

SOP-06: V5 Prediction Scoring

Operating SOP · Component #6 of the Operations Definition Register
Version 2.0 · 2026-08-10 · Owner: PPC Lead · Supersedes SOP-06 v1.0 (2026-07-24) in full · Thresholds and tolerances live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds

1. Verdict and Scope

VERDICT: V5 scores the prediction, never the outcome, and the difference is worth money. A system that grades outcomes teaches deciders to predict nothing falsifiable; a system that grades predictions produces a rule-level accuracy number that amends the framework itself under v4.0 Section 14 rule V5. This SOP governs the whole loop: the intake standard that makes a prediction scoreable, the Monday arrival sweep, realized-value assembly, the confound and void determination consumed from SOP-24 P3, the five-value outcome rubric, the four-value cause tag, filing to two instruments, the scoring-debt sweep, rule-level accuracy under rule V5, owner-level calibration, quarterly confidence calibration under rule V6, the override fold-in under rule V2, the human-versus-agent comparison that is the only evidence for widening an agent's drafting scope under rule V8, and the routing of a failing rule into a V3 amendment proposal.

Version 2.0 rebuilds v1.0 from 3 pages, 1,026 words, 0 step-level procedures, 5 failure modes and 0 worked examples to the 12-element SOP Build Kit v1.0 standard. It adds 14 procedures, 13 failure modes, 4 worked examples, 16 interface contracts and an operator ramp. It carries forward every threshold v1.0 traced to v4.0 and all eight canonical edge cases, which are account history that exists in no other document. It corrects three defects in v1.0: a four-value score rubric that dropped WRONG-DIRECTION, a confound citation pointing at the wrong SOP-24 procedure, and a set of scoring rules stated as adjectives rather than numbers.

1.1 The two objects this SOP keeps apart

A prediction is a written claim about a future measured value, filed at decision time in T6 Decision Log column **Expected Outcome Window** with its **Review Date**. An outcome is what the feeds report at that review date. V5 grades the first against the second and then asks a separate question: whose fault was any gap. Those are two axes, not one, and collapsing them is the failure this document exists to prevent. A push that reaches rank 8 because a competitor went out of stock is not a good prediction. A push that misses because the wrong bid landed is not a bad decision. v4.0 Section 14 rule V1 fixes the five fields a prediction must carry: metric, direction, magnitude, horizon, review date. v4.0 Section 13.2, the field **expected_outcome**, fixes the written form: named metric, direction, magnitude, horizon.

1.2 Why the rubric has five outcome values and not four

SOP-06 v1.0 used HIT, PARTIAL, MISS, VOID. Three other authorities use HIT, PARTIAL, MISS, WRONG-DIRECTION: the PPC Strategic Doctrine v1.0 Section 13 "Verification: the Loop that Makes Structure Learn", the Master PPC Data Architecture v1.0 Section 9 tab T6 Decision Log column definition, and the Account Rollup Workbook v1.0 tab A3 V5 Scoreboard, whose HIT, PARTIAL, MISS and WRONG-DIR columns are live formulas counting exactly those strings. v1.0 dropped a value the instruments count and added a value no instrument carries. Both halves of that are corrected here: WRONG-DIRECTION is restored because the scoreboard counts it, and VOID is retained because SOP-24 v2.0 P3 step 3 states that a confounded read is VOID and not a MISS and step 6 hands the void to this component by name. VOID has no column in T6 or A3 and no definition in v4.0; that gap is amendment item AM-2 and is worked in Section 6, WE-2.

1.3 Boundaries, with the owning component named for every excluded decision

Excluded decision	Owner	The seam, stated so it is owned
Every threshold, band, tolerance, gate, scenario and verdict	v4.0 (#8)	A numeral in this SOP without a v4.0 section beside it is a defect. Section 10.3 lists the numerals v4.0 does not carry; they are amendment items, never local values.
The confound and void cause classes, and the diagnosis score that names a loss	SOP-24	SOP-24 v2.0 P3 publishes five confound classes and hands each void to this SOP; this SOP consumes the class and never re-derives it.
Deployment accuracy, the T+1 diff, and the disposition of a value that landed wrong	#38	Per the Ledger row for #38, every staged set passes through it; this SOP routes EXECUTION-tagged rows there and grades no deployment.
The predictions themselves: what to change and what to expect	Every campaign-type SOP	Each owning SOP writes its own prediction inside its own procedure; this SOP screens the form at P1 and never authors the content.

Excluded decision	Owner	The seam, stated so it is owned
Push sizing, the required-clicks worksheet and the loss ceiling	SOP-23	This SOP scores a push prediction and returns systematic magnitude error to the worksheet owner; it never re-sizes a push.
Ladder rungs, step derivation, the discovered floor and the reversion rule	#28	This SOP scores a rung prediction against the recorded floor; a step size written here is a defect.
Incrementality test design, the register and the call	#35	This SOP scores the X12 capture prediction and consumes #35's void where a test rode a confounded window.
Automation rule registration, the weekly diff and rule retirement	#37	This SOP reads the rule change log only to separate an external change from an execution cause.
Recovery Mode entry, posture and close	SOP-26	This SOP scores predictions written inside the mode and never declares, runs or closes it.
Who holds which decision, and every escalation path	#3	Element 9 and Element 11 are drawn from the Authority Matrix v2.0 Section 1 rows by name; this SOP re-decides no cell.
Which units are terminal or LTSF-tiered	#27	Per Ledger conflict C6 the deciding document sits outside this corpus. A prediction on a terminal-tier unit scores on its own stated metric only, and this SOP infers no tier rule.
Envelope law, queue ranking and the funding release	Instrument layer	Per Ledger conflict C3, #31 carries no procedural document. This SOP names T0 zone H and Account Rollup tab A2 and cites no #31 procedure number.
Operator certification and the pass bar	#44	NOT BUILT, due 2026-09-25. This SOP supplies Element 11 as the input and asserts no certification standard.

2. Trigger Conditions

Every trigger below names the artifact that detects it and the maximum lag from signal to procedure start. Lag is measured from the moment the artifact is readable, not from the moment somebody notices.

Trigger	Detecting artifact	Maximum lag	Procedure
A review date has arrived	T6 Decision Log, column Review Date at or before today with column V5 Score reading Pending	Same Monday cascade	P2, then P3 to P7
A prediction is filed	T6 Decision Log, a new row written at deployment	At the staging step, before the row leaves for #38	P1 intake screen
A T+1 diff reports a mismatch	The #38 verification diff, Wednesday	1 working day	P6 cause tagging, EXECUTION branch
A confound is declared on a read	SOP-24 P3 handover carrying the class, or T5 Product Log column Trend-Exclusion Flag set across the read window	Same scoring pass	P4, then P5 VOID branch
The monthly batch date	The operating calendar (#4); v4.0 Section 14 rule V4 names monthly as the override review and stage review slot	Within the calendar month	P8, P9, P10, P12
Rule-level accuracy crosses its line	Account Rollup tab A3 V5 Scoreboard, hit rate by rule or scenario ID at 10 or more scored predictions	Same monthly batch	P9, then P14
The quarterly confidence read	v4.0 Section 14 rule V6 quarterly slot, read against E7 bands on A3	Within the quarter-close week	P11
Recovery Mode exits	SOP-26 exit log; v4.0 Section 14 rule V7 requires the R6 review to file as amendment proposals within one week	1 week from exit	P14
An agent drafting-scope request opens	The skills program docket under the 10-80-10 method (#7)	Before any scope change deploys	P13

Trigger	Detecting artifact	Maximum lag	Procedure
A prediction fails the intake screen	P1 checklist result	Same working session	Return to author; failure mode F5

3. Preconditions and Gates

Each gate below cites the T7 Manifest field group that carries it and states the observable that proves it passed. A gate that cannot be observed has not passed; it is unread.

Gate	Manifest ID	Rule, quoted to its v4.0 section	The observable that proves it passed
Decision log present	T7 row D7 Decision log + V5 scoring, source feed T6, refresh Weekly	v4.0 Section 14 rule V1: unscored actions do not accumulate	T7 row D7 reads PRESENT and T6 carries a row for every change deployed in the window, count matched against the #38 session record close count
G8 packet completeness	T7 PACKET COMPLETENESS SCORE cell	v4.0 Section 5 gate G8: DEGRADED packets, under 85% of gate-relevant fields, produce flags only. Section 3 family F15 sets FULL at or above 95%, WORKABLE 85-95%, DEGRADED under 85%	The completeness rollup stamping the cycle header reads at or above 85%. Below it, scores file as flags and the pass states what it could not see
G2 sample sufficiency	T7 rows A4 Statistical significance buckets (PRESENT) and B9 WoW deltas	v4.0 Section 5 gate G2: bid reads need 15 or more clicks in 7 days; CVR verdicts 100 or more clicks; CTR verdicts 1,000 or more impressions. WoW swings above 50% on thin samples are noise by definition	The realized window carries the sample the predicted metric requires. Short of it, the score is not a MISS; the read extends per gate G2
G1 validation quarantine	T7 rows feeding the K2 gate block; zero-flag columns compute mechanically	v4.0 Section 5 gate G1: impossible metrics quarantine the row, flag produced, no recommendation, data-refresh task opened same day	Zero zero-flag fires on the rows being scored. A quarantined row scores nothing and opens the refresh task instead
G10 event exclusion	T5 Product Log column Trend-Exclusion Flag; T2 Event Mode	v4.0 Section 5 gate G10: event weeks and the 14 days following are excluded from trend verdicts	The guard end date is computed explicitly and written into the score row, not judged by eye. Any overlap with the read window routes to P4
G9 price stability	T7 row D1 Per-variation margin (PARTIAL) and T5 event type price move	v4.0 Section 5 gate G9: a price change on the targeted variation within trailing 7 days marks CVR and ACOS reads PROVISIONAL	No price-move event inside the trailing 7 days of the review date, or the score carries the PROVISIONAL flag in the same words the engine computed, per v4.0 Section 13.3 rule J7
G12 MKL hygiene	T7 MKL refresh state	v4.0 Section 5 gate G12: a stale MKL past 30 days degrades the packet manifest	MKL refresh date within 30 days of the review date, and the scored term carries a syntax tag and a relevancy score
Trend comparability	T7 row A1 Four-period comparability (L30/P30/B30/L60), status MISSING, blocking all trend verdicts	v4.0 Section 3 family F7: trend classes IMPROVING, STABLE, DETERIORATING, VOLATILE, with VOLATILE at variance above 50% on adequate samples	Not currently passable. While A1 reads MISSING, a prediction whose metric is a trend class scores on a single-window read and the score row states the limitation. Logged AM-10
Rank series resolution	T7 rows D9 Daily rank series overlay (PARTIAL, head terms only) and B3 Rank pack (PARTIAL)	v4.0 Section 3 family F6: velocity classes IMPROVING at 2 or more positions per week toward target, FLAT at plus or minus 1, SLIPPING; ETA WITHIN horizon at or under the review horizon, default 3 weeks, BEYOND, NEVER at velocity at or under 0	The scored term appears in the daily series. Long-tail terms outside D9 score date accuracy at weekly resolution and the row says so
Deployment truth	The #38 T+1 verification diff for the deploying session	Per the Ledger row for #38, no deployment is complete until the next-day diff proves the values landed	A closed session record with zero unexplained mismatches covering the change the prediction was written against. Absent it, EXECUTION cannot be excluded and P6 withholds the JUDGMENT tag
Justification completeness	Master Action Record field presence, checked by code	v4.0 Section 12 rule P10: a verdict without a J-complete Master Action Record does not deploy, at any tier, in any mode. Section 13.3 rule J1: missing fields reject the record at the deploy step	The record carries expected_outcome and review_date populated. A row that reached T6 without them is a P1 intake failure caught late

4. Step-Level Procedures

Fourteen procedures. Time estimates are per unit named in the heading and assume the instruments are open: the product Command Workbook (tabs T5, T6, T7), the Account Rollup Workbook (tab A3), and the current #38 session record. Total weekly load at 31 filed predictions per product is 55-80 minutes for P2 to P8 combined.

P1. The prediction intake screen, run at staging, before the row leaves for #38

Per prediction, 60-90 seconds. Source: v4.0 Section 14 rule V1 and Section 13.2 field **expected_outcome**. This procedure is the cheapest one in the document and prevents the most expensive failure, because an unscorable prediction is not discovered until its review date, by which time the cycle that could have fixed it is closed.

1. Open the owning SOP's staged row and read the Master Action Record field **expected_outcome**. Confirm four elements are present as separate readable values, not as one sentence: the metric (named, and matching a column that exists in T2, T3 or T4), the direction (up, down, hold), the magnitude (a number or a bound, never "better"), and the horizon (a count of days or weeks).
2. Confirm the **review_date** field carries an actual calendar date. Compute it as the decision date plus the horizon and confirm it lands on a Monday, because the scoring pass runs inside the Monday cascade. A review date landing mid-week is moved forward to the next Monday and the horizon in the record is restated to match; the two must agree.
3. Confirm the conditions the prediction assumes are written: inventory band per v4.0 Section 5 gate G6, whether an event window is armed, the confidence score E7 with its three drivers per v4.0 Section 13.2 field **confidence**, and the benchmark source tier per v4.0 Section 12 rule P6. A prediction with no stated conditions cannot ever be scored VOID, because there is nothing for a confound to contradict.
4. Check the horizon against the confidence band. v4.0 Section 10 rule M10 states that LOW confidence, E7 under 50, shortens review horizons to 1 week and bars pushes. A LOW-confidence row carrying a 3-week horizon fails intake and returns to its author.
5. Write the row into T6 Decision Log with all eleven columns populated: **Date, Level, Target, Rule / Scenario Fired, Change (prior -> new), \$ Impact / wk, Expected Outcome Window, Review Date, V5 Score** set to Pending, **SI Conflict Note, Owner / Tier**. The tier value comes from v4.0 Section 12 rule P5.
6. Never write a prediction after the fact. A row whose **Date** is later than the deployment date it describes is retro-written and is failure mode F10. The T6 tab is append-only by design.

P2. The Monday arrival sweep, run first in the cascade

Per product, 5-8 minutes. Source: v4.0 Section 14 rule V1, which puts the scoring of last cycle's log as the next cycle's first step, and rule V4, which places the full audit weekly. The Master PPC Data Architecture v1.0 Section 9 states the same in one line for tab T6: the weekly audit's first step is scoring last week's rows.

1. Filter T6 Decision Log on **Review Date** at or before today AND **V5 Score** equal to Pending. This filtered set is the scoring worklist. Do not filter on **Date**; a long-horizon prediction filed six weeks ago arrives this week and a date filter hides it.
2. Count the worklist and record the count. Count separately the rows whose **Review Date** passed in an earlier cycle and which carry no documented gate G2 extension note from P3 step 3: that count is scoring debt and it reports every week whether or not it is zero, per failure mode F2. v4.0 sets no day count at which an arrived review becomes debt. This SOP therefore uses the cycle boundary, which is an artifact of the weekly cadence at v4.0 Section 14 rule V4 rather than a threshold, and logs the missing number at AM-4.
3. Sort the worklist by **\$ Impact / wk** descending. Scoring order follows dollars because the calibration value of a scored prediction rises with the size of the decision it graded, and because a truncated pass should truncate at the bottom.
4. Read T7 Manifest before reading any realized value. Record the completeness rollup and the status of rows A1, B3, D7, D8 and D9. Apply Section 3 gates. A pass that scores against a manifest it did not read is failure mode F13.
5. Open T5 Product Log and list every event dated inside the union of all read windows in the worklist. This single read serves P4 for the whole worklist and is cheaper than one read per row.

P3. Realized-value assembly, per prediction

Per prediction, 2-4 minutes; 5-7 minutes where the metric is rank and the term sits outside the daily series. The rule is that the realized value is pulled from the same feed and the same window definition that the prediction named, never from a different feed that happens to be closer to hand.

1. Identify the metric class from the **Expected Outcome Window** text and route to its feed: rank to T2 block K4 with the daily-series link, backed by T7 row D9; CVR, CTR, CPC, ACOS, CPA, orders, clicks, spend to T2 block K5 cross-type totals; placement-level values to T4; campaign-level budget and utilization to T3; TACOS and envelope to T0 and Account Rollup tab A0.
2. Build the realized window to the same length as the predicted horizon, ending on the review date. A 3-week prediction is read on 21 days of data, not on the last 7 because the last 7 are on screen.
3. Apply v4.0 Section 5 gate G2 to the realized window. If the window carries fewer than 15 clicks for a bid read, fewer than 100 clicks for a CVR read, or fewer than 1,000 impressions for a CTR read, stop. The outcome is not a MISS; the read extends and the row returns to the worklist with a new review date one cycle out. Record the extension so P8 can distinguish it from debt.

- Where the metric is rank and the term is inside T7 row D9 (head terms), read the daily series and record both the achievement date and the velocity in positions per day. Where the term is outside D9, read weekly resolution and mark the score row **date accuracy: weekly resolution**. Magnitude still scores exactly; only date precision degrades.
- Compute the realized value against the predicted magnitude and write both numbers into the score note. Never write a conclusion without both numbers beside it, per v4.0 Section 13.3 rule J2.

P4. The confound and void determination

Per prediction touched by an event, 3-5 minutes. Source: SOP-24 v2.0 P3, which is the procedure this SOP consumes by name, and v4.0 Section 5 gate G10. This SOP applies the determination; it does not define the classes.

- Take the five confound classes SOP-24 P3 step 4 publishes and test the read window against each: an event or deal window under gate G10, a competitor deal window under Section 8.7 rule S-C5, a price test in progress under gate G9, a listing change inside 14 days under Section 3 family F4, or a stockout on an adjacent variation shifting the traffic mix under Section 8.3 rule S-16.
- Compute the gate G10 guard end date explicitly: the last day of the event window plus 14 days. Write the computed date into the score row. SOP-24 P3 step 2 requires this be computed rather than judged by eye, and the reason is that a guard end estimated by eye is always earlier than the real one.
- Separate the two windows the confound can corrupt. A confound sitting on the decision date corrupts the diagnosis and belongs to SOP-24. A confound sitting inside the read window corrupts the score and belongs here. The same event can do one, both or neither; the score row says which.
- Where a confound falls inside the read window, the outcome is VOID and the row carries the class name and its v4.0 gate or rule reference. Where no confound falls inside, proceed to P5 with the read clean.
- Do not extend the loss ceiling. SOP-24 P3 step 7 re-dates the checkpoint to the guard end and records that v4.0 states no re-dating rule and no ceiling extension for the suppressed interval. A void score changes the calibration record; it does not authorize spend. A push whose ceiling is reached inside a suppressed interval stops at ceiling under v4.0 Section 8.4 rule S-R5 and any continuation is a fresh funding decision through Account Rollup tab A2. WE-2 computes what conflating the two costs.
- Where an incrementality test rode the same window, hand the void to #35 as well, per SOP-24 P3 step 6. A confounded read handed on as clean is failure mode F8.

P5. Outcome scoring: the five-value rubric

Per prediction, 1-2 minutes once P3 and P4 are done. This is the decision tree; it has five branches and is walked in order, top to bottom, first match wins.

Step	Test	Outcome	Why the order is fixed
1	Did a confound per P4 fall inside the read window?	VOID	Tested first because a contaminated read cannot be graded on any later test. SOP-24 P3 step 3: a VOID is not a MISS
2	Did the realized value move opposite to the predicted direction?	WRONG-DIRECTION	Tested before magnitude because a sign error is a different kind of error from a size error. Strategic Doctrine v1.0 Section 13 requires a written root cause for this value and for no other
3	Did a condition the prediction stated fail to hold, with the read still clean?	MISS	A stated condition that failed and did not produce the predicted behaviour falsifies the prediction even where the number landed close
4	Right direction, and magnitude and date both inside their tolerance?	HIT	All four elements of rule V1 must land: metric, direction, magnitude, date
5	Right direction, magnitude or date outside tolerance	PARTIAL	The residual branch. PARTIAL is the value that carries the most calibration signal, because a systematic direction of PARTIAL is a correctable assumption

Tolerance note, and it is a real constraint on this procedure. Steps 4 and 5 turn on the word tolerance. v4.0 uses that word 7 times, at Section 3 family F1, Section 3 family F16 metric E2, Section 4 rule X12, Section 9 lever rules, Section 10 rules M3 and M4, and Section 16, and quantifies it in none of them. Until v4.1 lands the bands, the scorer applies the prediction's own stated bound and nothing else: a prediction reading "rank 8 by day 14" is HIT only at rank 8 or better on or before day 14, and a prediction wanting a wider band states the band at P1. Where a prediction stated no bound, the outcome is not HIT and not PARTIAL; the row returns to its author as UNSCORABLE per failure mode F5. This is amendment item AM-1, and it is the single largest open item on this component.

P6. Cause tagging, on every non-HIT

Per non-HIT, 2-4 minutes. Four tags, closed, each with a fixed routing. VOID rows carry the confound class from P4 instead of a cause tag; every PARTIAL, MISS and WRONG-DIRECTION carries exactly one tag plus one line of evidence.

Cause tag	Definition and the evidence that names it	Routing, and what it does to calibration
JUDGMENT	The decision logic or its assumptions were wrong with the information available at decision time. Evidence: the record's own scenario line, showing which of the five X1 story elements carried the faulty read	Calibrates the decider. Counts in the owner-level read at P10 and in the rule-level read at P9
EXECUTION	The decision was right and landed wrong. Evidence: the #38 T+1 verification diff naming the delta, or an unlogged console edit surfaced by the diff	Routes to #38, never to the decider. Excluded from P10 entirely; appears in the #38 session review
WORLD	The world moved in ways the stated conditions could not name. Evidence: a T5 Product Log entry of event type competitor entry, or a documented platform change	Calibrates the condition writing, not the decider. Feeds P1's condition checklist, not the owner hit rate
DATA	The inputs were wrong or stale. Evidence: a T7 row reading STALE or MISSING for the field group the prediction depended on, or a gate G1 zero-flag fire, or a gate G9 PROVISIONAL mark	Routes to the feed owner and to gate discipline. Excluded from P10; accumulates as a manifest finding

1. Test EXECUTION first, because it is the only tag with an external artifact that either exists or does not. Open the #38 T+1 diff for the deploying session and compare intended value to landed value. A delta means EXECUTION; no delta clears it.
2. Where no T+1 diff exists for the window, do not tag JUDGMENT. EXECUTION has not been excluded, and a JUDGMENT tag applied without excluding it charges the decider for somebody else's error. The row files with the cause tag withheld and the owner-level read at P10 excludes it. This is interface I-2's non-arrival consequence.
3. Test DATA next against T7. Test WORLD next against T5. JUDGMENT is the residual and is only reached when the other three are excluded with evidence.
4. Write one line of evidence per tag. A tag without evidence does not file, per failure mode F4, because a cause series with untraced tags cannot support a rule review.

P7. Filing the score to both instruments

Per prediction, 30-45 seconds. Two writes, never one, because two different readers consume them.

1. Write the outcome value into T6 Decision Log column **V5 Score**, using the string exactly as spelled: HIT, PARTIAL, MISS, WRONG-DIRECTION, VOID. The Account Rollup tab A3 counts these with COUNTIFS formulas against literal strings; a spelling variant scores as zero in the rollup while looking correct in the product file.
2. Write the cause tag, the one-line evidence, and both numbers (predicted and realized) into the T6 Notes field or its equivalent adjacent cell.
3. Paste the row into Account Rollup tab A3, DETAIL block, columns **Product, Date, Target, Change (prior -> new), \$ Impact /wk, Review Date, V5 Score**. The summary block above computes **Predictions Logged, Scored, Pending**, per-value counts and **Hit Rate** from the detail rows. The Account Rollup READ ME names the Monday flow: Data Ops pastes the new T6 rows into A3 detail and the PPC Lead scores arrived reviews on A3.
4. Confirm the two instruments agree. A score present in T6 and absent in A3 is failure mode F7 and produces a rule-level accuracy number computed on a partial denominator.
5. VOID has no column in either instrument. Until AM-2 lands, a VOID files as the literal string VOID in the T6 **V5 Score** cell and is entered in the A3 detail block, where it is counted by none of the four summary formulas and therefore correctly excluded from **Hit Rate**. The scorer additionally records the VOID count in the pass note so the exclusion is visible rather than silent.

P8. The scoring-debt sweep and the monthly batch

Per product per month, 10-15 minutes. Source: v4.0 Section 14 rule V1, unscored actions do not accumulate, and rule V4, which places the monthly review slot.

1. Filter T6 on **V5 Score** equal to Pending AND **Review Date** in the past. Split the result into two lists: extended rows carrying a P3 step 3 gate G2 extension note with a new review date, and debt rows carrying nothing. Only the second list is debt.
2. Score the debt list in **\$ Impact / wk** order through P3 to P7. A prediction scored late is worth less than one scored on time, and worth far more than one never scored.
3. Report the debt count and the debt dollars as a single line in the monthly pass: rows, and the sum of their **\$ Impact / wk**. Both numbers report even at zero.
4. Where a review date passed and the underlying decision has since been superseded by a later change on the same target, score against the value at the original review date, not against today's state. The prediction was made about that date. Scoring it against a later state grades a different decision.

5. v4.0 sets no debt threshold. There is no number in v4.0 stating how many days past a review date a prediction may sit, or at what count debt escalates. Amendment item AM-4. Until it lands, the debt count reports and the PPC Lead decides; this SOP asserts no local trigger.

P9. Rule-level accuracy and the V5 rule review

Per monthly batch, 20-30 minutes across the account. Source: v4.0 Section 14 rule V5 and Section 3 family F15, metric E8.

1. On Account Rollup tab A3, group scored predictions by **Rule / Scenario Fired**, which T6 carries per row. v4.0 Section 3 family F16 defines E8 as prediction accuracy per rule ID and per owner; A3 carries hit rate by product, by rule or scenario ID, and by owner on a 13-week trend.
2. Compute accuracy per rule ID as HIT count divided by scored count. Exclude VOID rows from the denominator, since a VOID states that the rule was not tested. Record the excluded count beside the ratio.
3. Apply v4.0 Section 14 rule V5: any rule below 50% accuracy over 10 or more scored predictions enters review. Both conditions must hold. A rule at 20% on 4 scored predictions does not enter review; it accumulates.
4. For each rule entering review, pull its non-HIT rows and read the cause tag distribution. A rule whose misses are mostly EXECUTION is not a failing rule; it is a failing deployment path and routes to #38. A rule whose misses are mostly JUDGMENT on one shared assumption is a rule defect and routes to P14.
5. The review produces exactly one of three outputs, per rule V5: an amendment, a threshold change, or a documented defense of the rule. A rule that enters review and produces none of the three stays in review; it does not quietly leave.
6. Compute and record the account-wide hit rate alongside the per-rule numbers, because a rule at 45% inside an account running 42% is a different finding from the same rule inside an account running 70%.

P10. Owner-level calibration read

Per owner per month, 8-12 minutes. This is the read that calibrates deciders, human and agent, and it is deliberately narrower than the rule-level read.

1. Filter scored rows by **Owner / Tier**. Exclude every row tagged EXECUTION or DATA, and every VOID. What remains is the set the decider is answerable for: HIT, PARTIAL, MISS and WRONG-DIRECTION carrying JUDGMENT or WORLD.
2. Compute four numbers per owner: hit rate, PARTIAL rate, the direction of PARTIAL magnitudes (how many overshot the predicted magnitude and how many undershot), and the WRONG-DIRECTION count.
3. Read the PARTIAL direction as the primary signal. Systematically optimistic magnitudes are an assumption defect, not a skill defect, and the correction is to the assumption set rather than to the person. The account's recorded instance is velocity assumptions feeding SOP-23 push worksheets: the magnitude was consistently reached late, and the correction is a worksheet reference tightening returned to SOP-23, not a note in a review.
4. Read the VOID rate per owner as a separate line, computed on all rows including VOIDs. v4.0 Section 3 family F15 and Section 14 rule V5 give no VOID-rate ceiling, so this SOP states none: the number reports, the PPC Lead reads it against peers on the same decision classes, and every VOID must carry its documented confound class from P4 or it is not a VOID. Amendment item AM-5.
5. Minimum count. v4.0 sets the 10-or-more-scored-predictions bar at rule level only, under rule V5. It sets no minimum for an owner-level read. Amendment item AM-6. Until it lands, an owner-level read states its own denominator in the same line as its hit rate so the reader can weigh it.

P11. Confidence calibration, quarterly

Per quarter, 30-40 minutes. Source: v4.0 Section 14 rule V6 and Section 3 family F15, metric E7.

1. Group all scored predictions in the trailing quarter by the E7 confidence band recorded at decision time: HIGH at 75 or above, MEDIUM 50-74, LOW under 50, per v4.0 Section 3 family F15.
2. Compute hit rate per band. The expected ordering is HIGH above MEDIUM above LOW.
3. Apply rule V6: if LOW-confidence verdicts outperform their band or HIGH underperforms, the E7 weights recalibrate. E7 is a composite of sample size, benchmark source tier, manifest completeness and attribution cleanliness, per v4.0 Section 3 family F15; the recalibration names which of the four drivers carried the error.
4. Route the recalibration proposal to the v4.0 threshold owner through P14. E7 weights are v4.0 property; this SOP supplies the evidence and proposes, and per the Authority Matrix v2.0 row v4.0 threshold changes the CEO decides and the PPC Lead executes and reviews.
5. Check the second-order consequence before proposing. v4.0 Section 10 rule M10 makes E7 a sizing input: LOW halves lever step sizes and budget steps, shortens review horizons to 1 week and bars pushes. A recalibration that moves rows across the LOW line changes step sizes account-wide, so the proposal states how many rows in the trailing quarter would have crossed.

P12. The override fold-in, monthly

Per month, 15-20 minutes. Source: v4.0 Section 14 rule V2 and Section 3 family F15. This procedure closes the Ledger item recorded against this component as "Pending fold-in: V2 override-feedback loop".

1. Pull every row where a human overrode an AI-drafted verdict in the trailing month, with its logged reason. The Master PPC Data Architecture v1.0 Section 10 names the KDR column that carries the human confirmation stamp per the P5 tier.
2. Score the overrides the same way as any other prediction: the override carried its own expected outcome, and it is graded against it.
3. Compute the override rate and its direction of travel. v4.0 Section 14 rule V2 states that a falling override rate is the quality metric of the system. A rising override rate with high override hit rates means the drafting layer is behind the humans; a rising rate with low override hit rates means the humans are second-guessing a correct draft.
4. Fold validated overrides back into thresholds, scenarios or skills, per rule V2. A validated override is one that scored HIT while the drafted verdict it replaced would have scored otherwise on the same read. Route through P14.

P13. The human-versus-agent comparison and agent scope

Per scope request, 20-30 minutes. Source: v4.0 Section 14 rule V8 and the 10-80-10 method (#7). The scored decision log is the training corpus; this procedure is the evidence gate on widening what an agent may draft.

1. Hold the comparison to matched decision classes. An agent's hit rate on tail rows against a human's hit rate on oversight-class rows is not a comparison, because v4.0 Section 12 rule P5 puts every 500-or-more search-volume action, every gate failure, every structural change, every bid move above 25%, every budget move above \$50 per day and every LOW-confidence verdict into the human queue by construction.
2. Compute hit rate, PARTIAL rate and WRONG-DIRECTION count for both, on the same rule IDs, over the same window, with VOIDS excluded from both denominators.
3. Report the comparison to the PPC Lead with its denominators. Per the Authority Matrix v2.0 row Automation rules (#37) the PPC Lead decides and the incoming operator executes and reviews; agent scope changes that touch platform rules route through #37's register, and this SOP supplies evidence only.
4. No promotion threshold exists. v4.0 rule V8 names the training loop and sets no accuracy bar, sample size or margin for widening scope. Amendment item AM-8. Until it lands, this procedure produces the comparison and the decision is a named human's on the record.

P14. Amendment routing

Per amendment, 15-25 minutes. Source: v4.0 Section 14 rule V3 and rule V7.

1. Write the proposal in the form rule V3 requires: which rule is replaced, why, and the evidence. Oral re-derivations and parallel side-rules are process violations under rule V3, so a finding that lives only in a calibration conversation has not been filed.
2. Attach the scored evidence: the rule ID, the scored count, the hit rate, the cause tag distribution, and the specific rows. A proposal without its denominator cannot be weighed.
3. Route per the Authority Matrix v2.0 row v4.0 threshold changes: CEO decides, PPC Lead executes and reviews, every other role is informed.
4. Where the source is a Recovery Mode exit, file within one week of exit per v4.0 Section 14 rule V7, which requires every Recovery Mode exit to file its R6 review as amendment proposals within that window.
5. Track the proposal to a landed v4.1 section or to a documented rejection. An amendment item that neither lands nor is rejected is the same defect as an unscored prediction, one layer up.

5. Verdict Mapping

This SOP issues no keyword verdicts. It consumes them, and it holds every verdict to a predicted metric, which is why the mapping below runs the full closed vocabulary rather than a subset. Two separate closed vocabularies meet here and they are not interchangeable.

5.1 The two vocabularies

The verdict vocabulary is quoted verbatim from the Keyword Decision Record Template v3.0, sheet "Bands & Verdicts Reference", block "CLOSED VERDICT VOCABULARY (Framework Section 17): exactly one per Keyword Decision Record". Twenty-five strings. v4.0 Section 13.1 "The Closed Verdict Vocabulary" carries its own list of eighteen actions and states that a closed vocabulary is what makes the decision log queryable, the predictions scoreable per verdict type under rule V5, and the AI layer trainable. Per the Ledger row for #38, that document cross-maps the 25-string list to v4.0's 18-string Section 13.1 list explicitly, and the 18 / 25 / 26 count discrepancy is already an open item on the Ledger row for #46. This SOP quotes the KDR strings, which is what T6 and the KDR cycle file actually carry, and mints none.

The score vocabulary is HIT, PARTIAL, MISS, WRONG-DIRECTION, VOID, and its sources are named in 5.3. It is not a verdict vocabulary and never appears in a VERDICT cell.

5.2 Predicted-metric contract, per verdict string

Each verdict below binds its prediction to a metric class. A prediction whose metric does not match its verdict's contract fails the P1 intake screen, because a bid-down verdict predicting rank and a rank-push verdict predicting ACOS are both unscoreable against the rule that produced them.

Verdict string, quoted verbatim from KDR v3.0	The metric its prediction must name	The v4.0 section that sets the test, and the default horizon source
QUARANTINE: DATA REFRESH	Manifest or zero-flag state, not a performance metric	Section 5 gate G1. Scored on whether the refresh task closed and the flag cleared; a performance prediction on a quarantined row is unscorable
UNLOCK TRAFFIC	Clicks per day against the pre-change level	Section 9 rule L2, bids raised in 15-25% steps below 15 clicks; sufficiency read per Section 5 gate G2 at 15 clicks in 7 days
FIX BUDGET (TRUNCATED)	In-budget percentage of the day	Section 5 gate G3, in-budget below 70% marks every rate TRUNCATED. Horizon is one cycle, because the gate re-reads weekly
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	Cumulative clicks against the 50-click cap, and orders at the cap	Section 13.1, the FUND DELIBERATE TEST (50-CLICK CAP) entry. Scored on whether the cap was reached and what it returned, never on ACOS mid-test
FUND TEST (UNFUNDED PRIMARY)	Clicks delivered and orders returned inside the capped test	Section 5 gate G2 for the sufficiency read; Section 8.9 discovery scenarios for the success criterion
RANK PUSH: FUND SIZED PUSH	Organic rank against target rank, with the achievement date	Section 10 rule M2 for required clicks and the loss ceiling; Section 3 family F6 for velocity classes and for ETA WITHIN horizon at a default 3 weeks
RECOVERY PUSH	Organic rank against the pre-loss baseline rank	Section 7 rule D10 and Section 3 family F6: RECOVERY at best rank 10 or better in trailing 12 months, PROVEN at 5 or better. Cause class from SOP-24
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	Clicks per day against required clicks per day	Section 4 rule X4, the rank triad; Section 10 rule M2, required clicks equals DSTR divided by expected CVR
FIX CONVERSION FIRST; DEFER PUSH	CVR against its named benchmark	Section 5 gate G2, CVR verdicts need 100 or more clicks; Section 12 rule P6 for benchmark precedence; Section 5 gate G9 marks a read PROVISIONAL after a price move inside 7 days
PLACEMENT FIX FIRST	Top-of-search impression share, and placement-level CVR	Section 3 family F6, TOS percentage of 30 or above required on ranking campaigns; Section 10 rule M4 for the effective placement CPC test
RE-POINT VARIATION	Rank and orders on the backup variation, against the baseline snapshot	Section 5 gate G6 inventory bands; Section 9 rule L9, the re-point lever fires on YELLOW to RED same day
TAPER: RANK ACHIEVED, STEP BIDS DOWN	Rank held at or under target, and the CPC reached	Section 10 rule M3, bid down 10% per week from rank achieved; on slippage beyond tolerance restore one step and that value is the recorded minimum viable CPC
HOLD; DEFEND LIGHTLY AT RANK	Rank held, and the maintenance spend per week	Section 8.6 defensive scenarios; Section 10 rule M8 for the defense budget ceiling
DEFEND AT FLOOR	Rank held at the discovered floor CPC	Section 10 rule M3 for the floor definition and its quarterly re-test; Section 13.1 DEFEND AT FLOOR entry
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	CPC against max profitable CPC, and ACOS against its objective band	Section 3 family F1, max profitable CPC equals margin dollars times CVR, with CPC ratio bands at 0.8 SAFE, 0.8-1.0 AT-CEILING, 1.0-1.5 PUSH-ONLY, above 1.5 UNJUSTIFIABLE outside a sized push
INCREMENTALITY LADDER	Total capture, organic plus paid, across the ladder step	Section 4 rule X12: one ladder step down, and if total capture holds within tolerance for 2 weeks the ladder continues and the floor is recorded
ADD EXACT + RESTRUCTURE	Share of the term's clicks carried by the new exact	Section 4 rule X10, auto or broad carrying above 40% of a priority term's clicks is a structure failure; Section 3 family F8 structure states

Verdict string, quoted verbatim from KDR v3.0	The metric its prediction must name	The v4.0 section that sets the test, and the default horizon source
CONSOLIDATE DUPLICATES	The survivor campaign's efficiency against the pre-consolidation blended value	Section 9 rule L6, duplicate same-type same-objective campaigns consolidate to the winner. Strategic Doctrine v1.0 Section 13 requires structural changes to be held to the same standard as bid changes
GRADUATE FROM DISCOVERY	The new exact's CPA against the discovery source CPA	Section 9 rule L6 and Section 11 rule T6, graduation at 3 or more orders with the full unit deployed; Section 11 rule T1 for the transition
NEGATE (PRE-LOAD / REACTIVE / STEERING)	Wasted-ad-spend percentage, or the source term's click share after steering	Section 9 rule L5, the three negation modes. Bands from the KDR "Bands & Verdicts Reference" WAS percentage line
ENTER SBV ARBITRAGE	CVR-adjusted cost per SP-equivalent conversion on the substituted surface	Section 9 rule L8: the pivot opens when SP exact CPC is 1.5 times its ceiling or above and SBV CPC is 0.6 times the SP CPC or below
OPEN / ROTATE / EXIT CONQUEST	Cost per order on the conquest target against its objective band	Section 8.7 conquering scenarios; Section 3 family F1 objective Max-ACOS bands, conquest at 10-15%
STOP-LOSS: CONCEDE OR DEFER	Spend against the loss ceiling, and rank velocity at the stop	Section 8.4 rule S-R5, stop-loss fires on sufficient clicks with CVR at or above benchmark and rank velocity near zero across the review window
PAUSED-PENDING (REASON + RE-ENTRY)	The stated re-entry criterion, as a binary met or not met	Section 6 rule O5, Paused-pending is a formal objective carrying its blocking reason and re-entry criterion; Section 11 rule T1, re-enter only on the stated criterion
LEAVE TO HARVEST	Discovery-surface orders on the term, at the re-evaluation date	Section 6 rule O2 tail and Section 8.9 discovery scenarios. This verdict's re-evaluation date is the one most often missed; edge case EC7 records it

5.3 The score vocabulary and its sources

Score value	Source, quoted to its document and section	Definition applied here
HIT	Strategic Doctrine v1.0 Section 13; Data Architecture v1.0 Section 9 tab T6; Account Rollup v1.0 tab A3	Realized inside the stated bound on metric, direction, magnitude and date
PARTIAL	Same three sources	Direction right; magnitude or date outside the stated bound
MISS	Same three sources	A stated condition failed to produce the predicted behaviour, with the read clean
WRONG-DIRECTION	Same three sources. Strategic Doctrine v1.0 Section 13 adds the only value-specific obligation in the corpus: wrong-direction requires a written root cause	The realized value moved opposite to the predicted direction. Restored at v2.0; v1.0 dropped it
VOID	SOP-24 v2.0 P3 step 3, which states the read is VOID and not a MISS, and step 6, which hands the void to this component. No v4.0 definition and no instrument column exist. Amendment item AM-2	A confound per P4 fell inside the read window. Excluded from every accuracy denominator

5.4 Cause tags

JUDGMENT, EXECUTION, WORLD, DATA, carried forward unchanged from SOP-06 v1.0 Section 3, which is the only place in the corpus they are enumerated. Definitions and routing are at P6. This SOP owns them; SOP-24 owns the confound classes that make a VOID legal, per its v2.0 P3 step 4.

5.5 Strings this version does not use

SOP-06 v1.0 Section 5 recorded EC7 with the outcome string Scoring debt and EC8 with UNSCORABLE, and EC6 with working as designed. None of the three belongs to any closed vocabulary. At v2.0 they are retired as outcome values and survive as states: scoring debt is a P8 count, UNSCORABLE is a P1 and P5 return-to-author routing, and the process-versus-step distinction in EC6 is handled by scoring the two predictions separately. A locally minted score string is invisible to the Account Rollup tab A3 formulas exactly as a locally minted verdict string is invisible to the decision log.

5.6 The eight canonical edge cases, carried forward and re-scored

These are account history and exist in no other document in the library. Carried forward verbatim in substance from SOP-06 v1.0 Section 5, re-scored against the five-value rubric, with the two cases the rubric change affects marked.

Case	Score at v2.0	Why, and the consequence
EC1. Push predicted rank 22 to 8 by day 14; achieved day 13	HIT	Inside the stated bound on all four elements. Calibration records the velocity margin. Expanded as WE-1
EC2. Same push, achieved day 19	PARTIAL, cause WORLD or JUDGMENT	Direction right, date outside bound. A documented competitor deal mid-push tags WORLD; an undocumented drift tags JUDGMENT against the velocity assumption, and the worksheet reference corrects at SOP-23
EC3. Rank never moved, full traffic delivered	MISS, cause JUDGMENT	The stated condition, a validated CVR, was thin. The validation standard tightens for the class rather than the term
EC4. Event window landed mid-push, read impossible	VOID	Not a MISS. The confound is documented and the push re-predicts post-settle. Changed at v2.0: the void no longer implies a ceiling extension. Expanded as WE-2
EC5. Graduation predicted CPA at or under source; realized above, and the diff shows the wrong bid landed	MISS, cause EXECUTION	Routes to the #38 session review; the decider's calibration is untouched, which is the point of the second axis. Expanded as WE-3
EC6. Taper rung predicted hold; slipped; reverted per the ladder	Rung prediction: MISS. Process prediction: HIT	Changed at v2.0: v1.0 scored the process limb as "working as designed", which is not a score value. Two predictions, scored separately. Floor discovery is the goal, so a ladder that reverted did what rule M3 says it does
EC7. A LEAVE re-eval date passed unscored	Not a score. A P8 debt row	The unscored prediction is itself the finding, failure mode F2. The monthly batch sweeps it and the count reports
EC8. "Sales will improve"	Not a score. A P1 intake failure	Returned to author. Nothing files until it carries magnitude, date and conditions. This is the incident behind failure mode F5 and behind the entire P1 procedure
EC9. Bid reduced to close an ACOS gap; ACOS rose across the window with the read clean	WRONG-DIRECTION, cause JUDGMENT	New at v2.0. v1.0 had no branch for a sign error and such rows collapsed into MISS. Requires a written root cause per Strategic Doctrine v1.0 Section 13

6. Worked Examples

All four use the standard fictional base declared in the PPC Economics Worked-Example Set v1.0, because the anonymized live product economics spine is not attached to this session: AOV \$80.00, contribution \$28.00 per order, margin 35.0%, break-even ACOS 35.0%, Target ACOS 17.5%, Maximum Acceptable 26.3%, CVR 5.2% unless a case states otherwise. Values are illustrative and internally consistent. WE-1 and WE-2 are built on case E3 of that set so the sizing arithmetic reconciles to the training standard.

WE-1. The clean case: a sized push scored HIT, and what the margin is worth

The prediction, filed at deployment. Verdict RANK PUSH: FUND SIZED PUSH. Metric organic rank; direction toward 1; magnitude rank 8; horizon 14 days; review date day 14. Conditions stated: inventory GREEN per gate G6, no armed event window, CVR at or above 5.0%, confidence E7 = 78, band HIGH.

The sizing it was written against (Economics Worked-Example Set case E3). Rank 22, target 8. DSTR at rank 8 is 4.0 orders per day; current 1.6; gap 2.4. At CVR 5.0%, required clicks per day = $2.4 / 0.05 = 48.0$. At expected CPC \$2.30, planned spend per day = $48.0 \times \$2.30 = \110.40 . At velocity 1.0 position per day over 14 days, planned total = \$1,545.60. Planned paid recovery per day = $2.4 \times \$28.00 = \67.20 , so planned loss per day = $\$110.40 - \$67.20 = \$43.20$, and planned loss over 14 days = \$604.80 against a ceiling set at \$650.00.

P3, realized assembly. The term is inside T7 row D9, so the daily rank series is read. Rank 8 first recorded on day 13. Delivered clicks averaged 47.2 per day across 13 days = 613.6 clicks. Realized average CPC \$2.34. Realized spend = $613.6 \times \$2.34 = \$1,435.82$. Orders at CVR 5.0% = $613.6 \times 0.05 = 30.68$. Contribution recovered = $30.68 \times \$28.00 = \859.04 . Realized loss = $\$1,435.82 - \$859.04 = \$576.78$, which is 88.7% of the \$650.00 ceiling.

P4, confound. T5 Product Log carries no event dated inside days 1-13. No price move inside the trailing 7 days of the review date, so gate G9 does not fire. Read is clean; no VOID.

P5, scoring. Step 1 no confound. Step 2 direction correct, rank fell from 22 toward 1. Step 3 both stated conditions held: inventory GREEN throughout, realized CVR 5.0% at the stated floor. Step 4 magnitude reached, rank 8; date inside bound, day 13 against day 14. HIT.

P10, what the calibration records. Realized velocity = 14 positions / 13 days = 1.077 positions per day against the planned 1.0, a margin of +7.7%. At weekly resolution that is 7.54 positions per week, class IMPROVING per v4.0 Section 3 family F6, which sets IMPROVING at 2 or more positions per week toward target. The margin is recorded, not celebrated: one HIT with a 7.7% velocity margin is a single data point against the 10-scored-prediction bar rule V5 sets for rule-level review.

What happens next, and where this SOP stops. Rank 8 held from day 13 starts the 2-week at-target clock under v4.0 Section 11 rule T1, which flips the objective to post-objective and hands the term to the taper. This SOP files the score and stops; the flip is SOP-25's and the ladder is #28's.

WE-2. The edge case that goes wrong: a void read, and the ceiling it does not buy

Same prediction and same sizing as WE-1. Rank 22 to 8, horizon 14 days, ceiling \$650.00, planned loss \$43.20 per day.

What went wrong. On day 6 a confirmed deal window arms on the product and runs 4 days, closing end of day 9. T5 Product Log carries the event with **Trend-Exclusion Flag** set. Nobody re-reads the review date, and the scorer arrives on day 14 to score a push whose entire second week sat inside an event distortion.

P4, worked in the order SOP-24 P3 requires. Step 1: the confound sits inside the read window, not on the decision date, so it corrupts the score and not the diagnosis. Step 2: compute the guard end explicitly. v4.0 Section 5 gate G10 excludes event weeks and the 14 days following from trend verdicts, so the guard runs to day 9 + 14 = day 23. The day-14 review date sits 9 days inside the guard. Step 4: the class is the gate G10 event-window class, one of the five SOP-24 P3 step 4 publishes. The read is VOID. Step 5 of P4 is where the money is.

The arithmetic of the trap. At day 14 the push had run its full planned course: 14 days x \$110.40 = \$1,545.60 spent, 672 clicks, 33.6 orders at CVR 5.0%, \$940.80 recovered, realized loss \$1,545.60 - \$940.80 = \$604.80, which is 93.0% of the \$650.00 ceiling.

Re-dating the checkpoint to the guard end at day 23 adds 9 days. At \$43.20 loss per day that is 9 x \$43.20 = \$388.80 more, for a total of \$604.80 + \$388.80 = \$993.60, which is 152.9% of the ceiling. The push would breach its ceiling by \$343.60 while every day of the breach looked procedurally correct, because a re-dated checkpoint reads like an approved extension.

The correct handling, and the gap it exposes. The score files as VOID and is excluded from the owner's hit rate: the decider is protected from a read nobody could take, which is the entire purpose of the value. The ceiling does not move. SOP-24 P3 step 7 states plainly that v4.0 carries no re-dating rule and no ceiling extension for the suppressed interval. So the stop-loss under v4.0 Section 8.4 rule S-R5 stays live against the original \$650.00, the push stops at ceiling on day 15, and any continuation to day 23 is a fresh funding decision priced at \$388.80 and ranked in the M5 queue on Account Rollup tab A2 against every other ask. Additionally, v4.0 Section 8.12 rule S-E5 governs any rank reached inside an armed window: the achievement is provisional and the 2-week at-target clock under rule T1 starts only after the exclusion window closes, so even a rank-8 reading on day 12 would not start the taper clock until day 23.

Amendment item. The re-dating and ceiling-extension rule is absent from v4.0 and is logged at AM-3. SOP-24 v2.0 logs the same gap from its side at its AM-5. Two components hitting the same missing rule from opposite directions is the strongest evidence a v4.1 item can carry, and both references travel with the proposal.

WE-3. The second axis earning its keep: a MISS that is not the decider's

The prediction. Verdict GRADUATE FROM DISCOVERY on a term at 4 orders in the discovery surface, which clears the 3-order bar at v4.0 Section 11 rule T6. Metric CPA on the new exact; direction at or below; magnitude \$22.40, which is the discovery source CPA; horizon 3 weeks per the v4.0 Section 3 family F6 default review horizon; review date the third Monday out. Conditions: seed bid at Expected CPC \$2.80, steering negation live at source, CVR 12.5% carried from the discovery reads.

The arithmetic the prediction was built on. At CVR 12.5%, clicks per order = 1 / 0.125 = 8.0. At the staged seed bid of \$2.80, predicted CPA = 8.0 x \$2.80 = \$22.40, exactly at the source CPA, and \$5.60 under the \$28.00 contribution.

P3, realized. Realized CPA over the 3-week window, 22 orders: \$28.80. Working backwards at the same 8.0 clicks per order, that implies a realized CPC of \$28.80 / 8.0 = \$3.60.

P6, cause tagging in the order the procedure sets. EXECUTION is tested first. The #38 T+1 verification diff for the deploying session shows the keyword row landed at bid \$3.60 against a staged \$2.80. The delta is real and it is the whole gap. Note the size: \$3.60 / \$2.80 = 1.2857, a +28.6% bid delta, which exceeds the 25% bid-move line at v4.0 Section 12 rule P5 and was therefore human-confirmable at deployment and missed there.

P5 and P7, what files. Direction wrong on the CPA test, since predicted at or below \$22.40 and realized above it, but this is not WRONG-DIRECTION: the predicted direction was a bound, and the realized value moved against the bound because the input changed, not because the relationship reversed. Step 3 of the tree fires: a stated condition, seed bid at Expected CPC \$2.80, failed to hold. MISS, cause EXECUTION.

What it cost, and where the score goes. Overrun per order = \$28.80 - \$22.40 = \$6.40, which is 28.6% above prediction. Across 22 orders: \$140.80 of avoidable cost in one window. Against the \$28.00 contribution the term now loses \$28.80 - \$28.00 = \$0.80 per order, so the graduation that was \$5.60 per order profitable at the staged bid is \$0.80 per order unprofitable at the landed one, a swing of \$6.40. The row routes to #38's session review per interface I-15. The decider's hit rate is untouched and the row is excluded from the P10 owner read, which is the point of the axis: charging this to the decider would teach the graduation rule to predict conservatively against a deployment defect it does not control.

WE-4. Rule-level review: a scenario rule crosses the V5 line, and the defense that is only half a defense

The trigger. P9 groups scored predictions by **Rule / Scenario Fired** on Account Rollup tab A3. Scenario S-R3, the traffic-deficit fund verdict, has accumulated 14 scored rows across 5 products in the trailing 13 weeks: 5 HIT, 3 PARTIAL, 4 MISS, 1 WRONG-DIRECTION, 1 VOID.

The computation, with the denominator stated. VOID rows are excluded because a VOID states the rule was not tested. Scored denominator = 14 - 1 = 13. Accuracy = 5 / 13 = 38.5%. Both conditions of v4.0 Section 14 rule V5 hold: below 50%, and 10 or more scored predictions at 13. The rule enters review.

The cause distribution, read before any conclusion. Of the 8 non-HIT non-VOID rows, 1 is EXECUTION and routes to #38, 1 is DATA against a T7 row reading STALE, and 6 carry JUDGMENT. All 6 JUDGMENT rows fail on the same element: the required-clicks estimate derived from DSTR under v4.0 Section 10 rule M2 ran under the clicks the term actually needed. In all 6, the term had fewer than 8 weeks of push history, so the rule M2 amendment that corrects expected CVR and the click requirement by the term's observed Rank Response Elasticity (E4) had not yet engaged.

The defense, and why it is provisional. Excluding the 6 rows where the elasticity correction was structurally unavailable leaves 7 scored rows, and all 5 HIT rows survive the exclusion: 5 / 7 = 71.4%, well above the 50% line. That is a documented defense of the rule under rule V5, and it is not sufficient, because 7 is below the 10-scored-prediction bar the same rule sets. The rule stays in review until 10 elasticity-corrected predictions accumulate. Recording it as defended at 7 would be the same category error as scoring a CVR verdict on 80 clicks against the 100-click bar at v4.0 Section 5 gate G2.

The three outputs, per rule V5. One documented defense, filed provisional with its denominator. One threshold change proposed: rule M2's elasticity correction currently engages at 8 or more weeks of push history, and the evidence argues for an interim correction factor on terms below that bar, which is a v4.1 item and is logged at AM-7 rather than written into this SOP. One routing: the 1 EXECUTION row to #38, the 1 DATA row to the feed owner.

The second finding, which is about this procedure and not about S-R3. The 38.5% figure depends entirely on excluding the VOID from the denominator. Including it gives 5 / 14 = 35.7%, still below the line here, but on a rule sitting near 50% the choice of denominator decides whether a review opens at all. v4.0 Section 3 family F15 defines E8 as the percentage of predictions hit at horizon and does not state whether an untestable read sits in the denominator. That is amendment item AM-2 and it is the reason P9 step 2 requires the excluded count to be recorded beside the ratio.

7. Failure Modes and Escalation

Thirteen modes. F1 to F5 are carried forward from SOP-06 v1.0 Section 6 with their account-history incidents intact and their detection signals converted from adjectives to reads. F6 to F13 are new at v2.0.

Failure	Detection signal	Response	The incident that motivated it
F1. Outcome-bias scoring (carried from v1.0)	A HIT filed on a prediction whose stated conditions did not hold; a MISS filed on a row where T5 carries a documented confound inside the read window	Re-score through P5 in order. The rubric grades the prediction, not the luck. A HIT that violated its own conditions is a MISS at step 3	The pattern v1.0 recorded as the reason the rubric has two axes at all: scores were being written from whether the number looked good rather than from whether the claim held
F2. Scoring debt (carried from v1.0)	The P2 step 2 count: rows whose review date passed in an earlier cycle, carrying no gate G2 extension note, with V5 Score still Pending. The signal is the count rising month over month; no absolute trigger exists in v4.0, per AM-4	P8 batch sweeps in dollar order. The count and its dollars report every cycle at any value including zero. Persistent debt escalates to the PPC Lead as a cadence gap, not as an individual finding	EC7: a LEAVE TO HARVEST re-evaluation date passed unscored and the unscored prediction was itself the finding. Reinforced by the Transition Brief v1.2 week-1 instruction to confirm the V5 scoring backlog with the PPC Lead on day 1
F3. VOID abuse (carried from v1.0)	P10 step 4: an owner's VOID rate read against peers on the same decision classes, with every VOID checked for its documented confound class from P4. A VOID with no class named is the signal	The calibration read names it. A VOID without its documented confound is re-scored through P5 from step 2. v1.0 stated this as "far above peers", which is an adjective; the numeric ceiling is AM-5 and until it lands the read reports rates and denominators rather than firing a threshold	Recorded in v1.0 as a watch: the value that protects deciders from luck in both directions is also the value that pays a decider to find a confound

Failure	Detection signal	Response	The incident that motivated it
F4. Missing cause tags (carried from v1.0)	Any non-HIT, non-VOID row in T6 whose note carries no tag or no evidence line	Tag or the score does not file. Calibration is impossible without causes, and a rule review at P9 that cannot read a cause distribution cannot choose between amendment, threshold change and defense	Carried from v1.0 unchanged. It is the reason P6 step 4 requires one line of evidence per tag
F5. Vague predictions accepted (carried from v1.0)	A T6 row whose Expected Outcome Window carries no magnitude, no date, or no stated condition; found at scoring time rather than at intake	Return to author. Nothing files. The fix is at intake, which is why P1 runs at staging and not at review	EC8, the canonical instance: "Sales will improve" reached a decision record. It is unscorable on every axis and it is the single incident that produced the prediction standard
F6. WRONG-DIRECTION collapsed into MISS	Account Rollup tab A3 WRONG-DIR column reading zero across a full quarter while T6 notes describe realized values moving opposite to prediction	Re-score the affected rows at P5 step 2. A sign error and a size error have different corrections, and merging them hides the class that Strategic Doctrine v1.0 Section 13 singles out for a written root cause	The v1.0 rubric itself: it dropped the value that the T6 column definition, the A3 formulas and the Doctrine all carry. Every score filed under v1.0 is exposed to this
F7. Score filed in one instrument only	A row scored in T6 with no matching row in the Account Rollup tab A3 detail block, or an A3 row whose product-file counterpart still reads Pending	Complete both writes per P7. Reconcile the A3 Scored count against the T6 non-Pending count per product before the monthly batch closes	New. The A3 READ ME splits the Monday flow across two roles, Data Ops pasting and the PPC Lead scoring, and a split flow is where one-sided writes live
F8. A confounded read handed over as clean	A score filed HIT, PARTIAL or MISS on a window that T5 shows carried a trend-exclusion flag, or a #35 incrementality result consumed without its void state	Re-score through P4. Hand the void to #35 as well where a test rode the same window	SOP-24 v2.0 P3 step 6 states the principle in one line: a confounded read handed over as clean is worse than no read, because the consumer cannot detect the contamination and files a conclusion against it
F9. Accuracy computed on the wrong denominator	A P9 rule-level ratio filed without its excluded-VOID count beside it; or two ratios on the same rule differing by exactly the VOID count	Recompute with VOIDS excluded and the exclusion stated. Where the two denominators fall on opposite sides of the 50% line at rule V5, the rule enters review and the ambiguity files as evidence on AM-2	New, and worked in WE-4: S-R3 reads 38.5% on 13 or 35.7% on 14, and on a rule nearer the line the choice decides whether a review opens
F10. Retro-written predictions	A T6 row whose Date is later than the deployment date of the change it describes, or a prediction whose magnitude matches the realized value to an implausible precision	Void the row for calibration purposes and log it. v1.0 stated the rule directly: predictions file in T6 at decision time, never retro-written. T6 is append-only by design	Carried forward from the v1.0 prediction standard. A retro-written prediction produces a HIT that teaches nothing and inflates every denominator it enters
F11. Calibration read on too few rows	A P10 owner read reporting a hit rate without its denominator, or a rule review opened below the 10-scored-prediction bar at v4.0 Section 14 rule V5	State the denominator in the same line as every rate. Do not open a rule review below the bar; let the rule accumulate. Do not close one below it either, per WE-4	New. v4.0 sets a count bar at rule level and none at owner level, which invites reading a 2-of-3 owner rate as a finding
F12. Agent scope widened on an uncontrolled comparison	A P13 comparison whose two populations differ in decision class, in rule ID mix, or in window; or a promotion proposal citing a rate without a denominator	Rebuild the comparison on matched classes per P13 step 1. v4.0 Section 12 rule P5 routes the hardest decisions to humans by construction, so an unmatched comparison always flatters the agent	New. v1.0 named the human-versus-agent comparison as promotion evidence and set no control on how it is built
F13. Manifest-blind scoring	A scoring pass with no recorded T7 read; a trend-class prediction scored while T7 row A1 Four-period comparability reads MISSING; a WORLD tag applied while row D8 Competitor depth reads MISSING	Read T7 first per P2 step 4 and record the rollup. Score trend predictions on a single-window read with the limitation stated in the row. Do not tag WORLD on competitor grounds without a manually logged T5 competitor entry	New, and live: T7 currently carries A1 MISSING blocking all trend verdicts, D8 MISSING blocking competitor-context reads, D3 MISSING, and D9 PARTIAL at head terms only. SOP-24 v2.0 works the same D8 constraint from the diagnosis side

Escalation. F2 escalates to the PPC Lead when the debt count rises for 2 consecutive monthly batches, with the count and its dollars attached. F8 escalates on first occurrence, because the downstream consumer has already filed a conclusion. F10 escalates on first occurrence with attribution, because it is a record-integrity failure rather than a process gap. All other modes escalate on second occurrence within a quarter. Per the Authority Matrix v2.0 standing rule 3, the BM owner's review on verdict rows is the daily huddle and disagreements route up, not sideways.

8. Outputs and Artifacts

Artifact	Storage location	Retention	Downstream consumer
The filed score: outcome value, cause tag, evidence line, predicted and realized values	Product Command Workbook, tab T6 Decision Log, column V5 Score plus the adjacent note cell	Life of the product workbook; append-only, never overwritten	The next cycle's P2 sweep; the KDR cycle file, which per Data Architecture v1.0 Section 10 is the frozen snapshot V5 scores against
The account scoreboard rows	Account Rollup Workbook, tab A3 V5 Scoreboard, summary block and DETAIL block	13-week rolling trend per the tab definition, with detail retained beyond it	P9 rule-level review; P10 owner reads; the PPC Lead in the Monday flow; the BM owner in the daily huddle
The scoring-debt line: row count and summed \$ Impact / wk	The monthly batch note, filed with the A3 refresh	Rolling 12 months	PPC Lead, as the cadence read; failure mode F2 escalation
The rule-review docket: rule ID, scored count, hit rate, excluded VOID count, cause distribution, and one of amendment or threshold change or documented defense	Filed with the monthly batch and carried to the v4.1 amendment queue	Until the rule leaves review	The v4.0 threshold owner via P14; the Authority Matrix row v4.0 threshold changes
The quarterly confidence calibration: hit rate by E7 band with the driver named	Filed with the quarter-close batch	Rolling 8 quarters	The v4.0 threshold owner, since E7 weights and rule M10 step sizing are v4.0 property
The override fold-in read: override rate, direction of travel, validated overrides	Monthly batch note	Rolling 12 months	Threshold, scenario and skill owners per v4.0 Section 14 rule V2; closes the Ledger fold-in item on this row
The human-versus-agent comparison with both denominators	The skills program docket under the 10-80-10 method (#7)	Life of the agent scope grant	PPC Lead as the deciding role; #37 where the scope change touches a platform rule
EXECUTION-tagged rows with their T+1 diff references	Handed to the #38 session review	Life of the session record	#38, per the Ledger row for #38
Systematic PARTIAL-magnitude findings by rule	The monthly batch note, routed to the owning SOP	Until the assumption corrects	SOP-23 worksheet references; #28 rung assumptions; SOP-24 recovery sizing

The gate artifact. A scoring cycle closes only when the T6 non-Pending count equals the arrived-review count minus documented gate G2 extensions, and the Account Rollup tab A3 Scored count reconciles to it per product. An unreconciled cycle is open, not closed.

9. Skill-Conversion Block

Block	Specification, specific to this component
Input bindings	T6 Decision Log, all eleven columns, filtered on Review Date at or before today and V5 Score equal to Pending. T5 Product Log columns Date, Event Type, Affected Scope, Trend-Exclusion Flag. T7 Manifest rows A1, A4, B3, B9, D1, D7, D8, D9 plus the PACKET COMPLETENESS SCORE cell. T2 blocks K4 (rank, daily-series link, best rank) and K5 (cross-type totals, CTR, CVR, CPC, ACOS). T3 campaign rows for budget and utilization predictions; T4 for placement predictions; T0 zones C-F and Account Rollup tab A0 for envelope and TACOS predictions. The #38 T+1 verification diff for the deploying session. The #37 rule change log. SOP-24 P3 void handovers with their class references. Account Rollup tab A3, both blocks.

Block	Specification, specific to this component
Deterministic logic	The P1 six-point intake screen, including the rule M10 check that a LOW-confidence row cannot carry a horizon longer than 1 week. The P2 filter and the two counts, worklist and debt. The P3 metric-class routing table and the gate G2 sufficiency test with its extension branch. The gate G10 guard-end computation, last event day plus 14. The P5 five-branch decision tree in fixed order. The P6 tag ordering, EXECUTION then DATA then WORLD then JUDGMENT as residual, including the rule that no T+1 diff means no JUDGMENT tag. The P7 literal-string write and the two-instrument reconciliation. The P9 accuracy computation with VOID excluded and both conditions of rule V5 tested jointly. The P10 exclusion set. The P11 band grouping. The P13 matched-class filter.
Human checkpoints (drawn by row from the Authority Matrix v2.0 Section 1)	AI-executable with audit: the P2 sweep and both counts, P3 realized assembly, the P4 guard-end arithmetic, P5 scoring on clean unambiguous reads, the P8 debt sweep, the P9 and P10 and P11 computations, the P13 comparison arithmetic, and both P7 writes. Human-mandatory, by Authority Matrix row: row Weekly verdicts (established US), where the PPC Lead decides and the incoming operator executes within assigned scope, confirms every JUDGMENT tag, because a JUDGMENT tag is an assertion about a named person's or agent's reasoning. Row v4.0 threshold changes, where the CEO decides and the PPC Lead executes and reviews, takes every P14 amendment and every P11 recalibration proposal. Row Automation rules (#37), where the PPC Lead decides and the incoming operator executes and reviews, takes every P13 agent-scope change. Row Bulk deployments (#38), where the PPC Lead reviews and the incoming operator executes, takes every EXECUTION disposition. Row Incrementality tests, where the incoming operator decides and executes and the PPC Lead reviews, takes any void shared with #35. Row Code-red declaration + close, where the PPC Lead decides, takes scoring of predictions written inside Recovery Mode, since rule R3 tightens the P5 human-review thresholds to bid moves above 15% and budget moves above \$25 per day.
Cadence	Weekly Monday inside the cascade, first step, ahead of any new verdict: P2 to P7. At staging, per prediction: P1. Wednesday, on the T+1 diff: the P6 EXECUTION branch for the prior day's deployments. Monthly: P8, P9, P10, P12. Quarterly: P11. Within one week of a Recovery Mode exit: P14 per rule V7. Per scope request: P13.
Outputs	Drafted scores with cause tags and evidence lines queued for human confirmation on JUDGMENT rows; the two-instrument write; the debt count and its dollars; the rule-review docket with denominators; the owner calibration reads with their exclusion sets stated; the confidence-band table; the override fold-in; the agent comparison; F1 to F13 alarms, of which F2, F7, F9 and F13 are computable without human reading.
Training corpus role	Per v4.0 Section 14 rule V8, the scored decision log, override reasons and worked records are the training corpus for the AI skill files. This SOP is therefore both a consumer of agent drafts and the producer of the data those drafts are trained on, which is why P13 exists as a controlled comparison rather than as a rate on a dashboard.

10. Version Governance

10.1 Review triggers specific to this component

- Any change to the T6 Decision Log column set or to the Account Rollup tab A3 formula block. P7 writes literal strings that A3 counts with COUNTIFS; a column rename or a formula edit breaks scoring silently. This SOP increments before the next Monday cascade runs.
- Any change to SOP-24's confound classes. SOP-24 v2.0 Section 1.2 closes the list at seven cause classes and P3 step 4 publishes five confound classes; SOP-24 Section 10 states that a new confirmed class increments that SOP and propagates to its three named consumers, of which this is one.
- Landing of any AM item in Section 10.3 into v4.1. A tolerance band, a VOID definition or a debt threshold arriving in v4.0 replaces a stated gap here and changes P5, P8 or P9 directly.
- Any change to the closed verdict vocabulary. Section 5.2 binds all 25 KDR strings to a predicted-metric contract; a new or retired string leaves a verdict with no contract.
- Any change to E7's composition or to rule M10's step sizing. P11 exists to propose those changes and P1 step 4 enforces one of them.

6. A T7 manifest row moving to PRESENT where Section 3 currently records it as blocking. Rows A1, B3, D8 and D9 each currently constrain a scoring class; each arrival removes a stated limitation.
7. Quarterly regardless, alongside the rule V6 confidence calibration, because this SOP's own procedures are the instrument that grades every other SOP and an ungraded grader is the one gap the loop cannot detect.

10.2 Change note, v1.0 to v2.0

Supersedes v1.0 in full. Procedures 0 to 14. Failure modes 5 to 13, all five v1.0 modes carried with their incidents. Worked examples 0 to 4. Elements 2 of 12 to 12 of 12. Rubric corrected from four values to five: WRONG-DIRECTION restored from Strategic Doctrine v1.0 Section 13, the T6 column definition and the A3 formula block; VOID retained and its instrument gap logged. The confound citation corrected from SOP-24 P1 to SOP-24 v2.0 P3, since P1 at v2.0 is loss dating. The v1.0 statement that tolerances live in v4.0 corrected: v4.0 uses the word 7 times and quantifies it never, which is AM-1. Three v1.0 outcome strings retired per Section 5.5. Eight canonical edge cases carried forward, two re-scored, one added.

10.3 Amendment register: eleven items for v4.1

ID	The number this SOP needs and v4.0 does not contain	Where the gap bites, and the evidence travelling with it
AM-1	Tolerance bands per prediction class: the permitted magnitude deviation and date deviation that separate HIT from PARTIAL	P5 steps 4 and 5. v4.0 uses "tolerance" 7 times across Section 3 families F1 and F16, Section 4 rule X12, Section 10 rules M3 and M4, and Section 16, and quantifies it in none. Until it lands, P5 applies the prediction's own stated bound and returns unbounded predictions to their authors. Largest open item on this component
AM-2	VOID as a defined score value, and its treatment in the E8 denominator	Section 5.3, P7 step 5, P9 step 2, failure mode F9, WE-4. VOID appears zero times in v4.0. SOP-24 v2.0 P3 step 3 requires it. Neither T6 nor Account Rollup tab A3 carries a column for it
AM-3	The re-dating rule after a confound, and whether a suppressed interval extends a loss ceiling	P4 step 5 and WE-2, which prices one instance at \$388.80 against a \$650.00 ceiling, taking realized loss to 152.9% of ceiling. SOP-24 v2.0 logs the same gap from the diagnosis side at its AM-5; both references travel together
AM-4	The scoring-debt threshold: days past a review date at which a prediction is debt, and the count at which debt escalates	P8 step 5. v4.0 Section 14 rule V1 states that unscored actions do not accumulate and attaches no number. This SOP reports the count and asserts no trigger
AM-5	The VOID-rate ceiling per decider and decision class	P10 step 4 and failure mode F3. v1.0 stated it as "far above peers". No v4.0 band exists
AM-6	The minimum scored-prediction count for an owner-level calibration read	P10 step 5 and failure mode F11. v4.0 Section 14 rule V5 sets 10 or more at rule level and nothing at owner level
AM-7	An interim required-clicks correction for terms below the 8-week push-history bar in rule M2	WE-4. Rule M2's Rank Response Elasticity (E4) correction engages at 8 or more weeks; 6 of 8 non-HIT rows on S-R3 sat below that bar with no interim factor available
AM-8	The agent scope-promotion criterion: accuracy bar, minimum matched sample, and required margin over the human comparator	P13 step 4 and failure mode F12. v4.0 Section 14 rule V8 names the training loop and sets no bar
AM-9	A default review horizon per verdict class	P1 step 2 and Section 5.2. v4.0 Section 3 family F6 gives a 3-week default for ETA only, and Section 10 rule M10 a 1-week horizon under LOW confidence. Twenty-three of the 25 verdict strings have no stated default
AM-10	A scoring rule for trend-class predictions while four-period comparability is unavailable	Section 3 and failure mode F13. T7 row A1 reads MISSING and blocks all trend verdicts; v4.0 Section 3 family F7 defines the four trend classes and states no degraded-read procedure

ID	The number this SOP needs and v4.0 does not contain	Where the gap bites, and the evidence travelling with it
AM-11	Reconciliation of the closed verdict vocabulary counts, 18 in v4.0 Section 13.1 against 25 in the KDR Template v3.0 against 26 as cited by the Transition Brief v1.2 and the Ledger row for #8	Section 5.1 and 5.2, which bind a predicted-metric contract to each string. Already open on the Ledger row for #46 as "verdict-count discrepancy 18/25/26"; recorded here because a verdict with no contract is a prediction with no test

Owner: PPC Lead. Registered in the Supersession Map. Per v4.0 Section 14 rule V3, this document changes only by versioned amendment with a supersession note naming which rule was replaced, why, and the evidence.

11. Operator Ramp

Built from the Authority Matrix v2.0 Section 1 rows that name this component's decisions, one competence per row. The Matrix codes D as decides, E as executes, R as reviews, I as informed. Week 1 is observation and artifact familiarity, week 2 is supervised execution with the checkpoint holder reviewing, week 4 is unsupervised execution with the checkpoint holder auditing. The Transition Brief v1.2 already names two of these as live assignments: day 1 confirmation of the V5 scoring backlog with the PPC Lead, and execution of arrived-unscored reviews under supervision.

Authority Matrix row	Coding	Week 1: observe	Week 2: supervised	Week 4: unsupervised, audited
Weekly verdicts (established US)	PPC Lead D; incoming operator E within assigned scope; PPC Manager E; BM owner R	Sit the Monday P2 sweep read-only. Reproduce the worklist filter and the debt count from T6 and reconcile both to the PPC Lead's numbers. Name the eleven T6 columns and what each one is for	Run P2 to P7 on the assigned scope with the PPC Lead confirming every JUDGMENT tag before it files. Both instrument writes done by the operator; reconciliation checked by the Lead	Run the full pass on the assigned scope. Lead audits after filing: a sample of 5 rows re-walked through the P5 tree, and the T6 to A3 reconciliation checked. Gate: two consecutive weeks with arrived reviews scored on time and every miss root-caused in writing
Bulk deployments (#38)	PPC Lead R; incoming operator E; PPC Manager E	Read one Wednesday T+1 diff end to end alongside the deployment owner. Identify which rows in that diff would change a cause tag	Run the P6 EXECUTION branch on the prior day's deployments. Produce the EXECUTION-tagged handover to the #38 session review	Own the Wednesday EXECUTION pass. Audit is the #38 session record: every EXECUTION tag traceable to a named diff row, and no JUDGMENT tag filed on a window with no diff
Incrementality tests	Incoming operator D and E; PPC Lead R	Read SOP-24 v2.0 P3 and list its five confound classes with the v4.0 gate or rule that names each	Apply P4 to every worklist row touched by a T5 event, computing the gate G10 guard end explicitly. Lead checks each computed date	Own P4. Audit is the void register: every VOID carries its class and its computed guard end, and every void shared with a live incrementality test was handed to #35
v4.0 threshold changes	CEO D; PPC Lead E and R; all others I	Read the Section 10.3 amendment register and, for three items, state which procedure fails without the number	Draft one P14 amendment proposal in the rule V3 form: which rule, why, evidence, with the scored rows attached. Lead rewrites as needed	Draft the monthly proposal set unaided. Audit: does each proposal carry a rule ID, a scored count, a hit rate, an excluded-VOID count and a cause distribution. Decision stays with the CEO throughout; the operator never holds this D
Automation rules (#37)	PPC Lead D; incoming operator E and R; PPC Manager E	Read the #37 rule change log for one month and identify which T6 rows it would reclassify from JUDGMENT to an external change	Build one P13 comparison on matched classes with both denominators stated. Lead checks the matching before the numbers are read	Produce the comparison unaided; the scope decision remains the PPC Lead's. Audit: are the two populations matched on decision class, rule ID mix and window
Code-red declaration + close	PPC Lead D; incoming operator E; PPC Manager E	Read v4.0 Section 11 rules R1 to R6 and state how rule R3 changes the P5 human-review lines during the mode	Score a set of predictions written inside a Recovery window under supervision, applying the tightened lines	Score Recovery-window predictions unaided and file the rule V7 post-mortem input within one week of exit. Declaration and close remain the Lead's

Authority Matrix row	Coding	Week 1: observe	Week 2: supervised	Week 4: unsupervised, audited
Launch-stage verdicts + batches and Canada weekly block	Launch/CA operator D and E; PPC Lead R; BM owner R or I	Observe only. These scopes are not the incoming operator's under the Matrix	Observe only	Boundary competence: state why a launch-scope prediction is scored by the launch operator and what this SOP's P1 standard obliges regardless of scope

Vacancy. Authority Matrix v2.0 standing rule 5: an unfilled role's D reverts to the PPC Lead and its E to the PPC Manager until filled. A ramp interrupted by a vacancy pauses at the last completed week rather than skipping forward, because week 4 audit evidence cannot be produced by the auditor.

11.1 Live-defense questions

Ten questions. Each tests the reasoning behind a threshold or a rule ordering. None can be answered by reciting a number, and a correct number with the wrong reason is a fail.

1. Why does the P5 tree test VOID before it tests direction, rather than testing direction first and voiding afterwards? What would the account's hit rates look like if the order were reversed?
2. A push reached its target rank three days early, and the daily rank series shows the two competitors above it both went out of stock in week one. What is the score, and what is the argument for the other answer?
3. Why is EXECUTION excluded from the owner calibration read but included in the rule-level read? What would go wrong in each direction if that were reversed?
4. v4.0 rule V5 requires both below 50% accuracy and 10 or more scored predictions. Explain why a rule at 20% on 4 predictions is deliberately allowed to keep running, and what would break if it were not.
5. A confound voided a push's day-14 read and the guard end sits at day 23. The operator asks whether the loss ceiling extends by 9 days. Answer, and then explain why the honest answer and the operationally convenient answer differ.
6. Why does P6 forbid a JUDGMENT tag when no T+1 diff exists for the window, rather than defaulting to JUDGMENT as the residual? What does the default cost over a quarter?
7. An owner's VOID rate is triple every peer on the same decision classes. There is no v4.0 ceiling. What do you do, what do you not do, and what evidence would move you either way?
8. PARTIAL rows on one rule are 80% overshoot on magnitude. Why is that a correction to an assumption rather than to a person, and which document receives it?
9. Why does this SOP write the score into two instruments rather than one, given that the Account Rollup pulls from the product workbook? What class of error does the second write catch?
10. An agent's drafts score a higher hit rate than the human operator's over the same quarter. Give three reasons that comparison could be true and still be worthless.

12. Interface Contracts

Sixteen handoffs. Every row states the non-arrival consequence, because that is the clause the corpus has been missing and because this component sits downstream of every SOP that writes a prediction.

ID	Sending owner	Receiving owner	Artifact	Timing	What happens when it does not arrive
I-1	Every campaign-type SOP, via its own Section 4	This SOP, P1	The Master Action Record fields expected_outcome and review_date, written into T6 Decision Log	At staging, before the row leaves for deployment	The change is undeployable: v4.0 Section 12 rule P10 bars deployment of a verdict without a J-complete record, and Section 13.3 rule J1 rejects it at the deploy step. The row returns to its author as failure mode F5

ID	Sending owner	Receiving owner	Artifact	Timing	What happens when it does not arrive
I-2	#38	This SOP, P6	The T+1 verification diff for the deploying session, with zero unexplained mismatches	Wednesday, the day after deployment	EXECUTION cannot be excluded. P6 step 2 withholds the JUDGMENT tag, the row files with the cause tag open, and P10 excludes it from the owner read. Accuracy denominators shrink until the diff lands
I-3	SOP-24, P3	This SOP, P4	The void declaration with its confound class from the five SOP-24 P3 step 4 publishes	At the read, same scoring pass	The read scores as if clean, which is failure mode F8 and escalates on first occurrence. The consumer cannot detect contamination it was not told about
I-4	Data Ops	This SOP, P2	Refreshed T7 Manifest with the completeness rollup and row-level PRESENT / STALE / MISSING states	Monday 06:00 refresh, packet ready 09:00	Gate G8 applies: at under 85% of gate-relevant fields the pass produces flags rather than scores, and states what it could not see
I-5	Brand Management, via T5	This SOP, P4	T5 Product Log entries with Trend-Exclusion Flag and event type	Same day as the event	WORLD cannot be evidenced and no confound can be found. Non-HITs default toward JUDGMENT and the decider is charged for an unlogged world event. Data Architecture v1.0 Section 9 assigns T5 completeness to Brand Management and the reading obligation to PPC
I-6	#35	This SOP, P3	The incrementality test result and its total-capture read	At test close	INCREMENTALITY LADDER predictions are unscorable on their stated metric and extend under gate G2 rather than scoring MISS
I-7	#37	This SOP, P6	The rule change log covering the read window	Weekly, with the #37 diff	An automation-driven change cannot be separated from an execution error or from a decider's move. The cause tag files as open per I-2's rule
I-8	SOP-26	This SOP, P14	The Recovery Mode exit log with its recovered weekly dollars	At exit	The rule V7 one-week post-mortem clock has no start date. The amendment proposals miss the window and file late, which is the failure rule V7 exists to prevent
I-9	Data Ops	Account Rollup tab A3	New T6 rows pasted into the A3 DETAIL block	Mondays, after the product cascades close	A3 summary counts run on a stale detail block and every rule-level ratio is computed on a short denominator. Failure mode F7

ID	Sending owner	Receiving owner	Artifact	Timing	What happens when it does not arrive
I-10	This SOP, P7	T6 Decision Log and Account Rollup tab A3	The filed score: literal value string, cause tag, evidence line, both numbers	Same Monday pass	Rule V1 breaks at its own first step: unscored actions accumulate, the next cycle opens on an unclosed prior cycle, and scoring debt compounds under failure mode F2
I-11	This SOP, P6	#38 session review	EXECUTION-tagged rows with their diff references and the dollar consequence	With the monthly batch, or same week where the dollars warrant	The deployment path receives no signal, the same defect recurs, and the owning rule's accuracy carries a penalty that belongs elsewhere. WE-3 prices one instance at \$140.80 in a single window
I-12	This SOP, P10	SOP-23, SOP-24, #28	Systematic PARTIAL-magnitude findings by rule, with direction of error	Monthly	The assumption stays wrong and every future prediction on that rule inherits it. The account's recorded case is velocity assumptions feeding push worksheets
I-13	This SOP, P9 and P14	The v4.0 threshold owner, via the Authority Matrix row v4.0 threshold changes	The rule-review docket and the amendment proposals in the rule V3 form	Monthly, and within one week of a Recovery exit	v4.0 stops learning from its own evidence. Rule V5's stated purpose, that the framework is subject to its own evidence standard, goes unmet and failing rules keep firing
I-14	This SOP, P13	The skills program under the 10-80-10 method (#7)	The scored corpus and the matched human-versus-agent comparison with both denominators	Per scope request	Agent scope either freezes or widens on an uncontrolled rate, which is failure mode F12. Per v4.0 Section 14 rule V8 the scored log is the training corpus, so a gap here degrades the drafts as well as the promotion decision
I-15	This SOP, Element 11	#44 Operator onboarding and certification	The ramp table and the ten live-defense questions	On request; #44 is NOT BUILT, due 2026-09-25	No certification bar exists for this component and deployment authority is granted on impression. Per the Ledger row for #44, it consumes Element 11 from all SOPs and no document exists yet; this SOP asserts no pass bar of its own
I-16	This SOP, P12	Threshold, scenario and skill owners	The validated override set and the override-rate trend	Monthly	Rule V2's fold-in does not happen and the override rate loses its meaning as the system's quality metric. This is the fold-in item recorded as pending on the Ledger row for this component; P12 closes it

13. Reported Findings Beyond the Pre-Declared List

Anti-drift rule 5 requires a session to report conflicts it finds beyond those already registered. Per the Ledger accretion rule, conflict identifiers and Ledger versions are assigned at close-out and never inside a session, so these are reported as findings without IDs.

1. The score vocabulary is defined three ways across four documents. SOP-06 v1.0 Section 3 uses HIT, PARTIAL, MISS, VOID. The Strategic Doctrine v1.0 Section 13, the Data Architecture v1.0 Section 9 tab T6 column definition, and the Account Rollup Workbook v1.0 tab A3 all use HIT, PARTIAL, MISS, WRONG-DIRECTION. SOP-24 v2.0 P3 requires VOID. v4.0 contains none of the five strings. Resolved here by carrying five values and logging AM-2; reported for registration because it touches four documents and two instruments.
2. The Ledger row for this component cites SOP-24 P1 for the confound cause classes, and P1 is now loss dating. At SOP-24 v2.0 the confound and void determination is P3. The Ledger row for #35 already records the same correction applied to that component. This row's Must not restate cell needs the same edit and it is carried in the returned row below.
3. The attached Tier 2 copy of SOP-38 is v1.0 and the Ledger row for #38 names v2.0 as superseding it in full. This build cited #38 by Ledger row and one clause throughout and used no v1.0 procedure number. Reported so the Tier 2 routing list for this component is corrected before the next session.
4. Pack state is asserted three ways in this session's inputs. The version note records the Ledger at v2.5 with Wave 4 closed. The attached Ledger header reads v3.3, dated 2026-08-04, with Wave 1 returned and the Wave 2 gate met. The attached Ledger's own rows carry Wave 5A and Wave 5B returns for #20, #24, #28, #32, #33, #35, #37 and #42, and its rows reference conflicts through C23 while its Section 3 still lists seven. SOP-24 v2.0 Section 13 reports the same class of divergence from its own session, recording the Ledger at v2.8. This session treated the rows as current, minted no identifier, and reports the divergence.
5. v4.0 has no prediction-scoring section. Rule V5 sets a rule-level accuracy line and family F15 defines metric E8, and between them they specify a ratio without specifying how any single prediction is graded. Every scoring rule in the corpus lives in instruments and in this SOP. This is not a contradiction between documents and so is not a conflict; it is the structural reason nine of the eleven items in Section 10.3 exist, and it is reported so the v4.1 brief can treat them as one section rather than eleven scattered numbers.

Companions: Master PPC Decision Criteria System v4.0 (Section 3 families F1, F4, F6, F7, F8, F15, F16; Section 4 rules X1, X4, X10, X12; Section 5 gates G1, G2, G3, G6, G8, G9, G10, G12; Section 6 rules O2, O5; Section 7 rule D10; Section 8.4 rules S-R3, S-R5; Section 8.7; Section 8.9; Section 8.12 rule S-E5; Section 9 rules L2, L5, L6, L8, L9; Section 10 rules M2, M3, M4, M8, M10; Section 11 rules T1, T6, R3; Section 12 rules P5, P6, P10; Section 13.1 the closed verdict vocabulary; Section 13.2 the Master Action Record v2 output contract; Section 13.3 rules J1, J2, J7; Section 14 rules V1 to V8) · Keyword Decision Record Template v3.0, sheet "Bands & Verdicts Reference" (all 25 verdict strings) and sheet "Read Me" (the justification grammar and review tiers) · PPC Strategic Doctrine v1.0 principle 5 and Section 13 (the score vocabulary and the wrong-direction root-cause obligation) · PPC Authority Matrix v2.0 Section 1 (Elements 9 and 11, by row) and standing rules 3 and 5 · Master PPC Data Architecture v1.0 Sections 9, 10 and 11 (tabs T5, T6, T7, the KDR cycle file, Account Rollup tab A3) · PPC Command Workbook TEMPLATE v1.0 (T5, T6, T7 column sets) · Account Rollup Workbook v1.0 (tab A3 blocks and formulas, READ ME refresh flow) · PPC Economics Worked-Example Set v1.0 case E3 (the WE-1 and WE-2 sizing base) · SOP-24 Re-Ranking and Recovery v2.0 Sections 1.2 and P3 (the cause classes and the void determination) · SOP-12 Exact Campaign Operations v2.0 (the format exemplar) · #38 staging and T+1 verification, #35 incrementality, #37 rule register, #26 Recovery Mode, #23 push sizing, #28 ladder, #7 the 10-80-10 method, #44 onboarding, all by Ledger row and one clause · PPC Cross-Reference Ledger (component 6 row; conflicts C1 to C7 as declared, plus the findings above).

Inspiratek & Ecotero LLC | Confidential | Page 14 of 14