

SOP-13: Phrase Campaign Operations and the Per-Family Pilot

Version 2.0 · 2026-07-28 · Owner: PPC Lead · Supersedes SOP-13 v1.0 (2026-07-24) in full · Thresholds live in the Master PPC Decision Criteria System only; this SOP carries zero local thresholds. Threshold authority for this build: the uploaded Master PPC Decision Criteria System v3.0, operating as the v4.0 authority per the 2026-07-28 session ruling on PPC-DR-02; every citation below quotes that file by its actual section heading.

1. Verdict and Scope

VERDICT: Phrase is qualified discovery on proven roots, and it is the account’s one named structural gap at the match-type layer: the Criteria System’s F8 Structure and Registry family states that a primary syntax at 0% phrase is a standing finding, scenario S-H8 (Section 8.9, Discovery Scenarios) fires on a phrase layer at \$0 spend on a product with proven exact roots, the M6 structure targets (Section 10, Sizing Mathematics) set a 5-10% phrase spend destination, and the Data Architecture v1.0 Section 11 (A4 Structure targets) records the Bamboo 4pc product-wide \$0 phrase as a live one-cell finding. Version 2.0 rebuilds v1.0 from 4 procedures to 9, from 1 worked example to 3 (adding the mandatory edge case that goes wrong), from 5 failure modes to 9, and adds Elements 11 (operator ramp) and 12 (interface contracts), which were absent from all 25 shipped SOPs including the SOP-12 v2.0 exemplar. This SOP governs the full operating life of the two phrase surfaces: family phrase campaign creation, the dated per-family pilot that closes the S-H8 gap, the weekly operating pass, the stem-clustering read, the phrase-broad dedup wall, CAP SYNTAX receipt, event-window overlay, pause and retirement, and the stem graduation handoff.

1.1 The Two Phrase Surfaces

Every phrase term in the account lives on exactly one of two surfaces, and the surface determines its review depth and its exit path. **Surface 1 - the family phrase campaign:** one campaign per primary syntax family, Objective = Discovery, seeded with the family root as phrase plus validated modifier stems, budget-capped by design, never straddling two families (the family is the cap-enforcement unit per rule P4, Section 12, Precedence and Conflict Resolution). **Surface 2 - standing long-tail stem coverage:** validated stems whose per-term volume never clears the graduation bar remain in the family phrase campaign at a profitable bid as a deliberate end state under the LEAVE TO HARVEST verdict; graduating 20 micro-terms one by one is noise, and the stem held in phrase is the correct owner of that demand. Movement off Surface 2 happens only through the P4 stem routes; it is never a judgment call made mid-pass.

1.2 Session Rulings and Conflicts Declared Before Writing (anti-drift rule 5)

Four rulings and conflicts were declared before any content was written, so this build does not silently pick a side. **(a) PPC-DR-02, ruled this session:** the uploaded Criteria System v3.0 is the threshold authority the corpus cites as v4.0; this SOP cites it by its printed section headings. **(b) PPC-DR-04, provisional:** every verdict string in Section 5 below is quoted verbatim from the Keyword Decision Record Template v3.0, sheet “Bands and Verdicts Reference” (25 strings), the source the SOP Build Kit Element 5 authorizes alongside the Criteria System. The Criteria System’s own Section 13.1 (“The Closed Verdict Vocabulary”) lists 18 differently worded actions; the delta is logged in Section 13 of this SOP as amendment item AM-1 and the Section 5 mapping re-issues on the PPC-DR-04 ratification. **(c) Ledger conflict C3:** component #31 carries no procedural document; wherever this SOP receives budget law it references the instrument (T1 S3 spend-governance columns, the M5 queue per Criteria System Section 10) and never a #31 procedure number. **(d) Citation drift:** the uploaded Criteria System contains no Section 4.2 naming convention; per its own Section 17.2 supersession note the naming convention remains authoritative in the v1.0 Criteria System, which is outside this pack, so naming below references SOP-12 v2.0 Appendix A, the uploaded restatement of the governing convention.

1.3 Boundaries With Adjacent Components

Adjacent component	The seam, stated so it is owned
Criteria System (v3.0 operating as v4.0)	Decides every threshold, band, gate, scenario, verdict, and lever limit. This SOP executes; a threshold appearing here instead of there is a defect (SOP-12 v2.0 Section 1.3 rule, applied unchanged).
#30 Harvest and graduation	Decides which stems graduate. P4 routes candidates to #30; the GRADUATE FROM DISCOVERY verdict returns from #30 and SOP-12 stands the exact up. This SOP never issues a graduation.

Adjacent component	The seam, stated so it is owned
#29 Negation architecture	Owns negation doctrine, wall rules, and the wall audit. This SOP builds the prescribed walls at every structure event in the same working session and hands breach findings (F2) to the #29 wall audit.
SOP-14 Broad operations	Owns the shared weekly triage machinery (SOP-14 P3: the three candidate routes, the yield read, the head-term check, staging identical); this SOP runs it by reference and adds only the phrase-specific P4 and P5 reads. A change to SOP-14 P3 increments both SOPs (Section 10).
SOP-12 Exact operations	Owns exact standup. Stem graduations and every ADD EXACT + RESTRUCTURE found in phrase execute in SOP-12 P1/P2 with #29 walls; this SOP builds only the phrase-side steering negative.
#31 Budget and envelope (instrument)	Budget deltas, CAP SYNTAX, and pilot funding law arrive from the instrument layer (T1 S3 columns, the M5 queue per Criteria System Section 10). This SOP receives deltas and never originates them. Per conflict C3 no #31 procedure number is cited anywhere in this document.
#32 Deal and event windows	Owns the event plan. This SOP freezes phrase structure changes inside armed G12 windows (P7) and executes only what the event plan names.
#33 Sibling arbitration	Owns family term ownership. Sibling-owned roots are pre-loaded negatives in P1 per the declared S-F3 owner; a phrase campaign never seeds a root a sibling owns.
#38 Bulk deployment mechanics	Owns pull, staging format, upload, and T+1 verification. Every change this SOP makes exists first as a staged bulk set in #38 format.

2. Trigger Conditions

Per the Depth Standard, each trigger names its detecting artifact and the maximum allowed lag from signal to procedure start.

Trigger	Detecting artifact	Max lag	Procedure
Pilot standup	T2 K9 verdict row FUND TEST (UNFUNDED PRIMARY) with a dated release from the M5 queue (instrument)	Same day as the dated release; Tuesday default otherwise per #38	P2
Weekly cadence	Monday audit cascade reaches T2; search-term feed pull date = current week in T7	Same session as the Monday pass; deployment Tuesday per #38	P3 then P4 then P5
Pilot review gates	Paired T3 yield rows at day 14 and day 28 of the pilot window	Gate decision within 1 working day of the read landing	P2 steps 8-9
Family qualification	T2 K8 column "# Exact" at 1 or more on the family root with Validation State clean, plus the T1 family row live	P1 enters the next weekly verdict queue	P1
Stem graduation receipt	GRADUATE FROM DISCOVERY record returned by #30	Phrase-side steering negative built in the same working session as the SOP-12 standup (same day)	P9
CAP SYNTAX receipt	T1 row Cap Status = BREACH with the staged delta from the instrument layer	Budget rows staged in the same weekly session	P6
Event window arms	T2 "Event Mode" column set from the #32 event plan (gate G12, Criteria System Section 5)	Freeze effective the day the window arms	P7
WAS breach on a phrase row	T3 column "WAS %" above the Discovery limit (Criteria System Section 3, F11 Waste and Negation)	Negation pass in the same weekly session	P3 route via SOP-14 machinery; wall check P5

3. Preconditions and Gates

Each gate cites its T7 manifest ID and states the observable that proves it passed. A missing or stale field group suppresses the dependent execution class; the pass records what it could not see and continues on the rows it can (gate G8, Criteria System Section 5).

Manifest ID (T7 row)	Requirement	Observable that proves the pass
A2 Deployment state in-row	Current bulk state for every phrase campaign	T7 row A2 = PRESENT; T3 rows carry State, Daily Budget \$, Bid / AG Default \$ from the current-week BULK-US pull
A6 Objective + Expected CPC	Every phrase row tagged Discovery; a seed bid reference exists	T3 Objective Tag = Discovery on 100% of phrase rows. T7 A6 is PARTIAL: the account reference Expected CPC \$2.80 stands in until per-term values land, and every seed bid carries the PROVISIONAL note

Manifest ID (T7 row)	Requirement	Observable that proves the pass
B5 Coverage counts	Registry current per match type	T2 columns # Exact / # Phrase / # Broad / # Auto reconcile to the T3 registry rows for the family
B8 Economics chain	T0 named cells live (AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC)	T0 zone C populated; while T7 D1 reads PARTIAL (blended margin), the G9 PROVISIONAL tag rides every ceiling-referencing action (Criteria System Section 5, G9)
Search-term feed	SP + SB search term reports and STIS landed in the Raw Pull Pack (Data Architecture Section 12)	Pull date on the Raw Pull Pack sheets = current week. Feed absent: P3, P4, and P5 are suppressed for the week and the suppression is recorded in the pass output
Root validation	The family root carries validated exact structure before any phrase campaign stands up	T2 K8 # Exact at 1 or more on the root, Validation State clean, no G1 quarantine flag on the row
Gates in force	G2 sufficiency, G3 truncation, G12 event, and P4 syntax caps govern every read below	Bands per the KDR Bands Reference: clicks under 15 = no bid read; in-budget under 70% = every rate marked truncated; Event Mode clear or the plan names the change; T1 Cap Status inside cap

4. Step-Level Procedures

P1. Create a family phrase campaign (standing state) - 45-75 minutes per family

1. Confirm the family qualifies: T2 K8 # Exact at 1 or more on the family root with Validation State clean. An unproven root is broad's job (SOP-14); a phrase campaign on an unvalidated root is misfiled broad and is the F3 condition.
2. Confirm no family phrase campaign exists: filter T3 column Type = SP-P on the family's syntax tag. If one exists, extend it; a second phrase campaign on the same family is the F7-adjacent duplicate condition and routes to CONSOLIDATE DUPLICATES.
3. Confirm gates: G12 clear (T2 Event Mode blank or the event plan names this creation), G3 not relevant at creation, and the #33 check: no seed term sits on a root a sibling owns per the declared S-F3 assignment.
4. Name the campaign per SOP-12 v2.0 Appendix A with the family root in the keyword position and the objective segment = Disc; a name whose objective segment does not parse is the SOP-12 F6 class and renames ride alone.
5. Stage the campaign as bulk rows on the Sponsored Products Campaigns sheet per #38. Column values: Entity = Campaign then Ad Group then Keyword rows; Operation = Create; State = enabled; Targeting Type = Manual; Bidding Strategy = Dynamic bids - down only (Doctrine v1.0 Section 6.3: down-only is the default outside pushes); Daily Budget = the family discovery allocation share received from the instrument layer, capped by design.
6. One ad group. Seed the keyword rows: Keyword Text = the family root, Match Type = phrase, plus each validated modifier stem as its own phrase row where stem volume justifies dedicated coverage. Bid = the verdict-carried value; absent a per-term Expected CPC the account reference \$2.80 seeds it with the PROVISIONAL note in the reason string (manifest A6).
7. Pre-load negatives per #29 in the same staged set: Match Type = negativeExact rows for every graduated exact the family holds, brand terms, sibling-owned roots, and the taxonomy-known irrelevant classes (Criteria System Section 9, L5 pre-loaded mode).
8. Build the control-level dedup wall: Match Type = negativeExact rows for every phrase seed term, placed in the family BROAD campaign, in the same staged set, so broad explores beyond what phrase now covers and the two surfaces never buy the same query at two control levels. Record the wall list in the T3 row (Negation State column) so the #29 wall audit has its baseline. Skipping this step is the F2 root cause.
9. Update the registry: T2 # Phrase increments on every seeded term; the new T3 row carries Objective Tag, owner, and wall list; T1 column MT Cov Phrase % picks the campaign up through the syntax tag.
10. Log the creation in T6 with the prediction the driving verdict carries: metric, direction, magnitude, horizon, review date (rule V1, Criteria System Section 14). A creation without a logged prediction is incomplete and does not deploy.
11. Deployment routes through #38: Tuesday default; same day when the driving verdict carries a dated release. Same-day deployments still receive the full T+1 verification pass.

P2. The pilot (the dated build that closes the S-H8 gap) - standup 60-90 minutes; gates 20-30 minutes each

1. Scope: one phrase campaign in each primary family with proven structure, strongest family first. A family sitting in a chronic CONVERSION diagnosis is excluded until the diagnosis clears (Criteria System Section 8.2, S-Y5): pushing discovery into a family that cannot convert is waste with a process name.
2. Funding is envelope-neutral: the pilot budget reallocates from the family's existing discovery share and never expands the envelope (rule P2, Criteria System Section 12). Write the arithmetic into the standup record as three numbers: dollars stepped out of the family broad budget, dollars drawn from the family's unspent discovery allocation, and the unchanged family discovery total. WE-1 shows the computed form.
3. Duration and gates: a dated window with reads at day 14 and day 28. The standup T6 row logs the prediction: expected validated-term yield per \$100 of spend versus the family broad baseline, with the day-28 review date. A pilot without the dated window and the prediction row does not deploy (F1).
4. Standup: execute P1 in full, including step 8, the dedup wall.
5. Instrumentation: the T3 rows for the pilot campaign and the family broad campaign both carry the yield read every pass (Harvest Yield column), so the day-28 comparison is native to the workbook, never assembled at review.
6. Weekly during the window: P3 through P5 run normally; a daily pacing glance on the capped budget rides the standard alarms; no scale steps inside the window.
7. Day-14 gate, health only, 3 checks: (1) spend pacing to cap (in-budget healthy per the KDR Bands Reference budget-state bands), (2) negation cadence live (negatives added in both weekly passes), (3) zero wall breaches from P5. No scale or kill decision at day 14 unless a stop-loss condition fires.
8. Day-28 decision, issued as a verdict: SCALE (extend P1 to the remaining qualified families; the pilot converts to standing state), HOLD (extend the window once, with the specific open question named), or RETIRE (PAUSED-PENDING with a re-entry condition; the yield evidence files to T6 and the S-H8 gap is declared closed-by-evidence rather than closed-by-build). The decision rule reads yield and walls first; ACOS governs only through the Discovery objective band (Criteria System Section 6, O6). The rule's numeric tolerance for "comparable CPC" is not in the Criteria System and is logged as amendment item AM-3, not invented here.
9. The day-28 decision is a scored V5 row against the standup prediction, and the gate artifact is the decision table with both campaigns' yield reads side by side.

P3. Weekly operating pass - 30-45 minutes per family, inside the Monday cascade

1. Enter T2 only after T0 and T1 have been read; a pass that skipped the levels above is invalid by construction (Data Architecture Section 2, rule 1).
2. Score last cycle first: read the T6 rows whose review dates land this cycle and record their V5 scores. Unscored actions do not accumulate (rule V1).
3. Pull the window's search-term data: Ads console > Measurement and Reporting > Sponsored ads reports > Create report > Sponsored Products > Search term, summary, date range aligned to the 14-day pass window, filtered to the family phrase and family broad campaigns. Working columns: Campaign Name, Ad Group Name, Targeting, Match Type, Customer Search Term, Impressions, Clicks, Spend, 7 Day Total Orders (#), 7 Day Conversion Rate, Cost Per Click (CPC).
4. Run the shared triage machinery of SOP-14 P3 by reference on the phrase campaigns: the three candidate routes, the yield read, the head-term check, staging and deployment identical. Nothing about that machinery is restated here; a change there increments both SOPs.
5. Then run P4 (stem clustering) and P5 (dedup wall audit), the two phrase-specific reads.
6. Stage every resulting change per #38 with the verdict and scenario ID in the reason string; no scenario ID, no deploy. Tuesday deploy after the BM huddle; T+1 verification; a landed-value mismatch is the SOP-12 F10 class, handled per #38.

P4. Stem-clustering read (phrase-specific) - 20-40 minutes

1. Normalize every served term in the window: lowercase, strip terminal pluralization, sort tokens into a signature. This is the same normalized-token signature the Criteria System Section 3 F8 twin-set detection uses, applied at stem grain.
2. Group terms by modifier stem around the family root: size stems, attribute stems, use-case stems. Pool clicks and orders at stem level for the read; per-term samples at tail volume are thin by construction (gate G2 governs sufficiency; stem-level pooling itself has no explicit G2 clause and is logged as amendment item AM-4).
3. Route each stem through the decision tree below. Route by stem, never by individual term: a stem graduates once (F4 guard).
4. Append the stem map to the campaign's T3 notes in the fixed line format: stem | terms | clicks | orders | status | route | date. The next pass reads clusters, not raw terms.

P4 decision tree (4 branches; the Depth Standard requires a tree at 3 or more):

Condition read on the pooled stem	Route	Executed by
Stem validated (orders at the S-H1 graduation bar, Criteria System Section 8.9) AND stem volume clears the dedicated-coverage read	Graduate the stem as its stem phrase: candidate routed to #30; on the returned GRADUATE FROM DISCOVERY, SOP-12 stands the exact up and P9 builds the phrase-side steering negative the same session	#30 decides; SOP-12 executes; SOP-13 P9 walls
Stem validated AND long-tail (volume below the dedicated-coverage read)	LEAVE TO HARVEST: the stem stays in phrase at a profitable bid as standing coverage, noted per stem in the map. Deliberate, logged, re-read every pass	This SOP, Surface 2
Stem invalid (clicks at the reactive-negation bar with zero orders, F11 reactive mode)	Negative phrase on the stem in the family phrase campaign, staged this session	This SOP + #29 doctrine
A served string is a priority head term of the product (oversight class) with # Exact = 0 on the term	ADD EXACT + RESTRUCTURE issued at T2 per scenario S-R12 (Section 8.4, "Priority term (500+ SV, primary syntax) with zero exact campaigns"); SOP-12 executes; the head term's clicks are excluded from the phrase yield read because they were never discovery	T2 verdict; SOP-12 executes; see WE-3 and F6

P5. Dedup wall audit (new in v2.0) - 10 minutes

1. Join the window's served-term sets on Customer Search Term: family phrase campaign versus family broad campaign, from the same P3 step-3 report pull.
2. Any intersection is the F2 condition: the same query bought at two control levels with the account's own money. Build the negative-exact wall in broad on every intersecting term the same day; record the addition in the T3 wall list so the #29 wall audit carries it.
3. Mark both campaigns' rate reads on the intersecting terms confounded for the window and exclude those terms from the yield read; the double-bought spend reclassifies as structure waste in the audit, mirroring the treatment the Criteria System Section 8.9 gives the S-H4 unsteered leak.
4. Root-cause line in T6 names which session skipped its wall step; a recurring source escalates per Section 7.

P6. CAP SYNTAX receipt (new in v2.0) - 10-15 minutes

1. Trigger: the T1 family row reads Cap Status = BREACH and the instrument layer stages the family delta (conflict C3: instrument, not a #31 procedure).
2. Discovery reduces first in a capped family, phrase and broad proportionally to their current discovery spend shares; the proportional split is executed here, the cap decision is not.
3. Stage the Daily Budget reductions as #38 rows with the CAP SYNTAX verdict and the T1 finding ID in the reason string; T6 entry predicts the family share returning inside cap by the next weekly read.

P7. Event-window overlay (new in v2.0)

1. Arm on G12: the #32 event plan dates enter T2 Event Mode; every phrase row on the product carries the flag through the window plus the post-event guard.
2. Inside the window, phrase structure changes freeze (no seed additions, no campaign creations, no stem graduation standups) unless the event plan names them. Reactive and wall negation stays live: routing a wrong query out is never suspended for a deal.
3. Event-week rates are excluded from trend and yield verdicts through the guard period (gate G10, Criteria System Section 5); the pass reads affected rows for gates only. A pilot window never starts inside an armed window and pauses its day-count across one (F9 guards the breach case).
4. Post-window: seeds, bids, and the pilot clock resume from pre-event values as a staged, logged set per the #32 reconciliation.

P8. Pause, retirement, and re-entry (new in v2.0)

1. PAUSED-PENDING (REASON + RE-ENTRY) always carries both fields in T2 and T3; a pause without both is incomplete and returns to the pass.
2. RETIRE from the pilot (P2 day-28 path) pauses the campaign, dispositions every standing stem in the map (each either follows the family broad as ordinary discovery or negates), and states the re-entry condition, typically a family diagnosis change or a demand-signal change on the root.
3. Paused rows appear on the weekly pass as a re-entry check only: the pass tests the stated criterion and either re-enters or holds; it never re-litigates the pause.
4. Archive only past the governance period with no re-entry met; archiving is human-mandatory per the Authority Matrix and the T6 archive entry closes the campaign's history with its final stem map.

P9. Stem graduation handoff (new in v2.0) - same session as the SOP-12 standup

1. On the returned GRADUATE FROM DISCOVERY record from #30, hand SOP-12 the graduation packet: the stem phrase string, discovery CPC, pooled click and order evidence, and the served-term list under the stem.
2. Build the phrase-side steering negative in the same working session the exact goes live: negativeExact on the graduated string in the family phrase campaign, recorded in the T3 wall list. The graduated term never buys itself twice at two control levels (Criteria System Section 3, F8: the GRADUATED-UNSTEERED state is the leak this step closes from the phrase side).
3. T6 entry: the handoff line references the SOP-12 creation prediction rather than logging a second one; one graduation, one scored prediction.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds, and scenario logic live in the Criteria System. Every string below is quoted verbatim from the Keyword Decision Record Template v3.0, sheet "Bands and Verdicts Reference" (the closed vocabulary, 25 actions), the source the SOP Build Kit Element 5 authorizes; the delta against Criteria System Section 13.1 is amendment item AM-1 and this table re-issues on the PPC-DR-04 ratification. The table maps what this SOP does when each verdict lands on a phrase-relevant row; where execution belongs elsewhere, the boundary is stated so no verdict falls into a seam.

Verdict (closed vocabulary, verbatim)	Phrase execution
QUARANTINE: DATA REFRESH	No changes deploy on the row; the refresh task carries the staged set to the next clean cycle (gate G1).
UNLOCK TRAFFIC	Bid steps up per the L2 unlock mechanics (Criteria System Section 9) on the named phrase row; no rate judgments until G2 sufficiency.
FIX BUDGET (TRUNCATED)	Budget lever only per gate G3; bid reads resume the next untruncated day.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	No changes except budget assurance; the capped test runs to its click quota.
FUND TEST (UNFUNDED PRIMARY)	The pilot standup itself enters through this verdict (P2), envelope-neutral funding stated in the record.
RANK PUSH: FUND SIZED PUSH	No phrase execution; pushes run on exact (SOP-12, SOP-23 scope). A push verdict citing a phrase row is a routing defect reported to T2.
RECOVERY PUSH	No phrase execution; owned by the recovery scope (SOP-24 by number). Same routing-defect rule.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	No phrase execution; the required-clicks chain is a rank-objective instrument (Criteria System Section 10, M2).
FIX CONVERSION FIRST; DEFER PUSH	Phrase seeds on the affected family hold at maintenance; no new seeds while the family diagnosis routes to its fix owner; pilot exclusion per P2 step 1.
PLACEMENT FIX FIRST	Bids hold; placement programs belong to #28; this SOP deploys only campaign-level rows staged back.
RE-POINT VARIATION	No phrase-specific mechanics; the advertised-SKU edit executes per SOP-12 P10 pattern on the family campaign where it hosts the affected variation.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	No phrase execution; tapers live on exact (SOP-12 P4).
HOLD; DEFEND LIGHTLY AT RANK	No change this cycle on any phrase instance of the term.
DEFEND AT FLOOR	No phrase execution; floors are exact-side records.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	The bid half executes on the named phrase row in one staged step; the CAP SYNTAX half arrives from T1 and executes per P6.
INCREMENTALITY LADDER	No phrase execution; measurement design belongs to #35, ladders to #28.

Verdict (closed vocabulary, verbatim)	Phrase execution
ADD EXACT + RESTRUCTURE	The P4 branch-4 route: SOP-12 stands the exact up; this SOP builds the phrase-side steering negative same session and excludes the term from the yield read (WE-3).
CONSOLIDATE DUPLICATES	Two phrase campaigns on one family consolidate to the survivor per the SOP-12 P7 precedence, executed there; this SOP flags the condition (F7-adjacent).
GRADUATE FROM DISCOVERY	Receipt from #30: P9 in full, same session as the SOP-12 standup. Stems graduate as stem phrases per P4, never term by term.
NEGATE (PRE-LOAD / REACTIVE / STEERING)	Owned by #29 doctrine. This SOP executes the three states structurally: pre-load at P1 step 7, reactive at the P3/SOP-14 routes and P4 branch 3, steering at P1 step 8 (dedup wall) and P9 step 2.
ENTER SBV ARBITRAGE	No phrase execution; the SP-exact step-down belongs to SOP-12, the SBV build to #17-19.
OPEN / ROTATE / EXIT CONQUEST	No phrase execution; owned by #16.
STOP-LOSS: CONCEDE OR DEFER	The pilot day-28 RETIRE path, or any dated phrase test at expiry: P8 steps 1-2 with the yield evidence filed.
PAUSED-PENDING (REASON + RE-ENTRY)	P8 steps 1-3.
LEAVE TO HARVEST	The P4 branch-2 standing state: validated long-tail stems stay in phrase at a profitable bid, deliberate and noted per stem.

6. Worked Examples (illustrative values, internally consistent; live values come from T0 and T2)

WE-1 - The clean case: pilot standup, both gates, and the stem read (family “Bamboo”)

Setup. Family broad baseline yield: 1.27 validated terms per \$100 of discovery spend (the SOP-14 example read, carried on the paired T3 row). Pilot campaign seeded per P1: “bamboo sheets” as phrase plus the validated stems “deep pocket”, “split king”, and “cooling” as phrase rows. Funding, envelope-neutral per P2 step 2, written as three numbers: \$11/day stepped out of the family broad budget (\$25/day to \$14/day), \$7/day drawn from the family’s unspent discovery allocation, pilot budget \$18/day; the family discovery total moves from \$25/day spent to \$32/day spent inside the unchanged allocation. Standup T6 prediction: pilot yield at or above 1.5 validated terms per \$100 by day 28, versus the 1.27 baseline.

Day-14 health gate (P2 step 7). Spend \$247 of the \$252 cap-to-date = 98.0% pacing, in-budget healthy; negatives added in both weekly passes (cadence live); P5 join returns an empty intersection (zero wall breaches). All 3 checks pass; proceed, no scale or kill read.

Day-28 read (P2 step 8, with P4 applied). Spend \$498 of the \$504 cap = 98.8%; 259 clicks at CPC \$1.92; 13 orders, CVR 5.0%; 78 distinct served terms. Stem map: “split king” 3 terms pooled, validated at high volume, routes to #30 and graduates as the stem phrase; “deep pocket” 4 terms validated long-tail, LEAVE TO HARVEST standing coverage; “rv” 3 terms, 8 clicks, 0 orders, invalid, negative phrase “rv” staged. Validated terms routed: 7. Yield = $7 / \$498 \times 100 = 1.41$ per \$100: +11% over the 1.27 baseline at comparable CPC with clean walls, under the 1.5 prediction. Verdict at the gate: SCALE per the decision rule (above baseline, walls clean); the V5 row scores the standup prediction PARTIAL (right direction, short of magnitude) with the note that the “cooling” seed underperformed inside a family whose cooling demand is owned elsewhere per the #33 declaration, a seed-selection lesson the next family standup consumes.

WE-2 - The edge case that goes wrong: a mid-window seed without its wall (family “Satin”)

Setup. Pilot week 3. A new stem “deep pocket” was seeded into the family phrase campaign mid-window, and the session skipped P1 step 8: no same-session negative-exact wall in the family broad campaign.

Detection (P5). The weekly join on Customer Search Term finds “satin sheets queen deep pocket” served in BOTH campaigns inside the same 14-day window: phrase 42 clicks at CPC \$2.05 = \$86.10 spend, 1 order; broad 31 clicks at CPC \$1.71 = \$53.01, 1 order. Total: 73 clicks, \$139.11, blended CPC \$1.91, bought at two control levels with the account’s own money. This is the F2 condition.

Response. Same day: negativeExact on the term (and negativePhrase on the stem where the map shows siblings of the string) staged into broad and deployed; the addition writes to the T3 wall list for the #29 wall audit. Both campaigns’ rate reads on the term are marked confounded for the window; the \$139.11 reclassifies as structure waste in the audit. The yield read recomputes excluding the double-bought term: 5 validated on \$460 net = 1.09 per \$100 versus the 1.27 baseline. Because the decision rule reads yield and walls first (F5 guard), the day-28 outcome is HOLD with a one-window extension and the wall question named as closed, not RETIRE: the depressed yield was a wall artifact, not a demand verdict. T6 root cause names the mid-window session that skipped step 8; one recurrence escalates per Section 7.

WE-3 - The head term found in phrase: control precedes judgment

Setup. The family phrase campaign serves the head root itself, “bamboo sheets” (oversight class), at 66 clicks in the 14-day window = 4.7 clicks/day, 19% of the campaign’s clicks. T2 K8 on the term reads # Exact = 0 for this product: the same class of failure the June 30 review surfaced on the 6-piece head term carried by auto with zero exacts (Criteria System Section 16, Worked Pass 3).

Route (P4 branch 4). Scenario S-R12 (Section 8.4: a priority term with zero exact campaigns receives ADD EXACT + RESTRUCTURE before any performance verdict) fires at T2. SOP-12 executes the standup with #29 walls the same session; this SOP’s half is P9 step 2, the negativeExact on “bamboo sheets” in the family phrase campaign, staged in the same set, so the head term consolidates under exact control. The 66 clicks are excluded from the phrase yield read: a head term resident in a discovery surface was never discovery, and counting it would flatter the pilot. The T6 line references the SOP-12 creation prediction; this SOP logs no second prediction on the same event.

7. Failure Modes and Escalation

Failure	Detection signal	Response	Motivating account-history incident
F1. Pilot without a dated stop	Standup record missing window dates or the prediction row	Does not deploy; P2 step 3 is a precondition. The 50-click test discipline (Doctrine v1.0 Section 13) applies to every dated phrase test	None on file; the guard predates the failure by design
F2. Phrase-broad double-buying (wall missed)	P5 join returns a non-empty served-term intersection	Wall built same day per P5; reads confounded; spend reclassified structure waste; root cause names the session	The graduated-term-never-negated class: F8 GRADUATED-UNSTEERED (Criteria System Section 3) is the account’s named leak pattern at the exact seam; P5 is its phrase-side twin
F3. Phrase on an unproven root	T3 phrase campaign standing on a family whose root shows # Exact = 0 or Validation State unclear	Reclassify: pause; the root goes to broad discovery where it belongs (SOP-14)	None on file
F4. Term-by-term graduation noise	Multiple single-term graduation candidates from one stem in one pass	P4 stem routing enforced; the stem graduates once as the stem phrase	None on file
F5. Pilot judged on ACOS instead of yield	A kill recommendation citing ACOS while yield reads at or above baseline with clean walls	The decision rule reads yield and walls first; ACOS governs only through the Discovery band (Criteria System Section 6, O6)	The blended-number habit the Criteria System Section 1 verdict names account-wide
F6. Head term resident in phrase with no exact	P4 branch-4 read: an oversight-class string in the served set with T2 # Exact = 0	ADD EXACT + RESTRUCTURE per S-R12; SOP-12 executes; P9 walls; clicks excluded from yield (WE-3)	The 6-piece head term carried by auto with 0 exacts, June 30 review (Criteria System Section 16, Worked Pass 3): the same control-gradient failure on a different low-control surface
F7. Family straddle	One phrase campaign carrying seeds tagged to 2 syntax families (T2 Syntax Tags versus the T3 family assignment)	Split by family the same session; the family is the cap-enforcement unit (rule P4, Section 12) and a straddle makes T1 cap reads unreadable	The cooling-terms-inside-the-satin-tracker class (gate G12 rationale, Criteria System Section 5): demand tracked inside the wrong family
F8. Negation cadence lapse	T3 WAS % above the Discovery limit (F11 bands) on a phrase row for 2 consecutive passes	Reactive negation pass this session; if the breach source is off-taxonomy converting terms, the G12 route governs: MKL classification first, taxonomy expands deliberately or the term negates (S-H6)	The satin product taking clicks on cooling terms it cannot rank: the F11 pre-load failure named in the Criteria System Section 3
F9. Structure change inside an armed event window	T6 phrase structure entry dated inside a G12 window with no event-plan reference	Revert unless the event plan names it; deviation escalates per #32; trend and yield exclusions extend from the revert date; the pilot day-count correction is logged	None on file; guard mirrors SOP-12 F8

Escalation. The same failure mode twice in consecutive cycles on the same scope escalates to the PPC Lead with the T6 history attached. F2 recurrences additionally route through the #29 wall audit. Anything in the human-mandatory tier (pilot standup, both pilot gates, stem graduations via the #30 confirm, P8 archives) is confirmed per the Authority Matrix v2.0 before deployment, without exception.

8. Outputs and Artifacts

Artifact	Storage location	Retention	Downstream consumer
Stem map (stem terms clicks orders status route date)	T3 notes on the phrase campaign row, PPC Command Workbook for the product	Lives with the workbook; superseded maps remain in the row history	The next weekly pass; #30 (graduation candidates); #29 (wall baseline)
Pilot gate tables (day-14 health, day-28 decision with both yield reads)	T6 attachment on the gate rows	Append-only with T6	PPC Lead (gate decisions); SOP-06 (V5 scoring)
Staged bulk change sets (creations, negatives, budget deltas)	#38 staging location, #38 format, reason strings carrying verdict + scenario ID	Per #38, including the T+1 verification report	Data Operations / #38
T6 decision-log rows with predictions and review dates	T6 Decision Log tab, append-only	Retention horizon per the #39 fold-in (amendment pending there, not here)	SOP-06 V5 scoring; A3 verification scoreboard
Wall lists (pre-load, dedup, steering)	T3 Negation State column per campaign	With the campaign row; archived state closes with P8 step 4	#29 wall audit
Updated Ledger row for component #13	Section 14 of this document, five-column copy-paste form	Pasted into PPC_Cross_Reference_Ledger, version raised per its Section 6	The next build session

9. Skill-Conversion Block (Element 9 handoff)

Block	Specification (SOP-13-specific, not templated)
Input bindings (catalog IDs)	A2 deployment state; A5 threshold flags; A6 objective + Expected CPC (PARTIAL: \$2.80 account reference rides with the PROVISIONAL tag); B5 coverage counts (# Phrase, # Exact per term); B8 economics chain (G9 tag while D1 is blended); D7 decision log; plus the SP + SB search-term and STIS join from the Raw Pull Pack (Data Architecture Section 12). An absent search-term feed suppresses P3, P4, and P5 for the week; the skill never improvises around a missing feed.
Deterministic functions	Four, all closed-form: (1) stem signature = lowercase, de-pluralized, token-sorted string (the F8 twin-set normalization at stem grain); (2) yield = validated terms routed ÷ spend × 100; (3) wall-breach set = served-terms(phrase) intersect served-terms(broad) on Customer Search Term; (4) the P2 gate rule read over the paired T3 rows. The skill executes the P4 tree and the Section 5 mapping; it never computes its own thresholds and it version-pins the Criteria System per the V6 sync rule (Criteria System Section 14).
Human checkpoints (10-80-10, from the Authority Matrix v2.0 rows)	Per the row "Structure creation (all SOPs)": PPC Lead decides, the incoming operator and PPC Manager execute; per "Weekly verdicts (established US)": PPC Lead decides, operator executes assigned scope, BM owner reviews at the daily huddle (review, not veto). Human-mandatory in this SOP: pilot standup, both pilot gates, stem graduations (via the #30 confirm), P8 archives, every deployment confirm. Vacancy rule 5 of the Matrix applies: an unfilled role's D reverts to the PPC Lead, its E to the PPC Manager.
Cadence	Weekly with the Monday pass (P3-P5); daily pacing glance during a pilot window; wall-breach check on every search-term pull; same-session P9 on graduation receipt; G12 freeze same day the window arms.
Alarm specification	F1-F9 each carry a Section 7 detection signal and an evidence-line format (row, signal value, gate or rule reference). Routing: F2/F9 same day to the operator queue with the PPC Lead copied; F3/F4/F5/F8 into the weekly pass worklist; F1 blocks at staging QA; F6/F7 to the PPC Lead with T6 history.
Skill outputs	Drafted stem maps and candidate routes; the paired-campaign yield table; drafted bulk sets and T6 rows queued for human confirmation; F1-F9 alarms with evidence lines; the V5 scoring draft for phrase rows due this cycle.

10. Version Governance

Owner: PPC Lead. Named review triggers specific to this component: (1) any Criteria System version increment, at which point the Section 5 mapping re-validates before any pass runs on it; (2) the PPC-DR-04 vocabulary ratification, at which point Section 5 re-issues against the ratified string set and AM-1 closes; (3) any change to SOP-14 P3, because the shared machinery increments both SOPs; (4) the pilot day-28 gate, because SCALE converts P2 into the standing P1 state and this SOP increments to reflect it; (5) the D1 unblended per-variation margin landing (#41 due 2026-08-21), which retires the G9 PROVISIONAL tag riding every ceiling-referencing action here; (6) any build session that returns this component's Ledger row. Registered in the Supersession Map at publication. v2.0 change note: supersedes v1.0 in full; procedures 4 to 9 (adds P5 wall audit, P6 CAP SYNTAX receipt, P7 event overlay, P8 pause and retirement, P9 graduation handoff); verdict mapping expanded from 6 rows to the full 25-string closed vocabulary with boundary statements; worked examples 1 to 3 including the mandatory edge case; failure modes 5 to 9 with motivating incidents; the v1.0 WE funding arithmetic, which corrected itself mid-sentence in print, is restated cleanly in WE-1; Elements 11 and 12 added for the first time in the library.

11. Operator Ramp (Element 11)

Window	What the person taking over this SOP must demonstrate
Week 1	Reproduce the Economics Worked-Example Set cases E1 and E2 cold (the onboarding arithmetic gate). Shadow one full P3-P5 pass and independently produce the stem map for one family; the map must match the PPC Lead's on route assignments for every stem, with any disagreement argued from the P4 tree, not from preference. Name the four dedup wall placements of P1 from memory.
Week 2	Run the pass live under review: pull the search-term report to the P3 step-3 spec, run the P5 join, stage one negation set and one seed addition in #38 format with correct reason strings. Walk one graduation end to end verbally: candidate, #30 route, SOP-12 standup, P9 wall, single prediction.
Week 4	Own the weekly pass for assigned families. Run one pilot gate read (day-14 or day-28) and present the decision table with the verdict argued from the rule. Zero F2 findings attributable to own sessions across the window.
Live-defense questions	(1) Why does a validated long-tail stem stay in phrase instead of graduating, and which verdict string names that state? (2) Pilot yield 1.09 against a 1.27 baseline: what do you check before recommending RETIRE, and why? (3) A phrase-served string turns out to be an oversight head term with # Exact = 0: which scenario fires, who executes, and what is this SOP's half? (4) Broad and phrase both served the same query this week: name the failure, the same-day response, and what happens to the window's reads. (5) The family hits CAP SYNTAX: what does this SOP execute and what does it never originate, and where does the answer live given conflict C3? (6) A deal window arms mid-pilot: what freezes, what stays live, and what happens to the day-count?

12. Interface Contracts (Element 12)

For every upstream and downstream handoff: the sending owner, the receiving owner, the artifact, the timing, and what happens when it does not arrive. Built, per Ledger conflict C5, from the Ledger's Interfaces column plus the Data Architecture v1.0 artifact and timing definitions.

Direction and counterpart	Sender -> Receiver	Artifact	Timing	When it does not arrive
Upstream: Raw Pull Pack (#39 instrument)	Data Operations -> PPC operator	SP + SB search term reports and STIS sheets with current pull date	Monday 06:00 land; workbooks refreshed by 09:00 (Data Architecture Section 12)	P3-P5 suppressed for the week; suppression recorded in the pass output; no improvised pulls
Upstream: instrument budget law (C3: #31 has no SOP)	T1 S3 columns / M5 queue -> this SOP	Family discovery allocation; CAP SYNTAX deltas with finding IDs	Weekly with the cascade; same-session staging on a breach	No delta means no budget change: this SOP never originates one; a cap breach visible with no staged delta is reported to the PPC Lead same day
Upstream: #32 event plan	Deals owner + PPC Lead -> this SOP	Armed-window dates into T2 Event Mode	On window confirmation, before the window opens	An unarmed window discovered live is the F9 revert case; escalate per #32
Upstream: #33 ownership declarations	Account Rollup A1 -> this SOP	Declared family term owners (S-F3) for pre-load negatives	On declaration and at the monthly review	Absent a declaration on a contested root, the root is not seeded and the gap is logged; seeding into an undeclared conflict is the G11 class
Downstream: #30 harvest and graduation	This SOP (P4) -> #30	Routed stem candidates with pooled evidence in the stem map	Weekly, inside the Monday pass	#30 receives nothing and graduates nothing; a validated stem sitting unrouted 2 passes is a pass-quality finding on this SOP
Downstream: SOP-12 exact operations	This SOP -> SOP-12	ADD EXACT + RESTRUCTURE receipts (P4 branch 4) and graduation packets (P9 step 1)	Same day as the T2 verdict or the #30 confirm	SOP-12 not executing same day leaves the term double-exposed; the open item stays on the pass worklist flagged daily until the standup lands

Direction and counterpart	Sender -> Receiver	Artifact	Timing	When it does not arrive
Downstream: #29 negation architecture	This SOP -> #29	Wall lists (P1 steps 7-8, P9 step 2) and F2 breach findings	Same working session as each structure event; findings with the weekly pass	The #29 wall audit runs against an incomplete baseline; the missing list is rebuilt from the T3 row before the next audit
Downstream: #38 bulk mechanics	This SOP -> #38	Staged change sets with verdict + scenario ID reason strings	Tuesday default; same day on dated releases and F2 walls	Nothing deploys: #38 is the only door; a console edit found at T+1 is the SOP-12 F10 class
Lateral: SOP-14 broad operations	Shared -> shared	The P3 triage machinery (owned by SOP-14) and the dedup wall seam (owned here at P1 step 8 / P5)	Every weekly pass	A change to either side of the seam without the paired increment is a version-governance breach; both SOPs re-validate (Section 10 trigger 3)

13. Amendment Items Logged by This Build (the v4.1 queue; anti-drift rule 1)

ID	Item	Why it is logged instead of written
AM-1	Verdict vocabulary reconciliation: Criteria System Section 13.1 lists 18 actions; the KDR Template v3.0 Bands and Verdicts Reference and SOP-12 v2.0 Section 5 carry 25. This is PPC-DR-04.	Section 5 above uses the KDR set as the Build Kit authorizes and re-issues on ratification; the build does not pick the canon.
AM-2	Section-citation map: the corpus cites v4.0 Section 4.2 (naming) and 4.4 (bidding strategy); the uploaded Criteria System contains neither, and its Section 17.2 keeps the naming convention authoritative in the v1.0 Criteria System, outside this pack.	This SOP cites SOP-12 v2.0 Appendix A and Doctrine v1.0 Section 6.3 as the uploaded restatements; the canonical section map needs one ruling.
AM-3	Pilot decision-rule tolerance: "above baseline at comparable CPC" has no numeric tolerance in the Criteria System.	The rule is stated structurally (WE-1, WE-2 apply it); the tolerance is a threshold and thresholds are not invented here.
AM-4	Stem-level evidence pooling: G2 pools twin sets explicitly; pooling across a modifier stem is analogous but unstated.	P4 step 2 pools at stem grain as v1.0 did; the sufficiency clause needs a one-line G2 amendment to be canonical.
AM-5	Dedicated-coverage volume read for seeding a stem as its own phrase row (P1 step 6, P4 branch 1) has no stated bar.	Seeding judgment is currently human at the standup; the bar should live in the Criteria System, not here.

14. Ledger Row Returned by This Session (paste into PPC_Cross_Reference_Ledger, five-column Section 1 form)

#	Component and owning document	Owns exclusively	Must not restate: defer to	Interfaces, blockers and open findings
13	Phrase operations, including the per-family pilot. PPC_SOP_13_Phrase_Operations_and_Pilot_v2_0.pdf (supersedes v1.0 in full)	Phrase-match operations: family phrase creation (P1), the dated pilot (P2), the stem-clustering read and its 4 routes (P4), the phrase-broad dedup wall build and audit (P1 step 8, P5), CAP SYNTAX proportional execution on discovery (P6), phrase event overlay (P7), pause and retirement (P8), the phrase-side graduation handoff and steering wall (P9), and Surface-2 standing long-tail stem coverage under LEAVE TO HARVEST.	Graduation decisions: #30. Wall doctrine and audit: #29. Shared weekly triage machinery: SOP-14 P3. Exact standup and ADD EXACT + RESTRUCTURE execution: SOP-12. Thresholds and verdicts: the Criteria System. Budget law: the instrument layer per conflict C3. Staging: #38. Event plans: #32. Ownership declarations: #33.	Downstream: P4 candidates route to #30; graduation packets and S-R12 receipts to SOP-12 same day; wall lists and F2 findings to #29. Lateral: SOP-14 P3 coupling, both increment together. Open: AM-1 through AM-5 logged this build (AM-1 = PPC-DR-04); pilot standing state pending the day-28 gate; G9 PROVISIONAL rides until #41 lands (due 2026-08-21).

Companions: Master PPC Decision Criteria System v3.0, operating as the v4.0 threshold authority per the PPC-DR-02 session ruling (all thresholds and scenario logic) · Keyword Decision Record Template v3.0, Bands and Verdicts Reference sheet (verdict strings, band summaries) · PPC Strategic Doctrine v1.0 Sections 3.2, 3.4, 6.3, 13 (the why) · SOP-12 Exact Campaign Operations v2.0 (exemplar; Appendix A naming; P1/P7/P10 execution seams) · SOP-14 Broad, #29 Negation, #30 Harvest, #32 Events, #33 Arbitration, #38 Bulk Mechanics (boundaries, by number) · Master PPC Data Architecture v1.0 and the PPC Command Workbook TEMPLATE v1.0 (instrument mechanics, T-tab columns cited as printed) · PPC Cross-Reference Ledger v1.0 (component #13 row, superseded by Section 14 above on paste-in) · PPC Economics Worked-Example Set v1.0 (the ramp arithmetic gate) · PPC Authority Matrix v2.0 (Element 9 and 11 checkpoints).