PPC DEPARTMENT · OPERATING SOP · COMPONENT #14 OF THE OPERATIONS DEFINITION REGISTER

SOP-14: Broad Campaign Operations

Version 2.0 · 2026-08-07 · Owner: PPC Lead · Supersedes SOP-14 v1.0 (2026-07-24) in full · Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds. Built to the SOP Build Kit v1.0 twelve-element standard against the SOP-12 v2.0 depth floor.

1. Verdict and Scope

VERDICT: broad is the fenced discovery engine, the fence is the product, and the fence is built weekly or it is not built at all. A broad campaign is graded on cost per discovered converter, WAS percentage and graduation rate, never on its own ACOS, because v4.0 Section 6 rule O6 states that judging a campaign on another objective's metrics is a process violation. Broad carries the 15-20% match-type spend destination of the M6 structure targets (v4.0 Section 10, rule M6). Version 2.0 expands v1.0 from 6 procedures to 13, from a 7-row verdict mapping to the full 25-verdict closed vocabulary, from 1 worked example to 4, and from 5 failure modes to 12, adds Elements 11 and 12, corrects three citation defects carried in v1.0, and states for the first time that procedure P3 is an account-wide shared definition rather than a broad-only procedure. This SOP governs the full operating life of the three broad surfaces: family discovery standup, the modified-broad fence, the weekly three-route triage, defensive brand broad, graduation yield accounting, the negation-cadence audit, dated-test governance, CAP SYNTAX receipt, capped-by-design budget mechanics, event and recovery overlays, variation re-point, and restructure-or-retire.

1.1 The Three Broad Surfaces

Every broad keyword in the account lives on exactly one of three surfaces, and the surface determines its objective tag, its bid posture, its review depth and its exit path. Movement between surfaces is a staged, logged structure event, never a mid-pass judgment.

Surface	Structure and posture	Why it exists, and its exit
Surface 1 **Family discovery broad**	One campaign per syntax family, Objective = Discovery, seeded with the family root plus validated modifier roots, budget-capped by design, never straddling two families because the family is the cap-enforcement unit (v4.0 Section 12, rule P4). Bidding down-only per Doctrine v1.0 Section 6.3. Placement modifiers 0 on all three groups at creation.	Paying the algorithm to reveal demand the Master Keyword List has not mapped (Doctrine v1.0 Section 1, principle 1, the Control Gradient). Exit is graduation: winners leave upward through #30, losers leave sideways through negation, and the campaign itself exits only through P13.
Surface 2 **Modified broad**	The same family campaign carrying a required-word lock on a drifting keyword (+bamboo +sheets). The unfenced version pauses in the same change set so the two never serve simultaneously. Bid is carried from the paused keyword, never re-invented at deployment.	The middle rung of the control gradient when open broad drifts but phrase is too narrow (Doctrine v1.0 Section 3.3). Exit is either back down to open broad when the root-containment read recovers, or up to phrase coverage owned by SOP-13.
Surface 3 **Defensive brand broad**	One campaign per brand, Objective = Defensive, broad match on brand terms, floor bids at the discovered floor. Every generic discovery campaign carries steering negatives on brand terms so defense owns them exclusively.	Broad match captures misspellings and word-order variants that exact cannot enumerate. Judged on branded impression share held and incrementality, not on discovery yield (v4.0 Section 6, rule O6, Defensive metric set). Permanent state; the ladder re-runs to re-discover the floor.

1.2 P3 Is an Account-Wide Definition, Not a Broad Procedure

Three Ledger rows and three shipped v2.0 documents place the definition of the weekly three-route search-term triage inside this component. SOP-13 v2.0 Section 1.3 runs it by reference on phrase campaigns and adds only its stem-clustering and dedup reads. SOP-15 v2.0 P3 executes it on every auto group's keyword rows and adds only the two ASIN routes. SOP-29 v2.0 Section 1 owns the standards Routes B and C execute against and states that it never re-runs the triage. The consequence is a hard version-governance rule stated in Element 10: **a change to P3 increments SOP-13, SOP-14, SOP-15 and SOP-29 together.** P3 is therefore written below as a definition with numbered steps that another document can cite by step number, not as prose that only makes sense inside a broad campaign.

1.3 Rulings and Conflicts Declared Before Writing (anti-drift rule 5)

Six items were declared before any content was written, so this build does not silently pick a side.

Item	Statement	How this document proceeds
Version note	The Criteria System file's cover and footers read v3.0. That is a known stale header on the correct current file, reconciled in the pack's version note.	Cited throughout as v4.0 with its actual printed section headings. Not a finding, not an amendment item.
C9 **gate numbering**	v4.0 Section 5 assigns G9 to price stability, G10 to event exclusion and G12 to MKL hygiene. SOP-14 v1.0 Section 3, SOP-12 v2.0 Section 3, SOP-30 Section 3 and the Data Architecture use G9 for margin freshness, G10 for objective tagging and G12 for the event freeze.	Every gate is named by its v4.0 number and name, with the corpus alias given at first use in Element 3. Rides to v4.1 as SOP-29's AM-11; not re-logged here.
C10 **verdict vocabulary**	v4.0 Section 13.1 carries eighteen strings; the Keyword Decision Record Template v3.0 sheet "Bands & Verdicts Reference" carries twenty-five; Data Architecture v1.0 Section 2 rule 8 states twenty-six.	Element 5 quotes the KDR's twenty-five verbatim, matching exemplar SOP-12 v2.0 Section 5, and carries the v4.0 string as a stated cross-map on the two divergent rows so neither citation is dropped.
C12 **scenario ID**	Data Architecture v1.0 Section 6.1 labels the exact-live-with-source-unnegated filter S-H6; v4.0 Section 8.9 carries it as S-H4.	S-H4 is used for the unsteered-graduate condition throughout. S-H6 is used only for its own v4.0 condition, the off-taxonomy converter.
C3 **component #31**	Component #31 budget and envelope operations carries no procedural document; five documents cite it as though it had numbered procedures.	Wherever this SOP receives budget law it names the instrument and tab (T1 group S3 spend-governance columns, T0 zone H, the M5 queue per v4.0 Section 10). No #31 procedure number appears anywhere in this document.

Item	Statement	How this document proceeds
Naming citation	SOP-14 v1.0 P1 step 2 instructs the operator to name campaigns "per the 4.2 convention". The uploaded Criteria System contains no Section 4.2, and its Section 17.2 keeps the naming convention authoritative in the v1.0 Criteria System, which is outside this pack.	Naming references SOP-12 v2.0 Appendix A, the uploaded restatement of the governing convention. The dangling citation is logged as amendment item AM-6.

1.4 Boundaries With Adjacent Components

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, gate, scenario, verdict and lever limit. This SOP executes; a threshold appearing here instead of there is a defect.
#30 Harvest and graduation	Decides which terms graduate, pools evidence across all serving sources, and issues the verdict. P3 Route A delivers candidates; this SOP never validates a candidate and never issues a graduation.
#29 Negation architecture	Owns negation doctrine, the wall-level selection tree, the pre-load standard, the reactive execution standards and the monthly wall audit. This SOP produces Routes B and C with evidence lines and owns the weekly cadence read; it never sets a wall standard.
#12 Exact campaign operations	Owns exact standup. Every head term found living in broad produces an ADD EXACT + RESTRUCTURE receipt executed at SOP-12 P1 or P2; this SOP builds only the broad-side steering negative afterwards.
#13 Phrase operations	Owns the phrase surfaces, the stem-clustering read and the phrase-broad dedup wall decision. This SOP supplies the shared P3 machinery and takes the dedup wall's broad-side placement from #13's build.
#15 Auto operations, D5 split	Owns the four-group auto structure and the two ASIN routes. This SOP supplies the keyword-route definition #15 executes; auto group bids, budgets and verdicts are never touched here.
#28 Placement and bid-modifier	Owns ladder step derivation, exit-bid discovery and modifier sizing, including the brand-broad defense floor. This SOP deploys the campaign-level rows #28 stages back.
#16 PAT and conquest	Decides target opens, rotations and exits. Competitor-brand strings appearing in the Surface 3 served terms are incursion evidence routed to #16; this SOP never opens a target.
#33 Sibling arbitration	Owns family term ownership and the declaration workflow. A sibling-owned root is a pre-load negative here and is never seeded; the G11 check runs before any scale action on a shared root.
#32 Deal and event windows	Owns the event plan, ramp, caps and reconciliation. This SOP freezes broad structure changes inside armed windows and executes only what the event plan names.
#35 Incrementality testing	Designs and reads incrementality tests. The Surface 3 brand-broad step-downs execute here; the continue-or-restore call is #35's and v4.0's.
#38 Bulk deployment mechanics	Owns pull, staging format, upload and T+1 verification. Every change this SOP makes exists first as a staged bulk set in #38 format.
#37 Automation platform rules	Owns the rule register. This SOP sends exclusion-state notices for dated tests, modified-broad parallel states and event freezes; it never edits a rule.
Budget and envelope (instrument)	Budget deltas, CAP SYNTAX receipts and envelope law arrive from the instrument layer: T1 group S3 spend-governance columns, T0 zone H, and the M5 queue per v4.0 Section 10. This SOP receives deltas and never originates them (conflict C3).

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. The lag is measured from the artifact's timestamp, not from the moment a person read it; starting later than the stated lag is a failure mode, not a delay.

Trigger	Detecting artifact	Max lag	Procedure	Evidence the trigger fired
Weekly cadence	Monday audit cascade reaches T2 KEYWORD COMMAND; the T7 MANIFEST row for the search-term join shows the current week's refresh date	Same session; deployment Tuesday per #38	P3, then P5, then P7	Route counts A, B and C recorded per campaign in the T6 DECISION LOG, and a yield figure written to the T3 CAMPAIGN BOARD "Harvest Yield" column
New family standup	T1 SYNTAX BOARD row live for the family with "Gap KWs" above zero, and T2 K8 "# Broad" reading 0 across the family root	Enters the next weekly verdict queue	P1	A T3 row exists carrying Objective Tag = Discovery, the capped-by-design note, and a non-empty pre-load wall list
WAS breach on a broad row	T3 column "WAS %" above the Discovery limit of v4.0 Section 3, family F11 "Waste & Negation"	Negation pass in the same weekly session	P3 Routes B and C	Staged negation rows with complete evidence lines, and the recomputed WAS percentage recorded against the campaign
Root-containment drift	The P2 step 1 read: served terms containing the family root as a share of all served terms, this window against the prior window, both recorded on the T3 row	Modified-broad staged in the same weekly session	P2	Both window reads present on the T3 row with the delta stated in percentage points
Graduation approval returned	The SOP-30 P1 graduation log row reading APPROVE with this campaign named in its attached source list	Same day as the standup; the wall confirmation closes before the session does	P6	This campaign's negative list shows the entry, and the term is absent from its served terms at the next pull
Head-term detection	T2 campaign child rows, "Spend Share of KW %", read against v4.0 Section 4 rule X10 with the K8 "# Exact" count	Receipt to #12 the same day, outside the weekly rhythm	P3 step 6	The receipt exists in #12's queue and the term's clicks and spend are excluded from this campaign's yield base

Trigger	Detecting artifact	Max lag	Procedure	Evidence the trigger fired
Dated-test expiry	T1 group S3 "Dated Test / Stop-Loss" column reaching its expiry date, or the T2 K7 "Test Status" click quota filling	Verdict issued within 1 working day of the read landing	P8	The test row carries a closed verdict string and a scored T6 row against the standup prediction
CAP SYNTAX receipt	T1 row "Cap Status" reading BREACH with the staged delta from the instrument layer	Budget rows staged in the same weekly session	P9	The family's broad daily budget shows the reduced value at T+1 and the T1 share recomputes inside cap
Event window arms	T2 K2 "Event Mode" set from the #32 event plan (v4.0 Section 5, gate G10 event exclusion; corpus alias G12 event freeze)	Freeze effective the day the window arms	P11	No structure event dated inside the window without an event-plan reference in its T6 row
Recovery Mode entry	T0 zone A code-red flag, or the product's TACOS reading against its stage band per v4.0 Section 8.1, scenario S-A4	Same day as the entry log	P11 step 5	Discovery budgets show the sentinel level and the reduction total is logged as expected weekly savings
Inventory band change	T2 K2 "Inventory Band" flipping to YELLOW or RED on a targeted variation, with D2 naming the ranked backup	Same day on RED; next pass on YELLOW	P12	The advertised SKU on the affected campaign matches the D2-named backup at T+1
Yield failure	The T3 "Harvest Yield" column reading zero across the governance window with spend active	Evaluated in the same weekly pass the second zero lands	P13	A written evaluation exists covering seed refresh, negation audit and budget rebalance before any retirement is staged
Sibling declaration change	The #33 declaration record, a re-arbitration changing the owner, or a stewardship grant, landing in T2 K8 "Term Owner (Family)"	Same day the declaration date carries	P1 step 6, P3 step 3	The non-owner's broad campaign shows negative phrase on the owned root; the owner's campaign carries nothing for it
Verification return	The #38 T+1 verification report flagging a landed-value mismatch on a staged broad row	Redeploy the delta same day	P10 step 5	The redeployed delta verifies clean and the week's reads on affected rows are marked void

3. Preconditions and Gates

Every procedure below runs only after these preconditions read current in the T7 MANIFEST. A missing or stale field group suppresses the dependent execution class; the pass records what it could not see and continues on the rows it can (v4.0 Section 5, gate G8 packet completeness).

> **Gate naming, per conflict C9.** Gates are named below by their v4.0 Section 5 number and name, with the corpus alias in brackets at first use. The three that diverge: **G9** is price stability in v4.0 and margin freshness in the operating corpus; **G10** is event exclusion in v4.0 and objective tagging in the corpus; **G12** is MKL hygiene in v4.0 and the event freeze in the corpus. Where this document needs the corpus meaning it names the meaning in words as well as the number, so no instruction depends on which convention the reader learned first.

Manifest ID (T7 row)	Requirement before any procedure runs	Observable that proves it passed
Search-term join (Raw Pull Pack)	The Sponsored Products and Sponsored Brands search term reports plus STIS, joined to T2 campaign child rows per Data Architecture v1.0 Section 12. Broad is unreadable without it: an absent feed **suppresses P3 entirely** for the affected campaigns rather than degrading it, and no yield read is computed for the window.	T2 campaign child rows carry per-term-per-campaign splits with a non-zero "Spend Share of KW %" for the window, and the T3 "WAS %" column is populated. This feed has no dedicated T7 row, which is amendment item AM-1.
A2 Deployment state	Current bulk state in-row for every broad campaign; stale state fails freshness and the pass stops for those rows.	T7 row A2 = PRESENT; T3 rows carry State, Daily Budget $, and Bid / AG Default $ from the current-week pull.
A3 Pacing buckets	In-budget percentage present per campaign so the capped-by-design exemption of P10 can be distinguished from real truncation.	T3 "In-Budget Bucket" and "Truncation Flag (G3)" both populated; every capped-by-design row also carries the note in its T3 record.
A5 Threshold flags	Band and flag columns precomputed rather than re-derived at read time (Data Architecture v1.0 Section 2, rule 4).	T2 K5 flag columns present: ACOS band, CTR-versus-market, CVR-versus-market, zero-quarantine.
A6 Objective + Expected CPC	Every broad campaign tagged Discovery or Defensive; no untagged broad row proceeds (v4.0 Section 6, rule O1; corpus gate alias G10 objective tagging).	T3 "Objective Tag" populated on 100% of broad rows. T7 A6 is PARTIAL: the account reference Expected CPC rides with the provisional flag until the per-term value lands.
B5 Coverage counts	Registry counts per match type current, or creation legality and the head-term check are unreadable.	T2 K8 "# Exact", "# Phrase", "# Broad" and "# Auto" populated for every term the window served.
B8 Economics chain	T0 named cells live: AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC.	T0 zone C reads fresh; the v4.0 gate G9 price stability flag, and the corpus margin-freshness flag, both carry onto any ceiling-referencing action.
B9 WoW deltas	Prior-window figures present so the root-containment read of P2 and the cadence read of P7 have a comparison point.	T2 K5 "WoW Block" populated; the T3 row holds the prior window's yield and negation counts.
D2 Variation map	Targeted variation known with its ranked-backup chain, or P12 cannot execute.	T2 K2 "Targeted Variation(s)" and "Inventory Band" populated; unmapped rows carry the D2-pending flag and the campaign points at the preferred SKU.
D7 Decision log	The prior cycle's T6 rows whose review dates land this cycle are available for scoring; unscored actions do not accumulate (v4.0 Section 14, rule V1).	T6 rows for the due dates carry a V5 score before P3 step 1 completes.

Gates in force across every procedure

Gate	Effect on this SOP
G1 Validation quarantine	A quarantined row produces a flag and a data-refresh task, never a route. The term stays in its source untouched for the window.
G2 Sample sufficiency	Every route decision is checked against the v4.0 Section 5 read thresholds for the judgment being made. Below threshold the term stays in its source under a capped dated test rather than entering a negative list or a candidate queue.
G3 Budget truncation	A truncated campaign gets the budget lever before any bid conclusion. A deliberate discovery cap is not truncation; P10 step 1 states the distinction and the T3 note that carries it.
G6 Inventory adequacy	The band on the targeted variation governs. RED routes to P12 the same day; YELLOW freezes any spend increase on the family's discovery campaigns at the current level.
G7 Envelope and pacing	Campaign budgets reconcile to the product envelope. A discovery budget rise that breaches the envelope is refused at staging, not corrected after deployment.
G8 Packet completeness	A verdict whose gate fields are missing downgrades to a flag, and the pass states what it could not see. This is the gate that makes the search-term suppression rule enforceable rather than optional.
G9 Price stability [corpus: margin freshness]	A price change on the targeted variation inside the trailing window marks CVR and ACOS reads provisional; no conversion verdict is issued on a provisional read. The corpus margin-freshness flag rides the same column on ceiling-referencing actions.
G10 Event exclusion [corpus: objective tagging]	Event weeks and the following guard period leave the trend reads that feed P2 and P7. The corpus alias covers rule O1: an untagged broad campaign is unmanaged by definition and routes to assignment before it is analyzed.
G11 Sibling conflict	A shared root carrying two products' active discovery is a portfolio decision, not two product decisions. Nothing scales on a contested root until #33's declaration lands.
G12 MKL hygiene [corpus: event freeze]	Every analyzed term must exist in the current Master Keyword List with a syntax tag and a relevancy score; untagged terms route to MKL maintenance and taxonomy pre-negation, not to analysis. The corpus alias covers the armed-window structure freeze executed at P11.

4. Step-Level Procedures

Time estimates are planning figures for one product, not standards. Every step names the artifact it touches. Where a step needs a number, the number is read from v4.0 or from a named workbook cell at execution time and is never carried in this document.

P1. Create a family discovery broad campaign (35-50 minutes)

1. Confirm the family qualifies: the T1 SYNTAX BOARD row is live, "Gap KWs" reads above zero, and the family is not sitting in a chronic CONVERSION diagnosis. A family that cannot convert does not receive discovery spend (v4.0 Section 8.2, scenario S-Y5, which freezes bids at maintenance across a chronic conversion syntax).
2. Confirm creation is legal: read T2 K8 "# Broad" across the family root and the T3 registry. One broad campaign per syntax family. If the family campaign exists, extend it; a second campaign on the same family with no declared split is the F8 condition, not a creation.
3. Confirm gates: G6 GREEN on the D2-preferred variation, G12 clear or the event plan names this creation, G12 MKL hygiene passed on every seed term, and G11 clear on the family root against the #33 declaration.
4. Name the campaign per the convention restated in SOP-12 v2.0 Appendix A, with the family root in the keyword position, Objective = Discovery. The objective segment must parse or the T3 naming-compliance flag fires. Per section 1.3 above, the convention is cited there and not to a v4.0 section number.
5. Stage the campaign settings as bulk rows per #38: portfolio = the product's portfolio; bidding strategy down-only per Doctrine v1.0 Section 6.3, which sets down-only as the default everywhere outside ranking pushes; placement modifiers 0 on all three groups, because a modifier on a discovery campaign is paid visibility for traffic the campaign has not yet qualified.
6. Set the daily budget from the family's discovery allocation share held in the T1 group S3 spend-governance columns. Discovery budgets are capped by design. Write the capped-by-design note into the T3 row in the same session, because that note is what stops gate G3 from firing the budget lever on a deliberate cap (see P10 step 1).
7. Seed terms: the family root in broad, plus validated modifier roots where the family already has proven structure. Never seed a root a sibling owns per the #33 declaration; that root is a pre-load negative, not a seed.
8. Apply the pre-load negative set in the same session per #29's pre-load standard. Four classes belong in it: taxonomy-known irrelevant terms for the family, every already-graduated exact term as negative exact, all brand terms as steering negatives so defense owns them exclusively, and sibling-owned roots per the declared ownership. Record the list in the T3 "Negation State" column so the #29 wall audit has its baseline. Skipping this step is the F10 and F12 root cause.
9. Advertised SKU = the D2-preferred variation for the family, never the parent by default.
10. Update the registry: T2 K8 "# Broad" increments on every seeded root; the new T3 row carries Objective Tag, owner, capped-by-design note and wall list; the T1 column "MT Cov Broad %" picks the campaign up through the syntax tag.
11. Log the creation in T6 DECISION LOG with the discovery prediction the standup carries: expected validated terms routed per $100 of spend by the first review date, stated as a figure with its review date (v4.0 Section 14, rule V1). A creation without a logged prediction is incomplete and does not deploy.
12. Deployment routes through #38: Tuesday by default. Same-day deployments still receive the full T+1 verification pass.

P2. Modified-broad deployment, the fence step-up (20-30 minutes per keyword)

1. Compute the root-containment read on the campaign's served terms for the window: the count of served terms containing the family root, divided by the count of all served terms, expressed as a percentage. Record this window's read and the prior window's read on the T3 row, with the delta in percentage points. **The firing threshold for this read is not stated in v4.0 and is not invented here; it is amendment item AM-2.** What fires the procedure is the governed condition beside it: a WAS percentage above the Discovery limit of v4.0 Section 3, family F11 "Waste & Negation", read after the truncation check of step 2.
2. Run the lever order before choosing the fence: v4.0 Section 9, rule L1 places budget ahead of bid ahead of placement ahead of negation ahead of structure. Read the T3 "Truncation Flag (G3)" first. A truncated campaign's waste percentage describes its worst hours, so the breach routes to the budget lever and re-reads next clean cycle. Modified broad is a structure move and is therefore the last rung, not the first.
3. Decide fence or list. If the campaign's waste is concentrated in a small number of enumerable terms, Routes B and C of P3 handle it and no fence is built. If waste is distributed across many terms that share the property of not containing the required word, the term list is not the problem and the fence is: proceed.
4. Add the modified version of the drifting keyword with the required word locked (for example +bamboo +sheets). The bid is carried from the paused keyword, not re-invented at deployment; any later step obeys the limits of v4.0 Section 9, rule L2.
5. Pause the unfenced version in the same change set so the two never serve simultaneously. Two live versions of the same root at two fence levels is the F9 condition and it splits the window's data across both.
6. Carry the paused keyword's negatives to the ad-group level where they were keyword-scoped, and verify that the graduated-exact walls still apply to the new keyword. A fence change must not silently drop a wall.
7. Send the exclusion-state notice to the #37 register the same day: a rule that steps the paused keyword's bid, or re-enables it, undoes the fence.
8. Log the T6 prediction: the root-containment read expected at the next pull, stated as a figure with its review date. The next pass scores it.

P3. The weekly three-route search-term triage (30-45 minutes per product)

> **This procedure is the shared definition.** SOP-13 runs steps 1-9 on phrase campaigns and adds its stem read. SOP-15 runs steps 1-9 on every auto group's keyword rows and adds its two ASIN routes. SOP-29 applies its execution standards to whatever Routes B and C produce and does not re-run the sort. Cite steps by number. A change here increments all four documents (Element 10, trigger 3).

1. Enter T2 only after T0 and T1 have been read; a pass that skipped the levels above it is invalid by construction (Data Architecture v1.0 Section 2, rule 1). Score last cycle first: read the T6 rows whose review dates land this cycle and record their V5 scores. Unscored actions do not accumulate.
2. Pull the window's search-term data. Amazon Ads console: Measurement and Reporting, then Sponsored ads reports, then Create report, then Sponsored Products, then Search term, summary, date range aligned to the pass window, filtered to the campaigns in scope. Working columns: Campaign Name, Ad Group Name, Targeting, Match Type, Customer Search Term, Impressions, Clicks, Spend, 7 Day Total Orders (#), 7 Day Conversion Rate, Cost Per Click (CPC). Confirm the pull's row count against the T3 campaign list before sorting; a campaign missing from the pull is a suppression, not a clean campaign.
3. Apply the two guards before sorting anything. First, the G12 MKL hygiene guard: a served term with no syntax tag and no relevancy score in the current Master Keyword List routes to MKL maintenance, not into a route. Second, the G11 ownership guard: a served term on a root another sibling owns per the #33 declaration routes to the wall build, not into Route A.
4. **Route A, converters.** Every term whose evidence line shows in-window orders becomes a candidate row to #30 with a complete evidence line: clicks, orders, window dates, source campaign, CPC and spend. This SOP does not test the term against the graduation bar, because the bar reads pooled evidence across every source that served the term and a single-source read cannot see the others; the pooling and the decision are #30's, made against v4.0 Section 8.9 scenario S-H1 and the validation rules for the term's class. This SOP's job ends at the routing.

5. **Route B, proven waste.** Terms carrying spend with zero in-window orders stage as negative exact with the evidence line in the reason string. The firing rule is v4.0's and the numeral it needs is not printed in v4.0; see amendment item AM-3 and SOP-29's independently logged AM-1. Rows go to #29 P4 for the execution standards, including the converting-term guard and the sufficiency check. **A Route B row without a complete evidence line is returned to this SOP rather than completed downstream from the feed**, because a term arriving without evidence is exactly the term nobody checked.
6. **Route C, drift clusters.** Irrelevant term families that share a root, meaning wrong category, wrong material or wrong intent, stage as one negative phrase on the cluster root, one fence for many future terms. The evidence line carries the root, the member count, and the combined clicks and spend across members. The wall level and the reverse test are #29's; this SOP supplies the cluster and its evidence.
7. **Head-term check.** For every priority term the window served, read the T2 campaign child rows' "Spend Share of KW %" against the K8 "# Exact" count. A discovery campaign carrying a materially dominant share of a priority term's clicks with no live exact fires ADD EXACT + RESTRUCTURE before any performance verdict, per v4.0 Section 4 rule X10 and Section 8.4 scenario S-R12. The receipt goes to #12 the same day and the term leaves Route A: a coverage gap is not a graduation candidate. Its clicks and spend also leave the yield base of P5 (see WE-2).
8. **Off-taxonomy converters.** A converting term below the relevancy bar or outside the taxonomy routes per v4.0 Section 8.9 scenario S-H6 to the Master Keyword List classification decision first. Either the taxonomy expands deliberately with a logged syntax tag and a test plan, or the term negates. Incidental orders do not create a market position.
9. **Stage, deploy, verify, log.** Every change stages as a bulk row per #38 with the verdict string and the scenario ID in the reason string; a staged row missing its scenario ID does not deploy. Tuesday deploy after the brand-management huddle; T+1 verification; a landed-value mismatch is the SOP-12 F10 class handled per #38. Route counts A, B and C per campaign close the pass and feed P5 and P7.

P3 route decision tree

Condition read on the served term	Route	Executed by
No syntax tag or no relevancy score in the current MKL	MKL maintenance and taxonomy pre-negation (gate G12)	Brand Management owns the list; #29 places the pre-negation
Root owned by a sibling per the #33 declaration	Wall build on the owned root; never a candidate	#29 places the wall; #33 owns the declaration
Priority term with a dominant discovery click share and no live exact (X10, S-R12)	ADD EXACT + RESTRUCTURE receipt, same day; term excluded from the yield base	#12 stands the exact up; this SOP walls afterwards at P6
Converting term, relevancy at or above bar, inside the taxonomy	Route A candidate to #30 with the full evidence line	#30 pools, validates and decides
Converting term below the relevancy bar or off-taxonomy (S-H6)	MKL classification decision first, then expand or negate	Brand Management decides the taxonomy; #29 negates
Spend with zero in-window orders, term enumerable and isolated	Route B, negative exact, with the evidence line	#29 P4 standards; staged by this SOP
Spend with zero in-window orders, terms share a root and an irrelevance class	Route C, one negative phrase on the cluster root	#29 P1 tree and reverse test; cluster supplied by this SOP
Evidence below the v4.0 Section 5 gate G2 read thresholds for the judgment	No route. The term stays in its source under a capped dated test (P8)	This SOP

P4. Defensive brand broad operations (15-25 minutes per brand)

1. One brand-broad campaign per brand, Objective = Defensive. Broad match on brand terms captures the misspellings and word-order variants exact cannot enumerate.
2. Bids sit at the floor discovered by #28's ladder. This SOP deploys the campaign-level rows #28 stages back; it never derives a ladder step. Budget runs to the Defensive utilization target of v4.0 Section 3, family F10 "Budget & Pacing", which is the one objective where budget is filled before bids are touched (v4.0 Section 8.6, scenario S-D3).
3. Brand terms are steering-negatived out of every generic discovery campaign on the product and live only here. A brand term serving in a Surface 1 campaign is the F10 condition and is fixed the same day.
4. Weekly read, three figures only: branded impression share held, the campaign's ACOS against the Defensive band of v4.0 Section 6 rule O6, and the incursion list. Discovery yield is not read on this surface; grading a defensive campaign on graduation rate is the process violation rule O6 names.
5. Competitor brand strings appearing in this campaign's served terms are incursion evidence. Route them to #16 with the served-term line. This SOP never opens, rotates or exits a conquest target.
6. Where an incrementality test is running on branded terms, the dated step-downs execute here on the schedule #35 sets. The continue-or-restore call belongs to #35 and v4.0; this SOP stages the step and records the capture read at each step.
7. Log each step in T6 with its prediction and review date, exactly as a bid change on any other surface.

P5. Graduation yield computation and the yield ledger (10-15 minutes per product)

1. Define the numerator: validated terms routed to #30 in the window from this campaign. Terms routed as head-term receipts are excluded, because a coverage gap is a structure failure and crediting it as discovery output rewards the failure.
2. Define the denominator: the campaign's window spend, less the spend attributable to any term excluded from the numerator under step 1. The exclusion is what keeps a structure failure from suppressing the grade of a campaign that was doing its job (see WE-2, where the correction moves the read from 1.69 to 2.10).
3. Compute yield as validated terms routed divided by window spend, multiplied by 100, expressed per $100 of discovery spend to two decimal places.
4. Write the figure into the T3 "Harvest Yield" column beside the trailing baseline held on the same row. Two figures on one row is the whole instrument: a yield with no baseline beside it cannot be read as a trend.
5. Credit is per source, not per term: a term validated in this campaign and also in an auto group credits both sources, because both paid for the discovery. #30 P7 computes the same figure from its own side and the two must reconcile; a divergence is a data-quality finding, not a scoring dispute.
6. A zero yield with spend active is not an automatic failure in a single window; it is the trigger for P13 when it repeats across the governance window. One zero window with a rising root-containment read and a live negation cadence is a campaign working through a fence change.

P6. Graduation handoff and the source-side wall confirmation (10 minutes per graduation)

1. Trigger: the SOP-30 P1 graduation log row reading APPROVE with this campaign named in its source list. The graduation is one event with three parts, decided, stood up and walled, and it is incomplete until all three land; #30 owns the chain and #29 owns the wall standard.

2. This SOP's half is confirmation, not construction: verify that the negative exact for the graduated term appears in this campaign's negative list at the level #29 selected, in the same session as the standup.
3. Verify again at the next pull: the term is absent from this campaign's served terms. A list entry that exists while the term still serves is a breach even though the list looks clean, which is why both sides are checked.
4. Remove the term from the next window's yield numerator and its spend from the denominator: a graduated term is no longer discovery output, and leaving it in inflates the campaign's grade with work it already banked.
5. An approved graduation whose wall has not landed in this campaign by the end of the session is the F2 condition and the state v4.0 Section 3 family F8 names GRADUATED-UNSTEERED, with v4.0 Section 8.9 scenario S-H4 firing at the next pass. Per conflict C12, this filter is S-H4 and not S-H6.

P7. The negation-cadence audit (10 minutes per product)

The cadence is this component's exclusive ownership: #29 owns whether a wall is correct, this SOP owns whether the pass fired at all. The two questions have different failure signatures and a clean wall audit on a stalled cadence looks identical to a healthy campaign.

1. Read four figures per campaign from this pass and the prior one: negatives added at Route B, cluster roots added at Route C, member count behind those roots, and window spend removed.
2. Compute spend removed as the sum, across negated terms, of each term's in-window clicks multiplied by its CPC. Counts without dollars cannot be scored; the dollar figure is what makes the pass auditable and is what the T6 prediction is written against.
3. Read the cadence state. Negatives added at zero across two consecutive passes with spend active is the leak-by-design condition and fires F1. Name the cause explicitly as one of two: the feed was absent and the pass was suppressed under gate G8, or the feed was present and the pass was skipped. These are different failures with different owners and recording them as one loses the distinction.
4. Write the counts and the dollar figure into T6 DECISION LOG, filling Level, Target, Rule / Scenario Fired, Change, Expected Outcome Window and Review Date. The expected outcome is stated as the WAS percentage expected at the next read, which the following pass scores.
5. Two consecutive F1 conditions on the same campaign escalate to the PPC Lead with the T6 history attached, because a stalled fence is a silent budget leak that no alarm rings for.

P8. Dated-test governance on broad (15 minutes at open, 15 at close)

1. Every broad test carries three things at creation: a start date, a click quota or a stop-loss date, and a capped bid. A test missing any of the three does not deploy. Doctrine v1.0 Section 13 states the standard plainly: a test without a click quota, a capped bid and a dated stop-loss is not a test.
2. The click quota and the capped bid are read from the verdict that opened the test, which carries the v4.0 sizing. This SOP records them into the T1 group S3 "Dated Test / Stop-Loss" column and the T2 K7 "Test Status" field; it does not set them.
3. No lever moves inside the window except budget assurance under gate G3. A test whose bid was stepped mid-window has no readable result and the window restarts.
4. Send the exclusion-state notice to the #37 register at open, so no platform rule steps the capped bid.
5. At expiry the verdict executes as one of two closed strings. Where the test succeeded, the term scales into the standing campaign at the verdict-carried bid. Where it did not, STOP-LOSS: CONCEDE OR DEFER executes with the deferral condition named in words: what evidence would re-open this test, and when it would be re-read.
6. The close is a scored V5 row against the open prediction. A test that closes without scoring its own prediction has taught the system nothing, which is the entire failure rule V1 exists to prevent.

P9. CAP SYNTAX receipt execution on discovery (15-20 minutes)

1. Trigger: the T1 row "Cap Status" reading BREACH with the staged delta from the instrument layer. The cap decision, the M5 reallocation and the envelope arithmetic belong to the instrument layer per conflict C3; this SOP receives the delta and never originates it.
2. Broad budgets in the breaching family reduce first. The reason is stated once so the ordering is not re-argued each cycle: discovery spend is the most deferrable spend in a capped family, because its output is optional in any single window while a ranking push loses velocity and a defense campaign loses share the day it is starved.
3. Compute the reduction as an arithmetic result of the cap, not as a step size: the family's spend minus its capped allowance is the dollars to remove, and the removal lands on the family's broad daily budget first. **v4.0 states no per-step limit on a cap-driven budget reduction**, and the plus-or-minus 50% limit of Section 9 rule L2 is a bid limit, not a budget limit; the absence is amendment item AM-4 and the reduction is executed to the cap.
4. Where the removal would take the family's broad budget below the level at which the campaign can serve, the excess routes back to the instrument layer for reallocation across the family's other discovery surfaces rather than being absorbed silently. A campaign budgeted below its serving floor is a paused campaign with extra steps.
5. Re-read two shares after the change lands: the family's share against its cap on the T1 row, and the product's match-type broad share against the M6 destination of v4.0 Section 10 rule M6. A cap enforcement that also returns the structure split to band is a stronger result than either alone and is worth recording (see WE-4).
6. Log the reduction in T6 with the expected weekly dollars freed and the review date at which the family's share is re-read.

P10. Bid and budget change mechanics on capped-by-design campaigns (10 minutes per change set)

1. **The capped-by-design distinction, stated once.** A discovery campaign that spends its full budget every day is not truncated; it is capped, deliberately, and the cap is the instrument. The T3 row carries the capped-by-design note written at P1 step 6, and gate G3 reads that note before firing the budget lever. A campaign missing the note reads as truncated and will receive a budget rise it was never meant to have, which is the F11 condition.
2. Every bid or budget change exists first as a staged bulk row per #38 with prior value, new value, and a reason string carrying the verdict string and the scenario ID. No console edits outside the staged set.
3. Step sizes, floors and per-change limits are read from v4.0 Section 9 rule L2 at execution time; this SOP never restates them. A staged step exceeding an L2 limit fails staging quality control and returns to the pass.
4. The lever order of v4.0 Section 9 rule L1 governs every change set: budget where gate G3 is failing, then bid, then placement, then negation, then structure, then status. On a capped-by-design campaign the first rung is usually unavailable by design, which pushes discovery work toward negation and structure and is exactly the intended behaviour.
5. On a T+1 mismatch, redeploy the delta the same day, mark the week's reads on the affected rows void, and name a root cause. A console edit found by the verification pass names the actor and routes to change control at #37.

P11. Event-window overlay and the recovery-mode sentinel floor (10 minutes at arm, 10 at release)

1. Arm on the #32 event plan landing in the T2 K2 "Event Mode" column with its dates. Every broad row on the product carries the event flag through the window plus the guard period that follows it.

2. Inside the window, structure changes freeze unless the event plan names them: no creations, no fence changes, no retirements, no new dated tests. v4.0 Section 8.12 scenario S-E1 states the rule for the pre-event window directly, that no new experiments enter it. Budget assurance under gate G3 and the event plan's own staged actions remain live.
3. Reactive negation from Routes B and C continues inside the window and is not frozen: fencing waste is not an experiment. Route C cluster walls deploy normally.
4. Event-week rates leave the trend reads that feed P2 and P7 through the guard period, per gate G10 event exclusion. The weekly pass reads affected rows for gates only.
5. **Recovery Mode.** On entry, discovery budgets floor at sentinel levels rather than going to zero: the campaigns stay on, small, so the term histories and the negation lists stay live and the account is not rebuilding discovery from nothing at exit. The sentinel level itself is not a number this SOP owns; it arrives with the entry log from the instrument layer. Human-review thresholds tighten on entry and every discovery budget move is confirmed against the tightened tier.
6. At release, resume from the pre-window budgets and fence states unless the event reconciliation justifies new ones. The resumption is a staged, logged set, not a memory exercise.

P12. Variation re-point on broad (10 minutes)

1. Trigger: a RE-POINT VARIATION verdict, or gate G6 flipping RED on the targeted variation with D2 naming the ranked backup.
2. Stage the advertised-SKU edit to the D2-named variation and deploy the same day. A RED-driven re-point does not wait for Tuesday.
3. Log the re-point in T6 as a product event: the CVR baseline resets at re-point, the trend clock restarts, and the row is read for gates only until the v4.0 sufficiency window refills.
4. The yield baseline on the T3 row resets with it. Comparing a post-re-point yield against a pre-re-point baseline compares two different products' demand and produces a false trend.
5. On restock, re-point back and log it the same way. Where D2 carries no ranked backup for the family, the campaign holds on the preferred SKU with the D2-pending flag and the gap escalates the same day; this SOP does not choose a backup variation.

P13. Restructure or retire, pause, archive and re-entry (25-40 minutes)

1. Trigger: zero graduation yield across the governance window with spend active, or an ACOS outside the Discovery band of v4.0 Section 6 rule O6 with Routes B and C exhausted.
2. Evaluate in this order and record each answer, because retirement chosen before the first three have been tested is a diagnosis skip:
    - **Seed refresh.** Has the family produced new validated modifier roots since the last seeding? A campaign seeded once and never re-seeded is discovering against a stale map.
    - **Negation audit.** Is the fence the problem? Read the root-containment figures from P2 step 1 and the cadence figures from P7. A campaign whose cadence stalled is not a failed campaign; it is an unrun procedure.
    - **Budget rebalance.** Would the same dollars yield more in the family's other discovery surface? The comparison is the two campaigns' yield figures side by side on their T3 rows, which is why P5 step 4 requires the baseline to sit beside the read.
    - **Only then, retirement.**
3. Retirement is PAUSED-PENDING with both a reason and a named re-entry condition in T2 and T3. A pause without both is incomplete and returns to the pass. Paused rows appear on the weekly pass as a re-entry check only: the pass tests the stated criterion and either re-enters or holds. It never re-litigates the pause.
4. Where two broad campaigns exist on one family root with no declared split, consolidate rather than retire. Survivor precedence follows SOP-12 P7: the instance with the declared role, then the instance winning on the evidence read, then the instance with the longer clean history. The loser's negation lists carry to the survivor in full; a consolidation that drops a wall re-opens every term the wall fenced.
5. Archive only campaigns paused beyond the governance period with no re-entry condition met. Archiving is irreversible and sits in the human-mandatory tier; the archive entry in T6 closes the campaign's history with its final registry state and its final yield figure.
6. Every outcome of this procedure is a scored T6 prediction. A retirement predicts that the family's validated-term flow holds or improves at lower spend, and the next cycle scores it. A retirement that silently reduced discovery output is a wrong-direction score with a written root cause.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds and scenario logic live in v4.0. This table maps the complete 25-string closed vocabulary, quoted verbatim from the Keyword Decision Record Template v3.0 sheet "Bands & Verdicts Reference", to what this SOP does when each string lands on a broad-relevant row. Where execution belongs to an adjacent component the boundary is stated, so no verdict falls into a seam. Per conflict C10, two strings carry a cross-map to the differently worded v4.0 Section 13.1 "The Closed Verdict Vocabulary" action so neither citation is dropped.

Verdict string (KDR Template v3.0, verbatim)	Broad-campaign execution
QUARANTINE: DATA REFRESH	No route runs on the row and no negative stages against it. The term stays in its source for the window; the refresh task carries the row forward to the next clean cycle (gate G1).
UNLOCK TRAFFIC	The bid step executes per v4.0 Section 9 rule L2 on the named broad keyword; no rate judgment until gate G2 sufficiency. The budget lever is not the response on a capped-by-design campaign (P10 step 1).
FIX BUDGET (TRUNCATED)	Applies only where the T3 row does **not** carry the capped-by-design note. A deliberate discovery cap is exempt; treating it as truncation is the F11 condition.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	P8 step 3: no lever moves except budget assurance; the test runs to its click quota at the capped bid it opened with.
FUND TEST (UNFUNDED PRIMARY)	P8 opens the dated test on the relevant underfunded term at the capped bid the verdict names, with the click quota and stop-loss recorded in T1 group S3 and T2 K7. Never an uncapped bid rise.
RANK PUSH: FUND SIZED PUSH	No broad execution. Ranking runs on exact only, because rank credit must land on a chosen term at a chosen variation (Doctrine v1.0 Section 3.1). This verdict landing on a broad row is a routing defect; the receipt goes to #12.
RECOVERY PUSH	No broad execution; #12 stands up and funds the recovery. This SOP contributes the steering wall after the exact goes live, per P6.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	No broad execution. Required-clicks funding runs on the term's exact campaign at #12; feeding a discovery campaign to close a rank gap buys clicks the campaign chooses, not the clicks the plan sized.
FIX CONVERSION FIRST; DEFER PUSH	The family's discovery holds at current spend while the fix runs. No new seeds enter a family in chronic CONVERSION diagnosis (v4.0 Section 8.2, scenario S-Y5); P1 step 1 enforces this at standup.
PLACEMENT FIX FIRST	Bids hold; #28 executes the placement program. Broad campaigns carry modifiers at 0 by default (P1 step 5), so this rarely lands; where it does, this SOP deploys only the campaign-level rows #28 stages back.
RE-POINT VARIATION	P12 in full, same day on a gate G6 RED.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	No broad execution. Rank is not held in broad, so there is no taper ladder here; the exact's taper is #12 P4 and the ladder steps are #28's.
HOLD; DEFEND LIGHTLY AT RANK	Surface 3 only: the brand-broad campaign holds at its recorded floor with no change this cycle. Drift from the held position is what v4.0 watches, not this SOP.
DEFEND AT FLOOR	Surface 3 only: no bid change unless drift from the recorded discovered floor is detected, in which case the bid restores to the floor value in one staged step (P4 step 2).
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	The CAP SYNTAX half executes here first: broad budgets in the breaching family reduce first, per P9. The bid-down half executes on the named campaign per P10 against the ceiling value held in the T2 K6 economics block.
INCREMENTALITY LADDER	Surface 3 only, and only as execution: the dated step-downs deploy here on #35's schedule. Measurement design and the continue-or-restore call belong to #35 and v4.0.
ADD EXACT + RESTRUCTURE	Receipt routed to #12 the same day (P3 step 7). This SOP contributes the steering negative after the exact stands up, and excludes the term's clicks and spend from the yield base at P5 step 2.
CONSOLIDATE DUPLICATES	Two broad campaigns on one family root with no declared split consolidate per P13 step 4, applying the SOP-12 P7 survivor precedence. The loser's negation lists carry over in full.
GRADUATE FROM DISCOVERY	Issued by #30, never here. This SOP's half is P6: confirm the wall landed in this campaign, verify absence from served terms at the next pull, and remove the term from the next yield base. Cross-map per conflict C10: v4.0 Section 13.1 issues this action as "GRADUATE + STEER".
NEGATE (PRE-LOAD / REACTIVE / STEERING)	All three modes touch this SOP: pre-load at P1 step 8, reactive from P3 Routes B and C, steering confirmed at P6. Standards, wall levels and the reverse test are #29's; this SOP supplies routes with evidence lines and owns the cadence read at P7. Cross-map per conflict C10: v4.0 Section 13.1 issues this action as "CUT / NEGATE".
ENTER SBV ARBITRAGE	No broad execution; the format pivot belongs to #17-19 and the paired SP step-down to #12.
OPEN / ROTATE / EXIT CONQUEST	No broad execution. Competitor-brand strings in the Surface 3 served terms route to #16 as incursion evidence (P4 step 5); #16 decides.
STOP-LOSS: CONCEDE OR DEFER	P8 step 5: the dated test closes at its stop-loss with the deferral condition named in words, and the close is scored against the open prediction.
PAUSED-PENDING (REASON + RE-ENTRY)	P13 step 3: reason and re-entry condition both recorded in T2 and T3. The weekly pass tests the criterion and never re-litigates the pause.
LEAVE TO HARVEST	No structural change. The term keeps converting in broad below the graduation bar and is re-evaluated on the re-eval date held in #30's pool, which is where the aging rule lives. This SOP does not set or move that date.

6. Worked Examples

All four examples are illustrative. No live product economics spine was attached to this build session, so every example uses the standard fictional base: AOV $80.00, contribution $28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%. Live values come from the T0 zone C named cells. Every figure below is computed on the page and every arithmetic assertion was executed in code before publication.

WE-1 · The clean weekly pass: three routes, a WAS breach that fires the pass, and a yield read

Setup. Family broad "Bamboo", Surface 1, Objective = Discovery, daily budget $25.00 capped by design. Fourteen-day window: 9,800 impressions, 92 clicks, spend $158.24, 4 orders. Sixty-one served terms. Gates clear: inventory GREEN on the D2-preferred variation, no armed event window, Master Keyword List refreshed inside its cadence.

Derived reads. CPC = 158.24 / 92 = $1.72. CTR = 92 / 9,800 = 0.94%. CVR = 4 / 92 = 4.35%. Sales = 4 x $80.00 = $320.00. ACOS = 158.24 / 320.00 = 49.5%, which sits above the Discovery Max-ACOS band of v4.0 Section 3, family F1 "Product Economics", and is correctly *not* the grade: v4.0 Section 6 rule O6 assigns Discovery the metric set of cost per discovered converter, WAS percentage and graduation rate.

The read that fires the pass. Two terms carried the window's orders, taking 11 and 9 clicks: 20 clicks at $1.72 = $34.40 of converting spend. Spend on zero-order terms = 158.24 - 34.40 = $123.84. WAS = 123.84 / 158.24 = 78.3%, above the Discovery limit of v4.0 Section 3, family F11 "Waste & Negation", which fires the negation pass of Section 8.9 scenario S-H2. Truncation check first per rule L1: in-budget reads full against a deliberate cap and the T3 row carries the capped-by-design note, so gate G3 does not divert to the budget lever and the breach is real.

P3 routes over 61 served terms. Route A: both converting terms route to #30 with full evidence lines. "bamboo sheets split king" (11 clicks, 2 orders, CPC $1.72, source campaign named) and "bamboo sheet set deep pocket king" (9 clicks, 1 order). Neither reaches the graduation bar of Section 8.9 scenario S-H1 on this source alone; #30 pools the same terms from the auto close-match group and the first clears at three orders combined. This SOP does not make that test. Route B: four isolated terms with spend and zero orders, "bamboo shoots plant" (14 clicks), "bamboo flooring cleaner" (8), "bamboo pillow" (6), "bamboo cutting board" (5): 33 clicks at $1.72 = $56.76. Route C: one drift cluster on the root "crib" with three members ("crib sheets bamboo", "bamboo crib sheet", "crib bamboo bedding"), 9 clicks combined at $1.72 = $15.48, staged as one negative phrase. Head-term check: no single served term carries a dominant share of any priority term's clicks; clean.

P7 cadence arithmetic. Negatives added: 4 exact, 1 phrase root fencing 3 terms. Spend removed = 56.76 + 15.48 = $72.24. Projected next-window base = 158.24 - 72.24 = $86.00; remaining zero-order spend = 123.84 - 72.24 = $51.60; projected WAS = 51.60 / 86.00 = 60.0%. **One pass does not close a 78.3% WAS.** That is the finding, not a failure: the Discovery limit is reached by a cadence across passes, which is why v4.0's rule is a repeating pass and not a cleanup task.

P5 yield. 2 validated terms routed / $158.24 x 100 = **1.26 per $100** of discovery spend, written to the T3 "Harvest Yield" column beside the trailing baseline. The campaign is doing its job despite a 49.5% ACOS, because the two routed terms move converting demand up the control gradient where its economics improve.
Cross-document note: SOP-30 v1.0 and SOP-13 v2.0 cite a 1.27 baseline for this same window, computed on a spend figure rounded to $158; SOP-29 v2.0 WE-5 states the same window at $158.24, which yields 1.26. Reported as a conflict in the closing section; both computations are shown so neither citation is silently dropped.

T6 entries. Two routed with evidence lines; 4 negative exact and 1 negative phrase fencing 3 terms; $72.24 removed; yield 1.26; prediction, WAS at or below 60.0% at the next read with the crib cluster fenced; review date next pass.

WE-2 · The edge case that goes wrong: a head term found living inside the yield read

Setup. Family broad "Cooling", Surface 1, fourteen-day window: 7,400 impressions, 118 clicks, spend $237.18, 5 orders. CPC = 237.18 / 118 = $2.01. CTR = 118 / 7,400 = 1.59%. CVR = 5 / 118 = 4.24%. Sales = $400.00. ACOS = 237.18 / 400.00 = 59.3%.

What the operator did first, and why it was wrong. Four terms carried orders, so the operator routed four Route A candidates and closed the pass with a yield of 4 / 237.18 x 100 = 1.69 per $100. The pass was staged and would have deployed on that read.

What P3 step 7 caught. One of the four, "cooling sheets king", is a priority oversight-class term. Its T2 campaign child rows for the window read: this broad campaign 47 clicks, the auto loose-match group 19, the family phrase campaign 9, exact 0. Total 75 clicks on the term across all types. Broad's share = 47 / 75 = 62.7%, materially dominant, with K8 "# Exact" reading 0. That is v4.0 Section 4 rule X10, the auto-dominance read, resolving to Section 8.4 scenario S-R12, which requires ADD EXACT + RESTRUCTURE *before any performance verdict*. A term with no exact coverage is a structure failure. It is not a graduation candidate, and crediting it as discovery output would reward the failure that produced it.

The correction, computed. The head term leaves Route A: candidates fall from 4 to 3. Its clicks and spend leave the yield base: 47 clicks x $2.01 = $94.47 removed, base = 237.18 - 94.47 = $142.71. Corrected yield = 3 / 142.71 x 100 = **2.10 per $100**, against the 1.69 first computed, a difference of 0.41. **The campaign was being graded down for a failure that belonged to the account's structure, not to its discovery work.**

What still goes wrong, and how it is handled. The receipt goes to #12 the same day, Monday. The exact standup lands Thursday, because the funding release carried that date. The steering wall cannot be built until the exact is live, so this campaign keeps serving the head term for three days. That exposure is computed and recorded rather than ignored: 47 clicks over 14 days = 3.36 clicks per day, so three days is roughly 10.07 clicks at $2.01 = **$20.24 of double-buying**. The exposure window is written into T6 with the receipt so the #29 wall audit can see it, and the next window's yield base excludes the term from the wall date rather than from the window start.

Had the check been skipped. The pass would have closed at yield 1.69, the head term would have entered #30's queue as a graduation candidate, and #30 would have returned it as a rejection with the structure finding attached, one week later. The failure mode is F7, a yield read computed on a contaminated base, and its detection signal is exactly the mismatch between a Route A count and the head-term check.

WE-3 · Modified broad: choosing a fence over a list

Setup. Family broad "Satin", Surface 1, fourteen-day window: 125 clicks, spend $212.50, CPC = 212.50 / 125 = $1.70, 6 orders on 68 converting clicks (CVR 8.82% on the converting set). Spend on zero-order terms = $96.90, which is 57 clicks. WAS = 96.90 / 212.50 = 45.6%, above the Discovery limit of v4.0 Section 3, family F11, firing scenario S-H2.

Lever order, executed rather than assumed (rule L1). Truncation first: in-budget reads 96%, above the truncation line of v4.0 Section 5 gate G3, so the campaign's waste percentage describes its real behaviour and not its worst hours. The breach stands. Bid and placement are not the levers for a relevance problem. That leaves negation and structure.

The read that decides fence versus list. The 57 zero-order clicks are spread across 39 distinct terms, averaging 1.46 clicks each. Enumerating 39 negative exacts fences 39 strings against a term space that regenerates every window. The supporting evidence is the root-containment read of P2 step 1: served terms containing the family root were 44 of 61 last window (72.1%) and 33 of 68 this window (48.5%), a fall of 23.6 percentage points. **The term list is not the problem; the fence is.** Recorded honestly: the fall itself has no firing threshold in v4.0 and none is invented here (amendment item AM-2). What fires the procedure is the governed WAS breach; the containment read is the evidence that selects the response.

Execution. Add "+satin +sheets" as the modified version of the drifting keyword at the paused keyword's bid of $1.35, carried and not re-invented. Pause the unfenced version in the same change set so the two never serve simultaneously. Because the bid is unchanged, no step limit of v4.0 Section 9 rule L2 is engaged and the $0.50 floor is not approached. Keyword-scoped negatives move to the ad-group level; the graduated-exact walls are re-verified against the new keyword. Exclusion notice to the #37 register the same day, so no platform rule re-enables the paused version or steps the carried bid.

T6 prediction. Root containment at the next pull stated as a figure, with the WAS percentage expected alongside it, review date next pass. If containment recovers and WAS falls, the fence worked and the campaign continues on Surface 2. If containment recovers and WAS holds, the waste was never a drift problem and the pass returns to Routes B and C with the fence retained as a separate finding.

WE-4 · CAP SYNTAX receipt: discovery reduces first, and the structure split comes back into band

Setup. Product weekly tracked spend $1,840.00. The "Cooling" family is designated Secondary. Its spend reads $552.00, which is 552.00 / 1,840.00 = 30.0% of tracked spend, above the secondary reference share of v4.0 Section 3, family F13 "Syntax Rollups". The T1 row "Cap Status" reads BREACH and the instrument layer stages the delta. This SOP receives it and never originates it (conflict C3).

The arithmetic of the receipt. Capped allowance = 25.0% x $1,840.00 = $460.00. Excess = 552.00 - 460.00 = **$92.00** to remove. The family's spend splits exact $268.00, broad $168.00, phrase $71.00, auto $45.00, summing to $552.00.

Where the reduction lands, and why. Broad reduces first, because discovery output is optional in any single window while a ranking push loses velocity and a defense campaign loses share the day it is starved. Family broad falls 168.00 - 92.00 = $76.00 per week, which is 76.00 / 7 = **$10.86 per day**, down from $24.00. The reduction is 54.8% of the prior budget in a single step. That magnitude is an arithmetic result of the cap and not a step choice: v4.0 states no per-step limit on a cap-driven budget reduction, and the plus-or-minus 50% limit of Section 9 rule L2 is a bid limit. The absence is logged as amendment item AM-4 and the reduction is executed to the cap.

Serving-floor check (P9 step 4). At the family's blended discovery CPC of $1.72, $10.86 per day buys about 6.31 clicks per day. That is above the campaign's serving floor, so no dollars route back to the instrument layer for reallocation. Had it fallen below, the excess would return rather than being absorbed into a campaign budgeted below the level at which it can serve.

The two shares re-read after the change (P9 step 5). First, the family: broad is now 76.00 / 460.00 = 16.5% of the capped family, and the family sits at its 25.0% cap. Second, the product: total broad spend was $404.00, which is 404.00 / 1,840.00 = 22.0% of tracked spend and above the M6 broad destination of v4.0 Section 10 rule M6; after the removal it is 404.00 - 92.00 = $312.00, which is 312.00 / 1,840.00 = **17.0%** and inside the 15-20% band. The cap enforcement and the structure correction landed in one move, and the T6 entry records both, because a change that fixed two findings and reported one has under-reported its own value.

7. Failure Modes and Escalation

Twelve modes. Each carries its detection signal, its response, and the account-history incident that motivated it where one exists. The five incidents recorded in v1.0 are carried forward verbatim in substance, because they are institutional knowledge that no other document in the library holds.

Failure	Detection signal	Response	Motivating incident
F1. Negation cadence stalls	P7 step 3: negatives added reads zero across two consecutive passes with spend active	P3 executes immediately on the current window. Name the cause as one of exactly two: feed absent and pass suppressed under gate G8, or feed present and pass skipped. Two consecutive occurrences escalate to the PPC Lead with T6 history	Carried from v1.0 F1: the leak-by-design condition. A stalled fence is a silent budget leak and no alarm rings for it
F2. Graduated term still serving in broad	An approved graduation whose term appears in this campaign's served terms at the next pull; the v4.0 family F8 state GRADUATED-UNSTEERED, firing scenario S-H4	Wall built the same day at the level #29 selects; the #29 wall audit adds this campaign to its population; root cause names which session skipped its step	Carried from v1.0 F2. The term buys itself twice at two control levels while the exact fights the source it came from
F3. Head term carried by broad	P3 step 7: dominant discovery click share on a priority term with K8 "# Exact" reading 0 (rule X10, scenario S-R12)	ADD EXACT + RESTRUCTURE receipt to #12 the same day, outside the weekly rhythm; term leaves Route A; exposure window computed and recorded until the wall lands	Carried from v1.0 F3: the account's named failure class. A head term ranked by campaigns that choose their own traffic is unranked by control
F4. Drift explosion after a relevance-model shift	Root-containment share falling sharply window over window across several families at once, not one	P2 on the affected keywords; Route C clusters staged aggressively; the platform shift is entered in T5 PRODUCT LOG because it contaminates every trend read on the product, not just this campaign's	Carried from v1.0 F4. Semantic matching widened broad drift materially across the 2024-2026 period (Doctrine v1.0 Section 3.3)

Failure	Detection signal	Response	Motivating incident
F5. Discovery cannibalizing the family cap	Family discovery spend share above its allocation inside a capped family, read on the T1 row	The CAP SYNTAX receipt executes per P9: broad budgets reduce first, to the cap, with both shares re-read after	Carried from v1.0 F5. Discovery spending the primary family's rank velocity is the pattern the syntax caps exist to stop
F6. Pass run on a partial feed	The search-term pull's campaign count is short against the T3 campaign list, and routes were produced anyway	Void the window's routes for the missing campaigns and record the suppression under gate G8. Deployed negatives from a partial pull are reversed. A second consecutive miss escalates to Data Operations with the manifest history	New at v2.0. v1.0 stated that an absent feed suppresses P3 but gave no detection signal, so a partial feed read as a clean one
F7. Yield read on a contaminated base	A Route A count that includes a head-term receipt or an already-graduated term; or a yield figure whose denominator was not reduced by the excluded terms' spend	Recompute per P5 steps 1-3 with both figures shown, correct the T3 row, and re-score any T6 prediction written against the wrong figure	New at v2.0, from WE-2: the correction moved the read from 1.69 to 2.10 on one campaign in one window
F8. Second broad campaign on one family root	Registry read showing more than one live broad campaign against the same family root with no declared split	Consolidate per P13 step 4 with SOP-12 P7 survivor precedence; the loser's negation lists carry over in full before the pause	New at v2.0. Doctrine v1.0 Section 5.3 names thematic micro-splitting of broad as illegitimate because it destroys the data density that is broad's only virtue
F9. Fenced and unfenced versions serving together	Both the modified and the open version of one root returning served terms in the same window	Pause the open version immediately; the window's reads on both are void because the data split across two fence levels; root cause names the P2 step 5 omission	New at v2.0. The pause and the add must ride one change set; a two-session deployment leaves the gap open for a full window
F10. Brand terms serving in generic discovery	Brand strings appearing in a Surface 1 campaign's served terms	Steering negative placed the same day per #29; the brand term belongs only on Surface 3. Root cause names whether the P1 step 8 pre-load class was missing or the brand list changed without a re-attach	New at v2.0. Defense cannot hold a floor on terms a discovery campaign is also bidding for at discovery prices
F11. Capped-by-design read as truncated	A FIX BUDGET (TRUNCATED) verdict landing on a discovery campaign whose T3 row lacks the capped-by-design note	Restore the note, reverse the budget rise if it deployed, and re-read the campaign's rates at the intended cap. The missing note, not the verdict, is the defect	New at v2.0. v1.0 stated the exemption but placed the note nowhere, so the exemption depended on the reader remembering it
F12. Sibling-owned root seeded here	A seed term or served term on a root the #33 declaration assigns to another family product, with no stewardship grant	Remove the seed, place negative phrase on the owned root, and run the gate G11 check before any further scale action on the root. The declaration is #33's and is not re-argued in this pass	New at v2.0. Two family products buying one root bid against each other with the same money

Escalation. The same failure mode twice in consecutive cycles on the same scope escalates to the PPC Lead with the T6 history attached. F1 and F4 escalate on the second occurrence because both are silent; F3 and F12 escalate on the first, because an uncovered head term and a contested root both compound daily. Anything touching the human-mandatory tier of v4.0 Section 12 rule P5, which covers structural changes, gate failures and every deployment confirm, is confirmed per the Authority Matrix before deployment, without exception.

8. Outputs and Artifacts

Gates clear with artifacts, not assertions. Every row below names what is produced, where it is stored, how long it is held, and who consumes it downstream.

Event	Named artifact	Storage and retention	Downstream consumer
Weekly pass (P3)	Route A candidate list, one row per term with all serving sources attached and a complete evidence line; Route B and Route C staged sets with evidence lines in the reason strings; route counts per campaign	Candidate list and staged sets held with the #38 session record for the cycle; counts persist in T6 DECISION LOG, append-only	#30 P1 harvest pass consumes Route A; #29 P4 consumes Routes B and C; #38 consumes the staged set
Yield computation (P5)	The per-campaign yield figure with its trailing baseline beside it	T3 CAMPAIGN BOARD "Harvest Yield" column, current row; history in T6	#30 P7 reconciles against its own per-source credit; T1 and the audit grade the discovery surfaces
Cadence audit (P7)	Negatives added, cluster roots added, member count behind each root, and window spend removed in dollars	T6 DECISION LOG with the expected WAS figure and its review date	The #29 monthly wall audit reads the cadence beside its own wall correctness read; V5 scores the WAS prediction
Structure events (P1, P2, P13)	T3 row created or updated with Objective Tag, owner, capped-by-design note and wall list; T2 K8 registry count change; the T6 prediction row	T3 current row; T5 PRODUCT LOG entry for material structure events; T6 append-only	#29 wall audit takes the wall list as its baseline; #38 deploys; the next pass scores the prediction
Head-term receipt (P3 step 7)	The ADD EXACT + RESTRUCTURE receipt naming the term, the source campaign, the click share, the exact count, and the computed exposure window	#12's same-day queue; copy in T6 with the receipt date	#12 P1 or P2 stands the exact up; #29 walls afterwards; the exposure figure feeds the wall audit

Event	Named artifact	Storage and retention	Downstream consumer
Graduation confirmation (P6)	The wall confirmation line and the next-pull absence check for the graduated term	T3 "Negation State" column; the line is written back into #30's chain record	#30 cannot close its atomic chain without it; #29 audits both sides
Dated tests (P8)	The test record: start date, click quota or stop-loss date, capped bid, and at close the verdict string with its scored prediction	T1 group S3 "Dated Test / Stop-Loss"; T2 K7 "Test Status"; close scored in T6	The instrument layer's M5 queue reads closed tests as freed or committed dollars; V5 scores the close
CAP SYNTAX execution (P9)	The reduction record: excess computed, dollars removed, new daily budget, serving-floor check, and both re-read shares	T3 daily budget current; T1 row share recomputed; T6 with expected weekly dollars freed	The instrument layer reconciles the envelope; T1 confirms the family inside cap; the A4 structure-targets view confirms the M6 share
Every deployment	The staged bulk set in #38 format with verdict string and scenario ID in every reason string, and its T+1 verification report	Held with the #38 session record; mismatches recorded with their root cause	#38 owns the verification; a mismatch is the SOP-12 F10 class
Session close	The updated Ledger row for component #14 in copy-paste form	Pasted into the Cross-Reference Ledger, version raised per its Section 6	The next build session, which needs this component's ownership declaration and not this document

The weekly gate artifact for this component is the deployed change set carrying a complete evidence line on every negation row, its T+1 verification report, and the yield figure written beside its baseline. A pass that produced routes but no yield figure, or a yield figure with no baseline, has not closed.

9. Skill-Conversion Block (Element 9 handoff)

Block	Specification, specific to this component
Input bindings (catalog IDs and named cells)	A2 deployment state; A3 pacing and in-budget percentage plus the T3 truncation flag; A5 threshold flags; A6 objective tag and Expected CPC, riding provisional while the per-term value is absent; B5 coverage counts, specifically T2 K8 "# Exact", "# Phrase", "# Broad", "# Auto"; B8 economics chain from the T0 zone C named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC; B9 week-over-week deltas; D2 variation map with ranked backups; D7 decision log; the Master Keyword List with syntax tags and relevancy scores; the #33 term-ownership declaration; and the Sponsored Products and Sponsored Brands search-term reports plus STIS joined to T2 campaign child rows. **An absent search-term join suppresses P3, P5 and P7 for the week**; the skill records the suppression and never improvises around a missing feed.
Deterministic logic	Six closed-form computations and no more. (1) Route sort: the P3 decision tree read in table order, first match wins. (2) Root containment = served terms containing the family root divided by all served terms. (3) Yield = validated terms routed divided by window spend, times 100, on the base corrected per P5 steps 1-2. (4) Spend removed = the sum over negated terms of clicks times CPC. (5) WAS recomputation = remaining zero-order spend divided by the reduced base. (6) Head-term share = a campaign's clicks on a term divided by that term's clicks across all campaign types, read from the T2 child rows. The skill executes the tree and the Element 5 mapping; it computes no threshold, issues no verdict, and version-pins the Criteria System per the v4.0 Section 14 sync rule.
Human checkpoints (from the Authority Matrix v2.0 rows)	Per the row "Weekly verdicts (established US)": PPC Lead decides, incoming operator executes assigned scope, PPC Manager executes, brand-management owner reviews at the daily huddle, which is review and not veto. Per "Structure creation (all SOPs)": PPC Lead decides or per SOP scope, incoming operator and PPC Manager execute. Per "Bulk deployments (#38)": PPC Lead reviews, operators execute own scope. Per "Automation rules (#37)": PPC Lead decides, incoming operator executes and reviews. **Human-mandatory in this SOP:** campaign creation (P1), the fence decision (P2 step 3), every head-term receipt (P3 step 7), dated-test opens and closes (P8), CAP SYNTAX execution (P9), retirement and archive (P13), and every deployment confirm. **AI-executable with audit:** route sorting with evidence-line drafting, yield and cadence computation, root-containment reads, wall-confirmation checks, staged bulk drafting, and T6 drafting. Matrix rule 5 applies: an unfilled role's decision right reverts to the PPC Lead and its execution to the PPC Manager.
Cadence	Weekly with the Monday cascade for P3, P5 and P7. Same day on arrival for head-term receipts and graduation wall confirmations, outside the weekly rhythm. Same day on gate G6 RED for P12 and on an armed window for P11. Daily pacing glance on any live dated test. Monthly, the cadence figures hand to the #29 wall audit. Quarterly, the Surface 3 defense floor re-tests per #28.
Alarm specification	F1 to F12 each carry a Section 7 detection signal and an evidence-line format of row, signal value, and gate or rule reference. Routing: F3 and F12 same day to the operator queue with the PPC Lead copied, because both compound daily; F2, F9 and F10 same day into the change set; F6, F7, F8 and F11 into the weekly pass worklist; F1 and F4 to the PPC Lead with T6 history on the second occurrence; F5 to the instrument layer with the T1 share figures attached.
Skill outputs	Drafted Route A candidate rows with complete evidence lines; drafted Route B and Route C staged sets in #38 format ready for #29's standards; the yield figure with its baseline; the cadence figures with the dollar total; drafted T6 rows each carrying a prediction stated as a figure with a review date; F1 to F12 alarms with evidence lines; and the V5 scoring draft for broad rows due this cycle. Everything is drafted for confirmation at the tier the Authority Matrix sets; nothing deploys from the skill.

10. Version Governance

Owner: PPC Lead. Registered in the Supersession Map at publication.

Review triggers specific to this component. (1) Any Criteria System version increment, at which point the Element 5 mapping re-validates against the new string set before any pass runs on it. (2) **Any change to P3, which increments SOP-13, SOP-14, SOP-15 and SOP-29 together**, because three documents run the machinery this one defines and a fourth sets the standards its outputs execute against; a change on one side without the paired increment is a version-governance breach. (3) Ratification of the verdict-vocabulary conflict C10, at which point Element 5 re-issues against the ratified string set and amendment item AM-5 closes. (4) Any change to the #29 pre-load standard, which changes P1 step 8 without changing anything else in this document. (5) The landing of the unblended per-variation margin table, which retires the provisional flag riding every ceiling-referencing action here. (6) Any resolution of conflict C9, gate numbering, which changes how Element 3 names three gates. (7) Any build session that returns this component's Ledger row.

v2.0 change note. Supersedes v1.0 in full. Procedures expanded from 6 to 13, adding P5 yield computation as a procedure rather than a step, P6 graduation wall confirmation, P7 the cadence audit, P9 CAP SYNTAX execution, P10 capped-by-design change mechanics, P11 event and recovery overlays, and P12 variation re-point; the v1.0 P6 restructure-or-retire becomes P13 and gains the consolidation and archive branches. Verdict mapping expanded from 7 rows to the full 25-string closed vocabulary with boundary statements and two cross-mapped rows. Worked examples expanded from 1 to 4, including the mandatory edge case that goes wrong. Failure modes expanded from 5 to 12, with all five v1.0 incidents carried forward. Elements 11 and 12 added. Three citation defects corrected: the v1.0 naming citation to a Criteria System section that does not exist (now SOP-12 v2.0 Appendix A, logged as AM-6); the v1.0 verdict strings "CAP SYNTAX" and "GRADUATE FROM DISCOVERY" used outside a stated vocabulary, now quoted verbatim from the KDR with the C10 cross-map; and the v1.0 worked example's claim that two sub-bar terms "both meet validation", now corrected to routing with pooling and the decision left to #30. The v1.0 references to a "#31" allocation share are replaced by the instrument and tab per conflict C3.

Companions. Master PPC Decision Criteria System v4.0 (all thresholds and verdict issuance) · Keyword Decision Record Template v3.0, sheet "Bands & Verdicts Reference" (the closed vocabulary quoted in Element 5) · PPC Strategic Doctrine v1.0 Sections 1, 3.3, 3.4, 5.3, 6.3, 11, 12 and 13 (the why) · SOP-30 Harvest, SOP-29 Negation, SOP-12 Exact, SOP-13 Phrase, SOP-15 Auto, SOP-28 Placement, SOP-16 PAT, SOP-33 Sibling, SOP-32 Event, SOP-35 Incrementality, SOP-38 Bulk, SOP-37 Automation (boundaries) · Master PPC Data Architecture v1.0 and the PPC Command Workbook READ ME (instrument mechanics) · PPC Authority Matrix v2.0 (Elements 9 and 11) · PPC Cross-Reference Ledger (Element 12 and the ownership check).

11. Operator Ramp (Element 11)

Built from the Authority Matrix v2.0 rows that name this component's decisions: **"Weekly verdicts (established US)"** (PPC Lead decides; incoming operator executes assigned scope; PPC Manager executes; brand-management owner reviews), **"Structure creation (all SOPs)"** (PPC Lead decides or per SOP scope; incoming operator and PPC Manager execute), **"Bulk deployments (#38)"** (PPC Lead reviews; operators execute own scope), **"Automation rules (#37)"** (PPC Lead decides; incoming operator executes and reviews), **"Sibling declarations (S-series)"** (PPC Lead decides; incoming operator executes; brand-management owner reviews), and **"Code-red declaration + close"** (PPC Lead decides; incoming operator and PPC Manager execute), which governs the P11 sentinel floor. Each row becomes a demonstrated competence below. The checkpoint holder for review and audit is the PPC Lead, per the decide and review cells on those rows. Matrix rule 5 applies throughout: an unfilled role's decision right reverts to the PPC Lead and its execution to the PPC Manager until the role fills.

Window	What the person taking this over demonstrates
Week 1 **observation and artifact familiarity**	Reproduces cases E1 and E2 of the Economics Worked-Example Set cold on fresh illustrative inputs, matching to the cent, which is the onboarding arithmetic gate that document sets for every new operator. Shadows one full Monday pass end to end. Locates and reads, unaided: the T3 "Harvest Yield" column with its baseline, the T3 "Negation State" and capped-by-design note, the T2 K8 registry counts, one campaign's child rows with their "Spend Share of KW %", the T7 manifest row whose absence suppresses P3, and one prior T6 cadence entry with its dollar figure. Names the three broad surfaces and the objective tag each carries. Recites the four pre-load classes of P1 step 8 from memory.
Week 2 **supervised execution, checkpoint holder reviewing**	Runs the pass live with the PPC Lead reviewing every route before staging. Pulls the search-term report to the P3 step 2 specification with the named working columns. Produces all three routes with complete evidence lines and correct wall levels, applies both P3 step 3 guards, and computes the head-term check from the child rows rather than from the campaign total. Computes the yield on a corrected base and the cadence dollar figure, and stages one negation set in #38 format with reason strings carrying verdict and scenario ID. Walks one graduation end to end aloud: candidate, route to #30, standup at #12, wall at #29, confirmation here, yield-base correction. The pass bar for the week is zero rows forwarded to #29 without a complete evidence line.
Week 4 **unsupervised execution, checkpoint holder auditing**	Owns the broad scope across the assigned products: runs the pass, stages and deploys per #38, closes T+1, and executes one live structure event end to end, a P1 standup or a P2 fence, with the PPC Lead auditing the artifacts afterwards rather than gating them beforehand. Runs one dated-test close with the verdict argued from the rule. The audit samples the route counts, one negation evidence line, the yield base correction, and the cadence dollar figure. Two clean audit weeks with zero F1, F6 or F7 findings attributable to own sessions close the ramp.

Live-defense questions. These test the reasoning behind the thresholds, not recall of their values. An answer that quotes a number without naming what the number is protecting fails.

1. A broad campaign runs an ACOS well above the discovery band and routed three validated terms this window. Defend continuing it, using the objective-metric contract, and say what would change the answer.
2. Why does Route A route every converting term rather than only the terms that clear the graduation bar, and what specifically breaks if this SOP applies the bar itself?
3. A WAS breach lands on a campaign whose in-budget percentage is low. What do you check before staging a single negative, which rule orders that check, and why does the ordering exist?
4. Thirty-nine zero-order terms averaging one and a half clicks each. Argue the fence against the list, and name the read that decides it and the read that cannot decide it because it has no governed threshold.
5. A head term turns up carrying most of its own clicks in your broad campaign. What leaves Route A, what leaves the yield base, who executes, what is this SOP's remaining half, and what is the cost of the interval between the two?
6. The family hits its cap. Why does broad reduce before exact, what does this SOP execute, what does it never originate, and where does the answer live given that component #31 has no procedural document?
7. A deal window arms tomorrow and you have eleven Route B terms and one new dated test staged. What deploys, what freezes, and why is the answer different for those two?
8. Your negation cadence read zero this week. Name the two possible causes, say how you tell them apart from the artifacts alone, and explain why recording them as one failure loses something the account needs.

12. Interface Contracts (Element 12)

Built from the Cross-Reference Ledger's component-14 Interfaces cell, cross-checked against Master PPC Data Architecture v1.0 for artifact names, tab locations and timing, and against the paired interface rows already returned by the SOP-13, SOP-15 and SOP-29 sessions. Every row states what happens when the handoff does not arrive, because that is the clause the corpus was missing.

#	Interface	Sender to receiver	Artifact and timing	When it does not arrive
I-1	Search-term join	Data Operations to this SOP	Raw Pull Pack: Sponsored Products and Sponsored Brands search term reports plus STIS, joined to T2 campaign child rows. Monday 06:00 landing, workbook refresh by 09:00	P3, P5 and P7 suppress entirely for the affected campaigns; broad is unreadable without served terms and is never approximated from the campaign total. The pass records the suppression under gate G8; Routes A, B and C do not run and no yield figure is written. A second consecutive miss escalates to Data Operations with the manifest history
I-2	Pre-load sets at standup	#29 to this SOP's P1	The #29 pre-load table applied in the creation session and recorded in the T3 "Negation State" column as the wall baseline	The standup does not deploy. A discovery campaign created without its fence is the leak-by-design condition on day one; the session re-opens and the creating session is named
I-3	Route B terms and Route C clusters	This SOP P3 to #29 P4	One staged row per Route B term and one per Route C cluster root, each with a complete evidence line. Staged Monday, deployed Tuesday	#29 P4 suppresses for the affected campaigns rather than completing evidence from the feed; rows return here. Two consecutive passes at zero negatives with spend active is F1 here and F12 at #29

#	Interface	Sender to receiver	Artifact and timing	When it does not arrive
I-4	Route A candidate rows	This SOP P3 to #30 P1	One row per term with every serving source attached and a full evidence line of clicks, orders, window dates, source campaign, CPC and spend. Inside the Monday cascade	#30's harvest pass runs short and converting demand stays at the lowest control level another week. A validated term left unrouted for two passes is a pass-quality finding on this SOP, not on #30
I-5	Head-term receipts	This SOP P3 step 7 to #12	The ADD EXACT + RESTRUCTURE receipt with term, source campaign, click share, exact count and the computed exposure window. Same day, outside the weekly rhythm	The head term keeps being bought by a campaign that chooses its own traffic. The open item stays on the pass worklist flagged daily until the standup lands, and the exposure figure accrues in T6 so the cost of the delay is visible rather than assumed
I-6	Steering walls at graduation	#29 P3 to this SOP's campaigns	The negative exact for the graduated term placed in this campaign's list, same day as the standup, with the confirmation line written back to #30's chain record	The term is GRADUATED-UNSTEERED per v4.0 family F8 and scenario S-H4 fires at the next pass. This campaign re-buys a term the exact already owns; F2 opens here and the wall audit adds this campaign to its population
I-7	Yield and cadence reads	This SOP P5 and P7 to T3, T6 and the audit	The yield figure beside its baseline in the T3 "Harvest Yield" column; cadence counts and the dollar figure in T6, every pass	The discovery surfaces cannot be graded and #30 P7's per-source credit has no denominator to reconcile against. A missing yield figure voids the campaign's grade for the window rather than defaulting it to zero
I-8	Term-ownership declaration	#33 to P1 step 7 and P3 step 3	The declaration record, a re-arbitration changing the owner, or a stewardship grant, landing in T2 K8 "Term Owner (Family)", same day the declaration date carries	Nothing seeds and nothing scales on the contested root: gate G11 holds it. The root stays unseeded and the open declaration is flagged to #33 rather than resolved locally
I-9	Event-plan arming	#32 to P11	The armed-window declaration with its dates, landing in the T2 K2 "Event Mode" column before the ramp date	No dated test opens and no fence change stages inside the window; standing broad holds. A window observed live without a plan escalates to the deals owner per #32's own rules
I-10	Economics chain	The economics refresh to T0 named cells to this SOP	AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS and ExpectedCPC with their freshness state, before the Monday 09:00 packet-ready read	Ceiling-referencing actions hold with the provisional flag; the pass proceeds on routing and negation work and states what it could not price. Route sorting is unaffected, because no route decision references a ceiling
I-11	Variation map	Economics & Context Pack (D2) to P12 and P1 step 9	The keyword-to-variation map with ranked backups, current per its pack cadence	P12 cannot execute: RE-POINT VARIATION has no destination. Campaigns hold on the preferred SKU with the D2-pending flag, and a gate G6 RED on an unmapped row escalates the same day rather than waiting for the next pass
I-12	Budget delta and CAP SYNTAX receipt	Instrument layer (T1 group S3, T0 zone H, the M5 queue) to P9	The T1 "Cap Status" reading BREACH with the staged delta, in the weekly session	The family stays over cap another week and the breach is recorded on the T1 row with its date, so the duration of the over-cap period is visible. This SOP does not compute its own cap and does not reduce a budget without the receipt
I-13	Staged change sets	This SOP to #38	Bulk rows with verdict string and scenario ID in every reason string. Tuesday default; same day on dated releases and on gate G6 RED re-points	Nothing deploys: #38 is the only door. A console edit found at T+1 is the SOP-12 F10 class, the week's reads on the affected rows are void, and the actor is named
I-14	Exclusion-state notices	This SOP to the #37 register	Dated-test opens and closes, modified-broad parallel states, event freezes and recovery sentinel floors, same day as the change	Platform rules may fire on protected discovery scope: a capped test bid stepped, or a paused unfenced keyword re-enabled. The damage is reviewed at T+1 attribution and the missing notice is root-caused to the checklist step
I-15	Master Keyword List	Brand Management to the gate G12 check at P3 step 3	The current list with a syntax tag and a relevancy score on every term, refreshed inside its cadence	Off-taxonomy converters cannot be classified, so scenario S-H6 routing holds and those terms neither graduate nor negate; they stay in their source with the classification gap recorded. A list stale past its cadence degrades the packet manifest
I-16	Decision log and scoring	This SOP to T6, scored by V5	One T6 row per deployed change, each carrying a prediction stated as a figure with a review date	Unscored actions do not accumulate: the next cycle cannot score what was not predicted, and the campaign's history becomes a list of changes with no evidence that any of them worked

13. Amendment Items Logged by This Build (the v4.1 queue)

Six items. Each is a number or rule this SOP needed, could not trace to v4.0, and therefore did not write. None appears as a threshold anywhere in this document. Two overlap with items already logged independently by the SOP-29 session and are cross-referenced rather than duplicated, so the v4.1 owner does not resolve the same gap twice.

ID	The gap	What v4.1 owes, and why it matters here
AM-1	The search-term join has no dedicated T7 MANIFEST row. Data Architecture v1.0 Section 12 names the feed and its consumers; the manifest lists catalog groups A1 to E and the join sits across B5 and B9 without a row of its own.	A manifest row, because a feed whose absence suppresses three entire procedures must be declarable under gate G8. Today the suppression is real but the instrument cannot state it, so a suppressed pass and a clean pass look identical on the manifest.

ID	The gap	What v4.1 owes, and why it matters here
AM-2	The broad drift flag has no threshold. SOP-14 v1.0 Section 2 cites "the v4.0 flag" for a fall in the share of served terms containing the family root. v4.0 Section 3 family F9 defines a placement drift flag at a stated point movement, which is a different metric on a different object.	A firing threshold for root-containment fall, with its window and its point basis, or an explicit ruling that the WAS breach is the only governed trigger and the containment read is evidence only. P2 and WE-3 currently fire on the WAS breach and use containment as selecting evidence, which is defensible but is a local design standing in for a governed rule.
AM-3	The reactive-negation firing threshold, and the "price class" dimension that scopes it. Scenario S-H7 reads "zero-sale spend above threshold" without stating one, and rule L5 names spend-without-sales without a figure; v1.0 scoped the rule to "the term's price class", which is not a v4.0 dimension.	The clicks-at-zero-orders count and the spend figure that fire a Route B negative, plus either a price-class definition or an explicit redirect to margin-dollar class. Duplicates SOP-29's independently logged AM-1 and AM-2; recorded here to show the same gap binds Route B production as well as Route B execution, and to be resolved once.
AM-4	No per-step limit on a cap-driven budget reduction. Rule L4 sizes budget increases; rule L2's plus-or-minus 50% limit is a bid limit; scenario S-Y7 says reduce to cap and scenario S-A6 says reduce until the envelope holds, neither with a step bound.	Either a stated step bound on cap-driven reductions or an explicit statement that reductions to a cap are unbounded by design. WE-4 executes a 54.8% single-step cut because the cap arithmetic requires it, and an operator has no rule to check that against.
AM-5	Verdict vocabulary reconciliation. v4.0 Section 13.1 lists 18 actions; the KDR Template v3.0 lists 25; Data Architecture v1.0 Section 2 rule 8 states 26. This is Ledger conflict C10.	One ratified string set. Element 5 above uses the KDR's 25 as the Build Kit authorizes and as exemplar SOP-12 v2.0 does, and cross-maps the two divergent rows; the build does not pick the canon. Element 5 re-issues on ratification.
AM-6	Dangling naming citation. SOP-14 v1.0 P1 step 2 cites "the 4.2 convention"; the uploaded Criteria System has no Section 4.2, and its Section 17.2 keeps the naming convention authoritative in a document outside this pack.	A canonical section map for the naming convention and the bidding-strategy reference. P1 step 4 cites SOP-12 v2.0 Appendix A, the uploaded restatement, which is a correct pointer to a restatement rather than to the governing source. Overlaps SOP-13's AM-2; the same ruling closes both.

14. Ledger Row Returned by This Session

Paste into the Cross-Reference Ledger Section 1 and raise the version per its Section 6. Tab-separated, columns in the order: number, component, owning document, owns exclusively, must not restate, interfaces, source code, Ledger version, last updated, updated by build session.

```
14	Broad operations, including modified broad, negation cadence, graduation yield tracking	PPC_SOP_14_Broad_Campaign_Operations_v2_0.pdf (supersedes v1.0 in full; built 2026-08-07, Wave 4)	The three broad surfaces (family discovery, modified broad, defensive brand broad) and the objective tag each carries; family discovery standup (P1); the modified-broad fence decision and deployment (P2); THE WEEKLY THREE-ROUTE SEARCH-TERM TRIAGE (P3), which is the account-wide shared definition SOP-13, SOP-15 and SOP-29 run or execute against, including the two pre-sort guards, the route decision tree, the head-term check and the off-taxonomy route; defensive brand broad operations and incursion routing (P4); graduation yield computation, the corrected-base rule and the yield ledger (P5); the source-side wall confirmation half of the graduation chain (P6); THE NEGATION CADENCE AUDIT (P7), distinct from wall correctness which is #29's; dated-test governance on broad (P8); CAP SYNTAX execution ordering on discovery, broad first (P9); the capped-by-design G3 exemption and its T3 note (P10); event freeze and the recovery sentinel floor on discovery (P11); variation re-point on broad (P12); restructure-or-retire, consolidation, pause, archive and re-entry (P13). 13 procedures, 12 failure modes, 4 worked examples, 16 interface contracts each with a non-arrival consequence, the full 25-verdict mapping, Elements 11 and 12 added at v2.0. Zero local thresholds by design.	Thresholds, bands, gates, scenarios, levers, sizing and verdict issuance: v4.0 (#8) exclusively; a threshold appearing here instead of v4.0 is a defect. Graduation decisions, candidate pooling across sources, validation and the LEAVE TO HARVEST re-eval dates: #30 (this SOP delivers Route A candidates with evidence lines and never issues GRADUATE FROM DISCOVERY). Negation doctrine, wall levels, the reverse test, pre-load standards, reactive execution standards and the monthly wall audit: #29 (this SOP produces Routes B and C and owns only the cadence read). Exact standups and ADD EXACT + RESTRUCTURE execution: #12. Phrase surfaces, the stem read and the phrase-broad dedup decision: #13. The four-group auto structure and the ASIN routes: #15. Ladder derivation, exit-bid discovery, modifier sizing and the Surface 3 defense floor: #28. Target opens, rotations and exits: #16. Family term ownership: #33. Event plan, ramp, caps and reconciliation: #32. Incrementality test design and the continue-or-restore call: #35. Staging, upload and T+1: #38. The automation rule register: #37. Budget law, envelope, cap decisions and M5 allocation: the instrument layer (T1 group S3, T0 zone H, the M5 queue per v4.0 Section 10) per conflict C3; this SOP never cites a #31 procedure number. Interface contracts I-1 to I-16, each with a stated non-arrival consequence. Upstream: I-1 search-term join from Data Operations, Monday 06:00 (absence suppresses P3, P5 and P7 entirely per G8, never approximated); I-2 pre-load sets from #29 at standup (absence bars the standup); I-6 steering walls from #29 at graduation (absence leaves the term GRADUATED-UNSTEERED per v4.0 F8, S-H4); I-8 term-ownership declaration from #33 (absence holds the root unseeded under G11); I-9 event-plan arming from #32; I-10 economics chain to T0 named cells (provisional flag rides ceiling-referencing actions); I-11 D2 variation map (absence bars P12); I-12 CAP SYNTAX receipt from the instrument layer; I-15 MKL with syntax tags and relevancy from Brand Management (absence holds S-H6 routing). Downstream: I-3 Route B terms and Route C clusters with evidence lines to #29 P4; I-4 Route A candidate rows to #30 P1; I-5 head-term receipts to #12 same day with the computed exposure window; I-7 yield and cadence reads to T3 and T6; I-13 staged change sets to #38; I-14 exclusion-state notices to #37; I-16 T6 rows with dated predictions for V5. Coupled to #13, #15 and #29: P3 is the shared definition and a route-definition change increments all four documents. Open: AM-1 to AM-6 logged this build (AM-3 and AM-6 overlap SOP-29 AM-1/AM-2 and SOP-13 AM-2 and resolve once). Conflicts reported: C13 (new) the Bamboo window yield anchor, 1.27 on a rounded $158 in SOP-30 v1.0 and SOP-13 v2.0 against 1.26 on the $158.24 stated in SOP-29 v2.0 WE-5; C9, C10 and C12 confirmed as binding on this component, C10 resolved per the Ledger by quoting the KDR's 25 with cross-maps on the two divergent rows.	SOP-14 v2.0 Sections 1.2, 1.3, 4 (P3, P5, P7), 12, 13	1.9	2026-08-07	s9 W4 #14
```