

SOP-15: Auto Campaign Operations and the D5 Four-Way Split

Version 2.0 | 2026-08-05 | Owner: PPC Lead | Supersedes SOP-15 v1.0 (2026-07-24) in full | Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds | Rebuilt to the SOP Build Kit v1.0 twelve-element standard

1. Verdict and Scope

VERDICT: an auto campaign is four different products sharing one wrapper, and a blended auto produces one ACOS that hides four different decisions. Autos carry the 5-10% match-type spend destination of the M6 structure targets (v4.0, Section 10 "Layer 5: Sizing Mathematics (M1-M10)", rule M6) and 100% of the account's auto-side discovery flow: every close-match harvest candidate, every substitutes conquest signal, and the only standing indexing-health read the account owns. Version 2.0 expands v1.0 from 6 procedures to 12, from a 9-row verdict mapping to the full 25-verdict closed vocabulary, from 1 worked example to 4, and from 6 failure modes to 11, and adds the operator ramp and interface contracts the v1.0 standard never required. This SOP governs the D5 conversion of blended autos into the four-group structure, phase-staged deployment of the groups on new products, the weekly per-group pass with its keyword and ASIN routes, per-group bid, budget, and verdict execution, group-evidence hygiene for downstream consumers, the trawler and probe two-tier pattern, per-variation autos, the indexing-health read, event autos, and the automation-interface upkeep that keeps platform rules off protected auto scope.

1.1 The Four Targeting Groups and What Each Buys

The D5 split is the structural form of the four-group read defined in v4.0, Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)", Family F11 "Waste & Negation": each group is judged separately on cost per discovered converter, waste %, and graduation rate, which is also the Discovery objective-metric contract of rule O6 (Section 6 "Layer 1: Objective Assignment (O1-O8)"). Per the Strategic Doctrine v1.0 Section 4 "Auto Campaigns: Four Products in One Wrapper": **close match** buys queries closely matching the listing's indexed content, is the cleanest harvest source, and doubles as the indexing-health signal (P5); **substitutes** buys competitor detail pages and is conquest-lite without ASIN research, feeding the S-C1 target lists (v4.0, Section 8.7 "Conquesting Scenarios (S-C1 to S-C6)") through the P3.3 handoff; **loose match** buys semantically adjacent queries, carries the widest drift exposure, and runs the tightest waste control; **complements** buys adjacent-purchase pages, reads basket demand and inbound advertisers, and runs lowest, often event-only. Bid posture orders close highest, substitutes second, loose third, complements lowest per the same Doctrine section. A single blended auto collapses these four intents into one line; v4.0 Section 8.9 "Discovery Scenarios (S-H1 to S-H8)" scenario S-H3 orders the split read before any judgment, and the D5 structure makes S-H3 permanent.

1.2 Legitimate Multi-Auto Structures (and the One Illegitimate One)

Exactly four multi-auto structures are legal per Doctrine v1.0 Section 4, and each carries its own procedure here: (1) the **four-way split** itself, one campaign per targeting group per product, each with its own bid, budget, negation list, and verdict (P1, P2); (2) **trawler and probe** two-tier price discovery, a low-bid wide harvester beside a dated higher-bid validator on the same group, with shared negation lists mandatory at standup (P4; the shared-list rule executes per SOP-29 P4.1); (3) **per-variation autos** where variation demand diverges, each pointed at its D2-mapped variation (P10); (4) **dated event autos**, budget-capped, opened only inside a G12-armed window and closed with it (P6). The illegitimate fifth is a second blended auto created because the first "stopped working": that is a diagnosis skip, autos degrade from negation neglect and bid drift, and cloning resets the data without fixing either. A live auto matching none of the four legal patterns is failure mode F1 the day it is found.

1.3 Boundaries With Adjacent Components

Adjacent component	The seam, stated so it is owned
Criteria System v4.0	Decides every threshold, band, scenario, stage condition, and verdict. This SOP executes; it never decides. A threshold appearing here instead of v4.0 is a defect, and the numbers v4.0 does not yet carry are Appendix A amendment items, never local values.

#30 Harvest and graduation	Decides which terms graduate. This SOP delivers candidate rows with evidence lines and group tokens (P3 Route A) and never issues GRADUATE FROM DISCOVERY itself.
#29 Negation architecture	Owns negation doctrine, the pre-load tables, wall standards, shared-list management, and the wall audit. This SOP detects and routes auto-side negation needs (P3 Routes B, C, D) and applies #29's prescribed sets in the same session at every structure event.
#16 PAT and conquest	Owns target opens, flag tests, rotation, and exits. This SOP routes converting substitutes ASINs to #16 with evidence (P3.3) and ingests #16's failed-target list as negative-ASIN input; this SOP routes, #16 decides.
#12 Exact operations	Owns exact standups. ADD EXACT + RESTRUCTURE receipts from the P3 head-term check execute at SOP-12 P1 or P2; the source group contributes its steering wall after standup.
SOP-14 Broad operations	Owns the definition of the weekly three-route triage machinery (SOP-14 P3). This SOP executes that machinery on auto campaigns and adds the two ASIN routes autos uniquely need.
#37 Automation rules	Owns the rule register, reconciliation, and change control. Every ScaleInsights rule touching autos routes there; this SOP maintains the auto-side campaign states (parallel windows, capped-by-design rows, probes, event autos) that #37's exclusion lists consume (P11).
#38 Bulk mechanics	Owns pull, staging format, upload, and T+1 verification. This SOP produces staged sets in #38 format and consumes its verification report.
#32 Deal and event windows	Owns event plans, ramps, caps, and reconciliation. P6 event autos open only under an armed plan and close per #32's playbook.
#22 Launch-stage PPC	Owns the deployment sequence execution on new products. P2 staging honors the stage conditions v4.0 sets and the launch handoff #22 runs; stage-skipping that originates in launch routes back to #22 (F6).
#31 Budget and envelope	Envelope law lives in the instruments (T0, T1, Account Rollup A2) per open conflict C3; this SOP receives CAP SYNTAX deltas from T1 and never cites a #31 procedure number.
#34 Canada annex	Owns the Canada delta. Everything here applies in CA unless #34 overrides it.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. Lags use the same cadence terms v4.0 uses for gate responses (same day, same session) and the Monday cascade defined in Master PPC Data Architecture v1.0 Section 12 "Data Flow: Feed → Columns → Cadence".

Trigger	Detecting artifact	Max lag	Procedure
D5 conversion due	T3 Campaign Board: a live SP Auto row whose "Auto Group Split (D5)" cell reads blended; the D5 program row open in the Definition Register run ledger	Same week the row is found	P1 split conversion
Phase transition	T0 zone E stage read plus the v4.0 stage conditions clearing (Section 11 "Layer 6: Transitions, Stages, Modes (T1-T8) and Recovery Mode (R1-R6)", rule T2)	Next Monday pass	P2 staged standup of the newly eligible group
Weekly cadence	Monday cascade reaches the auto rows of T2 and T3; search-term data refreshed in the Raw Pull Pack (Monday 06:00 landing)	Same day (Monday)	P3 weekly pass, then P7, P8, P12
Two-tier test opened	A verdict in T2 K9 naming a capped test on a starved group (the KDR string CONTINUE TEST TO 50 CLICKS AT CAPPED BID or FUND TEST (UNFUNDED PRIMARY))	Same day as the funding release	P4 trawler and probe standup
Listing rewrite or monthly slot	T5 Product Log listing-change entry; otherwise the V4 monthly cadence slot (v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)")	Within the next weekly pass after the log entry	P5 indexing-health read

Event window armed	G12 armed by the #32 event plan; the window dates land in T2 K2 Event Mode	Per the event plan's ramp date	P6 event auto (dated, capped)
Group evidence requested downstream	#30's Monday harvest pass opens; SOP-30 Section 3 requires group-level auto evidence post-D5	Inside the same Monday pass	P9 group-token hygiene before routing
Structure state change	Any P1/P2/P4/P6 standup, close, or pause; the #38 staged set for it	Same day as the state change	P11 automation-interface upkeep
Verification return	#38 T+1 verification report flags a landed-value mismatch or an unattributed change on an auto row	Same day	F10 response via P11; redeploy per #38

3. Preconditions and Gates

Every procedure below runs only after these read current in the T7 Manifest (Master PPC Data Architecture v1.0 Section 9 "The Infrastructure Tabs: Logs and the Manifest (T5-T7)"). A missing or stale field group suppresses its dependent execution class; the pass records what it could not see and continues on the rows it can, per gate G8.

Precondition (T7 manifest row)	Requirement and the observable that proves it passed
D5 Auto four-group split	The targeting-group split is pullable for the window: the Sponsored Products Search Term Report carries the Targeting values close-match, loose-match, substitutes, complements, and T3's "Auto Group Split (D5)" cell resolves per campaign. Observable: T7 row D5 Status = PRESENT. When MISSING or STALE, group-level verdicts are suppressed for that window rather than approximated from the blend; the suppression is written into the pass record.
A2 Deployment state	Current bulk state in-row for every auto campaign (state, budget, default bid). Observable: T3 C2 cells populated with the current pull date.
A3 Pacing and in-budget	In-budget % and bucket per campaign. Observable: T3 C3 "In-Budget Bucket" resolved; G3 evaluation possible.
A6 Objective + Expected CPC	Every auto campaign carries Objective = Discovery (or the event tag under P6); UNTAGGED routes to assignment per O2 before execution (v4.0, Section 6 "Layer 1: Objective Assignment (O1-O8)", rule O1). Observable: T3 C1 Objective Tag populated, zero UNTAGGED auto rows.
A7 Eligibility	Advertised SKU eligible; a flip halts execution for the target same day per G4. Observable: T3 C2 Eligibility = Live.
B5 Coverage counts	Registry counts current (T2 K8 "# Auto" per term and the per-group campaign registry), or split-legality and head-term reads are unreadable. Observable: T2 K8 populated this cycle.
B8 Economics chain	T0 named cells live (AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC). The G9 PROVISIONAL tag carries onto any ceiling-referencing bid while margin is blended or stale. Observable: T0 zone C shows the margin-table version and a fresh G9 state.
D2 Variation map	Advertised SKU per auto campaign mapped, with ranked backups for P10 per-variation autos. Observable: T3 C1 "Targeted Variation + Band" resolved; unmapped rows carry the D2-pending flag and bar P10 standups.
ASIN-row identification	Search-term rows split into keyword queries versus ASIN matches (rows whose Customer Search Term is an ASIN string) before P3 routing. Observable: the pass worksheet carries the two-way split with row counts.
Own-catalog ASIN list + #16 target-status list	Both current for the P3 Route D negative-ASIN checks. Observable: list version dates inside the trailing week.

Gates in force. G1 quarantines impossible rows (flags only, no verdicts). G2 sample sufficiency governs every rate read: bid reads need at least 15 clicks (7d), CVR verdicts at least 100 clicks, CTR verdicts at least 1,000 impressions, with the 60-day fallback and EXTEND-class handling for thin rows, and WoW swings above 50% on thin samples read as noise (v4.0, Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass", gate G2); twin sets pool evidence. G3 marks rates TRUNCATED below 70% in-budget and forces the budget lever first on close and substitutes; loose and

complements rows carrying the T3 capped-by-design note are graded on the O6 discovery metrics at their cap instead (the Discovery utilization target is 70%+, not 100%, per v4.0 Section 3 Family F10 "Budget & Pacing" and the KDR v3.0 Bands & Verdicts Reference "Utilization targets" line; the exemption class itself is Appendix A item AM-4). G6 inventory bands govern advertised SKUs. G9 marks ceiling reads PROVISIONAL after a price change. G10 excludes event windows from trend verdicts. G12 freezes non-plan structure changes inside armed windows and requires every analyzed term to exist in the current MKL.

4. Step-Level Procedures

P1. The D5 conversion: blended auto to four-group structure *(60-90 min per product plus the parallel-window watch; planning estimate, not a threshold)*

1. **Inventory.** In T3, filter Type = SP Auto and read "Auto Group Split (D5)": every blended row lists with trailing 60-day spend, sales, and both negative lists (keyword and ASIN), pulled from the Sponsored Products bulk sheets in the Raw Pull Pack. The conversion order inside a product is by trailing spend, largest first.
2. **Create four campaigns,** one targeting group enabled in each, named with the group token in the KW/Root/Auto position per the naming convention restated in SOP-12 v2.0 Appendix A: Auto-Close, Auto-Sub, Auto-Loose, Auto-Comp. Objective = Discovery; bidding is down-only with capped test bids per the Doctrine v1.0 Section 2 control-gradient row for Discovery. In the bulk file, the four rows differ only in which Product Targeting Expression is enabled (close-match, substitutes, loose-match, complements) with the other three paused.
3. **Seed bids by group ordering** (close highest, substitutes second, loose third, complements lowest, per Doctrine Section 4): each seed derives from the blended auto's own group-level CPC history where the split report provides it, otherwise from the Expected CPC reference the driving verdict carries (T0 named cell). No bid is invented at creation; a standup with no evidence-carried seed returns to the pass as incomplete. Per-group reference seeds are Appendix A item AM-3 until v4.1 admits them.
4. **Split the budget** across the four by trailing group spend share from the same report. Loose and complements receive the capped-by-design note in their T3 rows at creation (Budget Verdict cell records the cap as the decision), which is what exempts them from FIX BUDGET (TRUNCATED) later (Section 3, AM-4).
5. **Migrate negatives, then pre-load fresh.** Keyword negatives move to Auto-Close and Auto-Loose; ASIN negatives move to Auto-Sub and Auto-Comp. Then apply the SOP-29 P1 pre-load table for the auto types in the same session: the broad keyword set to close and loose; failed-PAT-target ASINs to substitutes; own-catalog ASINs to complements (the self-cannibalization wall) and NOT to close or loose, which serve queries, not pages. A creation deployed without its pre-load is F1 in SOP-29 and F3 exposure here.
6. **Run the parallel window.** The four split campaigns run beside the legacy blended auto at capped budgets so coverage never gaps; at the window's end the legacy auto pauses as PAUSED-PENDING (REASON + RE-ENTRY) with reason = D5 split and re-entry = none. The window's duration and the pause day are Appendix A item AM-2; until v4.1 lands, the operating value is the one recorded in the product's own T6 conversion entry, and the daily watch below is what makes any value safe.
7. **Watch the window daily:** the two structures' combined spend against the pre-split baseline, and the legacy auto's state. A legacy auto spending past its recorded pause day is F5 and pauses immediately.
8. **Registry and log.** T2 K8 "# Auto" updates per term as the split lands; four T3 rows created with group tokens; one T6 entry carries the decomposition prediction: the blend's single ACOS will decompose into four group reads, and each group receives its own verdict from the closed vocabulary at the first post-split cycle. P11 runs the same day so #37's exclusion lists see the parallel-window rows.

P2. Phase-staged deployment on new products *(20-30 min per standup; planning estimate)*

1. The four groups never deploy simultaneously on a new product. Staging follows the v4.0 deployment sequence and stage conditions (Section 11, rule T2, which retains the Launch/Scale graduation criteria and the Phase 1-4 model): **close match** stands up at Phase 2 entry with broad, per the sequence; **substitutes** stands up at Phase 3 entry, the scale stage, under the organic-share and review conditions T2 carries; **loose and complements** stand up late and rarely: capped, often event-only, and only where close and substitutes sit at band.
2. Each standup runs P1 steps 2-5 for its single group. The T6 entry names the stage condition that authorized it; a standup entry without the cited condition is incomplete and does not deploy.

3. Skipping a stage requires a verdict, not a preference; deploying every group at once is the unpriced-risk pattern Doctrine v1.0 Section 2 names, and a group live without its authorizing condition is F6.

P3. Weekly pass: the SOP-14 P3 machinery plus the auto-specific routes *(25-40 min per product; planning estimate)*

1. **Enter after T0 and T1.** Product posture and family caps govern everything below them; score last cycle's due T6 rows first (V1: unscored actions do not accumulate).
2. **Pull and split.** From the Raw Pull Pack, open the Sponsored Products Search Term Report filtered to the product's auto campaigns, trailing 60 days with the 7-day and 30-day columns visible. Split every row into keyword queries versus ASIN matches (Customer Search Term values that are ASIN strings). Confirm each row carries its group token from the Targeting column (close-match, loose-match, substitutes, complements); rows without a resolvable token quarantine per P9.
3. **Run the keyword routes per SOP-14 P3 on every group's keyword rows.** Route A converters: any term at or above the graduation evidence bar becomes a candidate row to #30 with the full evidence line (clicks, orders, window, source campaign with group token, CPC); the graduation decision is #30's, made against v4.0 Section 8.9 scenario S-H1 (a discovery term reaching 3 or more orders graduates) and the validation rules for its price class. Route B proven waste: terms breaching the WAS limits (40% discovery, 10% elsewhere, per S-H2) stage negative exact with evidence lines, executed to SOP-29 P3 standards, including the converting-term guard. Route C drift clusters: wrong-intent clusters fence with negative phrase on the cluster root, one fence for many futures.
4. **Route D, the ASIN routes (auto-specific). P3.3 substitutes converters:** converting ASINs in Auto-Sub are conquest-lite validations; route them to #16 as PAT target candidates with the evidence line (clicks, orders, our price versus theirs where the competitor depth feed carries it; until that feed lands, the line notes its coverage). This SOP routes; #16 decides, and its S-C1 feed is exactly this list. **P3.4 negative-ASIN:** own-catalog ASINs appearing in complements or loose served terms stage negative-ASIN the same pass (the F3 wall); ASINs on #16's failed-target list stage negative-ASIN in Auto-Sub so the scout never re-buys a proven-bad page. Execution standards are SOP-29 P3's.
5. **Head-term check.** For every priority term served by an auto group, read the T2 campaign child rows' "Spend Share of KW %": auto or broad carrying more than 40% of a priority term's clicks with no live exact fires ADD EXACT + RESTRUCTURE before any performance verdict (v4.0, Section 4 "Layer X: The Interaction Layer", rule X10, and Section 8.4 "Ranking / Re-Ranking Keyword Scenarios (S-R1 to S-R16)", scenario S-R12). The receipt goes to SOP-12 the same day; the source group contributes its steering wall after standup per SOP-29 P2.
6. **Off-taxonomy converters.** A converting term below the relevancy bar or outside the taxonomy routes per S-H6 to the MKL classification decision first (gate G12): the taxonomy expands deliberately with a logged tag and test plan, or the term negates. Incidental orders do not create a market position.
7. **Stage, deploy, verify, log.** All changes stage as #38 bulk rows with verdict and scenario ID in the reason string; Tuesday deploy after the BM huddle; T+1 verification; every deployed change writes a T6 line with its prediction and review date. Route counts (A/B/C/D per group) close the pass and feed P12.

Decision tree for the routing step (any served-term row):

```
Served-term row (group token resolved? no → P9 quarantine)
|
+-- ASIN row?
|   +-- In Auto-Sub and converting ..... Route D1: evidence line → #16 (P3.3)
|   +-- Own-catalog ASIN (Comp/Loose) ..... Route D2: negative-ASIN this pass (F3 wall)
|   +-- On #16 failed-target list ..... Route D2: negative-ASIN in Auto-Sub
|   +-- Otherwise ..... hold; grade inside the group's 06 read
|
+-- Keyword row?
    +-- Orders at/above the S-H1 bar ..... Route A: candidate row + token → #30
    +-- Converting, off-taxonomy ..... S-H6: MKL decision first (G12)
    +-- Zero-sale past the S-H7 line ..... Route B: CUT / NEGATE (relevant primary
        |                               instead routes to the 50-click capped test)
    +-- WAS breach, proven waste ..... Route B: negative exact + evidence line
    +-- Wrong-intent cluster ..... Route C: negative phrase on the root
    +-- Priority term, auto share >40% ..... X10/S-R12: ADD EXACT + RESTRUCTURE → #12
    +-- Otherwise ..... stays; sufficiency per G2 governs the read
```


P4. Trawler and probe two-tier price discovery (30 min standup, 10 min per weekly read; planning estimate)

1. **Standup only through a verdict** opening the test on a named group: the trawler is the low-bid, wide, harvest-focused instance held at the L2 floor (\$0.50, v4.0 Section 9 "Layer 4: The Lever Library (L1-L10)", rule L2) or the group's evidence-carried low seed; the probe is the dated, higher-bid instance running to a click quota at the capped bid the verdict names (the KDR string CONTINUE TEST TO 50 CLICKS AT CAPPED BID is the usual carrier, and the 50-click capped test is the S-H7 relevant-primary route). The two must carry different bids by design.
2. **Shared negation lists are mandatory at standup** so the tiers cannot buy the same query twice; the list attaches per SOP-29 P4.1 (trawler and probe pairs share by rule) and lives in the T3 rows.
3. **At the probe's quota**, compute the read: probe CPC and CVR against the Max Profitable CPC (margin dollars times CVR, v4.0 Section 3 Family F1 "Product Economics") and the Expected CPC reference. A viable discovered CPC re-seeds the standing group bid through L2-compliant staged steps (no single change beyond 50%, \$0.50 floor); a failed probe closes STOP-LOSS: CONCEDE OR DEFER. Either way the probe campaign itself parks as PAUSED-PENDING (REASON + RE-ENTRY) with the test outcome as the reason.
4. **Log:** the probe's T6 chain carries one row per event (standup, quota reached, re-seed or stop-loss), each with its prediction, so V5 scores the test as a test.

P5. Indexing-health read (20 min per product per monthly slot; planning estimate)

1. Close match serves what Amazon believes the listing is. On the V4 monthly cadence slot, or within the next weekly pass after any T5-logged listing rewrite, export Auto-Close's served-term distribution (top terms by impressions, trailing 30 days) and compare it against the family's intended syntax map in T1 (S1 identity columns and the taxonomy tags).
2. Material divergence (the quantitative bar is Appendix A item AM-5; until v4.1, divergence is read against the syntax map qualitatively and stated with the served-term evidence) is an indexing or listing-content finding, not a bid problem: route it to the BM owner via the cross-framework handoff with the served-term evidence attached. Bids never fix indexing, per gate G5's logic.
3. The read files as a T6 product-level entry either way; a clean read is recorded, not skipped, exactly as the SOP-29 P6 audit records zero-breach months.

P6. Event autos (20 min per window; planning estimate)

1. Dated, budget-capped, opened only inside a G12-armed window with the expiry in the campaign name; sized per the event plan #32 owns (event sizing mathematics are v4.0 Section 10, rule M9); closed with the window per #32's playbook. Event-window rates never enter trend verdicts (gate G10).
2. An event auto never converts to a standing campaign silently; conversion requires a verdict, a fresh Objective tag, and a P1-style pre-load check, and a window campaign found live past its expiry is F11.

P7. Per-group bid and verdict execution mechanics (inside the weekly pass, 10-15 min; planning estimate)

1. Every bid or budget change exists first as a staged #38 row with prior value, new value, and the reason string carrying verdict and scenario ID. No console edits outside the staged set; a console edit found by verification is F10.
2. Step sizes, floors, and per-change limits read from L2 at execution time: traffic-unlock steps of 15-25% where clicks run under the G2 bid-read bar, changes capped at 50% per step, the \$0.50 floor, ladder steps of roughly 10% per week, and the placement-exit mechanic of bidding \$0.05-0.10 below a placement's blended CPC where a placement read exists. This SOP never restates the limits; a staged step exceeding L2 fails staging QA and returns to the pass.
3. The lever order is v4.0's L1: budget first where G3 is failing on close or substitutes, then bid, then placement, then negation, then structure, then status, with pause last after documented exhaustion. The auto-specific lever-order variant recorded in v1.0's boundary is not restated here; it rides as Appendix A item AM-1 per the Ledger's component-15 row.
4. One row receives one lever per cycle unless the verdict itself names a two-lever program; opposite trends never share an execution.

P8. Group budget and cap mechanics (inside the weekly pass, 10 min; planning estimate)

1. **FIX BUDGET (TRUNCATED)** applies to Auto-Close and Auto-Sub: below 70% in-budget, the budget lever runs first and the week's rate reads are marked TRUNCATED (gate G3). Auto-Loose and Auto-Comp rows carrying the capped-by-design note are graded at their caps on the O6 discovery metrics instead (Section 3, AM-4).

2. **CAP SYNTAX receipts** arrive from T1's cap enforcement (the instrument layer per conflict C3): auto budgets in the breaching family reduce with broad, loose and complements first, close and substitutes last, and the delta stages per #38 with the T1 reference in the reason string.
3. **S-H5 enforcement**: a group whose cost per discovered converter runs above 2 times the account norm for the group steps its budget down 25% (v4.0 Section 8.9, scenario S-H5); loose and complements are the usual offenders and close and substitutes retain budget priority, exactly as the scenario states. The step stages as one #38 row with S-H5 in the reason string.
4. **Recovery Mode overlay**: when a product is in Recovery, discovery budgets floor at sentinel levels, the autos stay on small (v4.0 Section 11, rule R3); the sentinel state is set by the product verdict and recorded in T3, and P8 never scales a sentinel row.

P9. Group-evidence hygiene: tokens, blends, and the downstream contract *(inside the weekly pass, 5-10 min; planning estimate)*

1. Post-D5, every piece of auto-sourced evidence that leaves this SOP carries its group token: candidate rows to #30 (whose Section 3 gate requires group-level auto evidence, never blended), ASIN handoffs to #16, and negation evidence lines to #29. The token is the campaign name's group segment plus the report's Targeting value; the two must agree.
2. A served-term row whose token cannot be resolved (report gap, renamed campaign, legacy blend still serving) quarantines: it routes to no downstream consumer, the T7 manifest note records the gap, and the row re-enters when the token resolves. Blended evidence validates nothing, and forwarding it anyway is F7.
3. Where a term earned evidence in two groups, the candidate row lists both sources with both tokens (one row per term, all sources attached, matching SOP-30 P1's dedupe rule), so the graduation walls land in every source.

P10. Per-variation autos *(25 min per standup; planning estimate)*

1. Legal only where variation demand genuinely diverges (size-specific or color-specific intent per Doctrine Section 4, case 3) and the D2 map resolves the variation with its ranked backup. Each per-variation auto points its advertised SKU at the D2-mapped variation, carries the variation in its name, and inherits the group token structure (a per-variation split is a refinement of a group, never a replacement for the D5 split).
2. Standup follows P1 steps 2-5 for the group being varied; the T6 entry names the divergence evidence (CVR-by-variation or demand split) that authorized it. G6 inventory bands govern per variation: RE-POINT VARIATION executes on RED per the D2 backup chain, same day, with the SOP-12 P10 mechanics as the execution reference.
3. A per-variation auto with no divergence evidence in its T6 entry is an F1-class structure and consolidates back.

P11. Automation-interface upkeep *(10 min at each structure event and inside the weekly pass; planning estimate)*

1. Every ScaleInsights rule that touches autos is inventoried and change-controlled in #37 (its register, Section 2 classes A-D); this SOP never edits a rule. What this SOP owns is the state feed: at every P1 parallel-window open and close, P4 probe standup and close, P6 event-auto open and close, and any manual-hold or sentinel state, the affected campaign names go to the #37 exclusion-list maintenance the same day (SOP-37 P3 carries exclusions with the state changes; the standup checklist carries the add).
2. The #38 T+1 attribution report is read every cycle for auto rows: a rule-fired change on a parallel-window, capped-by-design, probe, or event row is F10 here and F2 in #37; a change neither a session nor a rule claims escalates as an unlogged edit.
3. The weekly audit outranks inter-cycle automation on any shared target: the audit value wins and automation resumes from it (v4.0, Section 12 "Layer 7: Precedence and Conflict Resolution (P1-P10)", rule P3).

P12. Pass close and per-group yield reporting *(10-15 min; planning estimate)*

1. Per group, write the cycle's discovery reads into T3: cost per discovered converter, WAS %, Harvest Yield (converting terms found), Graduations (Qtr), and the "Auto Group Split (D5)" state. These are the O6 discovery metrics, and they, not group ACOS alone, grade the group.
2. Route counts close the pass: Route A candidates delivered to #30 (with tokens), Route B/C negatives staged to #29 standards, Route D handoffs to #16 and negative-ASIN counts, head-term receipts to #12. A pass artifact missing any route count is incomplete.

3. Discovery that never graduates anything is waste with a process name (Doctrine Section 2, Discovery row): a group at zero graduations across the quarter routes to the S-H5/S-H7 reads and, where the group is loose or complements, to the retirement question under the P8 cap logic; retirement itself is a verdict (PAUSED-PENDING (REASON + RE-ENTRY)), never a drift.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds, and scenario logic live in v4.0; the verdict strings below are quoted verbatim from the closed vocabulary in the Keyword Decision Record Template v3.0, Bands & Verdicts Reference sheet (the vocabulary v4.0 Section 13.1 "The Closed Verdict Vocabulary" governs). This table maps all 25 strings to what this SOP does when each lands on an auto-relevant row; where execution belongs to an adjacent component, the boundary is stated so no verdict falls into a seam.

Verdict (closed vocabulary, verbatim)	Auto-campaign execution
QUARANTINE: DATA REFRESH	No changes deploy on the row (G1); the refresh task carries the staged set to the next clean cycle; P9 holds any downstream routing from the row.
UNLOCK TRAFFIC	L2 unlock steps on the named group (close or substitutes in practice); no rate judgments until G2 sufficiency.
FIX BUDGET (TRUNCATED)	P8.1: budget lever only, on Auto-Close and Auto-Sub; capped-by-design loose and complements rows are exempt and grade on O6 at their caps.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	P4 probe (or an in-group capped test) runs to quota; budget assurance only; no rate verdicts mid-test.
FUND TEST (UNFUNDED PRIMARY)	P2 staged standup or P4 probe per the verdict's naming, stage condition cited in the T6 entry.
RANK PUSH: FUND SIZED PUSH	Never executes on an auto: autos cannot carry rank objectives (control gradient). The receipt routes to SOP-12; the auto side contributes the steering wall after standup.
RECOVERY PUSH	As above: exact-side execution at SOP-12; the auto group's history supplies evidence only.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	Exact-side execution; if the deficit read implicates auto walls (a steering negative choking a still-needed source), the wall question routes to #29 P7 removal control, never to a local delete.
FIX CONVERSION FIRST; DEFER PUSH	No auto scale on the affected term; the group holds its cap; the fix routes to the D8-named owner outside this SOP.
PLACEMENT FIX FIRST	Rare on autos; bids hold and the #28 program executes; this SOP deploys only campaign-level rows staged back.
RE-POINT VARIATION	P10.2 same day on the affected per-variation auto, per the D2 backup chain; baseline reset noted in T6.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	Exact-side; the auto group's only action is confirming its steering wall on the tapered term still stands (no removal on taper).
HOLD; DEFEND LIGHTLY AT RANK	No auto change; the group's read continues on its own O6 metrics.
DEFEND AT FLOOR	For autos, applies to Auto-Comp and trawler rows held at floor: no bid change; drift from floor restores in one staged step.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	Bid edit toward the ceiling on the named group per L2; the syntax-cap half arrives as a T1 delta and executes per P8.2.
INCREMENTALITY LADDER	Not an auto instrument; where a branded term still serves through an auto, the finding is a wall gap (route to #29), not a ladder.
ADD EXACT + RESTRUCTURE	P3.5 receipt to SOP-12 same day; the source group contributes its steering wall after standup; the restructure note names which group was carrying the term.
CONSOLIDATE DUPLICATES	For autos: an illegitimate clone or an evidence-free per-variation auto consolidates back to the legitimate structure (F1 response); survivor logic follows SOP-12 P7's precedence with the group token as the role test.

GRADUATE FROM DISCOVERY	Receipt of #30's decision on a Route A candidate: the steering negative lands in the source group(s) the same day per SOP-29 P2, and the group's Graduations (Qtr) count increments in T3.
NEGATE (PRE-LOAD / REACTIVE / STEERING)	P1.5 pre-loads at every standup; P3 Routes B, C, D reactive sets with evidence lines; steering entries at every graduation receipt. Standards and audit are #29's; negative-ASIN states are the auto-specific execution surface here.
ENTER SBV ARBITRAGE	No auto-side execution; the format build is #17-19's and the SP step-down is SOP-12's.
OPEN / ROTATE / EXIT CONQUEST	#16's verdict on targets this SOP sourced; the only auto-side execution is the EXIT echo: failed targets return on #16's list and stage negative-ASIN in Auto-Sub (P3.4).
STOP-LOSS: CONCEDE OR DEFER	P4 probe close or P6 event-auto close; freed budget returns to the envelope and reallocates at the instrument layer, never inside this SOP.
PAUSED-PENDING (REASON + RE-ENTRY)	P1.6 legacy pause at conversion; P4 probe park; P6 window close; any group retiring under the P12.3 read. Always with reason and re-entry condition; the weekly pass tests the criterion and never re-litigates the pause.
LEAVE TO HARVEST	The term stays on discovery economics in its source group at its capped bid; no exact stands up; the group keeps the read alive and the term re-enters through #30's pool per its re-eval date.

6. Worked Examples (illustrative values, internally consistent; live values come from T0 and T2)

WE-1 | The D5 decomposition: one blend, four verdicts. Illustrative on the standard fictional base (AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target 17.5%, Max Acceptable 26.3%); the Tier 3 live spine was not attached.

Setup. Legacy blended auto, trailing week before conversion: spend \$210.00, 7 orders, sales \$560.00, blended ACOS = $210.00 / 560.00 = 37.5\%$. P1 runs; first full post-split cycle reads, by group: Auto-Close spend \$92.00, 100 clicks, 5 orders, sales \$400.00; Auto-Sub \$58.00, 40 clicks, 1 order, \$80.00; Auto-Loose \$47.00, 50 clicks, 1 order, \$80.00; Auto-Comp \$13.00, 20 clicks, 0 orders, \$0.00. Totals: \$210.00 spend, \$560.00 sales, 37.5% blended: the identical number the old structure reported as one decision, now four rows.

Group reads. Close: $ACOS\ 92.00/400.00 = 23.0\%$, inside the Discovery Max-ACOS band of 15-30% (KDR v3.0 Bands & Verdicts Reference, "Objective Max-ACOS bands"); $CPC = 92.00/100 = \$0.92$; $CVR\ 5/100 = 5.0\%$; Max Profitable $CPC = 28.00 \times 0.050 = \1.40 , so \$0.92 sits at 0.66x ceiling, SAFE; 3 distinct converting terms, cost per discovered converter = $92.00/3 = \$30.67$ against an illustrative account norm of $\$24.00 = 1.28x$, under the 2x S-H5 line; one term reaches 3 orders and routes as a Route A candidate. Verdict on the group row: HOLD; DEFEND LIGHTLY AT RANK does not apply to autos, so the group holds under its O6 read with the candidate routed. Substitutes: $ACOS\ 58.00/80.00 = 72.5\%$, but the group's product is target evidence, not ACOS: 1 converting ASIN (9 clicks, 1 order) routes to #16 per P3.3, and 1 own-catalog ASIN found serving stages negative-ASIN. Loose: $ACOS\ 47.00/80.00 = 58.75\%$; $CPC\ 47.00/50 = \$0.94$ against a ceiling of $28.00 \times (1/50 = 2.0\%) = \0.56 , 1.68x ceiling; Route B takes 2 proven-waste terms to negative exact and Route C fences a wrong-material cluster with one negative phrase; cost per discovered converter $47.00/1 = \$47.00 = 1.96x$ norm: held, one step from the S-H5 trigger, and the T6 prediction says so. Complements: \$13.00 spend, 0 orders: the zero-sale terms route per S-H7; the group itself stays at its designed cap and floor bid, because the account keeps the read alive at minimum cost.

Log. The P1 decomposition prediction scores HIT (four distinct verdicts replaced one blend); four next-cycle predictions log per group. This example carries forward, recomputed on the standard base, the v1.0 Section 6 worked case.

WE-2 | A clean weekly pass with all four routes. Illustrative on the standard fictional base; evidence bars cited are v4.0's.

Route A. "bamboo sheets split king" in Auto-Close: 12 clicks, 3 orders in-window, CPC \$1.05. Three orders meets the S-H1 bar (v4.0 Section 8.9), so the candidate row goes to #30 with the evidence line: 12 clicks, 3 orders, trailing 30d, source Auto-Close (token close-match), CPC \$1.05. #30 decides; this pass only delivers.

Route B. Loose term "silk pillowcase" at 24 clicks, 0 orders, \$21.60 spend: proven waste against the discovery WAS limit (40%, S-H2); negative exact staged with the evidence line, executed to SOP-29 P3 standards; the converting-term guard confirms zero in-window orders before staging.

Route C. Six curtain-intent terms cluster in Auto-Loose: 31 clicks, 0 orders combined; one negative phrase on the root "curtain" fences the cluster; root and member count ride the evidence line.

Route D. Auto-Sub ASIN B0-prefixed row: 9 clicks, 1 order; our \$80.00 versus their \$86.40 (8.0% above us, from the available price source; the depth feed's absence is noted on the line). The row goes to #16 as a PAT candidate. One own-catalog ASIN appears in Auto-Comp served terms: negative-ASIN staged the same pass.

Head-term check. "bamboo sheets king" totals 68 cross-type clicks this window; Auto-Close carries 31 = 45.6%, above the 40% auto-dominance line (X10), and T2 K8 shows # Exact = 0: ADD EXACT + RESTRUCTURE fires per S-R12 before any performance verdict, receipt to SOP-12 same day, steering wall to follow the standup. Pass close: A=1, B=1, C=1 (6 terms fenced), D=2, head-term receipts=1; counts into P12.

WE-3 | The edge case: blended evidence after the split, plus an S-H5 breach. Illustrative on the standard fictional base.

The bad candidate. A candidate row reaches the pass worksheet reading "queen sheets deep pocket: auto, 16 clicks, 3 orders", no group token. Post-D5, that row cannot route: SOP-30's Section 3 gate requires group-level auto evidence, and forwarding the blend is F7 here and F3 there. P9 quarantines it and the operator re-pulls the Search Term Report by campaign. The split shows Auto-Close 7 clicks, 2 orders and Auto-Loose 9 clicks, 1 order: no single group meets the 3-order S-H1 bar. The blend would have graduated a term neither group validated. The corrected routing: the term stays on discovery economics (LEAVE TO HARVEST posture at source, with #30's pool carrying the re-eval), and the loose-side converter, sitting below the relevancy bar, routes S-H6 to the MKL decision first. The quarantine, the re-pull, and the corrected routes all write to the pass record.

The S-H5 breach. Same pass: Auto-Comp's month reads \$61.00 spend per discovered converter against the \$24.00 illustrative norm = 2.54x, above the 2x line, so the group's budget steps down 25% per S-H5: \$12.00/day becomes \$9.00/day, staged as one #38 row with S-H5 in the reason string, and close and substitutes retain their budgets per the scenario's priority clause. Prediction logged: cost per discovered converter back inside 2x within 4 weeks at the reduced cap, else the P12.3 retirement question opens.

WE-4 | Trawler and probe: a capped test that finds the price. Illustrative on the standard fictional base.

Standup. A verdict opens the two-tier test on Auto-Loose, suspected starved: the trawler holds the standing group at the L2 floor \$0.50; the probe stands up at the verdict-carried capped bid \$0.90 under CONTINUE TEST TO 50 CLICKS AT CAPPED BID, with the shared negation list attached at standup per SOP-29 P4.1 so the tiers cannot buy the same query twice.

Quota read. Day 11: probe at 50 clicks, spend \$45.00 ($CPC = 45.00/50 = \0.90), 3 orders, sales \$240.00, $CVR\ 3/50 = 6.0\%$, $ACOS\ 45.00/240.00 = 18.75\%$. Max Profitable CPC at the probe's $CVR = 28.00 \times 0.060 = \1.68 ; the \$0.90 discovered CPC sits at 0.54x ceiling and under the \$1.10 Expected CPC reference: viable.

Execution. The standing group re-seeds toward \$0.90 in L2-compliant steps: \$0.50 to \$0.75 this deployment (+50%, the per-change cap), \$0.75 to \$0.90 next cycle (+20%). The probe parks PAUSED-PENDING (REASON + RE-ENTRY), reason = quota reached, discovered CPC re-seeded, re-entry = a future two-tier verdict. Three T6 rows (standup, quota, re-seed step 1) each carry their prediction; a failed probe would instead have closed STOP-LOSS: CONCEDE OR DEFER with the freed cap returned to the envelope.

7. Failure Modes and Escalation

F1-F6 carry forward from v1.0 with their recorded incident motivations; F7-F11 are new at v2.0. Each row states the detection signal, the response, and the account-history incident where one exists.

Failure	Detection signal	Response	Motivating incident
F1. Second auto cloned because the first "stopped working"	A new blended auto with no group token and no driving verdict in T6; T3 diff	Pause the clone; run the actual diagnosis (negation neglect or bid drift); root-cause the creating session; consolidate per the Section 5 CONSOLIDATE DUPLICATES row	The diagnosis-skip pattern recorded in v1.0 Section 7; Doctrine Section 4 names cloning as the illegitimate case
F2. Blended read after the split exists	A verdict citing whole-auto ACOS where group rows are available in T3	Re-issue per group; post-split, the blended number is reporting, never a decision input (S-H3 made permanent)	Recorded in v1.0 Section 7; the account's pre-D5 state, one ACOS hiding four decisions
F3. Own-ASIN cannibalization	Own-catalog ASINs in complements or loose served terms (P3.4 scan)	Negative-ASIN same pass (the P1.5 wall); the #29 wall audit adds the campaign	Recorded in v1.0 Section 7; the self-cannibalization wall exists because the account paid to serve its own pages
F4. Loose drift explosion	Root-containing share of Auto-Loose served terms collapses after a relevance-model shift	Route C aggressive fencing; loose cap verified per P8; platform note to T5	Recorded in v1.0 Section 7; the 2024-2026 broad-match semantic widening Doctrine Section 3.3 documents
F5. Legacy auto never paused post-split	The blended auto spending past its recorded pause day (P1.7 daily watch)	Pause immediately as PAUSED-PENDING; reconcile the double-coverage window's spend in T6	Recorded in v1.0 Section 7: double coverage was buying the same queries at two structures
F6. Stage-skipping deployment	Substitutes, loose, or complements live on a product whose T6 standup entry cites no authorizing stage condition	Pause to the correct stage; the standup entry must cite the T2 condition; route to #22 when it originated in launch	Recorded in v1.0 Section 7; the unpriced-risk pattern Doctrine Section 2 names
F7. Token-less evidence forwarded downstream	A candidate or handoff row lacking its group token found by P9, or returned by SOP-30's F3 gate	Quarantine, re-pull by campaign, re-route on group-level evidence only; the forwarding session is named	New at v2.0; the WE-3 pattern, and the reason SOP-30 Section 3 carries its post-D5 gate
F8. Substitutes converters never handed to #16	Converting ASINs in Auto-Sub with no P3.3 handoff row in the pass artifact; #16's queue idle while conquest budget sits	Hand off same pass with evidence lines; a repeat routes as the sourcing-gap escalation SOP-16 Section 7 defines	New at v2.0; motivated by SOP-16's escalation clause for an empty flagged-candidate queue
F9. Auto carrying a head term unrouted	T2 child rows show auto share above 40% of a priority term's clicks with # Exact = 0 and no ADD EXACT + RESTRUCTURE receipt	P3.5 fires the receipt to #12 immediately; the pass reopens; weeks-at-state audit line stands	The account's named failure class: head terms carried entirely by auto with zero exacts (v4.0 Section 1 "Verdict First" and Section 16, Worked Pass 3, the six-piece pattern)
F10. Automation firing on protected auto scope	#38 T+1 attribution shows rule-fired changes on parallel-window, capped-by-design, probe, or event rows	P11 exclusion sync same day; the changes reviewed for damage; #37 F2 handling applies; audit value wins per P3 precedence	The stale-scope runaway recorded in SOP-37 Section 6's seeding pass (ladders found running on funded-loss campaigns)
F11. Event auto silently standing	A window-named campaign live past its expiry date with no conversion verdict in T6	Close per #32's playbook; conversion only through a fresh verdict, Objective tag, and pre-load check (P6.2)	New at v2.0; the P6 rule existed at v1.0 and this row gives it a detector

Escalation. F1 and F5 escalate on first recurrence (both are silent budget leaks). F3 recurring routes through #29's wall audit with session attribution. F6 routes to the launch handoff owner when it originates there. F7 recurring escalates to the PPC Lead with the pass records, because a token discipline that fails twice is a process gap, not a slip. F9 escalates on first occurrence exactly as SOP-30 escalates its head-term class: an uncovered head term compounds daily. F10 on push, event, or terminal-adjacent scope escalates same day per SOP-37's escalation clause. Anything in the human-mandatory tier (Section 9) is confirmed per the Authority Matrix before deployment, without exception.

8. Outputs and Artifacts

Event	Artifact set produced (gates clear with artifacts, not assertions), storage, retention, and downstream consumer
Weekly pass (P3, P7, P8, P12)	Per-group verdict rows in T2/T3; the staged #38 change set with verdict and scenario IDs; the deployed-and-verified set with its T+1 report; the route-count pass artifact (A/B/C/D, head-term receipts); candidate rows with tokens to #30; the P3.3 handoff list to #16; reactive negation sets with evidence lines to #29 standards; per-group yield reads into T3 (cost per discovered converter, WAS %, Harvest Yield, Graduations (Qtr)); T6 entries each carrying a prediction and review date. Storage: the product Command Workbook (T2, T3, T6) and the #38 staging store; retention per the workbook's append-only T6 rule, never purged. Consumers: #30, #16, #29, #38, the V5 scoring step, and the T0 zone I action summary.
D5 conversion (P1)	The conversion record: inventory line, four created T3 rows with group tokens and capped-by-design notes, migrated and fresh pre-load lists (recorded so the #29 audit has its baseline), the parallel-window watch lines, the legacy pause entry, and the T6 decomposition prediction. The conversion's gate artifact is the first post-split cycle showing four distinct group verdicts. Consumers: #29 (audit baseline), #37 (exclusion states), the Definition Register run ledger (D5 program row).
Standups and closes (P2, P4, P6, P10)	T3 rows created or parked with reasons and re-entry conditions; T6 entries citing the authorizing stage condition, test verdict, or window; shared-list attachments recorded (P4); the P11 state notices to #37. Consumers: #37, #32 (event reconciliation), the weekly pass.
Indexing-health read (P5)	The served-term distribution extract, the divergence read against the T1 syntax map, and the cross-framework handoff packet to the BM owner where divergence is material; a clean read logs as clean. Storage: T6 product-level entry with the extract attached. Consumer: the BM owner; PPC re-reads at the next slot.
Interface upkeep (P11)	Exclusion-state notices with campaign names and dates; the T+1 attribution review note; escalation packets for F10. Consumers: #37's register maintenance and the change-control record.

9. Skill-Conversion Block (Element 9 handoff)

Block	Specification
Input bindings (manifest IDs)	D5 targeting-group split (absence suppresses group verdicts for the window); A2 deployment state; A3 pacing and in-budget; A6 objective + Expected CPC; A7 eligibility; B5 coverage counts (T2 K8, per-group registry); B8 economics chain (T0 named cells with G9 state); D2 variation map; the ASIN-row identification on the search-term feed; the own-catalog ASIN list; the #16 target-status and failed-target lists; the ScaleInsights change log via the #38 T+1 attribution; T5 listing-change events (P5 trigger). An absent feed suppresses its dependent execution class per the manifest; the skill never improvises around a missing feed.
Deterministic logic references	v4.0 Section 8.9 scenarios S-H1 to S-H8 and the O6 discovery metric contract; gate logic G1, G2, G3, G6, G9, G10, G12; X10 and S-R12 for the head-term check; L1, L2, L5 lever rules; M6 and M9 references; SOP-14 P3 route definitions plus the two ASIN routes of this SOP; the group bid ordering per Doctrine Section 4; SOP-29 P1 pre-load tables and P4.1 shared lists; this SOP's Section 5 mapping. The skill executes mappings; it never computes its own thresholds, and it version-pins v4.0.

Human checkpoints (10-80-10, per the Authority Matrix v2.0 rows)	Per the "D5 conversions + auto operations" row (D and E: incoming operator; R: PPC Lead; vacancy per Section 2 rule 5): human-mandatory are every P1 conversion decision, every P2 staged standup, P4 probe standups and closes, P6 event-auto opens, P10 per-variation standups, group retirements, and every deployment confirm. Per the "Weekly verdicts (established US)" row, the pass's verdict executions run as E under the operator's assigned scope with the BM owner's R at the daily huddle. Per the "Bulk deployments (#38)" row, staging and deploy run as E in own scope. Per the "Automation rules (#37)" row, rule changes are never made here; P11 notices only. AI-executable with audit: the full P3 drafting per group, ASIN route lists, decomposition tables, token hygiene checks, S-H5 computations, F-mode detection, and T6 drafting.
Cadence	Weekly pass Monday with Tuesday deploy and T+1 verification; daily capped-budget and floor verification on loose, complements, and trawler rows; daily parallel-window watch during any P1 conversion; same-day P11 notices at every structure state change; P5 monthly and after listing rewrites; head-term receipts and RE-POINT execution same day.
Alarm specification	F1-F11 each carry the Section 7 detection signal, an evidence-line format (row, signal value, gate or rule reference), and a routing tier: F3, F5, F10 same day to the operator queue; F2, F4, F7, F11 into the weekly pass worklist; F1, F6, F8, F9 to the PPC Lead with the T6 history.
Skill outputs	Drafted per-group verdict rows, route lists and handoff packets, staged #38 sets, decomposition and yield tables, drafted T6 entries queued for confirmation, and F1-F11 alarms with evidence lines.

10. Version Governance

Owner: PPC Lead. Review monthly, and immediately on any v4.0 increment, at which point the Section 5 mapping and every S-H, O6, F10, F11, G2, G3, G12, L2, and X10 reference re-validates before any pass runs on it. Component-specific triggers beyond the monthly slot: (1) the D5 program completing across the portfolio, at which point P1 retires into a standing detection sweep and this SOP increments; (2) any Amazon change to the auto targeting groups or to the Search Term Report's Targeting tokens, which breaks the P9 token contract; (3) any #37 register change touching the auto-relevant rule classes, which re-runs the Section 5 and P11 checks; (4) the competitor depth feed landing, which upgrades the P3.3 evidence lines and is noted in the version history; (5) v4.1 resolving any Appendix A item, which retires that item's workaround the same increment. v2.0 change note: supersedes v1.0 in full; procedures expanded 6 to 12 (adds P7 bid mechanics, P8 budget and caps, P9 evidence hygiene, P10 per-variation autos, P11 automation interface, P12 yield close); verdict mapping expanded 9 rows to the full 25-string closed vocabulary with boundary statements; worked examples 1 to 4, recomputed on the standard base; failure modes 6 to 11 with incident attributions; Elements 11 and 12 added per the Build Kit standard; the v1.0 lever-order variant removed from the boundary text and parked at Appendix A per the Ledger's component-15 row. Registered in the Supersession Map; absorbs the scheduled D5 spec deliverable (the Completion Directive row closes against this document).

Companions: Master PPC Decision Criteria System v4.0 (all thresholds, stage conditions, and verdicts) | Strategic Doctrine v1.0 Sections 2, 3.4, 4 (the why) | Keyword Decision Record Template v3.0 (the closed vocabulary) | SOP-14 Broad (shared P3 machinery), SOP-30 Harvest, SOP-29 Negation, SOP-16 PAT, SOP-12 Exact, SOP-32 Events, SOP-37 Automation Rules, SOP-38 Bulk (boundaries) | Master PPC Data Architecture v1.0 and the Command Workbook READ ME (instrument mechanics) | PPC Cross-Reference Ledger (component #15 row).

11. Operator Ramp (Element 11)

Built from the Authority Matrix v2.0 rows that name this component's decisions: "D5 conversions + auto operations" (D, E: incoming operator; R: PPC Lead), "Weekly verdicts (established US)" (E: assigned scope), "Bulk deployments (#38)" (E: own scope), "Automation rules (#37)" (E, R: incoming operator, with D held by the PPC Lead), and "Event plans (#32)" (E). The ramp turns each row into a demonstrated competence; the checkpoint holder for review and audit is the PPC Lead per the R cell, and the vacancy rule (Matrix Section 2, rule 5) reverts D to the PPC Lead and E to the PPC Manager until the role fills.

Week	Demonstrated competence
Week 1: observation and artifact familiarity	Shadows one full Monday pass end to end. Locates and reads, unaided: the T3 "Auto Group Split (D5)" state per campaign, the T7 manifest row that suppresses group verdicts, the T2 K8 registry counts and child-row spend shares, one candidate row's complete evidence line with its group token, and one T6 chain for a prior probe. Reproduces the WE-1 decomposition arithmetic cold on fresh illustrative inputs, matching to the cent, in the manner the Economics Worked-Example Set's Use and Governance section requires of a new operator.

Week 2: supervised execution, checkpoint holder reviewing	Runs the P3 pass on one product with the PPC Lead reviewing every route before staging: all four routes produced with correct tokens, the head-term check computed from child rows, one Route B set staged to SOP-29 P3 standards with the converting-term guard applied, and the P12 yield close written. Drafts, and the PPC Lead confirms, one P4 probe standup with the shared list attached and one P11 exclusion notice. Zero forwarded rows without tokens across the week's work is the pass bar.
Week 4: unsupervised execution, checkpoint holder auditing	Owns the auto scope across the assigned products: runs the pass, stages and deploys per #38, closes T+1, and executes one live P1 conversion or P2 standup end to end (the D cell now exercised, with the PPC Lead auditing the artifacts after the fact rather than gating them before). The audit samples the route counts, one negation evidence line, the S-H5 computation, and the exclusion-state notices; two clean audit weeks close the ramp.

Live-defense questions (reasoning, not recall). (1) Why does a blended auto ACOS at band still fail the account, and what specifically does the four-group read expose that the blend cannot? (2) A substitutes ASIN converts twice: why is that a #16 candidate and not a graduation, and what would treating it as a graduation break? (3) Why does the parallel window exist at all, and what exact failure does pausing the legacy auto on day one risk? (4) Loose runs 60% ACOS and complements runs 85%: defend not pausing either, using the O6 metric set and the cap logic. (5) Why does the S-H5 step-down target loose and complements first while close and substitutes retain priority? (6) A rule steps a probe's bid up mid-test: name every contract that violated and the precedence rule that unwinds it. (7) Why must a graduation candidate's evidence be group-level after D5 when the blended total contains the same clicks and orders? (8) When is a second auto on the same product legitimate, and what is the tell that it is not?

12. Interface Contracts (Element 12)

Built from the Cross-Reference Ledger's component-15 Interfaces cell, cross-checked against Master PPC Data Architecture v1.0 for artifact names, tabs, and timing. Every row states what happens when the handoff does not arrive, because that is the clause the corpus was missing.

#	Interface	Sender → Receiver	Artifact and timing	Non-arrival consequence
I-1	Search-term feed with group split	Data Operations → auto operator	Raw Pull Pack: Sponsored Products Search Term Report with Targeting tokens; Monday 06:00 landing, workbook refresh by 09:00; T7 row D5 stamps PRESENT	Group verdicts suppressed for the window (never approximated from the blend); the pass records the suppression per G8; routing to #30/#16 holds; a second consecutive miss escalates to Data Operations with the manifest history
I-2	Pre-load sets at standup	SOP-29 (#29 standards) → this SOP's standups	The SOP-29 P1 pre-load table applied in the creation session; recorded in the T3 row as the wall baseline	The standup does not deploy: creation without pre-load is F1 in #29 and F3 exposure here; the session re-opens and the creating session is named
I-3	Failed-target list	#16 → this SOP (P3.4)	The failed-target register extract, same week as #16 declares exits	Auto-Sub keeps re-buying proven-bad pages (#16's F3 class); this SOP flags the gap to #16, stages the wall on receipt, and the wall audit notes the exposure window
I-4	Event plan arming	#32 → P6	The armed-window declaration with dates and caps, landing in T2 K2 Event Mode before the ramp date	No event auto opens (G12 bars non-plan structure); standing autos hold; a window observed live without a plan escalates to the deals owner per #32's own rules
I-5	Exclusion-state notices	This SOP (P11) → #37 register	Campaign names and states at every parallel-window, probe, event, sentinel, or hold change; same day as the change	Rules may fire on protected scope: F10 here, F2 in #37, changes reviewed for damage; a notice gap found by the T+1 attribution is root-caused to the missing checklist step

I-6	Economics chain	#41 refresh → T0 named cells → this SOP	AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC with the G9 freshness state, before the Monday 09:00 packet-ready read	Ceiling-referencing seeds and re-seeds hold (G9 PROVISIONAL); P4 quota reads defer their re-seed step; the pass proceeds on non-ceiling work and states what it could not price
I-7	Variation map	Economics & Context Pack (D2) → P10, P1	Keyword-to-variation map with ranked backups, current per its pack cadence	P10 standups barred; existing campaigns point at the preferred SKU with the D2-pending flag; RE-POINT VARIATION cannot execute and G6 RED on an unmapped row escalates same day
I-8	Candidate rows	This SOP (P3 Route A) → #30	One row per term, all sources attached, evidence line with group tokens; inside the Monday pass so #30's same-day pass consumes them	#30's harvest pass runs short and this pass is incomplete until the routes are logged; a token-blocked candidate travels as a quarantine note (P9), never as a blended row
I-9	Substitutes converter handoff	This SOP (P3.3) → #16 P1.1	The PAT candidate list with evidence lines (clicks, orders, price position where available), same pass	#16's queue starves at the top: the sourcing-gap escalation in SOP-16 Section 7 fires rather than standards dropping; a converter found unhandled in a later pass is F8 with the pass named
I-10	Reactive negation sets	This SOP (P3 Routes B, C, D) → #29 standards → #38	Negative exact, negative phrase, and negative-ASIN staged rows with evidence lines, Tuesday deploy	The WAS breach persists and compounds (E2-class quiet loss); the undeployed set carries forward as the next pass's first item and the deploy gap is noted in T6
I-11	Per-group yield reads	This SOP (P12) → T3 → the audit and A4 structure checks	Cost per discovered converter, WAS %, Harvest Yield, Graduations (Qtr), Auto Group Split (D5) state, every pass	Discovery grading goes blind: S-H5 is unenforceable that cycle and the M6 auto-share check loses its inputs; the pass artifact is incomplete until written
I-12	ScaleInsights auto-rule routing	This SOP → #37 register	Any observed auto-touching rule (new, edited, unregistered) reported to #37 the day it is seen	An unregistered live rule is #37's F1 and suspends on discovery; sitting on the observation makes this SOP the gap, so the report itself is the deliverable
I-13	Head-term receipts	This SOP (P3.5) → SOP-12	ADD EXACT + RESTRUCTURE receipt naming the carrying group, same day	The account's named failure class compounds daily (F9); escalation on first occurrence; the steering wall waits on the standup and the exposure window is logged
I-14	Decision log and predictions	This SOP → T6 → V5 scoring	One entry per deployed change with prediction, magnitude, horizon, review date; same day as deploy	Unscored actions do not accumulate (V1) and an unlogged change does not deploy at any tier (P10 completeness); the missing entry blocks the staged set at QA, not after it

Appendix A. v4.1 Amendment Items (the closing register; no number below appears as an operating value in this SOP)

Item	Gap, provenance, and the proposal carried to the threshold owner
AM-1	The auto lever-order variant. v1.0's boundary recorded a lever order (bids before placements before budget, pause last) that diverges from v4.0 Section 9 rule L1 (budget where truncated, then bid, then placement, then negation, then structure, then status). Carried per the Ledger component-15 row, which already parks it as a v4.1 item and declines to write it. Proposal: v4.1 either admits an auto-specific order with its rationale or retires the variant; until then L1 governs here (P7.3).
AM-2	D5 parallel-window duration and legacy-pause day. v4.0 carries no conversion-window standard. v1.0 deployed a 7-day window with the legacy pause on day 8; that operating value lives in each product's T6 conversion entry, not in this SOP's body. Proposal: admit the window and pause-day standard to v4.0's transition rules.

AM-3	Per-group bid seed references. v4.0 carries Expected CPC per keyword but no per-group auto reference for the four targeting groups. Until admitted, seeds derive from the blended auto's own group-level CPC history or the driving verdict's Expected CPC reference (P1.3, P2), and no seed is invented.
AM-4	The capped-by-design exemption class. G3 carries no exemption today; this SOP grades deliberately capped loose and complements rows on the O6 discovery metrics at their caps (grounded in the 70%+ discovery utilization target, v4.0 Section 3 Family F10, and the Doctrine's capped-discovery posture). Proposal: admit the exemption class and its marker (the T3 Budget Verdict cell recording the cap as the decision) to G3.
AM-5	Indexing-health divergence bar. P5 reads close-match served terms against the syntax map with no quantitative divergence threshold in v4.0. Proposal: admit a divergence measure and bar; until then the read is qualitative with evidence attached.