SOP-16: PAT and Conquest Operations

Version 2.0 · 2026-08-05 · Owner: PPC Lead · Supersedes SOP-16 v1.0 (2026-07-24) in full · Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds · Absorbs the scheduled PAT/conquest playbook deliverable

1. Verdict and Scope

VERDICT: conquest opens on evidence, never on ambition. A target earns spend only when it passes the v4.0 Section 8.7 entry test, row S-C1: "Candidate ASIN where we win ≥2 of price / rating / review count, or competitor flagged OOS." It is graded per target against the CPA ceiling of v4.0 Section 3, family F1 ("Sets the CPA ceiling (= margin $)") inside the Conquesting metric contract of Section 6, rule O6 ("Conquesting: CPA vs ceiling, win-probability inputs, ACOS 10-15"), and it exits the moment the flag basis reverses (S-C3) or the ceiling fails for 2 consecutive validated weeks (S-C2). The rule that carries the document is carried forward from v1.0 unchanged: **blended conquest ACOS is never a decision input**, because it averages targets that should not share a decision. Version 2.0 expands v1.0 from 7 procedures to 12, from a 6-row verdict mapping to the full scenario-wired mapping with quoted vocabulary strings, from 1 worked example to 4, and from 6 failure modes to 11, and it corrects the v1.0 open test: v1.0 opened on any single active weakness flag, which is looser than S-C1; v2.0 opens only on the S-C1 condition, and the two v1.0 flag classes that S-C1 does not admit as entry authorities (off-deal while we run one; conversion evidence alone) become queue-priority signals and ride to the v4.1 amendment list.

This SOP governs the whole operating life of product-targeting spend on both sides of the line: offense (conquest against competitor ASINs and refined categories) and defense (self-targeting on our own pages, complements, and virtual-bundle listings). It covers candidate intake from the three sources, flag-profile construction and the F12 win-probability classification, the open test and open mechanics, the weekly per-target read with its full decision tree, rotation and queue reallocation, exit and the failed-target register, category-refined standing conquest, the defense PAT standing set including virtual-bundle self-targeting, counter-conquest evaluation, reporting and cascade integration, and the event-window overlay.

1.1 The two structure classes of PAT spend

Every PAT dollar in the account lives in exactly one of two classes, and the class determines which read is legal. **Class 1, per-target conquest:** ASIN-targeted campaigns grouped by flag class, where every target is its own decision row and the blended number appears in reporting only. This class holds every offensive open and every counter-conquest response. **Class 2, refined-class conquest and the defense set:** category campaigns whose refinements (price floor above ours, rating ceiling below ours) enforce the S-C1 axes continuously, plus the standing self-targeting defense set. The category refinement is the one legal blend in conquest, because the refinement makes the class coherent; that exception is exactly why P7 exists to keep the refinements true. Defense self-targeting is graded on its own product (page occupancy for own-ASIN defense, basket lift for complements and virtual bundles), never on conquest economics and never on ACOS optics.

1.2 Boundaries with adjacent components

Adjacent component	The seam, stated so it is owned
Criteria System v4.0	Decides every verdict, threshold, band, scenario, and limit: Section 8.7 rows S-C1 to S-C6, Section 3 family F12 (BEATABLE / WINDOW), family F1 (the CPA ceiling), Section 6 rules O2, O6, O7, Section 5 gates, Section 9 levers, Section 13 vocabulary. This SOP executes; it never decides. A threshold appearing here instead of v4.0 is a defect.
SOP-15 Auto operations	Owns the substitutes auto that scouts candidates. SOP-15 P3.3 routes converting substitute ASINs here with evidence lines; SOP-15 routes, this SOP decides. Failed targets return to SOP-15's substitutes surface only as SOP-29 negative-ASIN entries.
SOP-29 Negation architecture	Owns walls, negative-ASIN mechanics, shared lists, and the wall audit. This SOP hands the failed-target register entries to SOP-29 the same week as every exit (SOP-29 Section 3 names "the failed-PAT-target list from #16" as a required input); SOP-29 P1 pre-loads it into every substitutes creation.
SOP-26 Mature/defend + code red	Owns incursion detection (P4) and the defense-stage posture including the annual defense-cost read (P2). SOP-26 P4 routes persistent incursors here for the P10 counter-conquest evaluation; this SOP returns the evaluation record. Defense PAT mechanics live here per the Ledger row for #26.

SOP-17/18/19 playbook	Owns SD contextual execution. The playbook's P7.2 runs this SOP's flag workflow unchanged and states it restates zero flag rules; the target list, flags, and open, rotate, exit calls arrive from this SOP.
SOP-30 Harvest	Owns keyword graduation. Keyword converters route to #30; ASIN converters route here via SOP-15 P3.3. No graduation logic is restated in this SOP.
#33 Sibling arbitration	Owns family term ownership. Conquest targets are ASINs, not shared terms, so G11 rarely binds; where a defense-set decision touches a family-owned surface, the declared owner is checked first.
#32 Deal and event window	Owns event plans, ramps, caps, and reconciliation. This SOP executes the S-C5 / S-E3 stand-downs inside armed windows and freezes non-plan structure per G12.
#38 Bulk mechanics	Owns staging format, upload, and T+1 verification. Every change this SOP produces exists first as a staged set per #38 and consumes the T+1 report.
#37 Automation rules	Owns the rule register. Any automation touching PAT bids is inventoried there; a rule change touching conquest increments this SOP's Section 5 mapping check.
Strategic Doctrine v1.0	Owns the why: Section 9 states the conquest doctrine this SOP executes ("conquest is opened on evidence, not ambition") and the defense-PAT and substitutes-scout roles.

2. Trigger Conditions

Verdict first: every trigger below names its detecting artifact and the maximum allowed lag from signal to procedure start. Where the corpus states no numeric lag, the lag is stated in the corpus's own cadence vocabulary (same day, same pass, same week); a clock-hour SLA for the reversal watch is amendment item A-5.

Trigger	Detecting artifact	Procedure invoked	Maximum lag
Candidate arrival	The SOP-15 P3.3 PAT candidate handoff list in the weekly pass outputs; the competitor tracker adding an ASIN (T2 column DR "Competitor Depth (D8)"); the P7 category scan	P1 intake, then P2 flag profile	Same weekly pass; profile complete before the next Monday cascade opens
Flag change on an open target	The daily reversal watch over T2 column DR source fields (price, rating, review count, stock, deal state) on every open target	P4 read for that target	Same day the feed shows the change; disposition staged same day
Weekly cadence	Monday cascade reaches conquest rows: T3 filtered to Type = PAT plus the per-target rows from the targeting join	P4 full per-target reads; P11 reporting	Inside the Monday cascade; deployment Tuesday per #38
Competitor deal window live	Deal-state field in the target's flag profile / T2 DR; T5 "Competitor entry" event rows	P4 branch 2 (S-C5 stand-down)	Same day; during armed events, on the daily event cadence per G7
WINDOW review date	The 2-week review date written at every OOS-class open per S-C1 ("competitor stockouts are time-boxed offensives (review in 2 weeks)")	P4 dated re-review of that target	On the dated day; a lapsed date is failure mode F4
Incursion escalation	SOP-26 P4 routes a persistent incursor with its response record	P10 counter-conquest evaluation	Inside the same weekly cycle as the escalation
Our reprice lands	T5 "Price move" event row on the product	P7 refinement refresh	Same week as the price move
Quarterly refresh date	The v4.0 Section 14 rule V4 quarterly slot on the cadence calendar	P7 refinement refresh; P8 defense floor re-check with the #28 ladder	Inside the quarterly window
VB status check result	The weekly virtual-bundle knocked-out check from the catalog owner	P9 pause or resume	Same week as the check
Event window arms	G12 arms on a confirmed deal window touching the product (#32 plan)	P12 event overlay	Same day the window arms

3. Preconditions and Gates

Verdict first: no open, scale, or per-target verdict runs until the gates below read current; a failing gate suppresses its dependent execution class and the pass records what it could not see (v4.0 Section 5, gate G8). Each gate names its manifest ID from the T7 Manifest of the PPC Command Workbook and the observable that proves it passed.

Gate (manifest ID)	Requirement and the observable that proves it
Competitor depth current (T7 row "D8 Competitor depth (price/rating/deal/OOS)")	Flag-profile source fields present with source and date per flag. Observable: T7 Status cell for D8 and per-flag date stamps on the profile. The T7 template ships this row MISSING with blocked classes "Conquest open/rotate/exit; defend-context reads": while MISSING, opens are suppressed to flags per G8. Where the row reads PARTIAL, a decision on partial flags with the coverage noted on the profile is legal; a decision on stale flags is not, and no OPEN deploys on a flag whose date predates the current read window. The feed is Ledger external dependency D4; its landing upgrades this gate and is a Section 10 version trigger. A numeric flag-age bar is amendment item A-4.
Per-target rows readable (T7 row "A8 PAT rows first-class")	The targeting join delivers per-target clicks, spend, and orders for every open target. Observable: the per-target read sheet populates one row per target from the Sponsored Products targeting data. The template ships A8 PARTIAL ("row-level PAT reads thin"): rows the join cannot deliver read as flags, never as approximations from the campaign blend, mirroring the SOP-15 Section 3 suppress-not-approximate rule for missing group splits.
Economics chain live (T7 row "B8 Row economics (BE chain)")	T0 named cells live: AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC. Observable: T2 column CO "CPA Ceiling $" and CM "Expected CPC $" populated on the product's rows. Per the SOP-12 v2.0 Section 3 precedent, the G9 PROVISIONAL tag carries onto any ceiling-referencing open while margin is blended or stale; the G9 gate-ID dual use is already registered as an open amendment item by the #41 build and is not re-opened here.
Sample sufficiency (G2)	v4.0 Section 5, gate G2: "Bid reads need ≥15 clicks (7d); CVR verdicts ≥100 clicks... Below threshold, fall back to the 60-day window; if still thin, the verdict is EXTEND TEST, never a rate judgment." Observable: T2 column V "Clicks 7/30/60d" and W "Sample State" on the target's row. S-C2's "validated weeks" means weeks that pass this gate.
Inventory adequacy (G6) and variation (D2)	The advertised SKU on every conquest and defense campaign is the D2-preferred variation, and its T2 column AD "Inventory Band" reads GREEN for any open or scale (v4.0 Section 8.3 row S-I1; RED blocks per gate G6). Observable: T3 column H "Targeted Variation + Band."
Eligibility (G4)	Advertised SKU eligible. Observable: T3 column O "Eligibility + Reason." A flip halts all execution for the target same day.
Event mode (G12)	T2 column AI "Event Mode" clear, or the #32 event plan names the change. S-C5 / S-E3 stand-downs are always live inside windows.
Prediction and horizon present (P10, O7, V1)	Every OPEN carries its per-target prediction (metric, direction, magnitude, horizon, review date per rule V1) and the O7 horizon date before it deploys. Observable: the T6 row exists with a review date. Per rule P10, a record missing these does not deploy; an open without them is failure mode F7.
Do-not-target exclusion list current	The standing exclusion list (brand-safety and policy exclusions) is checked before any open. This closes the Ledger's pending fold-in for #16 structurally; the list's owner and criteria are amendment item A-6. Observable: the profile carries an exclusion-check line.

4. Step-Level Procedures

P1. Candidate intake and queue maintenance (three sources)

Time: 20-30 minutes inside the weekly pass.

1. Open the conquest candidate queue for the product (the standing register block attached to the product's Command Workbook beside T3; one row per candidate ASIN). Columns: Candidate ASIN, Brand, Source (SUBSTITUTES / TRACKER / CATEGORY-SCAN), Evidence line, Date added, Profile state (PENDING / BUILT), Queue rank, Exclusion-check result.
2. **Source 1, substitutes converters:** take the SOP-15 P3.3 handoff list from the pass outputs. Each entry arrives with its evidence line (clicks, orders, our price vs theirs where the feed carries it). These are conquest-lite validated and enter at the top of the queue: proven leakage from their page to our listing is the strongest priority signal the system has. Priority is all it is: entry to spend still requires the P3 open test, because per v4.0 Section 8.7 row S-C1 the target lists "feed from the auto substitutes group and Keepa signals" while the entry condition is the win test.

3. **Source 2, the competitor tracker:** the tracked competitor set per family, read from T2 column DR "Competitor Depth (D8)" and the Keepa pull rows of the Raw Pull Pack. Any tracked ASIN not yet in the queue is added with source TRACKER.
4. **Source 3, category scans:** in the Ads Console, open Sponsored Products campaign creation to the product-targeting step, select the product's leaf category, and read the suggested-products list filtered to the product's price and rating window; record top sellers inside the window that are not yet queued, source CATEGORY-SCAN. This scan refreshes with every P7 refinement refresh.
5. Run the exclusion check (Section 3, last gate) on every new candidate; failures are marked EXCLUDED with the reason and never profiled.
6. De-duplicate against the failed-target register (P6): a register ASIN re-enters the queue only under the re-open rule (a new flag class per P6 step 4), and the queue row must name the new class.
7. Queue ranking: substitutes converters first, then tracker candidates whose profile already shows a BEATABLE read, then category-scan entries; ties break toward the larger evidence line. A queue with zero candidates whose profiles could pass P3 while conquest budget sits idle is a sourcing gap and escalates per failure mode F8, never a reason to lower the S-C1 bar.

P2. Flag-profile construction and F12 classification

Time: 5-8 minutes per candidate; batched inside the weekly pass.

1. For each PENDING candidate, build one flag profile with these fields, each flag carrying its source and date: Their price and ours with the delta %; their rating and ours; their review count and ours with the gap; stock state (IN / LOW / OOS); deal state (ON / OFF); conversion evidence (clicks and orders from the substitutes evidence line where present); coverage note (which fields the feed could not fill).
2. Classify per v4.0 Section 3, family F12: "we win a comparison on ≥2 of price / rating / reviews = BEATABLE; competitor OOS = WINDOW (time-boxed)." Write the class and the axes won into the profile. A profile that wins fewer than 2 of the 3 axes and shows no OOS is NOT-ENTERABLE and stays in the queue as a watch row.
3. Where the competitor depth row (T7 D8) is PARTIAL, a classification on the fields present is legal with the coverage noted; where it is MISSING, the profile is a flag, not an entry basis (Section 3, first gate).
4. Deal-state ON at profile time does not block classification; it blocks the open until the window ends, per v4.0 Section 8.7 row S-C5.
5. Profile state flips to BUILT; the profile is the gate artifact for any open on that target.

P3. The open test and OPEN mechanics

Time: 15-25 minutes per open, including staging.

1. **The test:** the target opens only when its profile satisfies v4.0 Section 8.7 row S-C1: "Candidate ASIN where we win ≥2 of price / rating / review count, or competitor flagged OOS." Anything else waiting in the queue is ambition, and an open without this basis is failure mode F2. A WINDOW-class open (OOS basis) writes its 2-week review date at open, per the same row: "competitor stockouts are time-boxed offensives (review in 2 weeks)."
2. **Gates:** Section 3 in full: D8 currency with per-flag dates inside the current read window; G6 GREEN on the D2-preferred advertised variation; G4 eligible; G12 clear or plan-named; economics chain live with the G9 state carried; the exclusion check passed.
3. **Structure:** conquest campaigns group by flag class per product, one ad group per class (BEATABLE-PRICE-LED, BEATABLE-RATING-LED, WINDOW), never one campaign per ASIN and never a mixed-class ad group; the class grouping is what keeps the per-target rows readable and the reporting coherent. The campaign name follows the v4.0 Section 4.2 convention with objective segment Conquest; naming non-compliance is the SOP-12 F6 pattern and renames ride alone.
4. **Objective:** per v4.0 Section 6, rule O2, "Competitor ASIN target → Conquesting." The T3 objective tag and the campaign-name objective segment are set in the same staged set.
5. **Targeting:** exact ASIN targeting on the candidate. Expanded targeting (their variations) is added only where the profile shows the flag basis holds family-wide (each variation's price or rating read carried in the profile); a flag proven on one child does not license the family.
6. **Entry bid:** seeded at the Expected CPC reference (T2 column CM; v4.0 Section 3, family F9 "Expected CPC (what the position should cost)") within the Section 9, rule L2 limits ("±50% per change, $0.50 floor"). v1.0 cited a "conservative entry bid per the v4.0 category CPC reference"; v4.0 contains no category CPC reference, so the seed is Expected CPC and a PAT-specific entry-bid rule is amendment item A-1. The bid is never invented; a driving verdict that carries a bid governs.
7. **Budget:** the class campaign carries a capped daily budget from the product envelope per the #31 instruments; conquest competes inside the envelope by rule M5 ordering and never by escalation (v4.0 Section 12, rule P2).

8. **Prediction and horizon:** the open logs its per-target prediction in T6: CPA at or under the ceiling (T2 column CO) by the review date, conditional on the flag basis holding, plus the O7 horizon date, since per v4.0 Section 6, rule O7 "every Ranking and Conquesting objective carries a horizon date" and renewal at the horizon requires "CPA within ceiling." No prediction, no deploy (rule P10; failure mode F7).
9. **Deploy:** stage per #38 with the scenario ID S-C1 in the reason string; Tuesday by default, same day where the driving signal is dated (a WINDOW open on a fresh stockout does not wait). T+1 verification returns per #38.

P4. The weekly per-target read

Time: 30-45 minutes for a product with 8-15 open targets.

1. Pull the per-target rows: the Sponsored Products targeting data joined to the campaign child rows of T2 (Data Architecture Section 6.1), trailing 7 days with the 30-day and 60-day windows behind it per G2's fallback. One row per target: clicks, spend $, orders, sales $, CPA, plus the current flag fields and their dates from the profile.
2. Score last cycle first: V5-score every T6 conquest row whose review date lands this cycle (v4.0 Section 14, rule V1: unscored actions do not accumulate).
3. Verify gates on every row before any metric read: G2 sufficiency (column V), the flag dates inside the read window, G6 band, G12 event flag. A failing gate executes the gate's response and defers the metric verdict for that row.
4. Re-run the F12 classification on current flag fields for every open target: the entry basis is re-tested at every read, because per v4.0 Section 8.7 row S-C3 "the leading indicator outranks the trailing one."
5. Walk each row through the decision tree below; write the disposition, the scenario ID, and the numbers into the read sheet; stage the resulting changes per #38 with scenario IDs in the reason strings.
6. The blended class line is computed for reporting only and is labeled REPORTING-ONLY on the sheet; a verdict citing it is failure mode F1. The one legal class read is the refined category campaign per P7.

> **P4 decision tree (per target, in order; first match wins):**
> **0.** Row fails G1 validation or G2 sufficiency → no rate verdict: quarantine flag or EXTEND TEST per gate text; budget assurance only.
> **1.** Entry basis gone: BEATABLE target now wins <2 of the 3 axes, or WINDOW target restocked or the 2-week review date arrives → exit this target at this review regardless of CPA (v4.0 8.7 rows S-C3 and S-C1 time-box); released spend to P5; register entry with reason FLAG-REVERSED or WINDOW-CLOSED.
> **2.** Target's deal state ON → stand down on that target for the window and stage the resume date (row S-C5: "conquesting into a deal price is buying the worst comparison of the month"); not an exit; the target's clock pauses.
> **3.** Basis holds · CPA ≤ ceiling (column CO) · ACOS inside the O6 10-15 band → continue; scale steps are legal per row S-C1 "Enter/scale conquest within the 10-15% band," inside the L2 limits. If the target-page CVR ≥ our own detail-page CVR, row S-C4 adds: scale within band and mine the target's page for the winning offer element; the insight routes to the listing team.
> **4.** Basis holds · CPA ≤ ceiling · ACOS above the band → no scale; one L2 ladder step down toward the CPC that prices the band; re-read next cycle.
> **5.** Basis holds · CPA > ceiling, first validated week → mark week 1 of the S-C2 clock on the row; one L2 step down; no exit yet, because row S-C2 exits on "CPA above ceiling for 2 consecutive validated weeks on a target ASIN."
> **6.** Basis holds · CPA > ceiling, second consecutive validated week → EXIT per S-C2 ("Exit that target; keep the campaign; rotate the target list"); P6 register entry and the SOP-29 handoff the same week.
> **7.** Refined category class row above band → P7 step 4 (row S-C6): tighten the refinement one step; a class that cannot reach band after two refinement passes exits.

P5. Rotation: queue reallocation on any exit

Time: 10 minutes per rotation event.

1. On any tree-branch-1, branch-6, or S-C6 exit, the released weekly spend is quantified on the read sheet (the exiting target's trailing 7-day spend).
2. The next queue candidate whose profile passes the S-C1 test opens per P3 in the matching flag-class ad group; the released spend is its starting reference. S-C2's own text is the instruction: "keep the campaign; rotate the target list."
3. If no queued candidate passes S-C1, the released budget returns to the product envelope per the #31 instruments and the queue state logs the F8 sourcing signal; the bar does not move.
4. Every rotation writes a T6 pair: the exit row scored against its open prediction, and the new open's prediction with its horizon.

P6. Exit and the failed-target register

Time: 10-15 minutes per exit including the handoff staging.

1. Every exit enters the failed-target register (the standing register block beside the queue) with: Target ASIN, Exit reason (CEILING-S-C2 / FLAG-REVERSED-S-C3 / WINDOW-CLOSED-S-C1 / REFINEMENT-EXIT-S-C6), Flag state at exit (the axes and values), Weeks open, Total spend and orders, Date.
2. **The SOP-29 handoff, same week:** CEILING exits are staged as negative-ASIN into the substitutes auto per SOP-29's standards, so the scout never re-buys a proven-bad page; SOP-29 Section 3 lists this register as a required input and its P1 table pre-loads "Failed-PAT-target ASINs (negative ASIN, from #16's list)" into every substitutes creation. A missed handoff is failure mode F3.
3. FLAG-REVERSED and WINDOW-CLOSED exits record the reversal, not a failure of the page: they do not stage a negative-ASIN, because the economics were bought against a weakness that is gone, and the target may re-enter the queue when a basis returns.
4. **The re-open rule:** a CEILING-exited target re-opens only on a NEW flag class: a target that failed while price-led is not blocked when its rating later collapses or it goes OOS; it is blocked from re-opening on the same class alone. The queue row for any re-entry names the new class and the register row it supersedes.

P7. Category-refined standing conquest (the one legal blend)

Time: 25-40 minutes per refresh; the weekly read rides inside P4.

1. One category campaign per product family with refinements set to structural weakness on the S-C1 axes: price floor above our current price, rating ceiling below our current rating. The refinement IS the flag test running continuously: the campaign can only serve against competitors weaker on both axes, which is what makes the class coherent enough to blend.
2. Band construction is mechanical from our own current values: the price floor sits one cent above our live price (T5-confirmed), the rating ceiling one displayed step below our live rating. Both values, their source reads, and the refresh date are recorded on the refinement band record.
3. Budget is capped by design (the T3 note set, per the SOP-15 precedent for capped-by-design rows); the class is graded on class CPA against the ceiling and the O6 band, the single exception to the per-target rule, and the exception is why the refinements must be right.
4. Above-band reads execute v4.0 Section 8.7 row S-C6: "Refine or exit: tighten the refinement set (price floor, rating floor) one step; a category campaign that cannot reach band after two refinement passes exits." The pass count lives on the band record.
5. **Refresh triggers:** the quarterly slot (v4.0 Section 14, rule V4) and, always, our own reprice the same week (the T5 price-move event): refinements aimed at last quarter's price position are failure mode F6, because a reprice down strands newly weak competitors outside a stale floor and a reprice up lets stronger ones in.
6. Every refresh re-runs the P1 category scan (source 3), stages the refinement edits per #38, and logs the T6 prediction: class CPA relative to ceiling and band by the next quarterly read.

P8. Defense PAT: own-ASIN and complement self-targeting

Time: 15-20 minutes weekly; the quarterly floor re-check rides the #28 ladder cycle.

1. **Own-ASIN self-targeting at floor, permanent:** our detail pages stop advertising competitors; the cheapest defense in the system per Doctrine v1.0 Section 9. Floor bids are the discovered floors from the #28 ladder, read per term from T4 Placement & Bid columns "Discovered floor $" and "Ladder state (step # + date)"; drift above floor without a verdict resets to floor in one staged step (the DEFEND AT FLOOR execution, as in SOP-12's mapping). The expansion trigger is v4.0 Section 8.6 row S-D2: "Branded IS slipping, or competitor ads appearing on our detail pages → Restore one ladder step; open/expand PAT self-targeting to occupy our own pages; the syntax defense budget (M8) is the ceiling."
2. **Complement self-targeting, capped:** our complementary products cross-sell on our own pages; graded on the basket read (units per order and AOV against the pre-standup baseline), never on ACOS optics (failure mode F5). Own-catalog negative-ASIN walls in the complements auto belong to SOP-29's P1 table and are not restated here.
3. **The cost of defense:** the annual defense-cost read belongs to SOP-26 P2 and is not run here; this SOP supplies the defense set's spend rows to it.
4. Weekly, the defense rows are read for occupancy (competitor ads present on our pages yes/no) and floor integrity; the S-D2 expansion executes the week the trigger shows.

P9. Virtual-bundle self-targeting

Time: 10-15 minutes weekly where VBs exist.

1. Runs only where virtual-bundle listings exist: PAT on our VB listings targeting our own component listings, lifting units per order and basket value on our own traffic. Each VB campaign is a T3 row (Type = PAT, Objective Tag = Defensive) whose column AP "Moves Received" carries the P9 pause and resume actions.
2. **Grading:** AOV and units-per-order lift against the pre-standup 4-week baseline, never on ACOS optics (failure mode F5): the tactic's product is basket size, and VB conversions frequently land on component listings, so the campaign row's own ACOS understates the tactic. v4.0 carries no numeric lift bar for this read; the metric standard is carried from v1.0 and a numeric bar is amendment item A-3.
3. **The knocked-out check, weekly:** the catalog owner's VB status check; a knocked-out VB (bundle broken by a component change) pauses its campaign the same week as PAUSED-PENDING with reason VB-KNOCKED-OUT and re-entry criterion VB-RESTORED. A check that does not arrive is treated as unknown, and unknown is not servable evidence: the campaign pauses pending the check (Element 12, row I-6).
4. The baseline clock annotates knocked-out gaps so lift reads never span a broken period.

P10. Counter-conquest evaluation

Time: 20-30 minutes per escalation.

1. Input: a persistent incursor escalated from SOP-26 P4 with its response record (the incursion held above the defense ceiling).
2. The incursor gets the full P2 flag profile and F12 classification like any candidate. The response is symmetric and the math decides it.
3. Classified BEATABLE or WINDOW: open per P3 in the matching class group, with the escalation reference in the T6 row; the target then lives under P4 like any other.
4. Classified NOT-ENTERABLE: the answer stays defensive, and the evaluation record returns to SOP-26 with the profile attached, because conquest against a competitor who wins the comparison is a donation (the v1.0 line, carried; failure mode F10 exists for the violation).
5. Every evaluation, opened or declined, is a written record: profile, classification, decision, and where declined, the counterfactual that would flip it (rule J5).

P11. Reporting, logging, and cascade integration

Time: 15-20 minutes inside the Monday cascade.

1. Enter the conquest rows only after T0 and T1 have been read; product posture and family caps govern everything below them (Data Architecture Section 2, rule 1).
2. The weekly gate artifact is the per-target read sheet: one row per open target with clicks, spend, orders, CPA vs the CO ceiling, flag fields with dates, F12 class, disposition, scenario ID; plus the refined-class row and the defense and VB rows with their own reads; plus the REPORTING-ONLY blended line.
3. Every staged change carries its scenario ID in the #38 reason string; no scenario ID, no deploy.
4. T6 rows per deployed change with predictions and review dates; the pass opens next cycle by scoring them (P4 step 2). Material structure events (class campaign creation, category refresh, VB pause) also write T5 product-log entries.
5. The BM owner's R on the verdict rows is the daily huddle per the Authority Matrix Section 2: review, not veto; disagreements route up, not sideways.

P12. Event-window overlay

Time: 10 minutes at arm and disarm; daily checks ride the event cadence.

1. Arm on G12: the confirmed window from the #32 plan flags every conquest row through the window plus the post-event guard.
2. Inside the window, competitor deal detection executes v4.0 Section 8.12 row S-E3, which routes to row S-C5: "stand down on that target for the window"; detection rides the daily event cadence (gate G7).
3. Competitor stockouts during events remain legal WINDOW opens where the event plan does not forbid the structure change; the 2-week S-C1 review date still applies and the read windows are excluded from trend verdicts per gate G10.
4. Structure changes beyond plan-named actions freeze per G12; stand-downs and reversal exits are never frozen, because they stop spend rather than start it.
5. Post-window, stand-downs resume by their staged resume dates and the reconciliation belongs to #32.

5. Decision Points: Verdict Mapping

Verdict first: all verdict issuance, thresholds, and scenario logic live in v4.0; the record-level verdict strings below are quoted verbatim from the Keyword Decision Record Template v3.0, sheet "Bands & Verdicts Reference," section "CLOSED VERDICT VOCABULARY," or from v4.0 Section 13.1 "The Closed Verdict Vocabulary," each with its source named. This SOP carries zero local thresholds. The conquest family is a single vocabulary string, "OPEN / ROTATE / EXIT CONQUEST"; the disposition (which of open, rotate, exit) lives in the scenario line and campaign moves of the record, which is also how the SOP-17/18/19 playbook P7.2 consumes it for SD contextual. The v1.0 strings "REFRESH REFINEMENTS" and "DEFENSE STATES" appear in neither vocabulary and are retired: both are procedure executions (P7, P8-P9), not verdicts.

Verdict string (source)	Execution in this SOP
"OPEN / ROTATE / EXIT CONQUEST" (KDR v3.0, Bands & Verdicts Reference)	The record-level verdict for every per-target and refined-class conquest row. Disposition OPEN executes P3 on an S-C1 basis; disposition ROTATE executes tree branch 1 (S-C3 reversal or S-C1 window close) plus P5 reallocation; disposition EXIT executes tree branch 6 (S-C2, two consecutive validated weeks above ceiling) plus P6. Deal stand-downs (S-C5), in-band scaling (S-C1), scale-and-mine (S-C4), and refine-or-exit (S-C6) ride the same record as campaign moves with their scenario IDs.
"DEFEND AT FLOOR" (KDR v3.0 reference; also v4.0 Section 13.1)	The standing state of the P8 defense set: own-ASIN self-targeting held at the discovered floors; drift restores to floor in one staged step; S-D2 expansion executes on its trigger.
"PAUSED-PENDING (REASON + RE-ENTRY)" (KDR v3.0 reference)	P9 knocked-out VB pause (reason VB-KNOCKED-OUT, re-entry VB-RESTORED) and any conquest campaign paused with a named reason and re-entry criterion.
"NEGATE (PRE-LOAD / REACTIVE / STEERING)" (KDR v3.0 reference)	Receipt-side only: the P6 register handoff stages failed-target negative-ASIN entries that SOP-29 owns and deploys; this SOP never restates wall rules.
"EXTEND TEST" (v4.0 Section 13.1)	Tree branch 0 on any open target below G2 sufficiency at its review: capped continuation to the sufficiency threshold; never a rate judgment.
"FLAG: DATA REPAIR" (v4.0 Section 13.1)	Any G1 quarantine or D8-MISSING state on conquest rows: flag produced, no recommendation, repair task with owner and date; opens suppressed while the manifest row blocks the class.

Boundary receipts: "ENTER SBV ARBITRAGE" and the brand-format rows belong to the SOP-17/18/19 playbook; "GRADUATE FROM DISCOVERY" belongs to #30; SOP-12's mapping states for conquest rows that no exact-campaign execution occurs "unless the verdict names a shared-term bid step." The quoted condition and verdict text of v4.0 Section 8.7 rows S-C1 to S-C6 used throughout this SOP: S-C1 "Candidate ASIN where we win ≥2 of price / rating / review count, or competitor flagged OOS → Enter/scale conquest within the 10-15% band; competitor stockouts are time-boxed offensives (review in 2 weeks)." S-C2 "CPA above ceiling for 2 consecutive validated weeks on a target ASIN → Exit that target; keep the campaign; rotate the target list." S-C3 "Win-probability degrades mid-flight (competitor price cut or review surge flips the ≥2-of-3 test) → Exit at the next review, not at the CPA lag: the leading indicator outranks the trailing one." S-C4 "Conquest CVR on target pages ≥ our own detail-page CVR → Scale within band and mine the target's page for the offer element winning the comparison." S-C5 "Competitor deal window live (deal state flag) → Stand down on that target for the window; resume at window end." S-C6 "Category targeting with refinements running above band → Refine or exit: tighten the refinement set (price floor, rating floor) one step; a category campaign that cannot reach band after two refinement passes exits."

6. Worked Examples

Verdict first: four examples, all illustrative and internally consistent; every value is computed on the page. No live product economics spine is attached to this build, so all four use the standard fictional base and say so: AOV $80.00, contribution $28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%. The conquest CPA ceiling equals contribution per order, $28.00 (v4.0 Section 3, family F1: "Sets the CPA ceiling (= margin $)"); the conquest band is ACOS 10-15% (Section 6, rule O6).

WE-1 · The clean weekly read: three targets, three different decisions, one misleading blend. Illustrative values on the standard fictional base; ceiling $28.00; band 10-15%.

T1 (substitutes source, 2 prior conversions in the evidence line). Profile: their price $89.60 vs ours $80.00 (+12.0%); rating 4.2 vs our 4.6; reviews 310 vs our 1,240. Wins price, rating, and reviews: 3 of 3, BEATABLE. Open 3 weeks. This week: 84 clicks at CPC $1.15 = spend $96.60; 9 orders (CVR 10.7%); sales 9 x $80.00 = $720.00; CPA $96.60 / 9 = $10.73; ACOS $96.60 / $720.00 = 13.4%. Tree branch 3: basis holds, CPA under ceiling, ACOS inside band: continue; a scale step is legal per S-C1 and stays inside the L2 limits.

T2 (tracker source, WINDOW class: opened 2 weeks ago on their stockout; today is the S-C1 review date). This week: 40 clicks at CPC $1.05 = spend $42.00; 4 orders; CPA $10.50; ACOS $42.00 / $320.00 = 13.1%. Flag check: RESTOCKED. Tree branch 1: exit at this review regardless of the passing CPA; the weakness that made their page leak is gone, and yesterday's CPA was bought against yesterday's stockout. Register entry WINDOW-CLOSED; no negative-ASIN (the page did not fail); released $42.00/wk to P5.

T3 (tracker source, opened BEATABLE on price + reviews). Flag re-check: rating recovered to 4.6, but price (+9.5%) and reviews (their 480 vs our 1,240) still win: 2 of 3, basis holds. This week: 82 clicks at CPC $2.50 = spend $205.00; 5 orders; CPA $41.00; ACOS $205.00 / $400.00 = 51.3%. Last week's validated read was CPA $33.60. Second consecutive validated week above the $28.00 ceiling: tree branch 6, EXIT per S-C2. Register entry CEILING-S-C2 ("their page converts its own traffic despite the price gap"); negative-ASIN staged to the substitutes auto per SOP-29 the same week; re-open blocked on the price-led class, open to a future rating or OOS class.

The blend that would have lied: total spend $96.60 + $42.00 + $205.00 = $343.60; total orders 18; blended CPA $343.60 / 18 = $19.09, comfortably under the $28.00 ceiling; blended ACOS $343.60 / $1,440.00 = 23.9%. The blended row reads "fine" while containing one in-band scale, one window-expiry exit, and one ceiling exit. The v1.0 case recorded the same shape at a blended $26.90. This is the per-target rule made visible, and it is the argument attached to failure mode F1. Rotation: $42.00 + $205.00 = $247.00/wk releases; the next BEATABLE queue candidate opens per P3 with its own prediction.

WE-2 · The edge case that goes wrong and gets handled: a stale flag buys two weeks of full price. Illustrative values on the standard fictional base; ceiling $28.00.

Setup. T5 opened as a WINDOW class on the competitor's stockout. The reversal watch lapses when the competitor-depth pull fails for the family (T7 row D8 flips PARTIAL); the target's flag date goes stale at 9 days. The competitor restocked on day 3.

What it cost. Week A: 55 clicks at CPC $1.20 = $66.00; 1 order; CPA $66.00; ACOS $66.00 / $80.00 = 82.5%. Week B: 48 clicks at CPC $1.25 = $60.00; 1 order; CPA $60.00; ACOS 75.0%. Two-week spend $126.00 against contribution recovered 2 x $28.00 = $56.00: net -$70.00 bought against a weakness that no longer existed. The account pattern this reproduces is v1.0's recorded quiet leak: paying full price against a recovered competitor.

Detection and handling. The P4 read runs its gate check first: the flag date (9 days) sits outside the read window and the WINDOW review date (2 weeks per S-C1) has arrived unserviced, which is the F4 signal. The read executes tree branch 1 immediately: exit at this review; the flag-state check needs no fresh feed to see that the time-box expired. Register entry WINDOW-CLOSED with the lapse noted; T6 scores the open's prediction WRONG-DIRECTION with the -$70.00 attached; the corrective rides the cadence: the reversal watch is daily on every open target, the flag-date column is mandatory on the read sheet, and a WINDOW open past its review date reads as expired at the next pass no matter what the feed says. No negative-ASIN: the page never failed; the process did.

WE-3 · Category refinement refresh after our reprice, then the S-C6 refine-or-exit read. Illustrative values on the standard fictional base; post-reprice AOV $72.00; ceiling $28.00; band 10-15%.

The refresh. Our price moves $80.00 to $72.00 (T5 price-move event). The standing category campaign's refinements read price floor $80.01+, rating ceiling 4.4 (ours 4.6). Stale bands are failure mode F6: competitors priced $72.01-$80.00 and rated at or below 4.4 are now weak on both S-C1 axes and invisible to the campaign. P7 executes the same week: price floor moves to $72.01; rating ceiling holds at one step below our unchanged 4.6; the band record logs both values, the T5 source, and the refresh date; the P1 category scan re-runs.

Read 1, post-refresh (validated). 240 clicks at CPC $0.95 = $228.00; 12 orders (CVR 5.0%); sales 12 x $72.00 = $864.00; class CPA $19.00 (inside the $28.00 ceiling); class ACOS $228.00 / $864.00 = 26.4%, above the 10-15% band. S-C6 executes, pass 1: tighten one step; the price floor steps to $75.00 so the class serves only against a wider price disadvantage.

Read 2 (validated). 170 clicks at CPC $0.70 = $119.00; 12 orders (CVR 7.1%); sales $864.00; class CPA $9.92; class ACOS $119.00 / $864.00 = 13.8%, inside the band. The class holds with 1 refinement pass spent; a second failed pass would have exited the campaign per S-C6. The lesson: the refinement is the flag test, and tightening it is the class-level equivalent of re-running F12 on every target at once.

WE-4 · Virtual-bundle defense: the basket read versus ACOS optics, and a knocked-out week. Illustrative values on the standard fictional base.

Setup. VB self-targeting live per P9. Pre-standup 4-week baseline on the family: 1.18 units per order, AOV $80.00. Post-standup validated 4-week read: 1.31 units per order (+0.13), AOV $88.80 (+$8.80 per order). Family orders 92/wk.

The two reads. Campaign-row optics: spend $84.00/wk (70 clicks at CPC $1.20); campaign-attributed orders 4 at $88.80 = $355.20; row ACOS $84.00 / $355.20 = 23.6%, far above the 10-15% defensive band, because VB conversions land heavily on component listings the row does not credit. A band-only read kills the campaign; that is failure mode F5. The recorded grading standard reads the basket instead: AOV lift $8.80 x 92 orders/wk = $809.60/wk of gross basket value against $84.00/wk of spend, a 9.6-to-1 gross ratio, with the honest caveat written on the read: this is a lift-versus-baseline comparison, not click attribution, and v4.0 carries no numeric lift bar for it (amendment item A-3). The inverse guard also holds: a VB row with flat units per order does not survive on a pretty row ACOS.

The knocked-out week. Week 5, a component price change breaks the bundle; the weekly VB status check flags KNOCKED-OUT. P9 executes the same week: the campaign goes PAUSED-PENDING (reason VB-KNOCKED-OUT, re-entry VB-RESTORED); the baseline clock annotates weeks 5-6 so no lift read spans the gap; the campaign resumes on the restore confirmation.

7. Failure Modes and Escalation

Verdict first: eleven modes, each with its detection signal and response; where an account-history pattern motivated the mode, the pattern is named. The v1.0 register recorded its six modes at pattern level without dated incidents; those six patterns (F1-F6 below) are the carried institutional knowledge, and no dates are invented for them.

Failure	Detection signal	Response and the recorded pattern
F1. Blended conquest read	A verdict or a staged change citing whole-campaign or whole-class conquest ACOS outside the P7 refined class	Re-issue per target; WE-1 is the argument (the blend reads $19.09 CPA while hiding a ceiling exit). Carried v1.0 pattern: the blended number that "said fine" over three different decisions.
F2. Unflagged open (ambition conquest)	A live target whose profile does not satisfy S-C1 (wins <2 of 3, no OOS), or has no profile	Exit or classify it honestly; the queue has candidates waiting and P1 sourcing, not the bar, is the fix. Carried v1.0 pattern, tightened: v1.0 policed "no flag"; v2.0 polices "no S-C1 basis."
F3. Failed target re-bought	A CEILING-register ASIN serving again in the substitutes auto's served terms	The SOP-29 wall gap: negative-ASIN staged now; the SOP-29 wall audit adds the campaign with session attribution. Carried v1.0 pattern: the scout re-buying a proven-bad page.

F4. Flag reversal unnoticed	Read rows with flag dates outside the read window; a WINDOW target past its 2-week review date unserviced	The flag check rides every read and the watch is daily on open targets; expired time-boxes read as expired without waiting for the feed. WE-2 quantifies one lapse at -$70.00 over two weeks. Carried v1.0 pattern: "paying full price against a recovered competitor is the quiet leak of conquest."
F5. Defense graded on ACOS optics	VB or complement rows judged on the campaign row's ACOS versus the 10-15% band	The basket metrics govern (units per order and AOV lift vs baseline); WE-4 shows a 23.6% row ACOS on a 9.6-to-1 gross basket ratio. Carried v1.0 pattern: the tactic's product is basket size.
F6. Stale refinements	Category band record unmoved after a T5 price-move event, or past the quarterly slot	P7 refresh now; WE-3 shows the stranded weak-competitor band a stale floor creates. Carried v1.0 pattern: refinements aimed at last quarter's price position.
F7. Open without prediction or horizon	A deployed conquest target with no T6 prediction row or no O7 horizon date	Rule P10: the record is incomplete and should never have deployed; write both now, date the horizon from the open date, and name the session. Per O7, objectives do not persist by inertia.
F8. Sourcing starvation misread as a standards problem	Zero queue candidates passing S-C1 while conquest budget sits idle for a cycle	Escalate as a sourcing gap: P1 sources get attention (substitutes flow, tracker coverage, category scan), never a lowered bar. Carried from the v1.0 escalation clause into a numbered mode.
F9. Conquest into a deal window	Spend on a target whose deal-state field reads ON, in or out of our own event windows	Stand down per S-C5 same day and stage the resume date; S-C5's own text is the reason: buying the worst comparison of the month.
F10. Counter-conquest without a basis	An open traced to a SOP-26 P4 escalation whose evaluation record shows NOT-ENTERABLE or is missing	Exit; the answer to an incursion we cannot out-compare stays defensive, and the evaluation record is written retroactively with the decline counterfactual. Carried v1.0 line: conquest against a stronger competitor is a donation.
F11. Class-blend leakage	A blended read applied to any per-target ad group by analogy to the P7 category exception	Re-issue per target; the exception is licensed by the refinement mechanics and by nothing else. The exception's existence is the reason this mode exists.

Escalation. F3 routes through the SOP-29 wall audit with session attribution. F8 escalates to the PPC Lead as a sourcing gap the cycle it is found. The same mode twice in consecutive cycles on the same scope escalates to the PPC Lead with the T6 history attached. Anything in the human-mandatory tier (every OPEN, every counter-conquest decision, every refinement band change, every pause) is confirmed per the Authority Matrix before deployment, without exception.

8. Outputs and Artifacts

Verdict first: gates clear with artifacts, not assertions. The weekly gate artifact is the per-target read sheet; an open's gate artifact is its flag profile. Retention policy lines are a pending #39 fold-in per the Ledger row for component 39; until that lands, the register and read sheets are kept without pruning because the P6 re-open rule needs the history.

Artifact	Named record and storage	Downstream consumer
Candidate queue with flag profiles	The conquest queue block beside T3 in the product Command Workbook; one row per candidate; profiles attached with per-flag source and date	P3 opens; the SOP-17/18/19 playbook P7.2 (SD contextual target list); the P10 evaluation
Open records	T6 Decision Log rows: prediction, O7 horizon, review date, scenario ID S-C1	V5 scoring next cycle; the Verification scoreboard (Account Rollup A3)
Per-target read sheet	Weekly, one row per open target plus the class, defense, and VB rows; the REPORTING-ONLY blend labeled; stored with the cycle's KDR snapshot	The Monday cascade and the BM huddle (R); the playbook P7.2; the SOP-26 P2 defense-cost read (defense rows)
Rotation and exit packets	Staged #38 sets with scenario IDs; paired T6 rows (scored exit, new open)	#38 deployment and T+1 verification; V5

Failed-target register	The register block beside the queue: ASIN, exit reason code, flag state at exit, weeks open, spend and orders, date; kept without pruning	SOP-29 P1 pre-load and P3.4 (negative-ASIN in substitutes, same week); P1 step 6 de-duplication; audits
Refinement band record	Price floor, rating ceiling, source reads, refresh date, S-C6 pass count; stored with the category campaign's T3 row	P7 refreshes; the monthly structure review (M6 context)
Defense and VB basket reads	Weekly occupancy and floor-integrity lines; VB units-per-order and AOV lift vs the annotated baseline	SOP-26 P2 (annual defense-cost read); P9 pause decisions
Counter-conquest evaluation records	Profile, F12 class, decision, counterfactual (rule J5)	SOP-26 P4 return path; the T6 log

9. Skill-Conversion Block

Verdict first: the AI layer drafts everything below its checkpoint line and decides nothing above it; the split follows the Authority Matrix row "PAT/conquest opens and exits" (Incoming operator D and E; PPC Lead R; BM owner I) and the P5 human-review tiers.

Block	Specification
Input bindings	The SOP-15 P3.3 handoff list; the Keepa competitor pull of the Raw Pull Pack into T2 column DR "Competitor Depth (D8)"; T7 rows "D8 Competitor depth," "A8 PAT rows first-class," "B8 Row economics" with their Status cells; T0 named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC; T2 columns CO "CPA Ceiling $," CM "Expected CPC $," AD "Inventory Band," V "Clicks 7/30/60d," AI "Event Mode"; the targeting join feeding the campaign child rows (Data Architecture Section 6.1); T3 columns D, H, O and the PAT type filter; the failed-target register; the VB status check; T5 event rows (price move, competitor entry, deal window); the do-not-target exclusion list. An absent feed suppresses its dependent execution class per the manifest; the skill never improvises around a missing feed.
Deterministic logic	v4.0 Section 3 family F12 classification (BEATABLE / WINDOW); Section 8.7 rows S-C1 to S-C6 exactly as the P4 tree orders them; Section 6 rules O2, O6, O7; gates G1, G2, G4, G6, G8, G12 with the G9 carry; Section 9 rule L2 limits on every bid step; the P6 register re-open rule (new flag class); the P7 band-construction mechanics; rules P10, V1, J5 on record completeness. No local thresholds; the skill version-pins v4.0 and re-validates the Section 5 mapping on any increment.
Human checkpoints	AI-executable with audit: queue maintenance and ranking (P1), flag profiles and F12 classification (P2), per-target read sheets and tree dispositions drafted (P4), reversal and staleness detection (F4 watch), register maintenance and the SOP-29 handoff staging (P6), refinement band drafts (P7), defense occupancy and basket reads (P8-P9), evaluation record drafts (P10). Human-mandatory, per the matrix row (Incoming operator decides and executes, PPC Lead reviews): every OPEN including WINDOW opens, every counter-conquest decision, every refinement band change, every pause, every deployment confirm.
Cadence	Daily: the reversal watch on every open target; WINDOW review-date checks; event-window deal detection per G7 when armed. Weekly: P1-P2 intake, the P4 reads inside the Monday cascade, P8-P9 defense and VB checks, deployment Tuesday per #38 with T+1 verification. Quarterly: P7 refresh on the V4 slot plus same-week on any T5 price move. Event-driven: P10 on SOP-26 P4 escalations; P12 at G12 arm.
Outputs	Drafted profiles and queue ranks; drafted read sheets with dispositions and scenario IDs; rotation and exit packets with paired T6 rows; refreshed band records; the SOP-29 handoff set; F1-F11 alarms, each carrying the detection signal, an evidence line (row, values, gate or scenario reference), and a routing tier: F4 and F9 same day to the operator queue; F1, F2, F6, F7, F11 into the weekly pass worklist; F3, F5, F8, F10 to the PPC Lead with the T6 history.

10. Version Governance

Verdict first: owner PPC Lead; review monthly and immediately on any v4.0 increment, at which point the Section 5 mapping and the P4 tree re-validate against Section 8.7 and the vocabulary before any pass runs on them. Component-specific triggers, each an increment on occurrence: (1) the competitor depth feed lands (T7 row D8 flips PRESENT; Ledger dependency D4 clears): P1-P2 coverage notes retire and this Section records the landing; (2) a new confirmed weakness class is admitted to v4.0: the flag-class list is living, and a class this SOP cannot classify under F12 increments both documents together; (3) the

SOP-17/18/19 playbook's SD contextual opens at structural depth on the same D4 landing: the shared flag-workflow interface (row I-5 below) re-verifies; (4) any Authority Matrix re-issue touching the "PAT/conquest opens and exits" row: Element 9 and Element 11 re-validate within the matrix's one-week re-issue window; (5) the VB program changes standup or knockout mechanics at the catalog owner: P9 re-validates; (6) amendment items A-1 to A-6 landing in a v4.1 release: the corresponding clauses replace their amendment references. Registered in the Supersession Map; supersedes v1.0 in full; absorbs the scheduled PAT/conquest playbook deliverable. v2.0 change note: procedures 7 to 12; verdict mapping rebuilt on quoted vocabulary strings with the v1.0 non-vocabulary strings retired; worked examples 1 to 4; failure modes 6 to 11; the open test corrected from any-single-flag to the S-C1 condition; Elements 11 and 12 added from primary sources per the program's conflict C5 resolution.

11. Operator Ramp

Verdict first: the ramp turns the Authority Matrix row "PAT/conquest opens and exits" (Incoming operator D, E; PPC Lead R; BM owner I) into demonstrated competence on a 1-2-4 week schedule; the checkpoint holder throughout is the PPC Lead, whose R on this row is the review seat, and the matrix's standing rule 5 (vacancy: D reverts to the PPC Lead, E to the PPC Manager) covers any gap. The Economics Worked-Example Set's governance clause applies: the incoming operator reproduces its computations cold before taking a scope; for this component the load-bearing cases are E2 (clicks to loss) and E4 (margin-dollar classes), plus this SOP's WE-1.

Week	Demonstrations, each observed by the PPC Lead (the matrix R seat)
Week 1: observation and artifact familiarity	Shadow one full Monday cascade through the conquest rows. Locate and read every artifact cold: the queue and a flag profile, the per-target read sheet, the failed-target register, the refinement band record, T7 rows D8 / A8 / B8, T2 columns DR / CO / CM / AD, the T6 conquest rows. Reproduce WE-1's arithmetic from the inputs without the answer sheet (all nine computations: three CPAs, three ACOS reads, the blend, the released spend, the S-C2 clock). State, unprompted, why the blend row is labeled REPORTING-ONLY.
Week 2: supervised execution	Build 3 flag profiles from live feed rows with correct F12 classes and coverage notes; draft one full P4 read sheet with tree dispositions and scenario IDs; draft one OPEN packet end to end (S-C1 basis, gates, structure, Expected CPC seed, prediction, O7 horizon) and one exit packet with the SOP-29 handoff staged. The PPC Lead reviews every item before anything deploys; every correction is written into the operator's copy of the tree.
Week 4: unsupervised execution under audit	Run the conquest rows of the cascade alone: intake, reads, dispositions, staging, T6 logging, the huddle presentation to the BM owner (I/R seats). The PPC Lead audits the week's artifacts after deployment: every disposition traces to a tree branch and scenario ID, every open carries its prediction and horizon, the register and handoff reconcile, zero F1/F2/F7 findings. Two consecutive clean audited weeks close the ramp; the closure record feeds #44 operator certification when that component lands (due 2026-09-25 per its Ledger row).

Live-defense questions (the PPC Lead asks these cold; the test is the reasoning behind the rule, not recall of its value): (1) A target's CPA is $10.50, well under ceiling, and its stockout just ended: what do you do and why is the passing CPA irrelevant? (2) Why is blended conquest ACOS banned on per-target groups yet legal on the refined category class, and what exactly licenses the exception? (3) Why does the CPA ceiling equal contribution dollars per order rather than a percentage, and what breaks if a percentage is copied across products (defend with the E4 case)? (4) A target failed on the price class in March; its rating collapses in July: may it open, and what must the queue row name? (5) SOP-26 escalates an incursor whose profile wins 1 of 3 axes: what is the answer and why is the alternative a donation? (6) A competitor opens a deal during our own event window: which two rules fire and in what order? (7) The competitor-depth manifest row reads MISSING: what may you still do this cycle, and what may you not?

12. Interface Contracts

Verdict first: every handoff below names its sending owner, receiving owner, artifact, timing, and the consequence when it does not arrive, because the non-arrival clause is what the corpus has historically omitted. Sources: the Cross-Reference Ledger row for component 16 (interfaces cell), corrected on one point in this build's returned row (the ASIN-converter route is SOP-15 P3.3, as both SOP-15 and this SOP state; the seed row attributed it to SOP-30), cross-checked against the Master PPC Data Architecture v1.0 for artifact names and locations.

Row	Direction and parties	Artifact and timing	Non-arrival consequence

I-1	Inbound; SOP-15 (PPC Lead scope) to this SOP	The P3.3 PAT candidate handoff list with evidence lines; every weekly pass	The queue thins from its top source; conquest continues on tracker and category-scan candidates at the unchanged S-C1 bar; a starved queue with idle budget escalates F8 as a sourcing gap, never as a standards question.
I-2	Inbound; Data Operations (Raw Pull Pack Keepa pull) to T2 column DR and the profiles	Competitor depth fields (price, rating, reviews, stock, deal) with dates; weekly pull, daily reversal watch on open targets; manifest row T7 "D8"	T7 D8 reads MISSING and the blocked-classes cell governs: "Conquest open/rotate/exit" suppresses to flags per G8; open targets read on available sources with coverage noted; WINDOW time-boxes still expire on their dates (WE-2); no new opens deploy. Ledger dependency D4; its landing is a Section 10 trigger.
I-3	Inbound; SOP-26 P4 (PPC Lead) to this SOP	Persistent-incursor escalation with the response record; event-driven, evaluated inside the same weekly cycle	No counter-conquest evaluation occurs; defense continues per SOP-26 inside its ceilings; nothing in this SOP self-originates an offensive answer to an incursion.
I-4	Outbound; this SOP (incoming operator) to SOP-29 (PPC Lead)	Failed-target register entries for negative-ASIN in the substitutes auto; same week as every CEILING exit; SOP-29 Section 3 names the list as a required input	The scout re-buys the proven-bad page (F3); the SOP-29 monthly wall audit detects the gap and attributes the session; the negative stages on discovery and the register backfills.
I-5	Outbound; this SOP to the SOP-17/18/19 playbook P7.2	Target list, flag states, and the open, rotate, exit calls for SD competitor contextual; with the weekly read sheet	SD competitor contextual holds closed at structural depth (its own stated posture while dependency D4 stands); own-page SD defense is unaffected; the playbook restates zero flag rules, so no parallel flag logic may appear there in the interim.
I-6	Inbound; catalog owner (Brand Management) to P9	The weekly virtual-bundle knocked-out check	Unknown is not servable evidence: VB campaigns pause pending the check the same week; the baseline clock annotates the gap; resume on the first passing check.
I-7	Inbound; Finance and the Economics Pack (T0 named cells) to the ceiling columns	AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC feeding T2 CO and CM; weekly refresh, monthly true-up; manifest row T7 "B8"	Ceiling-referencing opens carry the G9 PROVISIONAL state and defer where the driving read is a ceiling comparison; existing targets hold at current bids; the stale state is written on every affected record per rule J7.
I-8	Outbound; this SOP to #38 (bulk mechanics)	Staged change sets with scenario IDs; Tuesday cadence, same day for dated signals (fresh WINDOW opens, S-C5 stand-downs, reversal exits); T+1 verification returns	An unverified deployment is not live: the affected week's reads on those rows are void, the delta redeploys same day per #38's mismatch handling, and a console edit found outside the staged set routes to change control per the SOP-12 F10 precedent.
I-9	Outbound; this SOP to the Monday cascade and the BM huddle	The per-target read sheet and the T6 rows; weekly; the BM owner holds R via the daily huddle per the Authority Matrix standing rules	An unreviewed week is a process finding: the huddle flags it, the PPC Lead (R on opens and exits) audits the deployed set retroactively, and V5 scoring still runs next cycle because unscored actions do not accumulate.
I-10	Inbound; #32 (deals owner with the PPC Lead) to P12	The armed-window declaration and event plan; at G12 arm	Without a declared plan, no conquest structure changes occur inside a suspected window beyond stand-downs and exits; S-C5 and S-E3 remain live on the deal-state flags themselves, which do not depend on the plan's arrival.

13. Closing: v4.1 Amendment Items Logged This Build

Per the anti-drift protocol, every number this component needs that v4.0 does not contain is logged here and is not written into the SOP. Items already on the program's amendment register by earlier builds (the 18/25/26 closed-vocabulary count, the stale v3.0 cover labels, the G9 gate-ID dual use) are referenced where relevant and not re-opened.

1. **A-1, PAT entry-bid rule.** v1.0 cited a "conservative entry bid per the v4.0 category CPC reference"; no category CPC reference exists in v4.0. v2.0 seeds opens at Expected CPC (family F9) inside the L2 limits. A PAT-specific entry-bid rule is owed to v4.1 if the program wants one.
2. **A-2, additional entry flag classes.** v1.0 admitted "off-deal while we run one" and "proven conversion from their page" as opening flags; S-C1 admits neither as an entry authority (conversion evidence feeds target lists; deal state governs S-C5 stand-downs). If either is to open spend, v4.1 must add it to the S-C1 condition set.

3. **A-3, VB basket-lift bar.** The virtual-bundle defense read (units-per-order and AOV lift vs the pre-standup baseline) carries no numeric threshold anywhere in v4.0; the metric standard is carried from v1.0 without a bar.
4. **A-4, flag-age limit.** No numeric staleness window exists for competitor flag data; this SOP requires flag dates inside the current read window and blocks opens on stale flags without naming a day count.
5. **A-5, reversal-watch SLA.** Element 2 states detection lags in corpus cadence vocabulary (same day, same pass); a clock-hour SLA for the daily reversal watch on open targets is owed, mirroring the Element 2 lag item already logged by the #17/18/19 build.
6. **A-6, do-not-target exclusion list.** The brand-safety and policy exclusion check (Section 3) closes the Ledger's pending fold-in for #16 structurally; the list's owner, criteria, and change control are owed, coupled to #42 ad policy (due 2026-08-17).

Companions: Master PPC Decision Criteria System v4.0 (all thresholds and verdicts: Sections 3, 5, 6, 8.7, 9, 13) · Keyword Decision Record Template v3.0 (the closed verdict vocabulary) · PPC Strategic Doctrine v1.0 Section 9 (the why) · SOP-15 P3.3 (the scout), SOP-29 (walls and the register handoff), SOP-26 P2 and P4 (defense posture and incursions), SOP-17/18/19 P7.2 (SD contextual), #32, #37, #38 (windows, rules, staging) · Master PPC Data Architecture v1.0 and the PPC Command Workbook TEMPLATE v1.0 (instrument mechanics) · PPC Authority Matrix v2.0 (the decision seats) · PPC Cross-Reference Ledger (the component 16 row, updated by this build).