

SOP-17/18/19: SB, SBV, and SD Playbook

Version 2.0 · 2026-07-28 · Owner: PPC Lead · Supersedes SOP-17/18/19 v1.0 (2026-07-24) in full · One playbook, three formats, per the Cross-Reference Ledger row for component #17, which folds components #18 and #19 into this document · Thresholds live in the Master PPC Decision Criteria System (the attached criteria authority, cited by section heading throughout) and verdict strings in the KDR Template v3.0 closed vocabulary; this SOP carries zero local thresholds

1. Verdict and Scope

VERDICT: the non-SP formats are 3 distinct products sold under 2 console labels, and this playbook operates all 3 against the M6 structure targets of SP near 80%, SB 15-20%, and SD 5-10% of format spend. SB buys headline shelf on terms SP has already proven; SBV buys the same shopper cheaper whenever the CVR-adjusted inequality holds, and it is the standard response the criteria authority names to SP CPC pressure (Section 8.10, S-F3); SD buys pages and audiences and is the cheapest own-page defense surface. All 3 are graded click-attributed on the same economics chain as SP, never on view-attributed console numbers. Version 2.0 expands v1.0 from 6 procedures to 13, from 7 failure modes to 12, from 1 worked example to 4, from a 6-row local verdict list to the full 25-string closed vocabulary with format execution stated per string, and adds Elements 11 and 12, which no shipped SOP in the 23-document library carried before this build.

1.1 Three components, one playbook

The Ledger covers components 18 and 19 inside the row for 17 because the 3 formats share 1 measurement doctrine (click attribution only), 1 creative precondition gate (no launch on a placeholder asset), 1 monthly governance check (the M6 format split), and 1 event overlay. Splitting them would force 3 documents to restate those 4 shared mechanisms, which anti-drift rule 3 forbids. Each format section can split into a standalone SOP later without renumbering, per the v1.0 governance note carried forward.

1.2 The three products

SB (component #17) runs 3 declared roles, each its own campaign class: brand-defense headline on own brand terms, category headline on terms already proven in SP exact, and store spotlight for family cross-sell graded on the basket read. SB never discovers: a discovery term found in an SB list is the F1 condition, because the SP discovery layers (owned by SOP-13, SOP-14, SOP-15) do the same job at a lower CPC with better negation control. **SBV (component #18)** is the arbitrage instrument: the entry and exit test is the inequality $SBV \text{ effective CPC } \times (SP \text{ CVR} / SBV \text{ CVR}) < SP \text{ exact CPC}$, with its doctrine basis in Strategic Doctrine v1.0 Section 7 and its numeric trigger in the criteria authority Section 8.10 row S-F3. **SD (component #19)** is 2 products under 1 label: contextual page targeting (own-page defense at floor, competitor pages under SOP-16 flag logic unchanged) and audiences (views and purchases remarketing, plus the rare event-tied in-market buy), with SD Performance+ permitted only under the P9 guardrail cycle.

1.3 Boundaries with adjacent components

Adjacent component	The seam, stated so it is owned
Criteria authority (component #8)	Decides every threshold, band, scenario, and verdict. This SOP executes; a threshold written here instead of there is a defect per SOP-12 v2.0 Section 1.3.
#16 PAT and conquest	Owens the competitor weakness flag logic. SD competitor contextual runs SOP-16's open, rotate, and exit workflow unchanged; this SOP never restates a flag rule.
#12 Exact operations	On ENTER SBV ARBITRAGE, SOP-12 executes only the SP-exact ladder step-down, dated with the SBV go-live; the SBV build lives here. Stated identically in SOP-12 v2.0 Section 1.3.
#29 Negation architecture	Owens walls and negation doctrine. This SOP builds the prescribed walls at every SB keyword event in the same session and records them per #29 rules.
#38 Bulk-file mechanics	Owens pull, staging format, upload, and T+1 verification. Every change this SOP decides deploys as a #38 staged set and consumes the #38 verification report.
#32 Deal and event windows	Owens the event plan, ramp, caps, and reconciliation. This SOP deploys tentpole headline positions only when the event plan names them (P13).
#26 Mature and defend	Owens the defense set. SB brand-defense campaigns are members of that set; membership and escalation belong to #26, mechanics belong here.
#20 AMC operations	Owens instance setup and audience creation into SD. This SOP receives AMC-built audiences as SD inputs and executes them under P8 rules.
Strategic Doctrine v1.0	Owens the why: the SBV arbitrage basis (Doctrine Section 7) and format economics. Numbers stay in the criteria authority.
#42 Ad policy, moderation, claims	NOT BUILT, due 2026-08-17, coupled to this component per the Ledger. Until it lands, moderation rejections route per Element 12 row I-9.
Creative owners (outside PPC)	Video, store, and headline assets are preconditions supplied outside PPC. PPC parks candidacies on asset absence and never launches on a placeholder.

2. Trigger Conditions

Each trigger names its detecting artifact and the maximum allowed lag from signal to procedure start. Lags anchor to the cadence classes the criteria authority defines in Section 14 rule V4 (daily alerts and gate checks, weekly full audit, monthly reviews, quarterly re-tests); a lag stated as a cadence slot means the first occurrence of that slot after the signal.

Trigger	Detecting artifact	Maximum lag	Procedure invoked
Weekly cadence	Monday audit cascade reaches the T3 CAMPAIGN BOARD brand-format rows (Command Workbook tab T3, groups C4 and C6)	Same audit session	P10 weekly format read; P1, P2, P3, P7, P8 row reads inside it
Arbitrage candidacy	T2 K7 test-status field plus the creative asset register showing video status APPROVED on a term with qualified SP exact history	Next weekly pass after APPROVED lands	P4 probe entry
Probe qualification	Probe campaign clicks reach the G2 sufficiency bar (criteria authority Section 5, "Layer 0: Gates (G1-G12)")	Next weekly pass after the bar clears	P5 inequality test
Arbitrage re-test date	The dated re-test field on the open arbitrage record (T6 chain)	Same week the date lands; past-date is F3	P6 re-test
Remarketing pool signal	SD audience size read in the weekly SD bulk pull (Raw Pull Pack, SB / SBV / SD bulk and RAS sheets per Data Architecture v1.0 Section 12)	Next weekly pass	P8 audience entry evaluation
Performance+ expansion	Performance+ expansion log in the weekly SD pull	Weekly, zero skips; a skipped week is F4	P9 expansion review
Monthly split date	First Monday audit of the calendar month; A4 Structure targets tab of the Account Rollup Workbook	That session	P11 M6 split check
Event window arming	#32 event plan entering T2 with dates; event posture field PRE-EVENT opens 14 days out per criteria authority Section 3, family F2	Ramp actions per the plan's own dates	P13 event overlay
Moderation rejection	Ads Console moderation notification on any SB, SBV, or SD campaign	Same day to the PPC Lead	Element 12 row I-9 routing (interim, pending #42)
Defense escalation	#26 incursion or defense-set instruction referencing a brand format	Same session it arrives	P1 posture change; P7 own-page step

3. Preconditions and Gates

Every procedure below runs only after these read current. Manifest IDs are the T7 MANIFEST field groups of the PPC Command Workbook TEMPLATE v1.0; the observable that proves each gate passed is stated in the row. A missing or stale group suppresses the dependent execution class and the pass records what it could not see, per the manifest rule (criteria authority Section 5, gate G8 inside "Layer 0: Gates (G1-G12)").

Precondition	Manifest / register ID	Observable that proves it passed
Deployment state current	T7 group A2 (SP bulk) plus the SB / SBV / SD bulk and RAS sheets	T7 status PRESENT dated this Monday; T3 rows for every live brand-format campaign carry current state, budget, and bid.
Pacing and in-budget	T7 group A3	In-budget % populated on the format rows; a truncated row takes the budget lever first per gate G3 before any rate read.
Threshold flags precomputed	T7 group A5	Band and flag columns computed in-row; neither the operator nor the skill re-derives a threshold.
Objective and Expected CPC	T7 group A6	Every brand-format campaign carries exactly 1 objective tag in T3 group C1; an untagged row routes to assignment before any read (criteria authority Section 6, rule O1).
Economics chain live	T7 group B8 (T0 named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC)	Ceilings compute in-row; a stale margin table carries the PROVISIONAL tag onto any ceiling-referencing action.
Decision log live	T7 group D7 (tab T6)	Last cycle's due predictions carry V5 scores before new format actions stage; unscored actions do not accumulate.
Creative assets approved	Creative asset register (supplied outside PPC)	SB: store and headline assets live; SBV: video status APPROVED; absence parks the candidacy, and a launch on a placeholder is the F8 condition.
Brand Registry current	Ads Console brand state	Registry active for every brand advertised; an eligibility flip halts the affected campaigns same day per gate G4.
Attribution labels on reads	This SOP's read-row format	Every SD and SBV read row carries its attribution label; a decision citing a view-

	(Element 8)	attributed number is the F2 condition. The click-only rule is this playbook's standing measurement gate, carried from v1.0.
Event state known	T2 event flag from the #32 plan	Event exclusion honored: event weeks and the 14 days following are excluded from trend verdicts per the criteria authority's event exclusion gate (Section 5, G10 in the attached file; the same freeze is cited as G12 by SOP-12 v2.0 and the Ledger; the numbering divergence is logged as conflict C9 and amendment item A-8).
Human tier known	Criteria authority Section 12, rule P5	Structural changes, gate failures, and every 500+ SV term action route human-mandatory; Element 9 lists the format-specific checkpoint rows.

4. Step-Level Procedures

P1. SB brand-defense headline operations (weekly, 20 minutes)

1. Open the Ads Console at Sponsored Brands, filter to campaigns whose T3 C1 objective tag reads Defensive; cross-check the list against the #26 defense-set membership record. A defense campaign absent from the #26 set routes to #26 the same day; membership is theirs.
2. Keyword lists are exact-only own brand terms. Verify against T3 group C5 that every hosted keyword carries the brand root; a generic term inside a defense campaign is removed in the staged set and, where it is a proven category term, re-routed to P2 candidacy.
3. Bid posture is floor: read the campaign's click-attributed ACOS against the Defensive band of 10-15% from the criteria authority's objective bands (Section 3, family F1, "Objective Max-ACOS bands"; restated in Section 6, rule O6). Inside band with position held: no change. Above 15% with position held: 1 ladder step down per the L2 mechanics owned by #28, staged per #38.
4. Position evidence: read the headline position share for the brand terms from the weekly SB report. A slipping share with in-band ACOS restores 1 step, mirroring the criteria authority's defensive restore logic (Section 8.6, S-D2). The share bands themselves do not exist in the attached criteria file and are logged as amendment item A-6; until they land, direction (rising, flat, slipping across 2 reads) is the read, never a level judgment.
5. Walls per #29: brand terms stay negated out of generic discovery so defense owns them; verify the wall state field in T3 C5 reads STEERING for each defense term.
5. Every staged change logs a T6 prediction with a review date; the pass opens by V5-scoring last cycle's due rows.

P2. SB category headline operations (weekly, 30 minutes)

1. Candidacy rule: a term enters a category SB list only with proven SP exact history, defined as the term carrying qualified CVR and CPC evidence at the G2 sufficiency bar (at least 100 clicks for the CVR read per Section 5, gate G2) and a Conversions or post-objective state in T2. SB never discovers; F1 fires on any unproven or non-exact entry.
2. Creation: campaign per syntax family, exact-only list, headline and logo assets from the register, landing per the asset plan. Name per the v4.0 Section 4.2 convention as restated in SOP-12 v2.0 Appendix A, with the SB type segment; a non-parsing name is the naming defect class and renames ride alone.
3. Walls per #29 pre-load in the same session: brand terms excluded from category SB; the graduated-term steering logic mirrored so category SB and SP exact never fight for the same query without declared roles.
4. Bid: seed against the SP benchmark CPC for the same term from T2 group K5, since the SB premium must buy visible headline position. Read weekly on click-attributed ACOS against the SB objective band; SB rows are graded on their own band plus NTB depth, vCTR, and completion per T3 group C4 (Data Architecture v1.0 Section 7, "never judged on SP metrics").
5. Exit test: NTB% and ACOS both below bar for 2 validated weeks executes exit or re-creative per the criteria authority Section 8.10 row S-F2; the campaign steps down or pauses in the staged set with the reason string carrying S-F2.
5. T6 prediction per change; T3 registry updated (type counts in T2 K8 pick up SB presence per term).

P3. SB store-spotlight operations (weekly, 15 minutes)

1. Spotlight campaigns drive to Store pages for family cross-sell. Grade on the basket read: units per order and store dwell from the SB report, click-attributed sales against the campaign's own band, never last-click ACOS alone (carried from v1.0 P1.1c).
2. Store-page readiness is a pending fold-in named on the Ledger row for #17; until that fold-in lands, the precondition is binary: the Store page for the advertised family exists and is current per the brand owner. A stale or missing page parks the spotlight the way a missing video parks an SBV candidacy.
3. Changes stage per #38 with T6 predictions; the spotlight class carries the lowest budget priority of the 3 SB roles and is the first trimmed in a P11 rebalance.

P4. SBV candidacy and probe entry (per candidacy, 25 minutes)

1. Candidacy requires all 3: a term with proven SP exact performance (qualified CVR and CPC history at G2 sufficiency), an APPROVED video asset in the register, and gate state clean (no event freeze on the product, economics chain fresh).
2. Screen with the criteria authority's arbitrage trigger before spending: Section 8.10 row S-F3 names the head-term condition "SP exact CPC >=1.5x ceiling" together with "SBV CPC <=0.6x SP exact CPC", and Section 9 rule L8 defines the substitution lever. At candidacy the SBV CPC is an estimate from the console's suggested bid; the definitive test waits for the probe.
3. Open the probe as a dated, budget-capped campaign: 1 term, exact, video creative, cap and stop-loss date staged in the creation set. SBV CVR is unknown at entry, so the probe exists to measure it; no inequality verdict issues on day-3 optics.

4. Register the SP and SBV instances as a declared cross-format pair in T3 C5 with roles stated, so the registry never reads the pair as duplicates (the pair is 1 term posture per L8).
5. T6 entry: predicted SBV CVR band, click quota, cap, stop-loss date, review date.

P5. SBV inequality test and arbitrage registration (per test, 20 minutes)

1. Test only at qualified clicks: the CVR read requires at least 100 clicks on the probe per Section 5, gate G2. Below the bar at the stop-loss date, the probe closes as unqualified and the candidacy re-parks.
2. Compute the CVR-adjusted SBV cost: $\text{SBV effective CPC} \times (\text{SP CVR} / \text{SBV CVR})$. Compare against the SP exact CPC on the same term. The basis is Strategic Doctrine v1.0 Section 7 (the arbitrage instrument row) and the Economics Worked-Example Set case E6; the arithmetic convention (adjust before comparing, never compare raw CPCs across formats) is E6's lesson verbatim.
3. Decision tree:

P5 decision tree (3 branches). Adjusted SBV cost < SP exact CPC **AND** the S-F3 trigger held at entry: the arbitrage is OPEN. The keyword row carries the closed-vocabulary verdict **ENTER SBV ARBITRAGE**; #12 executes the SP-exact ladder step-down dated with the SBV go-live; the arbitrage record opens with both inputs, the computed adjusted cost, and a dated re-test. | Adjusted cost >= SP CPC with qualified data: the arbitrage is CLOSED at entry; the probe steps to pause as **PAUSED-PENDING (REASON + RE-ENTRY)** with re-entry stated as a materially lower SBV CPC or higher SBV CVR read; no SP step-down occurs. | Data unqualified at stop-loss: probe closes, spend recorded, candidacy re-parks; re-entry requires a new asset or a changed market read.

4. The re-test interval on an OPEN record is a dated field. The attached criteria file contains no interval value; this is amendment item A-3, and until it lands the record carries the interval the PPC Lead sets at registration, with the account's recorded case (a 6-week re-test) as the working reference. The obligation to re-test is absolute regardless of interval, because the inequality is a market state, not a permanent fact (v1.0 P2.4, carried forward).
5. T6 entry with the full computation in the reason string; human-mandatory confirmation per Element 9 before the pair deploys.

P6. SBV re-test, step-down, and exit (per re-test, 15 minutes)

1. On the dated re-test, pull current SBV CPC and CVR at qualified volume and recompute the adjusted cost exactly as P5 step 2.
2. Inequality still open: the record's re-test log appends the passing test, a new re-test date sets, and nothing else changes.
3. Inequality reversed: the SBV bid steps down 1 ladder step per the L2 mechanics, or pauses where the reversal margin exceeds a step's reach; the record closes with both tests on file; the SP exact continues untouched (the step-down at entry is not reversed by this SOP; any SP restore is a #12 decision on its own evidence). Creative fatigue signals (completion rate falling) execute exit or refresh per Section 8.10 row S-F4 even while the CPC side of the inequality holds.
4. A record past its re-test date is F3 and re-tests the day it is found.

P7. SD contextual operations (weekly, 30 minutes)

1. Own-page defense: own-ASIN contextual campaigns run at floor bids permanently as the cheapest page defense; the read is presence (impressions on own pages) at the Defensive band of 10-15% ACOS. A competitor incursion signal arriving from #26 P4 escalates the floor 1 step the same session.
2. Competitor pages: SOP-16's flag workflow runs unchanged: open only under a weakness flag, per-target reads, rotate on flag reversal, register failed targets. This SOP restates zero flag rules; the target list, flags, and open, rotate, exit calls arrive from #16, and the keyword-row verdict is the closed-vocabulary string **OPEN / ROTATE / EXIT CONQUEST**. #16 is upstream-blocked on the competitor depth feed (Ledger dependency D4), so competitor contextual stays structural until that feed lands; own-page defense is unaffected.
3. Blended contextual ACOS is never read; every verdict is per target class, mirroring the per-target doctrine.
4. Grading is click-attributed CPA against the ceiling, where the CPA ceiling equals contribution margin dollars per the criteria authority Section 3, family F1 (the margin row: "Sets the CPA ceiling (= margin \$)").

P8. SD audience operations (weekly, 25 minutes)

1. Views remarketing is the highest-yield entry class; purchases-based audiences serve complements and repurchase timing, capped; in-market and lifestyle conquest audiences run rarely and only event-tied per the #32 plan (postures per Doctrine v1.0 Section 8; entry economics per S-F5).
2. Pool-size gate: an audience campaign enters only above the pool floor, because a campaign on a trickle pool spends nothing and reads nothing (F6). The floor value does not exist in the attached criteria file and is amendment item A-1; until it lands, entry requires the PPC Lead's human-mandatory sign-off on the pool size, and the account's recorded working case (a 4,100-member views pool admitted) is the calibration reference, not a threshold.
3. Lookback windows for views and purchases audiences are named in Doctrine v1.0 Section 8 as doctrine, not as thresholds; the decision-grade window values are amendment item A-2. AMC-built audiences (cart abandoners, high-intent viewers) arrive from #20 and enter under the same pool gate.
4. Grading: SD in CPC mode is judged on click-attributed CPA against the ceiling per Section 8.10 row S-F5; SD in vCPM mode is judged on NTB% and assisted conversions, never raw ACOS, per row S-F6, and mode-aware evaluation is mandatory. The console's view-attributed numbers are recorded in reporting with their label and discarded from decisions (F2).
5. Every scale step re-checks marginal ACOS: a step is legal only while marginal ACOS stays at or under 1.5x average per Section 4 rule X11 and Section 10 rule M7.

P9. SD Performance+ guardrail cycle (weekly, 20 minutes)

1. Opt-in only, and only for remarketing audiences: Performance+ is acceptable for remarketing with a strict return floor, never for open prospecting on a mid-size budget (Doctrine v1.0 Section 8). The numeric guardrails (budget cap sizing, the return floor value, an expansion tolerance) do not exist in the attached criteria file and are amendment item A-4; until they land, every Performance+ campaign carries an explicit dollar cap set at entry, human-mandatory.
2. Exclusion lists load at entry: own catalog and previously failed targets.
3. The weekly expansion review is the gate artifact: read the expansion log, list what the mode added, and check the week's click-attributed CPA against the ceiling. A week without the review artifact is F4; a second consecutive miss escalates, because unbounded expansion is spend without an owner (carried from v1.0 escalation).
4. Suspension: cap breach or CPA above ceiling on the week's click-attributed read suspends the campaign in the staged set; re-entry requires the reloaded exclusion list and a re-set cap.

P10. Weekly format read pass (inside the Monday cascade, 45 minutes)

1. Enter T3 only after T0 and T1 have been read; a format pass that skipped product posture and family caps is invalid by construction (Data Architecture v1.0 Section 2, rule 1).
2. Score first: V5-score every T6 format row whose review date lands this cycle.
3. Filter T3 to Type in SB, SBV, SD. For each row, gates before metrics: eligibility, truncation, event flag, manifest status of the feeding groups.
4. Read each row against its own objective band only (T3 group C4), with attribution label verified on every SD and SBV number before it is believed.
5. Route by the P10 tree:

P10 routing tree (5 branches). Gate failing: execute the gate lever (budget on truncation, halt on eligibility) and defer the metric verdict, exactly as the SP pass does. | SB row failing S-F2 (NTB% and ACOS below bar 2 validated weeks): exit or re-creative per P2 step 5. | SBV row at or past re-test date: P6 now. | SD row citing a view-attributed number anywhere in its read: F2, re-issue click-attributed before any verdict. | Row clean and in band: no change; the WoW block reads next cycle.

3. Stage every resulting change per #38 with the verdict or scenario ID in the reason string; deploy Tuesday; consume the T+1 verification report and route mismatches to F10 same day.

P11. Monthly format-split check (first Monday of the month, 60 minutes)

1. Compute the month's spend shares for SP, SB, and SD from the A4 Structure targets tab of the Account Rollup Workbook, against the M6 targets of SP near 80%, SB 15-20%, SD 5-10% (criteria authority Section 10, rule M6; the same splits appear in Section 3 family F8).
2. Deviations are monthly findings to explain, never errors to auto-correct (M6's own clause). The response tree:

P11 drift tree (3 branches). All shares inside bands: record the check, no rebalance stages. | A format grew past its band top: that format re-justifies with its marginal yield read, meaning what the marginal dollar of that format bought in click-attributed contribution against the marginal SP dollar; a losing read stages a rebalance at keyword and campaign level (trims by verdict, spotlight class first per P3 step 3), never a blanket budget shift. | A format fell below its band floor: the gap is a finding with a named cause (asset absence, closed arbitrage, paused audiences) and a dated fill-or-explain owner; a silent shortfall is F7.

3. Split changes are deliberate staged deployments with T6 predictions; F7 is the split drifting on automation defaults with no rebalance decision on file.

P12. NTB reading rule (per read, 10 minutes)

1. The January 2026 attribution tightening re-baselined new-to-brand rates account-wide (Doctrine v1.0 Section 7 measurement notes). NTB comparisons run only within post-change windows; a trend drawn across the boundary manufactures a fake decline and is F5.
2. NTB is a strategy read (brand growth composition) and an S-F1, S-F2, and S-F6 grading input; it is never a bid input except where a verdict names it explicitly (carried from v1.0 P5.2).

P13. Event overlay (per the #32 plan)

1. Tentpole headline positions (SB and SBV) deploy only per the #32 event plan with the plan's caps; a headline deployment inside an armed window with no plan reference is the format-side twin of SOP-12's frozen-structure rule.
2. SD remarketing pools grow during windows; the post-window remarketing pass is the 1 deliberate settle-period action, named in the event plan itself, because it reads the window's pool while it is warm; this is the plan working, not an exclusion breach (carried from v1.0 P6.2).
3. Event weeks and the 14 days following are excluded from trend verdicts per the event exclusion gate; format reads inside the window run against the event projection only.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance and thresholds live in the criteria authority. Verdict strings below are quoted verbatim from the closed vocabulary in the Keyword Decision Record Template v3.0, sheet "Bands & Verdicts Reference", block "CLOSED VERDICT VOCABULARY (Framework Section 17)": 25 strings, exactly 1 per keyword record. Campaign-level moves ride the T3 campaign verdict subset (FUND / HOLD / TRIM, RAISE / CUT MULTIPLIER, NEGATION PASS, CONSOLIDATE, RETAG, PAUSE, ARCHIVE) per Data Architecture v1.0 Section 7. Record states used by this SOP's artifacts (an arbitrage record reading OPEN or CLOSED, a Performance+ campaign reading SUSPENDED) are artifact fields, not verdicts; the local verdict rows v1.0 Section 5 carried are retired by this rule.

Verdict (closed vocabulary, verbatim)	Brand-format execution
QUARANTINE: DATA REFRESH	No format changes deploy on the row; the refresh task carries the staged set to the next clean cycle.
UNLOCK TRAFFIC	SP-side verdict; no brand-format execution. A starved SBV probe takes the budget lever inside its cap, never a bid unlock.
FIX BUDGET (TRUNCATED)	Budget lever only on the truncated format campaign; rate reads resume the next untruncated day.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	An SBV probe below sufficiency continues to its click quota at the capped bid; no inequality verdict issues early (P5 step 1).
FUND TEST (UNFUNDED PRIMARY)	SP-side creation verdict owned by #12; if the funding note names a format probe, P4 executes it capped and dated.
RANK PUSH: FUND SIZED PUSH	Owned by #12 and #23. Formats never carry pushes; a push-tagged SB or SBV row is a tagging defect routed to retag.
RECOVERY PUSH	Owned by #12 and #24; no format execution.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	SP-side; no format execution.
FIX CONVERSION FIRST; DEFER PUSH	Formats on the affected term hold at current caps; no SB category additions for the term until the fix re-entry criterion clears.
PLACEMENT FIX FIRST	SP placement program owned by #28; no format execution.
RE-POINT VARIATION	Format campaigns advertising the affected SKU re-point in the same staged set #12 deploys, same day on inventory RED.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	Owned by #12 P4; no format execution. Post-taper, the term becomes a P2 category-SB candidate on its proven history.
HOLD; DEFEND LIGHTLY AT RANK	No format change; the cheapest-surface doctrine (S-F8) makes SBV the first instrument if presence later needs rebuilding.
DEFEND AT FLOOR	P1 and P7 floor postures execute it on defense SB and own-page SD; drift from floor restores in 1 staged step.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	Applies to a format campaign's bid where the row names it; the syntax-cap half executes at T1 per #31's instruments.
INCREMENTALITY LADDER	Measurement design owned by #35; this SOP deploys only the dated format-side spend steps the program names.
ADD EXACT + RESTRUCTURE	Owned by #12; no format execution until the exact is live and proven, after which P2 candidacy may follow.
CONSOLIDATE DUPLICATES	A same-format duplicate (2 SB campaigns on 1 family role, 2 SBV instances on 1 term) consolidates to the winner; the declared cross-format pair of P4 step 4 is never a duplicate.
GRADUATE FROM DISCOVERY	Owned by #30 into #12; no format execution. A graduated term earns P2 candidacy only after G2-qualified exact history exists.
NEGATE (PRE-LOAD / REACTIVE / STEERING)	Owned by #29; this SOP builds the prescribed SB walls same session (P1 step 5, P2 step 3) and records them.
ENTER SBV ARBITRAGE	P5 in full: pair registration, arbitrage record with dated re-test; #12 executes only the SP-exact ladder step-down, dated with the SBV go-live.
OPEN / ROTATE / EXIT CONQUEST	P7 step 2 executes it on SD competitor contextual targets under #16's flag logic unchanged.
STOP-LOSS: CONCEDE OR DEFER	An SBV probe at its stop-loss date without qualification closes per the P5 tree; freed budget returns to the envelope per #31's instruments, never inside this SOP.
PAUSED-PENDING (REASON + RE-ENTRY)	Executed on closed-at-entry probes, suspended Performance+ campaigns, and parked spotlights; every pause carries reason and re-entry criterion or it returns to the pass.
LEAVE TO HARVEST	No format instance is stood up; the term stays on discovery economics under #30, and SB never hosts it (F1).

6. Worked Examples

All 4 examples are illustrative and internally consistent, built on the standard fictional base because no live product economics spine is attached to this session: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Maximum Acceptable 26.3%. WE-1 and WE-2 carry forward the account cases recorded in v1.0 Section 6 as the institutional reference numbers.

WE-1 · SBV arbitrage lifecycle: entry, open, reversal, exit (clean case, carried from v1.0)

Setup. SP exact on the term: CPC \$2.10 at SP CVR 5.0%. Max Profitable CPC = $28.00 \times 0.050 = \$1.40$. The S-F3 trigger reads: SP CPC ratio = $2.10 / 1.40 = 1.50$, meeting the " $\geq 1.5x$ ceiling" condition at equality; console-suggested SBV CPC \$0.95 against the 0.6x bound of $0.6 \times 2.10 = \$1.26$, and $0.95 \leq 1.26$ holds. Video asset APPROVED. P4 opens the capped probe.

Qualified read (probe past 100 clicks). SBV CPC \$0.95, SBV CVR 3.2%. Adjusted SBV cost = $0.95 \times (5.0 / 3.2) = \1.48 per SP-equivalent

conversion, against the SP cost of \$2.10: a discount of $1 - (1.48 / 2.10) = 29.3\%$. The inequality is OPEN. The row carries **ENTER SBV ARBITRAGE**; the pair registers with declared roles; #12 steps the SP exact down 1 dated ladder step; the record opens with a 6-week re-test date (the working interval pending amendment A-3).

Week-6 re-test (P6). Video auction inflation: SBV CPC \$1.40 at CVR 3.0%. Adjusted cost = $1.40 \times (5.0 / 3.0) = \2.33 against \$2.10: REVERSED. The SBV bid steps down per L2 and the record closes with both tests on file; the SP exact continues untouched. Net teaching: the arbitrage bought roughly 6 weeks of the same shopper at a 29.3% discount and exited on arithmetic, not on sentiment.

WE-2 · SD audience read and the attribution trap (edge case, carried from v1.0 and extended)

Setup. Views remarketing pool at 4,100 members, admitted under the P8 human-mandatory pool sign-off (floor value pending amendment A-1). CPC mode, so the S-F5 grading applies: click-attributed CPA against the \$28.00 ceiling.

The trap. The console shows view-attributed CPA \$9.12. The week's click-attributed line: spend \$178.20, 9 orders, CPA = $178.20 / 9 = \$19.80$. An operator drafts a scale step citing the \$9.12 number: that is F2, and the read re-issues on \$19.80 before any verdict.

The correct read. \$19.80 against the \$28.00 ceiling holds; the scale step then passes the X11 marginal check: prior validated period spend \$118.80 on sales \$486.00, this period \$178.20 on \$648.00, marginal ACOS = $(178.20 - 118.80) / (648.00 - 486.00) = 59.40 / 162.00 = 36.7\%$, average ACOS = $178.20 / 648.00 = 27.5\%$, and $36.7\% \leq 1.5 \times 27.5\% = 41.3\%$, so 1 budget step is legal. The view-attributed \$9.12 is recorded in reporting with its label and discarded from the decision.

WE-3 · Performance+ expansion week gone wrong (edge case that gets handled)

Setup. Performance+ remarketing campaign, entry cap \$150.00 per week, exclusion lists loaded, weekly expansion review required (P9).

The miss. Week 1's expansion review is skipped (F4 fires on the missing artifact). Week 2's review finds the mode expanded its audience and spent \$212.00 against the \$150.00 cap, a breach of $212.00 / 150.00 - 1 = 41.3\%$; the week's click-attributed CPA rose from \$22.40 to \$31.50, which is $31.50 / 28.00 - 1 = 12.5\%$ above the \$28.00 ceiling.

The handling. Cap breach plus CPA above ceiling suspends the campaign in the staged set (P9 step 4); the record reads SUSPENDED with re-entry requiring the reloaded exclusion list and a re-set cap; the F4 escalation clock notes 1 missed review, and a second consecutive miss would escalate to the PPC Lead. The T6 root-cause line names the skipped review as the cause, not the mode: the guardrail existed and was not run.

WE-4 · Monthly split check with a rebalance (P11, both tree branches shown)

Prior month (in band, the carried v1.0 case). SP 78.4%, SB 16.2%, SD 5.4%: all inside the M6 bands; the check records and nothing stages.

This month (drift). Total format spend \$24,600.00: SP \$18,257.00, SB \$5,314.00, SD \$1,029.00. Shares: SP = $18,257 / 24,600 = 74.2\%$, SB = $5,314 / 24,600 = 21.6\%$, SD = $1,029 / 24,600 = 4.2\%$. SB sits above its 20% band top and SD below its 5% floor.

The marginal yield read. The last \$1,000.00 of SB spend bought \$310.00 of click-attributed contribution; the marginal SP dollar bought \$460.00 per \$1,000.00. SB does not re-justify its growth. Rebalance staged at campaign level: \$1,400.00 per month trims from the SB spotlight class first (P3 step 3), moving SB to $(5,314 - 1,400) / 24,600 = 15.9\%$ and SP to $(18,257 + 1,400) / 24,600 = 79.9\%$, both in band. The SD shortfall logs as a finding with its named cause (2 arbitrages closed in the period reduced SBV-adjacent SD retargeting) and a dated fill-or-explain owner; shares $79.9 + 15.9 + 4.2 = 100.0$. Both moves carry T6 predictions with next-month review dates.

7. Failure Modes and Escalation

Failure	Detection signal	Response, and the recorded account case where one exists
F1. SB discovering	Broad match or unproven terms found in an SB list at the P10 pass	Remove in the staged set; SB rides SP-proven winners only, because the discovery layers exist and are cheaper (v1.0 F1, carried).
F2. View-attributed decision	An SD or SBV verdict citing a view-attributed number	Re-issue click-attributed before any deploy; recurring F2 escalates with attribution of the drafting source. Recorded case: the \$9.12 view CPA against the \$19.80 click CPA in WE-2 (v1.0 Section 6, carried).
F3. Arbitrage never re-tested	An OPEN record past its re-test date at the P10 pass	Re-test the day it is found; the inequality is a market state. Recorded case: the week-6 reversal in WE-1, where a skipped re-test would have paid \$2.33 for a \$2.10 shopper indefinitely (v1.0 F3, carried).
F4. Performance+ expansion unreviewed	A week with no expansion review artifact	Review now; suspend if expansion ran unbounded; second consecutive miss escalates because unbounded expansion is spend without an owner (v1.0 F4 and escalation, carried).
F5. NTB trended across the boundary	A read comparing pre- and post-January-2026 NTB	Re-baseline inside post-change windows; the boundary rule exists for this (v1.0 F5, carried).
F6. Trickle-pool remarketing	An audience campaign running under the pool gate	Pause until the pool clears the gate; nothing reads on a trickle (v1.0 F6, carried).
F7. Split drift by default	P11 finds drift with no rebalance decision on file	P11 runs in full; the split is a decision, not a residue (v1.0 F7, carried).
F8. Launch on a placeholder asset	A live SB or SBV campaign whose register asset reads anything but APPROVED and live	Pause same day; the creative precondition is a gate, not a preference; root cause names the session that skipped the register check.

F9. Undeclared cross-format pair	Registry read shows SP and SBV live on 1 term with no declared pair roles in T3 C5	Declare the pair or consolidate; an undeclared pair reads as duplicates and corrupts the term's spend-share view.
F10. Landed values differ from staged	#38 T+1 verification mismatch on a format row	Redeploy the delta same day; the week's reads on affected rows are void; console edits outside the staged set route to change control per #37.
F11. Rebalance as blanket shift	A P11 rebalance deployed as campaign-level budget haircuts with no keyword-level verdicts in the reason strings	Revert and re-stage by verdict; structure targets rebalance at keyword level, never by blanket budget shifts (Doctrine v1.0 Section 12).
F12. Tentpole outside the plan	A headline deployment dated inside an armed window with no event-plan reference	Revert unless the plan names it; the deviation escalates per #32; exclusions extend from the revert date.

Escalation. The same failure mode in 2 consecutive cycles on the same scope escalates to the PPC Lead with the T6 history attached. Everything in the Element 9 human-mandatory list is confirmed before deployment, without exception.

8. Outputs and Artifacts

Event	Artifact set produced, storage, retention, and downstream consumer
Weekly pass (P10)	Updated T3 brand-format rows; per-format read rows with attribution labels; staged #38 change set with verdict or scenario IDs; the T+1 verification report consumed back; T6 entries each carrying a prediction and review date; V5 scores for last cycle's due rows. Storage: the product's Command Workbook; retention with the workbook's cycle history. Consumers: next cycle's V5 scoring, the A3 Verification scoreboard.
Arbitrage events (P5, P6)	The arbitrage record: both CPC and CVR inputs, the computed adjusted cost, the entry test, every re-test with dates, the paired campaign IDs and roles, state OPEN or CLOSED. Storage: T6 chain with the T3 pair declaration; retained for the life of the term. Consumers: #12 (the dated SP step-down), the V5 scorer, the next candidacy on the same term.
SD audience events (P8, P9)	Pool-size reads with sign-off, entry records with caps, the weekly Performance+ expansion review artifact (the gate artifact; an empty-change week still files one), suspension records with re-entry criteria. Storage: T6 plus the SD section of the weekly read pack. Consumers: the P10 pass, the F4 escalation clock, #20's audience feedback loop.
Monthly split check (P11)	The M6 comparison with shares to 0.1%, the marginal yield read where a format grew, any staged rebalance with its keyword-level verdict list, findings with owners for shortfalls. Storage: A4 Structure targets tab of the Account Rollup Workbook. Consumers: the PPC Lead's monthly review, #26 posture reads.
Event windows (P13)	The plan-referenced deployment set, the post-window remarketing pass record, the reconciliation inputs handed to #32. Consumers: #32 reconciliation, trend-exclusion bookkeeping.
NTB reads (P12)	Window-bounded NTB strategy reads with the boundary date stated on every comparison. Consumers: S-F1, S-F2, S-F6 grading; brand strategy reviews.

9. Skill-Conversion Block

Block	Specification, specific to this component
Input bindings (manifest and feed IDs)	T7 groups A2 (deployment state), A3 (pacing), A5 (threshold flags), A6 (objective and Expected CPC), B8 (economics chain), D7 (decision log); the SB / SBV / SD bulk and RAS sheets of the Raw Pull Pack feeding T3 rows and the D11 brand-format metrics (NTB%, vCTR, completion) per Data Architecture v1.0 Section 12; the creative asset register states; SD audience pool sizes and the Performance+ expansion log; the #16 flag feed (blocked on dependency D4); the #32 event plan dates; the #20 AMC audience handoffs. An absent feed suppresses its dependent execution class; the skill never improvises around a missing feed.
Deterministic logic references	The P5 and P6 inequality computation with the E6 arithmetic convention; the S-F3 and L8 trigger conditions; the S-F2, S-F5, S-F6 grading rules; the G2 sufficiency bars; the X11 and M7 marginal check; the M6 split comparison; the click-attribution gate; the closed verdict vocabulary from the KDR Template v3.0 reference sheet. The skill executes these by reference and computes zero thresholds of its own; it version-pins the criteria authority per the V-layer sync rules.
Human checkpoints (from Authority Matrix v2.0 rows)	Row "Weekly verdicts (established US)": PPC Lead decides, the incoming operator executes on assigned scope, so arbitrage entries and exits, SB exits, and Performance+ opt-ins and suspensions are drafted by the skill and confirmed at that row's D. Row "Structure creation (all SOPs)": every new SB, SBV, or SD campaign is a structure creation, decided at PPC Lead D. Row "PAT/conquest opens and exits" (incoming operator D, E; PPC Lead R): SD competitor contextual opens and exits ride this row because they run #16's flag logic. Row "Event plans (#32)": tentpole format deployments confirm at that row's D. Row "Bulk deployments (#38)": every deploy confirm. Row "v4.0 threshold changes": this component is I only; on any increment the Section 5 mapping re-validates before a pass runs.
Cadence	Weekly P10 read and P9 expansion review; re-tests on their record dates with a daily past-date alarm; monthly P11; pool watch weekly; moderation rejections same day; event actions per the #32 plan dates.
Alarm specification	F1 through F12 each carry the Section 7 detection signal, an evidence-line format (row, signal value, rule reference), and routing: F2, F3, F8, F10, F12 same day to the operator queue; F1, F5, F6, F9, F11 into the weekly worklist; F4 and F7 to the PPC Lead on recurrence.
Skill outputs	Drafted inequality tests with full computations; drafted read rows with attribution labels; drafted expansion review artifacts; the monthly split comparison with the marginal yield read; staged #38 sets queued for human confirmation; F1-F12 alarms

10. Version Governance

Owner: PPC Lead. Review triggers specific to this component: (1) any criteria authority version increment, at which point the Section 5 mapping and every cited section heading re-validate before a pass runs; (2) any Amazon change to SB, SBV, or SD capability or to attribution, which increments the affected format section before use, as the January 2026 NTB tightening did to P12; (3) the landing of component #42 Ad Policy, Moderation and Claims (due 2026-08-17), which retires Element 12 row I-9's interim routing; (4) the landing of the store-page readiness fold-in named on the Ledger row, which replaces P3 step 2's binary check; (5) the landing of dependency D4 (competitor depth feed), which activates P7 step 2 at full depth through #16; (6) resolution of any amendment item A-1 through A-8, which replaces the corresponding interim rule the same week. v2.0 change note: supersedes v1.0 in full; procedures 6 to 13; failure modes 7 to 12; worked examples 1 to 4; local verdict rows retired in favor of the closed vocabulary; Elements 11 and 12 added from primary sources per conflict C5's resolution. Registered in the Supersession Map as 3 components in 1 playbook; each section can split later without renumbering.

11. Operator Ramp

Built from the Authority Matrix v2.0 rows this component touches: "Weekly verdicts (established US)", "PAT/conquest opens and exits", "Event plans (#32)", "Bulk deployments (#38)", "Structure creation (all SOPs)", and "v4.0 threshold changes" (informed only). The checkpoint holder for weeks 2 and 4 is the D-holder of each row; per the Matrix vacancy rule, an unfilled role's D reverts to the PPC Lead and its E to the PPC Manager.

Week 1: observation and artifact familiarity

- Shadow 1 full P10 pass and 1 P9 expansion review. Demonstrate: locate every T3 brand-format row, name its objective tag and band, and state which T7 groups feed it.
- Reproduce the E6 arbitrage computation cold on 3 input sets, including 1 reversal, matching WE-1's arithmetic to the cent (the Economics Worked-Example Set governance requires reproducing the training cases before taking a scope).
- Read 1 open arbitrage record end to end and state its re-test date, its pair roles, and what closes it.

Week 2: supervised execution, checkpoint holder reviewing

- Run P10 with the PPC Lead reviewing every staged row before deploy (the "Weekly verdicts" row's D retained by the Lead).
- Draft 1 P5 inequality test and 1 P9 review artifact for the Lead's confirmation; draft 1 SD contextual rotate under a #16 flag with the "PAT/conquest opens and exits" checkpoint holder reviewing.
- Stage 1 change set per #38 end to end and reconcile its T+1 verification report under the "Bulk deployments" row's review.

Week 4: unsupervised execution, checkpoint holder auditing

- Own the weekly P10 and P9 cycle on the assigned scope; the Lead audits the T6 chain after deploy rather than before.
- Run 1 full P11 monthly check including the marginal yield read, presented to the Lead as a finished finding set.
- Zero F2, F3, or F4 events on the operator's scope across the 2 audited weeks is the pass bar; 1 event triggers a repeat of the week-2 supervision tier for the affected procedure.

Live-defense questions (reasoning, not recall)

- Why must the SBV CPC be CVR-adjusted before it is compared to the SP CPC, and what error does a raw comparison make when SBV CVR runs at 3.2% against SP's 5.0%?
- The console shows an SD CPA of \$9.12 and the ceiling is \$28.00. Why might scaling still be wrong, and which 2 gates do you check before the step?
- An arbitrage passed its entry test 9 weeks ago and has never re-tested. What is the worst case you might be paying right now, in dollars per SP-equivalent conversion, using WE-1's reversal numbers?
- Why does SB never discover, given that SB clicks convert? Name the 2 structural reasons the discovery layers beat it.
- A format grew to 21.6% of spend and its weekly ACOS is in band. Why is the split check still allowed to trim it, and what read decides?
- Why is a Performance+ week without a review artifact treated as a failure even when spend stayed under cap?

12. Interface Contracts

Every upstream and downstream handoff, with the clause the corpus has been missing: what happens when the handoff does not arrive. Sending and receiving owners are roles per the distribution standard.

ID	Interface	Sender	Receiver	Artifact and timing	Non-arrival consequence
I-1	SP benchmark evidence per term	PPC Lead (T2 owner)	This SOP (P2, P4, P5)	T2 K5 CPC and CVR blocks, current each Monday 09:00 per the Data Architecture assembly cadence	No arbitrage tests run and no P2 candidacies open this cycle; open records still re-test on their own stored SP inputs, flagged stale.
I-2	Creative assets	Creative and brand owners (outside PPC)	This SOP (P2, P3, P4)	Register entries: APPROVED video, live headline and store assets, on the asset plan's dates	Candidacy parks; nothing launches on a placeholder (F8); the parked list files as a P11 shortfall cause with the asset owner

					named.
I-3	Competitor weakness flags	#16 (blocked on dependency D4)	This SOP (P7 step 2)	Flagged target list with open, rotate, exit calls, weekly	SD competitor contextual stays closed at structural depth; own-page defense runs regardless; no target opens without a flag, ever.
I-4	Event plan	#32 owner (PPC Lead with the deals owner per the Matrix row)	This SOP (P13)	Dated plan with caps entering T2 at least at the 14-day PRE-EVENT boundary	No tentpole deploys; standing postures hold through the window; a window discovered armed without a plan freezes format changes and escalates to #32.
I-5	Walls	#29 rules, built by this SOP	#29 wall audit	Wall list recorded in T3 C5 in the same session as every SB keyword event	The SB addition does not deploy; a keyword without its wall record returns to the pass (the F2 double-buying class in SOP-12's registry is the account's known cost of skipping this).
I-6	Staging and verification	This SOP to #38, report back	#38, then this SOP	Staged bulk sets Tuesday; T+1 verification report Wednesday	Unstaged changes do not exist; a missing verification report voids the week's reads on affected rows and blocks the next staged set for those campaigns until reconciled (F10).
I-7	ENTER SBV ARBITRAGE step-down	This SOP (verdict and go-live date)	#12 (SP-exact ladder step)	The dated step-down in the same weekly staged cycle as the SBV go-live	If #12 does not execute, the pair economics misstate (the SP side keeps paying pre-arbitrage CPC); the P10 pass checks the step landed at T+1 and raises F10-class escalation to #12 if not.
I-8	AMC audiences	#20 (AMC cycle)	This SOP (P8)	Audience handoffs with size and definition, per the #20 cadence	SD audiences run on console remarketing pools only; the AMC-dependent audience classes stay unbuilt and the gap files as a P11 finding, not as improvised segments.
I-9	Policy, moderation, claims	#42 (NOT BUILT, due 2026-08-17)	This SOP (all creative-bearing campaigns)	Owed: the moderation and claims workflow	Interim: every moderation rejection routes same day to the PPC Lead, the campaign holds paused, and no resubmission occurs without Lead sign-off; this row retires when #42 lands (Element 10 trigger 3).
I-10	Format split rollup	This SOP (P11)	Account Rollup A4; PPC Lead monthly review; #26	The monthly M6 comparison with shares, yield read, and findings, first Monday	A missing check is F7 by definition; the Rollup's A4 tab shows the empty month and the next check runs on a 2-month window with the gap stated.
I-11	Defense set membership	#26	This SOP (P1, P7)	Defense-set membership and escalation instructions, weekly and on incursion	Existing defense campaigns hold their floor postures unchanged; no membership additions or removals occur locally; an unmatched incursion signal escalates back to #26.

Appendix A. v4.1 Amendment Items Logged by This Build

ID	Item, and the interim rule in force until it lands
A-1	SD remarketing pool-size floor: cited by v1.0 as "the v4.0 size floor" but no value exists in the attached criteria file. Interim: human-mandatory pool sign-off at entry (P8 step 2).
A-2	SD audience lookback windows as decision parameters: window values appear only in Strategic Doctrine v1.0 Section 8, the why layer. Interim: windows set at entry with the record stating the source.
A-3	SBV arbitrage re-test interval: the obligation exists, the interval value does not. Interim: interval set at registration, 6-week working reference from the recorded account case (P5 step 4).
A-4	SD Performance+ guardrail numerics (cap sizing, return floor, expansion tolerance): the mode appears only in Doctrine Section 8. Interim: explicit dollar cap at entry, human-mandatory (P9 step 1).
A-5	SD frequency cap values: the criteria file's S-F5 row names cap oversight with no values. Interim: platform defaults stand and are recorded on the entry record.
A-6	SB headline position-share bands: v1.0's position-share read carries no bands in the criteria file. Interim: direction across 2 reads, never a level judgment (P1 step 4).
A-7	Element 2 maximum-lag values where the corpus is silent: lags anchored to the Section 14 V4 cadence classes rather than invented as numbers; a dedicated lag table is proposed for v4.1.
A-8	Gate-numbering reconciliation (conflict C9): the attached criteria file's G9, G10, and G12 differ from the numbering SOP-12 v2.0, the Ledger, and

Appendix B. Weekly Format Pass Operator Card (P10 compressed; the SOP text governs)

1. T0 then T1 before T3, always. 2. Score last cycle's due T6 format rows first. 3. Filter T3 to Type in SB, SBV, SD; gates before metrics on every row (eligibility, truncation, event flag, manifest status). 4. Verify the attribution label on every SD and SBV number before believing it; a view-attributed number in a verdict is F2. 5. Each row grades against its own objective band only, plus NTB depth, vCTR, and completion where T3 C4 carries them. 6. SBV records at or past re-test date run P6 in the same session. 7. SB rows failing S-F2 for 2 validated weeks exit or re-creative. 8. Performance+ rows require this week's expansion review artifact on file. 9. Stage per #38 with verdict or scenario ID; no ID, no deploy; Tuesday deploy, T+1 verification, mismatch is F10 same day. 10. Every deployed change logs its T6 prediction with a review date.