

SOP-20: AMC Operations

PPC Department · Operating SOP · Component #20 of the Operations Definition Register
Version 2.0 · 2026-08-09 · Owner: PPC Lead · Supersedes SOP-20 v1.0 (2026-07-24) in full · Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds · Decision record: AMC Build-or-Defer Memo v1.0 (verdict: BUILD)

1. Verdict and Scope

VERDICT: AMC is the account's only deduped, user-level, cross-format evidence source, and it enters the system exclusively as a measurement layer with named consumers. It issues zero verdicts from the closed vocabulary, because the Master PPC Decision Criteria System v4.0 contains no AMC rule, band, or threshold today: every decision an AMC read touches is issued elsewhere, with the filed AMC result attached as evidence. Version 2.0 expands v1.0 from 6 procedures to 12, from 6 failure modes to 10, from 1 worked example to 4, adds Elements 11 and 12 (absent from every shipped SOP per conflict C5), and retires v1.0's 6-row verdict table: none of its strings appears in the closed vocabulary, so all 6 are reclassified as operating states pending v4.1 admission (AM-1). This SOP governs the entire operating life of the AMC layer: instance provisioning and account linking across both entity ad accounts, schema versioning and console reconciliation, the versioned 5-query template library (Q1-Q5) and its build order, the monthly Q1-Q3 cycle and its filing standard, the reading rules, audience construction at the platform minimums and the handoff of audience records into SD, refresh, decay watch and retirement, the quarterly Q5 dayparting evidence gate, ad-hoc named-evidence runs, and new-query governance. The economics of the layer: 0 license cost since the September 2025 self-service change, a 24-month lookback waiting in history, 5 template queries, and 2 platform constraints on audiences (2,000-user minimum, 500-user lookalike seed), all per Strategic Doctrine v1.0 Section 10 "AMC, DSP, and the Measurement Layer" (platform state verified July 2026).

1.1 Position: a measurement layer, never a decision layer

Strategic Doctrine v1.0 Section 10 fixes the layer's constitution: every AMC read enters the system as evidence attached to a verdict, never as a parallel reporting universe, and the workbook remains the decision surface. Operationally that means 3 things. First, no bid, budget, placement, negation, structure, or status change cites an AMC number unless a v4.0-issued verdict names that filed result as its evidence (F1 is the breach, and it escalates on first occurrence). Second, every query in the library exists because a named consumer SOP reads it into a named decision; a query without a consumer is retired, not kept (F3). Third, the layer's own outputs are records, not actions: filed result sets, audience records, gate answers, and log entries, each with a window, a query version, and a schema version, so the consumer can audit the evidence in the same 30 seconds v4.0 Section 13 demands of every Master Action Record.

1.2 Why this SOP issues zero closed-vocabulary verdicts

The closed vocabulary is held in two attached instruments: the Keyword Decision Record Template v3.0, "Bands & Verdicts Reference" sheet (25 strings) and v4.0 Section 13.1 "The Closed Verdict Vocabulary" (18 actions); the 18-versus-25 count divergence is a registered program finding (logged previously by the #41 and #46 build sessions) and is handled here by quoting the KDR strings per the standing rule and naming the v4.0 equivalent where the two differ. A full scan of both lists returns no AMC action of any kind. The nearest v4.0 touch points are structural, not vocabular: gate G2 in Section 5 supplies the per-cell sufficiency test this SOP applies to every read; Section 8.10 "Brand-Format Scenarios (S-F1 to S-F8)" issues the format decisions that consume Q1-Q3 evidence; Section 9 rule L10 defines the dormant dayparting lever whose evidence gate is Q5; Section 14 rule V4 supplies the monthly and quarterly calendar slots the cycle lands on. Precedent: SOP-43 (Ledger row #43) established the pattern for a component whose subject matter v4.0 does not yet regulate: issue nothing, run named operating states as register vocabulary, and log every needed number as a v4.1 amendment item rather than writing it locally. This SOP follows that pattern exactly; its 5 amendment items are listed in the closing section and cited as AM-1 to AM-5 throughout.

1.3 Boundaries With Adjacent Components

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, gate, scenario, and verdict. This SOP files evidence; it never decides. A decision number appearing in this SOP instead of v4.0 is a defect; the 5 numbers this SOP needs and v4.0 lacks are AM-1 to AM-5, not local rules.

Adjacent component	The seam, stated so it is owned
SB/SBV/SD Playbook (#17-19)	Owns all SD execution: audience campaign entry, caps, click-attributed grading against the ceiling, Performance+ guardrails, and every format verdict's execution. This SOP hands over audience records and filed result sets; it never creates, funds, or grades a campaign. Referenced by component number and clause only: the Tier 2 copy attached to this build is v1.0 while the Ledger row #17 records v2.0 as the superseding current version (registered conflict C8, third occurrence), so no playbook procedure number is load-bearing here.
#45 Dayparting (deferred)	Owns the deferral and its trigger. This SOP produces the quarterly Q5 evidence that answers the deferral question; it never deploys an hourly modifier. Per v4.0 Section 9 rule L10, dayparting stays dormant until the Marketing Stream feed (Catalog D12, Ledger dependency D5) lands with at least 4 weeks of hour-level evidence per campaign; Q5 evidence alone opens an evaluation, never a lever.
#21 DSP and Sponsored TV	Both remain deferred per the AMC Build-or-Defer Memo v1.0 Section 4, each with named re-evaluation triggers. Nothing in this SOP presumes either; AMC audiences deploy into SD without DSP, which is the memo's stated reason the deferral stands.
#35 Incrementality testing	Owns test design, run, and read. Q2 overlap evidence is observational context for format allocation; it never substitutes for a designed #35 test, and no incrementality claim cites Q2 alone.
#32 Event windows	Owns event plans and the armed-window state. This SOP consumes the window dates so reads honor v4.0 Section 5 gate G10 (event weeks and the 14 days following excluded from trend verdicts).
#39 Data feeds and manifest	Owns feed definitions, T7 field groups, and retention policy lines. AMC has no Catalog field-group row in T7 today; the assignment request is interface I-9 and an open item routed to #39, not a row this SOP invents.
#46 Access and permissions	Owns the instance seats. An AMC seat live without its Access Register row suspends on discovery per that SOP's discipline; the cycle halts until the seat registers (I-10).
#6 V5 prediction scoring	Owns scoring. The memo's 5 quarter-one predictions and every gate-answer prediction this SOP logs score through #6 per v4.0 Section 14 rules V1 and V5.

2. Trigger Conditions

Every trigger lands on a calendar slot defined by v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)", rule V4 (daily / weekly / monthly / quarterly), or fires on a named artifact. Numeric maximum-lag values in days from signal to procedure start exist nowhere in v4.0; the lag table is AM-5 (the same gap the #17-19 rebuild logged for its own Element 2). Until admitted, the binding lag is the calendar relationship stated per row, which is auditable from the artifacts named.

Trigger	Detecting artifact	Procedure	Maximum allowed lag
Monthly cycle	Query calendar rows in the AMC Query Register (Element 8)	P4	Filed inside the same V4 monthly slot, and before the M6 structure-target check reads Account Rollup tab A4, so the monthly format read consumes current evidence, never last month's.
Audience window advance	Refresh date column on the live audience record	P8	Same week as the recorded date; a live audience past its refresh date is F4 territory.
Weekly decay watch	Pool-size column on every live audience record	P8 step 1	Inside the weekly V4 slot; a sub-minimum pool notifies the playbook owner the same week.
Quarterly deep read	Query calendar quarterly row; V4 quarterly slot	P9	Gate answer filed before the quarterly framework review closes, so the deferral question is answered inside the quarter that asked it.
Named-evidence verdict	A T2 or T3 verdict draft citing AMC evidence	P10	Before the verdict deploys: per v4.0 Section 12 rule P10, a verdict without complete justification does not deploy, and an unfilled AMC citation is incomplete evidence.
New-query proposal	Written proposal entering the governance log	P11	Decision recorded before any build starts; nothing is queried first and justified later.

Trigger	Detecting artifact	Procedure	Maximum allowed lag
Platform or schema change	Schema version at any run differs from the Instance Record	P3 re-document, P12	Before the next filing; no result files under an undocumented schema.
Reconciliation fault	P2 monthly check: a format row absent or divergent beyond reason	P2 step 4	Data-refresh task opened the same day, matching the v4.0 Section 5 gate G1 same-day repair pattern.

3. Preconditions and Gates

Every procedure below runs only when these read current. AMC carries no Catalog field-group row in the T7 Manifest (the Master PPC Data Architecture v1.0 Section 9 group list runs A1-A8, B1-B11, C1-C8, D1-D12, E, and none is AMC), so until #39 assigns one, every filed result set carries its own manifest surrogate: window dates, query version, schema version, and UNREAD cell count on its header line. The G8 consequence still binds the consumers: a verdict whose named AMC evidence is missing or stale downgrades to a flag and states what it could not see.

Gate / precondition (source)	Requirement and the observable that proves it passed
Instance provisioned and linked (P1 record)	Both entity ad accounts linked; observable: the P2 reconciliation query returns SP, SB, and SD rows, each nonzero, for the trailing full month, and the Instance Record carries instance ID, linked account IDs, and schema version.
Backfill present (Doctrine v1.0 Section 10)	The 24-month lookback returns data at instance creation; observable: a P2 run windowed 12 months back returns populated rows.
G2 sample sufficiency (v4.0 Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass")	Quoted: "CVR verdicts ≥ 100 clicks; CTR verdicts $\geq 1,000$ impressions." Applied per query cell. A cell below the bar files as UNREAD (register state, AM-1) and receives no rate judgment; the v4.0 keyword-side analog of this state is its Section 13.1 action EXTEND TEST. Observable: the filed set's UNREAD count equals the count of sub-bar cells, checked mechanically.
G8 packet completeness (v4.0 Section 5)	The manifest surrogate line is present on every filed set; a consumer citing a set without one treats the evidence as missing. Observable: header line present, 4 fields populated.
G9 price stability (v4.0 Section 5)	A CVR-bearing read (Q2 conversion rates, Q5 hourly CVR) whose window spans a price change on the read's products carries PROVISIONAL, read from the T5 Product Log price-move rows. Observable: PROVISIONAL flag on the filed set naming the T5 row.
G10 event exclusion (v4.0 Section 5)	Quoted: "Event weeks and the 14 days following are excluded from trend verdicts." Q5 hour-cells and any trend read inside an armed #32 window are excluded, read from the T5 trend-exclusion flags. Observable: excluded-window list on the filed set. The corpus's G12 labeling of this flag in SOP-12 v2.0 and the Data Architecture T5 column is registered conflict C9; this SOP follows the v4.0 file's own gate IDs.
NTB boundary (Doctrine v1.0 Section 7 measurement notes)	NTB attribution tightened January 2026; NTB comparisons run only within post-2026-01 windows. Observable: every Q3 filing's window starts 2026-02-01 or later.
Audience platform minimums (Doctrine v1.0 Section 10)	Minimum audience 2,000 users; lookalike seeds 500. Platform constraints, not decision thresholds; the decision-side pool-entry floor and lookback windows remain absent from v4.0 (AM-2, same gap the #17-19 rebuild logged). Observable: the creation record shows the projected count against 2,000 and, where a lookalike is built, the seed count against 500, before any handoff.
Access current (#46)	Every seat touching the instance holds an Access Register row. Observable: the register's filtered view returns the seat; an unregistered live seat suspends per #46 and halts the cycle (I-10).

4. Step-Level Procedures

P1. Instance provisioning and account linking (one time; 2-3 hours plus propagation)

1. In the Amazon Ads Console for the primary entity account, open Measurement & Reporting, then Amazon Marketing Cloud, and create the self-service instance (0 license cost per Doctrine v1.0 Section 10). Record the instance ID and region.
2. Link both entity advertiser accounts (Ecotero LLC, Inspirattek LLC) to the instance from the instance's account-management screen. Observable: both appear in the linked-accounts list with status active.
3. Confirm backfill: run any template query windowed 12 months back; populated rows prove the 24-month lookback is live for the linked accounts.

4. Write the AMC Instance Record (Element 8): instance ID, region, linked account IDs, creation date, schema version string as displayed in the query editor, and the #46 seat list.
5. Register the seats with #46 the same day; the record is incomplete until the Access Register rows exist.

P2. Console reconciliation and schema documentation (45-60 minutes at setup; 15 minutes each monthly cycle)

1. Run the template attribution query for the trailing full calendar month, grouped by format (SP, SB with SBV, SD), returning attributed conversions and attributed sales per format.
2. Pull the same trailing month's per-format totals from the ad console (Sponsored Ads reports, campaign-type rollup, same attribution window as displayed).
3. Compare per format. AMC's modeled, deduped attribution will not equal console last-touch numbers, and the check is directional: every format present, nonzero, and the same order of magnitude, with the signed delta recorded on the Instance Record. A numeric divergence tolerance does not exist in v4.0 and is AM-4; until admitted, the recorded delta series is the evidence the tolerance will be set from.
4. Branch on the result: all formats present and in range: proceed and stamp the cycle reconciled. Any format absent, zero, or off by an order of magnitude: open a data-refresh task the same day (the v4.0 Section 5 gate G1 pattern), file nothing from this cycle, and notify consumers that the cycle is held (F7).
5. Document the schema version on every run; a changed version string triggers P3 re-documentation and a P12 NEW BASELINE before anything files.

P3. Build and maintain the template query library (build: 4-6 hours across the first quarter, in the memo's order; maintain: with P12)

Build order per the Build-or-Defer Memo v1.0 Section 5: Q1-Q3 in the first month, Q4 in the second, Q5 in the first quarterly cycle. Each query's register row carries: query ID, version, owner, named consumer SOP, the decision it feeds, window convention, and output table spec.

Query	Specification (dimensions, measures, window)	Named consumer and the decision it feeds
Q1 Path to conversion	By product family: ordered format-touch sequences preceding purchase; output: path share table with 5 rows (SP-only, SB-assisted, SBV-assisted, SD-assisted, multi-touch), share of converters and CVR per path; trailing full month.	The #17-19 playbook's monthly format-split read against the M6 structure targets (v4.0 Section 10 "Layer 5: Sizing Mathematics (M1-M10). Every Action Carries Numbers", rule M6: SP ~80%, SB 15-20%, SD 5-10%, deviations are findings to explain). Path evidence is what the marginal format dollar re-justifies with.
Q2 Cross-format overlap	Deduped users touched by 2 or more formats; measures: overlap share of converters, conversion rate of multi-format versus single-format users; trailing full month.	SB role grading under v4.0 Section 8.10 conditions S-F1 and S-F2 (SB judged on NTB% and brand value, never SP metrics); the SB incrementality question read on deduped users instead of format-siloed consoles.
Q3 NTB depth	New-to-brand share by format and term class; windows post-2026-01 only per the boundary gate; trailing full month.	Brand-growth composition strategy read for the playbook; NTB% inputs to the S-F1 condition. A strategy read, not a bid input, unless a verdict names it (P10).
Q4 High-intent audiences	Cart abandoners; detail-page viewers without purchase; per family; lookback window recorded on each audience record (no v4.0 window exists: AM-2); built as audiences only where P6 clears.	The playbook's SD audience entry: capped, click-attributed, ceiling-governed, per its own rules; this SOP hands the record only (I-3).
Q5 Hourly evidence	Conversions and CVR by hour of day, by family, 24 cells per family; per-cell G2 sufficiency; event windows excluded per G10; quarterly window.	The #45 dayparting deferral: the quarterly evidence gate (P9). Deployment evidence remains v4.0 Section 9 rule L10's Marketing Stream condition, which Q5 never substitutes for.

P4. Monthly cycle execution and filing (90-120 minutes per cycle)

1. Run the P2 step set first; a failed reconciliation holds the cycle (F7).
2. Run Q1, Q2, Q3 at current versions, trailing full month, per family.
3. File each result set with its manifest surrogate header: query ID and version, window start and end dates, schema version, row count, UNREAD cell count, PROVISIONAL and event-exclusion flags where they apply, and reading notes of 3 sentences or fewer per finding.
4. Write the T6 Decision Log entry (11 columns per Data Architecture v1.0 Section 9 "The Infrastructure Tabs: Logs and the Manifest (T5-T7)") for the cycle: date, level, target = the filed set IDs, rule reference = this SOP P4, expected outcome where a prediction is

logged, review date = next cycle.

5. Notify the playbook owner row that the cycle is filed, before the monthly M6 check reads A4; the consumer cites the filed result, never a screenshot (F9).

P5. Reading rules (applies to every read; 15 minutes per result set)

1. Per cell, apply G2: a CVR cell under 100 clicks or a CTR cell under 1,000 impressions is UNREAD. UNREAD cells are reported as UNREAD, never interpolated, never averaged into neighbors (F2).
2. Classify the read: strategy input (default) or named evidence (only when a verdict names the filed set per P10). No bid or budget change cites an AMC number without a v4.0 verdict naming it (F1).
3. Exclude armed event windows and the 14 days following from any trend statement (G10), reading the windows from the T5 flags (I-6).
4. Hold the NTB boundary: no Q3 comparison spans 2026-01 (F8).
5. Never trend across query versions: a changed query is a new version and a NEW BASELINE (P12, F5).

P6. Audience construction from Q4 (30-45 minutes per audience)

1. Pull the candidate pool's projected member count from the audience-creation screen; for lookalike builds, also the seed count.
2. Decision tree (3 branches):
 - Branch A: projected at or above 2,000 (and seed at or above 500 where a lookalike): create. An audience proposal that clears the minimums is created at its recorded window.
 - Branch B: projected under 2,000: do not create. The proposal is not created small; it is not created. Log the count and the shortfall; re-propose only when the pool grows or the window is legitimately widened, which is a new audience version.
 - Branch C: minimums clear but the window deviates from the register convention for that audience class: create only with the deviation recorded on the audience record; window values carry no v4.0 basis (AM-2), so the record is the audit trail the future threshold will be set from.
3. Name the audience **AUD-<family>-<queryID>v<version>-<window>-<YYYY-MM>**, so the name alone reconstructs its provenance.
4. Write the audience record (Element 8): name, query ID and version, window, projected count at creation, seed count where applicable, creation date, refresh date, status CREATED.
5. The creation decision is the D-holder's: PPC Lead per the Authority Matrix v2.0 Section 1 row "AMC cycle (#20)". The operator prepares; the D-holder decides (Element 9).

P7. Audience handoff into SD (15 minutes; execution belongs to #17-19)

1. Hand the audience record to the playbook owner row by component number: the playbook's SD audience-entry procedure governs the campaign build, caps, click-attributed grading against the T0 ceiling cells, and every subsequent format decision. This SOP launches nothing.
2. Record the handoff date and the receiving owner on the audience record; status remains CREATED until the playbook reports a hosting campaign (I-4), then LIVE.
3. Until an audience record exists and is handed over, the playbook's SD audiences run on console remarketing pools only, which is the standing non-arrival consequence its own Ledger row carries (I-3).

P8. Refresh, decay watch, and retirement (10 minutes weekly plus refresh events)

1. Weekly, read the current pool size of every LIVE audience. A pool under 2,000 notifies the playbook owner the same week; the hosting campaign pauses per that SOP's rules, and the audience status moves to DECAYED (F4 is this watch skipped).
2. On the recorded refresh date, re-project the pool: at or above 2,000, refresh and update the record; under 2,000, run the P6 tree (a widened window is a new version, not a silent edit).
3. Retire an audience whose pool stays under 2,000 across 2 consecutive refresh checks or whose consumer campaign closes: status RETIRED, retirement date and reason on the record. The record is never deleted; retention rides Element 8.

P9. Quarterly Q5 run and the dayparting evidence gate (2-3 hours per quarter)

1. Run Q5 per family at the current version for the trailing quarter, excluding armed event windows and their 14-day tails (G10).
2. Per family, apply G2 per hour-cell: count readable cells (CVR cells at 100 clicks or more) and UNREAD cells out of 24.
3. Compute the pattern read on readable cells only: pooled CVR per readable day-part against the readable remainder, stated as points of premium and relative percent. No interpolation across UNREAD cells, ever (F2).
4. Decision tree (3 branches):
 - Branch A: readable majority (13 or more of 24) and a pattern the D-holder judges material: the gate answer files as pattern present, and the PPC Lead may open a dayparting evaluation toward #45. Deployment still waits on v4.0 Section 9 rule L10: Marketing Stream curves with at least 4 weeks of hour-level evidence per campaign; Q5 evidence alone never sets a modifier. A numeric materiality bar does not exist in v4.0 and is AM-3; until admitted, the gate files the numbers and the call is the Matrix D-holder's.

- Branch B: readable majority, no material pattern: the deferral continues on evidence; the answer files with the computed premium.
 - Branch C: UNREAD majority: the deferral continues by insufficiency; the answer files with the UNREAD count. Insufficiency is an answer, not a gap to fill by interpolation.
5. File the gate answer with a V1 prediction where branch A opens an evaluation (metric, direction, magnitude, horizon), write the T6 entry, and route the answer to the #45 deferral owner (I-5); it scores through #6 next quarter.

P10. Ad-hoc named-evidence runs (30-60 minutes per request)

1. Trigger: a verdict draft in T2 or T3 cites AMC evidence. Run the named query at its current version for the window the verdict needs.
2. File the result set versioned and cross-referenced to the verdict's record ID; the Master Action Record's evidence line cites the filed set ID, never a screenshot or an unfiled number (F9).
3. Per v4.0 Section 12 "Layer 7: Precedence and Conflict Resolution (P1-P10)", rule P10, the verdict does not deploy until its justification is complete; an AMC citation without a filed set is incomplete by definition.
4. The run logs in the governance log with its requesting verdict, so the quarterly review sees which verdict classes actually consume AMC evidence.

P11. New-query governance (30 minutes per proposal; quarterly review inside the V4 slot)

1. A proposal enters only in writing: proposed query, the consumer SOP by number, and the decision it feeds, written before anything is built. Measurement without a consumer is a hobby, and it is declined as one.
2. Decision tree (3 branches):
 - Branch A: consumer named, decision named, no existing query covers it: ADMITTED (register state); the query builds at version 1 with its register row complete.
 - Branch B: consumer named but an existing query already feeds that decision: DECLINED with the pointer to the existing query; duplicates fragment baselines.
 - Branch C: no consumer or no decision named: DECLINED; the proposal may return when it has one.
3. Quarterly, review the library: a query whose named consumer has not cited a filing of it for 2 consecutive quarters retires (register state RETIRED; governance cadence carried from v1.0 P6.2, aligned to the V4 quarterly slot; it is library hygiene, not a spend threshold).
4. Admissions, declines, and retirements write to the governance log with the D-holder's sign-off per the Matrix row.

P12. Query versioning and re-baselining (10 minutes per change)

1. Any change to a query's logic, dimensions, measures, window convention, or underlying schema is a new version: increment the register row, record the change and date.
2. A new version is a NEW BASELINE: trend statements across versions are void, and any filing that trends across a version boundary re-files split at the boundary (F5). Cross-version trending is the AMC form of the NTB boundary error.
3. A platform schema change (P2 step 5) re-versions every affected query in one register pass before the next cycle files.

5. Decision Points: Verdict Mapping

Verdicts this SOP issues from the closed vocabulary: zero. The Keyword Decision Record Template v3.0 "Bands & Verdicts Reference" sheet holds 25 verdict strings and v4.0 Section 13.1 "The Closed Verdict Vocabulary" holds 18 actions; neither list contains an AMC action, so per anti-drift rule 1 this SOP mints none. v1.0's Section 5 table carried 6 local strings (RUN CYCLE / FILE; CREATE / REFRESH / RETIRE AUDIENCE; UNREAD; OPEN DAYPARTING EVALUATION; DEFER CONTINUES; ADMIT / DECLINE / RETIRE QUERY); all 6 are retired as verdicts at v2.0 and re-classed as operating states in the register vocabulary below, pending v4.1 admission under the v4.0 Section 14 rule V3 amendment protocol (AM-1). Two mappings remain and both run outward.

5.1 Closed-vocabulary verdicts that may name AMC filed evidence

Verdict (quoted verbatim, KDR v3.0 "Bands & Verdicts Reference")	AMC role and the execution boundary
ENTER SBV ARBITRAGE	Q1 path shares and Q2 overlap rates may inform the candidacy read on a term; the inequality, entry, and exit belong to the #17-19 playbook under v4.0 Section 8.10 (S-F3, S-F4) and Section 9 rule L8. This SOP's role ends at P10: run and file the named evidence before the verdict deploys. The v4.0 Section 13.1 equivalent string is FORMAT SUBSTITUTION (SBV/SD); the wording divergence rides the registered 18-versus-25 finding.
INCREMENTALITY LADDER	Q2 deduped-overlap evidence is context in the test-candidacy discussion; design, run, and read belong to #35, and no ladder cites Q2 alone. This SOP files the named evidence on request (P10).

Verdict (quoted verbatim, KDR v3.0 "Bands & Verdicts Reference")	AMC role and the execution boundary
QUARANTINE: DATA REFRESH	When P2 reconciliation fails, the affected consumer rows follow the quarantine path: no recommendation, repair task same day. The v4.0 Section 13.1 equivalent is FLAG: DATA REPAIR (G1, G8). This SOP opens the task and holds the cycle; the row-level verdict is issued where the row lives.

5.2 Operating states (register vocabulary, pending v4.1 admission per AM-1)

State	Where it is set and what it permits
RECONCILED / HELD (cycle)	P2 step 4. HELD files nothing and notifies consumers; RECONCILED releases P4.
FILED (result set)	P4 step 3 with the manifest surrogate header. Only FILED sets may be cited by any verdict (F9).
UNREAD (cell)	P5 step 1 on any sub-G2 cell. Permits nothing; forbids interpolation (F2).
CREATED / LIVE / DECAYED / RETIRED (audience)	P6-P8 on the audience record. DECAYED requires the same-week pause notice (F4).
ADMITTED / DECLINED / RETIRED (query)	P11 in the governance log with the D-holder sign-off.
DEFERRAL CONTINUES / EVALUATION OPEN (gate answer)	P9 branches. EVALUATION OPEN routes to #45 and still deploys nothing per L10.

6. Worked Examples

All 4 are illustrative and internally consistent; live values come from the instance and the T0 spine. Where product economics appear, the standard fictional base applies and is declared in the example.

WE-1. Clean monthly cycle: Q2 and Q3 file and the format read consumes them

Illustrative values. Window 2026-07-01 to 2026-07-31, schema v2.4.1, Q2 v1, Q3 v1; P2 reconciled (all 3 formats present, deltas recorded). Q2: 20,000 deduped converters; 3,680 touched 2 or more formats: overlap share = $3,680 / 20,000 = 18.4\%$. Multi-format converter CVR 8.1% versus single-format 5.2%: ratio = $8.1 / 5.2 = 1.56x$. Q3 (window entirely post-boundary): SB-touched converters 2,400, of which 984 new-to-brand: $NTB = 984 / 2,400 = 41.0\%$; SP-only converters 14,100, of which 3,384 new-to-brand: $NTB = 3,384 / 14,100 = 24.0\%$. Both sets file with 0 UNREAD cells at the family grain. Consumption: the playbook's monthly M6 read on Account Rollup A4 shows SP 78.4%, SB 16.2%, SD 5.4% (sum 100.0%), all inside the v4.0 Section 10 rule M6 targets; the S-F1 condition (NTB% healthy) cites the filed Q3 result by set ID. No rebalance is staged, and the reading note is 2 sentences. The layer did its one job: the split check ran on deduped evidence instead of format-siloed consoles.

WE-2. Clean audience lifecycle: creation, handoff, first read

Illustrative values on the standard fictional base (AOV \$80.00, contribution \$28.00, break-even ACOS 35.0%, Target 17.5%, Max Acceptable 26.3%). Q4 v1 returns a cart-abandoner pool for the satin family, 30-day window, projected 3,900 members: 3,900 against the 2,000 platform minimum = 95.0% headroom; P6 Branch A: create. Name: AUD-SATIN-Q4v1-30d-2026-08. Record written; the PPC Lead signs the creation per the Matrix #20 row D. P7 handoff same week; the playbook builds the capped SD entry under its own rules. First qualified read (the playbook's, cited here as the record trail): click-attributed CPA \$17.20 against the \$28.00 ceiling = 61.4% of ceiling; their verdict path, their execution. This SOP's artifacts at close: 1 audience record with status LIVE, 1 handoff line, 1 T6 entry with the V1 prediction "audience CPA at or under ceiling at first qualified read."

WE-3. Edge case: decay in flight, caught by the watch

Illustrative values. AUD-DCLR-Q4v1-30d-2026-06 created at projected 2,240. Week 9 decay watch: pool re-projects at 1,760: 240 under the 2,000 minimum = 12.0% short, a decline of 480 users = 21.4% from creation. Same-week notice to the playbook owner; the hosting campaign pauses per that SOP; status DECAYED, T6 entry written. Refresh attempt: widening the window to 60 days re-projects at 2,610, above minimum, but a widened window is a new audience version (P8 step 2), so AUD-DCLR-Q4v2-60d-2026-08 is created fresh with its own record, and v1 retires with reason "pool decay below platform minimum." What went wrong and got handled: nothing paused late, nothing trended across the window change, and the audience lineage is reconstructable from 2 records and their names.

WE-4. Edge case: quarterly Q5 with an UNREAD majority and an interpolation attempt

Illustrative values. Q5 v1, trailing quarter, DCLR family, 1 armed event window excluded per G10. Of 24 hour-cells, 8 clear the G2 CVR bar of 100 clicks (6 evening, 2 midday); 16 are UNREAD: $16 / 24 = 66.7\%$ UNREAD. Two near-miss cells sit at 87 and 41 clicks: both UNREAD, no exceptions. Readable pattern read: evening pooled 780 clicks, 42 orders: $CVR = 42 / 780 = 5.4\%$; midday pooled 240 clicks, 12 orders: $CVR = 12 / 240 = 5.0\%$; premium 0.4 points = 7.7% relative, computed on readable cells only. An analyst draft fills 4

UNREAD afternoon cells with the 5.0% midday rate to claim fuller coverage: F2 fires, the draft re-files with UNREAD marked, and the note quotes the gate: below the bar there is no rate judgment. Gate answer: P9 Branch C, DEFERRAL CONTINUES by insufficiency, filed with the 66.7% UNREAD count and the readable premium for the record. The dayparting question has an evidence-shaped answer, and the planning sheets stay empty until the gate opens.

7. Failure Modes and Escalation

Minimum detection-to-response contract per mode; the incident column carries the account-history motivation where one is recorded. v1.0 records no named, dated incidents; its single institutional trace (the dayparting planning sheets that existed before any hour-level evidence) carries into F6, and the remaining motivations cite their documented risk source rather than a manufactured history.

Failure	Detection signal	Response	Incident / motivation
F1. Parallel decision system	A bid, budget, or structure change citing an AMC number with no v4.0 verdict naming a filed set	Withdraw the change; route the question to the owning SOP; escalates on first occurrence: a parallel decision path is a governance breach	No recorded incident; constitution per Doctrine v1.0 Section 10 (evidence attached to a verdict, never a parallel reporting universe)
F2. Sub-minimum interpretation	A filed reading interpreting, averaging, or interpolating an UNREAD cell	Re-file with UNREAD marked; interpolation is manufacture; the note quotes G2	No recorded incident; risk named in the Build-or-Defer Memo v1.0 Section 2 (query misreads producing confident nonsense)
F3. Orphan queries	A library query with no consumer citation for 2 consecutive quarters at the P11 review	Retire per P11 step 3; measurement without a consumer is a hobby	No recorded incident; the memo's value test (every query has a named consumer) is the standing guard
F4. Audience decay unwatched	A LIVE audience record whose weekly pool read is under 2,000, or a skipped weekly read	Same-week notice to the playbook owner; hosting campaign pauses per that SOP; status DECAYED	No recorded incident; carried from v1.0 F4 as the named decay guard
F5. Unversioned query drift	A changed query filed under its old version; a trend statement spanning a version boundary	Re-version, NEW BASELINE, re-file split at the boundary; cross-version trends are void	No recorded incident; carried from v1.0 F5; the NTB boundary error generalized
F6. Dayparting on optics	A dayparting proposal, plan, or modifier draft without a filed Q5 gate answer	Decline at intake; the P9 gate is the only door, and deployment additionally waits on L10's Marketing Stream condition	Carried from v1.0 F6, the SOP's one institutional trace: dayparting planning sheets existed before any hour-level evidence did, and the rule is that the empty planning sheets stay empty until the gate opens
F7. Reconciliation divergence	P2 monthly check: a format row absent, zero, or off by an order of magnitude against console totals	Data-refresh task same day (G1 pattern); cycle HELD; nothing files; consumers notified	No recorded incident; new at v2.0, closing the gap that v1.0 checked linkage only at setup
F8. NTB boundary breach	A Q3 filing whose window starts before 2026-02-01, or any NTB trend spanning 2026-01	Re-file inside post-boundary windows; the comparison is void, not adjusted	No recorded incident; motivated by Doctrine v1.0 Section 7 (January 2026 tightening re-baselined NTB account-wide)
F9. Screenshot citation	A consumer document or record citing an AMC number with no filed set ID	Run P10, file, re-issue the citation; per v4.0 Section 12 rule P10 the verdict does not deploy until the evidence is complete	No recorded incident; new at v2.0; enforces the filing standard v1.0 stated but did not police
F10. Audience live without its record	A deployed audience with no creation record showing the 2,000 / 500 check, or a name that does not parse to query version and window	Suspend the handoff; write or reconstruct the record before the campaign resumes, or retire the audience	No recorded incident; new at v2.0, mirroring the #46 register discipline (an asset live without its row suspends on discovery)

Escalation: F1 escalates to the PPC Lead on first occurrence with the CEO informed per the Matrix #20 row (I column); F7 escalates same day to Data Operations; F2, F4, F5, F8, F9, F10 route through the monthly cycle review; F3 and F6 route through the quarterly review.

8. Outputs and Artifacts

Artifact	Content and storage	Retention	Downstream consumer
AMC Instance Record	Instance ID, region, linked accounts, schema version history, reconciliation delta series; 1 file in the AMC folder of the PPC document library, indexed by the Query Register	Life of the instance plus supersession per #1	P2 checks; #46 seat reconciliation; audits

Artifact	Content and storage	Retention	Downstream consumer
AMC Query Register	1 row per query: ID, version history, owner, consumer SOP, decision fed, window convention, calendar rows; the library's index	Permanent, append-only	P3-P5, P11, P12; the quarterly review; #44 (ramp reading)
Monthly Filed Result Sets (Q1-Q3)	Versioned files with the manifest surrogate header (window, query version, schema version, UNREAD count) and reading notes	Until superseded, minimum 24 months to match the platform lookback; the retention policy line is a #39 pending fold-in (I-9)	#17-19 monthly M6 read and S-F grading; any P10 named-evidence verdict
Audience Records	1 record per audience: name, provenance, counts against minimums, window, status history, handoff line	Permanent, append-only; RETIRED records kept	#17-19 SD entry (I-3); P8 decay watch; F10 audit
Quarterly Q5 Gate Answer	Readable and UNREAD counts, excluded windows, computed premium, branch taken, V1 prediction where applicable	Permanent	#45 deferral owner (I-5); #6 V5 scoring (I-8); the quarterly framework review
Governance Log	P11 admissions, declines, retirements; P10 run log with requesting verdicts	Permanent, append-only	Quarterly review; #1 run ledger
T6 Decision Log entries	11-column rows per Data Architecture v1.0 Section 9 for cycles, audience decisions, and gate answers	Permanent per T6 (append-only)	Account Rollup A3 scoreboard; #6 V5 scoring

9. Skill-Conversion Block

Block	Specification
Input bindings	The AMC instance (query execution surface and audience-creation screen with projected counts); the AMC Query Register and calendar rows; audience records with pool sizes and refresh dates; console per-format monthly totals for P2; T5 Product Log price-move rows and trend-exclusion flags (G9, G10 inputs); Account Rollup A4 for the consumption check; T6 for write-back; the #46 Access Register filtered view.
Deterministic logic	The V4 cadence calendar; the G2 per-cell test (CVR 100 clicks, CTR 1,000 impressions, quoted from v4.0 Section 5) and UNREAD filing; the 2,000 / 500 minimums check; the AUD naming grammar; the P12 versioning and NEW BASELINE rule; the 2-consecutive-quarter retirement check; the manifest surrogate header completeness check. Zero local thresholds: every number above quotes v4.0, Doctrine Section 10 platform facts, or a register convention, and the 5 absent numbers are AM-1 to AM-5, never computed locally.
Human checkpoints	Drawn from the Authority Matrix v2.0 Section 1 row "AMC cycle (#20)": CEO I, PPC Lead D, Incoming operator E, BM owner I; Standing Rule 5 applies on vacancy (D reverts to the PPC Lead, E to the PPC Manager). AI-executable with audit: query runs, filings with drafted reading notes, reconciliation computation, decay watch, Q5 counts and premium computation, drafted audience and T6 records. Human-mandatory (the D-holder): audience creation and retirement, query admission and retirement, the P9 gate answer and any EVALUATION OPEN, and any P10 evidence attachment to a verdict in the Human-mandatory review tier per v4.0 Section 13.2.
Cadence	Monthly: P2 plus Q1-Q3 (P4). Weekly: decay watch (P8 step 1). Quarterly: Q5 and the gate (P9), library review (P11 step 3). Event-driven: P10 runs, refresh dates, schema changes.
Outputs	Filed result sets, audience records, gate answers, governance log rows, drafted T6 entries queued for the D-holder, and F1-F10 alarms with evidence lines (set ID, cell or record, gate or rule reference).

10. Version Governance

Owner: PPC Lead. Review quarterly with the P11 library review, and immediately on any of these component-specific triggers:

1. An AMC platform change touching the audience minimums, the lookback, template availability, or the schema: P2/P3 re-document and P12 re-baseline before the next cycle, and this SOP increments if a procedure's field references move;
2. A v4.0 increment that admits any AM-1 to AM-5 item: Element 5 re-validates against the new vocabulary and the affected register states retire into real verdicts before the next cycle runs;
3. Ledger dependency D5 (Marketing Stream) landing: Q5's consumer relationship changes from deferral evidence to L10 deployment evidence, Element 12 row I-5 rewrites, and #45 opens its build;

4. A Build-or-Defer Memo v1.0 Section 4 re-evaluation trigger firing (an audience class hitting SD reach limits, a retargeting economics case, or a video pipeline plus material Q1 upper-funnel share): #21 reopens and this SOP's audience-deployment scope re-verifies;
5. The #17-19 playbook incrementing in a way that moves its SD audience-entry seam: Element 12 rows I-3 and I-4 re-pin;
6. The memo's 5 quarter-one predictions scoring through #6: any MISS routes to this review with its evidence.

Registered in the Supersession Map with the memo as the decision record.

11. Operator Ramp

Built from the Authority Matrix v2.0 Section 1: the rows that name this component's decisions are "AMC cycle (#20)" (CEO I, PPC Lead D, Incoming operator E, BM owner I) and, for the arithmetic standard, the Economics Worked-Example Set v1.0 "Use and Governance" rule that a new operator reproduces all 8 computations cold before taking a scope. The incoming operator holds E on this component: the ramp certifies execution, with every D remaining the PPC Lead's throughout and after the ramp.

Week	Demonstrated competences (checkpoint holder: PPC Lead per the Matrix #20 D cell)
Week 1: observation and artifact familiarity	Reads this SOP, the Build-or-Defer Memo v1.0, and Doctrine v1.0 Section 10; reproduces the 8 Worked-Example Set computations cold, with E6 (CVR-adjusted cross-format cost) and E8 (the TACOS = ACOS x paid share bridge) defended verbally since both bear directly on AMC reads; walks the Query Register and the last 2 filed cycles and names, for every filed set, its consumer SOP and the decision it fed; locates the G2 numbers in v4.0 Section 5 and the audience minimums in Doctrine Section 10 without prompting.
Week 2: supervised execution	Runs 1 full monthly cycle (P2 reconciliation, Q1-Q3 runs, filings with manifest surrogate headers and reading notes) with the PPC Lead reviewing every filing before release; runs 1 audience minimums check end to end (P6 tree, naming, record) with the creation decision staying with the D-holder; runs 1 week of decay watch and drafts the notice for a seeded sub-minimum case.
Week 4: unsupervised execution with audit	Runs the monthly cycle unsupervised; the PPC Lead audits the filed outputs after release against the manifest surrogate checklist (4 header fields, UNREAD counts correct, exclusion windows applied); prepares the quarterly Q5 computation pack (counts, premium, branch recommendation) for the D-holder's gate decision; passes the live-defense session below with 0 threshold values supplied by the examiner.

Live-defense questions (they test the reasoning behind the rules, not recall of the values):

- Why is a 90-click hour-cell UNREAD rather than nearly readable, and what does interpolating it manufacture?
- Why is an audience under the platform minimum not created small, and what would a trickle audience actually measure?
- Why does a query logic change void its own trend history, and what error class does trending across the change reproduce?
- Why are SD decisions click-attributed only, and what does the view-attributed number inflate (Doctrine Section 8)?
- Why can a clean Q5 pattern still not deploy an hourly modifier this quarter (L10's Marketing Stream condition)?
- Why does AMC evidence attach to verdicts instead of issuing them, and what breaks in V5 scoring and the decision log if a second verdict path exists?

12. Interface Contracts

Built from the Cross-Reference Ledger row for #20 (Interfaces cell) cross-checked against Master PPC Data Architecture v1.0 for artifact names and tab locations. Every row states the non-arrival consequence, because that is the clause the corpus has been missing (conflict C5).

ID	Sender, receiver	Artifact and timing	Non-arrival consequence
I-1	Amazon AMC platform to this SOP	Query results and schema, at every run	Cycle files late with the gap noted; consumers read the prior filed set marked STALE; no interpolation, no reconstruction.
I-2	This SOP to #17-19	Monthly filed Q1-Q3 sets, before the playbook M6 check reads Account Rollup A4	The split check runs on console-only evidence and records the absence; no rebalance may cite AMC that month.
I-3	This SOP to #17-19 playbook	Audience records for SD entry, on creation	SD audiences run on console remarketing pools only, the standing clause the Ledger row for #17 already carries.
I-4	#17-19 playbook to this SOP	Deployed-audience campaign status, weekly	The decay watch treats every live-unknown audience as LIVE (conservative) and the status gap logs to the monthly review.

ID	Sender, receiver	Artifact and timing	Non-arrival consequence
I-5	This SOP to the #45 deferral owner (PPC Lead)	Quarterly Q5 gate answer, inside the quarter	The deferral continues by default; no dayparting evaluation may open on optics; the planning sheets stay empty.
I-6	#32 to this SOP	Armed event windows via T5 trend-exclusion flags, at arming	G10 is unsatisfiable for the affected span: trend claims for it are suppressed to observations and marked as such on the filed set.
I-7	v4.0 owner (#8, PPC Lead) to this SOP	Admitted AMC rules at each v4.0 increment	Element 5 stays a zero-verdict mapping, the register states stand, and AM-1 to AM-5 remain open on the amendment list.
I-8	This SOP to #6 V5 scoring	The memo's 5 predictions at the first quarter, plus every V1 prediction on gate answers and audience creations, at review dates	Per v4.0 Section 14 rule V1, unscored actions do not accumulate; the quarterly review flags the scoring gap and holds the memo's scorecard open.
I-9	This SOP to #39 Data Operations	Request for an AMC field-group row in the T7 Manifest, plus the retention policy line; standing until assigned	Filed sets self-carry the manifest surrogate header; consumer verdicts naming missing AMC evidence downgrade to flags per G8.
I-10	#46 to this SOP	Access Register rows for every instance seat, on role fill and change	An unregistered live seat suspends on discovery per #46; the cycle halts until the seat registers.
I-11	This SOP to T6 and Account Rollup A3	Decision-log rows and their V5 scores, each cycle	The Monday audit's first step (scoring last cycle's rows) flags the gap; the affected predictions score late and the delay is recorded.

Closing: Amendment Items and Change Note

v4.1 amendment items logged this build (numbers this SOP needs and v4.0 does not contain; none written locally):

- AM-1: admission (or a standing register-vocabulary exception per rule V3) for the AMC operating states in Section 5.2, retiring v1.0's 6 local verdict strings permanently.
- AM-2: pool-entry floor and lookback windows for AMC and SD audiences as decision thresholds (the platform's 2,000 / 500 minimums are constraints, not policy; same gap the #17-19 rebuild logged).
- AM-3: the Q5 materiality bar: what CVR premium at what readable-cell coverage opens a dayparting evaluation.
- AM-4: the P2 reconciliation divergence tolerance between AMC attributed totals and console totals.
- AM-5: the Element 2 numeric lag table from signal to procedure start (mirrors the #17-19 log).

Open items routed to other owners (not v4.1): to #39: the T7 field-group assignment for AMC filed sets and the retention policy line (I-9). To the close-out: this component's Ledger row, returned with this build.

v2.0 change note: supersedes v1.0 in full; procedures expanded 6 to 12 (adds P2 reconciliation as a standing monthly check, splits reading rules, audience construction, handoff, and decay into P5-P8, adds P10 named-evidence runs and P12 versioning as standalone procedures); the Section 5 verdict table retired and rebuilt as a zero-verdict mapping with 3 outward evidence rows and a register-state table; failure modes expanded 6 to 10 (adds F7-F10); worked examples expanded 1 to 4; Elements 11 and 12 added from the Authority Matrix and the Ledger Interfaces column per conflict C5; v1.0's source attributions "the v4.0 conversion minimum" and "the v4.0 lookback windows" corrected: the former resolves to v4.0 Section 5 gate G2 and is carried, the latter resolves to nothing in v4.0 and rides AM-2.

Companions: Master PPC Decision Criteria System v4.0 (Sections 5, 8.10, 9, 10, 12, 13.1, 13.2, 14 cited herein) · Keyword Decision Record Template v3.0 (closed vocabulary) · AMC Build-or-Defer Memo v1.0 (decision record) · Strategic Doctrine v1.0 Sections 7, 8, 10 (the why and the platform facts) · SB/SBV/SD Playbook #17-19, current version per its Ledger row (primary consumer) · Master PPC Data Architecture v1.0 and the Command Workbook READ ME (instrument mechanics) · PPC Authority Matrix v2.0 (checkpoints) · Cross-Reference Ledger (ownership row #20).

Inspiratek & Ecotero LLC | Confidential | Page 11 of 11