

SOP-22: Launch-Stage PPC Operations and the Launch-to-KDR Handoff

Version 2.0 · 2026-08-07 · Owner: PPC Lead for PPC execution, jointly with the launch department for the plan this SOP executes inside · Supersedes SOP-22 v1.0 (2026-07-24) in full · Absorbs the scheduled Launch-to-KDR Handoff Map deliverable · Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds, and every launch-stage value v4.0 does not govern is named in Section 13 rather than written into the procedure

1. Verdict and Scope

VERDICT: launch-stage PPC buys an organic flywheel with sequenced, batched, deliberately unprofitable traffic, and it is the only stage in the system where the governing threshold set is incomplete, so the discipline that replaces the missing numbers is artifact verification: every gate in this stage passes on a named artifact held by a named role, never on a judgment that a condition "looks met." Two failure classes account for the stage's cost: deploying every campaign type at once, which converts a sequenced purchase into unpriced risk; and the handoff seam, where a product graduating with no co-signed packet arrives in the standing cycle blind and the standing system repays the launch spend to relearn what launch already knew. Version 2.0 expands v1.0 from 7 procedures to 13, from 6 failure modes to 12, from 1 worked example to 4, from a 5-row verdict table of locally invented strings to the full 25-verdict closed vocabulary quoted verbatim, and adds Element 11 operator ramp and Element 12 interface contracts, which no document in this library previously carried.

1.1 What this SOP owns

PPC execution inside the launch department's plan, across seven surfaces: the Phase 1 standup and its batch-1 deployment session; the honeymoon daily block; the batch progression pass and its envelope-neutral funding rule; the Phase 2 discovery entry and its condition-artifact verification; the product-type overlay that decides when placement premiums may be tested; the launch weekly audit interface; and the Launch-to-KDR Handoff Map that moves a graduating product into the standing decision system with a zero-day orphan window. The stage's closing artifact is the co-signed handoff line in the Command Workbook T6 Decision Log with the first standing cascade attached.

1.2 The phase-condition dependency, stated before it is relied on

SOP-22 v1.0 and the Cross-Reference Ledger row for component #22 both direct phase definitions and graduation conditions to "v4.0 Sections 3 and 9." In the current Criteria System v4.0, Section 3 is "Layer M: The Master Metric Criteria Dictionary (F1-F16)" and Section 9 is "Layer 4: The Lever Library (L1-L10)". Neither carries phase definitions, Phase 2 entry conditions, or launch-stage graduation conditions. Section 11 "Layer 6: Transitions, Stages, Modes (T1-T8) and Recovery Mode (R1-R6)" rule T2 states that "the Launch/Scale graduation criteria and Phase 1-4 model stand" and names the five dimensions they are measured on (profitability, organic sales %, review posture, CVR stability, TACOS vs CM2) without values. Section 17.2 "Supersession note (per V3)" assigns the stage and phase framework to the Master PPC Decision Criteria System v1.0 of March 2026, which is not held in this corpus.

This SOP therefore treats the phase and graduation frameworks the way SOP-27 is required to treat the LTSF tier declaration under conflict C6: it executes fully on what PPC does once a condition is declared met, it never infers the condition values, and it names the declaring artifact and its owner at every gate. Sixteen amendment items in Section 13 are the register of what v4.1 owes this component. Until they land, every phase and graduation gate in this document passes on the launch department's plan values with the source and version recorded in the phase-advance record, and the recorded source is what makes the decision auditable later.

1.3 Boundaries with adjacent components

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, gate, scenario, lever limit, sizing rule and verdict string. This SOP executes; it never decides. A threshold appearing in this SOP instead of v4.0 is a defect, and a launch-stage value v4.0 does not carry appears in Section 13, never in a procedure.
The launch department	Owns the launch plan, strategy, audits, and the keyword research producing the batch lists with their indexing status and product-type classification. This SOP owns PPC execution inside that plan and never authors the plan.
#23 Ranking push	Owns push sizing mathematics, the DSTR and required-clicks worksheet, pacing on funded pushes, and taper trigger execution. Launch runs batched mini-worksheets under its rules and never derives its own sizing method.
#12 Exact operations	Owns exact creation, naming, class assignment, taper execution and duplicate consolidation. Every batch and Halo exact stands up through SOP-12 P1 or P2.
#13 / #14 / #15	Own phrase, broad and auto mechanics. Phase 2 discovery stands up through them; this SOP decides only when the entry is verified, not how the campaigns are built.
#29 Negation architecture	Owns negation doctrine, wall-level selection and the wall audit. This SOP requires walls at every structure event in the same session and records the list; it never selects the level.
#30 Harvest and graduation	Decides which discovery terms graduate. Launch products enter its weekly pass from the first Monday of Phase 2 like any other product.
#28 Placement and modifiers	Owns ladder derivation, exit-bid discovery and modifier sizing. The P8 overlay decides when a placement test is legal to open; #28 sizes it.
#38 Bulk mechanics	Owns the staging contract, upload and the T+1 verification pass. Launch urgency never bypasses staging; a same-day deployment still receives the full pass.
#32 Deal and event windows	Owns the event plan, the armed-window declaration and reconciliation. Launches do not go live inside an armed window on the parent brand unless the plan names the launch.

Adjacent component	The seam, stated so it is owned
#33 Sibling arbitration	Owns the declaration on shared terms. A launch product sharing a head term with a live sibling is a portfolio decision before it is a launch decision.
#31 Budget and envelope	Owns the funding queue and envelope law. Per conflict C3 it has no procedural document, so this SOP references the instrument layer instead: T0 zone H and Account Rollup A2. No #31 procedure number is cited anywhere here.
#41 Margin table	Owns the per-variation margin refresh and freshness state. Launch runs on a provisional spine by design and receives the upgrade at handoff; it never authors margin values.
#36 Weekly audit	Receives the product at co-sign. From that moment the standing T0 to T2 cascade governs and this SOP is silent.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum lag between the signal appearing there and the procedure starting. A trigger detected later than its allowance is logged in T6 as a process miss with the delay in days, because an untracked late start is how a sequence silently becomes a simultaneous deployment.

Trigger	Detecting artifact	Max lag	Procedure
Launch scheduled	The launch department's dated plan with the batch-1 list attached	Same day	P1 readiness audit opens
Readiness clear	P1 checklist with every row evidenced, countersigned by the launch operator	2 days	P2 workbook pre-clone
Product go-live	Offer live and buyable in the console; batch-1 list indexed for the D2 preferred variation	Same day	P3 Phase 1 standup, day 0
Honeymoon open	The standup record with the dated honeymoon exit written into it	Daily	P4 honeymoon daily block
Batch mark reached	The batch calendar entry in T5 Product Log; the cadence is the launch plan's, recorded not derived	1 day	P5 batch progression pass
Progression closed	P5 log with terms held, continued and failed, plus available funding computed	1 day	P6 batch N+1 deployment
Phase 2 requested	Launch operator request with the condition artifact set attached	2 days	P7 condition verification
Product type declared	The launch plan's subjective or non-subjective classification	At standup	P8 overlay, rides every placement and discovery decision
Weekly audit lands	The launch department's weekly audit document	1 day	P9 interface exchange
G6 flips YELLOW or RED	T2 block K2 Inventory Band and Days to Stockout on the targeted variation	Same day on RED	P4 step 5 gate response, re-point through SOP-12 P10
Graduation claimed	Launch operator declaration with the graduation artifact set	2 days	P10 verification and packet
Packet complete	P10 packet checklist at 11 of 11 fields present	Same week	P11 workbook standup, P12 cascade and co-sign
Co-sign lands	The countersigned handoff line in T6 with the first cascade attached	30 days	P13 post-handoff calibration
T7 drops DEGRADED	T7 Manifest completeness below 85% of gate-relevant field groups	Same day	All verdict-issuing procedures downgrade to flags

3. Preconditions and Gates

Every gate cites its manifest field group from Master PPC Data Architecture v1.0 Section 9 tab T7, its identifier from Criteria System v4.0 Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass", and the observable that proves it passed. An assertion is not an observable. Per v4.0 Section 12 rule P1 gates outrank diagnosis and diagnosis outranks signals, so a failing gate stops the dependent procedure even when the launch calendar disagrees.

Gate naming. Conflict C9 records that the operating corpus and v4.0 assign different meanings to G9, G10 and G12. Every gate in this document is named by its v4.0 number and name, with the corpus alias given at first use, so no reader resolves the wrong gate.

Gate	Manifest ID	Requirement before the stage proceeds	Observable that proves it passed
G5 Indexing and listing placement	B1	Every batch-1 term indexed for the D2 preferred variation, with title or bullet placement on any term carrying a rank objective. No bid level ranks an unindexed term.	The indexing check result per term in T2 block K2 columns Indexed? and Listing placement, dated within the batch window. Backend-only terms carry a listing-queue ticket number.
G6 Inventory adequacy	D2	Launch cover on the D2 preferred variation sized against the full batching plan, not against batch 1 alone. RED is a projected stockout at or under 14 days.	T2 block K2 Inventory Band reading GREEN with Days to Stockout stated, plus the ranked backup chain present in the Keyword-to-Variation Map.

Gate	Manifest ID	Requirement before the stage proceeds	Observable that proves it passed
G9 Price stability corpus alias: margin freshness	D1	No price change on the targeted variation within the trailing 7 days, or every CVR and ACOS read in the window is marked PROVISIONAL. The launch economics spine is separately tagged PROVISIONAL by design until per-variation margin lands.	T5 Product Log carrying either no price-move row in the trailing 7 days or a dated row; T0 zone C Margin table version cell showing the version and its provisional state.
G12 MKL hygiene corpus alias: event freeze	C1-C8	Every batch term exists in the current Master Keyword List with a syntax tag and a relevancy score. An MKL stale past 30 days degrades the packet manifest.	T2 block K1 Syntax tags and Relevancy populated for every batch row, with the MKL refresh date in T7 inside 30 days.
G12 Event window v4.0 number: G10 event exclusion	-	The launch does not go live inside an armed window on the parent brand unless the event plan names it. Event weeks and the 14 days following are excluded from trend verdicts.	T5 Product Log deal-window row with its dates, cross-checked against the #32 event plan; absence of an armed window is itself the passing observable and is recorded as such.
G11 Sibling conflict	D6	Any batch term shared with a live sibling product carries a declared lead product before the launch term is funded. Two family pushes on one term bid against each other with the family's own money.	The ownership declaration row in Account Rollup A1 Family Boards, with the assignment basis stated; T2 block K8 Term owner (family) populated on every shared row.
G7 Envelope and pacing	A3	Campaign budgets reconcile to the declared launch envelope. Per v4.0 Section 8.1 scenario S-A6, budgets summing above 110% of the declared weekly envelope trigger mechanical reconciliation the same day.	T0 zone B Weekly PPC spend against the WeeklyEnvelope named cell, and the T3 Campaign Board daily budget column summing inside it.
G8 Packet completeness	All	The manifest reports present, stale and missing before any verdict is issued. Per v4.0 Section 3 family F15, DEGRADED is below 85% of gate-relevant field groups and produces flags only.	The T7 Manifest completeness rollup stamped on the cycle header, with the Verdict Classes Blocked column filled for every non-PRESENT row.
G2 Sample sufficiency	A4	Bid reads need at least 15 clicks over 7 days; CVR verdicts need at least 100 clicks; CTR verdicts need at least 1,000 impressions. Below threshold the verdict is a capped test, never a rate judgment.	T2 block K2 Clicks 7/30/60d and Sample state bucket per row. At launch this gate fails on most batch rows by construction, and the failure is the expected reading, not a defect.
G3 Budget truncation	A3	In-budget below 70% of the day marks every rate on the campaign TRUNCATED and the budget lever is evaluated before any bid conclusion.	T2 block K2 In-budget % on the worst hosting campaign, and the T3 In-budget bucket column.
G4 Eligibility	A7	The advertised SKU is eligible. An eligibility flip halts all lever activity for the target the same day and opens the listing task.	T3 group C2 Eligibility status and reason, checked at standup and in every daily block.
G1 Validation quarantine	A5	Impossible metrics quarantine the row: a flag is produced, no recommendation is issued, and a data-refresh task opens the same day.	T2 block K2 Validation state and the zero-flag quarantine column in block K5.

Gate interaction requiring escalation, not local resolution. v4.0 Section 3 family F2 sets the Launch TACOS reference band at 15-25% with deviation classes WITHIN at or under 1.2 times the band, ELEVATED at 1.2-1.5 times, BREACH at 1.5-2 times, and CODE RED above 2 times or above 1.5 times for 2 validated weeks. A product in its first weeks of launch carries almost no organic revenue, so its TACOS reads several multiples of 25% by construction and trips CODE RED mechanically. v4.0 Section 11 rule T8 sets Recovery Mode above stage defaults with no launch exemption. This SOP does not suspend a v4.0 rule on its own authority: the reading is recorded, flagged, and routed to the PPC Lead under the Authority Matrix v2.0 row "v4.0 threshold changes" where the CEO decides and the PPC Lead executes and reviews. Amendment item AM-16.

4. Step-Level Procedures

Thirteen procedures cover the stage from the readiness audit before go-live to the calibration pass 30 days after handoff. Time estimates assume the Raw Pull Pack has landed and the workbook is open. Every procedure writes to the workbook; one that leaves no row behind did not run.

P1. Pre-standup readiness audit

Time: 90 minutes. Runs day -5 to day -1. Owner: launch operator executes, PPC Lead reviews. Output: the countersigned readiness checklist.

1. Extract five fields from the launch plan into a working sheet: the batch-1 term list, the product-type classification (subjective or non-subjective), the declared envelope in dollars per week, the D2 preferred variation, and the dated honeymoon exit. A plan missing any of the five returns to the launch department; this SOP does not supply a missing plan field.
2. Run the G5 indexing check on every batch-1 term against the D2 preferred variation; record per term in T2 block K2 columns Indexed? and Listing placement. Terms reading Not indexed do not enter batch 1 at any bid: route to the Brand Management indexing owner and hold for a later batch, ticket reference into the T2 block K7 Situation text field.
3. Run the G12 MKL hygiene check: every batch-1 term carries a syntax tag and relevancy score in T2 block K1, and the T7 MKL refresh date is inside 30 days. An untagged term routes to MKL maintenance, not to a campaign.
4. Run the G6 read on the D2 preferred variation and its ranked backup chain; record days of cover against the full batching horizon, not batch 1 alone. YELLOW at readiness is a plan conversation before it is a PPC one.
5. Run the G11 sibling check against Account Rollup A1 Family Boards. Any term carrying a live objective on a sibling needs a declared lead product before it is funded here.
6. Run the event-window check against the #32 plan. Record it in T5 as a deal-window row even when no window exists, because a recorded absence is what makes the launch date defensible later.
7. Assemble the checklist with one evidence line per row, countersign with the launch operator, file it as the P1 artifact. Failing rows go to step 8.
8. Failing rows produce a dated blocker list with an owner each. The standup date moves or the affected terms drop out of batch 1, and the drop is logged before the standup proceeds.

P2. Workbook pre-clone and the provisional economics spine

Time: 60 minutes. Runs day -3. Owner: PPC Lead, with Data Operations wiring feeds. Output: a live Command Workbook with T0 zone C set and T7 stamped.

1. Clone the PPC Command Workbook template and name it for the product code per the template READ ME. The workbook exists before the campaigns do, so no launch decision is made outside an instrument.
2. Set the T0 zone C spine in cells C18 to C24: named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC, WeeklyEnvelope. Every ceiling references these, so they are set once and never duplicated downstream.
3. Derive, never type, the three chain values, per v4.0 Section 10 rule M1: BE_ACOS equals MarginPerOrder divided by AOV; TargetACOS is 50% of BE_ACOS; MaxAccACOS is 75%. A typed value in any of the three is a build defect.
4. Tag the margin version cell PROVISIONAL. A launch product has no realized per-variation margin, so the state rides every ceiling-referencing action until P11 upgrades it.
5. Load the batch-1 list and the full market keyword list into T2. The full list, not the funded slice: coverage state is a column, so unfunded demand is a filter rather than a research task later.
6. Set T7 status per Catalog field group. At launch, A1 four-period comparability and B9 week-over-week deltas both read MISSING because no prior periods exist. Record both with their blocked verdict classes rather than leaving them blank.
7. Stamp the T7 completeness rollup on the cycle header. Below 85% the launch proceeds but every downstream verdict issues as a flag per G8, and the standup record says so in plain words.

P3. Phase 1 standup, day 0

Time: 3 to 4 hours for a batch of 8 to 12 terms. Owner: launch operator decides and executes per Authority Matrix v2.0 row "Launch-stage verdicts + batches"; human-mandatory throughout. Output: the standup record, the T3 registry, the wall lists, and one T6 prediction row per term.

1. Confirm the P1 checklist is countersigned and the P2 workbook live. A standup without both is the F1 condition.
2. Deploy SP Exact single-keyword campaigns on batch-1 terms only, through SOP-12 P1. Fixed bids; one campaign, one keyword, one ad group; advertised SKU equal to the D2 preferred variation and never the parent by default.
3. Deploy the Halo exacts named in the plan through the same procedure, writing each retirement condition into its T3 row at creation. A Halo created without one is the F3 condition on its creation date, not on the day it is discovered.
4. Deploy nothing else: no broad, phrase, auto, brand format or conquest. Per v4.0 Section 6 rule O8 the launch stage carries a discovery-heavy mix target and Phase 2 is where that mix is earned. The sequence is the risk control.
5. Build the negation walls in the same session per #29 and record the list in each T3 row. That list is what the #29 audit measures against later.
6. Size each term with a mini-worksheet under SOP-23 P2 logic: target rank, DSTR for that rank, validated or modelled CVR, Expected CPC, required clicks per day, required budget per day, and its share of the loss budget. Version it into T2 block K7 columns DSTR, Required clicks/day, Required budget \$/wk and Loss ceiling \$/wk.
7. Assign one objective per term in T2 block K7 from the six-value dropdown: rank-sized terms take Ranking, capped tests take Discovery. An untagged row cannot be evaluated and enters the assignment worklist per v4.0 Section 6 rule O1.
8. Reconcile the sum of daily budgets against the WeeklyEnvelope named cell before staging. Per v4.0 Section 8.1 scenario S-A6 a sum above 110% of the declared envelope reconciles the same day, worst value first.
9. Set placement modifiers to zero on all three groups. The modifier is earned on evidence per #28, and at day 0 none exists.
10. Stage every campaign as bulk rows through #38, reason string carrying the objective and worksheet reference. Same-day deployment is permitted on the release date and still takes the full T+1 verification pass.
11. Write the plan's dated honeymoon exit into the standup record and into T5 as a targeting-change row, so the exit is a calendar fact rather than a memory.

12. Log one T6 prediction row per term: metric, direction, magnitude, horizon, review date. A term deployed without one cannot be scored, and an unscored action does not accumulate per v4.0 Section 14 rule V1.

P4. The honeymoon daily block

Time: 20 minutes per day. Owner: launch operator, PPC Lead informed. Output: a daily line in T6 for any intervention, and the gate-response record.

1. Read the four gates that can move in a day before any metric: G4 eligibility from T3 group C2, G6 inventory band and G3 in-budget from T2 block K2, G1 validation state. Gates before metrics, daily, without exception.
2. Read rank velocity and CVR per term from the DataRova daily series and T2 block K5 cross-type totals. Record; do not act yet.
3. Withhold efficiency tightening. Bids do not step down on ACOS optics while the window is open, because the window buys relevance signals that outlive it and a cut inside it discards the purchase. A plan rule, not a threshold; it holds to the dated exit written at standup.
4. Route a CVR materially below the launch model to the launch department as a listing or offer read before any lever moves. Per v4.0 Section 7 rule D6 the conversion diagnosis opens with the Price Position Index, and rule D8's funnel drop point names the owner. A conversion deficit is not a bidding problem at any stage.
5. Execute gate responses immediately and separately from the honeymoon rule. A G6 flip to RED runs the same-day taper and re-point under v4.0 Section 8.3 scenario S-I3 through SOP-12 P10, baseline snapshot written first. Per v4.0 Section 12 rule P1 gates outrank every lower layer, so a gate response is never withheld.
6. Execute the G3 budget lever same day on any truncation. A starved campaign's rates describe its worst hours, and at launch that misreading corrupts the batch model.
7. Write every intervention to T6 with its gate reference. Pacing and gate interventions do not change bids; bid decisions live in the progression pass and in verdict-carried ladders only.
8. At the dated exit, close the window in T5 and move to normal scenario governance. The exit is logged, not announced.

P5. Batch progression pass

Time: 2 hours. Runs at each batch mark on the launch plan's cadence. Owner: launch operator decides, PPC Lead reviews, BM owner reviews. Output: the progression log with terms held, continued and failed; the funding computation for the next batch.

1. Score last cycle first: read the T6 rows whose review dates land in this window and record V5 scores as HIT, PARTIAL, MISS or WRONG-DIRECTION. The pass opens here, not on new data.
2. Verify gates per row before any metric read: G2 sample state, G3 truncation, G6 inventory, G1 validation, G9 price stability. A failing gate executes its lever and defers the metric read on that row to the next window.
3. Classify every live term using the tree below. The classification, not the operator's impression, drives the funding computation.
4. Compute funding for the next batch as released budget plus unspent envelope. Released budget comes only from HELD terms and terms capped at maintenance by a fix-first verdict; unspent envelope is the declared envelope for the window minus actual spend. Batching is envelope-neutral by construction: no batch is funded by an envelope increase, which is a product decision with margin evidence and never emerges from a batch pass.
5. A term that reached target but has not held it for the interval v4.0 governs is not HELD and releases no budget. Per v4.0 Section 11 rule T1 the post-objective transition fires at 2 weeks at or under target rank. The launch hold read that would qualify a term earlier is AM-5 and is not substituted with a local number.
6. Record the progression log: terms held with released dollars, terms continued at full pressure, terms failed with the read and its routed owner, computed funding for batch N+1, and the step-1 scores.
7. Where fewer than half the terms delivered their required clicks, check the envelope and the per-term split before checking bids. The correction is fewer terms at full pressure, never more at lower pressure.

Batch term classification tree, applied in order; first match wins.

1. Does the row fail a gate? → **GATED**. Execute the gate lever, defer the metric read, release no budget.
2. Is the term at or under its target rank, and has it held there for the interval v4.0 Section 11 rule T1 governs? → **HELD**. Step pressure down within the v4.0 Section 9 rule L2 limits, release the freed budget to the next batch.
3. Did the term deliver at or above its required clicks per day with rank velocity intact? → **CONTINUED**. No change; the worksheet still governs; release no budget.
4. Did the term deliver its required clicks with CVR at or above benchmark and rank velocity at or near zero across the window? → **FAILED, over-competitive**. Route to v4.0 Section 8.4 scenario S-R5, stop-loss.
5. Did the term deliver at least 100 clicks with CVR below benchmark? → **FAILED, conversion**. Route to v4.0 Section 8.4 scenario S-R4, fix-first, cap at maintenance, release the difference.
6. Did the term deliver under its required clicks with CVR at or above benchmark? → **CONTINUED, traffic deficit**. Route to v4.0 Section 8.4 scenario S-R3, feed the named lever only.
7. Is the sample below the v4.0 Section 5 gate G2 line for the read the verdict needs? → **CONTINUED, capped test**. No rate judgment is legal; the test runs to its click quota.

P6. Batch N+1 deployment

Time: 2 to 3 hours. Runs within 1 day of P5 closing. Owner: launch operator; human-mandatory. Output: new T3 rows, wall lists, mini-worksheets, T6 predictions.

1. Take the batch list for this window from the plan in priority order. The order is the plan's; this SOP does not re-rank it.
2. Size the batch to the funding computed at P5 step 4, not to the list length. Terms that do not fit wait; batch size is an output of available funding, never an input.
3. Where funding supports fewer ranking-sized terms than the plan proposes, fund the priority subset at ranking size and place the rest at the capped test level per v4.0 Section 8.9 scenario S-H7, a dated 50-click capped test with success criteria. Fewer funded terms plus dated tests is legal; spreading the same money thinner across all of them is the F4 condition.
4. Run the full P1 gate set on every new term: indexing, MKL hygiene, inventory, sibling, event. Batch 3 receives the checks batch 1 received.
5. Stand up through SOP-12 P1 or P2 by class, build walls in the same session per #29, record the lists in T3.
6. Write a mini-worksheet per term into T2 block K7 and a prediction into T6, exactly as at standup.

7. Stage through #38 and record the deployment date. The window's progression log closes when T+1 verification returns clean.

P7. Phase 2 discovery entry, on verified conditions

Time: 2 hours for verification, plus the standup time the receiving SOPs require. Owner: launch operator decides, PPC Lead reviews. Output: the phase-advance record with its condition artifact set.

1. Receive the request with its artifact set. Verification runs on artifacts; a stated belief that conditions are met is not one.
2. Record the source of the condition values in the phase-advance record: which plan document, which version, which dated section. Per Section 1.2 the Phase 2 values are not carried by v4.0, so the recorded source is what makes the advance auditable and lets v4.1 reconcile it later.
3. Verify the T7 manifest first. Per v4.0 Section 12 rule P8 data-quality precedence beats every other consideration including urgency, and the correct urgent action on bad data is fixing it the same day. A DEGRADED manifest refuses the advance regardless of how the conditions read.
4. Verify each condition against its artifact and record the reading beside it. An unverifiable condition is a blocker with an owner and a date, never a rounded pass.
5. On a clean verification, stand up the Phase 2 surfaces through their owning SOPs: family broad through SOP-14 P1 at capped budgets and down-only bidding; auto close match through SOP-15 P2; phrase only where SOP-13's proven-root preconditions are met. This SOP verifies the entry; it does not build the campaigns.
6. Register with #30 harvest from the first Monday after entry. Launch products harvest like any other from day one of Phase 2; a discovery surface with no harvest destination is a slow leak by construction.
7. Set the discovery objective per v4.0 Section 6 rule O6, which judges Discovery on cost per discovered converter, waste percentage and graduation rate, never on another objective's metrics.
8. Where verification fails, refuse in writing with the failing condition, its artifact, its owner and the re-test date. A refused advance is a work item with a name on it; a deferred advance with no owner is how calendar impatience wins.

Phase 2 entry decision tree.

1. Does T7 read at or above 85% of gate-relevant field groups? If no → **REFUSE**, issue the data-repair task, re-test on the next refresh.
2. Is the condition value set present, sourced and dated in the launch plan? If no → **REFUSE**, escalate to the PPC Lead as an AM-1 dependency, re-test when the source lands.
3. Does every named condition verify against its artifact? If no → **REFUSE**, name the failing condition and its owner.
4. Is an event window armed on the parent brand? If yes and the event plan does not name this advance → **DEFER** to the window close plus the exclusion period.
5. All four clear → **ADVANCE**. Stand up through SOP-13/14/15, register with #30, write the phase-advance record.

P8. The product-type overlay

Time: rides every placement and discovery decision; no standalone session. Owner: launch operator. Output: the classification recorded once in T5 and applied per decision.

1. Take the classification from the plan and record it once in T5 as a targeting-change row. It is an input to this SOP, never a judgment made inside it.
2. On non-subjective products, follow the conversion: a top-of-search bid-adjustment test opens only after conversion evidence exists on the term. Non-subjective demand converts on specification match, so the premium is proven before it is paid. Until then modifiers stay at zero and #28 has nothing to size.
3. On subjective products, test top-of-search and rest-of-search together, because style-driven browsing converts across placements. The keyword universe is smaller, so discovery surfaces enter earlier within Phase 2 once P7 verifies the entry.
4. Apply v4.0 Section 3 family F9 to every placement read once evidence exists: a top-of-search premium survives only where its conversion beats rest-of-search inside the CPA ceiling, and a mix shift above 15 points across 4 weeks with no logged change routes to structural inspection before any performance explanation is accepted.
5. Waste-limit interaction: v4.0 Section 3 family F11 sets 10% standard and 40% discovery with no product-type split, while the KDR Template v3.0 "Bands & Verdicts Reference" sheet carries a product-type split v4.0 does not govern. Apply the v4.0 limits and log the divergence as AM-10.

P9. Launch weekly audit interface

Time: 60 minutes. Runs within 1 day of the launch audit landing. Owner: launch operator delivers, launch department chairs, BM owner reviews. Output: the keyword-level table and the staged actions.

1. The launch department's weekly audit is the governing review in this stage; PPC delivers into it and holds no competing review.
2. Deliver the keyword-level table per term with six columns: rank against target with history, clicks delivered against required, CVR against the launch model with the benchmark source named, spend against batch budget, gate state, and the current P5 classification.
3. Name the benchmark source on every CVR entry per v4.0 Section 7 rule D3, which ranks SQP brand-versus-market first, file reference targets second, account norm third, and lowers confidence on any fallback. A benchmark without a named source is not auditable.
4. Carry the flags honestly: every provisional read, fallback benchmark and manifest gap appears in the table in the words the workbook computed. Hiding a fallback is a process violation under v4.0 Section 13.3 rule J7.
5. Stage actions arising from the audit through #38 like any other. Launch urgency never bypasses the session record, and an action deployed outside staging cannot be verified at T+1.
6. Record the audit reference in T5 so the progression log and the audit trail reconcile later.

P10. Graduation verification and the handoff packet

Time: 3 hours. Runs within 2 days of the graduation claim. Owner: launch operator assembles, PPC Lead reviews. Output: the 11-field handoff packet.

1. Verify the graduation conditions against artifacts under the same rule as every other gate: record the source document, version and date of the values, because v4.0 Section 11 rule T2 names the five dimensions without values and the values come from the plan. AM-2.
2. Assemble the packet against this 11-field checklist and count completeness as fields present over fields required. Below 11 of 11 it does not proceed to P11.

#	Packet field	Content and the standard it must meet
1	Batch history with scored predictions	Every progression window with terms held, continued and failed, each prediction carrying its V5 score. Unscored rows are closed before handoff.
2	Per-term validated CVR	Only terms at or above the v4.0 gate G2 line of 100 clicks carry a validated CVR. Terms below it carry their click count and no CVR, so the standing system never inherits a fabricated benchmark.
3	Discovered CPC ranges per term	Observed CPC range with the window dates, feeding the Expected CPC reference the standing ladders step against.
4	Rank states against targets	Current rank, target rank, best rank with its date, and velocity, so the standing system can class each term under v4.0 Section 7 rule D10 as recovery or never-ranked.
5	Campaign registry with objectives	Every campaign with its type, objective tag, role and advertised SKU, matching the T3 Campaign Board exactly.
6	Negative and wall lists	Per campaign, the wall list recorded at each structure event, so the #29 wall audit has its baseline on day one of standing operation.
7	Open tests with dates	Every test with its click-count progress and review date. A test without a review date drifts unattended.
8	Halo retirement states	Each Halo exact with its written retirement condition and whether the condition is met.
9	Inventory and variation state	The D2 preferred variation, its ranked backup chain, current band, and any pre-OOS baselines written during the stage.
10	Realized economics	Blended validated CVR across qualifying terms, realized AOV, and the recomputed maximum profitable CPC, so the standing ceiling is a measured value rather than the launch model.
11	Launch audit trail reference	The pointer to the launch department's audit series, so a later question about the stage has a destination.

1. Where a field is short, name the missing element, its owner and its date. The packet completes before the handoff proceeds; it is not waived for schedule.
2. Recompute maximum profitable CPC from realized data per v4.0 Section 3 family F1, margin dollars times CVR. A launch model CVR that overstated reality leaves an overstated ceiling, and every ceiling-referencing row inherits the error.

P11. Workbook standup and reclassification

Time: 2 to 3 hours. Runs in the same week the packet completes. Owner: PPC Lead, launch operator executing. Output: the standing workbook with T0 upgraded and every campaign reclassified.

1. Upgrade the T0 zone C spine from the launch model to realized data: enter AOV and MarginPerOrder in cells C18 and C19 and the derived chain recomputes. Where per-variation margin now exists, clear the PROVISIONAL tag; where it does not, the tag carries forward and the standing system inherits an honest flag rather than a clean-looking guess.
2. Seed T2 with the launch keyword set plus the full market load, so the standing system sees unfunded demand as a filter from cycle one.
3. Import the launch campaign registry into T3, preserving objective tags, roles and advertised SKUs.
4. Retire Halo exacts whose conditions are met: PAUSED-PENDING with the charter as the stated reason and no re-entry, per v4.0 Section 6 rule O5. Pause, never delete; history stays readable.
5. Retag batch exacts by term state: a term continuing under a rank objective keeps Ranking with its target and horizon; a term at target transitions per v4.0 Section 11 rule T1; a capped term carries its fix-first re-entry criterion.
6. Move discovery campaigns into standing governance under SOP-13, SOP-14 and SOP-15 unchanged. No campaign is rebuilt at handoff: reclassification is tagging, not restructure, and a restructure here destroys the history the packet just captured.
7. Set T7 for the standing cadence: the MISSING rows for four-period comparability and week-over-week deltas begin filling from the first standing cycle, and the manifest says so rather than reading blank.

P12. First standing cascade and the co-sign

Time: 90 minutes inside the Monday cascade. Owner: PPC Lead and launch operator co-sign; human-mandatory for both. Output: the co-signed handoff line in T6.

1. Run the full T0 to T2 cascade on the first standing Monday: T0 posture, then T1 syntax caps, then T2 per-term verdicts. A T2 pass that skipped T0 and T1 is invalid by construction.
2. Issue verdicts from the closed vocabulary only, per Section 5. The first standing cycle is where launch vocabulary stops and standing vocabulary starts; mixing them is what makes a decision log unqueryable.
3. Log every verdict in T6 with its prediction and review date, as any standing cycle would.
4. Co-sign the handoff line: both owners sign the same T6 row, with the packet reference and first cascade attached.
5. Apply the seam rule: until the co-sign exists the launch owner holds the product, after it the standing owner does. There is no window in between by construction, and ownership never transfers by calendar, verbal handover, or the conditions being met.
6. Where the co-sign does not land, ownership stays and the delay escalates. Per Authority Matrix v2.0 standing rule 3 disagreements route up, never sideways. The escalation interval carried by v1.0 does not trace to v4.0 and is AM-14; until it lands the escalation runs on the authority path rather than a local number.

P13. Post-handoff calibration

Time: 90 minutes. Runs 30 days after co-sign. Owner: PPC Lead, launch department informed. Output: the calibration record and any amendment proposals.

1. Score the launch model against realized data on three lines: assumed CVR against blended validated CVR, assumed Expected CPC against discovered ranges, assumed rank velocity against realized velocity per term class.
2. Compute each delta as a percentage and hand all three to the launch department as inputs to the next plan. A launch model never scored stays wrong, and every later launch pays for it.

3. Read the first four standing cycles for orphan symptoms: terms with no objective tag, campaigns with no wall list, tests past review dates, launch predictions never scored. Each is an F6 finding naming the packet field that should have carried it.
4. Where a v4.0 rule produced a reading this stage could not act on, file an amendment proposal under v4.0 Section 14 rule V3, which requires versioned amendment with a supersession note and treats oral re-derivation as a process violation.
5. Close the launch record. The product is standing now and this SOP is silent until the next launch.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds and scenario logic live in v4.0. This SOP issues no verdict of its own and invents no string. Every string below is quoted verbatim from the Keyword Decision Record Template v3.0, sheet "Bands & Verdicts Reference", which carries the 25-item closed vocabulary. Per conflict C10 the v4.0 Section 13.1 "The Closed Verdict Vocabulary" list of 18 diverges on several rows; the resolution is to quote the KDR Template and carry the v4.0 string as a stated cross-map, so neither citation is dropped.

What v1.0 issued and v2.0 does not. SOP-22 v1.0 Section 5 mapped five strings in neither closed vocabulary: DEPLOY BATCH N, HOLD BATCH, PHASE ADVANCE, HALO RETIRE, HANDOFF. Those are gate outcomes with named artifacts, not verdicts. v2.0 expresses them that way: batch deployment is P6, the batch hold is a P5 tree classification, the phase advance is the P7 record, Halo retirement is P11 step 4, the handoff is the P12 co-sign. Conflict C17.

Verdict, quoted verbatim from the Keyword Decision Record Template v3.0 "Bands & Verdicts Reference" sheet	Launch-stage execution	v4.0 Section 13.1 cross-map
QUARANTINE: DATA REFRESH	No change deploys on the row; the refresh task opens same day with an owner; the staged set carries to the next clean window. On the manifest at P7 it refuses the phase advance outright.	FLAG: DATA REPAIR
UNLOCK TRAFFIC	Bids step up per v4.0 Section 8.4 scenario S-R1: 15-25% incrementally when clicks run under the G2 line with budget healthy and inventory GREEN. Common in week 1. No rate judgment until sufficiency.	-
FIX BUDGET (TRUNCATED)	Budget lever only, same day, per P4 step 6; the week's bid read is void. It outranks the progression read, because a truncated term has not been tested.	-
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	The test runs to its click quota at the capped bid; budget assurance only, no bid change. The standard verdict on terms below the G2 line, which at launch is most of them.	EXTEND TEST
FUND TEST (UNFUNDED PRIMARY)	P6 step 3 creation at the capped bid the verdict names, with test status and stop-loss date into T2 block K7. The mechanism v4.0 Section 8.9 scenario S-H7 provides for a relevant term that was underfunded.	FUND DELIBERATE TEST (50-CLICK CAP)
RANK PUSH: FUND SIZED PUSH	P3 or P6 standup at the worksheet's sizing through SOP-12 P1. The worksheet is SOP-23's, the campaign SOP-12's, the batch decision this SOP's.	RANK PUSH AT CONTROLLED CPC
RECOVERY PUSH	Legal only after a stockout inside the stage, sized toward the P4 step 5 baseline. A term with no baseline is never a recovery candidate; it is a cold start.	-
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	The named lever only, on the named campaign, never both in one window. P5 tree branch 6. The most common correct launch verdict and the most commonly mistaken for a conversion problem.	-
FIX CONVERSION FIRST; DEFER PUSH	Cap at maintenance, release the difference to the next batch, route the deficit by funnel drop point to the owner v4.0 Section 7 rule D8 names. In the honeymoon the routing runs but the cap waits for the dated exit.	Same string
PLACEMENT FIX FIRST	Bids hold; #28 executes. On non-subjective products it cannot issue before conversion evidence exists on the term, per P8 step 2.	-
RE-POINT VARIATION	SOP-12 P10 in full, same day, baseline snapshot written first. The CVR baseline resets and the term leaves the progression read until sufficiency refills. Worked in WE-2.	Same string
TAPER: RANK ACHIEVED, STEP BIDS DOWN	SOP-12 P4 in full. Inside launch it fires only where a term has held target for the interval v4.0 Section 11 rule T1 governs; below that it is not tapered and releases no budget.	TAPER TO FLOOR (POST-OBJECTIVE)
HOLD; DEFEND LIGHTLY AT RANK	No change this window. Correct on a term that has touched target but not held long enough to transition. It releases no budget, which is the point: a batch cannot be funded by a rank that has not proven stable.	CONTINUE; VERIFY MOMENTUM
DEFEND AT FLOOR	No bid change; drift from a recorded floor restores in one staged step. Rare inside launch, because floors are discovered post-taper.	Same string
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	Bid edit to the T2 block K6 Max profitable CPC value, single step. The syntax-cap half executes at T1 through the instrument layer. In the honeymoon the bid half waits for the dated exit; the cap does not.	Same string
INCREMENTALITY LADDER	Not issued inside launch: the test measures organic capture against paid spend and a launch product has no organic baseline. Available after co-sign under #35.	-
ADD EXACT + RESTRUCTURE	SOP-12 P1 or P2 for the uncovered term plus steering walls per #29. Inside launch it fires when a Phase 2 surface begins carrying a batch term's clicks, the failure v4.0 Section 4 rule X10 names above 40%.	Same string

Verdict, quoted verbatim from the Keyword Decision Record Template v3.0 "Bands & Verdicts Reference" sheet	Launch-stage execution	v4.0 Section 13.1 cross-map
CONSOLIDATE DUPLICATES	SOP-12 P7 in full, survivor selection human-mandatory. The usual launch cause is a Halo and a batch exact serving one term; the Halo's charter decides which is the duplicate.	-
GRADUATE FROM DISCOVERY	Received from #30 once registered at P7 step 6. Standup through SOP-12 same day, walls same session. Launch products graduate on the same bar as any other and receive no exception.	GRADUATE + STEER
NEGATE (PRE-LOAD / REACTIVE / STEERING)	Owned by #29. This SOP's role is structural: pre-load walls at the P3 standup, walls at every P6 event in the same session, steering walls the day a term graduates. Mode recorded per term in T2 block K8.	CUT / NEGATE
ENTER SBV ARBITRAGE	Not issued inside launch: the inequality compares a live SP exact CPC against a qualified SBV read, and a launch term has neither. Available after co-sign under #17-19.	FORMAT SUBSTITUTION (SBV/SD)
OPEN / ROTATE / EXIT CONQUEST	Not issued inside launch: conquest is a Phase 3-4 surface and the product has not yet earned the comparison it would conquest on. Owned by #16 when it opens.	-
STOP-LOSS: CONCEDE OR DEFER	The term steps to maintenance or PAUSED-PENDING per the verdict text; freed budget returns to the envelope and funds the next batch through P5 step 4, never through an escalation. P5 tree branch 4.	Same string
PAUSED-PENDING (REASON + RE-ENTRY)	Always carries a blocking reason and re-entry criterion in T2 and T3 per v4.0 Section 6 rule O5. The standard launch re-entry is the next batch window with released budget. Also the Halo retirement mechanism at P11 step 4, where the reason is the charter and there is no re-entry.	PAUSE (LEVERS EXHAUSTED)
LEAVE TO HARVEST	No exact stands up; the term runs on discovery economics under #30. Legal only after Phase 2 entry, because before it there is no discovery surface to leave it to.	Same string

One v4.0 verdict, PORTFOLIO CALL: ASSIGN LEAD, has no counterpart string in the Keyword Decision Record Template. Inside this stage the decision it names is executed as the G11 gate response at P1 step 5 and P6 step 4, with the declaration held by #33. Reported as an extension of conflict C10.

6. Worked Examples

No Tier 3 launch-stage economics spine is attached, and the attached Command Workbook B4 v1.0 is a mature-stage product whose spine does not transfer to a launch case. All four examples therefore run on the standard disclosed fictional base from the PPC Economics Worked-Example Set v1.0: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target 17.5%, Max Acceptable 26.3%, launch model CVR 5.2%, Expected CPC \$2.50. Values are illustrative and internally consistent; live values come from T0 cells C18 to C24. Instrument, column and tab references are real.

WE-1 · The clean case: standup sizing, the week-2 progression, and why batch 2 shrank

Setup. The plan delivers 8 batch-1 terms plus 3 Halo exacts, product type non-subjective, declared envelope \$210.00 per day. P1 clears on all six checks. P2 sets the spine: BE_ACOS = 28.00 / 80.00 = 35.0%; TargetACOS 17.5%; MaxAccACOS 26.3%. Max profitable CPC = 28.00 x 0.052 = \$1.46. Target CPC = 80.00 x 0.175 x 0.052 = \$0.73. Every launch bid decision sits in that spread, and a bid above \$1.46 is a subsidy a sized rank case must carry.

P3 sizing, per v4.0 Section 10 rule M2. Required clicks per day equals DSTR divided by expected CVR; required budget equals clicks times Expected CPC.

T1: DSTR 1.30 orders/day ÷ 0.052 = 25.0 clicks/day x \$2.50 = **\$62.50/day**

T2: DSTR 1.04 ÷ 0.052 = 20.0 clicks/day x \$2.50 = **\$50.00/day**

T3: DSTR 0.78 ÷ 0.052 = 15.0 clicks/day x \$2.50 = **\$37.50/day**

Ranked subtotal: 60.0 clicks/day, \$150.00/day

5 capped tests at 50 clicks x \$2.50 = \$125.00 each over a 14-day window = \$8.93/day each = \$44.64/day

3 Halo exacts at \$5.00/day = \$15.00/day

Total staged: \$209.64/day = 99.8% of the \$210.00 envelope. Breach line is 110% per v4.0 Section 8.1 scenario S-A6: clear.

Day-14 progression read (P5). Gates clear on all rows except one truncation on a test term, corrected the same day.

T1: 344 clicks, 18 orders, CPC \$2.52. Spend \$866.88. Sales 18 x \$80.00 = \$1,440.00. ACOS = 866.88 ÷ 1440.00 = **60.2%**, inside the Ranking envelope of 50-80% (v4.0 Section 3 family F1). CVR 18 ÷ 344 = 5.2%. Delivered 24.6 of 25.0 required clicks = 98.4%. Rank unranked to 14 against target 10. Loss = 866.88 - (18 x 28.00) = \$362.88 over 14 days = \$25.92/day, inside the batch loss budget. → **CONTINUED.**

T2: 277 clicks, 15 orders, CPC \$2.48. Spend \$686.96. Sales \$1,200.00. ACOS **57.2%**. CVR 5.4%. Delivered 19.8 of 20.0 = 99.0%. Rank 11 against target 12: at target, first window. → **HOLD; DEFEND LIGHTLY AT RANK**, not HELD.

T3: 129 clicks, 4 orders, CPC \$2.31. Spend \$297.99. Sales \$320.00. ACOS **93.1%**. CVR 4 ÷ 129 = 3.1% against the 5.2% launch model. Clicks 9.2 of 15.0 = 61.3%. → **FIX CONVERSION FIRST; DEFER PUSH.**

5 tests: 224 clicks, 7 orders, spend \$548.80, sales \$560.00, ACOS 98.0%.

3 Halos: \$210.00. **Total 14-day spend \$2,610.63 against a \$2,940.00 envelope: utilization 88.8%, unspent \$329.37.**

Why T2 does not release budget. T2 is at target and the temptation is to step it down and fund batch 2 with the saving. It does not qualify. v4.0 Section 11 rule T1 fires the transition to a post-objective posture at 2 weeks at or under target rank, and T2 has one window. The launch hold read that would qualify a term sooner is amendment item AM-5 and is not substituted with a locally chosen number. The correct verdict is HOLD; DEFEND LIGHTLY AT RANK and the released budget is zero.

Why T3 does. T3 caps at maintenance under its fix-first verdict, which is the capped test level of \$8.93/day against the \$21.29/day it was running, releasing \$12.36/day, or \$172.99 over the next 14 days.

Batch 2 funding = unspent \$329.37 + released \$172.99 = **\$502.36 over 14 days = \$35.88/day.**

4 terms at the capped test level: $4 \times \$8.93 = \$35.71/\text{day}$. Headroom \$0.17/day. **Batch 2 = 4 capped tests, zero ranking-sized terms.**

The lesson, and the reading that escalates. Batch 2 shrank from the plan's proposed ranking-sized set to four capped tests because exactly one term released budget and it was the failing one. The envelope is neutral by construction and the next batch is an output of what the last batch proved, not of what the plan hoped. Separately, 14-day TACOS reads \$2,610.63 against \$4,000.00 total revenue including \$480.00 organic, or **65.3%**: 2.6 times the v4.0 Section 3 family F2 launch band top of 25%, mechanically CODE RED. This SOP does not suspend the rule. The reading is recorded in T0 zone B, flagged, and routed to the PPC Lead under AM-16.

WE-2 · Edge case: inventory flips RED inside the honeymoon window

Setup. Day 9 of a 21-day honeymoon window. The daily block at P4 step 1 reads T2 block K2 and finds the D2 preferred variation at Days to Stockout 11, which is at or under the 14-day line and therefore RED under v4.0 Section 5 gate G6. Two of the eight batch-1 terms point at that variation: T1 at bid \$2.60 and T2 at bid \$2.45. The ranked backup carries Expected CPC \$2.35.

The conflict the operator will feel. The honeymoon rule says withhold lever moves; the gate says act today. The gate wins, and not narrowly: v4.0 Section 12 rule P1 places gates above diagnosis and diagnosis above signals, mechanically. The honeymoon rule suppresses *efficiency tightening*, a lever moved on ACOS optics. A re-point is a gate response to an inventory fact. Reading the honeymoon rule as a freeze on all activity is how a launch ranks into a stockout and burns the spend twice.

Execution, same day, per v4.0 Section 8.3 scenario S-13.

1. Write the pre-OOS baseline snapshot *before* anything moves: T1 rank 16, CVR 5.2%, CPC \$2.52; T2 rank 11, CVR 5.4%, CPC \$2.48. The snapshot is what makes these recovery candidates with proven ceilings later, rather than cold starts.

2. Re-point both ad-level targets to the ranked backup through SOP-12 P10, seeding each bid at the backup's Expected CPC per v4.0 Section 9 rule L9.

T1: $\$2.60 \rightarrow \$2.35 = -9.6\%$. T2: $\$2.45 \rightarrow \$2.35 = -4.1\%$.

Both inside the v4.0 Section 9 rule L2 limit of plus or minus 50% per change and clear of the \$0.50 floor. Both under the 25% bid-move line in v4.0 Section 12 rule P5, but both are structural changes, which are human-mandatory under the same rule regardless of size.

3. Taper the remaining terms on the RED variation to defense; no term on that variation scales while the band is RED.

What it costs the batch read. Per SOP-12 P10 step 3 the CVR baseline resets at re-point and the trend clock restarts, so T1 and T2 read for gates only until sufficiency refills. The day-14 pass classifies 6 terms, not 8, and the log records the exclusion with its cause. Sizing batch 2 on 8 terms of evidence when 2 had their clocks reset is how a launch model gets calibrated against numbers never measured.

Close. On restock, the reverse re-point runs with the recovery play under v4.0 Section 8.3 scenario S-15, sized against the proven ceiling in the snapshot with a 3-week ETA checkpoint. The pre-OOS baseline written on day 9 is the entire reason that recovery is funded at priority rather than treated as a cold start.

WE-3 · Edge case: the Phase 2 advance that is refused twice over

Setup. Week 6. The launch operator requests Phase 2 discovery entry, citing organic share and TACOS posture as met. P7 runs.

Step 3, the manifest, before any condition is read. T7 reports 24 gate-relevant field groups: 17 PRESENT, 4 PARTIAL, 3 MISSING.

Completeness = $17 \div 24 = 70.8\%$. The DEGRADED line is below 85% per v4.0 Section 3 family F15. 21 of 24 groups must read PRESENT to clear it, so **4 more must land**.

Per v4.0 Section 5 gate G8, a DEGRADED packet produces flags only. Per v4.0 Section 12 rule P8, data-quality precedence beats every other consideration including urgency.

Step 4, the conditions, refused on a second and independent ground. The Phase 2 entry values are not carried by the Criteria System v4.0, per Section 1.2 of this document. The launch plan attached to this launch states two of the three conditions without a version date, and the third is not stated at all. A condition without a sourced, dated value cannot be verified; it can only be agreed to, and agreement is not an artifact.

The wrong move, named so it is recognisable. The shortcut is to read the two dated conditions as met, treat the third as satisfied by them, and advance. That produces an entry with no record of what it was tested against: unreviewable, unscoreable, and unreconcilable when v4.1 lands the values. It also stands up discovery surfaces on a DEGRADED packet, so their first four weeks of search-term data arrive with no manifest behind them.

Execution. The advance is refused in writing on both grounds. Every affected row carries **QUARANTINE: DATA REFRESH**, cross-mapped to v4.0 Section 13.1 FLAG: DATA REPAIR. Three named data-repair tasks open with owners and dates. The missing condition value escalates to the PPC Lead as an AM-1 dependency. The re-test date is set to the refresh that lands the fourth field group. Batch progression continues unaffected: a refused phase advance stops Phase 2, not Phase 1.

Outcome. Entry clears 11 days later at completeness 87.5%, with all three conditions sourced and dated in the phase-advance record. The 11-day delay is visible, owned and cheap. The alternative was an unauditable advance that would have surfaced at the handoff, when it is neither.

WE-4 · Edge case: graduation verified, co-sign blocked, and the ceiling that moved

Setup. Week 14. Graduation conditions verify against the launch plan's values, with the source document, version and date recorded per P10 step 1. The packet assembles.

Packet check: **9 of 11 fields present = 81.8%**. Two short.

Missing field 2: three terms carry CVR figures computed on 61, 74 and 88 clicks, all below the v4.0 Section 5 gate G2 line of 100 clicks for a CVR verdict.

Missing field 6: wall lists absent for 2 of the 3 Halo exacts.

Why these two matter more than they look. A CVR computed on 61 clicks entering the standing system as a validated benchmark becomes the denominator of that term's ceiling, sizing and scenario call for months. A missing wall list leaves the #29 audit with no baseline on day one, so the first re-

bought graduate reads as a new defect rather than a known gap. Both are cheap now and expensive later, which is the profile of the failure this checklist catches.

Repair. The three thin terms carry forward as **CONTINUE TEST TO 50 CLICKS AT CAPPED BID** with their click counts and no CVR, so the standing system inherits an honest gap instead of a fabricated benchmark. The two wall lists rebuild from the #29 wall audit. Packet returns to **11 of 11 = 100.0%**.

The economics upgrade at P11 step 1. Six terms clear the 100-click line.

Clicks $344 + 277 + 129 + 112 + 104 + 101 = 1,067$. Orders $18 + 15 + 4 + 6 + 5 + 4 = 52$.

Blended validated CVR = $52 \div 1,067 = 4.9\%$ against the 5.2% launch model: **-6.3%**.

Max profitable CPC per v4.0 Section 3 family F1 = margin dollars x CVR:

launch model: $\$28.00 \times 0.052 = \mathbf{\$1.46}$

realized: $\$28.00 \times 0.04873 = \mathbf{\$1.36}$

The standing ceiling is 6.3% lower than the launch model assumed, and every ceiling-referencing row on the product inherits the corrected value at handoff.

The co-sign that does not land. Conditions were met on day 96. The standing owner has not countersigned by day 105. Under P12 step 5 the launch owner still holds the product, and both weekly passes have run on schedule because ownership never lapsed. The escalation interval carried by v1.0 does not trace to v4.0 and is amendment item AM-14, so the escalation runs on the authority path instead: the launch operator holds D and E on launch-stage verdicts and the PPC Lead holds R, and under Authority Matrix v2.0 standing rule 3 the disagreement routes up rather than sideways. The co-sign lands day 107.

Orphan window: zero days. Not because the handoff was fast, but because ownership was never contingent on it. That is the whole design of the seam rule: the co-sign transfers the product, and until it exists there is nothing to transfer, so there is nothing to drop.

7. Failure Modes and Escalation

Twelve failure modes, each with the signal that detects it in a named artifact, the response, and the account-history incident that motivated it where one exists. F1 to F6 carry forward from SOP-22 v1.0 with their incidents intact; F7 to F12 are new and draw their incidents from the patterns v4.0 records.

Failure	Detection signal	Response	Incident
F1. Everything-at-once deployment	Any broad, phrase, auto, brand-format or conquest campaign live before the P7 phase-advance record exists	Pause to the sequence same day; the campaigns return at Phase 2 through their owning SOPs. The P3 checklist is the control and its absence is the root cause every time.	The pattern the deployment sequence was written against: launching every type together, which reads as coverage and prices as unpriced risk because no surface can be attributed.
F2. Premature efficiency tightening	A bid step-down staged inside the honeymoon window whose reason string cites ACOS	Revert. The window buys relevance signals that outlive it. Gate responses are exempt and are distinguished by carrying a gate ID in the reason string.	Carried from v1.0. It repeats because launch ACOS looks alarming against every band an operator has internalised from mature products.
F3. Immortal Halo campaigns	A Halo exact live past the retirement condition in its T3 row, or with no condition in T3 at all	Retire per the charter at P11 step 4. A Halo created without a condition fails on its creation date; P3 step 3 is the control and the T3 field is the evidence.	Carried from v1.0: Halo exacts that outlived their purpose by months because nobody wrote their exit, and by then nobody could say what it should have been.
F4. Batch discipline lost to spread	Live term count above the plan's batch with per-term delivered clicks under the worksheet requirement across most rows	Re-batch to the priority set at full pressure. Fewer terms funded properly, never more funded thinly: thin pressure ranks none of them and produces readable evidence on none of them.	Carried from v1.0. The corpus records the same shape at product level as concentration without a plan, one term absorbing roughly a quarter of a product's clicks above break-even.
F5. Handoff orphan window	Graduation verified, no co-signed T6 line, no cascade run this week	The seam rule: the launch owner holds until the co-sign exists. Escalate the delay, never the ownership.	Carried from v1.0, and why the handoff map was a scheduled deliverable rather than a paragraph inside a launch plan.
F6. Launch data lost at handoff	A workbook standup missing packet fields: validated CVRs, CPC ranges, wall lists, pre-OOS baselines or scored predictions	The 11-field checklist re-runs and the handoff waits. A standing system starting blind repays the launch spend to relearn what launch already measured.	Carried from v1.0. Worked in WE-4, where two missing fields would have handed the standing system a benchmark computed on 61 clicks.
F7. Verdict strings invented at the seam	A T2 block K9 verdict, T6 reason string or staged reason string outside the 25-item closed vocabulary	Re-issue from the Section 5 table. A closed vocabulary is what makes the log queryable and predictions scoreable per verdict type; one invented string breaks both for every row carrying it.	SOP-22 v1.0 itself: its Section 5 mapped five strings present in neither closed vocabulary. Conflict C17.
F8. Phase advance on calendar	A phase-advance record with no condition artifact set, or with values whose source document and version are unrecorded	Revert the advance and hold the Phase 2 surfaces at current budgets. An advance whose test cannot be named cannot be reviewed or reconciled when v4.1 lands the values.	The gap this rebuild found: the phase framework is cited to v4.0 and is not in v4.0, so an advance recorded without its source is unauditably by construction.
F9. Launch term carried by a Phase 2 surface	Auto or broad carrying above 40% of a batch term's clicks per v4.0 Section 4 rule X10, read on the T2 campaign child rows	ADD EXACT + RESTRUCTURE where the exact is missing, steering walls per #29 where it exists. A term ranked by campaigns that choose their own traffic is unranked by control.	The head-term case v4.0 Section 8.4 scenario S-R12 was written for: a product's head term carrying all its traffic through auto with zero exact campaigns.
F10. Provisional economics treated as live	A ceiling-referencing action deployed while the T0 margin version cell reads PROVISIONAL with no provisional flag in its reason string	The flag rides every ceiling-referencing action to the P11 upgrade. Per v4.0 Section 13.3 rule J7 every provisional read appears in the record in the words the engine computed; hiding it is a process violation.	The account's standing condition: blended margin used where per-variation margin was owed, which sets ceilings on the wrong denominator across a whole product.

Failure	Detection signal	Response	Incident
F11. Sibling collision at launch	A batch term also carrying an active objective on a sibling with no lead declaration in Account Rollup A1	Neither side scales until the portfolio call assigns a lead and gives the second product a support or stand-down posture, per v4.0 Section 5 gate G11.	The divergence the corpus records at \$3.24 against \$1.48 CPC on the same term across two siblings, which the inspection order found to be settings archaeology rather than auction reality.
F12. Predictions logged and never scored	Launch-stage T6 rows past their review dates with the V5 score column empty at the P13 read	Score them late and record the delay. Unscored actions do not accumulate per v4.0 Section 14 rule V1, and a launch model never scored stays wrong while every later launch pays for it.	The pattern v4.0 Section 14 exists to end: changes made, effects never checked, the same argument re-run each cycle because last cycle's reasoning was never written down.

Escalation. F1 and F4 escalate to the launch owner and PPC Lead jointly on first occurrence, being plan breaches rather than tuning errors. F5 and F8 escalate to the PPC Lead with the artifact set. F7, F10 and F12 attach to the P13 read as standing audit lines. Any mode recurring in consecutive windows on the same scope escalates with its T6 history. Anything in the human-mandatory tier is confirmed under the Authority Matrix before deployment, without exception.

8. Outputs and Artifacts

Artifact	From	Storage location	Retention	Downstream consumer
Readiness checklist, countersigned	P1	Launch record, referenced from T5 Product Log	Held to the P13 close; no numeric schedule governed (AM-13)	P3, which does not open without it; the launch department's audit series
Command Workbook, cloned and spined	P2	The product's PPC Command Workbook, tabs T0 to T7	Permanent; it is the product's standing instrument	Every procedure here; the standing weekly cascade after co-sign
Standup record with mini-worksheets	P3	T2 block K7 fields, T3 registry rows, T6 prediction rows	T6 append-only; K7 updates in place, worksheet versioned	P5; SOP-23 worksheet calibration; packet field 1
Wall lists per campaign	P3, P6	The T3 row for each campaign	Live for the campaign's life; carried into packet field 6	The #29 monthly wall audit, which measures against them
Honeymoon daily lines	P4	T6, one line per intervention with its gate ID	Append-only	P5 context; the F2 detection read
Pre-OOS baseline snapshots	P4.5	T2 block K4 Pre-OOS baseline field	Held until the recovery play closes or the term retires	The recovery play under v4.0 Section 8.3 scenario S-15; packet field 9
Batch progression log	P5	T6, with per-term classification and the funding computation	Append-only	P6 sizing; packet field 1; P13 calibration
Phase-advance record	P7	T5 targeting-change row with the artifact set attached	Permanent; it is the stage's audit trail	#30 registration; the launch audit; v4.1 reconciliation when values land
Keyword-level audit table	P9	Delivered into the launch weekly audit; reference in T5	Per the launch department's own retention	The launch audit; the BM owner's review
Handoff packet, 11 fields	P10	Launch record, referenced from the T6 co-sign line	Permanent; it is the stage's closing evidence	P11; the standing owner; any later question about the stage
Co-signed handoff line	P12	T6, one row, both signatures, cascade attached	Append-only, permanent	The standing weekly cycle under #36; the Register run ledger
Calibration record	P13	Launch record plus T6; proposals route to the v4.1 register	Held to the next launch of the same product class	The next launch plan; SOP-23 worksheet assumptions; the v4.1 brief

Retention. The corpus governs retention by rule, not by schedule: T6 is append-only, paused rows are paused and never deleted so history stays readable, and the workbook is permanent. No numeric retention period for launch artifacts exists anywhere in the pack, and the Cross-Reference Ledger records retention policy lines as a pending fold-in against component #39. Stated as AM-13 rather than filled with a chosen number.

9. Skill-Conversion Block

Block	Specification
Input bindings	By Catalog field group: A3 pacing and in-budget for the P4 block and G3 responses; A4 sufficiency buckets for the P5 tree; A6 objective tag and Expected CPC for the mini-worksheets; A7 eligibility for G4; B1 indexing and listing placement for the P1 gate; B3 rank pack for progression and packet field 4; B5 coverage counts for the F9 read; B8 economics chain from T0 C18 to C24 for every ceiling; B9 week-over-week deltas, MISSING at launch by construction; C1 to C8 for MKL hygiene; D1 per-variation margin, PROVISIONAL through the stage; D2 variation map with backup chains for G6 and the L9 re-point; D7 decision log for V5 scoring. Plus three launch-only bindings the standing SOPs lack: the batch lists with indexing status, the product-type classification, and the dated honeymoon exit. An absent feed suppresses its dependent procedure under the manifest; the skill never improvises around a missing feed and never substitutes an account reference where a per-term value is owed.
Deterministic logic	The P5 seven-branch classification tree, evaluated in order, first match winning. The P5.4 funding computation: released budget plus unspent envelope, released drawn only from HELD terms and terms capped at maintenance. The P7 five-branch entry tree. The P10 11-field completeness count. The M2 worksheet arithmetic: required clicks equals DSTR divided by expected CVR, required budget equals clicks times Expected CPC. The P11 spine recomputation from realized data. Scenario matching, threshold comparison and verdict issuance are v4.0's, referenced by rule ID and never re-derived. The skill version-pins v4.0 and re-validates the Section 5 mapping on any increment.
Human checkpoints	By Authority Matrix v2.0 row. "Launch-stage verdicts + batches": launch operator decides and executes, PPC Lead reviews, BM owner reviews, CEO informed, so every batch deployment and phase advance is human-mandatory and cannot auto-apply at any confidence level. "Structure creation (all SOPs)": PPC Lead decides or the owning SOP's scope governs, launch operator executes in launch scope, so every P3 and P6 creation is human-confirmed. "Push funding (the M5 queue)": PPC Lead decides within the envelope, CEO above it, so a batch needing an envelope increase leaves this SOP. "Bulk deployments (#38)": PPC Lead reviews, launch operator executes in own scope. "Code-red declaration + close" read with standing rule 1, which splits the shared cell explicitly for launch products: the launch operator declares, the PPC Lead closes. "Event plans (#32)": PPC Lead decides with the deals owner, launch operator executes in launch scope. The P12 co-sign is dual human-mandatory by construction. Standing rule 5 applies throughout: an unfilled role's decision right reverts to the PPC Lead and its execution to the PPC Manager.
AI-executable with audit	Gate-state assembly and the daily gate read; worksheet computation from bound inputs; classification drafting against the P5 tree; funding-computation drafting; condition-artifact assembly and completeness counting for P7 and P10; packet assembly and field count; the P9 audit table; V5 score drafting on rows due; P13 delta computation; F1 to F12 alarm detection with evidence lines.
Cadence	Daily in the honeymoon window: the four-gate block plus rank and CVR reads. Daily throughout the stage on G4 and G6 alarms. At each batch mark: the full progression pass. Weekly: the audit interface delivery. Same day: every gate response, truncation and RED re-point. Same week as conditions verify: the handoff steps. 30 days post co-sign: the calibration read.
Alarm specification	Each of F1 to F12 carries its Section 7 signal, an evidence line naming the row, the signal value and the gate or rule reference, and a routing tier. Same day to the operator queue: F2, F9, F10, F11. Into the next progression worklist: F3, F4, F6, F7. To the PPC Lead with the artifact set: F1, F5, F8. To the P13 read: F12.
Skill outputs	Drafted worksheets; classifications with their tree branch cited; funding computations; condition and packet checklists with completeness counts; staged change sets in #38 format with verdict and scenario ID in every reason string; drafted T6 rows each carrying a prediction and review date; the P9 table; F1 to F12 alarms with evidence; the P13 deltas. Everything queues for human confirmation; nothing in this stage auto-applies.

10. Version Governance

Owner: PPC Lead for the PPC execution. The launch department owns the plan this SOP executes inside, and a change to either side of the P9 or P12 interface requires both owners. Registered in the Supersession Map at publication; supersedes SOP-22 v1.0 in full. The Completion Directive row for the Launch-to-KDR Handoff Map closes against procedures P10 to P12.

Review triggers specific to this component.

- Any v4.1 increment landing one of the sixteen Section 13 amendment items. Each landed item replaces a sourced-from-the-plan gate with a governed threshold, and the affected procedure re-validates before the next launch runs on it.
- Any v4.0 increment, at which point the Section 5 mapping re-validates against the new closed vocabulary before any batch pass runs against it.
- Any change to the launch department's plan template, because eight gates here take their values from it and a field that disappears silently disables the gate that read it.
- Any resolution of conflict C10, which would change how Section 5 quotes its strings.
- Any P13 calibration returning a launch-model delta beyond what the next plan can absorb, which signals that the sizing method rather than its inputs needs review.
- The first launch run after this version publishes: a read-through against the actual sequence, not a review of the document in the abstract.

11. Operator Ramp

Built from the Authority Matrix v2.0 rows that name this component's decisions; each row becomes a competence with a demonstration. Week 1 observes, week 2 executes supervised with the checkpoint holder reviewing, week 4 executes unsupervised with the checkpoint holder auditing. The checkpoint holder is named by the matrix row, not chosen.

Authority Matrix row	Week 1: observe	Week 2: supervised	Week 4: unsupervised, audited
Launch-stage verdicts + batches launch operator D and E; PPC Lead R	Sit the full progression pass. Reproduce the classification of every live term against the P5 tree and reconcile against the operator's, naming the branch for each disagreement.	Run the pass and draft the funding computation. The PPC Lead reviews every classification and the released-budget figure before batch N+1 stages.	Run the pass and deploy batch N+1. The PPC Lead audits the progression log after deployment and spot-checks three classifications against their evidence.
Structure creation (all SOPs) PPC Lead D; launch operator E in scope	Read three T3 rows end to end: objective tag, role, advertised SKU, wall list. Say what each field would have to read for the row to be a defect.	Stage a batch-term standup through SOP-12 P1 with walls in the same session. The PPC Lead confirms the wall list against #29 before deployment.	Stand up a full batch unsupervised. The PPC Lead audits the T3 registry and wall lists against the #29 audit baseline.
Push funding, the M5 queue PPC Lead D within envelope; CEO above	Recompute two mini-worksheets from their bound inputs and reach the same required clicks and required budget.	Size a batch to computed available funding and show the envelope reconciliation against the WeeklyEnvelope named cell.	Size and stage a batch inside the envelope. Correctly refuse and route a batch that would need an envelope increase rather than absorbing it.
Bulk deployments (#38) PPC Lead R; launch operator E in own scope	Read a staged set and a T+1 verification report. Identify every reason string missing a scenario ID.	Stage a launch deployment and read its verification report. Route one mismatch to its response.	Stage and deploy unsupervised, reconcile T+1, escalate any landed-value mismatch same day.
Code-red declaration + close standing rule 1: on launch products the launch operator declares, the PPC Lead closes	Read the T0 zone B TACOS line against the v4.0 launch band, state its deviation class, and say why the reading is expected at this stage.	Draft the AM-16 escalation on a live launch TACOS reading, stating what is escalated and what is not suspended.	Declare when the reading warrants it, and correctly decline to declare when the reading is the stage's normal shape.
Event plans (#32) PPC Lead D with deals owner; launch operator E	Read the event plan against the launch calendar and identify any overlap with the planned go-live.	Run the G12 event check at P1 step 6 and record the result, including a recorded absence.	Defer or proceed against an armed window, with the decision recorded and its basis named.

Live-defense questions

These test the reasoning behind thresholds, not recall of their values. An operator who can recite a band and cannot answer these cannot be left with a launch.

1. A batch term is at target rank after one progression window. Why does it release no budget, and what would have to be true for it to release budget?
2. A launch product reads TACOS at 2.6 times the top of its stage band. Why is that not a Recovery Mode declaration, and what is the correct action?
3. Inventory flips RED on day 9 of the honeymoon. The honeymoon rule says withhold lever moves. Which wins, and what general principle decides it?
4. Two batch terms both under-deliver clicks. One has CVR above the launch model, one has CVR at 3.1% on 129 clicks. Why do they receive different verdicts, and which lever is barred on the second?
5. Why does the packet carry a click count rather than a CVR on a term with 61 clicks, and what happens downstream if it carries the CVR instead?
6. Graduation conditions are met and the standing owner has not countersigned. Who owns the product, and why is that a rule rather than a courtesy?
7. A batch could run 12 terms at half pressure or 6 at full pressure inside the same envelope. Which, and what failure mode does the other choice produce?
8. The launch plan states a Phase 2 condition with no version date. Why does that refuse the advance even when the number looks reasonable?
9. Name a launch decision this SOP does not make, and name the document that makes it.
10. Max profitable CPC recomputes from \$1.46 to \$1.36 at handoff. Which two inputs moved it, and which rows on the product does the change touch?

12. Interface Contracts

Every upstream and downstream handoff this component holds, with sending owner, receiving owner, artifact, timing, and the consequence when it does not arrive. The non-arrival column is the clause the corpus has always been missing: without it a late input becomes an improvised substitute, and the improvisation is invisible in the record.

ID	Artifact	Sender	Receiver	Timing	Consequence of non-arrival
I-1	Launch plan: batch lists, indexing status, product type, envelope, dated honeymoon exit	Launch department	Launch operator	Day -5	P1 cannot open; the standup date moves. This SOP does not supply a missing plan field and does not infer a product type, batch list or envelope from another source.
I-2	Keyword-to-Variation Map with adequacy bands and ranked backups	Logistics and Catalog, via Economics and Context Pack tab 2	Launch operator	Day -5, monthly	G6 cannot pass and no ranking-sized term deploys. Campaigns may stand up at capped test level on the declared preferred SKU with the D2-pending flag set on every row.
I-3	Margin table with contribution per variation and its version stamp	#41, with Finance	PPC Lead, into T0 C18-C24	Day -3, monthly	The spine runs on the launch model with PROVISIONAL on every ceiling-referencing action. Absence never blocks the launch; an unflagged provisional read is the F10 condition.
I-4	Declared launch envelope and funding queue position	Instrument layer: T0 zone H and Account Rollup A2, per conflict C3	Launch operator	Day -3, weekly	No batch sizes. Sizing without a declared envelope is how a launch quietly becomes an envelope increase nobody approved.
I-5	Pre-load negative sets from the taxonomy	#29	Launch operator	At P3, same session	The standup does not close. A launch without pre-loaded negatives spends into known-irrelevant space from day 0, which the taxonomy already knew and the campaigns did not.
I-6	Mini-worksheet method and field set	#23	Launch operator	Standing reference	No term deploys at ranking size; the batch runs at capped test level only. This SOP never derives its own sizing method.
I-7	Wall lists at every structure event	#29	Launch operator, into T3	Same session as each P3 and P6 event	The structure event is incomplete and the term is treated as unsteered. The #29 audit has no baseline for that campaign and the first re-buy reads as a new defect.
I-8	Staged change sets in #38 format	Launch operator	#38	Same day on launch releases	Nothing deploys. A console edit outside the staged set cannot be verified at T+1 and voids the affected rows' reads for the window.
I-9	T+1 verification report	#38	Launch operator	Day after every deployment	Landed values are unconfirmed and the progression reads on unverified state. The deployment is treated as unlanded until the report returns.
I-10	Armed event-window declaration with dates	#32	Launch operator	Day -5 and on any arming	The G12 event check cannot pass and the launch date cannot be defended later. A recorded absence of a window is a valid arrival; silence is not.
I-11	Lead-product declaration on shared terms	#33, via Account Rollup A1	Launch operator	Day -5 and on any new shared term	G11 blocks funding on that term. It may run at capped test level pending the declaration; it may not be ranking-sized while two products could push it.
I-12	Exact campaign standups on batch terms	Launch operator	#12	At P3 and P6	Terms sit unfunded and the batch under-delivers. An exact stood up outside SOP-12 produces a campaign the registry does not know about.
I-13	Phase 2 entry verification with its artifact set	Launch operator	#13, #14, #15	At P7, within 2 days of request	No discovery surface stands up. The receiving SOPs do not build on a verbal advance and may refuse a standup without the record.
I-14	Harvest registration for the launch product	Launch operator	#30	First Monday after Phase 2 entry	Discovery runs with no graduation destination. Converting terms accumulate unclaimed and the product pays discovery prices for demand it has already proven.

ID	Artifact	Sender	Receiver	Timing	Consequence of non-arrival
I-15	Graduation issuances on launch discovery terms	#30	Launch operator	Weekly from Phase 2	Proven terms stay on discovery surfaces at uncontrolled CPC. Launch products get no exception from the graduation bar and none from its cadence.
I-16	Placement evidence once conversion exists on a term	Launch operator	#28	Post-conversion, per the P8 overlay	Modifiers stay at zero. On non-subjective products that is the correct resting state and costs nothing; on subjective products it delays a legal test.
I-17	T7 manifest status by field group	#39, via Data Operations	Launch operator	Weekly at packet-ready	G8 cannot be evaluated, so every verdict this window issues as a flag. An unread manifest is treated as DEGRADED, never as clean.
I-18	Raw Pull Pack with its period sheets	Data Operations	The Command Workbook	Monday 06:00 for a 09:00 packet	The progression pass runs on stale state and reads for gates only. At launch the four-period sheets read MISSING by construction; that is recorded, not worked around.
I-19	Keyword-level audit table	Launch operator	Launch department, BM owner	Within 1 day of the audit landing	The audit proceeds without the PPC read and the launch department decides on partial evidence. Two consecutive misses escalate to the PPC Lead.
I-20	Co-signed handoff line with packet and first cascade	Launch operator and PPC Lead jointly	The standing owner, into #36	Same week the packet completes	The launch owner keeps the product and keeps running it. There is no orphan window by construction; the delay escalates and ownership does not move.

13. The v4.1 Amendment Register for This Component

Sixteen values this SOP needs and v4.0 does not carry. None is written into a procedure above.

ID	Threshold or rule owed	What the procedure does until it lands
AM-1	Phase 1 to Phase 2 entry conditions with values	P7 verifies against the launch plan's values and records source, version and date in the phase-advance record.
AM-2	Launch graduation conditions with values on the five dimensions v4.0 Section 11 rule T2 names	P10 verifies against the launch plan and records the source the same way.
AM-3	Honeymoon window duration and exit rule	P3 step 11 writes the plan's dated exit into the standup record and T5; P4 closes on that date.
AM-4	Batch size limit and batch cadence	P5 runs on the plan's cadence recorded in T5; P6 sizes to computed funding rather than to a governed count.
AM-5	The launch hold read qualifying a term for a step-down before the standing transition interval	P5 step 5 falls back to v4.0 Section 11 rule T1 at 2 weeks at or under target rank. One window at target releases no budget.
AM-6	Launch envelope-setting rule for a product with no sales history	P2 takes the declared envelope from the plan; P3 step 8 reconciles under G7 and scenario S-A6.
AM-7	Launch batch loss ceiling, where v4.0 Section 10 rule M2 sizes from projected sales a launch term does not have	P3 step 6 carries the plan's batch loss budget as a stated input and marks it as such on every worksheet.
AM-8	Launch sufficiency substitute below the G2 lines	None is used. Terms below the lines receive capped-test verdicts and no rate judgment, which is why most launch verdicts are tests.
AM-9	The legal condition set for a Halo retirement charter	P3 step 3 requires a written condition and does not constrain its form. F3 detects a missing condition, not a wrong one.
AM-10	Product-type waste bands, carried by the KDR Template and not by v4.0	P8 step 5 applies the v4.0 Section 3 family F11 limits of 10% standard and 40% discovery, and logs the divergence.
AM-11	Stage TACOS band reconciliation across the three documents stating it differently	Every reference here uses the v4.0 Section 3 family F2 bands. Conflict C16.
AM-12	Whether launch batch exacts inherit the 30% top-of-search requirement before conversion evidence exists	P8 steps 2 and 3 run the overlay; modifiers stay at zero until evidence exists, and the interaction is flagged, not resolved.
AM-13	Retention schedule for launch artifacts	Element 8 states retention by rule: append-only logs, pause never delete, permanent workbook. No period is chosen.
AM-14	The co-sign escalation interval	P12 step 6 escalates on the Authority Matrix path with the delay counted in days and recorded, not against a chosen threshold.
AM-15	Provisional economics tolerance at launch, where the spine is provisional by design	P2 step 4 tags the spine PROVISIONAL and the tag rides every ceiling-referencing action to the P11 upgrade.
AM-16	Launch interaction with the CODE RED deviation class and Recovery Mode precedence	The reading is recorded, flagged, and routed to the PPC Lead under the Authority Matrix row owning v4.0 threshold changes. No local suspension.

Companions: Master PPC Decision Criteria System v4.0, Sections 3, 5, 6, 7, 8, 9, 10, 11, 12, 13 and 14 · Keyword Decision Record Template v3.0, sheet "Bands & Verdicts Reference" (the 25-item closed vocabulary quoted in Section 5) · PPC Authority Matrix v2.0 (Elements 9 and 11) · Master PPC Data Architecture v1.0, Sections 4, 6, 8 and 9 · PPC Command Workbook TEMPLATE v1.0 · PPC Strategic Doctrine v1.0, Sections 1, 2 and 3 · PPC Economics Worked-Example Set v1.0 (the base disclosed in Section 6) · SOP-12 v2.0, SOP-23, SOP-13, SOP-14, SOP-15, SOP-28, SOP-29, SOP-30, SOP-32, SOP-33, SOP-38, SOP-41 (boundaries) · PPC Cross-Reference Ledger (component #22 row; conflicts C1 to C18) · the launch department's plan and weekly audit frameworks.