

SOP-23: Ranking-Push Operations

Version 2.0 · 2026-08-08 · Owner: PPC Lead · Supersedes SOP-23 v1.0 (2026-07-24) in full · Every threshold lives in the Master PPC Decision Criteria System v4.0; this SOP carries zero local thresholds · Verdict strings are quoted verbatim from the Keyword Decision Record Template v3.0 Bands and Verdicts Reference, with the v4.0 Section 13.1 cross-map stated where the two diverge

1. Verdict and Scope

VERDICT: a ranking push is spend that carries five properties at once, and spend moving toward rank without all five is not a slow push, it is an unpriced subsidy that pays full auction cost for no governed outcome. The five are sized (a complete M2 block), dated (a horizon date and an ETA checkpoint), ceilinged (a weekly loss ceiling computed, not estimated), predicted (a V1 prediction with metric, direction, magnitude and horizon), and funded (an M5 queue release). This SOP owns the arithmetic that produces those five and the operating cadence that defends them from candidacy to post-mortem. It does not own campaign mechanics, queue law, ladder mechanics, recovery, or event overlays, each of which is named by number in Section 1.4.

What changed at v2.0. v1.0 was built to the superseded 10-element standard and measured 4 pages against a 10-18 page bar. v2.0 expands procedures from 7 to 13, failure modes from 7 to 13, and worked examples from 1 to 4; adds Element 11 operator ramp and Element 12 interface contracts, neither of which exists anywhere in the 23 shipped SOPs; and retires three verdict strings v1.0 used that belong to no closed vocabulary. Section 10 carries the full change note. The E3 sizing anchors and all seven v1.0 account-history incidents carry forward, now traced to source rather than asserted.

1.1 The five properties, stated as tests

A row that fails any test below is not a push and is governed by v4.0 Section 8.4 scenario S-R11 (CPC above 1.5 times ceiling with no sized push, no M2 record), which returns an immediate ladder toward ceiling rather than a smaller push.

Property	The test a row must pass	Where the standard comes from
Sized	T2 columns DSTR (orders/day), Required Clicks/Day, Required Budget \$/wk all populated from the same worksheet version.	v4.0 Section 10 rule M2, which requires every push to carry target rank, required clicks against actual, budget, ceiling, ETA checkpoint and stop-loss.
Dated	T2 columns Target Rank, ETA Checkpoint and Review Date carry dates, not blanks.	v4.0 Section 6 rule O7, under which every Ranking objective carries a horizon date and renews only on evidence.
Ceilinged	T2 column Loss Ceiling \$/wk computed as (ACOS above break-even) times projected sales, dated.	v4.0 Section 10 rule M2, loss-ceiling clause.
Predicted	A T6 Decision Log row exists with Expected Outcome Window and Review Date populated.	v4.0 Section 14 rule V1: unscored actions do not accumulate.
Funded	The term appears in T0 zone H as FUNDED with a dated release.	v4.0 Section 10 rule M5, which orders pushes by revenue potential at target rank divided by cost-to-close-gap.

1.2 The three candidacy classes

Class is assigned before sizing, because class changes the realism bar, the funding priority and the checkpoint. Classification is v4.0 Section 7 rule D10, which requires every FAR or UNRANKED term to be classed before objectives are assigned.

Class	Condition (v4.0 Section 3 family F6, best-rank row)	Consequence for this SOP
Cold	NEVER-RANKED: no best rank at or under 10 in the trailing 12 months and no pre-OOS baseline at or under 10.	Passes the full O3 realism gauntlet at P2 with no history credit. Sized on assumption, which is why its post-mortem calibration at P13 is the more valuable of the two.
Recovery	RECOVERY candidate: best rank at or under 10 within the trailing 12 months, or a pre-OOS baseline at or under 10. PROVEN where that rank is at or under 5.	Inherits a proven ceiling, skips the cold-start realism penalty and funds at better M5 priority per D10. The recovery push itself belongs to SOP-24 , which extends this document; this SOP classifies and hands over at P2 step 6.
Structure-blocked	Priority term at 500 or more search volume on a primary syntax with zero exact campaigns, or auto and broad carrying more than 40% of its clicks (v4.0 Section 4 rule X10).	v4.0 Section 8.4 S-R12 bars any performance verdict until the exact stands up. Candidacy is deferred, not declined: it re-enters at P1 after sufficiency per G2.

1.3 Why the sizing is the document

The account's standing failure is not badly executed pushes. It is spend that reaches break-even and keeps going with no arithmetic attached. v4.0 Section 3 family F16 records the motivating case: one term absorbing roughly a quarter of a product's ad clicks above break-even with no plan behind it. At \$6,000-8,000 per day of account spend, v4.0 Section 11 rule R1 puts a 2x stage-band breach at a four-figure daily leak. The P3 worksheet converts that pattern into either a funded decision or a bid-down verdict, with nothing in between, which is Strategic Doctrine v1.0 Section 1's third principle applied: economics set ceilings, never targets, and margin dollars rather than margin percent size every push.

1.4 Boundaries, with the owning component named for every excluded decision

Per the Cross-Reference Ledger, this component owns push sizing mathematics, DSTR and required-clicks, daily pacing on funded pushes, taper trigger execution, and the push precondition block. Everything below is referenced by number and one clause and is never restated here.

Excluded decision	Owner	The seam, stated so it is owned
Every threshold, band, scenario and verdict issuance	v4.0	A number appearing in this SOP that does not quote a v4.0 section is a defect. Section 13 of this document lists the numerals v4.0 does not yet carry; they are amendment items, not local values.
Campaign creation, bid and budget edits, taper execution at campaign level	SOP-12	This SOP issues the standup order and the taper trigger; SOP-12 P1 creates and P4 executes. A bid value written by this SOP is a defect.
Queue ranking, envelope law, budget release and return	#31	This SOP submits a sized worksheet and receives a dated release. It never ranks the queue and never moves budget. Per Ledger conflict C3, #31 has no procedural document, so this SOP cites the instrument (T0 zone H, Account Rollup A2) and never a #31 procedure number.
Ladder step derivation, exit-bid discovery, modifier sizing, the ceiling worksheet	#28	This SOP releases the taper trigger at P10 and stops. Every rung, every floor and every modifier value is #28's, and #28 is the authoritative ladder that SOPs 12, 23, 24, 25 and 26 all reference.

Excluded decision	Owner	The seam, stated so it is owned
Loss diagnosis, baseline restoration, recovery sizing, out-of-stock overlays	SOP-24	SOP-24 extends this document and inherits its procedures by reference. A rank that was held and then lost is SOP-24's from the moment the loss is dated, not this SOP's.
Event plans, caps, arming, hangover routing	#32	Windows amplify live pushes and never originate them. This SOP suspends and resumes at P12 on the armed window; the plan itself is #32's.
Family term ownership declaration and guard enforcement	#33	The push block is a SOP-23 precondition per the Ledger. This SOP holds candidacy until the declaration lands; it never adjudicates ownership.
Post-objective normalization after the taper	SOP-25	The taper exits into SOP-25. This SOP's last act on a successful push is releasing the trigger at P10 and scoring at P13.
Negation walls at every structure event	#29	Walls are built by #29 rules in the same working session as the standup; this SOP names the requirement and never the wall set.
Staging format, upload, and the T+1 verification pass	#38	Every change this SOP orders is staged, deployed and verified there. No console edit outside the staged set.
Which units are terminal or LTSF-tiered	#27	Per Ledger conflict C6 the deciding document sits outside this corpus. A CONSTRAINED earning-potential class bars ranking pushes outright under v4.0 Section 8.1 S-A10, and this SOP stops at that boundary.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. Where v4.0 states the lag, the section is cited. Where it does not, the cell says so and the item appears in Section 13; no lag is invented here.

Trigger	Detecting artifact and the observable	Maximum lag to procedure start	Procedure
Candidacy scan	T2 Keyword Command under the OVERSIGHT filter: rows where Organic Rank is worse than Target Rank and Objective reads Ranking or is blank.	One weekly cycle. v4.0 Section 14 rule V4 sets the weekly full audit as the cadence that produces Master Action Records.	P1
Funding release	T0 zone H M5 push queue: the entry state changes to FUNDED and carries a dated release.	Not stated in v4.0. Amendment item AM-6.	P6
Daily in-flight	T2 child rows for any campaign whose Objective reads Ranking with a live M2 block, read against T2 In-Budget % and the daily rank series.	Same day. v4.0 Section 14 rule V4 sets daily alerts, pacing and gate checks, with oversight-class checks during pushes.	P7
Weekly in-flight	The Monday audit cascade reaching T2 with a live push row present.	One weekly cycle per V4.	P8
Velocity stall	T2 Velocity (pos/wk) reading inside the FLAT band across the review window with T2 Actual Clicks/Day at or above Required Clicks/Day.	The review window closes at the ETA horizon, default 3 weeks per v4.0 Section 3 family F6.	P9
Early ETA warning	T2 ETA to Target (wks) crossing into BEYOND at current velocity.	Same weekly read. Family F6 requires the stop-loss conversation to open early, before the budget is gone.	P8
Target touched	T2 Organic Rank at or better than Target Rank on the daily series.	Same weekly read; the stability window starts from the touch date.	P10
Stability confirmed	T2 Organic Rank at or under Target Rank for 2 consecutive weeks (v4.0 Section 8.4 S-R6).	Same weekly read. The taper trigger was logged at standup and is not re-decided here.	P10
Horizon reached	T2 ETA Checkpoint date arriving with the target not yet held.	Same weekly read. v4.0 Section 6 rule O7 bars automatic continuation.	P11
Inventory band flip	T2 Inventory Band moving to YELLOW or RED on the targeted variation, or Days to Stockout falling to 14 or fewer.	Same day on RED (v4.0 Section 8.3 S-I3). YELLOW freezes at current spend with no same-day clause.	P12
Family conflict	T2 Family Conflict reading a live conflict on the term, or Sibling CPC View showing a second sibling with a push objective.	Not stated in v4.0. Gate G11 bars scale without naming a lag. Amendment item AM-7.	P12
Product posture flip	T0 zone A verdict banner entering push freeze or code red (v4.0 Section 8.1 S-A3, S-A4).	Same cycle on S-A3. S-A4 enters Recovery Mode, whose cadence is daily per v4.0 Section 11 rule R4.	P12
Event window armed	T2 Event Mode reading PRE-EVENT or IN-EVENT from the #32 event plan.	Arming is T-1 per the #32 plan; this SOP reads the state, it does not set it.	P12
Push close	Any exit: target held, stop-loss executed, concession logged, or candidacy withdrawn.	Not stated in v4.0. v1.0 asserted "at close" with no interval. Amendment item AM-8.	P13

3. Preconditions and Gates

Nothing in Section 4 runs until the manifest rows below read current on T7. Each row names the manifest field group by its Catalog ID, the gate it feeds by its v4.0 number, and the observable that proves the gate passed. Where the operating corpus uses a different gate number for the same idea, the corpus alias is given at first use per Ledger conflict C9; the v4.0 number governs.

Live manifest state, and what it costs this component. The T7 Manifest in the current Command Workbook template reads a packet completeness score of 0.571, below the 85% DEGRADED floor in v4.0 Section 5 gate G8. Four field groups this SOP depends on are not fully present: **D8 competitor depth is MISSING**, which degrades the DSTR input that v4.0 Section 3 family F14 defines as set by competitor velocity analysis; **D9 daily rank series is PARTIAL** on head terms only, which

T7 names as blocking push-window daily verification; **A6 objective and per-term Expected CPC is PARTIAL**, with T7 recording an account reference of \$2.80 in use until per-term values land; and **D1 per-variation margin is PARTIAL**, carrying the G9 PROVISIONAL tag onto every ceiling in this document. Under G8 a verdict whose gate fields are missing downgrades to a flag and the analysis states what it could not see. P3 step 9 enforces that statement on the worksheet itself.

Manifest ID	Gate	Requirement before any procedure runs	The observable that proves it passed
A4, A5	G1	No impossible-metric flag on the row: no CVR of zero with orders present, no ACOS without orders, no cross-table disagreement.	T2 column <code>Flag: Zero-Quarantine</code> is clear and <code>Validation State</code> is not QUARANTINE. A quarantined row produces a flag and a data-refresh task, never a push.
A4, A1	G2	Sample sufficiency for every rate the worksheet consumes: 15 or more clicks in 7 days for a bid read, 100 or more clicks for the CVR that sizes the push.	T2 <code>Clicks 7/30/60d</code> and <code>Sample State</code> . Below 100 clicks the CVR input falls back to the 60-day window; if still thin the verdict is EXTEND TEST and no push is sized. A1 reads MISSING on the live manifest, so four-period trend context is unavailable and single-window reads carry the flag.
A3	G3	The hosting campaign is not budget-truncated. In-budget below 70% of the day marks every rate on that campaign TRUNCATED and forces the budget lever before any bid conclusion.	T2 <code>In-Budget %</code> at or above 70. Corpus alias: the KDR Bands sheet labels 51-70 "strained"; G3's 70% line governs. A3 reads PARTIAL, so the day-of-week curve behind G7 is not yet available.
A7	G4	The advertised offer is eligible and unsuppressed.	T2 <code>Eligibility</code> reads eligible. A flip halts all execution for the target the same day and opens the listing task.
B1	G5	The term is indexed, and on a rank objective it is placed in title or bullets. No bid level ranks an unindexed term, and a push on a backend-only term funds the platform rather than rank.	T2 <code>Indexed?</code> reads yes and <code>Listing Placement</code> reads title or bullets. Backend-only routes to the listing queue before any push funding.
D1, D2	G6	The targeted variation reads GREEN. YELLOW freezes push spend at its current level; RED tapers to defense and re-points to the ranked backup the same day.	T2 <code>Inventory Band</code> and <code>Days to Stockout</code> , read against <code>Targeted Variation(s)</code> . GREEN requires days on hand above lead time plus 21 days at max-sales-day velocity. D2 reads PARTIAL, so the backup chain is confirmed on head terms only.
A3	G7	Campaign budgets reconcile to the product envelope and product pacing sits within 25% of the day-of-week curve.	T2 <code>Envelope Status</code> and T0 zone H available dollars. A push that does not fit the envelope is a queue decision at #31, not a smaller push here.
All	G8	The packet manifest reports present, stale or missing before analysis. Below 85% of gate-relevant fields the packet is DEGRADED and produces flags only.	T7 <code>Status</code> column and the completeness rollup stamped on the cycle header. See the note above: the live read is 0.571.
D1	G9	Price stability: no price change on the targeted variation inside the trailing 7 days, or the CVR and ACOS reads carry PROVISIONAL and no conversion verdict issues on them.	T5 <code>Product Log Event Type</code> filtered to price move, cross-read against T2 <code>Margin \$/Unit</code> . Corpus alias: SOP-12 v2.0, SOP-14, SOP-30 and the Data Architecture use G9 for margin freshness; v4.0 assigns G9 to price stability and both states are checked here.
D7	G10	Event weeks and the 14 days following are excluded from the trend reads that size a push.	T5 <code>Trend-Exclusion Flag</code> and T2 <code>Event Mode</code> . Corpus alias: the same corpus documents use G10 for objective tagging and G12 for the event freeze; v4.0 assigns G10 to event exclusion.
D6	G11	No second sibling carries an active push objective on the same term. Two siblings pushing one term is a portfolio decision, and neither push scales until the lead is assigned.	T2 <code>Family Conflict</code> and <code>Term Owner (Family)</code> , sourced from Account Rollup A1. D6 reads PARTIAL, and the Bamboo root ownership was left OPEN at cycle 1, which is a live instance rather than a hypothetical.
B11	G12	The term exists in the current Master Keyword List with a syntax tag and a relevancy score, and the list is not stale past 30 days.	T2 <code>Syntax Tags</code> and <code>Relevancy</code> populated. Corpus alias: G12 is used for the event freeze across the operating corpus; v4.0 assigns G12 to keyword-list hygiene.
B8, D1	Economics chain	T0 zone C is live: AOV, contribution per order, break-even ACOS, Target ACOS, Max Acceptable and the Expected CPC reference. Not a numbered gate; a precondition without which no ceiling in Section 4 computes.	T0 zone C margin table version and its fresh or provisional state. B8 reads PRESENT on blended margin; D1 reads PARTIAL, so every ceiling here carries PROVISIONAL until the per-variation table lands. This is external dependency D1 in the Ledger, open, and it caps this component by name.
B3, D9	Rank pack	Best rank with its date, pre-OOS baseline, velocity and ETA present for the term, and a daily rank series available for in-flight verification.	T2 <code>Best Rank (Hist) + Date</code> , <code>Pre-OOS Baseline</code> , <code>Velocity (pos/wk)</code> , <code>ETA to Target (wks)</code> . Both read PARTIAL on head terms only. A candidate outside head-term coverage cannot be classed at P2 step 1 and does not proceed.
B7	Placement splits	Placement-level splits present for the hosting campaign. Reads MISSING account-wide, which suppresses modifier increases at #28 and leaves S-R7 partially unexecutable. See WE-2.	T4 <code>Placement</code> and <code>Bid</code> rows carrying values rather than build markers. While MISSING, the top-of-search share requirement is measurable but not correctable by modifier.

4. Step-Level Procedures

Thirteen procedures cover the push lifecycle. Every step names the tab, the column and the value it reads or writes. A person who has never run a push executes these from the page. Time estimates assume the OVERSIGHT filter is already applied and the packet has refreshed.

P1. Candidacy assembly

Weekly, inside the Monday cascade. 25-40 minutes per product at roughly 100 oversight-class terms.

1. Open T2 Keyword Command. Confirm T0 and T1 have been read first: product posture and syntax caps constrain everything below them, and a T2 pass that skipped them is invalid by construction.
2. Score last cycle before assembling this one. Open T6 Decision Log, filter Review Date to this cycle, and write the V5 Score on every due row as HIT, PARTIAL, MISS or WRONG-DIRECTION. Unscored actions do not accumulate.
3. Apply the OVERSIGHT filter, then filter T2 to rows where Organic Rank is numerically worse than Target Rank. Rows with a blank Target Rank at 500 or more SV are unmanaged terms under v4.0 Section 6 rule O1 and route to target-setting before candidacy, not into it.
4. Drop rows failing the search-volume floor: v4.0 Section 6 rule O2 sets the ranking floor above 500 search volume. A term below it does not carry a Ranking objective regardless of its rank gap.
5. Drop rows whose Syntax Priority is secondary where the syntax already sits at its spend-share cap. v4.0 Section 12 rule P4 puts syntax caps above single-keyword strength, and a strong week never raises a cap mid-cycle.
6. Check Seasonal SV Index Dir. v4.0 Section 11 rule T4 times pushes into a rising index and never after the peak. A falling index is a defer with a dated re-entry, not a decline.
7. For each surviving row, run every Section 3 precondition and record which failed. **A failed precondition is the work item.** It is never a reason to size a smaller push, and it never routes to a bid change.
8. Write the surviving set into T2 Situation with the cycle date. Each survivor proceeds to P2. Candidates entering the queue without a completed P2 and P3 are the F1 condition.

P2. The O3 ranking-realism gauntlet

Per candidate, 8-12 minutes. Every clause must pass. Source: v4.0 Section 6 rule O3.

1. **Class the term first.** Read T2 Best Rank (Hist) + Date and Pre-00S Baseline. Best rank at or under 10 in the trailing 12 months, or a pre-OOS baseline at or under 10, classes the term RECOVERY; at or under 5 it is PROVEN. Otherwise NEVER-RANKED. Where B3 or D9 has no coverage for the term, the class cannot be established and the candidate stops here with the manifest gap recorded.
2. **Indexing and placement.** G5 passed per Section 3. Recorded, not assumed.
3. **Intent.** T2 Intent Type reads purchase. Research and comparison intent belong on cheaper surfaces with efficiency objectives; paying purchase prices for research traffic is the named error in v4.0 Section 3 family F14.
4. **Conversion.** T2 Brand CVR at or above Market CVR, with Benchmark Source recorded. Where the SQP row is broken, the fallback order is file reference target then account norm, and the fallback lowers the confidence score and routes the record to the human queue.
5. **Competitor credibility, both clauses.** Review Position Gap must not read FAR-BEHIND, which v4.0 Section 3 family F4 sets at under 50% of the top-5 median review count or 0.3 stars or more below it. Price Position Index must not read OUTPRICED, which family F5 sets above 1.25 on our price divided by median market price paid. BEHIND allows pushes on long-tail and mid-volume terms only while review velocity closes the gap. Note that D8 competitor depth reads MISSING on the live manifest, so this clause is evidenced on T2 Competitor Best Rank and Price Competitiveness Flag where present and flagged as unevidenced where not.
6. **Inventory.** G6 GREEN. YELLOW and RED route to v4.0 Section 8.3 S-I2 and S-I3 and are handled at P12, not by sizing around them.
7. **The THIN-margin clause.** Read T2 Margin \$/Unit. Below \$6 per unit the variation is THIN class, and O3 requires Expected CPC at or under margin dollars divided by 3. Compute it, write both numbers, and decline on failure. WE-3 works a live instance.
8. **Cost-to-close justified against revenue potential.** This clause needs the P3 worksheet to evaluate, so it is carried forward and resolved at P5 rather than guessed here.
9. Any clause failing closes candidacy with the failing clause named in T2 Situation. Set Objective to Paused-pending with the blocking reason and a re-entry criterion. Nothing exits the system silently.

P3. The sizing worksheet

Per surviving candidate, 15-20 minutes. This is the procedure the component exists for. Source: v4.0 Section 10 rule M2, with the chain definition in Section 3 family F14.

1. **Assemble the seven inputs** into T2: Organic Rank (current), Target Rank, DSTR (orders/day) at the target rank, validated CVR, Expected CPC \$, contribution per order from T0 zone C, and current daily orders on the term across organic and paid. Every input is read, never estimated at the keyboard.
2. **DSTR**. Set by competitor velocity analysis at the target rank, per family F14. It sizes the prize from market purchases, not from our current sales. Where D8 is MISSING the DSTR carries the unevidenced flag and the record says so.
3. **Order gap** = DSTR minus current daily orders. Write to T2.
4. **Required clicks per day** = order gap divided by expected CVR. Write to Required Clicks/Day. Where the term has 8 or more weeks of push history, correct the click requirement by the term's observed Rank Response Elasticity per the M2 amendment: evidence beats assumption. Below 8 weeks the elasticity correction is unavailable and the worksheet says so.
5. **Required spend per day** = required clicks times Expected CPC. Write the weekly figure to Required Budget \$/wk. Where A6 has no per-term Expected CPC, the account reference value in use is written into the worksheet as the assumption under test, not silently consumed.
6. **Duration**. Rank distance divided by the expected velocity for the term class. **AMENDMENT** v4.0 bands velocity only in positions per week and carries no per-class velocity reference for cold-push duration. Record the velocity assumption explicitly on the worksheet as an input under test, and log it against amendment item AM-4. Do not present it as a governed threshold.
7. **Total push budget** = required spend per day times duration.
8. **The loss ceiling, computed not estimated**. M2 defines the weekly loss ceiling as ACOS above break-even times projected sales, capped and dated. Compute the push ACOS as required spend per day divided by projected daily sales, subtract break-even ACOS from T0 zone C, multiply by projected weekly sales. Write to Loss Ceiling \$/wk with its date. **AMENDMENT** M2 requires the ceiling as a field and requires a stop-loss as a field, but no v4.0 scenario fires on the ceiling being reached; S-R5 fires on velocity, not on cumulative loss. See AM-1.
9. **State what the worksheet could not see**. Write into T2 Situation every manifest gap the sizing rode over: the D1 provisional margin, an unevidenced DSTR under D8, an account-reference Expected CPC under A6, an absent elasticity correction. Under G8 the analysis states what it could not see before it recommends anything, and confidence theater is a process violation under v4.0 Section 13.3 rule J7.
10. **Run the ceiling read**. Compute Max Profitable CPC as contribution times CVR and divide Expected CPC by it. At or under 0.8 the row is SAFE, 0.8-1.0 AT-CEILING, 1.0-1.5 PUSH-ONLY, above 1.5 UNJUSTIFIABLE outside a sized push. Above 1.5 inside a sized push, v4.0 Section 8.4 S-R10 makes the cheap-surface arbitrage check mandatory this cycle: where video CPC on the term is at or under 0.6 times the SP exact CPC, that surface enters for top-of-search presence and the SP exact steps down one ladder rung. The check is recorded on the worksheet whether or not it fires.
11. **Write the prediction**. Rank X to Y by the dated horizon, with the ETA checkpoint date. This is the row the push is scored against at P13 and it is written now, before any money moves.
12. **Version the worksheet** into the T2 plan block. It travels with the push, and every later decision in this document cites it by version rather than recomputing it.

P4. Confidence scoring and step modulation

Per sized candidate, 4-6 minutes. Source: v4.0 Section 10 rule M10 and Section 7 rule D9.

1. Compute the decision confidence score from its three stated drivers: sample size, benchmark source tier, and manifest completeness. Bands are 75 or above HIGH, 50-74 MEDIUM, below 50 LOW.
2. At LOW confidence, M10 **bars pushes outright**, halves every lever and budget step and shortens review horizons to one week. Under D9 the only legal verdicts at LOW confidence are extend-test, data-repair tasks and gate-driven actions. A LOW-confidence candidate does not proceed to P5.
3. At MEDIUM, standard step sizes apply. At HIGH, the upper end of every step range is permitted. Uncertainty buys smaller moves, never bigger ones.
4. With T7 completeness at 0.571 and benchmarks falling back below the SQP tier, the confidence composite is structurally depressed account-wide and more candidates fail this step than the rule anticipated. That is a data problem with a named owner, not a reason to relax the band.
5. Record the score and its three drivers on the record. Every Master Action Record carries it.

P5. Queue submission and the funding decision

Per candidate, 5-8 minutes. Queue ranking and envelope law belong to #31; this procedure is the submission and the receipt.

1. Compute the M5 ordering inputs from the worksheet: revenue potential at target rank divided by cost-to-close-gap. Write cost-to-close to T2 Cost to Close Gap \$.
2. Resolve the O3 clause carried forward from P2 step 8. Where cost-to-close is not justified against revenue potential at the M5 ordering, candidacy closes here with the concession logged.
3. Submit the worksheet to the queue surfaced in T0 zone H. Ties break by margin-dollar class and earning-potential class per M5. This SOP supplies the numerator and denominator; it does not rank the queue.
4. Check the product-level posture before expecting a release. Under v4.0 Section 8.1, S-A2 requires sign-off above the PPC Lead on new pushes while the product's total ad-cost ratio reads ELEVATED, S-A3 grants no new Ranking objectives at all during a push freeze, and S-A4 enters Recovery Mode, where R3 permits new Ranking objectives only on RECOVERY-class terms at HIGH confidence.
5. Also check S-A10: a CONSTRAINED earning-potential class bars ranking pushes outright and routes velocity requirements to the deals and pricing processes rather than to a PPC subsidy.
6. On release, read the dated release from T0 zone H and proceed to P6. Without a dated release there is no push, and spend moving toward rank in its absence is the F1 condition, not an informal start.

P6. Execution standup

Same day as the dated release. 20-30 minutes including staging. Campaign mechanics execute at SOP-12 P1.

1. Issue the standup order to SOP-12 carrying five values from the worksheet and nothing else: the dedicated single-keyword exact structure, the starting bid the worksheet carries, the daily budget as required weekly budget divided by 7, the objective tag Ranking, and the horizon date. **This SOP never writes a bid value of its own invention;** a creation with no verdict-carried bid returns as incomplete.
2. Confirm the bidding strategy is fixed. Strategic Doctrine v1.0 Section 6.3 sets fixed for ranking pushes because consistency of pressure is the point, and records that the April 2026 platform bidding change weights placement-level conversion probability far more aggressively, with category CPC inflation reported at 18-27% against Q1 baselines under dynamic raises.
3. Set placement modifiers at zero on all three groups unless placement-level evidence for this term already exists at #28. The modifier is earned rather than defaulted, and B7 reads MISSING account-wide, which suppresses increases.
4. Confirm walls are built in the same working session under #29 rules. A push whose source campaigns still serve the term buys the same click twice.
5. **Log the taper trigger at standup, on the day the push opens.** Write the achievement condition, the stability requirement reference and the ladder handoff into the T6 row now. Strategic Doctrine v1.0 Section 14 names spend continuing after rank achievement as an account failure class and names this logged trigger as the doctrine that prevents it.
6. Write the T6 Decision Log row: Rule / Scenario Fired, Change (prior -> new), \$ Impact / wk, Expected Outcome Window, Review Date, Owner / Tier. Tier is human-mandatory: v4.0 Section 12 rule P5 puts every 500-plus search-volume action and every structural change there.
7. Deployment routes through #38 with the T+1 verification pass. A landed-value mismatch is that component's five-branch disposition, and the push does not count as started until the diff proves it landed.

P7. Daily in-flight block

Daily on every live push, twice on deployment days. 6-10 minutes per push. Source: v4.0 Section 14 rule V4, daily alerts, pacing and gate checks with oversight-class checks during pushes.

1. **Pacing.** Read T2 In-Budget % and cumulative spend against the worksheet's daily line. Below 70% in-budget the campaign is truncated under G3, the budget lever is evaluated before any bid conclusion, and the week's bid read for the row is void.
2. **Rank.** Read the daily rank series and plot against the velocity line the worksheet assumed. D9 reads PARTIAL on head terms only, so a push outside that coverage has weekly rank resolution and the checkpoint arithmetic at P8 carries the coarser interval.
3. **Loss line.** Cumulative spend minus orders times contribution, compared against the dated weekly ceiling. Surface the number; the scenario call is v4.0's and is made at P8, not here.
4. **Gate flips.** Re-read eligibility, inventory band and event mode. Any flip routes to P12 the same day.
5. **No bid changes, ever, in this procedure.** Pacing interventions touch budget only. Bid decisions live in the weekly pass and in verdict-carried ladders. Every intervention writes a T6 line.

P8. Weekly in-flight read

Weekly per live push, 12-18 minutes. The scenario call is v4.0's; this procedure assembles the evidence in the order the rules require.

1. **Gates before metrics.** Re-run Section 3 on the row. A failing gate executes the gate's lever and defers the metric-driven verdict to the next clean cycle. v4.0 Section 12 rule P1 puts gates above diagnosis and diagnosis above signals.
2. **Write the five-element story line** with numbers, per v4.0 Section 4 rule X1: position, trajectory, economics, cause, plan. A line missing an element is an incomplete analysis and rule J2 rejects the record.
3. **Placement precedes bids.** Read T2 TOS IS %. Below the 30% floor that family F6 requires on ranking campaigns, v4.0 Section 8.4 S-R7 fires and the placement problem is addressed before any bid verdict, because most clicks off top-of-search means the campaign is not driving rank. Where B7 reads MISSING the modifier half is not executable and the condition is logged as a blocked verdict class rather than worked around.
4. **Run the rank triad,** v4.0 Section 4 rule X4, on CVR against benchmark, clicks per day against required, and rank velocity. The eight-cell cube resolves to four outcomes: at or above benchmark with clicks short and velocity flat is a traffic deficit and is funded; at or above benchmark with clicks met and velocity above zero is working and continues; at or above benchmark with clicks met and velocity near zero across the review window is over-competitive and the stop-loss fires; below benchmark in any combination is fix-first, because no bid level closes a conversion deficit.
5. **Read the ETA.** Family F6 classes it WITHIN at or inside the review horizon, default 3 weeks, BEYOND past it, and NEVER at velocity of zero or below. BEYOND triggers the stop-loss conversation early, before the budget is gone, and pre-stages the renewal decision at P11 for the horizon date.
6. **Check saturation on any scale step.** Marginal ACOS above 1.5 times average ACOS means the last dollars are buying worse traffic; freeze at the prior spend level and re-test in 2 weeks. This applies equally to budget increases, bid increases and modifier increases.
7. **Read the loss line against the ceiling** and record the percentage consumed. Where the ceiling is reached, see AM-1: no v4.0 scenario fires on it, so the row routes to the human-mandatory tier with the arithmetic attached rather than to an invented rule.
8. Write one verdict from Section 5, cite its scenario, and stage per #38. A staged row missing its scenario identifier does not deploy.

P9. Stop-loss execution

On the S-R5 condition. 10-15 minutes. Source: v4.0 Section 8.4 S-R5.

1. Confirm all three S-R5 conditions, not one: sufficient clicks and budget consumed per plan, CVR at or above benchmark, and rank velocity near zero across the review window, or observed elasticity near zero after adequate volume. A push failing on clicks is a traffic deficit and a different verdict entirely.
2. Confirm the placement check at P8 step 3 already ran. A stall diagnosed without the top-of-search read is the F4 condition, because a stall at full traffic on the wrong placement is not an auction verdict.
3. Issue STOP-LOSS: CONCEDE OR DEFER. The push objective ends here; the term re-classes and the budget returns to the envelope for M5 reallocation at #31.
4. **Concede** transitions the term to Conversions at the Max Profitable CPC ceiling. Compute it from contribution times realized CVR and route the bid step to SOP-12. Bid changes obey the lever limits at #28 and v4.0: no more than 50% per change with a \$0.50 floor.
5. **Defer** requires a new worksheet through the queue. There is no informal continuation, no reduced-rate version of the same push, and no carrying the same worksheet at a lower budget.
6. **AMENDMENT** v4.0 states no execution deadline for a fired stop-loss and no cool-down before the same term may be re-queued. v1.0 asserted same-day execution and an immediate re-queue path. Both are recorded as AM-2 and AM-9 rather than carried as thresholds.
7. Write the concession to T6 with the realized figures, and run P13 on the closed push.

P10. Achievement, stability and taper release

On the target being touched, then again at the stability confirmation. 8-12 minutes each. Source: v4.0 Section 8.4 S-R6 and Section 11 rule T1.

1. **Target touched** starts the stability measurement. Record the touch date in T6. Touching the target is not the achievement condition; holding it is.
2. **The stability requirement is 2 consecutive weeks at or under target rank**, per S-R6 and T1. Both rules state the same interval and it is the only stability figure in v4.0.
3. Hold the levers still through the window. The measurement reads rank at push pressure, and moving a bid or budget inside the window voids what is being measured. **AMENDMENT** The freeze itself is not governed by v4.0, and neither is the spend it authorizes. WE-1 computes what the window costs on a live sizing: \$1,545.60 of spend and \$604.80 of expected loss that the M2 block never sized, because M2 sizes to the horizon while S-R6 measures for two weeks past it. See AM-3.
4. **On confirmation, issue TAPER: RANK ACHIEVED, STEP BIDS DOWN.** The trigger logged at standup now executes. This is not a judgment call, and a rank-achieved row still carrying push-level spend is the F6 condition and the account's most expensive known failure class.
5. Hand off to SOP-12 P4 and the #28 ladder. This SOP supplies the achievement date, the push bid, and the worksheet reference. It supplies no rung, no step size and no floor.
6. Retag the objective from Ranking to the post-objective state per T1. Both the tag and the campaign-name objective segment update in the same staged set.
7. Run P13 on the closed push. Achievement does not exempt a push from its post-mortem.

P11. The horizon renewal decision

On the horizon date, per push. 15-20 minutes. Source: v4.0 Section 6 rule O7 and Section 8.4 S-R16.

1. Confirm the horizon date has arrived and the target is not held. Objectives do not persist by inertia, and renewal is a decision rather than a default.
2. Read the realized velocity across the whole push, not the last week, and the realized CVR and cost against the worksheet's assumptions.
3. **Re-size on evidence.** Remaining rank distance divided by realized velocity gives remaining days; times required spend per day gives the remaining cost to close. Where 8 or more weeks of push history exist, the elasticity correction applies. **AMENDMENT** Below 8 weeks, v4.0 does not state the re-sizing basis; using realized velocity directly is this SOP's stated method and is logged as AM-5.
4. **Apply the M5 test to the re-sized figure, not the original.** Compute re-sized cost-to-close per position against the figure at funding time. Renewal is legal only where the re-sized cost still clears M5 priority against the current queue.
5. Where it clears, renew with a new horizon date and a new worksheet version. Where it does not, transition to Conversions at ceiling and log the concession. WE-2 works an instance where cost-to-close inflates 3.25 times and the renewal fails.
6. Either outcome writes a T6 row with the arithmetic that produced it. A renewal without the re-sized number is the F12 condition.

P12. Suspension and resumption on gate flips

Same day on the flip. 10-20 minutes depending on the gate. Live pushes only.

1. **Inventory YELLOW.** v4.0 Section 8.3 S-I2 freezes the push at current spend with no scale steps, reviewed at the restock date. Rank preservation outranks rank progress.
2. **Inventory RED**, at 14 or fewer projected days to stockout: same day, taper to defense on the variation's terms, re-point ranking campaigns to the declared ranked backup, and write the pre-OOS baseline snapshot of rank, CVR and CPC per keyword. The snapshot is the recovery contract, and a push that stocks out without one leaves SOP-24 with no baseline to recover toward.
3. **Stockout despite the taper.** S-I4 moves push objectives to Paused-pending with re-entry on restock confirmed, and marks the trend windows for exclusion. The push is closed, not suspended, and P13 runs.
4. **Family conflict.** G11 bars scale on both siblings until the lead is assigned at #33. The push holds at current spend; it does not stop and it does not grow. WE-4 works an instance.
5. **Product posture.** S-A3 push freeze allows an existing push to continue only where velocity is above zero and the loss sits inside the ceiling; both conditions are tested and recorded, not assumed. S-A4 code red enters Recovery Mode, where R3 tightens human confirmation to any bid move above 15% and any budget move above \$25 per day, and caps step sizes at standard even at HIGH confidence.
6. **Event window armed.** The push continues and is amplified inside the window per the #32 plan. It is not originated there, and its in-window reads are governed by the plan's targets rather than by standing bands. G10 excludes the window and the 14 days following from the trend reads that would otherwise size or stop it.
7. Every suspension and resumption writes a T6 row naming the gate, the flip date and the resumption condition. A push suspended without a resumption condition is a push abandoned.

P13. Post-mortem, scoring and calibration

On every closed push without exception, 20-30 minutes. Source: v4.0 Section 14 rules V1 and V5.

1. Score the standup prediction HIT, PARTIAL, MISS or WRONG-DIRECTION against the horizon it named. Write to T6 V5 Score.
2. **Write the three calibration deltas separately**, because they have different owners and conflating them teaches nothing: realized velocity against assumed, realized CVR against assumed, and realized cost against budget. A push can miss on velocity while its CVR and cost assumptions were accurate, which is a sizing-reference problem rather than a forecasting problem.
3. Name which single assumption carried the error. Where the answer is the velocity reference, the finding attaches to amendment item AM-4, and repeated instances are the evidence that closes it.
4. Feed the realized rank movement per 100 incremental clicks into the elasticity series. At 8 or more weeks of history the series begins correcting the click requirement at P3 step 4 for this term class.
5. Where the push failed on a condition a missing feed made unfixable, attach the finding to the T7 row for that field group. WE-2 does exactly this against B7, and a feed with failed pushes attached to it argues its own priority better than a completeness score does.
6. Rule-level accuracy accrues per v4.0 Section 14 rule V5: any rule below 50% accuracy over 10 or more scored predictions enters review, and the review produces an amendment, a threshold change, or a documented defense.
7. **A closed push with no post-mortem is the F7 condition.** Calibration is the compounding asset of this SOP; without it the sizing assumptions stay wrong permanently and every future worksheet inherits the error.

5. Verdict Mapping

Every verdict issuance and every threshold behind it lives in v4.0. This table maps the verdicts that can land on a push-relevant row to what this SOP does when each arrives. Verdict strings are quoted verbatim from the Keyword Decision Record Template v3.0, Bands and Verdicts Reference sheet, closed verdict vocabulary, matching exemplar SOP-12 v2.0 Section 5. Where the v4.0 Section 13.1 string for the same state differs, it is carried in the cross-map column so neither citation is dropped. Zero local thresholds appear in this table.

Two states the KDR vocabulary cannot express. The KDR's twenty-five contains no member for a working push continuing unchanged, and no member for a family-level lead assignment. Both states arise in this component every cycle. For these two rows only, the string is quoted verbatim from **v4.0 Section 13.1 The Closed Verdict Vocabulary** with the source named in place, marked **v4.0** below. SOP-23 v1.0 filled the first gap by writing a bare **CONTINUE**, which is a member of neither vocabulary; that string is retired at v2.0. Two further v1.0 strings, **DIAGNOSIS BEFORE MONEY** and a bare **TAPER**, are retired for the same reason.

Verdict, quoted verbatim	v4.0 Section 13.1 cross-map	Scenario	What this SOP does
RANK PUSH: FUND SIZED PUSH	RANK PUSH AT CONTROLLED CPC	S-R2, S-R3	P2 gauntlet, P3 worksheet, P4 confidence, P5 submission. On the dated release, P6 standup and P7 arms the same day. The worksheet version travels with the push.
RECOVERY PUSH	No §13.1 member; v4.0 routes S-R13 into RANK PUSH AT CONTROLLED CPC	S-R13	P2 step 1 classes the term RECOVERY and hands to SOP-24, which owns loss diagnosis, baseline restoration and recovery-adjusted sizing. This SOP supplies the class, the proven ceiling and the pre-OOS baseline, then stops.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	No §13.1 member; S-R3 verdict text	S-R3	The named lever only, on the dedicated campaign, executed at SOP-12. Never both levers on one row in one cycle. Re-check the ETA in 2 weeks per S-R3.
CONTINUE; VERIFY MOMENTUM v4.0	Quoted from §13.1 in place	S-R2	No lever moves. Velocity and marginal checks are scheduled at P8. Where utilization sits below the 80% ranking target, S-R2 permits a single budget step of 20-30% and the marginal read runs at that step.
PLACEMENT FIX FIRST	No §13.1 member; S-R7 verdict text	S-R7	Bids hold. The modifier and budget program belongs to #28; this SOP deploys only the campaign-level rows #28 stages back. While B7 reads MISSING the modifier half is suppressed account-wide and the residual is logged as a blocked verdict class on T7.
FIX CONVERSION FIRST; DEFER PUSH	Identical string	S-R4	Cap at maintenance. The push does not fund. Route by funnel drop point with the price check first, and set Paused-pending with re-entry at CVR back to benchmark. No bid level closes a conversion deficit.
FIX BUDGET (TRUNCATED)	No §13.1 member; G3 response	G3	Budget lever only at P7 step 1. The week's bid read for the row is void and resumes on the next untruncated day.
UNLOCK TRAFFIC	No §13.1 member; S-R1 verdict text	S-R1	Applies pre-push at fewer than 15 clicks in 7 days with healthy budget and GREEN inventory; incremental bid steps of 15-25% to unlock traffic, and no rate judgments until sufficiency.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	TAPER TO FLOOR (POST-OBJECTIVE)	S-R6	P10 in full. The trigger logged at standup executes; the handoff goes to SOP-12 P4 and the #28 ladder. This SOP names no rung and no floor.
HOLD; DEFEND LIGHTLY AT RANK	No §13.1 member; the §13.1 near-neighbour is DEFEND AT FLOOR	S-R9	Applies where rank is at target but CVR sits below benchmark: hold at the minimum viable CPC while the conversion deficit routes to fix-first in parallel. Do not scale a rank held on weak conversion.
DEFEND AT FLOOR	Identical string	S-R9, S-D1	Post-taper state owned by SOP-25 and SOP-26. This SOP has already exited; listed so the verdict does not fall into a seam.
STOP-LOSS: CONCEDE OR DEFER	Identical string	S-R5, S-R16, S-Y9	P9 in full. Concede transitions to Conversions at ceiling; defer requires a new worksheet through the queue. Freed budget returns to the envelope and re-allocates at #31.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	Identical string	S-R11	The verdict for spend that looks like a push and is not one. Immediate ladder toward ceiling. A rank case may be written and sized afterwards, but aggression does not resume before it exists.
ADD EXACT + RESTRUCTURE	Identical string	S-R12	Candidacy defers at P1. Execution is SOP-12 P1 or P2 plus steering walls at #29. The candidate re-enters at the next cycle once sufficiency accrues on the new exact.
RE-POINT VARIATION	Identical string	S-I3, S-I6	P12 step 2 same day, with the pre-OOS baseline snapshot written before the swap. Execution at SOP-12; the backup chain is the declared variation priority flow.
ENTER SBV ARBITRAGE	FORMAT SUBSTITUTION (SBV/SD)	S-R10	Mandatory check at P3 step 11 whenever a live push runs above 1.5 times its ceiling. This SOP executes only the SP-exact single ladder step down; the video build belongs to components 17-19.
PAUSED-PENDING (REASON + RE-ENTRY)	PAUSE (LEVERS EXHAUSTED)	S-R15, S-I4	Every closed candidacy carries a reason and a re-entry criterion. A pause without both is incomplete and returns to the pass.
QUARANTINE: DATA REFRESH	FLAG: DATA REPAIR	G1, G8	No push sizes on a quarantined row. The refresh task opens the same day with an owner and a date.
PORTFOLIO CALL: ASSIGN LEAD v4.0	Quoted from §13.1 in place	G11, S-B2	Both sibling candidacies hold at P2 until #33 declares the lead. This SOP supplies each sibling's cost-to-close for the family M5 ordering and executes neither side of the declaration.
INCREMENTALITY LADDER	No §13.1 member; v4.0 carries the mechanism at X12	X12, S-A5	Post-objective, owned by SOP-26 and the measurement layer. Listed as a boundary: a push is never the vehicle for an incrementality test, because the test measures what spend above the floor buys and a push is spend below the rank.

6. Worked Examples

All four run on the standard training base, disclosed here because live per-variation economics are not attached: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%. This is external dependency D1 in the Ledger, open, and it caps this component by name. Every figure below was computed and asserted in code with float tolerances before it was written here. WE-1's anchors are carried forward unchanged from the Economics Worked-Example Set case E3 so that the two documents cannot drift apart.

WE-1 · A clean sized push, and the two governance gaps it exposes

Worksheet at P3. Term at organic rank 22 against a target of 8, distance 14 positions. DSTR at rank 8 reads 4.0 orders per day; current daily orders across organic and paid read 1.6. Order gap 2.4. Validated CVR 5.0% on 100-plus clicks. Expected CPC \$2.30.

Required clicks/day = $2.4 / 0.050 = 48.0$

Required spend/day = $48.0 \times \$2.30 = \110.40

Duration = 14 positions / 1.0 position per day = **14 days** (velocity is an assumption under test, per AM-4)

Total push budget = $\$110.40 \times 14 = \$1,545.60$

Paid orders/day = $48.0 \times 0.050 = 2.4 \rightarrow$ contribution recovered = $2.4 \times \$28.00 = \$67.20/\text{day}$

Expected loss/day = $\$110.40 - \$67.20 = \$43.20$, over 14 days = **\$604.80**

The ceiling, computed by the M2 formula rather than estimated. Push ACOS = $\$110.40 / (2.4 \times \$80.00) = 57.5\%$, inside the 50-80% ranking envelope. ACOS above break-even = $57.5\% - 35.0\% = 22.5$ points. Projected weekly sales = $2.4 \times 7 \times \$80.00 = \$1,344.00$. Weekly loss ceiling = $0.225 \times \$1,344.00 = \302.40 . Across the 14-day horizon that is \$604.80, which reconciles exactly to the expected-loss total above. The E3 case sets the operating ceiling at \$650, which sits \$45.20 above the computed figure with no governed rule for the margin; that rounding is amendment item AM-10.

Ceiling read at standup. Max Profitable CPC = $\$28.00 \times 0.050 = \1.40 . Expected CPC \$2.30 divided by \$1.40 is 1.64 times the ceiling, above the 1.5 line, so the row is UNJUSTIFIABLE outside a sized push and legal only inside this one. S-R10 therefore makes the cheap-surface check mandatory this cycle: the threshold is $0.6 \times \$2.30 = \1.38 , and video CPC on this term reads \$1.71, so the arbitrage does not open. Recorded on the worksheet either way.

Week 1 read at P8. Rank 14, so 8 positions in 7 days, 1.14 per day against the 1.0 line. Spend \$756.00 against the \$772.80 line, 97.8%, in-budget 100% so G3 is clean. Clicks 329, realized CPC \$2.30, orders 17 against 16.8 expected, CVR 5.17%. Clicks per day 47.0 against 48.0 required, 97.9%. ACOS 55.6%. Cumulative loss = $\$756.00 - (17 \times \$28.00) = \$280.00$, which is 92.6% of the one-week ceiling. Top-of-search share 41%, above the 30% floor. ETA: 6 positions remaining at 1.14 per day is 5.3 days, WITHIN the horizon. The rank triad returns working, so the verdict is **CONTINUE; VERIFY MOMENTUM**. No budget step is taken: utilization is already at 100%, so the S-R2 scale clause does not apply.

Day 13 and the collision. Rank 8 is touched on day 13 and the stability measurement starts. S-R6 requires 2 consecutive weeks at or under target, so confirmation lands on day 27. Two gaps become arithmetic at this point:

- **The window is unsized.** Holding levers still for 14 days at push pressure spends a further **\$1,545.60** and accrues a further **\$604.80** of expected loss. M2 sizes to the horizon; S-R6 measures for two weeks past it. Cumulative expected loss by day 27 reaches \$1,166.40 against a \$650 ceiling, exceeding it by **\$516.40**. This is AM-3.
- **Nothing fires on the ceiling.** Cumulative loss reaches \$650 on day 15.0, two days into the window. S-R5 requires velocity near zero and this push hit its target, so S-R5 does not fire. M2 requires the ceiling as a field and a stop-loss as a field, but v4.0 carries no scenario that fires on the ceiling being reached. The row routes to the human-mandatory tier with the arithmetic attached. This is AM-1, and it is the single most consequential gap this rebuild found.

Close. On day 27 rank holds at 8 or better across both weeks, TAPER: RANK ACHIEVED, STEP BIDS DOWN issues, and the trigger logged on day 0 executes into SOP-12 P4 and the #28 ladder. P13 scores the prediction HIT at day 13 against a day-14 horizon, and records the three deltas: velocity 1.08 per day realized against 1.0 assumed, CVR 5.06% against 5.0%, cost \$1,435 against \$1,545.60 budgeted. The velocity reference for this term class tightens on the evidence.

WE-2 · The edge case: every input right, the push still fails, and the failure is traced

Worksheet at P3. Term at rank 19 against a target of 6, distance 13. DSTR at rank 6 reads 5.0 orders per day; current 2.0. Order gap 3.0. Validated CVR 4.8%. Expected CPC \$2.60. Push bid \$2.74.

Required clicks/day = $3.0 / 0.048 = 62.5$

Required spend/day = $62.5 \times \$2.60 = \162.50

Duration = $13 / 0.62$ per day = 20.97 → planned **21 days**, which sets the horizon at the 3-week default review window

Total push budget = $\$162.50 \times 21 = \$3,412.50$

Push ACOS = $\$162.50 / (3.0 \times \$80.00) = 67.7\%$, inside the 50-80% envelope

Weekly loss ceiling = $(0.677 - 0.350) \times (3.0 \times 7 \times \$80.00) = \$549.50$, three-week total **\$1,648.50**

Max Profitable CPC = $\$28.00 \times 0.048 = \1.34 ; Expected CPC is 1.93 times ceiling, so S-R10 applies and the arbitrage threshold is $0.6 \times \$2.60 = \1.56

Week 1. Rank 18: 1 position in 7 days, 1.0 per week, which sits in the FLAT band. In-budget 100%, so this is not a truncation problem. Clicks 432, 61.7 per day against 62.5 required, 98.7%, so this is not a traffic deficit. CVR 4.86%, at benchmark, so this is not a conversion problem. Spend \$1,124.00 against the \$1,137.50 line, 98.8%. Cumulative loss \$536.00, which is 97.5% of the one-week ceiling. **Top-of-search share reads 21%, below the 30% floor.**

S-R7 fires and outranks the triad read: the placement problem precedes bid verdicts, because most clicks off top-of-search means the campaign is not driving rank. Verdict: **PLACEMENT FIX FIRST**. But T7 shows B7 placement splits MISSING, and #28 suppresses modifier increases account-wide until that feed lands. The campaign is not truncated, so the budget half of S-R7 has nothing to correct. **The verdict is issued and cannot be executed.** It is logged as a blocked verdict class against the T7 B7 row, and it becomes the named risk on the renewal decision.

Week 2. Rank 16: 2 positions, 0.29 per day. Cumulative clicks 866, orders 42, CVR 4.85%, spend \$2,251.00, cumulative loss \$1,075.00, which is 97.8% of the two-week ceiling. Top-of-search still 22%. ETA: 10 positions remaining at 0.29 per day is 35 days, landing on day 49 against a day-21 horizon, so the ETA reads BEYOND. Family F6 requires the stop-loss conversation to open now rather than at the horizon. Verdict: **CONTINUE; VERIFY MOMENTUM**, with the ETA flag raised and the renewal decision pre-staged for day 21.

Week 3, the horizon, at P11. Rank 15. Across the full push: 4 positions in 21 days, **0.190 per day realized against 0.62 assumed, 30.7% of plan.** Spend \$3,378.00 against \$3,412.50 budgeted, 99.0%. Clicks 1,299, orders 64, CVR 4.93%. Cumulative loss \$1,586.00, 96.2% of the three-week ceiling, still inside it.

Re-size on realized velocity: 9 positions remaining / 0.190 per day = **47.25 days**

Additional spend to close = $47.25 \times \$162.50 = \$7,678.13$

Additional expected loss = $47.25 \times \$78.50 = \$3,709.13$

Cost-to-close at funding = $\$3,412.50 / 13 = \262.50 per position

Cost-to-close re-sized = $\$7,678.13 / 9 = \853.13 per position, an inflation of **3.25 times**

M5 priority falls to **30.8%** of its value at funding time

Revenue potential at target rank has not changed, so the entire M5 movement sits in the denominator. The re-sized push does not clear the queue against the two candidates behind it. Verdict: **STOP-LOSS: CONCEDE OR DEFER**, conceding per S-R16 to Conversions at ceiling. Max Profitable CPC at realized CVR = $\$28.00 \times 0.0493 = \1.38 . The bid steps \$2.74 to \$1.38 in one move, a 49.6% reduction, inside the 50% per-change limit and above the \$0.50 floor. The \$162.50 per day budget line returns to the envelope and #31 re-allocates it.

Post-mortem at P13, and why this example matters. Prediction MISS. The three deltas separate cleanly: **CVR realized 4.93% against 4.80% assumed, 102.7% of plan; cost realized \$3,378.00 against \$3,412.50, 99.0% of plan; velocity realized 0.190 against 0.620, 30.7% of plan.** Two of the three assumptions were accurate to within 3%. The sizing chain was not wrong. The push failed because it never held top-of-search above the floor that v4.0 requires on a ranking campaign, so the clicks it correctly bought did not carry rank credit, and the one lever that would have corrected it was suppressed by a missing feed. The finding attaches to the T7 B7 row: a placement feed with a \$3,378.00 failed push attached to it argues its own priority more effectively than a completeness score of 0.571 does.

WE-3 · A THIN margin-dollar variation fails the realism gauntlet before any money moves

Candidate. A fitted-sheet term at organic rank 26 against a target of 8, on a variation carrying contribution of \$4.10 per unit at an AOV of \$13.90. Search volume clears the ranking floor, the term is indexed in title, intent is purchase and inventory reads GREEN, so the candidate arrives at P2 step 7 looking fundable.

Break-even ACOS = $\$4.10 / \$13.90 = 29.5\%$ · Target ACOS 14.8% · Max Acceptable 22.1%

Margin-dollar class: \$4.10 is below \$6.00, so the variation is **THIN**

O3 THIN clause: Expected CPC must be at or under margin dollars / 3 = $\$4.10 / 3 = \1.37

Expected CPC on the term reads **\$2.10**, over the cap by **\$0.73**

Max Profitable CPC = $\$4.10 \times 0.118 = \0.48 ; the CPC ratio is **4.34 times** ceiling, deep in the UNJUSTIFIABLE band

Clicks-to-loss = $\$4.10 / \$2.10 = 1.95$ clicks

Verdict: no push, and no smaller push. O3 fails on its final clause and candidacy closes at P2 step 9. The term takes PAUSED-PENDING (REASON + RE-ENTRY) with the reason recorded as the THIN clause and the re-entry criterion written as two alternatives with their arithmetic shown: **contribution at or above \$6.30 per unit** on the targeted variation, which is $3 \times \$2.10$ and also crosses THIN into STANDARD class, **or Expected CPC at or under \$1.37.**

What the operator must not do here. A 29.5% margin percentage looks unremarkable and invites a modest push instead. The binding constraint is margin dollars: at \$4.10 per unit the term tolerates fewer than two clicks per order, so no bid level makes the push work. v4.0 Section 4 rule X5 names this exactly. The failure is economic and routes to pricing and product, not to a bid.

WE-4 · Two siblings, one term: G11 holds both pushes before either is sized

Setup. A shared queen-size term appears as a push candidate on two family products in the same cycle. Sibling A runs an observed CPC of \$3.24 at organic rank 17. Sibling B runs \$1.48 at rank 6. The CPC ratio is **2.19 times**, above the 1.5 line that fires the structural inspection at v4.0 Section 4 rule X6. Neither instance carries an objective tag.

G11 fires at P2, before sizing. Two siblings carrying active push objectives on one term is a portfolio decision, not two product decisions, and neither push scales until the lead is assigned. Both candidacies hold. Per the Ledger, the push block is a SOP-23 precondition and the declaration belongs to #33.

The X6 inspection runs in its fixed order and finds the divergence is settings archaeology rather than auction reality: no objective tags on either instance; sibling A carrying a top-of-search modifier of +180% on a base bid of \$1.75, which authorizes an effective ceiling of $\$1.75 \times 2.80 = \text{\$4.90}$ per click, and under up-and-down bidding $\$1.75 \times 2.80 \times 2.00 = \text{\$9.80}$; sibling A on up-and-down where B runs fixed; and one legacy duplicate exact on A still bidding. The inflation is self-inflicted rather than market-driven, which routes to ladders and multiplier cuts rather than to acceptance of a higher Expected CPC.

Cost-to-close, supplied by this SOP for the family ordering. Family target rank 3, DSTR at rank 3 of 6.0 orders per day, common velocity assumption 0.5 positions per day.

Input	Sibling B	Sibling A
Current daily orders / order gap	3.6 / 2.4	1.1 / 4.9
Validated CVR	5.5%	5.1%
Required clicks/day	43.6	96.1
Expected CPC	\$1.60	\$2.90
Required spend/day	\$69.82	\$278.63
Distance / duration	3 positions / 6 days	14 positions / 28 days
Cost to close	\$418.91	\$7,801.57

Outcome. Revenue potential at target rank is materially equal, since it is the same term at the same target on one family, so the entire M5 ordering sits in cost-to-close, where A is **18.62 times** more expensive than B. The declaration returns **PORTFOLIO CALL: ASSIGN LEAD** with B as lead. B's candidacy proceeds to P3; A's closes with the concession logged and A holds Conversions at ceiling. A's structural repairs, being campaign-level, execute at SOP-12 and #28: consolidate the duplicate, cut the modifier to its evidence level, move to fixed bidding, and tag the objectives. The prediction carried on the declaration is family blended CPC on the term down 35% within 2 weeks at unchanged family rank capture.

Why this is the expensive case. Without G11 both pushes fund and the two products bid against each other with the same company's money in the same auction, raising the family's cost-to-close on a term one product would have won more cheaply. The divergence is recorded in the June 30, 2026 review and a family root ownership question remains open, so the gate protects a live condition rather than a theoretical one.

7. Failure Modes and Escalation

Thirteen modes. Each carries the signal that detects it, the response, and the account incident that motivated it where one exists. Incidents are sourced rather than asserted; the seven carried forward from v1.0 are marked.

ID	Failure	Detection signal	Response	Motivating incident
F1	Unfunded drip push	Sustained above-break-even spend on a term with no worksheet version in the T2 plan block and no FUNDED entry in T0 zone H.	S-R11 returns an immediate ladder toward ceiling. The row is sized properly through P2 to P5 or it is bid down. Nothing between the two is legal.	v1.0 The account's named pattern, recorded in v4.0 Section 3 family F16: one term absorbing roughly a quarter of a product's ad clicks above break-even with no plan behind it. Named again in the June 30, 2026 review as concentration without plans.
F2	Push into a stockout	Days to stockout below the push duration plus the stability window, at candidacy or in flight.	Candidacy blocked at G6. In flight, P12 executes the pre-OOS taper and the baseline snapshot. Ranking into a stockout burns the spend twice: once buying the rank, once losing it.	v1.0 The pre-OOS baseline exists in v4.0 Section 3 family F3 precisely because a head term fell from rank 7 across a stockout and was later mistaken for a cold start.
F3	Pushing an unindexed or backend-only term	G5 fails at P2 step 2: indexing reads no, or listing placement reads backend.	Route to the indexing or listing owner. The push waits. No bid level ranks an unindexed term and a push on a backend-only term funds the platform rather than rank.	v1.0 v4.0 Section 3 family F4 records the listing-coverage rollout as the routing mechanism, and syntaxes scoring below 60% coverage are listing projects before they are bidding projects.
F4	A stall answered with money	A bid or budget lift staged on a flat-velocity row before the placement read and the triad have run.	P8 order is enforced: gates, then placement, then the triad. The lift returns to the pass until the read exists. A stall at full traffic is a conversion or placement problem that money makes worse.	v4.0 Section 4 rule X2 names the mechanism: raising bids against adequate impression share buys more of the same rejection. WE-2 is the worked instance.
F5	Stop-loss deferred informally	The S-R5 conditions all met with the push still live in the following cycle, and no new worksheet in the queue.	P9 executes. Any continuation is a new sized worksheet through the queue, never the same push at a lower rate. Escalates on first occurrence.	v1.0 v4.0 Section 3 family F13 records the syntax-level version: rank spend that renews without movement is what the velocity engine exists to end.
F6	Taper never fires after rank is achieved	Rank at or better than target with push-level spend still live and weeks-at-verdict accruing.	P10 immediately. The trigger logged at standup executes. The weekly pass carries a standing audit line for this pattern, shared with SOP-12 F5.	v1.0 The account's most expensive known class. v4.0 Section 16 Worked Pass 1 records a head term held at rank 2 against a target of 3 for five consecutive weeks with the objective expired three weeks earlier and the spend never reduced.
F7	Worksheet never scored	A closed push with no V5 score and no calibration line in T6.	P13 runs late. Calibration is the compounding asset of this SOP, and an unscored push leaves every future worksheet inheriting the same wrong assumption.	v1.0 v4.0 Section 14 rule V1 states the principle directly: unscored actions do not accumulate.
F8	Two siblings pushing one term	T2 family conflict live, or the sibling CPC view showing a second push objective on the same term.	G11 holds both at P12 step 4 until #33 declares the lead. Neither push scales; neither stops.	The June 30, 2026 review recorded the same term at \$3.24 on one sibling and \$1.48 on another with no stated reason for either, and a family root ownership question was still open at the first cycle. WE-4 works it.
F9	Push sized on a blended margin figure read as final	The worksheet's ceiling cites a T0 economics block whose margin table carries the provisional state, with no flag on the record.	The PROVISIONAL tag carries onto every ceiling-referencing action and appears on the record in the words the engine computed. Confidence theater is a process violation under rule J7.	External dependency D1 is open in the Ledger and names SOP-23 push sizing as one of the two things it blocks. The live manifest shows a blended June 2026 figure in use.
F10	Rank credited to a push that auto is actually carrying	Auto and broad carrying more than 40% of the term's clicks, or zero exact instances, on a term with a Ranking objective.	S-R12 bars any performance verdict until the exact stands up and the sources are steering-negated. The push re-enters after sufficiency.	v4.0 Section 16 Worked Pass 3 records a head term at 74,000 search volume with zero exacts, two autos carrying 78% of its clicks at \$1.12, rank 22 against a target of 5. A head term ranked by auto is unranked by control.
F11	Push read on a truncated week	In-budget below 70% on the push campaign, with a bid conclusion staged from that week's rates.	G3 voids the week's bid read and forces the budget lever first. A starved campaign's rates describe its worst hours rather than its potential.	v4.0 Section 3 family F10 names truncation as the step-zero precondition the previous file generation could not see at all.
F12	Horizon passed with no renewal decision	The ETA checkpoint date in the past, target not held, push still live, and no T6 row carrying a re-sized cost-to-close.	P11 runs immediately. O7 bars automatic continuation, and S-R16 requires the re-sized figure to clear M5 before renewal. A renewal without the arithmetic is not a renewal.	v4.0 Section 6 rule O7 was written because ranking objectives in this account persisted by inertia; family F2 records that a ranking-stage designation expires and must be renewed on rank evidence.
F13	A push opened inside an armed event window	A standup dated inside an armed window with no reference to the #32 event plan.	Windows amplify live pushes and never originate them. The standup reverts unless the event plan names it, and the trend-exclusion window extends from the revert date.	SOP-32 states the rule from its own side, and #32's boundary in the Ledger names push origination as this component's, precisely so the seam has one owner.

Escalation. F5 escalates on first occurrence, because a push past its stop-loss conditions and still spending is the emergency class. F1 findings attach to the weekly pass as a standing audit line shared with SOP-12 F5. F6 and F8 escalate to the PPC Lead with the T6 history attached. Any mode occurring twice in consecutive cycles on the same scope escalates regardless of its individual routing. Every item in the human-mandatory tier is confirmed before deployment without exception, and Recovery Mode tightens that tier further under rule R3.

8. Outputs and Artifacts

Gates clear with artifacts rather than assertions. Each row names the record, where it is stored, how long it is kept, and who consumes it next.

Artifact	Produced by	Storage and retention	Downstream consumer
The sizing worksheet, versioned	P3	T2 plan block on the product's Command Workbook; retained for the life of the term, superseded versions kept for calibration.	#31 for the M5 ordering; #28 for the authorized CPC envelope that its ceiling worksheet tests against; SOP-24 as the base its recovery adjustments ride on; P13 for the calibration deltas.
The realism record	P2	T2 situation and plan fields with the cycle date; retained through the term's next candidacy.	The next cycle's P1, so a declined candidate is not re-assembled from scratch; the listing, pricing and review owners where the failing clause was theirs.
The confidence score with its three drivers	P4	Written onto the Master Action Record; frozen into the cycle KDR file.	The review-tier router at rule P5; the quarterly confidence calibration at rule V6.
The queue submission	P5	T0 zone H push queue, funded and queued entries with their precedence reason.	#31, which ranks and releases; Account Rollup A2, which holds the cross-product queue.
The standup record with the logged taper trigger	P6	T6 Decision Log, append-only.	SOP-12 for execution; P10, which executes the trigger rather than re-deciding it; #38, which stages and verifies.
Daily pacing and rank lines	P7	T6, one line per intervention; alarm evidence lines retained through the push close.	The F11 truncation audit; the weekly read at P8, which inherits the week's pacing history rather than re-deriving it.
The weekly in-flight verdict with its scenario	P8	T2 verdict block, mirrored into the cycle KDR file as the frozen snapshot.	#38 as the staged change set; the next cycle's V5 scoring.
The stop-loss or achievement record	P9 or P10	T6 with the realized figures attached.	#28 and SOP-12 P4 on a taper; #31 on a concession, for the returned budget; SOP-25 on the post-objective handoff.
The re-sized cost-to-close	P11	T2 cost-to-close field, versioned against the funding-time value.	#31 for the re-ordering; the M5 queue, which compares it against current candidates rather than against its own history.
The pre-OOS baseline snapshot	P12 step 2	T2 pre-OOS baseline fields; retained for at least the trailing 12 months that the recovery classification reads.	SOP-24, which recovers toward it; the D10 classification at P2 step 1 on any future candidacy for the term.
The post-mortem with three calibration deltas	P13	T6 with the V5 score; the deltas retained indefinitely as the elasticity series.	P3 step 4 on the next worksheet for the term class; rule V5 rule-level accuracy; the training corpus under rule V8; the T7 row for any feed a failure was traced to.

Gate artifacts. In flight, the daily pacing line is the standing artifact and its absence means the push was not operated. At close, the post-mortem is the gate artifact and a push cannot be marked closed without one.

9. Skill-Conversion Block

Block	Specification
Input bindings	By catalog ID: A1 four-period comparability, A3 pacing and in-budget, A4 sufficiency buckets, A5 threshold flags, A6 objective and per-term Expected CPC, A7 eligibility, B1 indexing and listing placement, B3 rank pack, B5 coverage counts, B7 placement splits, B8 economics chain, B11 market context, D1 per-variation margin, D2 keyword-variation map, D6 sibling view, D8 competitor depth, D9 daily rank series. Field-level: T2 plan block for the worksheet, T2 rank block for class and velocity, T2 economics block for ceilings, T0 zone C for the spine and zone H for the queue, T5 for attribution, T6 for predictions and scores, T7 for the manifest. An absent feed suppresses its dependent execution class; the skill never improvises around a missing feed and never substitutes an account reference for a per-term value without writing the substitution onto the record.
Deterministic logic	The P3 chain in order: order gap, required clicks, required spend, duration, total budget, and the M2 ceiling formula of ACOS above break-even times projected sales. The ceiling read of Expected CPC against contribution times CVR, resolved to four CPC-ratio bands. The O3 gauntlet as a conjunction where any false closes candidacy. The X4 triad as an eight-cell resolution to four outcomes. The ETA class from remaining distance over observed velocity. M5 ordering as revenue potential over cost-to-close. The confidence bands and their step modulation. Scenario selection, verdict issuance and every threshold are v4.0's; the skill version-pins v4.0 and re-validates Section 5 on any increment before running a pass on it. No local thresholds.
Human checkpoints	Executable with audit: candidacy assembly at P1, worksheet computation at P3, confidence scoring at P4, daily pacing and rank lines at P7, the evidence assembly at P8, the re-sized arithmetic at P11, and post-mortem calibration at P13. Human-mandatory , drawn from the Authority Matrix rows listed in Element 11: the funding decision at P5, the execution standup at P6, stop-loss execution confirmation at P9, the taper release at P10, the renewal decision at P11, and any suspension or resumption at P12. The matrix places the D on push funding with the CEO above the envelope and the PPC Lead within it, and every 500-plus search-volume action and every structural change sits in the human-mandatory tier regardless.
Cadence	Weekly for candidacy assembly and the in-flight read, inside the Monday cascade. Daily for pacing, rank and gate checks on every live push, twice daily on deployment days. Same day on inventory RED, eligibility flips and product posture changes. On the dated release for standup. On the horizon date for renewal. On close for the post-mortem. Recovery Mode replaces the weekly cadence with daily gate checks and twice-weekly mini-audits on the triage set.
Alarm specification	F1 to F13 each carry a detection signal from Section 7, an evidence-line format of row identifier, signal value and the gate or rule reference, and a routing tier. Same day to the operator queue: F2, F3, F9, F11. Into the weekly pass worklist: F4, F7, F10, F13. To the PPC Lead with T6 history: F1, F5, F6, F8, F12.
Outputs	Drafted worksheets with their manifest-gap statements, drafted realism records, pacing and loss-line entries, drafted weekly verdicts with scenario identifiers and five-element story lines, re-sized renewal arithmetic, drafted post-mortems with the three deltas separated, and F1 to F13 alarms with evidence lines. Every draft is queued for confirmation at its tier; nothing in this list deploys itself.

10. Version Governance

Owner: PPC Lead. **Review triggers specific to this component**, not a generic calendar:

- Any v4.0 increment touching Section 8.4, rule M2, rule M5, rule M10, rule O3, rule O7 or gate G6. Each of these is load-bearing in a named procedure, and the Section 5 mapping re-validates against the new verdict set before any pass runs on it.
- Any amendment item in Section 13 landing in v4.1.** AM-1 in particular changes P8 step 7 and P9 from a routed escalation into a fired scenario, which is a procedure change rather than a wording change.

- External dependency D1 landing. The per-variation margin table removes the PROVISIONAL tag from every ceiling in this document and makes the worked examples reproducible on live economics rather than the training base.
- Field group B7 landing. It makes S-R7 fully executable and closes the blocked verdict class that WE-2 traces a failed push to.
- Field group D8 landing. The DSTR input at P3 step 2 is defined as competitor velocity analysis, and it is currently unevidenced.
- Any change to SOP-24 P3, since that document extends this one and its recovery adjustments ride this worksheet.
- Any change to the #28 ladder mechanics, since P10 hands off into them.
- Quarterly at minimum, alongside the floor re-tests and family-level M5 that rule V4 places on that cycle.

v2.0 change note. Supersedes v1.0 in full. Procedures expanded 7 to 13, adding the O3 gauntlet as its own procedure, confidence modulation, the horizon renewal decision, and gate-flip suspension and resumption. Failure modes expanded 7 to 13. Worked examples expanded 1 to 4. Elements 11 and 12 added, neither having any model in the corpus. Three v1.0 verdict strings retired as members of no closed vocabulary: a bare CONTINUE, DIAGNOSIS BEFORE MONEY, and a bare TAPER. The v1.0 assertion that a breached loss ceiling executes a stop the same day without exception is retired as ungoverned and logged as AM-1 and AM-2. The v1.0 day-7 checkpoint is retired as a label and absorbed into the weekly read at P8, which rule V4 governs. All E3 sizing anchors and all seven v1.0 incidents carry forward, now sourced. Registered in the Supersession Map at publication.

11. Operator Ramp

Built from the Authority Matrix v2.0 rows that name this component's decisions. Each row becomes a competence the incoming operator demonstrates. Week 1 is observation and artifact familiarity, week 2 is supervised execution with the checkpoint holder reviewing, week 4 is unsupervised execution with the checkpoint holder auditing. Note the matrix vacancy rule: an unfilled role's decision right reverts to the PPC Lead and its execution to the PPC Manager until filled, so a ramp in progress does not leave a decision unowned.

Authority Matrix row	Rights on that row	Week 1: observe	Week 2: supervised	Week 4: unsupervised, audited
Push funding (the M5 queue)	CEO decides above the envelope; PPC Lead decides within it; incoming operator executes.	Sit the funding call on every candidate this cycle. Read T0 zone H and state, for each queued entry, which of the two decision rights applies and why.	Compute cost-to-close and the M5 ordering inputs for every candidate; the PPC Lead reviews the arithmetic before the call.	Own P2 through P5 end to end on assigned scope. The PPC Lead audits the ordering and the concessions rather than reproducing them.
Weekly verdicts, established US	PPC Lead decides; incoming operator and PPC Manager execute; brand-management owner reviews.	Shadow the full P8 read on two live pushes. Write the five-element story line for each without issuing a verdict.	Draft the weekly verdict with its scenario on every live push; the PPC Lead confirms before staging.	Issue weekly in-flight verdicts on assigned scope. The daily huddle review by the brand-management owner is a review, not a veto; disagreements route up.
Structure creation, all SOPs	PPC Lead decides or per SOP scope; incoming operator executes.	Watch one standup end to end, including the wall session and the T+1 verification return.	Issue the standup order at P6 with the five carried values; the checkpoint holder confirms before it reaches SOP-12.	Issue standups on assigned scope. Every creation on a funded push remains human-mandatory at deployment.
Bulk deployments	PPC Lead reviews; incoming operator and PPC Manager execute.	Read one full staged set and its T+1 diff. Identify which reason strings carry a scenario and which would fail staging.	Stage the change sets this SOP orders; the verification reconciliation is reviewed line by line.	Stage and reconcile unsupervised. A landed-value mismatch is routed same day rather than absorbed.
Sibling declarations, S-series	PPC Lead decides; incoming operator executes; brand-management owner reviews.	Read the family board and the open ownership question on at least one family root. State what G11 is currently blocking.	Supply both siblings' cost-to-close for a family ordering; the PPC Lead makes the call.	Detect and hold sibling conflicts at P12 step 4 without prompting, and supply the ordering inputs unasked.
Event plans	PPC Lead decides with the deals calendar owner; operator executes.	Read one armed window's plan and identify which live pushes it amplifies and under what allowance.	Execute P12 step 6 on an armed window under review: suspend the standing reads, not the push.	Operate live pushes through a full window and settle, unsupervised, escalating only plan deviations.
Code-red declaration and close	PPC Lead decides; incoming operator and PPC Manager execute.	Read the product posture rules and state, for each live push, what a push freeze and what a code red would each do to it.	Execute the P12 step 5 posture response on a product entering a freeze, with the velocity and ceiling tests recorded.	Apply the tightened Recovery Mode confirmation thresholds correctly without being reminded that they replaced the standing ones.
Launch-stage verdicts and batches	Launch operator decides and executes; PPC Lead reviews.	Identify which candidates on the board sit in launch scope and therefore route to a different decision holder.	Hand a launch-scope candidate over cleanly with its worksheet intact.	Route by scope without error. Executing a launch-scope push against the established-US decision right is a governance failure even where the arithmetic is correct.

Live-defense questions

These test reasoning behind the thresholds rather than recall of their values. The operator answers cold, with the workbook open and the criteria document closed.

1. A push has delivered every required click at the sized CPC, CVR is above benchmark, and rank has not moved in three weeks. Name the two things this could be and the order in which you must rule them out. Explain why raising the bid is wrong in both cases.
2. Two variations both read 32% ACOS. One is profitable and one is losing money. Explain how, and state which field on the worksheet tells you which is which before you look at anything else.
3. A push reaches its loss ceiling on day 15 having already touched its target rank on day 13. State what fires, what does not fire, and why this is currently an escalation rather than a rule.
4. Your worksheet assumed 0.8 positions per day and the push realized 0.3. CVR and cost both came in within 2% of plan. Which assumption failed, what do you change for the next worksheet on this term class, and what would you have to be careful not to change?
5. Explain why the stability window is two weeks and not one, in terms of what is being measured rather than what the rule says.
6. A term is a candidate on two sibling products in the same cycle. Explain what it costs the company to fund both, using the auction rather than the budget as your argument.
7. Top-of-search share reads 18% on a live push. The placement feed is not available. State what verdict issues, what part of it you can execute, and what you write down about the part you cannot.

8. A THIN-margin variation shows a 30% margin percentage, which looks healthy. Explain in one sentence why that number is nearly irrelevant to whether you can push the term, and name the number that is relevant.

12. Interface Contracts

Built from the Ledger's interfaces column for component #23, cross-checked against the Data Architecture for artifact names and tab locations. Every row states what happens when the handoff does not arrive, which is the clause the corpus has always been missing.

ID	Direction and counterpart	Artifact	Timing	Non-arrival consequence
I-1	Upstream from #11	The economics worked-example set as the arithmetic standard the worksheet follows, case E3.	Static; re-issues on any v4.0 change to sizing structure.	Without it, operators re-derive the sizing chain per push and the chain drifts between them. This SOP's P3 becomes the only written form of the arithmetic, which is a single point of failure the training layer was meant to remove.
I-2	Upstream from #41	The per-variation margin table with its freshness state, repointed into T0 zone C.	Monthly, plus event-driven refreshes, before the Monday packet-ready read.	Every ceiling in this SOP carries PROVISIONAL and every worked example runs on the training base. This is external dependency D1, open, and the Ledger names SOP-23 push sizing as one of the two things it blocks. Pushes still size, but no ceiling is defensible on live economics.
I-3	Upstream from #33	The family term-ownership declaration with the lead product and the support or stand-down posture.	On conflict detection; the family ordering runs quarterly.	G11 holds both sibling candidacies indefinitely. Neither push funds and the term stays unmanaged, which is worse than either sibling pushing alone. The push block is a SOP-23 precondition, so a stalled declaration stalls this component with no local workaround.
I-4	Upstream from #31 and the instrument layer	The dated funding release in the T0 zone H queue, and the envelope headroom figure.	Weekly within products; quarterly across the family.	No push starts. Sized worksheets accumulate unqueued and the operator faces pressure to start informally, which is precisely the F1 condition. Per Ledger conflict C3, #31 carries no procedural document, so a stalled release has no procedure to escalate against and routes to the PPC Lead by role.
I-5	Upstream from #32	The armed event window with its dates, affected campaign list, and any in-window allowance naming a live push.	T-14 at intake, T-1 at arming.	P12 step 6 cannot run. Live pushes are read against standing bands inside a window that distorts every input, producing stop-loss and scale verdicts on data v4.0 excludes from trend reads. The damage is silent because the numbers look readable.
I-6	Upstream from #39 and Data Operations	T7 manifest statuses for A1, A3, A6, B3, B6, B7, D1, D2, D8, D9.	Weekly, with the daily rank series refreshing daily.	P3 step 9 cannot state what the worksheet could not see, which turns a flagged sizing into a confident one. Under G8 that is the most expensive failure the system can produce: a confident recommendation on a silently incomplete packet.
I-7	Downstream to SOP-12	The standup order carrying structure, starting bid, daily budget, objective tag and horizon date; and later the taper trigger release.	Same day as the funding release; same day as stability confirmation.	A creation with no verdict-carried bid returns as incomplete and the release date passes unspent. On the taper side, non-arrival is F6: the push keeps running at push pressure on an achieved rank, which is the account's most expensive known failure class.
I-8	Downstream to #28	The achievement date, the push bid, and the worksheet's authorized CPC envelope, which the ceiling worksheet tests proposed modifiers against.	On taper release; on any in-flight modifier request.	The ladder starts from an unstated bid and the first rung is guessed, losing the floor the ladder exists to discover. In flight, #28 cannot test a modifier against a push envelope it does not have, so the ceiling worksheet falls back to the maintenance ceiling and blocks changes the push legitimately authorizes.
I-9	Downstream to SOP-24	The pre-OOS baseline snapshot of rank, CVR and CPC per keyword; the recovery classification; and the discovered floor and prior push CPC.	Same day on inventory RED; at close for the floor and CPC history.	SOP-24 has no baseline to recover toward and no bracket for its re-entry bid, so recovery is sized cold. The term misclassifies as NEVER-RANKED and loses the funding priority its proven ceiling earned, which is the misdiagnosis the recovery classification exists to prevent.
I-10	Downstream to SOP-25	The taper handoff: achievement date, held rank, and the worksheet reference.	On stability confirmation.	The normalization sequence starts without the achievement record, so the transition date is unknown and the retag never fires. The term reads as a live Ranking objective past its expiry and generates a false finding every cycle until corrected.
I-11	Downstream to #31	The released daily budget line on any concession or taper, and the re-sized cost-to-close on any renewal.	Same cycle as the verdict.	Freed budget stays notionally allocated to a push that has ended, so the queue under-reports available dollars and the next candidate is declined for budget that exists. This is the reallocation failure the opportunity-cost metric was defined to end.
I-12	Downstream to #38	Staged change sets carrying verdict and scenario identifier in every reason string.	Tuesday by default; same day on a dated release.	A staged row missing its scenario identifier does not deploy, so the change silently does not happen while the T6 row claims it did. The T+1 diff is the only thing that catches it, which is why the verification return is a contract and not a courtesy.
I-13	Downstream to T6 and the verification layer	Prediction rows at standup and at every in-flight verdict; V5 scores at every review date; the three calibration deltas at close.	At standup, weekly, and at close.	Unscored actions do not accumulate, so rule-level accuracy cannot compute and no rule in this SOP can ever be shown to be wrong. The framework loses the evidence standard it holds every other decision to.
I-14	Downstream to #29	The structure-event notice that a dedicated push campaign exists and needs its walls.	Same working session as the standup.	The source campaigns keep serving the term, so the push buys the same click twice at two CPCs and the rank credit is split across campaigns the worksheet never sized. The push then under-delivers against a plan that was arithmetically correct.

13. Amendment Register: numerals this SOP needs that v4.0 does not carry

Each item is a number or rule the procedures above require and that no attached document supplies. None is written into the SOP as a threshold. They are queued for v4.1 of the Criteria System, whose owner holds the decision right on threshold changes.

ID	The gap	Where it bites	What is needed, and the evidence available now
AM-1	No scenario fires on a loss ceiling being reached.	P8 step 7, P9, WE-1.	M2 requires the weekly loss ceiling as a field and a stop-loss as a field, but S-R5 fires on velocity near zero rather than on cumulative loss. A push can consume its full ceiling while velocity is healthy and no rule fires. Highest-priority item in this register. v1.0 filled the gap with a local rule that has no basis; that rule is retired here.

ID	The gap	Where it bites	What is needed, and the evidence available now
AM-2	No execution deadline on a fired stop-loss.	P9 step 6.	S-R5 states the stop-loss fires; it does not state by when it must be executed. v1.0 asserted same day. A push spending past a fired stop-loss is the F5 emergency class, so the interval matters.
AM-3	The stability window's spend posture is ungoverned, and its spend is unsized.	P10 step 3, WE-1.	S-R6 requires 2 consecutive weeks at target; M2 sizes only to the horizon. On the WE-1 sizing the window adds \$1,545.60 of spend and \$604.80 of expected loss beyond the plan, breaching a \$650 ceiling by \$516.40. Needed: whether the window is sized into the worksheet, and whether levers are frozen through it.
AM-4	No velocity reference for cold-push duration.	P3 step 6, every worked example.	Family F6 bands velocity in positions per week only; the duration calculation runs in positions per day. No per-term-class reference exists for a push with no elasticity history, and this assumption carried the entire error in WE-2. Needed: a reference by term class, or an instruction to treat duration as a checkpoint output rather than an input.
AM-5	No re-sizing basis below 8 weeks of push history.	P11 step 3.	M2's amendment corrects the click requirement by observed elasticity once 8 or more weeks exist. S-R16 requires a re-size at the horizon, which arrives at 3 weeks by default. Using realized velocity directly is this SOP's stated method and needs ratification or replacement.
AM-6	No maximum lag from funding release to standup.	Section 2, P6.	The release carries a date; nothing states how long the standup may trail it. A release that ages loses the queue position it won.
AM-7	No lag on the sibling-conflict hold.	Section 2, P12 step 4.	G11 bars scale until the lead is assigned but sets no clock on the assignment. Interface I-3's non-arrival consequence is an indefinite hold on both siblings, which is worse than either outcome of the decision.
AM-8	No interval for the post-mortem after close.	Section 2, P13.	V1 requires scoring; no rule states by when. F7 exists because post-mortems slip, and a mode with no clock cannot be audited.
AM-9	No cool-down before a conceded term re-enters the queue.	P9 step 5.	S-R5 permits defer with a new worksheet; nothing prevents the same term returning the following cycle with the same assumptions. Needed: either an interval or an explicit requirement that the re-entry cite what changed.
AM-10	No rule for setting the operating ceiling above the computed ceiling.	P3 step 8, WE-1.	The M2 formula returns \$604.80 on the E3 sizing and the case operates at \$650, a \$45.20 margin with no governed basis. Needed: round-up rule, or an instruction that the computed figure is the operating figure.
AM-11	No campaign-level pacing tolerance for a push.	P7 step 1.	G7 sets a 25% deviation alert against the day-of-week curve at product level. No campaign-level tolerance exists, and A3 currently reads PARTIAL with the day-of-week curve pending, so even the product-level test is not fully computable.
AM-12	No authorized CPC buffer under the 2026 auction.	P3 step 5, P6 step 2.	Strategic Doctrine v1.0 Section 6.3 instructs that push sizing carry a CPC buffer while the auction re-stabilizes and reports category inflation of 18-27% against Q1 baselines. That is a market observation, not an authorized numeral, and no buffer figure exists in v4.0. Sizing currently runs unbuffered.
AM-13	No push concurrency limit per product.	P5 step 3.	M5 orders the queue and the envelope caps total spend, but nothing limits how many pushes may run at once on one product. Concurrency affects the attribution of rank movement and the operator's daily capacity at P7.
AM-14	The closed vocabulary has no member for a working push continuing.	Section 5, P8.	The KDR's twenty-five contains no continue verdict and no lead-assignment verdict; v4.0 Section 13.1 carries both. This SOP quotes v4.0 in place for those two rows. The reconciliation belongs to the same release that resolves Ledger conflict C10.

Companions: Master PPC Decision Criteria System v4.0 (every threshold, band, scenario and verdict) · Keyword Decision Record Template v3.0 Bands and Verdicts Reference (the quoted verdict strings) · PPC Strategic Doctrine v1.0 Sections 1, 2 and 6 (the reasoning behind the rules) · PPC Economics Worked-Example Set v1.0 case E3 (the WE-1 anchors) · PPC Authority Matrix v2.0 (Elements 9 and 11) · Master PPC Data Architecture v1.0 and the Command Workbook (instrument mechanics) · SOP-12 (execution), SOP-24 (recovery), SOP-25 (post-objective), #28 (ladder), #29 (walls), #31 (queue and envelope), #32 (windows), #33 (family ownership), #38 (staging) · PPC Cross-Reference Ledger (component #23 row and conflicts C1 to C15).