

SOP-25: Transition and Maintenance Operations

Version 2.0 · 2026-08-08 · Owner: PPC Lead · Supersedes SOP-25 v1.0 (2026-07-24) in full · Thresholds, bands, scenarios and verdict issuance live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds · Rebuilt to the SOP Build Kit v1.0 twelve-element standard

1. Verdict and Scope

VERDICT: transition is the deliberate descent from push economics to sustainable economics without surrendering the rank the pushes bought, and it is a scheduled program with dates, not a drift. Version 2.0 expands v1.0 from 6 procedures to 14, from a 5-row verdict mapping to the full 25-string closed vocabulary, from 1 worked example to 4, and from 5 failure modes to 12, adds the operator ramp and the interface contracts the v1.0 standard never required, and corrects one v1.0 threshold that did not trace to v4.0. This SOP governs the product-level normalization the stage requires: stage-entry verification, the term-state inventory, the sequenced descent schedule, monthly milestone construction and the dual-limb checkpoint, weekly rung execution, the rank guard during descent, single-term slip reversion, cluster detection and schedule pause, accumulated-waste cleanup, the objective-mix rebalance, the event and Recovery overlays, and the exit handoff into SOP-26.

1.1 The Governing Rule: The Descent Is Priced Before It Runs

Everything in this document serves one rule. A transition is a dated schedule of named actions whose arithmetic is written before the first rung deploys, and a month is closed on two readings, not one: the actions that landed and the posture that moved. The Master PPC Decision Criteria System v4.0, Section 11 "Layer 6: Transitions, Stages, Modes (T1-T8) and Recovery Mode (R1-R6)", rule T1, sets the term-level transition: Ranking flips to post-objective at 2 weeks at or under target rank, then M3. Section 8.4 "Ranking / Re-Ranking Keyword Scenarios (S-R1 to S-R16)", scenario S-R6, states the same flip and names the cost of not running it: ranking spend that persists past this point is the definition of unmanaged spend. The PPC Strategic Doctrine v1.0, Section 1 "Verdict and First Principles", principle 5, supplies the reason the schedule is written down at all: structure that cannot be scored is opinion.

The account has the incident on file. v4.0 Section 16 "Three Worked Passes: The Full System on the Account's Real Failure Patterns", Worked Pass 1, is titled "the taper the account never runs": a head term at rank 2 against a target of 3, held five consecutive weeks, with the Ranking tag expired three weeks earlier and never transitioned, running CPC \$1.85 on a THIN \$4.10-margin variation at 48% ACOS against a 29.5% break-even. The layer walk found no failure to locate. The failure was a missing transition. This component exists so that finding is a scheduled event rather than a discovery.

1.2 Product-Level Normalization Versus Term-Level Taper

The distinction is the whole scope boundary and it is where v1.0 was thinnest. A taper is one term walking down a dated ladder. A transition is the product-level program that decides which terms taper, in what order, on what dates, funded from what, checked against what, and closed by what artifact. This SOP owns the program. It owns no rung mechanics at all.

Layer	What decides it	What this SOP does
The bar to taper a term	v4.0 Section 8.4, scenario S-R6, and Section 11, rule T1: rank at or under target for 2 consecutive weeks.	Reads the condition; never re-derives it, never waives it, never applies it early.
The size of one rung	v4.0 Section 10 "Layer 5: Sizing Mathematics (M1-M10). Every Action Carries Numbers", rule M3, inside the Section 9 "Layer 4: The Lever Library (L1-L10)" rule L2 limits.	Schedules the rung on a date. The step size is read, never chosen.
The exit target of a ladder	SOP-28 P5: the placement-CPC exit bid, being the blended placement CPC at the achieved position less the v4.0 decrement.	Records the target on the schedule row so the rung count is knowable in advance.
Which rung deploys, and how	SOP-12 P4 executes; #38 stages, uploads and verifies at T+1.	Issues the dated instruction and consumes the verification return.
The order and dating of the whole set	This SOP, P4.	The sequencing is the deliverable no adjacent component builds.

1.3 The Three Entry Routes and the One Exit

A product arrives in this stage by one of three routes, and each route hands over a different amount of prior work. Nothing enters by assertion: an entry without its artifact is not an entry.

Route	The condition, and where it is stated	What arrives with it
Taper handoff from a completed push	SOP-23 P6.3: hold confirmed, TAPER issues, handoff to SOP-12 P4 and the #28 ladder. The Cross-Reference Ledger's component-23 row records the seam: taper exits into SOP-25.	The push worksheet from T2 K7, the standup-logged taper trigger, the post-mortem calibration line, and the term's prior push CPC.
Stage renewal at the monthly slot	v4.0 Section 11, rule T7: monthly, every product's stage tag is confirmed or transitioned; expired Ranking-push tags without rank evidence auto-transition to Transition stage and its TACOS band, which mechanically triggers the Section 8.1 product-level posture checks.	A stage tag and a band, and nothing else. This is the thinnest entry and the one that most often finds unfinished push structure.
Objective expiry on the terms themselves	v4.0 Section 6 "Layer 1: Objective Assignment (O1-O8)", rule O7: every Ranking objective carries a horizon date and renews only on evidence; otherwise it transitions per T1 or exits per its stop-loss.	A worklist of terms whose objective horizon passed, each carrying its own T6 history.

The exit is single. v4.0 Section 11, rule T7, is the decision slot; SOP-26's Section 3 states the artifact it requires, which is this SOP's floor register, and its Section 2 names the trigger as the SOP-25 handoff. Operating mature without recorded floors is operating blind, so the register is the exit gate and not a courtesy.

1.4 Boundaries With Adjacent Components

Every excluded decision names the component that owns it. The Cross-Reference Ledger is the check.

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, scenario, gate, lever and verdict, including the stage bands and the transition condition. This SOP executes; it never decides a number. A threshold appearing here instead of v4.0 is a defect, and a number v4.0 does not carry is an Appendix A amendment item, never a local value.

#23 Ranking-push operations	Owns push sizing, the worksheet, funding, in-flight operation, achievement, hold confirmation and the post-mortem. Open funded pushes at stage entry continue under SOP-23 to their own conclusions; the descent schedule sequences around them, never through them.
#24 Re-ranking and recovery	Owns loss diagnosis, baseline restoration, recovery-adjusted sizing and the OOS overlays. A slip inside a controlled taper is not a recovery case. Failure mode F2 exists precisely because the two get confused, and the Ledger's component-25 row records that as the boundary's reason.
#28 Placement and modifier	Owns the authoritative ladder mechanics at its P5: the rung definition, the exit-target computation, the reversion rule, and the floor register schema. This SOP schedules rungs and consumes floors; it derives no step and computes no exit bid.
#12 Exact operations	Owns taper execution at P4, promotion and demotion between the two structure classes at P5 and P6, duplicate consolidation at P7, and the bid and budget change mechanics at P11. This SOP issues the dated instruction; SOP-12 deploys it.
#26 Mature and defend	Owns floor operations, the quarterly re-ladder, the defensive architecture, the code-red protocol and the monthly profit-harvest read. This SOP hands over the closed program record with the floor register and stops.
#29 Negation architecture	Owns negation doctrine, wall levels, the reactive execution standards and the monthly wall audit. The cleanup sweep at P10 is staged into #29's standards; this SOP sets no wall level and writes no evidence line.
#30 Harvest and graduation	Owns the weekly harvest pass. The harvest never stops during a descent: discovery caps tighten to the stage allocation and yield expectations hold. Structural maturity is not completeness of demand, and F4 is the mode that catches the confusion.
#38 Bulk mechanics	Owns pull, staging format, upload and the T+1 verification pass. This SOP produces staged rows in #38 format and consumes the verification report as each rung's proof of landing.
#32 Deal and event window	Owns the event plan, the armed-window declaration and reconciliation. Descent steps do not deploy inside an armed window, because the window distorts the rank read a taper depends on.
#33 Sibling arbitration	Owns the family ownership declaration. A descent on a family-shared term reads the declaration before it schedules, so a taper on one sibling is not read as a market signal by the other.
#31 Budget and envelope	Decides funding and every envelope-level budget verdict, including where freed descent dollars go. Per Ledger conflict C3 this component has no procedural document, so this SOP names the instrument and tab, T0 zone H and Account Rollup A2, and never cites a #31 procedure number.
#37 Automation rules	Owns the rule register and change control. A live automation rule touching a scheduled term is suspended for the descent window, because a rule and a ladder moving the same bid produce a read that scores nothing.

Depth note. This component fills the 10-18 page standard without padding because the descent spans four documents and the failure surface is the schedule, not the step. The v1.0 at 3 pages described the program; it did not make the program executable by a stranger, and it left the entry-artifact set, the milestone arithmetic, the reversion cost, the cluster basis and every non-arrival consequence unwritten. Those five gaps are where the account's post-objective spend persistence actually survives.

2. Trigger Conditions

Each trigger names the artifact in which the signal is detected and the maximum allowed lag from signal to procedure start. The calendar anchors (Monday 06:00 feed landing, 09:00 packet ready, Tuesday deploy after the BM huddle, T+1 verification) are the PPC Command Workbook TEMPLATE v1.0 READ ME's "CADENCE AND OWNERSHIP" line, and the cadence classes behind them are v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)", rule V4. v4.0 carries no numeric signal-to-start lag table for any component; where a lag below is expressed other than as a calendar anchor it is a planning figure, and the missing standard rides to v4.1 as Appendix A item AM-9.

Trigger	Detecting artifact (where the operator sees it)	Maximum lag, signal to procedure start	Procedure
Stage entry, taper handoff	SOP-23 P6.3 issues TAPER with the hold confirmed; the T2 K7 push worksheet closes and the T6 standup row carrying the logged taper trigger reaches its horizon.	The same week. A handed-over taper with no schedule row by the following Monday is failure mode F3.	P1 entry verification, then P2 and P4.
Stage entry, T7 renewal	The monthly stage review writes a Transition tag into T0 zone E "PRODUCT STATE VECTOR", dimension "Stage integrity", with its date.	The same week the tag lands.	P1 in full; the entry read is thinner and P2 carries more.
Objective horizon reached	T2 K7 "Review Date" on a Ranking-objective row passes with the objective unrenewed; the PUSHES LIVE filter view surfaces it against T2 K4 rank.	The pass in which the date lands.	P2 step 4 classifies the term; P4 schedules it.
Weekly cadence	The Monday cascade reaches T2 on a Transition-tagged product after T0 and T1 are read; the schedule's rungs dated this week are due.	Same day. Rungs stage Monday and deploy Tuesday per #38.	P6 rung execution, then P7.
Rank slip on a tapering term	The T2 K4 daily rank series (feed D9) reads a position worse than the rung's logged hold prediction; the T6 row for that rung scores MISS or WRONG-DIRECTION.	Same day. The rank guard is a daily read while any rung is live.	P8 single-term reversion.
Slip cluster	Two or more actively tapering terms slip against their predictions inside one read window, visible as two T6 rung rows scoring MISS in the same window.	Same day the second slip lands.	P9 schedule pause and cause re-read. The count itself is ungoverned; see AM-5.
Monthly milestone date	The milestone date on the program record in T5 PRODUCT LOG reaches the cycle date.	The pass in which the date lands. A milestone date passing unchecked is failure mode F3.	P11 dual-limb checkpoint.
Event window arming	#32 lands an armed-window declaration in T2 K2 "Event Mode" with its dates; the T5 "Trend-Exclusion Flag" sets.	Before the next scheduled rung inside the window.	P12 event overlay: the schedule holds, staging continues.
Recovery Mode entry	T0 zone A carries the CODE-RED verdict, or the product trips v4.0 Section 11, rule R1.	The same pass.	P12 Recovery overlay: tightened confirmation tiers per rule R3 apply to every rung.
Margin table change	T0 zone C "ECONOMICS BLOCK" shows a new margin table version; every ceiling in the file repoints.	The same cycle, per v4.0 Section 8.1, scenario S-A9.	P5 re-sizes the remaining schedule against the new ceilings.
Exit conditions read		The monthly slot.	P13 exit handoff to SOP-26.

The T7 monthly stage renewal reaches the product with the descent landed and the floor register populated in T2 K8.

3. Preconditions and Gates

Each gate below names its v4.0 identifier and function, the T7 MANIFEST field group whose status proves it can be read, and the observable an operator can point at to say the gate passed. Per Ledger conflict C9 the operating corpus uses a different number for three of these functions; the corpus alias is given at first use so no citation is dropped, and the numbering itself rides to v4.1 as AM-11. A gate that cannot be read is not a gate that passed: v4.0 Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass", rule G8, downgrades a verdict whose gate fields are missing to a flag, and the program states what it could not see.

Gate (v4.0 Section 5)	Manifest field group (T7)	What it does to this program	Observable that proves it passed
G1 Validation quarantine	A4, A5, B11	An impossible-metric row produces a flag and a data-refresh task, never a rung. A term whose evidence carries a quarantined row does not enter the schedule and does not exit it either: it holds at its current rung.	T2 K5 "Flag: Zero-Quarantine" reads Clean on every scheduled term.
G2 Sample sufficiency	A4, A1	Sets what the rung read can conclude, not whether the rung deploys. A rung's hold prediction is a rank read, which clears at its own evidence; the CVR and ACOS reads that price the next rung do not clear until their own thresholds do.	T2 K2 "Sample State" and "Clicks 7/30/60d". Below the CVR threshold the ceiling falls back to the T0 reference CVR and the rung record says so.
G3 Budget truncation	A3	A campaign under the in-budget threshold marks every rate TRUNCATED, and the budget lever is evaluated before any bid conclusion. A rung deployed on a truncated read is stepping a bid against the campaign's worst hours, which is failure mode F11.	T3 "Truncation Flag (G3)" and "In-Budget Bucket" clear on every campaign hosting a scheduled term.
G4 Eligibility	A7	An ineligible advertised SKU halts all lever work for the target. The scheduled rung holds with the listing task as its re-entry criterion; the schedule does not skip forward to the next term to keep the month's count.	T2 K2 "Eligibility" reads Live and T3 C2 "Eligibility + Reason" is clear.
G5 Indexing and listing placement	B1	Bars a rank-objective standup on an unindexed term. A descent is not a standup, so a tapering term is unaffected; a term the descent would re-point to a Ranking objective is.	T2 K2 "Indexed?" and "Listing Placement" read Yes and Title or Bullets where a rank objective is retained.
G6 Inventory adequacy	D2, F3 inputs	YELLOW freezes push spend at current level, which a descent already satisfies. RED routes the term out of this program entirely and into SOP-24 P4.1's pre-OOS taper, which is a different instrument with a different purpose: preserving rank memory through a stockout, not discovering a floor.	T2 K2 "Inventory Band" and "Days to Stockout" GREEN or YELLOW on every scheduled term, with the D2 map naming the ranked backup.
G8 Packet completeness	The T7 rollout	A DEGRADED manifest produces flags only. The program still runs the rungs it can read and records the suppressed classes; it never estimates a missing rank series and calls it a hold.	The T7 "PACKET COMPLETENESS SCORE" stamp on the cycle header, with the blocked verdict classes named in-row.
G9 Price stability corpus alias: G9 margin freshness	D1	A price change on the targeted variation inside the trailing window marks the CVR and ACOS reads PROVISIONAL. The rank read that governs a hold prediction is unaffected; the ceiling that prices the exit target is, and the schedule row carries the flag through.	T2 K2 and the T0 economics spine's G9 state. T0 C18:C24 (AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC) reads fresh or the record states PROVISIONAL.
G10 Event exclusion corpus alias: G12 event freeze	A1, T5	Event weeks and the days following are excluded from trend verdicts. No descent rung deploys inside an armed window, and a rank achieved inside one is provisional per v4.0 Section 8.12 "Event-Mode Scenarios (S-E1 to S-E6)", scenario S-E5: the T1 at-target clock starts only after the exclusion window. A term entering the schedule on an event spike is failure mode F6.	T2 K2 "Event Mode" and the T5 PRODUCT LOG "Trend-Exclusion Flag" clear on the affected window.
G11 Sibling conflict	D6	A family-shared term where a sibling carries an active push is a portfolio decision before it is a schedule row. Tapering one sibling while another pushes the same term produces a rank read neither product can interpret.	T2 K8 "Term Owner (Family)" reads a declared owner, not UNDECLARED, and T2 K2 "Family Conflict" is clear.
G12 MKL hygiene corpus alias: G10 objective tagging	B2 inputs via the Hub	Every scheduled term must exist in the current MKL with a syntax tag and a relevancy score. An off-taxonomy term does not get a descent schedule; it routes to the MKL classification decision, because tapering a term the product should not hold at all pays to discover a floor nobody wants.	T2 K1 "Syntax Tags" and "Relevancy" populated, and the MKL refresh inside its cadence.

Three standing preconditions beyond the gates. First, **the entry artifact set is complete**: a stage tag with its date, the term-state inventory of P2, and a live economics chain. An entry declared without the inventory is a stage tag with no program behind it. Second, **open funded pushes are identified and excluded** before the schedule is built; they run to their own conclusions under SOP-23 and enter this program only on that SOP's taper handoff. Third, **no automation rule is live on a scheduled term**: per the Ledger's component-28 row, #37 holds the ladder on automation rules, and a rule stepping a bid the schedule is also stepping makes every V5 score on that term unreadable.

4. Step-Level Procedures

Fourteen procedures. Time figures are planning estimates for scheduling the block; they are not decision numerals and nothing branches on them. Every navigation instruction names the actual tab, zone, block and column an operator opens in the PPC Command Workbook.

P1. Stage-entry verification and the descent mandate (20-30 min, once per entry; planning estimate)

1. **Name the entry route and produce its artifact.** Open T0 zone E "PRODUCT STATE VECTOR" and read the "Stage integrity" row. One of the three Section 1.3 routes must be evidenced: a SOP-23 P6.3 taper handoff with its T2 K7 worksheet closed, a T7 stage-renewal tag with its date, or a T2 K7 "Review Date" set of expired Ranking objectives. A stage tag with no route artifact is not an entry, and the correct action is to return it to the stage owner rather than to build a schedule on it.
2. **Read the entry posture in both currencies before anything else.** T0 zone B "SIX-NUMBER HEADER" gives weekly PPC spend, weekly sales and TACOS against the stage ceiling. Record the posture as a percentage and as a weekly dollar figure, because the schedule is built in dollars and only reported in percent. The percentage moves when sales move; the dollars move only when someone acts, and the second is what this program controls.
3. **Classify the deviation, do not judge it.** v4.0 Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)", Family F2 "Product State & Strategy", sets the Transition reference band at 7-12% and the deviation classes WITHIN at or under 1.2 times band, ELEVATED at 1.2-1.5 times, BREACH at 1.5-2 times, CODE RED above 2 times or above 1.5 times for 2 validated weeks. Compute the ratio and write the class on the program record. The class decides whether a Section 8.1 product-level scenario fires alongside the descent, not whether the descent runs: the descent runs because the stage says so.
4. **Read the gates strip before the schedule, not after.** T0 zone D carries twelve cells, G1 to G12, each PASS, FAIL or N-A with a one-line evidence note. A failed gate names the blocked action class. Record every FAIL and the terms it blocks; those terms enter the inventory with the block recorded and do not receive schedule dates until it clears.
5. **Read T7 MANIFEST and record the suppression.** Note the status of A1, A3, A4, B3, B7, D1, D2 and D9. A MISSING group does not stop the program; it suppresses the verdict classes named in that row's "Verdict Classes Blocked While Missing" cell, and the suppression is written into the program record. D9, the daily rank series, is the one absence that changes the program's shape rather than its scope: without it the rank guard at P7 cannot run daily and every rung's hold prediction scores late.
6. **Confirm the economics chain is live.** T0 zone C, cells C18:C24, carry AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC and WeeklyEnvelope as named cells. Every exit target and every ceiling in the schedule references them. A blended or stale margin table carries the G9 PROVISIONAL tag onto every ceiling-referencing row, and the record says so rather than presenting a provisional ceiling as a firm one.
7. **Write the mandate in one line with two numbers.** The mandate names the entry posture, the band ceiling in dollars, and the distance between them per week. That distance is the program's whole budget of work, and every later argument about whether a month was ambitious enough resolves against it.

P2. The term-state inventory (30-45 min per product; planning estimate)

1. **Enter T2 only after T0 and T1 have been read.** Product posture and family caps govern everything below them, and a T2 pass that skipped T0 and T1 is invalid by construction. Apply the FULL LIST filter view, not OVERSIGHT: a descent touches consolidated tail terms too, and the oversight filter hides them.
2. **Score last cycle first.** Open T6 DECISION LOG, filter to rows whose "Review Date" lands this cycle, and record the V5 score as HIT, PARTIAL, MISS or WRONG-DIRECTION. Unscored actions do not accumulate, per v4.0 Section 14, rule V1. Any rung row scoring MISS or WRONG-DIRECTION routes to P8 in this same pass.
3. **Classify every tracked term into exactly one of five states.** Read T2 K7 "Objective", K4 rank against target, and K8 "# Exact". The five states are: AT TARGET AND HELD (rank at or under target, hold confirmed, still carrying push economics); MID-PUSH (an open worksheet in K7 with a funded release); MAINTENANCE (post-objective already, at or near a recorded floor); DISCOVERY-ONLY (no exact instance, running in the discovery surfaces); and BLOCKED (a P1 step 4 gate failure). Every term lands in one state and the count of each goes on the program record.
4. **Test the AT TARGET AND HELD set against the transition condition itself.** v4.0 Section 8.4, scenario S-R6, requires organic rank at or under target for 2 consecutive weeks. Read the T2 K4 rank series, not the current week's cell. A term at target for one week is not yet a schedule row; a term at target for five weeks with the Ranking tag still live is the Worked Pass 1 case and is the highest-priority row in the schedule.
5. **Exclude the MID-PUSH set from the schedule and say so on the record.** These terms continue under SOP-23 to their own conclusions. Note each one's expected achievement date so P4 can sequence around it: a push landing in month 2 becomes a schedule row in month 3, and the schedule shows the placeholder rather than leaving a silent gap.
6. **Read the family ownership state on every candidate.** T2 K8 "Term Owner (Family)". Where the cell reads UNDECLARED on a family-shared term, gate G11 holds the row and it routes to #33 for the declaration before it takes a date. Where a sibling carries an active push on the same term, the row does not schedule at all this cycle: a taper on one side of an undeclared shared term is read by the other side as market movement.
7. **Record the ladder state each term arrives with.** T4 PLACEMENT & BID carries "Ladder State (step / date)", "Discovered Floor \$", "Placement-CPC Exit Bid \$" and "Next Step Date" per keyword and placement. A term with a recorded floor from a prior cycle does not restart at the top of the ladder; it probes from the floor. A term with no ladder state starts at step 0 and its rung count is unknown until P5 computes it.

P3. Objective-mix and structure-mix baseline (15-20 min per product; planning estimate)

1. **Read the objective mix as it stands.** T0 zone F "TACOS DECOMPOSITION" carries six rows: Ranking, Conversions, Defense, Discovery, Conquest and UNTAGGED, each with spend dollars and share of sales, summing to total TACOS. Record all six. The UNTAGGED row is read first: an untagged row cannot be graded at all per v4.0 Section 6, rule O1, and a descent that normalizes graded spend while untagged spend sits beside it has normalized the wrong half.
2. **Record the mix target gap honestly.** v4.0 Section 6, rule O8, sets stage mix targets for Launch, Ranking push and Mature and names none for Transition. The mix is therefore reported against its own trailing baseline and against the Mature target the product is descending toward, with the missing Transition target logged as AM-15. No target is invented for the intervening stage.
3. **Read the structure mix beside it.** Account Rollup A4 "Structure targets" carries the match-type spend split against the v4.0 Section 10, rule M6 targets. Deviations are findings to explain, not errors to auto-correct, and a descent is the correct occasion to explain them because push-era structure is exactly what M6 deviations record.
4. **Confirm discovery is still funded.** T1 SYNTAX BOARD, group S3, gives spend by objective per family. Discovery caps tighten to the stage allocation during a descent; they do not close. A transition product with no discovery yield reads is failure mode F4, and the check is the presence of the reads, not their level.
5. **Write the mix line onto the program record with the milestone path.** The mix is re-read at every monthly checkpoint beside the posture, so the record shows whether the descent moved the composition or only the total. A total that fell while the composition held is a smaller spend of the same shape, which is a lower cost and not a transition.

P4. The descent schedule build (45-60 min per product, once per entry; planning estimate)

1. **Take only the terms the inventory qualified.** The schedule's population is the AT TARGET AND HELD set that cleared P2 step 4, plus any MAINTENANCE term whose current bid sits above its recorded floor with no verdict authorizing the drift.

2. **Sequence, never batch.** Term-by-term ladders are sequenced within a syntax family and never started simultaneously inside one family, because simultaneous ladders confound the rank reads and a confounded read cannot find a floor. This rule is carried from v1.0 and from the structure conventions of SOP-28 P5; v4.0 carries no sequencing rule, and the gap rides to v4.1 as AM-4.
3. **Order within the family: strongest holds first, contested last.** Read T2 K4 "Rank gap" and "4-wk trend". A term holding well under target with a stable trend carries the least risk and releases budget earliest; a term sitting exactly at target with a slipping trend carries the most and goes last, cushioned by the budget the earlier ladders released. The ordering is a risk ordering, not a size ordering.
4. **Compute each term's rung count before dating it.** Read the term's current bid from T3 and its exit target from T4 "Placement-CPC Exit Bid \$". At the v4.0 Section 10, rule M3 step size of about 10% per week inside the Section 9, rule L2 limits of plus or minus 50% per change and a \$.50 floor, the rung count is the number of compounding steps that reaches the target. Writing the count in advance is what makes a missed schedule visible as a missed count rather than as a vague sense of slowness.
5. **Date every rung.** Each rung carries a deploy date, its ladder reference, its predicted hold position, and its review date. A rung without a date is not scheduled, it is intended, and undated descent is drift wearing a program's name.
6. **Attach the reversion budget to the schedule, not to the incident.** For each term, compute what one rung of reversion costs per week at current click volume: the delta between the rung's CPC and the prior rung's CPC, times weekly clicks. A schedule that has not priced its own reversions will read every reversion as a surprise.
7. **Version the schedule into the product's T5 PRODUCT LOG program record** with the milestone dates from P5 and the mandate from P1 step 7. The schedule is an artifact with a version, and a change to it is an entry, not an edit.

P5. Milestone construction: sizing the schedule into monthly landing points (25-35 min per product; planning estimate)

1. **Build the milestone from the actions, never the reverse.** Sum what the month's scheduled rungs, cleanups and normalizations remove per week, then compute the posture that arithmetic produces at held sales. A milestone chosen as a round percentage and then populated with whatever actions fit is failure mode F5 in advance.
2. **Size a taper rung's removal explicitly.** A rung's weekly removal is the CPC delta times the term's weekly clicks. State the click-hold assumption in the record: clicks are held constant across the rung for sizing, and the checkpoint reads the realized clicks against it. A ladder that loses clicks faster than it loses CPC is discovering its floor early, which the checkpoint should surface rather than absorb.
3. **Size a negation sweep's removal against the campaign's utilization, not its budget.** Removed spend leaves the envelope only where the hosting campaign is not budget-capped. Read T3 C3 "Utilization %" against the objective target. Where utilization sits below target the removed spend leaves; where the campaign is running at its budget the removed spend redistributes inside the same campaign and the sweep buys cleanliness, not savings. Both are legitimate; only one counts toward a milestone.
4. **Size a modifier normalization as the premium portion of placement spend.** Removing a modifier of m percent from a placement carrying S dollars of spend removes S times m divided by (100 plus m) at held placement clicks, because the modifier was funding exactly that fraction of the effective CPC. The modifier mechanics belong to SOP-28 P3.2; the arithmetic of what the month is buying belongs here.
5. **Size a duplicate consolidation as the difference, not the loser's spend.** The survivor absorbs the term's traffic at its own CPC. The month's removal is the pair's combined spend before consolidation less the survivor's spend after it, which is smaller than the loser's spend and is the only figure the checkpoint can verify.
6. **Re-size on any margin-table change.** v4.0 Section 8.1, scenario S-A9, recomputes every ceiling immediately on a margin move and puts every keyword whose CPC is above its new ceiling into the ladder queue this cycle. A descent schedule built against a superseded ceiling is aiming at the wrong floor.
7. **Halve every step under LOW confidence.** v4.0 Section 10, rule M10, halves all lever step sizes and budget steps at LOW confidence, shortens review horizons to one week, and bars pushes. A descent under LOW confidence is a slower descent, never a suspended one.

P6. Weekly rung execution (20-30 min per product per week; planning estimate)

1. **Open the schedule and take only the rungs dated this week.** Rungs are not pulled forward to catch up a missed month. A missed rung is a P11 finding with a cause, and re-dating it is an explicit schedule amendment.
2. **Re-read the gates on each due term before staging.** G3 truncation first: a rung on a truncated campaign is failure mode F11, and the budget lever runs first per v4.0 Section 9, rule L1's precedence. Then G6 inventory, G9 price stability, G10 event window and G11 family ownership.
3. **Take the step size from v4.0, never from the schedule's own memory.** The schedule carries the rung's date and its predicted hold; the step size is read at execution time from Section 10, rule M3, inside the Section 9, rule L2 limits. A staged step exceeding an L2 limit fails staging QA and returns to the pass.
4. **Issue one rung per term per cycle.** The one-lever rule applies: a term receives one rung, not a rung and a budget move, unless the verdict itself names both. Budget follows the bid down and never leads it down, per SOP-12 P4 step 3, so each step reads cleanly against the next.
5. **Retag the objective at the first rung, not at the last.** v4.0 Section 11, rule T1, flips Ranking to post-objective at the transition, and SOP-12 P4 step 4 retags Ranking to Conversions, or Defensive where the term is brand-adjacent, in the T3 objective tag and the campaign-name objective segment in the same staged set. A term walking down a ladder still tagged Ranking is graded against the wrong band for the whole descent, which is failure mode F10.
6. **Write the T6 entry before deployment, not after.** Each rung is a T6 row predicting rank hold at the reduced level, with metric, direction, magnitude, horizon and review date. A change without a logged prediction does not deploy at any tier, per v4.0 Section 12 "Layer 7: Precedence and Conflict Resolution (P1-P10)", rule P10.
7. **Stage per #38 and confirm at T+1.** Staged rows carry entity, prior value, new value, and a reason string carrying the verdict and the scenario identifier. Tuesday deploy after the BM huddle; the T+1 verification report confirms the landed value. A mismatch returns the same day rather than at next cycle.

P7. The rank guard during descent (10 min daily while any rung is live; planning estimate)

1. **Read the daily rank series on every actively tapering term** from T2 K4, sourced from the D9 daily feed, against the rung's logged hold prediction. The guard runs daily while a rung is live and stops when the term reaches its floor and holds. v4.0 Section 14, rule V4, sets daily checks for pushes and events and does not name a controlled descent; the daily cadence here is carried from v1.0 and the gap rides to v4.1 as AM-8.
2. **Read the position, then read the log.** A position worse than the prediction is not yet a slip verdict. v4.0 Section 7 "Layer 2: Diagnosis (D1-D10)", rule D5, consults the product log before PPC is credited or blamed for any swing. Open T5 PRODUCT LOG and check for a listing change, price move, OOS window, deal window or competitor entry inside the rung's window.
3. **Separate a slip from noise on the sample rule.** v4.0 Section 5, rule G2, states that week-over-week swings above its threshold on thin samples are noise by definition. A one-day position move on a low-volume term is not a slip; the guard reads the series, not the point.
4. **Route a confirmed single-term slip to P8 and a second confirmed slip in the same window to P9.** The routing is mechanical and the count is the only judgment involved, which is why the count is recorded rather than felt.
5. **Record a clean guard read as a read.** A rung that held is the evidence that the floor is lower than the current bid, and a program that only records its slips cannot show why it kept going.

P8. Single-term slip response: the one-rung reversion (15 min per case; planning estimate)

1. **Revert exactly one rung.** v4.0 Section 10, rule M3, restores one step on slippage beyond tolerance, and that restored value is the recorded minimum viable CPC. SOP-28 P5 executes the mechanics: revert one rung, hold two windows, then either re-probe or record the floor at the last held rung. Reverting two rungs discards the information the ladder just bought.
2. **Record the tolerance gap on the row.** M3 conditions the reversion on slippage "beyond tolerance" and v4.0 defines no tolerance figure anywhere. Every reversion decision in this document turns on it. The reversion is executed on the logged hold prediction being missed, the basis is stated on the row, and the missing standard rides to v4.1 as AM-6.
3. **Price the reversion in dollars per week** from the P4 step 6 figure: the CPC delta between the two rungs times weekly clicks. This is descent given back, it belongs in the month's checkpoint arithmetic, and it is the reason a reversion is a recorded event rather than a correction.
4. **Do not fund a recovery.** The default answer to a descent slip is the reversion. A recovery push during a controlled descent requires the SOP-24 gate in full: a named cause, RESOLVED, with a verification artifact. A live competitor deal window is a named cause that is not resolved; it closes by calendar. This is failure mode F2 and it is the single most expensive confusion in this document.
5. **Hold two windows, then decide.** The term either re-probes the reverted rung or records the floor at the last held rung, per SOP-28 P5. A reversion that is never re-probed leaves the term parked one rung above a floor nobody discovered, which is failure mode F7.
6. **Write the floor register entry when the floor lands.** Term, floor bid, date and the competitive context note describing what the market looked like when this floor held. The register schema is SOP-28's; populating it for this product is this program's output and SOP-26's entry requirement.

```
Rung's hold prediction missed on the daily read
|
|  --- T5 log shows a listing, price, OOS, deal or competitor event in the window?
|  |  --- Yes ..... attribution first (v4.0 Section 7, rule D5): the swing is
|  |  |  not the rung's. Hold the rung, record the cause, re-read
|  |  |  after the event's exclusion window.
|  |  --- No ..... continue
|  |
|  --- Series or single point? (v4.0 Section 5, rule G2)
|  |  --- Single point, thin sample .. noise by definition. No action. Guard continues.
|  |  --- Series ..... confirmed slip
|  |
|  --- Second confirmed slip in the same window on another tapering term?
|  |  --- No ..... P8: revert ONE rung, hold two windows, price the
|  |  |  give-back, then re-probe or record the floor.
|  |  --- Yes ..... P9: pause the schedule, run the cause classes.
|  |  |  Never: fund a recovery push (F2). The SOP-24 gate
|  |  |  needs a RESOLVED cause with an artifact.
```

P9. Cluster detection and schedule pause (30-45 min per occurrence; planning estimate)

1. **Pause the schedule on the cluster, before diagnosing it.** Slips clustering across two or more tapering terms inside one window are a market or product signal, not a coincidence of floors. Continuing to deploy rungs while the cause is unknown spends the schedule's remaining credibility on a hypothesis. The count of two is carried from v1.0; v4.0 carries no cluster rule and no count, and the gap rides to v4.1 as AM-5.
2. **Run the cause classes in SOP-24 P1's order, and only that order.** Date the loss precisely from the daily rank series, then check OOS event, price change, review shock, listing event, competitor surge, deal hangover, and taper too deep. The last class is the one this program owns, and it is checked last rather than first for the same reason the others exist: a program looking for its own fault finds it.
3. **Name one primary cause with its evidence.** UNKNOWN is a legal finding that routes to a deeper read, never to a default resumption and never to a default push.
4. **Route by cause class, not by instinct.** A cause outside PPC routes to its brand-management owner and the schedule stays paused until it resolves. A competitive cause with a calendar end holds the schedule to that calendar. Only "taper too deep" resumes the schedule directly, and it resumes it re-sized: the ladder inflected at a ladder date, so the exit targets on the affected family were wrong and P5 recomputes them.
5. **Escalate the re-read findings regardless of outcome.** A cluster pause that resolves cleanly is still the schedule's most informative event of the quarter, and the finding is what stops the next schedule from being built the same way.
6. **Resume with a re-probe, not with the next scheduled rung.** The paused terms re-probe the rung they reverted; the unaffected terms resume their dates. Resuming everything at once turns the pause into a batch, which is failure mode F1 arriving late.

P10. Accumulated-waste cleanup (45-90 min per cleanup class; planning estimate)

1. **Treat the descent as the scheduled occasion for push-era debris, not as an emergency sweep.** The push era tolerated waste the stage will not. Each cleanup class lands inside a named month per P5 so its saving is verifiable, rather than running continuously and never being attributable to anything.
2. **Reactive negation sweep beyond the weekly cadence.** Assemble the zero-sale and above-limit terms from the discovery surfaces and hand them to #29 P3 with their evidence lines. This SOP sets no wall level and writes no evidence line; it sets the occasion and counts the result. The relevance check always precedes the negative: a relevant non-converter routes to the fix queue, never to the negative list.
3. **Duplicate consolidation.** Read T2 K8 "Duplicate Flag" and the T3 role declarations. A second live exact instance without one of the four declared roles is SOP-12 P7 territory the day it is found. Survivor selection is human-mandatory and belongs to SOP-12; this program supplies the month and the expected release figure.
4. **Placement-modifier normalization on expired push justifications.** A modifier raised to fund a push whose objective has closed is a premium with no remaining reason. SOP-28 P3.2 steps it toward zero on the ladder cadence. v4.0 carries no horizon defining when a push-era justification is expired; the trigger used is the objective's own closure and the missing standard rides to v4.1 as AM-12.
5. **Quantify every cleanup class and feed the figure to P11.** A cleanup that cannot state its weekly dollars did not participate in the milestone, whatever else it achieved.
6. **Keep the harvest running throughout.** Discovery caps tighten to the stage allocation; discovery does not stop. #30's weekly pass continues and its yield reads continue landing in T3 C5 "Harvest Yield". Structural maturity is never completeness of demand.

P11. Monthly milestone checkpoint: the dual-limb read (30-40 min per milestone; planning estimate)

1. **Read the action limb first.** Open the #38 deployment records for the month and count the rungs that landed against the rungs that were dated, the cleanup classes that executed against those scheduled, and the reversions that occurred. The action limb is a count against a count and it does not admit interpretation.
2. **Read the posture limb second.** T0 zone B gives spend and TACOS. Compare against the milestone figure computed at P5.

3. **Fail the milestone if either limb fails.** A ratio at target with actions unlanded is milestone theater and the month did not happen; a ratio improved by sales growth while spend never normalized is the specific case, and it is detected by re-reading the same spend against the entry sales base. Actions landed with the posture unmoved is a sizing error and it routes to P5 rather than to more actions.
4. **Name the cause of any miss in one of two classes.** Actions unexecuted, or actions executed without effect. The two have different owners and different remedies, and a miss recorded without its class teaches the next month nothing.
5. **Record the miss rather than absorbing it.** v1.0 recorded a milestone "met within tolerance" at 0.1 percentage points over. No tolerance exists in v4.0 for a milestone read. This version states the realized figure, the gap, and the gap's weekly dollar value, and the missing standard rides to v4.1 as AM-3.
6. **Re-read the objective mix beside the posture** per P3 step 5, so the record shows whether the descent changed the composition or only the total.
7. **Reschedule explicitly.** A missed milestone produces a dated amendment to the schedule with the unexecuted actions named, not a quiet expectation that next month absorbs it. The schedule is the commitment, and F3 is the mode that catches a schedule treated as an aspiration.

P12. Event-window and Recovery-Mode overlays (10-15 min per window; planning estimate)

1. **Inside an armed event window the schedule holds and staging continues.** No descent rung deploys, because the window distorts the rank read the rung's hold prediction depends on. The rungs are staged, dated to the window's end, and the record carries the hold.
2. **Do not read event-window rank as a hold.** A rung whose window overlapped an event has no readable prediction; it re-reads after the exclusion period rather than scoring on distorted evidence.
3. **Refuse a term whose rank was achieved inside an event.** v4.0 Section 8.12, scenario S-E5, states that rank achieved during events is provisional and the at-target clock starts only after the exclusion window. A term entering the descent schedule on an event spike is failure mode F6, and the descent will discover a floor that the post-event market does not honour.
4. **In Recovery Mode, tighten the confirmation tier and keep descending.** v4.0 Section 11, rule R3, human-confirms bid moves above 15% and budget moves above \$25 per day and caps step sizes at standard even at HIGH confidence. A descent is the direction Recovery wants; what changes is who confirms each rung, not whether it deploys.
5. **Observe the mode precedence.** v4.0 Section 11, rule T8, orders Recovery Mode above Event Mode above stage defaults. A product in Recovery that enters an event runs the event at defense-and-proven-terms posture only.
6. **Record the overlay on the checkpoint.** A month run under an overlay and reported as an ordinary month makes the next month's V5 scores unreadable.

P13. Transition exit: the floor register handoff (30-40 min, once per exit; planning estimate)

1. **Verify the exit at the T7 monthly slot, never mid-month.** v4.0 Section 11, rule T7, is the stage renewal decision and it runs monthly against the rule T2 criteria. A stage advance declared from a single week's reading is failure mode F12.
2. **State the boundary honestly.** v4.0 Section 11, rule T2, refers the stage gates to the Launch and Scale graduation criteria and the Phase 1-4 model, which live in the Criteria System v1.0 of March 2026 and are not held in this corpus. This SOP therefore verifies what it can see, which is the posture against the band, the floor register's completeness, and the scored schedule, and it hands the stage decision to the PPC Lead with the BM owner rather than asserting criteria it cannot cite. The gap rides to v4.1 as AM-7.
3. **Complete the floor register before the handoff, not after.** Every term whose ladder closed carries a floor in T2 K8 with its date and its competitive context note. Terms whose ladders ended in reversion without a re-probe are listed as open, with their reverted rung and their reason. SOP-26 P1.1 resets drift above floor to floor, which it cannot do for a term whose floor was never recorded, and F8 is the mode that catches an exit declared on a partial register.
4. **Name the floors that sit at or above the ceiling.** Compare each recorded floor against the term's Max Profitable CPC from T2 K6. A floor above the ceiling is a position the product is renting rather than earning, and naming it on the handoff is what lets SOP-26 P5's harvest read and the v4.0 Section 4, rule X12 incrementality ladder decide whether it stays rented.
5. **Close the program record with its scored predictions.** Rungs deployed, V5 scores by class, reversions, cluster pauses, milestones met and missed with their causes, and the realized descent in weekly dollars against the P1 step 7 mandate.
6. **Route the freed dollars, do not spend them.** The released weekly spend posts to the M5 push queue at T0 zone H and Account Rollup A2 as funding supply. Per Ledger conflict C3 this SOP names the instrument and tab and never cites a #31 procedure number.

P14. Program close and carry-forward (20 min, once per exit; planning estimate)

1. **Write the calibration line.** Realized rung count against planned, realized removal against sized removal, realized reversion count against the priced reversion budget. These deltas are the next transition schedule's assumptions, and a schedule never calibrated stays wrong at the same size forever.
2. **Score the schedule as a whole, not only its rungs.** A program whose rungs mostly scored HIT while the posture missed its mandate was sequenced correctly and sized wrongly, and that is a different lesson from the reverse.
3. **Return the per-term floors to T4 and the register.** The floors re-test on the v4.0 Section 14, rule V4 quarterly cycle under SOP-26 P1.2, or earlier after a logged competitive shock.
4. **File the program record in T5 and leave the T6 chain open.** Rungs whose review dates fall after the exit still score, and the scores route to whoever holds the term under SOP-26.
5. **Hand back the amendment items this program generated** so the threshold owner sees which numbers the descent had to run without.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds and scenario logic live in v4.0. This table maps the complete 25-string closed vocabulary, quoted verbatim from the Keyword Decision Record Template v3.0 "Bands & Verdicts Reference" sheet and matching exemplar SOP-12 v2.0 Section 5, to what this SOP does when each verdict lands on a term inside a transition program. Where execution belongs to an adjacent component, the boundary is stated so no verdict falls into a seam. Per Ledger conflict C10 the vocabulary is stated at three different counts across the corpus; the resolution recorded there is followed, and the rows where v4.0 Section 13.1 "The Closed Verdict Vocabulary" carries a different string are cross-mapped in place so neither citation is dropped.

Verdict (KDR Template v3.0 closed vocabulary)	Transition and maintenance execution
TAPER: RANK ACHIEVED, STEP BIDS DOWN	The verdict this SOP exists to schedule. The term enters the descent schedule at P4 with its rung count, its dates and its exit target; P6 deploys one rung per cycle through SOP-12 P4; P7 guards it; P8 reverts it once on a confirmed slip. Cross-map: v4.0 Section 13.1 issues this decision under the string TAPER TO FLOOR (POST-OBJECTIVE), mapped there to S-R6 and S-D1. Both name one decision; the KDR string is the one written into T2 K9 because the workbook's data-validation list is built from it.
HOLD; DEFEND LIGHTLY AT RANK	The state a term occupies between its qualification at P2 step 4 and its scheduled slot at P4. No change this cycle. A term parked here past its scheduled date is failure mode F3, not a hold.

DEFEND AT FLOOR	The terminal state of a completed ladder and the state SOP-26 receives. Where drift above the recorded floor is detected without a verdict authorizing it, the bid restores to the floor value in one staged step at SOP-12. This SOP records the floor; SOP-26 P1.1 operates it thereafter.
INCREMENTALITY LADDER	On a term flagged for cannibalization, the descent doubles as the measurement: v4.0 Section 10, rule M3, states that on flagged terms the taper is the incrementality test, and Section 4, rule X12, sets the read as total organic plus paid capture holding for two windows before the ladder continues. This SOP schedules the rungs; the continue-or-restore call and the measurement design belong to #35 and v4.0.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	Distinct from a taper and often confused with one. This verdict steps a bid to the T2 K6 Max Profitable CPC because the economics failed; a taper steps a bid because the objective was achieved. A term arriving here mid-descent leaves the schedule and takes the single-step edit at SOP-12; the syntax-cap half executes at T1 through the instrument layer.
CONSOLIDATE DUPLICATES	P10 step 3 schedules it into a named month and states the expected release. Execution and survivor selection are SOP-12 P7's, and survivor selection is human-mandatory.
NEGATE (PRE-LOAD / REACTIVE / STEERING)	P10 step 2 sets the occasion for the cleanup sweep and counts its result. Wall level, evidence line and the reverse test belong to #29. This SOP never sets a wall level. Cross-map: v4.0 Section 13.1 issues the reactive branch under the string CUT / NEGATE.
PLACEMENT FIX FIRST	The descent holds on that term while #28 executes the placement program. A ladder stepped through an unresolved placement problem discovers a floor that describes one placement's economics and prices all three.
RECOVERY PUSH	No transition execution, and the boundary that failure mode F2 defends. A recovery on a term inside a controlled descent requires the SOP-24 gate in full: a named cause, verified RESOLVED with an artifact. Absent that, the answer to a slip is the P8 reversion. Note that v4.0 Section 8.4 carries the RECOVERY-class push at scenario S-R13, not S-R12; see Appendix A item AM-13.
RANK PUSH: FUND SIZED PUSH	No transition execution. A new push on a transition-stage product is a stage question before it is a keyword question: v4.0 Section 8.1, scenario S-A3, freezes new Ranking objectives at BREACH, and rule R3 bars them outside RECOVERY class in Recovery Mode. Where a push is legitimately funded, the term leaves the descent schedule and runs under SOP-23.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	No transition execution. The verdict lands on a live push and executes at SOP-12; a term receiving it is not in a descent.
STOP-LOSS: CONCEDE OR DEFER	No transition execution on the descent. Where a stop-loss closes an open push at stage entry, the term enters the P2 inventory as MAINTENANCE rather than as AT TARGET AND HELD, because it never reached target and has no rank to protect on the way down.
FIX CONVERSION FIRST; DEFER PUSH	The term caps at maintenance and leaves the schedule until the fix lands. A conversion deficit is not descended through: v4.0 Section 8.4, scenario S-R9, holds rank at the minimum viable CPC while the deficit routes to its owner, and a term at target on weak conversion is fragile and expensive to keep.
FIX BUDGET (TRUNCATED)	The rung holds. The budget lever runs first per v4.0 Section 9, rule L1, and bid reads resume the next untruncated day. Deploying a rung on a truncated read is failure mode F11.
RE-POINT VARIATION	The term exits the descent schedule the day G6 reads RED on the targeted variation. The pre-OOS taper at SOP-24 P4.1 is a different instrument with a different purpose, and the two are not interchangeable because one is discovering a floor and the other is preserving rank memory through a stockout.
PAUSED-PENDING (REASON + RE-ENTRY)	Used on a scheduled term held by a blocking gate: the schedule row stands, the rung holds, and the record carries the blocking gate as the reason and its clearance as the re-entry criterion. Nothing exits the schedule silently.
QUARANTINE: DATA REFRESH	No rung deploys on the row. The term holds at its current rung; the data-refresh task opens with an owner and a date. A ladder stepped on quarantined evidence has no readable prediction at any rung thereafter.
UNLOCK TRAFFIC	No transition execution. A term below the click-sufficiency read is not a descent candidate: there is no rank being held to walk down from.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	No transition execution. A dated test at a capped bid runs in a discovery surface and reaches the schedule only if it graduates and then achieves a rank.
FUND TEST (UNFUNDED PRIMARY)	No transition execution. Note that funding a test during a descent is legitimate: the descent normalizes achieved positions and does not close the product to new evidence, which is the same principle F4 enforces on discovery.
GRADUATE FROM DISCOVERY	No transition execution. The harvest continues throughout the descent per P10 step 6; a graduate enters exact under #30 and #12 and joins this schedule only when it later achieves a rank. Cross-map: v4.0 Section 13.1 issues this under GRADUATE + STEER.
ENTER SBV ARBITRAGE	No transition execution beyond one interaction: v4.0 Section 8.4, scenario S-R10, pairs the arbitrage with an SP-exact step down one ladder rung. Where that lands on a scheduled term, the rung is the arbitrage's, not the schedule's, and the schedule does not also deploy its own rung that cycle.
OPEN / ROTATE / EXIT CONQUEST	No transition execution. Conquest targets are #16's and carry their own band and their own exit test.
LEAVE TO HARVEST	No transition execution. The term runs on discovery economics under #30's rules and holds no rank this program could descend.
ADD EXACT + RESTRUCTURE	No transition execution, one precondition. A structure gap found during the P2 inventory routes to #12 and #29 before the term is scheduled: a descent on a term whose traffic is mostly carried by auto or broad is walking down a bid that controls a minority of the clicks, and the read it produces will describe the wrong campaigns. v4.0 Section 4, rule X10, and Section 8.4, scenario S-R12, order the structure repair before any performance verdict.

Zero local thresholds. Every numeral in this SOP that functions as a decision threshold is quoted from v4.0 with its section heading, or appears in the Appendix A amendment register. The program schedules a term on a gate, a rank condition and a recorded ladder state, and on nothing else. Where the corpus carries a number this SOP would otherwise need, and v4.0 does not, the number is not written: it is named as a gap.

6. Worked Examples

Four cases: one clean program from entry to third milestone, one that breaks mid-descent and is handled, one carrying two refusals, and one exit. The anonymized live product economics spine was not attached to this session, so all four run on the standard training base of the PPC Economics Worked-Example Set v1.0: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%, reference CVR 5.2%.

Values are illustrative and internally consistent; live values come from the T0 economics spine at C18:C24. Every figure below was recomputed programmatically before publication, with each line item rounded to the cent and every total taken as the sum of the rounded line items.

The chain used throughout: Max Profitable CPC = contribution x CVR = 28.00 x 0.052 = \$1.456, carried at \$1.46. Break-even ACOS = 28.00 / 80.00 = 35.0%. Target = 50% of break-even = 17.5%. Max Acceptable = 75% of break-even = 26.25%, carried at 26.3%. One M3 rung is a reduction of about 10%, so four rungs compound to 0.9 to the fourth power = 0.6561 and three rungs to 0.729.

WE-1 · The clean program: entry, three milestones, \$909.61 per week removed

Setup. The product enters Transition on the T7 monthly stage renewal: the Ranking-push tag expired without rank evidence and auto-transitioned. T0 zone B reads weekly PPC spend \$4,489.00 against weekly sales of \$33,500.00 from all sources.

P1 entry read. TACOS = 4,489.00 / 33,500.00 = 13.4%. The v4.0 Section 3, Family F2 Transition band is 7-12%, so the band ceiling of 12.0% equals 0.12 x 33,500.00 = \$4,020.00 per week and the distance to it is 4,489.00 - 4,020.00 = \$469.00 per week. The deviation class is 13.4 / 12.0 = 1.12 times band, which is WITHIN at or under 1.2 times: no Section 8.1 product-level review fires. This is the entry read most often misjudged. The product is above its band and below its deviation trigger simultaneously, and the descent runs because the stage says so, not because an alarm did. Gates strip: twelve PASS. Economics chain fresh, so no PROVISIONAL tag rides.

P2 inventory. 21 tracked terms: 7 AT TARGET AND HELD, 2 MID-PUSH with open worksheets, 9 MAINTENANCE, 3 DISCOVERY-ONLY. The 2 mid-push terms are excluded and their expected achievement dates are noted for month 3 placeholders. Of the 7 held terms, all clear the S-R6 condition of rank at or under target for 2 consecutive weeks on the T2 K4 series; one has held for 6 weeks with the Ranking tag still live, which is the Worked Pass 1 pattern and takes the first slot.

P4 and P5 schedule. Sequenced within the Bamboo family, strongest holds first. Month 1: 4 tapers plus the reactive negation sweep. Month 2: 3 tapers plus 2 duplicate consolidations. Month 3: modifier normalization on 5 expired push justifications.

Month 1 term	Bid CPC	Clicks/wk	Spend/wk	CPC after 4 rungs	Spend/wk after	Removed/wk
A (6-week hold)	\$2.42	96	\$232.32	\$1.5878	\$152.43	\$79.89
B	\$2.14	78	\$166.92	\$1.4041	\$109.52	\$57.40
C	\$1.96	64	\$125.44	\$1.2860	\$82.30	\$43.14
D	\$1.78	55	\$97.90	\$1.1679	\$64.23	\$33.67
Taper set			\$622.58		\$408.48	\$214.10

The negation sweep removes 23 zero-sale terms carrying \$148.20 per week across the two discovery campaigns. T3 C3 reads both at 61% utilization against the Discovery target, so neither is budget-capped and the removed spend leaves the envelope rather than redistributing inside it; the record states that test, because the same sweep on a capped campaign would have removed nothing from the total. Month 1 removes 214.10 + 148.20 = \$362.30 per week. Exit spend 4,489.00 - 362.30 = \$4,126.70, TACOS 4,126.70 / 33,500.00 = 12.3%. Both limbs pass: 16 rungs dated and 16 landed, sweep landed, posture 13.4% to 12.3%.

Month 2 term	Bid CPC	Clicks/wk	Spend/wk	CPC after 4 rungs	Spend/wk after	Removed/wk
E	\$2.28	88	\$200.64	\$1.4959	\$131.64	\$69.00
F	\$2.05	71	\$145.55	\$1.3450	\$95.50	\$50.05
G	\$1.92	59	\$113.28	\$1.2597	\$74.32	\$38.96
Taper set			\$459.47		\$301.46	\$158.01

The 2 duplicate consolidations: the losing instances carried \$214.90 per week combined and the surviving pair reads \$131.60 after consolidation, releasing 214.90 - 131.60 = \$83.30. Note that this is not the losers' spend: the survivors absorb the traffic at their own CPC, and the difference is the only figure the checkpoint can verify. Month 2 removes 158.01 + 83.30 = \$241.31. Exit spend 4,126.70 - 241.31 = \$3,885.39, TACOS 11.6%, inside the 7-12% band.

Month 3. The 5 expired push justifications carry \$986.00 per week of top-of-search spend at an average modifier of +45%. The premium portion is 986.00 x 0.45 / 1.45 = \$306.00 per week, and normalizing the modifiers to zero removes it at held placement clicks. Exit spend 3,885.39 - 306.00 = \$3,579.39, TACOS 10.7%.

Program result. 4,489.00 - 3,579.39 = \$909.61 per week removed, being 20.3% of entry spend and \$47,299.72 annualized at held sales. The posture landed inside the stage band. Whether the product advances to Mature is the T7 question at P13 and it is not this arithmetic's to answer.

WE-2 · The edge case: a slip cluster in month 2, and why the answer is not a push

Setup. The WE-1 program, one fact changed. In week 2 of month 2 the rank guard reads two positions worse than predicted on the same day: term E at rank 7 against a logged hold of 4-5, term F at rank 9 against a logged hold of 6-7. Both are on rung 4.

P7 guard, in order. The log is read before the ladder is blamed, per v4.0 Section 7, rule D5. T5 PRODUCT LOG carries a competitor entry dated inside the window: a deal price live on two competing ASINs on both terms. The sample rule at Section 5, rule G2, is applied and both moves persist across the series rather than one point, so both are confirmed slips rather than noise.

P9 cluster, because there are two. Two tapering terms slipping in one window is the cluster condition and the schedule pauses before the cause is settled. The count of two is carried from v1.0; v4.0 carries no cluster rule, the basis is stated on the row, and the gap rides as AM-5. The SOP-24 P1 cause classes run in order and return competitor surge as the primary cause with the log entry as its evidence, not taper too deep. That distinction matters: if the ladder had been the cause, the exit targets on the whole family would have been wrong and P5 would recompute them. It was not, so the targets stand.

P8 reversion, priced. Each term reverts exactly one rung per v4.0 Section 10, rule M3, and SOP-28 P5.

Term	Rung 4 CPC	Rung 3 CPC	Delta	Clicks/wk	Give-back/wk
E	\$1.4959	\$1.6621	\$0.1662	88	\$14.63
F	\$1.3450	\$1.4945	\$0.1494	71	\$10.61
Total					\$25.24

The refusal that defines this SOP. A recovery push was proposed on both terms: rank lost, revenue exposed, a funded worksheet available. It is refused. SOP-24's gate requires a named cause verified RESOLVED with an artifact, and a live competitor deal window is a named cause that is not

resolved; it closes by calendar, and the artifact does not exist yet. Funding a push into a competitor's deal price buys the worst comparison of the month, which is the same reasoning v4.0 Section 8.7, scenario S-C5, applies to conquering. This is failure mode F2 and the answer is the reversion.

Milestone arithmetic, restated. Month 2 planned to remove \$241.31 per week; the reversion returns \$25.24, so realized removal is \$216.07 and exit spend is $4,126.70 - 216.07 = \$3,910.63$, TACOS $3,910.63 / 33,500.00 = 11.7\%$ against an 11.6% milestone. The action limb passes: every dated rung landed and both consolidations executed. The posture limb misses by 0.1 percentage points, worth \$25.24 per week and \$1,312.48 annualized. v1.0 would have recorded this as met within tolerance. There is no tolerance in v4.0 for a milestone read, so this version records a miss, its cause class as actions executed without full effect, its dollar value, and the calendar on which the schedule resumes with a re-probe rather than with the next scheduled rung.

WE-3 · Two refusals: the milestone that met itself, and the rank that was not real

First refusal: the ratio moved and nothing else did. Same product, an alternative month 1. Seasonal demand lifts weekly sales from \$33,500.00 to \$36,800.00. Spend is unchanged at \$4,489.00 because none of the 16 dated rungs staged and the sweep never deployed: the #38 deployment record for the month is empty.

TACOS reads $4,489.00 / 36,800.00 = 12.2\%$, which beats the 12.3% milestone. On a single-limb read the month passes and the program reports progress. The dual read at P11 fails it on the action limb before the ratio is discussed at all: zero of 16 rungs landed and zero of one cleanup class executed. The diagnostic that makes it unarguable is one line: the same spend against the entry sales base of \$33,500.00 still reads 13.4%, exactly the entry posture. Nothing normalized. The \$362.30 per week that should have left the envelope did not, which is \$18,839.60 annualized, and when demand reverts the posture reverts with it while the program has spent a month of its schedule.

Why the dual read exists in this shape. A ratio has a numerator this program controls and a denominator it does not. Reading only the ratio hands the milestone to whichever of the two moved, and in a product with any seasonality the denominator moves more. This is failure mode F5, and the response is that the month is recorded as unlanded with its cause class as actions unexecuted, the schedule is amended with dates, and the unexecuted actions are named individually rather than carried forward as a total.

Second refusal: a term that reached target inside an event window. A term arrives at the P2 inventory reading rank 4 against a target of 5, held two consecutive weeks, which clears the S-R6 condition on its face. T2 K2 "Event Mode" shows an armed deal window covering both weeks and the T5 "Trend-Exclusion Flag" is set.

The term does not enter the schedule. v4.0 Section 8.12, scenario S-E5, states that rank gains inside an event are provisional and the T1 two-week at-target clock starts only after the exclusion window; the taper begins on confirmation, not on the event spike. Scheduling it anyway would walk a bid down against a position the post-event market never priced, and the ladder would discover a floor that dissolves the week the deal ends. The term is recorded as HOLD; DEFEND LIGHTLY AT RANK with the exclusion window as its re-entry date, and it re-qualifies on a post-window read. This is failure mode F6, and the cost of getting it wrong is not the rung: it is the floor register entry, which SOP-26 will operate against for a quarter.

WE-4 · The exit: a register with two open rows and two rented floors

Setup. Month 3 closes at 10.7% inside the 7-12% band. The T7 monthly stage renewal reaches the product. P13 runs.

The floor register. 13 terms carry recorded floors in T2 K8 with dates and competitive context notes. Mean recorded floor \$1.27, highest \$1.59, lowest \$0.94. Two terms, E and F from WE-2, ended in reversion with the re-probe still pending on the competitor windows; they are listed as open with their reverted rung and their reason, not omitted and not given a provisional floor.

The read SOP-26 needs and this SOP owes it. Each floor is compared against the term's Max Profitable CPC of \$1.46 from T2 K6. Eleven of the 13 sit under the ceiling. Two do not: \$1.59 and \$1.50. Those two positions are being rented rather than earned, and naming them on the handoff is what lets SOP-26's harvest read and the v4.0 Section 4, rule X12 incrementality ladder decide whether the rent is worth paying. An unannotated register hands over the same 13 numbers with that question invisible.

The scored program. 21 rungs deployed across three months and scored at their review dates: 14 HIT, 5 PARTIAL, 2 WRONG-DIRECTION, a hit rate of 66.7%. Both WRONG-DIRECTION scores are the WE-2 pair, which is the correct outcome: the ladder predicted a hold that a competitor's deal price broke, and the prediction was wrong for a reason the log records. One milestone met on both limbs, one missed on the posture limb by \$25.24 per week, one met on both.

The boundary, stated rather than crossed. The exit verifies the posture against the band, the register's completeness, and the scored schedule. It does not assert that the product meets the Mature-stage criteria, because v4.0 Section 11, rule T2, refers those criteria to the Phase 1-4 model of the Criteria System v1.0 of March 2026, which this corpus does not hold. The stage decision goes to the PPC Lead with the BM owner carrying the three artifacts this program can produce, and the missing criteria ride to v4.1 as AM-7. The released \$909.61 per week posts to the M5 queue at T0 zone H and Account Rollup A2 as funding supply, and per Ledger conflict C3 this SOP names the instrument and the tab rather than a #31 procedure number.

7. Failure Modes and Escalation

Twelve modes. Each carries the signal an operator can see in a named artifact, the response, and the account-history incident that motivated it where one exists. F1 to F5 are carried forward from v1.0 because the incidents behind them are institutional knowledge held nowhere else in the pack; F6 to F12 are new at v2.0 and close the seams the v1.0 depth standard left unwritten.

Failure	Detection signal	Response	Motivating incident
F1. Simultaneous family tapers	Two same-family ladders sharing a start window in the schedule, or two T6 rung rows on one family dated the same week.	Re-sequence per P4 steps 2-3 and re-date the later ladder. The rungs already deployed are not reverted; they are marked as an unreadable window and the family's exit targets are re-derived after the reads separate.	The v1.0 mode, carried. Confounded reads discover nothing: two ladders moving in one family produce a rank series neither ladder can claim, and the floor that emerges belongs to no term.
F2. Recovery push on a descent slip	A recovery standup staged on a term inside a controlled taper with no SOP-24 P2 restoration artifact in the record.	Blocked at the gate. The P8 reversion executes instead. The proposal is logged with its cause class so the pattern is visible if it recurs on the same operator or the same family.	The v1.0 mode, carried, and the reason the Cross-Reference Ledger's component-25 row names SOP-24 in its must-not-restate cell with the clause that F2 exists precisely because the two get confused. A slip inside a controlled taper is the ladder finding the floor. A recovery push fired before a cause is resolved buys rank the cause will take back.
F3. Schedule drift	A milestone date passing in T5 with scheduled actions	The missed-milestone record names the cause class, actions unexecuted or actions executed	The v1.0 mode, carried. The schedule is the commitment: a program with dates that nobody

	unexecuted, or a term sitting at HOLD; DEFEND LIGHTLY AT RANK past its scheduled slot.	without effect, and the schedule is amended with new dates and the unexecuted actions listed individually. Two consecutive misses escalate with the list attached.	enforces reads exactly like a program without dates, and the account's post-objective spend persistence is what that looks like after two quarters.
F4. Harvest stopped	Discovery yield reads absent in T3 C5 "Harvest Yield" on a transition-tagged product, or discovery spend at zero in the T0 zone F decomposition.	Discovery continues at the stage allocation. The caps tighten; they do not close. The gap routes to #30 with the cycles missed named.	The v1.0 mode, carried. Maturity of structure never means completeness of demand, and a descent is the moment the temptation is strongest because discovery is the easiest line to cut and the slowest to show what the cut cost.
F5. Milestone theater	A checkpoint passing on the posture limb with the #38 deployment record showing unlanded actions; detected mechanically by re-reading the month's spend against the entry sales base.	The month is recorded as unlanded whatever the ratio says, with the cause class named and the actions listed individually. P11 step 3 is the dual read and it is not a formality.	The v1.0 mode, carried, and reproduced in WE-3. A ratio improved by sales growth while spend never normalized is the specific case: the numerator is what this program controls and the denominator is what it does not, and reading only the quotient hands the milestone to whichever moved.
F6. Descent scheduled on an event-provisional rank	A schedule row whose qualifying at-target weeks fall inside a T5 trend-exclusion window, or a T2 K2 "Event Mode" flag on a term newly entered at P2 step 4.	The row leaves the schedule and holds with the exclusion window as its re-entry date. Rungs already deployed against the provisional rank revert and the term re-qualifies on a post-window read.	No prior incident recorded in the pack for this component. The mode is written from v4.0 Section 8.12, scenario S-E5, rather than from history: the ladder would discover a floor the post-event market never priced, and that floor then governs a quarter of SOP-26 operation. The first occurrence should be logged as the incident.
F7. Reversion never re-probed	A T4 "Ladder State (step / date)" showing a reverted rung with no subsequent probe date, past the two-window hold.	P8 step 5 executes: re-probe, or record the floor at the last held rung and close the ladder. A reversion is a hold with a decision attached, and the decision has a date.	New at v2.0. The reversion is the one step in the ladder with no forcing function: nothing breaks when it is left standing, the term simply sits one rung above a floor nobody discovered, and the schedule reports it as complete.
F8. Exit declared on a partial floor register	A P13 handoff where the count of recorded floors in T2 K8 is below the count of ladders the schedule closed, or where a floor carries no date or no context note.	The handoff holds. Open ladders are listed with their reverted rung and their reason rather than given provisional floors. SOP-26 P1.1 resets drift above floor to floor, which it cannot do for a term whose floor was never recorded.	New at v2.0, written from the seam. SOP-26's Section 3 states that operating mature without recorded floors is operating blind, and its F2 is floor rot; both presuppose a register this SOP is the sole producer of.
F9. Budget leading the bid down	A staged set where a campaign budget cut is dated before or with the rung on the same term, visible in the #38 change set.	Budget follows the bid down per SOP-12 P4 step 3 and never leads it. A budget cut ahead of the bid truncates the campaign, which marks every rate on it TRUNCATED under G3 and voids the rung's own read.	New at v2.0. The mechanism is self-concealing: the descent appears to accelerate because spend falls faster, and the reads that would have shown the truncation are the same reads the truncation invalidated.
F10. Objective never retagged	A term walking down a ladder with T3 "Objective Tag" still reading Ranking, or a mismatch between the T3 tag and the campaign-name objective segment.	Retag at SOP-12 P4 step 4 in the next staged set, both places together. The mismatch itself is SOP-12's F7 condition; what this SOP owns is that the retag was scheduled at the first rung and not the last.	New at v2.0. A term tagged Ranking is graded against the 50-80% ACOS envelope while it is actually being run toward a Conversions posture, so every band read during the descent is taken against the wrong band and the descent looks healthier than it is for its whole length.
F11. Rung deployed on a truncated read	A T3 "Truncation Flag (G3)" set on the hosting campaign in the same window as a deployed rung.	The budget lever runs first per v4.0 Section 9, rule L1; the rung re-dates to the next untruncated day. The already-deployed rung's hold prediction is voided rather than scored.	New at v2.0. A starved campaign's rates describe its worst hours, so the rung is stepping a bid against a distorted read and the floor it converges toward is the floor of a campaign that was never fully live.
F12. Stage advance declared on one reading	A Mature tag written into T0 zone E outside the T7 monthly slot, or an advance recorded without the three P13 artifacts.	The tag reverts pending the monthly renewal. The advance carries the posture read, the completed register and the scored schedule, or it does not happen.	New at v2.0. v4.0 Section 3, Family F2, notes that a mature product running a launch-shaped spend structure is a finding, not a preference; an advance declared early produces exactly that state, and it is then read as the next stage's problem rather than as this one's unfinished work.

Escalation. F2 escalates on first occurrence to the PPC Lead, because a recovery push funded against an unresolved cause spends real money on a hypothesis and the gate exists specifically to stop it. F3 escalates on the second consecutive occurrence with the unexecuted-action list attached. F5 escalates on first occurrence, because a month reported as landed that did not land corrupts every subsequent milestone the schedule is measured against. A P9 cluster pause escalates its re-read findings regardless of outcome. F12 escalates to the PPC Lead and the BM owner jointly, since the stage tag is the input to the next stage's entire posture. Everything touching the human-mandatory tier, which per v4.0 Section 12, rule P5, is every 500+ SV keyword action, every gate failure, every structural change, every bid move above 25% and every budget move above \$50 per day, is confirmed before deployment without exception; in Recovery Mode rule R3 tightens the bid tier to 15% and the budget tier to \$25 per day.

8. Outputs and Artifacts

Retention is stated below as the governing instrument's, because component #39 owns feed and manifest definitions and its retention policy lines are a pending fold-in. No retention period is invented here; the gap rides to v4.1 as AM-10.

Named record	Storage location	Retention	Downstream consumer and what they do with it
The program record (schedule, mandate, inventory, milestones)	T5 PRODUCT LOG, versioned, one entry per schedule version.	The product log's, as the stage's audit trail. Period ungoverned (AM-10).	P11 checks each milestone against it; P13 closes it; the operator ramp reads it as the week-1 artifact; the next transition on any product reads its calibration line as the sizing baseline.
The dated rung schedule	Inside the program record, with each rung's date, ladder reference, predicted hold and review date.	With the program record.	P6 executes from it weekly; SOP-12 P4 receives one dated instruction per rung; a missed date is F3's detection surface.
Staged rung sets	The #38 staged change set with prior value, new value, and the	Held with the #38 staged set.	SOP-12 P4 deploys; #38 returns the T+1 verification. The deployment record is the action limb of every P11 checkpoint, which is why the checkpoint reads it rather than asking whether the work was done.

	reason string carrying the verdict and scenario identifier.		
Monthly milestone verifications	The program record, one row per milestone: actions dated, actions landed, posture planned, posture realized, cause class on any miss.	With the program record.	The PPC Lead reads it at the monthly review; a second consecutive miss escalates from it; P14 computes the calibration line from the set.
Reversion records	The program record and T4 "Ladder State (step / date)".	With the T4 row for the term's life.	P8 step 5 re-probes or closes from it; F7 detects an unresolved one; the priced give-back enters the month's checkpoint arithmetic.
Cluster pause findings	The program record, with the SOP-24 P1 cause class, its evidence and its resolution calendar.	With the program record.	Escalated regardless of outcome; where the class is taper too deep, P5 re-sizes the affected family's exit targets from it.
The floor register	T2 K8 per term with floor bid, date and competitive context note; T4 "Discovered Floor \$".	Rolling, with the term's row.	The exit gate. SOP-26 P1.1 resets drift above floor to floor and P1.2 re-ladders one rung down per cohort quarterly; #28 P5 owns the register schema this populates.
Cleanup quantifications	The program record, one line per cleanup class with its weekly dollar figure and the utilization test behind it.	With the program record.	P11 counts them on the action limb; a class with no dollar figure did not participate in the milestone whatever else it achieved.
T6 decision-log entries	T6 DECISION LOG, one row per deployed rung with prediction, magnitude, horizon and review date.	Append-only; the log is never edited, only scored.	V5 scores each at its review date per v4.0 Section 14, rule V1; MISS and WRONG-DIRECTION scores route to P8 in the pass they land; #6 owns the scoring rubric.
Released-spend posting	T0 zone H "M5 PUSH QUEUE" and Account Rollup A2.	Rolling.	The queue treats verified released dollars as funding supply. This SOP posts the figure and names the instrument; the allocation decision is the envelope owner's, per Ledger conflict C3.
Suppression statement	The program record header, naming the T7 manifest groups that were STALE or MISSING and the classes they blocked.	With the program record.	The analysis states what it could not see before it recommends anything. A confident descent on a silently incomplete packet spends real money against a rank series nobody read.

The stage's gate artifact at exit is the closed program record with the floor register. A transition reporting a landed posture without a completed register has produced a cheaper month, not a transition.

9. Skill-Conversion Block

Block	Specification
Input bindings	T0 zone A verdict banner and stage tag; zone B six-number header (weekly spend, weekly sales, TACOS vs stage ceiling); zone C economics spine C18:C24 named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC, WeeklyEnvelope; zone D the twelve-cell gates strip; zone E "Stage integrity" and the state-vector rows; zone F the six-row TACOS decomposition including UNTAGGED; zone H the M5 queue. T1 SYNTAX BOARD group S3 spend governance, specifically "Spend Share Actual %", "Share vs Target (pp)", "Cap Status" and "Spend by Objective (\$)". T2 K1 identity (Syntax Tags, Relevancy, SV, Oversight Flag); K2 gates (Validation State, Sample State, In-Budget %, Eligibility, Indexed?, Inventory Band, Days to Stockout, Family Conflict, Event Mode); K4 rank (Organic rank daily series, Target rank, Rank gap, WoW delta, 4-wk trend, Velocity, OOS overlay flag); K6 economics (Max profitable CPC, Expected CPC, CPA ceiling); K7 plan (Objective, Loss ceiling, Review Date, Weeks-at-verdict); K8 registry (# Exact, Duplicate Flag, Term Owner (Family), Current bid vs suggested, TOS modifier %). T3 C2 deployment (Daily Budget, Bidding Strategy, Placement Mult TOS %), C3 pacing (Utilization %, In-Budget Bucket, Truncation Flag (G3)), C5 structure (Harvest Yield, Graduations (Qtr)), C1 "Objective Tag" and "Stage-Stack Role". T4 all six bid-mechanics columns, above all "Placement-CPC Exit Bid \$", "Ladder State (step / date)", "Discovered Floor \$" and "Next Step Date". T5 for the attribution check and the program record. T6 for scoring. T7 for suppression, above all group D9. An absent feed suppresses its dependent class per the manifest; the skill never improvises around a missing rank series and never scores a hold it could not read.
Deterministic logic	The five-state inventory classification; the S-R6 two-week at-target test read from the series rather than the cell; the mid-push exclusion; the G11 ownership hold; the family sequencing and the risk ordering within it; the rung-count computation from current bid to exit target at the M3 step size inside L2 limits; the reversion pricing; the milestone construction from actions with its four sizing forms (rung, sweep against utilization, modifier premium, consolidation difference); the dual-limb checkpoint with its entry-sales-base re-read; the slip-versus-noise test; the cluster count; the floor-versus-ceiling comparison at handoff. Every threshold is v4.0's and every rung mechanic is #28's. The skill schedules, sizes, stages and reconciles; it never sets a step size, never computes an exit bid, and never issues a stage tag.
Human checkpoints	Drawn from the PPC Authority Matrix v2.0 rows that name this component's decisions. AI-executable with audit: the term-state inventory, schedule drafting and sequencing, rung-count computation, milestone sizing, the weekly rung staging, the daily rank-guard read and its slip-versus-noise test, reversion pricing, cleanup quantification, checkpoint arithmetic, the floor-versus-ceiling comparison, and the suppression statement. Human-mandatory: the schedule sign-off itself (Matrix row "Weekly verdicts (established US)", D held by the PPC Lead); every cluster-pause re-read and its cause-class call; every stage advance and every stage-entry declaration (Matrix row "Structure creation (all SOPs)", D by the PPC Lead or per SOP scope); every deployment confirm (Matrix row "Bulk deployments (#38)", R by the PPC Lead); any descent on a launch-stage product (Matrix row "Launch-stage verdicts + batches", D and E held by the launch operator, which places those products outside this operator's scope until it transfers); any schedule row on a family-shared term pending a declaration (Matrix row "Sibling declarations (S-series)", D by the PPC Lead); and any suspension of an automation rule for a descent window (Matrix row "Automation rules (#37)", D by the PPC Lead). The v4.0 human-review tier applies on top and is tightened by rule R3 in Recovery Mode.
Cadence	Entry verification and schedule build once per stage entry, in the week the entry artifact lands. Weekly rung execution with the Monday cascade, after T0 and T1 and after last cycle's due T6 rows are scored, staged Monday and deployed Tuesday per #38. Daily rank guard while any rung is live. Monthly milestone checkpoint on its dated slot. Monthly objective-mix and structure-mix re-read beside it. Exit verification at the T7 monthly stage-renewal slot only. Recovery Mode tightens the confirmation tiers but does not change this cadence.
Alarm specification	F1 to F12 each carry a detection signal from Section 7 with an evidence-line format of row, signal value, and gate or rule reference. Routing tiers: F2, F5 and F12 same day to the PPC Lead, each escalating on first occurrence; F6, F9 and F11 same day to the operator queue because each voids a read that is still being scored; F1, F7 and F10 into the weekly pass worklist; F3, F4 and F8 into the monthly checkpoint and the exit with the reconciliation that detects them.
Skill outputs	The drafted program record with mandate, inventory and suppression statement; the drafted descent schedule with per-term rung counts, dates, predicted holds and priced reversions; drafted staged rung sets in #38 format; daily rank-guard reads with slip or noise classified; drafted reversion records with their give-back priced; cleanup quantifications with their utilization tests; drafted milestone checkpoints showing both limbs

10. Version Governance

Owner: PPC Lead. Review monthly, and immediately on any v4.0 increment, at which point the Section 5 mapping and every S-R6, S-R9, S-R10, S-R12, S-A1 to S-A9, S-E5, S-D1, T1, T2, T7, T8, R1, R3, O1, O7, O8, M1, M3, M6, M10, L1, L2, X12, X10, G2, G3, G6, G9, G10, G11, G12, D5, P5, P10, V1 and V4 reference re-validates before any schedule runs on it. Component-specific triggers beyond the monthly slot, none of them generic:

- Any change to the Transition TACOS band or its deviation classes.** The band is the destination every schedule in this document is sized toward, and it is currently stated at two values across the corpus. A change to v4.0 Section 3, Family F2 increments this SOP before any program is sized against it, and closing the corpus divergence increments it again.
- Any change to the S-R6 condition or the T1 transition rule.** These two state the same bar, at rank at or under target for 2 consecutive weeks. A change to one without the other increments this SOP and opens a conflict report, because the qualification test at P2 step 4 reads both.
- Any change to M3 taper mathematics or the L2 limits.** The rung count computed at P4 step 4, and therefore every schedule date downstream of it, is a function of the step size and the per-change limits. A step-size change re-dates every open schedule in the account.
- Any change to SOP-28's ladder mechanics at its P5.** This SOP schedules rungs it does not define. A change to the rung definition, the exit-target computation, the reversion rule or the floor register schema changes what this SOP can schedule and what it can hand over.
- Any change to SOP-24's cause classes at its P1.** P9 runs that list verbatim at every cluster pause. A new confirmed cause class increments SOP-24 and this SOP together.
- Any change to SOP-26's entry requirements.** The floor register is the exit artifact SOP-26 requires; a change to what SOP-26 needs at entry changes what P13 must complete before it hands over. The two documents are coupled at the register and increment together.
- The D9 daily rank feed changing status.** Ledger dependency reads on the daily series currently cap the rank guard's cadence and the granularity at which a slip can be separated from noise. A change in that feed's availability changes P7's shape, not merely its scope.
- v4.1 resolving any Appendix A item,** which retires that item's workaround in the same increment. AM-1, AM-3, AM-5, AM-6 and AM-7 each currently carry an operating convention this document states in place rather than a threshold it applies.

v2.0 change note. Supersedes v1.0 in full. Procedures expanded 6 to 14: P1 entry verification and P2 the term-state inventory are split from v1.0's single program build; P3 objective-mix baseline, P5 milestone construction with its four sizing forms, P8 the priced single-term reversion, P9 cluster handling as its own procedure with the SOP-24 cause-class routing, P11 the dual-limb checkpoint, P12 the event and Recovery overlays, P13 the exit handoff with the ceiling comparison, and P14 program close with the calibration line are new. Verdict mapping expanded from 5 rows to the full 25-string closed vocabulary with boundary statements, quoted verbatim from the KDR Template v3.0 and cross-mapped to v4.0 Section 13.1 on the three divergent rows this SOP touches. Worked examples 1 to 4, all arithmetic recomputed programmatically with float tolerances. Failure modes 5 to 12, with F1 to F5 carried forward with their v1.0 incidents intact. Elements 11 and 12 added per the Build Kit standard. **One v1.0 threshold corrected:** v1.0's worked example sized its descent against a Transition band of 8-12%, which traces to the KDR Template v3.0 "Bands & Verdicts Reference" sheet, the Command Workbook REFERENCE tab and Doctrine Section 2, not to v4.0. v4.0 Section 3, Family F2 carries the Transition band at 7-12% and is the sole threshold authority, so this version sizes against 7-12% and the divergence is reported as a conflict and logged as AM-1. **One v1.0 practice withdrawn:** v1.0 recorded a milestone as "met within tolerance" at 0.1 percentage points over. No milestone tolerance exists in v4.0; this version records the miss with its dollar value and logs the gap as AM-3. Registered in the Supersession Map; the Ledger row for component #25 is returned with this build and replaces the 2026-07-28 seed row.

Companions: Master PPC Decision Criteria System v4.0 (all thresholds, scenarios and verdict issuance) · Keyword Decision Record Template v3.0 (the closed vocabulary) · PPC Strategic Doctrine v1.0 Sections 1, 2 and 6.2 (the why) · SOP-23 Ranking-Push (the upstream taper handoff), SOP-24 Re-ranking and Recovery (the gate this SOP deliberately does not use on descent slips), SOP-26 Mature and Defend (the exit), SOP-28 Placement and Modifier (the authoritative ladder), SOP-12 Exact (rung execution, consolidation, demotion), SOP-29 Negation (the cleanup sweep), SOP-30 Harvest (discovery continues), SOP-32 (event plans), SOP-33 (ownership declarations), SOP-37 (rule suspension), SOP-38 (staging and T+1) · PPC Authority Matrix v2.0 (Elements 9 and 11) · Master PPC Data Architecture v1.0 and the PPC Command Workbook TEMPLATE v1.0 READ ME (instrument mechanics) · PPC Cross-Reference Ledger (component #25 row).

11. Operator Ramp

Built from the PPC Authority Matrix v2.0 rows that name this component's decisions. Seven rows apply. **"Weekly verdicts (established US)"**: D held by the PPC Lead, E by the incoming operator on assigned scope, E by the PPC Manager, R by the BM owner, which is where the weekly rung pass and the schedule sign-off sit. **"Structure creation (all SOPs)"**: D by the PPC Lead or per SOP scope, E by the incoming operator and the PPC Manager, which covers every stage-entry declaration and every stage advance. **"Bulk deployments (#38)"**: R by the PPC Lead, E by the incoming operator and the PPC Manager, which is every staged rung set. **"Launch-stage verdicts + batches"**: D and E held by the launch operator, R by the PPC Lead and the BM owner, which means a descent on a launch-stage product is not this operator's to schedule until scope transfers. **"Sibling declarations (S-series)"**: D by the PPC Lead, E by the incoming operator, R by the BM owner, which is why a family-shared schedule row waits on a declaration rather than proceeding on judgment. **"Automation rules (#37)"**: D by the PPC Lead, E and R by the incoming operator, which is the authority to suspend a rule for a descent window. **"Push funding (the M5 queue)"**: D by the CEO above envelope and by the PPC Lead within it, E by the incoming operator, which is where the freed dollars go at P13 and why this operator posts the figure rather than allocating it. The checkpoint holder for review and audit throughout is the PPC Lead per the D and R cells. The Matrix standing rules apply: the BM owner's R on verdict rows is the daily huddle, review and not veto, with disagreements routing up rather than sideways; and the vacancy rule reverts an unfilled role's D to the PPC Lead and its E to the PPC Manager until the role fills.

Week	Demonstrated competence
Week 1: observation and artifact familiarity	Shadows one full stage entry, or one full month of an open program where no entry is available, without executing. Locates and reads, unaided: T0 zone B's spend, sales and TACOS with the stage ceiling beside them, and states the posture in both percent and weekly dollars without prompting; T0 zone C's economics spine at C18:C24 and which cell prices an exit target; T0 zone D's twelve-cell gates strip and which action class each failing gate blocks; one term's T2 K4 rank series distinguished from its current-week cell; one T4 row's "Ladder State (step / date)", "Placement-CPC Exit Bid \$" and "Discovered Floor \$"; one closed ladder in the program record with its rung dates and its recorded floor; and the T7 row for the D9 daily rank feed with what its absence suppresses. Reproduces the WE-1 month 1 taper arithmetic cold on fresh illustrative inputs, matching to the cent, as the Economics Worked-Example Set's governance section requires of a new operator before taking a scope. States, from the artifacts alone and not from memory, which four components execute a single rung and which one of them owns the step size.
Week 2: supervised execution, checkpoint holder reviewing	Runs the weekly rung pass on one established US product with the PPC Lead reviewing every staged set before it deploys. The pass bar is five things, all of them checkable: every scheduled term re-gated before staging, with G3 truncation read first; one rung per term per cycle and no rung pulled forward to catch up a missed date; the objective retagged at the first rung, in both the T3 tag and the campaign-name segment; every rung carrying a T6 prediction written before deployment; and every deployed rung reconciled against the #38 T+1 return the following day. Drafts, and the PPC Lead confirms, one complete milestone checkpoint including the action limb read from the deployment record and the posture limb re-read against the entry sales base. Runs one rank-guard week end to end and classifies at least one position move as slip or noise with the series and the log entry cited. Zero rungs deployed on a truncated campaign across the week is the pass bar; a single one fails the week, because that is the F11 mechanism and it voids a read that is still being scored.

Week 4: unsupervised execution, checkpoint holder auditing	Owns the transition programs across the assigned established US products: builds a schedule from an entry, sequences it, sizes its milestones, runs the weekly rungs, guards them daily, and closes each month on both limbs. The D cell is now exercised on weekly verdicts within scope, with the PPC Lead auditing artifacts after the fact rather than gating them before. Executes at least one single-term reversion end to end with its give-back priced, and defends in writing why it was not a recovery case. The audit samples the rung-count computation on three terms against their exit targets, one milestone's dual read against the deployment record, the floor register's completeness against the ladders the schedule closed, and one cluster pause or, absent one, the operator's written cluster plan. Two clean audit months close the ramp. Launch-stage and Canada products stay outside this scope until the Matrix rows transfer, and the operator demonstrates that they know which rows those are and who holds D on each.
--	---

Live-defense questions. These test the reasoning behind the thresholds, not recall of their values. (1) The transition condition is stated twice in v4.0, at S-R6 and at rule T1. Name every route by which a term can satisfy it on the page and still not belong in this week's schedule. (2) A product enters at 13.4% against a 7-12% band and the deviation class reads WITHIN. Explain why a descent runs anyway, and what would be different if the class had read BREACH. (3) Defend sequencing over batching in dollars: what does a confounded rank read cost, and who pays it and when. (4) Two terms slip in one window. Argue the case for a recovery push, then dismantle it, and name the artifact whose absence decides it. (5) A milestone lands on the ratio and none of its actions deployed. Explain why the month failed, using only numbers the workbook already carries. (6) Why does the objective retag happen at the first rung rather than the last, when the term's economics do not change until the ladder ends. (7) A recorded floor sits at \$1.59 against a Max Profitable CPC of \$1.46. Explain what you are handing SOP-26, and what you would be hiding if you handed over the number without the comparison. (8) Budget follows the bid down and never leads it. Explain the failure that rule prevents, and why the failure conceals itself. (9) Which numbers in this SOP would you refuse to write if v4.1 never lands, and what do you do in the meantime.

12. Interface Contracts

Built from the Cross-Reference Ledger's component-25 Interfaces cell and the component-23, -24, -26, -28, -12 and -30 rows that name this component, cross-checked against Master PPC Data Architecture v1.0 and the PPC Command Workbook TEMPLATE v1.0 for artifact names, tabs and timing. Every row states what happens when the handoff does not arrive, because that is the clause the corpus was missing. Note on I-1, I-8 and I-13: the copies of SOP-23, SOP-24 and SOP-26 attached to this session are each at v1.0, so those seams are written from the Ledger rows and from each document's own Element 1 boundary statement rather than from procedure numbers that may move at their next increment.

#	Interface	Sender to receiver	Artifact and timing	Non-arrival consequence
I-1	Taper handoff from a completed push	SOP-23 P6.3 to this SOP P1	The closed T2 K7 push worksheet, the standup-logged taper trigger, the post-mortem calibration line and the prior push CPC, in the week the hold confirms.	The highest-consequence non-arrival in the document. The push simply keeps spending at push economics with its objective achieved, which is v4.0's Worked Pass 1 in progress and the account's named pattern. No alarm rings, because a healthy push and a finished push read identically on every metric except rank against target. Two consecutive weeks at or under target with no handoff escalates to the PPC Lead.
I-2	Stage tag and band	The T7 monthly stage renewal to this SOP P1	A Transition tag with its date written into T0 zone E "Stage integrity", at the monthly slot.	The product runs on its previous stage's TACOS band, so every posture read is taken against the wrong ceiling and the descent is never mandated. A product whose Ranking-push tag has expired with no renewal and no transition is the T7 auto-transition case and it routes to the stage owner rather than being assumed here.
I-3	Objective horizon expiries	v4.0 rule O7 via T2 K7 to this SOP P2	The set of Ranking-objective rows whose "Review Date" passed unrenewed, each with its T6 history, weekly.	Objectives persist by inertia, which rule O7 exists to prevent. The terms stay tagged Ranking and graded against the 50-80% envelope while carrying no live rank case, so they read acceptable indefinitely and never reach a schedule.
I-4	Ladder mechanics and exit targets	#28 P5 to this SOP P4 and P5	The rung definition, the step size reference, the reversion rule, and per term the placement-CPC exit bid in T4 "Placement-CPC Exit Bid \$", refreshed as position CPC data accumulates.	Rung counts cannot be computed, so the schedule can carry dates but not a destination and no operator can say how long a ladder is. The schedule is built to the term's Max Profitable CPC as a stated interim reference with the substitution recorded, and the affected rows are flagged for re-derivation the cycle the placement feed lands. Ledger dependency D3 is OPEN and caps the depth of this row.
I-5	Dated rung instruction	This SOP P6 to SOP-12 P4	One rung per term per cycle carrying the term, the prior bid, the new bid, the ladder reference, the predicted hold and the review date, staged Monday.	The rung does not deploy and the term holds at its current bid for another window at push economics. The schedule's date is missed, which is F3's detection surface, and the month's action limb fails at P11 with the specific rung named.
I-6	Staged change set and T+1 return	This SOP to #38, and #38 back to P6 step 7	Staged bulk rows with verdict and scenario identifier in the reason string, Tuesday deploy after the BM huddle, T+1 verification report back.	Without the staged set nothing deploys. Without the T+1 report the rung has no proof of landing, so its hold prediction cannot be scored and the month's action limb has a count it cannot verify. A landed-value mismatch returns the same day rather than at next cycle.
I-7	Daily rank series	The D9 feed to T2 K4 to this SOP P7	Daily organic rank per tracked term with the OOS overlay, refreshed daily.	The rank guard degrades from daily to weekly and a slip is discovered up to a week after it starts, by which time the next rung has usually deployed on top of it. The schedule slows deliberately to one rung per term per two windows while the feed is absent, and the record states the substitution rather than running the normal cadence blind.
I-8	Loss-diagnosis cause classes	SOP-24 P1 to this SOP P9	The ordered cause-class list, consumed at every cluster pause; and any change to that list at SOP-24's increment.	A cluster is diagnosed by whatever occurs to the operator, which reliably means the ladder is blamed first because it is the nearest change. The schedule then re-sizes exit targets that were never wrong, and the real cause runs untouched for another window.

I-9	The recovery gate	SOP-24 P2 to this SOP P8 step 4	The restoration verification artifact naming the cause and proving it RESOLVED, before any recovery spend deploys on a term inside a descent.	Nothing arrives, and nothing should: the absence is the gate working. The consequence of treating the absence as a formality is F2, which spends a funded worksheet on a cause that is still live and buys rank the cause takes back.
I-10	Product event log	Brand Management to T5 PRODUCT LOG to this SOP P7 and P9	Every listing change, price move, OOS window, deal window, competitor entry and eligibility flip, dated, with its trend-exclusion flag.	The attribution check at v4.0 Section 7, rule D5, cannot run, so every slip is attributed to the ladder by default. The guard holds the rung rather than reverting it, records that the log could not be read, and the reversion decision waits. Crediting or blaming PPC without the log is the single cheapest error to make and the hardest to detect afterwards.
I-11	Cleanup execution	This SOP P10 to #29 P3, SOP-12 P7 and #28 P3.2	The sweep candidate set with its month; the duplicate set with its expected release; the expired-justification modifier set. Each inside its scheduled month.	The month's cleanup class does not land and the milestone fails on the action limb with the class named. Push-era debris survives into the next stage, where SOP-26 finds it as structure that no longer matches the posture and has to remove it without the schedule that would have priced the removal.
I-12	Harvest continuity	#30 P12 to T3 C5 to this SOP P3 and P11	Per-source yield reads and "Graduations (Qtr)", every pass throughout the descent.	Discovery goes ungraded for the descent's duration and the temptation to cut it becomes unanswerable, because nothing shows what the cut costs. This is F4, and the check is the presence of the reads rather than their level.
I-13	Floor register handoff	This SOP P13 to SOP-26 P1	The completed register: per term the floor bid, its date, its competitive context note, and its comparison against the term's Max Profitable CPC, plus the open ladders listed with their reverted rung and reason. Delivered with the closed program record at the exit.	SOP-26 cannot operate. Its Section 3 states that operating mature without recorded floors is operating blind, its P1.1 resets drift above floor to floor and has no floor to reset to, and its quarterly re-ladder has no register to probe from and no diff to produce. Every term then re-discovers its floor from scratch a stage later, at the higher bid it was left at. The exit does not complete without this artifact, and F8 is the mode that enforces it.
I-14	Economics chain	#41 refresh to T0 C18:C24 to this SOP P5	AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS and ExpectedCPC with the G9 freshness state, before the Monday 09:00 packet-ready read.	Every ceiling and every exit target carries PROVISIONAL. Rungs continue, because a rung's hold prediction is a rank read and not a ceiling read, but the destination of each ladder is provisional and is re-derived at the next fresh refresh rather than left standing silently. A margin-table change mid-descent re-sizes the remaining schedule per scenario S-A9. Ledger dependency D1 is OPEN.
I-15	Event plan arming	#32 to this SOP P12	The armed-window declaration with dates and named actions, landing in T2 K2 "Event Mode" before the ramp date.	Rungs deploy inside a window whose rank reads are excluded from trend verdicts, so each one logs a prediction that cannot be scored and the ladder walks down on evidence it will later have to discard. A window observed live with no declaration escalates to the deals owner rather than being worked around here.
I-16	Family ownership declaration	#33 to this SOP P2 step 6	The declared term owner in T2 K8 "Term Owner (Family)" with its declaration date, before a family-shared term takes a schedule date.	Gate G11 holds the row. A descent scheduled on an undeclared shared term produces a rank read the sibling interprets as market movement and responds to with its own bid, so the two products bid against each other with the same money in opposite directions and neither ladder finds a floor.
I-17	Automation rule suspension	This SOP P1 to #37, and #37's confirmation back	The scheduled term list with the descent window, before the first rung; the rule-change log entry confirming suspension.	A rule and a ladder step the same bid in the same window. The V5 score on every rung of that term becomes unattributable, and the discovered floor is an artifact of two instruments rather than a market price. Per the Ledger's component-28 row, #37 holds the ladder on automation rules; without the confirmation the term leaves the schedule.
I-18	Released spend to the queue	This SOP P13 to T0 zone H and Account Rollup A2	The verified weekly released dollars, posted at the exit and at each milestone close.	The M5 queue under-reads its available funding and pushes that would have cleared the ranking go unfunded, so a successful descent quietly reduces the account's push capacity instead of increasing it. Per Ledger conflict C3 this SOP names the instrument and tab and never cites a #31 procedure number.
I-19	Decision log and predictions	This SOP to T6 to V5 scoring at #6	One entry per deployed rung with prediction, magnitude, horizon and review date, same day as deploy; scored at the review date.	Unscored actions do not accumulate and an unlogged change does not deploy at any tier. A ladder whose rungs carry no predictions produces a floor with no evidence that any step along the way was right, which makes the register a list of numbers rather than a record of findings.

Appendix A. v4.1 Amendment Items

Fifteen items. No number below appears as an operating value anywhere in this SOP. Each states the gap, where it was found, and the proposal carried to the threshold owner. Four duplicate items already logged by other build sessions and are marked so v4.1 folds each once.

Item	Gap	Provenance and proposal
AM-1	The Transition TACOS band is stated at two values	v4.0 Section 3, Family F2 carries Transition at 7-12%. The KDR Template v3.0 "Bands & Verdicts Reference" sheet, the PPC Command Workbook TEMPLATE v1.0 REFERENCE tab and the PPC Strategic Doctrine v1.0 Section 2 stage overlay all carry it at 8-12%. The same three instruments also diverge on Ranking push (10-18% against 12-15%), Launch (15-25% against "≤25% dated") and Harvest ("≤3%" against "≤4%"), and label the set "(M7)" where v4.0's M7 is marginal mathematics. This SOP uses v4.0's 7-12%; SOP-25 v1.0 used 8-12%. Proposal: reconcile the band table in v4.1 and re-issue the instrument reference sheets from it. Highest-priority item in this register: it is the destination every schedule in this document is sized toward. Reported as a new conflict with this build.
AM-2	The code-red trip is stated at two values	v4.0 Section 3, Family F2 sets CODE RED above 2 times band or above 1.5 times for 2 validated weeks, matching Section 11, rule R1. The KDR reference sheet and the Workbook REFERENCE tab both state "Code-red at >1.5× band" with no duration limb. P1 step 3 classifies the entry deviation and P12 reads the Recovery overlay, so both consume this. Proposal: align the instruments to v4.0. Reported as a new conflict with this build.
AM-3	Milestone tolerance	SOP-25 v1.0 recorded a milestone "met within tolerance" at 0.1 percentage points over. No tolerance for a milestone read exists anywhere in v4.0. P11 step 5 therefore records the realized figure, the gap and the gap's weekly dollar value, and declares no tolerance. Proposal: admit a milestone tolerance, or state that a milestone is met or missed with no band and that the dollar gap is the reported quantity.
AM-4	Family taper sequencing	The rule that ladders are sequenced and never simultaneous within a syntax family, and that strongest holds taper first, is carried from SOP-25 v1.0 and from SOP-28 P5 structure convention. v4.0 carries no sequencing rule for concurrent ladders at any level. P4 steps 2-3 apply it and state the basis. Proposal: admit a concurrency rule for ladders inside one family, or state that concurrency is permitted and that the resulting rank reads are treated as unattributable.
AM-5	Slip-cluster count	SOP-25 v1.0 set the schedule pause at slips clustering across 2 or more tapering terms in one window. v4.0 carries no cluster rule and no count. P9 step 1 applies the count of 2 and states the basis on the row. Proposal: admit a cluster definition, whether by count, by share of the active ladder set, or by family.
AM-6	Taper slip tolerance	v4.0 Section 10, rule M3, restores one step "on slippage beyond tolerance" and defines no tolerance anywhere. Every reversion in this document turns on it. P8 step 2 executes on the logged hold prediction being missed and states that basis. Proposal: admit a slip tolerance in positions, or state that the rung's own logged prediction is the tolerance by construction.
AM-7	Stage entry and exit criteria	v4.0 Section 11, rule T2, refers stage gates to "the Launch/Scale graduation criteria and Phase 1-4 model", which live in the Master PPC Decision Criteria System v1.0 of March 2026 and are not held in this corpus. Rule T7 supplies exactly one entry route, the expired Ranking-push tag, and no exit criteria for Transition to Mature. P13 step 2 therefore verifies what it can see and hands the stage decision up rather than asserting criteria it cannot cite. Proposal: fold the Phase 1-4 stage criteria into v4.1 so the stage boundaries are readable in the threshold authority.
AM-8	Rank-guard cadence during a descent	v4.0 Section 14, rule V4, sets daily checks for pushes and events and names no cadence for a controlled descent. P7 runs daily while a rung is live, carried from SOP-25 v1.0. Proposal: name the descent's cadence class in rule V4, or state that a live ladder inherits the push cadence.
AM-9	Element 2 lag table	The SOP Build Kit requires a maximum signal-to-start lag on every trigger. v4.0 Section 14, rule V4, sets cadence classes and no lag figures exist. Element 2 expresses lags as calendar anchors from the Command Workbook READ ME wherever possible. Proposal: admit a lag table. Duplicate of the SOP-17/18/19 session's item A-7 and the SOP-30 session's AM-12; fold once.
AM-10	Record retention	Element 8 requires a retention period per record. Ledger row 39 carries "Pending fold-in: retention policy lines" and no period exists for the program record, the schedule, the floor register or the cleanup quantifications. Proposal: admit retention lines at component #39. Duplicate of the SOP-29 and SOP-30 sessions' AM-13; fold once.
AM-11	Gate numbering	v4.0 Section 5 assigns G9 to price stability, G10 to event exclusion and G12 to MKL hygiene; the operating corpus uses G9 for margin freshness, G10 for objective tagging and G12 for the event freeze. Element 3 gives the corpus alias at first use so no citation is dropped. Proposal: reconcile the numbering in v4.1. Duplicate of Ledger conflict C9 and the SOP-29 and SOP-30 sessions' AM-11; fold once.
AM-12	Expired push justification horizon	SOP-28 P3.2 normalizes "a modifier whose justification expired" and neither that document nor v4.0 states when a push-era justification is deemed expired. P10 step 4 uses the objective's own closure as the trigger. Proposal: admit a horizon, or state that the justification expires with the objective it was raised for.
AM-13	The RECOVERY scenario identifier	v4.0 Section 8.4 carries the RECOVERY-class push at scenario S-R13 and carries S-R12 as ADD EXACT + RESTRUCTURE. The KDR Template v3.0 "Bands & Verdicts Reference" sheet cites the recovery zone as S-R12, and exemplar SOP-12 v2.0 Section 5 maps RECOVERY PUSH to S-R12. This SOP's F2 boundary turns on the recovery scenario being cited correctly. Proposal: correct the two instruments to S-R13. Analogous to Ledger conflict C12 and reported as a new conflict with this build.
AM-14	Concurrency of a descent and an open push on one product	P2 step 5 excludes MID-PUSH terms and sequences around them, carried from SOP-25 v1.0. v4.0 states no rule on whether a product may hold a funded push and a descent schedule simultaneously, and Section 8.1, scenario S-A3, only freezes new pushes at BREACH. Proposal: state whether a Transition-stage product may open or continue a funded push, and under what deviation class.
AM-15	Transition-stage objective-mix target	v4.0 Section 6, rule O8, sets stage mix targets for Launch, Ranking push and Mature and names none for Transition, while requiring the actual-versus-target mix as a monthly product finding. P3 step 2 therefore reports the mix against its own trailing baseline and the Mature target it descends toward. Proposal: admit a Transition mix target to rule O8, or state that the stage is measured on the direction of travel rather than on a target mix.

Appendix B. Transition Program Operator Card

P1 to P14 compressed; the SOP text governs.

1. No entry without its artifact: a route, a stage tag with a date, and an inventory. 2. Read the posture in percent and in weekly dollars; only the dollars are yours to move. 3. Classify the deviation, do not act on it: the descent runs because the stage says so. 4. Gates strip and manifest before the schedule, never

after; record what you could not see. 5. Five states, one per term; mid-push terms are excluded and their dates noted. 6. Test at-target on the rank series, not the current cell, and never on an event-window read. 7. Sequence within the family, strongest holds first, contested last, cushioned by what the earlier ladders released. 8. Compute every rung count before you date any rung, and price every reversion before you need one. 9. Build the milestone from the actions; a milestone chosen first and populated later is theater before it starts. 10. One rung per term per cycle; budget follows the bid down, never leads it; re-tag the objective at the first rung. 11. Prediction into T6 before deploy, always; stage Monday, deploy Tuesday, reconcile at T+1. 12. Guard daily: read the log before you blame the ladder, and the series before you call a slip. 13. One slip reverts one rung and holds two windows. Two slips pause the schedule and run the cause classes in SOP-24 P1's order. 14. The answer to a descent slip is the reversion. A recovery needs a cause verified RESOLVED with an artifact, and a live competitor window is not resolved. 15. Close every month on both limbs, and re-read the month's spend against the entry sales base before you believe a ratio. 16. Name every miss by cause class and re-date the schedule explicitly. 17. Exit only at the T7 slot, only with the register complete, and name the floors that sit above the ceiling. 18. Post the freed dollars to the queue; do not spend them. A program without its closed record and its register did not transition anything.