

# SOP-26: Mature and Defend Operations, Including the Code-Red Protocol

**Version 2.0** · 2026-08-08 · Owner: PPC Lead · Supersedes SOP-26 v1.0 (2026-07-24) in full · Every threshold, band, trip condition and verdict lives in the Master PPC Decision Criteria System v4.0 and the Keyword Decision Record Template v3.0; this SOP carries zero local thresholds and executes what those two documents decide.

## 1. Verdict and Scope

**VERDICT:** mature operations are floor operations. The product holds its earned position at the lowest CPC that still holds it, the difference between last quarter's price and this quarter's floor is harvested and re-queued, and blended efficiency above the mature band is treated as the emergency the Criteria System says it is rather than as a bad week. Version 2.0 expands v1.0 from 5 procedures to 14, from a 5-row verdict mapping to the full 25-verdict closed vocabulary, from 1 worked example to 4, and from 5 failure modes to 12, adds the operator ramp and the sixteen interface contracts, and makes three corrections of substance: the code-red trip is restated as the two-part test v4.0 actually carries rather than the single-limb 1.5x reading v1.0 used, the protocol's close is moved from day 14 to the two-consecutive-week exit test of v4.0 Section 11 rule R5, and four verdict strings v1.0 minted locally are retired in favour of the closed vocabulary.

This SOP governs the mature stage from stage entry to stage exit: the weekly mature posture read, floor maintenance and drift restoration, the quarterly re-ladder by term cohort, the standing defensive architecture and its membership register, the annual defense-cost read, the breach-band holding procedure, the code-red declaration and same-day freeze, the four-map diagnosis, the cut sequence ordered by loss dollars, the first re-read, the exit-test close with its structural escalation branch, incursion detection and the response ladder inside the defense ceiling, the monthly profit-harvest read that funds the queue, and the post-exit review that turns every recovery into an amendment proposal.

### 1.1 The Two Operating States and the Band That Separates Them

A mature product is in exactly one of two states at any moment, and the state is read from one number: product TACOS against the mature stage band. Criteria System v4.0 Section 3, family F2 "Product State & Strategy", sets the mature reference band at *Mature 4-7%* and classes deviation as *WITHIN ≤1.2x band · ELEVATED 1.2-1.5x · BREACH 1.5-2x · CODE RED >2x or >1.5x for 2 validated weeks*.

| State                 | Reading against the mature band                                           | Governing scenario                                                                                                          | What this SOP runs                                                                                                                                       |
|-----------------------|---------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Floor state</b>    | WITHIN (≤1.2x band) or ELEVATED (1.2-1.5x)                                | v4.0 Section 8.1, S-A1 healthy posture; S-A2 product review this cycle on ELEVATED or a deteriorating trend across 3+ weeks | P1 posture read, P2 floor maintenance, P3 quarterly re-ladder, P4 defensive architecture, P5 defense-cost read, P11 incursion response, P12 harvest read |
| <b>Breach band</b>    | BREACH (1.5-2x) for 1 validated week                                      | v4.0 Section 8.1, S-A3 push freeze                                                                                          | P13 breach hold. The protocol has not tripped; the freeze has. The second consecutive week above 1.5x converts this into a declaration.                  |
| <b>Code-red state</b> | CODE RED (>2x band in any validated week, or >1.5x for 2 validated weeks) | v4.0 Section 8.1 S-A4 (enter Recovery Mode) and Section 11 rules R1-R6                                                      | P6 declaration and freeze, P7 four-map diagnosis, P8 cut sequence, P9 day-7 re-read, P10 close or structural escalation, P14 post-mortem                 |

The distinction between the second and third rows is the single most consequential reading in this document and it is the one v1.0 collapsed. A product at 1.7x band in one validated week is in S-A3 push freeze and is not in Recovery Mode; declaring code red there runs a five-stage emergency protocol on a condition the Criteria System does not class as an emergency, and the cost is the R3 override set (tightened human-confirmation thresholds, discovery budgets floored at sentinel levels, no new Ranking objectives) applied to a product that has not earned it. A product at 1.6x band for a second consecutive validated week is in Recovery Mode whether or not anyone finds it dramatic. Neither reading is a judgement call.

### 1.2 Why the Mature Stage Carries an Emergency Protocol at All

Strategic Doctrine v1.0 Section 14 lists *Blended efficiency far above target on mature products, treated as normal* as one of the account's named failure classes, and the Criteria System states its own motivating case in Section 3 family F2: *The account's current ~20% blended TACOS against mature-stage products is a standing CODE RED and is why Section 11 exists*. Mature products carry the account's profit, so an efficiency failure there is not proportionally worse than the same failure at launch, it is categorically different: at launch the spend is buying a position that does not exist yet, and at maturity the spend is frequently re-buying a position the listing already holds. Section 3 family F16 names that condition E2, the incrementality and cannibalization estimate, and v4.0 Section 8.1 S-A5 fires *Taper-and-verify* on it with the burden of proof flipped onto the spend. That inversion is what makes floor operations the mature posture: every dollar above the discovered floor is presumed rented until an incrementality read proves otherwise.

### 1.3 Boundaries With Adjacent Components

Every decision excluded from this SOP names the component that owns it. The Cross-Reference Ledger row for component 26 is the check.

| Adjacent component                    | The seam, stated so it is owned                                                                                                                                                                                                                                                      |
|---------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Criteria System v4.0 (#8)</b>      | Decides the mature band, the deviation classes, the code-red trip conditions, the Recovery Mode rules R1-R6, every scenario, every lever limit and every verdict. This SOP executes; it never decides. A threshold written into this document instead of v4.0 is a defect.           |
| <b>#25 Transition and maintenance</b> | Owns the post-objective normalization sequence with its dates and ladder steps, and hands this SOP the stage entry. Ledger row 25 states the seam as "exits into SOP-26". This SOP begins at the handoff and never re-runs normalization.                                            |
| <b>#28 Placement and modifier</b>     | Owns the ladder mechanics every taper in the system references, the exit-bid discovery, the ceiling worksheet, the floor register format, and the bleeding-placement lever order. This SOP issues the re-ladder and bleed decisions; #28 supplies the mechanics and the step values. |
| <b>#12 Exact campaign operations</b>  | Executes campaign-level bid, budget, status and structure changes, including the P4 taper execution, the P6 demotion at floor, and the P7 duplicate consolidation this SOP's diagnosis surfaces.                                                                                     |
| <b>#14 Broad campaign operations</b>  | Owns defensive brand broad (P4) at defense floors and routes branded impression-share and incursion terms to this SOP. This SOP declares defense-set membership; #14 runs the brand-broad surface.                                                                                   |
| <b>#16 PAT and conquest</b>           | Owns defense PAT mechanics, own-ASIN and virtual-bundle self-targeting, and the counter-conquest evaluation on escalations this SOP raises at P11. This SOP escalates; #16 decides whether a counter-open is legal.                                                                  |
| <b>#17-19 SB, SBV and SD playbook</b> | Owns every non-SP format build. On a share-of-voice rebuild verdict this SOP names the terms and the defense budget; the format selection and creative belong there. Ledger row 17 records defense-set membership and escalation as this component's property.                       |

| Adjacent component               | The seam, stated so it is owned                                                                                                                                                                                                                                                                                                                                                                                                             |
|----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| #29 Negation architecture        | Owns negation doctrine, wall levels and the wall audit. The negation sweep inside the cut sequence is staged to #29's standards and executed under its rules; this SOP supplies the evidence lines.                                                                                                                                                                                                                                         |
| #30 Harvest and graduation       | Owns the graduation decision. The harvest never stops at maturity: a mature product still routes candidates and still receives graduates, and no code-red declaration suspends that pipeline.                                                                                                                                                                                                                                               |
| #35 Incrementality testing       | Owns test design and the continue-or-restore call. This SOP executes the ladder rung and records the floor; the read that decides whether capture held is #35's.                                                                                                                                                                                                                                                                            |
| #37 Automation rules             | Owns the rule register and change control. The automation check inside the four-map diagnosis reads #37's rule change log and hands back a suspension notice; this SOP never edits a rule.                                                                                                                                                                                                                                                  |
| #38 Bulk deployment              | Owns staging format, upload, and the T+1 verification diff. Every cut in this document is a staged, reasoned deployment in #38 format. Emergencies get faster queues, never lower standards.                                                                                                                                                                                                                                                |
| #32 Deal and event window        | Owns the event plan and the armed window. v4.0 Section 11 rule T8 orders the two modes: <i>Recovery Mode &gt; Event Mode &gt; stage defaults</i> , and a product in Recovery that enters an event runs the event at defense-and-proven-terms posture only.                                                                                                                                                                                  |
| #33 Sibling arbitration          | Owns the family ownership declaration. A mature product sharing a defended term with a sibling checks the declared owner before any restore step.                                                                                                                                                                                                                                                                                           |
| Budget envelope and the M5 queue | The released dollars this SOP verifies post to the M5 funding queue at the instrument layer: T0 zone H on the product Command Workbook and tab A2 of the Account Rollup Workbook. Per Ledger conflict C3 this SOP names the instrument and the tab and never cites a #31 procedure number; the v1.0 citation of "#31" at its P5 is one of the five dangling citations C3 records, and it is retired here and logged as amendment item AM-9. |
| Cross-framework read             | Pricing, competitive position and product decisions beyond PPC. A structural close at P10 hands the four maps out of PPC entirely. This SOP stops at that boundary and does not diagnose price or product.                                                                                                                                                                                                                                  |

## 1.4 What This SOP Does Not Contain

No numeral in this document is local. Where an operating question needs a number that v4.0 does not carry, the question is named in the amendment register at Section 10.3 and the procedure states the gap in place rather than filling it. Twelve such gaps are logged at this version. An operator who finds a threshold in this document that does not carry a v4.0 section heading beside it has found a defect and should raise it rather than execute it.

## 2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. Lag is measured from the moment the artifact carries the signal, not from the moment a person reads it, which is why the daily-cadence rows exist at all.

| Trigger                  | Signal and detecting artifact                                                                                                                                                                                        | Maximum lag      | Procedure and authority                                                                                                                                                                                                |
|--------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Stage entry              | The #25 normalization handoff arrives with its floor register and dated ladder state. Artifact: T2 block K8 discovered-floor values with dates, plus the T6 chain of the closing ladder steps.                       | One weekly cycle | P1 posture standup. A stage entry arriving without recorded floors is refused back to #25, because operating mature without recorded floors is operating blind and every downstream procedure here reads the register. |
| Weekly cadence           | The Monday cascade reaches T0 on a product carrying the Mature stage tag. Artifact: T0 zone B row 3, TACOS against stage band, and zone F, the TACOS decomposition by objective.                                     | Same cycle       | P1 weekly posture read, then P2 floor maintenance. Deployment Tuesday per #38 after the BM huddle.                                                                                                                     |
| Trip watch               | Daily product TACOS read against the mature band. Artifact: Account Rollup Workbook tab A0, the TACOS versus band column and the code-red flag cell.                                                                 | Same day         | P6 on a CODE RED reading per v4.0 R1; P13 on a BREACH reading. The A0 flag cell and the R1 trip are not the same test and the operator declares on R1 (see Section 3, gate note, and amendment item AM-1).             |
| Second-week confirmation | A product that read above 1.5x band last validated week reads above 1.5x again. Artifact: T6 decision log entry from the prior week's P13 hold, matched against this week's T0 zone B.                               | Same day         | P6 declaration. This is the limb that fires most often and it is the limb v1.0 could not see, because a single-week reading carries no memory.                                                                         |
| Quarterly re-ladder      | A term cohort reaches its re-test date. Artifact: T4 columns "Discovered Floor \$" and "Next Step Date"; v4.0 Section 10 rule M3 requires the floor <i>re-tested quarterly or after a logged competitive shock</i> . | One weekly cycle | P3 cohort re-ladder, executed through #28 P5 mechanics.                                                                                                                                                                |
| Competitive shock        | A logged shock on a defended term ahead of the quarterly date. Artifact: T5 product log, event type "Competitor entry", plus the T2 K8 competitor block (top-holder price, rating, reviews, deal and OOS flags).     | One weekly cycle | P3 out-of-cycle probe on the affected cohort only. M3 names the logged competitive shock as a re-test trigger equal to the calendar.                                                                                   |
| Incursion flag           | Branded impression share falls, or competitor brands appear in the brand campaign's served terms. Artifacts: T2 block K3 "Brand Impr Share" and "IS% Band" on branded rows, and the #14 P4 weekly brand-broad read.  | One weekly cycle | P11 response ladder inside the defense ceiling. v4.0 Section 8.6 S-D2 governs the restore step.                                                                                                                        |
| Share-of-voice failure   | Share of Voice composite falls below the PRESENT class on the product's top-10 terms. Artifact: the E5 composite defined in v4.0 Section 3 family F12, read on the defended set.                                     | One weekly cycle | P4 step 5 rebuild routing under v4.0 Section 8.10 S-F8, which names the cheapest surface first and SP CPC escalation last.                                                                                             |

| Trigger                     | Signal and detecting artifact                                                                                                                                                   | Maximum lag      | Procedure and authority                                                                                           |
|-----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|-------------------------------------------------------------------------------------------------------------------|
| <b>Cannibalization flag</b> | E2 fires on a branded or rank-3-or-better defended term. Artifact: T2 block K6 economics with the E2 estimate; v4.0 Section 8.6 S-D4.                                           | One weekly cycle | P2 step 6 hands the term to the incrementality ladder; the design and the continue-or-restore call belong to #35. |
| <b>Monthly harvest read</b> | The harvest-read date on the mature portfolio. Artifact: T6 decision log filtered to floor-state spend, against the prior-quarter baseline held in the same log.                | Two working days | P12 released-budget computation and verification, posting to T0 zone H.                                           |
| <b>Annual defense read</b>  | The defense-cost read date per primary syntax. Artifact: T1 group S7 "Defense budget \$", which v4.0 Section 10 rule M8 declares quarterly and which this read scores annually. | One weekly cycle | P5 defense-cost read. A defended set with no cost read on file is failure mode F4.                                |
| <b>Recovery exit</b>        | Two consecutive validated weeks within 1.2x band with no gate backlog, per v4.0 Section 11 rule R5. Artifact: T0 zone D gates strip plus two dated T0 zone B readings.          | Same cycle       | P10 close, then P14 post-mortem within one week of exit per R6.                                                   |

### 3. Preconditions and Gates

Every procedure below runs only after these preconditions read current in the T7 manifest. A missing or stale field group suppresses its dependent execution class and the pass records what it could not see; it never silently degrades a verdict. The code-red protocol is the deliberate exception and is stated as such at the end of this section.

#### 3.1 Manifest Preconditions

| Manifest field group (T7)            | Requirement before the dependent procedures run                                                                                                                                                                                                                                                                                                                                    |
|--------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>A1 four-period comparability</b>  | L30, P30, B30 and L60 present on the product. Without the four-period block the trip test reads a single window and cannot distinguish a spike from a trend, so P6 declarations are legal only on the current-week reading and the T6 entry records the single-window limitation. P3 re-ladder predictions are suppressed entirely: a rung with no trend context cannot be scored. |
| <b>A2 deployment state</b>           | Current bulk state in-row for every defended campaign. Stale state suppresses P2 drift restoration for those rows, because a restore against a stale bid value deploys the wrong delta.                                                                                                                                                                                            |
| <b>A3 pacing and in-budget</b>       | Present on every defensive campaign. v4.0 Section 9 rule L4 sets the defensive utilization target at 100%, and a defensive campaign under target takes the budget lever before any bid reading is accepted.                                                                                                                                                                        |
| <b>A6 objective and Expected CPC</b> | Every defended row carries exactly one declared objective. An untagged row on a mature product is an S-A8 stage-mismatch finding, not a defense row, and it routes to assignment before it is graded.                                                                                                                                                                              |
| <b>B7 placement splits</b>           | TOS, ROS and PP splits present for the campaigns in scope. The placement bleed map at P7 step 2 cannot be built without it; where it is MISSING the map is written as "not readable" and the cut sequence proceeds on the other three maps with the gap named. Ledger dependency D3 records this feed as OPEN.                                                                     |
| <b>B8 row economics</b>              | The T0 economics spine live at named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS and ExpectedCPC. Every ceiling in the cut sequence dereferences these cells. A blended or stale margin carries the G9 PROVISIONAL tag onto every ceiling-referencing action.                                                                                                       |
| <b>D6 sibling CPC view</b>           | Same-term CPC across siblings from Account Rollup tab A1. Absence suppresses the sibling limb of the structure scan at P7 step 3; the duplicate limb still runs.                                                                                                                                                                                                                   |
| <b>D7 decision log and V5</b>        | T6 live with review dates and scores. The second-week confirmation trigger reads last week's T6 entry; without it the two-week limb of the trip test is unreadable and the product is exposed to exactly the failure this protocol exists to prevent.                                                                                                                              |
| <b>D8 competitor depth</b>           | Top-holder price, rating, deal and stockout state. Absence suppresses the incursion cause classification at P11 step 2; the response ladder still runs on the impression-share signal alone with the coverage gap noted. Ledger dependency D4 records this feed as OPEN.                                                                                                           |

#### 3.2 Gates in Force

Gate numbering diverges between the Criteria System and the operating corpus. This is Ledger conflict C9, and per its resolution every gate below is named by its v4.0 name and number with the corpus alias given at first use. The alias is not a synonym and the two numbering schemes must not be mixed inside one reason string.

| Gate      | v4.0 Section 5 rule                               | The observable that proves it passed, on a mature product                                                                                                                                                                                                                                                                                |
|-----------|---------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>G1</b> | Validation quarantine                             | T2 flag column "Flag: Zero-Quarantine" reads Clean on every row entering a cut sequence. A quarantined row produces a flag and a data-refresh task, never a cut. v4.0 Section 12 rule P8 makes this beat urgency: <i>a quarantine or DEGRADED manifest beats every other consideration including urgency.</i>                            |
| <b>G2</b> | Sample sufficiency                                | T2 "Sample State" reads Full on any row receiving a rate-driven cut. Below sufficiency the row's verdict is EXTEND TEST and it leaves the cut list, however large its spend looks.                                                                                                                                                       |
| <b>G3</b> | Budget truncation                                 | T3 "Truncation Flag (G3)" clear on every campaign in scope, or the budget lever runs first. A truncated defensive campaign reading a high ACOS is describing its worst hours.                                                                                                                                                            |
| <b>G6</b> | Inventory adequacy                                | T2 "Inventory Band" GREEN or YELLOW on the defended variation. RED tapers to defense and re-points the same day, which outranks every mature-stage posture below it per v4.0 Section 12 rule P1.                                                                                                                                         |
| <b>G7</b> | Envelope and pacing                               | T0 zone H reconciles: the sum of campaign budgets against the declared weekly envelope. A release computed at P12 that does not reconcile to the envelope is not a release, it is an accounting error.                                                                                                                                   |
| <b>G8</b> | Packet completeness                               | T7 completeness score present and stated on the cycle header. A DEGRADED packet under 85% of gate-relevant fields produces flags only, and a code-red declaration on a DEGRADED packet carries the completeness score in its declaration record.                                                                                         |
| <b>G9</b> | Price stability<br>Corpus alias: margin freshness | No price change on the defended variation inside the trailing 7 days, or every CVR and ACOS reading in the diagnosis is marked PROVISIONAL and no conversion verdict issues on it. The corpus uses this number for margin-table staleness; both readings suppress the same class of action and the reason string states which one fired. |

| Gate       | v4.0 Section 5 rule                                                      | The observable that proves it passed, on a mature product                                                                                                                                                                                                                                                                    |
|------------|--------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>G10</b> | Event exclusion<br>Corpus alias: G12 event freeze                        | Event weeks and the 14 days following are excluded from trend verdicts. On a mature product this suspends the quarterly re-ladder, because a rung probed against event-inflated competition discovers a floor that will not hold. It never suspends a code-red declaration, and per rule T8 it never outranks Recovery Mode. |
| <b>G11</b> | Sibling conflict                                                         | Family ownership declared on any shared defended term before a restore step raises a bid on it. Two siblings restoring the same term is a portfolio decision at #33, not two product decisions.                                                                                                                              |
| <b>G12</b> | MKL hygiene<br>Corpus alias: the event freeze number, which is G10 above | Every defended term exists in the current MKL with a syntax tag and a relevancy score. A defended term that is not in the MKL is not a defense, it is an unexamined habit, and it routes to MKL maintenance rather than to the ladder.                                                                                       |

**The code-red protocol has one precondition and it is not a gate.** The protocol trips on the number and runs; that is what makes it a protocol rather than a proposal. What the gates do inside a declared code red is govern which cuts are legal, not whether the declaration stands: a G1 quarantine removes a row from the cut list and a G8 DEGRADED packet caps the verdicts to flags, but neither delays the declaration, the freeze, or the diagnosis. The one condition that must be true before declaring is that the reading is validated, which means the week is closed, the packet is present, and the sales figure is the SellerBoard read rather than an in-console estimate. v4.0 Section 11 rule R1 uses the word "validated" in both limbs of the trip and the word is doing work: an unvalidated week is not a week.

## 4. Step-Level Procedures

Fourteen procedures. P1 to P5 and P11 to P12 are floor-state work and run on the normal cadence. P13 is the breach-band hold. P6 to P10 and P14 are the code-red protocol and run only on a declaration. Time estimates assume the packet is present and the operator has run the procedure before; first executions run roughly double.

### P1. The weekly mature posture read (20 minutes per product)

1. Open the product Command Workbook and read T0 before anything else. v4.0 Section 8 sets the matrix precedence as *product (S-A) outranks syntax (S-Y) outranks inventory-keyword (S-I) outranks keyword objective matrices*, so a keyword verdict issued before the product posture is resolved is invalid by construction.
2. Score last cycle first. Read the T6 decision-log rows whose review dates land this cycle and record each V5 score as HIT, PARTIAL, MISS or WRONG-DIRECTION. Unsourced actions do not accumulate. On a mature product the rows that matter most are ladder rungs and release postings, because both make predictions that a single week either confirms or kills.
3. Read T0 zone B row 3: weekly PPC spend against weekly sales, giving product TACOS, beside the stage band ceiling. Compute the ratio to the band ceiling explicitly and write it into the cycle header. The ratio, not the raw percentage, is what every deviation class in v4.0 Section 3 family F2 is expressed in.
4. Classify the ratio into WITHIN, ELEVATED, BREACH or CODE RED against the F2 deviation classes. On CODE RED go to P6 the same day and stop reading. On BREACH go to P13. On ELEVATED continue this procedure and open the S-A2 product review named in step 5. On WITHIN continue.
5. ELEVATED only: run the four reads S-A2 names, in this order, and record each in T6 whether or not it finds anything. First, the E12 concentration read: T2 sorted by spend, the largest single term's share of product ad clicks and ad spend. Second, the E1 inflation decomposition on any term whose CPC has risen materially, which separates market inflation from self-inflicted inflation before aggressiveness is judged. Third, the X10 auto-dominance sweep: the share of any priority term's clicks carried by auto and broad. Fourth, re-verification of the top-10 spend terms against their declared objectives. New pushes require sign-off above the PPC Lead until the product reads WITHIN again.
6. Read T0 zone F, the TACOS decomposition by objective, and check that the mature mix matches the stage. v4.0 Section 6 rule O8 sets the mature target as *Mature: defense + conversions dominant, TACOS 4-7%*. A mature product carrying a discovery-heavy or ranking-heavy mix is an S-A8 stage-mismatch finding requiring a written objective-mix correction plan with a 2-week convergence date. The UNTAGGED row trending toward zero is the adoption metric; any dollar in it is unmanaged by definition.
7. Read T0 zone E, the product state vector, for the flywheel-health row. v4.0 Section 3 family F2 sets the mature organic-share reference at *Mature primary syntax organic share target ≥60-70%* and defines the RENTING flag as *ad share rising while organic share falls for 3+ weeks*. A RENTING reading fires S-A5 taper-and-verify, which routes every rank-3-or-better and branded term to the incrementality ladder at P2 step 6. RENTING overrides a good keyword-level ACOS: the product is paying for sales that organic rank should hold.
8. Read T0 zone D, the gates strip, and record which gates failed and which action class each failure blocks. Carry the failures into every procedure this cycle rather than re-reading them per row.
9. Write the posture line into the cycle header: state, ratio to band, mix status, flywheel status, gates failing. This one line is what the next twelve procedures inherit and it is what the BM huddle reviews.

### P2. Floor maintenance and drift restoration (25 minutes per product)

1. Apply the T2 OVERSIGHT filter view and add the defended set: rows whose objective tag reads Defensive, plus rows whose objective reads Conversions and which carry a recorded discovered floor in block K8. The defended set is a named list, not a filter result reinterpreted each week, and its membership is declared at P4.
2. For each row, read the current bid against the recorded discovered floor and its date. The floor is the ladder's product and the register is its memory; v4.0 Section 3 family F9 states the point directly: *the ladder state is the taper protocol's memory: without it every taper restarts from scratch*.
3. Where the current bid sits above the recorded floor with no verdict authorising the difference, issue DEFEND AT FLOOR and restore to the floor value in one staged step. Drift upward without a verdict is the most common quiet failure on a mature product because nothing alarms on it: the campaign still runs, the position still holds, and the premium is invisible until the harvest read finds it a month later.
4. Where the current bid sits below the recorded floor and the position has slipped, restore to the floor value and record the slip in T6. Where the position has not slipped, leave it and mark the row as a candidate for the next quarterly probe rather than recording a new floor mid-cycle. Floors are recorded by the ladder procedure at P3, never by an accident of drift.
5. Where the position is held and the bid is at floor, issue HOLD; DEFEND LIGHTLY AT RANK and change nothing. The row is read again next cycle against the week-over-week block. Doing nothing is a verdict here and it is logged as one.
6. Check the E2 cannibalization flag on every branded row and every row held at rank 3 or better. Where it fires, v4.0 Section 8.6 S-D4 routes the term to the incrementality test: one ladder step down, total capture watched for two weeks, the ladder continuing while capture holds. Issue INCREMENTALITY LADDER, stage the single step, and hand the measurement design and the continue-or-restore call to #35. This SOP owns the rung; it does not own the read.

7. Check the defensive utilization row on every defensive campaign in T3. v4.0 Section 8.6 S-D3 is unambiguous that budget precedes bids here: a defensive campaign under its utilization target takes the budget lever before any bid conclusion is accepted.
8. Stage every change as bulk rows per #38 with the four-segment reason string carrying the verdict verbatim, the v4.0 scenario ID, this SOP and procedure number, and the floor register reference. A staged row missing its scenario ID does not deploy.
9. Deploy Tuesday after the BM huddle. Read the #38 T+1 verification diff and redeploy any delta the same day.

### P3. The quarterly re-ladder by term cohort (40 minutes per cohort)

Competition moves floors in both directions, and a floor never re-tested is a price agreed with last quarter's market. The re-ladder is one probe rung down from the recorded floor, run per cohort so that a market-wide shift is legible as a pattern rather than as eleven unrelated slips.

1. Build the cohort. Group the defended set by syntax family and by the quarter in which each floor was recorded, reading T4 columns "Discovered Floor \$", "Ladder State (step / date)" and "Next Step Date". A cohort is one syntax family in one recording quarter; cohorts run separately so their results are comparable.
2. Check G10 event exclusion before probing. An armed event window or its 14-day guard suspends the probe for that cohort and the next date moves out by the guard period. A rung probed inside an event window discovers a floor that dissolves when the window closes.
3. Confirm the ceiling worksheet at #28 P2 for any cohort whose probe touches a modifier rather than a base bid. The effective-CPC arithmetic is multiplicative and the worksheet is mandatory before any modifier change; this SOP does not restate the arithmetic and does not skip it.
4. Probe one rung down from the recorded floor, at the M3 step size, within the L2 per-change limits and above the L2 floor value. The step size, the per-change cap and the bid floor are all v4.0 Section 9 rule L2 values read at execution time. Stage the whole cohort in one change set so the read window is common.
5. Log the hold prediction per term: position holds at the reduced level, with a named review date one read window out. A rung without a dated prediction is not a rung, it is drift, and it cannot be scored.
6. Hold one read window. Read at the review date: position held, or position slipped past the v4.0 tolerance.
7. On a hold, record the new lower floor in T4 with its date and the competitive context note describing what the market looked like when this floor held. The context note is what makes a future slip diagnosable.
8. On a slip, revert one rung to the prior recorded floor, hold two read windows, then either re-probe once or close the cohort at the prior floor. The reverted value is the recorded floor until a re-probe holds; a probe that failed does not lower the register.
9. Produce the register diff as the cohort's artifact: floors lowered, floors held, floors raised, each with its date and context note. The diff is what P12 reads to compute the released difference and it is the only evidence that the quarter's re-ladder actually ran.
10. Where a cohort slips as a whole rather than in individual terms, treat it as a competitive-shock reading rather than a bidding result and check the E1 inflation decomposition before re-probing. Market-driven inflation and self-inflicted inflation take opposite responses and the difference is measurable.

### P4. The standing defensive architecture and its membership register (30 minutes, monthly)

1. Declare the defended set in writing: term, surface, objective tag, floor bid, defense budget line, and the reason the term is defended. Ledger row 17 assigns defense-set membership and escalation to this component, so the declaration is this SOP's artifact and the format builds are not.
2. Brand broad: one brand-broad campaign per brand at defense floors, run by #14 P4. This SOP confirms membership and the floor value; #14 runs the surface and returns the weekly branded impression-share reading and any incursion terms.
3. Own-page defense: own-ASIN self-targeting at floor bids, permanently, per #16. Strategic Doctrine v1.0 Section 8 calls own-ASIN contextual *the cheapest page defense*, and the arithmetic behind that phrase is that the alternative is a competitor's ad on our own detail page at a price we do not set.
4. Shelf presence: brand-format placement on the head terms where the format split justifies it. Grade the defended set on the E5 Share of Voice composite, which v4.0 Section 3 family F12 bands as *DOMINANT >40% · PRESENT 15-40% · MARGINAL <15% of measured top-of-page real estate* and which it explicitly ties to the mature posture: the target is presence at minimum defense cost, not maximum spend.
5. Where the composite reads MARGINAL on the product's top-10 terms, route the rebuild under v4.0 Section 8.10 S-F8: cheapest surface first, at defense budgets, with SP CPC escalation as the last instrument rather than the first. Name the terms and the budget line; the format selection and the build belong to #17-19.
6. Check every defensive row against its own objective band. v4.0 Section 6 rule O6 grades Defensive on *IS held, ACOS 10-15, incrementality (E2)* and states plainly that judging a campaign on another objective's metrics is a process violation. A defensive campaign graded against the Conversions band will be cut when it is working.
7. Where defensive ACOS drifts above its band with impression share held, run the S-D5 auction-pressure check before touching a bid: if competitor conquesting intensified, hold and log the pressure period; if it did not, the drift is overbidding and the ladder resumes at P3 mechanics.
8. Where a term carries the defense-priority tag, apply S-D6: the tag protects it, and it is evaluated against the defense budget and the SOV target rather than against weekly ACOS bands. Protection is a declared, budgeted, quarterly-reviewed decision and not a standing exemption; a protected term with no quarterly review on file loses the protection.

### P5. The annual defense-cost read (60 minutes per primary syntax)

1. Read the declared defense budget from T1 group S7 "Defense budget \$". v4.0 Section 10 rule M8 defines it as *per primary syntax, defense budget \$ = (syntax organic revenue at risk x displacement probability) bounded by syntax contribution margin* and declares it quarterly. This procedure scores the declaration against a year of outcomes; it does not re-derive the number.
2. Sum trailing defense spend on the syntax across all surfaces: brand broad, brand exact, own-ASIN PAT, complement PAT, and brand-format placements. Cross-surface aggregation is mandatory. A defense budget scored against one surface understates its own cost and always reads well.
3. Compute branded revenue protected over the same period from the syntax's organic and branded sales lines, and state the displacement assumption explicitly as the assumption it is.
4. Divide: defense spend against branded revenue protected, expressed as a percentage, beside the M8 declared ceiling and beside the syntax contribution margin that bounds it.
5. Where the defended set includes terms carrying an E2 cannibalization flag, subtract the spend on those terms from the numerator and state the adjusted figure alongside the raw one. A defense-cost read that counts cannibalized spend as protection is measuring the wrong thing, and the incrementality reads at P2 step 6 are the evidence.
6. Where no read exists on file for the period, the finding is F4 and it is recorded as such. Protection has a price and an unmeasured price is not a decision, it is an assumption with a budget line.

## P6. Code-red declaration and the same-day freeze (45 minutes, same day, mandatory)

1. Confirm the trip against v4.0 Section 11 rule R1: *product TACOS >2x stage band any validated week, or >1.5x for 2 validated weeks, or by executive directive*. Write the arithmetic into the declaration: the week's spend, the week's sales, the resulting TACOS, the band ceiling, the ratio, and which limb fired. A declaration that does not show its ratio cannot be audited and cannot be reversed on evidence.
2. Confirm the week is validated: closed week, packet present, sales from the SellerBoard read rather than a console estimate. Where the packet is DEGRADED, declare anyway and carry the T7 completeness score inside the declaration record.
3. Declare in T5, the product log, as a dated event, and set the product state to Recovery Mode. R1 requires the entry to log the entry TACOS, the target band, and the projected weekly overspend in dollars. Compute the overspend as the week's actual spend minus the spend that the band ceiling would have permitted at the week's sales. That dollar figure is the mode's denominator and every action logged from here scores against it.
4. Freeze all non-defensive bid increases immediately, account-wide on the product, and all modifier increases with them. The freeze is not a cut; it stops the bleeding from receiving deliberate help while the diagnosis runs.
5. Apply the R3 temporary overrides in the same session and record that they are live: no new Ranking objectives except RECOVERY-class terms at HIGH confidence, discovery budgets floored at sentinel levels with auto left on and small, tightened human-confirmation thresholds on bid and budget moves, and step sizes capped at standard even where confidence is HIGH. The tightened confirmation thresholds are v4.0's numbers, read from R3 at execution time.
6. Set the R4 cadence: daily pacing and gate checks, twice-weekly mini-audits on the triage set, the weekly full cycle continuing unchanged. Recovery Mode is a faster and stricter version of the same system, not a different one, and the weekly cycle is what keeps it honest.
7. Confirm the declaration with the checkpoint holder before the freeze deploys. The Authority Matrix row "Code-red declaration + close" places the decision with the PPC Lead; on a launch-stage product the matrix standing rule splits the cell, with the launch operator declaring and the PPC Lead closing.
8. Notify the BM owner, who holds the informed position on this row, and the deals owner where an event window is armed or approaching. Rule T8 orders the modes: Recovery outranks Event, and a product in Recovery entering an event runs it at defense-and-proven-terms posture only.

## P7. The four-map diagnosis (3 to 4 hours of work; see the timing note below)

Four maps, built in this order, each an artifact rather than an impression. v4.0 Section 11 rule R2 fixes the triage order and this procedure is that order made executable. Nothing is cut until all four maps exist or are recorded as unreadable. **Timing.** R2 governs two of the four windows and no more: the gates sweep runs *same day*, and the top-20 spend terms each receive *a full Master Action Record this week*. v4.0 sets no completion deadline for the diagnosis as a whole, so this SOP states the two governed windows and does not invent a third. SOP-26 v1.0 carried a 48-hour diagnosis window that traces to no source; it is retired here and the gap is amendment item AM-12.

### Map 0. The gates sweep (R2 step 1)

1. Run quarantines, truncations, eligibility, MKL hygiene and inventory bands across every campaign on the product, same day. This is R2's first step and it precedes the spend maps because a truncated or quarantined row will otherwise be diagnosed as a spend problem it does not have.
2. Record each failure with the action class it blocks. A G4 eligibility flip found here halts execution for that target and opens the listing task the same day; a G1 quarantine removes the row from every subsequent map.

### Map 1. The waste map (R2 step 2)

1. Pull the top-20 spend terms per product in spend order. R2 requires each to receive a full Master Action Record this week, so the map is built at record grain rather than at summary grain.
2. For every term compute, in columns: window clicks, CPC, spend, CVR, orders, contribution recovered as orders multiplied by contribution per order, and the loss as spend minus contribution recovered.
3. Compute each term's own Max Profitable CPC as contribution per order multiplied by that term's own CVR, and place the observed CPC beside it. This is the per-click judgement surface and it is per term, never per product: two terms at the same CPC on the same product have different ceilings because they have different conversion rates.
4. Rank the map by weekly loss dollars, descending. Loss dollars outrank the ACOS column in triage, because a 90% ACOS on eleven dollars is a rounding error and a 41% ACOS on nine hundred is the emergency.
5. Flag the obvious-waste classes R2 names for immediate execution: zero-sale spend on irrelevant terms, GRADUATED-UNSTEERED leaks where an exact is live and the source was never negated, and duplicates. These execute without waiting for the rest of the sequence.
6. Sum the map's loss dollars and express them as a share of the declaration's weekly overspend. That share is the answer to the only question the close will ask: was this a term-level problem.

### Map 2. The placement bleed map

1. Read T4 per campaign per placement for the campaigns carrying the waste map's terms: spend, clicks, orders, CPC, CVR, and each placement's ACOS against the term's governing band.
2. Compute each placement's reading as a multiple of the campaign's blended reading. A term profitable at rest-of-search and ruinous at top-of-search is two decisions and one blended number hides both.
3. Mark any placement under the v4.0 traffic minimum as UNREAD. An UNREAD placement can receive a probe but never a premium, and it never enters the cut list on an inferred reading.
4. Where the B7 placement feed is MISSING per T7, write the map as not readable, name the gap in the declaration record, and proceed on the remaining maps. Ledger dependency D3 records this feed as open and the honest statement of the gap is the requirement.
5. Sum the bleed dollars and express them as a share of the weekly overspend, as in Map 1.

### Map 3. The duplicate and structure scan (R2 step 4)

1. Read T2 block K8 registry counts across the product for terms live in two or more same-type instances with no declared distinct role. The four legitimate roles are declared at SOP-12 Section 1.2 and anything outside them is a duplicate consolidating to the winner.
2. Run the S-R12 add-exact check in the same pass: any priority term with zero exact campaigns, or with auto and broad carrying more than the X10 share of its clicks, is a structure failure rather than a performance signal and it is repaired before it is judged.
3. Run the X6 sibling divergence check where the D6 feed is present: the same term across siblings at a CPC ratio above the v4.0 factor is a structural inspection in X6's fixed order, and settings archaeology is the presumption until proven otherwise.
4. Hand every consolidation to #12 P7 and every add-exact to #12 P1 or P2. Survivor selection is human-mandatory.

#### Map 4. The automation check

1. Read the #37 rule change log across the window and list every automated change that touched a defended or high-spend row.
2. Test each for ladder runaway: a rule that stepped a bid in one direction repeatedly against a stale target, without a human confirmation, is the pattern to find. Compare each affected row's current bid against its recorded floor and count the steps between them.
3. Where a runaway is found, suspend the rule pending #37 review and record the suspension notice back to #37. This SOP never edits a rule; it suspends and routes.
4. Attribute the runaway's dollars: the spend above floor on the affected rows across the window, expressed as a share of the weekly overspend.

Complete Maps 1 to 4 inside the week R2 allows for the Master Action Records, since the records are what the maps are built from. Close the diagnosis by reconciling the four shares against the declaration's overspend figure. Where the four maps do not account for the overspend, the residual is stated as a residual and named in the close. A diagnosis that reconciles to less than the overspend has found part of the problem, and saying so is what stops the day-14 close from being premature.

#### P8. The cut sequence, ordered by loss dollars (2 to 3 hours)

Cuts execute in loss-dollar order, not in the order they were found, and each cut obeys the standing lever precedence of v4.0 Section 9 rule L1: *budget (if G3 truncated) then bid then placement then negation then structure (graduation/consolidation) then status. Pause is last, after documented exhaustion of the others. Never skip forward.* Emergencies do not license lever-skipping; they license a faster queue.

1. Take the waste map in loss-dollar order. For each term, issue the verdict v4.0's scenario matrix returns and execute it. The common case on a mature product is BID DOWN TO PROFITABLE CPC; CAP SYNTAX, executed as a single-step bid edit to the row's own Max Profitable CPC value in T2 block K6. Two terms with opposite four-period trends never receive the same execution, and the trend context travels in the reason string.
2. Where a term's levers are documented as exhausted and the conditions of v4.0 Section 9 rule L7 are met, issue PAUSE (LEVERS EXHAUSTED) and enter it as PAUSED-PENDING with a reason and a named re-entry criterion. A pause without both is incomplete and returns to the pass.
3. Cut bleeding placements per #28 P6 in that SOP's lever order: the modifier first where one exists, then the base bid reduced against the bleeding placement's own CPC, then a placement split at persistent divergence. A flat bid cut across a campaign whose placement map shows one bleeder punishes the placements that work and it is the exact failure the placement read exists to prevent.
4. Land the negation sweep. Route the zero-sale and off-taxonomy terms from the waste map to #29 with their evidence lines, at the wall level #29 prescribes. This SOP supplies evidence and never selects the wall level.
5. Execute the structure repairs the scan found: consolidations to #12 P7, add-exact standups to #12 P1 or P2 with steering walls in the same session per #29.
6. Confirm each suspended automation rule with #37 and record the confirmation.
7. Stage every cut through #38 as a reasoned deployment carrying the code-red reference in the reason string, so the T+1 diff and every later audit can isolate the protocol's changes from the week's ordinary work.
8. Log each cut's expected weekly savings in T6 at the moment it is staged. R4 makes the running total against the entry overspend the mode's scoreboard, and a scoreboard assembled retrospectively is a story rather than a measurement.
9. Apply the tightened R3 confirmation thresholds to every cut. Cuts are still human-confirmed at the tightened levels; the speed comes from the queue, not from the standard.

#### P9. The first re-read, at the close of validated week 1 (45 minutes)

1. Read the week's TACOS and its ratio to band exactly as at P1 step 3, and place it beside the entry reading.
2. Read the direction, not the level. The test at day 7 is whether the trajectory has inflected, and v4.0 Section 4 rule X8 is the governing read: trajectory beats snapshot, and an improving reading continues its current treatment even at an ugly absolute level.
3. Not yet at band with a clear inflection is an acceptable day-7 reading and the protocol continues without deepening. Not yet turning is not acceptable and it deepens the cut sequence: return to the waste map, extend from the top-20 to the next spend decile, and re-run P8 on the extension.
4. Read the R4 scoreboard: expected weekly savings logged against realized change in weekly spend. A large gap between the two, with the cuts deployed and verified, means the diagnosis missed a driver and the four maps re-open rather than the cuts deepening blindly.
5. Confirm the freeze and the R3 overrides are still live. Recovery discipline lapsing quietly at day 7 is the most common way a protocol ends without closing.

#### P10. The close, or the structural escalation (60 minutes)

Day counts below are the R4 weekly cadence expressed on the operator calendar: day 7 is the close of validated week 1 after declaration, day 14 of week 2, day 21 of week 3. The readings are weekly validated readings, which is what both limbs of R1 and the R5 exit test are written in.

1. Read the day-14 reading and classify it. This SOP's v1.0 closed the protocol here on a single in-band reading. That is not the exit test. v4.0 Section 11 rule R5 states the exit as *TACOS within 1.2x stage band for 2 consecutive validated weeks, with no gate backlog. Exit is logged with the recovered weekly dollars.* A single in-band week is the first of two.
2. Where the day-14 reading is within 1.2x band, record it as week one of the exit test, hold the freeze and the R3 overrides, and re-read at day 21. Where day 21 is also within 1.2x band and the gate backlog is clear, the exit is legal.
3. Where the day-14 reading is approaching band but above 1.2x, the protocol continues on the R4 cadence and the exit clock has not started. Approaching is not within.
4. Where the reading is still elevated with no credible trajectory toward band, escalate STRUCTURAL. The problem is pricing, competition, or product wearing a PPC costume, and the four maps route out of PPC with the escalation. v4.0 Section 7 rule D6 sets the first check on the other side of that boundary: the Price Position Index precedes any listing spend, and money is not spent fixing images when the market is paying materially less for the same promise. This SOP stops at the routing and does not diagnose price or product.
5. On a legal exit, log it with the recovered weekly dollars: entry weekly spend against exit weekly spend at comparable sales, stated as a dollar figure and as a percentage. Clear the R3 overrides in one recorded action so that their end is as dated as their start.
6. Assemble the artifact chain and confirm it is complete before the close is accepted: the declaration with its ratio arithmetic, the four maps, the cut records with their staged sets and T+1 diffs, both re-reads, and the exit reading pair. A code red closes with its chain or it does not close.
7. Return the product to floor state and re-enter P1 on the normal cadence. The floors discovered during the cut sequence are recorded in T4 exactly as any other floor, with a context note stating they were discovered under Recovery conditions, because a floor found in an emergency is a floor found against a market that was not reacting to us yet.

### P11. Incursion detection and the response ladder (30 minutes)

1. Detect from two artifacts weekly: branded impression share and its band in T2 block K3 on branded rows, and competitor brands or ASINs appearing in the brand campaign's served terms, which arrive from the #14 P4 brand-broad weekly read.
2. Classify the cause before responding, using the D8 competitor block where the feed is present: a competitor entry, a deal window, a price move, or a review surge. Where D8 is MISSING the response still runs on the impression-share signal and the classification gap is recorded.
3. Respond inside the defense ceiling. v4.0 Section 8.6 S-D2 governs: *Restore one ladder step; open/expand PAT self-targeting to occupy our own pages; the syntax defense budget (M8) is the ceiling*. One step, not a matched escalation.
4. Hold the ceiling absolutely. A bidding war on our own brand terms is a tax a competitor levies at will and can raise at no cost to themselves, and the fight above the ceiling moves to conversion assets rather than to bids. Where the incursion persists at the ceiling, the answer is not a higher ceiling.
5. Escalate a persistent incursor to #16 P10 for counter-conquest evaluation, with the flag profile inputs this SOP holds. #16 decides whether a counter-open is legal on the evidence; unflagged, the answer stays defensive, because conquest against a stronger competitor is a donation.
6. Receive the counter-conquest evaluation record back from #16 and file it with the incursion record. The escalation is not closed until the evaluation returns, whichever way it decides.
7. Where the incursion coincides with a competitor deal window, apply S-C5 through #16 rather than through bids: stand down on that target for the window and resume at window end. Bidding into a competitor's deal price buys the worst comparison of the month.

### P12. The monthly profit-harvest read (50 minutes per product)

The mature portfolio is the funding queue's treasury and this read is how it pays in. Its whole discipline is in one distinction: a spend reduction is not a release until the position it was buying is proven to have held.

1. Compute current floor-state spend on the defended set across all surfaces for the month, from T6 filtered to floor-state rows and reconciled against T3 campaign spend.
2. Read the prior-quarter baseline for the same set from the same log. The baseline is a recorded figure, not a recollection, and where the set's membership changed between periods the comparison is made on the intersection with the difference stated.
3. Compute the raw difference. This number is a candidate, not a release.
4. Verify the position on the same set, on two axes. Rank: mean and per-term position against the prior period. Share: the E5 Share of Voice composite against its prior class. The release is real only where both held flat or better.
5. Where rank held but share fell a class, the release attributable to the terms whose share fell is withheld and routed to the S-F8 rebuild at P4 step 5 rather than posted. Withheld dollars are not a saving deferred, they are a saving that was not real, and posting them would fund a push out of money that was buying shelf presence.
6. Where both held, post the verified difference to the M5 funding queue at T0 zone H and, for family-level arbitration, at Account Rollup tab A2. Per Ledger conflict C3 the posting names the instrument and the tab; there is no #31 procedure to cite and citing one is the dangling-citation defect C3 records.
7. Write the posting into T6 with its verification evidence: the two axes, the periods compared, the withheld portion and its reason. The M5 queue funds pushes from landed savings, so a posting without verification evidence is an over-subscription waiting to happen.
8. Where the read finds floor-state spend rising against baseline with position unchanged, that is drift the weekly pass missed, and it routes back to P2 step 3 with the count of affected rows recorded as an F2 detection.

### P13. The breach-band hold (35 minutes)

The procedure that closes the gap between the instrument's flag and the protocol's trip. A product reading BREACH at 1.5-2x band for one validated week is not in Recovery Mode and is not to be treated as normal either. v4.0 Section 8.1 S-A3 defines the state exactly.

1. Confirm the classification: ratio to band ceiling between 1.5x and 2x, one validated week, and no prior week above 1.5x. A prior week above 1.5x makes this a declaration under R1's second limb and it goes to P6 instead.
2. Freeze new Ranking objectives on the product. No new push is granted while the product reads BREACH.
3. Test every live push against two conditions: rank velocity above zero, and spend inside its recorded loss ceiling. A push failing either does not continue; it takes the stop-loss its own sizing block already carries.
4. Cap CPCs above the 1.5x-ceiling line to the PUSH-ONLY band boundary pending review. The band structure is v4.0 Section 3 family F1: *≤0.8x ceiling SAFE · 0.8-1.0x AT-CEILING · 1.0-1.5x PUSH-ONLY · >1.5x UNJUSTIFIABLE outside a sized push*. Compute each row's ceiling from its own CVR and cap to 1.5 times that value, within the L2 per-change limit.
5. Run the E6 opportunity-cost reallocation: the expected value of the best unfunded push against the expected value of the worst-performing funded spend of equal size. A positive opportunity cost is a standing order to move budget rather than a suggestion.
6. Record the hold in T6 with the week's ratio and a dated re-read at the next validated week. This entry is what the second-week confirmation trigger reads; without it the two-week limb of the trip test cannot fire and the product can sit at 1.7x band indefinitely with every individual week looking like a bad week.
7. Where the next validated week reads at or above 1.5x, declare per P6 the same day. Where it reads below, the hold lifts and the product returns to the ELEVATED or WITHIN path at P1.

### P14. The post-exit review (90 minutes, within one week of exit)

1. Answer the R6 question in writing: which rules, had they been enforced, would have prevented entry. The question is deliberately about enforcement rather than about rule content, because most entries trace to a rule that existed and was not run.
2. Reconcile the four maps' shares against the entry overspend and state what the residual was and whether it was ever explained.
3. Score the cut records: expected weekly savings against realized, per cut, so that the next declaration sizes its cuts from evidence.
4. File the answers as V3 amendment proposals through the threshold-change path. The Authority Matrix places v4.0 threshold changes with the CEO deciding and the PPC Lead executing and reviewing, so this SOP's output is a proposal with evidence, never an edit.
5. Where a rule fired correctly and was overridden, log the override reason. Falling override rate is the quality metric of the system and an override with no reason is invisible to it.
6. Where the entry traced to a detection failure rather than a decision failure, raise it against this SOP's Element 2 trigger table at the next version increment. A protocol that trips late is a detection problem wearing a discipline costume.

## 5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds and scenario logic live in v4.0 and the closed vocabulary is quoted verbatim from the Keyword Decision Record Template v3.0, tab "Bands & Verdicts Reference", which carries the twenty-five strings the exemplar maps. Per Ledger conflict C10 the two divergent rows carry their v4.0 Section 13.1 string as a stated cross-map so neither citation is dropped. This table maps all twenty-five to what this SOP does when each lands on a mature-stage row, including the rows where the answer is that another component executes.

| Verdict (KDR Template v3.0, closed vocabulary)                                              | Mature-stage execution                                                                                                                                                                                                                                                                                       |
|---------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>QUARANTINE: DATA REFRESH</b>                                                             | The row leaves every map and every cut list this cycle, including inside a declared code red. v4.0 Section 12 rule P8 places data quality above urgency; the refresh task carries an owner and a date.                                                                                                       |
| <b>UNLOCK TRAFFIC</b>                                                                       | Rare at maturity and always a finding. A defended term below the sufficiency threshold is either mis-scoped as defense or has lost its position quietly; check the floor register and the position before stepping any bid up, because the freeze at P6 step 4 bars this verdict entirely inside a code red. |
| <b>FIX BUDGET (TRUNCATED)</b>                                                               | Budget lever only, executed at #12 P11 step 3. On a defensive campaign this is the S-D3 case and it runs before any bid conclusion is accepted.                                                                                                                                                              |
| <b>CONTINUE TEST TO 50 CLICKS AT CAPPED BID</b>                                             | No mature-stage execution beyond budget assurance. Inside a code red the R3 discovery floor applies and the test runs at sentinel level rather than stopping.                                                                                                                                                |
| <b>FUND TEST (UNFUNDED PRIMARY)</b>                                                         | Not issued from this SOP. Inside BREACH or Recovery the funding freeze bars it; on a floor-state product the funding decision is an M5 queue decision at the instrument layer.                                                                                                                               |
| <b>RANK PUSH: FUND SIZED PUSH</b>                                                           | Not this SOP's verdict. A mature product receiving a sized push is a stage question first: either the stage tag is wrong or the push is a re-ranking case at #24. Inside BREACH, S-A3 bars new Ranking objectives; inside Recovery, R3 bars them except RECOVERY-class terms at HIGH confidence.             |
| <b>RECOVERY PUSH</b>                                                                        | The one push class R3 permits inside Recovery Mode, and only on a D10 RECOVERY-class term at HIGH confidence. Execution belongs to #24 and the standup to #12; this SOP confirms only that the Recovery-mode condition is met.                                                                               |
| <b>TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS</b>                                             | Not a mature-stage verdict. Its arrival on a mature product means a rank case is live and the stage tag needs review at the monthly stage renewal.                                                                                                                                                           |
| <b>FIX CONVERSION FIRST; DEFER PUSH</b>                                                     | Cap the row at maintenance bids and hold. The fix routes to the funnel-drop-point owner outside this SOP. On a mature product this verdict frequently accompanies a chronic syntax-level conversion failure, where bids are frozen at maintenance across the whole syntax and further bid work is theatre.   |
| <b>PLACEMENT FIX FIRST</b>                                                                  | Bids hold; #28 executes the placement program and stages the campaign-level rows back. Inside a code red this is the second cut lever at P8 step 3 and it runs in #28's lever order.                                                                                                                         |
| <b>RE-POINT VARIATION</b>                                                                   | Executed at #12 P10 same day on a G6 RED reading. Inventory outranks mature posture: a defended floor on a variation about to go out of stock is defending a position that will not be servable.                                                                                                             |
| <b>TAPER: RANK ACHIEVED, STEP BIDS DOWN</b>                                                 | Arrives from #25 at stage entry rather than being issued here. Where it arrives mid-stage, the ladder mechanics are #28 P5 and the execution is #12 P4; this SOP records the resulting floor in the register.                                                                                                |
| <b>HOLD; DEFEND LIGHTLY AT RANK</b>                                                         | The standing state of the mature stage. No change this cycle; the week-over-week block is read next cycle; drift from the held position is the trigger, and the drift check is P2 step 3.                                                                                                                    |
| <b>DEFEND AT FLOOR</b>                                                                      | No bid change where the bid is at the recorded floor. Where drift from the floor is detected, the bid restores to the floor value in one staged step per P2 step 3.                                                                                                                                          |
| <b>BID DOWN TO PROFITABLE CPC; CAP SYNTAX</b>                                               | The workhorse verdict of the cut sequence. Bid edit to the row's own Max Profitable CPC value from T2 block K6, single step, within the L2 per-change limit. The syntax-cap half executes at T1 through the instrument layer and this SOP receives the resulting deltas.                                     |
| <b>INCREMENTALITY LADDER</b>                                                                | Issued at P2 step 6 on an E2-flagged branded or rank-3-or-better term. This SOP stages the dated step-down and records the floor; the measurement design and the continue-or-restore call belong to #35.                                                                                                     |
| <b>ADD EXACT + RESTRUCTURE</b><br>v4.0 Section 13.1 string: identical                       | Surfaced by the structure scan at P7 Map 3 and executed at #12 P1 or P2 with steering walls in the same session per #29. A mature product with an uncovered head term is carrying a launch-era gap into its profit stage.                                                                                    |
| <b>CONSOLIDATE DUPLICATES</b>                                                               | Handed to #12 P7. Survivor selection is human-mandatory per the Authority Matrix and is not delegated inside a code red.                                                                                                                                                                                     |
| <b>GRADUATE FROM DISCOVERY</b><br>v4.0 Section 13.1 cross-map: GRADUATE + STEER             | Owned by #30 and unaffected by mature-stage posture. The harvest never stops: a code-red declaration freezes bid increases, not the graduation pipeline, and a graduate arriving mid-protocol is stood up and walled the same day.                                                                           |
| <b>NEGATE (PRE-LOAD / REACTIVE / STEERING)</b><br>v4.0 Section 13.1 cross-map: CUT / NEGATE | Owned by #29. This SOP produces evidence lines from the waste map and never selects the wall level. The negation sweep is P8 step 4.                                                                                                                                                                         |
| <b>ENTER SBV ARBITRAGE</b>                                                                  | Relevant at maturity through the share-of-voice route: where the E5 composite reads MARGINAL, S-F8 names the cheapest surface first. This SOP names the terms and the defense budget; the format build belongs to #17-19 and the paired SP step-down to #12.                                                 |
| <b>OPEN / ROTATE / EXIT CONQUEST</b>                                                        | Owned by #16. This SOP escalates a persistent incursor at P11 step 5 and receives the counter-conquest evaluation record back; it never opens a target.                                                                                                                                                      |
| <b>STOP-LOSS: CONCEDE OR DEFER</b>                                                          | Executed on any push still live when BREACH or Recovery is entered and failing the velocity or loss-ceiling test at P13 step 3. Freed budget returns to the envelope and reallocates through the M5 queue at the instrument layer.                                                                           |
| <b>PAUSED-PENDING (REASON + RE-ENTRY)</b>                                                   | Executed at #12 P8 with a reason and a named re-entry criterion. Inside a code red the pause is the last lever after documented exhaustion of the others, not the first reach.                                                                                                                               |
| <b>LEAVE TO HARVEST</b>                                                                     | No exact instance is stood up; any live exact on the term is evaluated for demotion. The term runs on discovery economics under #30's rules regardless of the product's stage.                                                                                                                               |

### 5.1 Campaign-Level Verdicts

Campaign rows carry their own subset vocabulary, quoted from the KDR Template v3.0 "Campaign Decision Record" tab: *FUND / HOLD / TRIM BUDGET · RAISE / CUT MULTIPLIER · NEGATION PASS · CONSOLIDATE · RETAG OBJECTIVE · PAUSE (L7) · ARCHIVE*. On a mature product, TRIM BUDGET follows the bid down rather than leading it, so each ladder step reads cleanly; RAISE MULTIPLIER is barred entirely inside a declared code red by the P6 step 4 freeze; and RETAG OBJECTIVE is the verdict that catches a push-era campaign still tagged Ranking inside a mature stack, which is failure mode F5.

5.2 Verdict Strings Retired at This Version

SOP-26 v1.0 Section 5 issued four strings that appear in neither the KDR Template v3.0 closed vocabulary nor v4.0 Section 13.1: RE-LADDER (QUARTERLY), CODE-RED: FREEZE / CUT / DEEPEN / CLOSE / STRUCTURAL, INCURSION RESPONSE, and RELEASE TO QUEUE. A closed vocabulary is what makes the decision log queryable and predictions scoreable per verdict type, and a locally minted string is invisible to both. All four are retired. The operations they named survive as procedures rather than as verdicts: the quarterly re-ladder is P3 and issues DEFEND AT FLOOR or HOLD; DEFEND LIGHTLY AT RANK per term, the code-red stages are P6 to P10 and issue the cut verdicts above per row, the incursion response is P11 and issues DEFEND AT FLOOR inside the ceiling, and the release is P12 and is a posting rather than a verdict. Only DEFEND AT FLOOR carried forward from the v1.0 table unchanged.

6. Worked Examples

---

All four examples are built on the standard fictional training base, because the anonymized live product economics spine is not attached to this session: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%, CVR 5.2% unless a term's own rate is stated. Values are illustrative and internally consistent. Live executions read the T0 economics spine named cells instead, and the arithmetic is identical.

## WE-1. Code red from declaration to legal exit

**Setup.** A mature-tagged product, stage band 4-7%, band ceiling 7.0%. Week 1 closes at spend \$4,494.00 on sales \$42,000.00, giving TACOS 4,494.00 / 42,000.00 = 10.7% and a ratio to ceiling of 10.7 / 7.0 = 1.53x. That is BREACH, not CODE RED, and it was held at P13 with a dated re-read (this is WE-2). Week 2 closes at spend \$4,578.00 on sales \$42,400.00, giving 4,578.00 / 42,400.00 = 10.8% and a ratio of 1.54x. Second consecutive validated week above 1.5x: R1's second limb fires.

**P6 declaration.** Entry TACOS 10.8%, target band 4-7%, and the projected weekly overspend computed as actual spend minus the spend the ceiling would have permitted at this week's sales: 4,578.00 minus (0.07 x 42,400.00 = 2,968.00) = **\$1,610.00 per week**, or \$230.00 per day. That figure is the mode's denominator. Freeze deployed same day on all non-defensive bid and modifier increases; R3 overrides live and recorded; R4 cadence set; declaration confirmed by the PPC Lead per the Authority Matrix.

**P7 Map 1, the waste map.** Top-20 pulled; five terms carry the concentration. Each term's ceiling is its own contribution multiplied by its own CVR.

| Term  | Clicks/wk | CPC    | Spend      | CVR  | Orders | Contribution | Loss/wk         | Own ceiling |
|-------|-----------|--------|------------|------|--------|--------------|-----------------|-------------|
| W1    | 210       | \$2.40 | \$504.00   | 4.0% | 8.40   | \$235.20     | \$268.80        | \$1.12      |
| W2    | 160       | \$2.10 | \$336.00   | 4.5% | 7.20   | \$201.60     | \$134.40        | \$1.26      |
| W3    | 130       | \$2.80 | \$364.00   | 3.5% | 4.55   | \$127.40     | \$236.60        | \$0.98      |
| W4    | 95        | \$3.10 | \$294.50   | 3.0% | 2.85   | \$79.80      | \$214.70        | \$0.84      |
| W5    | 88        | \$2.25 | \$198.00   | 2.5% | 2.20   | \$61.60      | \$136.40        | \$0.70      |
| Total | 683       |        | \$1,696.50 |      | 25.20  | \$705.60     | <b>\$990.90</b> |             |

Map 1 accounts for 990.90 / 1,610.00 = **61.5%** of the weekly overspend. The v1.0 record of this incident reported 61%; the recomputation at record grain gives 61.5%.

**P8 cut sequence, and the L2 limit that shapes it.** The verdict on all five rows is BID DOWN TO PROFITABLE CPC; CAP SYNTAX, and the target is each row's own ceiling. Three of the five cuts cannot be executed in one step, because L2 caps a change at 50% of the current value: W1 at 2.40 to 1.12 is a 53.3% cut, W3 at 2.80 to 0.98 is 65.0%, W4 at 3.10 to 0.84 is 72.9%, and W5 at 2.25 to 0.70 is 68.9%. Only W2 at 2.10 to 1.26 (40.0%) lands in one step. The compliant program is therefore two staged sets:

| Term  | Step 1           | Cut % | Step 2           | Cut % | Wk-1 saving at constant clicks | Wk-2 additional |
|-------|------------------|-------|------------------|-------|--------------------------------|-----------------|
| W1    | \$2.40 to \$1.20 | 50.0% | \$1.20 to \$1.12 | 6.7%  | \$252.00                       | \$16.80         |
| W2    | \$2.10 to \$1.26 | 40.0% | none             |       | \$134.40                       |                 |
| W3    | \$2.80 to \$1.40 | 50.0% | \$1.40 to \$0.98 | 30.0% | \$182.00                       | \$54.60         |
| W4    | \$3.10 to \$1.55 | 50.0% | \$1.55 to \$0.84 | 45.8% | \$147.25                       | \$67.45         |
| W5    | \$2.25 to \$1.13 | 49.8% | \$1.13 to \$0.70 | 38.1% | \$98.56                        | \$37.84         |
| Total |                  |       |                  |       | <b>\$814.21</b>                | <b>\$176.69</b> |

The two columns sum to \$990.90, the map's full loss, and every proposed bid stays above the L2 floor value. The savings figures are computed at constant clicks, which is the conservative statement of the loss removed; realized click volume falls with the bid, so realized spend savings run higher and order counts run lower, and the T6 prediction states both. **This is why the day-7 test at P9 reads inflection rather than level:** on a compliant ladder, 82.2% of the waste-map saving lands in week 1 and the remainder cannot land until week 2.

**P7 Maps 2 to 4.** Product blended ACOS reads 4,578.00 / 13,080.00 ad-attributed sales = 35.0%. Product pages on two campaigns run at 2.3x that blended reading, or 80.5% ACOS, on \$620.00 of weekly spend: PP sales = 620.00 / 0.805 = \$770.19, contribution at 35.0% of sales = \$269.57, loss = 620.00 minus 269.57 = **\$350.43**, or 21.8% of overspend. Cut per #28 P6 lever order, the modifier first. The structure scan finds one undeclared duplicate pair, handed to #12 P7. The automation check finds one bid-ladder rule that stepped a defensive term up six times against a stale target: recorded floor \$0.62, current bid 0.62 x 1.10^6 = \$1.0984, deployed at cent granularity as \$1.10, a premium of \$0.48 per click, across four affected rows carrying 380 clicks per week = **\$182.40**, or 11.3% of overspend. Rule suspended pending #37 review.

**Reconciliation.** 990.90 + 350.43 + 182.40 = \$1,523.73 against \$1,610.00, leaving a residual of \$86.27, or 5.4%, which is stated as a residual rather than assigned. The four shares are 61.5% term waste, 21.8% placement bleed, 11.3% runaway rule, 5.4% unexplained. The v1.0 record of this incident stated 61% term waste, 22% placement bleed and "the remainder the runaway rule", which folded the residual into the rule's share; the split is separated here because an unexplained 5.4% is the part of a diagnosis most likely to return.

**P9 day 7.** Week 3 closes at spend \$3,780.00 on sales \$42,470.00 = 8.9%, ratio 1.27x, which is ELEVATED. The reading is not at band and it has clearly inflected from 1.54x, so the protocol continues without deepening. Scoreboard: expected weekly savings logged at staging were \$814.21 (week-1 waste steps) + \$350.43 (placement) + \$182.40 (rule) = \$1,347.04; realized reduction is 4,578.00 minus 3,780.00 = \$798.00, or 59.2% of expectation, consistent with the second bid steps landing in week 2 and the T+1 diff confirming mid-week.

**P10 day 14, and the correction to v1.0.** Week 4 closes at spend \$2,900.00 on sales \$42,650.00 = 6.8%, ratio 0.97x, which is WITHIN. Realized reduction from entry is 4,578.00 minus 2,900.00 = \$1,678.00, or 124.6% of the logged expectation: V5 scores the cut program HIT. **The protocol does not close here.** R5 requires two consecutive validated weeks within 1.2x band, so this is week one of the exit test, the freeze holds, and the re-read is dated at day 21. SOP-26 v1.0 closed its worked example at exactly this point on this exact reading, which is a one-week close against a two-week rule.

**P10 day 21, legal exit.** Week 5 closes at spend \$2,760.00 on sales \$42,900.00 = 6.4%, ratio 0.92x, second consecutive week within band, gate backlog clear. Exit logged with the recovered weekly dollars: entry spend \$4,578.00 against exit spend \$2,760.00 at sales that rose from \$42,400.00 to \$42,900.00, a reduction of **\$1,818.00 per week** (39.7%) on sales 1.2% higher, or \$94,536.00 annualized at this run rate. Band-ceiling spend at week-5 sales would be 0.07 x 42,900.00 = \$3,003.00, so the product exits \$243.00 under its ceiling. R3 overrides cleared in one dated action; artifact chain complete; P14 filed within one week.

## WE-2. The breach that is not a code red, and the flag that says it is

**Setup.** The week before WE-1. Spend \$4,494.00 on sales \$42,000.00 = 10.7%, ratio to the 7.0% ceiling of 1.53x. One validated week, no prior week above 1.5x. This is BREACH under the F2 deviation classes and it is S-A3, not S-A4.

**The instrument disagrees with the protocol.** Account Rollup tab A0 carries its code-red flag as TACOS greater than 1.5x band, so the cell lights this week. R1 does not fire, because neither of its limbs is met: the reading is not above 2x, and there is no prior validated week above 1.5x. **The operator declares on R1 and not on the cell.** Declaring here would apply the full R3 override set, including the discovery sentinel floor and the tightened confirmation thresholds, to a product the Criteria System classes one band below the emergency. The divergence is logged as amendment item AM-1 and the A0 cell is read as an early-warning prompt to run P13, which is a useful thing for it to be.

**P13 execution.** New Ranking objectives frozen. Two live pushes tested: one carries rank velocity of 0.6 positions per week and sits inside its loss ceiling, so it continues; the second reads velocity 0.0 across its review window and takes STOP-LOSS: CONCEDE OR DEFER, its budget returning to the envelope.

**The CPC cap arithmetic.** Every row above 1.5x its own ceiling caps to the PUSH-ONLY boundary. A defended term at CVR 5.2% carries Max Profitable CPC =  $28.00 \times 0.052 = \$1.46$ , so its 1.5x line is  $1.46 \times 1.5 = \$2.19$ . The row is running \$2.64, which is in the UNJUSTIFIABLE band outside a sized push, and it caps to \$2.19: a cut of  $(2.64 \text{ minus } 2.19) / 2.64 = 17.0\%$ , inside the L2 per-change limit and above the L2 floor. At 120 clicks per week the row's spend moves from  $120 \times 2.64 = \$316.80$  to  $120 \times 2.19 = \$262.80$ , a saving of \$54.00 per week.

**E6 opportunity cost.** The capped row still returns  $120 \times 0.052 = 6.24$  orders, worth  $6.24 \times 28.00 = \$174.72$  against \$262.80 of remaining spend: a weekly deficit of \$88.08 on the worst-performing funded spend of that size. The best unfunded push in the queue projects contribution above its cost, so the opportunity cost is positive and the reallocation is a standing order rather than a proposal.

**Where it goes.** The hold is written to T6 with the 1.53x ratio and a dated re-read. Seven days later the product reads 10.8% at 1.54x, the second-week trigger fires off that T6 entry, and WE-1 begins. Without the T6 entry the second limb of R1 is unreadable and the product sits above 1.5x indefinitely with each individual week looking like a bad week rather than a pattern.

## WE-3. A quarterly re-ladder that slips, reverts, and is confounded

**Setup.** A defended head term with a recorded floor of \$2.24 from a completed post-push ladder, and a cohort re-test date reached. Max Profitable CPC at the term's 5.2% CVR is  $28.00 \times 0.052 = \$1.46$ , so the floor carries a premium of  $2.24 \text{ minus } 1.46 = \$0.78$  per click. At 30 maintenance clicks per week that is \$23.40 per week, or \$3.34 per day.

**Is the premium justified.** The held position protects roughly 3.1 organic orders per day on the term, worth  $3.1 \times 28.00 = \$86.80$  per day of organic contribution against \$3.34 per day of premium: a ratio of 26.0 to 1. That ratio, not sentiment, is what HOLD; DEFEND LIGHTLY AT RANK means, and it is what the re-ladder re-tests each quarter.

**The probe.** One rung down at the M3 step size:  $2.24 \times 0.90 = \$2.02$ , a 9.8% step, inside the L2 per-change limit and above the L2 floor. Prediction logged: position holds at 4 or better through one read window, review dated.

**The slip.** At the review date the position reads 6 against a prediction of 4 or better. Per #28 P5 the rung reverts one step to \$2.24 and holds two read windows. The register is not updated: a probe that failed does not lower the recorded floor, and the floor stays at the last held rung.

**The confound, and why the revert was still right.** The T5 product log carries two competitor deal windows opening inside the slip window. Under X9 the product log is consulted before bids are blamed, so the slip is a market event rather than a result of the step, and the correct reading is that the probe was uninformative rather than that the floor is \$2.24. The revert was still the right execution, because the rung discipline is not a judgement about cause: it restores the position first and diagnoses after.

**The incrementality read is voided.** This term carries an E2 cannibalization flag, so under S-D4 the probe was doubling as the incrementality test, with total capture watched for two weeks. The deal windows contaminate exactly that window, so the capture read cannot be attributed and it is voided rather than reported. It is re-run outside the confound and the void is handed to #35 with the T5 references, because the continue-or-restore call belongs there and a confounded read handed over as clean is worse than no read.

**The re-probe.** Two read windows later, with both deal windows closed, the probe re-runs at \$2.02 and holds. The register updates: floor \$2.02, dated, context note recording that the first probe slipped inside two competitor deal windows and the second held after they closed. The new premium is  $2.02 \text{ minus } 1.46 = \$0.56$  per click, \$16.80 per week, \$2.40 per day, and the protection ratio moves from 26.0 to  $1 \text{ to } 86.80 / 2.40 = \mathbf{36.2 \text{ to } 1}$ . The quarter's saving on this one term is  $23.40 \text{ minus } 16.80 = \$6.60$  per week, or \$343.20 per year, and it is that figure, not the bid, that P12 will read next month.

## WE-4. A harvest read where the release is not real

**Setup.** Monthly harvest read on a defended set of 11 terms. Prior-quarter baseline floor-state spend \$3,120.00 per month; current floor-state spend \$2,704.00 per month; raw difference **\$416.00 per month**. On the v1.0 procedure this figure posts to the queue.

**Verification, axis 1: rank.** Mean position across the set moves from 3.2 to 3.1, and no individual term slipped past tolerance. Rank held, and slightly better.

**Verification, axis 2: share.** The E5 Share of Voice composite on the product's top-10 terms falls from 22% to 13%. Under the F12 bands that is a move from PRESENT (15-40%) to MARGINAL (below 15%). Share did not hold, and it did not hold in a way that crosses a class boundary.

**The split.** Three of the 11 terms carry the share loss and account for \$174.00 of the \$416.00 reduction. Verified release =  $416.00 \text{ minus } 174.00 = \mathbf{\$242.00 \text{ per month}}$ . The withheld \$174.00 is 41.8% of the raw difference. It is not deferred and it is not restored to bids: under S-F8 it routes to a presence rebuild at defense budgets on the cheapest surface first, with the terms and the budget line named here and the format build belonging to #17-19.

**What posting the whole figure would have cost.** At \$416.00 per month the queue would carry  $416.00 \times 12 / 52 = \$96.00$  per week of funding supply. The verified figure carries  $242.00 \times 12 / 52 = \$55.85$  per week. The difference of \$40.15 per week is a 71.9% over-subscription against the money that actually exists, and it would have funded a push out of dollars that were buying shelf presence. The push would then have been the visible failure, and the harvest read that mis-stated its funding would not have been.

**The posting.** \$242.00 per month to T0 zone H and to Account Rollup tab A2 for family-level arbitration, with the verification evidence written into T6: both axes, both periods, the withheld portion and its reason. Per Ledger conflict C3 the posting names the instrument and the tab; there is no #31 procedure number to cite.

## 7. Failure Modes and Escalation

Twelve modes. Each carries its detection signal, its response, and the account-history incident that motivated it where one exists. The incidents recorded in v1.0 Sections 6 and 7 are carried forward here because they cannot be reconstructed from any other document in the library.

| Failure                                               | Detection signal                                                                                                                                                                                     | Response, and the incident behind it                                                                                                                                                                                                                                                                                                                                                                                    |
|-------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>F1. Trip normalized</b>                            | A validated week meeting either R1 limb with no declaration in T5, or a T6 breach-hold entry with no dated re-read against it.                                                                       | Declare on discovery, run the protocol in full, and root-cause the silence separately from the spend. This is the named account pattern the protocol exists to kill, and the absence of a formal protocol was itself recorded in v1.0 as one of the account's named failure patterns. The escalation carries the days-silent count quantified, because the number of days is the finding, not the declaration.          |
| <b>F2. Floor rot</b>                                  | A cohort past its quarterly re-test date in T4 "Next Step Date", or P12 finding floor-state spend rising against baseline with position unchanged.                                                   | Re-ladder the cohort now; the register diff quantifies what the delay cost. Paying last year's price for this year's shelf is the whole of it, and the cost is invisible weekly and material annually: WE-3's single term carries \$343.20 a year between two adjacent rungs.                                                                                                                                           |
| <b>F3. Brand bidding war above the ceiling</b>        | A defense bid above the M8 syntax defense budget line on an incursion response, or two consecutive restore steps on one term.                                                                        | Restore the ceiling; route the fight per P11 steps 4 to 5. A bidding war on our own brand terms is a tax a competitor levies at will and raises at no cost to themselves, and above the ceiling the instruments are conversion assets, not bids.                                                                                                                                                                        |
| <b>F4. Defense theatre</b>                            | A declared defended set with no defense-cost read on file for the period, or a read computed on one surface.                                                                                         | Run P5 in full across all surfaces. Protection has a price, and an unmeasured price is an assumption with a budget line rather than a decision. Recorded in v1.0 as a standing pattern on this account.                                                                                                                                                                                                                 |
| <b>F5. Push-era structure surviving into mature</b>   | Dedicated push campaigns undemoted with floors long recorded, or a campaign still tagged Ranking inside a mature objective mix in T0 zone F.                                                         | #12 P6 demotions execute and the objective retags in the same staged set. Mature structure is tail-consolidated by default, and a mature product carrying a launch-shaped spend structure is an S-A8 finding rather than a preference. Recorded in v1.0.                                                                                                                                                                |
| <b>F6. The trip tested on the wrong number</b>        | A declaration record citing campaign ACOS, blended conquest ACOS, or ad-attributed ACOS instead of product TACOS against the stage band.                                                             | Void the declaration and re-test on TACOS. The Economics Worked-Example Set states the identity that makes this unambiguous: TACOS states stage posture and ACOS drives campaign verdicts, and applying one to the other's decision is the same mistake in two directions. A product can sit at 35.0% blended ACOS, exactly break-even, while its TACOS runs at 1.54x band, which is precisely WE-1's opening position. |
| <b>F7. Single-week close</b>                          | A code red closed on one in-band reading, or an exit logged with no second dated reading beside it.                                                                                                  | Re-open the protocol and hold to the R5 two-week test. SOP-26 v1.0's own worked example closed at day 14 on a single 6.8% reading, so this failure mode is written against the prior version of this document rather than against an operator.                                                                                                                                                                          |
| <b>F8. Recovery discipline suspended for an event</b> | An armed event window on a product in Recovery Mode with event-plan actions deploying outside the defense-and-proven-terms posture, or R3 overrides cleared before the R5 exit.                      | Restore the overrides same day and re-stage the event actions inside the posture. Rule T8 orders the modes and the ordering is not negotiable at the operator level: recovery discipline is not suspended for a deal.                                                                                                                                                                                                   |
| <b>F9. Undated or oversized rungs</b>                 | A ladder step in T4 with no date and no hold prediction, or a staged step exceeding the L2 per-change limit.                                                                                         | Re-stage the step at the compliant size with its date and prediction. WE-1 shows why this matters arithmetically rather than procedurally: four of the five cuts in that sequence need two steps to stay inside L2, and a single-step version of the same program would be rejected at staging QA and would arrive a week late having looked faster.                                                                    |
| <b>F10. Release posted without verification</b>       | A T0 zone H posting with no rank and share evidence in its T6 line, or a posting whose set membership changed between the compared periods with no intersection statement.                           | Withdraw the posting and re-run P12 steps 4 to 5. WE-4 is the arithmetic: 41.8% of a raw difference can be money that was buying shelf presence, and the queue funds pushes from landed savings, so the error surfaces as a failed push rather than as a failed read.                                                                                                                                                   |
| <b>F11. Runaway automation found late</b>             | A defended row whose current bid sits more than one rung above its recorded floor with no human confirmation in T6, or a #37 rule-change log entry touching a defended row with no matching T6 line. | Suspend the rule pending #37 review, restore the affected rows to floor, and attribute the dollars. In the v1.0 incident one rule stepped a defensive term six rungs against a stale target and the premium was carrying 11.3% of the week's overspend at the recomputation; nothing alarmed on it because each individual step was inside its own limit.                                                               |
| <b>F12. Post-mortem never filed</b>                   | A Recovery exit in T5 with no P14 review filed within one week, or a filed review with no amendment proposals attached.                                                                              | File it. Rule R6 exists because recovery that teaches nothing repeats, and rule V7 makes the filing an obligation rather than a courtesy. The specific failure to watch for is an exit whose review names only operator behaviour and no rule: an entry that traced to a detection gap is a trigger-table defect in this document and it is owed a version increment.                                                   |

**Escalation.** F1 escalates on discovery with the days-silent quantified. F6 and F7 escalate on discovery because both void a declaration or a close and therefore corrupt the artifact chain. A STRUCTURAL close at P10 step 4 escalates by definition, with the four maps attached and the residual stated. Any failure mode occurring twice in consecutive cycles on the same scope escalates to the PPC Lead with the T6 history attached. Everything in the human-mandatory tier (the declaration confirm, every pause, every structure change, survivor selection, rule suspensions, and the STRUCTURAL escalation itself) is confirmed per the Authority Matrix before deployment, without exception and without a Recovery-Mode discount.

## 8. Outputs and Artifacts

Gates clear with artifacts, not assertions. Each row names the record, where it is stored, how long it is held, and who consumes it downstream.

| Event                               | Named record                                                                                                                                              | Storage location                                                                       | Retention                                  | Downstream consumer                                                                                                                              |
|-------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------|--------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Weekly posture read (P1)</b>     | The posture line: state, ratio to band, mix status, flywheel status, gates failing                                                                        | Command Workbook cycle header, echoed to T0 zone A                                     | Append-only log life                       | The BM huddle; every keyword verdict issued this cycle inherits it; Account Rollup A0                                                            |
| <b>Floor maintenance (P2)</b>       | Staged restoration set with floor references, plus T6 rows carrying each prediction and review date                                                       | T6 decision log; staged set in #38 format                                              | Append-only log life                       | #38 for deployment and the T+1 diff; V5 scoring next cycle                                                                                       |
| <b>Quarterly re-ladder (P3)</b>     | The register diff: floors lowered, held, raised, each dated with its competitive context note                                                             | T4 columns "Discovered Floor \$", "Ladder State (step / date)", "Next Step Date"       | Life of the term in the register           | P12 harvest read; #28 as the ladder owner; #35 where a rung doubled as an incrementality step                                                    |
| <b>Defense declaration (P4)</b>     | The defended-set membership register: term, surface, objective tag, floor bid, defense budget line, reason                                                | T1 group S7 for the budget line; membership held with the product's defended set       | Life of the set, versioned per declaration | #14 P4 for brand broad; #16 for defense PAT; #17-19 for format placement and rebuilds                                                            |
| <b>Defense-cost read (P5)</b>       | Defense spend against branded revenue protected, with the displacement assumption and the cannibalization-adjusted figure stated                          | T6, filed against the syntax; the M8 declaration in T1 S7                              | Annual, superseded by the next read        | The quarterly M8 declaration; the BM owner's review row                                                                                          |
| <b>Code-red declaration (P6)</b>    | The declaration record: ratio arithmetic, limb fired, entry TACOS, target band, projected weekly overspend, packet completeness score, R3 override list   | T5 product log as a dated event; product state set to Recovery                         | Life of the product log                    | #37 (rule holds), #32 (event posture under T8), Account Rollup A0, the Authority Matrix checkpoint holder                                        |
| <b>Four-map diagnosis (P7)</b>      | Four maps plus the gates sweep, each with its share of overspend, and the reconciliation with its stated residual                                         | Attached to the declaration record in T5; source reads in T2, T3 and T4                | Life of the artifact chain                 | #28 (bleed map), #12 (structure scan), #29 (negation evidence lines), #37 (automation check), and the cross-framework read on a structural close |
| <b>Cut sequence (P8)</b>            | Cut records: staged sets carrying the code-red reference, each with expected weekly savings logged at staging, plus the T+1 diffs                         | T6 for the records; staged sets in #38 format                                          | Append-only log life                       | #38 verification; the R4 scoreboard; V5 scoring; P14's cut-record scoring                                                                        |
| <b>Re-reads and close (P9, P10)</b> | Day-7 and day-14 readings with ratios, the scoreboard line, and the exit pair with recovered weekly dollars                                               | T5 against the declaration; T6 for the scoreboard                                      | Life of the artifact chain                 | P14; Account Rollup A0; the amendment path on a structural close                                                                                 |
| <b>Incursion response (P11)</b>     | Incursion record: detection artifact, cause classification or the gap, ladder step taken, ceiling reference, and the returned counter-conquest evaluation | T6, cross-filed to the defended-set register                                           | Append-only log life                       | #16 P10 on escalation and back on return; #14 P4 for the brand-broad surface                                                                     |
| <b>Harvest read (P12)</b>           | The release posting with its verification evidence: both axes, both periods, withheld portion and reason                                                  | T0 zone H; Account Rollup tab A2 for family arbitration; evidence in T6                | Append-only log life                       | The M5 funding queue; #35 where an incrementality read informed the verification; family-level M5 quarterly                                      |
| <b>Breach hold (P13)</b>            | The hold record: week ratio, pushes tested and their outcomes, capped rows, E6 reallocation, dated re-read                                                | T6                                                                                     | Append-only log life                       | The second-week confirmation trigger reads this record directly; without it the R1 second limb cannot fire                                       |
| <b>Post-exit review (P14)</b>       | The R6 review: prevention answer, map reconciliation, cut scoring, override reasons, amendment proposals                                                  | Filed with the declaration record in T5; proposals routed to the threshold-change path | Life of the artifact chain                 | The v4.1 amendment brief; V7; this SOP's own Element 2 at the next increment                                                                     |

Retention: no rule in v4.0 sets a retention period for any of these records. The rows above state the log's own append-only life, which is a property of the instrument rather than a governed period. The gap is amendment item AM-8 and it duplicates items already logged by other build sessions, so v4.1 folds it once.

## 9. Skill-Conversion Block

| Block                      | Specification                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|----------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Input bindings</b>      | T0 zone B (weekly spend, weekly sales, TACOS against stage band), zone C economics spine named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC and WeeklyEnvelope, zone D gates strip, zone E state vector rows for flywheel health and stage integrity, zone F TACOS decomposition by objective, zone H envelope and M5 queue. T1 group S3 cap status and group S7 defense budget. T2 blocks K1 (oversight flag, coverage state), K2 (validation state, sample state, in-budget %, inventory band, event mode), K3 (brand impression share, IS% band), K5 (four-period trend, WoW block, zero-quarantine flag), K6 (margin per unit, BE ACOS, max profitable CPC, expected CPC, CPA ceiling), K8 (registry counts, duplicate flag, current bid, TOS modifier, sibling CPC view, competitor block). T3 utilization, in-budget bucket, truncation flag, objective tag, campaign verdict. T4 discovered floor, ladder state, next step date, placement verdict. T5 product log event types. T6 decision log with V5 scores. T7 manifest status per field group. Account Rollup A0 (code-red flag, TACOS against band) and A2 (cross-product M5 queue). The #37 rule change log. The #14 P4 brand-broad weekly read. An absent feed suppresses its dependent execution class per the T7 manifest and is stated; the skill never improvises around a missing feed. |
| <b>Deterministic logic</b> | The ratio computation and its classification into the four F2 deviation classes. The two-limb R1 trip test, which requires reading the prior week's T6 breach-hold entry and is therefore stateful across cycles. The overspend computation (actual spend minus band-ceiling spend at the week's sales). The per-term Max Profitable CPC computation and the L2 step decomposition where a target cut exceeds the per-change limit. The four-map construction and the share-of-overspend reconciliation with its residual. The cut ordering by loss dollars under the L1 lever precedence. The rung and reversion rules and the floor-recording rule (record at the last held rung, never at a failed probe). The release verification on both axes and the withheld-portion split. The R5 two-week exit test with the gate-backlog check. No local thresholds: every band, step size, limit and tolerance dereferences v4.0 at execution time and the skill version-pins v4.0.                                                                                                                                                                                                                                                                                                                                                                                                        |

| Block                                                                               | Specification                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
|-------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Human checkpoints</b><br>Drawn from PPC Authority Matrix v2.0, Section 1, by row | <b>AI-executable with audit:</b> the posture read and its classification, all four maps, the reconciliation, cut drafting in loss-dollar order with the L2 step decomposition, re-ladder rung drafting on schedule, drift detection against the floor register, the release computation and its verification arithmetic, the breach-hold record, the defense-cost read arithmetic, and the P14 reconciliation. <b>Human-mandatory:</b> the code-red declaration confirm and the close, per the matrix row "Code-red declaration + close" (PPC Lead decides, operator executes, PPC Manager executes, BM owner informed); every pause and every rule suspension, per "Automation rules (#37)" (PPC Lead decides, operator executes and reviews); every structural change and survivor selection, per "Structure creation (all SOPs)"; the STRUCTURAL escalation; every deployment confirm, per "Bulk deployments (#38)" (PPC Lead reviews, operator and PPC Manager execute); the release posting into the funding queue, per "Push funding (the M5 queue)" (CEO decides above envelope, PPC Lead decides within it); the incrementality continue-or-restore hand-off, per "Incrementality tests" (PPC Lead reviews, operator decides and executes); the counter-conquest escalation, per "PAT/conquest opens and exits" (PPC Lead reviews, operator decides and executes); and any shared-term restore step, per "Sibling declarations (S-series)" (PPC Lead decides, BM owner reviews). Weekly floor-state verdicts sit on the matrix row "Weekly verdicts (established US)": PPC Lead decides, the operator executes assigned scope, the BM owner reviews in the daily huddle, and that review is review rather than veto with disagreements routing up rather than sideways. On a launch-stage product the matrix standing rule splits the code-red cell: the launch operator declares and the PPC Lead closes. Under the vacancy rule an unfilled role's decision right reverts to the PPC Lead and its execution to the PPC Manager. Recovery Mode tightens the confirmation thresholds per R3 and the skill reads the tightened values rather than the standing ones while the mode is live. |
| <b>Cadence</b>                                                                      | Daily: the trip watch on every mature-tagged product against A0 and T0 zone B, plus pacing and gate checks. Weekly: the P1 posture read and P2 floor maintenance inside the Monday cascade, deployment Tuesday, T+1 verification Wednesday; the incursion scan; the F9 expired-rung scan. Monthly: the P12 harvest read; the P4 defended-set declaration review. Quarterly: the P3 re-ladder per cohort on its calendar; the M8 defense budget re-declaration. Annually: the P5 defense-cost read. Inside Recovery Mode the R4 cadence overlays all of the above: daily pacing and gate checks, twice-weekly mini-audits on the triage set, and the weekly full cycle continuing unchanged.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| <b>Alarm specification</b>                                                          | F1 to F12 each carry a detection signal from Section 7, an evidence-line format (row or product, signal value, the v4.0 rule or gate reference, and the artifact where the signal was read), and a routing tier. Same day to the operator queue: F1, F6, F7, F8, F11. Into the weekly pass worklist: F2, F3, F9, F10. To the PPC Lead with T6 history attached: F4, F5, F12, and any mode recurring in consecutive cycles.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>Skill outputs</b>                                                                | Drafted posture lines with classification and ratio; trip alarms with all four maps pre-drafted and reconciled; cut packets in loss-dollar order with L2-compliant step programs and expected weekly savings pre-logged; drafted register diffs; drafted release postings with verification evidence and the withheld split computed; breach-hold records with dated re-reads; F1 to F12 alarms with evidence lines; the V5 scoring draft for rows due this cycle; and the P14 reconciliation draft. Everything is drafted and queued for confirmation at the tier the matrix names; nothing in the human-mandatory set deploys from the skill.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |

## 10. Version Governance

Owner: PPC Lead, with the Brand Management counterpart per product on the defended-set declaration.

### 10.1 Review Triggers Specific to This Component

- Any v4.0 version increment, at which point four things re-validate before any pass runs: the mature stage band in family F2, the deviation classes and both limbs of the R1 trip, the Recovery Mode rule set R1 to R6 including the R5 exit test, and the Section 5 verdict mapping against the current closed vocabulary.
- Any change to the F2 stage band table, which changes the ratio every reading in this document is expressed in and therefore changes which products are in which state on the day it lands.
- Any change to the M8 defense budget formula, which is the ceiling P11 holds and the denominator P5 divides by.
- Any change to #28's ladder mechanics or step derivation, since P2, P3 and P8 all execute against them and the Ledger records SOPs 12, 23, 24, 25 and 26 as ladder consumers inheriting its changes.
- Every STRUCTURAL close at P10 step 4, which triggers a review of the four maps for coverage: a diagnosis that routed a problem out of PPC either found the boundary correctly or missed a PPC cause, and the maps are the only evidence of which.
- Every Recovery exit, whose P14 review may raise a detection gap against the Element 2 trigger table. A protocol that trips late is a trigger defect in this document.
- Any resolution of Ledger conflicts C9 or C10, which change gate naming in Element 3 and verdict strings in Element 5 respectively.
- Landing of the D3 placement feed, which converts the placement bleed map at P7 from a conditional artifact into a standing one, and of the D4 competitor depth feed, which converts the incursion cause classification at P11 step 2 from optional to standing.

### 10.2 Change Note, v1.0 to v2.0

Supersedes v1.0 in full. Procedures expanded from 5 to 14: P1 posture read, P4 defended-set declaration, P5 defense-cost read, P9 day-7 re-read and P10 close were implicit inside v1.0's P1 to P3 and are now separately executable; P13 breach hold and P14 post-exit review are new. Verdict mapping expanded from 5 rows to the full 25-verdict closed vocabulary with boundary statements, and four locally minted v1.0 strings retired (Section 5.2). Worked examples expanded from 1 to 4. Failure modes expanded from 5 to 12, with all five v1.0 modes and their incidents carried forward. Elements 11 and 12 added. Three corrections of substance: the trip restated as v4.0's two-part test rather than v1.0's single-limb 1.5x reading; the close moved from day 14 to the R5 two-consecutive-week exit test, which the v1.0 worked example failed; and the v1.0 P5 citation of a #31 procedure number retired to the instrument layer per Ledger conflict C3. Thresholds carried forward from v1.0 with their traces re-verified: the mature band 4-7% and the quarterly floor re-test. Incidents carried forward: the normalized trip, floor rot, the brand bidding war above ceiling, defense theatre, push-era structure surviving into mature, and the code-red incident with its five waste terms, its product-page bleed at 2.3x blended, its undeclared duplicate pair and its six-step runaway rule, all recomputed at record grain in WE-1.

### 10.3 Amendment Register: Twelve Items for v4.1

Each item is a numeral or a rule identity this SOP needed and v4.0 does not carry. None is written into the document.

| ID    | Item                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|-------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| AM-1  | The code-red trip is stated at two different tests. v4.0 family F2, Section 8.1 S-A4 and Section 11 R1 all carry the two-part test (above 2x band in a validated week, or above 1.5x for 2 validated weeks). The KDR Template v3.0 "Bands & Verdicts Reference" reads "Code-red at >1.5x band", Data Architecture v1.0 Section 11 tab A0 carries "Code-red flag (T6: TACOS > 1.5x band or spend untraceable to objectives)", and Strategic Doctrine v1.0 Section 14 reads "the code-red trip at 1.5x band forces the protocol". Highest-priority item in this register: it is the single threshold this component is built around, and the instrument currently flags a state the protocol does not recognise. |
| AM-2  | Three stage TACOS bands diverge between v4.0 family F2 and the KDR Bands sheet: Ranking push 10-18% against 12-15%, Transition 7-12% against 8-12%, Harvest at or below 3% against at or below 4%. Mature 4-7% agrees, which is why this component can operate while the item is open, but a product transitioning into or out of mature reads two different bands depending on which document is open.                                                                                                                                                                                                                                                                                                        |
| AM-3  | The KDR Bands sheet and its five worked calibration rows cite rule IDs from the superseded Decision Framework numbering: TACOS bands attributed to M7 and T6, concentration limits to M10, the recovery zone to S-R12, and objective-less spend to G10. Under v4.0 those IDs belong to marginal mathematics, the harvest-to-exact pipeline, confidence-weighted sizing, the add-exact scenario and event exclusion respectively. Distinct from conflicts C9 and C10, which cover gate numbers and verdict counts.                                                                                                                                                                                              |
| AM-4  | No numeral governs the incursion detection bar. Section 8.6 S-D2 fires on "Branded IS slipping" with no defined magnitude and no defined window, so P11 step 1 has a signal and no threshold.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| AM-5  | No numeral governs the day-7 inflection test. Rule R4 sets the cadence and P9 requires the trajectory to have inflected, but v4.0 defines no minimum improvement that counts as inflection, leaving the one judgement call in an otherwise arithmetic protocol.                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| AM-6  | No tolerance governs "held flat or better" in the harvest-read verification. P12 step 4 tests rank and share against the prior period and v4.0 sets no tolerance band on either axis, so the boundary between a real release and a withheld one is currently set per operator.                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| AM-7  | The annual defense-cost read has no governed pass or fail ratio. Rule M8 sizes the defense budget but sets no threshold at which defense spend against branded revenue protected is judged excessive, so P5 produces a number that nothing scores.                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| AM-8  | Record retention. No v4.0 rule sets a retention period for the code-red artifact chain, the floor register or the decision log. Duplicates items logged by prior build sessions; v4.1 folds once.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| AM-9  | The dangling #31 citation, per Ledger conflict C3. SOP-26 v1.0 P5 cited "#31" as though it carried numbered procedures; v2.0 retires the citation to the instrument layer (T0 zone H, Account Rollup A2). Logged as C3 requires, so the register records that one of the five citations C3 names has now been closed at its source.                                                                                                                                                                                                                                                                                                                                                                            |
| AM-10 | No numeral governs the residual tolerance at the diagnosis close. P7 requires the four maps to reconcile against the overspend and to state any residual, but v4.0 sets no share of overspend that the maps must account for before the diagnosis is complete.                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| AM-12 | No window governs the completion of the code-red diagnosis, and none governs delivery of the #37 rule change log into it. Rule R2 sets the gates sweep at "same day" and the Master Action Records at "this week", and stops there. SOP-26 v1.0 asserted a 48-hour diagnosis window with no traceable source; it is retired at v2.0 rather than carried, and P7 now runs to the two windows R2 does govern.                                                                                                                                                                                                                                                                                                    |
| AM-11 | Gate numbering, per Ledger conflict C9. Carried here because Element 3 executes against both numbering schemes; the resolution is v4.1's and this SOP applies the naming convention C9 prescribes in the meantime. Duplicates items already logged; v4.1 folds once.                                                                                                                                                                                                                                                                                                                                                                                                                                           |

## 11. Operator Ramp

Built from the PPC Authority Matrix v2.0 rows that name this component's decisions. Each row becomes a competence the incoming operator demonstrates, and the checkpoint holder named in the matrix is the person who reviews in week 2 and audits in week 4. The live-defense questions test the reasoning behind the thresholds, not recall of their values; an operator who can recite 4-7% and cannot say why a 1.7x week is not a code red has not passed.

### 11.1 Week 1: Observation and Artifact Familiarity

- Reproduce the posture read on three mature products without assistance: locate T0 zone B, compute the ratio to band ceiling, classify it into the four F2 deviation classes, and state which procedure the classification invokes. Demonstrated by matching the checkpoint holder's independent classification on all three.
- Trace the floor register end to end on one term: find the recorded floor in T4, the ladder chain that produced it in T6, and the context note that explains the market it was discovered against. State the term's current premium over its own Max Profitable CPC in dollars per click and per week.
- Read a closed code-red artifact chain and name each of its parts: declaration with ratio arithmetic, four maps with their shares, cut records, both re-reads, and the exit pair. Identify which part would be missing if the protocol had closed at day 14.
- Locate every artifact named in Element 8 in the live instruments. An operator who cannot find the artifact cannot produce it.
- Observe one full Monday cascade and one BM huddle. Attend the huddle as observer: the matrix places the BM owner's review on weekly verdicts, and the incoming operator needs to have seen a disagreement route up rather than sideways before running one.

### 11.2 Week 2: Supervised Execution

The operator executes with the checkpoint holder reviewing every output before it deploys. Scope follows the Transition Brief stack; the matrix rows below name the reviewer.

- Run P1 and P2 on assigned scope. Matrix row: "Weekly verdicts (established US)", where the PPC Lead decides and the operator executes assigned scope. Produce the staged restoration set and defend each drift restoration against the recorded floor.
- Draft one code-red diagnosis on a historical declaration, building all four maps from the archived packet and reconciling them against the recorded overspend. Compare the reconciliation to the original: a different residual is a finding to explain rather than an error to hide.
- Draft a cut sequence in loss-dollar order and decompose every cut into L2-compliant steps. Matrix row: "Bulk deployments (#38)", where the PPC Lead reviews and the operator executes. The checkpoint holder rejects any single-step cut exceeding the per-change limit, which is the fastest way to learn that the limit shapes the schedule.
- Execute one quarterly re-ladder cohort with the checkpoint holder confirming each rung and its dated prediction.
- Draft the monthly harvest read including the verification split, and state the withheld portion with its reason. Matrix row: "Push funding (the M5 queue)", where the PPC Lead decides within envelope and the CEO decides above it; the operator executes.
- Draft one incursion escalation to #16 and one incrementality hand-off to #35. Matrix rows: "PAT/conquest opens and exits" and "Incrementality tests", where the PPC Lead reviews and the operator decides and executes.

### 11.3 Week 4: Unsupervised Execution With Audit

The operator executes and deploys; the checkpoint holder audits after the fact against the artifact chain rather than before the fact against the draft. The human-mandatory set does not move: the declaration confirm, the close, pauses, rule suspensions, structure changes, survivor selection, the STRUCTURAL escalation and every deployment confirm remain confirmed per the matrix regardless of ramp stage.

- Own the weekly floor-state cycle on assigned scope end to end, including the T6 predictions and the V5 scoring of the prior cycle's due rows.
- Own the breach-hold procedure including the dated re-read that arms the second-week trigger. Audited specifically on whether the T6 entry exists, because that entry is the only memory the two-limb trip test has.
- Own the monthly harvest read and its posting, with the checkpoint holder auditing the verification evidence rather than the arithmetic.
- Declare a code red if one trips in scope, on the R1 test, with the confirm from the checkpoint holder. On a launch-stage product the matrix standing rule applies: the launch operator declares and the PPC Lead closes.
- Audit finding: any posture read whose classification the checkpoint holder would have made differently returns the operator to week-2 supervision on that procedure only, not on the whole scope.

### 11.4 Live-Defense Questions

1. A mature product reads 1.7x its band ceiling this week and read 1.1x last week. Is it a code red, and what runs? Answer to the rule, then say what the Account Rollup A0 cell says and why the operator does not act on it.
2. Why is the trip tested on TACOS and not on ACOS? Give a case where the two point in opposite directions, and say which decision each number governs.
3. A defended term's bid sits 30% above its recorded floor and the position is holding. What is the verdict and why is "it is working" not an argument?
4. You need to cut a term from \$3.10 to its \$0.84 ceiling. Write the step program and say why it takes two weeks. Then say what that implies for what the day-7 re-read is allowed to conclude.
5. The quarterly probe slips and the T5 log shows two competitor deal windows inside the read. What do you do with the rung, what do you record in the register, and what do you do with the incrementality read that was riding on it?
6. The harvest read shows \$416 of reduced floor-state spend with rank flat and share down a class. How much posts to the queue, and what happens to the rest? Say what breaks downstream if you post the whole figure.
7. A code red reads inside band at day 14. Why is that not a close, and what is the cost of closing there?
8. Defense spend on a syntax has doubled while branded revenue protected is flat. Name the failure mode, the procedure that detects it, and the one thing v4.0 does not tell you that you will have to escalate rather than decide.
9. An event window arms on a product in Recovery Mode. Which mode governs, what posture does the event run at, and where is that ordering written?
10. Name three things this SOP is forbidden to decide and say which component decides each.

## 12. Interface Contracts

Sixteen handoffs, split upstream and downstream. Every row states what happens when the handoff does not arrive, because that is the clause the corpus has always been missing. Sources: the Cross-Reference Ledger row for component 26 and the rows of every counterpart named, cross-checked against Master PPC Data Architecture v1.0 for artifact names and tab locations.

| ID  | Counterpart               | Artifact and timing                                                                                                       | Non-arrival consequence                                                                                                                                                                                                                                                                                                                                                           |
|-----|---------------------------|---------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| I-1 | #25, upstream             | Stage-entry handoff: the floor register with dated ladder state in T4 and the closing T6 chain, at stage transition       | Stage entry is refused back to #25 and the product stays in transition. Accepting entry without recorded floors makes P2 drift detection impossible (there is nothing to drift from), makes P3 unable to build a cohort, and leaves P12 with no baseline. This is the highest-consequence non-arrival on the upstream side because every other procedure here reads the register. |
| I-2 | Data Operations, upstream | The weekly packet: T0 zone B spend and sales from the SellerBoard read, landed Monday 06:00, workbooks refreshed by 09:00 | The week is not validated and no declaration issues on it. R1 uses "validated" in both limbs, so a missing or estimated sales figure suspends the trip test rather than degrading it. The breach-hold record still writes with the gap named, so the two-week clock is not silently reset.                                                                                        |
| I-3 | #41 economics, upstream   | The margin table refresh into the T0 economics spine named cells, monthly and on event                                    | Every ceiling in the cut sequence carries the G9 PROVISIONAL tag and every ceiling-referencing verdict states it. A cut computed against a stale margin cuts to the wrong number, and the direction of the error is unknown, so the cut deploys with the tag rather than waiting.                                                                                                 |
| I-4 | #14 P4, upstream          | Branded impression share and incursion terms from the brand-broad weekly read, each Monday cascade                        | P11 step 1 loses one of its two detection artifacts and runs on the T2 K3 impression-share reading alone. Competitor brands surfacing in served terms become invisible, which is the earlier of the two signals, so incursions are detected later and the response ladder starts from a worse position.                                                                           |
| I-5 | #28, upstream             | Ladder step derivation, exit-bid discovery and the ceiling worksheet result, on request per cohort and per bleed response | P3 probes and P8 placement cuts do not stage. This SOP owns the decision to re-ladder and the decision to cut a bleeder; it does not own the mechanics, and staging a rung without the derived step is inventing a threshold. The cohort's re-test date moves out and the delay is recorded as an F2 exposure rather than absorbed.                                               |
| I-6 | #37, upstream             | The rule change log covering the diagnosis window, on request at declaration; no delivery window is governed (AM-12)      | P7 Map 4 cannot be built and is recorded as unreadable in the declaration record. A runaway rule then survives the cut sequence and continues stepping bids against stale targets while the protocol runs, which is the WE-1 incident repeating inside the response to itself. The map is re-run when the log arrives, even mid-protocol.                                         |
| I-7 | #32, upstream             | The armed-window declaration and the event plan, at arming                                                                | G10 exclusion cannot be applied, so P3 probes may run inside an event window and discover floors that dissolve when it closes. Where the window is known but the plan is absent, P3 suspends for the affected cohorts rather than proceeding on an assumed window length.                                                                                                         |
| I-8 | #33, upstream             | The family ownership declaration on shared defended terms, current at declaration and before any restore step             | The G11 check fails and no restore step deploys on the shared term. Two siblings restoring the same term bid against each other with the family's own money, which is the divergence pattern the portfolio layer exists to catch. The term holds at its current bid and the escalation goes to #33.                                                                               |
| I-9 | #35, upstream             | The continue-or-restore call on an incrementality ladder, at the end of the two-week capture window                       | The rung holds where it is and the floor is not recorded, because a floor recorded without the capture read is a bid with a label. Where the window was confounded, this SOP hands the void and the confound references rather than a reading, and #35 re-runs; a confounded read passed over as clean is the worse failure.                                                      |

| ID   | Counterpart                           | Artifact and timing                                                                                                                                          | Non-arrival consequence                                                                                                                                                                                                                                                                                                                                |
|------|---------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| I-10 | #38, upstream                         | The T+1 verification diff on every staged set, the day after deployment                                                                                      | Deployed values are unproven and the week's readings on the affected rows are void. Inside a code red this is acute: the day-7 scoreboard compares expected savings against realized spend, and without the diff a shortfall cannot be separated into "the cut did not land" and "the cut landed and was not enough", which are opposite responses.    |
| I-11 | #16, downstream then upstream         | Escalation of a persistent incursor to #16 P10, same week as detection; the counter-conquest evaluation record returns                                       | Outbound: the incursion sits at the defense ceiling with no evaluation, and the standing answer stays defensive, which is the safe failure. Inbound: the incursion record does not close until the evaluation returns, either way. An escalation with no returned record is carried as open on the defended-set register rather than assumed declined. |
| I-12 | #17-19, downstream                    | Defense-set membership declaration and the S-F8 rebuild instruction naming terms and the defense budget line, at declaration and on a MARGINAL share reading | The presence rebuild does not happen and share continues falling below the PRESENT class while the harvest read keeps withholding the release. The two failures compound: the money is withheld correctly and is also not spent on the thing it was withheld for, so the position degrades while the budget sits.                                      |
| I-13 | #12, downstream                       | Structure repairs from P7 Map 3: consolidations to P7, add-exact standups to P1 or P2, demotions to P6, each with the code-red reference                     | The duplicate keeps splitting the term's data and double-paying the auction, and the uncovered head term keeps being carried by auto at an uncontrolled CPC. Both survive the cut sequence, so the diagnosis reconciles and the spend does not fall by the predicted amount, which surfaces at the day-7 scoreboard as an unexplained gap.             |
| I-14 | #29, downstream                       | Negation evidence lines from the waste map, staged in the same session as the cut sequence                                                                   | The negation sweep does not land, so zero-sale and off-taxonomy terms keep buying clicks at the reduced bid rather than at no bid. The bid-down verdict alone converts a large loss into a smaller loss, and the sweep is what converts it into no loss; a cut sequence without it under-delivers against its logged savings by the sweep's share.     |
| I-15 | Instrument layer, downstream          | The verified release posting to T0 zone H and Account Rollup tab A2, monthly at the harvest read                                                             | The M5 queue under-states available funding and queued pushes stay queued while the money exists. This is the quiet direction of the failure: an unposted release costs nothing visible and defers every push behind it, which is why P12 runs on a date rather than on prompting.                                                                     |
| I-16 | V5 and the amendment path, downstream | T6 rows with dated predictions on every deployed change; the P14 review with its amendment proposals within one week of a Recovery exit                      | Unscored actions do not accumulate: a cut program that is never scored cannot size the next one, so the following declaration re-derives its cuts from scratch. An unfilled P14 leaves the entry cause unexamined, and rule R6 states the consequence plainly: recovery that teaches nothing repeats.                                                  |

**Companions:** Master PPC Decision Criteria System v4.0 (family F2 stage bands and deviation classes, Section 8.1 S-A1 to S-A10, Section 8.6 S-D1 to S-D6, Section 8.10 S-F8, Section 9 levers L1 to L7, Section 10 M3 and M8, Section 11 T8 and R1 to R6, Section 12 P1 to P8, Section 13.1 closed vocabulary) · Keyword Decision Record Template v3.0 (the 25-verdict closed vocabulary and the campaign subset) · PPC Strategic Doctrine v1.0 Sections 6, 9, 12 and 14 (the why) · Master PPC Data Architecture v1.0 Sections 4, 6, 8, 9 and 11 (artifacts and tabs) · PPC Authority Matrix v2.0 Section 1 (Elements 9 and 11) · PPC Economics Worked-Example Set v1.0 cases E5 and E8 · SOP-25 (stage entry), SOP-28 (ladder mechanics and floors), SOP-12 (campaign execution), SOP-14 P4 (brand broad), SOP-16 (defense PAT and counter-conquest), SOP-17-19 (format rebuilds), SOP-29 (negation), SOP-30 (harvest), SOP-35 (incrementality), SOP-37 (rules), SOP-38 (staging and verification), SOP-32 (event windows), SOP-33 (family ownership) · PPC Cross-Reference Ledger (component 26 row and the interface list).