SOP-27: Terminal and LTSF-Tier PPC Rules

Version 2.0 · 2026-08-14 · Owner: PPC Lead with the LTSF Owner

Supersedes SOP-27 v1.0 (2026-07-24) in full · Closes Cross-Reference Ledger conflict C6 on the attachment of the LTSF Master Decision and Operations System v2.1 · Every threshold, band, scenario and verdict string lives in the Master PPC Decision Criteria System v4.0 and the Keyword Decision Record Template v3.0; this SOP carries zero local thresholds and executes what those two documents decide.

1. Verdict and Scope

> VERDICT: terminal-tier PPC executes a declaration it does not make. The LTSF Master Decision and Operations System v2.1 decides the SKU archetype, the risk tier, tier movement, the floor price, the salvage comparison and every terminal option. This SOP owns what PPC does on receipt of that declaration and owns nothing above it. On the day of the tag the objective function switches from rank and share to net recovery per unit moved against the declared alternative, and it switches back only when the LTSF system retracts the tag. Version 2.0 supersedes v1.0 in full. v1.0 stood at 2 pages and 7 rules because the deciding system sat outside this corpus, so tier criteria could not be read and were correctly not invented. That system is now attached, conflict C6 is closed, and rules R1 to R7 carry forward unchanged as the rule layer beneath 14 procedures, 25 verdict mappings, 4 worked examples, 12 failure modes, a 10-row operator ramp and 18 interface contracts.

1.1 Why this document is no longer 2 pages

SOP Build Kit v1.0 Section 4 "The Depth Standard V2.0" names this component as the honest example of a component that cannot fill 10 pages, at 2 pages by design. That statement was correct while C6 was open, because a document that stops at a boundary has only the boundary to describe. It is void at v2.0. With the LTSF Master v2.1 in the corpus the component acquires 4 archetype execution classes, 4 tier postures, 11 LTSF levers that PPC executes across 3 lever families, a code-red precedence seam, and a set of interfaces that were previously unwritable because the counterparty document could not be read. The depth is not padding; it is the material C6 was withholding.

1.2 The seam: decided above, executed inside

Every row on the left is read from the LTSF declaration and is never derived, recomputed, negotiated or second-guessed in PPC. Every row on the right is this SOP's and appears in no other document.

Decided by the LTSF Master v2.1 (read, never derived)	Executed by this SOP (owned here, nowhere else)
SKU archetype A / B / C / D, assigned by the owning Brand Manager and confirmed by the LTSF Owner per LTSF Section 5 "SKU Classification: The Four Archetypes"	The campaign-level execution class each archetype selects: which objective the product's terms carry, which levers open, which shut
Risk tier GREEN / YELLOW / RED / CRITICAL per LTSF Section 7.1 "Tier definitions and cadence"	The PPC posture each tier selects, and the posture change staged on the same day the tier changes
Tier movement per LTSF Section 7.3 "Tier movement rules"	Nothing. A tier move is an input event; PPC executes the new posture and proposes no move
The floor price per LTSF Section 4.2 "The floor price (clearance authority)"	The R2 allowable ad cost and maximum liquidation CPC computed from the declared floor, salvage and unit fields
The salvage comparison (removal net, liquidation net, disposal net) per LTSF Section 6.5 "Family 5: Terminal levers"	Nothing. PPC reads the winning alternative's per-unit value as the R2 subtrahend and prices no terminal option
Every terminal option and its execution, plus PO freezes and removal logistics	The lane-closure finding that routes back when arithmetic says PPC cannot clear the units
Discount depth, coupon depth, deal depth and the maximum combined stack	The bid and budget consequences of a depth change, and the CVR re-read that moves the R2 ceiling with it

1.3 The four archetypes as PPC execution classes

Archetype is a diagnosis of why the units are aged, and the diagnosis selects which PPC levers can work at all. Running the wrong class is the most expensive execution error available here, because every wrong week costs the full bracket rate on every unit.

Archetype	PPC execution class	LTSF Section 6.2 Family 2 levers that open	What PPC does not do
A — Fixable Demand	Rebuild traffic on a listing being repaired in parallel. P7. The paid side runs while Design and Brand Management shift the curve; PPC does not wait for the repair and does not substitute for it.	Competitor keyword replication; new discovery campaigns; break-even PPC scaling once the CVR read recovers	Does not lead with depth. A price lever on a broken listing buys the same rejection at a worse price, which is LTSF Principle 4
B — Variation Overstock	A child-ASIN targeting problem, executed on the D2 variation map and the re-point machinery this library already owns. P8.	Refined targeting on non-preferred children; new discovery campaigns scoped to those children	Does not touch the parent or the healthy children. No family-wide budget move, no shared coupon, no parent re-point
C — Structurally Dead	Clear what the ceiling buys and stop. P9. The 7-day decide window has a hard PPC consequence stated at 1.3.1 below.	Break-even PPC scaling on proven converters only, at the R2 ceiling, inside the dated liquidation campaign	Runs no experiments of any kind: no discovery campaign, no auto probe, no competitor replication build, no new keyword test
D — Aged Healthy	Speed, not exit. P10. The product sells and the margin holds; the batch is racing a bracket date.	Budget and bid increases on proven syntaxes; break-even PPC scaling with its spend ceiling at contribution above the floor price	Does not retag to liquidation and does not quarantine, because the units are not terminal. Rank objectives stay legal on the product's untagged terms

1.3.1 Archetype C and the 7-day window. LTSF Section 5 states of Archetype C: do not run further experiments, and decide within 7 days of classification. The PPC consequence is arithmetic, not attitude. Every discovery lever in LTSF Section 6.2 "Family 2: PPC velocity levers" carries a 7-14 day impact window, and LTSF Section 9.1 "Pre-registered 7-day tests" gives every lever a 7-day read before it can be scaled. A read window that equals or exceeds the decision window cannot return in time to change the decision, so the experiment is a pure charge with no decision value. On a 480-unit Archetype C batch in the 366-455 day bracket at the LTSF Section 3 rate of $6.90 per unit per month, one 7-day read window costs 480 x 6.90 / 30 x 7 = $772.80 and can change nothing. P9 therefore opens no test campaign on an Archetype C tag, and F7 is the failure mode when one appears.

1.4 The four risk tiers as PPC postures

The tier is declared, never read from PPC data. The posture below is what this SOP stages on the day the declaration lands or changes.

Tier	PPC posture	Structural state	Cadence this SOP runs
GREEN	Standing operations continue unchanged. The only PPC action is the proactive velocity read requested when a large batch approaches the pre-LTSF window.	No retag, no quarantine. Units remain in standing campaigns.	The weekly pass reads the flag and does nothing else. P1 step 6 records the read.
YELLOW	Velocity work live on the aged children or the aged batch, at the archetype's lever set. Push spend on the affected variation is frozen at its current level by v4.0 Section 5 gate G6 where inventory reads YELLOW for a stockout reason, and by this SOP's P13 where the freeze is LTSF-driven.	No retag to liquidation. Dedicated ad groups where the aged units sit on specific children.	Weekly, inside the Monday pass. P12 check line written whether or not the tier moved.
RED	Execution, not planning. The archetype's full lever set opens at the R2 ceiling. Archetype C and any terminal tag retag to liquidation under R1 and quarantine under R5.	Dated liquidation campaigns carrying the cliff date in the name for terminal-tagged units; dedicated ad groups for aged-batch units not yet tagged terminal.	Weekly minimum, with the daily pace watch armed inside the final 14 days before the declared cliff date.
CRITICAL	Emergency. R4 shuts everything that is not a proven converter at the R2 ceiling. The lane either clears at the ceiling or the finding routes back inside the same working session.	Quarantine complete. No terminal-tagged unit serves inside a standing campaign.	Twice weekly per LTSF Section 7.1, plus the daily Brand Management and PPC huddle line per LTSF Section 10.1. Daily pace watch is standing, not conditional.

1.5 Boundaries with adjacent components

Every decision excluded from this SOP names the component that owns it. The Cross-Reference Ledger row for component 27 is the check.

Adjacent component	The seam, stated so it is owned
LTSF Master Decision and Operations System v2.1 (outside the PPC component register)	Decides archetype, tier, tier movement, floor price, salvage comparison, discount depth and every terminal option. This SOP reads the declaration and executes. A tier criterion, a floor price or an archetype call written into this document is a defect.
Criteria System v4.0 (#8)	Decides every threshold, band, gate, scenario, lever limit and verdict string in the entire library. This SOP executes; it never decides. A threshold written here instead of v4.0 is a defect.
#12 Exact campaign operations	Executes campaign-level creation, bid, budget, status and structure changes, including the P10 re-point this SOP's Archetype B work depends on. This SOP issues the liquidation campaign specification; #12 stands the campaigns up.
#32 Deal and event window operations	Owns the window plan, its caps, its arming and its reconciliation. Its P12 computes a liquidation window under this SOP's R2, R3, R4 and R6, and contributes only the window plan around them.
#26 Mature, defend and code red	Owns the code-red declaration, the same-day freeze, the four-map diagnosis, the loss-ordered cut sequence and the exit test. The precedence seam where an LTSF RED or CRITICAL coincides with a code-red declaration is stated at 1.5.1 and worked at WE-4.
#25 Transition and maintenance	Owns the scheduled descent from push economics to sustainable economics and hands its exit into #26. Referenced here by Ledger row 25 only; a terminal tag is not a transition and does not enter its descent schedule. Not attached to this session.
#28 Placement and modifier operations	Owns ladder mechanics, exit-bid discovery and modifier sizing. R4 normalizes modifiers to zero on terminal-tagged entities, which is a state, not a ladder; no step is sized in this SOP.
#29 Negation architecture	Owns negation doctrine and wall levels. R5's quarantine is built with #29's walls in the same working session as the retag.
#33 Sibling and family arbitration	Owns shared-term ownership declaration. A terminal tag on one sibling is a family event: this SOP routes the finding and does not reassign the term.
#37 Automation platform rules	Excludes terminal-tagged entities from every automation rule and computes nothing about tiers. This SOP supplies the protected-scope state feed that tells it what to exclude.
#38 Bulk-file deployment mechanics	Owns staging format, upload and the T+1 verification pass. Every change in this SOP deploys through it.
#35 Incrementality testing	Owns test design and the continue-or-restore call. No incrementality test opens on a terminal-tagged entity: R4 removes the defense allocation the test would measure.
G6 inventory bands and the D2 re-point chain	Own adequacy bands and the ranked-backup chain. G6 measures stockout risk and reads GREEN on an overstocked variation; that is a correct reading and not an authorization to scale. LTSF risk is a separate declaration and this SOP is where the two are held apart.

1.5.1 The precedence seam with #26. v4.0 Section 11 "Layer 6: Transitions, Stages, Modes (T1-T8) and Recovery Mode (R1-R6)" rule T8 orders Recovery Mode above Event Mode above stage defaults. It does not name the terminal lane, so no precedence rule exists for the case where an LTSF RED or CRITICAL declaration coincides with a code-red declaration on the same product. This SOP states the operating consequence without inventing one. R5's structural quarantine is the mechanism: terminal-tagged units run in dated liquidation campaigns, which are not standing campaigns, so #26's cut sequence, which is ordered by loss dollars across standing campaigns, does not reach them. Cutting a liquidation campaign to reduce a code-red overspend raises the charge and forfeits the contribution, which is worked in dollars at WE-4. What does bind is v4.0 Section 11 rule R3: while Recovery Mode is active, every bid move above 15% and every budget move above $25 per day is human-confirmed, and that tightening applies to the liquidation lane exactly as it applies to everything else. The absence of a T8 clause naming the terminal lane is amendment item AM-2.

1.6 Rule identifiers, disambiguated

This SOP's own rules are R1 to R7 and are cited across the corpus without a document prefix; SOP-32 v2.0 P12 cites R2, R3, R4, R5 and R6 that way. v4.0 Section 11 uses R1 to R6 for Recovery Mode. The collision is real and both sets appear inside WE-4. Throughout this document, a bare R-number is this SOP's rule and a Recovery Mode rule is written with its document prefix as v4.0 R1 to v4.0 R6. The collision is amendment item AM-6.

1.7 Where every number in this document comes from

Three sources, and no fourth. First, v4.0, quoted with its section heading, which is the only threshold authority. Second, the LTSF declaration, whose fields are inputs to this SOP the way a CPC read is an input: the rate card, the tier definitions, the SLAs and the floor price are facts about the deciding system, not PPC thresholds, and they are cited to their LTSF section. Third, the worked examples, which are illustrative and say so. Any number that would need a fourth source is an amendment item in Element 10.5 and is not written into the procedure.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. v4.0 carries no interface-lag threshold of any kind, so every lag below is a cadence fact quoted from LTSF Section 9.2 "Execution SLAs", LTSF Section 7.1, or this SOP's own R5, and is marked with its source. Where neither document supplies a lag, the cell says so and the gap is amendment item AM-3.

Trigger	Detecting artifact	Maximum lag to procedure start	Procedure
New terminal or at-risk tag on a variation	The LTSF row in the LTSF Master Tracker, carrying archetype, tier, floor price, salvage value, remaining units and cliff date; mirrored into T2 Keyword Command block K2 "Targeted variation(s)" as the terminal flag	Same day. LTSF Section 9.2 sets 24 hours from decision to live for PPC campaign changes; R5 requires the quarantine inside the tag week	P1, then P2, P3, P4, P5, P6 in sequence
Tier change on an already-tagged SKU	Weekly auto-classification per LTSF Section 10.2 "The weekly cycle" step 3, landing in the LTSF row	Same day for a move into RED or CRITICAL; next Monday pass for a move down	P13, then the archetype procedure the new tier opens
Archetype change or first assignment	The archetype column added by LTSF Section 11.1 amendment 7, with date assigned and owner	Same day. An archetype change invalidates the live lever set	P1 step 4, then P7, P8, P9 or P10
Cliff-calendar update, including a bracket crossing	The cliff calendar columns added by LTSF Section 11.1 amendment 6: units per age bracket, the date the largest bracket crosses its next cliff, and the weekly velocity-to-cliff requirement	Same working session. R3's pace math is a function of the cliff date and is invalid the moment the date moves	P4 in full, then P12
Floor price, salvage value or depth change	The floor price column and salvage comparison columns per LTSF Section 11.1 amendments 2 and 3	Same working session. The R2 ceiling is a function of both	P3 recompute, then P4 re-run
Weekly expiry check	T6 Decision Log rows whose review date lands this cycle, plus the LTSF outcome log per LTSF Section 9.4	Monday, inside the standing pass, before any new verdict is read	P12
Daily pace watch inside the final 14 days before the cliff	Realized units cleared against the R3 required pace, read from T3 Campaign Board group C4 order counts on the dated liquidation campaigns	Daily, one read per working day. CRITICAL tier arms this unconditionally per LTSF Section 7.1	P12 step 5
Deal or event window touching a tagged variation	The armed-window declaration from #32 entering T2 with the event flag under v4.0 Section 5 gate G12	Before the window opens	P11 step 4, with #32 P12 owning the window plan
Code-red declaration on a product carrying a terminal tag	T0 Product Command zone A verdict banner reading CODE-RED, and the #26 declaration record	Same day as the declaration	P13 step 5 and the 1.5.1 seam
Lane cannot buy the required clicks at the R2 ceiling	The P4 traffic check output: required clicks per day against market CPC on the proven converter set	Same working session as the P4 run. LTSF Section 7.4 escalates any decision blocked more than 5 business days	P14
Tag retracted, or units cleared to zero	Remaining units reading zero in the LTSF row, or the tag column clearing	Next Monday pass	P14 step 4, closure record
Day-7 read on any deployed lever returns below half its logged prediction	The prediction and outcome log per LTSF Section 11.2, read against the T6 prediction	The day the read is clear. LTSF Section 9.1 replaces the lever rather than extending it	P12 step 4

3. Preconditions and Gates

Every procedure runs only after these read current. A missing or stale field group suppresses its dependent execution class and is recorded in T7 Manifest; the pass continues on the rows it can read and writes what it could not see, per v4.0 Section 5 "Layer 0:

Gates (G1-G12). Nothing Proceeds Until These Pass" gate G8.

Precondition	Manifest ID (T7 field group)	The observable that proves it passed	What is suppressed while it fails
The declaration is complete	Not a PPC manifest group. The LTSF row is the artifact; its presence is recorded in T7 as an external feed row	Seven fields present and dated: archetype, tier, floor price, salvage value per unit, remaining units, cliff date, and the winning terminal option. R2 and R3 are unreadable without all seven	Everything. P1 stops and returns the row to the LTSF Owner as incomplete. No partial execution, no estimated field
Converter evidence is sufficient	A4 sample sufficiency buckets	The proven converter set carries clicks at or above the v4.0 Section 5 gate G2 thresholds on the window the declaration covers; below them the verdict is CONTINUE TEST TO 50 CLICKS AT CAPPED BID and not a CVR read	The R2 maximum liquidation CPC, because it is a function of CVR. P3 holds and P4 cannot run
No quarantined rows in the affected set	Zero-flag columns feeding v4.0 gate G1	No impossible-metric flag on any row in the liquidation campaign set. A CVR reading zero against present orders quarantines the row	Every verdict on the quarantined row. The verdict is QUARANTINE: DATA REFRESH and a repair task opens the same day
Campaigns are not truncated	A3 pacing and in-budget	In-budget percent at or above the v4.0 gate G3 line on every dated liquidation campaign. A starved liquidation campaign misreads its own pace and its CPC	Every bid conclusion on the campaign. The budget lever executes first per v4.0 Section 9 rule L1
Advertised SKUs are eligible	A7 eligibility	Eligibility live on every advertised child in the liquidation campaign. An ineligible ad on terminal stock accrues bracket rate while buying nothing	All execution on that target, same day, per v4.0 gate G4
The variation map resolves	D2 variation map	Every term in the affected set maps to the specific aged child, not to the parent and not to the preferred child. Archetype B is unexecutable without this	P8 in full. The row carries the D2-pending flag and P8 does not start
Economics chain is fresh	D1 per-variation margin, B8 economics chain	T0 zone C named cells live with the margin table version dated inside the 45-day staleness clock; the v4.0 gate G9 state reads fresh rather than PROVISIONAL	Nothing is suppressed, but every ceiling-referencing action stages with the PROVISIONAL tag riding on it and is not executed clean
Price is stable on the affected children	D1 plus T5 Product Log price-move rows	No price change on the advertised child inside the trailing 7 days, per v4.0 gate G9. A depth change is a price change	Any CVR-derived conclusion, which means the R2 ceiling holds at its last value until a clean read returns. P3 step 6 states the hold
Event windows are known	Event-mode column in T2 block K2, fed by #32	The armed-window declaration present with its dates before the window opens, per v4.0 gate G12	Structure changes inside the window. #32 P12 governs the window; this SOP contributes R2 and R3 to it
Sibling conflicts are resolved	Family conflict state in T2 block K2	No sibling carrying an active push objective on a term the liquidation lane is buying, per v4.0 gate G11	Scale on the shared term until the portfolio call assigns a lead. The finding routes to #33
The keyword list is clean	MKL state feeding v4.0 gate G12 MKL hygiene	Every term in the liquidation campaign present in the current keyword list with a syntax tag and a relevancy score	Analysis of untagged terms, which route to maintenance rather than into the lane
Automation is excluded	D6 automation platform rules export, status OPEN per the Ledger	The exclusion list at #37 names every dated liquidation campaign. While D6 is OPEN this is confirmed by manual read, and the read is recorded	Nothing formally, but an unexcluded liquidation campaign is failure mode F9 and the manual read is the only control until D6 lands

4. Step-Level Procedures

Fourteen procedures. Time estimates are planning figures for one SKU at one tier and are not thresholds. Every staged change deploys through #38 and carries a reason string with its verdict and scenario ID.

P1. Read and validate the declaration (12 minutes per SKU)

1. Open the LTSF row for the SKU in the LTSF Master Tracker. Read the seven fields the declaration must carry, in this order: archetype (A / B / C / D with date assigned and owner, per LTSF Section 11.1 amendment 7), tier, floor price, salvage value per unit, remaining units by age bracket, the cliff date of the largest bracket, and the winning terminal option from the salvage comparison.

2. Any field blank or undated stops the procedure. Return the row to the LTSF Owner naming the missing field and the execution class it blocks. PPC computes nothing it can pull from that system and estimates nothing it cannot.
3. Copy the seven fields into T2 Keyword Command block K2 against the affected variation, and into the T5 Product Log as one dated event of type targeting change, with the LTSF row reference in the Link column. The Product Log entry is what makes every later CVR and CPC swing attributable, per v4.0 Section 7 "Layer 2: Diagnosis (D1-D10). Locate the Failure Before Touching a Lever" rule D5.
4. Route on archetype. The archetype selects the execution procedure and nothing else in this SOP overrides that selection.
5. Read the tier and set the posture per Section 1.4. Tier and archetype are independent: a RED Archetype A runs P7 at RED intensity, not P9.
6. Where the tier reads GREEN, record the read in T6 Decision Log with no staged change and stop. A GREEN tier is a real outcome of this procedure and consumes 3 of the 12 minutes.

```
P1 step 4 routing
|
+-- Archetype A (Fixable Demand) ......... P7 listing repair runs in parallel at Family 1; PPC rebuilds traffic
+-- Archetype B (Variation Overstock) .... P8 child-ASIN targeting on the D2 map; parent untouched
+-- Archetype C (Structurally Dead) ...... P9 proven converters only at the R2 ceiling; zero experiments
+-- Archetype D (Aged Healthy) ........... P10 break-even scaling on proven syntaxes; no retag, no quarantine
|
+-- Archetype blank or disputed .......... STOP. Return to the LTSF Owner. Section 5 of the LTSF Master
assigns it; this SOP does not classify and does not guess.
```

P2. The objective switch and the retag (R1) (15 minutes)

1. Trigger: a terminal tag, which is the Archetype C class and any unit the LTSF row marks terminal regardless of archetype. An at-risk tag without a terminal designation does not retag; it changes posture only, at P13.
2. Retag the affected terms to the liquidation objective in T2 block K7 "Objective". The objective field is mandatory under v4.0 Section 6 "Layer 1: Objective Assignment (O1-O8)" rule O1, and an untagged row is flagged UNMANAGED and cannot be evaluated at all.
3. Write the grading metric change into T2 block K7 "Situation text": net recovery per unit moved, which is salvage price net of fees minus ad cost per unit, compared against the declared alternative's per-unit value. Rank, share and defense metrics stop applying to these units the day of the tag.
4. Set the review date in block K7 to the cliff date, not to a standing cycle date. Every terminal PPC decision carries its expiry from the cliff calendar under R6, and LTSF Principle 5 states that no decision is open-ended.
5. Record the retag in T6 with its prediction: units cleared by the cliff date, the R3 required pace, and the review date. A retag without a logged prediction is incomplete and does not deploy, per v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)" rule V1.
6. The retag is human-mandatory. The Authority Matrix v2.0 Section 1 "The Matrix" row "Terminal/LTSF PPC lane" places D with the PPC Lead together with the LTSF Owner, and the incoming operator holds E.

P3. Compute the R2 allowable ad cost and the maximum liquidation CPC (10 minutes)

1. Read three declared values: net recovery per unit at the salvage price, the winning alternative's per-unit value, and the converter CVR for the proven set. The first two are declaration fields. The third is a PPC read from T2 block K5, subject to the v4.0 Section 5 gate G2 sufficiency line.
2. Allowable ad cost per unit = net recovery at the salvage price minus the alternative's per-unit value. Because the alternative is usually negative, the subtraction usually raises the ceiling: a charge avoided is value gained, and an ad that loses against price can still beat the alternative.
3. Maximum liquidation CPC = allowable ad cost per unit times CVR.
4. On Archetype D, where the units are not terminal and the product still sells at margin, the ceiling instead comes from the break-even PPC scaling lever in LTSF Section 6.2 "Family 2: PPC velocity levers", whose spend ceiling per unit is the contribution above the floor price. The declaration carries contribution at the current price and contribution at the floor price so PPC subtracts rather than models fees. The resulting CPC ceiling is the same object as v4.0 Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)" family F1's max profitable CPC, which is margin dollars times CVR; break-even scaling is the decision to operate at that ceiling rather than at the target line.
5. Write the ceiling and its date into T2 block K6 beside the standing max profitable CPC. Both values live in the row and are read apart. The standing value governs untagged terms on the same product.
6. Where gate G9 marks the reads PROVISIONAL because a price or depth change landed inside the trailing 7 days, the ceiling holds at its last dated value and every action staged against it carries the PROVISIONAL tag. It is not recomputed on a provisional CVR.

7. Sunk COGS appears in no column of this computation, per R7 and LTSF Principle 2. A bid computation citing acquisition cost is failure mode F2.

P4. Cliff pacing and the traffic check (R3) (20 minutes)

1. Required units per day = remaining units divided by days to the declared cliff date. Both are declaration fields; neither is a PPC estimate.
2. Required clicks per day = required units per day divided by the converter CVR.
3. Read market CPC on the proven converter set from T4 Placement & Bid, column "Expected CPC", and from the current suggested-bid values in the BULK-US Raw Pull Pack. Read it at the volume the lane needs, not at the volume it currently buys: a CPC that clears 40 clicks a day is not the CPC that clears 300.
4. Compare market CPC at the required volume against the P3 ceiling and route on the result.
5. Where the lane opens, size the daily budget as required clicks per day times market CPC, and stage it on the dated liquidation campaign. The budget is a derived figure and is re-derived every time the cliff date, the unit count, the floor price or the CVR moves.
6. Where the lane splits, state in T6 how many units PPC will clear and how many the terminal option takes, and route the second number to the LTSF Owner the same working session. A split is a normal outcome and is not an escalation.
7. Where the lane closes, run P14. Closing a lane on arithmetic is a correct outcome and not a failure, and it is the finding the LTSF system needs in order to move the price or execute the terminal option.

```
P4 step 4 traffic check
|
+-- Market CPC at required volume <= R2 ceiling
| -> LANE OPEN. Stage the full required-clicks budget. Arm P12 pace watch.
|
+-- Market CPC above ceiling, but a residual click volume clears at or under the ceiling
| -> LANE SPLIT. Buy what the ceiling buys; state the residual unit count;
| route the remainder to the LTSF Owner for the terminal option or a price move.
|
+-- No click volume clears at or under the ceiling
-> LANE CLOSED. P14. The finding routes back the same working session:
either the salvage price moves down, which raises CVR and re-runs this
procedure, or the terminal option executes on the whole pool.
```

P5. Structural quarantine (R5) (35 minutes for the first campaign on a SKU, 10 minutes per later variation)

1. Trigger: any retag under P2. The quarantine executes in the same week as the tag, without exception.
2. Specify the dated liquidation campaign and hand the specification to #12 for creation: one campaign per tagged variation, the cliff date carried in the campaign name so the expiry is readable from the campaign list, advertised SKU set to the specific aged child from the D2 map, and one ad group holding the proven converter set only.
3. Remove the tagged variation from the standing family structure the same week: pause the terms where the standing campaign advertises only that child, or negative-exact them where the standing campaign advertises several. #29 owns the wall level; this SOP supplies the term list and the reason.
4. Record the wall list in T3 Campaign Board group C5 "Negation state per mode" as the wall-audit baseline, exactly as any other structure event.
5. Set placement modifiers to zero on all three groups at creation. Modifiers are earned from placement evidence at #28, and R4 removes the evidence base on tagged entities.
6. Add the campaign to the #37 exclusion list so no automation rule touches it, and record the addition. While Ledger dependency D6 reads OPEN this is a manual read and the read is logged.
7. Mixed reads poison both sides: the standing product's economics and the liquidation's recovery math. A tagged variation still serving inside a standing campaign past the tag week is failure mode F3.

P6. Shut-off execution (R4) (15 minutes)

1. On terminal-tagged variations, stage the following as one set: rank objectives removed, funded pushes stopped, discovery spend closed, brand-defense allocation removed, and placement modifiers normalized to zero.
2. What remains: exact and auto close match on proven converters only, at the R2 ceiling. Nothing else serves.
3. Freed push budget does not return to the M5 queue from a liquidation lane. It returns to the LTSF recovery plan, because the dollars were funding a position the cliff calendar has already sold. State the amount and its destination in the same staged set.

4. A push or rank objective found live on a tagged variation executes R1 and R4 the same day and is failure mode F1. The spend was defending a position the calendar had already sold.

5. Where the product is not in Recovery Mode, this staged set is human-confirmed at the standard v4.0 Section 12 "Layer 7: Precedence and Conflict Resolution (P1-P10)" rule P5 tier. Where the product is in Recovery Mode, v4.0 R3 tightens the confirmation line and every step in this set is confirmed at the tightened line.

P7. Archetype A execution: rebuild traffic while the curve is repaired (40 minutes for the first build, 15 minutes per weekly pass)

1. Read the diagnosed defect from the declaration: click-through defect, conversion defect, or both. LTSF Section 5 assigns the diagnosis; PPC executes against it and does not re-diagnose the listing.

2. Do not lead with depth and do not wait for the repair. Both tracks run from day one. The listing sprint is Design's under the LTSF Section 9.2 48-hour SLA; the paid rebuild is this procedure's under its 24-hour SLA.

3. Competitor keyword replication, which LTSF Section 6.2 calls the single best-grounded targeting method available. Pull the reference competitors the declaration names: same price, same rating, comparable offer, inside the top 10,000 BSR. Pull their ranking and bidding terms from the Keepa competitor pull and the SQP export in the Raw Pull Pack. Stage the resulting exact set inside the dated liquidation campaign, seeded at the term's Expected CPC from T2 block K6, capped at the R2 ceiling.

4. Competitor keyword replication is a keyword lever and is not a conquest open. It buys the terms the competitor ranks for; it does not target the competitor's ASIN. PAT and conquest targets stay with #16 and enter only on that component's weakness flags.

5. New discovery campaigns, where the declaration authorizes a test budget. Fixed budget, capped bid, a dated stop-loss, and an expiry date before the cliff date rather than after it. A test that reads after the decision is made is a charge with no decision value.

6. Break-even PPC scaling opens only once the CVR read recovers past its declared benchmark. Scaling a defect multiplies the defect: LTSF Principle 4 states that a 40% discount on a listing converting far below category norm sells 6 units instead of 4 at worse economics, and the same arithmetic governs a budget increase.

7. Each lever carries a logged prediction with its basis tag before launch, per LTSF Section 8.2 "Velocity uplift estimates (the judgment half)": HIST from the product's own lever history, MARKET from the competitor-anchored ceiling, or TEST from a pre-registered 7-day test. An untagged multiplier renders as blank in the tracker and blocks the scenario from ranking, so PPC supplies the tag with the number.

8. Read at day 7 against the logged prediction. Below half the prediction the lever is replaced rather than extended, per LTSF Section 9.1 "Pre-registered 7-day tests".

P8. Archetype B execution: child-ASIN targeting (35 minutes for the first build, 10 minutes per weekly pass)

1. Read which children hold the aged units. The declaration names them; the LTSF Section 5 diagnostic signal is that specific colors or sizes hold 80% or more of the aged units against a parent that sells well.

2. Open T2 block K2 "Targeted variation(s)" and the D2 variation map, and list every live campaign whose advertised SKU is the preferred child while the aged units sit on a non-preferred child. That list is the whole problem: the aged children are usually excluded from campaigns entirely, so they receive no targeted spend at all.

3. Stage the re-point toward the aged child, executed at #12 P10 against the D2 map. Read the direction carefully: the G6 re-point moves spend away from a child that is thinning toward a ranked backup, and this re-point moves spend toward a child that is overstocked. Same machinery, opposite direction, and the D2 map holds both because it is a map of which variation a term lands on, not a map of which variation is scarce.

4. Where a term has no exact instance pointed at the aged child, the verdict is ADD EXACT + RESTRUCTURE and the instance stands up inside the aged child's dedicated ad group, with its walls built in the same session.

5. Refined targeting on the non-preferred children: dedicated ad groups per aged child, per-child bids computed from that child's own economics rather than the family blend. A child at a lower price point carries its own break-even, and reading it against the parent's blended line is the error v4.0 Section 8.8 "Conversions Scenarios (S-P1 to S-P8)" scenario S-P8 exists to catch.

6. New discovery campaigns scoped to the aged children only, at fixed test budgets, where the child's demand has never been mapped.

7. Do not touch the parent and do not touch the healthy children. No family-wide budget move, no shared coupon, no parent re-point. The healthy children subsidize the fix and their economics and price anchors stay intact; that is an LTSF hard rule, not a preference, and breaching it is failure mode F8.

8. Where the aged child shares a term with a sibling product, gate G11 holds and the finding routes to #33 before any scale step deploys.

P9. Archetype C execution: clear what the ceiling buys, then stop (20 minutes)

1. Retag and quarantine first: P2, then P5, then P6, in that order, inside the tag week.

2. Build the proven converter set from the lifetime search term reports and the SP and SB search term reports in the Raw Pull Pack: terms carrying orders at or above the v4.0 gate G2 sufficiency line on the window the declaration covers. Nothing else enters the ad group.

3. Set every bid at or under the R2 maximum liquidation CPC from P3. There is no ladder, no step size and no discovery from above: the ceiling is a hard line and the bid either clears the auction under it or it does not.

4. Open no experiments. No discovery campaign, no auto probe, no competitor replication build, no new keyword test, no placement test. Section 1.3.1 states the arithmetic: the read window exceeds the decision window, so the test cannot inform the decision and only accrues the bracket rate. An experiment opened on an Archetype C tag is failure mode F7.

5. Run P4. Archetype C lanes close on arithmetic more often than any other class, because the floor price is at or near salvage and the ceiling is therefore thin. Route the closure and do not respond by widening the term set.

6. Where the declaration escalates a new launch with no organic traction to a strategic viability review, PPC contributes the paid-side evidence and stops: the lifetime converter list, the untested-term map, and the tested-and-failed list. The exit, redesign or reposition call is not PPC's.

P10. Archetype D execution: break-even scaling against the bracket date (30 minutes for the first build, 12 minutes per weekly pass)

1. No retag and no quarantine. The units are aged, not terminal; the product sells and the margin holds. Rank objectives on the product's other terms are unaffected and continue at #12.

2. Read the weekly velocity-to-cliff requirement from the declaration's cliff calendar columns, and convert it to required units per day and required clicks per day under P4.

3. Identify the proven syntaxes from T1 Syntax Board group S5 "Four-quadrant diagnosis": the STRONG and VISIBILITY syntaxes carrying the affected size or child. Budget and bid increases on these are the fastest lever in the LTSF inventory, with a 1-3 day impact window.

4. Raise the ceiling from the target line to the break-even line under P3 step 4, and stage the bid step to the new ceiling on the proven syntaxes only. This is the whole of break-even PPC scaling: it moves no price, damages no anchor, stacks with no coupon, and costs zero margin dollars at the break-even line because the line is where contribution reaches zero.

5. Raise the daily budget to the required-clicks figure. Budget precedes bid whenever the campaign is truncated, per v4.0 Section 9 "Layer 4: The Lever Library (L1-L10). Order, Limits, Mechanics" rule L1.

6. Hold the guard: the realized net price must stay above the declared floor price at every step. Where a depth rung fires from the LTSF ladder, the net price falls, the guard tightens and the ceiling is recomputed at P3 rather than carried.

7. Log the prediction with its basis tag and read it at day 7 under LTSF Section 9.1.

8. Do not remove units the product will profitably sell. The play is speed. Where the pace math says the batch clears before its bracket date at the break-even ceiling, no depth rung and no terminal option is required and PPC says so in the P12 check line.

P11. The Family 3 to 5 levers that PPC executes (45 minutes for the first build, 10 minutes per later SKU)

1. Sponsored Display retargeting. Retarget detail-page viewers of the family with creative featuring the aged SKUs and their offer. Cheap incremental reach on shoppers already in market, with a 3-14 day impact window. Judged on cost per acquisition against the R2 allowable ad cost per unit, not on a standing efficiency band, and judged in its correct mode: v4.0 Section 8.10 "Brand-Format Scenarios (S-F1 to S-F8). SB, SBV, SD judged on their own physics" scenario S-F5 puts cost-per-click mode against the CPA ceiling and scenario S-F6 puts vCPM mode on new-to-brand and assisted conversions rather than raw efficiency. Scoring a vCPM buy on an efficiency band is structural error, not a bad week.

2. Brand Store clearance page fed by Sponsored Brands. One standing clearance page in the Brand Store aggregating every aged SKU, with a 1-2 week build window. One build serves every product's clearance at once, which is why it is built once and maintained rather than rebuilt per SKU. The page build is Design's; the Sponsored Brands traffic to it is this SOP's specification and #17-19's execution.

3. The membership rule for the clearance page is the tag, not the tier: a SKU enters when it carries a terminal or at-risk tag and leaves when the tag clears or the units reach zero. P12 checks membership weekly against the LTSF row so the page never advertises a SKU that has already cleared.

4. Where a deal or event window touches a tagged variation, #32 P12 owns the window plan and this SOP supplies R2 and R3 to it. The window's value on a terminal lane is the conversion move: a deal price raises CVR, which raises the R2 maximum liquidation CPC proportionally, which is how a lane that arithmetic closed at baseline conversion can open at the window price. Re-run P3 and P4 at the window CVR before the window opens, not after it closes.

5. Every other lever in LTSF Sections 6.3, 6.4 and 6.5 is executed outside PPC: coupons, Prime Exclusive Discounts, price drops, outlet deals, deals, brand tailored promotions, social codes, virtual bundles, cross-channel paths and every terminal lever. PPC reads their effect on CVR and net price and recomputes the ceiling; PPC does not set their depth.

P12. The weekly expiry check and the pace watch (R6) (20 minutes per SKU weekly, 5 minutes per daily read)

1. Run this before any new verdict is read on the SKU, in the same position the standing pass gives to scoring: v4.0 Section 14 rule V1 states that unscored actions do not accumulate.

2. Write the check line, which has exactly four numbers: realized recovery per unit moved, the declared alternative's per-unit value, realized pace in units per day, and the R3 required pace. The line is written whether or not anything changed.

3. Where the alternative wins on the first comparison, PPC stops that week. The objective function does not bend to finish what it started, and a lane that has become worse than removal is not improved by another week of spend. Route to P14.

4. Read every lever deployed on the SKU against its day-7 logged prediction. Below half the prediction the lever is replaced rather than extended, per LTSF Section 9.1. A coupon producing zero velocity change is a pure cost; the equivalent on the paid side is a discovery campaign at zero graduations.

5. Inside the final 14 days before the declared cliff date, and unconditionally at CRITICAL tier, read realized units cleared against the R3 pace once per working day from T3 group C4 order counts on the dated liquidation campaigns. A pace gap at day 10 is recoverable inside the ceiling; the same gap at day 3 is not, which is why the read is daily rather than weekly at the end.

6. Carry the SKU's line into the daily Brand Management and PPC huddle per LTSF Section 10.1 whenever a rung target or expiry date is at risk.

7. Re-derive the R3 pace whenever remaining units, the cliff date or the CVR moved. A pace figure carried from last week against a unit count that fell is a figure that overstates the remaining requirement and understates the ceiling.

P13. Tier posture change (15 minutes)

1. Trigger: the weekly auto-classification moves the tier, or an archetype change reclassifies the SKU. PPC executes the new posture and proposes no tier move; tier movement is decided at LTSF Section 7.3.
2. Route on the direction of the move.
3. On any move into RED or CRITICAL, stage the new posture the same day. On a move down, stage it in the next Monday pass, because a posture that relaxes has no cliff pressure behind it.
4. Write the posture change into T6 with the tier value on both sides of the move and the LTSF row reference. A posture change with no recorded prior tier cannot be scored.
5. Where a code-red declaration lands on the same product, the terminal lane continues inside its quarantine and every step in it is human-confirmed at the v4.0 R3 tightened line. Section 1.5.1 states the seam and WE-4 works it in dollars.

```
P13 step 2 routing
|
+-- GREEN -> YELLOW ....... open the archetype lever set on the aged children or batch; no retag
+-- YELLOW -> RED ......... full lever set at the R2 ceiling; retag and quarantine where terminal;
| arm the daily pace watch inside the final 14 days
+-- RED -> CRITICAL ....... R4 shut-off complete; twice-weekly cadence; daily pace watch standing;
| lane closes or clears inside the same working session
+-- Any move down ......... relax the posture at the next Monday pass; do not un-quarantine a
tagged variation on a tier move alone, because the tag, not the tier,
is what R5 quarantines
```

P14. Lane closure, handback and the closure record (25 minutes)

1. Trigger: P4 returns lane closed, or P12 returns the alternative winning, or the cliff date arrives, or remaining units reach zero.
2. Write the finding in the form the LTSF system can act on: required clicks per day, the R2 ceiling, market CPC at the required volume, the units PPC can clear at the ceiling, and the residual. A closure stated as a conclusion rather than as those five numbers cannot be acted on and is returned.
3. Set the affected rows to PAUSED-PENDING (REASON + RE-ENTRY), with the re-entry criterion written as the declaration change that would reopen the lane: a lower floor price, a deal-window CVR, or more days to the cliff. v4.0 Section 6 rule O5 requires the blocking reason, the re-entry criterion and a review date, and nothing exits the system silently.
4. Where units reached zero or the tag cleared, write the closure record: units cleared through the lane, total ad spend, realized recovery per unit against the alternative, and the predicted-versus-actual delta per lever for the LTSF outcome log. That delta is the basis tag HIST for the next exit program and is the reason the log exists.
5. Archive the dated liquidation campaign only after the closure record is written, and remove it from the #37 exclusion list and the Brand Store clearance page in the same session. An archived campaign left on the exclusion list hides a live automation gap.
6. Escalate rather than close where the decision has been blocked more than 5 business days awaiting an input, per LTSF Section 7.4 "Escalation thresholds". Blocked decisions accrue bracket rates like everything else.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds and scenario logic live in v4.0. Every string below is quoted verbatim from the Keyword Decision Record Template v3.0 sheet "Bands & Verdicts Reference", block "CLOSED VERDICT VOCABULARY (Framework Section 17):

exactly one per Keyword Decision Record", cross-read against v4.0 Section 13.1 "The Closed Verdict Vocabulary". This table states what each verdict does when it lands on a row inside the terminal lane. Twenty-five rows, because a verdict with no stated terminal-lane behaviour is a seam an operator falls into.

Verdict (closed vocabulary)	Terminal-lane execution
QUARANTINE: DATA REFRESH	No change deploys on the row and the repair task opens the same day. The R2 ceiling holds at its last dated value per P3 step 6. The repair ticket carries the SKU's bracket rate, because the charge accrues while the data is broken and the repair is therefore dated, not queued.
UNLOCK TRAFFIC	Bid step upward toward the R2 ceiling and never through it. Where the bid already sits at the ceiling, the verdict is a P4 re-run rather than a step: the lane is volume-constrained by arithmetic, not by bid timidity.
FIX BUDGET (TRUNCATED)	Budget lever only on the dated liquidation campaign and the week's bid read voids. On this lane the cost of truncation is measured in units not moved, which P12 records against the R3 required pace rather than against an efficiency band.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	Barred on Archetype C by Section 1.3.1. Legal on A, B and D, capped at the R2 ceiling and dated to expire before the declared cliff date, never after it.
FUND TEST (UNFUNDED PRIMARY)	As above, and the funding call sits with the LTSF Owner on a tagged SKU because a test budget is a charge against recovery rather than an investment against rank.
RANK PUSH: FUND SIZED PUSH	Does not apply to terminal-tagged entities. R1 removes rank objectives on the day of the tag and R4 shuts push spend. On an untagged term of the same product the verdict executes unchanged at #12.
RECOVERY PUSH	Does not apply to terminal-tagged entities, for the same reason. A pre-OOS baseline describes a position the cliff calendar has already sold.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	The terminal-lane required-clicks figure is R3's, derived from remaining units and days to cliff, not the standing DSTR figure. Where the R2 ceiling cannot buy them, the verdict routes to P14 instead of raising the bid.
FIX CONVERSION FIRST; DEFER PUSH	On Archetype A this is the standing state rather than a deferral: listing repair is LTSF Family 1 and runs in parallel under the 48-hour SLA. There is no push to defer, and PPC holds at maintenance on the affected terms while the repair lands.
PLACEMENT FIX FIRST	No terminal-lane form. R4 normalizes modifiers to zero on tagged entities, so there is no placement premium to correct. On untagged Archetype D terms the verdict executes at #28 unchanged.
RE-POINT VARIATION	The Archetype B core move, executed at #12 P10 against the D2 map, and pointed toward the aged child rather than away from a thinning one. P8 step 3 states the direction so the two uses of the same machinery are not confused.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	Does not apply. There is no rank objective on a tagged entity to achieve, and therefore no post-objective ladder to run.
HOLD; DEFEND LIGHTLY AT RANK	Does not apply. R4 removes brand-defense allocation on tagged entities.
DEFEND AT FLOOR	Does not apply on tagged entities, and the word floor is the trap. In this SOP a floor is the LTSF floor price, which is a price per unit. In the standing library a floor is a discovered minimum viable CPC, which is a bid. The two are different objects with the same name and confusing them is failure mode F5.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	The profitable CPC on a tagged entity is the R2 maximum liquidation CPC, not the standing max profitable CPC in T2 block K6. Both values sit in the row per P3 step 5 and the tagged row reads the R2 value. The syntax-cap half executes at the envelope instruments.
INCREMENTALITY LADDER	Does not apply. R4 removes the defense allocation the ladder measures, and the test design belongs to #35 in any case.
ADD EXACT + RESTRUCTURE	Legal on Archetype A and B where a proven converter carries no exact instance. The instance stands up inside the dated liquidation campaign or the aged child's dedicated ad group, never in the standing family structure. Barred on C as an experiment.
CONSOLIDATE DUPLICATES	Executes at #12 P7. A tagged term live in both a standing campaign and a liquidation campaign is not a duplicate awaiting consolidation; it is an R5 quarantine failure and routes to F3.

Verdict (closed vocabulary)	Terminal-lane execution
GRADUATE FROM DISCOVERY	Legal on A, B and D where discovery is open. The graduate lands inside the dated liquidation campaign with its walls built in the same session, and its expiry is the cliff date rather than a standing review date. Barred on C.
NEGATE (PRE-LOAD / REACTIVE / STEERING)	The steering wall is what makes R5's quarantine real rather than nominal. #29 owns level selection; this SOP supplies the term list and records the wall list in T3 group C5 at the retag.
ENTER SBV ARBITRAGE	Does not apply. Brand formats reach the terminal lane through P11 as the LTSF Family 2 and Family 3 levers and are judged on recovery per unit against R2, not on the SP-to-video cost inequality.
OPEN / ROTATE / EXIT CONQUEST	Does not apply here. Competitor keyword replication at P7 step 3 buys the terms a competitor ranks for and is a keyword lever; conquest targets a competitor ASIN and stays with #16 on that component's weakness flags. P7 step 4 states the distinction.
STOP-LOSS: CONCEDE OR DEFER	The terminal-lane form is lane closure at P14. Freed budget does not re-enter the push funding queue from a liquidation campaign; it returns to the LTSF recovery plan, because the dollars were never push dollars.
PAUSED-PENDING (REASON + RE-ENTRY)	Used at P14 step 3 for a lane closed on arithmetic. The re-entry criterion is written as the declaration change that would reopen the lane: a lower floor price, a deal-window conversion rate, or more days to the cliff.
LEAVE TO HARVEST	Does not apply. Terminal units have no harvest horizon. Every unit either clears before its cliff or takes the terminal option, and a term left to harvest on a tagged variation is spend with no expiry, which R6 exists to prevent.

Campaign-level verdicts follow the Keyword Decision Record Template v3.0 sheet "Campaign Decision Record" vocabulary: FUND / HOLD / TRIM BUDGET, RAISE / CUT MULTIPLIER, NEGATION PASS, CONSOLIDATE, RETAG OBJECTIVE, PAUSE (L7), ARCHIVE. On the terminal lane, RAISE MULTIPLIER never issues because R4 holds modifiers at zero; RETAG OBJECTIVE issues once, at P2; ARCHIVE issues once, at P14 step 5, and only after the closure record is written.

6. Worked Examples

Four cases. All values are illustrative and internally consistent, built on the standard fictional training base of the Economics Worked-Example Set v1.0 (AOV $80.00, contribution $28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Maximum Acceptable 26.3%, CVR 5.2%) except where a case declares its own child-level or converter-level values. No live product economics spine was attached to this session. Every declaration field is read from the LTSF row and is never derived here.

WE-1. RED tier, Archetype D: break-even PPC scaling against a bracket date

The declaration, read not derived. Archetype D (Aged Healthy). Tier RED. 624 aged units, all sitting in the 271-300 day bracket at the LTSF Section 3 rate of $5.45 per unit per month. The next rate step is the 366-455 day bracket at $6.90, and the largest bracket crosses it in 78 days. Declared floor price $58.40. Salvage value per unit negative $1.05. Contribution at the current price $28.00 per unit; contribution at the floor price $6.40 per unit. Terminal option on file: removal.

P1 and P13. Archetype D routes to P10. Tier RED opens the full lever set and arms the daily pace watch inside the final 14 days. No retag and no quarantine: the units are aged, not terminal, so the standing campaigns keep serving and the product's other terms keep their rank objectives at #12.

P4 pace. Required units per day = 624 / 78 = 8.0. Required clicks per day = 8.0 / 0.052 = 153.8.

P3 ceiling. This is Archetype D, so the ceiling comes from the break-even PPC scaling lever rather than from R2's salvage subtraction. Spend ceiling per unit is the contribution above the floor price, and the declaration carries both figures so PPC subtracts rather than models fees: 28.00 minus 6.40 = $21.60 of price headroom per unit, which confirms the guard has room and no depth rung is required at this level. The bid ceiling is the break-even line itself: contribution times conversion rate = 28.00 x 0.052 = $1.456, which is the same object as v4.0 Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)" family F1's max profitable CPC. The standing posture on this product buys at the target line, 80.00 x 0.175 x 0.052 = $0.728. Break-even scaling doubles the authorized click price from $0.73 to $1.46 and costs zero margin dollars at the line, because the line is where contribution reaches zero.

P4 traffic check. Market CPC on the STRONG and VISIBILITY syntaxes carrying the affected size reads $1.18 at the volume the lane needs. $1.18 is under the $1.456 ceiling with $0.28 of headroom per click, so the lane is OPEN. Daily spend = 153.8 x 1.18 = $181.54. Daily budget staged at that figure on the proven syntaxes only.

The week, computed. Units 8.0 x 7 = 56.0. Clicks 1,076.9. Spend 1,076.9 x 1.18 = $1,270.77. Sales 56.0 x 80.00 = $4,480.00. Lane ACOS = 1,270.77 / 4,480.00 = 28.4%. Read that number against the three lines: it is below break-even 35.0% and above Maximum Acceptable 26.3%, which is exactly the band the standing posture forbids and break-even scaling opens. Contribution recovered 56.0 x 28.00 = $1,568.00, so the lane nets $297.23 per week in cash while clearing at pace. Ad cost per unit = 1,270.77 / 56.0 = $22.69, leaving $5.31 of contribution retained per unit.

The full clear. 12,000 clicks over 78 days at $1.18 = $14,160.00 of spend returning 624 x $28.00 = $17,472.00 of contribution, net $3,312.00, and zero units cross the 366-day step.

Against the do-nothing baseline. At the standing lane's velocity of 3.1 units per day the batch clears 242 units in 78 days and 382 units cross into the $6.90 bracket. The step alone costs 382 x (6.90 minus 5.45) = $553.90 per month, and it recurs every month until the remainder clears. That number, not an efficiency band, is what authorizes the ceiling change.

The record. Verdict TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS on the proven syntaxes, with the required-clicks figure sourced from R3 rather than from DSTR. Prediction logged with basis tag HIST from the product's own prior budget-lift response. Review at day 7 per LTSF Section 9.1, and the pace watch arms at day 64.

WE-2. CRITICAL tier, Archetype C: the lane closes on arithmetic, and the experiment that would have made it worse

The declaration. Archetype C (Structurally Dead). Tier CRITICAL. 480 units in the 366-455 day bracket at $6.90 per unit per month. Cliff to the 456-plus bracket at $7.90 in 40 days. Salvage price nets $6.40 per unit; the winning alternative nets negative $3.10 per unit, which is the charge avoided. Converter CVR on the proven set reads 4.0%. This case follows the arithmetic convention of the Economics Worked-Example Set v1.0 case E7.

P2, P5, P6 inside the tag week. Terms retag to the liquidation objective, the variation leaves the standing family structure with its walls built in the same session, and R4 shuts rank objectives, pushes, discovery spend and brand-defense allocation, with modifiers normalized to zero. What remains is exact and auto close match on proven converters only.

P3 ceiling. Allowable ad cost per unit = 6.40 minus negative 3.10 = $9.50. That is more than the unit's own net recovery, and correctly so: every advertised sale also avoids a charge. Maximum liquidation CPC = 9.50 x 0.04 = $0.38.

P4 traffic check. Required units per day = 480 / 40 = 12.0. Required clicks per day = 12.0 / 0.04 = 300. Market CPC on the proven converters reads $0.55. The lane cannot buy 300 clicks at $0.38, so the lane does not open at full pace. Residual volume clearing at or under the ceiling reads 45 clicks per day, which is 45 x 0.04 = 1.8 units per day, or 72 units across the 40 days at a spend of 45 x 40 x $0.38 = $684.00, which is $9.50 per unit and exactly the ceiling by construction. Result: LANE SPLIT. PPC clears 72 units; 408 units route to the terminal option.

What goes wrong, and how it is handled. On the day the split is stated, the operator proposes a discovery campaign to find cheaper converters, on the reasonable-sounding argument that the current converter set is small and the ceiling is not the constraint everywhere. That proposal is failure mode F7 and Section 1.3.1 is the answer in numbers. Archetype C carries a 7-day decision window. Every Family 2 discovery lever carries a 7 to 14 day impact window and a 7-day read before it can be scaled. The read therefore returns at or after the decision, so it can change nothing, and it costs 480 x $6.90 / 30 x 7 = $772.80 of pure charge to run. The correct response is to route the split, not to widen the term set. The operator files the proposal as an observation on the closure record so the next Archetype C case inherits the reasoning rather than re-arguing it.

What the split is worth. Removal on the 408-unit remainder costs 408 x $1.17 = $477.36 once, against 408 x $7.90 = $3,223.20 per month if the units are held past the cliff. The removal path clears by a factor of 6.8, which is the LTSF standing test: the cheapest total-cost path wins and removal is not failure.

The record. Verdict on the residual set: BID DOWN TO PROFITABLE CPC; CAP SYNTAX, with the profitable CPC reading the R2 value of $0.38 rather than the standing max profitable CPC. Verdict on everything outside the residual set: PAUSED-PENDING (REASON + RE-ENTRY), re-entry criterion written as a declared salvage-price move that raises conversion, or a deal window that does the same. The finding routes to the LTSF Owner the same working session with the five P14 numbers attached.

WE-3. YELLOW tier, Archetype B: the aged children nobody was bidding on

The declaration. Archetype B (Variation Overstock). Tier YELLOW. The parent sells well. 900 aged units across the family, of which 738, or 82.0%, sit on three non-preferred child ASINs. Units are at day 226, inside the 211-240 day bracket at $1.00 per unit per month, so the current charge on the aged children is 738 x $1.00 = $738.00 per month. The next step is $1.50 at day 241 and the 271-day cliff at $5.45 lands in 45 days. Child economics declared separately from the family blend: child AOV $64.00, contribution $19.20 per unit, child conversion rate 4.6%.

P8 step 2, the finding. Every live exact campaign on the family advertises the preferred child. The D2 variation map confirms it: the three aged children appear as advertised SKUs on zero campaigns, so they receive no targeted spend at all. This is the ordinary Archetype B state and it is why the units aged.

Child economics, computed on the child and not on the blend. Break-even ACOS = 19.20 / 64.00 = 30.0%. Target = 15.0%, Maximum Acceptable = 22.5%. Child max profitable CPC = 19.20 x 0.046 = $0.883. Reading these children against the parent's 35.0% break-even would have overstated the ceiling by 5 percentage points, which is the error v4.0 Section 8.8 "Conversions Scenarios (S-P1 to S-P8)" scenario S-P8 exists to catch.

P4 pace. Days to the 271-day cliff = 271 minus 226 = 45. Required units per day = 738 / 45 = 16.4. Required clicks per day = 16.4 / 0.046 = 356.5. Market CPC on the child-specific terms reads $0.52, under the $0.883 ceiling, so the lane is OPEN. Daily spend = 356.5 x 0.52 = $185.39.

The window, computed. Across 45 days: spend $8,342.61, units 738, ad cost per unit $11.30, lane ACOS = 11.30 / 64.00 = 17.7%. That sits below Maximum Acceptable 22.5% and well below break-even 30.0%, so this lane is not even a break-even case; it is ordinary profitable spend that was simply never funded. Contribution 738 x $19.20 = $14,169.60, net after ads $5,826.99. Charge avoided by clearing before the cliff: 738 x (5.45 minus 1.50) = $2,915.10 per month.

What is not done. The parent price is not cut, no family-wide coupon is set, and the healthy children's campaigns are untouched. Breaching that is failure mode F8. Depth authority on the aged children, where any is needed, belongs to the LTSF Owner and not to PPC.

The record. Verdicts staged: RE-POINT VARIATION toward each aged child on terms already carrying an instance, and ADD EXACT + RESTRUCTURE where a proven converter carries no instance pointed at the aged child, each landing in that child's dedicated ad group with walls in the same session. Prediction basis tag MARKET, anchored on the preferred child's measured conversion at the same terms.

WE-4. A RED terminal lane inside a code-red declaration: the seam, priced

The setup. The WE-1 product, still running its 624-unit Archetype D lane at $1,270.77 per week, is declared code red on the same Monday. Product weekly spend $4,100.00 against weekly sales of $25,466.00 gives TACOS 16.1%. v4.0 Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)" family F2 sets the mature reference band at 4-7%, so 16.1% against the 7.0% ceiling is 2.30x band, above the greater-than-2x limb, and #26 declares under v4.0 Section 8.1 "Product-Level Scenarios (S-A). The posture every keyword inherits" scenario S-A4.

What binds the lane, and what does not. v4.0 Section 11 "Layer 6: Transitions, Stages, Modes (T1-T8) and Recovery Mode (R1-R6)" rule R3 tightens the confirmation line while Recovery Mode is active: every bid move above 15% and every budget move above $25 per day is human-confirmed. That applies to the liquidation lane exactly as it applies to everything else, so the lane continues and every step inside it is confirmed at the tightened line. R3's bar on new Ranking objectives is not binding here, because R4 removed rank objectives from these units on the day of the tag. Rule T8 orders Recovery Mode above Event Mode above stage defaults and names the terminal lane nowhere, which is amendment item AM-2.

The cut that looks right. #26's cut sequence is ordered by loss dollars across the product's standing campaigns. The liquidation lane is the largest single above-band line on the product, so an operator reading loss dollars alone would cut it first. Price the cut. Spend falls by $1,270.77 per week. Sales fall by $4,480.00 per week, so contribution falls by $1,568.00. The 624 units stay in FBA and accrue 624 x $5.45 = $3,400.80 per month, which is $793.52 per week. Net cash effect of the cut = 1,568.00 plus 793.52 minus 1,270.77 = $1,090.75 worse per week.

Why the metric rewards the wrong move. After the cut, product spend reads $2,829.23 against sales of $20,986.00, so TACOS reads 13.5%, which is 1.93x band. The cut moves the product out of the code-red limb and into breach, so the protocol's own exit test improves while the account loses $1,090.75 per week in cash. That is the entire reason R5's quarantine is load-bearing rather than tidy: dated liquidation campaigns are not standing campaigns, the cut sequence does not reach them, and the arithmetic above is why it must not.

The record. No verdict issues on the liquidation lane from the code-red diagnosis. The lane's own weekly line continues under P12. The finding, that a loss-ordered cut sequence run without the quarantine boundary would have destroyed $1,090.75 per week while improving the headline reading, files as an amendment proposal into the #26 post-exit review, which is where v4.0 Section 11 rule R6 says recovery lessons belong.

7. Failure Modes and Escalation

Twelve modes. F1 to F4 carry forward from v1.0 unchanged in substance. F5 to F12 are new at v2.0 and exist because the deciding system is now readable, which makes eight failure classes visible that a boundary-stopped document could not name.

Failure	Detection signal	Response, and the incident that motivated it
F1. Rank spend on terminal stock	A push objective or rank objective live on a tagged variation, read from T2 block K7 "Objective" against the terminal flag in block K2	R1 and R4 execute the same day. The spend was defending a position the calendar had already sold. Motivating class: the account's named pattern of spend continuing after the objective is achieved, listed in Strategic Doctrine v1.0 Section 14, here in its most expensive form because the position has no future value at all.
F2. Sunk COGS in the math	Any bid or ceiling computation citing acquisition cost, found in the T6 reason string or the P3 working	R7 and LTSF Principle 2. Recompute from the LTSF row. Motivating class: LTSF Section 2 records that displaying lifetime profit inclusive of COGS makes every option on an underwater SKU look like a loss, which produces paralysis rather than a decision.

Failure	Detection signal	Response, and the incident that motivated it
F3. Mixed reads	Terminal units still serving inside standing campaigns past the tag week, read from T3 group C1 advertised SKU against the tagged variation list	R5 quarantine executes and both reads re-baseline. Mixed reads poison the standing product's economics and the liquidation's recovery math at once, so neither number is usable until the separation lands.
F4. Expiry ignored	Liquidation spend live past a week in which the alternative won the R6 comparison, or past the declared cliff date	Stop the same day. The R6 check is the contract and the objective function does not bend to finish what it started. Motivating class: LTSF Section 9 records the prior operating pattern in which a non-performing lever ran for a month unexamined.
F5. The two floors confused	A bid staged at the LTSF floor price, or a price proposed at the discovered minimum viable CPC. Detected at staging QA: a bid value in the same order of magnitude as the product price	Reject the staged row and re-read Section 5. The LTSF floor is a price per unit and belongs to the LTSF system; the discovered floor is a bid and belongs to #28. Two objects, one word, and the collision is why this mode is written down.
F6. Ceiling carried after an input moved	An R2 ceiling in T2 block K6 whose date precedes the last change to floor price, salvage value, unit count, cliff date or conversion rate	P3 recompute the same working session, then P4 re-run. A ceiling is a function of five declared inputs and is stale the moment any one of them moves. The pace figure is stale with it.
F7. An experiment opened on an Archetype C tag	Any discovery campaign, auto probe, competitor replication build or new keyword test dated after an Archetype C classification	Close it the same day and route the intent to the closure record instead. Section 1.3.1 carries the arithmetic: on a 480-unit batch at $6.90 per unit per month one 7-day read window costs $772.80 and returns after the decision it was meant to inform. Worked at WE-2.
F8. Parent damage on an Archetype B fix	A price move, coupon or budget change touching the parent or a healthy child while the aged children are the tagged set	Revert the same day and re-stage against the aged children only. LTSF Section 5 states it as a hard rule: the healthy children subsidize the fix, so their economics and their price anchors are not spent on it.
F9. An automation rule reaching a liquidation campaign	A bid or budget change on a dated liquidation campaign appearing in the T+1 verification diff with no staged row behind it, or the campaign absent from the #37 exclusion list	Route the root cause to #37 change control the same day and restore the staged value. A rule tuned for efficiency bands will ladder a liquidation bid down toward a standing ceiling that does not apply, which closes the lane silently. Ledger dependency D6 reads OPEN, so the exclusion is confirmed by manual read until the export lands.
F10. Depth changed without a ceiling re-read	A coupon, deal or price move landing on a tagged child in T5 Product Log with no P3 recompute in the same cycle	Recompute at P3 and re-run P4. A depth change moves net proceeds and conversion at once, so it moves both limbs of the R2 ceiling. Running the old ceiling after a depth change either forfeits volume the new conversion rate would have bought or overspends against a net price that fell.
F11. Lane closure reported as a conclusion	A P14 handback in T6 that does not carry all five numbers: required clicks per day, the R2 ceiling, market CPC at required volume, units clearable at the ceiling, and the residual	Return it and rewrite. The LTSF system acts on the residual count and the ceiling, not on the word closed. A closure without its numbers reads as a PPC failure rather than as the arithmetic finding it is, and it invites a price argument instead of a price move.
F12. Quarantine released on a tier move rather than a tag change	A tagged variation returning to standing campaigns in the same cycle a tier moved down, with the terminal flag still set in block K2	Re-quarantine the same day. R5 quarantines the tag, not the tier. A CRITICAL SKU moving to RED is still terminal, and re-mixing its reads costs both sides of the separation a fresh baseline.

Escalation. The same failure mode twice in consecutive cycles on the same scope escalates to the PPC Lead with the T6 history attached. Anything in the human-mandatory tier at Element 9 is confirmed before deployment without exception. Separately, LTSF Section 7.4 "Escalation thresholds" escalates any decision blocked more than 5 business days awaiting an input directly to the LTSF Owner, because blocked decisions accrue bracket rates like everything else. The two escalation rules count different things, one in cycles and one in business days, and no rule reconciles them: amendment item AM-7.

8. Outputs and Artifacts

Gates clear with artifacts, not assertions.

Event	Named record	Location and retention	Downstream consumer
Declaration receipt (P1)	The seven declared fields copied against the variation, plus one dated T5 Product Log event of type targeting change carrying the LTSF row reference	T2 block K2 and T5 Product Log in the product's PPC Command Workbook; retained for the life of the product log	Every later attribution read on the product under v4.0 rule D5; #32 when a window touches the variation
Retag (P2)	Objective set to liquidation in block K7 with the grading metric in the situation text and the review date set to the cliff date, plus its T6 prediction row	T2 and T6 Decision Log; T6 retained as the scoring corpus	The V5 scoring pass; #37 for protected-scope exclusion; the LTSF outcome log
Ceiling and pace (P3, P4)	R2 allowable ad cost per unit, maximum liquidation CPC with its date, R3 required units and clicks per day, and the traffic-check result	T2 block K6 beside the standing max profitable CPC; recomputed rather than versioned	#32 P12, which builds its liquidation-window cap from the R2 ceiling times the R3 required pace
Quarantine (P5)	The dated liquidation campaign specification with the cliff date in the name, the wall list, the modifier-zero state, and the #37 exclusion-list entry	T3 Campaign Board groups C1 and C5; the wall list retained as the wall-audit baseline	#12 for creation; #29 for the wall audit; #37 for exclusion; #38 for staging
Shut-off (P6)	The staged set removing rank objectives, pushes, discovery spend and defense allocation, with the freed-budget amount and its destination stated	Staged through #38 in its format; the amount recorded in T6	The LTSF recovery plan, which receives the freed dollars; #38's T+1 verification
Weekly check (P12)	The R6 check line: realized recovery per unit, the alternative's per-unit value, realized pace, required pace. Written whether or not anything changed	T6, one row per SKU per cycle; retained for the life of the tag plus the closure record	The LTSF Owner's weekly cycle; the daily huddle where a target is at risk; Account Rollup A3 for V5 scoring
Daily pace read (P12 step 5)	Units cleared against required pace, one line per working day inside the final 14 days and standing at CRITICAL	T6; retained through closure	The daily Brand Management and PPC huddle
Lane closure (P14)	The five-number handback and the PAUSED-PENDING rows with their re-entry criteria	T6 and the LTSF row; retained as the decision record for the terminal option	The LTSF Owner, who moves the price or executes the terminal option on the residual
Closure record (P14 step 4)	Units cleared through the lane, total ad spend, realized recovery per unit against the alternative, and predicted-versus-actual per lever	T6 plus the LTSF prediction and outcome log per LTSF Section 11.2	The basis tag HIST source for the next exit program; the recovery incentive attribution under LTSF Section 10.4
Cycle snapshot	The tagged rows inside the generated KDR Cycle File, carrying block K9 verdict, scenario ID and expected outcome	KDR Cycle File, generated Monday afternoon as a values-only snapshot; frozen by construction	Next cycle's V5 scoring; the audit trail for what was true and what was decided on cycle date

9. Skill-Conversion Block

Block	Specification
Input bindings	The LTSF row fields (archetype, tier, floor price, salvage value per unit, contribution at price and at floor, remaining units by bracket, cliff date, winning terminal option). Catalog groups: A3 pacing and in-budget, A4 sample sufficiency, A7 eligibility, B8 economics chain, D1 per-variation margin, D2 variation map, D7 decision log. Instrument cells: T2 blocks K2, K5, K6 and K7; T3 groups C1, C4 and C5; T4 "Expected CPC"; T5 Product Log; T6; T7 Manifest. An absent feed suppresses its dependent execution class per Element 3; the skill never improvises around a missing feed and never estimates a declaration field.
Deterministic logic	R2: allowable ad cost per unit = net recovery at salvage price minus the alternative's per-unit value; maximum liquidation CPC = allowable ad cost times conversion rate. R3: required units per day = remaining units divided by days to cliff; required clicks per day = required units divided by conversion rate; the traffic comparison against market CPC at required volume. R6: realized recovery per unit against the alternative, and realized pace against required pace. The Archetype D ceiling substitution at P3 step 4. Every band, limit and verdict string is read from v4.0 and the Keyword Decision Record Template v3.0 at execution time and is never recomputed here.

Block	Specification
Human checkpoints, drawn by Authority Matrix v2.0 row	Row "Terminal/LTSF PPC lane" (D: PPC Lead with the LTSF owner; E: incoming operator; R: BM owner): the retag at P2, every lane closure at P14, and the quarantine execution at P5 are human-mandatory and are confirmed by the D holders jointly. Row "Structure creation (all SOPs)": the dated liquidation campaign specification. Row "Code-red declaration + close": the Section 1.5.1 seam, where the declaration is read-only to this SOP. Row "Event plans (#32)" (D: PPC Lead with the deals owner): any liquidation window. Row "Push funding (the M5 queue)": the freed-budget destination at P6 step 3. Row "v4.0 threshold changes" (D: CEO; operator is I): every amendment item in 10.5. Standing rule 5 applies throughout: an unfilled role's D reverts to the PPC Lead and its E to the PPC Manager until filled.
Executable with audit	P3 ceiling computation, P4 pace and traffic arithmetic, the P12 weekly check line and daily pace read, ceiling recomputation on any input change, staging of bid and budget rows inside the ceiling, T6 drafting, and the F1 to F12 alarm set with evidence lines.
Cadence	Weekly at every tier. Twice weekly at CRITICAL per LTSF Section 7.1. Daily pace read inside the final 14 days before the cliff and standing at CRITICAL. Same-day execution on a new tag, an archetype change, a move into RED or CRITICAL, and a cliff-calendar change. Day-7 read on every deployed lever. The daily Brand Management and PPC huddle line whenever a rung target or expiry date is at risk.
Alarm routing	F1, F3, F9 and F12 route to the operator queue the same day. F2, F5, F6, F10 and F11 enter the weekly pass worklist. F4, F7 and F8 route to the PPC Lead and the LTSF Owner jointly with the T6 history, because each of them costs bracket rate while it stands.
Outputs	Drafted staged change sets and drafted T6 rows queued for confirmation; the R6 check line per SKU per cycle; the five-number closure handback; the F1 to F12 alarm set; the predicted-versus-actual per lever feeding the LTSF outcome log.

10. Version Governance

Owner: PPC Lead jointly with the LTSF Owner. Six review triggers, each specific to this component.

1. Any version increment of the LTSF Master Decision and Operations System. R2 and R3 are functions of that document's declared columns, so a column change invalidates the formulas rather than merely dating them. This is the trigger that fires most often and it fires this document, not just its numbers.
2. Any change to the LTSF rate card in its Section 3, including a bracket boundary move. Every cliff date, pace figure and charge comparison in Elements 4 and 6 is anchored to it.
3. Any change to the LTSF tracker specification in its Section 11.1, because Elements 2, 3 and 12 name those columns as the detecting artifacts and a renamed column is a broken trigger.
4. Any v4.0 increment, at which point Element 5's mapping re-validates against the new verdict set before any pass runs on it, and the Section 1.5.1 seam re-reads against rule T8 in case the terminal lane has been named.
5. Any Cross-Reference Ledger version that rewrites the row for component 27 or the row of any counterparty named in Element 12, because an interface contract written against a superseded row is a contract with nobody.
6. Any re-issue of the PPC Authority Matrix, because every ramp competence in Element 11 is a matrix row and a moved letter moves a competence.

10.1 Change note, v1.0 to v2.0. Supersedes v1.0 in full. Conflict C6 closes on the attachment of the LTSF Master v2.1: v1.0 correctly stopped at the boundary and invented no tier criteria, and v2.0 crosses the boundary by reading rather than by inferring. Rules R1 to R7 carry forward unchanged in substance and keep their identifiers, because SOP-32 v2.0 P12 cites R2, R3, R4, R5 and R6 by number. Procedures expand from zero to 14. The verdict mapping expands from none to the full 25-string closed vocabulary. Worked examples expand from 1 to 4. Failure modes expand from 4 to 12. Elements 11 and 12 are added. The v1.0 worked example is preserved in substance inside WE-2, which extends it with the split outcome and the Archetype C experiment consequence that v1.0 could not state.

10.2 What this version does not claim. The LTSF Master Tracker V1.1, which SOP Build Kit v1.0 Section 3 routes to this component alongside the LTSF Master, was not attached to this session. Every tracker column named in Elements 2, 3 and 12 is taken from the LTSF Master's own Section 11.1 specification rather than from the tracker file, so a column that shipped under a different label than its specification will not match. That reconciliation is owed at the next build.

10.3 Registered in the Supersession Map at publication, against component #27 of the Operations Definition Register. Absorbs the scheduled terminal-ASIN rule-set deliverable, unchanged from v1.0.

10.4 Zero local thresholds. Every decision number in this document is one of three things and nothing else: a v4.0 value quoted with its section heading, a declared field or cadence fact from the LTSF Master v2.1 quoted with its section, or an illustrative worked-example value labelled as such in Element 6. Section 1.7 states the rule and Element 6 states the base.

10.5 The v4.1 amendment list

Ten items, each a number or rule this component needs and v4.0 does not carry. None is written into the document.

ID	Statement
AM-1	v4.0 Section 8.1 "Product-Level Scenarios (S-A). The posture every keyword inherits" scenario S-A10 puts CONSTRAINED, LTSF-burdened products on a profit-only posture and routes LTSF velocity requirements to deals and pricing rather than to PPC subsidy. Rule R2 authorizes allowable ad cost per unit that can exceed the unit's own contribution. Both are correct in their own frame, and v4.0 carries no rule for a unit that is in both. R5's quarantine is the operating separation this SOP uses; v4.0 should state it or reject it.
AM-2	v4.0 Section 11 rule T8 orders Recovery Mode above Event Mode above stage defaults and does not name the terminal lane. No precedence rule exists for an LTSF RED or CRITICAL coinciding with a code-red declaration. Section 1.5.1 states the consequence without inventing the rule; WE-4 prices it at $1,090.75 per week.
AM-3	v4.0 carries no interface-lag threshold of any kind. Every timing in Elements 2 and 12 is a cadence fact from LTSF Section 9.2, LTSF Section 7.1 or Data Architecture Section 12. No rule says how late is late.
AM-4	v4.0 carries no membership threshold for the proven converter set that R4 and P9 depend on: no order count, no click count, no window. This SOP uses the gate G2 sufficiency line as the nearest governed proxy and says so at Element 3, which is a substitution rather than an answer.
AM-5	v4.0 carries no LTSF dimension at all: no bracket rate, no floor price, no salvage value, no cliff date, no archetype and no risk tier. Every terminal number in this SOP is therefore a declaration field rather than a threshold, which is workable but leaves the threshold authority silent on the largest recurring charge in the portfolio.
AM-6	Rule-identifier collision. This SOP's rules are R1 to R7 and are cited without a prefix across the corpus, including SOP-32 v2.0 P12. v4.0 Section 11 uses R1 to R6 for Recovery Mode. Both appear in Section 1.5.1 and WE-4. Section 1.6 disambiguates in-document; the corpus needs one convention.
AM-7	Escalation rules count different things and no rule reconciles them. LTSF Section 7.4 escalates on dollars per product per month and on 5 business days of blockage. The PPC library escalates on the same failure mode recurring in consecutive cycles. A terminal lane sits under both.
AM-8	Cadence rules likewise. LTSF runs weekly with twice weekly at CRITICAL plus a mandatory daily huddle. v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)" rule V4 runs daily alerts, a weekly full audit, monthly and quarterly reviews, and Recovery Mode rule R4 runs daily with twice-weekly mini-audits. No rule states which cadence governs a terminal-tagged unit on a product that is not in Recovery Mode. This SOP runs the stricter of the two per Element 9 and states that as a working choice, not as a governed one.
AM-9	LTSF Section 9.1 replaces any lever returning below half its logged prediction at day 7. v4.0 Section 14 rule V5 reviews a rule falling below 50% accuracy over 10 or more scored predictions. The two use the same fraction on different objects, one on a single lever deployment and one on a rule's history, and an operator reading them together will conflate them.
AM-10	No aggregate pass mark exists for Element 11. Component #44, operator onboarding and certification, is not built, so competence recording stays with the checkpoint holder per 11.6 and no cross-SOP certification test exists. v4.0 carries no threshold for one.

11. Operator Ramp

VERDICT: an incoming operator holds this SOP when every Authority Matrix v2.0 row naming a terminal-lane decision can be executed at that operator's letter without the checkpoint holder changing a staged value, and when the operator can state why a number is not PPC's to compute before stating what the number is. The ramp is 10 matrix rows across 3 checkpoints: week 1 observation and artifact familiarity, week 2 supervised execution with review before deployment, week 4 unsupervised execution with a post-deployment audit. Week 4 moves nothing out of the human-confirmed tier. That line is fixed by v4.0 Section 12 "Layer 7: Precedence and Conflict Resolution (P1-P10)" rule P5 and, while Recovery Mode is active, tightened further by Section 11 rule R3. The competence this ramp exists to install above all others is the refusal: this document's most common failure is a PPC operator computing a floor price, a salvage value or an archetype that belongs to another system.

11.1 The Authority Matrix rows this SOP executes

Each row of the PPC Authority Matrix v2.0 Section 1 "The Matrix" below names a decision this component executes or receives. D decides, E executes, R reviews, I is informed.

Authority Matrix v2.0 row	D holder	Incoming operator	What SOP-27 does under this row
Terminal/LTSF PPC lane	PPC Lead with the LTSF owner; BM owner reviews	E	The whole document. This is the only matrix row whose D is held jointly with a role outside the PPC department, which is the structural fact behind Section 1.2: the seam is written into the authority table, not just into this SOP.
Weekly verdicts (established US)	PPC Lead	E, assigned scope	P12 inside the Monday pass, and Element 5 execution on any tagged row carrying a verdict this cycle.
Structure creation (all SOPs)	PPC Lead, or per SOP scope	E	The dated liquidation campaign specification at P5, and the ad-group builds at P8 and P9. The BM owner is informed on this row, so structure events are reported rather than negotiated.
Code-red declaration + close	PPC Lead	E	Nothing declarative. The declaration and its close are read-only here; P13 step 5 executes the tightened confirmation line and Section 1.5.1 holds the quarantine boundary.
Event plans (#32)	PPC Lead with the deals owner	E	Supplies R2 and R3 to #32 P12 and re-runs P3 and P4 at the window conversion rate before the window opens.
Push funding (the M5 queue)	CEO above envelope; PPC Lead within	E	P6 step 3 only: freed push budget is stated and routed to the LTSF recovery plan. This SOP originates no funding decision and receives none.
Bulk deployments (#38)	No D on this row; PPC Lead reviews	E	Every staged set in this document, and the T+1 reconciliation. Matrix standing rule 2 governs: E without D executes staged decisions only.
Sibling declarations (S-series)	PPC Lead	E	The gate G11 check before any scale step on a term the liquidation lane shares with a sibling. The declaration itself routes to #33.
Automation rules (#37)	PPC Lead	E and R	Supplies the protected-scope state feed listing every dated liquidation campaign, and reads the exclusion list back against it. F9 is the failure when the read is skipped.
v4.0 threshold changes	CEO	I	Nothing, and that is the point of its presence here. Ten amendment items sit in 10.5 because this component needs numbers v4.0 does not carry. An operator who writes one of them into a reason string has failed the ramp's central competence.

Scope and vacancy. Product scope for the incoming operator comes from the Transition Brief v1.2 stack per Authority Matrix standing rule 4, not from this document. Two matrix rows carry a different D and stay outside this ramp: "Launch-stage verdicts + batches" and

"Canada weekly block", both held by the launch and CA operator. If a checkpoint holder role is unfilled at any point, standing rule 5 applies: the unfilled role's D reverts to the PPC Lead and its E to the PPC Manager, and the ramp continues against those holders rather than pausing. The LTSF Owner has no vacancy clause in either document, which is a gap this SOP records rather than fills.

11.2 Week 1: observation and artifact familiarity

Nothing deploys under the incoming operator's name in week 1. The week produces nine written artifacts. The test is not whether the operator agrees with any decision; it is whether the operator can find every number this SOP consumes and say, for each one, which system owns it.

ID	Competence	Demonstration	Evidence and reviewer
W1.1	The seam, read from the source (row: Terminal/LTSF PPC lane)	Take one live LTSF row and mark each of its fields as decided-above or executed-here against Section 1.2. Compute nothing.	Annotated row; PPC Lead and LTSF Owner jointly. A sheet that marks a floor price as executed-here is returned.
W1.2	Declaration completeness (P1)	List the seven fields P1 requires and state, for each, the execution class that stops when it is blank.	Written list; LTSF Owner.
W1.3	The two ceilings in one row (P3)	Open T2 block K6 on a tagged variation and read the standing max profitable CPC and the R2 maximum liquidation CPC side by side, stating which governs which terms on that product.	Annotated block; PPC Lead.
W1.4	The rate card and the cliff (LTSF Section 3)	For one live batch, state the current bracket rate, the next bracket rate, the date the largest bracket crosses, and the monthly charge difference the crossing creates.	Four numbers with their bracket rows; LTSF Owner.
W1.5	Archetype literacy (Section 1.3)	For each of A, B, C and D, name the lever set that opens and the one thing PPC does not do. Propose no classification.	Four-row statement; BM owner, who holds R on the lane row.
W1.6	Quarantine as a structure, not a status (R5)	Read one dated liquidation campaign end to end: name in T3 group C1, wall list in group C5, modifier state, and the #37 exclusion entry. Name one standing campaign the wall fences and one it does not.	Structure read; PPC Manager, shared with the #29 owner.
W1.7	Authority literacy	For each of the ten rows in 11.1, state who holds D in this operator's scope and what standing rule 5 does if that holder is absent.	Row-by-row statement; PPC Lead.
W1.8	The pace math as a function of declared inputs (P4)	Recompute one live R3 pace figure from remaining units and days to cliff, then state which five inputs would invalidate it.	Recomputation with the five inputs named; PPC Lead.
W1.9	The check line (P12)	Sit one weekly R6 check and one daily pace read on a live tag, and state from the four numbers whether the lane is ahead of or behind its required pace.	Two dated notes; PPC Manager.

Week 1 closes when the checkpoint holders have recorded all nine as demonstrated. A competence not demonstrated repeats alongside the week 2 work; it does not block the ramp and it does not quietly disappear.

11.3 Week 2: supervised execution

The operator stages; the checkpoint holder reviews every staged row before it reaches #38. The supervision is on the artifact, not on the person: the reviewer reads the staged row and its reason string, and the reason string either carries a verdict and a scenario ID or it fails staging.

ID	Competence	Supervised task	The review test applied
W2.1	Same-week quarantine (P2, P5, P6)	Take one new terminal tag and stage the full sequence: retag, campaign specification, walls, shut-off, exclusion entry.	Every element present in one week, in order. A quarantine staged without its wall list fails, because a quarantine with no wall is a status change wearing a structure's name.

ID	Competence	Supervised task	The review test applied
W2.2	Ceiling computation (P3)	Compute the R2 ceiling on two SKUs, one terminal and one Archetype D, using the correct source for each.	The terminal SKU uses the salvage subtraction; the Archetype D SKU uses contribution above the floor. Using one method for both fails even where the arithmetic is clean.
W2.3	The traffic check and its three outcomes (P4)	Run the check on three live SKUs and route each to open, split or closed.	Market CPC read at the required volume, not at current volume. A closure routed as an escalation rather than as a finding fails.
W2.4	Archetype B direction (P8)	Stage one re-point toward an aged child against the D2 map.	The advertised SKU is the aged child, the bid is computed on the child's own economics, and nothing on the parent or the healthy children is touched.
W2.5	The Archetype C refusal (P9, F7)	Take one Archetype C tag and produce the converter set plus the list of levers not opened, with the charge cost of one 7-day read window stated.	The cost figure is computed from the batch's own bracket rate and unit count. A list of levers not opened with no number attached fails.
W2.6	Prediction and basis tag (LTSF Section 8.2)	Write the prediction for every lever staged this week: expected uplift, basis tag HIST, MARKET or TEST, and the day-7 read date.	Every lever tagged. An untagged multiplier renders as blank in the tracker and blocks the scenario, so an untagged prediction is not a weak prediction, it is no prediction.
W2.7	The five-number closure (P14, F11)	Write one lane-closure handback.	All five numbers present. A handback stating that PPC cannot clear the units, without the residual count and the ceiling, is returned unread.
W2.8	The refusal (row: v4.0 threshold changes)	Identify one request in the week that would place a number with no v4.0 or LTSF basis into this SOP or into a reason string, and route it as an amendment item instead of executing it.	The routed item names what the decision needs and why neither document answers it. This competence is demonstrated by declining, and an operator who declines nothing in week 2 has not demonstrated it.

Week 2 closes when the checkpoint holder's review changed no staged value on the final cycle of the week.

11.4 Week 4: unsupervised execution under audit

Week 3 repeats week 2 across the full assigned scope. From week 4 the operator deploys within scope without pre-review, with three permanent exceptions and one audit.

1. The P5 tier never becomes unsupervised. v4.0 Section 12 rule P5 places every 500-plus search-volume keyword action, every gate failure, every structural change, every bid move above 25%, every budget move above $50 per day and every low-confidence verdict in the human-confirmed tier permanently. Recovery Mode tightens it further under rule R3.
2. The Element 9 human-mandatory list stands. The retag, the quarantine execution and every lane closure are confirmed by the PPC Lead and the LTSF Owner jointly before deployment.
3. Declaration fields are never computed, at any tenure. This is not a review tier; it is a scope boundary, and no amount of demonstrated competence moves it.

ID	Competence	Audit evidence the checkpoint holder reads
W4.1	Ownership of the weekly lane for the assigned scope: declaration through ceiling through pace through staging through T+1.	The cycle's staged set, its T+1 verification report and its T6 rows read as one chain. A break anywhere in the chain is the finding.
W4.2	Routing every human-confirmed item to its holder before deployment.	The confirmation record against the deployed set. A P5-tier or Element 9 item deployed unconfirmed returns the scope to supervised execution the same day.
W4.3	Writing the R6 check line every cycle, including cycles where nothing changed.	One line per tagged SKU per cycle with all four numbers. A missing line on a quiet week is the failure, because a quiet week is exactly when a lane drifts past its expiry.
W4.4	Re-deriving the ceiling and the pace whenever a declared input moved.	The date on the K6 ceiling against the dates on the five inputs. F6 is the finding when the ceiling is older than its inputs.

ID	Competence	Audit evidence the checkpoint holder reads
W4.5	Holding the boundary: naming the owning system for every decision this SOP does not make.	Every item routed out this cycle names the LTSF Master by section or the counterparty component by number, with no restatement of that system's procedure.
W4.6	Producing the F1 to F12 alarm set with evidence lines and routing each per the Element 9 tiers.	Alarm list with row, signal value and rule reference on each line; routing matches the tier.

Audit findings return as competence gaps, not as reversals. A staged value the checkpoint holder would have chosen differently is a conversation; a computed declaration field or a missed confirmation is a stop.

11.5 Live-defense questions

Twelve questions, asked cold at the end of week 4. None asks for a threshold's value, and an answer reciting a number without its reasoning does not pass. Sources are named so the reviewer checks the answer rather than judging it.

Question	What a passing answer contains, and its source
Q1. A liquidation bid is authorized above the unit's own contribution. Why is that not a subsidy?	Because the comparison is against the alternative, which is negative: a charge avoided is value gained, so allowable ad cost is net recovery at salvage minus the alternative's per-unit value. An ad that loses against price can still beat the alternative. Source: rule R2 and LTSF Section 4.2.
Q2. The lane cannot buy the required clicks at the ceiling. Why is that a correct outcome rather than a PPC failure?	Because the arithmetic is the finding the deciding system needs. Closing the lane routes the choice to a price move or the terminal option, both of which recover more than spending above the ceiling would. Source: rule R3 and P14; worked at WE-2.
Q3. Why does sunk COGS appear in no column here, when it is real money already spent?	Because it is identical across every scenario for the same units, so it can never change which option wins, and displaying it makes every option on an underwater SKU look like a loss. Source: rule R7 and LTSF Principle 2.
Q4. An Archetype C SKU has a small converter set. Why is a discovery campaign the wrong answer?	Because the read window equals or exceeds the 7-day decision window, so the test cannot inform the decision it was opened for, and it costs the full bracket rate while it runs. Source: Section 1.3.1 and LTSF Section 5; failure mode F7.
Q5. Break-even PPC scaling doubles the authorized click price on an Archetype D batch. What makes that legal when the standing posture forbids it?	Because the ceiling being used is the max profitable CPC rather than the target line, and the objective has changed from efficiency to speed against a dated charge. The guard is the floor price: the realized net price stays above it at every step. Source: v4.0 Section 3 family F1 and Section 10 rule M1; LTSF Section 6.2.
Q6. A code red is declared on a product carrying a live liquidation lane. Why is the lane not the first thing cut?	Because the cut improves the headline reading and destroys cash: spend falls, contribution falls further, and the held units keep accruing. The quarantine is what keeps the cut sequence off the lane. Source: Section 1.5.1; priced at WE-4.
Q7. The tier moved from CRITICAL to RED. Why does the quarantine not relax?	Because R5 quarantines the tag, not the tier. A SKU that is still terminal still poisons both reads if it re-enters standing campaigns. Source: rule R5; failure mode F12.
Q8. A deal window opens on a tagged variation and the lane had been closed. Why might it now open?	Because the deal price raises conversion, which raises the maximum liquidation CPC proportionally, since that ceiling is allowable ad cost times conversion. The ceiling is re-run before the window, not after. Source: rule R2 and P11 step 4; the window plan is #32's.
Q9. You believe the declared floor price is too high and the lane would clear at a lower one. What is the correct action?	Route the finding with the five closure numbers and let the LTSF system move the price. PPC does not compute a floor and does not price against one it has not been given. Source: Section 1.2 and P14.
Q10. Why does the weekly check line get written on a week when nothing changed?	Because the check is the expiry contract: every terminal decision carries its expiry from the cliff calendar, and a quiet week is when a lane drifts past it unnoticed. Unscored actions also do not accumulate. Source: rule R6; v4.0 Section 14 rule V1.
Q11. The inventory gate reads GREEN on a variation carrying 624 aged units. Is the gate wrong?	No. That gate measures stockout risk and an overstocked variation is correctly not at risk of stockout. It is simply silent on the opposite failure, which is why the LTSF declaration is a separate input rather than a band on the same column. Source: v4.0 Section 5 gate G6; Section 1.5 of this SOP.
Q12. Your lane needs a number neither v4.0 nor the LTSF Master carries. What do you do with it?	Write it nowhere. Name what the decision needs, state that neither document answers it, and file it as an amendment item for the threshold owner; the row waits or executes the part that is governed. Source: v4.0 Section 14 rule V3; Authority Matrix row "v4.0 threshold changes", where the operator is informed and the CEO decides.

11.6 What this ramp cannot certify

Component #44, operator onboarding and certification, is not built. Until it exists there is no cross-SOP certification test and no pass mark, and neither v4.0 nor the LTSF Master carries a threshold for one. Two consequences hold. First, the checkpoint holders record each competence in 11.2 to 11.4 as demonstrated or not, per row, and that record is the certification; the aggregate pass mark is amendment item AM-10. Second, this section is written to be consumed whole by #44 when it lands, which is why every competence traces to a matrix row rather than to a house preference. One competence here has no counterpart in any other SOP's ramp: W1.1 and W4.5, the boundary against a system outside the PPC register, which #44 should carry forward as its own class.

12. Interface Contracts

> VERDICT: every handoff into and out of the terminal lane is an artifact with a named sending owner, a named receiving owner, a stated time and a stated consequence when it does not arrive. Eighteen contracts follow, 9 upstream and 9 downstream, drawn from the Cross-Reference Ledger row for component 27 and the rows of every counterparty naming #27, with artifact names taken from Master PPC Data Architecture v1.0 and, for the deciding system's own fields, from LTSF Section 11.1. The non-arrival rule is uniform: a handoff that does not arrive suppresses its dependent execution class and is recorded. It never lowers the bar for the decision and it never authorizes a substitute number. On this lane non-arrival also has a running cost, because the bracket rate accrues while the handoff is missing, so every consequence below states what is charged as well as what stops.

12.1 The five standing contract rules

ID	Rule
C1	The artifact is the interface. A verbal tag, a message or a recollection does not start a receiving procedure here. If the LTSF row does not carry the field, the handoff has not arrived, however confident both parties are.
C2	Non-arrival is recorded, not absorbed. The affected field group is marked in T7 Manifest with its status and the verdict classes it blocks, per v4.0 Section 5 gate G8.
C3	Suppression, never substitution. A missing declaration field suppresses its execution class. It does not license an estimated field, a smaller version of the action, or a value carried from a prior cycle without its date.
C4	Cite the Ledger row, not the counterparty's procedure number. Components rebuilt in a later wave move their procedure numbers, so each contract names its counterparty by component number and the Ledger row that declares its ownership. Where a procedure number appears it is this SOP's own or, for #32 and #12, one this session read directly in the attached document.
C5	No contract carries a maximum lag in hours or days, because v4.0 carries no interface-lag threshold. Times stated are cadence facts from LTSF Section 9.2, LTSF Section 7.1 and Data Architecture Section 12. A lag limit is amendment item AM-3.

12.2 Upstream contracts: what must arrive before this SOP acts

ID	Artifact and where it lands	Sending to receiving owner	Timing	When it does not arrive
I-1	The LTSF declaration: archetype, tier, floor price, salvage value per unit, remaining units by bracket, cliff date and the winning terminal option, in the LTSF row per LTSF Section 11.1 amendments 2, 3, 6 and 7	LTSF Owner to PPC operator	On classification, and on any weekly re-classification	P1 stops and nothing executes. The row returns naming the missing field. PPC computes no substitute, so the units continue accruing their bracket rate with no lane at all: on a 624-unit batch at $5.45 that is $113.36 per day of charge against a blocked decision, which is why LTSF Section 7.4 escalates blockage at 5 business days.

ID	Artifact and where it lands	Sending to receiving owner	Timing	When it does not arrive
I-2	Contribution at the current price and contribution at the floor price, per unit, for the Archetype D ceiling	LTSF Owner with Finance to PPC operator	With the declaration	P3 step 4 cannot run and break-even scaling does not open. The lane holds at the standing target line, which clears at the standing velocity and forfeits the pace. PPC does not model fees to fill the gap.
I-3	Cliff-calendar update: units per age bracket, the date the largest bracket crosses, and the weekly velocity-to-cliff requirement	Data Operations to PPC operator, via the LTSF row	Monday, with the LTSF weekly cycle refresh	The R3 pace holds at its last dated value and is marked stale in T6. A pace figure carried against a unit count that has fallen overstates the requirement and overspends; carried against a cliff date that has moved, it understates it and misses the date. F6 is the finding.
I-4	Refreshed product workbook and Raw Pull Pack: A3 pacing, A7 eligibility, A4 sufficiency, and the search-term joins feeding the converter set	Data Operations to PPC operator	Pack Monday 06:00; refresh complete 09:00	T7 marks the group stale or missing and dependent verdicts downgrade to flags per gate G8. The R6 check line still writes, using the last clean read with its date stated, because the expiry contract does not pause for a feed.
I-5	Per-variation margin and the D2 variation map with the aged children named	#41 owner and Brand Management to PPC operator	On the margin event and monthly	P8 does not start: Archetype B is unexecutable without knowing which child holds the units. The row carries the D2-pending flag and the aged children keep receiving no targeted spend, which is the state that aged them.
I-6	Armed-window declaration with its dates, entering T2 with the event flag	#32 owner to PPC operator	Before the window opens	The window is unrecorded and the ceiling is not re-run at the window conversion rate, so a lane that would have opened at the deal price stays closed and the units miss the one period in which the arithmetic worked. #32 P12 also cannot build its cap.
I-7	Code-red declaration record and the tightened confirmation line	#26 owner to PPC operator	Same day as the declaration	Steps in the lane deploy at the standard confirmation tier while Recovery Mode is active, which is a governance breach rather than an execution error. P13 step 5 cannot execute and Section 1.5.1's boundary is undefended against the cut sequence.
I-8	Declared family term owner where the lane buys a term a sibling also runs	#33 owner to PPC operator	Before any scale step on a shared term	Gate G11 stays unresolved and no scale step deploys on that term. The lane runs at reduced pace on its remaining terms and the shortfall is stated against the R3 requirement rather than absorbed.

ID	Artifact and where it lands	Sending to receiving owner	Timing	When it does not arrive
I-9	The #38 T+1 verification report on every staged set	#38 owner to PPC operator	T plus one after every deployment, including same-day deploys	Landed values are unconfirmed, so every read on the affected rows is void for the week and no verdict issues on them. F9 cannot fire without this report, which makes its absence more expensive than a mismatch inside it: an automation rule reaching a liquidation campaign would go undetected.

12.3 Downstream contracts: what this SOP owes, and to whom

ID	Artifact and where it lands	Sending to receiving owner	Timing	When it does not arrive
I-10	The R2 allowable ad cost per unit and maximum liquidation CPC, and the R3 required units per day, in T2 block K6	PPC operator to #32 owner	Before any liquidation window is planned	#32 P12 cannot build the window cap, which is the R2 ceiling expressed daily times the R3 required pace. The window then sizes from the standing lift index instead, which is the one sizing method its own P12 step 3 forbids on this class.
I-11	The five-number lane-closure handback: required clicks per day, the R2 ceiling, market CPC at required volume, units clearable, residual	PPC operator to LTSF Owner	Same working session as the P4 closure	The terminal option is not executed and the price is not moved, so the full pool keeps accruing at the bracket rate while the decision waits. On the WE-2 batch that is $110.40 per day. This is the most expensive non-arrival in the set because it stops the deciding system rather than the executing one.
I-12	Protected-scope state feed: every dated liquidation campaign, for the automation exclusion list	PPC operator to #37 owner	Same session as the quarantine, and on every archive	An automation rule ladders a liquidation bid toward a standing efficiency ceiling that does not apply and closes the lane silently. Nothing alarms, because the rule is doing what it was written to do. F9. Ledger dependency D6 reads OPEN, so the manual read is the only control.
I-13	The dated liquidation campaign specification: name carrying the cliff date, advertised child, converter ad group, modifier-zero state	PPC operator to #12 owner	Same week as the tag, per R5	The campaign is not created and the tagged units keep serving inside standing campaigns. Both reads are then mixed and neither the standing product's economics nor the recovery math is usable until a fresh baseline is set. F3.
I-14	The wall list on the new T3 row, as the wall-audit baseline	PPC operator to #29 owner	Same working session as the structure event	The wall audit has no baseline for the quarantine and cannot distinguish a missing wall from an unrecorded one, so every later audit on those terms produces a finding that costs a session to disprove.

ID	Artifact and where it lands	Sending to receiving owner	Timing	When it does not arrive
I-15	Freed push budget from the R4 shut-off, with the amount and its destination stated	PPC operator to the LTSF recovery plan, via the LTSF Owner	Same cycle as the shut-off	The dollars sit unallocated: they do not fund the recovery and they do not return to the push queue either, so the envelope reads healthy while nothing is buying anything. This is the quietest failure in the set because nothing overspends and nothing alarms.
I-16	The R6 weekly check line and the daily pace read	PPC operator to LTSF Owner and the daily huddle	Weekly on the Monday pass; daily inside the final 14 days and standing at CRITICAL	The rung target and the expiry date go unwatched, and the ladder's next rung does not fire on the date it was written for. LTSF Section 9.3 makes the rung fire automatic on a missed target, and this line is the evidence that the target was missed.
I-17	Predicted-versus-actual per lever, with the basis tag, into the prediction and outcome log	PPC operator to LTSF Owner and Data Operations	At the day-7 read and at closure	The next exit program has no HIST basis for any PPC lever, so its uplift estimates fall back to MARKET or TEST and it pays for evidence this program already bought. Incentive attribution under LTSF Section 10.4 also has no PPC record, since an unlogged action earns no credit.
I-18	Section 11 of this document as the ramp content for the certification test	PPC Lead to #44 owner	On the Ledger due date for #44	No cross-SOP certification exists. Competence recording stays with the checkpoint holders per 11.6 and the aggregate pass mark stays amendment item AM-10 rather than a number invented here.

12.4 Repeated non-arrival, and the contracts still owed

A single non-arrival is a logged suppression. The same interface failing twice in consecutive cycles escalates to the PPC Lead and the LTSF Owner jointly with the T6 history attached, and is carried to the counterparty as a Ledger-row disagreement rather than settled between operators: a session never edits another component's row. On this lane the escalation also carries the accrued charge for the blocked period, because that is the number that makes the delay concrete.

Owed	Statement
Interface lag limits	Every timing in 12.2 and 12.3 is a cadence fact, not a threshold. Neither v4.0 nor the LTSF Master says how late is late. Amendment item AM-3.
The LTSF Master Tracker V1.1	Routed to this component by SOP Build Kit v1.0 Section 3 and not attached to this session. Every tracker column named above comes from the LTSF Master's Section 11.1 specification rather than from the tracker file, so a column shipped under a different label will not match. Reconciliation owed at the next build.
A vacancy clause for the LTSF Owner	Authority Matrix standing rule 5 reverts an unfilled PPC role's D to the PPC Lead. The lane row's D is held jointly with the LTSF Owner, and neither document says what happens when that seat is empty. Recorded here, not filled here.
Counterparty procedure numbers	Contracts I-5, I-7 and I-8 cite Ledger rows rather than counterparty procedure numbers, per rule C4, because those components sit at versions later than the copies routed to this session. I-6, I-10 and I-13 cite procedures this session read directly in the attached SOP-32 v2.0 and SOP-12 v2.1.

Companions: LTSF Master Decision and Operations System v2.1 (the deciding system: archetypes, tiers, floor price, salvage comparison, lever inventory, cadence) · Master PPC Decision Criteria System v4.0 (all thresholds, gates, scenarios and verdicts) · Keyword Decision Record Template v3.0 (the closed verdict vocabulary) · PPC Strategic Doctrine v1.0 (the why layer) · PPC

Economics Worked-Example Set v1.0 case E7 (terminal arithmetic convention) · SOP-32 Deal and Event Window Operations v2.0 P12, SOP-12 Exact Campaign Operations v2.1, SOP-26 Mature and Defend v2.0, SOP-29, SOP-33, SOP-37, SOP-38, SOP-28 (boundaries) · Master PPC Data Architecture v1.0 and the PPC Command Workbook template (instrument mechanics) · PPC Authority Matrix v2.0 (Element 11) · PPC Cross-Reference Ledger, row for component 27 (Element 12) · Hanging Closet Organizer Full Liquidation Action Plan v1.0 (a worked application of this lane at program scale).

Inspiratek & Ecotero LLC | Confidential | SOP-27 v2.0