

PPC DEPARTMENT · OPERATING SOP · COMPONENT # 28 OF THE OPERATIONS DEFINITION REGISTER

SOP-28: Placement and Modifier Operations

Version 2.0 · 2026-08-08 · Owner: PPC Lead · Supersedes SOP-28 v1.0 (2026-07-24) in full · Every threshold, band, step size, limit and verdict lives in the Master PPC Decision

Criteria System v4.0 and the Keyword Decision Record Template v3.0; this SOP carries zero local thresholds and owns the mechanics those two documents' decisions execute through.

1. Verdict and Scope

VERDICT: the three placements are three products bought with one bid, and the modifier is the only instrument that buys them separately. Two rules bound

every operation in this document and neither bends: no modifier rises without placement-level conversion proof, and no modifier changes without the ceiling

worksheet. This SOP owns the mechanics that eleven other components' decisions execute through, which makes an error here a distributed error: SOP-12 P4,

SOP-14 P4, SOP-22, SOP-23 P10, SOP-24, SOP-25, SOP-26 P3 and P8, SOP-30, SOP-35, SOP-37 and SOP-38 all read a rung value, an exit bid, a floor or a

modifier from this document and none of them derives one.

What changed at v2.0. v1.0 was built to the superseded 10-element standard and measured 4 pages against a 10-18 page bar. v2.0 expands procedures from 8 to 14, failure

modes from 7 to 12, worked examples from 1 to 4; adds Element 11 operator ramp and Element 12 interface contracts, neither of which exists in any of the 23 shipped SOPs;

maps the full 25-string closed vocabulary plus the campaign subset in place of v1.0's 7-row table, of which six strings belonged to no closed vocabulary and are retired at

Section 5.2; and makes two corrections of substance. First, the v1.0 ladder worked example runs a first rung of 5.769% where v4.0 Section 10 rule M3 sets the taper step at

roughly 10% per week, and the format exemplar SOP-12 v2.0 already corrected the same ladder against this document; v2.0 adopts the compliant ladder and Section 6 WE-2

shows it. Second, three numerals v1.0 attributed to v4.0 are not in v4.0 at all: the modifier proof ratio, the per-placement traffic minimum, and the modifier step size. All three are

retired as ungoverned and logged in the Section 13 amendment register rather than carried. All seven v1.0 account-history incidents carry forward, now sourced rather than asserted.

1.1 The dependency that shapes this document, stated first

External dependency D3 in the Cross-Reference Ledger is the placement feed, and it is OPEN. The Ledger's own entry reads that the performance sections of SOP-28 stay

structural until it lands. The live T7 Manifest in the Command Workbook TEMPLATE v1.0 confirms it in-row: field group B7 Placement splits (TOS/ROS/PP) , source feed

TOS placement report , status MISSING , verdict classes blocked All T4 placement verdicts . A second group compounds it: B6 Bid intelligence (suggested

bid + gap) also reads MISSING , blocking Bid-gap sizing on ladders . Packet completeness stamps at 0.571, below the 85% DEGRADED floor of v4.0 Section 5 gate

G8.

The operating consequence is asymmetric and it is the single most important sentence an operator takes from this document. Modifier increases are suppressed account-

wide while B7 reads MISSING. Modifier cuts and normalizations remain legal throughout, because removing an unproven premium never needs proof. The ladder,

the ceiling worksheet, the exit-bid arithmetic, the floor register and the bleeding-placement lever order all run on base-bid and blended-CPC data that survives a B7 outage, so

this document is fully executable on its cut side and structurally executable on its raise side. Strategic Doctrine v1.0 Section 15 item 5 states the same posture from the doctrine

side: the placement doctrine activates fully when the placement report feed lands, and until then no modifier increases on unproven placements.

This is why brevity would be the wrong answer here rather than the honest one. The half of this component that a missing feed suppresses is one procedure of fourteen. The

half that eleven other components consume every week is unaffected and has never been written to depth.

1.2 The four decision surfaces this SOP owns

Surface What it decides Where the numbers come from

The ceiling worksheet Whether a proposed base bid and modifier pair is legal, computed before Effective CPC arithmetic per Strategic Doctrine v1.0 Section 6.1; the

the change is staged rather than after it is deployed. governing ceiling per v4.0 Section 10 rule M4, which sizes the multiplier so effective TOS CPC is at or under the minimum of max profitable CPC for the objective and Expected CPC plus tolerance.

The ladder The rung value, the rung date, the hold prediction, the reversion rule and v4.0 Section 10 rule M3 sets the step at roughly 10% per week; v4.0 Section

the point at which descent stops. 9 rule L2 sets the per-change limit at 50% and the bid floor at \$0.50.

The exit bid and the floor register Where the ladder is aiming and what it discovered when it stopped. v4.0 Section 3 family F9, blended placement CPC row: to leave an

underperforming placement, bid \$0.05-0.10 below its blended CPC. v4.0

Section 9 rule L2 carries the same decrement.

The bleeding-placement lever order Which lever moves first when one placement carries a term's failure and v4.0 Section 9 rule L1 lever precedence, with the placement decomposition

the others hold. classes of v4.0 Section 3 family F9 and the scenarios at Section 8.8 S-P5 and S-P6.

1.3 The three placement decomposition classes, and why blended reads cannot produce them

v4.0 Section 3 family F9 resolves every placement read to one of three verdicts, and the classes are per placement per term, never per campaign. TOS-EARNS: top-of-search

CVR materially above rest-of-search with top-of-search CPA at or under the ceiling, which justifies the multiplier up to the CPA ceiling.

TOS-VANITY: top-of-search CVR at or

below rest-of-search, which cuts the multiplier. PP-LEAK: product pages absorbing spend below product-page CVR norms, which exits the placement through the blended-CPC

mechanic. Family F9 states the reason the decomposition is mandatory in one line: the blended line hides that a term can be profitable at rest-of-search and ruinous at top-of-

search or the reverse. Section 6 WE-3 works a live instance where the blended ACOS reads 38.47% and looks survivable while one placement inside it runs at 147.92% and

carries a weekly loss of \$140.92.

AMENDMENT. Family F9 sets TOS-EARNS on top-of-search CVR being materially above rest-of-search, and v4.0 Section 9 rule L3 repeats the word without a number. No

factor exists anywhere in v4.0. SOP-28 v1.0 P3 asserted that a raise requires the CVR proof ratio to clear the v4.0 factor; there is no such factor to clear. This is amendment

item AM-1 and it is the highest-priority item in the Section 13 register, because it is the test the document's first absolute rule depends on.

1.4 Boundaries, with the owning component named for every excluded decision

Per the Cross-Reference Ledger row for component 28, this SOP owns ladder execution mechanics, arbitrage setup, the ceiling worksheet, bleeding-placement response, exit-

bid discovery, modifier sizing and the split-exact test. Everything below is referenced by number and one clause and is never restated here.

Excluded decision Owner The seam, stated so it is owned

Every threshold, band, step size, limit, v4.0 (#8) A numeral appearing in this SOP without a v4.0 section beside it is a defect. Section 13 lists the sixteen numerals

scenario and verdict issuance v4.0 does not yet carry; they are amendment items, never local values.

Campaign-level bid, budget, status and SOP-12 This SOP derives the rung value and the compliant base-and-modifier pair; SOP-12 P4 and P11 execute them. A

structure changes, including taper execution campaign edit made from this document is a defect.

and duplicate consolidation

Whether a taper happens, and when SOP-12 P4, #23, #24, #25, #26 The Ledger states it exactly: taper DECISIONS belong to those components and only the taper MECHANICS are

owned here. This SOP never decides that a term should descend; it decides how far each rung falls and where the descent stops.

#23

[page 2 of 18]

Push sizing, the authorized CPC envelope, This SOP receives the achievement date, the push bid and the authorized CPC envelope at taper release, and

achievement and stability confirmation tests proposed in-flight modifiers against that envelope. It supplies no push value.

The descent schedule, its dates and its #25 Per Ledger row 25 that component schedules rungs and consumes floors and derives no step. This SOP returns

sequencing the step value for each scheduled date.

Which cohorts re-ladder and when, and #26 Per Ledger row 26 that component issues the re-ladder and bleed decisions and this SOP supplies the

which placement bleeds get cut inside a mechanics and the step values. The floor register format is this component's and the register content is read by

code red both.

Incrementality test design and the continue- #35 Where v4.0 Section 4 rule X12 makes a rung double as an incrementality step, this SOP stages the rung and

or-restore call records the floor; the capture read and the call belong there.

Negation walls at every structure event, #29 A split-exact standup builds its walls in the same working session under #29 rules. This SOP names the

including on a placement split requirement and never the wall level.

Staging format, upload, the effective-CPC #38 Every rung and every modifier in this document is a staged, reasoned deployment. Per Ledger row 38 the P6

pre-upload scan and the T+1 verification diff scan there is an arithmetic check against this SOP's worksheet, never a re-decision.

The automation rule register and rule #37 Per Ledger row 37 the ladder holds are this component's. This SOP issues the hold notice naming which rows

change control are under a live ladder; it never edits a rule.

Family term ownership on a shared term #33 Gate G11 bars a restore step on a shared term until the lead is declared. This SOP holds; it never adjudicates

ownership.

The event plan and the armed window #32 Gate G10 excludes event weeks and the 14 days following from the trend reads a probe depends on. This SOP

reads the armed state; it does not set it.

The non-SP format build on ENTER SBV #17-19 This SOP supplies only the paired SP-exact single ladder step down; per Ledger row 17 the SBV go-live step-

ARBITRAGE down handoff routes through #12.

Budget envelope, the M5 queue and every Instrument layer Per Ledger conflict C3 this SOP names the instrument and the tab, T0 zone H and Account Rollup tab A2, and

budget release never cites a #31 procedure number.

Dayparting #45 v4.0 Section 9 rule L10 defines the lever at plus or minus 15% hourly modifiers on Marketing Stream curves with 4 or more weeks of hour-level evidence per campaign. T7 reads D12 Hourly dayparting MISSING and explicitly deferred. This SOP builds no hourly modifier and the reserved columns stay unbuilt.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. Where v4.0 states the lag, the section is cited. Where it does not, the

cell says so and the item appears in Section 13; no lag is invented here.

Trigger Detecting artifact and the observable Maximum lag to procedure start Procedure

Weekly cadence The Monday cascade reaches T4 Placement & Bid with One weekly cycle. v4.0 Section 14 rule V4 sets the weekly full P1

the T2 OVERSIGHT filter applied. audit as the cadence producing Master Action Records.

Modifier change request Any verdict from Section 5 proposing a modifier move, Same cycle. No modifier change deploys without the P2 P2, then P3 or P4

arriving in T2 block K9 or from an adjacent component. worksheet, so the worksheet is the gate rather than the clock.

Taper release TAPER: RANK ACHIEVED, STEP BIDS DOWN issued Same day as the release for the first rung derivation. Not P5, P6 anywhere in the system, arriving with the achievement stated in v4.0 as an interval. Amendment item AM-10.

date, the push bid and the authorized CPC envelope.

Scheduled descent rung A dated rung on the #25 descent schedule reaching its The scheduled date. v4.0 Section 10 rule M3 sets the step P5

date in T4 Next Step Date . interval at roughly one week between rungs.

Quarterly floor re-test A cohort reaching its re-test date, read from T4 One weekly cycle. v4.0 Section 10 rule M3 requires the floor P7

Discovered Floor \$ and Next Step Date . re-tested quarterly or after a logged competitive shock; Section

14 rule V4 places floor re-tests on the quarterly cycle.

Competitive shock on a defended term T5 Product Log Event Type reading competitor entry, One weekly cycle. Rule M3 names the logged competitive P7

cross-read against the T2 K8 competitor block. shock as a re-test trigger equal to the calendar.

Bleeding placement One placement failing its own band while the others hold, Not stated in v4.0. The weekly pass is the containing cadence P8

read on T4 Placement Verdict ; v4.0 Section 8.8 S- under V4. Amendment item AM-10.

P5 on a TOS-VANITY read, S-P6 on a PP-LEAK read.

Top-of-search share below floor T2 TOS IS % reading below 30 on a campaign carrying Same weekly read. S-R7 places the placement problem ahead P9

a Ranking objective; v4.0 Section 8.4 S-R7. of any bid verdict, so the lag is the pass itself.

Split candidacy Divergent placement economics proven on one term Not stated in v4.0, and the persistence bar carries no window P10 across qualified reads on T4, persistent rather than count. Amendment item AM-15.

single-window.

Bidding-strategy audit Any campaign carrying up-and-down bidding, read from Scheduled onto the governed quarterly cycle of v4.0 Section P11

T3 Bidding Strategy beside T3 14 rule V4 alongside the M3 floor re-tests. v4.0 states no

Placement Mult TOS % . cadence of its own for this audit; the ride-along is stated rather

than asserted as governed.

Placement drift T4 Drift Flag firing on a placement mix shift above Same weekly read. Family F9 routes drift to structural P12

15 points over 4 weeks with no logged change, per v4.0 inspection before performance explanations are accepted.

Section 3 family F9.

Margin recompute T0 zone C margin table version changing, which fires Same cycle. S-A9 puts every keyword whose CPC is above its P5, P2

v4.0 Section 8.1 S-A9. new ceiling into the ladder queue this cycle.

Feed state change T7 Manifest B7 or B6 status changing from MISSING. Same cycle. The suppression is lifted by a manifest read, not P13

by a memory.

Rung review date A T6 Decision Log row carrying a rung prediction The named review date. v4.0 Section 14 rule V1 makes the P14 reaching its Review Date . next cycle's first step the scoring of the last.

3. Preconditions and Gates

Nothing in Section 4 runs until the manifest rows below read current on T7. Each row names the manifest field group by its Catalog ID, the gate it feeds by its v4.0 number, and

the observable that proves the gate passed. Where the operating corpus uses a different gate number for the same idea, the corpus alias is given at first use per Ledger conflict

C9; the v4.0 number governs, and the two schemes are never mixed inside one reason string.

3.1 Live manifest state, and exactly what it costs this component

Five field groups this SOP depends on are not fully present, and their effects are not equivalent. B7 placement splits is MISSING, which makes every T4 performance

column unreadable and suppresses every modifier increase account-wide; it does not touch the ladder, which runs on base bid and blended CPC. B6 bid intelligence is

MISSING, which removes the suggested-bid reference and makes the family F9 auction-moved read, gap widening at or above 25% over 4 weeks, uncomputable. A6

objective and per-term Expected CPC is PARTIAL with an account reference of \$2.80 in use, which means the M4 minimum is computed against an account value rather

than a term value and the substitution is written onto the record. D1 per-variation margin is PARTIAL, carrying the G9 PROVISIONAL tag onto every ceiling in this

document. A8 PAT rows first-class is PARTIAL, which T7 records as blocking PAT ladder verdicts, so a rung on a PAT row is derived and staged with the coverage gap

named. Under G8 a verdict whose gate fields are missing downgrades to a flag and the analysis states what it could not see; P1 step 6 and P13 enforce that statement on the

read itself, and v4.0 Section 13.3 rule J7 makes hiding a fallback a process violation rather than a style choice.

Manifest ID Gate Requirement before any procedure runs The observable that proves it passed

A4, A5 G1 No impossible-metric flag on the row: no CVR of zero with orders present, no T2 Flag: Zero-Quarantine clear and Validation State not

ACOS without orders, no cross-table disagreement. quarantined. A quarantined row produces a flag and a data-refresh task, never a rung and never a modifier.

A4, A1 G2 Sample sufficiency for every rate the read consumes: 15 or more clicks in 7 T2 Clicks 7/30/60d and Sample State .

AMENDMENT. G2's

days for a bid read, 100 or more clicks for a CVR verdict, 1,000 or more minimums are stated at keyword and campaign grain. v4.0 carries no per-

impressions for a CTR verdict. Below threshold the window falls back to 60 placement minimum, and the UNREAD class this SOP operates therefore has

days; if still thin the verdict is EXTEND TEST. no governed bar. Amendment item AM-2. A1 reads MISSING, so four-period trend context is unavailable and single-window reads carry the flag.

A3 G3 The hosting campaign is not budget-truncated. In-budget below 70% of the day T3 Truncation Flag (G3) clear and In-Budget Bucket above 70%.

marks every rate on that campaign TRUNCATED and forces the budget lever Corpus alias: the KDR Bands sheet labels 51-70 strained; G3's 70% line

before any bid or placement conclusion. governs. A truncated week's placement read is void, not conservative: a starved campaign's placement mix describes its worst hours.

A7 G4 The advertised offer is eligible and unsuppressed. T2 Eligibility reads eligible. A flip halts all execution for the target the same day and opens the listing task.

D1, D2 G6 The targeted variation is not RED. RED tapers to defense and re-points the T2 Inventory Band and Days to Stockout read against Targeted

same day, which outranks every placement posture below it per v4.0 Section 12 Variation(s) . A rung staged onto a variation about to stock out is

rule P1. discovering a floor for a position that will not be servable. D2 reads PARTIAL, so the backup chain is confirmed on head terms only.

A3 G7 Campaign budgets reconcile to the product envelope. A modifier raise is a T2 Envelope Status and T0 zone H available dollars. A raise that does not

spend increase and must fit the envelope before it is legal. fit the envelope is a queue decision at the instrument layer per C3, not a smaller

modifier here.

All G8 The packet manifest reports present, stale or missing before analysis. Below T7 Status column and the completeness rollup stamped on the cycle header.

85% of gate-relevant fields the packet is DEGRADED and produces flags only. The live read is 0.571. Every verdict in this document issued under a

DEGRADED packet carries the score in its record.

D1 G9 Price stability: no price change on the targeted variation inside the trailing 7 T5 Event Type filtered to price move, cross-read against T2 Margin \$/ days, or every CVR and ACOS read in the placement decomposition carries Unit . Corpus alias per C9: SOP-12 v2.0 Section 3, SOP-14, SOP-30 and the PROVISIONAL and no conversion verdict issues on it. Data Architecture use G9 for margin freshness; v4.0 Section 5 assigns G9 to price stability. Both states are checked here and the reason string names which one fired.

D7 G10 Event weeks and the 14 days following are excluded from the trend reads a T5 Trend-Exclusion Flag and T2 Event Mode . Corpus alias per C9: probe depends on. No modifier change and no probe rung deploys inside an the operating corpus uses G12 for the event freeze and G10 for objective armed window except where the event plan names it. tagging; v4.0 assigns G10 to event exclusion. A rung probed inside a window discovers a floor that dissolves when the window closes.

D6 G11 Family ownership is declared on any shared term before a restore step raises a T2 Family Conflict and Term Owner (Family) , sourced from bid on it. Account Rollup A1. D6 reads PARTIAL. Two siblings restoring one term is a portfolio decision at #33, not two placement decisions.

B11 G12 The term exists in the current Master Keyword List with a syntax tag and a T2 Syntax Tags and Relevancy populated. Corpus alias per C9: G12 is relevancy score. used for the event freeze across the operating corpus; v4.0 assigns G12 to keyword-list hygiene. A defended floor on a term that is not in the list is a floor with no owner.

B7 Placement Placement-level splits present for the campaigns in scope, at TOS, ROS and PP T4 performance columns carrying values rather than the B7 build marker.

read grain. Reads MISSING account-wide. While MISSING: every placement is UNREAD, every modifier increase is suppressed, and P1 returns a structural read only.

Cuts and normalizations stay legal. This is Ledger dependency D3, OPEN, and it caps this component by name.

B6 Bid Suggested bid and bid gap present per row. T2 Suggested Bid / Gap % carrying values. Reads MISSING. The family reference F9 auction-moved signal, gap widening at or above 25% over 4 weeks, cannot compute, so a rung that slips cannot be separated from market movement on this evidence alone and P7 step 8 routes to the T5 log instead.

B8, D1 Economics T0 zone C live: AOV, contribution per order, break-even ACOS, Target ACOS, T0 zone C margin table version and its fresh or provisional state. B8 reads chain Max Acceptable and the Expected CPC reference. Not a numbered gate; a PRESENT on blended margin; D1 reads PARTIAL, so every ceiling in this precondition without which no ceiling in Section 4 computes. document carries PROVISIONAL until the per-variation table lands. This is Ledger dependency D1, OPEN.

[page 4 of 18]

4. Step-Level Procedures

Fourteen procedures. P1 to P4 are the weekly read and the modifier program. P5 to P7 are the ladder, the floor and the re-test, which are the mechanics eleven other components consume. P8 to P12 are the response procedures. P13 is the feed-suppression procedure and P14 is scoring. Every step names the tab, the column and the value it reads or writes. A person who has never run a placement pass executes these from the page. Time estimates assume the packet has refreshed and the operator has run the procedure before; first executions run roughly double.

P1. The weekly placement read (T4)

Weekly, inside the Monday cascade. 20-30 minutes per product at roughly 100 oversight-class terms across three placements.

1. Confirm T0 and T1 have been read first. Product posture and syntax caps constrain everything below them, and a T4 pass that skipped them is invalid by construction
per v4.0 Section 8, where product scenarios outrank syntax which outrank keyword which outrank format.
2. Score last cycle before reading this one. Open T6 Decision Log, filter Review Date to this cycle, and write the V5 Score on every due rung and every due modifier
change as HIT, PARTIAL, MISS or WRONG-DIRECTION. On this SOP the rows that matter most are rungs, because each one carries a hold prediction that a single week either confirms or kills.
3. Read the T7 Manifest row for B7 before opening T4. If it reads MISSING, the placement performance columns are not readable and this procedure runs its structural half only: steps 4 and 5 are skipped, step 6 runs in full, and P13 governs the cycle. If it reads PRESENT or STALE, continue.
4. For each oversight-class keyword, read T4 one row per placement across TOS, ROS and PP: Impressions , Clicks , CVR , CPC \$, CPA \$, Spend Share % , Sales Share % , IS % .
5. Compute the three reads each placement needs, per placement, never per campaign.
 - The proof read. Each placement's CVR against the campaign blended CVR and against rest-of-search CVR, which is the input the P3 raise gate consumes. Record the ratio as computed; the bar it must clear is AM-1 and is not yet governed.
 - The ceiling read. Each placement's own max profitable CPC, computed as contribution per order from T0 zone C multiplied by that placement's own CVR, with the observed CPC beside it. This is per placement because two placements at the same CPC on the same term have different ceilings when they convert differently.
 Resolve the ratio into the v4.0 Section 3 family F1 bands: at or under 0.8 SAFE, 0.8-1.0 AT-CEILING, 1.0-1.5 PUSH-ONLY, above 1.5 UNJUSTIFIABLE outside a sized push.
 - The band read. Each placement's ACOS against the term's governing objective band from v4.0 Section 6 rule O6, which is the objective the row inherits and not the campaign's blended band.
6. Mark UNREAD any placement below sufficiency, and record why. An UNREAD placement can receive a probe under a verdict but never a premium. AMENDMENT. v4.0
carries no per-placement traffic minimum, so this SOP applies the G2 keyword-grain minimums to the placement row and states the substitution on the record rather than treating the placement as sufficient by default. Amendment item AM-2.
7. Route the findings. Proof ratios feed the P3 eligibility test. Any placement resolving to TOS-VANITY or PP-LEAK feeds P8. Drift flags feed P12. Rows carrying an open ladder feed P5 on their dated step. Rows at a recorded floor with no live verdict take HOLD; DEFEND LIGHTLY AT RANK and change nothing, which is a verdict and is logged as one.
8. Write the cycle line into the T4 header: rows read, rows UNREAD with the reason, placements resolving to each of the three F9 classes, open ladders and their next step dates, and the manifest groups that were missing. This line is what P14 scores against.

P2. The ceiling worksheet, mandatory before any modifier change
Per proposed change, 8-12 minutes. Source: v4.0 Section 10 rule M4 for the governing ceiling; Strategic Doctrine v1.0 Section 6.1 for the effective-CPC arithmetic. No modifier change
of any direction deploys without this worksheet, including a cut.

1. Read the current settings from T3: Bid / AG Default \$, Placement Mult TOS % , Placement Mult PP % , and Bidding Strategy . All four are required; a worksheet built without the bidding strategy is computing the wrong ceiling.
2. Compute the current effective CPC ceiling for every placement the change touches, using the three-factor form: base bid multiplied by one plus the placement modifier, multiplied by one plus the dynamic raise. Under fixed and dynamic down-only the dynamic raise is zero. Under dynamic up-and-down the platform can add up to a further doubling at top of search. Platform parameters, not decision numerals. The 0 to 900% modifier range and the up-and-down raise at top of search are Amazon

platform mechanics recorded in Strategic Doctrine v1.0 Section 6.1; v4.0 carries neither, and their absence from the threshold authority is amendment item AM-11. They are the domain of the arithmetic, not a threshold this SOP applies.

3. Compute the proposed effective CPC ceiling the same way, for every placement, not only the one being changed. A base-bid move changes all three ceilings at once and that is the fact most single-placement reasoning misses.

4. Identify the governing ceiling. At maintenance it is the term's own max profitable CPC, contribution multiplied by that placement's CVR, from T2 Max Profitable CPC

\$. During a funded push it is the authorized CPC envelope carried on the #23 worksheet, received at interface I-4. Where neither is available the worksheet stops and

the change does not stage; a modifier sized against an unknown ceiling is the failure the worksheet exists to prevent.

5. Apply the M4 minimum: the proposed effective TOS CPC must sit at or under the lower of max profitable CPC for the objective and Expected CPC plus tolerance.

AMENDMENT. M4 states a tolerance and sets no value, so the second limb of the minimum is not computable; this SOP applies the first limb, records that the second

was unavailable, and logs amendment item AM-4. A6 reads PARTIAL with an account reference of \$2.80 in use, so where the second limb is applied at all it is applied

against an account value and the substitution is written onto the record.

6. Return the verdict. Where the proposed effective ceiling sits at or under the governing ceiling, the worksheet CLEARS and the change proceeds to P3 or P4. Where it does not, the change is BLOCKED as staged.

7. On a BLOCKED raise, compute the compliant pair before returning it. A raise can ride a base-bid cut: raise the modifier, cut the base, and hold the net effective ceiling inside the governing ceiling. Compute the base value at which the pair clears, state both numbers, and stage them as one set. The pair deploys together or not at all,

because deploying the modifier ahead of the base cut authorizes exactly the click the worksheet blocked. Section 6 WE-1 works the arithmetic.

8. Attach the worksheet to the change as its gate artifact. A modifier change staged without an attached worksheet fails staging QA at #38 and returns to the pass. The

worksheet records: current and proposed base, current and proposed modifier per placement, bidding strategy, both effective ceilings, the governing ceiling with its

source, the M4 limbs applied and the limbs unavailable, and the CLEARS or BLOCKED verdict.

[page 5 of 18]

P3. Modifier raise operations

Per raise, 10-15 minutes. Source: v4.0 Section 9 rule L3 and Section 10 rule M4; Strategic Doctrine v1.0 Section 6.2. Currently suppressed account-wide: B7 reads MISSING and

P13 governs. The procedure is stated in full so that it is executable on the day the feed lands and so that the suppression has a defined shape rather than an absence.

1. Confirm the read is legal before the raise is considered. v4.0 Section 9 rule L3 permits a placement adjustment only after G2 sufficiency and a D1 diagnosis. A raise on

an UNREAD placement is the F2 condition regardless of how the CVR looks.

2. Test the proof condition. The placement's CVR must exceed rest-of-search CVR on read-qualified traffic, and the placement's CPA must sit at or under the CPA ceiling

from T2 CPA Ceiling \$. AMENDMENT. Family F9 requires the CVR to be materially above and sets no factor. Record the computed ratio and route the raise to the

human-mandatory tier until AM-1 lands, rather than clearing it against an invented bar.

3. Run the P2 worksheet. It is not optional and it is not a formality: Strategic Doctrine v1.0 Section 6.1 puts the triple multiplication as the reason the worksheet exists, and

Section 6 WE-1 shows a \$2.20 base with a 150% modifier under up-and-down authorizing an \$11.00 click, which is 7.55 times a \$1.456 ceiling.

4. Confirm the envelope. A raise increases spend, so G7 applies: the raise fits the product envelope or it is a queue decision at the instrument layer, T0 zone H, and not a smaller modifier here.

5. Check saturation before staging. v4.0 Section 4 rule X11 applies to top-of-search multiplier increases exactly as it applies to bid and budget increases: where marginal

ACOS runs above 1.5 times average ACOS on the line, freeze at the prior spend level and re-test in 2 weeks. Read T2 Marginal ACOS against ACOS .

6. Apply the confidence modulation. v4.0 Section 10 rule M10 halves every lever step at LOW confidence, permits standard steps at MEDIUM and the upper end of every range at HIGH. AMENDMENT. v4.0 carries no modifier step size to modulate. SOP-28 v1.0 attributed a top-of-search pattern to v4.0 Section 8; no such pattern exists there, and it is retired rather than carried. Amendment item AM-3. Until it lands, a raise is staged at the value the ceiling worksheet's compliant pair returns, which is an arithmetic output rather than a step size, and the derivation is written onto the record.

7. Log the prediction before staging: placement share shift and the efficiency consequence, with a named review date. Strategic Doctrine v1.0 Section 13 holds a modifier raise to the same standard as a bid change: a modifier raise predicts placement-level CVR capture.

8. Stage through #38 with a reason string carrying the verdict verbatim, the v4.0 scenario ID, this SOP and procedure number, the worksheet result, and the computed proof ratio. A staged row missing its scenario ID does not deploy.

9. Route the confirmation. Every raise deployment is human-mandatory: v4.0 Section 12 rule P5 places every 500-plus search-volume action and every bid move above 25% in that tier, and a modifier raise moves the effective ceiling by more than 25% in almost every case the worksheet clears.

P4. Modifier cut and normalization operations

Per cut, 6-10 minutes. Legal throughout the B7 outage: removing an unproven premium never needs proof. Source: v4.0 Section 9 rule L3 and Section 8.8 S-P5.

1. Identify the cut class. A decomposition cut follows a TOS-VANITY read under S-P5 and cuts the multiplier so the term earns at rest-of-search, re-read in 2 weeks. A bleed cut is the first lever of P8 and cuts the bleeding placement's premium. A normalization steps a modifier toward zero because its justification expired: the push closed, the deal window ended, the rank was achieved and tapered.

2. For a normalization, prove the justification expired rather than assuming it. Read T3 Objective Tag against the modifier's originating T6 row: a modifier raised for a Ranking push on a term now tagged Conversions is a premium outliving its reason. AMENDMENT. v4.0 sets no horizon after which an expired justification must be normalized. Amendment item AM-16, which duplicates the expired-justification horizon logged by the SOP-25 session; v4.1 folds them once.

3. Run the P2 worksheet on the cut as on any other change. A cut changes the effective ceiling and the worksheet is the record of what it changed it to; a cut deployed without one leaves no trace of the ceiling before the cut, which is what the next raise would have to be sized against.

4. Compute the effect before staging: the placement's spend at the current effective ceiling against its spend at the proposed one, at constant clicks. State it as the conservative figure it is, because realized click volume falls with the ceiling, so realized spend reduction runs higher and realized order counts run lower. Both directions go into the prediction.

5. Where the cut is a normalization to zero, stage it as a single step. The L2 per-change limit of 50% governs bid values; a modifier moving to zero is not a bid change and the limit does not bind it. Where the cut is partial, the residual modifier must still clear the worksheet at the new value.

6. Stage through #38 with the reason string carrying the verdict, the scenario ID, the worksheet result and the expiry evidence for a normalization.

7. Write the T6 row with a dated prediction. A cut with no prediction cannot be scored, and an unscored cut means the next raise on the same placement starts from no evidence.

P5. The ladder: authoritative rung mechanics

Per rung, 8-12 minutes. Source: v4.0 Section 10 rule M3 for the step, Section 9 rule L2 for the limits. This is the procedure SOP-12 P4, SOP-14, SOP-22, SOP-23 P10, SOP-24, SOP-25, SOP-26 P3 and P8, SOP-30 and SOP-35 all execute against. It decides no descent; it prices every one.

1. Confirm the descent decision arrived from its owner and is not being made here. The arriving artifact names the term, the starting bid and the reason: from #23 at taper
release it carries the achievement date, the push bid and the authorized CPC envelope; from #25 it carries the scheduled date; from #26 it carries the cohort and the
probe instruction; from #35 it carries the incrementality step. Absent that artifact there is no rung to derive and the request returns to its origin.
2. Read the starting bid from T2 Current Bid \$ and confirm it against the arriving artifact. A mismatch is an F11 condition: something moved the bid between the
decision and the derivation, and the rung is derived from the live value with the discrepancy recorded, never from the stale one.
3. Derive the rung. The step is roughly 10% per week per M3, applied to the current value: the rung equals the current bid multiplied by 0.90. Compute it, do not estimate it,
and write it to T4 Ladder State (step / date) with the date.
4. Test the rung against the L2 limits before staging. The change must sit inside 50% of the current value, which a 10% step always does, and the resulting bid must sit at
or above the \$0.50 floor. A derived rung below \$0.50 is truncated to \$0.50 and the truncation is recorded: the ladder has reached a limit rather than a floor, and those are
different findings.
5. Test the rung against the exit target from P6. Where a full 10% step would carry the bid below the exit bid, the rung lands exactly on the exit bid instead and is the
terminal rung of the descent. A terminal rung is smaller than 10% by construction and that is compliant, not an exception: the ladder is aiming at the exit bid, and
overshooting it discards the floor the ladder exists to find.
6. Log the hold prediction with the rung. The prediction is that position holds at the reduced level, with a named review date one read window out. v4.0 Section 14 rule V1
is the authority: unscored actions do not accumulate. A rung without a dated prediction is not a rung, it is drift, and it is the F4 condition.
7. Hold one read window. No second rung deploys before the first is scored. Moving two rungs inside one window makes the slip unattributable to either.
8. Read at the review date. Position held, or position slipped past tolerance. AMENDMENT. M3 sets the reversion on slippage beyond tolerance and states no tolerance
value. Amendment item AM-5, which duplicates the taper slip tolerance logged by the SOP-25 session; v4.1 folds them once. Until it lands the slip is recorded with its
magnitude in positions and routed to the human-mandatory tier rather than resolved against an invented bar.
9. On a hold, record the new lower value as the current ladder state and derive the next rung at its date.
10. On a slip, revert one rung to the prior held value and hold. AMENDMENT. v4.0 sets no hold period for a reverted rung and no re-probe count. SOP-25 and SOP-26 v2.0
both operate a two-window hold and a single re-probe; both figures are ungoverned. Amendment items AM-8 and AM-9. This SOP applies the corpus practice, states
that it is corpus practice rather than a v4.0 threshold, and logs both.
11. A failed probe never lowers the register. The reverted value is the recorded floor until a re-probe holds. This is the single rule that makes the register mean something
across quarters, and violating it is the F8 condition.
12. Stage every rung through #38 as its own change set carrying the next step's date in the reason string, so the T6 chain reads as a dated ladder and V5 scores each rung
independently.

[page 6 of 18]

P6. Exit-bid discovery and the floor register

Per descent, 10-15 minutes at derivation, 5 minutes at each update. Source: v4.0 Section 3 family F9 and Section 9 rule L2 for the decrement; Strategic Doctrine v1.0 Section 6.2 for the register's purpose.

1. Compute the exit bid at the start of the descent, not at its end. The ladder needs a target before its first rung or it is descending toward nothing. The exit bid is the
blended placement CPC at the achieved position minus the decrement of \$0.05-0.10 that v4.0 family F9 and rule L2 both carry. Write it to T4 Placement-CPC Exit

Bid \$.

2. Select the decrement inside the governed band and record the selection. AMENDMENT. v4.0 sets the band and no rule for which end of it to use. The Economics

Worked-Example Set v1.0 case E5 and SOP-28 v1.0 both operate \$0.07. This SOP uses the corpus value, states that it is a selection inside a governed band rather

than a governed value, and logs amendment item AM-6.

3. Re-derive the exit bid as the position's CPC data accumulates. The blended placement CPC at the achieved position moves as the market moves, and an exit target

computed at week one and never updated is aiming at last month's auction. Re-derive at each rung and record whether the target moved.

4. Record the floor when the ladder stops. The register entry carries four things: term, floor bid, date, and the competitive context note describing what the market looked

like when this floor held. The context note is what makes a future slip diagnosable, and Strategic Doctrine v1.0 Section 6.2 names it as part of the register rather than as

commentary.

5. Write the entry to T4 Discovered Floor \$ with its date and set Next Step Date to the quarterly re-test. Instrument note: T4 carries no column for the context

note. Until one exists the note is written to the T6 row that recorded the final rung and the T6 date is carried in the register entry, so the note is retrievable rather than

lost. This is reported as a finding at Section 13.

6. Confirm the register entry exists before the descent is marked complete. v4.0 Section 3 family F9 states the consequence directly: the ladder state is the taper protocol's

memory, and without it every taper restarts from scratch. A completed ladder with no register entry is the F7 condition and the entry is reconstructed from the T6 rung

chain the day it is found.

7. Hand the entry downstream. #25 receives it at the transition exit, #26 reads it for drift restoration and cohort construction, #12 restores to it on a drift detection, and #35

reads it as the defended value where the descent doubled as an incrementality test.

P7. The quarterly floor re-test

Per cohort, 25-40 minutes. Source: v4.0 Section 10 rule M3, which requires the floor re-tested quarterly or after a logged competitive shock, and Section 14 rule V4, which places floor

re-tests on the quarterly cycle. The decision to re-ladder a cohort is #26's; the probe mechanics are this SOP's.

1. Receive the cohort from #26 with its member terms and its recording quarter. This SOP does not build cohorts and does not select which ones run.

2. Check G10 before probing. An armed event window or its 14-day guard suspends the probe for that cohort and the next date moves out by the guard period. Record the

suspension with the new date; a suspended probe with no new date is a probe abandoned.

3. Run the P2 worksheet for any cohort member whose probe touches a modifier rather than a base bid. The effective-CPC arithmetic is multiplicative and the worksheet is

mandatory on a probe exactly as on any other change.

4. Derive the probe rung: one step down from the recorded floor at the M3 step size, inside the L2 per-change limit and above the \$0.50 floor. Stage the whole cohort in

one change set so the read window is common and a market-wide shift is legible as a pattern rather than as unrelated slips.

5. Log the hold prediction per term with a named review date. A probe without a dated prediction cannot be scored and produces no register movement in either direction.

6. Read at the review date and apply P5 steps 8 to 11 unchanged. A hold records the new lower floor with a fresh context note. A slip reverts and the register does not

move.

7. Produce the register diff as the cohort's artifact: floors lowered, floors held, floors raised, each with its date and context note. The diff is what #26 reads to compute the

released difference at its harvest read, and it is the only evidence that the quarter's re-test actually ran.

8. Where a cohort slips as a whole rather than in individual terms, treat it as a competitive reading rather than a bidding result and route to the v4.0 Section 7 rule D7

inflation decomposition before re-probing: market-driven inflation and self-inflicted inflation take opposite responses. Note the live constraint: the family F9 auction-moved signal, bid gap widening at or above 25% over 4 weeks, needs B6 which reads MISSING, so the decomposition runs on the T5 product log and the competitor block rather than on the bid-gap series, and the record says so.

P8. The bleeding-placement response, in lever order

Per bleeding placement, 15-20 minutes. Source: v4.0 Section 9 rule L1 lever precedence and rule L3; scenarios S-P5 and S-P6 at Section 8.8. The order is not a preference and never inverts.

1. Confirm one placement is failing while the others hold. A term failing at every placement is not a placement problem and routes to the verdict its own scenario returns;

the placement read is what separates the two and P1 step 5 is where it happened.

2. Confirm the gates first. A truncated week under G3 voids the placement read entirely, because a starved campaign's placement mix describes its worst hours; the budget

lever runs and the placement conclusion defers to the next untruncated week. This is the F12 condition.

3. Lever 1, the premium goes first. Cut the bleeding placement's modifier where one exists, per P4. This is the cheapest available correction and the only one that

touches nothing else: a modifier is placement-specific by construction, so cutting it cannot damage the placements that work. Run the P2 worksheet, stage, and hold one

read window.

4. Lever 2, the base bid against the bleeder's own CPC. Where the placement still bleeds after the premium is gone, reduce the base bid to the bleeding placement's

own blended CPC minus the F9 decrement. Before staging, compute what the reduced base does to the effective ceiling of every other placement on the campaign,

because the base is shared: the exit that removes the bleeder also caps the earner. State both figures in the record. Where the earner's verified weekly contribution at

risk exceeds the bleeder's weekly loss removed, this lever is not taken and step 5 governs.

5. Lever 3, the split, at persistent divergence. Where the divergence persists and lever 2 would cost more than it recovers, the answer is P10: the base bid stops being

shared. This is the structural correction and it sits behind bid and placement in the L1 precedence exactly as L1 orders it.

6. Never a flat bid cut across the campaign. v4.0 Section 9 rule L3 and Strategic Doctrine v1.0 Section 6.2 both name it: a flat cut punishes the placements that work

and frequently leaves the bleeder serving, because the bleeder often carries a modifier that the flat cut does not touch. Section 6 WE-3 prices exactly this: a 30% flat cut

on the worked campaign leaves the product-page bleeder with \$0.155 per click of unused headroom while stripping \$0.29 per click from the placement carrying the campaign's best per-click return.

7. Write the T6 row naming the lever taken, the levers not taken, and why. A bleed response that records only what was done cannot be audited against the order it was

required to follow.

[page 7 of 18]

P9. Placement arbitrage standup and the rest-of-search inverse (the M4 pattern)

Per standup, 20-30 minutes. Source: v4.0 Section 10 rule M4; Strategic Doctrine v1.0 Section 6.2. The raise half is suppressed while B7 reads MISSING; the base-siting half and

the inverse are executable throughout.

1. Confirm the entry condition. The forward pattern requires a proven top-of-search CVR premium per P1, a top-of-search impression-share deficit on a term whose

objective values visibility, and a clear P2 worksheet at the target. v4.0 Section 3 family F6 sets the deficit line for Ranking campaigns at top-of-search share below 30,

and Section 8.4 S-R7 fires there.

2. Site the base bid at rest-of-search economics: the base is set so that rest-of-search and product pages stay cheap, which means at or under the rest-of-search

placement's own max profitable CPC. Compute it from contribution multiplied by rest-of-search CVR.

3. Size the modifier so that base multiplied by one plus the modifier reaches the target top-of-search CPC and no further. The target is the M4 minimum from P2 step 5.

The base keeps two placements cheap; the modifier alone funds the premium placement. That separation is the whole of the pattern.

4. Confirm the bidding strategy before deploying. Under up-and-down the platform adds its own raise on top of the modifier, which means the sized ceiling is not the

authorized ceiling. Strategic Doctrine v1.0 Section 6.3 sets fixed for ranking pushes and down-only as the default everywhere else; an arbitrage standup on up-and-down

is sized against a ceiling that is not the one in force.

5. Run the inverse where it applies. On a term already holding strong organic position, the rest-of-search harvest runs base at rest-of-search economics with the modifier at

zero, buying incremental shelf below the fold while organic holds the top. This standup needs no proof ratio because it raises nothing, so it stays legal while B7 is

MISSING.

6. Log the placement-share prediction with a review date, and re-read the standup at the normal cadence like any other structure. Both directions of the pattern are

structures, not settings.

7. Where the raise half is suppressed, issue the verdict, execute the executable half, and log the blocked half against the T7 B7 row per P13. Section 6 WE-4 works an

instance where three separate blocks stack on one raise and the only executable action is a bidding-strategy correction.

P10. Placement-split exact candidacy

Per candidacy, 15-25 minutes for the test; the standup itself executes at SOP-12. Source: Strategic Doctrine v1.0 Section 5.2 case 2, which the Ledger row for component 9 records

this SOP as implementing.

1. Test candidacy against the doctrine's own bar: divergent placement economics proven on one term across qualified reads, persistent rather than single-window.

AMENDMENT. Neither the doctrine nor v4.0 states how many windows make a divergence persistent. Amendment item AM-15.

Record the number of qualified windows

observed and route the candidacy to the human-mandatory tier rather than clearing it against an invented count.

2. Confirm the divergence is economic rather than volumetric. Two placements with different click counts are not divergent economics; two placements with different CVR

against the same contribution are. Compute both placements' own max profitable CPC and show the separation.

3. Confirm the split is the cheapest correction available. A split creates a second campaign, a second budget line and a second set of walls, so it is justified only where the

shared base bid is the binding constraint. Where a modifier cut resolves the divergence, the modifier cut is the answer and the split is not.

4. Declare the two roles before the campaigns exist. The top-of-search campaign runs the arbitrage structure of P9: modest base, modifier funding the premium. The rest-

of-search campaign runs base at rest-of-search economics with the modifier at zero. Both roles are declared in T3 so that the pair never reads as duplicates: this is one

of the four legitimate duplicate roles that SOP-12 Section 1.2 carries, and an undeclared pair is that document's P7 territory the day it is found.

5. Hand the standup to SOP-12 P1, run twice, with the declared roles, the sited base bids and the sized modifiers. This SOP supplies the values; it creates nothing.

6. Confirm walls are built in the same working session under #29 rules. A split whose source campaigns still serve the term buys the same click twice at two prices.

7. Log the standup prediction: the combined efficiency of the pair against the pre-split blended read, with a review date. Each campaign grades on its own placement's

economics thereafter, and grading either against the pre-split blend is a process violation under v4.0 Section 6 rule O6.

P11. The bidding-strategy and effective-ceiling audit

Per product, 30-45 minutes, scheduled onto the governed quarterly cycle alongside the M3 floor re-tests. Source: Strategic Doctrine v1.0 Section 6.3.

1. Filter T3 to campaigns carrying up-and-down bidding, and beside each read Placement Mult TOS % and Bid / AG Default \$. The combination of up-and-down

and a live modifier is the audit's target, because that is the combination where the authorized ceiling and the intended ceiling diverge most.

2. Run the P2 worksheet on every row found, whether or not a change is proposed. The audit's product is a list of authorized ceilings, and a row with no worksheet has an unknown one.

3. Rank the findings by the gap between the authorized effective ceiling and the governing ceiling, in dollars per click, multiplied by the row's weekly clicks. That product is the bounded weekly exposure. State it as bounded rather than as realized: the dynamic raise applies per auction rather than to every click, so the figure is a ceiling on the exposure and not a saving.

4. Apply the doctrine posture as the correction: fixed for ranking pushes, down-only as the default everywhere else, up-and-down only where a term carries months of validated CVR and explicit ceiling math. A strategy change is a cut in ceiling terms and needs no placement proof, so it stays legal while B7 is MISSING.

5. Record the April 2026 auction context on any row whose CPC rose materially, and route it through the v4.0 Section 7 rule D7 inflation decomposition before any aggressiveness verdict. Strategic Doctrine v1.0 Section 6.3 reports category CPC inflation of 18-27% against Q1 baselines under dynamic raises; that is a market observation carried in the doctrine, not an authorized numeral, and no v4.0 threshold consumes it.

6. AMENDMENT. Doctrine 6.3 instructs that push sizing carry a CPC buffer while the auction re-stabilizes and states no buffer figure; v4.0 carries none either. Amendment item AM-12, which duplicates the item logged by the SOP-23 session; v4.1 folds them once. Sizing that consumes this SOP's ceilings currently runs unbuffered and the record says so.

7. Issue the ladder-hold notice to #37 naming every row under a live ladder or a live worksheet, so that automation does not step a bid this SOP is stepping. v4.0 Section 12 rule P3 resolves the collision in the audit's favour; the notice is what prevents it arising.

P12. Drift detection and the structural-inspection handoff

Weekly with P1, 8-12 minutes per product. Source: v4.0 Section 3 family F9 drift row and Section 4 rule X6.

1. Read T4 Drift Flag , which fires on a placement mix shift above 15 points over 4 weeks with no logged change. The qualifier carries the weight: a mix that moved because someone moved it is a deployment, and a mix that moved on its own is a finding.

2. Confirm the absence of a logged change against T6 and T5 before treating the shift as drift. A shift matching a staged set is not drift and closes here.

3. Route unexplained drift to structural inspection before any performance explanation is accepted. Family F9 states the routing directly, and v4.0 Section 4 rule X6 carries

the fixed inspection order, whose second item is placement multipliers and effective top-of-search CPC and whose third is bidding strategy. Those two items are this

SOP's to supply.

4. On a sibling divergence, supply the effective top-of-search CPC for each sibling instance rather than the base bids. Two siblings at the same base bid with different

modifiers under different bidding strategies are not priced the same, and the base-bid comparison hides it. #33 makes the ownership call; this SOP supplies the

arithmetic that shows whether the divergence is settings archaeology or auction reality.

5. Check G11 before any restore step on a shared term. Two siblings restoring the same term bid against each other with the family's own money, and the restore holds until the lead is declared.

6. Write the finding to T6 whether or not it produced a change. A drift flag that resolved to a logged deployment is evidence the log is working, and recording it is how that stays true.

[page 8 of 18]

P13. Operating under a suppressed placement feed

Every cycle while T7 reads B7 MISSING. 10-15 minutes per product. This procedure exists because the alternative to a defined suppression is an undefined one, and an undefined suppression degrades into raises made on assumption.

1. Read the T7 row for B7 at the start of the cycle and stamp its status on the T4 header. The suppression is lifted by a manifest read and never by recollection.
2. Mark every placement UNREAD for the cycle and state it. An UNREAD placement is not a placement performing at zero and is not a placement performing at the

campaign blend; it is a placement whose performance is unknown, and the three are different inputs to different decisions.

3. Run the cut side in full. P4 normalizations, P8 lever 1 modifier cuts where a bleed is visible from campaign-grain evidence, P5 rungs, P6 exit-bid derivation, P7 probes,

P11 strategy corrections and P12 drift routing all run unchanged. None consumes B7.

4. Suppress the raise side entirely. No modifier increase stages while B7 is MISSING, on any row, under any verdict, at any confidence. A raise staged during the outage is the F10 condition and it escalates on discovery.

5. For every verdict that arrives requiring a raise, issue the verdict, execute the executable half, and log the blocked half as a blocked verdict class against the T7 B7 row

with the term, the campaign, the verdict string, and the evidence that would have supported it. The log is the feed's business case: a placement feed with named failed

pushes and blocked verdicts attached to it argues its own priority more effectively than a completeness score of 0.571 does.

6. Carry the same statement onto every record the cycle produces. Under G8 the analysis states what it could not see before it recommends anything, and under v4.0

Section 13.3 rule J7 hiding the gap is a process violation rather than a presentation choice.

7. On the cycle where B7 flips to PRESENT, do not release the backlog at once. Run P1 in full to establish the first clean read, hold one window, and only then work the

blocked-verdict log in order of the weekly loss dollars each entry carries. A backlog released against a first read has no trend context, and A1 four-period comparability

reads MISSING, so the first clean read is a single window by construction.

P14. Scoring, calibration and the register diff

Weekly for rungs and modifier changes, monthly for the register diff. 15-25 minutes. Source: v4.0 Section 14 rules V1 and V5.

1. Score every due T6 row: rung hold predictions, modifier raise predictions of placement-share shift, cut predictions of spend reduction at constant clicks, and split standup

predictions of combined efficiency against the pre-split blend. Write HIT, PARTIAL, MISS or WRONG-DIRECTION.

2. Separate the two error classes on any missed rung, because they have different owners. A rung that slipped because the market moved is a competitive finding routing

to the T5 log and the D7 decomposition. A rung that slipped because the step was too large is a mechanics finding routing to this document. Conflating them teaches

nothing and produces a ladder that is either too timid or too fast forever.

3. Produce the monthly floor-register diff across the product: floors lowered, held, raised, and floors whose re-test date has passed without a probe. The last group is the

finding, and it is #26's failure mode F2 detected from this side.

4. Compute rule-level accuracy per v4.0 Section 14 rule V5: any rule below 50% accuracy over 10 or more scored predictions enters review, and the review produces an

amendment, a threshold change, or a documented defense. On this SOP the rule most likely to trip it is the M3 step size itself, and that is the intended behaviour.

5. Attach every finding traced to a missing feed to the T7 row for that field group, with its dollar consequence. This is the same mechanism P13 step 5 uses, and the two

logs are one artifact read two ways.

6. A rung with no score is the F4 condition and a completed ladder with no register entry is the F7 condition. Both are detected here, and both are repaired from the T6

chain rather than re-derived from scratch.

5. Decision Points: The Full Verdict-to-Execution Mapping

Every verdict issuance, threshold and scenario behind it lives in v4.0. Verdict strings are quoted verbatim from the Keyword Decision Record Template v3.0, sheet Bands &

Verdicts Reference , which carries the twenty-five strings the exemplar SOP-12 v2.0 Section 5 maps. Per Ledger conflict C10 the v4.0 Section 13.1 string is carried as a

stated cross-map on the divergent rows so neither citation is dropped. Zero local thresholds appear in this table.

Verdict (KDR Template v3.0, closed vocabulary) v4.0 Section 13.1 Placement and modifier execution

cross-map

QUARANTINE: DATA REFRESH FLAG: DATA REPAIR The row leaves every placement read, every ladder and every worksheet this cycle. G1 is mechanical and v4.0

Section 12 rule P8 places data quality above urgency. A rung already staged on the row is pulled from the

change set.

UNLOCK TRAFFIC No 13.1 member; S-R1 Bid steps of 15-25% per v4.0 Section 9 rule L2, executed at SOP-12. This SOP reviews the placement modifiers

verdict text at the same time per S-R1 and returns any that are suppressing serving, but it raises none while B7 is MISSING.

FIX BUDGET (TRUNCATED) No 13.1 member; G3 Budget lever only at #12. The week's placement read for the campaign is void, not conservative, and resumes on

response the next untruncated day. This is the F12 condition when ignored.

CONTINUE TEST TO 50 CLICKS AT CAPPED BID EXTEND TEST No modifier change and no rung. The capped bid is the test's own instrument; a placement adjustment inside a

running test changes what the test measures.

FUND TEST (UNFUNDED PRIMARY) FUND DELIBERATE Modifiers at zero on all groups at standup. v4.0 Section 9 rule L3 bars a placement adjustment before G2

TEST (50-CLICK CAP) sufficiency, and a test has none by definition. The modifier is earned rather than defaulted.

RANK PUSH: FUND SIZED PUSH RANK PUSH AT Modifiers at zero unless placement-level evidence for this term already exists here. This SOP receives the

CONTROLLED CPC authorized CPC envelope from #23 at I-4 and tests every in-flight modifier request against it rather than against the maintenance ceiling.

RECOVERY PUSH No 13.1 member; v4.0 As above, with one addition: where a pre-OOS ladder state and floor exist for the term, they are supplied to #24

routes S-R13 into RANK as the recovery bracket. A recovery re-entry bid seeded without the prior floor is sized cold on a term with a PUSH AT proven ceiling.

CONTROLLED CPC

TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS No 13.1 member; S-R3 S-R3 names the top-of-search multiplier and budget as its levers. While B7 is MISSING the multiplier half is not

verdict text executable, so the budget half runs alone at #12 and the modifier half logs as a blocked verdict class per P13 step 5.

FIX CONVERSION FIRST; DEFER PUSH Identical string No modifier moves. A conversion deficit is not corrected by buying a better placement, and raising a modifier

against it buys more of the same rejection per v4.0 Section 4 rule X2.

PLACEMENT FIX FIRST No 13.1 member; S-R7 The verdict this component exists for. Bids hold at #12 while this SOP runs the program: P1 read, P2

verdict text worksheet, P9 arbitrage sizing where the pattern applies, P8 lever order where a bleed is present. While B7 is MISSING the top-of-search share deficit remains measurable and the modifier correction does not, so the verdict

issues and the executable half is named. Section 6 WE-4 works exactly this state.

RE-POINT VARIATION Identical string Executed at #12 P10 same day. This SOP re-derives the exit bid against the backup variation's own placement

CPCs rather than carrying the primary's, because a re-point resets the CVR baseline and the ceiling with it.

TAPER: RANK ACHIEVED, STEP BIDS DOWN TAPER TO FLOOR P5 in full from the arriving achievement date, push bid and authorized envelope, P6 for the exit target and the

(POST-OBJECTIVE) register entry. This SOP names every rung and every floor; it decides no taper. Execution is #12 P4.

HOLD; DEFEND LIGHTLY AT RANK

[page 9 of 18]

No 13.1 member; the No rung and no modifier change. The row is read again next cycle against its recorded floor. Doing nothing is a

13.1 near-neighbour is verdict here and is logged as one, because an unlogged hold is indistinguishable from a row nobody read.

DEFEND AT FLOOR

DEFEND AT FLOOR Identical string No change where the bid sits at the recorded floor. Where drift above the floor is detected with no verdict

authorising it, this SOP supplies the floor value and #12 restores to it in one staged step. The restore is a return to a recorded value, not a new decision, and needs no worksheet.

BID DOWN TO PROFITABLE CPC; CAP SYNTAX Identical string The bid edit is #12's, to the row's own max profitable CPC from T2 K6, inside the L2 per-change limit. Where the

target cut exceeds 50% of the current value this SOP returns the compliant multi-step program with its dates, because a single-step version fails staging QA and arrives a week late having looked faster.

INCREMENTALITY LADDER No 13.1 member; v4.0 This SOP stages the single rung and records the floor; the two-week capture read and the continue-or-restore

carries the mechanism call belong to #35. Where the window was confounded, the void is handed on with the T5 references and no at X12 floor is recorded, because a floor recorded without the capture read is a bid with a label.

ADD EXACT + RESTRUCTURE Identical string Standup at #12 P1 or P2 with walls at #29. This SOP supplies modifiers at zero at creation and nothing else; the

new exact has no placement history and therefore no earned premium.

CONSOLIDATE DUPLICATES No 13.1 member; Handed to #12 P7. Where the losing instance held placement evidence the survivor lacks, this SOP receives it

carried inside ADD with the consolidation note and sizes the survivor's modifiers from it rather than rediscovering them, which is the EXACT + evidence-preservation clause of #12 P7 step 4 read from this side.

RESTRUCTURE

GRADUATE FROM DISCOVERY GRADUATE + STEER Owned by #30. This SOP supplies the exit-bid reference that seeds the new exact where a related floor exists on

the root, and modifiers at zero otherwise.

NEGATE (PRE-LOAD / REACTIVE / STEERING) CUT / NEGATE Owned by #29. No placement execution. A negation changes which terms serve, not how their placements are

priced, and re-reading the placement split after a wall lands is a P1 item rather than a response.

ENTER SBV ARBITRAGE FORMAT This SOP derives the single SP-exact ladder step down that v4.0 Section 9 rule L8 pairs with the video entry,

SUBSTITUTION (SBV/ dated to the SBV go-live. The step is one rung at the M3 size, not a cut to floor: the pair is one managed posture

SD) and the SP side steps rather than exits. The video build belongs to #17-19.

OPEN / ROTATE / EXIT CONQUEST No 13.1 member; Owned by #16. Placement execution on PAT rows runs unchanged here, with the T7 A8 PARTIAL coverage gap

carried at S-C1 to S-C6 named on any PAT rung, because T7 records PAT ladder verdicts as blocked while that group is partial.

STOP-LOSS: CONCEDE OR DEFER Identical string On a concession the term transitions to Conversions at ceiling and this SOP supplies the step program from the

push bid to that ceiling, multi-step where the L2 limit requires it. On a defer, the ladder state and any discovered floor are preserved for the re-entry rather than reset.

PAUSED-PENDING (REASON + RE-ENTRY) PAUSE (LEVERS Executed at #12 P8. The ladder state and the floor register entry survive the pause: a term that re-enters at a re-

EXHAUSTED) derived bid has discarded a quarter of ladder evidence, which is the F7 condition arriving by a different route.

LEAVE TO HARVEST Identical string No exact instance and therefore no placement program. Discovery-surface placement modifiers are the hosting

campaign's and follow its own objective band under O6.

5.1 Campaign-level verdicts

Campaign rows carry their own subset vocabulary, quoted from the KDR Template v3.0 Campaign Decision Record sheet: FUND / HOLD / TRIM BUDGET, RAISE / CUT

MULTIPLIER, NEGATION PASS, CONSOLIDATE, RETAG OBJECTIVE, PAUSE (L7), ARCHIVE. Two of these are this SOP's own output. RAISE MULTIPLIER issues only

from P3 with all three conditions met and the P2 worksheet clear, and it is suppressed account-wide while B7 reads MISSING. CUT MULTIPLIER issues from P4 and from P8

lever 1 and is legal throughout. TRIM BUDGET follows a bid or modifier cut rather than leading it, so that each ladder step reads cleanly against a stable budget line.

5.2 Verdict strings retired at this version

SOP-28 v1.0 Section 5 issued seven strings. One, the placement scenario reference, is a scenario pointer rather than a verdict and is carried in Section 4 instead. The

remaining six appear in neither the KDR Template v3.0 closed vocabulary nor v4.0 Section 13.1 and are retired here: RAISE MODIFIER (proof + worksheet) , NORMALIZE

MODIFIER , ENTER ARBITRAGE / ROS HARVEST , TAPER (from any SOP) , RE-LADDER (QUARTERLY) and PLACEMENT-SPLIT . A closed vocabulary is what makes the

decision log queryable and predictions scoreable per verdict type under v4.0 Section 14 rule V5, and a locally minted string is invisible to both. The operations they named

survive as procedures rather than as verdicts: the raise is P3 and issues RAISE MULTIPLIER at campaign grain, the normalization is P4 and issues CUT MULTIPLIER, the

arbitrage and harvest standups are P9 and issue PLACEMENT FIX FIRST at keyword grain, the taper is P5 and executes TAPER: RANK ACHIEVED, STEP BIDS DOWN, the

quarterly re-test is P7 and issues DEFEND AT FLOOR or HOLD; DEFEND LIGHTLY AT RANK per term, and the split is P10 and issues ADD EXACT + RESTRUCTURE. RE-

LADDER (QUARTERLY) was retired independently by the SOP-26 session at its Section 5.2; this is the same string retired at its mechanics owner.

5.3 A fourth vocabulary, at placement grain

The Master PPC Data Architecture v1.0 Section 8 and the Command Workbook TEMPLATE v1.0 T4 tab both carry a placement verdict column with the values PREMIUM

EARNED, CUT MULTIPLIER and EXIT PLACEMENT. Only CUT MULTIPLIER is a member of the KDR campaign subset; PREMIUM EARNED and EXIT PLACEMENT are

members of no closed vocabulary in the corpus. This SOP writes the T4 column with those values because they are the instrument's data-validation list and the column has to

be filled, and it issues its verdicts from the closed vocabularies at keyword and campaign grain. The divergence is reported as finding N3 at Section 13.2 as an extension of C10

and C18 rather than resolved here.

6. Worked Examples

All four run on the standard fictional training base, disclosed here because the anonymized live per-variation product economics spine is not attached to this session and Ledger

dependency D1 is OPEN: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%, blended CVR 5.2% unless a

placement's own rate is stated. Live executions read the T0 zone C named cells instead and the arithmetic is identical. Every figure below was computed and asserted in code

with float tolerances before it was written here. WE-1 and WE-2 carry forward the anchors of SOP-28 v1.0's single worked example so the two versions reconcile, with the one

arithmetic correction Section 10.2 records.

[page 10 of 18]

WE-1 · The ceiling worksheet: a blocked raise, the compliant pair, and the warning case

Setup. A maintenance term at contribution \$28.00 and CVR 5.2%. The proposal arriving at P2 is a top-of-search modifier of plus 25% on the current base bid of \$1.30, on a campaign running dynamic down-only.

Governing ceiling (max profitable CPC) = $28.00 \times 0.052 = \$1.456$

Current effective TOS ceiling = $1.30 \times (1 + 0.25) \times (1 + 0.00) = \1.625

1.625 is above 1.456 by = \$0.169 per click

VERDICT: BLOCKED as staged.

The compliant pair, computed rather than negotiated. The raise can ride a base-bid cut. Solving for the base at which the pair clears: the base must sit at or under 1.456

divided by 1.25, which is \$1.1648. Taking \$1.15 as the staged value:

Proposed effective TOS ceiling = $1.15 \times (1 + 0.25) = \1.4375

Headroom under the governing ceiling = $1.456 - 1.4375 = \$0.0185$ per click

Base cut required = $(1.30 - 1.15) / 1.30 = 11.54\%$

At 30 clicks per week: authorized spend falls from $1.625 \times 30 = \$48.75$ to $1.4375 \times 30 = \$43.13$

The pair stages as one change set and deploys together or not at all. Deploying the modifier on the Tuesday and the base cut on the Thursday authorizes exactly the \$1.625

click the worksheet blocked, for two days, on every auction the term enters. The base cut of 11.54% sits inside the L2 per-change limit of 50% and the resulting bid of \$1.15

sits above the \$0.50 floor, so both halves are stageable.

Why the worksheet is mandatory rather than advisable: the warning case. Strategic Doctrine v1.0 Section 6.1 states the triple multiplication. Take a \$2.20 base with a

plus 150% top-of-search modifier on a campaign left on dynamic up-and-down:

Authorized effective TOS CPC = $2.20 \times (1 + 1.50) \times (1 + 1.00) = \11.00

Ratio to the \$1.456 ceiling = 7.55 times

v4.0 Section 3 family F1 band: above 1.5 times = UNJUSTIFIABLE outside a sized push

No operator proposes an \$11.00 click. The \$11.00 click is proposed by a \$2.20 bid, a modifier set once for a push that ended, and a bidding strategy nobody re-read. Each

of the three is individually defensible and the product of the three is not, which is why the arithmetic runs before the change and not after the invoice. This is failure mode F6

and it escalates on discovery.

What issues. At campaign grain, RAISE MULTIPLIER on the compliant pair, human-mandatory per v4.0 Section 12 rule P5 because the effective ceiling moves more than

25%. Under the live manifest this raise does not stage at all: B7 reads MISSING, so the proof ratio behind the plus 25% has no read-qualified traffic to compute from,

and P13 governs. The base cut and the strategy correction in the warning case both stage normally, because both are cuts.

WE-2 · The ladder end to end: exit target, rung, slip, confound, re-probe, floor

Setup. A push closes at rank achieved. #23 hands over at interface I-4: achievement date, push bid \$2.60, and the authorized envelope. The blended placement CPC at the achieved position reads \$2.31.

Exit bid = blended placement CPC - decrement = $2.31 - 0.07 = \$2.24$

Rung 1 = $2.60 \times 0.90 = \$2.34$ (10.00% step)

Rung 2 = 2.34 to the exit bid = \$2.24 (4.27% terminal rung)

Total descent 2.60 to 2.24 = 13.85%

The v1.0 correction. SOP-28 v1.0's worked ladder ran \$2.60 to \$2.45 as its first rung, which is a 5.77% step against the roughly 10% per week that v4.0 Section 10 rule M3

sets, and it then ran four rungs where two reach the same target. The format exemplar SOP-12 v2.0 already carries the corrected ladder on this exact anchor, at \$2.60 to

\$2.34 then to the \$2.24 exit reference, and records in its own change note that v1.0 showed a first step past the M3 size. The authoritative ladder document was the one out

of step with the exemplar that cites it. v2.0 adopts the compliant ladder.

Rung 1 deploys. \$2.60 to \$2.34, dated, with the hold prediction that position holds at the reduced level and a review date one week out. At the review date position holds.

The value is recorded as the current ladder state and rung 2 is derived at its date.

Rung 2 slips. \$2.34 to \$2.24 deploys and position slips past tolerance at the review date. Per P5 step 10 the rung reverts one step to \$2.34 and holds. The register does

not move: a probe that failed does not lower the recorded floor, so the floor stays at \$2.34 pending a re-probe.

The confound. The T5 Product Log carries two competitor deal windows opening inside the slip window, event type competitor entry. v4.0 Section 4 rule X9 puts the

product log ahead of any attribution to bids, so the correct reading is that the probe was uninformative rather than that the floor is \$2.34. The revert was still the right

execution: the rung discipline restores position first and diagnoses after, and it is not a judgement about cause.

The re-probe. Two read windows later, with both windows closed, \$2.24 re-probes and holds. The register updates: floor \$2.24, dated, with the context note recording that

the first probe slipped inside two competitor deal windows and the second held after they closed. That note is what makes the next slip on this term diagnosable in thirty seconds instead of a week.

What the floor is worth, and what it costs.

Weekly spend at the push bid, 30 maintenance clicks = $2.60 \times 30 = \$78.00$

Weekly spend at the discovered floor = $2.24 \times 30 = \$67.20$

Weekly reduction = \$10.80 (annualized \$561.60)

Premium the held rank still pays = $2.24 - 1.456 = \$0.784/\text{click}$

Premium per week at 30 clicks = \$23.52 (\$3.36/day)

Organic contribution the rank protects = $3.1 \text{ orders/day} \times 28.00 = \$86.80/\text{day}$

Protection ratio = $86.80 / 3.36 = 25.8$ to 1

That ratio is what HOLD; DEFEND LIGHTLY AT RANK means arithmetically, and it is what P7 re-tests each quarter. Reconciliation note. The Economics Worked-Example

Set v1.0 case E5 carries this case at \$23.52 per week and \$3.36 per day, which reproduces exactly at full precision from a \$1.456 ceiling. SOP-26 v2.0 WE-3 carries the same case at \$23.40 and \$3.34, which is what the same arithmetic returns from a ceiling rounded to \$1.46 before the subtraction. E5 governs and the difference is \$0.12

per week of rounding order, not a disagreement about the floor; it is recorded at Section 13.2 finding N6 so the two documents reconcile at the next increment rather than teaching an operator that the arithmetic is approximate.

[page 11 of 18]

WE-3 · A bleeding placement, the lever order, and what the flat cut would have cost

Setup. One exact campaign on one term. Base bid \$1.80, top-of-search modifier 0%, product-page modifier plus 25%, dynamic down-only so the dynamic raise is zero.

Effective ceilings: top of search and rest of search at \$1.80, product pages at 1.80 multiplied by 1.25, which is \$2.25. Week 1 placement read at P1, with each placement's own ceiling computed from contribution multiplied by that placement's own CVR:

Placement Clicks CPC Spend CVR Orders Contribution Net ACOS Own ceiling / ratio

TOS 150 \$1.55 \$232.50 6.4% 9.60 \$268.80 +\$36.30 30.27% \$1.792 / 0.86 AT-CEILING

ROS 220 \$1.05 \$231.00 4.5% 9.90 \$277.20 +\$46.20 29.17% \$1.260 / 0.83 AT-CEILING

PP 130 \$1.42 \$184.60 1.2% 1.56 \$43.68 -\$140.92 147.92% \$0.336 / 4.23 UNJUSTIFIABLE

Blended 500 \$1.2962 \$648.10 4.212% 21.06 \$589.68 -\$58.42 38.47% Hides all three rows

What the blended line hides. The campaign reads 38.47% ACOS, which against a 35.0% break-even looks like a row needing a modest correction. Inside it, one placement runs at 147.92%, which is 3.85 times the blended reading, absorbs 28.48% of the spend and returns 7.41% of the sales. The other two both sit inside the AT-CEILING band and are the reason the blend looks survivable. This is v4.0 Section 3 family F9's PP-LEAK class and Section 8.8 S-P6 is the governing scenario.

Lever 1: the premium goes first. The product-page modifier cuts from plus 25% to zero at P4, which drops the product-page effective ceiling from \$2.25 to \$1.80 and touches neither other placement, because a modifier is placement-specific by construction. No proof is required to remove an unproven premium. Week 2 read on product pages: 74 clicks at \$1.18, spend \$87.32, CVR 1.4%, orders 1.036, contribution \$29.01.

PP net, week 1 = -\$140.92

PP net, week 2 after the modifier cut = -\$58.31

Weekly loss removed by one placement-specific lever = \$82.61

Campaign net, week 1 = -\$58.42

Campaign net, week 2 = +\$24.19 (swing \$82.61)

Lever 2, and why it is tested rather than taken. Product pages still bleed at \$58.31 per week. The L2 exit mechanic sets the exit bid at the bleeding placement's own

blended CPC minus the decrement: 1.18 minus 0.07, which is \$1.11. But the base bid is shared. Cutting the base from \$1.80 to \$1.11 is a 38.33% cut that also caps top of

search at \$1.11 against the \$1.55 it is currently paying, putting the whole of that placement's verified \$36.30 per week of net at risk to remove \$58.31 of loss. The bounded

net gain is \$22.01 per week and it can be negative. P8 step 4 requires that comparison to be computed before the lever is taken; here it fails the test and step 5 governs,

which routes to P10.

The flat cut, priced. Failure mode F3 is a flat campaign-wide bid cut taken instead of the placement work. Suppose a 30% flat cut, \$1.80 to \$1.26:

TOS effective ceiling falls 1.80 to 1.26, against \$1.55 being paid = \$0.29/click stripped from the earner

PP effective ceiling falls 2.25 to 1.575, against \$1.42 being paid = \$0.155/click of headroom still unused

Result: the placement carrying +\$36.30/wk is capped below its paid CPC.

The placement carrying -\$140.92/wk keeps serving with headroom to spare.

That inversion is the exact failure the placement read exists to prevent, and it is not a matter of degree: the flat cut damages the earner first and the bleeder last, because

the bleeder is protected by the very modifier the flat cut does not touch. v4.0 Section 9 rule L3 and Strategic Doctrine v1.0 Section 6.2 both name it, and this is what the

naming costs when it is ignored.

WE-4 · The arbitrage that cannot be sized: three blocks on one raise, and the one action that executes

Setup. A term on a Ranking objective. Base bid \$1.15, top-of-search modifier 0%, bidding strategy dynamic up-and-down inherited from a launch-era build. Blended CVR

4.9%, rest-of-search CVR 4.5%. Sufficiency met at 214 trailing clicks against the G2 CVR bar of 100. In-budget 100%, so G3 is clean and there is no budget half to correct.

Inventory GREEN. T2 TOS IS % reads 18, which sits in the KDR 12-25 band and below the 30 that v4.0 Section 3 family F6 requires on ranking campaigns. v4.0 Section

8.4 S-R7 fires: the placement problem precedes any bid verdict.

P11 runs before P9, and finds an open authorization. The audit reads the bidding strategy before the modifier is proposed:

Governing ceiling = $28.00 \times 0.049 = \$1.372$

Authorized under up-and-down = $1.15 \times (1 + 0.00) \times (1 + 1.00) = \2.30

Ratio to ceiling = $2.30 / 1.372 = 1.68$ times (UNJUSTIFIABLE band)

Authorized under fixed = $1.15 \times (1 + 0.00) = \1.15

Ratio to ceiling = $1.15 / 1.372 = 0.84$ times (AT-CEILING band)

Bounded weekly exposure closed, at 63 clicks = $63 \times (2.30 - 1.15) = \72.45

The exposure is stated as bounded rather than as a saving. The dynamic raise applies per auction rather than to every click, so \$72.45 is a ceiling on what the open

authorization could cost in a week, not what it did cost. Overstating it as a saving would be the same category of error the flat cut makes in WE-3.

The arbitrage the operator wants, and the three blocks on it. The M4 pattern would site the base at rest-of-search economics and size the modifier to reach the target

top-of-search CPC. The base is already close: the rest-of-search ceiling is 28.00 multiplied by 0.045, which is \$1.26, and \$1.15 sits at 0.91 of it. Sizing the modifier requires

the M4 minimum, and every input to it fails:

1. The proof ratio has no bar. Family F9 requires top-of-search CVR materially above rest of search and states no factor. Amendment item AM-1.

2. The first limb of the minimum has no input. Max profitable CPC for the objective at top of search needs the top-of-search CVR, which comes from B7. T7 reads B7

MISSING.

3. The second limb has no tolerance. M4 sets the second limb at Expected CPC plus tolerance and states no tolerance value. Amendment item AM-4. Compounding it,

T7 reads A6 PARTIAL with an account reference of \$2.80 standing in for the per-term Expected CPC, so even a governed tolerance would be applied to an account value.

What issues, and what does not. At keyword grain the verdict is PLACEMENT FIX FIRST, quoted from the KDR closed vocabulary and cited to S-R7. The executable half

is the bidding-strategy correction to fixed, which is a cut in ceiling terms and needs no placement proof, staged through #38 with the worksheet attached. The blocked half is

logged against the T7 B7 row per P13 step 5, carrying the term, the campaign, the verdict string, the 12-point top-of-search share gap, and the three named blocks. There

is no budget half: in-budget reads 100%, so the other lever S-R7 names has nothing to correct either.

Why this example is in the document. The verdict is correct, the diagnosis is correct, the measurement is available, and the correction is not executable. That state is not

a failure of the pass and it is not a reason to substitute a smaller modifier chosen by feel. It is a \$72.45-per-week bounded exposure closed and a 12-point share deficit left

standing with a named owner, and the log entry is what turns dependency D3 from a line in the Ledger into a business case. The SOP-23 session recorded the same

mechanism from the other side, attaching a failed push of \$3,378.00 to the same B7 row.

[page 12 of 18]

7. Failure Modes and Escalation

Twelve modes. Each carries the signal that detects it, the response, and the account-history incident that motivated it where one exists. Incidents are sourced rather than

asserted; the seven carried forward from SOP-28 v1.0 are marked, and each now names the document that records it.

ID Failure Detection signal Response Motivating incident

F1 Modifier raised without A raise deployed with no computed proof Normalize the modifier to zero at P4 the day it is v1.0. Strategic Doctrine v1.0 Section 14 lists it as a

placement-level proof ratio in the reason string, or with a ratio found, then re-raise properly at P3 if the proof named account failure class and Section 6.2 states

computed on traffic below sufficiency. exists. The normalization does not wait for the the doctrine that prevents it: prove before you pay,

proof, because the premium is being paid and top-of-search premium CVR is common but not

meanwhile. universal, with product-page-strength products

inverting it.

F2 An UNREAD placement A premium live on a placement below Normalize to zero and probe traffic first if the v1.0. v4.0 Section 5 gate G2 bars a rate judgment

treated as a read sufficiency, or a placement verdict issued placement matters. An UNREAD placement can below sufficiency at any grain, and Section 9 rule L3

on a row whose T4 performance columns receive a probe under a verdict; it can never permits a placement adjustment only after G2 carry the build marker rather than values. receive a premium. passes. The per-placement bar itself is amendment

item AM-2, which is why this mode is detected on

the value rather than on a threshold.

F3 A flat bid cut taken on a single- A whole-campaign bid reduction staged Reverse to the P8 lever order: the bleeding v1.0. v4.0 Section 9 rule L3 and Strategic Doctrine

placement bleed where the T4 read shows one placement placement's modifier first, then its base against v1.0 Section 6.2 both name the flat cut as the failure

failing and the others inside their bands. its own blended CPC with the shared-base cost the placement read exists to prevent. Section 6

computed, then the split. Restore the working WE-3 prices it: a 30% flat cut strips \$0.29 per click

placements' economics in the same staged set. from the earner and leaves the bleeder \$0.155 per

click of unused headroom.

F4 Undated or unpredicted rungs A T4 Ladder State entry with no date, Re-stage the rung with its date and prediction. A v1.0. v4.0 Section 14 rule V1 states the principle:

or a rung deployed with no T6 hold descent without dates is not a ladder, it is drift unscored actions do not accumulate. v4.0 Section 3

prediction and no review date. wearing a ladder's name, and nothing about it family F9 states the consequence: the ladder state is can be scored or reversed on evidence. the taper protocol's memory, and without it every taper restarts from scratch.

F5 Premiums outliving their A modifier live past the event that P4 normalization on the ladder cadence. v1.0. Strategic Doctrine v1.0 Section 14 names

justification authorized it: a closed push, an ended Premiums are rented and the rent stops with the post-objective spend persistence as an account

deal window, an achieved and tapered reason. The scan runs weekly rather than on failure class, and v4.0 Section 16 Worked Pass 1 rank. Detected by reading T3 Objective request, because nothing alarms on a premium records the instance: a head term held at rank 2

Tag against the modifier's originating T6 that is simply still there. against a target of 3 for five consecutive weeks with row. the Ranking objective expired three weeks earlier and the spend never reduced.

F6 Up-and-down bidding beneath The P11 audit finding a campaign on Worksheet immediately, then the doctrine v1.0. Strategic Doctrine v1.0 Section 6.1 records the

a live modifier with no dynamic up-and-down carrying a non-zero posture: fixed for ranking pushes, down-only arithmetic that makes it expensive: a \$1.50 bid with

worksheet top-of-search modifier with no P2 elsewhere, up-and-down only with months of a plus 300% top-of-search modifier under up-and-

worksheet on file. validated CVR and explicit ceiling math. down authorizes a \$12.00 click. Section 6 WE-1

Escalates on discovery, because it is an open works the \$11.00 variant. The click is never ceiling breach in waiting rather than a breach proposed; it is assembled from three individually that has happened. defensible settings.

F7 Floors unrecorded, or lost at a A completed ladder with no T4 Reconstruct the entry from the T6 rung chain the v1.0. v4.0 Section 3 family F9 is the source. The

pause Discovered Floor \$ entry; or a term day it is found, with the context note written from downstream cost is recorded in the Ledger row for

re-entering from PAUSED-PENDING at a the T5 log for the window in question. The component 26, which names the stage entry without

re-derived bid rather than at its recorded register is the ladder's product; losing it discards recorded floors as the highest-consequence non-

floor. the quarter of evidence that produced it. arrival on its upstream side: drift detection has nothing to drift from and the quarterly re-ladder has no register to probe from.

F8 A failed probe lowering the A T4 floor value written on a rung that Restore the register to the last held rung and re- v4.0 Section 10 rule M3 sets the recorded value at

register slipped, or a re-test date advanced after a date the re-test. The reverted value is the the restored step rather than at the attempted one.

reverted probe as though the probe had recorded floor until a re-probe holds, and the SOP-26 v2.0 WE-3 works the confounded instance

held. failed probe is recorded as a probe, not as a on this document's own anchor: a probe slipping

floor. inside two competitor deal windows, reverted, and

re-probing successfully after the windows closed,

where writing the slip into the register would have

raised the term's floor on evidence that was never

valid.

F9 A rung staged outside the L2 A staged step exceeding 50% of the Re-stage as a compliant program with each step SOP-26 v2.0 WE-1 is the recorded instance: in a

limits current value, or landing below the \$0.50 dated. Where the target requires more than one code-red cut sequence, cuts to the rows' own

bid floor, or a multi-step target compressed step, return the whole program at once so the ceilings required two staged steps each because the

into one step. schedule is visible: a single-step version fails single-step versions breached the L2 per-change

staging QA at #38 and arrives a week late limit, which is why 82.2% of that sequence's saving having looked faster. landed in week 1 and the remainder could not land until week 2.

F10 A modifier increase staged Any staged row raising a placement Pull the row from the change set. If it deployed, Ledger dependency D3 is OPEN and states that the while B7 reads MISSING multiplier during the suppression, detected normalize it the same day and record the performance sections of this SOP stay structural at the #38 pre-upload scan or in the T+1 breach. Escalates on discovery, because the until the feed lands. The cost of operating around it diff. raise was made on an assumption presented as rather than through it is recorded from the other a read and the assumption is now in the log as side: the SOP-23 session traced a failed push of evidence. \$3,378.00 to the same B7 row, where the correct verdict issued and the one lever that would have executed it was suppressed.

F11 Automation and the ladder A T6 rung chain interrupted by a rule-fired Suspend the rule pending #37 review, restore SOP-26 v2.0 WE-1 records the instance at record stepping the same bid change on the same row, or a current bid the row to its ladder state, and re-issue the P11 grain: one bid-ladder rule stepped a defensive term more than one rung from the recorded step 7 hold notice. v4.0 Section 12 rule P3 up six times against a stale target, from a recorded ladder state with no human confirmation. resolves the collision in the audit's favour: the floor of \$0.62 to a deployed \$1.10, a premium of audit value wins and automation resumes from \$0.48 per click across four rows carrying 380 clicks it. per week, or \$182.40 and 11.3% of that week's overspend. Nothing alarmed on it because every individual step sat inside its own limit.

F12 A placement verdict or a rung staged from Void the week's placement read for the v4.0 Section 3 family F10 names truncation as the a week where T3 Truncation Flag campaign and run the budget lever first. The step-zero precondition the previous file generation

[page 13 of 18]

A placement decision taken on (G3) fired, or on a row whose T2 Flag: read resumes on the next untruncated day; the could not see at all, and v4.0 Section 5 gate G3

a truncated or quarantined Zero-Quarantine is set. rung does not deploy from voided evidence. states the mechanism: a starved campaign's rates

week describe its worst hours rather than its potential. On

placement data the distortion is worse than on

blended data, because truncation ends the day at

whichever placement was still serving.

Escalation. F6 and F10 escalate on discovery, both because they are open authorizations rather than closed losses: F6 is a ceiling breach waiting for an auction and F10 is a

premium already being paid on unread traffic. F1 recurring escalates with the raiser attributed. F11 escalates to the PPC Lead with the T6 chain and the #37 rule identity

attached. Any mode occurring twice in consecutive cycles on the same scope escalates regardless of its individual routing. Everything in the human-mandatory tier is confirmed

per the Authority Matrix v2.0 before deployment, without exception.

8. Outputs and Artifacts

Gates clear with artifacts rather than assertions. Each row names the record, where it is stored, how long it is kept, and who consumes it next.

Artifact Produced by Storage and retention Downstream consumer

The weekly placement read P1 T4 Placement & Bid, one row per oversight keyword per P3 eligibility, P8 candidacy, P12 drift routing; #26 as the bleed evidence its

placement, with the cycle line on the tab header; append- cut sequence orders by loss dollars; the T7 row for any group the read could

only log life. not see.

The ceiling worksheet P2 Attached to the change it gates, referenced from the T6 row; #38 as the input its pre-upload effective-CPC scan checks against; #23

retained for the life of the campaign so the prior authorized where the change is tested against a push envelope; the next raise on the ceiling is recoverable. same placement, which is sized from the last cleared worksheet rather than from memory. This is the modifier change's gate artifact.

The staged modifier change P3, P4 Staged set in #38 format carrying verdict, scenario ID, this #12 for the campaign-level deployment; #38 for staging and the T+1 diff; V5

SOP and procedure number, worksheet result and proof scoring at the named review date.
ratio; T6 row with the prediction.

The dated rung P5 T4 Ladder State (step / date) and Next Step #12 P4 and P11 for execution; #25 as the priced step for its scheduled date;

Date ; one T6 row per rung, each independently scoreable, #26 P3 and P8; #35 where the rung doubles as an incrementality step; #37 append-only. as the hold notice's subject.

The exit bid P6 T4 Placement-CPC Exit Bid \$, re-derived at each P5 step 5, which lands the terminal rung on it; #23, whose taper handoff is

rung with the prior value retained in the T6 chain. aimed at it; #30, which seeds graduated exacts from it where a related floor exists on the root.

The floor register entry P6 T4 Discovered Floor \$ with its date and Next Step #25 at the transition exit; #26 for drift restoration, cohort construction and

Date ; the competitive context note carried on the final the harvest-read baseline; #12 for the drift restore value; #14 as the Surface rung's T6 row until T4 has a column for it. Retained for the 3 defense floor; #35 as the defended value. This is the taper's gate life of the term in the register. artifact: a descent is not complete without it.

The register diff P7, P14 Produced per cohort at the quarterly probe and monthly #26, which reads it to compute the released difference at its monthly harvest

across the product; filed against the T4 rows it moved. read and which cannot verify a release without it; P14 for the passed-re-test-date finding.

The bleed response record P8 T6, naming the lever taken, the levers not taken, and the The #26 cut sequence, which executes placement cuts in this SOP's lever

computed comparison that decided between them. order; the F3 audit line; V5 scoring on the lever's prediction.

The standup record P9, P10 T3 with both declared roles for a split; T6 with the combined- #12 P1 for the two creations; #29 for the walls in the same session; the

efficiency prediction against the pre-split blend. registry, so the pair never reads as duplicates at #12 P7.

The audit finding list P11 Ranked by bounded weekly exposure, filed with the #37, which holds its rules off the named rows; #12 for the strategy

quarterly cycle; the ladder-hold notice issued from it. corrections; the BM owner's review row.

The blocked-verdict log P13 Written against the T7 B7 row with term, campaign, verdict Data Operations and #39 as the feed's business case; P13 step 7 as the

string, supporting evidence and dollar consequence; work order on the day B7 flips; #26 and #23, which attach their own findings retained until the feed lands and the backlog is worked. to the same row.

The scored rung chain P14 T6 with V5 scores and the two separated error classes; Rule-level accuracy under v4.0 Section 14 rule V5, where the M3 step size

retained indefinitely as the step-size evidence series. is the rule most likely to enter review; the training corpus under rule V8; the v4.1 amendment brief.

Retention. No rule in v4.0 sets a retention period for the floor register, the worksheet set or the rung chain. The rows above state the instrument's own append-only life, which is

a property of the file rather than a governed period. The gap is amendment item AM-13 and it duplicates items already logged by the SOP-25, SOP-26, SOP-29 and SOP-30

sessions; v4.1 folds them once.

9. Skill-Conversion Block

Block Specification

Input bindings By catalog ID: A1 four-period comparability, A3 pacing and in-budget, A4 sufficiency buckets, A5 threshold flags, A6 objective and per-term Expected CPC, A7 eligibility, A8 PAT rows first-class, B6 bid intelligence, B7 placement splits, B8 economics chain, B11 market context, D1 per-variation margin, D2 keyword-variation map, D6 sibling CPC view, D7 decision log and V5, D12 hourly dayparting (reserved, unbuilt). Field-level: T4 all twenty columns as the primary surface; T3 Bid / AG Default \$, Placement Mult TOS % , Placement Mult PP % , Bidding Strategy , Truncation Flag (G3) ; T2 Max Profitable CPC \$, Expected CPC \$, CPA Ceiling \$, Current Bid \$, Suggested Bid / Gap % , TOS Modifier % , TOS IS % , Marginal ACOS , Family Conflict , Term Owner (Family) ; T0 zone C for the economics spine and zone H for envelope headroom; T5 for the attribution read at P7 step 8 and P12 step 2; T6 for rungs, predictions and scores; T7 for the manifest. An absent feed suppresses its dependent execution class asymmetrically and the skill encodes the asymmetry: B7 absent suppresses every raise and no cut, and the skill never substitutes a campaign blended CVR for a placement CVR, an account Expected CPC for a per-term value, or a prior week's placement mix for an unread one, without writing the substitution onto the record.

Deterministic logic The effective-CPC form as a three-factor product, computed for every placement on the campaign rather than for the one being changed. The M4 minimum with each limb's availability tested and recorded separately. The compliant-pair solve: given a target modifier and a governing ceiling, return the base at which the pair clears, or return BLOCKED where no base above the \$0.50 floor clears it. The per-placement ceiling as contribution multiplied by that placement's own CVR, resolved into the four F1 CPC-ratio bands. The rung derivation as current value multiplied by 0.90, tested against the L2 per-change limit and the \$0.50 floor, and truncated to the exit bid where a full step would overshoot it. The multi-step decomposition where a target cut exceeds the L2 limit, returned as a dated program rather than as a single step. The exit bid as blended placement CPC minus the decrement, re-derived per rung. The reversion rule, including the invariant that a failed probe never lowers the register. The three F9 decomposition classes. The drift test at 15 points over 4 weeks with the logged-change check. The bounded-exposure computation at P11 step 3. Scenario selection, verdict issuance and every threshold are v4.0's; the skill version-pins v4.0 and re-validates Section 5 on any increment before running a pass. No local thresholds, and the sixteen numerals at Section 13 are absent from the skill exactly as they are absent from this document.

Human checkpoints Executable with audit: the full T4 read at P1, worksheet computation at P2, the compliant-pair solve, rung derivation on schedule, exit-bid re-derivation, the F5 expired-justification scan, the drift test, the bounded-exposure ranking at P11, register diff assembly, and V5 scoring drafts. Human-mandatory, drawn from the Authority Matrix v2.0 rows listed in Element 11: every raise deployment and every arbitrage or split standup, per the matrix rows "Structure creation (all SOPs)" and

[page 14 of 18]

"Bulk deployments (#38)"; every re-ladder probe on a defended head term, per "Weekly verdicts (established US)"; every rule suspension arising from F11, per "Automation rules (#37)"; the incrementality rung handoff, per "Incrementality tests"; any restore step on a shared term, per "Sibling declarations (S-series)"; and every proof-ratio clearance until AM-1 lands, because the bar it would be cleared against does not exist. v4.0 Section 12 rule P5 places every 500-plus search-volume action, every gate failure, every structural change and every bid move above 25% in the tier regardless. Cadence Weekly: the T4 read, the F5 expired-justification scan, drift detection, and rungs falling due on their dates inside the Monday cascade, deployed Tuesday per #38 with the T+1 diff Wednesday. Per rung date rather than per week where a descent schedule from #25 sets other dates. Monthly: the floor register diff. Quarterly: the

P7 floor re-tests on the governed V4 cycle, with the P11 bidding-strategy audit riding the same cycle. Same day: a G4 eligibility flip, a G6 RED read, a margin-table version change firing S-A9, and a taper release arriving from #23. Every cycle while B7 reads MISSING: P13.

Alarm specification F1 to F12 each carry a detection signal from Section 7, an evidence-line format of row identifier, signal value, the v4.0 rule or gate reference, and the artifact where the signal was read, plus a routing tier. Same day to the operator queue: F6, F10, F11, F12. Into the weekly pass worklist: F2, F4, F5, F8. To the PPC Lead with the T6 chain attached: F1, F3, F7, F9, and any mode recurring in consecutive cycles.

Outputs Drafted ceiling worksheets with both effective ceilings and the limb-availability statement; drafted compliant pairs; drafted rung schedules with dates and hold predictions; drafted multi-step programs where L2 requires them; drafted exit bids and register entries with their context notes; drafted bleed responses in lever order with the shared-base comparison computed; drafted register diffs; drafted blocked-verdict log entries with their dollar consequences; the P11 finding list ranked by bounded exposure with the ladder-hold notice; F1 to F12 alarms with evidence lines; and the V5 scoring draft for rows due this cycle. Every draft is queued for confirmation at the tier the matrix names; nothing in this list deploys itself, and the skill stages no modifier increase while T7 reads B7 MISSING.

10. Version Governance

Owner: PPC Lead.

10.1 Review triggers specific to this component

- Any v4.0 increment touching Section 9 rules L1, L2, L3 or L8, or Section 10 rules M3, M4, M7 or M10, or Section 3 family F9, or Section 8.8 scenarios S-P5 and S-P6, or Section 8.4 S-R7. Each is load-bearing in a named procedure and the Section 5 mapping re-validates against the new verdict set before any pass runs on it.
- Any amendment item at Section 13 landing in v4.1. AM-1 in particular converts P3 step 2 from a recorded ratio routed to a human into a mechanical test, which is a procedure change rather than a wording change. AM-3 does the same for P3 step 6.
- Dependency D3 landing. The placement feed converts P1 steps 4 and 5 from suppressed to standing, retires P13, and makes P3 executable for the first time since this component was written. It also opens the blocked-verdict backlog, which P13 step 7 governs.
- Dependency D1 landing. The per-variation margin table removes the G9 PROVISIONAL tag from every ceiling in this document and makes the worked examples reproducible on live economics rather than on the training base.
- Field group B6 landing. It restores the family F9 auction-moved signal and makes the P7 step 8 decomposition computable from the bid-gap series rather than only from the product log.
- Any Amazon placement-control change: a new placement group, a change to the modifier range, or a change to the dynamic-raise behaviour increments this SOP before use, because all three sit inside the effective-CPC form that every worksheet computes.
- Any change to #23's push worksheet or its authorized CPC envelope, since P2 step 4 dereferences it, and any change to #25's descent schedule or #26's cohort construction, since P5 and P7 execute against both.
- Resolution of Ledger conflicts C9 or C10, which change gate naming in Element 3 and verdict strings in Element 5 respectively.
- Quarterly at minimum, alongside the floor re-tests and the bidding-strategy audit that ride the same governed cycle.

10.2 Change note, v1.0 to v2.0

Supersedes v1.0 in full. Procedures expanded from 8 to 14: the raise and the cut are separated into P3 and P4 because they have opposite proof requirements and opposite legality under the current feed state; exit-bid discovery and the floor register are separated from the ladder into P6 because the register is consumed by four components that never run a rung; the quarterly re-test becomes P7 in its own right; drift detection becomes P12; and P13 and P14 are new. Failure modes expanded from 7 to 12, with all seven

v1.0 modes and their incidents carried forward and now sourced to the document that records each. Worked examples expanded from 1 to 4. Elements 11 and 12 added, neither having any model in the corpus. Verdict mapping expanded from 7 rows to the full 25-string closed vocabulary with the v4.0 Section 13.1 cross-map and the campaign subset, and six locally minted v1.0 strings retired at Section 5.2.

Three corrections of substance. First, the v1.0 ladder worked example is corrected to M3 compliance: its first rung of \$2.60 to \$2.45 is a 5.77% step against the roughly 10% per week M3 sets, and the exemplar SOP-12 v2.0 had already corrected the same ladder against this document at \$2.60 to \$2.34 then to the \$2.24 exit reference. Section 6

WE-2 carries the compliant version and the anchors otherwise reconcile with v1.0 and with the Economics Worked-Example Set case E5. Second, three numerals v1.0 attributed to v4.0 do not exist there and are retired rather than carried: the modifier proof ratio (v1.0 P3.1, "the v4.0 factor"), the per-placement traffic minimum (v1.0 Section 3, "the v4.0 traffic minimum PER PLACEMENT"), and the modifier step size (v1.0 P3.1, "the +25% top-of-search pattern and its variants live there"). All three are now AM-1, AM-2 and AM-3. Third, v1.0's two-window hold and single re-probe after a reverted rung are retained as corpus practice and re-labelled as ungoverned rather than presented as mechanics, at AM-8 and AM-9.

Thresholds carried forward from v1.0 with their traces re-verified: the \$0.05-0.10 placement-CPC exit decrement (v4.0 Section 3 family F9 and Section 9 rule L2); the roughly 10% per week taper step, the 50% per-change limit and the \$0.50 bid floor (rule L2 and rule M3); the quarterly floor re-test or logged competitive shock (rule M3); the three-factor effective-CPC form and the 0 to 900% modifier range and up-and-down raise (Strategic Doctrine v1.0 Section 6.1, platform parameters, now logged at AM-11 as absent from v4.0); the doctrine posture of fixed for pushes and down-only as default (Doctrine Section 6.3); and the split-exact legitimacy test (Doctrine Section 5.2 case 2).

Incidents carried forward: the raise without proof, the premium on an unread placement, the flat cut on a bleed, the undated descent, the immortal premium, the up-and-down ceiling breach, and the unrecorded floor. Registered in the Supersession Map at publication.

11. Operator Ramp

Built from the PPC Authority Matrix v2.0 Section 1 rows that name this component's decisions. Each row becomes a competence the incoming operator demonstrates, and the checkpoint holder named in the matrix is the person who reviews in week 2 and audits in week 4. Note the matrix vacancy rule: an unfilled role's decision right reverts to the PPC Lead and its execution to the PPC Manager until filled, so a ramp in progress never leaves a rung unowned. Scope assignments follow the Transition Brief v1.2 stack, which the matrix records.

Authority Matrix row Rights on that row Week 1: observe Week 2: supervised Week 4: unsupervised, audited Weekly verdicts (established US) PPC Lead decides; incoming Shadow the full P1 read on two Run P1 on assigned scope and draft Own the weekly placement pass on operator and PPC Manager products. For every placement every placement verdict with its assigned scope. The brand-management execute; brand-management resolving to a class, state which of the scenario. The PPC Lead confirms owner's review in the daily huddle is owner reviews; CEO three F9 classes it is and which before staging. Every UNREAD mark is review rather than veto; disagreements informed. scenario governs, without issuing a defended out loud. route up, not sideways. verdict.

Bulk deployments (#38) PPC Lead reviews; incoming Read one staged rung set and its T+1 Stage rungs and modifier cuts; the Stage and reconcile unsupervised. A operator and PPC Manager diff. Identify which reason strings carry checkpoint holder reviews the landed-value mismatch on a rung routes execute; launch operator a scenario ID and which would fail reconciliation line by line and rejects same day rather than being absorbed executes own scope. staging. any single-step cut exceeding the L2 into the next rung. per-change limit, which is the fastest

way to learn that the limit shapes the schedule rather than the ambition.

Structure creation (all SOPs)

[page 15 of 18]

PPC Lead decides or per Read one completed placement-split Run the P10 candidacy test on one live Run candidacy and hand over standups

SOP scope; incoming pair in T3 and state both declared divergence and produce the sited base on assigned scope. Every raise operator, launch operator and roles, then find the four legitimate bids and sized modifiers for both roles. deployment and every standup remains

PPC Manager execute. duplicate roles at SOP-12 Section 1.2 The checkpoint holder confirms before human-mandatory at deployment and say why this pair is not a duplicate. the standup reaches #12. regardless of ramp stage.

Automation rules (#37) PPC Lead decides; incoming Read one rule-change log against a Issue the P11 step 7 ladder-hold notice Own the hold notice each cycle. Audited

operator executes and live ladder and identify which rows the for the cycle and reconcile it against specifically on whether the notice reviews; PPC Manager rule would touch. State what a runaway the rule register. Draft one F11 preceded the rungs rather than followed

executes; CEO informed. looks like in the T6 chain. suspension with the dollars attributed. them, because a notice issued after a rule fires is a report rather than a control.

Incrementality tests PPC Lead reviews; incoming Trace one completed incrementality Stage one incrementality rung and Stage rungs and hand off unsupervised.

operator decides and rung end to end: the step, the two- hand the measurement design to #35 Audited on whether a confounded executes; CEO and brand- week capture window, and where the with the confound references. The PPC window was passed over as clean, which

management owner informed. continue-or-restore call was made. Lead reviews the handoff, not the is the worse failure than reporting no

arithmetic. read at all.

Sibling declarations (S-series) PPC Lead decides; incoming Read one family board and compute Supply the X6 items 2 and 3 evidence Detect and hold shared-term restore

operator executes; brand- each sibling's effective top-of-search for one live divergence: multipliers, steps at G11 without prompting, and

management owner reviews; CPC from base, modifier and bidding effective CPCs and strategies. The supply the effective-CPC evidence

CEO informed. strategy. State why the base-bid PPC Lead makes the ownership call at unasked. comparison is the wrong one. #33.

Event plans (#32) PPC Lead decides with the Read one armed window and identify Suspend an affected cohort's probes at Operate ladders and probes through a

deals calendar owner; which probes and which modifier P7 step 2 under review, and re-date full window and resume unsupervised, operator executes; brand- changes it suspends and for how long, each one. Suspending without re- escalating only plan deviations. management owner reviews. including the 14-day guard. dating is the finding.

v4.0 threshold changes CEO decides; PPC Lead Read Section 13 of this document and, Draft one amendment entry from a live Raise amendment items as proposals

executes and reviews; for three of the sixteen items, state instance, with the evidence attached with evidence. An operator who writes a

everyone else informed. what the procedure does today in the rather than the opinion. missing threshold into a worksheet rather absence of the numeral. than into this register has misunderstood

the whole document, and this is the audit that catches it.

11.1 Live-defense questions

These test the reasoning behind the thresholds rather than recall of their values. The operator answers cold, with the workbook open and the criteria document closed. An

operator who can recite the \$0.05-0.10 decrement and cannot say why the ladder aims below the blended CPC rather than at it has not passed.

1. A campaign reads 38% blended ACOS against a 35% break-even and looks like a modest correction. Explain how the same campaign can contain a placement losing \$140 a week, and name the one number you compute per placement that makes it visible in one line.
2. Why does the exit bid sit below the blended placement CPC rather than at it? Say what bidding exactly at the blended CPC would buy you, and what the decrement is actually purchasing.
3. You cut a bleeding placement's modifier and it still bleeds. The obvious next move is to cut the base bid to the bleeder's own CPC. State the arithmetic you must run before taking that lever, and say what result would stop you taking it.
4. A modifier raise and a base-bid cut are staged as one set. The deploy window is tight and someone proposes pushing the modifier today and the base cut on Thursday.

Explain, in dollars per click, what those two days cost.

5. A rung slips. Explain why the register does not move, and say what would be true about next quarter's probe if it did.
6. Two siblings run the same term at the same base bid. Explain how they can still be paying materially different prices per click, and name the three settings you would read in order.
7. Top-of-search share reads 18% against a 30% floor and the placement feed is not available. State what verdict issues, what part of it you can execute, what you write down about the part you cannot, and why writing it down is worth doing rather than an admission of defeat.
8. You need to cut a bid from \$3.10 to \$0.84. Write the step program and say why it takes two weeks. Then say what that implies about what a day-7 read of the cut is allowed to conclude.
9. Nobody proposes an \$11.00 click, yet the audit finds them. Explain how one gets assembled, and name which of the three contributing settings is the one most likely to have been correct when it was set.
10. Name three things this SOP is forbidden to decide, and say which component decides each.
12. Interface Contracts

Twenty-two handoffs, split upstream and downstream. Every row states what happens when the handoff does not arrive, because that is the clause the corpus has always been missing. Sources: the Cross-Reference Ledger row for component 28 and the rows of every counterpart named, cross-checked against Master PPC Data Architecture v1.0 for artifact names and tab locations. Substitution stated: SOP-25 was attached to this session at v1.0 while the Ledger records it at v2.0, and SOP-24, SOP-32 and SOP-33 are being rebuilt concurrently in Wave 5A. The seams for all four are therefore written from their Ledger rows rather than from procedure numbers, per the Ledger's own v2.5 correction-register guidance that a Wave 5A session should expect those numbers to move.

12.1 Upstream

ID Counterpart Artifact and timing Non-arrival consequence

I-1 Data Operations and Field group B7, placement splits at TOS, ROS and PP grain, from The highest-consequence non-arrival on the upstream side, and it is the live state.

#39 the TOS placement report; weekly, landed with the Raw Pull Pack Every placement is UNREAD, every T4 performance column is unreadable, and every

Monday 06:00. modifier increase is suppressed account-wide. P1 runs its structural half only and P13 governs the cycle. Cuts, rungs, exit bids, floors and probes are unaffected, which is why the suppression is stated as asymmetric rather than as an outage. This is Ledger dependency D3, OPEN.

I-2 Data Operations and Field group B6, suggested bid and bid gap per row; weekly. The family F9 auction-moved signal, gap widening at or above 25% over 4 weeks, cannot

#39 compute. A cohort that slips as a whole cannot be separated into market movement and self-inflicted overshoot on bid evidence, so P7 step 8 falls back to the T5 product log and the competitor block and the record says so. Reads MISSING.

I-3 #41 economics The per-variation margin table refreshed into the T0 zone C Every ceiling in this document carries the G9 PROVISIONAL tag and every worked

named cells; monthly and on event, before the Monday packet- example runs on the training base. Worksheets still clear and rungs still derive, but no

ready read. ceiling is defensible on live economics, so a BLOCKED verdict cannot be argued against on the numbers. Ledger dependency D1, OPEN.

I-4 #23 At taper release: the achievement date, the push bid, and the The ladder starts from an unstated bid and the first rung is guessed, which discards the

worksheet's authorized CPC envelope. On any in-flight modifier floor the ladder exists to discover. In flight, this SOP cannot test a proposed modifier

request during a live push: the same envelope. against a push envelope it does not have, so the ceiling worksheet falls back to the maintenance ceiling and blocks changes the push legitimately authorizes.

[page 16 of 18]

I-5 #25 The dated rung instruction from the descent schedule, per term, Rungs have step values and no dates, which is the F4 condition arriving from outside.

with its scheduled date; weekly through the descent. The descent stalls at whatever value the last rung reached, and because that value is not

a discovered floor it enters no register, so the stall is invisible until the next quarterly re-test finds a term with no re-test date.

I-6 #26 The cohort re-ladder instruction with its member terms and Floors go unre-tested past their quarterly date, which is that component's floor-rot failure

recording quarter, quarterly; and the bleeding-placement cut mode detected from this side and quantified by the register diff that did not run. Inside a

decision inside a cut sequence, same day on a declaration. code red the placement lever of the cut sequence has no derived step, so the sequence

either skips the lever or takes a flat cut, which is F3.

I-7 #35 The continue-or-restore call at the end of the two-week capture The rung holds where it is and no floor is recorded, because a floor recorded without the

window on any rung that doubled as an incrementality step. capture read is a bid with a label. Where the window was confounded this SOP hands the

void and the T5 references rather than a reading; a confounded read passed over as clean is the worse failure.

I-8 #32 The armed-window declaration with its dates and affected G10 exclusion cannot be applied, so probes run inside windows and discover floors that

campaign list, at arming; and any in-window allowance naming a dissolve when the window closes, and modifier changes deploy against event-inflated

modifier change. placement mixes. Where the window is known but the plan is absent, P7 step 2 suspends the affected cohorts rather than proceeding on an assumed window length.

I-9 #33 The family ownership declaration on shared terms, current before G11 bars the restore and the term holds at its current bid. Two siblings restoring one term

any restore step. bid against each other with the family's own money, which is the divergence pattern the portfolio layer exists to catch. The escalation goes to #33 and the hold has no local workaround.

I-10 #38 The T+1 verification diff on every staged rung and modifier set, the Deployed values are unproven, and on a ladder that is acute: a rung that slipped and a

day after deployment. rung that never landed produce the same reading at the review date and take opposite responses. Without the diff the reversion rule fires on evidence that may describe a deployment failure, and the register records a floor that was never tested.

I-11 #37 The rule change log covering the ladder window, and confirmation Automation steps the same bid the ladder is stepping and the T6 chain becomes

that named rows are held; on request at the hold notice. unscorable, because no rung's outcome can be attributed to the rung. This is F11, and

the recorded instance ran six unconfirmed steps against a stale target before anything alarmed.

I-12 #22 Placement evidence accumulated during launch, handed over at The standing owner rediscovers placement economics the launch already paid for, and

the standing-owner co-sign. the first modifier decision on the product is made with no history at a point where the term is most likely to need one. Ledger row 22 records this handoff from its side as its I-16.

12.2 Downstream

ID Counterpart Artifact and timing Non-arrival consequence

I-13 #12 The derived rung value with its date, the compliant base-and- SOP-12 P4 has no step to execute and the taper stalls at push pressure on an achieved

modifier pair, and the multi-step program where L2 requires one; rank, which is that document's F5 and the account's most expensive known failure class.

same cycle as the verdict, same day on a dated release. A pair delivered without both halves is worse than none: deploying the modifier alone

authorizes exactly the click the worksheet blocked.

I-14 #25 Rung mechanics, step values and exit targets for each scheduled The descent schedule has dates and no step sizes, so either the schedule stalls or a step

descent date; weekly through the descent. is derived locally, and a locally derived step is a threshold invented in a document whose Ledger row explicitly states it derives no step. Ledger row 25 records this as its I-4,

capped by dependency D3.

I-15 #26 The floor register format and its content, the register diff per cohort Drift restoration has no floor to reset to, the quarterly re-ladder has no register to probe

and per month, and the bleeding-placement lever order; quarterly from, and the monthly harvest read has no baseline against which a spend reduction can

for the diff, on request for the lever order. be verified as a real release. Ledger row 26 names the missing register as the highest-consequence non-arrival on its upstream side.

I-16 #23 The ceiling worksheet result on any in-flight modifier request The push either runs without the modifier the sizing assumed or takes it untested, and in

tested against the push envelope; same cycle as the request. the second case the authorized CPC exceeds the envelope the push was funded against,

which converts a sized push into an unsized one without anyone changing the worksheet.

I-17 #38 Staged modifier and rung sets carrying verdict, scenario ID, A staged row missing its scenario ID does not deploy, so the rung silently does not

procedure number, worksheet result and the next step's date in happen while the T6 row claims it did, and the next rung derives from a value that never

the reason string; Tuesday by default, same day on a dated landed. The T+1 diff is the only thing that catches it, which is why I-10 is a contract rather

release. than a courtesy. Ledger row 38 records the pre-upload effective-CPC scan as an arithmetic check against this SOP's worksheet, so an absent worksheet leaves that scan

with nothing to check.

I-18 #35 The staged incrementality rung and, on completion, the recorded There is no measurable step, so the test has no treatment and the capture read measures

floor as the defended value; on the test's schedule. nothing. A test designed against a step that was never staged returns a clean-looking null

that will be cited later as evidence the spend was incremental.

I-19 #14 The Surface 3 defense floor and the exit-bid reference for Brand broad runs at an undiscovered defense floor, which means the defense budget is

defensive brand broad; on discovery and at each re-test. spent at a price nobody tested. Per Ledger row 14 that component defers ladder

derivation, exit-bid discovery, modifier sizing and the Surface 3 defense floor here and derives none of them.

I-20 #30 The exit-bid reference and any related root floor for seeding a Seed bids are set without the ladder reference, so a graduate on a root where a floor is

graduated exact; same day as the graduation. already known starts above it and has to descend to a value the register already held.
Per

Ledger row 30 that component defers ladder derivation and exit-bid discovery here.

I-21 #37 The ladder-hold notice naming every row under a live ladder or a Automation and the ladder fight over the same bid. v4.0 Section 12 rule P3 resolves the

live worksheet; each cycle, ahead of the rungs it protects. collision after the fact in the audit's favour; the notice is what prevents it arising, and a

notice issued after the rungs is a report rather than a control.

I-22 T6 and the verification Dated rung predictions at every step, modifier-change predictions Unsourced actions do not accumulate, so rule-level accuracy on the M3 step size cannot

layer at #6 at every raise and cut, V5 scores at every review date, and the compute and this SOP's central mechanic can never be shown to be wrong. Every

separated error classes at P14 step 2. downstream component would keep executing a step size that has never been tested against its own predictions.

13. Amendment Register and Conflicts

13.1 Amendment register: sixteen numerals this SOP needs that v4.0 does not carry

Each item is a number or a rule identity the procedures above require and that no attached document supplies. None is written into this SOP as a threshold. They are queued

for v4.1 of the Criteria System, whose owner holds the decision right on threshold changes per the Authority Matrix v2.0 row for v4.0 threshold changes, where the CEO decides

and the PPC Lead executes and reviews.

[page 17 of 18]

ID The gap Where it bites What is needed, and the evidence available now

AM-1 No factor defines a placement CVR proof P3 step 2, P9 step 1, WE-4. Highest-priority item in this register. v4.0 Section 3 family F9 sets TOS-EARNS on top-of-search

ratio. CVR being materially above rest of search, and Section 9 rule L3 repeats the word without a number.

This is the test the document's first absolute rule depends on, so today every raise routes to a human because there is no bar to clear mechanically. SOP-28 v1.0 P3 asserted a v4.0 factor that does not exist; it is retired here.

AM-2 No per-placement traffic minimum. Section 3, P1 step 6. v4.0 Section 5 gate G2 sets 15 clicks in 7 days for a bid read, 100 clicks for a CVR verdict and 1,000

impressions for a CTR verdict, all at keyword and campaign grain. The UNREAD class this SOP operates has no governed bar. This SOP applies the G2 minimums to the placement row and states the substitution; v1.0 asserted a v4.0 per-placement minimum that does not exist.

AM-3 No modifier step size. P3 step 6, P4 step 5. v4.0 sets bid ladder steps at roughly 10% per week and budget steps at 20-30%, and sets no step size

for a modifier at all. Section 8.4 S-R3 names the top-of-search multiplier as a lever without a magnitude.

v1.0 attributed a plus 25% top-of-search pattern to v4.0 Section 8; no such pattern is there. Until it lands a raise is staged at the ceiling worksheet's compliant-pair output, which is an arithmetic result rather than a step.

AM-4 The M4 tolerance has no value. P2 step 5, WE-4. Section 10 rule M4 sizes the multiplier so effective TOS CPC is at or under the minimum of max

profitable CPC for the objective and Expected CPC plus tolerance. The second limb is not computable, so the minimum reduces to one limb and the reduction is recorded on every worksheet.

AM-5 The M3 slippage tolerance has no value. P5 step 8, P7 step 6. Rule M3 reverts one step on slippage beyond tolerance and states no tolerance. Every slip currently

routes to the human-mandatory tier with its magnitude in positions attached. Duplicates the taper slip tolerance logged by the SOP-25 session; v4.1 folds once.

AM-6 No rule selects within the \$0.05-0.10 exit P6 step 2, WE-2. Family F9 and rule L2 both set the band and neither states which end to use. The Economics Worked-

decrement band. Example Set case E5 and SOP-28 v1.0 both operate \$0.07. A five-cent spread on a floor is \$0.05 per

click of permanent premium or discount on every maintenance click the term ever takes.

AM-7 No product-page CVR norm. P1 step 5, P8 step 1. Section 8.8 S-P6 defines PP-LEAK as product pages absorbing spend below product-page CVR

norms; no norm exists in v4.0 at any grain. WE-3's product-page row is classed on its own ceiling ratio of 4.23 rather than against a norm, which works where the leak is severe and does not where it is marginal.

AM-8 No hold period for a reverted rung. P5 step 10, P7 step 6. Rule M3 reverts one step and stops. SOP-25 and SOP-26 v2.0 both operate a two-read-window hold.

The figure is corpus practice and this SOP applies it as corpus practice with the label attached, not as a v4.0 mechanic.

AM-9 No re-probe count after a reverted rung. P5 step 10, P7 step 6. SOP-26 v2.0 P3 carries a single re-probe then closure at the prior floor. v4.0 carries no count, so

nothing governs whether a term re-probes once, three times, or every quarter until it holds.

AM-10 No lag from a taper release or a bleed Section 2, P5 step 1, P8. The taper release carries a date; nothing states how long the first rung may trail it, and a rank held at

detection to the procedure that answers it. push pressure for the interval is spending at push economics with no objective. The same absence

covers the bleed response. Duplicates the Element 2 lag-table item logged by the SOP-25, SOP-30 and SOP-17/18/19 sessions; v4.1 folds once.

AM-11 No authorized modifier range and no P2 step 2, P11, WE-1, WE-4. Strategic Doctrine v1.0 Section 6.1 records the platform's 0 to 900% modifier range and its up-to-

dynamic-raise ceiling in v4.0. doubling dynamic raise at top of search. These are Amazon platform mechanics rather than decision numerals, and this SOP treats them as the domain of the arithmetic. They nonetheless appear in every ceiling computed in this document and in SOP-23's own worked examples, and the threshold authority carries neither. Needed: either their admission to v4.0 as platform constants or an explicit statement that the doctrine holds them.

AM-12 No CPC buffer figure for the 2026 auction. P11 step 6. Doctrine Section 6.3 instructs that push sizing carry a CPC buffer while the auction re-stabilizes and

reports category inflation of 18-27% against Q1 baselines. That is a market observation, not an authorized numeral. Sizing that consumes this SOP's ceilings currently runs unbuffered. Duplicates the item logged by the SOP-23 session; v4.1 folds once.

AM-13 Record retention. Section 8. No v4.0 rule sets a retention period for the floor register, the worksheet set or the rung chain. The register in particular is read across quarters by four components, so its retention is not a filing question.

Duplicates items logged by the SOP-25, SOP-26, SOP-29 and SOP-30 sessions; v4.1 folds once.

AM-14 Gate numbering. Section 3. Carried here because Element 3 executes against both numbering schemes at G9, G10 and G12. Duplicates Ledger conflict C9; v4.1 folds once. This SOP applies the naming convention C9 prescribes in the meantime.

AM-15 No persistence bar for split candidacy. P10 step 1, Section 2. Doctrine Section 5.2 case 2 requires divergent placement economics proven across qualified reads,

persistent rather than single-window, and states no window count. A split creates a second campaign, a second budget line and a second wall set, so the bar deciding when it is justified is structural rather than cosmetic.

AM-16 No horizon after which an expired modifier P4 step 2, F5. A modifier raised for a push that closed is a premium with no reason, and v4.0 sets no interval after

justification must be normalized. which it must go. F5 exists because nothing alarms on a premium that is simply still there. Duplicates the expired-justification horizon logged by the SOP-25 session; v4.1 folds once.

13.2 Conflicts found beyond the registered set

The Ledger register runs to C18 at v2.5 and the build brief pre-declares C1 to C7. The eleven registered conflicts beyond the brief's list were not re-discovered here; C3, C9,

C10, C17 and C18 bind this component and are applied as their resolutions prescribe. Seven further findings are reported below. They are labelled N1 to N7 rather than C19

onward, because the Ledger's v2.5 numbering rule states that a concurrent session returns its findings and the close-out assigns every C-series identifier. Wave 5A

has four sessions open at once and this is exactly the collision that rule was written after.

ID Finding Statement

N1 The dayparting lever carries two IDs. v4.0 Section 9 assigns L10 to dayparting and L8 to format substitution. Master PPC Data Architecture v1.0 Section 8 cites the

reserved hourly columns as "(L8)", and the Command Workbook TEMPLATE v1.0 T7 Manifest row D12 Hourly dayparting reads "L8 dayparting (reserved, do not build)". Same shape as C9, C12 and C17: an identifier divergence rather than a content one. It matters here because L8 is the SBV arbitrage lever this SOP executes the paired step-down for, so a skill reading L8 from the instrument would hold a dayparting flag against an arbitrage instruction.

N2 The placement verdict set is cited to the wrong Master PPC Data Architecture v1.0 Section 8 states the T4 placement verdict values "per S-P4", and the Command Workbook T7

scenario. Manifest row B7 lists the blocked classes as "All T4 placement verdicts; S-P4; M4 arbitrage". v4.0 Section 8.8 S-P4 is profitable

headroom with a budget step of 20-30%. The placement decomposition scenarios are S-P5 for TOS-VANITY and S-P6 for PP-LEAK. Two instruments point the placement verdict at a budget scenario.

N3 A fourth verdict vocabulary exists at placement The Data Architecture Section 8 and the Command Workbook T4 tab both carry Placement Verdict with the values

grain. PREMIUM EARNED, CUT MULTIPLIER and EXIT PLACEMENT. Only CUT MULTIPLIER is a member of the KDR Template v3.0

campaign subset; PREMIUM EARNED and EXIT PLACEMENT belong to no closed vocabulary in the corpus. An extension of C10

[page 18 of 18]

and C18 rather than a replacement: C10 resolves which vocabulary governs at keyword grain, C18 records that it is short by two members, and this records a whole grain the closed vocabularies do not cover.

N4 The instrument carries two modifier columns Strategic Doctrine v1.0 Section 6.1 states that all three placement groups are independently adjustable. The Command Workbook

where the doctrine carries three placement T3 Campaign Board carries Placement Mult TOS % and Placement Mult PP % and no rest-of-search column. A rest-of-

groups. search adjustment therefore cannot be staged from the instrument, cannot be audited in the T3 row, and would not appear in the

P2 worksheet's inputs. Every arbitrage structure in P9 sites its base at rest-of-search economics precisely because that placement has no modifier, so the omission is currently benign and would stop being benign the day a rest-of-search adjustment is used.

N5 The floor register has no column for its context Strategic Doctrine v1.0 Section 6.2 makes the competitive context note part of the register entry, and SOP-26 v2.0 P3 requires it on

note. every recorded floor because it is what makes a future slip diagnosable. Command Workbook T4 carries

Discovered Floor \$ and Next Step Date and no note column. P6 step 5 routes the note to the final rung's T6 row as a workaround, which keeps it retrievable and makes it a second lookup rather than a register field.

N6 One worked-example anchor is carried at two The Economics Worked-Example Set v1.0 case E5 states the maintenance premium on the \$2.24 floor at \$23.52 per week and

figures across three documents. \$3.36 per day. SOP-26 v2.0 WE-3 states the same case at \$23.40 and \$3.34. E5 reproduces exactly at full precision from a \$1.456

ceiling; SOP-26's figures reproduce from a ceiling rounded to \$1.46 before the subtraction. The difference is \$0.12 per week of rounding order and nothing operational turns on it, but the register value it describes is this component's artifact and a worked example that disagrees with itself across documents teaches an operator to distrust the arithmetic everywhere else. E5 governs; correct SOP-26 at its next increment.

N7 Two consumers restate the ladder mechanics Ledger rows 25 and 26 both defer rung mechanics, step derivation, exit-bid computation, the reversion rule and the floor register

this component owns. schema to #28, and both documents then carry those mechanics procedurally: SOP-25 v1.0 at its P5 and SOP-26 v2.0 at its P3

steps 4 to 8. The restatements are currently faithful, which is the problem: nothing detects the moment they stop being. Both restate the ungoverned two-window hold and single re-probe of AM-8 and AM-9 as though they were mechanics, which is how an

ungoverned figure acquires the appearance of a standard by repetition. Recommended resolution: both cite this SOP's P5 by reference and carry no step content, which is what their own Ledger rows already say they do.

Companions: Master PPC Decision Criteria System v4.0 (Section 3 families F1, F6, F9, F10 and F12; Section 4 rules X2, X6, X9, X11 and X12; Section 5 gates G1 to G12; Section 6 rules O6

and O8; Section 7 rule D7; Section 8.1 S-A3 and S-A9, Section 8.4 S-R6, S-R7 and S-R10, Section 8.5 S-M2 and S-M5, Section 8.6 S-D1, Section 8.8 S-P1, S-P5 and S-P6, Section 8.10 S-F3

and S-F8; Section 9 levers L1 to L4, L8 and L10; Section 10 rules M3, M4, M7 and M10; Section 12 rules P1, P3, P5 and P8; Section 13.1 closed vocabulary and Section 13.3 rule J7; Section 14

rules V1, V4 and V5) · Keyword Decision Record Template v3.0, sheets "Bands & Verdicts Reference" and "Campaign Decision Record" (the quoted verdict strings) · PPC Strategic Doctrine v1.0

Sections 5.2, 6.1, 6.2, 6.3, 13, 14 and 15 (the reasoning behind the rules and the platform arithmetic) · PPC Economics

Worked-Example Set v1.0 case E5 (the WE-2 anchors) · PPC Authority

Matrix v2.0 Section 1 (Elements 9 and 11) · Master PPC Data Architecture v1.0 Sections 4, 6, 7, 8, 9 and 12 and the PPC Command Workbook TEMPLATE v1.0 tabs T0, T2, T3, T4, T5, T6 and

T7 (instrument mechanics) · SOP-12 (campaign execution and the format exemplar), SOP-14 (defensive brand broad), SOP-22 (launch handover), SOP-23 (push envelope and taper release),

SOP-24 (recovery brackets), SOP-25 (descent schedule), SOP-26 (cohorts, floors and the cut sequence), SOP-29 (walls on splits), SOP-30 (graduation seeds), SOP-32 (armed windows),

SOP-33 (family ownership), SOP-35 (incrementality), SOP-37 (rule holds), SOP-38 (staging, the effective-CPC scan and the T+1 diff) · PPC Cross-Reference Ledger v2.5 (the component 28

row, dependency D3, and conflicts C1 to C18).