SOP-29: Negation Architecture Operations

Operating SOP, Component #29 of the PPC Operations Definition Register | Version 2.0 | 2026-08-07 | Owner: PPC Lead | Supersedes SOP-29 v1.0 (2026-07-24) in full | Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds

1. Verdict and Scope

VERDICT: negation is traffic routing, not cleanup. The negative lists ARE the fences that make the control gradient real, and every one of them is versioned structure carrying an evidence line, not maintenance debris. A campaign's negative list is part of its structure and is versioned with it, per Strategic Doctrine v1.0 Section 11. This SOP governs the three negation states, wall-level selection, the pre-load standard at every structure event, the steering walls that close the graduation chain, the reactive execution standards applied to the weekly routes, negative-ASIN standards, shared-list management, sibling and brand walls, the monthly wall audit, removal control, negation-state reconciliation, and the event and recovery overlays.

THE ONE GOVERNING SENTENCE FOR THIS COMPONENT. Every negative carries five things: a declared state (pre-load, reactive or steering, per v4.0 Section 9 rule L5), a wall level (negative exact, negative phrase or negative ASIN), an evidence line or event reference, a named owner, and a verification line proving it landed in the list it was staged against. A negative missing any of the five is structure debt. A negative nobody can explain is itself a finding, and it is the finding this SOP exists to prevent.

Version 2.0 supersedes v1.0 in full. Procedures expand 7 to 12, failure modes 5 to 12, worked examples 2 to 5, and the Section 5 mapping from a 4-row local table to the full 25-verdict closed vocabulary. Elements 11 and 12 are added; neither existed in v1.0 nor in any shipped SOP. Thirteen v4.1 amendment items are logged at Section 10.3, all of them thresholds v1.0 asserted or implied without a traceable source.

1.1 The Three Negation States

The states are v4.0's, defined at Section 9 rule L5 and carried as a per-term field at Section 3 family F11. They are not severity levels and not a sequence. They are three jobs fired by three events, and an operator who confuses them places the right negative at the wrong moment.

State	What fires it	When it is placed	Evidence it carries	v4.0 issuing reference
PRE-LOAD	Taxonomy-known irrelevance: the MKL already knows the class is wrong for this product (wrong material, wrong size class, wrong intent).	In the creation session, before the campaign has served an impression.	An event reference, not a performance line: creating SOP, creation date, and the taxonomy class fenced. Pre-load has no clicks to cite by definition.	Section 9, L5 (pre-loaded mode). Section 8.9, S-H2, "pre-load negatives from the taxonomy at launch so irrelevant syntaxes never spend".
REACTIVE	Proven waste on served evidence: a WAS% breach against the objective limit, or spend without sales on a term the relevance check has already failed.	Staged Monday in the weekly pass, deployed Tuesday per #38.	A full performance line: clicks, orders, window dates, source campaign, CPC, spend, and the relevance-check result.	Section 9, L5 (reactive mode). Section 3, F11 (WAS% limits). Section 8.9, S-H2 and S-H7.
STEERING	A control-level change that must not double-buy: a graduation, a sibling ownership declaration, a brand-defense assignment, or a phrase-to-broad dedup.	The same day the control change lands. Steering is mechanical at graduation, not discretionary.	The event reference plus the full source list: every campaign that served the term, not only the one that validated it.	Section 9, L5 (steering mode). Section 3, F8 (steering-negation state). Section 8.9, S-H1 and S-H4.

The state is written per term on the T2 campaign child rows in the column headed **Negation State on Term (child)** and rolled to the campaign in T3 group C5 under **Negation State (pre/reactive/steering)**. Permitted values are the three states plus **none**. A term whose exact instance is live while its sources read **none** is the GRADUATED-UNSTEERED condition of v4.0 Section 3 family F8, and it is the most expensive readable state this SOP is responsible for: the source keeps buying the term at uncontrolled CPC while the exact fights it with our own money.

1.2 The Three Wall Levels

Three negative match types exist and each fences a different shape of future traffic. Choosing between them carries a cost either way, which is why it is P1 rather than a footnote.

Wall level	What it fences	Correct use	The failure it causes when misapplied
Negative exact	Exactly one served term and its close variants. One term, one wall point.	Every steering wall. Every Route B reactive term. Every graduated exact held out of its sources.	Under-fencing: a drift family of 40 terms needs 40 wall points, the operator places 3, and the other 37 keep spending.
Negative phrase	Every current and future served term containing the root, in order. One wall point, an unbounded set of futures.	Route C drift clusters, sibling-owned roots, brand roots held out of generic discovery, and the stem wall at stem graduation.	Over-fencing: a root chosen by string prefix rather than taxonomy silently fences owned terms. Failure mode F5, computed at WE-3.
Negative ASIN	One product detail page, on the product-targeting surface only: auto substitutes, auto complements, and PAT.	Own-catalog pages held out of complements (the self-cannibalization wall) and failed conquest targets held out of substitutes.	Applied to auto close or auto loose, where it changes nothing: those groups serve queries, not pages, so the wall is a no-op that reads as coverage.

1.3 What a Wall Is Not

Four boundaries, each of which has produced a real defect class here. A wall is **not a bid lever**: v4.0 Section 9 rule L1 places negation fourth, after budget, bid and placement, so an operator reaching for a negative to fix a cost problem has skipped three levers. It is **not a cleanup pass**: negatives are placed at named events on named evidence, never swept. It is **not a relevance judgment made at execution time**: relevance is decided upstream in the MKL, and v4.0 Section 3 family F11 states the guard directly, that the relevance check always precedes the negative and a relevant non-converter routes to the fix queue. It is **not reversible without a case**: removal runs through P10 with a named case class and evidence.

1.4 Boundaries With Adjacent Components

Read against the PPC Cross-Reference Ledger v1.7 row 29. Each seam is stated so no negation decision falls into a gap and no component absorbs another's scope.

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, gate, scenario, lever and verdict, including the three L5 modes, the WAS% limits and the reactive rules this SOP executes against. This SOP sets execution standards and never sets a number; a threshold appearing here instead of v4.0 is a defect.
#30 Harvest and graduation	Decides which terms graduate and assembles the source list. This SOP owns the wall half of the atomic chain and returns the verification line to #30's chain record. The chain spans both documents; a change to either half increments both.
#14 Broad operations	Owns the weekly three-route search-term triage that generates reactive candidates. This SOP owns the standards those routes execute against and never re-runs the triage.
#13 Phrase operations	Owns the phrase-broad dedup wall build and audit and the phrase-side graduation handoff, per Ledger row 13. This SOP supplies the wall standard; the dedup decision is #13's.
#15 Auto operations, D5 split	Owns the four-group structure and detects and routes candidates at P1.5 and P3 Routes B-D. This SOP owns the prescribed pre-load sets those routes apply in the same session, per Ledger row 15.
#12 Exact campaign operations	Owns exact standups. This SOP prescribes the walls SOP-12 builds at every P1, P2, P5 and P7 structure event, same session, recorded in the T3 row.
#16 PAT and conquest	Decides target opens, rotations and exits and declares the failed-target list. This SOP places the resulting negative ASINs and never decides that a target has failed.
#33 Sibling arbitration	Owns candidacy detection, the composite and the ownership declaration. This SOP builds and moves the walls the declaration orders, same day, and holds the non-owner's defense posture intact.
#38 Bulk-file deployment	Owns staging format, pull, transcription, validation, upload, T+1 verification and rollback. This SOP owns what the wall is and where it goes; #38 owns proof that it landed.
#37 Automation rules	Owns the rule register and change control on the automation deployment path. Negation staged here never routes through a rule without #37's register entry.
#32 Deal and event window	Owns the event plan and its named actions. This SOP freezes non-plan wall changes inside an armed window per P12 and executes only what the plan names.
#31 Budget and envelope	COMPLETE by absorption; no procedural document exists (Ledger conflict C3). Where this SOP needs the envelope it names the instrument and tab (T0 zone H, Account Rollup A2), never a #31 procedure number.

2. Trigger Conditions

Every trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. A trigger firing without its procedure starting inside the stated lag is a failure mode, not a delay, and the lag is measured from the artifact's timestamp rather than from the moment a person read it.

Trigger	Detecting artifact	Max lag	Procedure	Evidence the trigger fired
Structure event	The staged creation block in the #38 change set from SOP-12 P1/P2, SOP-13 P1, SOP-14 P1 or SOP-15 P2, carrying its pre-load set as a unit per #38 P1 step 3.	Same session. The pre-load never deploys later than the campaign it fences.	P2 pre-load standard, applied through the P1 tree.	The T3 row for the new campaign shows a populated list attachment and a non-empty negation-state field.
Graduation approval	The SOP-30 P1 graduation log row reading APPROVE with its full source list attached per SOP-30 P1.1.	Same day, walls placed and verification line returned.	P3 steering walls, all sources.	Each source's negative list shows the entry and the line is written back into #30's chain record.
Weekly reactive routes	The SOP-14 P3 Route B and Route C output, delivered with evidence lines into the Monday cascade at T2.	Staged Monday, deployed Tuesday. Two consecutive passes at zero negatives with spend active is F12.	P4 reactive execution standards.	The staged set carries one row per Route B term and one per Route C cluster root, each with a complete reason string.
ASIN route	The SOP-15 P3.3 substitutes handoff, the #16 failed-target declaration, and own-catalog appearances in the served-ASIN feed for complements and loose.	Same pass as detection.	P5 negative-ASIN standards.	Complements and substitutes show the ASIN entries; the own-catalog scan returns zero own ASINs serving.
Shared-list edit	A version increment on any named shared list, or an F4 divergence alarm between two campaigns sharing a discovery role.	Same day for the re-attach; the divergence source is named before the session closes.	P6 shared negative lists.	The list version log shows the addition with its evidence line and every attached campaign reads the same version.

Trigger	Detecting artifact	Max lag	Procedure	Evidence the trigger fired
Sibling declaration	The SOP-33 P4 declaration record, a P6 re-arbitration changing the owner, or a P6.3 stewardship grant.	Same day the declaration date carries. A wall state disagreeing with a declaration date is F8 here and F3 at #33.	P7 sibling and family walls.	The non-owner's discovery campaigns show negative phrase on the owned root; the owner's carry nothing for it.
Brand assignment	A brand-defense plan change from the SOP-14 P4 structure, or a new brand entering the catalog.	Same session as the defense structure change.	P8 brand-term walls.	Brand terms return zero appearances in the served terms of every generic discovery campaign at the next pull.
Monthly audit	The calendar slot registered against component #4 (amendment AM-4: the slot itself is ungoverned).	The audit completes inside the month it opens; findings stage the day they are found.	P9 monthly wall audit.	The artifact exists with its points-checked count whether or not it found a breach. Zero-breach months are recorded, never skipped.
Removal request	Any proposal to delete a negative, from any source, including an automation rule proposing it.	No removal deploys in the session it is proposed.	P10 removal control.	The removal log carries the case class, the evidence and the confirming role.
Registry reconciliation	The T2 campaign child rows and T3 group C5 read against the served-term feed at each weekly pull.	Same pass. A GRADUATED-UNSTEERED row closes in the cycle it is found, per v4.0 Section 8.9 S-H4.	P11 negation-state reconciliation.	Zero rows reading exact-live with source state **none**.
Event window arms	The #32 event plan entering T2 with its dates, arming the event gate.	Same day the window arms.	P12 event overlay.	The armed-window record lists which wall classes remain live and which are frozen.
Recovery mode entry	A product entering Recovery Mode per v4.0 Section 8.1 S-A4 and Section 11 rules R1-R6.	The pass runs on the recovery cadence from the day of entry.	P12 recovery overlay.	The recovery triage record names the negation pass among its completed steps.

3. Preconditions and Gates

Every procedure runs only after these read current in the T7 Manifest. A stale or missing field group suppresses its dependent wall class rather than degrading it: this SOP builds no wall against a list it cannot prove is current, because a wall action against a stale list is a wall against the wrong world. The pass records what it could not see and continues on the classes it can.

3.1 Input Preconditions

Manifest field group	Requirement before any wall action	Observable that proves it passed
B5 Coverage counts	Registry counts per match type current in T2 block 8, so the graduated-term population and the source list are readable. Feed: campaign registry, weekly.	T7 row B5 reads PRESENT and the T2 #Exact / #Phrase / #Broad / #Auto columns reconcile to the T3 campaign count for the product.
Search-term feed	Served-term rows present for every campaign in scope with the per-campaign split. Absence suppresses P4 entirely for the affected campaigns.	The T2 campaign child rows carry a non-null Spend Share of KW % for the window, which computes only when the join landed.
Served-ASIN feed	Product-target served rows present for substitutes and complements, or P5 suppresses for those groups. Post-D5, evidence carries the group token and is never blended.	The T3 row for each auto group shows a populated Auto Group Split (D5) link and the child rows carry group tokens.
Graduated-term list	From the registry, current to the last completed harvest pass. This is the population P9 audits and the list P3 walls against.	The graduation log's last row date is inside the current cycle and its count matches the trailing registry delta.
Brand-term list	Current per brand across all five brands, including variants and misspellings the exact list cannot enumerate.	The list version date is inside the current month and P8's zero-appearance check returns clean.
S-series declarations	The SOP-33 declaration register current, each with its date and named re-arbitration triggers.	Every root appearing in two or more siblings' campaigns has a declaration row, or an open S-series finding inside its deadline.
Own-catalog ASIN list	Complete across both entities and all five brands, including child variations, since a child ASIN self-cannibalizes exactly as a parent does.	The list count equals the catalog count in the Economics and Context Pack for the current month.
Failed-target list	Declared by #16 per its exit decisions. This SOP places what #16 declares and never infers a failure from CPA alone.	Each entry carries the #16 exit verdict reference and its declaration date.
T3 list attachments	Each campaign's attached shared lists named and versioned in its T3 row, because list attachment is campaign structure.	Every campaign sharing a discovery role reads the same list name and version string.

3.2 Gates in Force

GATE NUMBERING NOTICE. The Criteria System v4.0 Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass" and the operating corpus number three gates differently. This SOP names every gate by its v4.0 name and number and gives the corpus alias at first use, so an operator cannot execute the wrong check from the wrong document. Reported as conflict C9 and amendment AM-11; not resolved inside this SOP.

Gate (v4.0 Section 5)	What it blocks here	Observable that proves it passed
G8 Packet completeness	The governing gate for this SOP. The manifest reports present, stale or missing before analysis; a wall class whose input list is missing is suppressed and stated, never built on assumption. v4.0 sets DEGRADED at under 85% of gate-relevant fields, producing flags only.	T7 completeness is recorded on the pass header and every suppressed wall class is named in the pass record.
G12 MKL hygiene **(the corpus calls the event freeze G12)**	Pre-load in full. Every analyzed term must exist in the current MKL with a syntax tag and relevancy score; untagged or irrelevant tracked terms route to MKL maintenance and taxonomy pre-negation, not to analysis. v4.0 sets refresh at bi-weekly and degrades the manifest past 30 days stale.	The MKL version date is inside 30 days and the pre-load set resolves to tagged taxonomy classes rather than free text.
G10 Event exclusion **(the corpus number for this is G12)**	Non-plan wall changes inside an armed window, and the reading of event-week rates as reactive evidence. v4.0 excludes event weeks and the 14 days following from trend verdicts.	The armed-window record in T2 shows the dates and the staged set contains zero non-plan wall rows dated inside them.
G2 Sample sufficiency	Any reactive negative staged on a sample below the v4.0 read thresholds: 15 or more clicks in 7 days for bid reads, 100 or more for CVR verdicts, 1,000 or more impressions for CTR verdicts. Below threshold the verdict is a test continuation, never a rate judgment. Twin sets pool evidence.	Every Route B row's evidence line shows a click count at or above the applicable v4.0 threshold for the judgment being made.
G1 Validation quarantine	Any wall staged on a quarantined row. Impossible metric combinations quarantine the row, produce a flag and no recommendation, and open a data-refresh task the same day.	Zero staged negation rows point at a term whose T2 zero-quarantine flag is set.
G11 Sibling conflict	Sibling wall placement while a contested root sits undeclared. Two siblings pushing one term is a portfolio decision, not two product decisions.	Every root with two or more sibling appearances has a declaration row, or an open finding inside its deadline with the #33 P2 guards in force.
G3 Budget truncation	Nothing here directly, and that is the point: it is named so it is not misapplied. A truncated campaign's WAS% describes its worst hours, and v4.0 rule L1 puts the budget lever ahead of negation, so a truncated campaign gets its budget lever before its breach is read as a negation case.	The T3 Truncation Flag (G3) reads clear on every campaign whose WAS% is being read as reactive evidence this pass.
G9 Price stability **(the corpus uses G9 for margin-table freshness; both are named)**	Reactive reasoning that leans on conversion economics while a price change inside 7 days marks CVR and ACOS reads provisional, or while the margin table carries a provisional tag.	The staged row's reason string carries the provisional tag where either condition holds, so the read is auditable rather than silently confident.

4. Step-Level Procedures

Twelve procedures. Time estimates assume one operator working from the Command Workbook on a single product of roughly 280 tracked terms with the week's feeds landed; multiply by product count for a full account pass. Every staged row leaves this SOP in the #38 P1 format, one row per change, reason string built as **[SOP-29][verdict][state and wall level][evidence reference]**.

P1. Select the wall level

3 minutes per wall decision

Runs inside every other procedure that places a negative. Stated first because the wall-level choice is where both the under-fencing and over-fencing failure classes originate, and because v1.0 left it implicit.

DECISION TREE. Four branches.
Step 1. Is the target a product detail page rather than a query? If yes, the wall is **negative ASIN** and it goes to P5. It is placed only in auto substitutes, auto complements, and PAT campaigns; placing it in auto close or auto loose is a no-op that reads as coverage. If no, continue.
Step 2. Does one specific served term need fencing, with no expectation of unenumerated siblings? If yes, the wall is **negative exact**. Every steering wall is negative exact, without exception, because a graduated term is one term and phrase-walling it would fence the modifier space the source is supposed to keep discovering. If no, continue.
Step 3. Does a root exist that (a) appears in every member of the cluster, (b) appears in no term this product or any sibling owns, and (c) survives the reverse test below? If yes, the wall is **negative phrase** on that root. If no root satisfies all three, fall back to negative exact on each member and log the cluster

as unbounded so the next pass re-tests it.
Step 4. The reverse test, mandatory before any negative phrase deploys. Query the trailing window's served terms for the proposed root across every campaign the wall will touch. Count the terms that contain it. Subtract the cluster members. Any remainder is a term the wall will fence that nobody intended to fence: if the remainder is not zero, the root is wrong and the operator returns to step 3 with a longer root. The remainder count goes in the reason string whether it is zero or not, so the audit can see the test was run.

1. Read the candidate from its source: a graduation record, a Route B term, a Route C cluster, a declaration, or a brand list entry.
2. Walk the tree above. Record the branch taken and the reverse-test remainder in the staged row's reason string.
3. Where the tree returns negative phrase, record the fenced member count as well, so P9 can later compute what the single wall point is actually holding.
4. Where two operators would plausibly disagree on the root, the longer root wins. Under-fencing costs the spend on terms that leak; over-fencing costs contribution on terms already earning, and the second is larger and harder to detect.

P2. Pre-load standard at every structure event

12-20 minutes per campaign created

Executes in the creation session, staged as part of the creating SOP's settings block per #38 P1 step 3. A creation deployed without its pre-load is F1 here and in the creating SOP. The sets below are prescribed here and applied by the creating SOP; that split is the Ledger boundary for rows 12-15.

Campaign being created	Prescribed pre-load set	Why this set and not another
Family broad **Family phrase**	Every graduated exact in the family, negative exact, one wall point each. All brand terms across both entities, negative exact on the enumerated terms plus negative phrase on each brand root. Sibling-owned roots per the current declarations, negative phrase. Known-irrelevant taxonomy classes for the family, negative phrase on each class root.	Discovery surfaces are the entry point for every drift class; the fence is built before the campaign serves, not after it has taught the algorithm what to buy.
Auto close **Auto loose**	The same keyword-side set as family broad. Own-catalog ASINs are **not** negated here.	Close and loose serve queries, not pages. An own-ASIN negative in these two groups fences nothing and produces a false coverage read at the next audit.
Auto substitutes	Failed-PAT-target ASINs from the #16 declared list, negative ASIN. Own-catalog ASINs, negative ASIN. No keyword-side pre-load beyond the brand set.	Substitutes is the conquest signal source per v4.0 Section 3 family F11; walling the targets #16 has already exited stops the group re-proposing them as fresh candidates.
Auto complements	Own-catalog ASINs, negative ASIN, in full including child variations. Failed-target ASINs, negative ASIN.	Complements buys pages bought alongside ours, which by construction includes our own catalog. This is the self-cannibalization wall and it is the one pre-load with no keyword analogue.
Exact campaigns	No keyword pre-load; exact serves what it is. Brand campaigns exclude competitor-brand terms only where the defense plan names them.	An exact campaign cannot drift, so a pre-load set on it is either inert or an error waiting for an audit to find.
PAT campaigns	Own-catalog ASINs, negative ASIN. Failed-target ASINs from prior rotations on the same target list.	Conquest on our own pages is the same money on both sides of the auction.

1. Open the creating SOP's staged creation block. Read the campaign type and the syntax family from the name segments.
2. Pull the current lists from Section 3.1: graduated terms for the family, brand terms, sibling-owned roots, own-catalog ASINs, plus the #16 failed-target list where the type calls for it. Check each version date before use; a stale list suppresses its own line of the pre-load, and the suppression is written to the pass artifact header naming its **T7 MANIFEST** row rather than worked around.
3. Resolve the known-irrelevant taxonomy classes from the MKL, not from recall. This is the G12 dependency: an untagged class cannot be pre-loaded because there is nothing to name in the reason string.
4. Walk P1 on each entry. Graduated terms and brand terms resolve to negative exact; brand roots, sibling roots and taxonomy classes resolve to negative phrase; ASIN entries resolve to negative ASIN.
5. Stage every entry as its own row in the creation block, so it deploys as a unit and a partial pre-load shows as a row-count delta rather than an absence.
6. Write the list attachment into the new campaign's T3 row and set its negation-state field to **pre-load**. A T3 row with an empty attachment field is the F1 detection signal.
7. Record the pre-load row count against the prescribed count for the type. The two must match or the difference is named before the creation deploys.

P3. Steering walls at graduation

6-10 minutes per graduated term

The coupled half of the SOP-30 P3 atomic chain. #30 decides and stands up, this SOP walls, and neither half is complete alone. A standup whose walls are not verified the same day is F2 here and F1 at #30.

1. Receive the term and its **full** source list from the SOP-30 P1.1 queue row: every campaign that served the term in the evidence window, not only the one whose evidence validated it. A term validated in broad and also served by auto close is one graduation with walls in both.
2. Verify the list against the served-term feed before placing anything: query the term across every campaign on the product and on declared siblings and compare the returned set to what #30 supplied. A campaign in the feed but absent from the list is added, and the discrepancy returns to #30 as a queue-assembly finding.
3. Place negative exact for the term in each source. Sources are typically family broad, family phrase, auto close, auto loose, and any sibling discovery campaign carrying the root per the current declaration.

4. Check each placement against the campaign's ad group structure. A campaign with two or more ad groups takes the negative at the level covering all of them: a keyword-level negative on a multi-group campaign leaves the other groups serving. This is the list-scope error class, and the one breach that passes a naive list check.
5. Stem graduations wall differently per SOP-30 P4: negative phrase on the stem in the phrase source, fencing every member, plus negative exact in broad. Members are not walled individually, because they are the standing coverage the phrase campaign exists to hold.
6. Stage all wall rows in one change set with the graduation record reference in every reason string, so the chain reads as one event rather than as scattered negatives.
7. Verify same day: each source's list shows the entry at the level step 4 selected. Write the verification line into #30's chain record; the pass is not complete until it exists.
8. Verify again at the next pull: the term is absent from all sources' served terms. A list entry that exists while the term still serves is a breach even though the list looks clean, which is why P9 checks both sides.

P4. Reactive execution standards on the weekly routes

25-40 minutes per product per pass

SOP-14 P3 sorts the window's served terms into three routes and hands Routes B and C here with their evidence lines. This SOP does not re-run the triage or re-decide a route; it applies the standards below and stages the result.

1. **Route B, proven waste.** One negative exact per term. The evidence line is mandatory: clicks, orders, window dates, source campaign, CPC, spend, relevance-check result. A row arriving without one returns to #14 rather than being completed here from the feed, because the term reaching this SOP without evidence is exactly the term nobody checked.
2. **Route C, drift clusters.** One negative phrase per cluster root, selected through the P1 tree with its reverse test executed and its remainder recorded. The evidence line carries the root, the member count, and the combined clicks and spend across members.
3. **The converting-term guard.** Never negate a term that converted in-window without a diagnosis note; conversion inside the window overrides the click rule pending the read. v4.0 owns the rule at Section 3 family F11 and this step is the execution guard that makes it enforceable. A staged negative whose evidence line shows in-window orders is held, not deployed; the hold is F3.
4. **The truncation check.** Before reading a WAS% breach as a negation case, read the T3 Truncation Flag (G3). A truncated campaign's waste percentage describes its worst hours, and v4.0 rule L1 places budget ahead of negation, so the breach routes to the budget lever first and re-reads next clean cycle.
5. **Sufficiency.** Each row's click count is checked against the v4.0 Section 5 G2 read thresholds for the judgment being made. Below threshold the term stays in its source under a capped, dated test rather than entering a negative list.
6. Stage every row in the #38 P1 format with reason strings built to this SOP's pattern. Deploy Tuesday with the pass.
7. Compute and record the pass arithmetic: terms fenced, clusters fenced, member count behind each cluster, window spend removed, and the resulting WAS% per campaign. The spend-removed figure is what makes the pass auditable; counts without dollars cannot be scored.
8. Write the counts and routes into **T6 DECISION LOG**, filling Level, Target, Rule / Scenario Fired, Change (prior -> new), Expected Outcome Window and Review Date. State the prediction as a WAS% figure with a date, never as a direction.

P5. Negative-ASIN standards

10-15 minutes per product per pass

1. **Own-catalog appearances.** Scan the served-ASIN feed for complements and loose each pass; any own ASIN returned stages as negative ASIN in that campaign the same pass. Include child variations, which self-cannibalize exactly as parents do and are what the catalog list most often misses.
2. **Failed conquest targets.** Place what #16 declares, when #16 declares it, into auto substitutes and into the PAT campaigns on the same target list. This SOP never infers a failure from a CPA read; the exit decision is #16's per Ledger row 16.
3. **Group discipline.** Negative ASINs go into substitutes, complements and PAT only. Read the group from the T3 column **Auto Group Split (D5)**. Post-D5 the four groups are separate structures per #15, so one placed on the wrong group is inert rather than merely misplaced.
4. **Evidence.** The reason string for an own-catalog wall carries the ASIN, the parent, the clicks and spend it drew in the window, and the catalog-list version. For a failed target it carries the #16 exit verdict reference and the declaration date.
5. Own-catalog walls are permanent structure and are not re-litigated at each audit. A failed-target wall is removable only through P10, on a #16 re-entry declaration.

P6. Shared negative-list management

8-12 minutes per pass, plus 20 minutes per new list

1. **Sharing is by rule, not convenience.** Two structures share a list where they share a discovery role: the trawler and probe pair at standup per SOP-15 P4, and the family broad plus family phrase pair, which share the graduated, brand and sibling sets. Doctrine v1.0 Section 11 states the basis: campaigns sharing a discovery role share a list so they cannot buy the same mistake twice.
2. **Naming and versioning.** Every shared list carries a name identifying its family and role and a version string that increments on every addition or removal. Additions log with their evidence lines exactly as standalone negatives do.
3. **One edit, everywhere.** A shared list is edited once and lands in every campaign it is attached to. Editing a campaign's own list to add something the shared list should carry is the divergence source, and it is how F4 begins.
4. **Attachment is structure.** It lives in the T3 row and is read at every creation and audit. A campaign that lost an attachment through a rebuild reads clean on its own list while missing everything the shared list holds.
5. **Divergence check.** Each pass, compare the entry counts and version strings of every pair sharing a role. A mismatch names the diverged campaign, the divergence count, and the direct edits that bypassed the list. The tolerance that should fire this alarm is ungoverned; see amendment AM-7.

P7. Sibling and family walls

10 minutes per declaration, same day

1. Open the SOP-33 declaration register and read the row for the root: owner, declaration date, composite table, named re-arbitration triggers. The non-owner's discovery campaigns receive negative phrase on the owned root; the owner's campaigns carry nothing for it.
2. **The non-owner's defense posture is protected, and protecting it is this SOP's job.** Ownership is of the root, not of the modifier space that names the sibling. The non-owner holds exact defense at floor on its own model-specific variants, meaning the root plus its size, count, or model

terms. Over-walling the non-owner out of its own name is failure mode F9 here and F4 in #33, and the reverse test in P1 step 4 is what catches it before deployment: run the test against the non-owner's own exact list as well as its served terms.

3. On a re-arbitration changing the owner, the walls move the same day the declaration changes. The old owner's discovery gains the wall; the new owner's loses it through P10 removal control with the declaration as the case evidence. Both movements stage in one change set so the family cannot sit half-walled.
4. On a temporary stewardship grant under SOP-33 P6.3, the wall moves with the stewardship and carries the reversion date in its reason string. A stewardship wall without a reversion date in the string is F10.
5. Twin sets carry shared negatives per v4.0 Section 8.11 S-B4 and Section 9 rule L6: one managed term, one bid posture, shared negatives. A twin pair whose lists diverge is read as one finding, not two.

P8. Brand-term walls

15 minutes per brand per month

1. Brand terms live in defense and nowhere else. Every generic discovery campaign across the product's family carries the brand set: negative exact on the enumerated brand terms, negative phrase on each brand root to catch the variants and misspellings the enumeration cannot reach.
2. The five brands are walled across both entities, not per entity: a brand term leaking into a sibling brand's discovery campaign spends against the family from the inside.
3. The defensive brand-broad campaign per SOP-14 P4 is the one structure that holds brand terms, and it carries competitor-brand negatives only where the defense plan names them. This SOP places what the plan names; the plan is #14's and #17-19's.
4. Competitor brands appearing in the brand campaign's served terms are conquest context and route to #16. They are not negated here on sight, because an incursion term is evidence before it is waste.
5. The check that proves the wall: zero brand-term appearances in the served terms of every generic discovery campaign at the pull. A brand term appearing anywhere generic is a finding regardless of what the negative lists show.

P9. The monthly wall audit

90-140 minutes per product

The audit is the artifact that makes every other procedure checkable. It runs on the calendar slot registered against component #4; the slot itself carries no v4.0 basis and is logged as amendment AM-4.

1. **Assemble the population.** Every term graduated in the trailing quarter with its full source list. The sampling floor is ungoverned (AM-3), so until it is, the audit runs the full trailing-quarter population rather than a sample: a sample without a governed rule is a judgment call wearing a statistic's clothing.
2. **Check both sides per wall point.** Side one: the list shows the entry at the level P3 step 4 selected. Side two: the source's served terms do not contain the term. Absent from the list is a breach; present in the list and also in served terms is equally a breach, because list scope errors happen and only side two catches them.
3. **Brand check.** Zero brand terms in any generic discovery campaign's served terms.
4. **Own-ASIN check.** Zero own-catalog ASINs in complements' or loose's served rows; any appearance stages a negative ASIN the same day and carries its window spend into the finding.
5. **Sibling check.** Every declaration's wall state matches its declaration date. Compute the lag in days for any mismatch, because the lag is what the fix is scored against.
6. **Shared-list check.** Every pair sharing a role reads the same list version.
7. **Compute the breach rate** as breaches divided by wall points checked and record it whether or not it is zero, alongside the **T7 MANIFEST** completeness score for the cycle. Zero-breach months are recorded, never skipped: an audit with no artifact is indistinguishable from an audit that did not run.
8. **Attribute every breach to its session**, the creation or graduation session that should have placed the wall. Attribution turns a breach list into a cause; without it the list repeats next month.
9. Stage all fixes the same day. The audit artifact is the points-checked count, the breach list with session attribution, the fixes staged, and the named root causes.

P10. Removal control

15 minutes per removal, never same-session

A removal requires a named case class. There are exactly four.
Case 1, sibling re-arbitration. The declaration changed and the outgoing non-owner's wall comes down. Evidence: the new declaration record with both composites.
Case 2, brand-strategy change. A brand term moves scope, or a competitor-brand negative is lifted because the defense plan changed. Evidence: the plan revision.
Case 3, erroneous negation with evidence. The wall fenced something nobody intended, and the evidence is the served-term or contribution collapse that proves it. Worked example WE-3 is this case.
Case 4, #16 target re-entry. A failed target is re-declared as a candidate and its negative ASIN comes down. Evidence: the #16 re-entry declaration.
A failed graduate is not a removal case. Per SOP-30 P6, a graduated exact that fails in exact gets a diagnosis, and the source walls stay regardless of the exact's outcome. Re-opening the source would buy the same clicks at less control.

1. No removal deploys in the session it is proposed. The case class is named and evidence attached first, which is what separates a controlled removal from a fence breach with a plausible story.
2. Removals log identically to additions: entity, wall level, case class, evidence, confirming role, date.
3. Removal is human-mandatory in every case; see Section 9 for the checkpoint and Section 11 for who holds it.
4. An unexplained removal found in the audit is reinstated pending review and escalates on first occurrence. The escalation clock is ungoverned; see amendment AM-8.

P11. Negation-state reconciliation

20-30 minutes per product per pass

1. Open **T2 KEYWORD COMMAND**, apply the OVERSIGHT filter view, and expand the campaign child rows. Read the column **Negation State on Term (child)**, which carries one of pre-load, reactive, steering or none, beside **Spend Share of KW % (child)**.
2. **The unsteered-graduate scan.** Filter for rows where an exact instance is live and a source campaign's state reads **none**. This is the GRADUATED-UNSTEERED condition of v4.0 Section 3 family F8 and the scenario at v4.0 Section 8.9 S-H4, which closes the leak in the same cycle it is found and reclassifies the source's spend on the term as structure waste in the audit. Note that the Data Architecture v1.0 Section 6.1 labels this filter S-H6; that is conflict C12 and the correct scenario ID is S-H4.
3. **The auto-dominance cross-check.** Where a priority term shows auto or broad carrying more than 40% of its clicks per v4.0 Section 4 rule X10, the finding is structure failure rather than a performance signal. It routes to #12 as ADD EXACT + RESTRUCTURE, with this SOP contributing the steering wall once the exact stands up; this SOP does not stand up the exact.
4. Reconcile against the **T3 CAMPAIGN BOARD** group C5 column **Negation State (pre/reactive/steering)**: a campaign reading clean at C5 while any child row reads **none** is a rollup defect, not a clean campaign.
5. Record the reconciliation counts in the pass artifact: rows scanned, unsteered rows found, rows closed this cycle, rows carried with a named reason.

P12. Event-window and recovery overlays

10 minutes to arm, 5 minutes per pass while armed

1. **Arming.** When the #32 event plan enters T2 with its dates, classify every wall class as frozen or live for the window and record the classification in the armed-window record.
2. **Live inside a window:** steering walls at graduation, because a graduation deploying without its wall is a leak that the window would amplify; own-catalog negative ASINs, because self-cannibalization is worst at event traffic volumes; and any wall the event plan names.
3. **Frozen inside a window:** reactive negation from the weekly routes and Route C cluster walls. Event-week rates are excluded from trend verdicts per v4.0 Section 5 G10, so event-week waste evidence is not evidence, and a phrase wall placed on event traffic fences a demand shape that will not exist in two weeks.
4. **Post-window.** The frozen reactive queue re-reads against post-window evidence rather than deploying from the queue as staged. The 14-day exclusion following the window applies to the re-read.
5. **Recovery Mode.** A product in Recovery Mode per v4.0 Section 8.1 S-A4 and Section 11 rules R1-R6 runs the negation pass on the recovery cadence rather than the weekly one, and the pass is a named step in the recovery triage record. Recovery tightens the human-review tiers per v4.0 rule R3, so the checkpoints in Section 9 tighten with it.
6. **Mode precedence** is v4.0's at Section 11 rule T8 and Section 12 rule P7: Recovery Mode over Event Mode over stage defaults. A product in Recovery that enters an event runs the event under recovery constraints, and this SOP's freeze list follows that ordering rather than the calendar.

5. Decision Points: The Full Verdict-to-Wall Mapping

All verdict issuance, thresholds, bands and scenario logic live in v4.0. This table maps the complete 25-verdict closed vocabulary of the Keyword Decision Record Template v3.0, tab "Bands & Verdicts Reference", to what this SOP does when each verdict lands. Verdict strings are quoted verbatim. Ten verdicts produce a wall action and are mapped individually. Fifteen produce none, and are grouped below by the reason they produce none, because knowing why a verdict is inert is what stops an operator reaching for a negative on a bid problem.

> **VOCABULARY SOURCE NOTICE.** The Criteria System v4.0 Section 13.1 "The Closed Verdict Vocabulary" carries an eighteen-item list whose strings differ from the KDR's on the two rows this SOP most depends on: v4.0 issues "CUT / NEGATE" where the KDR issues "NEGATE (PRE-LOAD / REACTIVE / STEERING)", and "GRADUATE + STEER" where the KDR issues "GRADUATE FROM DISCOVERY". The Data Architecture v1.0 Section 2 rule 8 and Ledger row 8 both state twenty-six. This SOP quotes the KDR's twenty-five, matching exemplar SOP-12 v2.0 Section 5, and carries the v4.0 string as a stated cross-map on the two divergent rows so neither citation is dropped. Reported as conflict C10.

5.1 The Ten Verdicts That Produce a Wall Action

Verdict, quoted verbatim from the KDR Template v3.0	Negation-architecture execution
NEGATE (PRE-LOAD / REACTIVE / STEERING) **v4.0 Section 13.1 cross-map: "CUT / NEGATE"**	The verdict this SOP owns end to end. The parenthetical names which of the three states is issued and therefore which procedure runs: pre-load to P2, reactive to P4, steering to P3 or P7. The mode is recorded on the row per v4.0 Section 3 family F11.
GRADUATE FROM DISCOVERY **v4.0 Section 13.1 cross-map: "GRADUATE + STEER"**	P3 in full, same day, in every source campaign attached to the term by SOP-30 P1.1, with the verification line returned to #30's chain record. The v4.0 string names the steering half explicitly, which is the half this SOP owns.
ADD EXACT + RESTRUCTURE	The steering half executes once #12 has stood the exact up: negative exact in every source that was carrying the term, same day as the standup. This is the resolution path for the auto-dominance read at v4.0 Section 4 rule X10 and for scenario S-R12.
CONSOLIDATE DUPLICATES	Carry every negation list the paused instance held that the survivor lacks, in the same change set as the pause. A survivor missing the loser's fences inherits its traffic and none of its protection; failure mode F11.
OPEN / ROTATE / EXIT CONQUEST	On EXIT only: the exited target enters the #16 failed-target list and this SOP places negative ASIN in substitutes and in the PAT campaigns on that list, per P5. Opens and rotations produce no wall.
FUND TEST (UNFUNDED PRIMARY)	No wall on the tested term. Where the test stands up a new campaign, P2 pre-load applies to that campaign in the creation session.
RANK PUSH: FUND SIZED PUSH	No wall on the pushed term. Where the push creates a dedicated campaign, P2 applies; where it follows a graduation, the P3 walls already stand and are re-verified rather than rebuilt.
RECOVERY PUSH	As the sized push, plus the pre-OOS wall set is re-verified before the push funds. Walls placed before an outage are the ones most often found stale at restock.

Verdict, quoted verbatim from the KDR Template v3.0	Negation-architecture execution
RE-POINT VARIATION	No keyword wall. The own-catalog negative-ASIN set re-checks against the newly targeted variation the same day: the backup's own ASIN must be walled out of complements exactly as the primary's was.
PAUSED-PENDING (REASON + RE-ENTRY)	No wall required, since a paused term cannot serve. Add one where the paused term sits in a campaign that also hosts phrase or broad ad groups, because those groups can still serve it.

5.2 The Fifteen Verdicts That Produce No Wall, Grouped by Reason

Reason class	Verdicts, quoted verbatim from the KDR Template v3.0	Why no wall, and what the operator does instead
A different lever owns it	"UNLOCK TRAFFIC" · "FIX BUDGET (TRUNCATED)" · "PLACEMENT FIX FIRST" · "BID DOWN TO PROFITABLE CPC; CAP SYNTAX" · "TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS"	v4.0 Section 9 rule L1 places negation fourth, after budget, bid and placement, so a negative on any of these rows means a lever was skipped. Two specifics: a truncated campaign's WAS% describes its worst hours, so a reactive case on it is held per P4 step 4 and re-read the next untruncated cycle; and a traffic deficit caused by an over-broad wall fencing the term's own traffic routes to P10 case 3, never to a bid lift.
A wall would destroy the evidence or the demand	"QUARANTINE: DATA REFRESH" · "CONTINUE TEST TO 50 CLICKS AT CAPPED BID" · "FIX CONVERSION FIRST; DEFER PUSH" · "INCREMENTALITY LADDER" · "STOP-LOSS: CONCEDE OR DEFER" · "LEAVE TO HARVEST"	Each row is a term the account is still buying information about, or still earning from. G1 blocks any wall on a quarantined row and pulls one already staged. FIX CONVERSION FIRST is the most common source of an incorrect reactive negative: a relevant non-converter routes to the fix queue and never to the negative list, per v4.0 Section 3 family F11. A wall inside a live incrementality measurement contaminates it, so it is held with its reason string until the measurement closes. A conceded term is not a fenced term; walling it prevents the deferred re-entry the verdict names. LEAVE TO HARVEST is converting demand below the graduation bar, and walling it removes that demand from the only surface funding it.
The walls already stand and are read at the audit, not at the verdict	"TAPER: RANK ACHIEVED, STEP BIDS DOWN" · "HOLD; DEFEND LIGHTLY AT RANK" · "DEFEND AT FLOOR" · "ENTER SBV ARBITRAGE"	Steady-state postures on terms the account already controls. Steering walls placed at graduation hold through a taper regardless of the exact's bid path, because the term's economics are the term's economics per SOP-30 P6. Existing sets are read for drift at the monthly audit, not re-litigated at each verdict. Where a defended term is brand, P8 confirms it is walled out of every generic discovery campaign so defense actually owns it. The SBV pivot changes which surface buys the term, not whether it is bought.

Campaign-level note: the T3 group C6 subset vocabulary carries **NEGATION PASS** as a campaign verdict. It is a routing instruction to run P4 on the named campaign, not a keyword verdict, and it never substitutes for a term-level verdict from the closed vocabulary above.

6. Worked Examples

Five examples: three clean cases and two that go wrong and get handled. No live operator economics are attached to this build, so all five use the standard fictional base declared in the PPC Economics Worked-Example Set v1.0: AOV $80.00, contribution $28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Maximum Acceptable 26.3%, CVR 5.2% unless stated. Every figure below is computed on the page; nothing is asserted that is not calculated.

WE-1 (clean). Pre-load failure found late: the satin product tracking cooling terms

The incident this example encodes. The Criteria System v4.0 Section 3 family F11 records it directly: a satin product tracking cooling terms it cannot rank and taking clicks on them is a pre-load failure plus an MKL hygiene failure, because the taxonomy knew and the campaigns did not. This is the account's named pre-load failure class and the reason P2 runs in the creation session rather than after the first read.

Setup. Satin 4pc, auto loose campaign, Objective Discovery, 28-day window. Campaign window spend $612.00. The served-term feed returns 6 cooling terms drawing 118 clicks at CPC $1.42.

> Cooling-term spend = 118 x $1.42 = **$167.56**
> Of the 6 terms, 5 carry zero orders across 104 of those clicks: 104 x $1.42 = **$147.68** of spend that bought nothing
> One term returned 1 order at $80.00, so contribution recovered = $28.00, and the window net on the cooling class = $167.56 - $28.00 = **$139.56 lost**
> Campaign zero-order spend across all served terms = $271.32, so WAS% = 271.32 / 612.00 = **44.33%**, above the Discovery limit of 40% at v4.0 Section 3 family F11, which fires the negation pass at Section 8.9 scenario S-H2
> The cooling class alone is 147.68 / 271.32 = **54.43%** of the campaign's entire waste
> Run rate = $167.56 / 28 days = **$5.98 per day**, or $5.98 x 365 = **$2,184.26 per year** if nobody places the wall

Wall selection (P1). Step 1: these are queries, not pages, so not negative ASIN. Step 2: six terms exist today and the taxonomy says the class is wrong, which means unenumerated siblings are certain, so not negative exact. Step 3: the root "cooling" appears in all six, and the satin line owns no cooling term. Step 4, the reverse test: served terms containing "cooling" across the three campaigns in scope total 6; cluster members total 6; remainder **0**. The root is safe.

Execution. One negative phrase on "cooling" in three keyword-side sources: satin family broad, satin auto close, satin auto loose. Three wall points fence six current terms and every future cooling term the taxonomy has not enumerated. The alternative, six negative exacts per source, is 18 wall points fencing only what already happened. The G12 MKL hygiene route runs in parallel: the cooling class is tagged in the MKL as out-of-taxonomy for the satin line so the pre-load set for the next satin creation carries it automatically.

Artifact. Three staged rows, one evidence line each carrying root, member count 6, combined clicks 118, combined spend $167.56, reverse-test remainder 0. Prediction logged: cooling-root appearances in the satin served terms fall to 0 at the next pull, and the campaign's WAS% falls below the 40% Discovery limit.

WE-2 (clean). Steering walls at graduation, all sources

Setup. SOP-30 approves "bamboo sheets split king" and hands the queue row here with three attached sources.

> Family broad: 11 clicks at CPC $1.72 = $18.92, 2 orders
> Auto close: 6 clicks at CPC $1.53 = $9.18, 1 order
> Sibling 6-pc broad: 4 clicks at CPC $1.66 = $6.64, 0 orders
> Combined: 21 clicks, 3 orders, $34.74 of source spend. Three orders clears the graduation bar at v4.0 Section 9 rule L6 and Section 8.9 scenario S-H1
> Blended source CPC = $34.74 / 21 = **$1.65**; the exact's seeded bid, carried by the graduation verdict and owned by #30 and #12 rather than by this SOP, is **$1.30**
> Cost of skipping the walls for one window = $34.74 - (21 x $1.30) = **$7.44**, or $7.44 x 52 = **$386.88 per year** of paying a 27% premium to buy traffic the account already owns at exact

Execution (P3). Step 2, the source-list verification, queries the term across every campaign on the product and on declared siblings and returns the same three campaigns #30 supplied: the list is confirmed rather than assumed. Step 4 checks ad group structure: the family broad carries two ad groups, so its negative exact is placed at the level covering both. The sibling 6-pc broad receives the wall because the declaration assigns the bamboo root to the 4pc; the 6-pc's own count-modifier variants are untouched, per P7 step 2.

Verification. Day+1: all three lists show the entry at the correct level. Next pull: the term is absent from all three sources' served terms. Both lines are written back into #30's chain record, which is what completes the atomic chain. Three wall points, zero breaches.

WE-3 (goes wrong, handled). An over-broad phrase root fences owned demand

Setup. Bamboo family broad, weekly pass. Route C hands over a 2-term drift pair, "cooler bag sheets" and "cooler king sheets". The operator stages negative phrase on the root **"cool"**, reasoning that it fences both members with one wall point. The P1 step 4 reverse test is not run.

What the wall actually fenced. The root "cool" is contained in "cooling", which is a live owned root: the bamboo family broad legitimately served 3 cooling terms carrying 8 orders in the trailing 28-day window.

> Served terms on the campaign: 61 before the wall, 38 after = **23 terms fenced** against 2 intended
> Clicks: 92 before, 41 after = **51 clicks removed**
> Contribution carried by the 3 owned cooling terms = 8 orders / 28 days x $28.00 = **$8.00 per day**
> The wall deployed Tuesday and was caught at the following Monday pass, 6 days live: 6 x $8.00 = **$48.00 of contribution lost**

Why T+1 verification did not catch it. The #38 T+1 diff proves the staged value landed, and it did land, exactly as staged. Deployment accuracy and wall correctness are different questions, and this is the case that shows the seam: #38 owns proof of landing and this SOP owns whether the thing that landed was right. The detection came from the P11 reconciliation read on served-term collapse, which is failure mode F5.

Response. P10 case 3, erroneous negation with evidence. The served-term collapse and the contribution arithmetic above are the evidence. Removal is human-mandatory and does not deploy in the session it is proposed, so the negative phrase "cool" comes down in the next change set and negative exact goes on the 2 actual drift terms in the same set. Running the reverse test on the correct root: served terms containing "cooler" total 2, cluster members 2, remainder **0**.

Root cause and prediction. The cluster root was chosen by string prefix rather than by taxonomy, which is P1 step 3 condition (b) skipped and step 4 not run. Prediction logged: served-term count returns to the 58-61 band within 2 windows and the next monthly audit returns zero owned-root fences. The reverse-test remainder now appears in every staged phrase row's reason string whether it is zero or not, so the audit can see the test was run rather than inferring it from the outcome.

WE-4 (goes wrong, handled). The monthly wall audit finds five breaches

Setup. Audit date 2026-08-03. Population: 14 terms graduated in the trailing quarter, sources varying 3 to 4 per term, giving **48 wall points**.

> Checked 48, clean on both sides 45, breaches **3**: breach rate = 3 / 48 = **6.25%**
> Breaches 1 and 2: both terms graduated in the session of 2026-07-13 where auto loose was missed as a source. Side one caught them, the entry being absent from the list
> Breach 3: the list shows the entry and the term still serves. The campaign carries 2 ad groups and the negative sits at keyword level in ad group 1 only. Side one reads clean; **only side two catches it**, which is why P9 step 2 checks both
> Own-ASIN check: 1 own child ASIN in complements, 22 clicks at CPC $0.86 = **$18.92** over 30 days, annualizing to 18.92 / 30 x 365 = **$230.19** spent against sales the catalog already owned
> Sibling check: a declaration changed 2026-07-21 with the wall still on the outgoing non-owner = **13 days** of lag
> Findings: 3 breaches + 1 own-ASIN + 1 sibling lag = **5**. Brand and shared-list version checks clean

Attribution and response. Two root causes, not five findings. Breaches 1 and 2 attribute to one creation session that took the source list from the graduation record without the P3 step 2 verification against the served-term feed. Breach 3 attributes to a wall-level selection that did not check ad group structure, which is P3 step 4, and is why that step exists explicitly. The sibling lag routes to #33 as its F3 with P7 executing the move. All five fixes stage the same day. Artifact: 48 points checked, 5 findings, 5 fixes, 2 named root causes, breach rate 6.25%. Prediction: zero breaches from post-fix sessions at the next audit.

WE-5 (clean). A reactive pass with the converting-term guard held

Setup. Bamboo family broad, Objective Discovery, budget capped by design. Window: 92 clicks at CPC $1.72 = $158.24 spend. SOP-14 P3 hands over 6 Route B terms and one Route C cluster.

> Pre-pass zero-order spend across all served terms = $91.16, so WAS% = 91.16 / 158.24 = **57.61%**, above the 40% Discovery limit and firing S-H2
> Route B as delivered: 6 terms. One of the 6 shows 1 order inside the window, so the converting-term guard at P4 step 3 **holds it**: the negative is not staged, a diagnosis note opens, and the term re-reads next cycle. **5 terms proceed**, carrying 34 clicks: 34 x $1.72 = **$58.48**
> Route C: cluster root "crib", 3 member terms, 9 clicks: 9 x $1.72 = **$15.48**. Reverse test: served terms containing "crib" = 3, members = 3, remainder **0**
> Total fenced this pass = $58.48 + $15.48 = **$73.96** per window, or x 52 = **$3,845.92 per year**
> Post-pass campaign spend = $158.24 - $73.96 = **$84.28**; remaining zero-order spend = $91.16 - $73.96 = **$17.20**
> Post-pass WAS% = 17.20 / 84.28 = **20.41%**, back inside the 40% Discovery limit

Staged output. 6 rows: 5 negative exact plus 1 negative phrase. The single phrase wall does the work of 3 exact walls today and fences every future crib term without another pass. The held term is the pass's most important row: it converted, so under v4.0 Section 3 family F11 the relevance check routes it to the fix queue rather than the negative list, and staging it would have removed converting demand on a click-count rule.

Prediction. Next window WAS% holds at or below 25% with the crib cluster fenced; the held term returns a diagnosis verdict at its review date. Both are scored at the named dates rather than inspected.

7. Failure Modes and Escalation

Twelve failure modes. Each carries the signal that detects it, the response, and the account-history incident that motivated it where the corpus records one. Where no incident exists the column says so rather than inventing one, because a fabricated incident is worse than an absent one: it teaches an operator to believe a pattern that never happened.

Failure	Detection signal	Response	Motivating incident
F1. Creation without pre-load	A new campaign whose T3 list attachment is empty or whose staged pre-load row count is below the P2 prescribed count for its type.	Pre-load same day through P2 in full. The creating session is named. The creating SOP's own F-line notifies so the miss is recorded on both sides of the seam.	The satin product tracking cooling terms it cannot rank and taking clicks on them, recorded at v4.0 Section 3 family F11 as a pre-load failure plus an MKL hygiene failure: the taxonomy knew, the campaigns did not. Computed at WE-1.
F2. Wall-less standup, chain broken	The P3 verification line is absent from #30's chain record, or the term appears in a source's served terms after graduation.	Walls same day. The harvest pass reopens and does not close until the line is written. The session is root-caused, not just the wall.	The GRADUATED-UNSTEERED state, defined at v4.0 Section 3 family F8 as a leak in which the source keeps buying the term at uncontrolled CPC while the exact fights it. Recurring instances route through the P9 audit with session attribution.
F3. Converting term negated without diagnosis	A staged negative whose own evidence line shows in-window orders and carries no diagnosis note.	Hold the negative before deployment; run the read; route to the fix queue if relevance holds. The P4 step 3 guard exists for exactly this row.	None recorded in the corpus as a dated incident. The rule is preventive and derives from v4.0 Section 3 family F11, which states that a relevant non-converter routes to the fix queue and never to the negative list.
F4. Shared-list drift	Two campaigns sharing a discovery role return different entry counts or different version strings at the P6 divergence check.	Re-attach the named list; find the divergence source, which is almost always direct edits bypassing the list; the bypassing session is named.	None dated. The structural basis is Doctrine v1.0 Section 11: two campaigns sharing a discovery role share a list so they cannot buy the same mistake twice. The alarm tolerance is ungoverned, amendment AM-7.
F5. Wall built at the wrong level	Served-term count or adjacent tail clicks collapse on a campaign window over window with no other structural change; or a phrase row staged with a non-zero reverse-test remainder.	P10 case 3, erroneous negation. Replace the over-broad root with negative exact on the actual members; restore serving; the remainder now rides in every phrase row's reason string.	Recorded in the corpus as SOP-12 v2.0 failure mode F9, "Wall built at the wrong level: source-campaign impressions collapse on adjacent tail terms after a wall build", routed to the #29 wall audit for root cause. Computed at WE-3.
F6. List-scope error	The negative list shows the entry and the term still appears in served terms. Detected only by the second side of the P9 check.	Re-place the negative at the level covering every ad group on the campaign. Audit the creating session's other walls on multi-group campaigns in the same pass.	None dated in the corpus. Found by construction at WE-4 breach 3, and it is the reason P3 step 4 and P9 step 2 both exist as explicit steps rather than as assumptions.
F7. Uncontrolled removal	The P9 audit finds a negative absent from a list with no matching entry in the removal log.	Reinstate pending review; escalate on first occurrence. An uncontrolled removal is a silent fence breach and it is the only finding here that escalates without a repeat.	None dated. The escalation clock is ungoverned, amendment AM-8.
F8. Sibling wall not moved at re-arbitration	A declaration date and its wall state disagree. The lag is computed in days at the audit.	P7 executes the move same day. The declaration workflow at #33 adds the wall step to its checklist so the lag cannot recur silently.	Recorded at SOP-33 v1.0 failure mode F3, "Walls lagging declaration", and at SOP-29 v1.0 F5. The Bamboo-root ownership call between the 4pc and the 6pc sat open through the first cycle per the Data Architecture v1.0 Section 11 A1 note. Computed at WE-4.
F9. Over-walling the non-owner	The non-owner's own model-specific variants, meaning the root plus its size, count or model terms, return zero impressions after a sibling wall lands.	Restore the P7 step 2 posture: ownership is of the root, never of the modifier space that names the sibling. The non-owner keeps exact defense at floor on its own variants.	Recorded at SOP-33 v1.0 failure mode F4. This SOP is the execution side: the reverse test at P1 step 4 runs against the non-owner's own exact list, which is what prevents it before deployment.
F10. Stewardship wall without a reversion date	A wall placed under a temporary stewardship grant whose reason string carries no reversion date, or a stewardship wall still standing past the owner's restock plus its recovery.	Revert per the dated record. Silent conversion of stewardship into ownership is not a path to ownership.	Recorded at SOP-33 v1.0 failure mode F5, "Immortal stewardship". The wall is the mechanism by which a stewardship becomes permanent in practice, which is why the date lives in this SOP's reason string.

Failure	Detection signal	Response	Motivating incident
F11. Survivor inherits traffic without the fences	After a duplicate consolidation, the survivor's negative list is missing entries the paused instance held.	Carry the missing entries in the same change set as the pause. A survivor missing the loser's fences inherits the traffic and none of the protection.	None dated as a negation incident. The consolidation mechanics at SOP-12 v2.0 P7 step 3 require carrying every negation list the loser held; this failure mode is the audit check that the requirement was met.
F12. Negation cadence stalls	Negatives added equal zero across two consecutive passes on a campaign with active spend.	P4 executes immediately. The cause is named as one of two things and never left ambiguous: the search-term feed was absent, or the pass was skipped. Those have different fixes.	Doctrine v1.0 Section 3.3 states the condition directly: broad campaigns without an active negation cadence are leaking by design. The leak is silent because no alarm fires on a stalled pass, only on its consequences weeks later.

Escalation. F7 escalates on first occurrence. F1 and F2 escalate when the same creating or graduating session produces a second instance, with the session attributed and the audit history attached. F12 escalates to the PPC Lead on the second consecutive stalled pass, because a stalled fence is a budget leak that reports nothing. F8 escalates jointly with #33, since the declaration and the wall are one event held by two documents. Every escalation carries the wall points involved, the window spend at stake, and the session attribution; an escalation without a dollar figure returns to the sender.

8. Outputs and Artifacts

Gates clear with artifacts, not assertions. Each record below names where it lives, how long it is held, and who consumes it. Formal retention periods are not governed anywhere in the corpus; the periods stated are the operational minima the audit population rule and the V5 scoring cycle already require, and the gap is logged as amendment AM-13.

Named record	Storage location	Retention	Downstream consumer
Pre-load checklist per creation	The T3 CAMPAIGN BOARD row for the new campaign: list attachment plus the Negation State field. The staged rows sit in the #38 creation block.	The campaign's operating life plus the governance period after archive.	P9 audit baseline; the F1 check in SOP-12, SOP-13, SOP-14 and SOP-15; the operator ramp week 1 read.
Wall verification line per graduation	The SOP-30 chain record, mirrored into T6 DECISION LOG as the graduation's negation half.	Trailing four quarters minimum, since the audit population is the trailing quarter and trend needs four.	SOP-30 P3 step 3 chain completion; P9 audit; V5 prediction scoring at the review date.
Reactive change set with evidence lines	The #38 session record, with the deployed rows reflected in T6 and the counts in the pass artifact.	Held with the #38 session record.	#38 P6 T+1 diff; V5 scoring; the WAS% re-read at the following pass.
Versioned shared lists with addition and removal logs	The list object itself, named and versioned; attachment recorded in each T3 row.	Full version history, permanent. A list's history is the only way a divergence source is ever found.	P6 divergence check; P9 shared-list check; every creation attaching the list.
Negative-ASIN register	T3 group C5, with the served-ASIN scan output attached to the pass artifact.	Permanent for own-catalog entries; for the campaign's life for failed-target entries.	#16 target rotation decisions; P9 own-ASIN check; the RE-POINT VARIATION re-check.
Monthly wall audit artifact	The audit report: points checked, breach list with session attribution, fixes staged, root causes, breach rate.	Rolling 12 months minimum so the breach rate reads as a trend rather than a snapshot.	PPC Lead review; SOP-30 F1 root-causing; this SOP's own Section 10 review trigger; the operator ramp week 4 audit.
Removal log	The removal register: entity, wall level, case class, evidence, confirming role, date.	Permanent. This is the record that makes F7 detectable at all.	P9 audit; escalation; #33 re-arbitration evidence trail.
Negation-state reconciliation counts	T2 KEYWORD COMMAND child rows and the T3 group C5 rollup; counts in the weekly pass artifact.	Weekly, with the workbook's cycle history.	P11 unsteered-graduate scan; the S-H4 close; the skill layer's alarm feed.
Suppression record	The pass artifact header, naming every wall class suppressed by a missing or stale input and the manifest row that caused it.	With the pass artifact.	T7 MANIFEST completeness read; G8 evidence; the next pass's re-attempt list.
v4.1 amendment register	Section 10.3 of this document.	Until v4.1 publishes and each item is either folded in or formally declined.	The threshold owner; the CEO decision on v4.0 threshold changes per the Authority Matrix row.

The **weekly gate artifact** for this component is the deployed negation change set with a complete evidence line on every row and its T+1 diff attached. The **monthly gate artifact** is the wall audit report, including in months where it found nothing. The **event gate artifact** for a graduation is the verification line, which is what closes the atomic chain across this document and #30.

9. Skill-Conversion Block

Block	Specification
Input bindings	T2 KEYWORD COMMAND campaign child rows, specifically Negation State on Term (child) and Spend Share of KW % (child). T3 CAMPAIGN BOARD group C5, specifically Negation State (pre/reactive/steering), Duplicate / Self-Overlap, Harvest Yield, Graduations (Qtr) and Auto Group Split (D5), plus C4 WAS % and C3 Truncation Flag (G3). The registry graduated-term list from T2 block 8. The brand-term list. The S-series declaration register. The own-catalog ASIN list from the Economics and Context Pack. The #16 failed-target list. The served-term feed and the served-ASIN feed per campaign with group tokens. The MKL with syntax tags and relevancy scores. T7 MANIFEST status per field group. An absent feed suppresses its dependent wall class per G8; the skill never improvises around a missing feed and never completes a partial evidence line from the feed.
Deterministic logic	The P1 wall-level tree including the reverse test and its remainder computation. The P2 pre-load table keyed by campaign type. The P3 all-sources rule and the ad group level check. The P4 execution standards including the converting-term guard, the truncation check, and the G2 sufficiency check. The P5 group-discipline rule. The P6 version and attachment comparison. The P7 same-day wall move and the non-owner posture test. The P9 two-sided audit checklist and the breach-rate computation. The P11 unsteered-graduate filter, which is exact-live and source-state-none, matching v4.0 Section 8.9 scenario S-H4. Every threshold consumed is v4.0's and is read at execution time by section reference; the skill version-pins v4.0 and never caches a number.
Human checkpoints	**AI-executable with audit:** pre-load set drafting from the four lists, wall drafting from #30's source lists, the P1 tree walk including the reverse test, reactive staging from delivered routes, evidence-line assembly, list version diffs, the full monthly audit with two-sided breach detection and attribution, the reconciliation scan, and all arithmetic in the pass artifact. **Human-mandatory:** every P10 removal without exception; every sibling wall move and every stewardship wall; every F3 hold disposition; every reverse test returning a non-zero remainder, since that is the F5 decision point; and the deployment confirm itself, which is #38's checkpoint. Recovery Mode tightens these per v4.0 rule R3.
Cadence	Event-driven same session for creations. Same day for graduations, declarations, stewardship grants and #16 declarations. Weekly with the Monday pass for reactive routes, ASIN scans, shared-list divergence and reconciliation; deployment Tuesday per #38. Monthly for the wall audit on the #4 calendar slot. A served-term breach scan runs at every pull, because a graduated term reappearing in a source is the cheapest signal in the system to read and the most expensive to miss.
Alarm specification	F1 through F12 each carry a detection signal from Section 7, an evidence-line format of campaign or term, signal value, gate or rule reference, and wall points involved, and a routing tier. Same day to the operator queue: F1, F2, F5, F6. Into the weekly pass worklist: F3, F4, F9, F11, F12. To the PPC Lead with audit history attached: F7 on first occurrence, F8 jointly with #33, F10 with the reversion record. Every alarm carries the window spend at stake; an alarm without a dollar figure does not route.
Skill outputs	Drafted pre-load sets and wall sets queued with their event references; drafted reactive change sets in #38 P1 format with complete reason strings; the monthly audit artifact with breach rate and session attribution; the reconciliation counts; list version diffs; F1-F12 alarms with evidence lines; and the drafted T6 entries with predictions stated as figures with review dates. Everything is drafted for confirmation at the tier Section 11 assigns; nothing deploys from the skill.

10. Version Governance

10.1 Ownership and review triggers

Owner: PPC Lead. This component's review triggers are specific to what it holds and are not a monthly calendar habit:

- **Any v4.0 increment** touching Section 9 rule L5, Section 3 families F8 or F11, Section 8.9 scenarios S-H1 through S-H8, or Section 5 gates G2, G8, G11 or G12. Those are the exact references this SOP executes against, and a change to any of them re-validates Sections 3, 4 and 5 before the next pass runs.
- **Any change to the SOP-30 chain definition.** The atomic chain spans both documents and a change to either half increments both. This SOP does not increment quietly when #30 moves.
- **Any change to the #38 P1 staging format.** Every negative leaves here in that format; a schema change is coordinated at #38 first and lands here before the next session.
- **Any D5 structural change at #15** that adds, removes or redefines an auto targeting group, because the P2 pre-load table is keyed by group and a fifth group would have no prescribed set.
- **Any new brand or entity** entering the catalog, which changes the P8 brand set across both entities.
- **Two consecutive monthly audits with a rising breach rate**, which is evidence that a procedure is wrong rather than that an operator was careless.
- **Publication of v4.1**, at which point every item in Section 10.3 is either folded in and removed from this register or formally declined with a reason.

Registered in the Supersession Map at publication. Coupled to SOP-30: the atomic chain spans both documents. Coupled to SOP-14: the routes this SOP executes against are defined there and SOP-13 and SOP-15 run the same machinery, so a route-definition change increments all four.

10.2 What v2.0 carried forward from v1.0, and what is new

Thresholds carried forward: none, because none existed. SOP-29 v1.0 contains zero decision numerals. It defers every threshold to v4.0 by reference and states so in its own header. There was nothing in that category to carry, and the build brief's expectation that there was is corrected here rather than satisfied by manufacturing one.

Incidents carried forward: none from v1.0, because it records none. Its Section 6 is labelled illustrative and its F1-F5 carry detection signals without dated events. The negation-relevant incidents in this document are drawn from where the corpus actually records them: the satin cooling-terms pre-load failure from v4.0 Section 3 family F11 (F1, WE-1); the GRADUATED-UNSTEERED leak from v4.0 Section 3 family F8 (F2); the wrong-level wall from SOP-12 v2.0 failure mode F9 (F5, WE-3); the lagging and over-broad sibling walls from SOP-33 v1.0 failure modes F3, F4 and F5 (F8, F9, F10); the open Bamboo-root ownership call from the Data Architecture v1.0 Section 11 tab A1 (F8); and the leaking-by-design condition from Doctrine v1.0 Section 3.3 (F12).

Structure carried forward: none. v1.0's 10-element, 4-page structure is superseded per conflict C1. The structural exemplar is SOP-12 v2.0.

Substantively new in v2.0: P1 wall-level selection with the reverse test, which did not exist in any form and is the direct answer to the F5 class; P5 as a standalone ASIN procedure rather than a clause inside the reactive routes; P8 brand walls; P11 reconciliation, which turns the negation-state column from a record into a scan; P12 event and recovery overlays; failure modes F3, F6, F7, F11 and F12; four of the five worked examples; the full 25-verdict mapping; and Elements 11 and 12 in their entirety.

10.3 The v4.1 Amendment Register

Thirteen items. Each is a number or rule this SOP needs, could not trace to v4.0, and therefore did not write. None appears as a threshold anywhere in this document.

ID	The gap	What v4.1 owes, and why it matters here
AM-1	The reactive-negation firing threshold. v4.0 Section 8.9 scenario S-H7 reads "Zero-sale spend above threshold on an irrelevant term" without stating the threshold; Section 9 rule L5 names "spend-without-sales" without a figure.	The clicks-at-zero-orders count and the spend-at-zero-orders figure that fire a Route B negative. SOP-14 P3 Route B and this SOP's P4 both execute "the v4.0 reactive-negation rules" that v4.0 does not numerically state, so the routes are delivered today on a rule nobody can quote.
AM-2	"Price class" is not a v4.0 dimension. SOP-14 P3 Route B and SOP-29 v1.0 P3.1 scope the reactive rule to "the term's price class". v4.0 Section 3 family F1 defines margin-dollar class (THIN under $6, STANDARD $6-15, RICH over $15); family F11 defines WAS% limits with no price dimension.	Either a price-class definition or an explicit redirect to margin-dollar class. The two are not interchangeable: a high-price low-margin variation and a low-price high-margin variation land in opposite classes under the two readings, and the reactive rule fires differently on each.
AM-3	The wall-audit population rule. v1.0 P6.1 sampled "all terms graduated in the trailing quarter, minimum" with no source; v4.0 Section 14 rule V4 does not name this audit or its population.	The population or sampling rule. P9 runs the full trailing-quarter population by default, which is defensible but is a local choice standing in for a governed one.
AM-4	The wall-audit calendar slot. v1.0 set "first Monday of the month". v4.0 Section 14 rule V4 lists the monthly items and this audit is not among them.	The slot, registered against component #4, which owns the operating calendar per Ledger row 4. Every SOP Element 2 trigger is supposed to land on a slot defined there.
AM-5	Negative-phrase cluster bounds. No document states a minimum member count for a Route C cluster, a maximum breadth for a cluster root, or a served-term-collapse tolerance that voids a phrase wall after deployment.	The minimum members and the collapse tolerance. This is the gap that produced WE-3; the P1 reverse test is this SOP's local guard, and a guard is not a threshold.
AM-6	The own-catalog negative-ASIN rule has no rule ID. Doctrine v1.0 Section 11 states it as doctrine; v4.0 carries no L-rule, no S-scenario and no threshold for self-cannibalization walls.	A rule ID and its scope: which auto groups, which formats, whether child variations are in scope, whether the wall is permanent. P2 and P5 currently execute a doctrine statement.
AM-7	Shared-list divergence tolerance. v1.0 failure mode F3 detects "diverged lists" with no figure.	The entry-count or version-lag difference that fires F4 here. Without it the check is a comparison with no trigger.
AM-8	Removal-control escalation clock. v1.0 F4 escalates an uncontrolled removal "immediately" with no reinstatement window or review deadline.	The hours to reinstate and the review window. F7 escalates on first occurrence, which is the right posture; the clock is what makes it enforceable.
AM-9	Wall-breach escalation rate. v1.0 escalated on "repeated F1 from one session pattern" with no repeat count or rate.	The breach rate or per-session repeat count that escalates. P9 computes a breach rate; nothing in the corpus says which number is bad.
AM-10	Steering-wall verification lag. SOP-30 P3.4 makes the chain atomic same day; v4.0 carries no maximum interval from wall deployment to verification line.	The maximum hours from deployment to the verification line. "Same day" is the coupled document's rule, not a governed threshold, and it is the interval during which a GRADUATED-UNSTEERED state is live and unrecorded.
AM-11	Gate numbering divergence. v4.0 Section 5 assigns G9 to price stability, G10 to event exclusion, G12 to MKL hygiene. SOP-12 v2.0 Section 3, SOP-14 Section 3, SOP-30 Section 3 and the Data Architecture use G9 for margin freshness, G10 for objective tagging, G12 for the event freeze.	Either v4.1 restates the three gates to match operating usage or the corpus re-numbers. Material here because pre-load depends on MKL hygiene and the P12 overlay on the event freeze, and they carry the same number under the two schemes. Reported as conflict C9.
AM-12	WAS% boundary sense and sub-bands. v4.0 Section 3 family F11 reads "≤10% everywhere except Discovery ≤40%". The KDR Bands tab reads "Standard <10% · Discovery <40%" and adds product-type sub-bands (10% non-subjective, 10-15% semi, 15-20% subjective) that v4.0 does not carry.	The inclusive-or-exclusive resolution, and whether the sub-bands govern. Bedding and sleep products span the subjective and semi-subjective classes, so the sub-bands would move the firing point on most of this catalog.
AM-13	Record retention. No document in the corpus states a retention period for any PPC operating record.	Retention per record class. Section 8 states operational minima derived from the audit population and the V5 scoring cycle, which is a working answer standing in for a governed one.

11. Operator Ramp

Built from the PPC Authority Matrix v2.0 Section 1, taking the rows that name this component's decisions and turning each into a competence the incoming operator demonstrates. Week 1 is observation and artifact familiarity. Week 2 is supervised execution with the checkpoint holder reviewing every output before it deploys. Week 4 is unsupervised execution with the checkpoint holder auditing after the fact. The matrix's standing rules apply throughout: one D per row per scope; an E without a D executes staged decisions only; the BM owner's R on verdict rows is the daily huddle, meaning review and not veto; and under the vacancy rule an unfilled role's D reverts to the PPC Lead and its E to the PPC Manager.

Authority Matrix row	Who holds what here	Week 1: observe	Week 2: supervised	Week 4: unsupervised, audited
Structure creation (all SOPs)	PPC Lead D or per SOP scope; incoming operator E; PPC Manager E; BM owner I.	Read the last 10 creation sessions' T3 rows. State, for each, which P2 set applied and whether the row count matched the prescribed count.	Draft the pre-load set for every creation in the week. The checkpoint holder compares to the P2 table before the creation block deploys.	Pre-load sets deploy with the creation. The wall audit at month end is the audit, and the operator's sessions are attributable in it.
Weekly verdicts, established US	PPC Lead D; incoming operator E within assigned scope; PPC Manager E; BM owner R at the huddle.	Sit the Monday pass. Trace three Route B terms from the SOP-14 output through to a staged row and read each evidence line aloud against P4 step 1.	Stage the full reactive set. Every row reviewed before Tuesday deployment, with particular attention to the converting-term guard and the truncation check.	Stage and deploy within scope. The T+1 diff and the WAS% re-read at the next pass are the audit.
Sibling declarations, S-series	PPC Lead D; incoming operator E; Launch/CA operator I; BM owner R.	Read the declaration register. For two live declarations, locate the walls in the non-owner's campaigns and confirm the owner's campaigns carry nothing.	Execute one wall placement and one wall move under review, including the non-owner defense-posture check at P7 step 2.	Execute same-day moves on declarations. The lag computation at the next audit is the audit.
D5 conversions and auto operations	PPC Lead R; incoming operator D and E.	Read the four P2 auto pre-load sets and state why own-catalog ASINs appear in complements and substitutes and not in close or loose.	Draft the pre-load sets for one D5 conversion; the checkpoint holder confirms group discipline before deployment.	Own the auto pre-load sets. The own-ASIN check at the monthly audit is the audit.
PAT and conquest opens and exits	PPC Lead R; incoming operator D and E.	Read the #16 failed-target list and trace two entries to their negative ASIN placements.	Place negative ASINs on one declared exit under review.	Place on declaration. The audit checks that no exited target reappears as a substitutes candidate.
Bulk deployments, #38	PPC Lead R; incoming operator E; PPC Manager E.	Read three closed session records and locate the negation rows, their reason strings and their T+1 diff lines.	Build the staged negation set to the #38 P1 format. The upload action itself stays with its #38 holder.	Build and hand off. A held row returned for an incomplete reason string is the audit signal.
Event plans, #32	PPC Lead D with the deals owner; incoming operator E; PPC Manager E; BM owner R.	Read one closed event window's armed record and state which wall classes were frozen and why.	Classify wall classes for one arming window under review.	Arm and hold the freeze. Zero non-plan wall rows dated inside the window is the audit.
Code-red declaration and close	PPC Lead D; incoming operator E; PPC Manager E.	Read the recovery triage record from the last code-red and locate the negation pass step.	Not exercised in week 2 unless a code-red is live; if one is, execution is fully supervised.	Run the negation pass on the recovery cadence with the tightened review tiers of v4.0 rule R3.
v4.0 threshold changes	CEO D; PPC Lead E and R; operator I.	Read Section 10.3 and state, for three items, what this SOP does today in the absence of the number.	Add any new gap found in the week's work to the register, with its source-search evidence.	Maintain the register. Writing a number rather than logging a gap is the failure this row exists to prevent.

11.1 Live-Defense Questions

Asked at the end of week 4. These test the reasoning behind the rules, not recall of values. An operator who can quote every threshold and fail these will place correct-looking walls in wrong places.

1. A term shows 40 clicks and zero orders in the window and the relevance check passes. Why does it not get a negative, what does it get instead, and which v4.0 family states the rule?
2. A graduation record lists two source campaigns. The served-term feed shows the term in four. What do you do, and why is the discrepancy itself a finding rather than a correction?
3. You can fence a 12-term drift family with one negative phrase or twelve negative exacts. Argue both sides, then state which error is more expensive and why the asymmetry decides the default.
4. A campaign's WAS% reads 61% and its truncation flag is set. Explain the precedence order that governs your next move and what you are not allowed to conclude from that 61%.
5. A graduated exact is failing while its source walls stand. An operator proposes removing the walls to let the source recover the traffic. Give the argument against, and name the SOP and procedure that owns it.
6. Your monthly audit returns 48 points checked and zero breaches. Why does the artifact still get written, and what would a reviewer learn from it that a skipped audit would have hidden?
7. A sibling declaration changed nine days ago and the wall has not moved. Name the two failure modes this triggers, in which two documents, and why the same event is held by both.
8. An own-catalog ASIN appears in complements' served rows and the campaign is profitable. Explain why the wall goes in anyway.

9. The event window arms tomorrow and you have eleven Route B terms staged. What deploys, what freezes, and what happens to the frozen queue after the window closes?
10. You need a reactive threshold that v4.0 does not contain. Describe exactly what you do, and what you must not do, including what goes into the staged row in the meantime.

12. Interface Contracts

Sixteen contracts, built from the Cross-Reference Ledger v1.7 row 29 Interfaces cell and cross-checked against the Master PPC Data Architecture v1.0 for artifact names and tab locations. Every row states what happens when the handoff does not arrive, because a handoff with no stated non-arrival consequence is an assumption rather than a contract, and that clause is the one the corpus has always been missing.

ID	Sending owner	Receiving owner	Artifact	Timing	Consequence when it does not arrive
I-1	SOP-30, PPC Lead	This SOP, operator E	Graduation approval row with the full source list attached per SOP-30 P1.1.	Same day as approval.	No steering walls are placed; the term enters GRADUATED-UNSTEERED per v4.0 Section 3 family F8 and the source keeps buying at uncontrolled CPC. P11 catches it next pass as scenario S-H4 and F2 fires with session attribution. The chain is incomplete at #30's F1 too, so both documents record it.
I-2	This SOP, operator E	SOP-30, PPC Lead	The wall verification line, one per source, written into #30's chain record.	Same day as the walls.	#30's harvest pass cannot close: the chain is atomic and incomplete without the line. F2 fires here and the pass reopens rather than being marked complete with a gap.
I-3	SOP-14, operator E	This SOP, operator E	Route B terms and Route C clusters with complete evidence lines, from the weekly three-route triage.	Monday with the cascade; deployment Tuesday.	P4 suppresses for the affected campaigns rather than running on partial data. Two consecutive suppressed passes with spend active is F12, the leak-by-design condition, and the cause is named as feed-absent or pass-skipped, never left ambiguous.
I-4	SOP-13, operator E	This SOP, operator E	Phrase-broad dedup wall requests and phrase-side steering handoffs per Ledger row 13.	Weekly with the pass.	Phrase seed terms keep serving in the family broad, so the account buys the same demand at two control levels. The standing dedup wall reads absent at the next audit and attributes to the phrase side.
I-5	SOP-15, operator D and E	This SOP, operator E	Auto Routes B through D detections with group tokens, per SOP-15 P1.5 and P3.	Weekly with the pass.	Auto-group waste stays unfenced and, post-D5, blended evidence validates nothing. This is the receiving end of SOP-15's own I-10, so a gap appears in both registers rather than in neither.
I-6	This SOP, operator E	SOP-12, SOP-13, SOP-14, SOP-15	The prescribed pre-load set for the campaign type being created, staged as part of the creation block.	Same session as creation.	The campaign deploys unfenced and teaches the algorithm what to buy before the first read. F1 fires here and in the creating SOP; SOP-15 records it as its own F1 and F3.
I-7	SOP-16, operator D and E	This SOP, operator E	The failed-target declaration with its exit verdict reference.	On declaration.	Exited targets stay unwalled in substitutes and are re-proposed as fresh candidates by the group that surfaced them, so conquest spend recurs on targets already judged. This SOP does not infer an exit: without the declaration there is no wall.
I-8	SOP-33, PPC Lead	This SOP, operator E	The declaration record, re-arbitration outcome, or stewardship grant with its reversion date.	Same day as the declaration date.	The family runs undeclared or half-walled; F8 fires here and F3 at #33 on the same event. Under a stewardship with no reversion date, F10 fires and the stewardship silently trends toward ownership.
I-9	This SOP, operator E	SOP-33, PPC Lead	Same-day wall confirmation per placement or move.	Same day.	#33's gate artifact is the declaration record with the wall confirmation attached, so it stays incomplete: the declaration reads as decided while traffic still routes the old way.
I-10	This SOP, operator E	SOP-38, operator E	The staged negation change set in #38 P1 format, one row per change with a complete reason string.	Monday stage, Tuesday deploy.	Nothing lands. A row with an incomplete reason string is held at #38 and returned here rather than completed by the transcriber, so the wall waits a full cycle.
I-11	SOP-38, operator E	This SOP, operator E	The T+1 diff report covering the negation rows.	The morning after any deployment.	The walls are unverified, and a wall not proven to have landed is indistinguishable from one that failed; A2 stays stale so the next cascade reads a world that does not exist. Note the seam: the diff proves landing, not correctness. WE-3 is the case where it passed and the wall was still wrong.
I-12	Data Operations	This SOP, operator E	The search-term feed and the served-ASIN feed with per-campaign splits and D5 group tokens, from the SP and SB search term reports and STIS in the Raw Pull Pack.	Monday 06:00, with the workbook refresh completing by 09:00.	P4 and P5 suppress entirely for the affected campaigns per G8, recorded in the pass artifact header with the manifest row that caused it. Negation is unreadable without served terms, and it degrades to nothing rather than to guesswork.

ID	Sending owner	Receiving owner	Artifact	Timing	Consequence when it does not arrive
I-13	Brand Management	This SOP, operator E	The current MKL with syntax tags and relevancy scores, and the current brand-term list across both entities.	Bi-weekly per the v4.0 Section 5 G12 refresh cadence; stale past 30 days degrades the manifest.	P2 cannot resolve taxonomy classes, so that line of every pre-load set suppresses and the creation deploys part-fenced. This is the mechanism behind the satin cooling-terms incident: the taxonomy knew and the campaigns did not.
I-14	SOP-32, PPC Lead with the deals owner	This SOP, operator E	The event plan entering T2 with its window dates and its named actions.	Same day the window arms.	Reactive walls deploy inside an armed window against event-week evidence that v4.0 Section 5 G10 excludes from trend verdicts, fencing a demand shape that will not exist in two weeks. The fence outlives the window.
I-15	This SOP, operator E	SOP-37, PPC Lead	Notice of any negation class routed through a platform automation rule rather than through a bulk session.	Before the rule arms.	An automation rule places or removes negatives outside change control and the P9 audit later finds walls with no session record. That reads as F7, an uncontrolled removal, escalating on first occurrence against an operator who did nothing wrong.
I-16	This SOP, operator E	T6 Decision Log, PPC Lead	One entry per deployed change with its prediction stated as a metric, direction, magnitude, horizon and review date.	With every deployment.	V1 and V5 cannot score the action, so unscored actions accumulate and the rule-level accuracy read at v4.0 Section 14 rule V5 has nothing to compute here. The audit can still find breaches; nothing can tell whether the walls did what they claimed.

Companions. Master PPC Decision Criteria System v4.0 (all thresholds, gates, scenarios, levers and verdict issuance) · Keyword Decision Record Template v3.0, tab "Bands & Verdicts Reference" (the closed verdict vocabulary quoted in Section 5) · PPC Strategic Doctrine v1.0 Sections 3.1-3.4, 4 and 11 (the why) · PPC Authority Matrix v2.0 Section 1 (Elements 9 and 11) · Master PPC Data Architecture v1.0 Sections 6, 6.1, 7 and 9 (T2 child rows, T3 group C5, T5-T7) · PPC Command Workbook TEMPLATE v1.0 (T3 CAMPAIGN BOARD, T6 DECISION LOG, T7 MANIFEST) · PPC Economics Worked-Example Set v1.0 (the arithmetic base declared in Section 6) · PPC Cross-Reference Ledger v1.7 row 29 (ownership boundary and Element 12 source) · SOP-30 Harvest and Graduation (the coupled half of the atomic chain) · SOP-14 Broad Operations (the routes), SOP-13 Phrase, SOP-15 Auto and D5 Split, SOP-12 Exact (the creation events), SOP-33 Sibling Arbitration (the declarations), SOP-16 PAT and Conquest (the ASIN lists), SOP-38 Bulk-File Deployment Mechanics (the execution surface), SOP-32 Event Window, SOP-37 Automation Rules.