

SOP-30: Harvest and Graduation Operations

Version 2.0 · 2026-08-07 · Owner: PPC Lead · Supersedes SOP-30 v1.0 (2026-07-24) in full · Thresholds, bands, scenarios and verdict issuance live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds · Rebuilt to the SOP Build Kit v1.0 twelve-element standard

1. Verdict and Scope

VERDICT: graduation is the single most valuable recurring event in the account, because it moves converting demand up the control gradient where its economics improve, and this SOP owns that decision. Version 2.0 expands v1.0 from 7 procedures to 12, from a 5-row verdict mapping to the full 25-string closed vocabulary, from 1 worked example to 4, and from 5 failure modes to 12, adds the operator ramp and the interface contracts the v1.0 standard never required, and corrects two v1.0 verdict strings that were paraphrases rather than closed-vocabulary quotations. This SOP governs the weekly harvest pass over candidates routed from every discovery source, evidence qualification against the sufficiency gate, validation and disposition, class assignment and family routing, the atomic same-day execution chain and its verification, stem graduation, the head-term fast path, management and aging of the LEAVE TO HARVEST pool, failed-graduate handling, dual-source and twin-set and sibling-shared graduations, the event and Recovery overlays, and per-source yield reporting.

1.1 The Chain Rule: One Event, Three Parts

Everything in this document serves one rule. A graduation is one event with three parts, decide, stand up, and wall, and it is incomplete until all three land and the third is verified. The Master PPC Decision Criteria System v4.0, Section 11 "Layer 6: Transitions, Stages, Modes (T1-T8) and Recovery Mode (R1-R6)", rule T6, states the unit: graduations batch weekly, and each graduation deploys an exact campaign with an objective tag, a seed bid at Expected CPC, steering negation at source, a registry entry, and a review date, and a graduation missing any element is incomplete. The PPC Strategic Doctrine v1.0, Section 3.4 "The keyword lifecycle across match types", names the price of breaking it: the most expensive form of structure debt is the graduated term never negated at source, which then buys itself twice at two different control levels. v4.0 Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)", Family F8 "Structure & Registry", gives that state its name, GRADUATED-UNSTEERED, and Section 8.9 "Discovery Scenarios (S-H1 to S-H8)", scenario S-H4, orders it closed in the same cycle. This SOP owns the decide half of the chain and the assembly of the source list; SOP-29 owns the wall half.

The three parts are executed by three components, which is why the verification line exists rather than an assumption: this SOP decides, SOP-12 stands the exact up (P1 or P2), and SOP-29 places the steering walls (P2). A pass is not closed by a decision; it is closed by SOP-29 returning a wall verification line for every source campaign attached to every approved term.

1.2 The Four Candidate Sources and What Each Delivers

Nothing enters this pass except through a routed candidate row carrying an evidence line. The pass does not discover; it adjudicates.

Source	What arrives, and who owns the route	What the evidence line must carry	What this SOP does not take from it
SOP-14 P3 Route A	Converters from family broad and, by the same shared machinery, from family phrase. SOP-14 owns the definition of the weekly three-route triage.	Clicks, orders, window, source campaign, CPC, per the route that produced it.	The routing decision itself, and the negation cadence that produces Routes B and C.
SOP-13 P4 stem routes	A validated high-volume stem, with its member long-tail set named.	The stem string, the member set, and the pooled member evidence.	Stem clustering, the dated pilot, and the phrase-side standing coverage under LEAVE TO HARVEST.
SOP-15 P3 Route A	Auto-side converters, one row per term per targeting group, each carrying its group token.	The group token in addition to the standard fields. Post-D5 evidence is group-level, never blended.	The D5 split, per-group verdicts, and the two ASIN routes.
Head-term receipts	ADD EXACT + RESTRUCTURE receipts from any source's head-term check.	The carrying campaign, its share of the term's clicks, and the absence of a live exact.	The head-share computation, which each source runs on its own served terms.

ASIN converters are not candidates and never enter this pass. SOP-15 P3.3 routes converting substitutes ASINs to SOP-16 directly, with evidence, and SOP-16 decides. The Cross-Reference Ledger's component-16 row records this explicitly and corrects an earlier attribution of the route to this component; the correction is honoured here.

1.3 The Three Dispositions and the One Bar

Every candidate row leaves this pass in exactly one of three states, and the bar that separates them belongs to v4.0, not to this document. v4.0 Section 8.9, scenario S-H1, sets it: a discovery search term reaching 3 or more orders graduates to exact with a declared objective and is negated at source so traffic consolidates under control. v4.0 Section 9 "Layer 4: The Lever Library (L1-L10)", rule L6, restates it as the structure lever, and Section 11, rule T1, restates it as the objective transition. Three independent statements of the same bar in the threshold authority is the strongest signal in the corpus that the bar is not negotiable at execution time.

Disposition	Condition	Closed-vocabulary verdict issued
Approve	Orders at or above the S-H1 bar, relevancy and taxonomy clear per gate G12, evidence group-level where the source is an auto.	GRADUATE FROM DISCOVERY. v4.0 Section 13.1 "The Closed Verdict Vocabulary" issues the same decision under the string GRADUATE + STEER; the cross-map is stated in Section 5 and neither citation is dropped.
Hold	Converting below the S-H1 bar, relevant, and still earning inside its source's discovery economics.	LEAVE TO HARVEST
Route out	Evidence contradicts relevance or taxonomy membership, whatever the order count.	The gate G12 route runs first; its terminal states are a logged taxonomy expansion or NEGATE (PRE-LOAD / REACTIVE / STEERING).

1.4 Boundaries With Adjacent Components

Every excluded decision names the component that owns it. The Cross-Reference Ledger is the check.

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, scenario, gate, lever and verdict, including the graduation bar itself. This SOP executes; it never decides a number. A threshold appearing here instead of v4.0 is a defect, and a number v4.0 does not carry is an Appendix A amendment item, never a local value.
#14 Broad operations	Owens the weekly three-route triage machinery and the definition of Route A. This SOP consumes Route A output; it never re-runs the triage and never re-derives an evidence line.
#13 Phrase operations	Owens stem clustering, the four stem routes, the phrase-side graduation handoff, and Surface-2 standing long-tail stem coverage. This SOP receives a routed stem and decides whether it graduates.
#15 Auto operations	Owens the D5 four-group structure, per-group verdicts, and group-evidence token hygiene. This SOP refuses blended auto evidence post-D5 and returns the row rather than validating it.
#12 Exact operations	Owens the exact standup on both classes (P1 single-keyword, P2 consolidated tail), and owns promotion and demotion between the two classes (P5, P6). This SOP assigns the class a graduate enters; it never executes a promotion.
#29 Negation architecture	Owens negation doctrine, the wall-level selection tree, the steering wall standards, and the monthly wall audit. This SOP hands over the approved term with its full source list and consumes the wall verification line; without that line this pass cannot close.
#16 PAT and conquest	Owens target opens, flag tests, rotation and exits. ASIN converters route to #16 through SOP-15 P3.3, not through this SOP.
#33 Sibling arbitration	Owens the family ownership declaration and the composite that produces it. This SOP reads the declaration to assemble the source list and to satisfy gate G11 before a shared term's chain runs.
#38 Bulk mechanics	Owens pull, staging format, upload and the T+1 verification pass. This SOP produces staged rows in #38 format and consumes the verification report as the chain's proof of landing.
#32 Deal and event window	Owens the event plan, the armed-window declaration and reconciliation. This SOP stages normally inside an armed window and deploys only what the plan names.
#31 Budget and envelope	Decides funding, CAP SYNTAX and every envelope-level budget verdict. Per Ledger conflict C3 this component has no procedural document, so this SOP names the instrument and tab (T0 zone H, Account Rollup A2) and never cites a #31 procedure number.
#28 Placement and modifier	Owens ladder step derivation and exit-bid discovery. A graduate's seed bid comes from the verdict and the Expected CPC reference, never from a ladder this SOP computes.

Depth note. This component fills the 10-18 page standard without padding because the chain spans three documents and the failure surface is the seam, not the step. The v1.0 at 3 pages described the pass; it did not make the pass executable by a stranger, and it left the source-list assembly, the dual-source credit rule, the pool aging behaviour, and every non-arrival consequence unwritten. Those four gaps are where the account's graduated-unsteered leaks actually occur.

2. Trigger Conditions

Each trigger names the artifact in which the signal is detected and the maximum allowed lag from signal to procedure start. The calendar anchors (Monday 06:00 feed landing, 09:00 packet ready, Tuesday deploy after the BM huddle, T+1 verification) are the Command Workbook READ ME's "CADENCE AND OWNERSHIP" line, and the cadence classes behind them are v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)", rule V4. v4.0 carries no numeric signal-to-start lag table for any component; where a lag below is expressed in hours rather than as a calendar anchor it is a planning figure, and the missing standard rides to v4.1 as Appendix A item AM-12.

Trigger	Detecting artifact (where the operator sees it)	Maximum lag, signal to procedure start	Procedure
Weekly cadence	The Monday cascade reaches T2 KEYWORD COMMAND after T0 and T1 are read; candidate rows from SOP-14 P3, SOP-13 P4 and SOP-15 P3 are present in the pass worklist with evidence lines.	Same day. The pass starts once the T7 MANIFEST completeness stamp is read and the T6 rows due this cycle are scored.	P1 queue assembly, then P2 to P5, then P12 close.
Head-term receipt	An ADD EXACT + RESTRUCTURE receipt arriving from any source's head-term check, visible as a T2 K9 "Action Source (Finding IDs)" entry naming the carrying campaign.	Same day on arrival, outside the weekly rhythm. A receipt that waits for Monday is failure mode F5.	P7 immediately; the validation queue is skipped by construction.
Stem route	SOP-13 P4 routes a validated stem with its member set into the pass worklist.	The same weekly pass in which it arrives.	P6 stem graduation.
Re-evaluation date	T2 K7 "Review Date" on a LEAVE TO HARVEST row reaches or passes the cycle date; the PAUSED-PENDING and pool filter views surface them.	The pass in which the date lands. An entry passing its date untouched is failure mode F4.	P8 pool re-evaluation, inside the weekly pass.
Wall verification return	SOP-29 returns a verification line per source, or the T+1 report from #38 shows a wall point missing or a graduated term still serving in a source's served terms.	Same day. A chain open past T+1 reopens the pass.	P5 step 6; failure mode F1 on a miss.
Registry conflict	T2 K8 shows "# Exact" at 2 or more for one term with the "Duplicate Flag" set and no declared role in the T3 row.	The same pass, before any further graduation on that term.	P1 step 5 dedupe; the consolidation itself executes at SOP-12 P7.
Failed graduate	A graduated exact's T2 K9 "Last Cycle Prediction Score" returns MISS or WRONG-DIRECTION against the graduation prediction at its review date.	The pass in which the score lands.	P9 diagnosis. The source walls stay regardless of the outcome.
MKL refresh	The bi-weekly MKL refresh changes a tracked term's syntax tag or relevancy score, visible in T2 K1 "Relevancy" and "Syntax Tags".	The next pass. A term whose relevancy fell below bar cannot graduate on evidence gathered before the change.	P3 step 4 re-tests gate G12 on every candidate and every open pool entry.

Trigger	Detecting artifact (where the operator sees it)	Maximum lag, signal to procedure start	Procedure
Event window arming	#32 lands an armed-window declaration in T2 K2 "Event Mode" with its dates.	Before the next staging step inside the window.	P11 event overlay; staging continues, deployment follows the plan.
Recovery Mode entry	T0 zone A carries the CODE-RED verdict, or the product enters Recovery per v4.0 Section 11, rule R1.	The same pass.	P11 Recovery overlay: tightened confirmation tiers per rule R3 apply to every chain.

3. Preconditions and Gates

Each gate below names the v4.0 identifier and function, the T7 MANIFEST field group whose status proves it can be read, and the observable an operator can point at to say the gate passed. Per Ledger conflict C9 the operating corpus uses a different number for three of these functions; the corpus alias is given at first use so no citation is dropped, and the numbering itself rides to v4.1 as AM-11. A gate that cannot be read is not a gate that passed: v4.0 Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass", rule G8, downgrades a verdict whose gate fields are missing to a flag, and the pass states what it could not see.

Gate (v4.0 Section 5)	Manifest field group (T7)	What it does to this pass	Observable that proves it passed
G1 Validation quarantine	A4, A5, B11	An impossible-metric row produces a flag and a data-refresh task, never a candidacy. A candidate whose evidence line contains a quarantined row does not enter the queue.	T2 K5 "Flag: Zero-Quarantine" reads Clean on every row in the evidence line. The Bamboo 4pc file's broken SQP rows, CVR fields reading zero with orders present, are the canonical case this gate was written for.
G2 Sample sufficiency	A4, A1	Sets what may be read from the evidence, not whether the term graduates. The S-H1 bar is an order count and clears at its own count; the CVR and CTR verdicts that would price the seed bid do not clear until their own thresholds do.	T2 K2 "Sample State" and "Clicks 7/30/60d". Below the CVR threshold the ceiling read falls back to the product reference CVR from T0, and the record says so.
G3 Budget truncation	A3	A truncated source campaign's rates describe its worst hours. The order count still counts; the CPC in the evidence line is marked and does not seed a bid on its own.	T3 "Truncation Flag (G3)" and "In-Budget Bucket" on every source campaign in the source list.
G4 Eligibility	A7	An ineligible advertised SKU halts all lever work for the target. No standup deploys against it; the candidate holds with the listing task as its re-entry criterion.	T2 K2 "Eligibility" reads Live and T3 C2 "Eligibility + Reason" is clear on the receiving campaign.
G5 Indexing and listing placement	B1	Bars a rank-objective standup on an unindexed term. A graduate entering on a Conversions or Discovery-graduate objective is unaffected; a graduate entering on Ranking is not, and routes to the indexing queue first.	T2 K2 "Indexed?" and "Listing Placement" read Yes and Title or Bullets for rank-objective graduates.
G6 Inventory adequacy	D2, F3 inputs	RED on the targeted variation blocks the standup and re-points ranking campaigns to the declared backup the same day. A graduate whose D2-preferred variation is RED holds; the walls do not.	T2 K2 "Inventory Band" and "Days to Stockout" with the D2 map naming the ranked backup.
G8 Packet completeness	The T7 rollout	A DEGRADED manifest produces flags only. The pass still runs on the rows it can read and records the suppressed classes; it never approximates a missing source's evidence.	The T7 "PACKET COMPLETENESS SCORE" stamp on the cycle header, with the blocked verdict classes named in-row.
G9 Price stability corpus alias: G9 margin freshness	D1	A price change on the targeted variation inside the trailing 7 days marks the CVR and ACOS reads PROVISIONAL. The order count is unaffected; the ceiling comparison carries the flag through to the record.	T2 K2 and the T0 economics spine's G9 state. The economics chain (AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC in T0 C18:C24) reads fresh or the record states PROVISIONAL.
G10 Event exclusion corpus alias: G12 event freeze	A1, T5	Event weeks and the 14 days following are excluded from trend verdicts. Graduations still stage; standups deploy outside the frozen scope unless the event plan names them.	T2 K2 "Event Mode" and the T5 PRODUCT LOG "Trend-Exclusion Flag" on the affected window.
G11 Sibling conflict	D6	A term where two siblings carry active push objectives is a portfolio decision before it is a graduation. The declared owner determines whose exact receives the graduate and which sibling campaigns join the source list.	T2 K8 "Term Owner (Family)" reads a declared owner, not UNDECLARED, and T2 K2 "Family Conflict" is clear.
G12 MKL hygiene corpus alias: G10 objective tagging	B2 inputs via the Hub	Every candidate must exist in the current MKL with a syntax tag and a relevancy score. An untagged or below-bar term routes to MKL maintenance and taxonomy pre-negation, not to graduation, whatever its order count. The satin product's cooling terms are the motivating case.	T2 K1 "Syntax Tags" and "Relevancy" populated, and the MKL refresh inside its bi-weekly cadence. An MKL past 30 days degrades the manifest and this gate reads accordingly.

Two standing preconditions beyond the gates. First, every candidate row carries an evidence line, and a candidate without one does not enter the pass: the pass adjudicates evidence and cannot manufacture it. Second, post-D5 auto evidence is group-level and carries its group token; a blended auto row validates nothing, which is failure mode F3 and the reason SOP-15 P9 quarantines token-blocked rows at source rather than forwarding them.

4. Step-Level Procedures

Twelve procedures. Time figures are planning estimates for scheduling the Monday block; they are not decision numerals and nothing branches on them. Every navigation instruction names the actual tab, block and column an operator opens.

P1. Candidate queue assembly and dedupe (15-25 min per product; planning estimate)

1. Enter only after T0 PRODUCT COMMAND and T1 SYNTAX BOARD have been read. Product posture and family caps govern everything below them, and a T2 pass that skipped T0 and T1 is invalid by construction. Read T0 zone D, the twelve-cell gates strip, and T0 zone C, the economics block, before anything else: the seed bid every approval carries is priced against T0 C18:C24.
2. Score last cycle first. Open T6 DECISION LOG, filter to rows whose "Review Date" lands this cycle and whose "Level" reads Keyword, and record the V5 score in the "V5 Score" column as HIT, PARTIAL, MISS or WRONG-DIRECTION. Unsourced actions do not accumulate (v4.0 Section 14, rule V1). Graduation rows scoring MISS or WRONG-DIRECTION route to P9 in this same pass.
3. Read the manifest before the queue. Open T7 MANIFEST and record the status of A1, A4, B5, B8, D1, D2 and D6. A MISSING group does not stop the pass; it suppresses the verdict classes named in that row's "Verdict Classes Blocked While Missing" cell, and the suppression is written into the pass artifact.
4. Assemble one row per TERM, not one row per arrival. Collect every routed candidate from SOP-14 P3 Route A, SOP-13 P4, SOP-15 P3 Route A, and any head-term receipt. Where the same search term arrives from two or more sources, it is ONE candidate row with all serving sources attached, not two candidacies. The same term validated in broad and in auto-close is one graduation with walls in both sources; treating it as two is failure mode F2 and produces the duplicate the registry will find next month.
5. Attach the full source list, which is the deliverable SOP-29 cannot build for itself. For each candidate open the T2 keyword parent row and click the outline control to expand its campaign child rows. Record every campaign whose child row shows non-zero clicks on the term in the window, reading the child "Camp Type (child)", "Camp Objective (child)" and "Spend Share of KW %" columns. The source list is every campaign that SERVED the term, not only the campaign that routed it. A source list assembled from the routing arrival alone is the single most common origin of a graduated-unsteered leak.
6. Extend the source list across the family. Read T2 K8 "Term Owner (Family)". Where the term is family-shared and an ownership declaration exists, the non-owner's discovery campaigns join the source list because they will need the sibling wall. Where the cell reads UNDECLARED, gate G11 holds the chain and the term routes to #33 for the declaration before approval; the candidate stays in the queue with the block recorded.
7. Dedupe against live structure. Read T2 K8 "# Exact" and "Duplicate Flag". A term already carrying a live exact is not a graduation candidate: it is either already graduated and unsteered (route to the S-H4 close, and the steering wall is the whole remedy), or it is a duplicate with no declared role (route to SOP-12 P7 for consolidation). Neither case adds a second exact.
8. Rank the queue by the loss the delay costs, not by arrival order: head-term receipts first (they leave the queue immediately under P7), then dual-source terms with the largest combined window spend, then single-source terms. The ordering exists because a pass that runs out of confirmation time should have spent it on the largest uncontrolled spend.

P2. Evidence qualification (10-20 min per product; planning estimate)

1. Read the evidence line as it arrived. The pass does not recompute clicks or orders from the raw pull; it reads what the routing SOP produced, because the route owns the definition. A line missing any of clicks, orders, window, source campaign or CPC returns to its source rather than being completed here.
2. Check the group token on every auto-sourced row. Post-D5, an auto row without a resolvable group token is quarantined at source per SOP-15 P9 and must not appear here. If one does, return it and record the exposure: blended evidence validates nothing, which is failure mode F3.
3. Pool a dual-source term's evidence for the order count. Clicks, orders and spend sum across the attached sources; the CPC is recomputed as pooled spend divided by pooled clicks and is labelled a blended discovery CPC, never carried forward as a bid.
4. Pool twin-set evidence for the sufficiency read. v4.0 Section 5, rule G2, states that twin sets pool evidence across the set for sufficiency; T2 K1 "Twin-Set ID" identifies the set. v4.0 does not state whether the S-H1 order count itself pools across a twin set, so the pass records the twin partner's evidence beside the candidate's and does not add the counts together; the gap rides to v4.1 as AM-15.
5. Apply gate G2 to what the evidence can price, not to whether it graduates. The order bar clears at its own count. The CVR read that prices the ceiling does not clear until its own threshold does; below it, the ceiling is computed from the product reference CVR in T0 and the record carries the fallback and its confidence effect, per v4.0 Section 7 "Layer 2: Diagnosis (D1-D10)", rules D3 and D9.
6. Mark every evidence line whose source campaign carries the G3 truncation flag. The CPC in a truncated line is not a price signal, so the seed reference falls to Expected CPC alone and the record says which input was unusable.

P3. Validation and disposition (20-35 min per product; planning estimate)

1. Run gate G12 first, before the order count is read. Open T2 K1 and confirm the term carries a current syntax tag and a relevancy score at or above bar. A term below bar or off-taxonomy routes to the MKL classification decision whatever its order count: the taxonomy either expands deliberately, logged with a syntax tag and a test plan, or the term negates. This is v4.0 Section 8.9, scenario S-H6, and it is the one place where meeting the graduation bar does not graduate a term. Three incidental sales do not create a market position.
2. Read the pooled order count against the S-H1 bar. At or above, the candidate is an approval subject to steps 3 and 4. Below, it is a hold.
3. Test relevance against the evidence, not against the tag alone. A term whose orders landed at a price point or variation the product does not serve contradicts its own relevancy score; that contradiction routes to step 1's MKL decision and the reason is logged on the row.
4. Test the gate set that blocks execution rather than candidacy: G4 eligibility, G5 indexing where the graduate will carry a Ranking objective, G6 inventory on the D2-preferred variation, and G11 family ownership. A candidate failing one of these is approved on evidence and held on execution, with the blocking gate and its re-entry criterion named. Approval and execution are separate states and the record shows both.
5. Write the disposition into T2 K9 "VERDICT" using the closed-vocabulary string, with the fired scenario identifiers in "Scenario ID(s)". A staged row without a scenario identifier does not deploy.
6. Log the reason on every non-approval. A hold carries its accumulated evidence and a review date; a route-out carries the gate that fired and the owner it went to. Nothing exits this pass silently.

```

Candidate row (evidence line complete? no ..... return to source SOP)
|
+-- G12: in current MKL, tagged, relevancy at/above bar?
|   +-- No ..... S-H6 route: MKL classification decision
|   |   first (G12). Terminal states: logged
|   |   taxonomy expansion, or NEGATE
|   |   (PRE-LOAD / REACTIVE / STEERING).
|   +-- Yes ..... continue
|
+-- Auto-sourced row without a group token? ..... F3: return to SOP-15 P9. No verdict.
|
+-- Pooled orders at/above the S-H1 bar?
|   +-- No, and term is relevant and converting ... LEAVE TO HARVEST; pool entry with
|   |   accumulated evidence + review date (P8)
|   +-- No, and zero-sale past the S-H7 line .... not this pass: the reactive route owns it
|   |   (SOP-14 P3 Route B / SOP-15 P3 Route B)
|   +-- Yes, orders contradict relevance ..... step 3: back to the S-H6 route
|   +-- Yes ..... GRADUATE FROM DISCOVERY, then:
|       +-- G4 / G5 / G6 / G11 all clear? ..... approve and execute: P4, then P5
|       +-- One or more failing ..... approve on evidence, hold on execution;
|           name the gate and its re-entry criterion

```

P4. Class assignment and family routing (10-15 min per product; planning estimate)

1. Read the oversight flag, do not judge it. T2 K1 "Oversight Flag" carries Y or N. The criterion behind it is v4.0's, not this SOP's: Section 12 "Layer 7: Precedence and Conflict Resolution (P1-P10)", rule P5, makes every 500+ SV keyword action human-confirmed, and Section 6 "Layer 1: Objective Assignment (O1-O8)", rule O2, sets the SV floor above 500 for a Ranking objective. The Data Architecture and the KDR Template add a second limb, top spend decile, which v4.0 does not carry; where a candidate qualifies only on that limb the assignment is recorded as resting on the ungoverned limb and rides to v4.1 as AM-3.
2. Assign the class. Oversight flag Y assigns the single-keyword class and the standup runs at SOP-12 P1. Oversight flag N assigns the consolidated tail class and the standup runs at SOP-12 P2. The class is read from the flag, never from the strength of the week; assigning on judgment is failure mode F12.
3. Assign the syntax family from the term's root using the family map in T2 K1 "Syntax Tags". The family determines which consolidated campaign receives a tail graduate and which caps govern it, because the family is the cap-enforcement unit.
4. Read the receiving family's cap state before routing a tail graduate. Open T1 SYNTAX BOARD, S3 group, and read "Cap status" and "Share vs target (pp)". A family already over cap still receives the graduate, because a graduation is a structure change and not a spend increase, and the term was already being bought at a lower control level. The cap consequence is a budget verdict and executes at T1 through the instrument layer, never inside this pass.
5. Carry the objective. A tail graduate enters on the receiving campaign's dominant tail objective; an oversight-class graduate enters on the objective the verdict declares. No graduate enters untagged: an untagged row is UNMANAGED and cannot be evaluated at all.
6. Carry the seed bid on the record. v4.0 Section 11, rule T6, seeds at Expected CPC, read from T2 K6 "Expected CPC \$". Record the ceiling beside it: Max Profitable CPC from T2 K6, which is margin dollars times CVR per v4.0 Section 3, Family F1 "Product Economics", and Section 10 "Layer 5: Sizing Mathematics (M1-M10)", rule M1. Where the seed sits above the ceiling, the record states it and names the objective that permits it; where no objective permits it, the graduate enters at the ceiling and the difference is logged. The bid is never invented at standup: SOP-12 P1 step 7 returns a creation whose verdict carried no bid.
7. Write class, family, objective, seed bid and ceiling onto the graduation record before standup. A standup instruction missing any of the five is incomplete and SOP-12 returns it.

```

Approved term
|
+-- T2 K1 Oversight Flag = Y? ..... Class 1 single-keyword ..... standup at SOP-12 P1
|   (qualifies only on the top-spend-decile limb? record the
|   ungoverned basis on the row; AM-3)
|
+-- Oversight Flag = N
|   +-- Family consolidated campaign exists? ..... extend it ... standup at SOP-12 P2
|   +-- No family campaign exists ..... SOP-12 P2 creates one; never straddle
|   |   two families in one campaign
|   +-- Receiving ad group already at its size cap . SOP-12 P2 opens a second ad group in
|       the SAME campaign, never a second
|       campaign on the family

```

P5. The atomic execution chain (20-30 min per approval batch; planning estimate)

1. Issue the verdict. Write GRADUATE FROM DISCOVERY into T2 K9 with the scenario identifiers S-H1 and, where the head-term path applied, S-R12. Human confirmation is mandatory on every approval and every rejection at this step, per the Authority Matrix v2.0 row "Weekly verdicts (established US)" and v4.0 Section 12, rule P5.
2. Hand the standup to SOP-12 the same day, carrying class, family, objective, seed bid and ceiling. This SOP does not create campaigns, set placement modifiers or write campaign names.
3. Hand the term and its FULL source list to SOP-29 in the same session. The source list is every campaign attached at P1 step 5 and step 6, not only the campaign that validated the term. SOP-29 owns which wall level each source receives; this SOP owns that no source is missing from the list.
4. Confirm the registry moved. T2 K8 "# Exact" increments and a T3 CAMPAIGN BOARD row exists or is extended, with C5 "Graduations (Qtr)" incrementing on the receiving campaign. A registry that did not move means the standup did not land.
5. Write the T6 entry with the graduation prediction before deployment, not after. The prediction names the metric, direction, magnitude and horizon: the exact's CPA against its blended discovery CPA at the review date, with the CPC ceiling as the constraint. A creation without a logged prediction does not deploy at any tier, per v4.0 Section 12, rule P10.
6. Close the chain on the verification line, not on the deployment. SOP-29 returns one line per source confirming the negative is present in that source's list; #38's T+1 report confirms the landed values on the standup. The chain is closed when every source has returned its line AND the term is absent from every source's served terms in the following window. A standup whose walls are not verified is failure mode F1 and the pass reopens.
7. Where the chain cannot close inside the window because a source's wall could not be placed, record the open wall point, its reason and its owner in the pass artifact. An open wall point is a quantified exposure, not a formality: it is the term's blended discovery CPC less the exact's seed bid, multiplied by the clicks the source continues to take.

P6. Stem graduation (15 min per stem; planning estimate)

1. Receive the routed stem from SOP-13 P4 with its member long-tail set named. The stem construct itself belongs to #13; this SOP decides only whether the stem graduates.
2. Read the pooled member evidence against the S-H1 bar. The stem's evidence is the sum across its member terms in the window, because the stem is what the members share and no single member need clear the bar alone.
3. Graduate the stem ONCE, as the stem phrase into exact. Member long-tail terms stay in phrase as standing coverage under LEAVE TO HARVEST, which is #13's Surface-2 responsibility, not a disposition this pass re-opens each week.
4. Specify the wall pattern on the record: the stem as negative phrase in the phrase source, fencing every member, plus negative exact in broad. Members are not walled individually. The pattern is SOP-29's to execute; naming it on the record is this SOP's, because the source list is what SOP-29 consumes.
5. Log the member count fenced beside the wall count, so the pass artifact shows how many future terms two wall points closed.
6. Record the standing gap: v4.0 contains no stem construct and therefore no rule making the graduation once-only. The once-only rule is carried from #13 and #29's structure conventions and rides to v4.1 as AM-14.

P7. Head-term fast path (20-30 min per receipt, same day; planning estimate)

1. Trigger: an ADD EXACT + RESTRUCTURE receipt from any source's head-term check. The condition is v4.0's: Section 4 "Layer X: The Interaction Layer. How Metrics Combine Into the Story of a Keyword", rule X10, reads auto or broad carrying more than 40% of a priority term's clicks as a structure failure rather than a performance signal, and Section 8.4 "Ranking / Re-Ranking Keyword Scenarios (S-R1 to S-R16)", scenario S-R12, orders ADD EXACT + RESTRUCTURE before any performance verdict on a priority term with zero exact campaigns.
2. Run the chain immediately, on arrival, outside the weekly rhythm. The validation queue is skipped by construction: a head term is not a discovery candidate proving itself, it is a coverage gap that has already proven itself and is being bought at the wrong control level every day it waits. A receipt sitting in the weekly queue is failure mode F5 and escalates on first occurrence.
3. Assemble the source list exactly as P1 step 5, then run P4 class assignment. A priority term at 500+ SV carries the oversight flag and enters the single-keyword class.
4. Execute P5 in full: startup at SOP-12, walls at SOP-29, registry, prediction, verification line. The steering negation closes the source's leak the same day the exact goes live.
5. Record the exposure the receipt closed: the carrying campaign's share of the term's clicks, its CPC, and the exact's seed bid. The account's named failure class is a head term ranked by campaigns that choose their own traffic; the six-piece satin case, all head-term traffic through auto with zero exacts, is the standing example, and the 6pc auto-dominance finding surfaces as a one-cell read in Account Rollup A4.

P8. LEAVE TO HARVEST pool management (10-20 min per product; planning estimate)

1. The pool is a list that is reviewed, not a set of terms rediscovered each week. Open T2 with the PAUSED-PENDING filter view and the K7 "Review Date" column; every LEAVE entry carries its accumulated evidence, its source list as known at entry, and a review date.
2. At re-evaluation, re-read gate G12 first. An entry whose relevancy fell below bar at the last MKL refresh leaves the pool through the S-H6 route regardless of accumulated orders. Evidence gathered under an old taxonomy does not survive the taxonomy changing.
3. Test the accumulated window against the S-H1 bar. Clearing it graduates the entry through P4 and P5 in this pass. Not clearing it holds the entry with its evidence carried forward and a new review date.
4. Record the aging gap honestly on every hold. v4.0 carries no aging rule and no re-evaluation interval for this pool; the interval used is an operating value that lives in the entry's T6 row, not a threshold in this document, and the missing standard rides to v4.1 as AM-4. The v1.0 of this SOP cited "the v4.0 aging rule"; no such rule exists in the attached file, and this version does not repeat the citation.
5. Report the pool as a distribution, not a count: entries, age bands, and the oldest entry's age, written into the pass artifact. A pool reported only as a number hides the rot, and a rotting pool is failure mode F4. The alarm bar for that distribution is also ungoverned and rides as AM-5.
6. Re-attach the source list at re-evaluation rather than trusting the entry's stored one. Sources change between windows: a campaign paused since entry is no longer a wall point, and a campaign that began serving the term since entry is a new one.

```
Pool entry reaching its review date
|
+-- G12 re-read: still tagged, relevancy at/above bar?
|   +-- No ..... S-H6 route: MKL decision; entry leaves the pool
|   +-- Yes ..... continue
|
+-- Accumulated orders at/above the S-H1 bar?
|   +-- Yes ..... GRADUATE FROM DISCOVERY: P4 class, P5 chain,
|   |               source list re-attached at today's state
|   +-- No ..... continue
|
+-- Still converting at all in the trailing window?
|   +-- Yes ..... LEAVE TO HARVEST holds; evidence carried forward;
|   |               new review date; disposition logged (AM-4 interval)
|   +-- No, and spend continues .... not this pass: the reactive negation route owns a
|   |               zero-sale term (SOP-14 P3 Route B)
|   +-- No, and spend has stopped .. entry closes with its reason; re-entry requires a
|   |               new-window evidence line (standard ungoverned, AM-10)
```

P9. Failed-graduate handling (15 min per case; planning estimate)

1. Trigger: a graduation prediction scoring MISS or WRONG-DIRECTION at its review date in T6, meaning the source CVR did not transfer to the exact.
2. The verdict is diagnosis, not un-graduation. The term's economics are the term's economics, and re-opening the source would buy the same clicks at less control for more money. This is the whole reason the rule exists.
3. The source walls STAY regardless of the exact's outcome. A failed graduate is not a removal case; SOP-29 P7 removal control admits four named case classes and this is not one of them. Un-walling a failed graduate re-creates the double-buy the graduation was built to end.
4. Route the exact to its own verdict path. A graduated exact that fails follows SOP-12's execution of whatever v4.0 issues on it: a bid step toward the ceiling, a taper, or a pause. That verdict is not this SOP's to issue and the exact does not return to the harvest queue.
5. Run the attribution check before the diagnosis is written. Open T5 PRODUCT LOG and confirm whether a listing change, price move, OOS window, deal window or competitor entry sits inside the graduation window; v4.0 Section 7, rule D5, consults the product log before PPC is credited or blamed for any swing.
6. Record the failure against the source, not only the term: a source whose graduates repeatedly fail is producing evidence that does not transfer, and that is a yield-quality finding for the discovery SOP that owns it. The horizon over which a graduate is called failed is ungoverned and rides to v4.1 as AM-9.

P10. Dual-source, twin-set and sibling-shared graduations (10 min inside the pass; planning estimate)

1. Dual-source: one row, all sources attached, one graduation, walls in every source. The evidence pools for the order count (P2 step 3) and every attached source appears in the list handed to SOP-29.
2. Twin sets: read T2 K1 "Twin-Set ID". A twin set is one managed term with one bid posture and shared negatives, per v4.0 Section 3, Family F8, and Section 8.11 "Sibling & Portfolio Scenarios (S-B1 to S-B6)", scenario S-B4. A graduating twin carries its partner's source campaigns into the wall list, because a wall on one spelling that leaves the other serving fences nothing.
3. Sibling-shared terms: gate G11 holds the chain until T2 K8 "Term Owner (Family)" names a declared owner. The owner's exact receives the graduate; the non-owner's discovery campaigns join the source list and receive the sibling wall.
4. Where the declaration changes after a graduation has landed, the walls move with it. The declaration workflow is #33's and the wall move is SOP-29 P5's; this SOP's part is that the graduation record names the declaration and its date, so a later re-arbitration can find every wall it needs to move.
5. Credit the yield to each source that served the term. This convention has no v4.0 basis and rides to v4.1 as AM-7; until it lands, it is applied consistently and stated on every yield read so the reads remain comparable across products.

P11. Event-window and Recovery-Mode overlays (5-10 min per window; planning estimate)

1. Inside an armed event window, graduations stage normally. The evidence is still evidence and the queue does not pause; what changes is deployment, which follows the #32 event plan and lands outside the frozen scope unless the plan names the standup.
2. Event-window evidence is read for the order count only. Event weeks and the 14 days following are excluded from trend verdicts per gate G10, so an event-window CVR or CPC does not price a seed bid; the seed falls to the pre-window Expected CPC reference and the record carries the exclusion flag.
3. Rank gains inside an event are provisional per v4.0 Section 8.12 "Event-Mode Scenarios (S-E1 to S-E6)", scenario S-E5. A graduate entering on a Ranking objective during a window carries the provisional state on its record so the downstream taper clock is not started on an event spike.
4. In Recovery Mode, the tighter confirmation tiers of v4.0 Section 11, rule R3, apply to every chain in the pass: bid moves above 15% and budget moves above \$25 per day are human-confirmed, and step sizes cap at standard even at HIGH confidence. Graduations are structure repairs and sit inside R2's fourth triage band, so they continue rather than freeze.
5. Record the overlay on the pass artifact. A pass run under an overlay and reported as an ordinary pass makes the next cycle's V5 scores unreadable.

P12. Pass close: yield, registry reconciliation and the T6 chain (15-20 min per product; planning estimate)

1. Compute the per-source yield: validated graduations credited per \$100 of the source campaign's window spend. Read spend from T3 C4 "Spend \$ (L30)" scoped to the window and write the read into T3 C5 "Harvest Yield", with "Graduations (Qtr)" incrementing. These reads grade the discovery SOPs, which is why they are computed here and not there.
2. Record that v4.0 sets no floor for the yield read. Section 6, rule O6, names graduation rate as one of the three Discovery objective metrics but carries no band for it, so the read is reported against its own trailing baseline and the missing band rides to v4.1 as AM-8.
3. Reconcile the registry. For every approval, T2 K8 "# Exact" moved by exactly one and no term shows a count of 2 or more without a declared role in its T3 row. A count that moved by two is failure mode F2 and the dedupe at P1 step 4 failed.
4. Reconcile the walls. Every approved term has one verification line per attached source. Count wall points checked against wall points confirmed and write both numbers; an equal count is the passing state and a difference is an open exposure, not a rounding note.
5. Write the pass artifact: the graduation log, one row per candidate, carrying term, sources, evidence, disposition, class, family, seed bid, ceiling, wall points and their verification state, prediction, and review date. This log is the pass; a pass without it did not happen.
6. Close the T6 chain. Every deployed change carries its prediction and review date, and the next cycle's P1 step 2 scores it. Where a suppression or an overlay applied, the artifact names it and the manifest state that caused it.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds and scenario logic live in v4.0. This table maps the complete 25-string closed vocabulary, quoted verbatim from the Keyword Decision Record Template v3.0 "Bands & Verdicts Reference" sheet and matching exemplar SOP-12 v2.0 Section 5, to what this SOP does when each verdict lands on a harvest-relevant row. Where execution belongs to an adjacent component, the boundary is stated so no verdict falls into a seam. Per Ledger conflict C10 the vocabulary is stated at three different counts across the corpus; the resolution recorded there is followed, and the two rows where v4.0 Section 13.1 "The Closed Verdict Vocabulary" carries a different string are cross-mapped in place so neither citation is dropped.

Verdict (KDR Template v3.0 closed vocabulary)	Harvest and graduation execution
GRADUATE FROM DISCOVERY	The decision this SOP owns. P3 validation, P4 class and family, P5 the atomic chain. Issued only after gate G12 clears and the pooled order count reaches the S-H1 bar. Cross-map: v4.0 Section 13.1 issues this decision under the string GRADUATE + STEER, mapped there to S-H1 and T6. Both name one decision; the KDR string is the one written into T2 K9 because the workbook's data validation list is built from it.
LEAVE TO HARVEST	P8 pool entry with accumulated evidence, source list and a review date. No exact is stood up; the term keeps converting in its discovery source under that source's economics. Deliberate, tracked and aged, never a silent non-decision.
ADD EXACT + RESTRUCTURE	P7 head-term fast path, run on arrival. The validation queue is skipped; the chain is identical. The restructure note names which source campaign was carrying the term.
NEGATE (PRE-LOAD / REACTIVE / STEERING)	The steering half of every graduation, owned by #29. This SOP delivers the term and its full source list and consumes the verification line; it never sets a wall level. The pre-load and reactive states never originate here.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	The candidate is not yet a candidate row: it stays in its source at the capped test bid until the quota is reached. This pass takes no action and does not open a pool entry. v1.0 of this SOP carried this string truncated; the full string is restored here.
FUND TEST (UNFUNDED PRIMARY)	A relevant primary term that was underfunded rather than unproven. The capped, dated test is created at SOP-12; this SOP records that the term is under test and defers candidacy to the quota, so it is not re-adjudicated weekly.
QUARANTINE: DATA REFRESH	No candidacy and no disposition on the row. The data-refresh task opens with an owner and a date; the candidate re-enters the queue on the next clean pull with its evidence re-read, not re-used.
CONSOLIDATE DUPLICATES	Detected at P1 step 7 when the registry shows a live exact with no declared role. Consolidation executes at SOP-12 P7; this pass adds no second exact and records the duplicate as the reason the candidacy closed.
RANK PUSH: FUND SIZED PUSH	No harvest execution. A graduate may carry a Ranking objective at standup, but push funding is the M5 queue's and the push sizing is #23's; this SOP hands over the class, family, objective and seed bid only.

Verdict (KDR Template v3.0 closed vocabulary)	Harvest and graduation execution
RECOVERY PUSH	No harvest execution. A recovery-class term is not a discovery candidate: it has a proven ceiling and its own funding path.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	No harvest execution. The verdict lands on a live exact and executes at SOP-12; a graduate reaching it has already left this pass.
UNLOCK TRAFFIC	No harvest execution. Applies to a live campaign below the click-sufficiency read; a discovery term at that volume is not a candidate.
FIX BUDGET (TRUNCATED)	No harvest execution. The truncation flag on a source campaign marks its CPC unusable as a seed reference (P2 step 6); the budget lever itself is the instrument layer's.
FIX CONVERSION FIRST; DEFER PUSH	No harvest execution. A conversion deficit on a graduated exact routes by funnel drop point outside this SOP. It never re-opens the graduation.
PLACEMENT FIX FIRST	No harvest execution. Placement analysis and modifier sizing are #28's; a graduate's modifiers start unearned at standup per SOP-12 P1 step 8.
RE-POINT VARIATION	No harvest execution, one precondition: a graduate's advertised SKU is the D2-preferred variation. Where G6 reads RED before standup the graduate holds; where it flips after standup the re-point executes at SOP-12 P10.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	No harvest execution. The taper is the exact's lifecycle, executed at SOP-12 P4.
HOLD; DEFEND LIGHTLY AT RANK	No harvest execution. A held-at-rank term is post-graduation by definition.
DEFEND AT FLOOR	No harvest execution. Brand terms live in defense and are steering-negated out of generic discovery, which means they are never harvest candidates in the first place.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	No harvest execution. Where a cap breach lands on the family receiving a tail graduate, the graduation still routes (P4 step 4) and the budget half executes at T1 through the instrument layer.
INCREMENTALITY LADDER	No harvest execution. Measurement design belongs to #35 and the ladder to #28.
ENTER SBV ARBITRAGE	No harvest execution. A format substitution acts on a live term's surface, not on a discovery candidate.
OPEN / ROTATE / EXIT CONQUEST	No harvest execution. ASIN converters are not graduations and reach #16 through SOP-15 P3.3, never through this pass.
STOP-LOSS: CONCEDE OR DEFER	No harvest execution on the term. Where a stop-loss closes a dated test in a discovery source, this pass records that the candidacy path closed with it, so the term is not re-queued from stale evidence.
PAUSED-PENDING (REASON + RE-ENTRY)	Used on an approval held by a blocking gate: the approval stands on evidence, the execution holds, and the record carries the blocking gate as the reason and its clearance as the re-entry criterion. Nothing exits the system silently.

Zero local thresholds. Every numeral in this SOP that functions as a decision threshold is quoted from v4.0 with its section heading, or appears in the Appendix A amendment register. The pass rejects a candidate on a gate, an order count or a relevancy tag, and on nothing else. Where the corpus carries a number this SOP would otherwise need, and v4.0 does not, the number is not written: it is named as a gap.

6. Worked Examples

Four cases: one clean pass, one that breaks and is handled, one that refuses evidence and refuses a bar-clearing term, and one that graduates something no single member could. The anonymized live product economics spine was not attached to this session, so all four run on the standard training base of the PPC Economics Worked-Example Set v1.0: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%, reference CVR 5.2%. Values are illustrative and internally consistent; live values come from the T0 economics spine at C18:C24. Every figure below was recomputed programmatically before publication.

The chain used throughout: Max Profitable CPC = contribution x CVR = 28.00 x 0.052 = \$1.456, carried at \$1.46. Break-even ACOS = 28.00 / 80.00 = 35.0%. Target = 50% of break-even = 17.5%. Max Acceptable = 75% of break-even = 26.25%, carried at 26.3%.

WE-1 · The clean weekly pass: one dual-source graduation, one tail graduation, one hold

Setup. Monday cascade, T0 and T1 read, T7 completeness stamped, T6 rows due this cycle scored. Three candidate rows arrive: two from the family broad campaign's SOP-14 P3 Route A, one of which also arrives from Auto-Close's SOP-15 P3 Route A carrying its group token.

P1 queue assembly. Term A "bamboo sheets split king" arrives twice. It is ONE candidate row with two sources attached, not two candidacies. Expanding A's T2 campaign child rows shows a third serving campaign, the sibling 6-piece broad, at non-zero clicks; T2 K8 "Term Owner (Family)" reads a declared owner, so gate G11 clears and the sibling campaign joins the source list as a wall point. A's source list closes at 3. Term B "bamboo sheet set deep pocket king" is single-source, list of 1. Term C "split king bamboo sheet set" is single-source, list of 1. Registry read: all three at # Exact = 0, so none is a duplicate or an unsteered graduate.

P2 evidence qualification. A pools: broad 11 clicks at CPC \$1.72 = \$18.92, Auto-Close 6 clicks at CPC \$1.12 = \$6.72. Pooled 17 clicks, 3 orders, \$25.64. Blended discovery CPC = $25.64 / 17 = \$1.51$. Discovery CPA = $25.64 / 3 = \$8.55$ against \$28.00 contribution. Sales $3 \times \$80.00 = \240.00 , discovery ACOS = $25.64 / 240.00 = 10.7\%$. Gate G2: at 17 clicks the CVR read does not clear its threshold, so the ceiling is computed from the T0 reference CVR and the record carries the fallback and its confidence effect. B: 9 clicks at \$1.68 = \$15.12, 3 orders, CPA \$5.04, ACOS $15.12 / 240.00 = 6.3\%$. C: 14 clicks at \$1.50 = \$21.00, 2 orders, CPA \$10.50, ACOS $21.00 / 160.00 = 13.1\%$.

P3 disposition. Gate G12 first on all three: current MKL, syntax tag Bamboo, relevancy at or above bar, clean. A at 3 pooled orders and B at 3 orders both reach the S-H1 bar: GRADUATE FROM DISCOVERY. C at 2 orders does not: LEAVE TO HARVEST, pool entry opened with accumulated evidence, its source list, and a review date. Gates G4, G5, G6 and G11 clear on A and B, so both are approved AND executable in the same pass.

P4 class and seed. A carries the oversight flag at SV 720: single-keyword class, standup at SOP-12 P1. B does not, at SV 210: consolidated tail class, standup at SOP-12 P2, family Bamboo, extending the existing family campaign from 14 terms to 15. Seed bids from T2 K6 Expected CPC: A at \$1.40, B at \$1.30. Ceiling recorded beside each at \$1.46. A's seed sits \$0.06 under the ceiling while its blended discovery CPC sat \$0.05 over it, which is the graduation argument in one line: the same demand moves from \$1.51 per click to \$1.40 per click, \$0.11 per click, \$1.87 on this window's 17 clicks, and the control to hold it there. At the seed and the pooled CVR the projected CPA is $1.40 / 0.1765 = \$7.93$ against the \$8.55 discovery baseline.

P5 chain, same day. A: standup at SOP-12 P1, walls handed to SOP-29 for all 3 sources, T2 K8 "# Exact" 0 to 1, T3 row created, T6 prediction written: exact CPA at or below \$8.55 within 30 days at CPC at or below the \$1.46 ceiling. B: standup at SOP-12 P2, wall handed for 1 source, T3 "Graduations (Qtr)" increments on the consolidated campaign, T6 prediction written. Four wall points staged in total. Tuesday deploy per #38; T+1 verification returns clean on both standups; SOP-29 returns 4 verification lines; the following window's served-terms scan shows both terms absent from all 4 sources. Chains closed.

P12 close. Yield: family broad window spend \$158.00 with 2 credited = 1.27 validated graduations per \$100; Auto-Close window spend \$92.00 with 1 credited (A, dual-source credit) = 1.09 per \$100. Both reads land in T3 C5 "Harvest Yield". Registry reconciled: two counts moved by exactly one each, no term at 2 or more without a declared role. Walls reconciled: 4 points checked, 4 confirmed. Pool: 1 entry, age 0. Pass artifact is the graduation log with two verified chains and one logged hold.

WE-2 · The edge case: a source that was never on the list, found at T+1

Setup. Same pass as WE-1, one fact changed. Term A was also served by Auto-Loose, whose search-term row carried no resolvable targeting token and was therefore quarantined at SOP-15 P9 rather than forwarded. The quarantine worked exactly as designed: it stopped a blended row from becoming evidence. What it did not do is tell this pass that a fourth campaign was serving the term, because a quarantined row produces no candidate arrival.

What the pass saw. A's source list closed at 3 of the 4 campaigns actually serving the term, a completeness of 75%. Nothing in the pass was wrong: the registry read clean, the three walls were correctly specified, SOP-29 returned 3 verification lines, and #38's T+1 report on the standup was clean. The chain reported closed on complete evidence that was itself incomplete.

Detection. The served-terms scan in the following window shows "bamboo sheets split king" serving in Auto-Loose. This is the F1 detection signal in its second form: not an absent verification line, but a graduated term appearing in a source after graduation. The exact and Auto-Loose are now both buying the term, which is the GRADUATED-UNSTEERED state named in v4.0 Section 3, Family F8.

Sizing the exposure, because it is a number and not a formality. Auto-Loose CPC on the term \$1.94 against the exact's seed of \$1.40: a double-buy premium of \$0.54 per click. At the 4 clicks Auto-Loose took in the exposure day, \$7.76 of spend carried \$2.16 of premium. Left uncaught at that rate the leak runs 28 clicks and \$15.12 per week, \$786.24 across a year, on one term. That is the arithmetic behind the doctrine's claim that the most expensive form of structure debt is the graduated term never negated at source.

Response. Wall staged the same day and deployed; the pass reopens and does not close until the fourth verification line returns. Root cause is named precisely, because naming it wrongly is how it recurs: the failure is not a skipped wall step in this pass, it is a source list assembled from routing arrivals rather than from the campaign child rows. P1 step 5 exists for this reason. Two findings route out: to SOP-15 P9 that a token gap silently removed a campaign from a downstream source list, and to SOP-29's wall audit to add Auto-Loose as an attached source for the term.

The incident this reproduces. SOP-29's monthly wall audit found 2 breaches across 41 wall points on 12 graduated terms, a breach rate of 4.9%, and both breaches traced to one creation session where Auto-Loose was missed as a source. The same source, the same mechanism, twice. That is why P1 step 5 reads campaign child rows rather than trusting the arrival, and why P12 step 4 counts wall points checked against wall points confirmed rather than reporting a pass as clean.

WE-3 · Two refusals: blended evidence, and a term that clears the bar and still does not graduate

First refusal: blended auto evidence. A candidate arrives for “cooling bamboo sheets” from a product whose D5 conversion has not completed. The evidence line reads 41 clicks, 3 orders, \$63.55 spend, blended CPC \$1.55, blended CPA \$21.18. Against \$28.00 contribution the CPA reads comfortable and the order count reaches the bar. The pass returns the row anyway, unvalidated, because post-D5 the evidence contract is group-level.

What the blend was hiding. The four targeting groups split as follows.

Group	Clicks	Orders	Spend	CPC	CPA	vs \$1.46 ceiling
Close match	12	2	\$13.44	\$1.12	\$6.72	\$0.34 under
Loose match	21	1	\$40.11	\$1.91	\$40.11	\$0.45 over
Substitutes	5	0	\$8.05	\$1.61	none	\$0.15 over
Complements	3	0	\$1.95	\$0.65	none	\$0.81 under
Blended	41	3	\$63.55	\$1.55	\$21.18	\$0.09 over

Close match produced 2 orders in 12 clicks at \$6.72 per order. Loose match produced 1 order in 21 clicks at \$40.11 per order, which is \$12.11 above contribution. Seeding a graduate off the blend at \$1.55 would set the seed \$0.43 above the close-match evidence that actually earned the graduation, and would price the term against traffic that converted once in twenty-one clicks. The row returns to SOP-15 P3 for group-level re-issue; the pass records the return, not a verdict, and the exposure window is logged. This is failure mode F3, and it is the reason the group token is a precondition rather than a preference.

Second refusal: the bar is met and the term still does not graduate. “Cooling satin pillowcase” arrives from the same product's loose-match auto with 19 clicks, 3 orders, CPC \$1.36, spend \$25.84, CPA \$8.61, ACOS $25.84 / 240.00 = 10.8\%$. On performance alone this is the strongest candidate in the queue: the order count reaches the S-H1 bar, the CPA sits \$19.39 under contribution, and the ACOS sits well inside every band.

P3 step 1 refuses it before the order count is read. T2 K1 shows the term tagged to the Cooling syntax on a Bamboo product with a relevancy score below bar. Gate G12 fires and v4.0 Section 8.9, scenario S-H6, routes the term to the MKL classification decision first: either the taxonomy expands deliberately, logged with a syntax tag and a test plan, or the term negates. No exact is stood up on either path this pass. The disposition, its gate, and its owner go on the row; the decision itself belongs to Brand Management, not to this pass.

Why the refusal is correct and not conservative. Three incidental sales do not create a market position. The motivating incident is in the corpus: the satin product tracking cooling terms it cannot rank and taking clicks on them, which the threshold authority classes as a pre-load failure plus an MKL hygiene failure, because the taxonomy knew and the campaigns did not. Graduating this term would have built a permanent, budgeted, walled structure around a category the product does not compete in, and every future pass would have read the resulting exact as evidence that the category was chosen.

WE-4 · The stem, the pool, and a cap that does not stop a graduation

Setup. SOP-13 P4 routes the stem “split king” with a member set of 6 long-tail terms, 4 of which served in the window.

Member term	Clicks	Orders	CPC	Spend
split king sheets bamboo	9	1	\$1.55	\$13.95
split king bamboo sheets soft	8	2	\$1.48	\$11.84
split king cooling bamboo set	7	1	\$1.61	\$11.27
split king sheet set deep	7	0	\$1.48	\$10.36
Stem pooled	31	4	\$1.53	\$47.42

P6 stem graduation. The highest single member reaches 2 orders and no member clears the S-H1 bar alone. The stem pools to 4 orders across 31 clicks and \$47.42, CPC \$1.53, CPA \$11.86, sales 4 x \$80.00 = \$320.00, ACOS 47.42 / 320.00 = 14.8%. The stem is what the members share, so the stem is what graduates, once, as the stem phrase into exact at a seed of \$1.35 from T2 K6 against the \$1.46 ceiling. Member long-tail terms stay in phrase as standing coverage under LEAVE TO HARVEST, which is SOP-13's Surface-2 responsibility and is not re-adjudicated here each week. Wall pattern named on the record for SOP-29: the stem as negative phrase in the phrase source, fencing all 6 members, plus negative exact in broad. Two wall points fence six terms and every future term the stem will produce, which is why members are not walled individually. The record also carries the standing gap: v4.0 contains no stem construct, so the once-only rule is carried from SOP-13 and SOP-29 structure convention and rides to v4.1 as AM-14.

P8 pool re-evaluation in the same pass. The pool opens at 7 entries, 3 of them due.

Entry	Accumulated evidence at re-evaluation	Disposition
deep pocket king sheets	Window 1: 11 clicks, 2 orders, CPC \$1.54, \$16.94. Window 2: 6 clicks, 1 order, CPC \$1.53, \$9.18. Accumulated 17 clicks, 3 orders, \$26.12, CPC \$1.54, CPA \$8.71.	Clears the S-H1 bar on accumulated evidence. GRADUATE FROM DISCOVERY. Source list re-attached at today's state, not the state stored at entry.
split king sheets bamboo cooling	Still 2 orders after the second window.	LEAVE TO HARVEST holds, evidence carried forward, new review date. The interval used is an operating value in the entry's T6 row, not a threshold here; v4.0 carries no aging rule and the gap rides as AM-4.
cooling king sheet set	Relevancy fell below bar at the bi-weekly MKL refresh.	Leaves the pool through the S-H6 route regardless of accumulated orders. Evidence gathered under an old taxonomy does not survive the taxonomy changing.

Pool closes at 5 entries, oldest 4 windows. The distribution is written into the artifact, not just the count, because a pool reported as a number hides the rot that failure mode F4 detects.

The cap that does not stop the graduation. “Deep pocket king sheets” carries SV 340, no oversight flag, so it routes to the tail class in the Deep Pocket family. T1 S3 reads that family at 27.4% of tracked spend against its 25.0% secondary cap, a breach of 2.4 percentage points. The graduate routes anyway. A graduation is a structure change, not a spend increase: the term was already being bought, at a lower control level, at a higher CPC, out of the same envelope. Refusing the graduation would leave the spend exactly where it is and remove the only instrument that can price it. The cap consequence is a budget verdict and executes at T1 through the instrument layer, and per Ledger conflict C3 this SOP names the instrument and tab rather than a #31 procedure number.

P12 close. Family phrase window spend \$71.00 with 1 credited = 1.41 validated graduations per \$100, into T3 C5. Two graduations, two chains, three wall points staged, pool distribution reported, four T6 predictions written with review dates.

7. Failure Modes and Escalation

Twelve modes. Each carries the signal an operator can see in a named artifact, the response, and the account-history incident that motivated it where one exists. F1 to F5 are carried forward from v1.0 because the incidents behind them are institutional knowledge held nowhere else in the pack; F6 to F12 are new at v2.0 and close the seams the v1.0 depth standard left unwritten.

Failure	Detection signal	Response	Motivating incident
F1. Wall-less standup (chain broken)	A verification line absent at P5 step 6, or a graduated term appearing in a source's served terms in the following window.	Wall staged and deployed same day; the pass reopens and does not close until every line returns; the exposure is priced in dollars per click and logged, not noted.	The SOP-29 monthly audit's 2 breaches across 41 wall points, both traced to one creation session where Auto-Loose was missed as a source. Recurring F1 routes through that audit with session attribution.
F2. Double graduation across sources	T2 K8 “# Exact” reads 2 for one term with the Duplicate Flag set and no declared role in the T3 row.	The dedupe at P1 step 4 failed: one term, one row, all sources attached. Consolidation executes at SOP-12 P7; this pass fixes the queue assembly that produced the second candidacy.	The same term validated in broad and in auto arriving as two routed candidates in one pass, which is the normal arrival pattern and therefore the standing risk.

Failure	Detection signal	Response	Motivating incident
F3. Validation on blended auto evidence	An auto-sourced evidence line without a resolvable group token, post-D5.	Return to source for group-level re-issue; no verdict is written on the row. Blended evidence validates nothing.	The four-group blend hiding a close-match CPA of \$6.72 behind a loose-match CPA of \$40.11, reproduced in WE-3: the blend reads acceptable and the seed derived from it is \$0.43 wrong.
F4. Rotting LEAVE pool	The pool age distribution shows entries past their review date untouched; the PAUSED-PENDING filter view surfaces them.	P8 executes on every overdue entry in this pass; the cadence gap that let them age is root-caused rather than absorbed.	The pool is the one part of this SOP with no forcing function: nothing breaks when it is skipped, which is precisely why it gets skipped. The alarm bar for the distribution is ungoverned and rides as AM-5.
F5. Head-term receipt queued	An ADD EXACT + RESTRUCTURE receipt sitting in the weekly candidate queue rather than having run on arrival.	P7 immediately. Escalates on first occurrence, not on recurrence: the queue is for discovery candidates proving themselves, never for coverage gaps that have already proven themselves.	The account's named failure class: a head term carried entirely by auto with zero exacts, ranked by campaigns that choose their own traffic. The six-piece satin case surfaces as a one-cell auto-dominance finding in Account Rollup A4.
F6. Graduation on a term failing gate G12	A graduation record whose term carries no current syntax tag, a relevancy below bar, or an MKL entry older than the refresh cadence allows.	The standup reverses and the walls stay pending the MKL decision; the term routes through S-H6. A structure built on an ungoverned taxonomy is harder to remove than the spend it was built to control.	The satin product tracking cooling terms it cannot rank and taking clicks on them: the taxonomy knew and the campaigns did not. WE-3 reproduces the refusal that prevents it.
F7. Seed bid invented at standup	SOP-12 returns a creation whose driving verdict carried no bid, or a T3 row whose starting bid matches no value on the graduation record.	The record is completed at P4 step 6 with Expected CPC and the ceiling, and the standup re-runs. A bid invented at creation time has no prediction attached to it and cannot be scored.	SOP-12 P1 step 7 exists as a return condition, which means the upstream omission has happened. This SOP is the upstream.
F8. Source list truncated to the validating source	A wall count lower than the campaign child-row count showing non-zero clicks on the term in the window.	Re-assemble per P1 step 5 from child rows, re-issue the list to SOP-29, and reopen the chain. Distinct from F1: the wall was not missed, the source was never on the list.	WE-2 in full. The pass reported clean on evidence that was itself incomplete, and the leak ran at \$0.54 per click until the served-terms scan caught it.
F9. Standup deployed inside an armed event window without the plan naming it	A T6 structure entry dated inside a G10 window with no event-plan reference in its row.	Revert unless the event plan names it; the deviation routes to #32 and the trend-exclusion window extends from the revert date. Staging inside the window was always legal; deploying outside the plan was not.	No prior incident recorded in the pack for this component. The mode is written from the seam rather than from history, and the first occurrence should be logged as the incident.
F10. Yield credited wrongly	Per-source graduation counts in T3 C5 that do not reconcile to the pass artifact's chain count, or a dual-source term credited once.	Recompute from the graduation log; state the crediting convention on the read. Yield grades the discovery SOPs, so a wrong credit misgrades another component's work.	The crediting convention has no v4.0 basis (AM-7), so it can drift between passes without anything detecting it except this reconciliation.
F11. Graduation deployed without a prediction	A T6 row for a graduation with an empty "Expected Outcome Window" or "Review Date".	Block the staged set at QA rather than after deployment. An unlogged change does not deploy at any tier, and an unscored action does not accumulate.	The verification loop is the only mechanism that tells this SOP whether its graduations were right. A graduation without a prediction removes the term permanently from that loop.

Failure	Detection signal	Response	Motivating incident
F12. Class assigned on judgment	A single-keyword standup on a term whose T2 K1 oversight flag reads N, or a tail routing on a term whose flag reads Y.	Re-assign per P4 step 2 and move the term through SOP-12 P5 or P6 rather than re-creating it. The boundary between classes is the flag, never preference.	A strong week is the standing temptation, and the syntax caps that a strong week would breach are portfolio law. Precedence is explicit: caps outrank single-keyword strength.

Escalation. F5 escalates on first occurrence, because an uncovered head term is the account's named failure class compounding daily. F1 and F8 escalate on recurrence within the same source or the same session pattern, with the wall audit's session attribution attached. F6 escalates on first occurrence to Brand Management and the PPC Lead jointly, because it is a taxonomy decision that PPC executed without authority. Everything touching the human-mandatory tier, which is every approval, every rejection, every class assignment on an oversight-class term, and every deployment confirm, is confirmed per the Authority Matrix before deployment, without exception.

8. Outputs and Artifacts

Retention is stated below as the governing instrument's, because component #39 owns feed and manifest definitions and its retention policy lines are a pending fold-in. No retention period is invented here; the gap rides to v4.1 as AM-13.

Named record	Storage location	Retention	Downstream consumer and what they do with it
The graduation log (the pass artifact)	The KDR Cycle File generated Monday from T2 blocks K1 to K9, one frozen row per candidate.	The KDR cycle file's, as the audit trail. Period ungoverned (AM-13).	V5 scoring next cycle reads it as what was true and what was decided on cycle date; the SOP-29 wall audit reads it as the population of wall points to sample; the operator ramp reads it as the week-1 artifact.
The standup instruction	The staged #38 change set, plus the T2 K9 "Campaign Moves (prior -> new)" cell.	Held with the #38 staged set.	SOP-12 P1 or P2 executes it the same day. Carries class, family, objective, seed bid and ceiling; incomplete on any of the five and SOP-12 returns it.
The source list	Issued to SOP-29 with the approval; recorded on the T3 row of each source campaign as its wall baseline.	Held with the T3 row for the campaign's life.	SOP-29 P2 places one wall per source and returns one verification line per source. It is the only artifact in the chain that SOP-29 cannot build for itself.
Wall verification lines	Returned into the chain record inside the graduation log.	With the graduation log.	This SOP's P5 step 6 closes the chain on them. The monthly wall audit reconciles points checked against points confirmed.
Per-source yield reads	T3 CAMPAIGN BOARD, C5 "Harvest Yield" and "Graduations (Qtr)".	Rolling, with the T3 row.	Grades the discovery SOPs (#13, #14, #15) against the Discovery objective-metric contract; rolls into Account Rollup A4 structure targets; no v4.0 floor exists for the read (AM-8).
The LEAVE pool report	T2 rows under the PAUSED-PENDING filter view, with K7 "Review Date" and the accumulated evidence.	Until the entry graduates, negates or closes.	Next cycle's P8 works the due entries; the age distribution is the F4 detection surface.
T6 decision-log entries	T6 DECISION LOG, one row per deployed change with prediction, magnitude, horizon and review date.	Append-only; the log is never edited, only scored.	V5 scores each at its review date; MISS and WRONG-DIRECTION scores route to P9 in the pass they land.
Rejection and route-out records	The graduation log row, carrying the gate that fired and the owner it routed to.	With the graduation log.	Brand Management works the MKL decisions; the re-entry evidence standard for a rejected candidate is ungoverned (AM-10).
Suppression statement	The pass artifact header, naming the manifest groups that were STALE or MISSING and the classes they blocked.	With the graduation log.	The analysis states what it could not see before it recommends anything. A confident verdict on a silently incomplete packet is the most expensive failure this system can produce.

The pass gate artifact is the graduation log with every approved chain showing its verification line. A pass reporting approvals without verification lines has produced decisions, not graduations.

9. Skill-Conversion Block

Block	Specification
Input bindings	Candidate rows with evidence lines from SOP-14 P3 Route A, SOP-13 P4 and SOP-15 P3 Route A, the last carrying group tokens. T2 K1 identity (Twin-Set ID, Syntax Tags, Relevancy, SV, Oversight Flag, Coverage State); K2 gates (Validation State, Clicks 7/30/60d, Sample State, Eligibility, Indexed?, Listing Placement, Targeted Variation, Inventory Band, Family Conflict, Event Mode); K5 cross-type totals and the zero-quarantine flag; K6 economics (Max Profitable CPC, Expected CPC, CPA Ceiling); K7 plan (Review Date, Test Status); K8 registry (# Exact, Duplicate Flag, Term Owner (Family)); K9 verdict block. The campaign child rows, specifically "Camp Type", "Camp Objective", "Spend Share of KW %" and "Negation State on Term", which are what the source list is built from. T3 C4 spend for the yield denominator and C5 for the yield reads. T5 product log for the P9 attribution check. T6 for scoring. T7 manifest for suppression. T0 C18:C24 economics spine for every ceiling. The D2 variation map and the #33 ownership declaration. An absent feed suppresses its dependent class per the manifest; the skill never improvises around a missing feed and never completes an evidence line it did not receive.
Deterministic logic	Queue assembly with the one-row-per-term dedupe and all-sources attachment; source-list construction from campaign child rows rather than from routing arrivals; group-token presence test on auto rows; evidence pooling across sources; the G12-before-order-count ordering in P3; the S-H1 bar comparison; the class read from the oversight flag; the seed-and-ceiling pairing; the chain checklist and its verification reconciliation; the pool re-evaluation tree; yield computation and per-source crediting; the P12 registry and wall reconciliations. Every threshold is v4.0's and every wall level is SOP-29's. The skill routes, stages, computes and reconciles; it never sets a threshold and never selects a wall level.

Block	Specification
Human checkpoints	Drawn from the Authority Matrix v2.0 rows that name this component's decisions. AI-executable with audit: queue assembly and dedupe, source-list construction, evidence pooling, gate reads, disposition drafting, class-assignment drafting, chain checklists, verification reconciliation, pool aging reports, yield computation, T6 drafting, and the suppression statement. Human-mandatory: every APPROVE and every route-out confirmation and every class assignment on an oversight-class term (Matrix row "Weekly verdicts (established US)", D held by the PPC Lead); every head-term fast-path standup and every structure creation (Matrix row "Structure creation (all SOPs)"); every deployment confirm (Matrix row "Bulk deployments (#38)"); graduations on launch-stage products (Matrix row "Launch-stage verdicts + batches", D and E held by the launch operator); and any chain on a family-shared term pending a declaration (Matrix row "Sibling declarations (S-series)"). The v4.0 human-review tier applies on top: every 500+ SV keyword action, every gate failure, every structural change, and every LOW-confidence verdict is human-confirmed regardless of who drafted it.
Cadence	Weekly with the Monday pass, after T0 and T1 and after last cycle's due rows are scored. Head-term receipts same day on arrival, outside the weekly rhythm. Chain verification same day as standup, with the served-terms confirmation at the following window. Pool re-evaluation every pass on entries reaching their date. Pool distribution reported every pass. Yield reads every pass. Recovery Mode tightens the confirmation tiers but does not change this cadence.
Alarm specification	F1 to F12 each carry a detection signal from Section 7 with an evidence-line format of row, signal value, and gate or rule reference. Routing tiers: F1, F3, F5, F6 and F8 same day to the operator queue, with F5 and F6 escalating on first occurrence; F2, F7, F9, F11 and F12 into the weekly pass worklist; F4 and F10 into the pass close with the reconciliation that detects them.
Skill outputs	The drafted graduation log queued for confirmation, one row per candidate with disposition, class, family, seed bid, ceiling and source list; drafted standup instructions in #38 staging format; drafted source lists issued to SOP-29; verification reconciliation reports; the pool distribution report; per-source yield reads; drafted T6 entries with predictions and review dates; the suppression statement; and F1 to F12 alarms with evidence lines.

10. Version Governance

Owner: PPC Lead. Review monthly, and immediately on any v4.0 increment, at which point the Section 5 mapping and every S-H1, S-H4, S-H6, S-R12, T6, T1, L6, O6, G2, G11, G12, X10, P5 and M1 reference re-validates before any pass runs on it. Component-specific triggers beyond the monthly slot, none of them generic:

- Any change to the S-H1 bar or its evidence window.** The bar is the single number this entire document turns on, and it is stated three times in v4.0 (S-H1, L6, T1). A change to any one of the three without the other two increments this SOP and opens a conflict report.
- Any change to the atomic chain definition on either side.** This SOP is coupled to SOP-29: the chain spans both documents, and a change to either half increments both. A change to SOP-29's wall-level selection tree changes what the source list must carry.
- Any change to a candidate route's definition.** SOP-14 owns the three-route triage, SOP-13 owns the stem routes, SOP-15 owns the auto routes and the group-token contract. A route-definition change alters what arrives here and increments this SOP alongside the route owner.
- D5 completion across the portfolio.** When every product's autos carry the four-group structure, the F3 return condition becomes a detection sweep rather than a standing precondition, and P2 step 2 changes shape.
- Any change to the oversight-class definition.** P4 reads the flag; a change to what sets the flag changes every class assignment. This includes v4.1 admitting or rejecting the top-spend-decile limb (AM-3).
- The D2 variation map landing at full depth.** Ledger dependency D2 is OPEN and currently caps the depth at which a graduate's advertised variation and its ranked backup can be specified.
- v4.1 resolving any Appendix A item,** which retires that item's workaround in the same increment.

v2.0 change note. Supersedes v1.0 in full. Procedures expanded 7 to 12 (adds P2 evidence qualification, P4 class assignment as its own procedure with the seed-and-ceiling pairing, P10 dual-source and twin-set and sibling handling, P11 event and Recovery overlays, P12 pass close with the two reconciliations). Verdict mapping expanded from 5 rows to the full 25-string closed vocabulary with boundary statements, quoted verbatim from the KDR Template v3.0 and cross-mapped to v4.0 Section 13.1 on the one divergent row this SOP issues. Worked examples 1 to 4, all arithmetic recomputed programmatically. Failure modes 5 to 12 with incident attributions. Elements 11 and 12 added per the Build Kit standard. Two v1.0 verdict strings corrected: CONTINUE TEST TO 50 CLICKS restored to CONTINUE TEST TO 50 CLICKS AT CAPPED BID, and REJECT (logged reason), which appears in no closed vocabulary in the corpus, replaced by the S-H6 route and its two terminal states. The v1.0 citation of a "v4.0 aging rule" for the LEAVE pool is withdrawn: no such rule exists in the attached file, and the gap is logged as AM-4 rather than repeated. Registered in the Supersession Map; the Ledger row for component #30 is returned with this build and carries a seed-row scope correction.

Companions: Master PPC Decision Criteria System v4.0 (all thresholds, scenarios and verdict issuance) · Keyword Decision Record Template v3.0 (the closed vocabulary) · PPC Strategic Doctrine v1.0 Sections 3.1, 3.4, 5.1, 11, 13 (the why) · SOP-29 Negotiation (the coupled half of the chain), SOP-12 Exact (standup and class movement), SOP-13, SOP-14, SOP-15 (the candidate routes), SOP-16 (ASIN converters, via SOP-15), SOP-33 (ownership declarations), SOP-32 (event plans), SOP-38 (staging and T+1) · PPC Authority Matrix v2.0 (Elements 9 and 11) · Master PPC Data Architecture v1.0 and the Command Workbook READ ME (instrument mechanics) · PPC Cross-Reference Ledger (component #30 row).

11. Operator Ramp

Built from the PPC Authority Matrix v2.0 rows that name this component's decisions. Five rows apply. "Weekly verdicts (established US)": D held by the PPC Lead, E by the incoming operator on assigned scope, E by the PPC Manager, R by the BM owner. "Structure creation (all SOPs)": D by the PPC Lead or per SOP scope, E by the incoming operator and the PPC Manager, E by the launch operator on launch and CA scope. "Bulk deployments (#38)": R by the PPC Lead, E by the incoming operator and the PPC Manager. "Launch-stage verdicts + batches": D and E held by the launch operator, R by the PPC Lead and the BM owner, which means graduations on launch-stage products are not this operator's to decide until scope transfers. "Sibling declarations (S-series)": D by the PPC Lead, E by the incoming operator, R by the BM owner, which is why a family-shared graduation waits on a declaration rather than proceeding on judgment. The checkpoint holder for review and audit throughout is the PPC Lead per the D and R cells. The Matrix vacancy rule reverts an unfilled role's D to the PPC Lead and its E to the PPC Manager until the role fills.

Week	Demonstrated competence
Week 1: observation and artifact familiarity	Shadows one full Monday pass end to end without executing. Locates and reads, unaided: one candidate row's complete evidence line and its group token; the campaign child rows under a T2 keyword parent and the "Spend Share of KW %" column that the source list is built from; T2 K8 "# Exact", "Duplicate Flag" and "Term Owner (Family)"; the T7 manifest rows that suppress a verdict class; one closed chain in the graduation log with its verification lines; and one LEAVE pool entry with its accumulated evidence and review date. Reproduces the WE-1 pooling and yield arithmetic cold on fresh illustrative inputs, matching to the cent, as the Economics Worked-Example Set's governance section requires of a new operator before taking a scope. Names, from the artifacts alone and not from memory, which three components execute the three parts of one graduation.
Week 2: supervised execution, checkpoint holder reviewing	Runs the pass on one established US product with the PPC Lead reviewing every disposition before staging. The pass bar is four things, all of them checkable: every candidate deduped to one row per term with all sources attached; every source list assembled from campaign child rows and reconciled against the child-row count; gate G12 read before the order count on every row; and every approval carrying class, family, objective, seed bid and ceiling before it leaves for SOP-12. Drafts, and the PPC Lead confirms, one full atomic chain including the source list issued to SOP-29 and the T6 prediction, then closes it on the returned verification lines. Runs one pool re-evaluation including at least one hold with its aging disclosure. Writes the pass artifact including the suppression statement. Zero approvals with an unreconciled source list across the week is the pass bar; a single one fails the week, because that is the F8 mechanism.
Week 4: unsupervised execution, checkpoint holder auditing	Owens the harvest pass across the assigned established US products: runs the pass, issues dispositions, stages per #38, closes T+1, and reconciles walls and registry at P12. The D cell is now exercised on weekly verdicts within scope, with the PPC Lead auditing artifacts after the fact rather than gating them before. Executes at least one head-term fast path on arrival and one F-line response end to end. The audit samples the source-list reconciliation on three chains, one wall verification set, the yield reads against the graduation log, and one pool distribution. Two clean audit weeks close the ramp. Launch-stage and Canada graduations stay outside this scope until the Matrix rows transfer, and the operator demonstrates that they know which rows those are.

Live-defense questions. These test the reasoning behind the thresholds, not recall of their values. (1) The graduation bar is stated three times in v4.0, at S-H1, L6 and T1. Why would you refuse to graduate a term that reaches it, and name every route by which that can happen. (2) A term arrives from broad and from Auto-Close in the same pass. Explain why that is one graduation and not two, and name exactly what breaks if it is treated as two. (3) Why is the source list this SOP's deliverable rather than SOP-29's, given that SOP-29 places the walls. (4) A graduated exact fails: its CVR did not transfer from the source. Defend leaving the source walls in place, in dollars. (5) A tail graduate routes into a family already 2.4 percentage points over its secondary cap. Defend routing it anyway. (6) Why does the pass read gate G12 before the order count rather than after, when reading it after would be faster on most rows. (7) A candidate's evidence line carries 41 clicks and 3 orders from an auto campaign and no group token. The CPA reads \$21.18 against \$28.00 contribution. Name what you cannot know and what you would seed wrongly if you proceeded. (8) The stem pooled to 4 orders and no member reached 3. Explain why the stem graduates once rather than the two members that converted, and what the wall pattern buys. (9) Which numbers in this SOP would you refuse to write if v4.1 never lands, and what do you do in the meantime.

12. Interface Contracts

Built from the Cross-Reference Ledger's component-30 Interfaces cell and the component-29, -12, -13, -14, -15 and -16 rows that name this component, cross-checked against Master PPC Data Architecture v1.0 for artifact names, tabs and timing. Every row states what happens when the handoff does not arrive, because that is the clause the corpus was missing. Note on I-5 and I-6: the Ledger records component #29 at v2.0 and its interface contracts I-1 and I-2 as the two halves of this chain; the copy of SOP-29 attached to this session is the superseded v1.0, so the seam below is written from the Ledger row rather than from that document's procedure numbers.

#	Interface	Sender to receiver	Artifact and timing	Non-arrival consequence
I-1	Broad and phrase candidate rows	SOP-14 P3 Route A to this SOP P1	One candidate row per term with the evidence line (clicks, orders, window, source campaign, CPC), inside the Monday pass.	The pass runs on the sources that did arrive and records the gap; converting broad terms keep buying themselves at discovery control. Two consecutive passes with spend active and zero Route A arrivals is a stalled negotiation cadence at the source and escalates to the PPC Lead, not a lowered bar here.
I-2	Stem routes	SOP-13 P4 to this SOP P6	The stem string with its member set named and pooled member evidence, inside the Monday pass.	P6 does not run this pass; member terms continue as standing phrase coverage and no harm accrues beyond deferral, since the stem is not being double-bought. The deferral is logged so the stem is not re-discovered as new next pass.
I-3	Auto candidate rows with group tokens	SOP-15 P3 Route A to this SOP P1	One row per term per targeting group, each carrying its group token, inside the Monday pass.	Rows without tokens are refused, not approximated (F3). A token-blocked row travels as a quarantine note, so this pass learns that a campaign served the term even when the evidence cannot be used, which is the gap WE-2 turns into a leak when it is missing.
I-4	Head-term receipts	Any source's head-term check to this SOP P7	An ADD EXACT + RESTRUCTURE receipt naming the carrying campaign and its share of the term's clicks, same day.	The account's named failure class compounds daily: the head term stays ranked by campaigns that choose their own traffic. Escalates on first occurrence at either end, whether the receipt was never sent or was sent and queued.
I-5	Approval plus the full source list	This SOP P5 to #29	The approved term with every attached source campaign, issued in the same session as the standup instruction.	#29 cannot build this list for itself, so the walls it places are the walls it was told about: the term goes GRADUATED-UNSTEERED at every omitted source and the exact fights its own discovery campaigns. This is the highest-consequence non-arrival in the document.
I-6	Wall verification lines	#29 to this SOP P5 step 6	One line per source confirming the negative is present, plus the following window's served-terms confirmation.	The chain cannot close and the pass stays open. The graduation is not reported as complete, the open wall point is priced in dollars per click, and F1 fires.

#	Interface	Sender to receiver	Artifact and timing	Non-arrival consequence
I-7	Standup instruction	This SOP P5 to SOP-12 P1 or P2	Class, family, objective, seed bid and ceiling, same day as the verdict.	SOP-12 returns an incomplete instruction rather than inventing a bid (F7). The graduate does not stand up, the walls are not placed against a live exact, and the term stays at discovery control for another window.
I-8	Registry counts and child rows	T2 K8 and the campaign child rows to this SOP P1	Current per-match-type counts and per-campaign spend share on the term, refreshed with the Monday 09:00 packet.	Dedupe and source-list construction are both unreadable: the pass cannot tell a graduation from an unsteered graduate from a duplicate. Approvals suspend on affected terms and the manifest state is recorded, per the completeness gate.
I-9	Economics chain	#41 refresh to T0 C18:C24 to this SOP P4	AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS and ExpectedCPC with the G9 freshness state, before the Monday 09:00 packet-ready read.	Seed and ceiling carry PROVISIONAL. Approvals continue, because the graduation bar is an order count and not a ceiling read, but the seed is recorded as provisional and re-priced at the next fresh refresh rather than left standing silently.
I-10	Variation map	Economics and Context Pack (D2) to this SOP P4	Keyword-to-variation map with ranked backups, current per its pack cadence.	The standup points at the preferred SKU with the D2-pending flag; a graduate whose variation is unmapped cannot be re-pointed if inventory turns, so a G6 flip on that row escalates same day. Ledger dependency D2 is OPEN and caps the depth of this row.
I-11	MKL with syntax tags and relevancy	Brand Management to this SOP P3 step 1	The bi-weekly MKL refresh, landing before the pass.	Gate G12 cannot be read and no graduation is issued: a structure built on an ungoverned taxonomy costs more to remove than the spend it was meant to control. An MKL past 30 days degrades the manifest and the pass states the degradation rather than proceeding on stale tags.
I-12	Family ownership declaration	#33 to this SOP P1 step 6 and P10	The declared term owner in T2 K8, with its declaration date.	Gate G11 holds the chain on family-shared terms. The candidate stays queued with the block recorded rather than graduating into a sibling conflict; a term graduated without the declaration will need its walls moved at the first re-arbitration and nobody will know which walls those are.
I-13	Staged change set and T+1 return	This SOP to #38, and #38 back to P5 step 6	Staged bulk rows with verdict and scenario identifier in the reason string, Tuesday deploy after the BM huddle, T+1 verification report back.	Without the staged set nothing deploys and the pass holds. Without the T+1 report the chain has no proof of landing and cannot close; a mismatch returns to the pass the same day rather than at next cycle.
I-14	Per-source yield reads	This SOP P12 to T3 C5 to the audit and Account Rollup A4	Validated graduations per \$100 of window spend, plus Graduations (Qtr), every pass.	The discovery SOPs go ungraded for the cycle: a broad campaign at band that graduates nothing reads as healthy, which is the slow leak the doctrine names. The M6 structure check loses an input and the pass artifact is incomplete until written.
I-15	Event plan arming	#32 to this SOP P11	The armed-window declaration with dates and named actions, landing in T2 K2 Event Mode before the ramp date.	Staging continues, deployment does not: a standup inside an unarmed-but-live window is F9. A window observed live with no plan escalates to the deals owner rather than being worked around here.
I-16	Decision log and predictions	This SOP to T6 to V5 scoring	One entry per deployed change with prediction, magnitude, horizon and review date, same day as deploy.	Unscored actions do not accumulate, and an unlogged change does not deploy at any tier. The missing entry blocks the staged set at QA rather than after it; a graduation without a prediction leaves the verification loop permanently (F11).

Appendix A. v4.1 Amendment Items

Fifteen items. No number below appears as an operating value anywhere in this SOP. Each states the gap, where it was found, and the proposal carried to the threshold owner. Four duplicate items already logged by other build sessions and are marked so v4.1 folds each once.

Item	Gap	Provenance and proposal
AM-1	The graduation evidence window	v4.0 Section 8.9, scenario S-H1, sets the bar at 3 or more orders and names no window; Section 9 rule L6 and Section 11 rule T1 restate the bar and also name no window. Gate G2 defines 7d, 30d and 60d buckets for rate reads only. A term reaching 3 orders across 60 days and one reaching 3 across 7 days are different terms and currently graduate identically. Proposal: admit an evidence window to S-H1, or state explicitly that the count is window-agnostic and why. Highest-priority item in this register: it is the single bar the entire document turns on.
AM-2	Price class versus margin-dollar class	SOP-14 P3, SOP-29 and SOP-30 v1.0 all validate "at the term's price class". The phrase appears nowhere in v4.0, which carries margin-dollar class THIN under \$6, STANDARD \$6-15, RICH above \$15 at Section 3 Family F1. Proposal: retire the phrase or define it. <i>Duplicate of the #29 session's AM-2 per Ledger row 29; fold once.</i>
AM-3	Oversight-class second limb	v4.0 carries the 500+ SV limb at Section 12 rule P5 and Section 6 rule O2. Master PPC Data Architecture v1.0 Section 6 and the KDR Template v3.0 Read Me both add "or top spend decile", which v4.0 does not carry. P4 class assignment reads the flag, so an ungoverned limb silently decides which SOP-12 procedure runs. Proposal: admit the second limb with its computation, or remove it from the instrument documents.
AM-4	LEAVE pool aging rule and re-evaluation interval	SOP-30 v1.0 cited "the v4.0 aging rule". No aging rule exists in the attached v4.0. The interval currently lives in each entry's T6 row as an operating value. Proposal: admit a re-evaluation interval and an aging rule stating whether accumulated evidence resets or the candidate closes.
AM-5	Pool distribution alarm bar	Failure mode F4 detects a rotting pool from the age distribution, and no threshold exists for when a distribution is rotting. Proposal: admit a pool-age or pool-size bar to the Discovery objective-metric contract at rule O6.
AM-6	Same-day atomicity clock	The chain's same-day completion requirement is carried by PPC Strategic Doctrine v1.0 Section 3.4. v4.0 Section 11 rule T6 says graduations batch weekly and names no completion clock; Section 8.9 scenario S-H4 says the unsteered leak closes "same cycle". Doctrine is not a threshold authority. Proposal: admit the clock to T6, at same day or same cycle, and reconcile the two phrasings.
AM-7	Dual-source yield credit convention	v4.0 rule O6 grades Discovery on graduation rate and sets no rule for crediting a term validated in two sources. Credit each, credit the validator, or split, are three different gradings of the same work. Proposal: admit the convention so per-source yield reads remain comparable across products.
AM-8	Graduation-rate floor	Rule O6 names graduation rate as a Discovery metric and carries no band for it, unlike the WAS limit beside it. The yield read is therefore reported against its own trailing baseline with no external bar. Proposal: admit a floor, or state that the metric is trend-only by design.
AM-9	Failed-graduate horizon	No v4.0 rule states how long a graduated exact runs before its failure to transfer is called. P9 currently triggers on the graduation prediction's own review date, which each record sets individually. Proposal: admit a horizon so the trigger is uniform across records.
AM-10	Re-entry evidence for a routed-out candidate	SOP-30 v1.0 required "new-window evidence" for a rejected term to re-enter candidacy. That standard has no v4.0 basis and is not written here. Proposal: admit a re-entry evidence standard, or state that a routed-out term re-enters on the ordinary route with its reason visible.
AM-11	Gate numbering	v4.0 Section 5 assigns G9 to price stability, G10 to event exclusion and G12 to MKL hygiene; the operating corpus uses G9 for margin freshness, G10 for objective tagging and G12 for the event freeze. Element 3 gives the corpus alias at first use so no citation is dropped. Proposal: reconcile the numbering in v4.1. <i>Duplicate of Ledger conflict C9 and the #29 session's AM-11; fold once.</i>
AM-12	Element 2 lag table	The Build Kit requires a maximum signal-to-start lag on every trigger. v4.0 Section 14 rule V4 sets cadence classes and no lag figures exist. Element 2 expresses lags as calendar anchors from the Command Workbook READ ME wherever possible. Proposal: admit a lag table. <i>Duplicate of the SOP-17/18/19 session's item A-7; fold once.</i>
AM-13	Record retention	Element 8 requires a retention period per record. Ledger row 39 carries "Pending fold-in: retention policy lines" and no period exists for the graduation log, the pool, or the source lists. Proposal: admit retention lines at component #39. <i>Duplicate of the #29 session's AM-13; fold once.</i>
AM-14	The stem construct	The word stem appears nowhere in v4.0. The graduate-once rule, the pooled-member evidence basis, and the two-point wall pattern are structure conventions carried from SOP-13 and SOP-29. Proposal: admit the stem construct and its graduation rule, or state that stems graduate under S-H1 on pooled evidence with the pooling rule named.
AM-15	Twin-set pooling and the order bar	Gate G2 pools twin-set evidence across the set for sufficiency. v4.0 does not state whether the S-H1 order count also pools across a twin set. P2 step 4 therefore records the partner's evidence beside the candidate's and does not add the counts. Proposal: state whether the order bar pools across a twin set.

Appendix B. Weekly Harvest Pass Operator Card

P1 to P12 compressed; the SOP text governs.

1. T0 then T1 before T2, always; read the T0 gates strip and the economics spine first. 2. Score last cycle's due T6 rows before adjudicating anything new; MISS and WRONG-DIRECTION route to P9 this pass. 3. Read T7; a MISSING group suppresses its named classes and the suppression goes in the artifact. 4. One row per TERM, all sources attached; two arrivals of one term is one graduation. 5. Build the source list from the campaign child rows, not from the routing arrival; reconcile the count. 6. Gate G12 before the order count, on every row, always. 7. Auto row with no group token: return it, do not validate it. 8. At or above the S-H1 bar and clean on G12: GRADUATE FROM DISCOVERY. Below and converting: LEAVE TO HARVEST with a review date. Off-taxonomy: the S-H6 route. 9. Class from the oversight flag, never from the week; family from the root; seed from Expected CPC with the ceiling recorded beside it. 10. Chain the same day: standup to SOP-12, term plus FULL source list to SOP-29, registry moved, prediction written before deploy. 11. Close on the verification lines, not on the deployment; count points checked against points confirmed. 12. Yield per source per \$100, registry reconciled, pool reported as a distribution, artifact written. A pass without the artifact did not happen.