

SOP-32: Deal and Event Window Operations

Version 2.0 | 2026-08-08 | Owner: PPC Lead jointly with the deals calendar owner | Supersedes SOP-32 v1.0 (2026-07-24) in full | Rebuilt to the SOP Build Kit v1.0 twelve-element standard | Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds and logs 12 ungoverned values as Appendix A amendment items | Absorbs the scheduled deal-window playbook deliverable and the scheduled annual-seasonality fold-in

1. Verdict and Scope

VERDICT: an event window is a priced distortion of four numbers at once, so it gets its own plan, its own caps, its own reads, and a dated exclusion period before standing governance touches the data again. Traffic, conversion, CPC and price all move inside the window, and the price move alone marks every in-window CVR and ACOS read PROVISIONAL under v4.0, Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass", gate G9. That single fact is the document: a window generates the most data of any 7 days in the quarter and the least usable data, and the entire job of this SOP is to buy the volume without letting the distortion enter a standing verdict. Version 2.0 expands v1.0 from 7 procedures to 13, from a 5-row verdict mapping written in five locally invented strings to the full 25-string closed vocabulary plus the 2 v4.0-only strings and the 10-string campaign-grain subset, from 1 worked example to 4, and from 7 failure modes to 12, and adds the operator ramp, the 18 interface contracts, and the annual seasonality procedure the v1.0 standard never required.

This SOP governs the PPC consequences of a confirmed deal or event window from calendar intake to scored close: classification, the event plan and its prediction, cap construction, modifier and placement posture, push amplification and the in-window allowance, arming with its freeze list and baseline snapshot, the in-window daily block, cap-hit and stop decisions, close and revert, the settle period and the quarantine of window data, post-event scoring and class-factor calibration, liquidation windows on terminal-tagged units, and the annual seasonality calendar that times pushes into demand rises the deals calendar does not carry. Two rules carry the whole document and both are v4.0 rules, not local ones: every campaign touching the window carries an explicit cap sized by rule M9, and window-distorted data never feeds a standing verdict, per gate G10.

1.1 Why the Window Is a Distortion and Not a Good Week

Four values move together inside a window and each moves a different decision surface. **Price** falls, which cuts contribution per order and therefore lowers the break-even chain the whole product is judged against (v4.0, Section 10 "Layer 5: Sizing Mathematics (M1-M10)", rule M1: break-even ACOS is margin divided by AOV, Target is 50% of break-even, Max Acceptable is 75%). A 15% discount does not cut the ceiling by 15%; it cuts contribution by the full discount and the ceiling with it, which is why the window's Max Profitable CPC is usually *below* the baseline ceiling even though traffic is higher. **Conversion** rises, which raises Max Profitable CPC in the other direction (margin dollars times CVR, v4.0 Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)", Family F1 "Product Economics"), so the two moves partly cancel and the net is an arithmetic question the plan answers before the window opens, never an intuition. **Traffic** rises, which exhausts unchanged budgets by early afternoon and marks every rate on the affected campaigns TRUNCATED under gate G3. **CPC** rises because every competitor is bidding the same days. A verdict written on any one of those four in isolation is wrong in a direction the operator cannot predict, which is the reason gate G10 exists and the reason the settle period is not optional politeness.

1.2 The Five Event Classes This SOP Operates

Five classes carry different plan shapes, and the class is declared at intake because it selects the intake lead time, the cap basis, the modifier posture, and the scoring form. The classes are institutional record carried forward from v1.0 Section 1: **lightning deals** (hours, single variation, no ramp, the cap is the whole control); **7-day deals** (the standard case, full plan, worked at WE-1); **tentpoles** (Prime Day, BFCM: multi-day, multi-product, twice-daily cadence, the only class where a first-time window has no own-history factor, worked at WE-2); **price-discount events** (a standing price move rather than a badge, so the price limb of Section 1.1 applies and the traffic limb largely does not); and **liquidation deals** on terminal-tagged units, where the objective function is component #27's and not the stage's, worked at WE-4. v4.0 carries no closed list of deal classes and the Master PPC Data Architecture v1.0, Section 9 "The Infrastructure Tabs: Logs and the Manifest (T5-T7)" carries "Deal window" as one single value in the T5 Product Log's closed event-type list, so the class is recorded in the plan and not in the instrument. That gap is Appendix A item AM-10.

1.3 Boundaries With Adjacent Components

Every excluded decision names the document that owns it. Boundaries written against components rebuilding concurrently in Wave 5A (#24, #28, #33) cite the Cross-Reference Ledger row rather than a procedure number, because those numbers move at their rebuild and the Ledger v2.5 correction register records roughly 40 owed corrections of exactly this class across six documents.

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, gate, scenario, lever limit, sizing rule and verdict string. This SOP executes; it never decides. A threshold appearing here instead of v4.0 is a defect, and a value v4.0 does not carry goes to Appendix A, never into a procedure.
Pricing and deals calendar (P2.3, P3.5)	Decide whether a deal runs, at what price, and whether its margin is approved. This SOP never re-litigates the deal price; it plans around a confirmed window. Neither process document is held in this corpus, which caps intake depth at what the calendar entry delivers (Section 12, I-1).
#23 Ranking-push operations	Owns push sizing, the M2 worksheet, funding, standup order, stop-loss and taper trigger release. Windows AMPLIFY live pushes and never originate them: a push that opens inside a window is F8 here and violates v4.0 Section 8.12 "Event-Mode Scenarios (S-E1 to S-E6)", scenario S-E1. This SOP supplies the armed-window declaration; #23 decides what its push does with it.
#24 Re-ranking and recovery	Owns loss diagnosis, baseline restoration and recovery sizing. The deal-hangover cause class belongs there and resolves by calendar. This SOP hands over only at settle expiry and only where post-settle rank is still below the armed baseline. Cited per Ledger row 24; procedure numbers move at its Wave 5A rebuild.
#15 Auto operations	Owns event autos end to end. This SOP arms the window they may open inside and closes it; the standup, the expiry naming and the no-silent-conversion rule are SOP-15 P6 and P6.2, verified against SOP-15 v2.0 as attached.
#12 Exact campaign operations	Owns exact structure. Its P9 event overlay freezes structure changes inside an armed window and executes only what the event plan names. This SOP produces the plan that names them.
#27 Terminal and LTSF-tier rules	Owns the liquidation objective switch, the R2 allowable-ad-cost ceiling, R3 cliff pacing and the R6 weekly expiry check. A liquidation window computes under those rules; this SOP contributes only the window plan, its caps and its scoring. The LTSF Master Decision and Operations System is outside this corpus per conflict C6 and this SOP stops at that boundary.
#28 Placement and modifier operations	Owns the ceiling worksheet, ladder rungs, exit-bid discovery and modifier sizing, including the normalization that steps an event premium back out. This SOP decides only whether a premium is tested inside a window and books its removal on the revert checklist. Cited per Ledger row 28; procedure numbers move at its Wave 5A rebuild.

#33 Sibling and family arbitration	Owns the ownership declaration and guard enforcement on contested roots. A window does not adjudicate ownership by accident: an undeclared contested root inside the affected set is noted in the plan and the G11 portfolio call runs there. Cited per Ledger row 33.
#26 Mature and defend, code red	Owns Recovery Mode execution and the code-red trip. Mode precedence is v4.0's, not this SOP's: Recovery outranks Event Mode, so a product in Recovery entering a window runs the window at defense-and-proven-terms posture only (Section 11 rule T8, Section 12 rule P7).
#38 Bulk-file deployment	Owns staging format, upload, the processing report and the T+1 verification diff. Every arming set, cap change, allowance raise and revert set stages in #38 format with the event reference in the reason string, and this SOP consumes the verification diff.
#37 Automation platform rules	Owns the rule register and change control. This SOP never edits a rule; it supplies the armed-window campaign list so #37's exclusion handling keeps platform rules off frozen scope for the duration.
#35 Incrementality testing	Reads the event exclusion when it voids or defers a test window. This SOP publishes the exclusion dates; #35 decides what its test does about them.
#31 Budget and envelope	Envelope law lives in the instrument layer (T0 zone H, Account Rollup A2) per open conflict C3; this SOP never cites a #31 procedure number. Window caps sum inside the product envelope and the envelope decision is not made here.
#34 Canada annex	Owns the Canada delta. Everything here applies in CA unless #34 overrides it; the CA deals calendar arrives on the same interface with the same non-arrival consequence.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. Lags use the cadence terms v4.0 uses (same day, same session, daily) and the Monday cascade defined in Master PPC Data Architecture v1.0, Section 12 "Data Flow: Feed to Columns to Cadence": Raw Pull Pack lands Monday 06:00, workbooks refresh by 09:00, the audit runs Monday, deployment is Tuesday after the BM huddle, and daily checks run on pushes and events per v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)", rule V4.

Trigger	Detecting artifact	Max lag	Procedure
Calendar confirmation	A confirmed entry on the deals calendar (Economics and Context Pack Tab 4 "Deal calendar link (P3.5)") carrying class, dates and deal price	Same day the entry is confirmed, and not later than the intake date the class carries	P1 intake and classification, then P2 plan
Pre-event window opens	T2 K2 "Event Mode" resolves to PRE-EVENT, 14 days before window start (v4.0 Section 3, Family F2 "Product State and Strategy", Event posture)	Same day the state flips	P2 to P5 complete before P6 can run
Plan review due	The drafted plan with its prediction, held for the joint sign-off the Authority Matrix v2.0 "Event plans (#32)" row requires	Not later than 48 hours before arming	P2 step 7 sign-off
Arming mark	T-1 relative to window start; the plan carries the arming date	Same day	P6 arming, freeze list and baseline snapshot
Window live	T2 K2 "Event Mode" reads IN-EVENT; T3 C2 state shows the plan's staged values landed	Daily; twice daily on tentpoles	P7 in-window block
Cap reached on any affected campaign	T3 C3 "In-Budget Bucket" and the pacing read against the event projection; gate G7 event cadence	Same day, inside the block that found it	P8 cap-hit decision
Pacing deviation past the line	Spend against the event projection deviating more than 15% (v4.0 Section 8.12, scenario S-E2)	Same day	P8 correction on the named branch
Competitor deal detected on a conquest target	The competitor deal-state flag in T2 K8 "Competitor Depth (D8)" (manifest group D8; see Section 3)	Same day	P7 step 5 stand-down per S-E3 and S-C5
Window close	The window end timestamp on the calendar entry	Same day	P9 revert checklist
Settle expiry	The exclusion period ending 14 days after the window per gate G10	Inside the next Monday pass	P11 scoring; standing governance resumes
Terminal tag inside a planned window	An LTSF terminal or at-risk tag landing on a variation in the affected set	Same day	P12 liquidation window under #27's rules
Seasonal SV Index turning	T2 K1 "Seasonal SV Index Dir" rising on a syntax, outside any deal window (v4.0 Section 8.12, scenario S-E6)	Inside the next Monday pass	P13 seasonality timing
Verification return	The #38 T+1 verification diff flags a landed-value mismatch on an arming, raise or revert row	Same day	Redeploy the delta per #38; F7 response if it is a revert row

3. Preconditions and Gates

Every procedure below runs only after these read current in the T7 Manifest (Master PPC Data Architecture v1.0, Section 9 "The Infrastructure Tabs: Logs and the Manifest (T5-T7)"). A missing or stale field group suppresses its dependent execution class rather than degrading it silently, per gate G8; the plan records what it could not see and proceeds on what it can.

Precondition (T7 manifest field group)	Requirement and the observable that proves it passed
The confirmed calendar entry	Class, start and end timestamps, deal price and deal fee present. Observable: the Economics and Context Pack Tab 4 deal-calendar row resolves and the T5 Product Log carries a "Deal window" entry with its dates. A window with no calendar row does not arm, because this SOP plans around a confirmed decision and never makes one.
A2 Deployment state	Current bulk state in-row for every campaign in the affected set: state, daily budget, bid, bidding strategy, placement multipliers. Observable: T3 C2 cells populated with the current pull date. Without it the baseline snapshot is a guess and F6 is already true.

A3 Pacing and in-budget	In-budget percentage and bucket per campaign, and the day-of-week spend curve the deviation read compares against. Observable: T3 C3 "In-Budget Bucket" and "Pacing Deviation Flag" resolved. The Command Workbook template records A3 as PARTIAL with the day-of-week curve pending, so where the curve is absent the in-window pacing read runs against the plan's own projection only and the plan records the substitution.
A7 Eligibility	Every advertised SKU in the affected set eligible, with the deal badge live at arming. Observable: T3 C2 "Eligibility + Reason" reads live. G4 halts every lever on a target the day its ad reads ineligible or suppressed and opens the listing fix task the same day; advertising an unbuable offer into event traffic spends into a wall at event CPC.
B8 and D1 Economics chain	T0 zone C named cells live (AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC) with their margin-table version. Observable: T0 zone C shows the version and the G9 state. The deal price is itself a price change, so gate G9 marks every in-window CVR and ACOS read PROVISIONAL for the duration and for the trailing 7 days after the price restores, and no conversion or fix-first verdict issues on those reads.
Inventory cover (G6 band and days to stockout)	Cover clears projected window velocity plus the settle period on every variation in the affected set. Observable: T2 K2 "Inventory Band" GREEN and "Days to Stockout" exceeding window days plus settle days. A deal into a stockout converts paid momentum into a recovery case and is the worst available outcome of any procedure in this document; YELLOW freezes push spend at current level and RED tapers and re-points the same day (gate G6, Section 8.3 "Inventory Scenarios (S-I)", S-I2 and S-I3).
D8 Competitor depth	Competitor price, rating, review count, deal state and stockout flags current for every conquest target inside the affected set. Observable: T2 K8 "Competitor Depth (D8)" populated inside the trailing week. The manifest carries D8 as MISSING in the template, so while it is MISSING the S-E3 stand-down cannot be detected mechanically, the plan states the blind spot, and the finding downgrades to a flag per gate G8 rather than to silence.
D9 Daily rank series	Daily organic rank on every amplified term, so the in-window rank watch and the post-settle read compare against the same series. Observable: T2 K4 "Organic Rank" carries a daily series link. The manifest carries D9 as PARTIAL at head terms only, so amplified terms outside the head set are read weekly and the plan records which terms carry which cadence.
E Composite indices (the Event Lift Index)	The Event Lift Index for this product and this event class, computed from our own history (v4.0 Section 3, Family F16 "Category E: Extended Intelligence Metrics", metric E9). Observable: the class factor resolves with the number of prior windows behind it. Where no prior window exists, rule M9's first-time path applies and the borrowed category default is an Appendix A gap, not a number this SOP writes.
Family conflict state	Every root in the affected set either carries a declared family owner or carries the guard that applies while it is undeclared. Observable: T2 K2 "Family Conflict" resolved and T2 K8 "Term Owner (Family)" populated. Gate G11 makes two siblings pushing one term a portfolio decision, not two product decisions.
MKL hygiene	Every term in the affected set exists in the current MKL with a syntax tag and a relevancy score (gate G12). Observable: the MKL refresh date inside 30 days. An untagged term inside an amplified set buys event traffic the taxonomy never claimed.

Gates in force, by v4.0 name and number, with the corpus alias given at first use. Conflict C9 records that the operating corpus numbers three gates differently from v4.0; this SOP names every gate as v4.0 names it. **G1** quarantines impossible rows to flags with no recommendation. **G2** sample sufficiency governs every rate read: bid reads need at least 15 clicks (7d), CVR verdicts at least 100 clicks, CTR verdicts at least 1,000 impressions. **G3** marks every rate on a campaign TRUNCATED below 70% in-budget and forces the budget lever before any bid conclusion, which is the standing state of a capped campaign in the back half of an event day. **G4** eligibility. **G6** inventory adequacy. **G7** envelope and pacing: campaign budgets reconcile to the product envelope, non-event pacing deviation above 25% against the day-of-week curve alerts the same day, and *during events pacing is checked daily against the event projection*, which is the cadence clause this SOP runs on. **G8** packet completeness. **G9** price stability, which a deal price fires by definition. **G10** event exclusion, the gate that carries the whole settle rule: event weeks and the 14 days following are excluded from trend verdicts, and event-window performance is evaluated only against the event projection and the Event Lift Index, never against baseline thresholds. The operating corpus calls this exclusion "G12", including the Data Architecture's own T5 column label "Trend-exclusion window flag (G12)"; the v4.0 number is G10 and this SOP uses it. **G11** sibling conflict. **G12** MKL hygiene, which in v4.0 carries no event clause at all.

The freeze has no gate. The in-window freeze on structure changes, main-image changes and un-named modifier moves is a standing account rule enforced across the corpus and attributed in five documents to "G12". v4.0's G12 is MKL hygiene and v4.0's G10 excludes window data from trend verdicts; neither carries a freeze clause. The freeze therefore operates in this SOP as a plan state declared at P6 and enforced by the components that hold the levers (#12 P9 on exact structure, #15 P6 on event autos, #37 on platform rules), and its absence from v4.0 is Appendix A item AM-1, the highest-priority item in this register. Nothing about the freeze changes operationally; what changes is that this document stops citing a gate number that does not carry it.

4. Step-Level Procedures

Time estimates are planning figures for scheduling the operator's week. They are not thresholds and nothing in v4.0 governs them.

P1. Event intake and classification (25-40 min per window; planning estimate)

1. **Pull the calendar entry.** Open the Economics and Context Pack, Tab 4 "Deal calendar link (P3.5)", and read the confirmed row: event name, class, start timestamp, end timestamp, deal price, deal fee, affected ASINs and variations. A row missing the end timestamp or the deal price is not confirmed and returns to the deals calendar owner the same day; this SOP does not infer either.
2. **Classify.** Assign one of the five Section 1.2 classes. The class selects the intake lead time, the cap basis, the modifier posture and the scoring form, and it is written into the plan header. A window carrying a terminal-tagged variation in its affected set is a liquidation window regardless of what the calendar calls it, and routes to P12.
3. **Set the intake date.** Standard classes intake at the opening of the PRE-EVENT window, which v4.0 sets at 14 days before start (Section 3, Family F2, Event posture: NONE / PRE-EVENT (14d window) / IN-EVENT / POST-EVENT (14d)). Tentpoles intake earlier at the lead time the A-to-Z System Map v0.9 Section 6 records for this program (event plans at T-14 and T-30); v4.0 carries no tentpole lead time, so the operating value lives in the calendar and the map, and the gap is Appendix A item AM-2.
4. **Read the product posture before the window, not during it.** Open T0 for every affected product and record four states into the plan header: the stage tag and its TACOS band (zone B row 3), whether the product is in Recovery Mode (zone A verdict and the code-red flag on Account Rollup A0), the gate strip G1-G12 (zone D), and the earning-potential class (zone E). A product in Recovery runs the window at defense-and-proven-terms posture only, because Recovery outranks Event Mode under v4.0 Section 11 "Layer 6: Transitions, Stages, Modes (T1-T8) and Recovery Mode (R1-R6)", rule T8, restated as precedence in Section 12 "Layer 7: Precedence and Conflict Resolution (P1-P10)", rule P7. A CONSTRAINED earning-class product carrying above-break-even spend runs profit-only and its velocity requirements route to deals and pricing rather than to PPC subsidy (Section 8.1 "Product-Level Scenarios (S-A)", scenario S-A10).
5. **Recompute the window economics.** From T0 zone C, take AOV, MarginPerOrder and the margin-table version. Compute the window's own chain at the deal price: window AOV equals the deal price; window contribution equals MarginPerOrder minus the discount in dollars; window break-even ACOS equals window contribution divided by window AOV; window Target is 50% of that and window Max Acceptable 75%, per rule M1. Compute window Max Profitable CPC as window contribution times the expected in-window CVR (Family F1). Record all six values in the plan. v4.0 does not state whether window ceilings are computed on deal-price or baseline economics, so this SOP computes both and states both in the plan; the governance gap is Appendix A item AM-11.
6. **Check inventory against the whole exposure.** Days of cover must exceed window days plus the 14-day settle on every variation in the affected set, read from T2 K2 "Days to Stockout". Cover that clears the window but not the settle is a YELLOW-class finding that caps the plan rather than blocking it; cover that does not clear the window returns the entry to the deals calendar owner with the shortfall quantified in units.
7. **Log intake.** One T5 Product Log entry with Event Type "Deal window", the dates, the affected scope and the trend-exclusion flag set. The log entry is what every later attribution check reads before crediting or blaming PPC for a swing (v4.0 Section 7 "Layer 2: Diagnosis (D1-D10)", rule D5).

```
Intake classification (one confirmed calendar entry)
|
+-- Affected set contains an LTSF terminal or at-risk tag? ... yes -> LIQUIDATION -> P12
|                               (objective function is #27's; R2/R3 govern)
+-- Duration under 24 hours, single variation? ..... yes -> LIGHTNING -> P2 short form, cap is the control,
|                               no ramp, no modifier test
+-- Multi-day, multi-product, category-wide traffic event? .... yes -> TENTPOLE -> P2 full form at the tentpole lead
|                               time; twice-daily P7; first-window path at P3.4
+-- Price move without a deal badge or event traffic? ..... yes -> PRICE-DISCOUNT -> P2 full form, price limb only:
|                               recompute ceilings (P1.5), skip the uplift cap basis
+-- Otherwise ..... 7-DAY DEAL -> P2 full form, the standard case
```

P2. The event plan and its prediction (60-90 min for a 7-day deal, 2-3 hours for a tentpole; planning estimate)

1. **Build the affected-campaign set.** Filter T3 Campaign Board to the product's live rows and select every campaign whose advertised SKU or ASIN appears on the calendar entry. Record campaign name, type, objective tag, current daily budget, current bid or ad-group default, and current placement multipliers TOS and PP from T3 C2. This list is the plan's scope and nothing outside it is touched inside the window.
2. **Add the halo scope.** Family and sibling campaigns not carrying the discounted SKU still lift when the family's traffic rises. Add them as a separately labelled halo block, because halo spend rides the same caps but the halo campaigns are not selling the discounted offer and their ceilings are the baseline ones from P1 step 5, not the window ones. v4.0 carries no rule defining halo membership; the plan states its membership test and the gap is Appendix A item AM-8.
3. **Name the live pushes the window will amplify.** From T2 K7, list every term on the affected products carrying a Ranking objective with a funded worksheet. Each gets an in-window allowance line at P5. Pushes are amplified, never originated: scenario S-E1 states that no new experiments enter the window.
4. **Read the contested roots.** For every root in the affected set, check T2 K2 "Family Conflict" and K8 "Term Owner (Family)". An undeclared contested root is noted in the plan with the guard that applies while it is undeclared, and the declaration is #33's to make. A window does not settle ownership by spending through it.
5. **Write the caps (P3), the modifier posture (P4) and the allowances (P5)** into the plan as three tables. The plan is a single document and the single change authority inside the window: an action not named in it does not deploy inside the window.
6. **Write the prediction.** Four values with a named horizon: units sold across the affected set, spend against total caps, the post-settle rank delta on the amplified terms measured against the P6 baseline, and the halo lift in orders. This is the row V5 scores at settle expiry, and a prediction without all four is incomplete (v4.0 Section 14, rule V1; Section 12, rule P10).
7. **Write the stop conditions.** The plan names its own stop conditions in numbers, because inside the window standing bands do not apply and the plan is the only thing that can stop it. v4.0 carries S-E2's 15% pacing-deviation correction trigger but no stop threshold, so the stop condition is a plan value with a named owner and the gap is Appendix A item AM-6.
8. **Sign off jointly.** The Authority Matrix v2.0 Section 1 row "Event plans (#32)" holds D with the PPC Lead together with the deals owner, E with the incoming operator, E with the launch and CA operator in their own scope, E with the PPC Manager, and R with the BM owner at the daily huddle. Standing rule 1 of that matrix requires shared D cells to name their split: on this row the PPC Lead signs the PPC side and the deals owner signs the calendar and price side, and neither signs alone. Unsigned plans do not arm.

P3. Cap construction, the overspend guard (20-30 min per window; planning estimate)

1. **Establish the baseline.** Per campaign in the affected set, take trailing daily spend from T3 C4 over a clean 14-day window that contains no prior event days and no logged price move (T5 confirms). A baseline drawn across a prior window inflates every cap downstream of it.

2. **Size the cap.** Per v4.0 Section 10, rule M9: event budget per priority term equals baseline daily spend times the Event Lift Index times event days. Applied at campaign grain, the daily cap is baseline daily spend times the class Event Lift Index, and the window budget is that figure times event days. The Event Lift Index is this product's own measured multiplier for this event class (Family F16, metric E9), not a category rule of thumb.
3. **Write the caps as ceilings, not targets.** Reaching a cap is a decision point (P8), never a silent raise and never a silent stop. An affected campaign with no cap in the plan is failure mode F1, and inside a tentpole it is the highest-cost single defect this SOP can produce, because event auctions are exactly where a dynamic-bidding campaign runs away.
4. **The first-window path.** Where no prior window of this class exists on this product, rule M9's own clause applies: first-time events borrow the category default and log the variance. v4.0 carries no category-default table, and the program's open-decisions register carries this as item PPC-DR-35 with tentpole factors recorded as having no basis. This SOP therefore does not write a first-window factor. The plan runs at the baseline with its named allowance (P5) as the only expansion authority, states the missing factor on its face, and the gap is Appendix A item AM-4. The exposure this creates is quantified at WE-2 so the cost of the gap is visible rather than argued.
5. **Reconcile to the envelope.** The summed window caps plus unaffected campaign budgets must sit inside the product weekly envelope on T0 zone H. Product envelope outranks keyword ambition (Section 12, rule P2); an over-subscribed window returns to the plan for reduction by the plan's own priority order, not by an across-the-board haircut, and never by an envelope raise made inside this SOP.
6. **Stage the caps** as #38 rows with prior value, new value, and a reason string carrying the event reference and the plan version. Deployment is at arming (P6), not before: caps deployed early spend event budget on non-event days.

P4. Modifier and placement posture inside the window (20-30 min per window; planning estimate)

1. **Default posture is no change.** A placement premium is earned from placement-level conversion evidence, never from the expectation that event traffic converts better. v4.0 Section 9 "Layer 4: The Lever Library (L1-L10)", rule L3 admits a premium only on a TOS-EARNS decomposition and cuts it on TOS-VANITY; rule M4 sizes any multiplier so that effective top-of-search CPC stays at or below the lower of Max Profitable CPC for the objective and Expected CPC plus tolerance.
2. **Where a premium is tested,** the ceiling worksheet is #28's and the plan records three things: the effective top-of-search CPC the multiplier produces, the window Max Profitable CPC from P1 step 5 that it is checked against, and the dated removal that goes onto the revert checklist at P9. A premium with no booked removal is F7 waiting to happen.
3. **Read the window ceiling, not the baseline ceiling.** Because the deal price cuts contribution, the window Max Profitable CPC is frequently below the baseline value, so a multiplier that was legal last week can be illegal inside the window at the same bid. This is the single most common arithmetic error in event planning and it is computed at P1 step 5 precisely so P4 cannot repeat it.
4. **Bids hold.** The plan changes budgets and, where evidenced, multipliers. Bid moves inside a window are the P8 exception path only, and never on a rate read: bid changes on window-distorted reads are F2.
5. **Hourly boost removal.** Where any hourly boost or daypart adjustment is live from a prior window, it is removed at arming. Component #45 is a dated deferral and its register line assigns event-window boost removal to this component until the Marketing Stream hourly feed lands, so the removal is this SOP's item and appears on the P6 freeze list.

P5. Push amplification and the in-window allowance (15-25 min per live push; planning estimate)

1. **Confirm the push is already live and funded.** Read T2 K7 for the term: Objective Ranking, Required Budget \$/wk and Loss Ceiling \$/wk populated, ETA Checkpoint dated. A term without a funded worksheet is not a push and does not receive an allowance; funding a push inside a window is F8.
2. **Write the allowance.** Each amplified push gets a named allowance expressed as a dollar or percentage expansion of its window cap, available on a deliberate call at P8 and on no other basis. v4.0 carries no event-allowance rule: rule L4's single-step budget expansion of 20-30% is a performance-and-utilization rule for standing operation, not an event authority. The allowance is therefore a plan value and the gap is Appendix A item AM-5.
3. **State what the window does and does not change about the push.** The worksheet's loss ceiling, ETA checkpoint and stop-loss stay #23's and are not suspended by the window. What the window changes is the click price and the conversion rate the worksheet assumed, so the post-window recalibration at P11 hands #23 the realized figures for its next worksheet.
4. **Record the rank read cadence.** Amplified terms carrying a daily rank series are read daily inside the window; terms outside the head set are read at the weekly cadence and the plan says which is which, because manifest group D9 is PARTIAL at head terms only.
5. **Rank gained inside the window is provisional.** Scenario S-E5 states that the two-week at-target clock starts only after the exclusion window, so a term that touches target rank on day 3 of a deal has not started its clock and no taper issues on the event spike. The taper trigger release stays #23's decision on the post-settle read.

P6. Arming at T-1: the freeze list, the baseline snapshot and deployment (30-45 min per window; planning estimate)

1. **Set Event Mode.** T2 K2 "Event Mode" flips to IN-EVENT on every affected keyword row and T3 rows carry the window reference. Downstream components read this cell as the armed-window declaration, so the flip is the arming act.
2. **Publish the freeze list.** Name every frozen class explicitly: structure changes (creations, consolidations, promotions, demotions, archives, campaign or ad-group state changes) except those the plan names; main-image and A-plus changes on affected ASINs; modifier moves outside the plan; and platform-rule actions on the affected campaigns. The list goes to #12 for its P9 overlay, to #15 for event autos, to #37 for its exclusion handling and to #38 for the session pre-flight freeze check. The freeze itself has no v4.0 gate and rides as Appendix A item AM-1, which changes the citation and not the practice.
3. **Take the baseline snapshot.** Per affected term: organic rank from T2 K4, trailing CVR and CPC from T2 K5, current bid from T2 K8, and the price state from T0 zone C. Per affected campaign: daily budget, bid, placement multipliers TOS and PP, and 14-day trailing daily spend from T3 C2 and C4. Store the snapshot as a dated block against the plan. Every post-event read compares to this snapshot and to nothing else. Without it the entire post-event story is unreadable, which is failure mode F6 and the reason arming is a procedure rather than a state change.
4. **Deploy the arming set** through #38: caps, any planned modifier changes, any event autos #15 stands up, and the state changes the plan names. Every row carries the event reference in its reason string. Deployment is same-day at arming and still takes the full T+1 verification pass; a landed-value mismatch on an arming row is a same-day redeploy.
5. **Confirm buyability at arming** and again at window open: offer live, buy box held, deal badge showing, the variation orderable. Verified with the artifact, not asserted.
6. **Write the T6 entries.** One row per deployed change with prior value, new value, the plan reference as the rule or scenario ID, the expected outcome, and the review date set at settle expiry. Unscored actions do not accumulate (rule V1).

P7. The in-window daily block (15-20 min per block; twice daily on tentpoles; planning estimate)

1. **Pacing against the projection, per campaign.** Read cumulative spend against the plan's own daily projection, not against the day-of-week curve: gate G7's event clause checks pacing daily against the event projection, and scenario S-E2 sets the correction trigger at a deviation above 15%. Record the deviation as a number in the block; "on pace" without a figure is not a read.
2. **Units and orders against the projection.** The pacing read alone cannot separate the two branches S-E2 defines, because over-pacing with matching sales lift and over-pacing without it take opposite actions. Both numbers are read in the same block or neither is.

3. **Buyability and badge.** Offer live, buy box held, badge showing. A badge that drops mid-window is a calendar-side event and routes to the deals owner the same day with the affected hours quantified.
4. **Rank watch on the amplified terms.** Daily where D9 carries the series. The window is a discounted rank push and its velocity is part of the plan's return, so the rank line is recorded even on days when no decision follows from it.
5. **Competitor deal check on conquest targets.** A competitor deal live on a target means standing down on that target for the window and resuming at window end, per scenario S-E3 routing to Section 8.7 "Conquering Scenarios (S-C1 to S-C6)", scenario S-C5: conquering into a competitor's deal price buys the worst comparison of the month. While manifest group D8 reads MISSING this check cannot run mechanically and the block records that it could not.
6. **Halo lift.** Record family campaign spend and orders daily; halo spend rides the same caps and an uncapped halo is the same F1 defect one level out.
7. **No reactive bid surgery.** In-window ACOS against standing bands is not a decision input. The plan's own targets govern, gate G9 has already marked the reads PROVISIONAL, and gate G3 has marked every capped campaign's rates TRUNCATED. Three separate v4.0 rules say the same thing about the same numbers, which is why F2 is a withdrawal and not a debate.

P8. Cap-hit decisions and the plan stop condition (10-15 min per decision; planning estimate)

1. **A cap hit is a decision, and silence is not one.** When a campaign reaches its cap, the block records the hour, the cumulative spend, the cumulative units, the deviation against the projection, and one of three outcomes: RAISE within the plan's named allowance, HOLD at cap, or the plan's stop condition. Each writes a T6 row with its evidence line.
2. **RAISE requires the allowance and the worksheet.** The raise is legal only up to the allowance the plan named at P5 or P2, and the worksheet check computes projected cost per order at the raised cap against the window contribution from P1 step 5. A raise whose cost per order clears the window contribution is funded; one that does not is a HOLD regardless of how the units are pacing.
3. **HOLD is the default on the under-performing branch.** Over-pacing without matching sales lift is not a funding case, it is a leak, and scenario S-E2 directs truncating it the same day rather than raising into it.
4. **Under-pacing on priority terms** takes the other S-E2 branch: budget or top-of-search steps toward the projection, sized inside rule L2's limits of 50% per change and the \$0.50 floor and inside rule M4's effective-CPC ceiling.
5. **Human confirmation.** Rule P5 puts every budget move above \$50 per day, every bid move above 25% and every gate failure in the human-mandatory tier, and the Authority Matrix "Event plans (#32)" row holds every cap RAISE with the plan's two D-holders. A raise inside a tentpole is confirmed before it deploys, without exception.

```

Cap reached on an affected campaign (P7 block found it)
|
|-- Is the campaign in the plan's affected set with a written cap?
|   |-- no ..... F1: cap it now from the plan's basis; root-cause the plan omission
|   |-- yes -> continue
|
|-- Units pacing at or above the plan projection?
|   |-- no (over-pacing without matching lift) ..... S-E2 truncate branch: HOLD at cap, find the leak the same
|   |                                               day (placement, irrelevant term, runaway bid), stage the fix
|   |-- yes -> continue
|
|-- Does a named allowance exist for this campaign in the plan?
|   |-- no ..... HOLD at cap; the shortfall is quantified in the block and the plan
|   |                                               gap is named at P11 scoring. No unallowanced raise, ever.
|   |-- yes -> continue
|
|-- Worksheet check: projected cost per order at the raised cap <= window contribution per order?
|   |-- no ..... HOLD at cap; record the arithmetic that refused the raise
|   |-- yes -> continue
|
|-- Confirmed by both plan D-holders per the Authority Matrix row?
|   |-- no ..... HOLD until confirmed; the window does not wait for a decision it
|   |                                               can make later, but it does wait for the authority
|   |-- yes ..... RAISE to the allowance ceiling, stage one #38 row per campaign
|   |                                               with the event reference, log the T6 evidence line
|
|-- Separately, every block: spend past the plan's stop threshold WITH units under it?
|   |-- yes ..... PLAN STOP: execute the plan's own named stop the same day inside
|   |                                               the window; this is the one in-window action that reduces scope

```

P9. Close and revert (30-45 min per window; planning estimate)

1. **Run the revert checklist** at the window end timestamp, not on the next Monday. Caps restore to their P6 baseline values; any allowance raise steps out; event modifiers step out through #28's normalization; event autos close per SOP-15 P6.2, which requires a fresh verdict, a fresh objective tag and a pre-load check before any event auto becomes a standing campaign; any hourly boost restored for the window is removed; and any campaign state changed for the window returns to its baseline state.
2. **Deploy the revert as one staged set** per #38 with the event reference, and read the T+1 verification diff against the P6 baseline snapshot line by line. The window's gate artifact at close is the executed revert checklist with its verification diff attached, not the operator's recollection that it was done.
3. **Quantify anything that did not revert.** A premium or a cap living past close is F7, and the correction carries the number of days it lived and the spend above baseline it produced. Premiums are rented by the window and returned with it.
4. **Do not lift the freeze at close.** The freeze on structure changes runs through settle, because a structure change deployed into a settling row destroys the only clean comparison the post-event read has.

P10. The settle: exclusion marks and the quarantine of window data (10-15 min to set, then standing; planning estimate)

1. **Set the marks.** Every affected row in T2 carries the event exclusion mark for the window days plus the 14 days following, per gate G10. Standing verdicts do not fire on marked rows: the weekly pass reads them for gates only, exactly as SOP-12 P9 step 3 does on its side.
2. **Quarantine the data from trailing averages.** A 7-day window plus a 14-day settle removes 21 days from every trailing series, so the four-period trend columns on T2 K5 (ACOS L30, P30, B30, L60) cannot produce a trend verdict on these rows until the clean days rebuild. The pass states which rows are unreadable rather than reading them thinly, because a 30-day window holding 21 excluded days leaves 9 days, which will not clear gate G2's sufficiency bars on most terms.
3. **The hangover is expected and reads against the baseline.** The post-deal dip compares to the P6 baseline snapshot and never to the window peak. A recovery push staged during settle is F5: the hangover cause class resolves by calendar and belongs to #24, and only a post-settle position still below the armed baseline is a case at all.

4. **Publish the exclusion dates** to #35, which reads them when it voids or defers an incrementality test, and to the weekly audit so the exclusion is visible on the pass rather than remembered.

P11. Post-event scoring and class-factor calibration (45-60 min per window; planning estimate)

1. **Score at settle expiry, not at close.** The honest measure of a window is the post-settle position against the armed baseline; the window peak measures nothing that persists.
2. **Compute the event record** across five lines: units against predicted units; spend against total caps including any allowance raises; the post-settle rank delta per amplified term against the P6 snapshot; halo lift in orders; and the deal fees the calendar entry carried. v4.0 carries no rule for how deal fees enter the event profit-and-loss or the contribution used for the window loss line, so the fee is reported as its own line and not folded into contribution, and the gap is Appendix A item AM-12.
3. **Score the prediction** HIT, PARTIAL, MISS or WRONG-DIRECTION on the T6 row, per the scoring vocabulary the Data Architecture carries at T6. A window scoring PARTIAL states which limb hit and which missed, in numbers.
4. **Calibrate the class factor.** Realized spend multiplier equals realized daily spend divided by the P3 baseline; realized unit lift equals realized units divided by predicted units. The Event Lift Index for this product and class moves toward the realized figures, exactly as push worksheets calibrate from realized velocity. The class playbook is a compounding asset: the second window of a class is priced from the first, and the fifth is priced well.
5. **Return the recalibration** to the components that assumed the old figures: to #23 for the amplified pushes' next worksheet, to #28 where a premium was tested, and to the deals calendar owner as the PPC-side read on the class.
6. **Route the post-settle read.** The three-branch decision tree below closes the window.

```

Post-settle read on each amplified term (measured against the P6 baseline snapshot)
|
+-- Position at or better than baseline?
|   +-- yes, and at or better than target rank for the required period after the exclusion window
|   |   ..... The at-target clock now starts per S-E5 and T1; the taper decision is #23's
|   +-- yes, but not at target ... Standing governance resumes on the row at the next pass; no event action
|
+-- Position below baseline?
|   +-- The named cause is the window itself and no other log event intervened
|   |   ..... Recovery case: hand to #24 with the P6 snapshot, the window record and the
|   |   |   settle dates attached. #24 dates the loss and names the cause; this SOP does
|   |   |   not diagnose. Cited per Ledger row 24.
|   +-- Another T5 log event lands inside the settle (price move, listing change, OOS, competitor entry)
|       ..... Attribution is contested: the handoff to #24 carries both events and states
|           that the window is one candidate cause, not the finding (rule D5, rule X9)

```

P12. Liquidation windows on terminal-tagged units (40-60 min per window; planning estimate)

1. **The objective function switches and this SOP does not switch it.** Component #27 owns the retag to liquidation, the R2 allowable-ad-cost ceiling, R3 cliff pacing, R4 what shuts off, R5 the structural quarantine into dated liquidation campaigns and R6 the weekly expiry check. This SOP contributes the window plan around them.
2. **Take the inputs, do not derive them.** The LTSF row supplies salvage-parity price, the alternative's per-unit value, remaining units and the cliff date. The LTSF Master Decision and Operations System is outside this corpus per conflict C6, so PPC computes nothing it can pull from that system and stops at the boundary where the tier decision lives.
3. **Build the cap from the ceiling, not from the lift index.** On a liquidation window the binding number is #27's R2 ceiling expressed as a daily figure: allowable ad cost per unit times the R3 required units per day. That is the cap. The Event Lift Index sizes standing windows; it does not outrank a recovery ceiling.
4. **The window's value is the CVR move.** A deal price raises conversion, which raises the R2 maximum liquidation CPC proportionally, which is precisely how a lane that arithmetic closed at baseline conversion can open at the window price. Where the lane still cannot buy the required clicks at the ceiling, the finding routes back to the LTSF system rather than becoming a subsidy, and closing a lane on arithmetic is a correct outcome.
5. **Predict in units against the cliff schedule,** not in rank or share: units cleared, remaining units at close, and the revised units-per-day the remaining cliff days require. Rank, share and defense metrics do not apply to these units.
6. **Settle differently.** Gate G10's exclusion protects trend verdicts, and terminal units have no trend verdicts to protect. The settle on a liquidation window is #27's R6 weekly check resuming, not a 14-day quarantine of standing governance.

P13. The annual seasonality calendar and pre-peak timing (3-4 hours for the annual build, 10 min per weekly read; planning estimate)

1. **Deal windows and demand seasonality are different machinery** and only the first has a calendar. This procedure absorbs the scheduled seasonality fold-in registered against this component: the index machinery, the pre-peak timing rule and the peak-readiness checklist.
2. **The weekly read.** On the Monday pass, read T2 K1 "Seasonal SV Index Dir" per syntax on every product. A rising index on a syntax gives its candidate pushes priority in the M5 funding order and pushes proposed after the peak are deferred to the next cycle year with the finding logged, per v4.0 Section 8.12 scenario S-E6 and Section 11 rule T4: pushes time into a rising index, never after the peak.
3. **Why timing beats budget here.** A push funded into a rising index buys rank while demand is climbing and holds it through the peak; the same dollars spent after the peak buy the same rank into falling demand and the position depreciates while it is being paid for. The index is a standing input to the O3 realism check and to M5 allocation, so the timing decision is made in the funding queue and not in this SOP.
4. **The annual calendar.** Build one calendar per year covering the category's demand curve alongside the deals calendar's confirmed windows, so that push funding, inventory cover and event planning are read against the same twelve months. For the bedding category the peak sits in the fourth quarter, which means the calendar is built before the fourth quarter opens rather than inside it. v4.0 carries the timing rule but no calendar cadence or build date, so the date lives in the program cadence and the gap is Appendix A item AM-9.
5. **Peak readiness.** Before the peak opens, confirm four states per product and record them: inventory cover clearing the peak plus the settle at peak velocity, the margin table refreshed so ceilings are not provisional through the highest-spend weeks of the year, the Event Lift Index seeded for every class the peak will run, and the affected-set list drafted for the tentpoles already on the calendar. A peak entered with a provisional margin table prices every ceiling in its highest-volume weeks on a stale number.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds and scenario logic live in v4.0. The verdict strings below are quoted verbatim from the closed vocabulary in the Keyword Decision Record Template v3.0, "Bands & Verdicts Reference" sheet, which is the resolution conflict C10 records for the three different counts the corpus carries. Two strings that v4.0 Section 13.1 "The Closed Verdict Vocabulary" carries and the KDR twenty-five does not are quoted from v4.0 with their source named, per the resolution of conflict C18. Two rows carry a v4.0 cross-map because the two documents name the same action differently, per C10.

v1.0's five verdict strings are retired. SOP-32 v1.0 Section 5 mapped ARM EVENT PLAN, CAP HIT: RAISE (per allowance) / HOLD, PLAN STOP CONDITION, REVERT + SETTLE and SCORE EVENT. None appears in the closed vocabulary, in v4.0 Section 13.1, or in the campaign-grain subset. They are plan states and procedure names, not verdicts, and this version records them as such: arming is P6, the cap-hit call is P8, the stop is the plan's own named condition at P2 step 7, the revert is P9 and the score is P11. Retiring them removes five strings that a keyword decision record could never legally carry.

Verdict (KDR v3.0 closed vocabulary, verbatim)	Execution inside or around an event window
QUARANTINE: DATA REFRESH	No changes deploy on the row (gate G1). Inside a window the refresh task still runs same day, because a quarantined row inside an armed set means the cap adherence read for that campaign is unverifiable and the block records it.
UNLOCK TRAFFIC	The under-pacing branch of scenario S-E2 where clicks run under the projection with rank velocity intact: rule L2 unlock steps on the named campaign, inside the 50% per-change limit. Never issued on a window-distorted rate read.
FIX BUDGET (TRUNCATED)	The standing state of a capped campaign in the back half of an event day. Inside the window the budget lever is bounded by the cap and the allowance, so the response is the P8 decision, not an uncapped raise. Outside the window it executes normally.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	An open test entering a window pauses its click quota rather than counting event clicks toward it: event clicks are not comparable to the test's baseline. The test resumes at settle expiry with the in-window clicks excluded from the quota.
FUND TEST (UNFUNDED PRIMARY)	Not executed inside a window. Scenario S-E1 bars new experiments entering the window; the funded test stands up at settle expiry.
RANK PUSH: FUND SIZED PUSH	Not originated inside a window (F8). A push funded before arming is amplified per P5 with its allowance; a push funded during the window deploys at settle expiry unless the plan itself named it at P2 step 3.
RECOVERY PUSH	Not executed during settle (F5). The hangover resolves by calendar and belongs to #24; only a post-settle position still below the P6 baseline routes there, per the P11 tree.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	Executes on an amplified push inside the window when clicks run under the worksheet line with rank velocity intact, bounded by the cap and the named allowance. The named lever only, never both levers on one row in one cycle.
FIX CONVERSION FIRST; DEFER PUSH	Suppressed inside the window and through settle: gate G9 marks in-window CVR reads PROVISIONAL and no fix-first or conversion verdict issues on a provisional read. A conversion deficit visible in the window is re-read on clean days after settle expiry.
PLACEMENT FIX FIRST	#28 executes; this SOP deploys only the campaign-level rows staged back and books the removal on the revert checklist where the change was made for the window.
RE-POINT VARIATION	Executes same day even inside a frozen window: gate G6 RED outranks the freeze, because re-pointing to the declared backup is what stops the window spending into a stockout. The re-point is added to the revert checklist so the baseline comparison stays honest.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	Not issued on an event spike. Scenario S-E5 starts the at-target clock only after the exclusion window, so the taper decision waits for the post-settle read and belongs to #23 when it comes.
HOLD; DEFEND LIGHTLY AT RANK	No window action. The row holds its cap and its posture and is read for gates only through settle.
DEFEND AT FLOOR	Branded and floor-held rows keep their floors inside the window; event traffic raises branded volume without changing the defended value. Drift from the recorded floor restores in one staged step at settle expiry, not inside the window.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	Suppressed on window-distorted rows through settle (F2 when staged anyway). The window Max Profitable CPC from P1 step 5 is the in-window ceiling reference; the standing ladder resumes on clean days.
INCREMENTALITY LADDER	No window execution. #35 reads the published exclusion dates and voids or defers its test; a ladder step inside a window measures the window, not the increment.
ADD EXACT + RESTRUCTURE	A structure change, therefore frozen unless the plan names it. The receipt holds and executes at #12 on the first clean day, and the exposure window is logged so the delay is visible.
CONSOLIDATE DUPLICATES	Frozen inside the window: consolidating mid-window destroys both halves of the comparison. Executes at settle expiry at #12.
GRADUATE FROM DISCOVERY (v4.0 Section 13.1 cross-map: GRADUATE + STEER)	Frozen inside the window; #30's issuance continues and the standup queues. Event-window discovery evidence is read after settle because event conversion is not the term's standing conversion.
NEGATE (PRE-LOAD / REACTIVE / STEERING) (v4.0 Section 13.1 cross-map: CUT / NEGATE)	Reactive negation on proven in-window waste is the one structure-adjacent action the freeze admits, because it reduces scope rather than changing it. Standards are #29's; the negative carries the event reference so the wall audit can date it.
ENTER SBV ARBITRAGE	Not opened inside a window: a format entry is a new experiment under S-E1. An arbitrage already live continues under its own posture and its rates are excluded with everything else.
OPEN / ROTATE / EXIT CONQUEST	#16's verdict. The event-side execution is the S-E3 and S-C5 stand-down: a target running its own deal is stood down for the window and resumes at window end. Opens are frozen; exits execute normally.
STOP-LOSS: CONCEDE OR DEFER	Executes on an amplified push whose worksheet ceiling is breached inside the window: the window does not suspend #23's stop-loss. Freed budget returns to the envelope at the instrument layer, never inside this SOP.
PAUSED-PENDING (REASON + RE-ENTRY)	The parking form for anything the window defers: frozen structure receipts, paused tests, stood-down conquest targets. Always with reason and a re-entry condition, and for window-driven parks the re-entry condition is dated at settle expiry.
LEAVE TO HARVEST	The term stays on discovery economics inside its source campaign at its capped bid; no exact stands up inside the window.

CONTINUE; VERIFY MOMENTUM (quoted from v4.0 Section 13.1; absent from the KDR twenty-five per C18)	The standing state of an amplified push pacing to plan inside its caps: no changes, the daily block continues, and the velocity check is scheduled. This is the most frequently issued in-window state and the KDR sheet has no member for it.
PORTFOLIO CALL: ASSIGN LEAD (quoted from v4.0 Section 13.1; absent from the KDR twenty-five per C18)	Fires where gate G11 catches two siblings amplified on one term inside one window. The call is #33's; this SOP holds both campaigns at cap until it lands rather than letting the window arbitrate ownership by spend.

5.1 The Campaign-Grain Subset

Caps, reverts and allowances are campaign-grain actions, and the Master PPC Data Architecture v1.0, Section 7 "Level 3 · Campaign Board (Tab T3)" carries a ten-value campaign verdict vocabulary at column group C6 that this SOP writes directly. These are the strings that land in T3, quoted verbatim from the Command Workbook REFERENCE tab validation list.

Campaign verdict (T3 C6, verbatim)	When this SOP writes it
FUND BUDGET	Cap deployment at arming (P6) and any confirmed allowance raise (P8). The reason string carries the event reference and the plan version.
HOLD	The default cap-hit outcome where no allowance exists or the worksheet refuses the raise, and the standing state of every affected campaign pacing to plan.
TRIM BUDGET	The S-E2 over-pacing-without-lift branch and the plan's own stop condition. The only in-window action that reduces scope.
RAISE MULTIPLIER	An evidenced event premium sized by rule M4 against the window ceiling, with its dated removal booked on the revert checklist at the same moment it is raised.
CUT MULTIPLIER	The revert of an event premium at close (P9) and the response to a TOS-VANITY decomposition found inside the window.
NEGATION PASS	Reactive negation on proven in-window waste, executed to #29 standards with the event reference on the evidence line.
CONSOLIDATE	Frozen inside the window; written at settle expiry when a duplicate found during the window is resolved.
RETAG OBJECTIVE	Written when an event auto converts to a standing campaign, which requires the fresh verdict, fresh objective tag and pre-load check SOP-15 P6.2 sets.
PAUSE	Event autos and dated event campaigns at close (P9), and stood-down conquest targets for the window's duration.
ARCHIVE	Never inside a window or its settle; archiving is irreversible and destroys the baseline comparison.

6. Worked Examples

All four examples use the standard fictional base because the Tier 3 live product economics spine was not attached to this build: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%, baseline CVR 5.2%. Values are illustrative and internally consistent; live values come from T0 zone C and T2. Max Acceptable is quoted at the corpus base figure of 26.3%, which is 75% of 35.0% presented to one decimal; the exact product is 26.25%.

WE-1 | A 7-day deal, clean case: the window ceiling falls while the traffic rises

Intake (P1). Class 7-day deal, discount 15%. Deal price = $80.00 \times 0.85 = \$68.00$. The discount comes off contribution in full, so window contribution = $28.00 - 12.00 = \$16.00$ per order. Window break-even ACOS = $16.00 / 68.00 = 23.5\%$; window Target = 11.8%; window Max Acceptable = 17.6%. Measured conversion uplift for this class on this product is 1.35x, so in-window CVR = $5.2\% \times 1.35 = 7.02\%$. Window Max Profitable CPC = $16.00 \times 0.0702 = \$1.12$, against the baseline ceiling of $28.00 \times 0.052 = \$1.46$. *The window lowers the per-click ceiling from \$1.46 to \$1.12 while raising traffic.* Both figures go in the plan header per AM-11, because v4.0 does not state which basis governs.

Caps (P3). Affected set of three campaigns with clean 14-day baselines: hero exact \$28.00/day, family broad \$16.00/day, sibling halo \$8.00/day, total \$52.00/day. The product's Event Lift Index for this class, from two prior windows, is 2.2. Daily caps per rule M9 = $28.00 \times 2.2 = \$61.60$, $16.00 \times 2.2 = \$35.20$, $8.00 \times 2.2 = \$17.60$, total \$114.40/day; window budget = $114.40 \times 7 = \$800.80$. The plan's named allowance is +20% on the hero campaign only, available on a deliberate call.

Prediction (P2 step 6). 640 units across the affected set; spend inside caps; post-settle rank on the hero term at +3 or better against the baseline; halo lift positive. Plan units per day = $640 / 7 = 91.4$.

Day 2 cap hit (P7, P8). The hero campaign reaches its \$61.60 cap at 14:00. Even-curve spend projection through 14:00 = $61.60 \times (14 / 24) = \35.93 , so the spend deviation is $(61.60 - 35.93) / 35.93 = 71.4\%$, far past the 15% line scenario S-E2 sets. Units through day 2 = 205 against a projection of $2 \times 91.4 = 182.9$, which is 12.1% ahead. Lift matches spend, so this is not the truncate branch. Worksheet check on the raise: raised cap = $61.60 \times 1.20 = \$73.92$ /day; at the in-window hero CPC of \$1.05 that buys $73.92 / 1.05 = 70.4$ clicks/day; at 7.02% CVR that is 4.94 orders/day; cost per order = $73.92 / 4.94 = \$14.96$ against the window contribution of \$16.00, clearing by \$1.04 per order. The CPC of \$1.05 sits at $1.05 / 1.1232 = 0.93x$ the window ceiling, inside the AT-CEILING band Family F1 defines at 0.8-1.0x. Both plan D-holders confirm; the raise stages as one #38 row with the event reference. Verdict written to T3: FUND BUDGET.

What was not done. The in-window ACOS on the hero campaign at the raised cap reads $14.96 / 68.00 = 22.0\%$, which is above the window Max Acceptable of 17.6% and would read as a bid-down case against any standing band. No bid moved. Gate G10 excludes the read, gate G9 has marked it PROVISIONAL because the price changed, and gate G3 has marked the campaign's rates TRUNCATED because it hit cap. The plan's targets governed and three separate rules agreed.

Close and settle (P9, P10). Rank ran 12 to 7 during the window. Caps revert to \$28.00 / \$16.00 / \$8.00, the raise steps out, no event autos to close, no premium was tested so none to normalize. Exclusion marks set for the 7 window days plus 14 settle days, which removes 21 days from every trailing series on the affected rows.

Score (P11). Total caps deployed = 6 days $\times$ \$114.40 + 1 day at the raised total of \$126.72 = \$813.12. Realized spend \$769.00, which is $769.00 / 813.12 = 94.6\%$ of caps. Units 601 against 640 predicted = 93.9%, short by 6.1%. Post-settle rank 9 against the baseline 12 = +3, exactly at prediction. Score PARTIAL: the rank limb HIT, the units limb MISSED by 6.1%. Calibration: realized daily spend = $769.00 / 7 = \$109.86$ against the \$52.00 baseline = 2.11x, so the class Event Lift Index trims from 2.20 toward 2.11 and the unit-side assumption trims 6.1%. The deal fee from the calendar entry is reported as its own line and is not folded into the \$16.00 contribution, per AM-12.

WE-2 | A first-time tentpole: what the missing category default actually costs

Setup. The product enters its first BFCM window. Five campaigns in the affected set with a combined clean baseline of \$146.00/day. No prior window of this class exists on this product, so the Event Lift Index is unseeded. Rule M9's first-time path directs borrowing the category default; v4.0 carries no category-default table and the program's open-decisions register carries the tentpole factor as having no basis at item PPC-DR-35. This SOP therefore writes no factor. The plan runs the caps at the baseline, states the missing factor on its face, and carries a named dollar allowance as the only expansion authority (AM-4, AM-5).

Day 1, first block at 11:00 (P7). Cumulative spend across the affected set \$138.70. Even-curve projection through 11:00 = $146.00 \times (11 / 24) = \66.92 , so the deviation is $(138.70 - 66.92) / 66.92 = 107.3\%$. Units through 11:00 = 96 against a plan-day projection of 118 units, whose 11-hour share is $118 \times (11 / 24) = 54.1$, so units run 77.5% ahead. Spend and lift are both far ahead: the S-E2 branch is the deliberate raise, not the truncate.

The arithmetic that makes the case (P8). Observed spend rate = $138.70 / 11 = \$12.61/\text{hour}$. The \$146.00 cap exhausts at $146.00 / 12.61 = 11.58$ hours, which is 11:35, leaving 12.42 hours or 51.8% of the day with the affected set dark on the highest-conversion day of the year. At the observed rate the forfeited spend would be $12.42 \times 12.61 = \$156.62$. Observed conversion efficiency = $96 \text{ orders} / \$138.70 = 0.692 \text{ orders per dollar}$, so the forfeited orders = $156.62 \times 0.692 = 108$. At the window contribution of \$16.00 per order that is $108 \times 16.00 = \$1,728.00$ of contribution not bought, on day 1 of a 2-day event, from a single missing table value.

The decision. The plan's named allowance covers a raise of \$156.62 for the day, taking the deployed cap to $146.00 + 156.62 = \$302.62$, staged as one #38 row per campaign with the event reference and confirmed by both D-holders per the Authority Matrix row before it deploys. Rule P5 puts it in the human-mandatory tier independently, because the move exceeds \$50 per day. The second block runs at 18:00 per the tentpole cadence, which itself has no v4.0 basis and rides as AM-7.

What this example is for. It is not an argument for a factor. It is the quantified cost of not having one, computed so that AM-4 arrives at the threshold owner with a number attached instead of a preference. Until v4.1 lands, the plan allowance plus the S-E2 daily correction plus gate G7's daily event cadence are the three bounds that keep the exposure to one day of a known size.

WE-3 | The edge case: an excluded read that got a verdict, and a hangover that was not a case

Setup. The WE-1 window has closed. The hero campaign's window record: spend \$421.00, clicks 401 at a \$1.05 CPC, 28 orders at 6.98% CVR, sales at the deal price = $28 \times 68.00 = \$1,904.00$.

What went wrong. The Monday pass after close staged a scale step on the hero row. The workbook's economics block (T2 K6) references T0's named cells, which still carry the baseline AOV \$80.00 and contribution \$28.00 because the price has restored. The pass therefore read ad cost per order = $421.00 / 28 = \$15.04$ against a \$28.00 contribution and saw \$12.96 per order of headroom. The true in-window surplus was $16.00 - 15.04 = \$0.96$ per order. The headroom was overstated by $12.96 / 0.96 = 13.5$ times. The same rows read two different ACOS figures depending on the join: against deal-price sales, $421.00 / 1,904.00 = 22.1\%$; against list-price sales of $28 \times 80.00 = \$2,240.00$, $421.00 / 2,240.00 = 18.8\%$, which flatters by 3.3 percentage points and reads inside Max Acceptable.

The response (F9). The staged step is withdrawn before deployment. The row carries the gate G10 exclusion mark for the window days plus 14, so no trend verdict was legal on it in the first place; the read that produced the verdict was excluded data. Had it deployed at the L4 single-step 25% on a post-window baseline of \$28.00/day, that is \$7.00/day across the 14 remaining settle days, or \$98.00 of spend carrying no scoreable prediction. The withdrawal is logged with its own T6 line naming the exclusion mark that should have blocked it, because the mark existed and the pass did not read it.

The second failure (F5). On settle day 3 the hero term reads rank 14 against the baseline 12 and a recovery push is staged. It is blocked at the gate: the post-deal dip is expected, it reads against the P6 baseline and not against the window peak of rank 7, and the hangover cause class resolves by calendar. The drift is quantified in the T6 note rather than acted on.

The read that was a case (P11). At settle expiry on day 15 after close, rank reads 13 against the baseline 12: still below baseline, so this is a recovery case. The T5 log carries no other event inside the settle, so the window is the named candidate cause and the handoff to #24 carries the P6 baseline snapshot, the window record and the settle dates. #24 dates the loss and names the cause; this SOP does not diagnose, and the handoff cites Ledger row 24 rather than a procedure number because those numbers move at its Wave 5A rebuild.

The cost of the exclusion arithmetic. A 7-day window plus a 14-day settle removes 21 days. The 30-day trailing column on T2 K5 therefore holds 9 clean days after the window, which on a term running 401 clicks in 7 event days and roughly 20 clicks per clean day yields about 180 clicks: enough for a bid read at gate G2's 15-click bar, short of the 100-click bar for a CVR verdict on most weeks. The pass states which rows are unreadable rather than reading them thinly.

WE-4 | A liquidation window: the deal price opens a lane that arithmetic had closed

Setup, inputs taken from the LTSF row and not derived here. 480 units tagged terminal with the cliff 40 days out. Salvage price nets \$6.40 per unit; the LTSF alternative nets -\$3.10 per unit. Under #27's rule R2 the allowable ad cost per unit = $6.40 - (-3.10) = \$9.50$, and under R3 the required pace = $480 / 40 = 12.0$ units per day. At the baseline converter CVR of 4.0% the maximum liquidation CPC is $9.50 \times 0.04 = \$0.38$, the required clicks are $12.0 / 0.04 = 300$ per day, and the converters' market CPC reads \$0.55: the lane cannot buy 300 clicks at \$0.38 and closes on arithmetic, which is a correct outcome and routes back to the LTSF system.

What the window changes. A 7-day liquidation deal raises the converter CVR to 6.5%. The R2 ceiling moves with it: maximum liquidation CPC = $9.50 \times 0.065 = \$0.62$, now above the \$0.55 market CPC. Required clicks fall to $12.0 / 0.065 = 184.6$, call it 185 per day. The lane opens because conversion moved, not because anyone bid more.

The cap (P12 step 3). The binding number is the ceiling, not the lift index: daily cap = allowable ad cost per unit $\times$ required units per day = $9.50 \times 12.0 = \$114.00$. At \$114.00 and a \$0.55 CPC the campaign buys $114.00 / 0.55 = 207.3$ clicks per day, which at 6.5% CVR clears 13.47 units per day against the required 12.0. Cost per unit = $114.00 / 13.47 = \$8.46$ against the \$9.50 allowable, inside by \$1.04.

The prediction (P12 step 5), in units against the cliff. Across 7 window days the campaign clears $13.47 \times 7 = 94.3$, call it 94 units. Remaining units fall from 480 to 386 with 33 cliff days left, so the required pace falls from 12.0 to $386 / 33 = 11.7$ units per day. The window buys 0.3 units per day of slack against the cliff. That is the honest return and it is what the plan predicts; it is not a rank number, a share number or an ACOS number, because none of those applies to these units under #27's rule R4.

Settle (P12 step 6). No 14-day quarantine of standing governance runs here, because gate G10 protects trend verdicts and these units have none. What resumes is #27's weekly expiry check comparing realized recovery per unit against the alternative. The week the alternative wins, PPC stops that week, and the window does not buy an exception to that.

7. Failure Modes and Escalation

F1 to F7 carry forward from v1.0 Section 7 with their recorded incident motivations. F8 to F12 are new at v2.0. Each row states the detection signal, the response, and the account-history incident that motivated it where one exists.

Failure	Detection signal	Response	Motivating incident
F1. Uncapped campaign inside an armed window	A campaign in the affected set with no cap row in the plan; T3 C3 shows no cap against a live window reference	Cap now from the plan's own basis; quantify the uncapped hours and the spend above the baseline curve; root-cause the plan omission at P11 scoring	Recorded in v1.0 Section 7 as the overspend-guard failure, worst in tentpoles. Event auctions are where dynamic spend runs away, which is why the cap and not the bid is this document's primary instrument
F2. Bid surgery on window-distorted reads	A standing-band verdict staged against in-window data: any bid or ladder row whose evidence dates inside the window or the settle	Withdraw before deployment; the plan's targets govern inside the window; re-read on clean days after settle expiry. Three v4.0 rules independently bar the read: G9 PROVISIONAL, G3 TRUNCATED, G10 excluded	Recorded in v1.0 Section 7. The account's pattern is a post-window bid-down on rows whose ACOS was a price artifact, and WE-3 quantifies the mirror-image error at 13.5 times overstated headroom
F3. Structure or main-image change inside the window	The P6 freeze list checked against the T5 Product Log and the T6 decision log; a structure entry dated inside the window with no plan reference	Revert where reversible; quantify the read damage in excluded days; root-cause the session that deployed it; extend the exclusion window from the revert date rather than from the window close	Recorded in v1.0 Section 7; the standing main-image freeze inside active windows is an account rule the Transition Brief v1.2 records as riding with every confirmed window. SOP-12 v2.0 carries the mirror failure as its own F8
F4. Event autos live past close	A window-named campaign live past its expiry date with no conversion verdict logged	Close per SOP-15 P6.2; conversion to a standing campaign requires a fresh verdict, a fresh objective tag and a pre-load check. Dated names exist to make this visible from the campaign list alone	Recorded in v1.0 Section 7 and carried as SOP-15's own F11, which gives the rule its detector on the auto side
F5. Hangover panic	A recovery push, bid raise or budget raise staged on an affected row during the settle period	Block at the gate; the hangover reads against the P6 baseline and resolves by calendar; only a post-settle position still below baseline is a case, and it routes to #24 with the snapshot attached	Recorded in v1.0 Section 7. Spending to arrest an expected dip pays twice: once for the window and once for the settle it was always going to have
F6. No baseline snapshot	The P6 arming record missing rank, CVR, CPC, bid, budget or multiplier values for any affected row	Reconstruct from the Raw Pull Pack and the T3 pull dates where possible; where it cannot be reconstructed the gap is named in the P11 score and the window's rank limb is recorded as unscorable rather than assumed	Recorded in v1.0 Section 7. Without the snapshot the entire post-event story is unreadable and the class factor cannot calibrate, which costs the next window as well as this one
F7. Revert missed	Caps, allowance raises, premiums, boosts or state changes live past close; the #38 T+1 diff against the P6 baseline snapshot	Re-run the revert checklist; quantify the days lived and the spend above baseline; premiums are rented by the window and returned with it	Recorded in v1.0 Section 7 as the same class SOP-28 carries on its side: a premium that outlives its evidence is paid for indefinitely because nothing re-asks whether it earned its level
F8. A push originated inside the window	A T2 K7 row gaining a Ranking objective with a funded worksheet dated inside the window, or a new dedicated campaign standing up against a window-dated verdict with no plan reference	Withdraw and re-date to settle expiry unless the plan named it at P2 step 3; scenario S-E1 bars new experiments entering the window. A push sized on event conversion is sized on a number that ends with the window	New at v2.0. The boundary is stated in v1.0 Section 1 (pushes are amplified, never originated) but v1.0 gave it no detector, so nothing found the violation
F9. A trend verdict fired on an excluded row	Any verdict on a row carrying the gate G10 exclusion mark whose evidence window overlaps the window or the settle; the weekly pass gate-check column	Withdraw before deployment; log the withdrawal with the mark that should have blocked it; where it deployed, void the week's reads on the affected rows and quantify the spend that carried no scoreable prediction	New at v2.0, from the WE-3 pattern. The exclusion is a gate the pass must read, and a mark that exists but is not read is the same as no mark at all
F10. A Recovery-Mode product run at full event posture	An affected product whose T0 zone A verdict or Account Rollup A0 code-red flag reads Recovery, carrying a plan with discovery, conquest or unproven-term scope	Reduce the plan to defense and proven terms the same day per rule T8; the reduction is a plan revision, not an in-window improvisation; the code-red owner is informed per the Authority Matrix code-red row	New at v2.0. Recovery outranks Event Mode in v4.0 and neither v1.0 nor its plan template read the mode before building the affected set
F11. Competitor-deal stand-down missed on a conquest target	A conquest target carrying a live competitor deal flag with the campaign still bidding; T2 K8 "Competitor Depth (D8)" where the group is present	Stand down on that target for the window and resume at window end per S-E3 and S-C5; where manifest group D8 reads MISSING the block records that the check could not run, and the finding downgrades to a flag per gate G8 rather than being dropped	New at v2.0. Conquesting into a competitor's deal price buys the worst comparison of the month at the highest CPC of the month, and the detection depends on a feed the Ledger carries as open dependency D4
F12. A window armed into insufficient cover	Days to stockout on any affected variation below window days plus settle days at projected window velocity, read at P1 step 6 or breached in flight	At intake, return the entry with the shortfall in units. In flight, gate G6 governs: YELLOW freezes push spend at current level, RED tapers to defense and re-points to the declared backup the same day, and the re-point outranks the freeze	New at v2.0 as a named failure with a detector; v1.0 stated the rule in its preconditions without one. A deal into a stockout converts paid momentum into a recovery case and forfeits the deal fee as well as the spend

Escalation. F3 escalates the same day: a freeze breach inside a live window damages a comparison that cannot be rebuilt. F1 found mid-tentpole escalates immediately with the uncapped hours and the spend above the baseline curve quantified. F12 escalates the same day to the deals calendar owner as well as the PPC Lead, because the remedy may be a calendar decision rather than a PPC one. F9 recurring escalates to the PPC Lead with the pass records, because an exclusion discipline that fails twice is a process gap rather than a slip. F10 escalates to the code-red owner named on the Authority Matrix "Code-red declaration + close" row. Anything in the human-mandatory tier of Section 9 is confirmed per the Authority Matrix before deployment, without exception.

8. Outputs and Artifacts

Gates clear with artifacts, not assertions. Every row names the record, where it is stored, how long it is kept and who consumes it.

Event	Artifact set produced	Storage and retention	Downstream consumer
Intake and classification (P1)	The classified calendar read: class, dates, deal price, deal fee, affected ASINs and variations; the product posture header (stage, mode, gate strip, earning class); the window economics block with both bases computed; the inventory cover read against window plus settle	The plan document against the calendar entry; one T5 Product Log row with Event Type "Deal window" and the trend-exclusion flag set. T5 is the closed event list and is retained for the life of the product	Every later attribution check (rule D5); #24 at any post-settle recovery handoff; the deals calendar owner where the entry returns
The event plan (P2 to P5)	The affected-campaign table with baselines; the halo block with its membership test; the amplified-push list with allowances; the cap table; the modifier posture with any booked removal; the contested-root notes; the four-line prediction; the named stop conditions; the two signatures	The plan document, versioned, stored against the calendar entry and referenced by every staged row in the window. Retained through the following year's window of the same class so the next plan is built from the last one	#12, #15, #16, #23, #26, #29, #33, #37, #38 all read the plan; the BM owner reviews it at the daily huddle
Arming (P6)	The freeze list with its four named classes; the baseline snapshot per term and per campaign; the deployed arming set with its #38 T+1 verification diff; the buyability confirmation; the T6 rows with prediction and review date	T2 K2 Event Mode and T3 window references carry the live state; the snapshot stores as a dated block against the plan; T6 is append-only and never purged	#12 P9, #15 P6, #37 exclusion handling, #38 session pre-flight; the P11 score reads the snapshot
In-window blocks (P7, P8)	One dated block per day, twice daily on tentpoles, each carrying: per-campaign spend against projection with the deviation as a number, units and orders against projection, buyability and badge state, rank lines on amplified terms, the competitor-deal check or the note that it could not run, halo spend and orders, and every cap-hit decision with its evidence line and worksheet arithmetic	T6 decision log rows for every decision; the block record against the plan. T6 retained per the workbook's append-only rule	The P11 score; #23 for amplified pushes; the PPC Lead on any escalation; V5 scoring at the review dates
Close (P9)	The executed revert checklist with each line marked and dated; the revert staged set; the #38 T+1 diff read line by line against the baseline snapshot; the quantified exception list for anything that did not revert	Against the plan; the diff stores with the #38 session record	#28 where a premium normalized; #15 where an event auto closed; #37 where exclusion states lift; the P11 score. The executed checklist is the window's gate artifact at close
Settle (P10)	The exclusion marks on every affected T2 row with their start and end dates; the published exclusion date range; the settle-period T6 notes recording any drift observed and not acted on	T2 row marks for the settle duration; the published range against the plan and in the weekly pass header	#35 voiding or deferring tests; the weekly audit; #12 and every component reading rows for gates only
Scoring (P11)	The event record on five lines (units, spend against caps, post-settle rank delta per term, halo lift, deal fees); the V5 score with the limb detail; the class-factor calibration showing realized against assumed on both the spend and unit sides; the recalibration returns	T6 with the V5 score written against the arming prediction; the class factor updates the product's Event Lift Index series, retained across windows because the series is the asset	#23 next worksheet; #28 where a premium was tested; the deals calendar owner; Account Rollup A3 verification scoreboard; the next window of this class
Liquidation windows (P12)	The R2 ceiling and R3 pace read at the window price; the daily cap derived from the ceiling; the units-against-cliff prediction; the units cleared and revised pace at close	Against the plan and the dated liquidation campaign record #27 owns	#27's weekly expiry check; the LTSF owner; the deals calendar owner where the lane still cannot clear
Seasonality (P13)	The weekly index read per syntax per product; the annual demand calendar alongside the confirmed-window calendar; the peak-readiness confirmation on its four states	The annual calendar as a standing document, re-issued yearly; the weekly read as a line in the Monday pass record	The M5 funding queue at T0 zone H and Account Rollup A2; #23 candidacy timing; inventory planning; the deals calendar owner

9. Skill-Conversion Block

Block	Specification
Input bindings (manifest field groups and named cells)	The confirmed deals-calendar entry (Economics and Context Pack Tab 4) with class, dates, deal price and deal fee; A2 deployment state for the affected set (T3 C2); A3 pacing and in-budget with the day-of-week curve where present (T3 C3); A7 eligibility and the deal-badge state; B8 and D1 economics chain from the T0 zone C named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC with the margin-table version and G9 state; G6 inventory band and days to stockout per affected variation (T2 K2); D8 competitor depth for conquest targets in the set (T2 K8); D9 daily rank series on amplified terms (T2 K4); the E composite group carrying the product's Event Lift Index per class; T2 K7 push worksheets for amplified terms; T2 K2 Family Conflict and K8 Term Owner for contested roots; the T5 Product Log; the #38 T+1 verification diff; the LTSF row where a terminal tag is present. An absent feed suppresses its dependent execution class per the manifest and the skill states what it could not see; it never improvises around a missing feed.
Deterministic logic references	v4.0 Section 8.12 scenarios S-E1 to S-E6 and Section 8.7 scenario S-C5; rule M9 cap sizing and rule M1 the break-even chain; rules L1, L2, L3, L4 and M4 for lever order, limits and multiplier sizing; gates G1, G2, G3, G4, G6, G7, G8, G9, G10, G11 and G12 by their v4.0 numbers; rules T3, T4 and T8 with rule P7 for mode precedence; rules P2, P5 and P10 for precedence, review tier and record completeness; rules V1, V4 and V5 for prediction, cadence and scoring; #27 rules R2, R3, R4 and R6 on liquidation windows; this SOP's Section 5 mapping and its 5.1 campaign-grain subset. The skill executes mappings, never computes its own thresholds, and version-pins v4.0.

Human checkpoints, drawn from the Authority Matrix v2.0 rows that apply	Per the <i>"Event plans (#32)"</i> row (D: PPC Lead with the deals owner; E: incoming operator, launch and CA operator in own scope, PPC Manager; R: BM owner): human-mandatory are plan sign-off by both D-holders, every cap RAISE, every execution of a plan stop condition, tentpole arming, and any plan revision inside a live window. Per <i>"Weekly verdicts (established US)"</i> (D: PPC Lead; E: assigned scope; R: BM owner at the daily huddle): the resumption of standing governance at settle expiry and the withdrawal of any verdict staged on an excluded row. Per <i>"Push funding (the M5 queue)"</i> (D: PPC Lead within envelope, CEO above it): any in-window allowance that would take the affected set past the product envelope. Per <i>"Bulk deployments (#38)"</i> (E: own scope; R: PPC Lead): arming, raise and revert sets stage and deploy as E. Per <i>"Code-red declaration + close"</i> (D: PPC Lead): the F10 posture reduction on a Recovery-Mode product. Per <i>"Terminal/LTSF PPC lane"</i> (D: PPC Lead with the LTSF owner): every liquidation window plan. Per <i>"Sibling declarations (S-series)"</i> (D: PPC Lead): the G11 hold on a contested root inside a window. Per <i>"Structure creation (all SOPs)"</i> : any structure change the plan names inside a frozen window. Standing rule 1 of the matrix requires the shared D on the event-plan row to name its split, which Section P2 step 8 does; standing rule 5 reverts an unfilled role's D to the PPC Lead and its E to the PPC Manager. AI-executable with audit: plan drafting from the calendar entry, cap arithmetic, the window economics block, baseline snapshot assembly, pacing and deviation lines, cap-hit worksheet computation, revert checklist drafting, exclusion mark setting, the event record and the class-factor calibration, and F1 to F12 detection with evidence lines.
Cadence	Intake at the class lead time; plan drafted and signed not later than 48 hours before arming; arming at T-1; daily in-window blocks with a second block on tentpoles; cap-hit decisions same day inside the block that found them; revert at the close timestamp; exclusion marks held for window days plus 14; scoring at settle expiry inside the Monday pass; the seasonality index read weekly and the annual calendar built before the category peak opens.
Alarm specification	F1 to F12 each carry the Section 7 detection signal, an evidence-line format (campaign or term, signal value, the gate or scenario reference, and the window reference), and a routing tier: F1, F3, F10 and F12 same day to the operator queue and the named escalation; F2, F5, F9 and F11 into the block record and the next pass worklist; F4, F6, F7 and F8 to the PPC Lead with the T6 history attached.
Skill outputs	Drafted event plans with caps, allowances, prediction and stop conditions; drafted freeze lists and baseline snapshots; staged #38 sets for arming, raises and reverts; dated in-window block records with deviations computed; drafted cap-hit worksheets; the revert checklist; exclusion mark sets with published dates; the scored event record with class-factor calibration; and F1 to F12 alarms with evidence lines. Every draft is queued for the confirmation tier its Authority Matrix row requires.

10. Version Governance

Owner: PPC Lead jointly with the deals calendar owner, because the plan interface spans both and neither signs alone. Review after every tentpole and quarterly otherwise, and immediately on any v4.0 increment, at which point the Section 5 mapping and every S-E, S-C5, M9, M1, G3, G6, G7, G9, G10, G11, G12, T3, T4, T8, L2, L3, L4 and M4 reference re-validates before any plan is written against it. Class factors version with their P11 calibrations and the version travels on the plan.

Component-specific review triggers beyond the quarterly slot. (1) Any change to the deals calendar's own fields or cadence, which changes what intake receives and therefore what P1 can classify. (2) A v4.1 release resolving any Appendix A item, which retires that item's workaround in the same increment; AM-1 retires the freeze's gate-less status and AM-4 retires the first-window path at P3 step 4. (3) The first tentpole run on any product, because the first window of a class is the one that seeds the Event Lift Index and the plan shape for every window after it. (4) Any Amazon change to deal mechanics, badge behavior or the deal-fee structure, which changes the P11 event record and the P1 economics block. (5) The landing of manifest group D8, which converts the F11 competitor stand-down from a stated blind spot to a mechanical check and closes the corresponding line in Section 3. (6) The landing of the Marketing Stream hourly feed, which moves hourly boost handling from this component to #45 and retires P4 step 5. (7) Any rebuild of #24, #28 or #33, which re-validates the three boundary rows written from Ledger rows rather than from procedure numbers.

v2.0 change note. Supersedes v1.0 in full. Procedures expanded 7 to 13, adding P3 cap construction as its own procedure, P4 modifier and placement posture, P5 push amplification and the allowance, P8 cap-hit decisions with its decision tree, P10 the settle as a procedure rather than a clause, and P13 the annual seasonality calendar, which absorbs the scheduled seasonality fold-in registered against this component. Verdict mapping rebuilt from five locally invented strings to the full 25-string KDR closed vocabulary plus the two v4.0-only strings per conflict C18, plus the ten-string campaign-grain subset this SOP writes into T3; v1.0's five strings are retired and mapped to their procedures. Worked examples expanded 1 to 4 and recomputed on the standard fictional base, with the v1.0 7-day case carried forward at WE-1 and its spend-against-caps figure corrected from 96% to 94.6% once the allowance raise is included in the cap total. Failure modes expanded 7 to 12 with detectors added for the two rules v1.0 stated without one. Elements 11 and 12 added per the Build Kit standard, with 18 interface contracts each carrying a non-arrival consequence. Gate citations corrected throughout to v4.0 numbering per conflict C9: the event exclusion is G10 and not G12, and the in-window freeze is recorded as carrying no v4.0 gate at all rather than borrowing one. Thresholds carried forward from v1.0 that trace to v4.0: the 14-day pre-event window, the 14-day post-event exclusion, the 15% in-window pacing correction line, the cap formula, the 70% truncation line, the inventory bands and the 50% and \$0.50 lever limits. Incidents carried forward from v1.0 Section 7: F1 to F7 with their motivations intact. Absorbs the scheduled deal-window playbook deliverable, whose Completion Directive row closes against this document; see the delivery note on that row's Definition of Done. Registered in the Supersession Map.

Companions: Master PPC Decision Criteria System v4.0 (all thresholds, gates, scenarios and verdicts) | Keyword Decision Record Template v3.0, "Bands & Verdicts Reference" (the closed vocabulary) | PPC Strategic Doctrine v1.0 Sections 2, 12, 13 (the why) | PPC Authority Matrix v2.0 (Elements 9 and 11) | Master PPC Data Architecture v1.0 and the PPC Command Workbook READ ME (instrument mechanics) | SOP-23 Ranking Push, SOP-24 Recovery, SOP-15 Auto Operations, SOP-12 Exact, SOP-27 Terminal and LTSF, SOP-28 Placement, SOP-33 Sibling Arbitration, SOP-26 Mature and Code Red, SOP-29 Negation, SOP-35 Incrementality, SOP-37 Automation, SOP-38 Bulk (boundaries) | PPC Cross-Reference Ledger (component #32 row) | P2.3 Pricing and P3.5 Deals Calendar (the deciding processes, neither held in this corpus).

11. Operator Ramp

Built from the Authority Matrix v2.0 rows that name this component's decisions: **"Event plans (#32)"** (D: PPC Lead with the deals owner; E: incoming operator, launch and CA operator in own scope, PPC Manager; R: BM owner), **"Weekly verdicts (established US)"** (E: assigned scope), **"Bulk deployments (#38)"** (E: own scope), **"Push funding (the M5 queue)"** (E), **"Code-red declaration + close"** (E), **"Terminal/LTSF PPC lane"** (E), **"Sibling declarations (S-series)"** (E) and **"Structure creation (all SOPs)"** (E). Each row becomes a demonstrated competence. The checkpoint holder for review and audit is the PPC Lead per the R and D cells, with the deals calendar owner co-reviewing everything on the plan interface; the vacancy rule at matrix Section 2 standing rule 5 reverts an unfilled role's D to the PPC Lead and its E to the PPC Manager until the role fills. Note that the matrix's standing rule 4 records the Completion Directive as still naming the PPC Lead on two rows the Transition Brief v1.2 assigns to the incoming operator; the brief controls, and this ramp is written to the brief.

Week	Demonstrated competence
Week 1: observation and artifact familiarity	Shadows one full window end to end where the calendar provides one, or reconstructs the last completed window from its artifacts where it does not. Locates and reads unaided: the deals-calendar entry and its five fields; T2 K2 "Event Mode" on an affected row; the T5 Product Log "Deal window" entry and its trend-exclusion flag; a P6 baseline snapshot block and what each line will be compared against; the T3 C3 pacing columns and which of them the day-of-week curve is missing from; and one prior window's scored T6 row with its class-factor calibration. Reproduces the WE-1 window economics block cold on fresh illustrative inputs, matching to the cent: deal price, window contribution, window break-even, window Target, window Max Acceptable and window Max Profitable CPC, and states which of the six moved in the opposite direction to the traffic. Reads the Authority Matrix "Event plans (#32)" row aloud and names who signs which half.
Week 2: supervised execution, checkpoint holder reviewing	Drafts one complete event plan on a live or historical calendar entry with the PPC Lead and the deals owner reviewing every table before anything stages: affected set with clean baselines, halo block with its membership test stated, amplified-push list with allowances, cap table computed from rule M9 with the Event Lift Index and its window count named, modifier posture with any booked removal, contested-root notes, the four-line prediction and the named stop conditions. Drafts the freeze list and the baseline snapshot and has both confirmed. Runs three consecutive in-window blocks under review, each carrying its deviation as a number on both the spend and the units line. Drafts one cap-hit worksheet and takes it to a HOLD or a RAISE recommendation with the arithmetic on the page; the confirmation is the checkpoint holder's. Pass bar for the week: zero blocks recording "on pace" without a figure, and zero plan tables with a cap that cannot be traced to a baseline and a factor.
Week 4: unsupervised execution, checkpoint holder auditing	Owns the window scope on the assigned products: intake, plan, arming, blocks, cap-hit decisions inside the named allowance, revert, settle marks and scoring, with the PPC Lead auditing artifacts after the fact rather than gating them before. The two D-cells stay with the plan sign-off and the cap RAISE, which are not delegated. The audit samples the baseline snapshot completeness, one cap computation traced to its baseline window and factor, the revert checklist against the T+1 diff line by line, the exclusion marks and their published date range, and the class-factor calibration arithmetic. Executes one full post-settle read to its correct branch, including the handoff packet to #24 if the branch calls for it. Two clean audited windows, or two clean audited months where no window falls in the period, close the ramp.

11.1 Live-Defense Questions

These test the reasoning behind the thresholds, not recall of their values.

1. A 15% discount on an \$80.00 product raises conversion by a third. Does the maximum profitable CPC go up or down, and why is the answer not obvious from either number alone?
2. Three separate v4.0 gates independently bar a bid change on an in-window row. Name all three and explain what each one is actually protecting, since they are not protecting the same thing.
3. The window peak shows rank 7 and the post-settle read shows rank 13 against a baseline of 12. Which of those three numbers is the measurement, and what would be lost by scoring on either of the others?
4. A push touches its target rank on day 3 of a 7-day deal. Why does the at-target clock not start, and what would starting it cost the account two weeks later?
5. Why is a cap a ceiling on autonomous spend rather than a target, and what specifically goes wrong in an event auction if a campaign is left uncapped rather than merely under-budgeted?
6. The plan says caps revert at close but the freeze runs through settle. Defend the asymmetry: why would lifting the freeze at close be worse than lifting the caps at close?
7. A first-time tentpole has no Event Lift Index. Explain why the correct response is a named allowance rather than a borrowed factor, and what the allowance buys that a guessed factor does not.
8. A product in Recovery Mode has a confirmed deal window. Name the precedence rule, and explain why a discount is exactly the situation where the temptation to suspend recovery discipline is strongest and the case for it is weakest.
9. Two siblings are both amplified on one root inside one window. What does the window do while ownership is undeclared, and why is spending through it the one option that is never neutral?
10. A liquidation window opens a lane that arithmetic had closed at baseline. Which input moved, which did not, and why does that distinction decide whether the lane stays open after the window?
11. The settle removes 21 days from a 30-day trailing column. Name the second gate this collides with and explain what an operator must write on the pass instead of a verdict.
12. The freeze on structure and main-image changes is enforced everywhere in this corpus and carries no v4.0 gate. Explain why that is an amendment item rather than a reason to stop enforcing it, and what would change in this document if v4.1 admitted it.

12. Interface Contracts

Built from the Cross-Reference Ledger's component-32 Interfaces cell and from the rows of the components that name #32 in their own Interfaces cells, cross-checked against Master PPC Data Architecture v1.0 for artifact names, tabs and timing. Every row states what happens when the handoff does not arrive, because that is the clause the corpus was missing.

#	Interface	Sender to receiver	Artifact and timing	Non-arrival consequence
I-1	Confirmed window	Deals calendar (P3.5) to P1	The calendar entry with class, start and end timestamps, deal price, deal fee and affected ASINs, landing in the Economics and Context Pack Tab 4 at or before the class lead time	No window is planned and none arms: this SOP plans around a confirmed decision and never infers one. A window that runs without an entry is unplanned spend at event CPC and escalates to the deals owner the same day with the hours quantified. Neither the calendar nor the pricing process document is held in this corpus, so the entry is also the only specification of its own fields
I-2	Deal price and margin approval	Pricing (P2.3) to P1 step 5	The approved deal price with its margin position, on the same entry	The window economics block cannot be computed, so every ceiling in the plan is baseline-only and every cap is unpriced. The plan states this and arms only on the deals owner's written confirmation that the price is approved; it does not estimate the margin
I-3	Economics chain	#41 refresh to T0 named cells to P1	AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC with the margin-table version and G9 freshness state, before the Monday 09:00 packet-ready read	Every window ceiling carries the G9 PROVISIONAL tag; the plan computes with the stale values, marks the ceiling lines provisional on its face, and the peak-readiness check at P13 step 5 fails until the refresh lands. A peak entered on a provisional table prices its highest-volume weeks on a stale number
I-4	Inventory cover and bands	Economics and Context Pack (D2, G6) to P1 step 6	Days to stockout and the adequacy band per affected variation with the ranked backup chain, current per the pack's daily inventory cadence	The window cannot be cleared for arming: cover against window plus settle is the one precondition with no workaround, and F12 is the failure it prevents. An unmapped variation additionally bars the RE-POINT VARIATION response, so a RED landing mid-window has no backup to move to
I-5	Competitor deal state	Data Operations (D8) to P7 step 5	Competitor price, rating, review count, deal state and stockout flags into T2 K8, weekly, for every conquest target in the affected set	The S-E3 and S-C5 stand-down cannot be detected mechanically. The block records that the check could not run and the finding downgrades to a flag per gate G8; conquest targets in the set are stood down by hand where the deals owner or the tracker reports a competitor window. Ledger dependency D4 remains open
I-6	Daily rank series	Data Operations (D9) to P5, P7, P11	Daily organic rank per amplified term into T2 K4, retained through the window and the settle	The in-window rank watch drops to the weekly cadence and the post-settle read compares two weekly points instead of a series, which cannot date an inflection. The plan records which terms carry which cadence rather than implying a series it does not have
I-7	Live push registry and worksheets	#23 to P5	The list of funded live pushes on affected products with their loss ceilings, ETA checkpoints and stop-loss conditions, read from T2 K7 before the plan is drafted	No allowance is written for that push and the window amplifies it without one, so a cap hit on a push campaign resolves to HOLD every time. The push's own stop-loss stays #23's and is not suspended, so the exposure is bounded, but the window forfeits the amplification it was planned for
I-8	Armed-window declaration	P6 to #12, #15, #16, #22, #23, #26, #29, #37	Event Mode set on T2 K2 with the window dates, plus the freeze list, published at T-1 and readable by every consuming component. A recorded absence of a window is a valid arrival; silence is not	Downstream components run standing operation into a live window: structure changes deploy inside the freeze (F3 here, and the mirror failure in each consuming SOP), event autos cannot legally open at all, and platform rules fire on frozen scope. The consuming component escalates the missing declaration rather than assuming no window exists
I-9	Event autos	P6 and P9 to #15 P6	The armed window with its dates and caps at open, and the close instruction at the window end timestamp	No event auto opens, which is the safe direction. At the other end, a missing close instruction leaves a dated auto live past expiry: SOP-15 F11 detects it on the auto side and F4 here on the window side, and the two detectors exist because one of them will find it first
I-10	Exclusion state to automation	P6 and P9 to #37	The affected campaign list with the window dates at arming and the release at close, same day as each	Platform rules may fire on frozen scope: bids and budgets move inside the window with no plan reference, the changes are unattributed on the #38 T+1 diff, and the read damage is unbounded because the rule kept firing. This SOP reports the state; it never edits a rule
I-11	Staged sets	P3, P6, P8, P9 to #38	Arming, cap, allowance-raise and revert rows in #38 staging format with the event reference and plan version in every reason string, deployed same day at arming and close	The plan exists and the account does not run it: caps are not live, the freeze check has nothing to check against, and the window runs on baseline budgets into event traffic. A revert set that does not deploy is F7 with its days quantified
I-12	Verification diff	#38 to P6 and P9	The T+1 diff against the staged values, read line by line against the P6 baseline snapshot at close	Landed values are unverified, so the window's own record cannot be trusted and the P11 score is computed on assumed rather than confirmed deployment. The score records the gap rather than reporting a clean run

I-13	Placement premium and its normalization	P4 to #28, and back at P9	The ceiling worksheet for any tested premium before it deploys, and the dated normalization on the revert checklist at close	Without the worksheet no premium is tested, which is the correct default. Without the normalization the premium lives past the window as rented real estate nobody re-asks about, which is F7 here and the same failure class on #28's side. Cited per Ledger row 28
I-14	Contested-root declaration	#33 to P2 step 4	The family ownership declaration or the guard in force while undeclared, read before the affected set is finalized	Both siblings hold at cap and the G11 portfolio call is requested; the window does not adjudicate ownership by spend. An undeclared root amplified on both siblings bids the family against itself with the family's own money at the highest CPC of the quarter. Cited per Ledger row 33
I-15	Post-settle recovery handoff	P11 to #24	The P6 baseline snapshot, the window record, the settle dates and the post-settle position, handed over at settle expiry where position sits below baseline	The loss goes undiagnosed and the position drifts while the cause is unnamed; the hangover cause class has an owner precisely so this does not become a PPC bid problem. Where the handoff is refused for lack of a snapshot, that is F6 arriving late. Cited per Ledger row 24
I-16	Exclusion dates to test design	P10 to #35	The published exclusion range covering window days plus 14, at arming and confirmed at close	An incrementality test reads a window as its treatment effect and returns a false result that then justifies a ladder step or a restore. #35 voids or defers on receipt; without the dates it has nothing to void against
I-17	Liquidation tier declaration	LTSF system via #27 to P12	The LTSF row carrying salvage-parity price, the alternative per-unit value, remaining units and the cliff date, at the tag	The window cannot be planned as a liquidation window and PPC computes nothing it can pull from that system. The plan stops at the boundary and the entry returns to the LTSF owner; the LTSF Master document is outside this corpus per conflict C6
I-18	Scored event record and class factor	P11 to #23, #28, the deals owner and Account Rollup A3	The five-line event record, the V5 score with its limb detail, and the recalibrated Event Lift Index, at settle expiry	The next window of this class is priced from the same assumption that just missed, and the class factor never converges: the playbook stops compounding and every window is a first window. #23's next worksheet carries stale velocity and CVR assumptions for the amplified terms

Appendix A. v4.1 Amendment Items

No number below appears as an operating value in this SOP. Each item names the gap, its provenance, the operating path used until v4.1 lands, and the proposal carried to the threshold owner.

Item	Gap, provenance, operating path and proposal
AM-1	The in-window freeze has no v4.0 gate. Highest priority in this register. The freeze on structure changes, main-image changes and un-named modifier moves inside an armed window is enforced across the corpus and attributed to "G12" by SOP-32 v1.0, SOP-12 v2.0 Section 3 and its P9, SOP-15 v2.0, SOP-23 v1.0 and the Transition Brief v1.2. v4.0's G12 is MKL hygiene and carries no event clause; v4.0's G10 excludes window data from trend verdicts and carries no freeze clause. This is the event-side instance of the identifier divergence registered as conflict C9. Operating path: the freeze is declared as a plan state at P6 and enforced by the components holding the levers, unchanged in practice. Proposal: admit the armed-window freeze as a named gate clause with its four scope classes stated, and reconcile the corpus alias in the same release.
AM-2	Tentpole intake lead time. v4.0 Family F2 carries the 14-day PRE-EVENT window and nothing longer. The program's registered cadence carries event plans at T-14 and T-30, the longer figure for tentpoles. Operating path: the tentpole lead time lives on the calendar entry and in the program cadence, not in this SOP's body. Proposal: admit a class-conditional intake lead time to the F2 event-posture row.
AM-3	Per-class settle lengths. Gate G10 and Family F2 carry a single 14-day post-event exclusion for every class. v1.0 ran a per-class settle and its worked case used 7 days on a 7-day deal, which is shorter than the gate. Operating path: G10's 14 days governs every class in this version, including where v1.0 ran shorter, because the gate is the threshold authority. Proposal: either admit per-class settle lengths with their evidence or state explicitly that G10's 14 days is universal so the plan template stops carrying a field the gate overrides.
AM-4	First-window category default lift factors. Rule M9 directs first-time events to borrow the category default and log the variance; v4.0 carries no category-default table. The program's open-decisions register carries this as item PPC-DR-35 with tentpole factors recorded as having no basis and a proposed interim treatment awaiting ratification. Operating path: no factor is written; the plan runs at baseline with a named dollar allowance as the only expansion authority, bounded by the S-E2 15% correction and the gate G7 daily event cadence. WE-2 quantifies the exposure this creates at \$1,728.00 of contribution on one day of one event. Proposal: ratify PPC-DR-35 and admit the resulting table to M9.
AM-5	The in-window allowance. The plan's named cap-raise allowance has no v4.0 basis. Rule L4's single-step 20-30% budget expansion is a performance-and-utilization rule for standing operation and does not address event authority. Operating path: the allowance is a plan value, written at P2 or P5, confirmed by both D-holders before any raise deploys, and bounded by the P8 worksheet check against window contribution. Proposal: admit an event-allowance rule, or state that L4 governs inside windows and bound it explicitly.
AM-6	The plan stop condition. v4.0 carries S-E2's 15% pacing-deviation correction trigger but no stop threshold: there is no rule stating at what point a window stops rather than corrects. Operating path: the stop condition is a plan value written at P2 step 7 with a named owner, expressed in numbers on both the spend and the units limb. Proposal: admit a stop-condition threshold expressed as a multiple of the projection on each limb.
AM-7	Tentpole in-window cadence. Rule V4 carries daily pacing and gate checks during pushes and events, and gate G7 carries daily event pacing. The twice-daily tentpole block that v1.0 ran and this version retains has no v4.0 basis. Operating path: the class carries its cadence in the plan. Proposal: admit an event-class cadence to V4, since a tentpole day compresses a normal week of spend into hours.
AM-8	Halo scope definition. No v4.0 rule defines which family and sibling campaigns count as halo for an event, yet halo spend rides the caps and is scored on the prediction. Operating path: the plan states its own membership test at P2 step 2 and records it. Proposal: admit a halo membership rule, or assign the boundary to #33 in the Ledger so one document owns it.
AM-9	Annual calendar cadence and build date. Rule T4 and scenario S-E6 govern the timing rule (pushes time into a rising Seasonal SV Index, never after the peak) but carry no calendar cadence and no build date. The scheduled fold-in registered against this component asks for the annual calendar dated before the fourth quarter opens. Operating path: P13 step 4 states the rule and the program cadence carries the date. Proposal: admit the annual calendar to the V4 cadence rule or to component #40 when it lands.
AM-10	Event-class taxonomy. v4.0 carries no closed list of deal classes, and the Master PPC Data Architecture v1.0 Section 9 T5 closed event list carries "Deal window" as a single value, so the class cannot be recorded in the instrument and cannot be filtered on. Operating path: the five classes are declared in the plan header and carried forward from v1.0 as institutional record. Proposal: admit the class list and add it as a second column on the T5 event row so class-level scoring becomes a filter rather than a reconstruction.
AM-11	The basis for window ceilings. Scenario S-A9 requires ceilings to recompute immediately from a new margin table and gate G9 marks price-changed reads PROVISIONAL, but no rule states whether a temporary deal price governs the window's ceilings or the baseline economics do. The two answers differ materially: on the WE-1 figures the window ceiling is \$1.12 against a baseline of \$1.46, so the same bid is legal on one basis and not the other. Operating path: P1 step 5 computes both and the plan carries both, with the window basis governing in-window decisions and the baseline governing the halo block. Proposal: state the basis explicitly in the M1 chain.
AM-12	Deal-fee treatment. The calendar entry carries a deal fee and no v4.0 rule states whether it enters the event profit-and-loss, the contribution used for the window loss line, or neither. Operating path: the fee is reported as its own line in the P11 event record and is not folded into contribution, so no ceiling in this SOP is computed on a number that includes it. Proposal: state the fee treatment in the M1 chain and in the M9 event-sizing rule so the window's true cost per unit is one number rather than two.