

SOP-34: The Canada Marketplace Annex

PPC DEPARTMENT · OPERATING SOP · COMPONENT #34 OF THE OPERATIONS DEFINITION REGISTER

Version 2.0 · 2026-08-11 · Owner: the CA operator under the PPC Lead · Supersedes SOP-34 v1.0-structure (2026-07-24) in full · Rebuilt to the SOP Build Kit v1.0 twelve-element standard · Thresholds live in the Master PPC Decision Criteria System v4.0 only; this annex carries zero local thresholds and logs 15 ungoverned values as Appendix A amendment items · Absorbs the scheduled marketplace-expansion trigger clause registered against this component in the SOP Program Gap Register v1.0, Section 4

1. Verdict and Scope

VERDICT: Canada runs the same decision system on CAD-native economics at a weekly cadence, and this annex is complete when every CA parameter carries a definition, a source, a derivation, a scaling relation and a named consumer, not when its value is populated. The logic transfers whole; every number re-parameterizes; and the single structural finding of this rebuild is that the Master PPC Decision Criteria System v4.0 carries no marketplace dimension at all, so a CA value is not a number waiting in a column that exists, it is a column that does not yet exist. Version 2.0 expands v1.0-structure from 0 step-level procedures to 13, from a 6-row PENDING-CA table to a 14-row parameter register in which each row states source, derivation, scaling relation and consumers, from 0 worked examples to 4, from 4 failure modes to 12, and adds the operator ramp, the 16 interface contracts and the marketplace-expansion procedure the v1.0 standard never required.

1.1 The Unresolved Dependency and What It Prevents

This component is held on dependency D7 and conflict C7: the CA account packet, 5 items per v1.0 Section 7, has not landed. This build proceeds to the boundary rather than stopping, on the component owner's scope direction that the annex is a process document whose job is to establish the mapping and decision structures rather than to carry a snapshot of CA data that would be stale on arrival. That direction supersedes the Cross-Reference Ledger's C7 resolution ("hold SOP-34 out of the build queue") for structure, and does not supersede it for values.

No CA value is written anywhere in this document. Every number in this document is either quoted from v4.0 with its section heading, computed in a worked example from the standard fictional training base, or listed in Appendix A as an amendment item.

The specific decisions the missing packet prevents, and what each waits on:

Decision this annex cannot make	Packet item it waits on	What runs in the meantime
Any CA ceiling: break-even ACOS, Target, Max Acceptable, Max Profitable CPC per product	Item 1, per-product CAD economics	P3 builds the spine structure and the named cells; the CA workbook computes nothing until the cells hold CAD-native values. No CA campaign runs before P3 closes (F5).
The CA volume floor register: the numeric floors that replace v4.0 Section 6, rule O2's search-volume floors	Item 2, CA keyword export with volumes and ranks	P4 fixes the derivation (the measured CA-to-US demand ratio per syntax family) and the review cadence. WE-1 computes the method on illustrative volumes so the arithmetic is proven before the data lands.
The CA campaign registry state and the consolidation plan	Item 3, CA campaign registry as-is	P5 fixes the structure rule and the promotion test; the as-is registry decides only how much consolidation work exists, never whether the rule applies.
CA conquest target lists and the CA competitor depth read	Item 4, CA competitor set with price and rating states	Flag-obvious cases only per Section 4.4; the block records that the mechanical check could not run, per gate G8.
CA event class factors and the pre-peak calendar	Item 5, 12 months of CA seasonality by family	P11 runs CA windows at the CA baseline with a named allowance and no borrowed factor, on the same first-window path SOP-32 uses.
The CA mature band, the CA code-red trip and the CA TACOS stage bands	Items 1 and 2 together, then a v4.1 release	Nothing. These are the three CA values with no operating substitute: the code-red protocol at #26 P3 trips on a multiple of a band that does not exist in Canada. Appendix A items AM-6 and AM-7, and the only rows in this annex that block rather than degrade.

1.2 Why the Parameter and Not the Value Is the Deliverable

A CA keyword list, a CA competitor set and a CA volume table all change continuously, so a document gated on their current values is stale the week after it publishes and useless in the week it is needed. What does not change is the specification: what the parameter is, which source supplies it, how it is derived from that source, how it stands relative to the US equivalent, and which decisions consume it. A parameter specified to those five points is executable the day its value arrives, by an operator who was not present when it was specified. That is the test this document is built to pass, and P2 is where it is applied row by row.

The inverse failure is the one the corpus already carries. v1.0 Section 3 marked six parameter classes PENDING-CA with a one-line rule each and no source, no derivation and no consumer list, which is why v1.0 could be read in five minutes and executed in none of them. A blank with a label is not a specification; it is a deferral wearing one.

1.3 What This Annex Owns, and the Standing Transfer Rule

This annex owns the Canada delta and nothing else. The standing rule, unchanged from v1.0 and restated because everything below depends on it: all operating SOPs 12 through 38 apply in Canada exactly as written unless a row in this document explicitly overrides them. An override is a row in Section 2, 3, 4 or 5 of this annex naming the US rule and the CA treatment. Silence is not an override, and an operator who cannot find an override row executes the US procedure. Marketplace splits are a doctrine requirement rather than a convenience: the PPC Strategic Doctrine v1.0, Section 5.3 "Multiple phrase and broad campaigns" names marketplace splits (US against CA) as always separate campaigns, so no CA campaign shares a structure, a budget or a data pool with a US campaign.

1.4 Boundaries With Adjacent Components

Every excluded decision names the document that owns it. Boundaries against components whose procedure numbering may move at their own rebuild cite the Cross-Reference Ledger row rather than a procedure number.

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, gate, scenario, lever limit, sizing rule and verdict string, in every marketplace. This annex executes; it never decides. v4.0 currently carries no marketplace dimension, so every CA-specific value is an Appendix A item and never a number written here.
CA pricing and strikethrough owner	Decides the CA price, the strikethrough state and whether a CA promotion runs. This annex reads the decision and executes its PPC consequences; it never re-litigates a price, and it never converts a US price into a CA one.
CA inventory and stock owner	Decides CA cover, restock timing and the CA backup chain. This annex executes gate G6 responses on what that owner publishes.
#12 Exact campaign operations	Owns exact structure, creation, the weekly operating pass and taper execution. In Canada the same procedures run against the CA workbook at the CA cadence; the only CA delta is which terms qualify for a dedicated campaign (Section 4.5) and how often the pass runs (Section 4.6).
#13 Phrase and #14 Broad operations	Own discovery structure and the weekly search-term triage machinery. The CA search-term pass at P9 executes SOP-14's three-route triage unchanged; what changes is the window it reads and the accumulation rule, never the routes.
#15 Auto operations	Owns auto structure and the D5 targeting-group split end to end, including in Canada; its Ledger row defers the Canada delta to this annex, and this annex's delta is cadence and volume qualification only.
#16 PAT and conquest	Owns candidate intake, target list construction and the open, rotate and exit calls. In Canada the full machinery activates when CA volume clears the scaled criteria; until then only flag-obvious cases run (Section 4.4).
#23 Ranking push operations	Owns push sizing, the M2 worksheet, funding order, standup, stop-loss and taper release. A CA push uses the same worksheet with CA DSTR, CA velocity and CA Expected CPC references; no CA push is sized on a US reference (F5).
#26 Mature, defend and code red	Owns floor operations, the quarterly re-ladder and the code-red protocol at its P3. The protocol transfers unchanged and trips on the v4.0 multiple of the mature band; the CA band does not exist, which is AM-7 and the reason CA code red is currently declared on the PPC Lead's call rather than mechanically.
#27 Terminal and LTSF-tier rules	Owns the liquidation objective switch and the R-series rules, computed against the CA fee schedule and CA salvage values. The LTSF Master Decision and Operations System is outside this corpus per conflict C6; this annex stops at that boundary and infers no tier rule.
#28 Placement and modifier operations	Owns the ceiling worksheet, ladder rungs, exit-bid discovery and modifier sizing. Ladder mechanics transfer unchanged; the CA ladder prices against CA blended CPC and the CA ceiling, never a converted US one. Cited per Ledger row 28.

Adjacent component	The seam, stated so it is owned
#29 Negation and #30 Harvest	Own wall logic and the atomic graduation chain. Both transfer unchanged. The graduation bar itself is v4.0's and does not scale (Section 4.9); what changes in Canada is the window over which the bar is reached.
#32 Deal and event window operations	Owns the event plan, cap construction, arming, the in-window block, revert, settle and scoring, in every marketplace. The CA deals calendar arrives on its own interface with its own class list; this annex supplies the CA calendar seam and the CA class factors, and #32 decides what the window does.
#33 Sibling and family arbitration	Owns the ownership declaration and guard enforcement on contested roots. Family boards are per marketplace: a CA sibling contest is arbitrated on CA composites and never on the US declaration for the same root. Cited per Ledger row 33.
#35 Incrementality testing	Owns test design, the register and the call. CA test design waits on CA volumes clearing the read minimums; this annex publishes the CA volume state and #35 decides whether a CA test is possible.
#37 Automation platform rules	Owns the rule register and change control. CA campaigns enter the register as their own scope with their own values; this annex never edits a rule and never lets a US rule address a CA campaign (F6).
#38 Bulk-file deployment	Owns staging format, upload, the processing report and the T+1 verification diff. Every CA change stages in #38 format with the marketplace key in the reason string, and this annex consumes the verification diff.
#31 Budget and envelope	Envelope law lives in the instrument layer (T0 zone H, Account Rollup A2) per open conflict C3; this annex never cites a #31 procedure number. The CA envelope is its own envelope and never a share of the US one.

1.5 Why This Document Is Longer Than the Delta It Describes

Most of Canada is the US system, and a document that said only that would be one page and would be executed by nobody. The length here is not the delta; it is the specification of the delta's parameters to the five points named at Section 1.2, plus the weekly block written as procedures a person who has never run it can execute. Every section that restates a US procedure has been cut to a reference; where this document runs long it is because it is specifying a parameter, arguing the scaled-floor case that is the annex's reason to exist, or stating a non-arrival consequence the corpus has never carried.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. Lags use the cadence terms v4.0 uses at Section 14 "Layer 9: Verification and Learning (V1-V8)", rule V4 (daily, weekly, monthly, quarterly), read against the Monday cascade defined in Master PPC Data Architecture v1.0, Section 12 "Data Flow: Feed to Columns to Cadence": the Raw Pull Pack lands Monday 06:00, workbooks refresh by 09:00, the audit runs Monday, deployment is Tuesday after the BM huddle. The CA cadence is weekly-only by design and the reasoning is at Section 4.2; the daily limb of V4 does not fire in Canada except where a gate response is same-day by its own rule.

Trigger	Detecting artifact	Max lag	Procedure
The CA account packet lands, in whole or in part	Delivery from the Head of Amazon Operations against the 5-item list at Section 4.12; each item is received against its own row in the parameter register	Same week; the conversion runs in the next Monday block	P12 packet-landing conversion, then P1 to P5 in order for whatever the arriving items unblock
CA workbook not yet instantiated on a product with live CA spend	No CA Command Workbook against a product carrying CA campaigns in the CA registry	Same week the spend is found	P1 instantiation, and the spend is recorded as ungoverned in the block until P1 closes
A CA parameter row has a value but no source, derivation or consumer	The P2 parameter register's own completeness column	Same session it is found	P2 step 6 reconciliation; the value is suspended from use until the row closes (F7)
Weekly CA block due	The Monday cascade slot after the CA feeds refresh	Same day; a block run late is run late and logged, never skipped (F3)	P6 to P10 in order, one session
Two consecutive CA blocks missed	No CA session record for two consecutive weeks in the CA T6	Same day the second gap is found	Escalation to the PPC Lead per Section 7, F3; the block runs before any CA deployment resumes

Trigger	Detecting artifact	Max lag	Procedure
CA price or strikethrough change	The CA pricing owner's published change, mirrored to the CA T5 Product Log as event type "Price move"	Same day	P8 step 2; gate G9 marks CA CVR and ACOS reads PROVISIONAL for the trailing 7 days per v4.0 Section 5, gate G9
CA cover falls or a CA variation goes out of stock	The CA inventory owner's cover read into CA T2 K2 "Inventory Band" and "Days to Stockout"	Same day, because gate G6's RED response is same-day by its own rule	P8 step 3; G6 responses execute at CA scope
CA campaign truncated	CA T3 C3 "In-Budget Bucket" reading below the v4.0 gate G3 line	Inside the block that found it	P7 pacing and truncation sweep
A CA term crosses its scaled volume floor in either direction	The CA volume refresh against the CA floor register built at P4	Inside the next Monday block	P4 step 5 re-qualification, then P5 where the structure class changes
A CA syntax family's measured demand ratio moves past its recorded review band	The quarterly ratio re-measure at P4 step 6	Inside the quarter's first block after the re-measure	P4 full re-derivation; every floor moves together or none does
A confirmed CA deal or event window	The CA deals calendar entry with class, dates and CA deal price	Same day the entry is confirmed, and not later than the class lead time	P11 CA event handling, which hands the plan to #32
CA verification return	The #38 T+1 verification diff flagging a landed-value mismatch on a CA row	Same day	Redeploy the delta per #38; F9 response where the row is a revert or a gate-driven row
A v4.1 release carrying any CA parameter	The Criteria System increment notice	Before the next CA block runs	P2 step 7 retirement of the affected Appendix A item, then Section 10 revalidation
A new marketplace is proposed	An expansion decision from the CEO or the Head of Amazon Operations	Before any spend in that marketplace	P13 marketplace instantiation; expansion never starts undefined

3. Preconditions and Gates

Every CA procedure runs only after these read current in the CA T7 Manifest (Master PPC Data Architecture v1.0, Section 9 "The Infrastructure Tabs: Logs and the Manifest (T5-T7)"). A missing or stale field group suppresses its dependent execution class rather than degrading it silently, per v4.0 Section 5, gate G8; the block records what it could not see and proceeds on what it can. The CA manifest is its own manifest: a PRESENT status on the US manifest says nothing about Canada, and reading one for the other is the mechanism behind F6.

Precondition (CA T7 manifest field group)	Requirement and the observable that proves it passed
The CA parameter register (this annex, P2)	Every row carries definition, source, derivation, scaling relation, consumers and a status. Observable: the register's completeness column reads complete on every row whose consumer is exercised this cycle. A parameter with a value and no source is suspended from use the session it is found; a parameter with a source and no value suppresses only the decisions listed in its consumers cell.
B8 and D1 economics chain, CAD-native	CA T0 zone C named cells live in CAD (AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC, WeeklyEnvelope) with their CA margin-table version. Observable: CA T0 zone C shows the CA margin-table version and the gate G9 state, and the cells' currency label reads CAD. No CA ceiling computes from a converted US value, ever (F2).
A2 CA deployment state	Current bulk state in-row for every CA campaign: state, daily budget, bid, bidding strategy, placement multipliers. Observable: CA T3 C2 cells populated with the current CA pull date. Without it there is no baseline for any CA change and no diff for #38 to verify against.
A3 CA pacing and in-budget	In-budget percentage and bucket per CA campaign. Observable: CA T3 C3 "In-Budget Bucket" and "Pacing Deviation Flag" resolved. The Command Workbook template records A3 as PARTIAL with the day-of-week curve pending, and at CA volume a day-of-week curve needs more weeks to stabilise than in the US, so while it is absent the pacing read runs against the CA weekly envelope only and the block records the substitution.
A7 CA eligibility	Every advertised CA SKU eligible with the CA offer live. Observable: CA T3 C2 "Eligibility + Reason" reads live. Gate G4 halts every lever on a target the day its ad reads ineligible or suppressed and opens the listing fix task the same day.

Precondition (CA T7 manifest field group)	Requirement and the observable that proves it passed
CA inventory cover (G6 band and days to stockout)	Cover and band per CA variation from the CA inventory owner, with the CA backup chain named. Observable: CA T2 K2 "Inventory Band" and "Days to Stockout" resolved on every CA-advertised variation. Gate G6 governs unchanged: GREEN scales, YELLOW freezes push spend at current level, RED tapers to defense and re-points to the declared backup the same day.
The CA volume feed (packet item 2 and its refresh)	CA search volume per term per syntax family, current inside the refresh window the CA floor register names. Observable: CA T2 K1 "SV" populated with the CA source and pull date. While it is MISSING, no term is qualified or de-qualified against a floor and the block records the whole family as unqualified rather than defaulting to the US floor (F1).
The CA MKL and syntax taxonomy (G12)	Every CA term exists in the CA Master Keyword List with a syntax tag and a relevancy score. Observable: the CA MKL refresh date inside the window v4.0 Section 5, gate G12 sets. Taxonomy tags are shared with the US taxonomy by contract; divergent taxonomies across marketplaces would make the demand ratio at P4 uncomputable, because the two sides would not be counting the same family.
CA competitor depth (D8 equivalent)	CA competitor price, rating, review count and stockout flags for CA conquest candidates. Observable: CA T2 K8 populated inside the trailing week. The manifest carries D8 as MISSING in the template and no CA equivalent exists at all before packet item 4, so the conquest machinery stays at flag-obvious cases and the block records the blind spot; the finding downgrades to a flag per gate G8 rather than to silence.
CA rank series	CA organic rank per CA priority term. Observable: CA T2 K4 "Organic Rank" carries a CA source and a pull date. The US daily rank series (D9) is a US feed and says nothing about CA rank; a CA push read on a US rank series is F5.
The marketplace key	Every CA packet row, registry row and staged bulk row carries a marketplace value. Observable: the key resolves on the CA T3 registry and on every #38 staged row's reason string. The corpus carries this as open decision PPC-DR-20 with the recommendation to tag every packet and registry key from day one so later inclusion is non-breaking; without the key CA spend mixes into US totals silently, which is F6.

Gates in force, by v4.0 name and number. All twelve gates at v4.0 Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass" apply in Canada unchanged, because a gate is a procedure and procedures transfer. What the CA operator must hold separately is which gates carry an absolute bar and which carry a demand-relative one, since only the second class re-parameterizes:

- Absolute, and identical in Canada. G1 validation quarantine, G2 sample sufficiency, G3 budget truncation, G4 eligibility, G5 indexing and listing placement, G8 packet completeness, G9 price stability, G10 event exclusion, G11 sibling conflict, G12 MKL hygiene. G2 is the important case and the one most often got wrong: its click and impression bars are statistical, not commercial. A conversion-rate estimate on 40 clicks has the same width in Toronto as in Denver, so the bar does not scale down with the marketplace. What changes at CA volume is the time it takes to reach the bar, which is why the CA read runs on the 60-day window as standard rather than as a fallback (Section 4.9).
- Demand-relative, and re-parameterized in Canada. G6 inventory adequacy, whose band inputs are CA cover against CA velocity, and G7 envelope and pacing, whose envelope is the CA envelope. Neither gate's rule changes; both consume CA inputs and would be meaningless on US ones.
- The floors that are not gates. The search-volume floors at v4.0 Section 6, rule O2 (above 500 for Ranking, above 1000 for Market Share) are objective-assignment criteria rather than gates, and they are the single most consequential US-scaled value in the system for Canada. They are demand-relative by construction and they are the subject of P4 and WE-1. v4.0 states them as absolute numbers with no marketplace conditioning, which is Appendix A item AM-3.

The oversight-class problem, stated once. v4.0 Section 12 "Layer 7: Precedence and Conflict Resolution (P1-P10)", rule P5 places every action on a keyword above 500 search volume in the human-mandatory review tier. At CA volumes almost no term reaches that figure, so the tier that exists to force human review would empty out precisely in the marketplace with the thinnest instrumentation and the least history. This annex therefore runs the whole CA block in the human-mandatory tier: every CA verdict is confirmed at the weekly sign-off the Authority Matrix v2.0 "Canada weekly block" row requires. That is stricter than P5 and never looser, which is the only safe direction to resolve an under-specified tier. The SV-keyed tier definition is Appendix A item AM-13.

4. Step-Level Procedures

Time estimates are planning figures for scheduling the CA operator's week. They are not thresholds and nothing in v4.0 governs them. P1 to P5 are the standing-up sequence, run once per product and re-run when a parameter moves. P6 to P10 are the weekly block, run every Monday as one session. P11 to P13 are the conditional procedures.

4.1 P1. Instantiate the CA Command Workbook (45-70 min per product; planning estimate)

1. Clone the template. Copy PPC Command Workbook TEMPLATE v1.0 and name it per the template's own CLONE PROCEDURE step 1, with the marketplace token in the code position so the CA file is never confused with its US sibling on a file list. The clone carries all ten tabs: T0 Product Command, T1 Syntax Board, T2 Keyword Command, T3 Campaign Board, T4 Placement and Bid, T5 Product Log, T6 Decision Log, T7 Manifest and REFERENCE.
2. Set the marketplace key. Write the marketplace value into the workbook header and into the registry key column on T3, so every row this file produces is attributable to Canada downstream. This is the step that prevents F6, and it is the first step rather than the last because a row created without the key is found later by reconciliation rather than by construction.
3. Leave the economics spine empty and say so. The template's CLONE PROCEDURE step 2 sets T0 zone C named cells (AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC, WeeklyEnvelope) from the product packet. Until CA packet item 1 lands there is no CA packet to set them from, so the cells stay empty, the workbook's gate strip at T0 zone D reads the G9 state as blocked on the CA margin table, and no ceiling in the file computes. Do not seed the cells with converted US values to make the formulas resolve; a resolved formula on a converted input is F2 with the evidence hidden.
4. Load what exists. Data Operations loads the CA MKL into T2 and the CA campaign registry into T3, deleting the template's example rows per CLONE PROCEDURE step 3. Where packet items 2 and 3 have not landed, T2 and T3 stay empty and the T7 manifest names both as MISSING with the verdict classes each blocks.
5. Build the CA manifest honestly. T7 is populated from the CA feed reality, not copied from the US manifest. Every CA field group gets its own status, its own source feed and its own "Verdict Classes Blocked While Missing" cell. The CA completeness score is computed on the CA manifest and stamped on the CA block header. A CA workbook reporting the US completeness score is a false readiness signal and the most expensive single defect this procedure can produce.
6. Register the file. The CA workbook is recorded against the product in the Account Rollup Workbook A0 row set as its own row, not folded into the product's US row. Family boards at A1 are per marketplace for the same reason: a CA sibling grid built from US CPCs compares two auctions that never met.

4.2 P2. Build and Maintain the CA Parameter Register (2-3 hours for the first build, 15-25 min per weekly reconciliation; planning estimate)

This is the annex's central instrument and the procedure the rest of the document depends on. The register is one row per CA parameter, and a row is complete when all six columns are written, whether or not the value column holds a number.

1. Open the register with the six columns fixed: parameter, definition, source, derivation, scaling relation to the US equivalent, consumers, and status. The value is a seventh column and it is the only one allowed to read pending.
2. Write the definition in the terms the consuming rule uses. If the consumer is v4.0 rule M2 clicks-to-rank sizing, the definition is written as M2 reads it, so the substitution at use is mechanical rather than interpretive.
3. Name the source to the artifact. Not "the CA packet" but the packet item number, the feed and the tab or column it lands in. A source that names a document rather than a field is not a source; it is a hope.
4. Write the derivation as arithmetic. The derivation states how the value is computed from the source, in an expression an operator can execute. Where the derivation is a measurement rather than a formula, it names the measurement window and the refresh cadence.
5. Write the scaling relation. Three classes only, and the class determines the treatment for the life of the parameter. Class T (transfers): the US value applies unchanged in Canada because the quantity is a ratio, a bar or a procedure constant that carries no marketplace content. Class N (native): the value is computed natively from CA inputs and bears no relation to the US value at all. Class R (ratio-derived): the value is the US value scaled by a measured CA-to-US relation, and the row names which relation and where it is measured. Assigning a class is the substantive act; a parameter with no class is unusable because nobody knows what to do with it when the US number changes.
6. List the consumers by rule ID and procedure. Every decision that reads this parameter. This is the cell that makes a pending value survivable: it converts "we do not have that number" into a named list of decisions that are suspended, which is the difference between a stated blind spot and a silent one.
7. Reconcile weekly and retire on v4.1. The weekly block reads the register before it reads anything else and confirms three things: no row has a value without a source, no row's consumers include a decision the block made this week while the row was pending, and every row whose value arrived since the last block has been through step 6. When a v4.1 release lands a CA parameter, the register row's status flips to governed, the Appendix A item retires in the same increment, and Section 10 re-validation runs.

The register at v2.0 publication, with the class assigned on every row and the value pending on every row that has one:

Parameter	Class	Source	Derivation and scaling relation	Consumers (suspended while pending)
CA per-product economics spine	N	Packet item 1, into CA T0 zone C named cells	Computed natively in CAD: CA AOV minus CA landed cost minus CA referral fee minus CA fulfillment fee minus CA returns provision. No relation to the US value; the US spine is not an input. Refresh monthly with the CA margin-table version stamped.	v4.0 rule M1 break-even chain; every ceiling; M2 push budgets; M4 effective-CPC ceilings; #26 floor economics; #27 CA terminal math
CA break-even, Target and Max Acceptable ACOS	T	Derived in-workbook from the row above	Rule M1 transfers whole: break-even is margin divided by AOV, Target is 50% of break-even, Max Acceptable is 75%. The ratios are procedure constants and carry no marketplace content; only their inputs are CA.	Every CA verdict on an ACOS read; the CA ladder; #26 CA floors
CA Expected CPC reference	N	Packet items 2 and 3, then the CA auction as observed in the CA bulk	CA auction levels per category, measured from CA delivered CPC by syntax family. The US reference is never a proxy and never a starting point: the two auctions have different depth, different seasonality and different currency.	v4.0 rule L2 bid references; M2 required budget; M4 multiplier sizing; ladder seeds at #28; O3 realism
CA search-volume floors (the floor register)	R	Packet item 2, refreshed on the CA volume feed	The v4.0 Section 6, rule O2 floors scaled by the measured CA-to-US demand ratio per syntax family, computed at P4 as CA family SV divided by US family SV on the matched term set. Re-measured quarterly; every floor in a family moves together.	O2 objective assignment; O3 realism; the oversight-class filter; P5 structure class; #23 push candidacy; #16 conquest candidacy
CA demand ratio per syntax family	N	P4 measurement on packet item 2 against the US T2 for the matched family	Measured, not assumed. The ratio is the annex's one derived quantity and it is recomputed rather than inherited; v1.0's characterisation of CA volume as roughly an order of magnitude under the US is institutional record, not a divisor.	The floor register; P5 structure class; CA pool minimums; #35 CA test feasibility
CA sample-sufficiency bars	T	v4.0 Section 5, gate G2, unchanged	Transfers with no scaling. The bars are statistical, so the confidence interval on a rate does not widen or narrow with the marketplace. What CA volume changes is the window needed to reach them, computed at Section 4.9.	Every CA rate verdict; the read-limit flag; #35 test windows
CA DSTR and rank velocity references	N	Packet item 2 ranks, then the CA rank series	CA daily sales-to-rank at CA volumes, measured on CA competitors. A US DSTR applied to a CA push over-sizes the push by the ratio of the two markets and prices the whole worksheet wrong.	Rule M2 clicks-to-rank; #23 push sizing, funding and ETA; #24 CA recovery sizing
CA TACOS stage bands	N	Packet item 1 plus the CA stage declaration; awaiting v4.1	CA stage bands computed from CA contribution and CA stage posture. No derivation from the US bands is defensible, because the bands encode a market's paid-share structure rather than a product's arithmetic.	Rule T2 stage gates; T7 stage renewal; rule R1 Recovery trigger; #25 and #26 stage postures

Parameter	Class	Source	Derivation and scaling relation	Consumers (suspended while pending)
CA mature band and code-red trip multiple	N	The CA TACOS band row above; awaiting v4.1	The trip is a multiple of the CA mature band, and the multiple itself is v4.0's and transfers. Without the band the multiple has nothing to multiply.	#26 P3 code-red protocol; the CA T0 zone A verdict banner; Account Rollup A0 code-red flag
CA structure and format targets	R	Packet item 3 plus the CA registry as it develops	Rule M6's splits are destinations rather than quotas; in Canada the exact and consolidated split moves toward tail-consolidation at the rate CA volume qualifies terms, so the target is stated as a direction with a review date rather than a band until CA volume supports one.	Rule M6 structure review; P5 structure class; the monthly CA structure finding
CA event class factors (Event Lift Index)	N	Packet item 5, then CA window history	Measured per CA event class from CA windows only. The Boxing Day class has no US counterpart, and US Prime windows differ in CA in both dates and depth, so a US factor is not a first-window default here.	Rule M9 event sizing; #32 cap construction and scoring; P11
CA competitor set and depth	N	Packet item 4, refreshed weekly once the CA feed exists	CA competitors on CA listings at CA prices. Overlap with the US set is coincidental and is never assumed.	O3 credibility check; #16 CA conquest; gate G6 context reads; #26 CA incursion detection
CA envelope	N	The CA stage band and CA sales, into CA T0 zone H	Computed from the CA stage TACOS ceiling on CA revenue. The CA envelope is never a share of the US envelope, and CA pushes compete only against CA pushes in the CA M5 queue.	Gate G7; rule P2 precedence; rule M5 allocation; every CA budget verdict
CA read minimums for SD pools and AMC	R	The demand ratio row plus the CA registry	Pool and instance minimums scale with the demand ratio because they are volume-qualification bars rather than statistical ones. Re-checked semi-annually per Section 4.4.	SD audience build; #20 CA AMC instance decision; #35 CA test feasibility

Nine of these fourteen rows carry class N. That is the finding, not an accident of drafting: most of the Canada delta is not a translated US number but a natively computed one, which is why the annex's structure is worth more than the values it is waiting for, and why the FX shortcut is not a shortcut but a different and wrong answer (WE-2).

4.3 P3. Build the CAD-Native Economics Spine (60-90 min per product at packet landing; planning estimate)

1. Take CA inputs only. From CA packet item 1: CA list price, CA landed cost per unit including inbound to CA fulfilment, the CA referral rate for the category, the CA fulfilment fee for the size tier, and the CA returns provision. Every input is a CA figure sourced from a CA document. If an input is only available as a US figure, it is not an input and the row stays pending.
2. Compute the chain in CAD. CA contribution equals CA AOV minus the four cost lines. CA break-even ACOS equals CA contribution divided by CA AOV. CA Target is 50% of break-even and CA Max Acceptable is 75%, per v4.0 Section 10 "Layer 5: Sizing Mathematics (M1-M10)", rule M1. CA Max Profitable CPC equals CA contribution times CA conversion rate, per v4.0 Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)", Family F1 "Product economics". Worked at WE-2.
3. Write the values into the named cells, and nothing else into them. CA T0 zone C carries the CAD values with the currency label and the CA margin-table version. Every derived ceiling in the CA workbook repoints automatically, exactly as the template's clone procedure describes for a US product.
4. Record the exchange rate nowhere in the decision chain. FX belongs in the reporting layer that consolidates CAD results into a group total. It is not an input to any cell that a verdict reads. The whole of F2 is this one step, and WE-2 quantifies what happens when it is skipped.
5. Set the CA margin staleness clock. Gate G9 marks reads PROVISIONAL on a price change; the margin table carries its own staleness horizon per the Economics and Context Pack's margin-table rule. The CA table gets its own clock, because a fresh US table

says nothing about the CA fee schedule, which changes on its own calendar.

6. Publish the spine to its consumers. The CA spine is the input to #23 push worksheets, #28 ladder ceilings, #26 floor economics and #27 CA terminal math. Publishing means the block record names the version and the date, so a consumer can tell which spine its arithmetic was built on.

4.4 P4. Derive the CA Volume Floor Register (2-3 hours for the first build, 30-45 min per quarterly re-measure; planning estimate)

This is the annex's reason to exist and the procedure WE-1 works end to end.

1. Match the term sets. For each syntax family, take the CA term set from packet item 2 and the US term set for the same family from the US T2. Matching is by syntax tag, not by string, because the same family carries different strings in two markets. Terms that exist in only one market stay in the count on their own side; excluding them would bias the ratio toward the shared head.
2. Measure the demand ratio. The family's CA-to-US demand ratio equals total CA family search volume divided by total US family search volume across the matched sets. Record the ratio, the two totals, the term counts and the measurement date. Ratios are per family and never per account: the families do not travel together, and one blended account ratio would over-floor the strong families and under-floor the weak ones.
3. Scale the floors. The CA floor for each objective equals the v4.0 Section 6, rule O2 floor for that objective multiplied by the family's measured ratio. Record both the US figure and the CA figure so the derivation is auditable from the register row alone. Round to the nearest whole search-volume unit and record the rounding.
4. State what did not scale. The sample-sufficiency bars at v4.0 Section 5, gate G2 do not scale and are recorded on the same page as the floors so the two are never confused. This is the mirror-image error of F1 and it is F8: an operator who has just learned that US floors are too high for Canada will reach for the same divisor on the click bars, and that produces verdicts on samples that cannot support them.
5. Qualify the terms. Every CA term is marked against the CA floors: qualified for a Ranking objective, qualified for Market Share, or below both. The qualification drives objective assignment at O2, structure class at P5, and the oversight filter the weekly block reads. A term crossing a floor in either direction between refreshes is re-qualified in the next block and its structure follows at P5.
6. Re-measure quarterly and on a category shift. The ratio is a measurement of two markets and both move. Re-measure at the quarter, and immediately where a CA category event, a CA listing expansion or a US category shift makes the last measurement unrepresentative. When the ratio moves, every floor in that family moves together in one staged change; floors that drift individually stop being a register and become a collection of exceptions.

...

CA term qualification (one term, one CA volume refresh)

```
|
+-- CA volume feed present and inside its refresh window?
| +-- no ..... No qualification runs. Whole family recorded unqualified in the block;
| | US floors are NOT substituted (F1). Feed gap named at gate G8.
| +-- yes -> continue
|
+-- Term carries a CA syntax tag and relevancy score in the CA MKL (gate G12)?
| +-- no ..... Route to CA MKL maintenance; no objective, no structure, no spend.
| +-- yes -> continue
|
+-- CA volume at or above the family's scaled Market Share floor?
| +-- yes ..... Qualified for a Market Share objective, subject to O2's syntax-priority
| | cap: secondary-syntax terms cannot carry Market Share.
| +-- no -> continue
|
+-- CA volume at or above the family's scaled Ranking floor?
| +-- yes ..... Ranking candidate. O3 realism runs next on CA inputs: CA indexing,
| | CA CVR against a CA benchmark, CA competitor credibility, CA
| | inventory GREEN. Realism failure returns it to Discovery, not to a push.
| +-- no -> continue
|
+-- Term converting inside the CA consolidated structure?
+-- yes ..... LEAVE TO HARVEST at CA discovery economics; re-tested each refresh.
```

+- no Discovery or Paused-pending per O2's tail, with the re-entry criterion written as the floor it must cross.

...

4.5 P5. Instantiate CA Structure at CA Volume (2-4 hours for the first build per product; planning estimate)

1. Start consolidated. The CA default is the consolidated-tail-first posture carried forward from v1.0 Section 4, and the reasoning is the Strategic Doctrine v1.0, Section 5.1 two-class doctrine: dedicated single-keyword campaigns buy budget isolation and clean rank attribution, and consolidated syntax-exact campaigns buy data density. At CA volume the second consideration dominates for almost every term, because per-term isolation at tail volume produces campaigns that never accumulate enough conversions to leave learning. The class boundary is the oversight flag rather than preference, and in Canada the oversight flag is set by the P4 floor register.
2. Open a dedicated CA campaign only on a qualified term. A CA term gets its own campaign when it clears the scaled Ranking floor and carries a funded push or a declared distinct role per the Doctrine's Section 5.2 legitimacy test. Everything else lives in a consolidated CA syntax-exact campaign, ad group size per SOP-12's consolidation rules.
3. Never clone the US structure. A 1:1 clone of a US product's campaign set into Canada creates campaign counts the CA volume cannot feed and is F4. The CA registry is built from the CA qualified set upward, not from the US registry downward.
4. Build the walls in the same session. Negation walls per #29 at every CA structure event, exactly as in the US: the graduated CA exact is negative-exact in the CA phrase, broad and auto campaigns the same session. CA walls never reference US campaigns and US walls never reference CA ones.
5. Set the format posture. CA runs SP-heavier until CA-scale asset economics exist for SB and SD, per v1.0 Section 3, and SD audiences stay largely inactive because CA remarketing pools rarely clear their read minimums; both are re-checked semi-annually against the P2 register's read-minimum row. This is a posture with a review date, not a permanent exclusion.
6. Stage everything through #38. Every CA campaign creation, state change, budget and bid stages in #38 format with the marketplace key and the CA plan reference in the reason string, deploys on the CA block's deployment slot, and is verified against the T+1 diff.
7. Promote and demote on the floor register, not on impression. A consolidated CA term that crosses the scaled Ranking floor and passes O3 realism is promoted to its own campaign; a dedicated CA campaign whose term falls below the floor for two consecutive refreshes is demoted back into the consolidated campaign at taper. Both moves are logged with a prediction and scored at their review dates.

4.6 P6. The Weekly CA Block, Part 1: Cascade-Lite (25-40 min; planning estimate)

The CA block is one session per week and it is the contract. The cadence is weekly and not daily because CA volumes accrue decision-grade evidence weekly: a daily read on data that produces a handful of clicks a day manufactures reactions to noise, and v4.0 Section 5, gate G2 already says that swings on thin samples are noise by definition. The daily limb of rule V4 fires in Canada only where a gate response is same-day by its own rule, which is G4 eligibility, G6 RED and G9 on a price move.

1. Score last week first. Open the CA T6 Decision Log and score every CA row whose review date has arrived: HIT, PARTIAL, MISS or WRONG-DIRECTION, per v4.0 Section 14, rule V1. Unsourced actions do not accumulate, and at CA volume the log is short enough that there is no excuse for a backlog.
2. Read the parameter register. P2 step 7's three checks, before any metric is read. A block that reads metrics before it reads the register can spend the whole session computing against a suspended parameter.
3. Read the CA manifest. T7 statuses and the CA completeness score, stamped on the block header. Name every field group reading MISSING or STALE and the verdict classes it blocks this week.
4. T0 Product Command. Read the CA verdict banner, the six-number header, the CA economics block with its margin-table version and G9 state, the CA gate strip G1-G12, and the CA envelope at zone H. Any RED component forces the AT RISK styling and names the blocked action class; the CA block does not proceed past a G1 quarantine or a G6 RED without executing that gate's response first.
5. T1 Syntax Board. Read the CA syntax strip: spend share against target, cap status, four-quadrant diagnosis with weeks-in-diagnosis, and coverage. Syntax caps are portfolio law and they transfer unchanged; the CA cap is a share of the CA envelope.
6. T2 Keyword Command, oversight view. Apply the CA oversight filter, which in Canada is the P4 qualified set rather than a raw search-volume threshold. Verify gate columns before any metric read: sample state, in-budget, eligibility, indexing, inventory band, family conflict, event mode.

4.7 P7. The Weekly CA Block, Part 2: Pacing and Truncation Sweep (15-25 min; planning estimate)

1. Read in-budget per CA campaign from CA T3 C3 and bucket it against the v4.0 Section 5, gate G3 line. Every rate on a campaign below that line is marked TRUNCATED and the budget lever is evaluated before any bid conclusion is allowed on that campaign. Record the percentage as a number; a pacing line without a figure is not a read.
2. Read utilisation against the objective target, not against a single account figure. The targets are objective-conditional at v4.0 Section 9 "Layer 4: The Lever Library (L1-L10)", rule L4 and one utilisation target across objectives is a category error.

3. Reconcile to the CA envelope. Summed CA campaign budgets against the CA weekly envelope at CA T0 zone H, per gate G7. Over-subscription resolves by the CA M5 rank order, never by an across-the-board reduction and never by borrowing from the US envelope.
4. Stage the budget responses inside rule L4's step limits, one lever per row per cycle. Where the CA day-of-week curve is absent per Section 3, the pacing read runs against the CA weekly envelope only and the block records the substitution rather than implying a curve it does not have.

4.8 P8. The Weekly CA Block, Part 3: The Pricing, Strikethrough and Stock Interfaces (10-20 min; planning estimate)

1. Read, do not decide. This annex executes the PPC consequences of CA pricing and stock decisions and makes neither. The three interfaces are I-7, I-8 and I-9 at Section 12, each with its non-arrival consequence.
2. Price and strikethrough. Take the CA pricing owner's current state and any change dated inside the trailing 7 days. A CA price change fires gate G9: CA CVR and ACOS reads are PROVISIONAL for the trailing 7 days and no fix-first or conversion verdict issues on them. Log the change in CA T5 as event type "Price move". At CA volume a 7-day provisional window removes most of a usable reading period, so the block states which rows are unreadable this week rather than reading them thinly.
3. Stock. Take CA cover and the CA band per variation from the CA inventory owner. Gate G6 responses execute at CA scope the same day: YELLOW freezes CA push spend at its current level, RED tapers to defense and re-points CA ranking campaigns to the declared CA backup. A CA re-point needs a CA backup chain; where the CA variation map has no backup, the re-point response is unavailable and the block records that the only remaining response is the taper.
4. Route what is not PPC. A CA listing suppression, a CA buy-box loss or a CA offer problem routes to its owner the same day with the affected spend quantified, and the affected targets carry gate G4's halt until it clears.

4.9 P9. The Weekly CA Block, Part 4: The CA Search-Term Pass at Scale (25-40 min; planning estimate)

1. Run SOP-14's triage unchanged. The three routes are SOP-14 P3's and this annex does not restate them: converters route to the #30 graduation queue, proven waste takes negative exact, drift clusters take negative phrase on the cluster root. The routes transfer because they are procedures.
2. Read on the window the sample supports, not on the calendar. v4.0 Section 5, gate G2 sets the bars and directs the fallback to the 60-day window when the shorter window is thin. In Canada the 60-day window is the standing read window for rate verdicts rather than the exception, because at CA click rates almost no term reaches the CVR bar inside 7 days. The weekly block still runs weekly; what it does weekly is accumulate, route and stage, and what it does on the 60-day window is issue rate verdicts.
3. Compute the read-limit line and apply it. A term that cannot reach gate G2's CVR bar inside the 60-day window at its current click rate cannot produce a legal rate verdict at all, and carrying it as an extended test forever is the failure that F1 describes. The line is arithmetic on v4.0's own two numbers rather than a new threshold: the CVR bar divided by the number of weeks in the 60-day window. Terms below the line are flagged read-limited in the block record with their current click rate, and they are governed by structure and by gates rather than by rate verdicts until their click rate rises. Computed at WE-1.
4. Hold the graduation bar where v4.0 put it. Graduation runs at the order count v4.0 Section 9, rule L6 sets, and that count does not scale with the marketplace: three orders is three orders of evidence. What changes is that a CA term reaches it over a longer window, so the CA harvest queue is built by accumulation across passes rather than found inside one pass. A CA graduation that fires on a scaled-down order count is F8.
5. Stage the negatives with the evidence line in the reason string, carrying the marketplace key, the route and the window the evidence came from. A CA negative whose evidence line does not name its window cannot be audited later, because the window is the contested part of every CA read.

4.10 P10. The Weekly CA Block, Part 5: Verdicts, Staging and the Session Record (25-40 min; planning estimate)

1. Write one verdict per CA oversight row from the closed vocabulary at Section 5, with its scenario ID and the eight-element justification grammar the Keyword Decision Record Template v3.0 "Read Me" sheet sets out. The CA record is a Master Action Record like any other, and v4.0 Section 13.2 "Master Action Record v2: the output contract" governs its fields unchanged.
2. State the flags honestly in the record. Every fallback benchmark, every provisional read, every read-limited term and every manifest gap appears in the record's flags in the words the workbook computed. At CA volume a record with no flags is more likely to be incomplete than clean.
3. Apply the one-lever rule and the lever order at rule L1: budget where gate G3 is failing, then bid, then placement, then negation, then structure, then status. One row receives one lever per cycle unless the verdict itself names two.
4. Stage every change through #38 with the marketplace key, the verdict, the scenario ID and prior and new values in the reason string. Deploy on the CA block's deployment slot after sign-off, and read the T+1 verification diff line by line. A mismatch is a same-day redeploy of the delta.
5. Take the sign-off. The Authority Matrix v2.0 "Canada weekly block" row holds D and E with the launch and CA operator and R with the PPC Lead, with the CEO and the BM owner informed. Per Section 3's oversight-class note, every CA verdict is in the human-mandatory tier, so the sign-off is the block's gate rather than a formality.

6. Write the session record. One artifact per block: the manifest state and completeness score, the parameter-register check, the gate findings, the read-limited list, the verdicts with their predictions and review dates, the staged set and its diff result, and what the block could not see. The block is the contract and the session record is the evidence it was honoured; a week with no session record is F3 regardless of what was deployed.

4.11 P11. CA Events, Deals and Seasonality (45-90 min per window; planning estimate)

1. The CA calendar is its own calendar. CA windows carry CA classes, including the Boxing Day class that has no US counterpart, and CA Prime windows differ from US ones in dates and depth. A CA window is never planned from the US calendar entry for the same period.
2. Hand the plan to #32 and take back the plan's CA inputs. Event plan construction, cap sizing, arming, the in-window block, revert, settle and scoring are #32's and are not restated here. What this annex supplies is the CA class list, the CA economics spine at the CA deal price, the CA affected set and the CA class factor where one exists.
3. Run the first CA window of every class with no borrowed factor. Rule M9 sizes an event from the product's own measured lift history and directs first-time events to the category default; there is no CA category-default table and a US factor is not one. The CA window therefore runs at the CA baseline with a named dollar allowance as the only expansion authority, exactly as SOP-32 handles an unseeded class, and the gap is Appendix A item AM-11.
4. Compute the window ceiling on the CA deal price. A CA deal price cuts CA contribution in full and therefore cuts the CA window ceiling, while CA conversion rises; the net is arithmetic that the plan answers before the window opens. The method is #32's and the inputs are the CA spine's.
5. Score and seed the CA class factor. The scored CA window is what seeds the CA Event Lift Index for that class. Until CA packet item 5 lands there is no CA seasonality history at all, so the first CA window of each class is the seed and its scoring is worth more than its result.
6. Read CA seasonality separately from US seasonality. Rule T4 times pushes into a rising seasonal index and never after the peak; the index is a CA index. A CA push funded against a US seasonal read buys CA rank into whatever the CA curve is actually doing, which on several bedding families is a different month.

4.12 P12. Packet-Landing Conversion: Structure to Running System (1 day for the full 5 items; planning estimate)

1. Receive item by item, not as a package. Each of the 5 CA packet items unblocks a named set of rows in the P2 register and a named set of decisions. Partial arrival is normal and useful: item 1 alone converts the economics spine and unblocks every ceiling, which is more than half the annex's blocked surface.
2. Item 1, per-product CAD economics. Run P3 in full. Set CA T0 zone C, stamp the margin-table version, clear the G9 blocked state, and publish the spine to its consumers.
3. Item 2, the CA keyword export with volumes and ranks. Run P4 in full: match the sets, measure the ratios, scale the floors, qualify the terms, and record what did not scale. Then run P5 for the structure the qualification implies.
4. Item 3, the CA campaign registry as-is. Load into CA T3, reconcile against the P5 qualified structure, and produce the consolidation and closure list as staged #38 rows. The as-is registry is an input to the work list, never a justification for the structure it describes.
5. Item 4, the CA competitor set. Load into CA T2 K8, then re-open the CA conquest and defence surfaces that Section 4.4's flag-obvious limitation held closed, per #16's own procedures.
6. Item 5, 12 months of CA seasonality by family. Build the CA annual demand calendar alongside the CA deals calendar, and seed the CA class factors where the history supports one. Where 12 months of CA history does not cover a class, the first CA window of that class remains the seed.
7. Reconcile the register and close the version. Every arriving value goes through P2 step 6 before use. When all five items have landed and every register row reads governed or has a named v4.1 item, this annex increments and the Definition Register's PARTIAL status on component #34 closes. The first weekly block on the converted system runs the following Monday.

4.13 P13. Marketplace Instantiation: The Expansion Trigger (2-3 hours per new marketplace; planning estimate)

This procedure absorbs the scheduled fold-in registered against this component in the SOP Program Gap Register v1.0, Section 4: the annex doubles as the template for any future marketplace, and a new marketplace instantiates the annex before any spend, so expansion never starts undefined.

1. The trigger is the expansion decision, not the first campaign. The moment a new marketplace is decided, this procedure opens. Spend in an uninstantiated marketplace is ungoverned spend by definition and there is no threshold below which that is acceptable, because the problem is the absence of a decision surface rather than the size of the loss.
2. Copy this annex and re-answer its four questions per marketplace: what transfers unchanged, what re-parameterizes, what stays deliberately simpler, and what the packet must contain to convert structure into system. The four-part shape is the transferable asset; the answers are not.

3. Build the parameter register first, per P2, before any workbook is cloned. The register is what makes the new marketplace's blanks survivable.
4. Name the marketplace packet and its owner with a date, on the model of the CA 5-item packet and its owner at interface I-1. A packet with no named owner is the state that held this component at v1.0-structure for the whole of the production run.
5. Set the marketplace key before the first row exists. Retro-fitting a marketplace key across a live registry is a schema change touching every table, which is exactly the outcome the open decision on marketplace scope was written to avoid.
6. Declare the cadence from the volume, not from habit. The CA weekly cadence is a consequence of CA volume against gate G2, not a preference. A larger marketplace may earn the US daily-plus-weekly rhythm on day one, and a smaller one may not earn weekly.

5. Decision Points: The Full Verdict-to-Execution Mapping

All verdict issuance, thresholds and scenario logic live in v4.0. The 25 keyword verdict strings below are quoted verbatim from the closed vocabulary in the Keyword Decision Record Template v3.0, "Bands & Verdicts Reference" sheet, which is also the data-validation source list on the PPC Command Workbook TEMPLATE v1.0 REFERENCE tab. Two further strings that v4.0 Section 13.1 "The Closed Verdict Vocabulary" carries and the KDR twenty-five does not are quoted from v4.0 with their source named. Nothing is retired from v1.0 of this annex, because v1.0 carried no verdict strings at all: its Section 5 was a transfer-status table, not a verdict mapping, which is the single largest reason it could not be executed.

The standing rule for this column. Every verdict transfers to Canada unchanged in meaning. What the CA column below states is the CA-specific execution condition: which CA input the verdict consumes, which CA window it may be read on, and where a CA verdict is suspended because its parameter is pending rather than because the verdict does not apply.

Verdict (KDR v3.0 closed vocabulary, verbatim)	CA execution condition
QUARANTINE: DATA REFRESH	Fires on gate G1 in the CA workbook exactly as in the US. At CA volume the commonest cause is a join gap rather than an impossible metric, so the refresh task names the CA feed and not the row. No CA change deploys on a quarantined row.
UNLOCK TRAFFIC	Rule L2's step on a CA row under the bid-read click bar, inside the 50% per-change limit and the floor rule L2 sets. Legal at CA volume and frequently correct there, because a CA term below the click bar is usually under-fed rather than under-performing.
FIX BUDGET (TRUNCATED)	The P7 sweep response wherever CA in-budget sits below the gate G3 line. Executes on the CA envelope only. Where a CA discovery campaign is capped deliberately, the CA T3 row records the deliberate cap so the budget lever is not fired against a designed constraint.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	Executes unchanged; the click quota is the verdict's own and does not scale. At CA click rates the quota takes materially longer to fill, so the test carries its expected fill date computed from the row's CA click rate and is reviewed on that date rather than weekly.
FUND TEST (UNFUNDED PRIMARY)	Executes on a CA primary-syntax term with no CA coverage. The cap and the date are the verdict's; the bid seeds at the CA Expected CPC reference and is suspended while that parameter is pending, because a test funded at a US bid reference tests the bid rather than the term.
RANK PUSH: FUND SIZED PUSH	Suspended until the CA DSTR, CA velocity and CA Expected CPC references exist. Rule M2 needs all three, and a push sized on US references over-sizes by the ratio of the two markets. This is the largest single decision class the CA packet unblocks.
RECOVERY PUSH	Same suspension as the row above and the same reason; the recovery class is the sizing rule applied to a term with a proven CA ceiling, and the proven ceiling must be a CA rank history.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	Executes on a live CA push where CA clicks run under the worksheet line with CA rank velocity intact. The named lever only, never both levers on one row in one cycle.
FIX CONVERSION FIRST; DEFER PUSH	Requires a CA conversion read against a CA benchmark. Suspended on any row inside gate G9's provisional window after a CA price change, and suspended on any read-limited row, because a conversion verdict on a sample that cannot reach the gate G2 bar is the verdict this annex exists to prevent.
PLACEMENT FIX FIRST	#28 executes; this annex stages only the CA campaign-level rows it returns. CA placement work waits on CA placement data, which is a CA feed and not a US one.

Verdict (KDR v3.0 closed vocabulary, verbatim)	CA execution condition
RE-POINT VARIATION	Executes same day on gate G6 RED and outranks any freeze, because re-pointing to the declared backup is what stops CA spend running into a CA stockout. Requires a CA backup chain; where the CA variation map has no backup the response is unavailable and the block records that the taper is the only remaining lever.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	Rule M3's ladder transfers whole, including its step size and the recorded minimum viable CPC. The CA taper reads a CA rank series; a taper stepped against a US rank series is F5 and the account cannot see the position it is giving up.
HOLD; DEFEND LIGHTLY AT RANK	The standing state of a held CA position. At CA volume this is the correct verdict on more rows than in the US, because holding is what a row with insufficient evidence for any other action should do.
DEFEND AT FLOOR	CA branded and floor-held rows sit at their recorded CA floors from the CA register, discovered by #28's ladder against CA blended CPC. A US floor imported into Canada is a price from another auction.
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	The ceiling is the CA Max Profitable CPC from the CA spine, so this verdict is suspended entirely until P3 closes. It is also barred on any read whose CVR sample sits under gate G2's bar, which at CA volume means it is read on the 60-day window and not on the week. Worked at WE-4.
INCREMENTALITY LADDER	#35 owns the test. In Canada the feasibility question comes first: where CA volume cannot clear the read minimums the correct output is that no CA test is possible this cycle, stated with the numbers, rather than a test that returns a result nobody can trust.
ADD EXACT + RESTRUCTURE	Executes at #12 against the CA registry. In Canada the receipt fires on the P4 floor crossing rather than on a raw head-share read, so the structure follows qualification rather than impression.
CONSOLIDATE DUPLICATES	The commonest CA structure verdict during packet-landing conversion, because the as-is CA registry at packet item 3 is where a cloned US structure would show up. Executes at #12; the consolidation list is produced at P12 step 4.
GRADUATE FROM DISCOVERY (v4.0 Section 13.1 names the same action GRADUATE + STEER)	Fires at the order count rule L6 sets, unchanged. What changes in Canada is the window over which the count accumulates, so the CA graduation queue is carried across passes rather than found in one. A CA graduation on a reduced order count is F8.
NEGATE (PRE-LOAD / REACTIVE / STEERING) (v4.0 Section 13.1 names the same action CUT / NEGATE)	All three modes execute in Canada per #29. The evidence line carries the marketplace key and the window the evidence came from, because at CA volume the window is the contested part of every waste read.
ENTER SBV ARBITRAGE	Requires CA SP and CA SBV cost and conversion reads on the same terms. Suspended while CA format economics are pending; the CA posture is SP-heavier until those reads exist, per Section 4.5 step 5.
OPEN / ROTATE / EXIT CONQUEST	#16's verdict. In Canada it runs on flag-obvious cases only until CA packet item 4 lands, and the block records that the mechanical competitor check could not run rather than implying it was clean.
STOP-LOSS: CONCEDE OR DEFER	Executes on any dated CA test or CA push reaching its stop-loss condition. The condition is the worksheet's and is not relaxed for a thinner marketplace; freed CA budget returns to the CA envelope at the instrument layer.
PAUSED-PENDING (REASON + RE-ENTRY)	The parking form for CA rows blocked on a pending parameter. The re-entry condition is written as the specific parameter and its packet item, so the row re-enters mechanically when the value lands rather than when someone remembers it.
LEAVE TO HARVEST	The correct verdict on a converting CA term below its scaled floor: it stays in the CA consolidated structure at CA discovery economics with no dedicated campaign, and it is re-tested against the floor register at every refresh.
CONTINUE; VERIFY MOMENTUM (quoted from v4.0 Section 13.1; absent from the KDR twenty-five)	The standing state of a healthy CA line inside its caps: no changes, the next block continues, and the velocity check is scheduled at the CA cadence. At CA volume this and HOLD; DEFEND LIGHTLY AT RANK are the two most frequently correct outputs of a block.

Verdict (KDR v3.0 closed vocabulary, verbatim)	CA execution condition
PORTFOLIO CALL: ASSIGN LEAD (quoted from v4.0 Section 13.1; absent from the KDR twenty-five)	Fires where gate G11 catches two CA siblings on one CA root. The call is #33's and is made on CA composites; a US ownership declaration on the same root does not settle a CA contest, because the two families face different competitors at different prices.

5.1 The Campaign-Grain Subset

Budgets, multipliers and structure states are campaign-grain actions. The ten strings below are quoted verbatim from the PPC Command Workbook TEMPLATE v1.0 REFERENCE tab CAMPAIGN VERDICT validation list, and land in CA T3 column group C6 per Master PPC Data Architecture v1.0, Section 7 "Level 3 · Campaign Board (Tab T3)".

Campaign verdict (CA T3 C6, verbatim)	When the CA block writes it
FUND BUDGET	The P7 response on a truncated CA campaign inside the CA envelope, and any staged CA budget increase at P10. The reason string carries the marketplace key.
HOLD	The default on any CA campaign pacing to plan and on any campaign whose evidence does not clear a gate this week. At CA volume this is the most frequently written campaign verdict and that is correct rather than passive.
TRIM BUDGET	The response to a CA campaign over-subscribing the CA envelope by the M5 rank order, and to a CA dated test at its stop-loss.
RAISE MULTIPLIER	Suspended until CA placement evidence exists. A CA multiplier raised on US placement decomposition is a US conclusion applied to a different auction.
CUT MULTIPLIER	Executes where a CA multiplier survives from a cloned structure with no CA evidence behind it. During packet-landing conversion this is a routine finding rather than an exception.
NEGATION PASS	The P9 staged output, executed to #29 standards with the CA evidence line and its window named.
CONSOLIDATE	The P5 and P12 structure response wherever CA campaign count exceeds what the CA qualified set supports. This is the F4 remedy.
RETAG OBJECTIVE	Written when a CA campaign's objective no longer matches the term's qualification after a floor refresh, and whenever a CA campaign is found carrying no objective at all, which blocks evaluation under rule O1.
PAUSE	Executes with the reason and the re-entry criterion, per rule L7 and rule O5. In Canada the commonest re-entry criterion is a floor crossing or a pending parameter landing.
ARCHIVE	Rare and deliberate. Archiving destroys the history a thin marketplace can least afford to lose, so a CA campaign is paused with a re-entry condition unless the product itself is exiting Canada.

6. Worked Examples

All four examples use the standard fictional training base because the Tier 3 live product economics spine was not attached to this build and because no CA economics exist to use: AOV \$80.00 USD, contribution \$28.00 USD per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%, baseline CVR 5.2%. All CA figures are illustrative and internally consistent, and every one of them is computed on the page rather than asserted. Live values come from CA T0 zone C and the CA T2 once the packet lands. The corpus base figure of 26.3% for Max Acceptable is 75% of 35.0% presented to one decimal; the exact product is 26.25%.

WE-1 | The scaled-floors problem, worked: what US volume floors do to a CA family, and what the scaled read produces instead

This is the annex's central argument and until this version it had no worked demonstration.

Setup. One CA syntax family, 40 tracked terms, illustrative monthly CA search volume. The four largest terms carry 620, 410, 355 and 290; the remaining 36 terms carry 2,505 combined with none above 190. Family total CA search volume = 620 + 410 + 355 + 290 + 2,505 = 4,180.

Applying the US floors as written. v4.0 Section 6 "Layer 1: Objective Assignment (O1-O8)", rule O2 grants a Ranking objective only on a term with search volume above 500, and a Market Share objective only above 1,000. Against the CA family: exactly 1 of 40 terms clears the Ranking floor, and 0 of 40 clear the Market Share floor. That one term carries 620 of 4,180 = 14.8% of the family's demand, which means 85.2% of the family is structurally ineligible for any objective above Discovery. This is not a judgement that the demand is not worth pursuing; it is an arithmetic accident of applying an absolute number to a market roughly eleven times smaller.

The click side, which is worse. The family runs an illustrative 96 clicks per week: 22, 14, 11 and 9 on the four head terms and 40 across the remaining 36. v4.0 Section 5, gate G2 needs at least 15 clicks in the 7-day window for a bid read, so 1 of 40 terms carries a legal 7-day bid read. Under the US operating pattern of weekly rate verdicts, the other 39 return the same output every week forever: insufficient evidence. That is F1 in its exact form. Everything is UNREAD, so nothing is decided, so nothing changes, so nothing becomes readable.

The scaled read, step by step. The matched US term set for the same family carries an illustrative 46,200 monthly search volume. The measured demand ratio is $4,180 / 46,200 = 0.0905$, that is 1 unit of CA demand for every 11.05 units of US demand on this family. The scaled Ranking floor is $500 \times 0.0905 = 45.24$, recorded in the CA floor register at 45 with the rounding noted. Against that floor: the four head terms qualify, and of the remaining 36 terms an illustrative 14 sit at or above 45 (1,930 combined) with 22 below (575 combined, checking to $1,930 + 575 = 2,505$). Qualified terms = $4 + 14 = 18$ of 40, carrying $1,675 + 1,930 = 3,605$ of $4,180 = 86.2\%$ of the family's demand.

What the two reads produce, side by side. US floors: 1 term governed, 14.8% of family demand inside the objective system, 39 terms with no objective and no route to one. Scaled floors: 18 terms governed, 86.2% of family demand inside the objective system, and 22 terms correctly left in Discovery and harvest because they are genuinely thin even at CA scale. The scaled read is not more permissive in general; it is 5.7 times more permissive on volume qualification and it changes nothing else.

What did not scale, and why that matters more than what did. The gate G2 bars are statistical and do not move. A conversion-rate estimate on 40 clicks carries the same width in either market, so the CVR bar stays where v4.0 put it. What CA volume changes is the time to reach it. The 60-day fallback window that gate G2 already names is $60 / 7 = 8.571$ weeks, so the click rate required to reach the CVR bar inside that window is $100 / 8.571 = 11.67$ clicks per week. Against the family: the head term at 22 per week reaches 188.6 clicks in 60 days and clears; the second at 14 reaches 120.0 and clears; the third at 11 reaches 94.3 and does not; the fourth at 9 reaches 77.1 and does not. So 2 of 40 terms can carry a legal CVR verdict and 38 are read-limited. The read-limit line of 11.67 clicks per week is arithmetic on two v4.0 numbers and is not a new threshold; it is written into the block record, never into a rule.

The lesson. Two different classes of number look identical on the page and behave in opposite ways. Demand floors are relative to a market and must be scaled or they exclude the market. Sufficiency bars are relative to a confidence interval and must not be scaled or they license verdicts the evidence cannot support. The annex's job is to keep those two straight, and F1 and F8 are the two directions of getting it wrong.

WE-2 | CAD-native economics against the FX shortcut, and what the shortcut costs

Setup, CA-native. Illustrative CA figures for one product, computed natively per P3: CA AOV CAD \$109.99; CA landed cost CAD \$41.20; CA referral fee at 15% = $109.99 \times 0.15 = \text{CAD } \16.50 ; CA fulfilment fee CAD \$8.85; CA returns provision at 3.5% of AOV = $109.99 \times 0.035 = \text{CAD } \3.85 . CA contribution = $109.99 - 41.20 - 16.50 - 8.85 - 3.85 = \text{CAD } \39.59 per order, a margin of $39.59 / 109.99 = 36.0\%$.

The CA chain, per rule M1. CA break-even ACOS = 36.0%. CA Target = 50% of break-even = 18.0%. CA Max Acceptable = 75% of break-even = 27.0%. At an illustrative CA conversion rate of 4.6%, CA Max Profitable CPC = $39.59 \times 0.046 = \text{CAD } \1.82 , and the CA target CPC = $109.99 \times 0.18 \times 0.046 = \text{CAD } \0.91 . Those two figures are the whole CA bidding envelope for this product and neither was derived from a US number.

The FX shortcut, run for comparison. Take the US contribution of \$28.00 and the US AOV of \$80.00 and convert at an illustrative rate of 1.3600 CAD per USD: contribution CAD \$38.08 on AOV CAD \$108.80. Break-even ACOS = $38.08 / 108.80 = 35.0\%$, unchanged from the US, because conversion preserves the ratio. CA Max Profitable CPC on this basis = $38.08 \times 0.046 = \text{CAD } \1.75 .

The first error: the level. The shortcut understates contribution by $39.59 - 38.08 = \text{CAD } \1.51 per order, which is 3.8% of the true CA figure. It does so through two mistakes that partly cancel: the CA price is not the converted US price (CAD \$109.99 against CAD \$108.80) and the CA fee schedule is not the US one. The shortcut's answer was close for reasons that will not repeat, which is worse than being clearly wrong, because nothing in the output signals that it should be re-derived.

The second error, and the one F2 names: the ceiling tracks the currency. Move the illustrative rate from 1.3600 to 1.4100, a 3.68% move. The FX-derived contribution becomes $28.00 \times 1.41 = \text{CAD } \39.48 and the FX-derived CA Max Profitable CPC becomes $39.48 \times 0.046 = \text{CAD } \1.82 , a 3.68% rise. Nothing about the Canadian business changed: the same CA price, the same CA fees, the same CA conversion rate, the same competitors. The CAD-native ceiling stayed at CAD \$1.82 across both dates because it never consumed a rate at all.

What the swing costs. On an illustrative 240 CA clicks per week across the affected set, a ceiling moving CAD \$1.75 to CAD \$1.82 moves the authorised spend band by $240 \times 0.07 = \text{CAD } \15.46 per week, or $13 \times 15.46 = \text{CAD } \200.93 per quarter, on zero business change. The dollar figure is small and it is not the damage. The damage is that every ladder step, every recorded floor from rule M3 and every push worksheet re-prices against a reference that moved for no operational reason, so the discovered minimum viable CPC becomes untestable: the ladder cannot tell a competitive shift from a currency tick. A floor that cannot be re-tested is not a floor.

The lesson. FX is a reporting layer. The moment it enters a decision cell, the product's economics acquire a second driver that no PPC lever can respond to, and the account starts making bidding decisions about the Canadian dollar.

WE-3 | A clean CA weekly block

Setup. One CA product, post-conversion, running six CA campaigns against a CA weekly envelope of CAD \$420.00 and the CA spine from WE-2 (contribution CAD \$39.59, break-even 36.0%, Target 18.0%, Max Acceptable 27.0%, Max Profitable CPC CAD \$1.82).

P6, cascade-lite. Four CA T6 rows carry review dates this week and are scored first: two HIT, one PARTIAL, one MISS. The CA manifest reads the CA competitor group MISSING and the CA day-of-week curve absent, so the header records that conquest verdicts and curve-based pacing are blocked this week. CA T0 shows the gate strip clear and the envelope at zone H.

P7, pacing. CA spend for the week reads CAD \$381.50 against the CAD \$420.00 envelope = 90.8% envelope utilisation. Campaign CA-EX-01, the hero exact, budgeted CAD \$14.00 per day, spent CAD \$92.40 = CAD \$13.20 per day = 94.3% utilisation and in-budget at 96%: healthy against gate G3's line, and at or above the ranking utilisation target, so no budget lever fires. Campaign CA-BR-01, the family broad, reads in-budget 58%: below the gate G3 line, so every rate on it is marked TRUNCATED, the budget lever is evaluated before any bid conclusion, and the verdict written to CA T3 is FIX BUDGET (TRUNCATED). No bid on CA-BR-01 is touched this cycle, which is the one-lever rule and rule L1's order working together.

P8, interfaces. The CA pricing owner reports no change inside the trailing 7 days, so no row carries the gate G9 provisional mark. The CA inventory owner reports all CA variations GREEN with cover clear. Both are recorded as received; a recorded absence of change is a valid arrival and silence is not.

P9, the search-term pass. 41 CA served terms in the window. Route B takes two terms with spend and no orders at their price class, staged as negative exact with the evidence line naming the 60-day window. Route C takes a 3-term drift cluster on one irrelevant root, staged as one negative phrase. Route A finds no term at the graduation bar this week, which is the expected CA outcome: at CA click rates the bar is reached across passes, so the two candidates nearest it are carried in the CA harvest queue with their running order counts and their expected arrival dates rather than being graduated early or dropped.

P10, verdicts and record. Eleven CA oversight rows receive verdicts. Six read HOLD; DEFEND LIGHTLY AT RANK or CONTINUE; VERIFY MOMENTUM, which at CA volume is the correct majority. One reads FIX BUDGET (TRUNCATED). Two read NEGATE (PRE-LOAD / REACTIVE / STEERING) from the P9 routes. Two read PAUSED-PENDING (REASON + RE-ENTRY), both blocked on the CA Expected CPC reference with that parameter named as the re-entry condition. Eight rows carry the read-limited flag with their click rates. The staged set goes out through #38 with the marketplace key, the T+1 diff returns clean, and the session record carries the manifest state, the two blind spots, the eleven verdicts with their predictions and review dates, and the read-limited list.

What was not done. No bid moved on any read-limited row, no CA conquest verdict was written while the CA competitor group read MISSING, and no rate verdict was written on the 7-day window. Three separate gate rules agreed on the same restraint and the block wrote down that they had.

WE-4 | The edge case: a CA verdict issued on a window that could not carry it

Setup. The same CA product and spine. Campaign CA-EX-02 carries one CA term. The week reads: spend CAD \$58.90, 19 clicks, 1 order.

What the 7-day read appears to show. $CPC = 58.90 / 19 = \text{CAD } \3.10 . $CVR = 1 / 19 = 5.26\%$. Sales at CA AOV = $1 \times 109.99 = \text{CAD } \109.99 , so $ACOS = 58.90 / 109.99 = 53.6\%$, which is 1.49 times the CA break-even of 36.0% and twice the CA Max Acceptable of 27.0%. Ad cost per order = CAD \$58.90 against CA contribution of CAD \$39.59, a loss of CAD \$19.31 on the order. On the face of it this is an unambiguous BID DOWN TO PROFITABLE CPC; CAP SYNTAX, and the operator stages it.

Why it is not legal (F5 and F8 together). Two gates bar the read independently. Gate G2 requires at least 100 clicks for a CVR verdict, and the ACOS figure above is a CVR-driven rate read carried by a single order: the whole verdict rests on whether that one order was 1-in-19 or 1-in-40. Gate G3 marks every rate on the campaign TRUNCATED, because CA-EX-02 ran at 62% in-budget, below the gate's line, so the budget lever is evaluated before any bid conclusion is allowed. The staged step is withdrawn before deployment and the withdrawal is logged with the gate that should have blocked it.

What the legal window shows. On the 60-day window the same term reads: spend CAD \$402.60, 141 clicks, 9 orders. $CPC = 402.60 / 141 = \text{CAD } \2.855 . $CVR = 9 / 141 = 6.38\%$. Ad cost per order = $402.60 / 9 = \text{CAD } \44.73 against contribution CAD \$39.59, a loss of CAD \$5.14 per order. The direction agrees with the 7-day read and the magnitude does not: $19.31 / 5.14 = 3.76$ times overstated. A verdict sized on the short window would have been sized 3.76 times too hard.

The correct sequence. Budget first, per rule L1 and gate G3: the truncation is the finding this cycle and the response is the budget lever, not the bid. The rate verdict waits for the next block on the 60-day window with the truncation cleared. When it comes, the ceiling is CA Max Profitable CPC computed on the 60-day CVR: $39.59 \times 0.0638 = \text{CAD } \2.53 . Stepping the CPC from CAD \$2.855 to CAD \$2.53 is an 11.5% move, well inside rule L2's per-change limit, so the ladder step is a single legal step rather than a cut.

What the withdrawal saved. Had the 7-day read driven the sizing, the ceiling would have been computed on the 7-day CVR: $39.59 \times 0.0526 = \text{CAD } \2.08 . That is CAD \$0.44 below the correct ceiling, a bid set 17.5% too low. The term runs $141 / 8.571 = 16.45$ clicks per week, so over the following four weeks the wrong bid would have governed roughly 65.8 clicks at a level that risks a CA position the account cannot currently see, because the CA rank series is a CA feed and the US daily rank series says nothing about it. The cost of the error is not the immediate spend; it is a position given up on evidence that never existed, in the one marketplace with no daily series to catch it.

The lesson. At CA volume the temptation runs in both directions in the same session: scale the floors down so that anything can be governed, and scale the sufficiency bars down so that anything can be decided. The first is correct and the second is F8. The 60-day

window is not a delay; it is the shortest window in which the verdict is real.

7. Failure Modes and Escalation

F1 to F4 carry forward from v1.0 Section 8 with their detection and response intact; F1 and F2 are the annex's two named patterns and the reason it exists. F5 to F12 are new at v2.0. Each row states the detection signal, the response, and the motivating pattern where one is recorded.

Failure	Detection signal	Response	Motivating pattern
F1. Raw US thresholds applied to CA volumes	Everything reading UNREAD in a CA block, or triggers that never fire across consecutive blocks with CA spend active; any CA qualification run against an unscaled floor	Run P4 and qualify against the scaled floor register. Where the CA volume feed is absent, record the family as unqualified rather than substituting the US floor. Quantify the excluded demand as a share of family search volume, as WE-1 does	Carried from v1.0 Section 8 as the annex's stated reason to exist. WE-1 quantifies it at 85.2% of one family's demand excluded by a single unscaled number
F2. FX-converted margins inside the decision chain	CA T0 zone C cells traceable to a US value times a rate; CA margins moving with the currency between refreshes on an unchanged CA price and fee schedule	Rebuild the spine natively per P3 from CA inputs only. Void every ceiling computed from the converted values and re-derive; FX returns to the reporting layer	Carried from v1.0 Section 8. WE-2 quantifies the two limbs: contribution understated by CAD \$1.51 per order, and a ceiling moving 3.68% on a currency tick alone
F3. The skipped weekly CA block	No CA session record for a week in the CA T6	Run the block late and log the gap with what it could not reconstruct. Two consecutive gaps escalate and no CA deployment resumes until a block has run	Carried from v1.0 Section 8. The block is the contract; a skipped week is not a scheduling detail because nothing else reads the CA account at all
F4. US structure cloned 1:1 into Canada	CA dedicated campaign counts far above what the CA qualified set at P4 supports; CA campaigns whose terms sit below every scaled floor	Consolidate per P5 and the Doctrine's two-class boundary; stage the consolidation list through #38. Structure follows qualified volume, never the US registry	Carried from v1.0 Section 8. Per-term isolation at tail volume produces campaigns that never leave learning, which is the Doctrine's own reason for the consolidated class
F5. A CA decision made on a US input	Any CA verdict, bid, floor, DSTR, benchmark or rank read whose source is a US feed, US series or US column; the CA record's source field naming a US artifact	Withdraw before deployment; re-derive on the CA input or suspend the decision with the parameter named as the re-entry condition. Where it deployed, void the affected reads and quantify the spend that carried no valid basis	New at v2.0. This is the generalised form of F2 and it is the failure the parameter register exists to make visible: every register row's source column is the detector
F6. CA spend mixing into US totals	A CA row with no marketplace key; a US total that reconciles only when CA spend is included; a platform rule scoped to a campaign set that spans both marketplaces	Set the key at source per P1 step 2 and re-key the affected rows; re-scope the rule at #37. Every account total is restated for the affected periods rather than corrected forward silently	New at v2.0, from the corpus's own open decision on marketplace scope, which recommends tagging every packet and registry key from day one precisely so a retrofit is never needed
F7. A CA parameter with a value and no source	The P2 register's completeness column on any row where the value column is populated and the source, derivation or consumers cell is not	Suspend the value from use the session it is found, and every decision in the consumers cell with it. The row closes by naming the source or the value is removed; a number nobody can re-derive is not a parameter	New at v2.0. This is how a placeholder becomes a threshold: someone needs a number, writes a plausible one, and the next operator finds it in a register and assumes it was governed
F8. Sample-sufficiency bars scaled down for Canada	Any CA rate verdict on a sample below the gate G2 bar; a CA graduation on a reduced order count; a CA floor register carrying a scaled click or impression bar	Withdraw the verdict and re-read on the window that reaches the bar. Restate the CA floor register's own statement of what did not scale, per P4 step 4	New at v2.0 and the mirror image of F1. The operator who has just learned that US floors are too high for Canada reaches for the same divisor on the click bars; WE-4 quantifies the result at 3.76 times overstated
F9. A CA change deployed without its verification diff	A CA staged set with no #38 T+1 diff read against it, or a diff read in aggregate rather than line by line	Read the diff and redeploy the delta the same day. A CA block that cannot show its diff cannot show that it changed what it says it changed, and the week's predictions become unscorable	New at v2.0. At CA volume a single unverified change is a larger share of the account's total change set than the same error in the US
F10. A CA event run on US class factors or a US calendar	A CA window plan carrying an Event Lift Index sourced from US history, or a CA window dated from the US calendar entry for the same period	Run the window at the CA baseline with a named allowance as the only expansion authority, per P11 step 3. Score it and seed the CA factor from the result	New at v2.0. The Boxing Day class has no US counterpart at all, which makes the absence of a CA factor structural rather than a data gap that time will close

Failure	Detection signal	Response	Motivating pattern
F11. Pricing, strikethrough or stock decided inside this annex	A CA block record containing a price decision, a promotion decision or a stock decision rather than a read of one	Return the decision to its owner the same day with the PPC consequence quantified. This annex executes consequences; a PPC document that starts deciding prices has acquired a second job and will do neither well	New at v2.0. The boundary is stated in v1.0 Section 1 without a detector, so nothing found the violation
F12. Packet items used before the register reconciles them	A CA decision citing an arriving packet value whose register row has not been through P2 step 6	Suspend the decision until the row closes. Landing day is exactly when a value is most likely to be used before anyone has written down where it came from	New at v2.0. The conversion at P12 is a single busy day and the register check is the only thing standing between a fast conversion and an unauditable one

Escalation. F1 and F8 escalate to the PPC Lead on discovery with the affected verdict count and the demand or evidence share quantified, because both are systematic rather than incidental: they misgovern a class of rows, not a row. F2 escalates the same day to the PPC Lead and the finance seam, since a converted spine invalidates every ceiling downstream of it. F3 escalates on the second consecutive gap. F6 escalates to the PPC Lead and the Head of Amazon Operations, because restating account totals is a reporting event and not a PPC one. F11 escalates to the owner of the decision that was taken, not to PPC. Anything in the human-mandatory tier is confirmed at the Authority Matrix "Canada weekly block" sign-off before deployment, without exception.

8. Outputs and Artifacts

Gates clear with artifacts, not assertions. Every row names the record, where it is stored, how long it is kept and who consumes it.

Event	Artifact set produced	Storage and retention	Downstream consumer
Instantiation (P1)	The CA Command Workbook with all ten tabs, the marketplace key set, and the CA T7 manifest built from CA feed reality with a CA completeness score	The CA workbook, one per product, versioned; the CA row in Account Rollup A0 and the CA family board at A1	Every later CA procedure; the PPC Lead at review; the Definition Register, which cannot close #34 until instantiation exists per product
The parameter register (P2)	One row per CA parameter with definition, source, derivation, scaling class, consumers, status and value; the weekly reconciliation result	A standing document against this annex, versioned on every change, retained across annex versions because the derivation history is the audit trail	Every CA decision; the v4.1 amendment owner, who receives Appendix A from it; the incoming CA operator, for whom it is the fastest route into the component
Economics spine (P3)	CA T0 zone C named cells in CAD with the CA margin-table version and the currency label; the CA staleness clock	CA T0, refreshed monthly with the version stamped; prior versions retained so a past verdict can be re-read on the spine it was written against	#23 push worksheets; #28 ladder ceilings; #26 CA floors; #27 CA terminal math; every CA ceiling in the workbook
Floor register (P4)	Per syntax family: the matched term counts, both search-volume totals, the measured ratio, the scaled floors with their rounding, the qualified term list, and the explicit statement of what did not scale	Against the annex, re-issued at every quarterly re-measure with the prior measurement retained; the qualification state lands in CA T2 K1	O2 objective assignment; P5 structure class; #23 push candidacy; #16 conquest candidacy; #35 CA test feasibility; the CA oversight filter the block reads
Structure instantiation (P5)	The CA registry with objective tags and roles; the consolidation and closure list; the wall list per structure event; the staged #38 set	CA T3 for the registry; the staged sets with the #38 session record; T6 rows carrying the structure predictions	#12 and #29 at execution; #38 at deployment; the monthly CA structure finding at rule M6
The weekly CA block (P6-P10)	One dated session record carrying: the manifest state and CA completeness score, the parameter-register check, the scored prior-cycle rows, the gate findings, pacing and truncation reads as numbers, the interface receipts, the search-term routes with counts, the read-limited list with click rates, every verdict with its prediction and review date, the staged set, and the T+1 diff result	CA T6 for the decision rows, append-only; the session record against the CA workbook, retained for the life of the product	Next week's scoring pass; the PPC Lead at sign-off; Account Rollup A3 verification scoreboard; #38 at deployment; every component receiving a CA routed candidate
Interface receipts (P8)	The CA pricing and strikethrough state, the CA stock state and band per variation, each recorded as received with its date, including where the report is that nothing changed	In the session record and, where a change occurred, as a CA T5 Product Log row under its event type	Gate G9 provisional marks; gate G6 responses; every later CA attribution check, which reads the log before crediting or blaming PPC
CA events (P11)	The CA class declaration, the CA affected set, the CA spine at the deal price, and the scored CA window with its class-factor seed	Against the CA event plan #32 holds; the class factor accumulates as a CA series	#32 for the plan and its scoring; #23 for any amplified CA push; the CA deals calendar owner; the next CA window of that class

Event	Artifact set produced	Storage and retention	Downstream consumer
Packet conversion (P12)	The per-item conversion record: what landed, which register rows closed, which Appendix A items retired, and which decisions resumed	Against the annex and the Definition Register row for #34	The Definition Register, which closes the PARTIAL status on the fifth item; the Ledger row for #34; the PPC Lead
Marketplace instantiation (P13)	The new marketplace's four answers, its parameter register, its named packet with a dated owner, and its marketplace key	A new annex document for that marketplace, registered in the Supersession Map	The CEO and Head of Amazon Operations at the expansion decision; every operating SOP, which then applies there as written

9. Skill-Conversion Block

Block	Specification
Input bindings (CA manifest field groups and named cells)	The CA parameter register with its status and consumers columns; the CA T0 zone C named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC and WeeklyEnvelope with the CA margin-table version and G9 state; A2 CA deployment state at CA T3 C2; A3 CA pacing and in-budget at CA T3 C3; A7 CA eligibility; the CA inventory band and days to stockout at CA T2 K2; the CA volume feed at CA T2 K1 with its pull date; the CA rank read at CA T2 K4; the CA MKL refresh date and syntax tags; the CA competitor block at CA T2 K8 where it exists; the CA T5 Product Log; the CA T6 Decision Log; the CA T7 Manifest; the CA registry at CA T3; the #38 T+1 verification diff; the CA pricing and CA inventory owner reports; the CA deals calendar; and the marketplace key on every one of them. An absent CA feed suppresses its dependent execution class per the manifest, and the skill states what it could not see. The skill never reads a US feed to fill a CA binding, and never substitutes a US value for a pending CA parameter: both are hard stops, not degradations.
Deterministic logic references	v4.0 gates G1 to G12 by their v4.0 numbers, with the absolute and demand-relative split at Section 3 of this annex; rule O1 objective requirement, rule O2 assignment and its search-volume floors as the scaled inputs to P4, rule O3 realism on CA inputs, rules O5 to O7; rules D1 to D10 for diagnosis order and benchmark hierarchy; rules L1 lever order, L2 bid limits, L4 budget steps, L5 negation modes, L6 structure and L7 status; rules M1 the break-even chain, M2 clicks-to-rank, M3 the taper ladder, M4 effective-CPC ceilings, M5 allocation, M6 structure targets and M9 event sizing; rules T1, T2, T4 and T8; rules P1, P2, P5 and P10; rules V1, V4 and V5; the Section 5 mapping and the 5.1 campaign-grain subset of this annex; and P2's three scaling classes, which are this annex's own logic and the one place the skill decides anything. The skill executes mappings, never computes its own thresholds, and version-pins v4.0 and the parameter register together, because a CA verdict is only reproducible if both versions are known.

Block	Specification
Human checkpoints, drawn from the Authority Matrix v2.0 rows that apply	Per the "Canada weekly block" row (D and E: the launch and CA operator; R: the PPC Lead; I: CEO and BM owner): human-mandatory are the weekly block sign-off in full, every CA verdict inside it per the Section 3 oversight-class note, every CA structure change, and every CA parameter admitted to the register. Per "v4.0 threshold changes" (D: CEO; E and R: PPC Lead; I: the CA operator): the CA operator is informed and never decides a threshold, which is why every CA gap in this document is an Appendix A item. Per "Structure creation (all SOPs)" (E: launch and CA operator in own scope): CA campaign creation, consolidation and closure execute as E inside the signed block. Per "Bulk deployments (#38)" (E: own scope; R: PPC Lead): CA staged sets deploy as E and the T+1 diff is read before the block closes. Per "Event plans (#32)" (E: launch and CA operator in own scope): CA window execution is E; the plan's D stays with the PPC Lead and the deals owner. Per "Push funding (the M5 queue)" (E: launch scope; D: PPC Lead within envelope, CEO above it): any CA push funding decision routes up, and no CA push is funded inside the block. Per "Launch-stage verdicts + batches" (D and E: launch and CA operator): a CA product still in launch stage runs its launch verdicts on this same row. Per "Code-red declaration + close" (D: PPC Lead): a CA code red is declared by the PPC Lead on judgement while the CA band is pending, and the CA operator executes. Standing rule 5 of the matrix reverts an unfilled role's D to the PPC Lead and its E to the PPC Manager, which is the operative state whenever the CA seat is vacant. AI-executable with audit: cascade-lite drafting, the parameter-register completeness check, the ratio measurement and floor scaling at P4, term qualification, the pacing and truncation sweep, the search-term triage draft with all three routes and their evidence lines, read-limit computation and flagging, verdict drafting with the justification grammar, staged-set drafting in #38 format with the marketplace key, the diff read, and F1 to F12 detection with evidence lines.
Cadence	Weekly: the full CA block, one session, on the Monday cascade slot after the CA feeds refresh, with deployment on the CA slot after sign-off. Same day regardless of the block: gate G4 eligibility halts, gate G6 RED responses, gate G9 marks on a CA price change, and any #38 diff mismatch. Monthly: the CA margin-table refresh and the CA structure finding at rule M6. Quarterly: the demand-ratio re-measure and floor re-derivation at P4 step 6, the recorded-floor re-tests at rule M3, and the SD pool and AMC read-minimum re-check. Semi-annually: the format-posture re-check. On event: the CA window procedures at P11. On arrival: P12 per packet item. On decision: P13 per new marketplace.
Alarm specification	F1 to F12 each carry the Section 7 detection signal, an evidence-line format (CA campaign or term, the signal value, the gate or rule reference, the register row where a parameter is implicated, and the marketplace key) and a routing tier: F1, F2, F5 and F8 to the PPC Lead the session they are found, because each misgoverns a class of rows rather than one row; F3 and F6 to the PPC Lead and the Head of Amazon Operations on the second occurrence; F4, F7, F9, F10 and F12 into the block record and the next block's worklist; F11 to the owner of the decision that was taken, the same day.
Skill outputs	Drafted CA workbook instantiations with the manifest built from feed reality; drafted and reconciled parameter-register rows with the scaling class proposed and flagged for confirmation; computed demand ratios, scaled floor registers and qualified term lists with the did-not-scale statement attached; drafted CA structure and consolidation lists; the full weekly CA block record with every read as a number; drafted CA verdicts with predictions, review dates and honest flags; staged #38 sets with the marketplace key; the read-limited list; the CA harvest queue with running order counts and expected arrival dates; CA event class-factor seeds; and F1 to F12 alarms with evidence lines. Every draft is queued for the confirmation tier its Authority Matrix row requires, which for this component is the whole block.

10. Version Governance

Owner: the CA operator under the PPC Lead, per the Authority Matrix v2.0 "Canada weekly block" row, which holds D and E with the launch and CA operator and R with the PPC Lead. Review quarterly at the demand-ratio re-measure and immediately on any v4.0 increment, at which point the Section 5 mapping and every gate, rule and scenario reference in this annex re-validates before the next CA block runs. The parameter register versions with the annex and travels with it.

Component-specific review triggers beyond the quarterly slot.

1. Any CA packet item landing. Each of the five items closes a named set of register rows and retires a named set of Appendix A items; the annex increments on each, rather than waiting for all five, because the value of item 1 arriving is not deferred by item 5 being late.
2. A v4.1 release carrying the CA parameter column. This is the release that converts the majority of this document from specification to execution, and it retires AM-1 through AM-13 in one increment. Every register row's status flips and every suspended decision class re-opens in the same block.
3. A measured demand ratio moving past its recorded review band on any family. The floors move together in one staged change; a register that has drifted family by family is no longer a register.
4. A CA fee-schedule or CA referral-rate change. The CA spine re-derives at P3 and every CA ceiling repoints. This is a CA calendar event and does not follow the US fee calendar.
5. A CA marketplace structural change: a new CA program, a CA-only ad format, or a change in CA deal mechanics, each of which changes what P8, P9 or P11 receives.
6. The first CA window of any class, because it seeds that class's factor and the plan shape for every window after it.
7. Any rebuild of #12, #14, #23, #26, #28, #32, #33 or #38, which re-validates the boundary rows at Section 1.4 that are written from Ledger rows rather than from procedure numbers.
8. An expansion decision for any new marketplace, which runs P13 and produces a sibling annex; the CA annex is then reviewed for anything the new instantiation exposed as CA-specific that had been written as general.

v2.0 change note. Supersedes v1.0-structure in full. Procedures expanded from 0 to 13: the annex previously carried no step-level procedures at all and was one of the three documents the program index recorded as a skeleton rather than a thin document. The v1.0 six-row PENDING-CA table becomes the 14-row parameter register at P2, with each row carrying definition, source, derivation, scaling class, consumers and status, so that a pending value suspends a named list of decisions instead of leaving a blank. The three scaling classes (transfers, native, ratio-derived) are new and are the mechanism by which the annex distinguishes what must scale from what must not. Verdict mapping is new: v1.0 carried no verdict strings, and v2.0 maps the full 25-string closed vocabulary from the Keyword Decision Record Template v3.0 plus the two strings v4.0 Section 13.1 carries and the KDR does not, plus the ten-string campaign-grain subset this annex writes into CA T3. Worked examples expanded from 0 to 4, computed on the standard fictional training base, with WE-1 supplying the scaled-floors demonstration the annex's central argument had never carried. Failure modes expanded from 4 to 12, with F1 to F4 carried forward intact and F5 to F12 added, including F8 as the explicit mirror image of F1. Elements 11 and 12 added per the Build Kit standard, with 16 interface contracts each carrying a non-arrival consequence. Absorbs the scheduled marketplace-expansion trigger clause registered against this component, at P13. Two corrections of substance to v1.0: its Section 2 attributed the in-window freeze to gate G12 and named G10 as the objective gate, where v4.0's G12 is MKL hygiene and G10 is event exclusion, so this version names every gate as v4.0 names it and records the corpus alias divergence in Appendix A; and its Section 3 cited a CA parameter column in v4.0, which v4.0 does not contain and which the program's own amendment queue lists as pending, so that citation is retired and becomes Appendix A item AM-1.

The register's status on component #34 stays PARTIAL until the fifth packet item lands and P12 step 7 closes. Registered in the Supersession Map.

Companions: Master PPC Decision Criteria System v4.0 (all thresholds, gates, scenarios and verdicts) | Keyword Decision Record Template v3.0, "Bands & Verdicts Reference" (the closed vocabulary) | PPC Strategic Doctrine v1.0, Sections 5.1 and 5.3 (the two-class structure doctrine and the marketplace-split rule) | PPC Authority Matrix v2.0 (Elements 9 and 11) | Master PPC Data Architecture v1.0 and the PPC Command Workbook TEMPLATE v1.0 READ ME (instrument mechanics and the clone procedure) | SOP-12 Exact, SOP-13 Phrase, SOP-14 Broad, SOP-15 Auto, SOP-16 Conquest, SOP-23 Push, SOP-24 Recovery, SOP-26 Mature and Code Red, SOP-27 Terminal, SOP-28 Placement, SOP-29 Negation, SOP-30 Harvest, SOP-32 Events, SOP-33 Arbitration, SOP-35 Incrementality, SOP-37 Automation, SOP-38 Bulk (all apply in Canada as written unless a row here overrides) | PPC Cross-Reference Ledger (component #34 row) | the CA pricing, CA inventory and CA deals owners (weekly interfaces).

11. Operator Ramp

Built from the Authority Matrix v2.0 rows that name the launch and CA operator: "Canada weekly block" (D and E: launch and CA operator; R: PPC Lead; I: CEO and BM owner), "Launch-stage verdicts + batches" (D and E), "Structure creation (all SOPs)" (E in own scope), "Bulk deployments (#38)" (E in own scope), "Event plans (#32)" (E in own scope), "Push funding (the M5 queue)" (E in launch scope), "v4.0 threshold changes" (I) and "Automation rules (#37)" (I). Each row becomes a demonstrated competence. The checkpoint holder for review and audit is the PPC Lead per the R cell. The matrix's vacancy rule at Section 2 standing rule 5 reverts an unfilled role's D to the PPC Lead and its E to the PPC Manager until the seat fills, which is the operative state during any gap in the CA seat and is therefore part of what the incoming operator is told on day one. The matrix's standing rule 4 records that the Completion Directive still names the PPC Lead on two rows the Transition Brief v1.2 assigns to the incoming operator; the brief controls, and this ramp is written to the brief.

Note on the two rows the CA operator holds as I rather than D or E. Being informed on v4.0 threshold changes and on automation rules is not a lesser competence; it is the boundary that keeps this component honest. The CA operator is the person most exposed to missing numbers and therefore the person most tempted to write one. Week 1 tests that they know which side of that line they stand on.

Week	Demonstrated competence
Week 1: observation and artifact familiarity	Shadows one full CA weekly block end to end, or reconstructs the last completed block from its session record where no live block falls in the week. Locates and reads unaided: the CA parameter register and, for any three rows, the source, derivation, scaling class and consumer list without prompting; the CA T7 Manifest with its CA completeness score, and names three verdict classes currently blocked and why; CA T0 zone C and the currency label on every named cell; the CA floor register with one family's measured ratio and the scaled floors it produced; the CA T5 Product Log and what a "Price move" row does to the following week's reads; and one prior CA session record with its read-limited list. Reproduces the WE-2 CA spine cold on fresh illustrative CA inputs, matching to the cent: CA contribution, CA break-even, CA Target, CA Max Acceptable, CA Max Profitable CPC and CA target CPC, then states what each figure would have been under the FX shortcut and why the difference is not a rounding question. Reads the Authority Matrix "Canada weekly block" row aloud and states what the CA operator may not decide.
Week 2: supervised execution, checkpoint holder reviewing	Runs three consecutive CA weekly blocks with the PPC Lead reviewing every table before anything stages: the register check, the manifest read with the completeness score on the header, the pacing and truncation sweep with every in-budget figure recorded as a number, the interface receipts including the ones reporting no change, the three-route search-term pass with an evidence line and a named window on every staged negative, and the verdicts with predictions and review dates. Measures one syntax family's demand ratio end to end at P4, derives the scaled floors, produces the qualified term list, and writes the did-not-scale statement in their own words. Drafts one structure recommendation from that qualification and takes it to a CONSOLIDATE or a dedicated-campaign recommendation with the arithmetic on the page. Pass bar for the week: zero pacing lines recorded without a figure, zero rate verdicts written on a sample under the gate G2 bar, and zero staged rows without the marketplace key.
Week 4: unsupervised execution, checkpoint holder auditing	Owns the CA scope: the weekly block in full, structure creation and consolidation inside the qualified set, CA staged deployment and its diff, and the CA event execution where a window falls in the period, with the PPC Lead auditing artifacts after the fact rather than gating them before. The sign-off on the block itself is not delegated, because it is the D and R pair that the Authority Matrix row holds and it is also the only human-mandatory check the whole CA surface gets. The audit samples: one register row traced from a live verdict back to its source artifact; one floor traced to its measured ratio and its two search-volume totals; the read-limited list against the click rates that produced it; the T+1 diff against the staged set line by line; and the session record's blind-spot statement against the manifest. Two clean audited blocks in a row, plus one correctly refused verdict (a decision the operator declined to make because its parameter was pending, with the parameter and its packet item named), close the ramp. The refusal is a pass condition rather than a nice outcome, because in this component the ability to not decide is the competence being certified.

11.1 Live-Defense Questions

These test the reasoning behind the thresholds, not recall of their values.

1. The search-volume floor for a Ranking objective scales for Canada and the click bar for a conversion verdict does not. Both are numbers in the same document about the same keyword. Explain what makes them different kinds of number, and what goes wrong in each direction if the distinction is missed.
2. A CA product's contribution can be computed two ways and both produce a plausible figure. Name the two, state which one moves when nothing about the Canadian business has changed, and explain why that movement makes a recorded bid floor untestable rather than merely inaccurate.
3. A CA term reads 53.6% ACOS on 19 clicks and one order, against a break-even of 36.0%. Two separate gates independently bar the obvious verdict. Name both, and explain what each one is protecting, since they are not protecting the same thing.
4. The CA cadence is weekly and the US cadence is daily plus weekly. Defend the difference from the evidence rather than from workload, and name the specific gate rule that makes a daily CA read a worse read rather than merely a more expensive one.
5. The graduation bar does not scale for Canada but the volume floor does. A term that would graduate in the US in one week takes a CA term two months. Explain why holding the bar is correct, and what the CA harvest queue must carry so that the longer window does not become a lost term.
6. A parameter row in the register has a value and no source. Explain why that is treated as more dangerous than the same row with a source and no value, and what the next operator would do with each.

7. The annex assigns every parameter one of three scaling classes. Take the CA Expected CPC reference and the CA break-even ratio chain, state which class each carries, and explain why the classes are opposite for two quantities that appear in the same bid calculation.
8. v4.0's human-mandatory review tier keys on a search-volume figure that almost no CA term reaches. State what this annex does about that, and explain why the resolution runs in the stricter direction rather than the looser one.
9. A CA sibling contest on a shared root already has a US ownership declaration for the same root. Explain why that declaration does not settle the CA contest, and what the CA block does while the CA call is outstanding.
10. The first CA window of a class has no lift factor and no CA history to borrow one from. Explain why running the window at the CA baseline with a named allowance is a better answer than borrowing the US factor for the nearest equivalent class, and what the window buys the account even if it under-performs.
11. This annex says all operating SOPs apply in Canada as written unless a row here overrides them, and it says silence is not an override. Explain what an operator does when they cannot find an override row for a case that plainly feels Canadian, and why that rule is safer than its opposite.
12. Half of this document specifies parameters whose values do not exist. Explain what makes that document useful rather than a deferral, and name the single test that decides whether any given row of it has met that bar.

12. Interface Contracts

Built from the Cross-Reference Ledger's component-34 Interfaces cell and from the rows of the components that name #34 in their own Interfaces cells, cross-checked against Master PPC Data Architecture v1.0 for artifact names, tabs and timing. Every row states what happens when the handoff does not arrive, because that is the clause the corpus was missing. A recorded absence is a valid arrival on every row below; silence is not.

#	Interface	Sender to receiver	Artifact and timing	Non-arrival consequence
I-1	The CA account packet, 5 items	Head of Amazon Operations to P12	Items 1 to 5 per Section 4.12, each against its own register rows, delivered item by item rather than as a package	The annex stays at structure and the Definition Register keeps #34 at PARTIAL. Every consequence below follows from this one row, which is why it is I-1: the component is not blocked on analysis, it is blocked on five deliveries with one named owner
I-2	CA per-product economics (packet item 1)	Head of Amazon Operations with the finance seam, to P3	CA list price, CA landed cost, CA referral rate, CA fulfilment fee and CA returns provision per product, in CAD, monthly with a version stamp	No CA ceiling computes and no CA bid has a legal reference, so every CA campaign runs without a ceiling or does not run. The annex's answer is that it does not run: a CA campaign live without a CA spine is spending against an unknown break-even, and the block records it as ungoverned spend rather than estimating the margin
I-3	CA keyword export with volumes and ranks (packet item 2)	Data Operations to P4	CA search volume and CA rank per term per syntax family, with the matched US family totals available for the ratio, refreshed on the CA volume cadence	No demand ratio, so no scaled floors, so no qualification. The block records every family as unqualified and does not substitute US floors, because substituting them is F1 and produces a marketplace where nothing is ever readable. Structure holds at consolidated and no CA promotion runs
I-4	CA campaign registry as-is (packet item 3)	Data Operations to P5 and P12	The current CA campaign, ad group and target inventory with budgets, bids and states, once at conversion and thereafter through the CA bulk	The consolidation work list cannot be produced, so any cloned CA structure survives undetected and F4 persists. CA spend continues on a structure nobody has reconciled, and the first evidence of the problem is a campaign count that CA volume cannot feed

#	Interface	Sender to receiver	Artifact and timing	Non-arrival consequence
I-5	CA competitor set (packet item 4)	Data Operations to P4 realism and #16	CA competitor price, rating, review count and stockout state for CA conquest and defence candidates, weekly once the feed exists	The O3 credibility check runs without competitor evidence, so CA push realism is assessed on rank and economics alone; conquest stays at flag-obvious cases; and the block records that the mechanical check could not run. The finding downgrades to a flag per gate G8 rather than being dropped
I-6	CA seasonality, 12 months by family (packet item 5)	Head of Amazon Operations to P11	CA demand by family by month, once at conversion and re-issued annually	No CA annual calendar and no CA class factors, so every CA window is a first window and rule T4's timing rule has no CA index to time against. CA pushes are funded on CA volume alone with no seasonal read, which on a fourth-quarter category is the most expensive month to be blind in
I-7	CA price and strikethrough state	CA pricing owner to P8 step 2	Current CA price, current strikethrough state, and any change dated inside the trailing 7 days, weekly before the CA block and same day on a change	Gate G9 cannot fire, so CA CVR and ACOS reads carry no provisional mark and a price-driven swing is read as a PPC result. The block's response is to treat the state as unknown and suppress conversion verdicts for the week rather than assume stability, and the gap escalates to the pricing owner on the second consecutive week
I-8	CA cover, bands and the backup chain	CA inventory owner to P8 step 3	Days to stockout and adequacy band per CA variation with the ranked CA backup, on the CA inventory cadence	Gate G6 cannot be evaluated, which is the one precondition with no workaround: CA spend continues into an unknown cover position and a stockout converts paid momentum into a CA recovery case. An unmapped CA variation additionally bars the re-point response, so a RED landing mid-week has no backup to move to
I-9	The CA deals calendar	CA deals and promotions owner to P11 and #32	Confirmed CA windows with class, dates and CA deal price, at or before the class lead time	No CA window is planned and none arms. A CA window that runs without an entry is unplanned spend at event CPC in a marketplace with no class factor to size it, and it escalates to the deals owner the same day with the hours quantified
I-10	The CA feed refresh	Data Operations to P6	CA bulk, CA search-term rows and CA rank into the CA workbook before the CA block's slot, weekly	The block runs on stale data or does not run. The annex's answer is that it runs and states its staleness: the manifest carries the last refresh date, the CA completeness score falls, and every verdict class the stale groups block is named. A block that skips because the feed is late is F3 wearing a reason

#	Interface	Sender to receiver	Artifact and timing	Non-arrival consequence
I-11	CA staged sets	P5 and P10 to #38	CA creations, budgets, bids, multipliers, negatives and state changes in #38 staging format with the marketplace key, the verdict, the scenario ID and prior and new values in every reason string, on the CA deployment slot	The block decided and the account did not change: verdicts sit undeployed, predictions cannot be scored, and next week's block reads a state that does not match its own log. At CA volume one undeployed set is a large share of the week's total change
I-12	CA verification diff	#38 to P10 step 4	The T+1 diff against the CA staged values, read line by line	Landed values are unverified, so the week's record cannot be trusted and every prediction is scored against an assumed deployment. This is F9, and the block records the gap rather than reporting a clean run
I-13	The CA parameter register and its Appendix A items	P2 to the PPC Lead as v4.0 threshold owner	The register with every pending row, and Appendix A, at every annex increment and immediately on any new gap found in a block	The gaps stay invisible to the only role that can close them, and the annex's pending rows harden into permanent workarounds. This is the row that decides whether this document is a specification or an excuse, and it is the reason Appendix A is addressed to a named role rather than to the file
I-14	The CA marketplace key	P1 to #37, #38 and the Account Rollup	The key on every CA registry row, every staged row and every rollup read, set at instantiation and carried thereafter	CA spend mixes into US totals, platform rules address CA campaigns from US scope, and the family boards compare two auctions that never met. This is F6, and its correction is a restatement of account totals rather than a forward fix
I-15	CA scope declarations to the cross-cutting components	P5 and P10 to #33 and #16	The CA qualified set, the CA contested roots and the CA conquest candidates, with the marketplace key, at instantiation and on every floor refresh	#33 arbitrates a CA root on US composites or not at all, and #16 builds CA target lists from US candidates. Both siblings then hold at their current posture and the CA contest is settled by whichever campaign spends more, which is the one outcome that is never neutral
I-16	Marketplace instantiation	CEO or Head of Amazon Operations to P13	The expansion decision with the marketplace named, before any spend in it	Spend opens in an ungoverned marketplace with no parameter register, no key, no cadence and no packet owner, and the first governing document is written after the first mistake. The whole of P13 exists to make the arrival of this row precede the arrival of the spend

Appendix A. v4.1 Amendment Items

No number below appears as an operating value in this annex. Each item names the gap, its provenance, the operating path used until the amendment lands, and the proposal carried to the threshold owner. Items AM-1 to AM-13 are addressed to the Master PPC Decision Criteria System; AM-14 and AM-15 are addressed to the documents named in their rows, because the gap is not v4.0's to close.

Item	Gap, provenance, operating path and proposal
AM-1	The CA parameter column does not exist in v4.0. Highest priority in this register. SOP-34 v1.0 Section 3 cites a scaling rule in "v4.0's CA column" and its Section 1 states that v4.0 remains the threshold authority "with a CA parameter column pending". The attached v4.0 contains no CA column, no marketplace column and no marketplace-conditional threshold anywhere in Sections 3 to 15. The program's own remaining-work list carries the Canada parameter column as part of the v4.1 amendment batch, which confirms it as pending rather than missing. Operating path: every CA value lives in this annex's P2 register with a status, and no CA threshold is written into any operating document. Proposal: admit a marketplace dimension to v4.0 as a conditional column on every threshold that carries marketplace content, and populate the CA cells from the P2 register at the same release.
AM-2	v4.0 carries no marketplace dimension at all. Every band, floor, benchmark and bar in v4.0 is single-marketplace by construction, and nothing in the document states which marketplace. This is the structural form of AM-1 and it is why a CA value is not a number waiting in an empty cell. Operating path: this annex's three scaling classes at P2 step 5 substitute for the missing dimension by classifying each threshold as transferring, native or ratio-derived. Proposal: state the marketplace assumption explicitly in v4.0 Section 1, and mark each threshold's class in the Dictionary so the classification is v4.0's rather than an annex's.
AM-3	Search-volume floors are stated as absolute numbers. v4.0 Section 6, rule O2 grants objectives on search volume above fixed figures with no marketplace conditioning, and those figures exclude 85.2% of one illustrative CA family's demand at WE-1. Operating path: P4's measured demand ratio scales the floors per syntax family, with both the source figure and the scaled figure recorded in the floor register. Proposal: state the floors as demand-relative with a named scaling basis, or admit a per-marketplace floor table fed from the ratio measurement.
AM-4	Sample-sufficiency bars carry no statement that they are absolute. v4.0 Section 5, gate G2 sets click and impression bars with no marketplace conditioning, which is correct, but the document nowhere says so, and an operator scaling floors in the same session has nothing on the page to stop them scaling these too. Operating path: P4 step 4 requires the floor register to state what did not scale, and F8 detects the error. Proposal: add one sentence to G2 stating that the bars are statistical and do not vary by marketplace, which is the cheapest amendment in this register and prevents the most expensive class of error in it.
AM-5	No rule states the read window for a marketplace whose click rates cannot reach the bars weekly. G2 names a 60-day fallback as an exception; in Canada it is the standing case. Operating path: Section 4.9 runs the 60-day window as standard for CA rate verdicts, computes the read-limit line from G2's own figures, and flags terms below it rather than carrying them as extended tests indefinitely. Proposal: admit a read-limited state to the Dictionary with its own handling, so a term that cannot reach sufficiency has a defined posture rather than an endlessly renewed test.
AM-6	CA TACOS stage bands do not exist. Rule T2 makes stage select the TACOS band, and the bands as carried are US bands. Operating path: none. CA stage posture is declared by the PPC Lead on judgement and recorded, which is a decision without a threshold and is stated as such in the block. Proposal: admit CA stage bands computed from CA contribution, or state explicitly that the US bands apply in CA with evidence.
AM-7	The CA mature band and the CA code-red trip multiple do not exist. #26's P3 protocol trips on the v4.0 multiple of the mature band; the multiple transfers and the band does not exist, so the trip cannot be computed in Canada. Operating path: CA code red is declared by the PPC Lead on judgement with the entry figures recorded, and the protocol then runs unchanged. Proposal: admit the CA mature band with AM-6 in the same release, since the trip is derived from it.
AM-8	No CA Expected CPC reference and no rule for deriving one. Rules L2, M2 and M4 all consume Expected CPC, and v4.0 carries no method for establishing it in a marketplace with no history. Operating path: the reference is measured from delivered CA CPC by syntax family once packet items 2 and 3 land, and every decision consuming it is suspended until then. Proposal: admit a cold-start derivation for Expected CPC in a new marketplace, which is also what P13 will need for the next one.

Item	Gap, provenance, operating path and proposal
AM-9	No CA DSTR or rank-velocity references. Rule M2 sizes every push from DSTR and expected conversion; both are market-specific and neither exists for Canada. Operating path: CA pushes are suspended, which is the largest single decision class this annex currently holds closed. Proposal: admit CA references with packet item 2, and state in M2 that DSTR is marketplace-scoped so the substitution can never be made silently.
AM-10	No CA structure or format targets. Rule M6's splits are US splits and the CA posture is deliberately SP-heavier and consolidation-first. Operating path: the CA target is stated as a direction with a review date rather than a band, per the P2 register row. Proposal: admit marketplace-conditional structure targets, or state that M6 is a US-only check and name what replaces it elsewhere.
AM-11	No CA event class list and no CA class factors. Rule M9 sizes events from measured lift and directs first-time events to a category default; there is no CA default table and the Boxing Day class has no US counterpart. Operating path: the first CA window of each class runs at CA baseline with a named allowance and seeds the factor from its own scoring, per P11 step 3. Proposal: admit a CA class list with AM-6, and state that class factors are marketplace-scoped so a US factor can never be borrowed by default.
AM-12	The CA cadence has no basis in v4.0. Rule V4's cadence is daily plus weekly plus monthly plus quarterly, written for a marketplace whose volumes support a daily read. The CA weekly-only block is correct on the evidence and ungoverned on the page. Operating path: Section 4.6 states the cadence and its reasoning, and the same-day gate responses are carried as exceptions. Proposal: admit a volume-conditional cadence to V4, which P13 will need for every future marketplace whose volume differs again.
AM-13	The human-mandatory review tier keys on a US search-volume figure. Rule P5 places every action on a keyword above a fixed search volume in the human-mandatory tier; almost no CA term reaches it, so the tier empties in the marketplace least able to afford it. Operating path: this annex places the whole CA block in the tier, which is stricter than P5 and never looser. Proposal: express the tier's keyword trigger as an oversight-class membership rather than a raw volume figure, so it follows the scaled floors automatically in every marketplace.
AM-14	Addressed to Master PPC Data Architecture v1.0 and the Command Workbook TEMPLATE v1.0. The architecture specifies one workbook per product with no marketplace key, and the Raw Pull Pack is named for a single marketplace's bulk. There is therefore no schema position for the marketplace value that F6 depends on. Operating path: P1 step 2 writes the key into the workbook header and the T3 registry key column by convention, which works and is not enforced by the instrument. Proposal: admit the marketplace as a first-class key on the workbook, the registry, the manifest and the rollup, per the recommendation already recorded in the program's open decision on marketplace scope.
AM-15	Addressed to the Cross-Reference Ledger. The component-34 row records the blocker and the pending CA parameter column but carries no interface list, so Element 12 above was built from this annex's own procedures and the rows of components naming #34, rather than from a Ledger Interfaces cell. Operating path: the 16 rows at Section 12 are returned with this build's Ledger row so the cell fills from the document that owns the seams. Proposal: adopt the returned Interfaces cell at the next Ledger version.