

SOP-35: Incrementality Testing Operations

PPC Department · Operating SOP · Component #35 of the Operations Definition Register

Version 2.0 · 2026-08-08 · Owner: the incoming operator decides and executes with the PPC Lead reviewing, per PPC Authority Matrix v2.0 Section 1 row "Incrementality tests" · Supersedes SOP-35 v1.0 (2026-07-24) in full · Every tolerance, window, band, gate and verdict lives in the Master PPC Decision Criteria System v4.0 and the Keyword Decision Record Template v3.0; this SOP carries zero local thresholds and executes what those two documents decide.

1. Verdict and Scope

VERDICT: incrementality asks the one question that ultimately justifies mature and defensive spend: what happens to TOTAL orders when this spend steps down. It answers with controlled step-downs reading total (organic plus paid) capture, never paid orders, per v4.0 Section 4 rule X12. Version 2.0 expands v1.0 from 6 procedures to 13, from a 6-row verdict table to the closed vocabulary quoted verbatim, from 2 worked examples to 4, and from 6 failure modes to 12, and adds the operator ramp and the twelve interface contracts. It makes three corrections of substance.

First, all six v1.0 verdict strings are retired: none appears in the Keyword Decision Record Template v3.0 closed vocabulary (Section 5.2).

Second, the v1.0 release seam, "released budget posts to #31 via the SOP-26 P5 pipeline", is retired per Ledger conflict C3 and re-seamed to SOP-26 P12 release verification posting to the instrument layer at T0 zone H and Account Rollup tab A2; at SOP-26 v2.0, P5 is the annual defense-cost read and the release pipeline is P12.

Third, two v1.0 citations of v4.0 objects that do not exist there, the "v4.0 baseline window" and the class-calendar cadences, are retired and logged as v4.1 amendment items rather than carried (AM-2, AM-5).

This SOP governs candidate intake and the test register (P1), the testing calendar (P2), baseline construction (P3), test design and its prediction (P4), ownership and gate clearance with the freeze (P5), rung staging (P6), the capture read at each rung (P7), confound handling and voids (P8), the continue-or-restore call back to SOP-26 (P9), conclusion and the finding record (P10), release integration (P11), scoring and the quarterly backlog review (P12), and the annual defense-read feed (P13). The reader is the operator holding the incrementality scope per the Authority Matrix; a person who has never run a test executes P1-P13 from the page.

1.1 The Two Test Classes

Every test in this document is one of two classes, and the class decides the step mechanics.

Class 1, rung tests, cover rank-sensitive entities: category terms holding rank, and every E2-flagged term at rank 1-3. The step is one #28-derived rung per v4.0 Section 10 rule M3 (a step of about 10% per week, inside the Section 9 rule L2 limits: 50% cap per change, \$0.50 bid floor), because category rank burns while it is measured, and an off-cliff test on a rank-sensitive term destroys the position it is reading (failure mode F2, carried from v1.0).

Class 2, branded zero tests, cover branded defense classes only: these may test to zero directly, because brand-term position does not decay like category rank and the zero read is the honest one for the class. The zero test executes as a dated campaign state action staged per #38 and deployed at #12, never as a bid edit; whether the zero step needs an explicit exemption from the L2 per-change cap is ungoverned in v4.0 and is amendment item AM-6.

Both classes read the same thing: total capture against a filed baseline, over the 2-week window per rung that v4.0 Section 4 rule X12 and Section 8.6 scenario S-D4 fix ("watch total capture 2 weeks"). The 2-week capture window is the one governed duration in this document.

1.2 Why the Test Exists

Strategic Doctrine v1.0 Section 13 states the method in one sentence: step spend down on a term and read whether total (organic plus paid) orders hold, which is the only honest test of whether a dollar is buying incremental demand or renting demand already owned. v4.0 Section 3 family F16 builds the detector: the E2 cannibalization estimate compares expected organic capture at current rank against paid share of sales on the term, and the flag fires when paid share exceeds expected paid share at that rank plus tolerance. Rule X12 closes the loop with the sentence this whole document executes: spend above the discovered floor on flagged terms is by definition rented, not earned. Section 8.1 scenario S-A5 sets the posture consequence on a RENTING mature product: taper-and-verify, with the burden of proof flipped onto the spend. The test is not a research exercise; it is how the burden of proof gets discharged, in dollars, on a date.

1.3 Boundaries With Adjacent Components

Every decision excluded from this SOP names the component that owns it. The Cross-Reference Ledger row for component 35 is the check.

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides the test mechanic (X12), the flag (family F16, E2), the scenarios (S-D4, S-A5, S-Y10), the capture window, the ladder mathematics (M3), the lever limits (L2), every gate, and every verdict. This SOP executes; it never decides. A tolerance or window written into this document instead of v4.0 is a defect. The capture tolerance itself does not exist in v4.0 and is amendment item AM-1, never a local number.
#28 Placement and modifier	Owns rung derivation, ladder mechanics, exit-bid discovery, and the floor register format that T4 implements. This SOP requests rungs and records floors through #28's format; it never derives a step. Referenced per its Ledger row (v2.0, Wave 5A); v1.0 procedure numbers are not cited.
#12 Exact campaign operations	Executes every campaign-level change a test needs: rung deployment confirmations, the zero test's dated state action, restore deployments, and RE-POINT VARIATION on a G6 RED.
SOP-24 P3 (#24)	Owns the confound and void determination and publishes the five confound classes this SOP consumes by name: event or deal window (G10), competitor deal window (S-C5), price test (G9), listing change inside 14 days (family F4), adjacent-variation stockout (S-I6). This SOP hands windows in and receives classifications; it never re-derives the classes. The v1.0 citation "SOP-24 P1" is updated to P3 per the attached SOP-24 v2.0, which states it is the procedure #35 consumes by name.
SOP-26 (#26), three seams	Inbound: its P2 step 6 issues INCREMENTALITY LADDER on an E2 flag, stages the rung, and hands the measurement design and the continue-or-restore call here (its interface I-9). Outbound one: NON-INCREMENTAL findings route to its P12 release verification, which alone posts to the instrument layer. Outbound two: its P5 annual defense-cost read consumes this SOP's P13 evidence package for the cannibalization-adjusted figure.
#32 Deal and event windows	Owns event plans, arming, and settle marks. This SOP applies gate G10 exclusion against #32's declarations and the T5 log; it never schedules around an assumed window length. Referenced per its Ledger row (v2.0, Wave 5A).
#33 Sibling arbitration	Owns the family ownership declaration. Gate G11: no test arms on a shared term without the declaration, because a contested term reads the family's interference, not incrementality.
#38 Bulk deployment	Owns staging format, upload, and the T+1 verification diff. Every rung, zero step, and restore in this document is a staged, reasoned deployment in #38 format, and no window starts until its rung is proven landed.
#37 Automation rules	Owns the rule register and change control. This SOP declares frozen test scope as an exclusion state and reads the rule change log; it never edits a rule. A rule firing on frozen scope voids the affected read and suspends pending #37 review.
#6 V5 prediction scoring	Owns the scoring rubric. This SOP files predictions per rule V1, scores them per V5 inside P1, and hands voids with their class references so a VOID is never filed as a MISS.
Instrument layer (M5 queue)	Releases post to T0 zone H and Account Rollup tab A2 through SOP-26 P12 verification only. Per Ledger conflict C3 there is no #31 procedure to cite, and this SOP never posts directly (failure mode F10).

1.4 What This SOP Does Not Contain

No numeral in this document is local. Where an operating question needs a number that v4.0 does not carry, the question is named in the amendment register at Section 10.3 and the procedure states the gap in place rather than filling it. Eight such gaps are logged at this version, and the first of them, the capture tolerance for X12, is the single number this component is built around. An operator who finds a tolerance, a window length, or a cadence in this document without a v4.0 section heading beside it, or without an AM reference declaring it an input under test, has found a defect and raises it rather than executing it.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. v4.0 states no per-trigger lag table for this component; every lag below is anchored to a governed cadence (Section 14 rule V4's daily, weekly, and quarterly slots) and the absence of finer-grained lags is amendment item AM-7.

Trigger	Signal and detecting artifact	Maximum lag	Procedure and authority
E2 flag handoff	SOP-26 P2 step 6 stages INCREMENTALITY LADDER on a branded or rank 1-3 term. Artifact: the staged rung in #38 format plus the T6 entry with its review date, per SOP-26 interface I-9.	One weekly cycle	P1 intake, P4 design. v4.0 Section 8.6 S-D4 governs the test; the rung is theirs, the read is this SOP's.
Product RENTING flag	T0 zone E flywheel row: ad share rising while organic share falls for 3+ weeks on a Mature product. Artifact: the F2 directional flag read at SOP-26 P1 step 7.	One weekly cycle	P1 batch entry under v4.0 Section 8.1 S-A5: tests on all rank 1-3 and branded terms of the product.
Syntax RENTING	Ad dependency rising against the declared organic-share target on a primary syntax. Artifact: T1 syntax board organic-share columns against the declared target.	One weekly cycle	P1 batch entry under v4.0 Section 8.2 S-Y10: tests on the syntax's rank-held terms; the M8 defense budget binds the syntax until organic share recovers.
Quarterly floor re-test on a flagged term	T4 "Next Step Date" reaches its date on an E2-flagged term. Artifact: T4 columns "Discovered Floor \$", "Ladder State (step / date)", "Next Step Date".	One weekly cycle	P1 intake. v4.0 Section 10 rule M3: on E2-flagged terms the taper doubles as the incrementality test (X12); Section 14 rule V4 sets floor re-tests quarterly.
Structural re-open	A T5 Product Log event on a tested entity: price move, listing / image / A+ change, or competitor entry. Artifact: T5 event types (closed list of 8) with affected keywords.	One weekly cycle	P1 re-open. The old finding described the old world; carried from v1.0 Section 2.
Capture-window expiry	A rung's 2-week window closes. Artifact: the T6 review date on the rung's prediction row.	Same cycle	P7 read, then P9 call where SOP-26 holds the rung. X12 fixes the window; V1 fixes the review date.
Confound signal	A T5 entry, a #32 arming notice, or a #37 rule-log entry lands inside an open window. Artifacts: T5 trend-exclusion flag, the #32 armed-window declaration, the #37 rule change log.	Same day	P8 void protocol through SOP-24 P3. Gate G10 computes the guard end; the read restarts clean.

3. Preconditions and Gates

Every procedure below runs only after these preconditions read current in the T7 manifest. A missing or stale field group suppresses its dependent read class and the record states what it could not see; it never silently degrades a conclusion.

3.1 Manifest Preconditions, With the Live State

The T7 Manifest in the current Command Workbook template reads a packet completeness score of 0.571, below the 85% DEGRADED floor that v4.0 Section 5 gate G8 sets, so conclusions on affected rows cap at flags until the packet recovers, and every design states the manifest state on its face. The field groups this component depends on, with their live status:

Manifest field group (T7)	Live status	Requirement and the suppression consequence
A1 Four-period comparability (L30/P30/B30/L60)	MISSING	Trend context for baselines and reads. While MISSING, every baseline states the single-window limitation on its face and the X8 trajectory read at P7 Branch C runs on the WoW block only.
B8 Row economics (BE chain)	PRESENT (blended)	The T0 zone C named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC. D1 per-variation margin reads PARTIAL, so every ceiling-referencing statement in a design or finding carries the G9 PROVISIONAL tag per the corpus margin-freshness alias.
B11 Market context (SQP shares)	PRESENT	The organic-capture estimate limb of the E2 machinery (expected organic capture from SQP click curves, family F16). One row on benchmark fallback lowers E7 confidence per rule D3 and the record names the source.
D7 Decision log and V5	PRESENT	T6 with review dates and scores. The capture-window expiry trigger reads T6; without it no window closes on schedule and predictions cannot score.

Manifest field group (T7)	Live status	Requirement and the suppression consequence
D9 Daily rank series	PARTIAL (head terms)	The rank axis of every rung test. Rank-sensitive rung tests are legal only where D9 covers the term; elsewhere the rank axis reads the weekly K4 block and the design states the coarser grain.
D8 Competitor depth	MISSING	The competitor deal-window confound class is visible only where T5 carries a manually logged competitor entry. Every read states this coverage gap while D8 is MISSING (Ledger dependency D4, OPEN).
B9 WoW deltas	PARTIAL	The week-over-week block that substitutes for A1 trend context. Fills from cycle 2 per the manifest note.

3.2 Gates in Force

Gate numbering diverges between the Criteria System and parts of the operating corpus; this is Ledger conflict C9, and per its resolution every gate below is named by its v4.0 Section 5 name and number with the corpus alias given at first use. The two numbering schemes are never mixed inside one reason string.

Gate (v4.0 Section 5)	The observable that proves it passed, on a test entity
G1 Validation quarantine	T2 zero-flag column reads Clean on every row in test scope. A quarantine fire mid-test removes the row and voids its window with no confound class; the verdict is QUARANTINE: DATA REFRESH and the refresh task carries an owner and a date.
G2 Sample sufficiency	T2 "Sample State" adequate for any rate read the design consumes. The capture read itself is order-count arithmetic against a baseline, but a CVR statement inside a design or finding obeys the 100-click limb.
G6 Inventory adequacy	T2 "Inventory Band" GREEN or YELLOW on the targeted variation. RED aborts the test same day and routes RE-POINT VARIATION to #12; the OOS is simultaneously a SOP-24 P3 confound for adjacent reads.
G8 Packet completeness	The T7 score stated on the design and every conclusion. DEGRADED (the live state, 0.571) caps conclusions on affected rows to flags per v4.0 Section 12 rule P8.
G9 Price stability (corpus alias: margin freshness)	No price change on the targeted variation inside the trailing 7 days, read from T5 filtered to price moves, or the test start moves past the price window and every CVR-adjacent read carries PROVISIONAL. Both readings of G9 suppress the same class of statement and the record says which fired.
G10 Event exclusion (corpus alias: the G12 event freeze)	No armed or guarded event window overlaps a baseline or read window: event weeks and the 14 days following are excluded from trend reads. Read from #32 declarations and the T5 trend-exclusion flag; the guard end is computed to a date, never eyeballed.
G11 Sibling conflict	Family ownership declared per #33 on any shared term before arming. No declaration, no test; the escalation goes to #33.
G12 MKL hygiene	Every tested term exists in the current MKL with a syntax tag and relevancy score. An untested untagged term routes to MKL maintenance first; a test on an unexamined habit produces a finding about nothing.

4. Step-Level Procedures

Thirteen procedures. P1-P2 run weekly inside the Monday cascade. P3-P6 run per test. P7-P9 run per rung window. P10-P11 run per conclusion. P12 runs weekly for scoring and quarterly for the backlog. P13 runs annually per primary syntax. Time estimates assume the packet is present and the operator has run the procedure before; first executions run roughly double.

P1. Candidate intake and the test register (25 minutes, weekly)

1. Open the test register (schema at step 8, held with the product Command Workbook beside T6) and score last cycle first: read the T6 rows whose review dates land this cycle and record each V5 score as HIT, PARTIAL, MISS or WRONG-DIRECTION. Unscored actions do not accumulate, per v4.0 Section 14 rule V1.
2. Intake source 1, E2 flags: filter T2 to rows whose block K6 E2 estimate fires the cannibalization flag (paid share above expected paid share at rank plus tolerance, v4.0 Section 3 family F16). On branded and rank 1-3 terms the flag arrives through SOP-26 P2 step 6

with the staged rung and the string INCREMENTALITY LADDER; enter the row with the handoff reference. The rung is theirs; the design and the read are this SOP's.

3. Intake source 2, product RENTING: a product whose T0 zone E flywheel row fires RENTING (family F2) enters under S-A5 as a batch: every rank 1-3 and branded term of the product, each its own register row with the S-A5 reference.
4. Intake source 3, syntax RENTING: S-Y10 on a primary syntax enters the syntax's rank-held terms as a batch, and the register records that the M8 defense budget is the binding syntax ceiling until organic share recovers. The ceiling's number is T1 group S7's; this SOP records the constraint and #26 with the instrument layer enforces it.
5. Intake source 4, defensive classes: every class in the SOP-26 P4 defended-set declaration with no capture evidence on file in the trailing year enters as a candidate, because SOP-26 P5 subtracts E2-flagged spend from its numerator and an unmeasured protection price is an assumption with a budget line. This is the v1.0 candidate class carried forward as the companion to SOP-26's F4.
6. Intake source 5, quarterly floor re-tests: T4 "Next Step Date" reaching its date on an E2-flagged term enters per the M3 doubling sentence (Section 10 rule M3: on E2-flagged terms the taper doubles as the incrementality test).
7. Intake source 6, structural re-opens: a T5 event (price move, listing / image / A+ change, competitor entry) on a concluded entity re-opens its test with the T5 reference, because the old finding described the old world. Carried from v1.0 Section 2.
8. Register row schema: entity (term or campaign class); class (rung or zero); source reference (scenario ID, SOP-26 handoff, T4 date, or T5 event); entry date; state (QUEUED / DESIGNED / ARMED / READING / VOID-RESTART / CONCLUDED-UNACTED / CLOSED); owner; review date. A row without a source reference does not enter: the register is flag-driven and trigger-driven, never curiosity-driven. The state names are register vocabulary, not verdict strings, and their formal admission is amendment item AM-8's companion note.

P2. The testing calendar and slotting (15 minutes, weekly, same pass)

1. Slot QUEUED rows against the two sequencing rules. One variable per entity: an entity under test freezes everything else on it (P5), so a second test on the same entity queues behind the first. One test per family: two open tests sharing an A1 Family Boards grid confound each other and answer nothing twice (v1.0 F3 carried; the family grain is the Account Rollup A1 grid).
2. Priority inside the queue: SOP-26 handoffs first, because a staged rung is already waiting on a design; then S-A5 and S-Y10 batches; then M3 doublings on their T4 dates; then defensive-class candidates; then re-opens. Ties break by weekly spend at the tested level, descending, because the finding value is proportional to the spend under question.
3. Governed anchors: floor re-tests run quarterly per v4.0 Section 14 rule V4, and protected terms review quarterly per Section 8.6 S-D6. Ungoverned cadences: v1.0 carried a class calendar (semi-annual on branded classes, annual per term cohort) tracing to no v4.0 rule; those cadences are retired and logged at AM-5. Until v4.1 lands a cadence, entry is by flag, batch scenario, quarterly M3 date, defensive-class gap, or structural re-open, and the P12 backlog review is the guard against never-testing.
4. Nothing tests continuously: an entity past its designed windows without a conclusion is failure mode F6, and the calendar closes what curiosity keeps open. Carried from v1.0.
5. Write the slot decisions into the register with dates. The register is the artifact the P12 review audits.

P3. Baseline construction (20 minutes per test)

1. The baseline artifact files before any rung deploys and carries: total orders per week on the entity from all sources; the paid split (orders, spend, clicks, CPC) by campaign from the T2 child rows (block K5 columns); the organic rank series (block K4, with the daily-series link where D9 covers the term); share context from block K3 (SQP brand purchase share, benchmark source named per rule D3); and the economics snapshot from the T0 zone C named cells with the margin table version and G9 state.
2. Total-orders source by class. Product-scope zero tests read weekly product orders from the SellerBoard figure in T0 zone B: the validated number, never a console estimate. Term-scope rung tests read term total capture: paid orders from the K5 cross-type totals plus the organic estimate from the E2 machinery (expected organic capture at current rank from SQP click curves, family F16), with the estimate named as an estimate and its benchmark source stated.
3. The clean-window check: the baseline window contains no event week and no post-event guard. Read T5 for deal windows, cross-check #32 settle marks, and apply gate G10 (event weeks plus the 14 days following are excluded from trend reads). A baseline built on an event week inherits the distortion permanently; filing one is failure mode F7. Where recent history is wall-to-wall events, the baseline waits.
4. Baseline length: v4.0 defines no baseline window for incrementality tests, and v1.0's citation of "the v4.0 baseline window" traced to nothing; it is retired and logged at AM-2. The baseline states its own window on the artifact as an input under test, built from whole validated weeks. Live constraint: A1 four-period comparability reads MISSING, so every baseline states the single-window limitation on its face.
5. Prediction precondition: no rung deploys until the P4 prediction is filed against this baseline. A baseline without a prediction is an observation, not a test.

P4. Test design and the prediction (25 minutes per test)

1. One variable. The design names exactly one campaign or one term class as the stepped variable; everything else on the entity freezes for the duration. Two moving pieces produce an unattributable read.

2. Class selection per Section 1.1: rank-sensitive entities design as rung tests at the M3 step size inside the L2 limits; branded defense classes may design as zero tests; everything else designing to zero is F2, restored on discovery.
3. The capture window per rung is 2 weeks, per v4.0 Section 4 rule X12 and Section 8.6 S-D4.
4. The prediction, per v4.0 Section 14 rule V1 and the Section 13.2 record contract: named metric (total orders on the entity), direction (holds), magnitude (the design's stated hold bound, plus a stated rank bound on rung tests), horizon (the 2-week window), review date. The hold bound is the prediction's magnitude and it is stated on the design as an input under test: v4.0 sets no capture tolerance for X12, and that gap is AM-1, the highest-priority item in the amendment register. The bound is never presented as a threshold, never reused across designs without restatement, and every conclusion states the bound it was read against. A conclusion citing a bound its design never stated is failure mode F12.
5. The design packet: entity and class; baseline reference; the stepped variable with prior and proposed values; the rung schedule with dates; the prediction; the confound watch list (the five SOP-24 P3 classes, named); the restore rule quoted from X12 (on a capture drop, restore and record the defended value); and the abort conditions (G6 RED, G1 quarantine, G4 eligibility flip).
6. Zero-test specifics: the step is a dated campaign state action (budget to zero or paused for the window), staged per #38 and executed at #12, with the resume value recorded so the restore is mechanical. It is a designed test step, not PAUSE (LEVERS EXHAUSTED): levers are not exhausted, the test is running.

P5. Ownership, gate clearance, and the freeze (15 minutes per test)

1. G11 first on any shared term: family ownership declared per #33 before arming. Two siblings testing or defending the same term is a portfolio decision, and a contested term reads the family's interference rather than incrementality. No declaration, no arming; the escalation goes to #33 and the flag owner's rung holds where it is.
2. Read the remaining gates from the T0 zone D strip and the T2 flag columns, per Section 3.2: G1 Clean on every row in scope; G2 for any rate statement; G6 GREEN or YELLOW; G8 stated on the design; G9 clear or the start date moves past the price window; G10 clear across every planned window; G12 MKL current on the term.
3. Declare the freeze: no bid, budget, modifier, negation, or structure change on the frozen scope for the duration, other than the designed steps. The freeze list writes into the register row and the T6 entry.
4. Exclusion notice to #37: declare the frozen scope to the automation rule register so no platform rule steps a bid inside the test. This SOP maintains the state feed only and never edits a rule. A rule change touching frozen scope mid-test voids the affected read and the rule suspends pending #37 review.
5. Arm the register row (state ARMED) with the first rung date.

P6. Rung staging and deployment (10 minutes per rung)

1. The rung value arrives from #28 mechanics; for zero tests the rung is the designed state action. This SOP never derives a step: staging a rung without the derived value is inventing a threshold, and the window moves out instead.
2. Stage per #38 with the four-segment reason string: the verdict string verbatim (INCREMENTALITY LADDER), the v4.0 scenario reference (S-D4, or S-A5 / S-Y10 for batch entries, or the M3 doubling), this SOP and procedure number, and the test register reference. A staged row missing its scenario reference does not deploy.
3. Deploy on the rung date and read the #38 T+1 verification diff the next day. A rung that did not land did not start its window: re-date the window from the landed day. An unlanded rung read as if live is a staging failure, not a test result.
4. Write the rung into T6 with the V1 prediction fields, and move the register row to READING.

P7. The capture read at each rung (20 minutes per read)

1. At the window close, compute total capture against baseline: total orders on the entity across the window's validated weeks, beside the baseline figure, as a level and as a percentage. Paid orders and spend read alongside as the mechanism, never as the answer: paid orders falling is the test working, and a finding citing paid metrics as the result is failure mode F1.
2. On rung tests, read the rank axis across the window (K4, daily link where D9 covers the term). A capture hold with rank sliding past the design's stated rank bound is a hold on borrowed time and classifies as a drop for the restore rule.
3. Run the P8 confound sweep before classifying anything. A read classifies only after the SOP-24 P3 determination returns clean.
4. The decision tree, four branches:

Branch	Condition	Resolution
A. Capture holds	Total capture inside the stated bound and rank inside its bound.	The spend at the tested level was non-incremental. The rung holds and the ladder proceeds to the next rung, per X12: the ladder continues and the floor is recorded. The held rung registers in T4 through #28's format with the test reference.

Branch	Condition	Resolution
B. Capture drops	Total capture past the bound, or rank past its bound.	Incremental spend confirmed. Restore one rung per X12: if capture drops, restore and record the defended value. The incremental floor records at the restored level in T4 with the test reference, and the row's standing verdict becomes DEFEND AT FLOOR.
C. Ambiguous	Inside the bound but deteriorating across the window (X8 trajectory).	v4.0 defines no extension allowance, and v1.0's single-extension rule traced to nothing (AM-4). The design's own review date governs; the ambiguity states on the record; where the trajectory runs against the prediction, the conservative execution is Branch B's restore, per X8's one-step-conservative rule.
D. Void	A confound inside the window per P8.	The read restarts clean after the G10 guard; the register row states VOID-RESTART with the class. The rung itself does not move: a void has no evidence in either direction.

5. Score the rung's prediction in T6 at the read. A VOID is a VOID, never a MISS: misfiling teaches the calibration series that a sound assumption was wrong. SOP-24 P3 carries the rule; this SOP consumes it.

6. Each subsequent rung repeats P6-P7 until a Branch B restore, a designed stop, or the L2 floor value; the ladder never steps below it.

P8. Confound watch and the void protocol (10 minutes per read, plus same-day handling)

1. The watch runs across every open window, not only at the close: T5 entries in the five classes SOP-24 P3 publishes; the #37 rule log against frozen scope; #32 arming notices.
2. Live constraint: D8 competitor depth reads MISSING, so the competitor deal-window class is visible only where T5 carries a manually logged competitor entry, and every read states the coverage gap while D8 is MISSING.
3. On a confound inside an open window, the read is VOID per SOP-24 P3, which owns the determination: this SOP hands the window references and receives the classification. The rung does not move on a void; position restores only on Branch B evidence.
4. Compute the guard end explicitly (the event week plus the 14 days per G10) and re-date the window from the guard end. The register row states VOID-RESTART with the class and the T5 references.
5. Hand the void to #6 where a prediction was riding the window, so the score files as VOID.
6. Never hand a confounded read forward as clean: SOP-26 interface I-9 names it the worse failure, because the consumer cannot detect the contamination. Failure mode F8, with the account's dated incident behind it (WE-3).

P9. The continue-or-restore call (15 minutes per call)

1. For tests riding an SOP-26 rung (the P2 step 6 handoff or the P3 quarterly doubling), the call returns to SOP-26 at the end of the 2-week capture window, per its interface I-9: CONTINUE (capture held; the ladder proceeds and SOP-26 records the floor at the held rung) or RESTORE (capture dropped; SOP-26 restores and records the defended value).
2. The call carries: the window dates, the baseline reference, the capture arithmetic, the rank axis, the bound it was read against, and the SOP-24 P3 determination.
3. On a confounded window the call is the void with its class references, never a reading. SOP-26 holds the rung and the floor is not recorded, per its I-9 clause; the re-dated window returns here.
4. The call lands the same cycle the window closes, on the V4 weekly anchor. A late call leaves the counterpart's rung undated past its review, which is its F9 exposure; the absence of a governed per-call lag is inside AM-7.

P10. Conclusion and the finding record (20 minutes per test)

1. A test concludes with the full rung history and exactly one finding class: NON-INCREMENTAL AT THE TESTED LEVEL (capture held at the final held level; the finding value per week is the pre-test spend minus the held-level spend) or INCREMENTAL FLOOR (a Branch B restore; the floor value with its test reference). The class names are register vocabulary, not verdict strings; the executing verdicts are Section 5's.
2. The finding files with: entity; class; rung history with dates and T+1 confirmations; capture arithmetic per rung; the bound(s) read against; the confound record; the finding value in dollars per week; and the acted consequence (the P11 routing or the T4 record).
3. INCREMENTAL FLOOR: the floor records in T4 ("Discovered Floor \$", "Ladder State (step / date)", "Next Step Date") through #28's register format, with a context note naming the test and the market state, because a floor without context is undiagnosable at the next slip. The row's standing verdict is DEFEND AT FLOOR, and the next quarterly date sets per V4.
4. NON-INCREMENTAL: the finding value routes per P11. The entity's standing state after conclusion is its stage verdict, commonly HOLD; DEFEND LIGHTLY AT RANK on a held mature term.

5. Restoration verification where a restore executed: the #38 T+1 diff on the restore, plus the next validated week's rank read against the pre-rung position. A restore without verification is a bid with a label.
6. The register row moves to CONCLUDED-UNACTED until the P11 routing or the T4 record lands, then CLOSED. CONCLUDED-UNACTED is a dated state the P12 review hunts, not a resting place.

P11. Release integration (10 minutes per finding)

1. A NON-INCREMENTAL finding is a candidate release, not a posting. It hands to SOP-26 P12, whose verification (rank and share axes, the withheld-portion split) decides what is real; the verified difference posts to the instrument layer at T0 zone H and Account Rollup tab A2. Per Ledger conflict C3 there is no #31 procedure to cite, and v1.0's routing through the SOP-26 P5 pipeline is retired: at SOP-26 v2.0, P5 is the annual defense-cost read and the release pipeline is P12.
2. This SOP never posts to the queue directly, and the reason is arithmetic rather than etiquette: SOP-26's recorded harvest case computes a raw monthly release of which 41.8% was money buying shelf presence, and an unverified posting over-subscribes the M5 queue and surfaces later as a failed push. Failure mode F10.
3. For findings on entities outside the mature defended set, the finding routes through the product's harvest-read owner with the same verification demand: the M5 queue funds from landed savings only, per Strategic Doctrine v1.0 Section 12.
4. Write the routing into T6 and the register row.

P12. Scoring, calibration, and the quarterly backlog review (45 minutes quarterly, plus weekly scoring inside P1)

1. Weekly: score due predictions per V1 inside P1 step 1. Rule-level accuracy accrues per V5: a design pattern below 50% accuracy over 10 or more scored predictions enters review, and the review produces an amendment, a change proposal, or a documented defense.
2. Calibration: a MISS on a hold prediction (the design predicted non-incremental and the spend was real) is the calibration doing its job on the next candidate ranking. Carried from v1.0's second worked case; the register's priority order re-weights toward classes whose predictions miss.
3. Quarterly, in the V4 quarterly framework-review slot: sweep the register for CONCLUDED-UNACTED rows and for entities past their designed windows without a conclusion (F6). Each unacted finding quantifies in dollars per week and the total is the review's headline number. F5 escalates here with the unacted findings quantified, carried from v1.0; WE-4 computes a quarter of theater at \$2,099.50 on three rows.
4. Reconcile the register against T4 (every INCREMENTAL FLOOR finding has its floor row) and against T0 zone H (every routed release reached SOP-26 P12 or is dated in its queue).

P13. The annual defense-read feed (30 minutes annually per primary syntax)

1. SOP-26 P5 subtracts E2-flagged spend from the defense-cost numerator and states the adjusted figure beside the raw one, with the incrementality reads as the evidence. This procedure packages that evidence per defended syntax: the trailing-year findings on its terms (class, tested level, value, dates), the current T4 floors with test references, and the open register rows.
2. The package lands before the P5 read date; the trigger is SOP-26's Element 2 row and the artifact is this SOP's. A defended syntax with no capture evidence on file re-enters P1 as a candidate class, closing the loop v1.0 named: protection whose price must be proven, not assumed.

5. Decision Points: Verdict Mapping

All verdict issuance, thresholds and scenario logic live in v4.0, and every string below is quoted verbatim from the Keyword Decision Record Template v3.0, tab "Bands & Verdicts Reference", the closed vocabulary of 25 strings. v4.0 Section 13.1 carries an 18-string list under the same closed-vocabulary rule; the divergence is the known count conflict recorded across prior build sessions, and where a KDR string has no identical Section 13.1 string the cross-map states it. This table maps the strings that land on or arise from tested scope; the remaining closed-vocabulary strings do not intersect this component and their owners are named in the Ledger.

Verdict (KDR Template v3.0, closed vocabulary)	Execution in this SOP
INCREMENTALITY LADDER (v4.0 Section 13.1 cross-map: no identical string; nearest is DEFEND AT FLOOR, defined "Branded/held term at discovered floor; incrementality-tested where flagged")	The entry verdict. Issued by the flag owner (SOP-26 P2 step 6 on S-D4; batch entries under S-A5 and S-Y10), never by this SOP. This SOP receives it, designs at P4, reads at P7, and returns the P9 call.
DEFEND AT FLOOR	The Branch B outcome. The restored value records as the defended floor in T4 with the test reference; this SOP issues the restore instruction inside the P9 call or the P10 finding, and the rung owner (#28 mechanics, #12 deployment) executes.
TAPER: RANK ACHIEVED, STEP BIDS DOWN (v4.0 Section 13.1 cross-map: TAPER TO FLOOR (POST-OBJECTIVE))	The state a Branch A continuation rides where the test doubles an M3 taper: the ladder proceeding is this string's territory and its owner executes; this SOP supplies the capture evidence that keeps it legal on an E2-flagged term.

Verdict (KDR Template v3.0, closed vocabulary)	Execution in this SOP
HOLD; DEFEND LIGHTLY AT RANK	The standing state after a NON-INCREMENTAL conclusion on a held mature term: no change, the week-over-week block reads next cycle. Boundary row; SOP-26 owns the posture.
QUARANTINE: DATA REFRESH (v4.0 Section 13.1 cross-map: FLAG: DATA REPAIR)	A G1 fire inside test scope removes the row and voids its window with no confound class; the refresh task carries an owner and a date. Data quality beats the test calendar, per v4.0 Section 12 rule P8.
RE-POINT VARIATION	A G6 RED mid-test aborts the test the same day; execution at #12 per rule L9. The OOS is simultaneously a SOP-24 P3 confound class for adjacent reads.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	Boundary row. The click-quota discipline governs traffic tests (Doctrine Section 13); an incrementality window is dated per X12, not click-counted. This string never issues from this SOP, and conflating the two test types is a design defect.
STOP-LOSS: CONCEDE OR DEFER	Boundary row. A test is not a push and takes no stop-loss; a live push inside proposed test scope is a one-variable violation and the test waits or the scope narrows. Push economics belong to their owners.
PAUSED-PENDING (REASON + RE-ENTRY)	Boundary row. An entity blocked from testing (a chronic conversion failure per rule D2, a fix running per rule D8) waits under its own O5 state, and the test re-opens on the O5 re-entry criterion.

5.1 Campaign-Level Verdicts

Campaign rows carry the subset vocabulary quoted from the KDR Template v3.0 "Campaign Decision Record" tab: FUND / HOLD / TRIM BUDGET · RAISE / CUT MULTIPLIER · NEGATION PASS · CONSOLIDATE · RETAG OBJECTIVE · PAUSE (L7) · ARCHIVE. Inside a frozen test scope, RAISE / CUT MULTIPLIER and TRIM BUDGET are barred except as designed steps, NEGATION PASS defers to the window close unless it is a #29-prescribed pre-load riding a structure event outside the frozen scope, and the zero test is a designed budget or status action rather than PAUSE (L7), because levers are not exhausted while the test is running.

5.2 Verdict Strings Retired at This Version

SOP-35 v1.0 Section 5 issued six strings that appear in neither the KDR Template v3.0 closed vocabulary nor v4.0 Section 13.1: OPEN TEST, RUNG HOLDS / NEXT RUNG, RESTORE + RECORD FLOOR, EXTEND WINDOW (once), VOID READ, and RELEASE TO QUEUE / DEFEND FLOOR. A closed vocabulary is what makes the decision log queryable and predictions scoreable per verdict type, and a locally minted string is invisible to both. All six are retired. The operations they named survive as procedures rather than as verdicts: opening a test is P3-P5 under the inbound INCREMENTALITY LADDER, the rung branches are P7's tree, the restore is Branch B issuing DEFEND AT FLOOR, the extension had no source and is AM-4, the void is the P8 protocol consuming SOP-24 P3, and the release is the P11 routing into SOP-26 P12, which is a handoff rather than a verdict.

6. Worked Examples

All four examples are built on the standard fictional training base, because the anonymized live product economics spine is not attached to this session: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Max Acceptable 26.3%, CVR 5.2% unless a term's own rate is stated. Values are illustrative and internally consistent. Live executions read the T0 zone C named cells instead, and the arithmetic is identical.

WE-1. A branded zero test, clean: the finding routes, it does not post

Setup. A brand exact defense campaign spends \$84.00 per week: CPC \$0.75, 112 clicks per week, 31 paid orders per week (a 27.7% branded CVR, which is why brand terms are the classic candidate). Baseline per P3: product total orders 212 per week from the SellerBoard figure in T0 zone B, brand-term organic position 1 held, no event week or guard inside the window, single-window limitation stated because A1 reads MISSING.

Design per P4. Class 2, zero test: the branded class may test to zero (Section 1.1). The step is a dated campaign state action staged per #38 and executed at #12, resume value \$0.75 recorded. Window: 2 weeks per X12. Prediction filed per V1: total orders hold within the design's stated bound of a 3.0% fall (an input under test pending the AM-1 tolerance), brand-term position holds at 1.

Read per P7. Weeks 1-2 at zero: paid orders fall by 31 per week, which is the mechanism working, not the answer. Total orders read 210 per week: a fall of 2 orders, which is $2 / 212 = 0.94\%$, inside the stated 3.0% bound. Position holds at 1. The brand's organic presence caught 29 of the 31 paid orders, an organic absorption of $29 / 31 = 93.5\%$. The P8 sweep returns clean.

Finding per P10-P11. NON-INCREMENTAL AT THE TESTED LEVEL: finding value \$84.00 per week, \$4,368.00 per year at 52 weeks. The finding routes to SOP-26 P12 for verification on both axes before any posting; nothing posts from this SOP. Prediction scores HIT in T6. The lesson: the test's product is a routed finding with an acted consequence, and the routing is one line because the boundary is one line.

WE-2. A rank-3 rung test that breaches: the spend was real, and the floor is the proof

Setup. A category term held at organic rank 3 fires the E2 flag (family F16: paid share above expected paid share at rank 3 plus tolerance). Recorded floor from a completed ladder: \$1.90. Term CVR 5.2%, so Max Profitable CPC = $28.00 \times 0.052 = \$1.46$ (K6), and the bid sits \$0.44 per click above its profit ceiling, which is exactly what the flag is asking about. Baseline: 41 total orders per week (14 paid, 27 organic estimate per K3), 269 paid clicks per week ($269 \times 0.052 = 13.99$ paid orders), paid spend $269 \times 1.90 = \$511.10$ per week.

Design. Class 1, rung test. One M3 rung from #28: $1.90 \times 0.90 = \$1.71$, a 10.0% step, inside the L2 per-change cap and above the \$0.50 floor. Prediction: total orders hold within the design's stated bound of a 5.0% fall, rank holds at 3-4; window 2 weeks; the bound is an input under test per AM-1.

Read. Week 2: total orders 37.3 per week, a fall of 41.0 to $37.3 = 3.7$ orders, which is $3.7 / 41.0 = 9.0\%$, past the 5.0% bound; rank sliding toward 5, past the rank bound. Both limbs of Branch B. The P8 sweep returns clean, so the read stands.

The arithmetic that decides it, at constant clicks (the conservative statement). The rung's spend saving: $269 \times (1.90 \text{ minus } 1.71) = 269 \times 0.19 = \51.11 per week. The contribution lost: $3.7 \text{ orders} \times \$28.00 = \$103.60$ per week. The rung loses \$52.49 per week net, and the ratio $103.60 / 51.11 = 2.03$ says the tested spend was buying real orders at roughly two dollars of contribution per dollar of bid saving.

Resolution per Branch B. Restore one rung to \$1.90; the incremental floor records at \$1.90 in T4 with the test reference and a context note; the standing verdict is DEFEND AT FLOOR. The \$0.44 per click premium over the \$1.46 ceiling is now a measured price, not an assumption: stepping below \$1.90 costs \$2.03 in contribution per \$1.00 saved. Restoration verifies on the #38 T+1 diff and the next validated week's rank read. Prediction scores MISS, which is the calibration doing its job on the next candidate ranking, carried from v1.0.

WE-3. The confounded window: a void handled, and the incident behind failure mode F8

Setup. A defended head term with recorded floor \$2.24 reaches its quarterly T4 date; the term carries an E2 flag, so the SOP-26 P3 probe doubles as the incrementality test per the M3 sentence. The rung: one step to \$2.02 (a 9.8% step). The capture read rides the same 2-week window as the rank hold prediction.

The confound. Inside the window, T5 carries two competitor deal-window entries on top-10 holders. Per P8 the windows hand to SOP-24 P3, which classifies the confound (competitor deal window, the S-C5 class) and returns VOID for both the rank read and the capture read. The rung reverts on the counterpart's slip discipline; the capture read is not attributed, because deal-priced competition contaminates exactly the demand the test measures.

The handling. The P9 call returns the void with its class references and T5 links, never a reading: SOP-26 holds the rung and no floor records, per its I-9 clause. The guard end computes per G10 (event weeks plus 14 days), the window re-dates, and the register row states VOID-RESTART. The prediction files as VOID, not MISS.

The clean re-run. Both deal windows closed, the probe re-runs at \$2.02 and capture holds inside the stated bound: Branch A, the ladder continues, and the floor records at \$2.02 with a context note naming the first window's confound. The premium over the \$1.46 ceiling moves from $2.24 \text{ minus } 1.46 = \0.78 to $2.02 \text{ minus } 1.46 = \0.56 per click; at 30 maintenance clicks per week that is \$23.40 falling to \$16.80, a saving of \$6.60 per week, \$343.20 per year on one term.

The incident. This is the account's one dated incident touching this component, recorded in the SOP-26 v2.0 rebuild's worked cases: the read was voided and re-run clean, and the interface clause "a confounded read handed over as clean is worse than no read" exists because of it. Failure mode F8 carries it.

WE-4. The quarterly backlog: what test theater costs, in dollars per week

Setup. The P12 quarterly review sweeps the register and finds three rows aged in CONCLUDED-UNACTED across a full 13-week quarter. Row 1: the WE-1 class branded finding, \$84.00 per week non-incremental, never routed to SOP-26 P12. Row 2: a defensive complement-PAT class finding, \$47.50 per week non-incremental at the tested level, unrouted. Row 3: an INCREMENTAL FLOOR finding whose T4 record was never written, leaving the bid at the pre-test \$2.10 against a proven defended value of \$1.85 on 120 clicks per week: an unproven premium of $0.25 \times 120 = \$30.00$ per week re-accumulating.

The arithmetic. Weekly theater cost: $84.00 + 47.50 + 30.00 = \161.50 . Over the quarter: $161.50 \times 13 = \$2,099.50$. The two unrouted releases ($84.00 + 47.50 = \$131.50$ per week) are the quiet direction: the M5 queue understates its funding supply and queued pushes stay queued while the money exists, which is the same failure direction SOP-26's release interface names. The unwritten floor is the loud direction eventually: \$30.00 per week of spend with a proof on file that it buys nothing.

Resolution. Act the backlog: route rows 1-2 to SOP-26 P12, write row 3's T4 record and restore the bid through #28 mechanics and #38 staging. Escalate per F5 with the \$2,099.50 quantified. The lesson, carried from v1.0 in its own words: the test's product is the moved budget, and findings never left un-acted are the point.

7. Failure Modes and Escalation

Twelve modes. Each carries its detection signal, its response, and the account-history pattern or incident that motivated it where one exists. F1-F6 carry forward from v1.0 with their patterns; F7-F12 are new at v2.0.

Failure	Detection signal	Response, and the pattern or incident behind it
F1. Paid-orders read	A finding, call, or conclusion citing paid metrics as the answer.	Re-read on totals. Paid orders falling is the mechanism, not the result: it is the test working. Carried from v1.0 as the component's first-listed pattern.
F2. Off-cliff on a rank-sensitive term	A zero test outside the branded class, or a staged step past the L2 per-change cap.	Restore immediately and estimate the rank cost the same day: rank burns while it is measured, which is why the rung ladder exists. Carried from v1.0 with its same-day escalation.
F3. Concurrent family tests	Two open register rows sharing an Account Rollup A1 family grid.	Sequence them; confounded tests answer nothing twice. Carried from v1.0; the family grain is the A1 grid.
F4. Event contamination	A rung or read overlapping a T5 deal window, a #32 settle mark, or the G10 guard.	Void per P8 and restart at the computed guard end. Carried from v1.0, now executed through SOP-24 P3 rather than locally.
F5. Test theater	CONCLUDED-UNACTED rows aging; findings with no T4 record and no P11 routing.	Act the backlog and escalate quarterly with the unacted findings quantified. Carried from v1.0; WE-4 prices one quarter at \$2,099.50 on three rows.
F6. Permanent testing	An entity past its designed windows with no conclusion in the register.	Conclude on the evidence held; the calendar closes what curiosity keeps open. Carried from v1.0.
F7. Dirty baseline	A baseline window overlapping a T5 event entry or a G10 guard.	Rebuild before any rung deploys. A distorted baseline poisons every downstream read of the test, permanently, because every branch compares against it. New at v2.0.
F8. Confounded read passed as clean	A call or finding with no P8 sweep on the record, or a T5 event inside a window the record does not carry.	Void retroactively, notify every consumer, re-run outside the confound. The dated incident: a quarterly probe on a defended head term whose capture window contained two competitor deal windows; the read was voided and re-run clean (WE-3), and SOP-26 I-9 names the pass-as-clean variant the worse failure. New at v2.0, incident carried.
F9. Restore skipped after a breach	A Branch B read with the bid still at the rung one cycle later.	Restore now, verify per P10 step 5, and log the delay's contribution cost from the breach arithmetic (WE-2's rate: \$52.49 per week net on that term). New at v2.0.
F10. Direct queue posting	A T0 zone H entry sourced to this SOP with no SOP-26 P12 verification evidence attached.	Withdraw and route through P12. The over-subscription arithmetic is SOP-26's recorded harvest case, where 41.8% of a raw release was money buying shelf presence. New at v2.0; replaces the retired v1.0 #31 seam.
F11. Shared-term test without a declaration	A register row arming on a term whose T2 family conflict state is undeclared.	Hold arming and escalate to #33 per G11. A contested term reads the family's interference, not incrementality. New at v2.0.
F12. Bound drift	A conclusion citing a bound its design never stated, or a design reusing a prior bound without restatement.	Re-classify against the design's own bound or void the read. The AM-1 gap is escalated, never filled per operator: an improvised tolerance is a local threshold wearing a prediction costume. New at v2.0.

Escalation. F2 escalates the same day with the rank cost estimated. F8 escalates on discovery to every consumer of the affected read, because a contaminated conclusion may already be executing. F10 escalates on discovery: queue integrity is account-wide. F5 escalates at the quarterly review with dollars quantified. Any mode recurring in consecutive cycles on the same scope escalates to the PPC Lead with the register history attached. Everything in the human-mandatory tier (every test arming, every zero test, every restore instruction, every finding routing, and every void notification) is confirmed per the Authority Matrix before deployment.

8. Outputs and Artifacts

Gates clear with artifacts, not assertions. Each row names the record, where it is stored, how long it is held, and the downstream consumer. No v4.0 rule sets a retention period for any of these records; the rows state the instrument's own append-only life and the gap rides in AM-8, which duplicates items already logged by prior sessions and folds once at v4.1.

Event	Named record	Storage location	Retention	Downstream consumer
Intake and slotting (P1-P2)	The test register: entity, class, source reference, state, dates, owner	Register block held with the product Command Workbook beside T6	Append-only log life	P12 backlog review; the weekly cascade; SOP-26 handoff reconciliation
Baseline (P3)	Baseline artifact: totals, paid split, rank series, share context, economics snapshot, clean-window statement	Filed with the register row; sources T0 zone B/C, T2 K3-K5	Life of the test chain	Every P7 read; #6 scoring context
Design (P4-P5)	Design packet with the V1 prediction, freeze list, confound watch list, abort conditions	Register row plus the T6 prediction entry	Life of the test chain	#38 staging; #37 exclusion state; the checkpoint holder
Rung (P6)	Staged set with the four-segment reason string; the T+1 diff	#38 format; T6	Append-only log life	#38 verification; the window clock
Read (P7-P8)	Capture read with branch classification; void records with class and T5 references	T6; register row state	Life of the test chain	P9 call; #6 (voids); SOP-24 P3 round-trip
Call (P9)	The continue-or-restore call with window dates, arithmetic, bound, determination	Returned on the SOP-26 I-9 channel; copy in T6	Life of the test chain	SOP-26 P2/P3 rung execution and floor recording
Conclusion (P10)	The finding record: class, rung history, values, acted consequence	Register row; T4 for floors through #28's format	Life of the term in the register	SOP-26 P12 (releases); #28 and the T4 register (floors); P13 package
Routing (P11)	The release handoff with the finding value and evidence	T6; SOP-26 P12 intake	Append-only log life	SOP-26 P12 verification; T0 zone H and Account Rollup A2 after verification
Review (P12)	The quarterly backlog review with the theater total in dollars per week	T6, filed against the register	Append-only log life	The PPC Lead; the v4.1 amendment path where a gap recurs
Defense feed (P13)	The evidence package per defended syntax	Filed against T1 group S7's syntax	Annual, superseded by the next package	SOP-26 P5 cannibalization-adjusted figure

9. Skill-Conversion Block

Block	Specification
Input bindings	T0 zone B (weekly validated spend, sales, orders from the SellerBoard read), zone C economics named cells AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC with the margin version and G9 state, zone D gates strip, zone E flywheel row. T1 organic-share columns and group S7 "Defense budget \$". T2 blocks K3 (SQP shares, benchmark source), K4 (rank series, daily link, pre-OOS baseline), K5 (cross-type totals, WoW block, zero-flag), K6 (E2 estimate, Max Profitable CPC, Expected CPC), K2 (inventory band, event mode, family conflict state). T4 "Discovered Floor \$", "Ladder State (step / date)", "Next Step Date". T5 event types with the trend-exclusion flag. T6 decision log with review dates and V5 scores. T7 statuses per field group. The #32 armed-window declarations. The #37 rule change log against frozen scope. The #38 T+1 diff. The SOP-26 I-9 handoff channel. An absent feed suppresses its dependent read class per the T7 manifest and is stated; the skill never improvises around a missing feed.
Deterministic logic	Register state machine (QUEUED through CLOSED) and the two sequencing rules. Baseline assembly with the clean-window check and the G10 guard-date computation. The class check (rung against L2 limits; zero only on branded defense classes). The V1 prediction contract with the bound stated as an input under test. The capture computation per window and the four-branch classification, including the rank-axis drop rule and the X8 conservative rule on Branch C. The void consumption from SOP-24 P3 with window re-dating. The call format and its deadline on the V4 weekly anchor. The finding classes and their routing split (P12 handoff against T4 record). The backlog quantification in dollars per week. No local thresholds: every band, limit, window and tolerance dereferences v4.0 at execution time, the one ungoverned value (the capture bound) is read from the design packet only, and the skill version-pins v4.0.

Block	Specification
Human checkpoints (drawn from PPC Authority Matrix v2.0, Section 1, by row)	AI-executable with audit: register upkeep and scoring, baseline assembly, design drafting, capture arithmetic, branch classification drafts, void detection drafts, call drafting, finding drafts, backlog quantification, the P13 package. Human-mandatory: every test arming and every zero test, per the row "Incrementality tests" (the incoming operator decides and executes, the PPC Lead reviews, the CEO and BM owner are informed); every staged deployment confirm, per "Bulk deployments (#38)" (the PPC Lead reviews, the operator executes); any test touching a shared defended term, per "Sibling declarations (S-series)" (the PPC Lead decides, the BM owner reviews) through the #33 declaration; and the release routing's downstream posting, which is never this skill's action, per "Push funding (the M5 queue)" (the CEO decides above envelope, the PPC Lead within it). Weekly reads inside assigned scope sit on "Weekly verdicts (established US)": the PPC Lead decides, the operator executes assigned scope, the BM owner reviews in the daily huddle, review not veto. Under the vacancy rule an unfilled role's decision right reverts to the PPC Lead and its execution to the PPC Manager.
Cadence	Weekly: P1-P2 inside the Monday cascade, scoring first; window closes read per their T6 dates; calls the same cycle. Daily during open windows: the P8 confound watch against T5, #32 notices, and the #37 log. Quarterly: the P12 backlog review in the V4 framework-review slot; M3 doublings on their T4 dates. Annually: the P13 package per defended syntax ahead of the SOP-26 P5 date.
Alarm specification	F1-F12 each carry a detection signal from Section 7, an evidence-line format (entity, signal value, the v4.0 rule or gate reference, the artifact where it was read), and a routing tier. Same day to the operator queue: F2, F8, F10, plus any G1/G6 abort. Into the weekly worklist: F1, F4, F7, F9, F11, F12. To the PPC Lead with register history: F3, F5, F6, and any mode recurring in consecutive cycles.
Skill outputs	Drafted register updates and scores; baseline artifacts; design packets with predictions; staged rung sets in #38 format; capture reads with branch drafts; void records; drafted calls; finding records with routing drafts; the quarterly backlog review with the theater total; the P13 package; F1-F12 alarms with evidence lines. Everything queues for confirmation at the tier the matrix names; nothing in the human-mandatory set deploys from the skill.

10. Version Governance

Owner: the incoming operator under the PPC Lead's review, per the Authority Matrix row and its standing rule 4 (the matrix records the Transition Brief's assignment on incrementality).

10.1 Review Triggers Specific to This Component

Any v4.0 increment, at which point four things re-validate before any test arms: rule X12 and its capture window, scenario S-D4 and the S-A5 / S-Y10 batch conditions, the M3 doubling sentence, and the Section 5 gate set this document reads.

The landing of the AM-1 capture tolerance in v4.1, which converts every design bound from an input under test to a governed value and re-opens every open design the day it lands.

Any change to #28's ladder mechanics or register format, since every rung and every floor record here executes through them.

Any change to SOP-24 P3's published confound classes (the closed list of five), which this SOP consumes by name.

Any change to SOP-26's P2 step 6 handoff, its I-9 call contract, or its P12 verification, which are the three seams this component lives on.

The landing of the D8 competitor depth feed (Ledger dependency D4) or full D9 rank coverage, which convert two stated coverage gaps into standing reads.

Every quarterly P12 review that finds theater, whose quantified total is evidence about this document's own routing discipline.

10.2 Change Note, v1.0 to v2.0

Supersedes v1.0 in full. Procedures expanded from 6 to 13: the register and calendar (P1-P2), gate clearance with the freeze and the #37 exclusion (P5), staging with T+1 (P6), the call (P9), release integration (P11), the backlog review as a dated procedure (P12), and the defense feed (P13) were implicit or absent in v1.0's P1-P6 and are now separately executable. Verdict mapping moved from six locally minted strings to the closed vocabulary with cross-maps; all six v1.0 strings retired (Section 5.2). Worked examples expanded from 2 to 4; failure modes from 6 to 12 with all six v1.0 patterns carried. Elements 11 and 12 added. Corrections of substance: the release seam re-pointed from "#31 via SOP-26 P5" to SOP-26 P12 and the instrument layer per Ledger conflict C3; the SOP-24 citation updated from P1 to P3 per the attached v2.0; the untraced "v4.0 baseline window" and the class-calendar cadences retired to

AM-2 and AM-5. Thresholds carried forward from v1.0 with traces re-verified: the 2-week capture window (X12, S-D4), the one-rung step and restore rule (M3, X12), the branded-class zero design (carried as this SOP's own design property, with its L2 question logged at AM-6). Incidents and patterns carried: the six v1.0 failure patterns by name, the two v1.0 worked cases recomputed at full grain (WE-1, WE-2), and the confounded-probe incident recorded in the SOP-26 v2.0 rebuild (WE-3, F8).

10.3 Amendment Register: Eight Items for v4.1

ID	Item
AM-1	The capture tolerance. Rule X12 reads "if total capture (organic + paid) holds within tolerance for 2 weeks" and S-D4 repeats the window; no tolerance numeral exists anywhere in v4.0. This is the single number this component is built around, and until it lands every design states its own bound as an input under test (P4 step 4). Highest-priority item in this register.
AM-2	The baseline window. v1.0 cited "the v4.0 baseline window"; no such window exists (the only baselines v4.0 defines are the pre-OOS snapshot of family F3 and the T5 recovery baseline). Retired; each baseline states its own window as an input under test until v4.1 defines one.
AM-3	The rank-slip tolerance. Rule M3 reads "on slippage beyond tolerance, restore one step" with no numeral, so the rank bound on rung tests is design-stated. Likely duplicates items from the ladder-owning sessions; v4.1 folds once.
AM-4	The ambiguous-read rule. v1.0 allowed one window extension on an inside-tolerance-but-trending read; no v4.0 rule defines the state or the allowance. Retired; Branch C runs on X8's conservative rule and the design's own review date.
AM-5	Standing re-test cadence per candidate class. v1.0 carried semi-annual on branded classes and annual per term cohort with no v4.0 basis; retired. The governed anchors are quarterly floor re-tests (V4, M3) and the quarterly S-D6 protection review; every other cadence is a v4.1 question.
AM-6	The zero-test step-size exemption. The branded zero test steps past the L2 per-change cap by construction; whether the cap binds a designed status or budget action, or an explicit exemption is owed, is ungoverned. The zero test is carried as this SOP's own design property meanwhile, restricted to branded defense classes.
AM-7	Lag table. No per-trigger and per-call maximum lags exist for this component; everything anchors to the V4 cadence slots. Echoes the Element 2 lag gap logged by a prior format-playbook session; v4.1 folds once.
AM-8	Record retention, plus the register-state vocabulary (QUEUED / DESIGNED / ARMED / READING / VOID-RESTART / CONCLUDED-UNACTED / CLOSED) which is register vocabulary outside the closed verdict vocabulary and is owed formal admission. Retention duplicates items logged by prior sessions; v4.1 folds once.

11. Operator Ramp

Built from the PPC Authority Matrix v2.0 rows that name this component's decisions. The governing row is "Incrementality tests": the CEO informed, the PPC Lead reviewing, the incoming operator deciding and executing, the BM owner informed. The PPC Lead is therefore the checkpoint holder who reviews in week 2 and audits in week 4. Supporting rows: "Weekly verdicts (established US)" for reads inside assigned scope, "Bulk deployments (#38)" for staging, "Sibling declarations (S-series)" for the G11 check, and "Push funding (the M5 queue)" for why the posting is never this operator's action. The live-defense questions test the reasoning behind the rules, not recall of values: an operator who can recite the 2-week window and cannot say why a capture hold with rank sliding is a drop has not passed.

11.1 Week 1: Observation and Artifact Familiarity

- Reproduce the capture computation on one closed test from its artifacts alone: baseline, rung history, window arithmetic, branch classification. Demonstrated by matching the recorded classification without seeing it first.
- Trace one T4 floor to its test: the register row, the design bound, the Branch B read, the restore's T+1 diff, and the context note. State the term's premium over its own Max Profitable CPC in dollars per click and per week.
- Locate every artifact named in Element 8 in the live instruments: the register, T6 prediction rows, T4 floor columns, T5 event types, the T7 statuses this component depends on. An operator who cannot find the artifact cannot produce it.
- Walk the five SOP-24 P3 confound classes from that document and name the detecting artifact for each, including the D8 gap.

- Observe one full Monday cascade and one BM huddle in which a continue-or-restore call lands, as observer: the matrix places the BM owner's review on weekly verdicts, and the operator needs to have seen review-not-veto before running one.

11.2 Week 2: Supervised Execution

The operator executes with the PPC Lead reviewing every output before it deploys, per the "Incrementality tests" row's R.

- Design one rung test and one branded zero test end to end (P3-P5), each with its V1 prediction and stated bounds. The reviewer rejects any design whose bound is unstated or inherited, which is the fastest way to learn F12.
- Run one capture read on a historical window and classify it; a different branch than the record is a finding to explain, not an error to hide.
- Draft one void with the SOP-24 P3 round-trip and the G10 guard-date computation shown.
- Stage one rung per #38 with the four-segment reason string. Matrix row: "Bulk deployments (#38)", the PPC Lead reviewing, the operator executing.
- Draft one finding and its P11 routing, and defend why the posting is not this operator's action. Matrix row: "Push funding (the M5 queue)".
- Run one G11 check on a shared term against the #33 declaration. Matrix row: "Sibling declarations (S-series)", the PPC Lead deciding, the BM owner reviewing.

11.3 Week 4: Unsupervised Execution With Audit

The operator decides and executes per the matrix D and E; the PPC Lead audits after the fact against the artifact chain. The human-mandatory set does not move at any ramp stage.

- Own the register and calendar weekly, including scoring, and every open window's read and call on assigned scope.
- Own one full test from intake to CLOSED, including the routing or the T4 record. Audited specifically on whether the P8 sweep exists on every read, because the sweep's absence is the F8 signature.
- Own one void end to end, including consumer notification.
- Run one quarterly-style backlog sweep on the live register with the dollars quantified.
- Audit finding: any read whose branch the PPC Lead would have classified differently returns the operator to week-2 supervision on P7 only, not on the whole scope.

11.4 Live-Defense Questions

1. Why does the test read total orders and not paid orders, and what exactly is happening when paid orders fall during a clean test?
2. Why may a branded defense class test to zero while a rank-3 category term may not, and what burns in the second case?
3. Capture holds inside the bound but rank slides past its bound. Which branch, and why is "capture held" not an argument?
4. What does a VOID protect, in both directions, and what does misfiling one as a MISS teach the calibration series?
5. Where does the hold bound on your design come from today, and what are you required to do because v4.0 does not carry it?
6. Why does a NON-INCREMENTAL finding not post to the queue from this SOP, and what breaks downstream if it does? Give the recorded percentage.
7. On a quarterly M3 doubling: whose rung is it, whose read is it, and whose floor record is it? Name the documents.
8. What re-opens a concluded test, and why is a finding dated?
9. Two competitor deal windows land inside your open window and D8 is MISSING. Walk the protocol, including how you knew about the windows at all.
10. Name three things this SOP is forbidden to decide and say which component decides each.

12. Interface Contracts

Twelve handoffs, split upstream and downstream. Every row states what happens when the handoff does not arrive, because that is the clause the corpus has always been missing. Sources: the Cross-Reference Ledger row for component 35 and the rows of every counterpart named, cross-checked against Master PPC Data Architecture v1.0 for artifact names and tab locations.

ID	Counterpart	Artifact and timing	Non-arrival consequence
I-1	SOP-26, upstream	The INCREMENTALITY LADDER handoff: the staged rung, the T6 entry with its review date, and the term's E2 evidence, at P2 step 6 or the P3 doubling.	Nothing arms on the term here, and an SOP-26 rung dated without a design on this side is drift on theirs: the P1 weekly reconciliation catches the mismatch and returns it, because a rung running without a capture design is exactly the unmeasured taper the flag exists to prevent.

ID	Counterpart	Artifact and timing	Non-arrival consequence
I-2	#28, upstream	Rung derivation and the floor register format, on request per design and per restore.	The rung does not stage: deriving a step here is inventing a threshold. The window moves out, the delay records on the register row, and the counterpart's re-test date inherits the slip.
I-3	SOP-24 P3, upstream	The confound and void determination with the class, per read and on demand mid-window.	No read classifies: the branch decision waits, because a read classified without the sweep is the F8 signature. The window holds open with the gap stated rather than closing on an unswept read.
I-4	#32, upstream	Armed-window declarations and settle marks, at arming.	G10 cannot apply prospectively: windows get scheduled into events they could have avoided. The retrospective T5 scan still voids the affected reads, but every void that a declaration would have prevented is recorded as an I-4 miss, and affected windows re-date.
I-5	#33, upstream	The family ownership declaration on any shared term, before arming.	The test does not arm (G11) and the escalation goes to #33; the flag owner's rung holds where it is. Two siblings reading the same term measure each other, not incrementality.
I-6	#41 economics, upstream	The margin refresh into the T0 zone C named cells, monthly and on event.	Every ceiling-referencing statement in designs and findings carries the G9 PROVISIONAL tag. Capture arithmetic itself still runs (it is order counts against a baseline), so tests continue with the tag stated rather than waiting.
I-7	Data Operations, upstream	The weekly validated packet: SellerBoard totals into T0 zone B, Monday 06:00, workbooks by 09:00.	The week is not validated and no window closes on it: windows extend to the next validated week rather than reading console estimates, and the extension records on the register row. An estimate-based capture read is a dirty baseline's twin.
I-8	#38, both directions	Staged rung sets out; the T+1 verification diff back, per rung and per restore.	Without the diff the rung is unproven and its window does not start (P6 step 3). At a later shortfall the record cannot separate "the rung did not land" from "the rung landed and capture moved", which are opposite responses.
I-9	#37, both directions	The frozen-scope exclusion state out at arming; the rule change log back, continuously across open windows.	A platform rule stepping a bid inside frozen scope voids the affected read and the rule suspends pending #37 review. An exclusion state that never lands means every window on that scope is exposed to silent contamination, and the register row carries the exposure until it lands.
I-10	SOP-26, downstream	The continue-or-restore call at each capture-window close, on its interface I-9 channel.	Stated on their side and honored here: the rung holds and no floor records. On this side a missed call deadline leaves their rung undated past its review; the register carries the due date and a late call is a weekly-cascade finding, not a private slip.
I-11	SOP-26 P12 and the T4 register, downstream	NON-INCREMENTAL findings to P12 with evidence; INCREMENTAL FLOOR records into T4 through #28's format, at conclusion.	The two F-classes by direction: an unrouted release leaves the M5 queue understating its supply while queued pushes wait (the quiet failure), and an unwritten floor leaves a proven premium re-accumulating (the loud one eventually). WE-4 prices both; P12's quarterly sweep is the backstop.

ID	Counterpart	Artifact and timing	Non-arrival consequence
I-12	#6, V5 and SOP-26 P5, downstream	Predictions and scores into T6 weekly; voids with class references as they occur; the P13 evidence package annually per defended syntax.	Unscored actions do not accumulate and the next design is uncalibrated; a void filed as a MISS corrupts the series in the other direction. Without the P13 package, SOP-26 P5's cannibalization-adjusted figure degrades to the raw number with the gap stated, and the defense-cost read measures the wrong thing.

Companions: Master PPC Decision Criteria System v4.0 (Section 3 families F2 and F16, Section 4 rules X8, X9 and X12, Section 5 gates, Section 8.1 S-A5, Section 8.2 S-Y10, Section 8.6 S-D1 to S-D6, Section 9 rule L2, Section 10 rules M3 and M8, Section 12 rule P8, Section 13.1 and 13.2, Section 14 rules V1, V4 and V5) · Keyword Decision Record Template v3.0 (the closed vocabulary and the campaign subset) · PPC Strategic Doctrine v1.0 Sections 12 and 13 (the why) · Master PPC Data Architecture v1.0 Sections 4, 6, 8, 9 and 11 (artifacts and tabs) · PPC Authority Matrix v2.0 Section 1 (Elements 9 and 11) · PPC Economics Worked-Example Set v1.0 case E5 (the premium arithmetic WE-3 re-uses) · SOP-24 P3 (voids), SOP-26 (handoff, release verification, defense read), SOP-28 (rungs and floors), SOP-12 (execution), SOP-32 (windows), SOP-33 (ownership), SOP-37 (rules), SOP-38 (staging and T+1) · PPC Cross-Reference Ledger (component 35 row and the interface list).

Inspiratek & Ecotero LLC | Confidential | Page 12 of 12