

SOP-37: Automation Platform Rules

PPC Department · Operating SOP · Component #37

Version 2.0 · 2026-08-09 · Owner: PPC Lead · Supersedes SOP-37 v1.0 (2026-07-24) in full

Thresholds live in the Master PPC Decision Criteria System v4.0 only. This SOP carries zero local thresholds. Every inventoried rule parameter in this document is recorded platform state carrying a reconciliation outcome, never an authority.

1. Verdict and Scope

VERDICT

A rule is a standing decision executing at machine speed, and it is legitimate only while four things stay true at once: it is registered with an owner, it reconciles to a named v4.0 rule or carries a dated documented exception, its scope excludes every entity whose objective function is not the one the rule assumes, and every change it makes is attributable in the T+1 diff. A rule missing any of the four is suspended on discovery, not tolerated until the next pass. The account has already paid for the failure this SOP exists to prevent: one target-ACOS ladder walked a defensive term up 6 consecutive steps against a target left stale by a retag, and the code-red four-map diagnosis found it, not the automation layer that owned it.

This SOP governs the rule register and its schema; rule intake and registration; the reconciliation of every live rule against v4.0; rule-class guardrails and the scope-exclusion state machine; change control; the weekly runaway watch; suspension and re-enable; the rank-based taper rule and its single-rung boundary with #28; the precedence arbitration between rule action and staged manual action; portfolio-object governance as the second budget-enforcement surface; the sync cadence at every v4.0 increment; the quarterly full pass; automation posture inside Recovery Mode and code red; and platform capability, seat and export changes.

1.1 The dependency that shapes this document, stated first

External dependency D6 in the Cross-Reference Ledger Section 2 is the automation platform rules export, and it is OPEN. The Ledger's effect line reads that the rule register cannot be reconciled against v4.0 without it. That is correct about the mechanism and wrong as a reason to hold this component at v1.0 depth, and the distinction decides how the document is written.

D6 blocks exactly one thing: a machine-readable diff. While it is OPEN, P1 runs as a human transcription of the live rule list, which carries transcription risk and a detection latency bounded by the operator's pass rather than by a scheduled job. It blocks nothing else. Reconciliation outcomes are per-rule judgments against v4.0 sections an operator reads directly; scope exclusion runs off campaign-state feeds the workbook already carries; change control is a T6 discipline; precedence is decided by the #38 T+1 attribution report, which exists today; portfolio governance runs off console objects and the envelope instruments. Thirteen of fourteen procedures are fully executable at D6 OPEN. One is executable at reduced assurance and says so in the row.

Manifest state at build, read from the live T7 Manifest of the Command Workbook TEMPLATE v1.0: A1 four-period comparability reads MISSING, blocking all trend verdicts, which is why the P6 runaway watch compares against a single-window expected volume. A3 pacing buckets reads PARTIAL with the day-of-week curve pending, which is why P10 collisions surface as truncation rather than as pacing deviation. A6 objective plus per-term Expected CPC reads PARTIAL with Expected CPC held at an account reference, which caps the precision of every ladder-priced rule. A8 PAT rows first-class reads PARTIAL, so SBV-PAT reconciliation carries the coverage gap in the row. Packet completeness stamps at 0.571, below the DEGRADED floor of gate G8, so every outcome recorded this cycle carries its manifest state beside it.

1.2 What changed at v2.0

v1.0 was built to the superseded 10-element standard and measured 3 pages against a 10-18 page bar. v2.0 expands procedures from 7 to 14, failure modes from 6 to 12 and worked examples from 1 to 4, adds Elements 11 and 12, neither of which exists in any of the 23 shipped SOPs including the format exemplar, and maps the full 25-string closed verdict vocabulary in place of v1.0's 5-row table of register states.

Four corrections of substance.

First, v1.0 P5 sized the runaway signature against "the v4.0 consecutive-step reference," which does not exist: v4.0's consecutive-period language governs weeks in scenario conditions, never steps in a ladder. It is retired as ungoverned and logged as AM-2 rather than carried.

Second, v1.0 reconciled B-class pause thresholds against "the v4.0 reactive thresholds per ad type," where scenario S-H7 conditions on a threshold v4.0 does not state and family F11 gives waste as a percentage; those parameters become documented exceptions pending AM-1.

Third, v1.0 treated the B-class pause as a question about the click count, when the prior question is the lever order: rule L1 places status last and rule L7 requires a log naming the exhausted levers, so a rule that pauses on a click count satisfies neither. That is F7 at v2.0 and was invisible at v1.0.

Fourth, v1.0 P7 suspended rule action for a staged decision's read window where rule P3 has automation resuming from the audit value; the difference is reported in Section 13.1. All six v1.0 failure modes and both account-history incidents carry forward, now sourced rather than asserted.

1.3 The rule register: four classes, one schema

Every live rule appears in the register with nine fields: name, class, trigger condition, action, bounds, scope, owner, reconciliation outcome with its date, and last-synced v4.0 version. A live rule found outside the register is F1 and suspends on discovery. The four classes below are the inventory as seeded from the live platform in 2026-07, restated here as the objects P3 reconciles, not as thresholds this document sets.

Class	Inventoried trigger and action (recorded state)	Governing v4.0 surface and standing guardrail
A. Inch-up (traffic assurance)	Target bid raised by a fixed \$0.05 where clicks run under 3, per ad type (SP, headline, SBV, SBV-PAT, SD), each bounded: SP roughly \$0.30-1.30, the video and display variants \$0.30-1.00.	Section 5 gate G2 makes a bid read need at least 15 clicks in 7 days, and Section 8.4 scenario S-R1 answers the under-15-click case with bid steps of 15-25% per Section 9 rule L2. The rule's direction agrees; its step form does not, and the divergence is arithmetic, not stylistic (WE-1). Scoped OUT of push, event, terminal and settling entities, and out of any entity deliberately held at a bid floor.
B. Pause-at-zero (waste control)	Pause plus bid floor at zero orders over a 10-day window: SP at 30 or more clicks, headline at 30, SBV at 15, SBV-PAT at 15, SD at 12.	Section 9 rule L1 places status last after documented exhaustion of budget, bid, placement, negation and structure; rule L7 requires the log entry naming the exhausted levers. Section 8.9 scenario S-H7 conditions on zero-sale spend above a threshold v4.0 does not state. The class therefore carries a lever-order defect (F7) and an ungoverned parameter set (AM-1) at once. The converting-term guard owned by #29 rides every rule-fired pause.
C. Target-ACOS step ladders	Step ladders toward target ACOS at 5, 10, 15, 20, 25, 30 and 35 percent tolerances, split by order state: the 1-or-more-order family (3 or more clicks, maximum bid \$2.50) and the zero-order family (3 or more clicks, top-of-search share capped at 30, maximum bid \$1.20).	THE scope-sensitive class. Section 6 rule O6 makes each objective judged only on its own metric set, so an efficiency ladder on a ranking, event or liquidation objective is a category error, not a tuning error. The maximum-bid fields are flat dollars where Section 10 rule M1 derives the ceiling per variation as contribution multiplied by CVR, which is why a flat cap can sit above the product's own ceiling (WE-2). Stale targets after a retag are the runaway root cause (F3).
D. Rank-based taper	At target rank with the platform's hold condition met, the keyword steps down.	Section 8.4 scenario S-R6 fires the objective transition only after organic rank sits at or under target for 2 consecutive weeks, and Section 8.12 scenario S-E5 makes rank achieved inside an event window provisional until the exclusion window closes. The rule executes rung one of the #28 ladder and flags the term; rungs two onward depend on hold predictions the rule cannot make, and the step size is Section 10 rule M3's roughly 10% per week, derived at #28 and never here.

1.4 Boundaries, with the owning component named for every excluded decision

Per the Cross-Reference Ledger row for component 37, this SOP owns the rule register, the sync-check cadence and change control so automation never drifts from the criteria. Everything below is referenced by number and one clause and is never restated here.

Excluded decision	Owner	The seam, stated so it is owned
Every threshold, band, scenario, limit and verdict issuance	v4.0 (#8)	Rules implement v4.0 and never rival it. A numeral in this document without a v4.0 section beside it is either recorded platform state carrying a reconciliation outcome or a defect. Section 13 lists the thirteen numerals v4.0 does not yet carry.
The manual deployment path, staging format, upload and the T+1 verification diff	#38	This SOP consumes the T+1 attribution report as its precedence evidence and its unlogged-edit detector. A rule-fired change that the diff cannot attribute is the #38 F4 class arriving on this SOP's desk.

Excluded decision	Owner	The seam, stated so it is owned
Ladder mechanics: rung derivation, exit-bid discovery, the floor register, modifier sizing	#28	Per the Ledger row for 28, ladder holds are that component's and it issues the hold notice naming which rows are frozen. The D-class rule executes rung one and hands off; it never derives a rung.
Runaway diagnosis inside a declared code red	SOP-26 P3	That protocol's four-map diagnosis reads this SOP's rule-change log as its automation check. This SOP owns the standing weekly watch outside code red and the suspension execution inside it.
Terminal and LTSF-tier entities	SOP-27	Terminal handling belongs to that lane exclusively. This SOP excludes tagged entities from B and C classes and computes nothing about which units or tiers, per conflict C6: the deciding system sits outside this corpus.
Auto-side rule scope state and the per-group protected surfaces	SOP-15 P11	That SOP maintains the exclusion state feed and never edits a rule; this SOP registers, reconciles and suspends.
Negation doctrine, wall levels, shared lists, the converting-term guard	#29	A rule-fired pause or negative applies #29's standards unchanged. This SOP routes; it does not standard.
Envelope law, queue ordering, CAP SYNTAX and envelope raises and cuts	The instrument layer (T0 named cells, T1 caps, Account Rollup A2) per conflict C3	Component #31 carries no numbered procedures, so P10 reconciles portfolio caps against the instrument and tab, and the dangling-citation gap is logged as an amendment item rather than cited as a procedure.
Event-window arming, ramps, caps and reconciliation	#32	This SOP excludes armed and settling entities; it never decides what an event plan permits.
Automation-platform seats, credentials and their rotation on departure	#46	That SOP owns the seat, this SOP owns the rules. A seat change that removes a rule owner triggers P7 here within the same session.
Billing continuity and the outage response	#43	That SOP requests this SOP's rule log to clear its triage branch 4. This SOP owns the post-restoration rule-read suspension.
Which terms graduate, and the exact standups that follow	#30 and #12	A rule never creates, retags or consolidates structure. Structure is the last lever in the L1 order and it is human territory.

2. Trigger Conditions

Each trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. Lag is measured from the artifact's own timestamp, not from when someone read it.

Trigger	Detecting artifact	Maximum lag	Procedure
Weekly register diff	The platform's live rule list read against the register, inside the Monday cascade after the packet-ready read	Same working day as the packet-ready read	P1
New rule requested	A rule-intake row raised by any operator, or a rule found live and unregistered by the P1 diff	1 working day to registration or suspension	P2, P7
v4.0 version increment	The version stamp on the Criteria System file changing against the register's last-synced field	1 working day to open the pass, 5 working days to close it	P11, P3
Quarterly full pass	The register's last full-pass date reaching 13 weeks, per v4.0 Section 14 rule V4 quarterly cadence	The scheduled week	P12, P3
Campaign state change	A push standing up or closing, an event window arming or settling, a terminal tag landing, a manual hold being placed: read from the T2 gate columns and the T3 campaign state	Same day as the state change	P4

Trigger	Detecting artifact	Maximum lag	Procedure
Runaway signature	Rule-fired change counts per entity per window in the platform change log, read against expected volume	Same working day as the weekly read	P6
Rule edit without a record	The register diff against the T6 change records	Same working day	P5, P7
Rule vs manual collision	The #38 T+1 verification diff attributing a change to a rule on an entity carrying a staged decision from the same cycle	Same day as the T+1 report	P9
Rank-target reached	The platform hold condition firing on a D-class rule, cross-read against organic rank in T2 block K4	Same day for the flag; the rung derivation waits on #28	P8
Code red declared	The declaration written to T5 by SOP-26 P3, same day as the trip	Same day	P13
Recovery Mode entry or exit	The product's Recovery state per v4.0 Section 11 rules R1 and R5	Same day	P13
Portfolio object created or edited	The console portfolio list read against the portfolio section of the register	Same working day as the weekly read	P10
Billing or platform outage restored	The restoration verification from #43, or the platform's own incident close	Same day as restoration	P14 step 4
Seat or credential change	The #46 access reconciliation diff naming an automation-platform seat	Same day	P14
Platform capability change	The vendor's release note or a new rule type appearing in the console	5 working days to register the class or decline it in writing	P14, P2

3. Preconditions and Gates

Each gate below names the manifest field group it reads and the observable that proves it passed. A missing or stale group suppresses its dependent execution class; the pass records what it could not see and continues on the rows it can, per v4.0 Section 5 gate G8.

Gate or precondition	Manifest ID and source	The observable that proves it passed
Register currency	The register's own last-synced fields	Last-synced version equals the version on the attached Criteria System file and the last full-pass date is inside 13 weeks. A register synced to a stale v4.0 is a register of guesses, and no outcome recorded against it is admissible.
Live rule list readable	Ledger dependency D6, automation platform rules export, status OPEN	A dated console read of the live rule list, transcribed by entity and initialised, exists for this cycle. The row carries the manual-read flag and states F1 latency at one cycle.
Deployment state current	A2 deployment state in-row, T7 PRESENT	Current bid, budget and state for every entity carrying a live rule, timestamped inside this cycle. Stale A2 stops P9 arbitration, because prior-value evidence is unreadable.
Attribution report present	A2 plus the #38 T+1 verification diff	A T+1 diff exists for every deployment since the last pass with zero unexplained mismatches, so every change carries a source of session or rule. Absent, P6 and P9 do not run.
Threshold flags computed	A5 threshold flags in-row, T7 PRESENT	Gate and band flag columns compute on the affected T2 rows, so reconciliation reads flags and never re-derives a threshold.
Objective and Expected CPC	A6, T7 PARTIAL	Exactly one declared objective in T2 block K7 on every entity carrying a C-class or D-class rule. Expected CPC is held at the account reference until per-term values land, so ladder-priced rules carry the A6 PARTIAL flag.

Gate or precondition	Manifest ID and source	The observable that proves it passed
Economics chain live	B8 chain, T0 named cells	AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC, with the D1 margin version and its G9 state. Maximum profitable CPC computes per variation per rule M1. G9 PROVISIONAL carries onto every ceiling-referencing outcome, so a cap approved on provisional margin is dated and re-checked at the next refresh.
Coverage counts current	B5 coverage counts, T7 PRESENT	Registry counts per match type are current, so a rule scoped to a campaign class can be proved to cover what the operator thinks it covers.
G1 validation quarantine	T2 zero-flag columns	No quarantined row inside a rule's scope; rule action on a quarantined entity suspends until the data-refresh task closes.
G3 budget truncation	K2 in-budget %, worst hosting campaign	At or above 70% of the day on every entity a C-class rule prices. Below it every rate is TRUNCATED and a rule-fired bid change on that day is void by construction.
G6 inventory adequacy	Inventory band and days to stockout, T2 block K2, computed from the D2 map	Band GREEN on the targeted variation for any rule that can raise a bid; YELLOW and RED entities leave raise scope the same day the band flips.
G9 price stability	The price change log in T5, trailing 7 days	No price change on the targeted variation inside 7 days; inside the window no conversion-driven outcome is recorded as a finding.
G10 event exclusion	Event posture in T2 block K2 and the #32 armed-window record	Event weeks and the 14 days following are excluded from trend verdicts; a D-class taper firing inside them is acting on a provisional read per scenario S-E5 and is F9.
G12 MKL hygiene	The MKL refresh date	MKL current inside 30 days; a stale list past 30 days degrades the manifest and rules on untagged terms route to maintenance rather than reconciliation.
G8 packet completeness	T7 completeness %	Stated, not assumed. Below 85% the packet is DEGRADED and every outcome this cycle is a flag carrying its completeness stamp. At build the stamp read 0.571.
Mode posture known	The product's Recovery and event state	Per rule T8 and rule P7, Recovery outranks Event and both outrank stage defaults. Mode is read before any re-enable, because Recovery tightens the confirm tier under rule R3.

4. Step-Level Procedures

P1. The weekly register diff

25-40 minutes at 24 registered rules. Runs inside the Monday cascade, after the packet-ready read and before any T2 verdict work, because a rule that moved a bid overnight changed the prior value every verdict this cycle is measured against.

1. Open the register at its current version beside the platform's live rule list. While dependency D6 is OPEN there is no export, so this is a transcription read: work top to bottom by ad type, never from memory of last week.
2. Transcribe each live rule's nine fields into a dated working copy. Match rules by name plus trigger, never by name alone, because a rule edited in place keeps its name.
3. Produce the three-way diff: live and registered, live and not registered, registered and no longer live.
4. Live and not registered is F1: suspend in the console the same session, then route to P2 or to a written retirement. A rule nobody registered is a decision nobody made.
5. Registered and no longer live: confirm a deletion was staged in T6. A deletion without a record is F4 and the row stays open with its owner named.
6. Field mismatch on a live registered rule: apply the decision tree below.
7. Write the diff artifact: date, rules read, three counts, every mismatch with its disposition, the manual-read flag while D6 is OPEN, and the operator's initials.

P1 step 6 decision tree, field mismatch on a live registered rule:

1. Is there a T6 change record dated on or before the platform's last-change date naming this rule and this field? If yes the register is stale, not the platform: update the row, carry the reconciliation outcome forward only if the changed field is not the one the outcome rests on, and stop.
2. If no record exists, is the change inside the rule's registered bounds and scope? If yes it is a silent edit inside bounds, F4: revert to the registered value pending review, reconstruct the record from console history, route to P5.
3. If no record exists and the change sits outside registered bounds or scope, suspend the rule the same session and route to P7. Damage assessment runs on every entity the rule touched since the platform's last-change date, using the #38 T+1 diffs for that window.
4. If console history cannot name an actor, escalate to the PPC Lead with the entity list and route the seat question to #46 the same day. An unattributable edit is an access finding before it is an automation finding.

P2. Registering a rule, and the pre-go-live reconciliation

20-30 minutes per rule. Runs on intake and on every unregistered rule P1 finds.

1. Collect the nine register fields from the requester. A request missing bounds or scope is returned, not completed by the registrar: the operator who wants the rule is the one who knows what it must not touch.
2. Name the v4.0 rule the proposed rule implements, by section and rule ID. A requester who cannot name the governing section is proposing to invent a threshold, and the request is declined at intake.
3. Run P3 before go-live. A rule enters at AGREES, at ADJUST with the adjusted value already applied, or at DOCUMENTED EXCEPTION with an owner and a date. There is no provisional live period.
4. Set scope exclusions per P4 before enabling. A rule enabled first and scoped afterward has already run against the entities it was meant to avoid.
5. Set bounds so the rule's largest possible action stays inside rule L2: 50% of the current value per change, floor \$0.50. Wider bounds are out of bounds before the rule fires once.
6. Check the confirm tier. Rule P5 makes bid moves above 25%, budget moves above \$50 per day, actions on 500-plus search-volume keywords, gate failures, structural changes and LOW-confidence verdicts human-confirmed. A rule whose action class can cross those lines registers only with the crossing entities excluded, because a rule cannot produce a confirmation.
7. Write the T6 entry with the v4.0 citation and a prediction carrying metric, direction, magnitude, horizon and review date. Rule V1 makes an unscored action one that does not accumulate, and a standing decision accumulates fastest.
8. Enable, with the first review at the next weekly pass rather than the quarterly one.

P3. Reconciliation against v4.0

10-15 minutes per rule; 4-6 hours for a full 24-rule pass. Runs at every increment (P11), quarterly regardless (P12), and at registration (P2).

Three outcomes exist and only three. AGREES: the rule implements the v4.0 threshold as written, at the same value, on the same condition, inside the same limits; record section, rule ID, both values and the date. ADJUST: the rule diverges and moves to v4.0's value this session, staged per P5, prior and new values recorded. DOCUMENTED EXCEPTION: the divergence is deliberate, written, owned and dated, or v4.0 carries no counterpart at all; the first needs a reason and a review date, the second needs an amendment item in Section 13. Divergence with none of the three is F5.

Per-class checks, in this order:

1. Lever order first, every class. Which lever does the rule pull, and where does rule L1 place it: budget, bid, placement, negation, structure, status. A rule pulling a later lever without evidence the earlier ones were pulled is F7 however well its numbers reconcile. A correctly parameterised rule on the wrong lever does more damage than a badly parameterised one on the right lever.
2. A-class step form. Compare the fixed dollar step against rule L2's 15-25% at both registered bounds. Outside the band at either bound: ADJUST where the platform accepts a percentage step, DOCUMENTED EXCEPTION with AM-4 where it does not. Record the computed percentages, not the judgment (WE-1).
3. B-class parameters. Scenario S-H7 conditions on zero-sale spend above a threshold v4.0 does not state and family F11 gives waste as a percentage, so the per-ad-type click counts reconcile to nothing: DOCUMENTED EXCEPTION, AM-1. The #29 converting-term guard rides every rule-fired pause exactly as it rides a manual one.
4. C-class scope. Ladders run only where efficiency is the declared objective per rule O6. Read the objective tag on every entity in scope, not the campaign name. A mismatch is F2, not an outcome.
5. C-class ceiling. Compare the flat maximum bid against the entity's own maximum profitable CPC per rule M1, contribution multiplied by CVR. Above 1.5 times that ceiling the rule can carry a bid into the band scenario S-R11 calls unjustifiable outside a sized push, so no exception is available: ADJUST to the tightest ceiling in scope, or narrow the scope until the cap is legal (WE-2).
6. D-class transition condition. The platform's hold condition is not v4.0's test: scenario S-R6 requires organic rank at or under target for 2 consecutive weeks and scenario S-E5 makes event-window rank provisional. Record the divergence and enforce it through scope (P4) and the P8 flag rather than through the rule's own condition.
7. Record every outcome with its manifest state. An outcome recorded on a DEGRADED packet is a flag, not a finding, per gate G8.

P4. Scope-exclusion state management

5 minutes per state change, same day; 20 minutes for the weekly reconciliation of the whole set.

1. Standing exclusions from B and C classes at all times: campaigns with an open push worksheet, event campaigns and settling rows inside their window and guard, entities carrying a terminal or LTSF tag per SOP-27, and any entity under a declared manual hold. A-class adds one: any entity deliberately at a bid floor, because a floor is a discovered value under rule M3 and a rule that lifts it has erased the discovery.
2. The reason is one sentence: pushes run on loss lines, events on caps, terminal on recovery mathematics, holds on a human's stated reason, and an efficiency ladder judges all four against an objective none of them carries, which rule O6 calls a process violation.
3. Maintain exclusions as states change, same day, in both directions: a push standing up adds its campaigns on the standup checklist and a push closing removes them; a window arming adds and settling removes at the end of the guard; a terminal tag adds the week it lands.
4. Verify in the platform, not in the register. The register records intent, the console holds state, and only the console stops the rule. Read the scope back after every edit.
5. Weekly, reconcile the exclusion set against the T2 gate columns and T3 states in both directions: every entity that should be excluded is, and every entity in the set still has a reason. A stale exclusion is a rule silently not running where it should.
6. Apply SOP-15 P11 exclusion-state notices for protected auto scope the same session and confirm receipt. That SOP maintains the feed and never edits a rule; the edit is this procedure's.

P5. Change control

10-15 minutes per change.

1. Creations, edits and deletions are staged decisions carrying what a bid change carries: reason string, owner, version, and a T6 entry with a prediction and review date.
2. The reason string follows the #38 staging contract: component tag, driving outcome or verdict, governing reference, evidence reference. A reason string that cannot name a v4.0 section or a reconciliation outcome is an invented threshold with a rule wrapped around it.
3. Apply in the console, read the value back, record the read-back. The register is the rule world's deployment log and is only true if the read-back matches.
4. Increment the register version on every change, so outcomes recorded against a version can be re-read later against the rules live at the time.
5. A rule edit found without its record is F4: revert to the registered value pending review, then reconstruct the record from console history and the #38 T+1 diff for the window.
6. Deletions keep their row, marked retired with date and reason. A retired row is how the next operator learns which rules the account has already tried and why they stopped.

P6. The weekly runaway watch

20-30 minutes, immediately after P1 on the same change-log read.

1. Pull rule-fired change counts per entity per window for the trailing 7 days. While A1 four-period comparability reads MISSING, the expected-volume comparison is single-window and the artifact says so rather than implying a trend read the manifest cannot support.
2. The signature is consecutive same-direction steps on one entity inside one window. v4.0 carries no consecutive-step count for ladders, so this procedure detects and escalates the pattern and asserts no passing number: the flag count is AM-2, and until it lands the operator flags every entity at three or more same-direction rule-fired steps in the window and records the count as a detection convention in the artifact, not as a threshold.
3. Suspend that entity's rule action pending the read. Suspend the entity, not the rule, unless two or more entities on one rule flag in the same window.
4. Read the cause in order: retag since the target was set, objective change, a margin move recomputing the ceiling under scenario S-A9, an excluded state P4 missed, a G1 quarantine.
5. The classic cause is a stale target after a retag. The retag checklist at SOP-12 P4 and P6 carries the rule-target update step, which kills the cause at source rather than catching it weekly here.
6. Quantify before re-enabling: bid at window start against bid now, premium per click, clicks in the window, dollars. A runaway with no dollar figure will be re-litigated; one with a figure will not (WE-3).
7. Route the suspension packet to the PPC Lead the same day, the entity to the weekly pass for its correct verdict, and the root cause to whichever checklist should have prevented it. A second runaway from one cause escalates the checklist gap, not the rule.

P7. Suspension and re-enable

10 minutes to suspend, 20-30 minutes to re-enable.

1. Execute in the console and record in the register the same session, with the trigger named: F1, F2, F3, F4, F5, F7 or a P6 runaway.

2. State the scope of the freeze: this entity, this rule, or this class. Class-wide is reserved for a defect in the governing logic rather than in one instance of it.
3. Re-enable requires four things present at once: root cause named and fixed at source, reconciliation re-recorded against the current v4.0 version, scope exclusions verified in the console, and the damage record closed with its dollar figure.
4. Re-enable is human-mandatory. The Authority Matrix automation-rules row gives the PPC Lead the decision and the operator the execution, and rule R3 tightens confirmation inside Recovery Mode to bid moves above 15% and budget moves above \$25 per day.
5. A rule re-enabled twice for one root cause does not get a third: it retires, and the need is met by a staged manual procedure until the cause is structurally closed.

P8. The rank-based taper rule and its single-rung boundary

10 minutes per firing.

1. The rule executes rung one and flags the term to the ladder process. It never walks a term down unattended, because rungs two onward depend on hold predictions the rule cannot make.
2. The rung value is derived at #28 from rule M3, roughly 10% per week, inside rule L2's limits. This SOP checks that the rule stopped after one, not the arithmetic.
3. Verify the transition condition against v4.0 rather than the platform: scenario S-R6 requires organic rank at or under target for 2 consecutive weeks. A firing on a shorter hold is recorded and the term returns to the weekly pass without a rung.
4. Verify event state: scenario S-E5 makes rank gained inside a window provisional and starts the clock only after the exclusion window closes. A firing inside an armed window or its 14-day guard is F9 and reverts.
5. Write the flag with entity, rung executed, date and rank evidence. From that point the term belongs to #28 and to whichever component owns the descent decision.
6. If the term sits inside a scheduled descent owned elsewhere, revert the rung and let the schedule run. Two descents on one term compound into a step rule L2 does not permit.

P9. Precedence: rule action against staged manual action

15 minutes per collision.

1. The governing rule is P3: the weekly audit outranks the automation platform between cycles, the audit value wins where both touched the same target, and automation's inter-cycle actions are reviewed each audit.
2. Detection is the #38 T+1 diff, which attributes every change to its source. A change neither a session nor a rule claims escalates as that SOP's unlogged-edit class.
3. Restore the audit value first and investigate second. The audit value is the one carrying a J-complete record under rule P10; the rule-fired value is not.
4. Then decide the rule's status on that entity for the read window. The read-window suspension is this SOP's implementation mechanism and is not stated in rule P3, which has automation resuming from the audit value. Until AM-13 settles it the operating instruction is to suspend the rule on that entity for the stated read window, because a rule that resumes immediately can re-move the value inside the window meant to measure it, which is F6.
5. Record the collision in T6 with the conflict flag set, so scoring can separate the entity's outcome from the rule's interference.
6. Where collisions recur on one rule the fix is scope, not vigilance: the entities the weekly pass touches every cycle are the entities the rule should not.

P10. Portfolio-object governance

30-45 minutes at the first pass, 10 minutes weekly thereafter.

Console portfolios are the second surface on which a budget can be capped, and until this procedure existed nobody owned them. A portfolio cap suppresses spend on campaigns whose own budgets are healthy, so the symptom is campaign-level truncation with no campaign-level cause.

1. Register every portfolio object: name, member campaigns, cap type and value, owner, and the product envelope it sits inside. A portfolio with no registered owner is suspended from carrying a cap, which means the cap is removed and the envelope governs alone.
2. Creation is human-mandatory and follows the Authority Matrix structure-creation row: the PPC Lead decides or delegates per SOP scope, the operator executes.
3. Reconcile every cap against the envelope instruments rather than a #31 procedure, per conflict C3: the product's weekly envelope named cell on T0, the family caps on T1, the cross-product queue on Account Rollup A2. The line is arithmetic: portfolio caps inside a product, expressed weekly, against that product's declared weekly envelope.
4. Three outcomes. Cap at or above the envelope: the envelope binds, the object is inert, record it. Cap below the envelope: the portfolio binds and the envelope is not the operative constraint, which is a finding routed to the envelope owner with the suppressed dollars computed (WE-4). Cap below the envelope and set to enforce it: two surfaces are enforcing different numbers and one is wrong, so raise the cap to the envelope or cut the envelope to the cap in one staged change and never leave both live.

5. A portfolio containing campaigns from more than one product cannot reconcile by construction and is restructured: an envelope is a product-level object.
6. v4.0 carries no portfolio-cap tolerance and no alert bar. Gate G7 governs campaign budgets summing to the envelope and pacing deviation above 25% against the day-of-week curve; it does not name the portfolio object. That is AM-8, and until it lands divergence is reported at its dollar value rather than tested against a bar.
7. Weekly, diff the console portfolio list against the register exactly as P1 diffs the rule list. An unregistered portfolio is the same class of finding as an unregistered rule.

P11. Sync at every v4.0 increment

4-6 hours across 24 rules, opened within 1 working day of the increment and closed inside five.

1. Read the change note and list the rules it touches: thresholds, bands, lever limits, scenario conditions, and any change to the closed verdict vocabulary.
2. Run P3 on every rule, not only the ones the note names. A rule can be made non-compliant by a change it cites indirectly, and the register's value is that this is checked rather than assumed.
3. Record the new last-synced version on every row, including rows whose outcome did not change. A row without the current version is unproven, not unchanged.
4. Where the increment closes an amendment item raised here, move the affected rules out of DOCUMENTED EXCEPTION and close the Section 13 entry with the version that closed it.
5. Where the increment creates a divergence, ADJUST the same session where the platform permits it and suspend where it does not. A rule is not left running against a superseded number while a platform limitation is discussed.
6. Write the pass artifact: version from and to, rules read, outcome counts, adjustments with prior and new values, suspensions, amendments opened and closed.

P12. The quarterly full pass

5-7 hours, once per quarter regardless of increments, per the quarterly cadence of rule V4.

1. Run P3 on every rule, P4 on the entire exclusion set, P10 on every portfolio object.
2. Age every documented exception. An exception past its review date without renewal is a divergence, and it becomes ADJUST or the rule suspends. v4.0 sets no maximum life, which is AM-10; until it lands the review date is the control.
3. Read the quarter's suspensions as a set and look for the pattern rather than the incidents: one root cause twice is a checklist gap, one class three times is a class needing narrower scope or retirement.
4. Re-verify bounds against current economics. Ceilings recompute when the margin table moves under scenario S-A9, so a cap legal last quarter can be illegal this quarter without anyone touching the rule.
5. Score the quarter's rule predictions in T6 per rule V1 and report the register's own accuracy beside the account's. A standing decision never scored is a habit.
6. Produce the quarterly artifact and register it as the baseline the next quarter diffs against.

P13. Automation posture inside code red and Recovery Mode

45-90 minutes on declaration day, then 15 minutes daily while the mode is live.

1. On a declaration written to T5, deliver the rule-change log for the affected product to SOP-26 P3 the same day. That protocol's four-map diagnosis reads this log as its automation check, and its absence blocks one of four maps.
2. Freeze rule-fired increases on the product the same day. The declaration freezes non-defensive increases, and a rule that keeps raising is the freeze failing silently.
3. Where the diagnosis names a runaway rule, suspend per P7 and return the review inside the protocol's window rather than at the next weekly pass.
4. Apply the rule R3 tightening to every rule still live: bid moves above 15% and budget moves above \$25 per day are human-confirmed and step sizes cap at standard even at HIGH confidence. A rule whose action class can exceed those lines narrows its scope or suspends, because it cannot produce a confirmation.
5. Honour mode precedence per rules T8 and P7: Recovery outranks Event and both outrank stage defaults, so a product in Recovery entering an event window keeps the Recovery posture on its rules.
6. On exit, restore the pre-mode posture as a staged change with a record, not by memory, and re-verify scope before re-enabling anything.

P14. Platform capability, seats, exports and outage recovery

Variable: 30 minutes for a seat change, 2-3 hours for a new rule class.

1. A new rule type in the console is registered as a class or declined in writing within 5 working days. An unclassified capability is a rule nobody has decided about, which is how unregistered rules arrive.

2. A new class enters only with its governing v4.0 surface named, its guardrails written and its standing exclusions set. Where no v4.0 surface governs the action, the class is declined and the gap is logged: a capability is not an authority.
3. Seat and credential changes route from #46. A departure or role change touching an automation seat triggers, the same day, reassignment of every row that seat owns and P7 suspension of any rule whose owner is now vacant. Under the Matrix vacancy rule an unfilled role's decisions revert to the PPC Lead and execution to the PPC Manager, so a vacant owner is a temporary state with a named holder, never an empty field.
4. After a billing or platform outage is restored, suspend rule action on the affected scope until one clean full day of data has landed. Rules read the window the outage distorted, and an outage day looks like a collapse in clicks and orders, which is the input the pause and ladder classes react to hardest. Gate G10 excludes event weeks and carries no outage exclusion, which is AM-9; until it lands the suspension window is stated in the T6 record and the affected dates carry the trend-exclusion flag #43 raises.
5. When the D6 export lands, P1 converts from transcription to a scheduled diff: the manual-read flag retires, F1 detection moves from one cycle to one day, and the export schema registers so a schema change is itself a P14 event.

5. Verdict Mapping

All verdict issuance, thresholds and scenario logic live in v4.0 and in the Keyword Decision Record Template v3.0's closed vocabulary. This SOP issues no keyword verdicts. What it decides is narrower and it is the question the account keeps getting wrong: which verdicts a standing rule may execute at all. The table below maps the full 25-string closed vocabulary, quoted verbatim, to the automation disposition and its governing surface.

Verdict (closed vocabulary, verbatim)	Rule may execute	Governing surface and the seam
QUARANTINE: DATA REFRESH	No	Gate G1 quarantines the row; rule action on the entity suspends until the refresh task closes.
UNLOCK TRAFFIC	Yes, A-class	Section 8.4 scenario S-R1 with the 15-25% steps of rule L2. The fixed-dollar step is the divergence WE-1 computes.
FIX BUDGET (TRUNCATED)	No	Gate G3 makes budget the first lever, but envelope decisions sit at the instrument layer per conflict C3. No budget-class rule is registered.
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	No	A dated test with a click quota is a staged human decision; the entity is excluded for its duration.
FUND TEST (UNFUNDED PRIMARY)	No	Funding is an M5 queue decision. Rules never fund.
RANK PUSH: FUND SIZED PUSH	No	Push scope is a standing exclusion (P4); a rule inside a funded push is F2 the day it fires.
RECOVERY PUSH	No	As above, and the recovery sizing belongs to the recovery component.
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	No	The required-clicks arithmetic is Section 10 rule M2 sizing, which no rule computes.
FIX CONVERSION FIRST; DEFER PUSH	No	Section 8.4 scenario S-R4 forbids bidding a conversion deficit; the entity enters the exclusion set.
PLACEMENT FIX FIRST	No	Placement is #28's lever and its gate is placement-level conversion proof. No modifier-class rule is registered.
RE-POINT VARIATION	No	Targeting change, human-mandatory, driven same day by gate G6.
TAPER: RANK ACHIEVED, STEP BIDS DOWN	Rung one only, D-class	Scenario S-R6 sets the condition, rule M3 sets the step, #28 derives the rung; the rule fires one rung and flags (P8).
HOLD; DEFEND LIGHTLY AT RANK	No	The verdict's content is that nothing changes this cycle.
DEFEND AT FLOOR	No, entity excluded	The floor is a discovered value under rule M3; drift restores in one staged step (WE-3).
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	Bid half only, C-class	The bid half moves toward the rule M1 ceiling; the syntax-cap half is an envelope action no rule executes.
INCREMENTALITY LADDER	No	A measured test; a rule stepping inside the window destroys the read.

Verdict (closed vocabulary, verbatim)	Rule may execute	Governing surface and the seam
ADD EXACT + RESTRUCTURE	No	Structure sits last before status in rule L1 and is human-mandatory under rule P5.
CONSOLIDATE DUPLICATES	No	Survivor selection is human-mandatory.
GRADUATE FROM DISCOVERY	No	Decided at #30, executed at #12 same day with walls the same session.
NEGATE (PRE-LOAD / REACTIVE / STEERING)	Routing only	#29 owns doctrine, wall level and the converting-term guard; a rule surfaces candidates and places nothing.
ENTER SBV ARBITRAGE	No	A cross-format posture change; the SP-exact step-down executes at #12.
OPEN / ROTATE / EXIT CONQUEST	No	Owned by #16, on flag evidence no rule holds.
STOP-LOSS: CONCEDE OR DEFER	No	The stop-loss reads velocity and ETA against a sized plan; its owner concedes.
PAUSED-PENDING (REASON + RE-ENTRY)	No	The verdict requires a reason and a re-entry criterion, which is a sentence no rule writes. This is the string the B-class pause approximates, and the approximation is F7.
LEAVE TO HARVEST	No	A pool-management decision at #30.

The campaign-level subset and the register's own state vocabulary

Campaign-level verdicts in the closed subset (FUND BUDGET, HOLD, TRIM BUDGET, RAISE MULTIPLIER, CUT MULTIPLIER, NEGATION PASS, CONSOLIDATE, RETAG OBJECTIVE, PAUSE, ARCHIVE) are executed by their owning components; no registered rule executes any of them today. Separately, this SOP operates six register states that belong to no closed vocabulary and are not verdicts: REGISTERED, SUSPENDED, RETIRED, AGREES, ADJUST and DOCUMENTED EXCEPTION, plus the scope states EXCLUDED and IN SCOPE. They describe rules, not keywords, and they are declared here rather than smuggled in as verdict strings. Their formal admission is amendment AM-7.

6. Worked Examples

All four use the standard fictional training base and say so: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Maximum Acceptable 26.3%, CVR 5.2% unless stated. Live values come from the T0 economics spine per product. Rule parameters shown are the inventoried platform state of Section 1.3.

WE-1. The pass re-run to the v2.0 standard: 24 rules, and why AGREES fell from 21 to 6

Setup. The register carries the 24 rules seeded from the live inventory: 5 A-class inch-up (one per ad type), 5 B-class pause-at-zero (one per ad type), 13 C-class ladders (7 in the one-or-more-order family at a \$2.50 maximum bid, 6 in the zero-order family at \$1.20), 1 D-class taper. The v1.0 seeding pass recorded 21 AGREES, 1 ADJUST, 1 DOCUMENTED EXCEPTION, 1 SUSPENDED.

A-class, step form. The action is a fixed \$0.05 raise; rule L2 expresses the unlock step as 15-25% of the current value and scenario S-R1 is the scenario it serves. At the SP lower bound \$0.30 the step computes to 0.05 divided by 0.30, which is 16.7%: inside the band. At the upper bound \$1.30 it computes to 0.05 divided by 1.30, which is 3.8%: far below it. The step stays inside the band only between \$0.30 and \$0.333, since 0.05 divided by 0.15 is \$0.333, a span of \$0.033 inside a \$1.00 band, or 3.3% of the rule's own range. The consequence is a slow rule, not a dangerous one: reaching L2's smallest permitted step at \$1.30 needs 0.195 divided by 0.05, which is 3.9 firings, so 4 weekly firings deliver what one governed step delivers. Outcome for all 5: DOCUMENTED EXCEPTION, safe-side, dated, AM-4. A divergence acting more slowly than v4.0 permits is documentable; one acting faster is not.

B-class, lever order before parameters. The action pauses and sets a bid floor. Rule L1 places status last and rule L7 requires a log naming the exhausted levers, which no rule writes, so the class as configured cannot be legal. The resolution is a split, not a suspension: the bid-floor half is a bid-lever action and stays, the pause half is removed and its candidates route to the weekly pass as a worklist. Outcome for all 5: ADJUST. The click counts of 30, 30, 15, 15 and 12 over a 10-day window reconcile to nothing and ride in the row as AM-1.

C-class, scope then ceiling. The scope read finds the zero-order ladder live on 2 campaigns carrying open push worksheets: the account's recorded incident, F2 rather than an outcome, exclusion applied the same day and the window's changes read for damage. The ceiling check then splits the class. Maximum profitable CPC on the standard base is 28.00 multiplied by 0.052, which is \$1.46. The zero-order family's \$1.20 cap sits under it: 6 rules AGREES, carrying the G9 provisional flag. The one-or-more-order family's \$2.50 cap sits at 2.50 divided by 1.46, which is 1.71 times the ceiling, above the 1.5 multiple at which scenario S-R11 calls the spend unjustifiable outside a sized push: 7 rules ADJUST, worked in WE-2.

D-class. The platform fires on its own hold condition; scenario S-R6 requires 2 consecutive weeks at or under target and scenario S-E5 makes event-window rank provisional. Outcome: DOCUMENTED EXCEPTION, enforced through scope and the P8 verification, since the platform does not express the two-week test.

Result. 24 rules: 6 AGREES, 12 ADJUST, 6 DOCUMENTED EXCEPTION; one ADJUST rule also suspends pending its root-cause fix (WE-3); 2 scope exclusions land the same day.

The lesson. AGREES fell from 21 to 6 without a single rule changing, because the two passes asked different questions. v1.0 recorded AGREES where a rule's direction matched v4.0's intent. v2.0 records it only where two numbers are identical, on the same condition, inside the same limits, pulling the lever v4.0 puts in that position. Agreement in spirit is the state a register drifts through on its way to being wrong.

WE-2. The flat cap that reaches past the ceiling

Setup. Standard base. A C-class ladder in the one-or-more-order family, maximum bid \$2.50, scoped across three variations of one product.

The ceiling, per variation. Rule M1 derives maximum profitable CPC as contribution multiplied by CVR. Variation 1: 28.00 multiplied by 0.052 gives \$1.46. Variation 2, contribution \$24.60 at the same CVR: 24.60 multiplied by 0.052 gives \$1.28. Variation 3, contribution \$31.40 at CVR 6.1%: 31.40 multiplied by 0.061 gives \$1.92. One rule, one cap, three ceilings.

What the cap authorises on variation 1. At \$2.50 with CVR 5.2%, cost per order is 2.50 divided by 0.052, which is \$48.08 against \$28.00 of contribution: a loss of \$20.08 per order. At 40 clicks per week the rule spends \$100.00 to buy 2.08 orders, recovering \$58.24 of contribution: a net loss of \$41.76 per week and \$2,171.52 per year if nobody reads the row. The efficiency reading agrees, as it must: sales are 2.08 multiplied by \$80.00, which is \$166.40, so ACOS is 100.00 divided by 166.40, which is 60.1%, and 60.1 divided by 35.0 is 1.72 times break-even, matching the 1.71 CPC multiple to rounding. When the two readings disagree the input is wrong.

Disposition. DOCUMENTED EXCEPTION is not available: it would authorise spend scenario S-R11 forbids outside a sized push, and no rule sizes one. Outcome ADJUST, with two legal endings. Cap to the tightest ceiling in scope, \$1.28, compliant everywhere and costing variation 3 its headroom, since 1.92 minus 1.28 over 1.92 is 33.3% of the ceiling it could profitably pay. Or split the rule per variation so each cap equals its own ceiling, at the cost of three register rows where there was one. This account takes the first this cycle and the second at the next quarterly pass, and records the choice with the 33.3% figure attached so the trade was priced rather than assumed.

The lesson. A flat cap is not a ceiling; it is a number that happens to sit above or below one, and which it changes every time the margin table moves under scenario S-A9.

WE-3. The runaway that never breached its band (the recorded incident, worked)

Setup. A branded defensive term, discovered floor \$1.10 recorded per rule M3, CVR 18.0% on branded traffic, AOV \$80.00. The term is retagged to Defensive. A C-class ladder keeps it in scope with the target ACOS left at its pre-retag value and walks the bid up 6 consecutive weekly steps at the rule's configured 10% increment.

The walk. \$1.10 to \$1.21, \$1.33, \$1.46, \$1.61, \$1.77, \$1.95. The end bid is 1.10 multiplied by 1.10 to the sixth power, which is \$1.95, and the premium over the recorded floor is \$0.85 per click. At 30 maintenance clicks per week that standing premium is \$25.50 per week, or \$1,326.00 per year. The premium accumulated during the walk is the sum of the six weekly premiums, 0.11 plus 0.23 plus 0.36 plus 0.51 plus 0.67 plus 0.85, which is \$2.73 per click, or \$81.90 at 30 clicks per week.

Why nothing alarmed. At the floor, cost per order is 1.10 divided by 0.18, which is \$6.11, so ACOS is 6.11 divided by 80.00, which is 7.6%. At the end of the walk, cost per order is 1.95 divided by 0.18, which is \$10.83, so ACOS is 13.5%. The Defensive objective band under rule O6 is 10-15%. The runaway ended inside the band, so no efficiency alarm could fire: on a defended term the band is not the control, because scenario S-D1 makes the discovered floor the control and states that branded terms never scale. The rule spent 77% more per click than the floor while reading healthy on every flag the account had, which is why the code-red four-map diagnosis found it and the automation layer that owned it did not.

Handling. The entity's rule action suspends the same day (P6 step 3). The cause is named: the objective changed and the rule's target did not. The bid restores to \$1.10 in one staged step, since drift from a recorded floor restores rather than re-ladders. The retag checklist at SOP-12 P4 and P6 gains the rule-target update step, and P4 here adds branded defended terms at a recorded floor to the A-class standing exclusion. The 10% step is recorded as platform state, not as compliance: rule M3's roughly 10% per week governs tapers downward and v4.0 carries no governed upward step or repeat count, which is AM-3 and AM-2.

The lesson. A rule inside its band can still be outside its mandate, and the number that would have caught this was a floor nobody was diffing.

WE-4. The edge case: truncation with no campaign-level cause

Setup. Illustrative values in the shape of the T0 six-number header. A product with a declared weekly envelope of \$4,259 carries 6 campaigns inside one console portfolio. The Monday cascade reads in-budget at 64.9% on all 6, below the 70% floor of gate G3, so every rate on them is TRUNCATED and no bid conclusion is legal this week.

The wrong first hypothesis. The instinct is a runaway budget rule. P1 and P6 run first and find nothing, correctly: the register carries no budget-class rule at all and the change log shows zero rule-fired budget actions on those entities. Two procedures return clean and the symptom persists, which is the moment the diagnosis widens rather than repeats.

What it was. P10 reads the console portfolio list against the register and finds the portfolio unregistered, carrying a cap of \$500 per day. The envelope is \$4,259 per week, which is 4,259 divided by 7, or \$608.43 per day. The cap suppresses 608.43 minus 500.00, which is \$108.43 per day, or \$759.00 per week, which is 759 divided by 4,259, or 17.8% of the declared envelope. Two budget surfaces were live and enforcing different numbers, which is the silent disagreement P10 exists to end.

The second thing that goes wrong. The correct fix, raising the cap to \$608 per day or cutting the envelope to \$3,500 per week, cannot deploy this week: the product sits inside an armed event window owned by #32, where only changes the event plan names deploy, and mode precedence under rule T8 is not negotiable for convenience. So the cap holds, the suppression is recorded at its dollar value for the window, the 6 campaigns' bid reads are void for the week under gate G3, and the change stages for the day the window closes with the envelope owner named as decider. Registration itself does not wait, because registering an object is not a change to it. Residue. The T6 entry carries the suppressed dollars, the void bid reads, and the staged change with its deploy date. v4.0 carries no portfolio-cap tolerance and no alert bar, so the finding is reported at \$759.00 per week rather than tested against a threshold, which is AM-8.

The lesson. When two procedures return clean and the symptom persists, the missing surface is usually one nobody registered, and an unregistered object is invisible in exactly the way an unregistered rule is.

7. Failure Modes and Escalation

Failure	Detection signal	Response
F1. Unregistered rule (carried from v1.0)	The P1 weekly diff finds a live rule absent from the register	Suspend the same session, then register per P2 or retire in writing. While D6 is OPEN the detection latency is one cycle, not one day, and the row says so.
F2. Rules on excluded scope (carried; the account's recorded incident)	Rule-fired changes on push, event, terminal, settling or held entities, read from the change log against the T2 gate columns	Exclusion applied the same day; the window's changes reviewed for damage against the #38 T+1 diffs. The recorded instance is the zero-order ladder found live on 2 campaigns with open push worksheets. Escalates same day on push or terminal scope, because it spends against a funded plan or a cliff.
F3. Stale target after a retag (carried; the runaway root cause)	A runaway signature at P6, or a target older than the entity's current objective tag	Suspend the entity's rule action; restore the recorded floor or the correct target in one staged step; the retag checklist at SOP-12 P4 and P6 carries the rule-target update so the cause dies at source (WE-3).
F4. Silent rule edit (carried)	The P1 register diff against the T6 change records, or a #38 T+1 mismatch nobody's record explains	Revert to the registered value pending review; reconstruct the record from console history; an unattributable actor routes to #46 the same day as an access finding.
F5. Undocumented divergence (carried)	A rule value off v4.0 with no AGREES, ADJUST or exception recorded, found at P3	ADJUST to v4.0's value, or write the exception with owner and review date. Tolerance is not a state.
F6. Double-touch (carried)	The #38 T+1 diff shows a rule change and a staged change on one entity inside one read window	P9 precedence: restore the audit value, then suspend the rule on that entity for the read window. Recurrence on one rule is fixed by narrowing scope, not by watching harder.
F7. Lever-order violation (new)	A rule whose action pulls a later lever in the rule L1 order than its evidence supports; the standing instance is a pause fired on a click count	Split or suspend the action per WE-1: the bid half stays, the status half routes to the weekly pass as a worklist. A pause needs the rule L7 log naming the exhausted levers, which a rule cannot produce.
F8. Ceiling-blind cap (new)	A flat maximum bid above 1.5 times the tightest maximum profitable CPC in the rule's scope, found at P3 or after a margin move under scenario S-A9	ADJUST the cap to the tightest ceiling in scope, or split the rule per variation. No documented exception is available, because the exception would authorise spend scenario S-R11 forbids (WE-2).
F9. Event-window firing (new)	A D-class taper or a C-class step dated inside an armed window or its 14-day guard, read against the #32 window record and gate G10	Revert the change; the entity enters the exclusion set for the window plus guard; rank gained inside the window is provisional under scenario S-E5 and starts no transition clock.

Failure	Detection signal	Response
F10. Rule firing through a failed gate (new)	A rule-fired change on an entity carrying a G1 quarantine, a G3 truncation, a G6 YELLOW or RED band, or a G9 provisional price flag	Void the change and the week's read for the row; suspend rule action on the entity until the gate clears. Gates outrank every signal beneath them per Section 12 rule P1, and a rule is a signal.
F11. Post-outage firing (new)	Rule-fired pauses or step-downs clustering on the day after a billing or platform outage, cross-read against the #43 restoration record	Suspend rule action on the affected scope until one clean full day has landed; reverse pauses fired on outage data; flag the dates for trend exclusion. v4.0 excludes event weeks under gate G10 and carries no outage exclusion, which is amendment AM-9.
F12. Portfolio-cap collision (new)	Campaign-level truncation with no campaign-level cause, and a portfolio cap below the product envelope, found at P10	Register the object, compute the suppressed dollars, and reconcile the two surfaces to one number in a single staged change. Never leave both live (WE-4).

Escalation. F2 on push or terminal scope escalates the same day. F7 and F8 escalate on discovery, because both describe a class rather than an instance and both were live before they were named. A second occurrence of any failure from the same root cause escalates the checklist gap rather than the rule, per the v1.0 standard carried forward. Any suspension open longer than one full cycle escalates to the PPC Lead with the damage record attached. Anything touching the human-mandatory tier, which is rule creation, re-enable, documented exceptions, portfolio creation and any suspension on push scope, is confirmed per the Authority Matrix before it lands.

8. Outputs and Artifacts

Gates clear with artifacts, not assertions. Every artifact below has a named location, a retention, and a downstream consumer that reads it rather than a folder it rests in.

Artifact	Content and storage	Retention	Downstream consumer
The rule register	Nine fields per rule plus the portfolio section; the automation section of the workbook layer, versioned continuously	Permanent; retired rows kept	Every procedure here; SOP-26 P3 code-red automation check; #38 P6 external-change classification
Weekly diff artifact	Date, rules read, three counts, mismatches with dispositions, manual-read flag, initials	13 weeks rolling, then the quarterly pass supersedes	The weekly pass; the PPC Lead; the quarterly pattern read at P12
Runaway watch report	Flagged entities, step counts, causes, suspensions, damage in dollars	13 weeks rolling	PPC Lead same day; SOP-26 P3 during code red
Reconciliation record	Per rule: outcome, v4.0 section and rule ID, both values, date, manifest state	Permanent, by version	The v4.1 amendment brief; the quarterly pass; #1 governance
Change records (T6)	Rule creations, edits, deletions with reason string, owner, prediction and review date	Permanent, append-only	V5 scoring; #38 attribution; the audit
Suspension packet	Trigger, scope of freeze, root cause, damage record, re-enable conditions	Permanent	PPC Lead; the checklist owner whose gap it names
Exclusion set	Entity list by exclusion class, with the state that put each there	Live, reconciled weekly	SOP-15 P11; #32; SOP-27 lane; the weekly pass
Portfolio register section	Object, members, cap, owner, envelope reconciliation and its dollar delta	Permanent	The envelope owner at the instrument layer; #38
Sync pass artifact	Version from and to, rules read, outcome counts, adjustments, suspensions, amendments opened and closed	Permanent, one per increment	#1 Supersession Map; every SOP whose Element 10 re-validates on the same increment
Quarterly full-pass artifact	The complete register state, exception ages, prediction scores, the quarter's pattern read	Permanent	The next quarter's diff baseline; #44 onboarding

The system's standing gate artifact is the register at its last sync. The weekly artifact is the diff plus the runaway report. A week with neither did not happen, whatever the console shows.

9. Skill-Conversion Block

Block	Specification
Input bindings	The platform live rule list and change log (Ledger dependency D6, OPEN: transcription until it lands); the register with its version and last-synced fields; A2 deployment state; A3 pacing buckets (PARTIAL); A5 threshold flags in-row; A6 objective and Expected CPC (PARTIAL); A8 PAT rows (PARTIAL); B5 coverage counts; B8 economics chain with the D1 margin version and its G9 state; T2 block K2 gate columns (validation, in-budget, eligibility, inventory band, envelope, cap, event mode); T2 block K7 objective; T3 campaign state and objective tags; T5 product log for price and outage events; T6 decision log including the conflict flag; T7 manifest completeness; the #38 T+1 attribution report; the console portfolio list. An absent feed suppresses its dependent procedure and the skill states which, rather than improvising a substitute.
Deterministic logic	The three-way P1 diff and its step-6 tree; the three reconciliation outcomes and their per-class check order, lever first; the L2 bound test at both registered bounds; the ceiling test, cap against contribution multiplied by CVR, flagged above 1.5 times; the exclusion state machine across push, event, terminal, hold and floor states; the runaway signature as consecutive same-direction steps per entity per window; the P9 precedence resolution; the P10 envelope arithmetic, portfolio caps summed weekly against the declared envelope; the exception ageing test at P12. The skill computes none of its own thresholds and version-pins v4.0 in every reconciliation record.
Human checkpoints (Authority Matrix rows)	The row for automation rules (#37) reads: PPC Lead decides, incoming operator executes and reviews, PPC Manager executes, CEO and BM owner informed. So: AI-executable with audit are the P1 transcription and diff, the reconciliation drafting, the runaway detection and its damage arithmetic, the exclusion reconciliation from state feeds, the P10 envelope arithmetic, and every artifact draft. Human-mandatory are rule creation and re-enable (PPC Lead decides), documented exceptions and their renewals, suspensions on push scope, portfolio creation per the structure-creation row, and the P13 posture change on a code-red declaration, which the code-red row gives the PPC Lead to close. The bulk-deployment row governs anything this SOP stages through #38. Under the vacancy rule an unfilled owner's decisions revert to the PPC Lead and execution to the PPC Manager.
Cadence	Weekly: the P1 diff and the P6 runaway watch inside the Monday cascade, before T2 verdict work. Same day: exclusion updates on every state change, P9 collisions on the T+1 report, P13 on a declaration. Per increment: the P11 full sync within one working day, closed inside five. Quarterly: P12 regardless of increments. Inside Recovery Mode the confirm thresholds tighten per rule R3 and the exclusion reconciliation runs daily rather than weekly.
Outputs	Drafted register updates and diff artifacts queued for confirmation; drafted reconciliation rows with their v4.0 citations and both values; runaway packets with entity, step count, cause hypothesis and damage in dollars; exclusion set deltas; portfolio reconciliation lines with the dollar delta; F1 to F12 alarms with evidence lines.
Alarm specification	Each alarm carries the entity or rule, the signal value, the governing gate or rule reference, and a routing tier. Same day to the operator queue: F1, F2, F9, F10, F11. Into the weekly worklist: F3, F4, F5, F6, F12. To the PPC Lead on discovery with the damage record: F7 and F8, because both name a class rather than an instance.

10. Version Governance

Owner: PPC Lead. The register versions continuously and independently of this document; this document reviews on the triggers below, each specific to this component rather than to the calendar.

- Any v4.0 increment, which re-validates Section 1.3's class table, the Section 5 mapping and every reconciliation outcome. This SOP increments before the next weekly pass runs against the new version.
- Any change to the closed verdict vocabulary, which re-validates Section 5 string by string.
- Ledger dependency D6 landing, which rewrites P1 from transcription to a scheduled diff and retires the manual-read flag, the one-cycle latency statement and part of Section 1.1.
- Any automation-platform capability change that adds a rule class, a new action type, or a change to how bounds are expressed, per P14.
- Any Authority Matrix re-issue touching the automation-rules row, the structure-creation row or the code-red row, which re-writes Elements 9 and 11.

6. The arrival of a standalone budget and envelope procedure for component #31, which would replace the instrument-layer reconciliation in P10 with a procedure citation and close part of conflict C3.

7. Two suspensions from the same root cause in one quarter, which is evidence that a procedure here failed rather than a rule.

Registered in the Supersession Map. v2.0 change note: supersedes v1.0 in full; procedures expanded 7 to 14, adding P2 registration, P7 suspension and re-enable, P10 portfolio governance, P11 sync, P12 quarterly pass, P13 code-red posture and P14 platform, seat and outage handling; failure modes expanded 6 to 12; worked examples expanded 1 to 4; the 5-row local table replaced by the full 25-string closed-vocabulary mapping; Elements 11 and 12 added, neither of which exists in any shipped SOP; the ungoverned consecutive-step reference retired; the B-class lever-order defect named and resolved; the flat-cap ceiling test added; the portfolio-object fold-in absorbed from the Gap Register.

11. Operator Ramp

Built from the Authority Matrix rows that name this component's decisions. Week 1 is observation and artifact familiarity, week 2 is supervised execution with the checkpoint holder reviewing, week 4 is unsupervised execution with the checkpoint holder auditing. The checkpoint holder is the PPC Lead, who holds the D on the automation-rules row; the incoming operator holds E and R on that same row, which is the competence this ramp builds.

Week	Demonstrates	Evidence the checkpoint holder accepts
Week 1	Artifact familiarity and the register's shape. Reads the register cold and states, per rule, its class, its governing v4.0 section, its reconciliation outcome and its scope exclusions. Observes one full P1 diff and one P6 runaway watch. Locates every input this SOP binds: the change log, A2, the T+1 report, the T2 gate columns, T6, T7 and the console portfolio list.	A written read of 5 rules chosen by the checkpoint holder, each naming the v4.0 section by heading and both values, with no outcome asserted from memory. A stated list of which procedures degrade while D6 is OPEN and by how much.
Week 2	Supervised execution on the automation-rules row's E. Runs P1 end to end including the step-6 tree, P4 on one live state change, P5 on one edit, and P6 including the damage arithmetic. Drafts one P3 reconciliation across a full class. Observes one P9 collision resolution and one P10 portfolio read.	The week's diff artifact and runaway report, reviewed line by line before they are filed. One reconciliation set the checkpoint holder can accept without changing an outcome. One correctly refused action: an edit, exclusion or re-enable the operator declined to make because it sits above the E without D line, since under the matrix's standing rule an E without D executes staged decisions only.
Week 4	Unsupervised execution with audit. Owns the weekly P1, P4, P5, P6 and P9 cycle and the P10 weekly portfolio diff. Runs a full P3 pass across all four classes and drafts the amendment items it produces. Executes a suspension per P7 and prepares the re-enable packet for the PPC Lead's decision. Carries the P13 posture actions if a declaration lands.	Four consecutive weekly artifact sets audited after filing with zero unexplained mismatches against the #38 T+1 diffs. One suspension packet complete with root cause, damage in dollars and re-enable conditions. One amendment item written to the standard of Section 13: the gap named, the v4.0 section that should carry it named, and no invented number.

Live-defense questions. These test the reasoning under the thresholds, not recall of their values. A correct answer names the rule and explains why it binds; a recited number without its reason is a fail.

1. A rule's cap is \$1.20 and the term's maximum profitable CPC is \$1.46. The cap is under the ceiling. Give two separate reasons that rule may still be wrong, and say which you would check first and why.
2. A defended term reads 13.5% ACOS inside a 10-15% band while its bid sits 77% above its recorded floor. Which number is the control on that term, and what does the other one tell you?
3. A rule pauses a keyword at zero orders after 30 clicks. Name the two v4.0 rules that decide whether that pause is legal at all, and explain why the click count is the second question rather than the first.
4. The Monday audit sets a bid and the rule moves it back that night. Which value stands, what exactly does the T+1 diff prove, and what does the register change next week so the question does not recur?
5. A portfolio cap and a product envelope disagree by \$759 a week. Explain why that is a finding rather than a setting, who decides it, and why registering the object does not wait for an armed event window to close.
6. The register records AGREES on a rule. State precisely what that claim asserts and three things it does not assert.
7. Explain why a rule that acts more slowly than v4.0 permits can hold a documented exception, while one that acts faster cannot.
8. Ledger dependency D6 is OPEN. Name the one procedure that degrades, state by how much, and explain why the other thirteen do not.
9. A margin table refresh lands and no rule was edited. Name every register field that may have become false, and the section that makes it so.

10. You find a rule firing correctly, inside bounds, on the right lever, against a target that matches v4.0, on a campaign carrying an open push worksheet. What is the finding, and why is it not a reconciliation outcome?

12. Interface Contracts

Every row states the sending owner, the receiving owner, the artifact, the timing, and what happens when it does not arrive. The last column is the one the corpus has always been missing.

ID	Sender to receiver	Artifact	Timing	On non-arrival
I-1	Automation platform, via Data Operations, to this SOP	Live rule list and change log; the machine-readable export is Ledger dependency D6	Weekly, before the Monday packet-ready read	D6 is OPEN today, so P1 runs as a dated console transcription with the manual-read flag set and F1 detection stated at one cycle. If even the console read cannot be taken, no reconciliation outcome is recorded that week and the register carries an unread-week marker.
I-2	SOP-15 P11 to this SOP	Exclusion-state notices for protected auto scope	Same day as any auto-side state change	The affected auto groups stay in B and C scope by default and F2 becomes likely; the weekly P4 reconciliation against the T3 states is the fallback detector, one cycle slower.
I-3	SOP-15 to this SOP	Per-group automation rule reporting	Weekly	The auto-side runaway read runs on account-level change counts only, so a per-group runaway is detectable only after it moves the group total. Stated in the runaway artifact.
I-4	#38 to this SOP	T+1 verification diff with every change attributed to session or rule	The morning after every deployment	P6 and P9 do not run: rule-fired and human changes cannot be separated, F4 cannot be distinguished from platform lag, and precedence under rule P3 is unenforceable. Rule re-enables freeze until the diff resumes.
I-5	#28 to this SOP	Ladder hold notice naming which rows are frozen	Same day as issue	The D-class rule can fire a rung onto a term already inside a scheduled descent, compounding two steps into one week. P8 step 6 is the manual check that catches it late.
I-6	#28 to this SOP	The rung value for a flagged term	Inside the ladder cycle	The P8 flag holds and no rung deploys. This SOP never derives a rung to fill the gap.
I-7	This SOP to SOP-26 P3	The rule-change log for the affected product, and any suspension executed under the protocol	Same day as the code-red declaration	One of the protocol's four diagnosis maps cannot be built, and a runaway root cause stays unexcluded while the cut sequence runs. The protocol proceeds with the gap named in its record.
I-8	#32 to this SOP	Armed-window record with dates, and the settling rows at window close	At arming, and at the end of the guard period	Entities cannot be excluded on time and F9 becomes likely. The standing default is to exclude the product's entities from B and C classes on any unconfirmed window signal until the record arrives.

ID	Sender to receiver	Artifact	Timing	On non-arrival
I-9	SOP-27 and the tier-deciding system to this SOP	Terminal and LTSF tags with their dated scope	Same week as the tag	Terminal entities keep efficiency rules against a liquidation objective, which contradicts the objective function outright. The fallback is the weekly T2 read, and the row carries the gap. Per conflict C6 this SOP infers no tier criteria of its own.
I-10	#43 to this SOP	Billing incident and restoration records with the outage window	Same day as incident and as restoration	F11 goes undetected: pauses and step-downs fired on outage-distorted data stand as if they were evidence. The affected dates carry no trend-exclusion flag and the following week's reads inherit the distortion.
I-11	#46 to this SOP	Access reconciliation diff naming automation-platform seats and any credential rotation	Weekly, and same day on any departure or role change	Register rows keep owners who no longer hold seats. P14 step 3 cannot run, and the vacancy rule applies by default: decisions revert to the PPC Lead and execution to the PPC Manager until the diff lands.
I-12	The instrument layer (T0 envelope named cell, T1 caps, Account Rollup A2) to this SOP	Declared weekly envelope and family caps for P10	Weekly	Portfolio caps cannot be reconciled. The standing response is that unreconcilable caps are removed rather than guessed, so the envelope governs alone and the portfolio object becomes inert.
I-13	The Economics and Context Pack to this SOP	D1 margin table with its version and G9 freshness state	Monthly and on event	Ceilings are stale, every cap approval carries the G9 provisional flag, and F8 risk rises silently because a cap that was legal can become illegal without anyone touching the rule.
I-14	This SOP to the weekly pass and T6	Register states, exclusion set deltas, suspensions and the diff artifact	Weekly, before T2 verdict work	The pass reads bids without knowing which entities a rule touched overnight, so prior values are unreliable and the cycle's verdicts inherit the uncertainty. This is the interface whose absence caused the account's recorded runaway to survive six weeks.
I-15	This SOP to #44 operator onboarding, due 2026-09-25	Element 11 ramp content and the live-defense set	On #44's build	The certification carries no automation module, and operators take the automation-rules E without a demonstrated competence behind it.
I-16	This SOP to #1 governance and the Supersession Map	Sync and quarterly pass artifacts, and every amendment item raised	Per increment and per quarter	The v4.1 amendment brief is assembled without this component's thirteen items, and the same gaps are rediscovered by the next session that touches a rule.

13. Amendment Register: Items for v4.1

Every item below is a number or a rule this SOP needed and v4.0 does not contain. None is written into the procedures above. Each names the section that should carry it.

ID	The gap	Where it belongs and why it matters here
AM-1	Per-ad-type zero-order pause parameters: the click count and the window length that make a pause legal	Section 8.9 scenario S-H7 conditions on zero-sale spend above a threshold it does not state; Section 3 family F11 gives waste as a percentage. The five B-class rules therefore run on numbers with no source. Highest priority in this register.
AM-2	A consecutive-step count that defines a runaway	v4.0's consecutive-period language governs weeks in scenario conditions, never steps in a ladder. v1.0 cited a reference that does not exist; P6 now detects the pattern and records its own convention rather than asserting a threshold.
AM-3	A governed upward step size and a per-window repeat cap for rule-fired increases	Section 9 rule L2 gives 15-25% for the unlock case and a 50% per-change limit; Section 10 rule M3 gives roughly 10% per week for tapers downward. Nothing governs repeated same-direction raises, which is the exact shape of the account's recorded runaway.
AM-4	Admission of a fixed-dollar step form, or a rule that steps are expressed in percent only	Section 9 rule L2 is expressed in percent; the platform's inch-up is a fixed increment. WE-1 shows the divergence is 16.7% at one bound and 3.8% at the other.
AM-5	A requirement that any platform maximum-bid field be derived from the entity's own ceiling	Section 10 rule M1 derives the ceiling per variation. A flat cap is a different object and its legality changes whenever the margin table moves under Section 8.1 scenario S-A9.
AM-6	Whether a rule-fired change owes its own justification record	Section 12 rule P10 states that a verdict without a J-complete Master Action Record does not deploy, at any tier, in any mode. A standing rule deploys many changes against one registration record, and v4.0 does not say whether that satisfies P10 once or fails it every time.
AM-7	Formal admission of register-state vocabulary	REGISTERED, SUSPENDED, RETIRED, AGREES, ADJUST, DOCUMENTED EXCEPTION, EXCLUDED and IN SCOPE describe rules, not keywords, and sit outside the closed verdict vocabulary of Section 13.1. They are declared in Section 5 rather than used as verdicts.
AM-8	A portfolio-cap tolerance against the product envelope, and its alert bar	Section 5 gate G7 reconciles campaign budgets to the envelope and alerts on pacing deviation above 25% against the day-of-week curve. It does not name the console portfolio object, which is a second enforcement surface (WE-4).
AM-9	An outage exclusion window for rule reads and for trend verdicts	Gate G10 excludes event weeks and the 14 days following. No equivalent exists for a billing or platform outage, though the distortion is larger and its direction is the one the pause and ladder classes react to hardest.
AM-10	A maximum life for a documented exception	Section 14 rule V3 governs how the framework changes, not how long a divergence may stand. P12 ages exceptions against their own review dates, which is a local control where a governed one belongs.
AM-11	The automation-platform export schema (Ledger dependency D6)	Not a v4.0 item but a dependency specification: the field list that makes P1 a scheduled diff instead of a transcription, and the schema-change trigger that makes it a P14 event.
AM-12	A human-confirm tier for the standing decision, not only for the action	Section 12 rule P5 confirms actions above 25% on bids and \$50 per day on budgets. A rule is one decision that will produce many such actions, and no tier governs the decision itself.

ID	The gap	Where it belongs and why it matters here
AM-13	Whether a staged manual decision suspends rule action for its read window	Section 12 rule P3 states the audit value wins and automation resumes from it after upload. Resuming from a value is not the same instruction as standing off it, and the difference is the double-touch failure. Reported in full below.

13.1 Conflicts found beyond the pre-declared set

Three findings. No conflict identifier and no Ledger version is minted here: identifiers and versions are assigned at close-out, never inside a build session.

1. v4.0 rule P3 against SOP-37 v1.0 P7, on what happens after a collision. Rule P3 says the audit value wins and automation resumes from it after upload. v1.0 P7 says a staged manual decision suspends rule action on that entity for the decision's read window. Resuming from a value and standing off it are different instructions, and the gap between them is exactly where the double-touch failure lives: a rule that resumes immediately can re-move the audited value inside the window that was supposed to measure it. Resolution taken: v4.0 governs the precedence question, so the audit value is restored first, unconditionally; the read-window suspension is carried as this SOP's implementation mechanism and declared as such at P9 step 4, pending AM-13.
2. Two D-number namespaces collide across the corpus. The Cross-Reference Ledger Section 2 numbers external dependencies D1 to D7, where D6 is the automation platform rules export that gates this component. The Metric Expansion Catalog numbering used throughout the Master PPC Data Architecture numbers field groups D1 to D12, where D6 is the twin-set and sibling view. D1 and D2 happen to coincide in both schemes, which makes the collision harder to see rather than easier, and D3, D5 and D6 do not. This component is the one most exposed to it, because its own blocking dependency is called D6 and its manifest reads a different D6 every Monday. Recommended fix at the source: prefix the Ledger's dependency identifiers so the two spaces cannot be confused.
3. The Ledger's header version narrative disagrees with its own row versions and with the pack's state note. The attached Ledger header carries v2.9 and describes that version as the Wave 1 return of component #43 with the Wave 2 gate met, while Section 1 already carries rows returned through Wave 5A at versions up to v2.9, including the row that closes Wave 5A. The pack's version note separately describes the pack state as Ledger v2.5 with Wave 4 closed and Wave 5A open. All three cannot be current at once. This does not change any boundary this session read, since every row used here was read at its own stated version, and it is reported because a session that trusted the header narrative rather than the rows would take a different view of which components are safely citable.

13.2 Already-registered items this build re-encountered

Reported once, not re-opened: the gate-numbering divergence, where the operating corpus uses G12 for the event-window freeze while v4.0 Section 5 assigns G10 to event exclusion and G12 to MKL hygiene, is already registered as an open conflict by the SB, SBV and SD playbook build; this document cites v4.0's own numbering throughout. The closed-vocabulary count divergence across 18, 25 and 26 strings is already registered by the #41 and #46 builds; this document maps the 25-string vocabulary of the Keyword Decision Record Template v3.0, quoted verbatim. Component #31's dangling procedure citations remain conflict C3, and P10 reconciles to the instrument and tab rather than to a #31 procedure number.

Companions. Master PPC Decision Criteria System v4.0 (all thresholds, bands, scenarios and verdict issuance) · Keyword Decision Record Template v3.0 (the closed verdict vocabulary quoted in Section 5) · PPC Authority Matrix v2.0 (Elements 9 and 11) · Master PPC Data Architecture v1.0 and the PPC Command Workbook TEMPLATE v1.0 (manifest IDs, tabs and named cells) · PPC Cross-Reference Ledger (boundaries and Element 12) · PPC Strategic Doctrine v1.0 (the why layer) · SOP-38, SOP-28, SOP-26, SOP-27, SOP-15, SOP-29, SOP-32, SOP-43, SOP-46 (boundaries, by number and one clause).

Inspiratek & Ecotero LLC | Confidential | Page 18 of 18