

SOP-38: Bulk-File Deployment Mechanics

Version 2.0 · 2026-08-07 · Owner role: PPC Lead · Supersedes SOP-38 v1.0 (2026-07-24) in full · Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds and logs eleven ungoverned operating numerals as v4.1 amendment items rather than writing them.

1. Verdict and Scope

VERDICT: deployment is the only step in the system where a correct decision can still cost money, because every other component is judged on what it decided and this one is judged on what actually landed. This SOP governs the path from a staged verdict to a value live in the advertising console and proven live by a next-day diff: the staging contract every other component's output must satisfy, session pre-flight, the bulk pull and its pre-state snapshot, transcription by entity ID, the eight pre-upload validations, the effective-CPC ceiling scan, upload and processing-report reading, delta classification and session close, the T+1 verification diff, mismatch disposition, rollback, the compressed same-day path, console single-value edit control, and the session record with its state write-back. Two rules outrank everything below them. First: a deployment session closes only when success count equals staged count, or when every difference between the two carries a written row, cause, and disposition. Second: no deployment is complete until the T+1 diff proves the values landed, because an unverified deployment and a failed deployment are indistinguishable from inside the workbook.

Version 2.0 change basis. v1.0 measured 1,808 words across 8 procedures, 6 failure modes, and 1 worked example, against the SOP Build Kit v1.0 standard of 12 elements at 10 to 18 pages and against the exemplar floor of 6,367 words, 12 procedures, 10 failure modes, and 3 worked examples set by SOP-12 v2.0. v2.0 expands procedures 8 to 14, failure modes 6 to 12, worked examples 1 to 4, and the verdict mapping from a 4-row session-control table to the full 25-verdict closed vocabulary with a stated cross-map to the 18-string list in v4.0 Section 13.1 "The Closed Verdict Vocabulary". Elements 11 and 12 are added; neither existed in v1.0 nor in any of the 23 shipped SOP documents.

The standing rule. The workbook is the source of truth and the bulk file is transport, never record. A value's authority comes from T2 K9 and the T6 log, not from the sheet that carried it to the console, and a bulk file retained as though it were a decision record is a second source of truth by another name. This rule governs every procedure below and is the reason P14 writes A2 from verified values rather than from the file that was uploaded.

1.1 Why This Component Carries No Decision Rights

The PPC Authority Matrix v2.0 Section 1 row "Bulk deployments (#38)" is the only operating row in the matrix with no D cell: PPC Lead holds R, the incoming operator holds E, the launch and CA operator holds E on its own scope, the PPC Manager holds E, and neither the CEO nor the BM owner appears. That absence is deliberate and it is the whole design of this component. Matrix Standing Rule 2 states it directly: an E without a D "executes staged decisions only; the #38 precedence rule applies." A transcriber who disagrees with a staged value routes it back to its owner and never edits it, at any tier, in any mode. The corollary sets the failure surface this SOP defends: because this component decides nothing, every defect it can produce is an accuracy defect, and accuracy defects are silent until money has already moved through them.

1.2 The Three Deployment Paths and Which One Owns Which Change

Path	What travels it	Owning component	Proof of landing
Manual bulk session	Every staged change set produced by the operating SOPs: bids, budgets, states, placement multipliers, negatives, creations, target and SKU swaps	This SOP, P1 to P14	P9 T+1 diff with zero unexplained mismatches
Automation rule firing	Standing rules in the automation platform register: inch-up, pause-at-zero, target-ACOS ladders, rank-based taper rung one	#37, rule register and change control	P9 diff attributes the change to the rule; #37 P7 precedence resolves the overlap
Console single-value edit	Genuine urgencies only, one value at a time, never a batch	This SOP, P13, with same-day back-fill into a session record	P9 diff plus the back-filled session record; an unmatched value is F4

A change that travels no path is not a change, it is an incident: the P9 diff finds a value no session record and no rule log explains, and F4 opens with an attribution question rather than a correction.

1.3 Boundaries With Adjacent Components

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, gate, scenario, lever limit, and verdict string. This SOP transports and proves; it never decides. A threshold appearing in this SOP instead of v4.0 is a defect, per SOP-12 v2.0 Section 1.3.
#37 Automation platform rules	Owens the rule register, the reconciliation states, scope exclusions, and rule change control. This SOP owns the manual path and supplies the T+1 attribution that #37 P7 uses to separate rule-fired from session-fired changes.
#12, #13, #14, #15, #16, #17-19, #29	Owens what changes and why. Each hands this SOP a staged change set meeting the P1 contract and consumes the P9 diff. This SOP never re-derives their values and never re-reads their evidence.
#28 Placement and bid modifiers	Owens ladder step derivation, exit-bid discovery, and modifier sizing. This SOP receives the resulting values and runs the P6 ceiling scan against them as an arithmetic check, never as a re-decision.
Budget envelope (instrument layer)	Envelope law lives in the instrument: T0 zone H and Account Rollup A2 per Ledger conflict C3. This SOP names the instrument and tab and never cites a #31 procedure number, because #31 carries no procedural document.
#32 Deal and event window	Owens the event plan, its ramp, caps, and reconciliation. Inside an armed window this SOP deploys only what the plan names, and every such row carries the event reference in its reason string.
#6 V5 prediction scoring	Owens the scoring rubric. This SOP supplies execution-cause routing: when a scored prediction misses because a value did not land, the miss is an execution cause and returns here, not to the deciding SOP.
#39 Data feed and manifest	Owens feed definitions and the T7 manifest field groups. This SOP reads T7 status and never re-defines a field group.
#43 Ad-account billing continuity	Owens billing failure detection and the INCIDENT-OPEN broadcast. That broadcast freezes new deployments account-wide until #43 P8 passes; this SOP holds rather than working around a frozen account.
#46 Access and permissions	Owens the Access Register and the unrecognized-actor response. This SOP's T+1 attribution report is the real-time detector #46 P10 consumes.

#34 Canada marketplace annex	Owns the Canada delta. The staging contract, validations, and verification standard in this SOP apply unchanged to the CA block; only the marketplace field and the deploying role differ.
------------------------------	--

2. Trigger Conditions

Every trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. Lag values in this table are calendar facts drawn from the Master PPC Data Architecture v1.0 Section 12 "Assembly cadence" and from the same-day language of the v4.0 gates that fire them; none is a decision threshold and none is invented here. Where the corpus states no lag, the cell says so and the item appears in Section 10.3.

Trigger	Detecting artifact	Maximum lag to procedure start	Procedure	Source of the lag
Weekly deployment	KDR Cycle File generated Monday PM with its Master Action Queue; T2 K9 verdict rows carrying campaign_moves	Tuesday of the same cycle, after the BM huddle	P2 to P9, full session	Data Architecture v1.0 Section 12 "Assembly cadence"
Graduation chain	GRADUATE FROM DISCOVERY verdict in T2 K9 with the #30 graduation record attached	Same day as the verdict	P12 compressed session	SOP-12 v2.0 Section 2, graduation-arrival row
Funding release	M5 queue entry flips to FUNDED with a dated release in T0 zone H	Same day as the release	P12 compressed session	SOP-12 v2.0 Section 2, funding-release row
Inventory RED	T2 K2 inventory band flips to RED on a targeted variation; D2 names the ranked backup	Same day	P12, with the RE-POINT VARIATION reason string	v4.0 Section 5 gate G6 "Inventory adequacy": RED "re-points ranking campaigns to the ranked backup the same day"
Eligibility flip	T3 C2 eligibility status changes with its reason code	Same day	P12 for the halt rows; the listing task is not this SOP's	v4.0 Section 5 gate G4 "Eligibility": the listing fix task opens same day
Event window action	#32 event plan names a dated action inside an armed window	The date the plan names	P12, event reference mandatory in every reason string	v4.0 Section 5 gate G10 "Event exclusion" and the #32 plan
Code-red action	Recovery Mode entry logged with its entry TACOS and projected weekly overspend	Same day, at the tightened confirmation tier	P12 under P2 step 6 authority re-read	v4.0 Section 11 Recovery Mode R1 and R3
Verification	Any closed upload session in the session record	The next day, unconditionally	P9, then P10 on any mismatch	This SOP; the T+1 cadence is the component's own mechanic
Rollback need	A P10 deployment-error classification, or a V5 score of WRONG-DIRECTION in T6 naming execution rather than judgment as the miss cause	Same day as the classification	P11	v4.0 Section 14 rule V1 scoring plus this SOP's P10
Account freeze	#43 INCIDENT-OPEN broadcast	Immediate stop; no new session opens	P2 step 2 refuses to open the session	Ledger row 43 interface I-6, "freezes new deploys account-wide until P8 passes"
Held-row return	A row held at P5 or P6 re-enters the staged set with a corrected value from its owner	The next scheduled session	P1 re-validation, then the normal path	No governed lag exists in v4.0; logged as AM-4

3. Preconditions and Gates

Each precondition names its manifest field group by Catalog ID as carried in the Command Workbook T7 MANIFEST tab, and states the observable that proves it passed. A precondition proved by assertion rather than by artifact has not passed. Where a field group reads PARTIAL or MISSING in T7, the dependent row class does not deploy and the session record states what it could not see, per v4.0 Section 5 gate G8 "Packet completeness".

Manifest ID	Field group and its T7 source feed	Requirement before the session opens	Observable that proves it passed
A2	Deployment state in-row (bid, budget, state); SP bulk sheets from the Raw Pull Pack	Current console state present for every affected entity. A2 is the single field group this SOP both consumes and writes.	T7 status reads PRESENT with a refresh timestamp later than the last closed session's T+1 diff
A5	Threshold flags in-row; computed in the workbook	Band and flag columns computed, so the ceiling scan at P6 reads flags rather than re-deriving thresholds	T7 status PRESENT; T2 K6 max profitable CPC and Expected CPC populated on every bid row in the staged set
A6	Objective plus per-term Expected CPC; objective owners and the T0 reference	Every staged row's entity carries exactly one declared objective, because the reason string cannot be completed without it	T3 C1 objective tag non-blank on every affected campaign; UNTAGGED rows return to their owner before staging closes
A7	Eligibility in-row; SP bulk	No affected target sits ineligible or suppressed	T3 C2 eligibility column clear across the affected-campaign list
B5	Coverage counts per match type; the campaign registry	Registry counts current, so a creation row can be proved legal and a negative can be proved to have a live parent	T2 K8 counts refreshed in the same pull as A2
B8	Row economics, the break-even chain; derived from the T0 named cells	T0 economics spine live: AOV, MarginPerOrder, BE_ACOS, TargetACOS, MaxAccACOS, ExpectedCPC	T0 cells C18 to C24 populated and the margin table version stamped
D1	Per-variation margin; Finance	Margin basis known for every bid row that the ceiling scan will test	T7 status with its G9 state recorded; a PROVISIONAL margin carries onto every ceiling-referencing row in the session record

D7	Decision log and V5 scoring; T6 in the workbook	T6 reachable and append-only, because every deployed row writes a prediction row here	T6 accepts the session's rows and the prior session's rows are present and scored or dated
-----------	---	---	--

Gates in force at session open

Gate	v4.0 name and rule	Effect on this session
G1	Section 5 "Validation quarantine": impossible metrics quarantine the row, flag produced, no recommendation	A quarantined row's staged change does not deploy. It is not held for a value correction; it is returned, because the evidence under it is void.
G6	Section 5 "Inventory adequacy": GREEN scales, YELLOW freezes push spend at current level, RED tapers and re-points the same day	YELLOW blocks any staged row that raises push spend. RED admits only taper, re-point, and defense rows, on the same-day path.
G7	Section 5 "Envelope and pacing": campaign budgets reconcile to the product envelope; pacing deviation above 25% versus the day-of-week curve alerts same day	A staged budget set whose sum breaks the envelope does not deploy as a set. It returns whole; this SOP never deploys the profitable half.
G8	Section 5 "Packet completeness": DEGRADED packets, under 85% of gate-relevant fields, produce flags only	A DEGRADED manifest closes the session to verdict-driven rows. Gate-driven rows and data-repair rows remain deployable, per v4.0 Section 12 rule P8.
G9	Section 5 "Price stability": a price change on the targeted variation within the trailing 7 days marks CVR and ACOS reads PROVISIONAL	The PROVISIONAL flag rides into the reason string of every affected row so the next cycle's scoring reads it. It does not by itself block deployment.
G10	Section 5 "Event exclusion": event weeks and the 14 days following are excluded from trend verdicts	Inside an armed window only plan-named rows deploy. The corpus alias for this freeze is G12; both names point at the same freeze, per Ledger conflict C9.
G12	Section 5 "MKL hygiene": every analyzed term exists in the current MKL with a syntax tag and relevancy score	A staged row on an untagged term returns to MKL maintenance rather than deploying, because its verdict was issued outside the taxonomy.

Gate-name discipline. Ledger conflict C9 records that v4.0 Section 5 and the operating corpus assign G9, G10, and G12 to different subjects. This SOP names every gate by its v4.0 number and v4.0 name and gives the corpus alias at first use. v1.0 of this SOP additionally used "G2" for bulk-file freshness, which matches neither the v4.0 assignment, sample sufficiency, nor the corpus alias used by SOP-12 v2.0 Section 3, manifest freshness. That third usage is retired here: file freshness is a P3 precondition of this SOP, not a gate, and it carries no G-number until v4.1 assigns one. The gap is logged as AM-11.

4. Step-Level Procedures

Fourteen procedures. Time estimates assume a session of 20 to 40 staged rows on one marketplace and are execution estimates for a trained operator, not governed limits. A first unsupervised session runs roughly double.

P1. The staging standard: the contract every component's output meets

Time: not a session step. This is the contract validated at P5 step 1; a producing SOP that breaks it discovers the break here, one cycle late.

1. **One row per change.** Never one row per keyword, never one row per decision. A verdict that moves a bid and a budget produces two rows. The columns, in order: marketplace | entity type | campaign ID | ad group ID | keyword or target ID | entity name (reference only, never the match key) | current value | new value | change type | reason string | staging owner role | review date | confirmation tier.
2. **Entity type** is one of: campaign, ad group, keyword, product ad, product target, placement, negative keyword, negative target. **Change type** is one of: bid, budget, state, placement modifier, bidding strategy, negative-add, negative-remove, target swap, SKU swap, creation.
3. **The reason string format**, four bracketed segments in fixed order: [SOP-##][verdict string, verbatim][scenario or gate reference][evidence reference]. Worked instance: [SOP-12][TAPER: RANK ACHIEVED, STEP BIDS DOWN][S-R6 to T1][T6 2026-08-04 step 2 of 4]. A row with a missing or malformed segment does not deploy; it returns to its owner. This is not a house-style preference: v4.0 Section 12 rule P10 "Justification completeness" states that a verdict without a J-complete Master Action Record does not deploy at any tier in any mode, and that completeness is checked mechanically before the deploy step accepts the record.
4. **The verdict segment is quoted verbatim from the closed vocabulary and is never paraphrased, abbreviated, or case-folded.** The operative list is the 25-string vocabulary held in the Keyword Decision Record Template v3.0 "Bands & Verdicts Reference" sheet, mirrored exactly in the Command Workbook REFERENCE sheet as the KW VERDICT data-validation list. Verbatim matters mechanically here and nowhere else in the library: the reason string is the only field the T+1 attribution can match on, and a paraphrase makes a landed change unattributable. Compound verdicts are legal where the workbook READ ME admits them; both strings appear, joined by a plus sign, each verbatim.
5. **Creations carry their full settings block as one unit**, staged as a group and never split across sessions: campaign name per the naming convention, portfolio, daily budget, bidding strategy, ad group, advertised SKU, keyword or target with match type, starting bid, placement modifiers, and the pre-load negative set the creating SOP prescribes. A creation that arrives without its negative set is incomplete at P5 and returns whole.
6. **Confirmation tier is stamped by the producing SOP, not decided here.** Per v4.0 Section 12 rule P5 "Human-review tiers", the following are human-confirmed: every 500+ SV keyword action, every gate failure, every structural change, every bid move above 25%, every budget move above \$50 per day, and every LOW-confidence verdict. The routine tail auto-applies with spot checks. Inside Recovery Mode the tightened values of v4.0 Section 11 rule R3 replace them: every bid move above 15% and every budget move above \$25 per day is human-confirmed.
7. **The staged set is a set.** It carries a session identifier, a row count by change type, the producing components, and the T7 completeness score at the time of staging. A staged set without its count block cannot be validated at P5 step 1 and does not open a session.

P2. Session pre-flight: scope, authority, and freeze check

Time: 10 to 15 minutes. Runs before any file is downloaded, because a session that should not have opened costs more than one that opens late.

1. **Read the freeze conditions first, in this order:** an open #43 INCIDENT-OPEN broadcast, an unresolved delta from a prior session on any affected scope, an armed event window touching any affected product, and a Recovery Mode entry on any affected product. The first three block or narrow the session. The fourth does not block it; it re-tiers it.
2. **An unresolved delta freezes new sessions for the affected scope only.** Scope means the campaign IDs carried in the open delta log, not the product and not the marketplace. Rows outside that ID list proceed in the same session.
3. **Assemble the affected-campaign list** from the staged set: the distinct campaign IDs across all rows. This list is the session's boundary for every later step. An entity that is not on it is out of scope, and a change that touches one is scope contamination, failure mode F10.
4. **Read T7 MANIFEST** for the eight field groups in Section 3. Record the completeness score in the session record. If the score places the packet in DEGRADED per v4.0 Section 5 gate G8, the session narrows to gate-driven and data-repair rows and the narrowing is written down.
5. **Confirm the huddle on the weekly path.** The BM owner's R on the verdict rows is exercised in the daily huddle per Authority Matrix v2.0 Standing Rule 3, which is review and not veto; a disagreement routes up rather than holding the session.
6. **Re-read the confirmation tier of every row** against v4.0 Section 12 rule P5, or against Section 11 rule R3 where a Recovery Mode entry is live on that row's product. Rows whose tier is human-mandatory are grouped and their confirmations are collected before the upload step, never during it. The confirming role is the one the Authority Matrix names for the decision class, not the operator running the session.
7. **Open the session record** with: session identifier, date and clock time, marketplace, operator role, staged row count by change type, producing components, T7 score, freeze conditions checked and their state, and the affected-campaign list.

P3. Pull and pre-state snapshot

Time: 10 to 20 minutes, feed-dependent.

1. **Download fresh bulk sheets** scoped to the affected entities and to a window that covers every row's current value. Never edit a previously downloaded file. The reason is arithmetic rather than tidiness: the current values a transcription diffs against are the values that were true at download, and every hour of age is an hour in which a rule, a co-operator, or the platform moved one of them.
2. **Record the download clock time in the session record.** The corpus governs no maximum file age, so the operator records the age at upload and the session record carries it. The absent limit is logged as AM-2 and is not invented here.
3. **Snapshot the pre-state.** Save the untouched fresh download, timestamped, under the session identifier, before any cell is edited. This file is the rollback reference at P11 and it is the only artifact that makes rollback an execution rather than a reconstruction. A session whose snapshot was taken after the first edit has no rollback path and is aborted at this step.
4. **Filter working views** to the affected-campaign list from P2 step 3. Every later count is computed inside this filter, so a filter set wrongly produces counts that agree with each other and disagree with reality.
5. **Confirm the pull covers every staged row.** A staged row whose entity is absent from the pull is a registry disagreement, not a missing filter: the entity was archived, moved, or never existed. It returns to its owner at this step, before transcription, and it is the cheapest place in the whole procedure to catch failure mode F9.

P4. Transcription by entity ID with ordering discipline

Time: 1.5 to 3 minutes per row for value changes; 8 to 12 minutes per creation block.

1. **Match entities by ID, never by name.** Names duplicate and the stage-stack deliberately creates near-identical ones: v4.0 Section 6 rule O4 makes dual objectives on one keyword legal and explicitly tagged, so a hot Ranking exact and a cool Conversions exact on the same term are both correct and both

similarly named. Name matching in that structure lands the change on the wrong twin and the error is invisible until the next cycle reads two campaigns behaving unlike their objectives.

2. **Order the upload in three blocks, in this sequence.** Block 1, creations, so that keywords, targets, and negatives have live parents to attach to. Block 2, value modifications: bids, budgets, placement modifiers, bidding strategy. Block 3, state changes, last. The sequence is a dependency rule, not a preference: a pause executed before a dependent edit on the same entity voids the edit, and the platform reports the voided edit as a success.
3. **Within block 1, order creations parent-first:** campaign, then ad group, then product ad, then keyword or target, then negatives. A negative staged against a keyword ID that block 1 has not yet created rejects.
4. **Transcribe current values from the staged row, then reconcile them against the fresh pull.** Where the staged current value and the pulled current value disagree, that row stops. The world moved between staging and this session, so the verdict was issued against a state that no longer exists, and the row returns to its owner for a re-read. It is not adjusted to the new current value, because the new current value may be exactly what the verdict was reacting to.
5. **Transcribe new values exactly as staged, to the staged precision.** Bid values carry two decimal places. A value that requires rounding to enter the sheet is a staging defect and returns.
6. **Leave every column the staged row does not name untouched.** A bulk sheet carries columns this session has no business in; an unintended edit to one of them uploads as successfully as an intended one.

P5. Pre-upload validation: the eight checks

Time: 15 to 25 minutes. Every check is recorded with its result in the session record; a check without a recorded result did not run.

1. **Count check.** Transcribed rows equal staged rows, reconciled by change type and not only in total. A total that matches while two change types are transposed is the failure this per-type reconciliation exists to catch.
2. **Reason-string completeness.** Every row carries four bracketed segments. Checked mechanically per v4.0 Section 13.3 rule J1, which states that field presence is checked by code and not by reviewers.
3. **Verdict-string verbatim check.** Every reason string's verdict segment matches a member of the 25-string KDR vocabulary character for character. Near-misses fail: a truncated, re-cased, or re-punctuated string is a failed check and not a warning, because the T+1 attribution and the V5 scoring both match on it.
4. **Spot check.** A sample of rows verified end to end against the workbook: right entity ID, right current value, right new value, right reason string, right confirmation tier. The corpus governs no sample size and no sampling rule, so the operator records the count sampled and the selection basis in the session record. The absent sampling rule is logged as AM-1; v1.0 of this SOP asserted a five-row sample with no source and that number is not carried forward.
5. **Ceiling scan.** Every bid row and every placement-modifier row tested per P6. Any row above its ceiling with no sizing block is HELD and returned to its owner with the arithmetic attached. It is never repaired inside the transcription, because repairing it would be this component making a decision.
6. **Lever-limit check.** Every bid row tested against v4.0 Section 9 rule L2 "Bids": a maximum of plus or minus 50% per change and a floor of \$0.50. A staged step outside those limits fails staging quality assurance and returns. This SOP checks the limit; #28 and the deciding SOP own the step size inside it.
7. **Scope diff.** No transcribed row touches an entity outside the affected-campaign list. Run as a set comparison rather than by eye.
8. **Encoding pass on created names.** Every name in block 1 survives a round trip through the upload template with its segments intact, because the naming convention is what the audit filters parse and a name that loses a separator is a campaign that cannot be graded. The corpus states no character-set rule for the reason string or the name segments; the gap is logged as AM-9.

P6. Ceiling scan arithmetic

Time: 2 to 4 minutes per bid row, or seconds where the workbook flag columns are populated. Read the flag; do not re-derive the threshold.

1. **Compute the effective CPC the staged row authorizes,** using the identity in PPC Strategic Doctrine v1.0 Section 6.1: effective CPC ceiling equals base bid multiplied by one plus the placement modifier, multiplied by one plus the dynamic raise. Doctrine 6.1 records that up-and-down bidding can add up to 100% at top of search. The triple multiplication is why a modest-looking bid can authorize a click several times its own value.
2. **Read the ceiling it is tested against** from T2 K6: max profitable CPC, defined in v4.0 Section 3 family F1 as margin dollars multiplied by CVR, and Expected CPC, defined in family F9 as what the position should cost.
3. **Classify the ratio** using the pre-computed CPC ratio bands in v4.0 Section 3 family F1 "Product Economics": at or below 0.8 times ceiling is SAFE; 0.8 to 1.0 times is AT-CEILING; 1.0 to 1.5 times is PUSH-ONLY; above 1.5 times is UNJUSTIFIABLE outside a sized push.
4. **Apply the disposition by band.** SAFE and AT-CEILING proceed. PUSH-ONLY proceeds only where the row's reason string references a sizing block, per v4.0 Section 10 rule M2, which requires every push to carry target rank, required clicks versus actual, budget, weekly loss ceiling, ETA checkpoint, and stop-loss. UNJUSTIFIABLE is HELD and returned regardless of what the reason string claims, because v4.0 Section 8.4 scenario S-R11 classes CPC above 1.5 times ceiling with no M2 record as spend that must ladder toward ceiling before aggression resumes.
5. **Placement-modifier rows are tested on the same identity.** v4.0 Section 10 rule M4 sizes the multiplier so that effective top-of-search CPC stays at or below the lesser of max profitable CPC for the objective and Expected CPC plus a tolerance. v4.0 names that tolerance and does not number it, so this SOP tests against the lesser of the two hard values and records the untested tolerance margin. The missing numeral is logged as AM-8.
6. **A PROVISIONAL margin carries into the result.** Where T7 D1 reads PARTIAL or the G9 clock is live, every ceiling computed in this scan is provisional, the session record says so, and the flag rides into the reason string of each affected row.

P7. Upload and processing-report reading

Time: 5 minutes to upload; 10 to 20 minutes to read the report properly.

1. **Confirm the human-mandatory confirmations are in hand** before the upload action. The upload action itself is human-mandatory in every session, without exception: it is the moment a decision becomes money.
2. **Upload the three blocks in the P4 step 2 order.** Where the platform accepts one file, the block order is the row order inside it.
3. **Read the processing report completely, at row level.** Partial success is the dangerous case and the reason this step is written as a procedure rather than a line. A report skimmed for a success banner hides silently skipped rows, and a silently skipped row is a decision that was made, confirmed, staged, validated, and never executed, while every record in the system says it was.
4. **Classify every non-success into one of three classes.** Rejected rows: the platform names a cause; fix the cause and re-upload inside the same session. Silently skipped rows: no error is named and the row simply did not take; these are found by comparing the success count to the attempted count, never by reading the error list. Format errors: fix the sheet, never the intent, and never fix a format error by changing a value.
5. **Re-upload corrected rows inside the same session** and read the second report with the same discipline. Each re-upload increments an attempt counter in the session record.
6. **Never re-upload the whole file to fix a subset.** A whole-file re-upload re-applies rows that already landed, and where a rule fired in between, it overwrites the rule's change without attribution.

P8. Delta classification and session close

Time: 10 to 20 minutes. Three or more branches, so the decision tree governs.

Decision tree. Compare success count to attempted count.

- **Branch A, success equals attempted.** The session closes on the upload side. Proceed to P9; the session is not complete, only its upload half is.
- **Branch B, success below attempted, cause named by the platform.** Fix the cause per P7 step 4 and re-upload in this session. If the fix requires a value change, the row is not fixable here: it returns to its owner and becomes a delta.
- **Branch C, success below attempted, no cause named.** Identify the missing rows by set difference between transcribed and landed. Each becomes a delta row carrying: row identifier, entity ID, staged value, cause where determinable, disposition, and owner. A delta with no determinable cause is logged as undetermined rather than guessed.
- **Branch D, an error cascade: multiple rejects sharing one cause.** ABORT the session. Do not fix rejects individually inside a cascade, because a shared cause means the staged set or the pulled file is wrong at a level above the row. Fix the cause, then restart from P3 with a fresh pull.
- **Branch E, a stale-file discovery mid-session or scope contamination found after upload.** ABORT and go to P11 rollback for anything that landed, then restart from P3.

1. **The session closes only at success equals staged, or with every difference logged** as a delta row with cause and disposition. There is no third close state and no partial close.
2. **An unresolved delta freezes the affected scope** per Section 3, for the campaign IDs in the delta log only.
3. **Held rows are not deltas.** A row held at P5 or P6 never entered the attempted count. It is logged separately with its hold reason and its owner, and it returns through the next cascade with a corrected value.

P9. T+1 verification diff

Time: 20 to 35 minutes. Runs the day after every deployment, unconditionally, including after a session that reported every row successful.

1. **Pull fresh** on the affected-campaign list. This is a second pull, not a re-read of the P3 file.
2. **Diff row by row: intended value against deployed value**, for every row in the attempted set. Not a sample. The diff is cheap because it is mechanical, and it is the only artifact in the system that distinguishes a deployed decision from a decided one.
3. **Match on entity ID and change type**, then read the reason string from the session record to attribute the row. Where the landed value matches and the reason string is present, the row is verified.
4. **Classify every mismatch per P10** before any correction is staged.
5. **Scan for values the session did not stage.** Any changed value on an affected entity that no session row and no rule log claims is an unattributed change: F4 opens, and it opens as an attribution question routed to #46, not as a value correction.
6. **The verification artifact is the diff report with zero unexplained mismatches.** Explained is not the same as matched: a mismatch classified, dispositioned, and dated is explained. A mismatch noticed is not.
7. **Write the A2 deployment state** from the verified values so the next cascade reads truth rather than intent, and close the session record.

P10. Mismatch classification and disposition

Time: 5 to 15 minutes per mismatch. Four branches, so the decision tree governs. Classification precedes correction in every branch.

Decision tree. For each mismatch, ask the questions in this order and stop at the first yes.

- **Branch 1, does the #37 rule change log show a firing on this entity inside the window?** Then the class is **external change, rule-attributed**. v4.0 Section 12 rule P3 governs: the weekly audit outranks the automation platform between cycles, the audit value wins, and automation resumes from it after upload. Restore the audited value, and raise the scope exclusion with #37 rather than re-restoring it every week.
 - **Branch 2, does another session record or a console-edit log claim this entity?** Then the class is **external change, human-attributed**. Identify the acting role, reconcile the two intents, and route the unlogged half to F4.
 - **Branch 3, does the landed value equal the staged value under a platform transformation?** Rounding to platform precision, a currency or percentage format change, or a modifier expressed differently are transformations rather than mismatches. The corpus states no tolerance band for this test, so the operator records the raw and transformed values and the class as transformed rather than asserting a tolerance. The absent tolerance is logged as AM-5.
 - **Branch 4, none of the above, and the landed value differs from the staged value.** Then the class is **deployment error**, and it is this component's. Stage the correction as a new mini-session through P1 to P9 in full. The week's performance reads on the affected rows are void, because the entity did not run the values the decision assumed.
 - **Branch 5, the landed value equals the pre-state value and no error was reported.** Then the row did not take: the class is **silent skip found at T+1**, which means P8 missed it. Redeploy the row and record the P8 miss, because a silent skip that survives to T+1 is a validation failure and not a platform failure.
1. **Platform lag is a sub-class of branch 3, not a default.** Where a value is expected to settle, the row is re-checked once on a later pull and the re-check is dated in the session record. v4.0 states no lag tolerance; the number of days is not invented here and is logged as AM-3.
 2. **A deployment error triggers a root-cause line** naming which validation would have caught it: count check, verbatim check, spot check, ceiling scan, lever limit, scope diff, or none of them. Where the answer is none of them, the validation set is short by one check and that becomes a version-governance trigger under Section 10.
 3. **Every classified mismatch writes to T6** with its class, disposition, and the date it was closed.

P11. Rollback

Time: 30 to 50 minutes. Rollback is a deployment and receives the full standard.

1. **Trigger.** A P10 branch-4 deployment error, or a V5 score of WRONG-DIRECTION in T6 whose root cause names execution rather than judgment. A wrong judgment is not rolled back here; it is re-decided by its owning SOP and arrives as a new staged set.
2. **Restore from the P3 snapshot**, not from memory and not from the workbook's prior-value column. The snapshot is the only record of what the console actually held before the session touched it.
3. **Stage the rollback as its own session** with the reason string [SOP-38][ROLLBACK][session reference][defect class from P10], then run P5 validation, P7 upload, P8 close, and P9 verification in full. Urgency does not shorten this; it shortens the queue ahead of it.
4. **State rollbacks record the elapsed window in T6.** Un-pausing an entity that was paused for a period does not restore the entity to its pre-pause condition: the paused hours are absent from the next read, and v4.0 Section 5 gate G2 "Sample sufficiency" will judge that read on the clicks it actually contains. Recording the window is what stops the next cycle from reading a short sample as a performance change.
5. **Restore budgets alongside states** where the paused entities carried them, and record the forgone spend across the paused window so the envelope reconciliation at T0 zone H reads correctly.
6. **A rollback does not re-litigate the decision.** It returns the account to the state the decision was made against, and the decision returns to its owner intact.

P12. Same-day and emergency compressed path

Time: 35 to 60 minutes for a set of 1 to 8 rows. Three or more branches on entry, so the decision tree governs.

Decision tree at entry.

- **Is a G1 quarantine or a DEGRADED manifest live on the affected rows?** Then stop. v4.0 Section 12 rule P8 "Data-quality precedence" states that a quarantine or DEGRADED manifest beats every other consideration including urgency, and that the correct urgent action on bad data is fixing the data the same day.
 - **Is a Recovery Mode entry live on the product?** Then re-tier every row to the R3 values before staging: human confirmation at bid moves above 15% and budget moves above \$25 per day.
 - **Is an event window armed?** Then only rows the event plan names may travel this path, and each carries the event reference in its reason string.
 - **Otherwise**, proceed with the compressed session.
1. **The compressed session is the full standard with a smaller set, not a reduced standard.** Stage per P1, pre-flight per P2 steps 1 to 3 and 6, pull and snapshot per P3, transcribe per P4, validate per P5 with every one of the eight checks run and recorded, upload per P7, close per P8, verify per P9 the next day.
 2. **Nothing about urgency removes a step.** The ceiling scan in particular survives compression, because same-day rows are disproportionately push and event rows and those are exactly the rows whose modifiers stack.
 3. **The spot check scales to volume rather than disappearing.** On a set of 1 to 8 rows the spot check and the count check converge: every row is verified end to end and the session record says so.
 4. **The T+1 verification is not compressible.** A same-day session that skipped its next-day diff is an open session, and its scope is frozen under Section 3 until the diff exists.

P13. Console single-value edit control and back-fill

Time: 2 minutes to edit; 10 to 15 minutes to back-fill correctly. The asymmetry is the point.

1. **A console single-value edit is permitted for a genuine urgency only** and only where the urgency cannot wait for P12. It is one value on one entity. Two values are a bulk session with a small file.
2. **Log the edit to T6 at the moment of the edit**, not at the end of the day: the T6 row carries date, level, target, the rule or scenario that fired, the prior and new value, the expected outcome window, the review date, and the acting role and tier. The workbook T6 tab is append-only, so the entry cannot be tidied later.
3. **Back-fill a session record the same day**, containing the same fields a bulk session would carry: session identifier, affected entity, staged equivalent row with its four-segment reason string, and the confirmation tier that applied.
4. **The edit still receives a T+1 diff.** A console edit is not exempt from verification; it is the deployment path with the weakest evidence trail and therefore the one that most needs the diff.
5. **An unlogged console edit is failure mode F4** regardless of the size of the value or the seniority of the role that made it. The detection is mechanical at P9 step 5 and the response is attribution, not correction.

P14. Session record, state write-back, and retention

Time: 10 to 15 minutes.

1. **Close the session record** with: staged count, held count with reasons and owners, attempted count, success count, delta log, upload attempt count, processing reports, T+1 diff report, mismatch classifications with dispositions, and the clock times of the P3 pull and the P7 upload.
2. **Write A2** from verified values only. A2 written from intended values is the single most damaging shortcut available in this procedure, because every downstream gate and every next-cycle current-value reconciliation reads it as truth.
3. **Write T6 predictions.** Every deployed row carries the prediction its verdict issued: metric, direction, magnitude, horizon, and review date, per v4.0 Section 14 rule V1, which states that unscored actions do not accumulate. This SOP does not author the prediction; it carries it from the staged row into T6 so that V5 can score it.
4. **Route execution causes.** Where a prior cycle's prediction scores MISS or WRONG-DIRECTION and the cause is that the value did not land, the score routes to this component and not to the deciding one.
5. **Retain the artifact set.** The corpus states no retention period for session records, pre-state snapshots, or processing reports, and the Ledger row for component #39 carries "retention policy lines" as a pending fold-in. Until a period is set, artifacts are retained and the absence is logged as AM-6. No local period is written here.

5. Verdict Mapping: What This Component Does With Each Verdict

All verdict issuance, every threshold, and every scenario condition live in v4.0. This table maps the full 25-string closed vocabulary to what this component does when a row carrying that verdict reaches the staged set. The vocabulary quoted is the Keyword Decision Record Template v3.0 "Bands & Verdicts Reference" list, mirrored exactly by the Command Workbook REFERENCE sheet KW VERDICT validation list, and matching the exemplar SOP-12 v2.0 Section 5. The cross-map column gives the v4.0 Section 13.1 "The Closed Verdict Vocabulary" string where one exists verbatim, so neither citation is dropped; where the cell reads "no verbatim counterpart", v4.0 Section 13.1 carries no character-identical string and the KDR string is the one written into the reason string.

Zero local thresholds in this table. This component holds no verdict of its own. Its session-control decisions are DEPLOY, HOLD ROW, ABORT SESSION, and ROLLBACK, and all four are execution states rather than verdicts: none of them changes what a value should be.

Verdict, quoted verbatim from the KDR v3.0 closed vocabulary	Deployment treatment in this component	v4.0 Section 13.1 cross-map
QUARANTINE: DATA REFRESH	No row deploys. The staged set for the quarantined entity is returned, not held, because the evidence under it is void per gate G1. A data-repair row may deploy in the same session; a value row may not.	FLAG: DATA REPAIR
UNLOCK TRAFFIC	Bid row; P6 ceiling scan and P5 lever-limit check both apply. The step size is #28's and the deciding SOP's; the plus or minus 50% limit is tested here.	No verbatim counterpart
FIX BUDGET (TRUNCATED)	Budget row, deployed ahead of any bid row on the same entity in the same session, because a bid read through a truncated day is unreadable per gate G3.	No verbatim counterpart
CONTINUE TEST TO 50 CLICKS AT CAPPED BID	Usually a no-change row. Where it stages a budget-assurance change only, that row deploys; a bid change carrying this verdict is a staging defect and returns.	EXTEND TEST
FUND TEST (UNFUNDED PRIMARY)	Creation block or a bid and budget pair at the capped test value the verdict names. Block 1 ordering applies to the creation.	FUND DELIBERATE TEST (50-CLICK CAP)
RANK PUSH: FUND SIZED PUSH	Creation or a bid and budget lift. P6 admits a PUSH-ONLY band ratio only because the reason string references the M2 sizing block; without that reference the row is HELD.	RANK PUSH AT CONTROLLED CPC
RECOVERY PUSH	As above. The pre-OOS baseline reference travels in the evidence segment so the T+1 diff and the later V5 score read against the same baseline.	No verbatim counterpart
TRAFFIC DEFICIT: FEED TO REQUIRED CLICKS	One lever per row per cycle: a bid row or a budget row, never both, unless the verdict text itself names a two-lever program.	No verbatim counterpart
FIX CONVERSION FIRST; DEFER PUSH	Bid and budget rows capping the entity at maintenance. The fix itself is not a PPC row and never enters a bulk session.	FIX CONVERSION FIRST; DEFER PUSH
PLACEMENT FIX FIRST	Placement-modifier rows only; bid rows on the same entity in the same session are a staging defect. P6 step 5 tests the modifier.	No verbatim counterpart
RE-POINT VARIATION	SKU swap at ad level, deployed same day on the P12 path when gate G6 reads RED. Block 2 row; a re-point does not wait for the weekly session.	RE-POINT VARIATION
TAPER: RANK ACHIEVED, STEP BIDS DOWN	One ladder step per session, with the next step's date carried in the evidence segment so T6 reads as a dated chain and each step scores independently.	TAPER TO FLOOR (POST-OBJECTIVE)
HOLD; DEFEND LIGHTLY AT RANK	No row. A staged row carrying this verdict is a staging defect and returns to its owner.	CONTINUE; VERIFY MOMENTUM
DEFEND AT FLOOR	No row unless the T+1 diff or a fresh pull shows drift from the recorded floor, in which case one bid row restores the floor value in a single step.	DEFEND AT FLOOR
BID DOWN TO PROFITABLE CPC; CAP SYNTAX	Bid row to the T2 K6 max profitable CPC value, single step, tested against the plus or minus 50% limit. The syntax-cap half arrives separately as envelope deltas from the instrument layer.	BID DOWN TO PROFITABLE CPC; CAP SYNTAX
INCREMENTALITY LADDER	Dated step-down rows, one step per session. The continue-or-restore call is not made here; this component deploys the step and returns the T+1 proof.	No verbatim counterpart
ADD EXACT + RESTRUCTURE	Creation block plus the steering negatives, staged as one unit and uploaded in P4 block order so the negatives attach to live parents. Splitting the unit across two sessions is failure mode F9's most common origin.	ADD EXACT + RESTRUCTURE
CONSOLIDATE DUPLICATES	A state row pausing the loser, a budget row moving its budget to the survivor, and negative-add rows carrying the lists the survivor lacks. State rows go last per P4 block 3.	No verbatim counterpart
GRADUATE FROM DISCOVERY	Creation block plus steering negatives at source, on the same-day P12 path. The graduation record's discovery CPC seeds the bid the verdict carries; this component does not seed it.	GRADUATE + STEER
NEGATE (PRE-LOAD / REACTIVE / STEERING)	Negative-add rows in P4 block 1 after their parents. Every negative carries its mode in the evidence segment. Parent liveness is proved at P3 step 5, which is where F9 is caught cheaply.	CUT / NEGATE
ENTER SBV ARBITRAGE	Only the SP-exact ladder step-down row travels a bulk session under this SOP's contract, dated with the go-live of the paired format. The paired build belongs to its own component.	FORMAT SUBSTITUTION (SBV/SD)
OPEN / ROTATE / EXIT CONQUEST	Product-target add, bid, or state rows as the verdict names. Target IDs, never target names: conquest targets are the entity class where name matching fails most quietly.	No verbatim counterpart
STOP-LOSS: CONCEDE OR DEFER		STOP-LOSS: CONCEDE OR DEFER

	Bid and budget rows to maintenance, or a state row to paused, as the verdict text names. Freed budget returns to the envelope at the instrument layer, not through a bulk row.	
PAUSED-PENDING (REASON + RE-ENTRY)	State row, P4 block 3, last in the upload. The reason and re-entry criterion live in T2 and T3, not in the bulk file; the bulk file carries only the state.	PAUSE (LEVERS EXHAUSTED)
LEAVE TO HARVEST	No row. Any live instance on the term is handled by its owning SOP and reaches this component, if at all, as a separate staged row with its own verdict.	LEAVE TO HARVEST

Why the cross-map column exists. Seven of the 25 KDR strings appear character-identical in v4.0 Section 13.1. Eighteen do not, and eleven v4.0 strings have no KDR counterpart. Ledger conflict C10 registers the divergence and resolves it in favour of the KDR's 25, matching SOP-12 v2.0 Section 5. This component is the one that has to live with the divergence mechanically, because the reason string carries the verdict verbatim and the T+1 attribution matches on it. Until v4.1 reconciles the two lists, the KDR string is what is written and the v4.0 string is what is cross-mapped, and the choice is recorded here rather than left to the operator. Logged as AM-10.

6. Worked Examples

All four examples are illustrative. No anonymized live product economics spine is attached to this build, so every example uses the standard training base declared in the PPC Economics Worked-Example Set v1.0: AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Maximum Acceptable 26.3%, CVR 5.2% unless stated otherwise. Every figure below is computed on the page; nothing is asserted that is not calculated.

WE-1 · The clean weekly session: 27 rows, four producing components, everything lands

Setup. Tuesday session. Staged set: 27 rows across four components. 15 bid changes, 5 budget changes, 4 negative-adds, 2 creations, 1 state change. Sum: $15 + 5 + 4 + 2 + 1 = 27$.

P2 pre-flight (11 minutes). No billing incident open. No unresolved delta on any affected campaign ID. No armed event window. No Recovery Mode entry, so v4.0 Section 12 rule P5 tiers apply rather than the tightened Section 11 rule R3 values. T7 completeness score recorded. Affected-campaign list: 14 campaign IDs. Confirmation tiers re-read: 4 bid rows exceed the 25% move threshold, 2 budget rows exceed the \$50 per day threshold, and both creations are structural changes, so $4 + 2 + 2 = 8$ rows are human-mandatory and the remaining $27 - 8 = 19$ sit in the auto tier with spot checks. The 8 confirmations are collected before the upload step, not during it.

P3 pull (14 minutes). Fresh pull at 09:40, scoped to the 14 campaign IDs. Snapshot saved untouched under the session identifier at 09:44, before the first edit. Working views filtered to the affected list. All 27 staged entities present in the pull.

P4 transcription (38 minutes). Matched by campaign ID, ad group ID, and keyword ID throughout. Ordering: block 1 carries the 2 creation settings blocks and then the 4 negatives that attach to them; block 2 carries the 15 bid rows and 5 budget rows; block 3 carries the single state change. Current values reconciled against the pull row by row: 27 of 27 agree, so no row stops here.

P5 validation (19 minutes). Count check by change type: 15, 5, 4, 2, 1, matching the staged set exactly. Reason strings: 27 of 27 carry four bracketed segments. Verbatim check: 27 of 27 verdict segments match a KDR string character for character. Spot check: 6 rows sampled, selection basis recorded as the 4 human-mandatory bid rows plus both creations, all 6 clean. Scope diff: clean, no row outside the 14 IDs. Encoding pass on the 2 created names: both round-trip with segments intact. Lever-limit check on the 15 bid rows: the largest move is \$1.40 to \$1.87, which is $(1.87 - 1.40) / 1.40 = 33.6\%$, inside the plus or minus 50% limit of v4.0 Section 9 rule L2, and above the 25% human-confirmation threshold of Section 12 rule P5, which is why it was already tiered human-mandatory at P2. The largest budget move is \$45.00 to \$102.00 per day, a rise of \$57.00 per day, above the \$50 per day threshold and likewise already tiered.

P6 ceiling scan (9 minutes). Economics from T0: contribution \$28.00 on AOV \$80.00, so break-even ACOS = $28.00 / 80.00 = 35.0\%$, Target = 17.5%, Maximum Acceptable = 26.3%. Term CVR 5.2%, so max profitable CPC = $28.00 \times 0.052 = \$1.46$ and Target CPC = $80.00 \times 0.175 \times 0.052 = \0.73 . One bid row reads \$1.40 base on a campaign carrying a top-of-search modifier of +45% under fixed bidding, so the dynamic raise is zero: effective top-of-search CPC = $1.40 \times 1.45 \times 1.00 = \2.03 . Against the \$1.46 ceiling that is $2.03 / 1.46 = 1.39$ times, which lands in the PUSH-ONLY band of v4.0 Section 3 family F1 (1.0 to 1.5 times). The row's reason string references its M2 sizing block, so it proceeds. The scan records the ratio rather than the adjective: the row is not "aggressive", it is authorizing \$0.57 per click above the ceiling, funded by a sized push.

P7 upload and P8 close (21 minutes). 8 human confirmations in hand. Three blocks uploaded in order. Processing report read at row level: 27 success, 0 rejected, 0 skipped. Attempted 27, success 27, so P8 branch A applies and the upload half closes with no delta log.

P9 verification, Wednesday (26 minutes). Second pull on the 14 campaign IDs. Row-by-row diff: 27 of 27 landed as intended. Scan for unstaged values on affected entities: none. A2 written from the verified values. T6 carries 27 prediction rows with their review dates. Session record closed: 27 staged, 27 deployed, 27 verified, 0 held, 0 deltas, pull-to-upload age 3 hours 6 minutes recorded.

What the example demonstrates. A clean session still produces eleven artifacts and still takes roughly 2 hours 18 minutes across two days. The cost of the standard is visible; the cost of skipping it is not, which is why WE-3 exists.

WE-2 · The edge case: a ceiling breach, an orphaned negative, and a silent skip in one session

Setup. Tuesday session. 31 staged rows. Same economics base, except that the term in row 12 carries a validated CVR of 4.1% rather than 5.2%, so its own ceiling differs from the product's blended figure.

P6 catches the ceiling breach. Row 12 stages a bid of \$2.40 on a campaign carrying a top-of-search modifier of +90% under up-and-down bidding, which PPC Strategic Doctrine v1.0 Section 6.1 records as adding up to 100% at top of search. Effective top-of-search CPC = $2.40 \times 1.90 \times 2.00 = \9.12 . The ceiling for this term is max profitable CPC = $28.00 \times 0.041 = \$1.15$. The ratio is $9.12 / 1.15 = 7.94$ times ceiling, far above the 1.5 times boundary that v4.0 Section 3 family F1 classes as UNJUSTIFIABLE outside a sized push, and the row carries no M2 reference. **Row 12 is HELD** and returned to its owner with the arithmetic attached. It is not repaired here: reducing \$2.40 to a compliant value would be this component choosing a bid, and the correct instrument is v4.0 Section 8.4 scenario S-R11, which ladders toward ceiling under the deciding SOP. Staged 31, held 1, so attempted = $31 - 1 = 30$.

P3 step 5 did not catch the orphan, and P7 did. Row 19 is a negative-exact pointing at a keyword ID whose parent campaign was archived 6 days earlier. The entity was present in the pull as an archived row, so the P3 presence test passed on a technicality. The upload rejects it with a named cause. Fix: the negative re-points to the live twin campaign by ID, taken from the T2 K8 registry rather than by searching names. Re-uploaded in the same session; success. The root-cause line names the staged row as built against a registry read older than the archive event, and the P3 presence test is recorded as passing on presence rather than on state, which is the check-design gap this example exists to surface.

P8 finds the silent skip by counting, not by reading errors. First processing report: 28 success against 30 attempted, with exactly 1 named rejection. $30 - 28 - 1 = 1$ row unaccounted for. Set difference between transcribed and landed identifies it: a keyword bid row whose keyword state reads ARCHIVED, which

the platform skipped without naming. After the row 19 fix re-uploads, success reaches 29 against 30 attempted, leaving 1 delta. The delta row is logged with entity ID, staged value \$1.74, cause "keyword archived between staging and upload", disposition "returned to owner for re-read", and its owner role. Reconciliation: 29 deployed + 1 held + 1 delta = 31 staged.

Session close and freeze. The unresolved delta freezes new sessions for that single campaign ID under Section 3, not for the product and not for the marketplace. The other 13 campaign IDs in the session remain open for the next cycle.

P9 verification, Wednesday. 29 of 29 deployed rows landed as intended. Zero unexplained mismatches. The held row and the delta row are carried in the session record as open items with owners, so the next cascade reads them as work rather than as history.

What the example demonstrates. Three defects, three different detectors, and none of them was the error list. The ceiling breach was caught by arithmetic, the orphan by a platform rejection the operator read at row level, and the silent skip by subtraction. A session that reads only the error list finds one of the three.

WE-3 · T+1 mismatch classification: all four branches in one diff, and the price of the one that hides

Setup. 24 rows deployed Tuesday, all reported successful. Wednesday pull and diff. 20 rows match exactly. 4 mismatch, and each takes a different P10 branch: $20 + 2 + 1 + 1 = 24$.

Branch 3, two rows: platform transformation. Two budget rows read the prior day's figure. Both are classed as settling, re-checked once on a later pull, and both resolve to the staged values. The session record dates the re-check and records the class as transformed rather than asserting a governed lag window, because v4.0 states none. Logged against AM-3.

Branch 1, one row: external change, rule-attributed. A bid reads \$1.19 against \$1.31 staged. The #37 rule change log shows a target-ACOS ladder step fired overnight on that entity. v4.0 Section 12 rule P3 governs: the weekly audit outranks the automation platform between cycles, the audit value wins, and automation resumes from it after upload. The value restores to \$1.31 in a mini-session, and the entity is raised with #37 for a scope exclusion rather than being restored again every Wednesday. Restoring without raising the exclusion is how a one-row mismatch becomes a standing weekly cost.

Branch 4, one row: deployment error, and this is the expensive one. A bid staged at \$1.62 landed at \$16.20: a decimal transposition in transcription, which uploaded successfully because \$16.20 is a valid bid. Against the term's max profitable CPC of $28.00 \times 0.052 = \$1.46$, the landed value is $16.20 / 1.46 = 11.13$ times ceiling. The entity took 38 clicks in the day before detection. Spend at the landed bid: $38 \times 16.20 = \$615.60$. Orders at CVR 5.2%: $38 \times 0.052 = 1.976$. Contribution recovered: $1.976 \times 28.00 = \$55.33$. Net loss for the day: $615.60 - 55.33 = \$560.27$. Had the staged value landed, spend would have been $38 \times 1.62 = \$61.56$ against the same \$55.33 of contribution, a net loss of \$6.23. The difference, $560.27 - 6.23 = \$554.04$, is the cost of this single row, and it equals $38 \times (16.20 - 1.62) = \554.04 exactly, as it must.

Why the ceiling scan did not catch it. The scan at P6 ran against the staged value of \$1.62, which is $1.62 / 1.46 = 1.11$ times ceiling and sits legally in the PUSH-ONLY band with its M2 reference. The transposition happened at transcription, downstream of the scan and upstream of the upload, and the only checks positioned to catch it are the P5 spot check, if that row is sampled, and the P9 diff, which catches it always. The root-cause line names the spot check as the validation that would have caught it, and the sample did not include the row. That is the direct operational argument for the sampling rule logged as AM-1: an unsampled row's only defence is a diff one day later, and one day at 11.13 times ceiling is what one day costs.

Disposition. The correction is staged as a new mini-session through P1 to P9 in full, not edited in the console, because a console edit here would leave the largest financial event of the week with the weakest evidence trail. The week's performance reads on that entity are void: it did not run the values the decision assumed. T6 carries the correction, the void note, and the root-cause line.

WE-4 · Rollback of a state deployment scored WRONG-DIRECTION on an execution cause

Setup. Tuesday 10:20, a state change paused 3 campaigns as part of a consolidation, and the pause was deployed correctly and verified at T+1. Friday 10:20, the V5 read in T6 scores the originating prediction WRONG-DIRECTION and the root cause names execution rather than judgment: the consolidation's survivor was never given the loser's negative lists, so the survivor served queries the paused campaigns had been fencing.

Trigger test. The cause is execution, so P11 applies. Had the root cause named judgment, the row would return to its owning SOP for a re-decision and would reach this component later as a new staged set. This distinction is the whole of P11 step 1 and it is the one that keeps rollback from becoming a second opinion.

Restore from the snapshot. The three campaign states restore from the P3 snapshot taken Tuesday at 09:44, not from the workbook's prior-value column. The three campaigns carried daily budgets of \$34.00, \$27.00, and \$18.00, summing to $34.00 + 27.00 + 18.00 = \79.00 per day, and those budgets restore alongside the states so the envelope reconciliation at T0 zone H reads correctly.

Record the elapsed window. Tuesday 10:20 to Friday 10:20 is 72 hours. That window is written to T6 because the next read on these entities will contain 72 fewer hours of clicks, and v4.0 Section 5 gate G2 "Sample sufficiency" will judge it on the clicks it actually holds. Without the window note, the next cycle reads a short sample as a performance change. Forgone spend across the window at the restored budgets: $3 \times 79.00 = \$237.00$, recorded so the envelope read is not credited with a saving that was an outage.

Run the full standard. The rollback is staged with the reason string [SOP-38][ROLLBACK][session 2026-08-04-US-01][P10 branch 4], validated through all eight P5 checks, uploaded, closed at 3 attempted and 3 success, and verified at T+1 with 3 of 3 landed. The missing negative lists are not added in this session: they are a #29 staged set and they arrive through the normal path with their own reason strings.

What the example demonstrates. Rollback is the one procedure where the temptation to work quickly is highest and the evidence requirement is unchanged. The snapshot is what makes it a 40-minute execution instead of a reconstruction, and the snapshot only exists because P3 step 3 was run on a day when nothing had gone wrong.

7. Failure Modes and Escalation

Twelve failure modes. Each carries a detection signal, a response, and the account-history incident that motivated it where the pack holds one. Where the pack holds none, the cell says so rather than supplying a plausible one, because a fabricated incident is worse than an absent one: it invites an operator to reason from a case that never happened.

Failure	Detection signal	Response	Motivating incident, where the pack holds one
F1. Stale-file editing	Current-value mismatches clustering at P4 step 4, or a pull-to-upload age in the session record that exceeds the session's own duration by a wide margin	ABORT under P8 branch E, re-pull, restart from P3. Rows that already landed go to P11.	v1.0 records this as a class and names the P3 snapshot rule as its answer. No dated incident is held in the pack, and none is supplied here. The absent freshness limit is AM-2.
F2. Partial success unread	Success count below attempted count with no delta log, or a session record whose only upload evidence is a summary line	Reconcile immediately by set difference. The processing report is read at row level in every session, and the session record carries the three counts separately so the arithmetic is visible.	None dated in the pack. The class is carried forward from v1.0 F2 unchanged.
F3. Name-based matching landing on the wrong twin	A change landing on a stage-stack sibling: two campaigns on one term whose behaviour stops matching their objective tags	P11 rollback of the wrong-entity change, then redeploy by ID. ID matching is enforced at P4 step 1 rather than recommended.	v1.0 records the twin case as the reason for the rule, without particulars. The structural cause is documented: v4.0 Section 6 rule O4 makes dual-objective instances legal and similarly named, so the pack explains why the class exists even though it holds no dated instance.
F4. Unlogged console edit	P9 step 5: a changed value on an affected entity that no session record and no #37 rule log claims	Attribution first, correction second. Identify the acting role, back-fill the record, and route to #46, whose P10 uses this report as its real-time detector for unrecognized-actor response. Recurrences escalate with attribution attached.	The pack holds a dated adjacent case: the 2026-08-01 billing failure that #43 was built around, recorded in the Ledger row for component #46. It is a different failure class and is cited only as the reason the attribution channel to #46 exists.
F5. Verification skipped	No T+1 diff artifact exists for a closed session, or A2 carries a refresh timestamp older than the session's upload	Run it late. The session is not closed until the diff exists and A2 is stale until then, which means every downstream current-value reconciliation is reading intent as if it were state.	None dated in the pack. WE-3 quantifies the exposure at \$554.04 for a single undetected row over a single day.
F6. State-change ordering voiding an edit	An edit reported successful whose entity was paused earlier in the same upload; the T+1 diff shows the pre-state value on a row the report called successful	P4 block ordering enforced: creations, then values, then states. The voided edit redeploys in a new session; it is not re-uploaded into the same file.	None dated in the pack. Carried forward from v1.0 F6 unchanged.
F7. Reason string incomplete or verdict not verbatim	P5 step 2 and step 3 fail counts above zero; or a T+1 diff row that cannot be attributed to a verdict because the string does not match the vocabulary	Return the row to its owner. Do not repair a reason string inside a deployment session: the string is the producing SOP's justification, and rewriting it here creates a justification this component invented.	The pack holds the structural cause rather than an incident: Ledger conflict C10 records the vocabulary stated at three counts, 18 in v4.0 Section 13.1, 25 in the KDR and the workbook validation list, and 26 in the Data Architecture rule 8 and Ledger row 8. Seven strings are common to the two lists. An operator quoting from the wrong list produces an unattributable row.
F8. Ceiling breach repaired instead of held	A transcribed bid or modifier differing from its staged value with no returned-and-restaged row behind it; or a session record with a ceiling-scan result of pass on a row whose arithmetic fails	The row is HELD and returned with the arithmetic attached. This component does not choose bids. A repaired breach is this component making a decision, which the Authority Matrix row for #38 gives it no D cell to make.	The pack holds a dated platform driver: PPC Strategic Doctrine v1.0 Section 6.3 records Amazon's late-April 2026 update to the dynamic bidding engine and CPC inflation of 18 to 27% against Q1 baselines under dynamic raises. That is the environment in which stacked modifiers stopped being theoretical.
F9. Orphaned child row: parent archived between staging and upload	Upload rejection naming a parent that does not exist, or a P3 step 5 presence test that passed on an archived entity	Re-point the child to the live parent by ID from the T2 K8 registry, re-upload in the same session. Where the parent has no live twin, the row returns to its owner as a structural question rather than a transcription one.	None dated in the pack. WE-2 works the case and exposes the check-design gap: the P3 presence test passes on presence, not on state.
F10. Scope contamination	P5 step 7 scope diff finds a transcribed row outside the affected-campaign list, or the T+1 diff finds a changed entity that was never in the session's scope	ABORT under P8 branch E and roll back anything outside scope. A contaminated session cannot be partially trusted, because the mechanism that admitted one out-of-scope row admitted it silently.	None dated in the pack. The class is new at v2.0; v1.0 carried the scope diff as a validation step with no matching failure mode.
F11. Rule and manual double-touch inside one window	P10 branch 1 firing repeatedly on the same entity across consecutive cycles, which is a scope-exclusion gap rather than a series of coincidences	Restore the audited value per v4.0 Section 12 rule P3, then raise the scope exclusion with #37 rather than restoring weekly. A second firing on the same entity after an exclusion request escalates.	The pack holds the interlock rather than an incident: #37 P7 states that a staged manual decision suspends rule action on that entity for the decision's read window, and that the T+1 attribution is what separates the two sources. Note that #37 P5 cites a "v4.0 consecutive-step reference" that does not exist in v4.0; the runaway signature is therefore unnumbered on both sides. Logged as AM-7.
F12. Session frozen on an unresolved delta past its escalation point	A delta row open in the session record with no disposition and no owner action, while the affected campaign IDs remain frozen	Escalate to the PPC Lead with the processing report attached. The freeze is scope-limited by design, so the cost of leaving it open is bounded but it is not zero: the frozen IDs miss a full cycle of verdicts.	None dated in the pack. v1.0 asserted a 24-hour escalation clock with no source; the clock is not carried forward and the gap is AM-4.

Escalation. F4 escalates on second occurrence with attribution attached. F8 and F10 escalate on first occurrence, because both mean a validation ran and was overridden rather than a validation being absent. The same failure mode twice in consecutive cycles on the same scope escalates to the PPC Lead with the T6

history attached, matching the exemplar's standard. Anything touching the human-mandatory tier is confirmed per the Authority Matrix before deployment, without exception, and the upload action itself is human-mandatory in every session.

8. Outputs and Artifacts

Gates clear with artifacts, not assertions. Every artifact below is named, located, and consumed by someone; an artifact with no downstream consumer is not produced.

Artifact	Storage location	Retention	Downstream consumer and what it does with it
Session record (staged counts, held rows, delta log, clock times, freeze checks, T7 score)	Deployment session folder, keyed by session identifier	Not governed; retained pending AM-6	The next session's P2 step 2 freeze check reads open deltas from it; #6 reads it for execution-cause routing on missed predictions
Pre-state snapshot (untouched timestamped bulk download)	Same session folder, alongside the record	Not governed; retained pending AM-6	P11 rollback restores from it; it is the only artifact that makes rollback an execution rather than a reconstruction
Processing report or reports, one per upload attempt	Same session folder	Not governed; retained pending AM-6	P8 delta classification; F12 escalation attaches it; #43 triage branches 2 to 4 read it
T+1 diff report with zero unexplained mismatches	Same session folder, dated to the verification day	Not governed; retained pending AM-6	The gate artifact for any deployment. Consumed by #12 F10, #29, #13, #14, #16 as their landed-value proof; by #37 P7 for source attribution; by #46 P10 as the unrecognized-actor detector; by #43 for incident triage
A2 deployment state, written from verified values	Command Workbook T3 CAMPAIGN BOARD group C2, refreshed into T7 as PRESENT	Live, overwritten each cycle	Every next-cycle current-value reconciliation at P4 step 4; every gate in Section 3 that reads current state
T6 decision-log rows, one per deployed change with its prediction	Command Workbook T6 DECISION LOG, append-only	Log life	V5 scoring at the named review date, per v4.0 Section 14 rule V1; the Account Rollup A3 verification scoreboard
Held-row register with hold reason and owner role	Session record, carried forward until closed	Until the row redeploys or is formally declined	The producing SOP's next pass; a held row that never returns is a decision quietly abandoned
Delta log with row, cause, and disposition	Session record	Until every row is dispositioned	P2 step 2 of the next session; the freeze it drives is scoped to its campaign IDs
Root-cause lines on deployment errors	T6 alongside the correcting row	Log life	Section 10 version governance: a root cause naming no existing validation is a trigger to add one

9. Skill-Conversion Block

Block	Specification
Input bindings	Staged change sets meeting the P1 contract from every producing component, carrying entity IDs, current and new values, change type, four-segment reason string, staging owner role, review date, and confirmation tier. The fresh bulk pull scoped to the affected-campaign list. A2 deployment state. T2 K6 max profitable CPC and Expected CPC for the ceiling scan. T2 K8 registry counts for parent-liveness and twin resolution. T3 C1 objective tags and C2 deployment fields. T7 manifest status for the eight field groups in Section 3. The KDR v3.0 verdict vocabulary as a fixed string list for the verbatim check. The #37 rule change log. The #43 incident-broadcast state. T6 for prediction write-back.
Deterministic logic	The P1 four-segment reason-string grammar and its field-presence test. The P4 ID-match resolution and the three-block ordering with parent-first sequencing inside block 1. The P5 eight checks, of which seven are pure set or string operations and one, the spot check, is a sampling procedure with no governed rule. The P6 ceiling identity from Doctrine 6.1 and the four CPC ratio bands from v4.0 family F1. The L2 lever-limit test. The P8 close condition, success equals attempted or every difference logged, and its five branches. The P10 five-branch mismatch classifier, evaluated in order with a stop at first match. The P5 tier assignment against v4.0 rule P5, or rule R3 where a Recovery entry is live.
Human checkpoints	Drawn from PPC Authority Matrix v2.0 Section 1. Row "Bulk deployments (#38)": PPC Lead R, incoming operator E, launch and CA operator E on own scope, PPC Manager E, no D cell. Matrix Standing Rule 2 makes every E a staged-decision execution only. AI-executable with audit: transcription drafting, all eight P5 validations, the P6 arithmetic, processing-report parsing, delta classification, the full T+1 diff, and the P10 classification draft. Human-mandatory in every session: the upload action itself, ABORT and ROLLBACK decisions, and held-row dispositions. Human-mandatory per row: the v4.0 rule P5 tiers, replaced by the rule R3 tiers where Recovery Mode is live on that row's product, confirmed by the role the Matrix names for that decision class and not by the operator running the session. Row "Code-red declaration + close": PPC Lead D, so the emergency path's authority re-read at P12 routes there. Row "Automation rules (#37)": PPC Lead D, so an F11 scope-exclusion request routes there rather than being settled inside a deployment session.
Cadence	Weekly Tuesday session after the huddle, following the Monday cascade. Same-day compressed sessions on graduation chains, funding releases, inventory RED, eligibility flips, event-plan actions, and code-red actions. T+1 diff the morning after every deployment without exception, including same-day and console-edit deployments. A standing scan at every pull for values without session records, which is the F4 detector. Freeze-condition re-read at the open of every session.
Outputs	Drafted transcriptions with their validation results queued for the human upload; the ceiling-scan result per bid and modifier row with the ratio stated; parsed processing reports with classified deltas; T+1 diff reports with classified mismatches and dispositions; A2 state writes from verified values only; T6 prediction rows carried from the staged set; F1 to F12 alarms with evidence lines naming the row, the signal value, and the rule or gate reference; execution-cause routing packets to #6.
What the skill must refuse to do	Compute a bid, repair a ceiling breach, complete a reason string, choose between two closed vocabularies, invent a spot-check sample size, assert a platform-lag tolerance, or close a session at success below attempted without a delta log. Each refusal maps to a numbered amendment item in Section 10.3, so the skill's silence is traceable rather than arbitrary.

10. Version Governance

Owner role: PPC Lead, per the Authority Matrix row for bulk deployments, where the Lead holds R and no role holds D.

10.1 Review triggers specific to this component

- Any change to the advertising platform's bulk-sheet schema.** A schema change increments this SOP before the next session runs, because column positions and accepted value formats are the substance of P4 and P5.
- Any change to the P1 staging contract.** The contract binds every producing SOP, so a change here is coordinated with all of them before it is written. A producing SOP that changes its output format coordinates the change here first, not after.

3. **Any v4.0 increment.** The Section 5 verdict mapping re-validates against the new verdict set before any session runs on it, and the rule P5 tier values in P1 step 6 and the rule R3 values in P12 re-read at the same time.
4. **Any resolution of Ledger conflict C10.** When the closed vocabulary is reconciled to one list, Section 5 collapses to that list and the cross-map column retires.
5. **A deployment-error root cause naming no existing validation.** If a P10 branch-4 error was catchable by no check in P5, the validation set is short by one and this SOP increments to add it. This is the only trigger in the list that fires from operational evidence rather than from a document change.
6. **Any of the eleven amendment items in Section 10.3 landing in v4.1.** Each landing replaces a recorded gap with a governed numeral, and the affected procedure step is rewritten to cite v4.1 by section heading.

10.2 Change note, v1.0 to v2.0

Supersedes v1.0 in full. Procedures expanded 8 to 14: P2 session pre-flight, P6 ceiling-scan arithmetic, P8 delta classification, P10 mismatch classification, P13 console edit control, and P14 session close are new or split out from compressed v1.0 steps. Failure modes expanded 6 to 12: F7 reason-string and verbatim failure, F8 ceiling breach repaired, F9 orphaned child row, F10 scope contamination, F11 rule and manual double-touch, and F12 frozen session are new. Worked examples expanded 1 to 4. Verdict mapping expanded from a 4-row session-control table to the full 25-verdict vocabulary with a v4.0 Section 13.1 cross-map. Elements 11 and 12 added. Three local numerals asserted in v1.0 with no source are retired rather than carried: the five-row spot check, the plus-one-day platform-lag re-check, and the 24-hour session-freeze escalation clock. Each is logged in Section 10.3. The v1.0 use of "G2" for bulk-file freshness is retired, since it matches neither the v4.0 assignment nor the corpus alias. The Ledger's pending fold-in for this component, "rollback and restore snapshot procedure", is closed by P3 step 3 and P11.

10.3 Open v4.1 amendment items

Eleven numerals this component needs and v4.0 does not contain. None is written into the procedures above.

ID	The missing numeral	Where it bites and why it was not invented
AM-1	Pre-upload spot-check sample size and selection rule	P5 step 4. v4.0 rule P5 names "spot checks" for the routine tail with no count and no sampling basis. WE-3 shows the cost of an unsampled row at \$554.04 for one day. v1.0 asserted five rows with no source.
AM-2	Maximum bulk-file age between pull and upload	P3 step 2. The freshness rule exists in prose in v1.0 and in this SOP; the clock does not exist anywhere. Until it lands, the age is recorded rather than bounded.
AM-3	Platform-lag re-check window at T+1 mismatch	P10 branch 3. v1.0 asserted one day, once. Overlaps the item SOP-29 logged as its AM-10, verification lag; the two should be reconciled as one numeral rather than two.
AM-4	Escalation clock on an unresolved delta, and the return lag for a held row	P8 step 2, F12, and the held-row trigger in Section 2. v1.0 asserted 24 hours with no source.
AM-5	Landed-value tolerance for platform transformation	P10 branch 3. Without it, rounding and format transformations are classified by operator judgment, which is the one thing this component is built to remove.
AM-6	Retention period for session records, pre-state snapshots, processing reports, and diff reports	Element 8 and P14 step 5. The Ledger row for component #39 already carries "retention policy lines" as a pending fold-in; this is the same gap seen from the deployment side.
AM-7	The consecutive-step reference for the runaway signature	F11. SOP-37 P5 and its Element 9 both cite "the v4.0 consecutive-step reference"; no such reference exists in v4.0. The signature is therefore unnumbered on both the automation side and the deployment side.
AM-8	The tolerance term in the placement-multiplier ceiling	P6 step 5. v4.0 rule M4 sizes the multiplier against "Expected CPC + tolerance" and family F9 repeats the phrase; neither numbers it. The scan currently tests against the two hard values and records the untested margin.
AM-9	Character-set and separator rule for reason strings and created campaign names	P5 step 8. The naming convention parses only if segments survive upload, and no rule states which characters survive.
AM-10	Which closed vocabulary a reason string carries, and which string is canonical on the 18 rows where the two lists diverge	P1 step 4 and Section 5. Ledger C10 registers the divergence and resolves the citation question; it does not resolve the machine-attribution question, which is this component's.
AM-11	A gate number for deployment-file freshness, or a ruling that it is not a gate	Section 3. Three different subjects currently travel under G2 across v4.0, SOP-12 v2.0, and SOP-38 v1.0. This SOP treats freshness as a P3 precondition and claims no G-number.

Registered in the Supersession Map. **Companions:** Master PPC Decision Criteria System v4.0 (all thresholds and verdicts) · Keyword Decision Record Template v3.0 (the closed vocabulary) · PPC Authority Matrix v2.0 (Element 9 and Element 11 checkpoints) · Master PPC Data Architecture v1.0 and the Command Workbook READ ME (T2, T3, T6, T7 mechanics) · PPC Strategic Doctrine v1.0 Sections 6.1 to 6.3 (the ceiling identity and the 2026 auction context) · PPC Economics Worked-Example Set v1.0 (the arithmetic base) · SOP-12 v2.0 (the format exemplar and the staging producer) · SOP-37 (the other deployment path) · PPC Cross-Reference Ledger (ownership boundaries).

11. Operator Ramp

Built from the PPC Authority Matrix v2.0 rows that name this component's decisions. The ramp exists because the incoming operator holds E on the "Bulk deployments (#38)" row from day one and the row carries no D cell, which means the competence being certified is accuracy under a standard rather than judgment about values.

Week 1: observation and artifact familiarity

The checkpoint holder is the PPC Lead, who holds R on the deployment row. The operator observes two full sessions and demonstrates the following, all of which are reads rather than executions.

1. **States why the deployment row has no D cell** and what Matrix Standing Rule 2 means in practice: an E without a D executes staged decisions only. The operator can name what they would do with a staged value they believe is wrong, which is return it, and why editing it would be a role violation rather than a helpful correction.
2. **Locates every artifact in Section 8 cold:** the session record, the pre-state snapshot, both processing reports, the T+1 diff, T3 group C2 where A2 is written, T6 where predictions land, and T7 where the eight field groups in Section 3 report their status.
3. **Reproduces the ceiling identity from PPC Strategic Doctrine v1.0 Section 6.1 without the document open** and computes, on the training base of the PPC Economics Worked-Example Set v1.0, $\text{max profitable CPC} = 28.00 \times 0.052 = \1.46 and $\text{Target CPC} = 80.00 \times 0.175 \times 0.052 = \0.73 . The operator then states the four CPC ratio bands and their source: SAFE, AT-CEILING, PUSH-ONLY, and UNJUSTIFIABLE at 0.8, 1.0, and 1.5 times, per v4.0 Section 3 family F1 "Product Economics". Naming the source is half the competence; a band recited without it is a number the operator cannot defend when it changes.
4. **Reads a closed session record and states what it proves and what it does not.** Specifically: that a success count equal to the attempted count proves the platform accepted the rows and proves nothing about what landed.
5. **Recites the P4 three-block ordering and the reason for it,** which is dependency, not tidiness.
6. **Names the 25-string vocabulary's location** and the fact that it diverges from v4.0 Section 13.1 on 18 of 25 rows.

Week 2: supervised execution

The operator drafts and the PPC Lead confirms every step. The operator holds E; the Lead's R is exercised in real time rather than after the fact.

1. **Runs P2 pre-flight unaided,** including the four freeze conditions in order, and correctly scopes an unresolved delta's freeze to campaign IDs rather than to the product.
2. **Transcribes a set of 15 or more rows by ID** with the three-block ordering, and correctly stops a row whose staged current value disagrees with the fresh pull rather than adjusting it.
3. **Runs all eight P5 checks and records all eight results,** including a recorded result for the spot check that states the count sampled and the selection basis, since no rule governs it.
4. **Holds at least one row correctly.** The Lead seeds a ceiling breach into the practice set. The competence is holding it and attaching the arithmetic, not detecting it.
5. **Reads a processing report at row level** on a set containing one silent skip, and finds the skip by subtraction rather than by reading the error list.
6. **Assembles the confirmation set before the upload,** correctly tiering rows against v4.0 rule P5, and correctly re-tiering a product under Recovery Mode to the rule R3 values.
7. **Runs the T+1 diff and classifies every mismatch through the P10 branches in order,** stopping at first match rather than reasoning to the most likely cause.
8. **The Lead confirms the upload action in every session this week.** That confirmation is permanent, not a training measure: the upload action is human-mandatory in every session at every experience level.

Week 4: unsupervised execution with audit

The operator runs sessions alone. The PPC Lead audits the closed session record and the T+1 diff rather than the execution.

1. **Three consecutive sessions close with zero unexplained mismatches** and with every P5 check carrying a recorded result.
2. **One session is deliberately aborted** under P8 branch D or E during the audit period, and the abort is correct: the operator has demonstrated that stopping is available, since an operator who has never aborted has not shown they can.
3. **One rollback is executed end to end** from the snapshot, with the elapsed window and forgone spend both recorded.
4. **One deployment error is root-caused** to a named validation, or to none, with the "none" case raised as a Section 10.1 trigger rather than absorbed.
5. **Held rows and deltas from the operator's sessions are all dispositioned,** which tests follow-through rather than execution: a held row that never returns is the quietest failure this component produces.

Live-defense questions

These test the reasoning behind the thresholds, not recall of their values. The Lead may open v4.0 during the defense; the operator answers with the reasoning and cites where the number lives.

1. A staged bid row would move a bid by 24%. Recovery Mode is live on the product. Which confirmation tier applies, and why do the two rule sets carry different numbers rather than one?
2. A row's effective top-of-search CPC computes to 1.2 times max profitable CPC. Explain why that can be correct and what single reference in the reason string decides it.
3. The processing report reads 30 of 30 successful. Name two distinct defects that are still possible, and the check that would find each.
4. A T+1 diff shows a bid \$0.12 below the staged value. Walk the P10 branches in order and state what evidence you would need at each to stop there.
5. Why does this SOP hold a ceiling-breach row instead of reducing the bid to the ceiling, given that reducing it would be arithmetically correct and faster?
6. A verdict string in a reason string reads "TAPER TO FLOOR (POST-OBJECTIVE)". It is a valid v4.0 Section 13.1 string. Does the row deploy? Explain what the two vocabularies do to the T+1 attribution.
7. The pre-state snapshot was taken after the first two rows were transcribed. What is now impossible, and at what point in the session should that have stopped it?
8. A budget row and a bid row are staged on the same truncated campaign in one session. Which deploys first, and what would reading the bid through the truncated day have cost the next cycle?
9. You paused three campaigns on Tuesday and un-paused them on Friday. What must be written to T6 beyond the state change, and which gate would misread the next cycle without it?
10. Name a numeral you use in this procedure that v4.0 does not govern, and state what you do instead of using a number.

12. Interface Contracts

Built from the Cross-Reference Ledger's Interfaces column for component #38 and from the rows of the components that name #38, cross-checked against Master PPC Data Architecture v1.0 for artifact names and tab locations. Every row states what happens when the handoff does not arrive, because that is the clause the corpus has been missing.

12.1 Upstream: what arrives here

ID	Sending owner	Artifact and its location	Timing	Non-arrival consequence
I-1	#12 exact operations, PPC Lead or assigned operator	Staged change set meeting P1, from T2 K9 verdicts and their campaign_moves	Monday PM, for Tuesday	Exact-relevant rows do not deploy this cycle. The gap belongs to #12 P3 step 7; this component does not construct rows to fill it.
I-2	#13 phrase operations	Staged change set meeting P1	Monday PM, for Tuesday	Phrase rows do not deploy. Where a phrase-broad dedup wall was expected in the same cycle as an exact standup, the standup deploys without its wall and the exposure is recorded for #29.
I-3	#14 broad operations	Staged change set meeting P1	Monday PM, for Tuesday	Broad rows do not deploy; the weekly triage routes that depend on them defer one cycle.
I-4	#15 auto operations and the D5 split	Staged change set meeting P1, including per-group budget rows	Monday PM, for Tuesday	Auto rows do not deploy. A pending steering negative that does not deploy leaves the term GRADUATED-UNSTEERED per v4.0 family F8, which is a live leak rather than a deferred change, and it is recorded as such.
I-5	#16 PAT and conquest	Staged change set meeting P1, target IDs mandatory	Monday PM, for Tuesday	Conquest rows do not deploy; open targets continue at their current bids. #16 records the deferral against its own weekly read.
I-6	#29 negation architecture	Staged change set meeting P1, every negative carrying its declared mode	Monday PM, for Tuesday, and same day on structure events	Walls do not land. #29 cannot close its atomic chain, and the source campaigns keep buying terms their exacts are fighting. Recorded and escalated on a second consecutive cycle.
I-7	#17, #18, #19 brand formats	Staged change set meeting P1, plus the SP-exact step-down row dated to the paired go-live	Monday PM, or the dated go-live	The paired posture breaks: one side of a format substitution live and the other not. The step-down row is held rather than deployed alone.
I-8	#28 placement and bid modifiers	Derived ladder steps and modifier values inside the producing SOP's staged rows	Monday PM, for Tuesday	Ladder rows deploy at their last governed value or not at all; this component never derives a step. A ladder that misses a dated step has its chain broken in T6 and V5 scores the gap.
I-9	Instrument layer, T0 zone H and Account Rollup A2	Envelope deltas as staged budget rows	Monday PM, for Tuesday	Budget rows do not deploy as a set. Per gate G7 the set returns whole; this component never deploys the profitable half of an envelope reconciliation.
I-10	#32 deal and event window	Event plan naming dated actions, entering T2 with its dates	Before the window arms	Inside an armed window with no plan, nothing deploys on the affected products. The freeze is the safe state, so non-arrival is conservative rather than blocking.
I-11	Data Operations	Raw Pull Pack landed and the product workbooks refreshed	Monday 06:00 and 09:00 per Data Architecture v1.0 Section 12	The session cannot open: P3 has nothing fresh to pull against and P4 step 4 cannot reconcile current values. The session defers rather than running on an old file.
I-12	#39 data feed and manifest	T7 MANIFEST status for the eight field groups in Section 3	With the Monday refresh	Completeness is unknown, which under gate G8 is treated as DEGRADED rather than as PRESENT. The session narrows to gate-driven and repair rows.
I-13	#37 automation platform rules	Rule change log covering the deployment window	Continuous; read at P10	P10 branch 1 cannot be evaluated, so a rule-fired change classifies as unattributed and opens F4 against a human who did not act. The ambiguity escalates the same day rather than being resolved by assumption.
I-14	Authority Matrix holder for the decision class	Human confirmation stamp for every row in the human-mandatory tier	Before the P7 upload action	Those rows do not deploy. The rest of the session proceeds; a missing confirmation narrows a session, it does not abort one.
I-15	#43 ad-account billing continuity	INCIDENT-OPEN or INCIDENT-CLEAR broadcast	Immediate on incident	An INCIDENT-OPEN with no matching clear leaves every session closed. Arrival of the open state is what freezes; absence of the clear is what keeps it frozen, and neither is worked around.

12.2 Downstream: what leaves here

ID	Receiving owner	Artifact and its location	Timing	Non-arrival consequence
I-16	Every producing component: #12, #13, #14, #15, #16, #17-19, #29	T+1 diff report with zero unexplained mismatches, in the session folder	The morning after every deployment	The producing SOP cannot prove its verdict executed. #12 F10 cannot fire, #29 cannot close its wall verification, and every affected row's weekly read is being taken against an unproven state.
I-17	Next cycle's cascade, via the workbook	A2 deployment state written from verified values into T3 group C2 and stamped in T7	At P9 step 7	The next cascade reconciles current values against intent rather than state, which corrupts P4 step 4 for every component at once. This is the highest-consequence non-arrival in the table.
I-18	#6 V5 prediction scoring	Execution-cause routing packet: scored misses whose cause is a value that did not land	At the review date of the affected T6 rows	Execution misses are scored as judgment misses. The deciding SOP's rule-level accuracy under v4.0 rule V5 falls for a reason that is not its own, and the wrong rule enters review.
I-19	#37 automation platform rules	Attribution report separating session-fired from rule-fired changes, plus scope-exclusion requests from repeat F11 firings	Weekly, with the diff	#37 P7 precedence cannot be applied and its F6 double-touch cannot be detected. A rule and a session keep overwriting one entity indefinitely.

I-20	#46 access and permissions	Unattributed-change lines from the P9 step 5 scan	With every diff	#46 P10 loses its real-time detector for unrecognized-actor response, and an unauthorized change is indistinguishable from an unlogged authorized one.
I-21	#43 ad-account billing continuity	T+1 diff and A2 state for triage branches 2 to 4	On request during an incident	#43 cannot separate a billing-caused serving stop from a deployment-caused one, and its triage branches stall on the ambiguity.
I-22	T6 decision log, consumed by V5 and the Account Rollup A3 scoreboard	One prediction row per deployed change, carried from the staged row	At P14 step 3	Unscored actions do not accumulate per v4.0 rule V1: the change happened and the system has no record that predicts anything about it, so it can never be scored.

Appendix A. The Staged Row, Field by Field

Restated from P1 for operator convenience; P1 governs. A row missing any field returns to its owner before the session opens.

Field	Content	Why the session cannot run without it
marketplace	The account and marketplace the row belongs to	The CA block runs the same standard with a different deploying role; a mixed-marketplace file uploads to one of them
entity type	One of eight values in P1 step 2	Determines which of the three upload blocks the row belongs to
campaign ID, ad group ID, keyword or target ID	Platform IDs, populated to the depth the entity type requires	The match key at P4 step 1. Names are reference only
entity name	Human-readable reference	Never the match key. Present so a held row can be discussed without an ID lookup
current value	The value at staging time	Reconciled against the fresh pull at P4 step 4; a disagreement stops the row
new value	The value the verdict carries, at staged precision	Transcribed exactly; a value needing rounding is a staging defect
change type	One of ten values in P1 step 2	The count check at P5 step 1 reconciles by type, not only in total
reason string	Four bracketed segments in fixed order	Checked at P5 steps 2 and 3; it is the only field the T+1 attribution can match on
staging owner role	The role a held row returns to	A held row with no owner is a decision quietly abandoned
review date	The date the prediction is scored	Carried into T6 at P14 step 3 so V5 can score it
confirmation tier	Auto, spot-check, or human-mandatory	Stamped by the producing SOP against v4.0 rule P5 or rule R3; re-read at P2 step 6, never re-decided here

Appendix B. Deployment Session Operator Card

P1 to P14 compressed; the SOP text governs.

1. Freeze check in order: billing incident, open delta, armed event window, Recovery Mode. **2.** Affected-campaign list from the staged set; it is the boundary for everything after. **3.** T7 read and score recorded. **4.** Tiers re-read; human confirmations collected before the upload, never during. **5.** Fresh pull, clock time recorded, snapshot saved untouched before the first edit. **6.** Match by ID. Order: creations, then values, then states. **7.** Current value disagrees with the pull: the row stops, it is not adjusted. **8.** Eight checks, eight recorded results: count by type, reason-string completeness, verdict verbatim, spot check with its basis stated, ceiling scan, lever limit, scope diff, encoding. **9.** Ceiling breach with no sizing reference: HELD and returned with the arithmetic. Never repaired here. **10.** Upload is human-mandatory in every session. Read the report at row level. **11.** Success below attempted: find the missing rows by subtraction, not by reading errors. **12.** Close at success equals staged, or log every difference with row, cause, disposition, owner. **13.** T+1 diff every time, on everything, including console edits. **14.** Classify before correcting: rule, human, transformation, deployment error, silent skip, in that order. **15.** A2 from verified values only. T6 gets one prediction row per deployed change.

SOP-38 Bulk-File Deployment Mechanics v2.0 · Component #38 of the Operations Definition Register · Thresholds live in the Master PPC Decision Criteria System v4.0 only · Zero local thresholds; eleven ungoverned numerals logged at Section 10.3 rather than written · Inspiratek & Ecotero LLC