

SOP-42: Ad Policy, Moderation and Claims

PPC Department · Operating SOP · Component #42 of the Operations Definition Register

Version 1.0 · 2026-08-09 · First build (no prior version) · Owner: PPC Lead, Creative seam per the Gap Register charter · Due 2026-08-17, delivered ahead of date · Coupled to components 17-19

Thresholds live in the Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds. v4.0 contains zero moderation, ad-policy, or claims rules today, so this SOP issues zero verdicts from the closed vocabulary; every numeric control this component needs is logged on the v4.1 amendment register in Section 13, never written locally.

1. Verdict and Scope

VERDICT: this SOP is the account's single owner of ad-creative policy compliance, moderation handling, and claims review, the three duties the CrossReference Ledger row for #42 names and no shipped SOP performs. Until this document, moderation rejections routed interim to the PPC Lead with no procedure behind the route (Ledger row #17, interface I-9, which retires on this SOP landing), and the standing company claims rule for bamboo textile copy ("viscose derived from bamboo" or "rayon derived from bamboo", never raw fiber claims, per the PPC SOP Program Gap Register v1.0 Section 3 #42 entry) was enforced on listings, packaging, and A+ but on zero ad surfaces. Ad headlines and video scripts are exactly where a claims violation recurs, at 15-20% of mature ad spend on the SB side alone (Strategic Doctrine v1.0 Section 7 heading). This SOP closes that gap with 1 standing instrument (the Claims and Policy Register), 10 procedures, 8 failure modes, and 11 interface contracts, and it does so while issuing zero closed-vocabulary verdicts, because v4.0 carries no rule class for this domain: the SOP-43 v1.0 precedent, applied a second time.

Scope, stated as three lanes.

- Lane 1, pre-submission claims compliance: every SB headline, SBV script, SD headline, custom image, and Store spotlight copy passes the P2 checklist against the Claims and Policy Register before any submission, so a claims violation is caught at \$0 spent rather than in moderation or in an audit.
- Lane 2, the moderation lifecycle: detection of rejections and policy-driven serving stops, root-cause classification into 5 closed classes, correction, resubmission, the 2-rejection escalation trip from the component charter, and the appeal path.
- Lane 3, standing policy state: the restricted-term list, the keyword-level policy blocks, the do-not-target exclusion content that couples to SOP-16, platform and company policy-change response, and the monthly audit line the charter sets: zero live ad assets carrying non-compliant claims.

Depth note per the Build Kit. The Gap Register scoped this component at 3-4 pages under the superseded 10-element standard (conflict C1). The Build Kit v1.0 standard governs this build. Procedure count lands at 10 against the exemplar's 12: SOP-12 governs the account's highest-frequency surface and #42 governs a narrower compliance surface; 10 is every genuine procedure this component owns, and nothing here is padded to reach a count.

1.1 Boundaries With Adjacent Components

Adjacent component	The seam, stated so it is owned
Criteria System v4.0 (#8)	Decides every threshold, band, scenario, and verdict in the library. It carries zero moderation or claims rules today (verified against the attached file, Sections 3-14); therefore this SOP issues zero closed-vocabulary verdicts, and every needed numeric control is a Section 13 amendment item per the rule amendment protocol in v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)", rule V3.
#17-19 SB/SBV/SD playbook	Owns creative execution: format roles, builds, the creative precondition gate, S-F2/S-F4 scenario execution, and the SBV arbitrage workflow. This SOP supplies the P2 claims PASS its precondition gate consumes and takes over its interim moderation route (its row I-9). Per its Ledger row, this SOP restates zero creative-execution steps.
#16 PAT and conquest	Owns target opens, rotation, and exits. This SOP owns the policy-restricted portion of the exclusion-list content (P8), answering the open owner question SOP-16's row logged as its AM-6, coupled to #42. #16 decides and executes targeting; this SOP supplies the screen.
#29 Negation architecture	Owns negative placement, wall standards, and shared-list rules. This SOP rules restricted-term content; #29 places the resulting negatives in its surfaces per its own rules.

Adjacent component	The seam, stated so it is owned
#12 Exact operations	Owns SP campaign mechanics. Receives nothing from this SOP except incident freeze notices on affected targets; on "ENTER SBV ARBITRAGE" (Keyword Decision Record Template v3.0, Bands & Verdicts Reference), a rejected SBV creative suspends the paired step-down until approval lands, and the step-down itself stays SOP-12's per its row.
#32 Deal and event windows	Owns the event plan, arming, ramp, and reconciliation. This SOP supplies the creative-readiness check at arming (P9) and same-day in-window rejection handling; the plan and its caps stay #32's.
#38 Bulk mechanics	Owns staging, upload, and T+1 verification. This SOP reads the T+1 report as detection evidence only and stages nothing through it except corrected-state confirmations already owned by the affected SOP.
#37 Automation rules	Owns the rule register. No platform rule may act on an asset inside an open moderation incident; this SOP feeds the exclusion state, #37 owns the rules, the SOP-15 P11 pattern.
#43 Billing continuity / #46 Access	Account-level billing failures are #43's exclusively. Access to moderation surfaces and the role notification inbox is provisioned per #46's Access Register; this SOP consumes, never grants.
Catalog, listings, A+ (Brand Management)	Listing-side claims enforcement stays with its current owners. This SOP covers ad creative only; a claims defect found on a listing routes to BM and is not handled here.
Legal / company standing rules	The company claims rules (the bamboo wording rule and its class-mates) are set outside PPC. This SOP enforces them on ad surfaces and never interprets a new claims class on its own; an unregistered claims class routes to the PPC Lead with the Creative seam for a ruling.

2. Trigger Conditions

VERDICT: 8 triggers, each with a detecting artifact and a maximum lag. The same-day lags borrow the alert standard v4.0 already sets for the adjacent signal class (Section 3 "Layer M: The Master Metric Criteria Dictionary (F1-F16)", family F4 "Listing, Creative & Review State": eligibility flips alert same day; Section 5 "Layer O: Gates (G1-G12). Nothing Proceeds Until These Pass", G4: the fix task opens same day). A formal lag table for this component is owed to v4.1 (Section 13, AM-4), the same gap the playbook logged for itself.

Trigger	Detecting artifact	Maximum lag to procedure start
New creative submission upstream	The playbook build session presents the asset set for P2 before submission	Zero by construction: P2 is a precondition of submission, not a response to it
Weekly cadence	T3 CAMPAIGN BOARD refreshed from the SB/SBV/SD bulk (Data Ops lands the Raw Pull Pack Monday 06:00; workbooks complete by 09:00 per Data Architecture v1.0 Section 12)	Same session: P3 runs inside the Monday cascade
Moderation rejection or policy serving stop	Role notification inbox; T3 C2 State and Eligibility + Reason columns; #38 T+1 report on creative-carrying launches	Same day on any asset attached to a live push, an armed event window, or an arbitrage pairing; next Monday pass otherwise (interim rule pending AM-4)
Account-level policy notice or warning	Ads console / seller account policy notification to the role inbox	Same day, to P7, always: an account-level notice is never a weekly item
Platform or company claims-rule change	Amazon policy update notice; Creative seam or Legal notice on a company rule	Register update before the next submission; live-asset sweep opens same week (P7)
Event window arming	#32 event plan declares the armed window into T2 (the SOP-12 P9 arming pattern)	P9 readiness check at arming; daily in-window checks for the window duration
Register ruling request	A P2 hold (unregistered claims class) or a P8 referral files a ruling request line in the register change log	Ruled before the affected submission proceeds; the PPC Lead with the Creative seam rules open requests inside the week at latest, so a held asset never ages silently
Monthly audit slot	The monthly review slot per v4.0 Section 14 "Layer 9: Verification and Learning (V1-V8)", rule V4	P10 runs in the monthly slot; the #40 cadence document assigns the calendar day when it lands (AM-8)

3. Preconditions and Gates

VERDICT: 5 preconditions, each with the observable that proves it passed. A missing precondition suppresses its dependent procedure class and is recorded, never worked around: the G8 packet-completeness pattern (v4.0 Section 5) applied to a compliance surface.

Precondition (manifest ID)	Requirement and passing observable
A2 Deployment state	T3 C2 State column current from this week's SB/SBV/SD bulk join. Observable: T7 MANIFEST row A2 reads PRESENT with this Monday's refresh date. Missing: P3 runs on the notification inbox only and the pass records the degraded read.
A7 Eligibility	Eligibility + Reason current per campaign in T3 C2. Observable: T7 row A7 PRESENT, this week. Missing: rejection detection falls back to inbox plus console spot-read; the pass records it.
Claims and Policy Register current	This SOP's own instrument (P1), version-stamped, last-review date inside the monthly cycle. Observable: the register header version cell and review date. Missing or stale: P2 cannot issue PASS; submissions hold, which is the designed fail-safe direction. The register carries no Catalog manifest ID yet; the Definition Register entry is owed (AM-9).
Access provisioned (#46)	The executing role holds ads-console access including campaign manager and the role notification inbox. Observable: the #46 Access Register filtered view shows the role's row live for both entity accounts.
Event-plan visibility (#32)	Armed windows readable where the account reads them: the confirmed window dates entered into T2 by the #32 plan (the SOP-12 P9 arming pattern). Observable: the armed-window record with dates present for any product carrying event creative. Missing: P9 cannot run its arming check and reports exactly that to the #32 owner; the arming decision is #32's to hold or take.
Gates in force (v4.0 Section 5)	G4: an ineligible or suppressed ad halts all lever action for the target and the fix task opens same day; no procedure below bids, budgets, or judges rates on an asset inside an open incident. G10: delivery-outage windows this SOP flags are excluded from trend verdicts. Observable: the T5 PRODUCT LOG row with its Trend-Exclusion Flag set (P4 step 6).

4. Step-Level Procedures

VERDICT: 10 procedures. P1 stands up the single instrument everything else reads; P2 is the gate that catches violations at \$0 spent; P3-P6 are the moderation lifecycle from detection to appeal; P7-P8 hold the standing policy state current; P9-P10 are the event overlay and the charter audit line. Every procedure names its artifacts, columns, and navigation so a first-time operator executes from the page; every decision inside them is either a register ruling routed to a human checkpoint or a mechanical check, never a judgment call dressed as a step.

P1. Build and maintain the Claims and Policy Register (first build 90 min; weekly upkeep 15 min)

1. Create the register as one workbook, four tabs, stored in the instrument library beside the Command Workbook set, named **Claims_and_Policy_Register_v1_0.xlsx**, version in the header cell, never the filename beyond major versions.
2. Tab 1, CLAIMS REGISTER, one row per claim, columns: Claim ID (CL-001 up); Claim text as used (verbatim); Class, closed list of 6: fiber-content / performance / pricing-promo / comparative / health-adjacent / certification; Brands and products covered; Approved wording (verbatim); Banned wording (verbatim); Substantiation source + date; Ad surfaces where live (SB headline / SBV script / SD headline / custom image / Store spotlight); Status, register vocabulary of 4: APPROVED / RESTRICTED / BANNED / PENDING-SUBSTANTIATION; Review date; Owner.
3. Seed row CL-001, the standing company rule from the Gap Register charter: approved wording "viscose derived from bamboo" and "rayon derived from bamboo"; banned wording: any raw-fiber claim, including "bamboo fiber", "100% bamboo", and "bamboo fabric" used as fiber content; class fiber-content; status APPROVED for the two approved strings, BANNED for the raw-fiber strings; surfaces: all five. This row is the reason the component exists and it is never edited without a Creative seam + PPC Lead ruling.
4. Seed the performance class: every cooling, temperature, or sleep-quality claim used in ad copy enters with its substantiation source or enters as PENDING-SUBSTANTIATION, and a PENDING claim never ships (P2 step 5).
5. Seed the pricing-promo class: discount and deal claims are legal in ad copy only inside a platform-confirmed deal window (P9 consumes this row).

6. Seed the comparative class as an empty class with a standing rule row: no comparative claim and no competitor naming ships without a named PPC Lead + Creative seam ruling recorded on its own row; the class starts empty on purpose, because every entry is a ruling, never a habit.
7. Seed the certification and health-adjacent classes: any certification named in ad copy (fabric standard, safety mark) carries the certificate reference and expiry date as its substantiation source; a lapsed certificate flips the row to BANNED on its expiry date mechanically, which is why the date is a column and not a note. Health-adjacent sleep claims follow the performance class rules with the Creative seam ruling mandatory on first registration.
8. Tab 2, SUBMISSION LOG, append-only, one row per P2 run: Date; Asset ID; Campaign name per the naming convention; Surfaces; Claim IDs touched; Checker role; Result PASS / RETURNED; Returned-for Claim ID.
9. Tab 3, MODERATION INCIDENTS, append-only, one row per rejection or policy stop: Incident ID (MI-001 up); Date; Asset; Campaign; Rejection reason verbatim; Root-cause class (P4); Rejection count on this asset; State, register vocabulary: OPEN / CORRECTED / APPEALED / CLOSED; Case ID; T5 row link; Closure date.
10. Tab 4, RESTRICTED LIST (P8): restricted copy terms; keyword-level policy blocks (terms the account will not bid for policy reasons); do-not-target ASIN exclusions (siblings and declared partner ASINs, per the Gap Register fold-in to #16); each row with basis, date, and the receiving SOP (#29 or #16) plus receipt date.
11. Weekly upkeep inside P3: review dates due this week are confirmed or re-ruled; every change writes a change-log line (date, row, prior, new, ruler role).
12. The register is upstream of every asset, always: wording changes land in the register first and in assets second, so there is exactly one source of approved language (the design-rule-2 pattern from Data Architecture v1.0 Section 2).

P2. Pre-submission claims checklist (10-20 min per submission set)

1. Trigger: any playbook build or edit that creates or changes creative text or imagery: SB headline, SBV script and on-screen text, SD headline, custom image overlay text, Store spotlight copy. The playbook presents the asset set; its creative precondition gate consumes this procedure's result and is not restated here (Ledger row #17).
2. Inventory every text element of the set into a working list: headline strings, script lines, image overlay text, display URLs to Store pages.
3. Exact-string scan of every element against Tab 1 banned wording. Any hit: result RETURNED with the Claim ID; the asset goes back to the Creative seam; no partial passes on a set that shares a campaign.
4. Fiber-content check on every bamboo-range asset regardless of scan result: the copy states one of the two approved wordings wherever fiber content is stated or implied; imagery implying raw-fiber content (bamboo stalks presented as the fabric source) routes to the Creative seam for a ruling and the asset holds.
5. Claims-class walk on every remaining claim-bearing element, the decision tree:
 - Claim matches an APPROVED register row: pass the element.
 - Claim matches a BANNED or RESTRICTED row: RETURNED, cite the row.
 - Performance claim with no register row: enter it as PENDING-SUBSTANTIATION; RETURNED until the substantiation source lands and the row is ruled APPROVED.
 - Pricing-promo claim: pass only when P9 confirms a platform-confirmed deal window covering the flight dates; otherwise RETURNED.
 - Comparative claim or competitor naming: RETURNED to the Restricted List ruling path; comparative copy never ships without a named ruling.
 - Fiber-content, health-adjacent, or certification claim with no row: hold the asset; open a register ruling with the Creative seam and the PPC Lead; never pass by judgment.
 - Descriptive non-claim text (color, size, set count): pass; add a register note only when the checker had to think about it, which is the signal a class ruling is owed.
6. Write the SUBMISSION LOG row. PASS releases the set to the playbook's submission step; RETURNED names the failing Claim ID and the set does not submit.
7. The checklist is per submission event, not per week: an asset edited after a PASS re-runs P2 in full, because moderation re-reviews edited assets and so does this SOP.

P3. Weekly moderation status pass, inside the Monday cascade (15 min in a clean week)

1. Run after the workbook refresh completes (by 09:00 per Data Architecture v1.0 Section 12), and after T0/T1 are read, per the cascade order every weekly pass obeys.
2. Open T3 CAMPAIGN BOARD; filter Type to SB, SBV, SD. Read two C2 columns per row: State, and Eligibility + Reason.
3. Flag every row whose state is a moderation state (in review, rejected) or whose eligibility reason names a policy cause. Cross-check the role notification inbox for moderation and policy mail since the last pass; the inbox catches what a stale bulk misses, and the bulk catches what a misrouted mail misses.
4. Console verification on every flagged row, per entity account and per format: in the ads console, open Sponsored Brands (or Sponsored Display) from the campaign-type filter, open the flagged campaign, open its Ads view, and read the per-creative status and

the stated reason. SBV creatives carry their status on the video creative entry; SD creatives on the ad entry. The console status is the authoritative current state; the bulk is last pull's state; a disagreement means the event happened inside the week and the console wins.

5. Check the console notification center (the bell) and the account health / policy notices surface on both entity accounts for anything the mail route dropped; a notice found here but not in the inbox is an I-2 routing defect logged the same day.
6. Every rejected or policy-stopped asset opens P4 in the same session. Every asset sitting in review since before the prior pass opens a status check through the case path if it is push-coupled or window-coupled, and a log line otherwise.
7. Confirm closure evidence on incidents in CORRECTED state: T3 State back to delivering, or the T+1 report from #38 confirming the corrected launch landed.
8. Run the P1 step 9 register upkeep.
9. Write the pass line to the SUBMISSION LOG: date, rows read, incidents opened, incidents closed. A zero-incident week still logs the pass; the empty line is the passing state, the SOP-12 P9 deviation-log convention.

P4. Rejection intake, root-cause classification, and routing (20-40 min per incident)

1. Open the MODERATION INCIDENTS row: capture the rejection reason verbatim, the asset, the campaign, the timestamp, and the rejection count on this asset (Tab 3 history, exact asset match). Capture from the source with the fuller text: the console per-creative status detail usually carries the reason string and the policy area name; the notification mail sometimes carries a policy link the console omits; record both when both exist. Screenshot the console detail into the incident folder named by Incident ID, because rejection text is not exportable and the case path (P6) needs it verbatim.
2. Check the trip first: rejection count 2 on one asset routes to P6 before any further correction attempt, per the component charter (Gap Register v1.0 Section 3, #42 entry; the numeric trip rides the v4.1 register as AM-3, not as a local threshold).
3. Classify root cause into the closed class list of 5, the decision tree:
 - (a) Claims content: the reason names a claim, a superlative, a health or safety statement, or restricted language. Route to P5, this SOP's own correction path.
 - (b) Creative-technical: specs, resolution, logo, text-over-image ratio, video length or audio. Route to the Creative seam for the asset fix; this SOP tracks the incident and re-runs P2 on the fixed asset before resubmission, because a technical fix can introduce a claims change.
 - (c) Landing mismatch: the ad conflicts with the Store page or detail page it lands on. Route to the page owner via Brand Management; the ad-side incident holds OPEN until the page or the ad is aligned; which one changes is the page owner's call with the PPC Lead.
 - (d) Policy category: the product or format is restricted as a category matter. Route to the PPC Lead: contest through P6 with the written policy text, or retire the asset class and record the retirement in the register.
 - (e) Unclear or generic reason: one clarification request through the case path (P6 step 2 mechanics), then classify on the response; an unanswered clarification defaults to class (a) treatment, the conservative branch.
4. Set incident state OPEN and start the daily check cadence: while any incident is OPEN, the role inbox and the asset's console status are read daily, not weekly.
5. Freeze notice: the affected target takes no lever action while the incident is OPEN, the G4 pattern (v4.0 Section 5); notify the owning SOP's operator on push-coupled or arbitrage-coupled targets the same day.
6. Attribution write: on any delivery-affecting incident, write the T5 PRODUCT LOG row the same day, event type "eligibility flip" from the closed list of 8 (Data Architecture v1.0 Section 9), detail naming the moderation event, Trend-Exclusion Flag set. This is what lets the attribution rules (v4.0 Section 4 "Layer X", X9; Section 7, D5) see the stop instead of blaming bids. A dedicated moderation event type is owed to the T5 closed list (AM-6).

P5. Claims correction and resubmission (30-60 min per incident)

1. Correct the copy strictly to the APPROVED wording in the register row. No improvisation at the asset: a claim with no approved wording gets one drafted by the Creative seam, ruled by the PPC Lead, written to the register, and only then written into the asset (P1 step 10 order, always).
2. Re-run P2 on the corrected asset in full; log the new SUBMISSION LOG row referencing the incident ID.
3. Resubmit through the owning playbook build procedure (Ledger row #17; not restated here). Set incident state CORRECTED.
4. Verify on the daily check: approval landed, T3 State delivering, and on relaunch-class corrections the #38 T+1 report confirms the landed values.
5. Close: state CLOSED with closure date; the closure line notes what register change, if any, the incident produced, because an incident that teaches the register nothing repeats.
6. If the resubmission is rejected, return to P4 step 2: that is the second rejection on the asset and the trip fires.

P6. Appeal and case escalation (30 min to open; daily checks while open)

1. Eligibility to appeal, any one of: the 2-rejection trip fired; the rejection is class (e) after classification; or the rejection is judged wrong against the written policy text by the PPC Lead. An unchanged asset is never resubmitted a second time without a case: the

identical-resubmission loop is the F2 pattern, it burns the appeal path and invites account-level flags.

2. Open the support case carrying: asset ID and campaign; the rejection reason verbatim; the policy section quoted; the register Claim ID with its substantiation source and date. A case that cannot cite its substantiation is not opened; it goes back to P5.
3. Record the Case ID in the incident row; state APPEALED; daily checks.
4. A case unresolved at the next weekly pass escalates to the PPC Lead with the incident history; a case touching an armed event window escalates same day (P9).
5. Write the outcome back to the register: an upheld rejection turns the disputed wording BANNED account-wide with the case cited as basis; an overturned rejection records the ruling as precedent on the claim row. Either way the register moved, which is the point of the path.

P7. Policy and claims rule-change response, platform and company (30-90 min per change event)

1. Intake, two sources, one flow: (branch 1) a platform policy update or account-level policy notice arrives at the role inbox; (branch 2) the Creative seam or Legal changes a company claims rule. Both are same-day intakes; an account-level warning is read by the PPC Lead the day it arrives, without exception.
2. Assess which register classes and rows the change touches; draft the register edits; the PPC Lead rules them with the Creative seam.
3. Update the register first (P1 step 10 order), with the change-log line naming the source notice.
4. Open the live-asset sweep the same week: pull the live SB/SBV/SD asset inventory (P10 step 2 mechanics) and check every asset against the changed rows only; hits open P4-class incidents with correction staged same day.
5. Route restricted-list consequences per P8 where the change adds blocked terms or targets.
6. Log the change event in T6 DECISION LOG with the sweep result and a review date; an account-level notice additionally writes a T5 row, because it is an account event the attribution layer must see.

P8. Restricted-term and exclusion-list maintenance (20 min monthly; 60 min quarterly review)

1. This SOP owns the content and the ruling of Tab 4; it executes none of it in campaign surfaces. Three routes, each with a receipt:
 - Restricted copy terms: enforced by P2 exact-string scan; no external route.
 - Keyword-level policy blocks: routed to #29 under its shared-list rules the same week as the ruling; #29 places, this SOP records the receipt date.
 - Do-not-target ASIN exclusions (siblings and declared partner ASINs never conquered, the Gap Register fold-in to #16): routed to #16 the same week; #16 screens its opens against it; receipt recorded. This resolves the open owner question in SOP-16's row (its AM-6): the owner is this SOP, the PPC Lead ruling.
2. Additions come from: P2 comparative-claim returns; P4 class (d) retirements; P7 change events; #16 or #29 referrals; the partner and sibling declarations from #33's family records.
3. Quarterly, in the v4.0 Section 14 V4 quarterly review class: walk the full list; every row is re-based or retired; a list change without its receipt from #16 or #29 inside the week is the F8 condition.
4. Every list change writes the change-log line and, where a live target or term is affected, a T6 entry with the routing receipt as evidence.

P9. Event-window creative readiness and in-window handling (20 min at arming; 10 min daily in-window)

1. At #32 arming of a confirmed window: pull the event plan's creative-carrying line items; verify every asset is APPROVED in the register, live-eligible in T3 C2, and, for pricing-promo copy, that the deal is platform-confirmed for the exact flight dates. Report the check result to the #32 owner at arming (interface I-8); the arming decision stays #32's.
2. Pricing-promo assets submit only after the deal is platform-confirmed: submitting against an intended-but-unconfirmed deal is the WE-2 failure and the register carries it as a standing rule row after that incident class first fires.
3. Inside the window plus the 14-day pre and post posture windows the account already tracks (v4.0 Section 3, family F2, Event posture row), the moderation check cadence is daily, and any rejection is a same-day P4 regardless of the asset's coupling.
4. An in-window rejection that cannot be corrected and re-approved before its flight date routes to the event plan's fallback the same day: the format-role fallback belongs to the playbook and #32 (their rows); this SOP's job is that the fallback decision is taken on day 0, not discovered at reconciliation.
5. No new claims wording enters the register during an armed window except through an incident: the S-E1 no-new-experiments posture (v4.0 Section 8.12 "Event-Mode Scenarios (S-E1 to S-E6)") applied to copy.

P10. Monthly claims compliance audit, the charter audit line (45-75 min monthly)

1. Runs in the monthly slot (v4.0 Section 14 V4 monthly class; calendar day assigned by #40 when it lands, AM-8).
2. Pull the live-asset inventory: T3 filtered to SB/SBV/SD with State delivering, plus the Store spotlight and custom-image asset list from the Creative seam, because a delivering campaign is the surface but the asset is the unit.
3. For every live asset, map every claim-bearing element to an APPROVED register row. The audit line, from the charter: zero live ad assets carrying non-compliant claims. Any miss opens a P4-class incident with correction staged same day, even where the platform

- approved the asset: platform approval never overrides the company claims rule (the F3 logic).
4. Cross-check the SUBMISSION LOG against the inventory: a live asset with no P2 row is the F4 condition and gets a retroactive P2 the same day plus a root cause naming the skipped build step.
 5. Confirm register hygiene: review dates current, PENDING-SUBSTANTIATION rows aged out or resolved, incident rows in terminal states.
 6. Write the audit result line to the register and the monthly compliance note to Account Rollup A4 beside the playbook's format rollout (Data Architecture v1.0 Section 11), so the monthly structure review reads format spend with its compliance line attached. Repeated misses on one source escalate to the PPC Lead with the log history, the SOP-12 escalation convention.

5. Verdict Mapping

VERDICT: this SOP issues zero verdicts from the closed vocabulary, and that is a designed property, not a gap in the build. The closed vocabulary is defined in v4.0 Section 13.1 "The Closed Verdict Vocabulary" and carried operationally in the Keyword Decision Record Template v3.0, Bands & Verdicts Reference sheet; a full-text search of the attached v4.0 for moderation, ad-policy, claims, or creative-compliance rules returns none, so no closed-vocabulary verdict can lawfully originate here (anti-drift rule 1). The states this SOP writes (register: APPROVED / RESTRICTED / BANNED / PENDING-SUBSTANTIATION; incidents: OPEN / CORRECTED / APPEALED / CLOSED; checklist: PASS / RETURNED) are register vocabulary in the SOP-43 v1.0 precedent, pending v4.1 admission or a formal standing exception (AM-2). Zero local thresholds appear in this SOP.

Three closed-vocabulary verdicts touch this scope without originating here, each quoted verbatim with its source:

Verdict (quoted verbatim)	Interaction with this SOP
"PAUSED-PENDING (REASON + RE-ENTRY)" (KDR Template v3.0, Bands & Verdicts Reference)	Where a moderation stop is the blocking reason on a term or campaign, the verdict is issued by the owning SOP under v4.0; this SOP supplies the reason string and the re-entry criterion: approval landed and delivery confirmed (P5 steps 4-5).
"ENTER SBV ARBITRAGE" (KDR Template v3.0, Bands & Verdicts Reference)	A rejected SBV creative suspends the entry until approval: the SP-side step-down does not execute against a dark SBV surface. Execution split stays per the Ledger rows for #12 and #17-19; this SOP supplies only the incident state.
Scenario S-F2 re-creative branch and S-F4 creative refresh (v4.0 Section 8.10 "Brand-Format Scenarios (S-F1 to S-F8)")	When v4.0 issues an exit-or-re-creative outcome, the replacement creative enters through P2 before submission. The scenario decision is v4.0's; the compliance pass is this SOP's.

6. Worked Examples

VERDICT: 3 examples: one clean lifecycle priced end to end, one edge case where the escalation trip fires inside an armed window, and one audit-line catch of a claim the platform approved. Together they exercise P2 through P10 and every register state at least once, and each one ends with the register or the rule set having learned something, because an incident that changes nothing repeats.

All three examples are illustrative and internally consistent, built on the standard fictional base (AOV \$80.00, contribution \$28.00 per order, break-even ACOS 35.0%, Target 17.5%, Max Acceptable 26.3%) and the Economics Worked-Example Set v1.0 E6 format values (SP exact CPC \$2.10 at CVR 5.0%; SBV CPC \$0.95 at CVR 3.2%). Live values come from T0 and the live plans. Every figure below is computed on the page.

WE-1 · Clean case: SBV rejection on a raw-fiber claim, detected, corrected, priced

Setup. Tuesday 14:20, Amazon rejects a live SBV asset on the bamboo range: video overlay reads "100% bamboo softness", a raw-fiber claim, register row CL-001 BANNED wording. Daily budget on the SBV line \$35.00. The SBV runs as the cheap side of an open arbitrage pairing per the E6 pattern.

Execution. Wednesday 09:00 daily-check detection (an incident-cadence read; lag 18h40m = 18.7h). P4: class (a) claims content, first rejection on the asset, no trip. T5 row written, Trend-Exclusion Flag set; freeze notice to the SP-side operator: the paired step-down holds. P5: overlay corrected to the approved wording "viscose derived from bamboo"; P2 re-run PASS; resubmitted Wednesday 11:00. Approved Friday 09:30 (platform review 46.5h); delivery restart confirmed in T3 the same morning; incident CLOSED.

The price of the dark window. Delivery dark Tuesday 14:20 to Friday 09:30 = 2.80 days. Undeployed SBV spend = 2.80 x \$35.00 = \$98.00. Expected SBV clicks forgone = 98.00 / 0.95 = 103.2; expected orders = 103.2 x 0.032 = 3.30. Cost per order: SBV = 0.95 / 0.032 = \$29.69; SP exact = 2.10 / 0.05 = \$42.00; arbitrage premium = 42.00 - 29.69 = \$12.31 per order. With SP absorbing the demand during the dark window, excess acquisition cost = 3.30 x 12.31 = \$40.62, or 40.62 / 2.80 = \$14.51 per dark day.

The lesson. Had detection waited for the Monday weekly pass (Tuesday 14:20 to Monday 09:00 = 5.78 days dark before correction even starts), the same arithmetic scales to $40.62 \times 5.78 / 2.80 = \83.85 of premium plus the correction days on top. The daily incident cadence is worth \$14.51 per day on one mid-sized pairing, which is why P4 step 4 exists and why the lag table is owed to v4.1 as a rule, not a habit (AM-4).

WE-2 · Edge case: second rejection inside an armed event window; the trip fires and the fallback is priced

Setup. A 2-day deal window (Friday-Saturday) is armed by #32. The SB category headline asset carries the pricing claim "Sale: 30% off". P9 passed at arming on the plan as drafted. Wednesday 10:00, the asset comes back rejected for the second time: rejection 1 (Monday) was class (b), a logo spec; the corrected asset's rejection 2 names the pricing claim, because the deal is not yet confirmed platform-side even though it is booked internally.

Execution. P4 step 2: rejection count 2 on one asset, the charter trip fires, route to P6 with no third blind resubmission. Case opened Wednesday 11:00 carrying the deal booking confirmation and the register pricing-promo row. Same-day escalation to the PPC Lead under P9 step 4, window minus 2 days. Fallback decision taken Wednesday, not Friday: if the case does not clear by Thursday 17:00, the event runs SP-side per the event plan's format fallback (owned by #32 and the playbook; referenced, not restated).

The exposure, priced. Event plan SB budget \$600.00 per day $\times$ 2 days = \$1,200.00. At illustrative event-window figures (SB CPC \$1.10, event CVR 6.5% against the 5.2% base): planned clicks = $1,200.00 / 1.10 = 1,090.9$; expected orders = $1,090.9 \times 0.065 = 70.9$; expected contribution = $70.9 \times \$28.00 = \$1,985.20$; expected net = $1,985.20 - 1,200.00 = \$785.20$. That \$785.20 is what a dark SB slot puts at risk across the window, and it is on the table Wednesday because the trip forced the case path instead of a third rejection on Thursday night.

Resolution. The deal state syncs platform-side Thursday 14:00; the case clears 16:00; the asset approves; Friday 06:00 arming re-check reads APPROVED and delivering. Incident CLOSED with the precedent line: pricing-promo assets submit only after the deal is platform-confirmed. The register gains the standing rule (P9 step 2) and the incident taught the system something it keeps.

WE-3 · Edge case: the audit line catches an approved-but-non-compliant claim

Setup. P10 monthly audit. An SD line has been delivering 24 days with the headline "cools you 3x faster", a performance claim with no register row: it entered through a build session that skipped P2 (the F4 condition), and moderation approved it, which settles nothing about compliance (the F3 logic).

Execution and exposure accounting. The audit line fails: 1 live non-compliant asset against a required zero. Incident opened; correction staged same day: claim withdrawn from the copy pending substantiation, asset resubmitted with approved descriptive copy. Exposure recorded in the incident row for any downstream ruling: 24 days $\times$ \$22.00 per day = \$528.00 of spend carried on the claim; at CPC \$1.15, clicks = $528.00 / 1.15 = 459.1$; at CVR 4.0%, orders = 18.4; attached contribution = $18.4 \times \$28.00 = \515.20 . Root cause: the build session that skipped P2, named in the log; second occurrence from the same source escalates.

The lesson. Moderation is a floor set by the platform for the platform. The audit line is the company's own bar, it is absolute (zero, from the charter), and it is the only mechanism that catches what moderation happily approves.

6.1 Computation summary (every figure above, one table; asserted in code at QA)

Figure	Computation	Value
WE-1 dark window	Tue 14:20 to Fri 09:30 = 67.17h / 24	2.80 days
WE-1 undeployed spend	$2.80 \times \$35.00$	\$98.00
WE-1 clicks / orders forgone	$98.00 / 0.95; \times 0.032$	103.2; 3.30
WE-1 cost per order, SBV / SP	$0.95 / 0.032; 2.10 / 0.05$	\$29.69; \$42.00
WE-1 premium per order / window / per day	$42.00 - 29.69; 3.30 \times 12.31; / 2.80$	\$12.31; \$40.62; \$14.51
WE-1 weekly-pass counterfactual	$40.62 \times 5.78 / 2.80$	\$83.85
WE-2 window budget	$\$600.00 \times 2$	\$1,200.00
WE-2 clicks / orders / contribution	$1,200.00 / 1.10; \times 0.065; \times \28.00	1,090.9; 70.9; \$1,985.20
WE-2 expected net at risk	$1,985.20 - 1,200.00$	\$785.20
WE-3 spend on the claim	$24 \times \$22.00$	\$528.00
WE-3 clicks / orders / contribution	$528.00 / 1.15; \times 0.040; \times \28.00	459.1; 18.4; \$515.20

7. Failure Modes and Escalation

VERDICT: 8 failure modes. This is a first build, so account-history incidents do not exist yet for most rows; where the motivating condition is structural rather than historical, the row says so plainly instead of inventing history.

Failure	Detection signal	Response and motivation
F1. Rejection sits undetected	T3 C2 State scan at P3; role-inbox daily reads during incidents; spend flat at zero delivery on a format row	P3/P4 same session. Motivation: structural, and documented: the playbook's row I-9 routed rejections interim to the PPC Lead precisely because no detection owner existed before this SOP.
F2. Unchanged resubmission loop	SUBMISSION LOG shows the same asset resubmitted with identical wording after a rejection	Block; route to P6. An identical resubmission burns the appeal path and invites account-level policy flags; the trip at 2 rejections (charter, AM-3) exists to make this mechanically impossible to drift into.
F3. Approved-but-non-compliant claim live	P10 audit line; P7 sweep after a rule change	Correct same day even though approved and delivering; platform approval never overrides the company claims rule. Motivation: the charter's origin, the FTC-anchored bamboo rule enforced everywhere except ads.
F4. Register bypass (no P2 row on a live asset)	P10 cross-check of SUBMISSION LOG vs live inventory	Retroactive P2 same day; root cause names the skipped build step; a recurring source escalates to the PPC Lead with the log history.
F5. Event asset rejected inside an armed window	P9 daily in-window check; #32 pacing anomaly on the format line	Same-day P4; fallback decision same day per P9 step 4; trend exclusion per G10. WE-2 is the walked case.
F6. Silent delivery stop misread as performance	A format row's WoW spend collapse with no T5 event row behind it	Write the missing T5 row; mark the affected reads excluded; the root cause is the P4 step 6 write that did not happen. This is the X9/D5 attribution discipline applied to moderation.
F7. Rule change not propagated to live assets	Register change-log date with no same-week sweep line behind it (P7 step 4)	Run the sweep immediately; incidents for hits; the change-event T6 entry is incomplete until the sweep line exists.
F8. Restricted-list drift	Quarterly P8 review diff; a routed change with no #16/#29 receipt inside the week	Re-route with receipts same week; an exclusion ruled here but not landed in the executing SOP protects nothing. Motivation: structural; SOP-16's AM-6 recorded the ownerless state this SOP ends.

Escalation. The same failure mode twice in consecutive cycles on the same scope escalates to the PPC Lead with the log history attached. Any account-level policy notice, any F5, and any fired 2-rejection trip escalates same day. Anything in the human-mandatory tier (Section 9) is confirmed before it executes, without exception.

8. Outputs and Artifacts

VERDICT: every procedure closes with named artifacts in named locations, because gates clear with artifacts, not assertions. The register and its three logs are append-only and permanent: a compliance function whose history can be edited has no history.

Event	Artifact set produced (gates clear with artifacts, not assertions)	Storage, retention, consumer
P1 register build and upkeep	Claims_and_Policy_Register_v1_0.xlsx, four tabs, version-stamped, change-logged	Instrument library beside the Command Workbook set; permanent, append-only logs; consumed by the playbook precondition gate, #16, #29, P10, and #44 (Element 11 content)
P2 checklist runs	SUBMISSION LOG rows, one per run, PASS / RETURNED with Claim IDs	Register Tab 2; permanent; consumed by P10 cross-check and by the playbook build session
P3-P6 incidents	MODERATION INCIDENTS rows with class, states, case IDs, closure notes; T5 PRODUCT LOG rows on delivery-affecting events; freeze notices to owning SOPs	Register Tab 3 and workbook T5; permanent; consumed by X9/D5 attribution reads and by the weekly pass of every affected SOP
P7 change events	Register change-log lines; sweep result; T6 DECISION LOG entry with review date; T5 row on account-level notices	Register and workbook T6; permanent; consumed by the monthly review and the v4.1 amendment owner

Event	Artifact set produced (gates clear with artifacts, not assertions)	Storage, retention, consumer
P8 list changes	Tab 4 rows with basis and receipt dates; T6 entries on live-target effects	Register Tab 4; permanent; consumed by #29 and #16 in their surfaces
P9 / P10	Arming check result to #32; monthly audit line with the zero-or-miss count; monthly compliance note to Account Rollup A4	Register and A4; permanent; consumed by #32 arming and the monthly structure review

9. Skill-Conversion Block

VERDICT: the automatable core is the P2 exact-string scan, the two classification trees, and the P10 diff; every wording ruling, every appeal, and every account-level response stays human. The split below is this component's own 10-80-10 line, drawn from the Authority Matrix rows that govern in the absence of a #42 row.

Block	Specification, specific to this SOP
Input bindings	T3 CAMPAIGN BOARD C2 State and Eligibility + Reason (from the SB/SBV/SD bulk + RAS sheets, Monday 06:00 feed per Data Architecture v1.0 Section 12); T7 MANIFEST rows A2 and A7; the role notification inbox; Claims_and_Policy_Register_v1_0.xlsx all four tabs; the #32 event plan dates; the #46 filtered access view. An absent feed suppresses its dependent class per Section 3 and the pass records it; the skill never improvises around a missing feed.
Deterministic logic	P2 step 3 exact-string scan of asset text against Tab 1 banned wording (fully automatable); the P2 step 5 class tree and the P4 step 3 class tree as written; the P10 diff of live inventory vs SUBMISSION LOG; the P8 receipt-date check. The skill drafts results; it computes no thresholds, and every ruling on wording lives with humans.
Human checkpoints (from Authority Matrix v2.0, Section 1 "The Matrix" and Section 2 Standing Rules)	The Matrix carries no #42 row (AM-5); interim rights follow the documented route and the vacancy rule: register status rulings (APPROVED / BANNED), appeal filings, class (d) retire-or-contest calls, and every account-level notice response are PPC Lead decisions, per the playbook I-9 interim route and Standing Rule 5. Checklist and pass execution follow the "Structure creation (all SOPs)" row pattern: operators execute, the Lead decides scope. Event-window fallbacks follow the "Event plans (#32)" row: PPC Lead decides with the deals owner; launch/CA operator and PPC Manager execute; BM owner reviews. The BM owner's review is the daily huddle per Standing Rule 3: review, not veto.
Cadence	P3 weekly inside the Monday cascade; daily inbox and status reads while any incident is OPEN and through armed windows plus their 14-day postures; P10 monthly; P8 quarterly; P2 on every submission event; account-level notices same day, always.
Alarms	F1-F8, each with its Section 7 detection signal, an evidence line (asset, campaign, signal value, register or log reference), and routing: F1/F5 and every trip fire same day to the PPC Lead; F2/F6/F7 to the operator queue same session; F3/F4/F8 into the weekly pass worklist with monthly escalation on recurrence.
Skill outputs	Drafted incident rows and classifications queued for confirmation; drafted P2 scan results with PASS blocked pending human release on any RETURNED history; the P10 audit diff; F1-F8 alarms with evidence lines; drafted T5/T6 entries.

10. Version Governance

VERDICT: this SOP re-versions on policy events, not on a calendar. Owner: PPC Lead. Review triggers, specific to this component:

1. any v4.0 increment that admits moderation, claims, or ad-policy rules, at which point every AM item in Section 13 re-checks against the new text and this SOP's register vocabulary either enters the closed system or retires;
2. any Amazon ads policy change notice touching a register class (P7 both detects the event and triggers this review);
3. any change to a company standing claims rule from Legal or the Creative seam;

4. any version increment of the #17-19 playbook, the coupled component, at which point interfaces I-3, I-4, and I-5 re-validate against its new row;
5. the Authority Matrix re-issue that adds the #42 row (AM-5), at which point Section 9's interim rights retire;
6. closure of the first 5 live moderation incidents, which reviews the P4 class list against observed rejection reasons.

Registered against component #1 document governance; the register instrument itself is owed a Definition Register entry (AM-9). v1.0 change note: first build; no prior version; nothing carried forward.

11. Operator Ramp

VERDICT: the ramp is built from the Authority Matrix v2.0 rows that govern this scope in the absence of a #42 row: "Structure creation (all SOPs)" (operator executes, Lead decides), "Event plans (#32)" (Lead decides with the deals owner; operator and Manager execute; BM owner reviews), the "Weekly verdicts (established US)" huddle convention (Standing Rule 3: the BM owner's R is review, not veto; disagreements route up), and Standing Rule 5's vacancy rule for the interim decision rights. Certification content hands to #44 on its due date (interface I-10).

Week 1, observation and artifact familiarity. Read this SOP and the register cold; reproduce from memory the 6 claim classes, the 4 register statuses, the 5 rejection root-cause classes, and register row CL-001's approved and banned wordings verbatim. Sit the Monday P3 pass as observer; walk T3 C2 and the T7 A2/A7 rows; trace one closed incident end to end through Tabs 2-3, the T5 row, and the T6 entry.

Week 2, supervised execution. Run P2 on one live submission set and P3 on one Monday pass with the PPC Lead reviewing every PASS and RETURNED call before release. Classify one rejection (live or synthetic) through the P4 tree with the Lead confirming the class and the routing. Draft, but do not file, one P6 case package and have the Lead grade it against P6 step 2's required contents.

Week 4, unsupervised with audit. Run P2, P3, and P4 unsupervised for the week with the PPC Lead auditing the SUBMISSION LOG and incident rows at week end. Run one full P10 audit end to end, including the A4 note. Rulings, appeals, and account-level notices remain with the Lead regardless of ramp stage: those are decision rights, not competences, and they move only when the Matrix row lands.

Competence certified	Demonstrated by	Governing Authority Matrix v2.0 row
Checklist execution (P2) and pass execution (P3)	Week 2 supervised runs; week 4 unsupervised with log audit	"Structure creation (all SOPs)": operator E, Lead D on scope
Incident classification and routing (P4-P5)	One classified incident confirmed by the Lead; one solo by week 4	"Weekly verdicts (established US)" huddle convention; Standing Rule 3 (BM owner review, not veto)
Case package construction (P6)	One graded package, week 2; filing stays with the Lead	Interim per the playbook I-9 route and Standing Rule 5 (vacancy rule), pending the #42 Matrix row (AM-5)
Event-window conduct (P9)	Walk WE-2 aloud: the trip, the case, the day-0 fallback	"Event plans (#32)": Lead D with deals owner; operator and PPC Manager E; BM owner R
Audit-line execution (P10)	One full monthly audit end to end, week 4	"Structure creation (all SOPs)" pattern; escalations to the Lead

Live-defense questions (reasoning, never threshold recall):

- Why does platform approval settle nothing about compliance, and what single mechanism catches the difference?
- Why is an unchanged resubmission worse than a slow correction?
- Why does the register update before the asset, in every case, with no exception for urgency?
- Why does a mid-week delivery stop get a T5 row the same day, and which two v4.0 rules consume that row?
- Why does the restricted list route to three different components instead of executing here, and what breaks in the ownership model if this SOP places one negative itself?
- Why is the 2-rejection trip written as a charter control on the v4.1 list rather than as a number in this SOP?

12. Interface Contracts

VERDICT: 11 contracts, each with sender, receiver, artifact, timing, and the non-arrival consequence, because the missing consequence clause is the defect class this rebuild program exists to end. Sources: the Ledger row for #42 and the rows for #17-19, #16, #43, and #44, cross-checked against Data Architecture v1.0 for artifact names and timings.

ID	Sender to receiver, artifact	Timing	When it does not arrive
I-1	Data Operations to this SOP: SB/SBV/SD bulk + RAS state columns landing in T3 C2	Monday 06:00 feed; workbook by 09:00	P3 runs on the notification inbox alone; T7 marks A2/A7 stale; the pass records the degraded read and the week's detection claim is downgraded, the G8 pattern.
I-2	Platform to the PPC Manager role inbox: moderation and policy notifications	Continuous	Notifications landing in a personal inbox is the F1 standing condition; re-route same day via #46 provisioning; incidents found late carry the lag in their rows.
I-3	Creative seam to this SOP: asset text and imagery for P2	Before any playbook submission	The submission does not proceed: the playbook's creative precondition gate has no PASS to consume. Fail-safe direction is hold, never ship-unchecked.
I-4	This SOP to the playbook build session: P2 result, PASS or RETURNED with Claim ID	Same session as the check	The playbook holds the launch; a launch that proceeded anyway is an F4 finding at the next P10.
I-5	This SOP to the playbook: moderation-rejection handling (P4-P6), replacing its interim route	On this SOP landing, its row I-9 retires	If an incident gets no response inside one day, the interim route (PPC Lead direct) resumes as the standing fallback; the gap is logged as an F1 line.
I-6	This SOP to #29: keyword-level policy blocks and restricted-term rulings for wall placement	Same week as the ruling, with receipt	Walls run stale; the #29 monthly wall audit is the backstop detector; the missing receipt is the F8 condition here regardless of what #29 finds.
I-7	This SOP to #16: do-not-target exclusion content (siblings, declared partners, policy-blocked targets)	Same week as the ruling, with receipt	#16 opens targets unscreened; the S-C1 open test does not check policy exclusions, which is exactly why its AM-6 coupled to this component. Missing receipt is F8.
I-8	This SOP to #32: event creative-readiness check result at arming (P9)	At window arming; daily in-window updates	The window arms on unverified creative and carries WE-2 class exposure; the arming-checklist line that would hold the arm belongs in #32 and is proposed as a fold-in (AM-7).
I-9	This SOP to the workbook T5 PRODUCT LOG: event rows on delivery-affecting incidents	Same day as the incident	F6: the attribution layer blames bids for a policy stop; trend verdicts on affected rows are polluted until the row is backfilled and exclusions applied.
I-10	This SOP to #44 operator onboarding: Element 11 ramp and certification content	By #44's due date, 2026-09-25	#44 builds this component's certification blind; the consequence is recorded on #44's side per its row (it consumes Element 11 from every SOP).
I-11	This SOP to Account Rollup A4: the monthly compliance note beside the playbook's format rollup	Monthly, with P10	The monthly structure review reads format spend with no compliance line; the miss surfaces at the next P10 and the note backfills with both months covered.

13. Closing: v4.1 Amendment Register and Standing Findings (this build)

VERDICT: 10 amendment items, zero locally written numbers. Per v4.0 Section 14 rule V3, every rule this component needs that v4.0 does not carry is logged here for the threshold owner, never written into the body as a local number.

Item	Content
AM-1	v4.0 carries zero moderation, ad-policy, or claims rules. The whole rule class is owed: rejection-response requirements, appeal limits, incident severity classes, and the compliance gate's place in the gate order.
AM-2	Register vocabulary admission: APPROVED / RESTRICTED / BANNED / PENDING-SUBSTANTIATION, incident states OPEN / CORRECTED / APPEALED / CLOSED, and checklist results PASS / RETURNED enter the closed-vocabulary system or receive a formal standing exception (the SOP-43 AM-2 pattern, second instance).

Item	Content
AM-3	The escalation trip at 2 rejections on one asset: charter value from the Gap Register v1.0 Section 3 #42 entry, operating as a charter control pending v4.0 admission.
AM-4	The Element 2 lag table: same-day lags currently borrow the F4/G4 same-day alert standard; the moderation-specific maxima need a v4.0 basis (the playbook logged the same gap for itself).
AM-5	Authority Matrix v2.0 carries no #42 row. Interim rights per the playbook I-9 route and Standing Rule 5; the row is owed at the next Matrix re-issue.
AM-6	The T5 PRODUCT LOG closed event-type list of 8 has no moderation or policy event type; interim logging rides the "eligibility flip" type with detail text; a dedicated type is owed (the SOP-43 billing-outage pattern, second instance).
AM-7	Fold-in proposal to #32: an arming-checklist line that holds arming pending the I-8 creative-readiness result.
AM-8	Calendar slots for the P10 monthly audit and the P8 quarterly review belong to #40 Monthly, Quarterly and Annual Cadence, due 2026-09-18.
AM-9	The Claims and Policy Register is owed a Definition Register entry under component #1.
AM-10	Ownership resolution for the Ledger close-out: this SOP takes the policy-restricted exclusion-list content, answering SOP-16's open AM-6. SOP-16's row is not edited by this session; the disagreement-free resolution is reported for close-out per the accretion rule.

Appendix A. Operator Cards (compressed; the SOP text governs)

Weekly card (P3)

1. Monday, after 09:00 refresh, after T0/T1.
2. T3 filtered to SB/SBV/SD; read C2 State and Eligibility + Reason.
3. Inbox sweep since last pass; console verification on every flag, both accounts.
4. Every rejection or policy stop: P4 this session.
5. In-review beyond one pass: status check if push- or window-coupled.
6. Confirm CORRECTED incidents delivering (T3 or the T+1 report).
7. Register review dates due this week: confirm or re-rule.
8. Log the pass line; a zero-incident week still logs.

Incident card (P4-P6)

1. Capture verbatim + screenshot; count rejections on the asset.
2. Count reads 2: P6, no third blind resubmission, ever.
3. Classify into the 5 classes; route per the tree.
4. Freeze notice on the target same day; no lever action inside an open incident (G4 pattern).
5. Delivery affected: T5 row same day, Trend-Exclusion Flag set.
6. Daily checks while OPEN.
7. Correct to register wording only; register first, asset second.
8. Re-run P2 in full; resubmit through the playbook.
9. Verify approval and delivery; CLOSED with the lesson line.
10. Window-coupled anything: same-day escalation, fallback decided day 0.

Appendix B. Incident Evidence-Line Format (feeds the Section 9 alarms)

Every F1-F8 alarm and every incident row carries one evidence line in this fixed format, so the skill layer parses it and a reviewer audits it in seconds:

[Incident or F-ID] | [entity account] | [campaign name per convention] | [asset ID] | [signal: column or source + value] | [register or rule reference] | [state] | [next action + date]

Worked instance from WE-1:

MI-014 | Ecotero | EBS4-SBV-bamboo sheets king-HSV-Video-Rank-B4-KING-WHITE | VID-2231 | Console creative status: Rejected, claims policy | CL-001 BANNED wording | CORRECTED | verify approval, daily check, close on delivery.

A line missing its register reference is incomplete and does not close its incident.

Companions: Master PPC Decision Criteria System v4.0 (all thresholds; zero exist for this domain today) · Keyword Decision Record Template v3.0 (the closed vocabulary this SOP issues nothing from) · PPC SOP Program Gap Register v1.0 Section 3 #42 entry (the charter) · Cross-Reference Ledger rows #42, #17-19, #16, #43 (boundaries) · Master PPC Data Architecture v1.0 (T3, T5, T6, T7, A4) · PPC Authority Matrix v2.0 (interim rights) · Strategic Doctrine v1.0 Sections 7-8 (format context) · Economics Worked-Example Set v1.0 (arithmetic conventions).

Inspiratek & Ecotero LLC | Confidential | Page 10 of 10