SOP-45: Dayparting Deferral Control

Operating SOP, Component #45

v1.0 · First build · Owner: PPC Lead · Threshold authority: Master PPC Decision Criteria System v4.0 only; this SOP carries zero local thresholds · Roles, never names

1. Verdict and Scope

VERDICT: dayparting is deferred, not undecided, and this document governs the deferral rather than the practice. v4.0 Section 9 "Layer 4: The Lever Library (L1-L10). Order, Limits, Mechanics" defines L10 as "Dayparting (new, dormant until data): plus or minus 15% hourly modifiers, deployed only on Marketing Stream curves with 4 weeks minimum of hour-level evidence per campaign (Catalog D12). Until the feed lands, no dayparting decisions are made on guesswork." The feed has not landed. The Command Workbook T7 MANIFEST row "D12 Hourly dayparting" reads Source Feed "Marketing Stream", Refresh "Deferred", Status "MISSING". Master PPC Data Architecture v1.0 Section 12 "Data Flow: Feed to Columns to Cadence" lists nine live feeds and Marketing Stream is not one of them. The Cross-Reference Ledger records external dependency D5 "Marketing Stream hourly feed" as OPEN. Therefore this SOP authorizes zero hourly modifiers, issues zero verdicts of its own, and exists to do four things: hold the deferral in a stated, dated, weekly-checked condition; stop dayparting-shaped actions from entering the account through side doors while the lever is dormant; convert the deferral into an execution build the day the feed lands; and retire the slot permanently if the evidence says the pattern does not exist.

1.1 Why this document is 11 pages and not 18, and why that is correct

The SOP Build Kit v1.0 sets 10-18 pages and the SOP-12 v2.0 exemplar floor is 12 procedures and 4 worked examples. This component cannot reach that floor honestly, and the reason is recorded here rather than solved with padding. The Cross-Reference Ledger row for #45 states the ownership in seven words: "The deferral and its trigger condition." Its must-not-restate cell states the rest: "Nothing operational until the trigger fires." v4.0 supplies exactly two numbers for the lever itself, the plus or minus 15% modifier bound and the 4-week evidence window, and no third. It does not state the bucket granularity, the step cadence, the re-read interval, the hold test, or the exit rule. An execution SOP written today would invent all five, which anti-drift rule 1 forbids and which the v4.1 amendment list in Section 13 records instead. The precedent is in the corpus: the Ledger registers component #21 "DSP and Sponsored TV deferrals" as "COMPLETE as a closed decision ... No procedure exists and none is owed."

What this document does carry at full depth is the part that is executable today and is currently unwritten anywhere: the ten standing procedures that keep a dormant lever from leaking into the account, the arithmetic that decides whether any hour-level claim is readable at all, ten failure modes drawn from the seams this component touches, and the activation gate. That content is real, it runs every Monday, and no other document in the 46-component set owns it.

1.2 Boundaries, stated by component number and one clause

Adjacent component	The seam, stated so it is owned
Criteria System v4.0	Decides every threshold, band, scenario, and verdict, including L10's plus or minus 15% bound and its 4-week window. This SOP executes the deferral; it never sizes an hourly modifier. A threshold appearing in this SOP instead of v4.0 is a defect.
#12 Exact campaign operations	Owns intraday funded-push pacing and the hourly truncation alarm on push campaigns. Hour-level reading of spend already exists in the account and it belongs there. This SOP consumes those alarms as a confound list and never duplicates the pacing procedure.
#28 Placement and bid-modifier operations	Owns the ceiling worksheet and every modifier operation. An hourly modifier is a fourth multiplicand in that worksheet, and the arithmetic is theirs. Referenced by number and clause only per Ledger conflict C8; the copy attached to this session is superseded.
#37 Automation platform rules	Owns the rule register, the weekly diff, and change control for every automation-platform rule. A time-conditioned rule is registered and reconciled there. This SOP supplies one input to that register: the exclusion reason. Referenced by number and clause only per Ledger conflict C8.
#20 AMC operations	Owns the AMC audience and query lifecycle and produces the hourly-pattern evidence named in Strategic Doctrine v1.0 Section 10. This SOP consumes that evidence as an intake artifact and never writes or runs a query.
#39 Data feed and manifest operations	Owns feed definitions, manifest field groups by ID, and the T7 layer. The D12 row's definition and its status transitions are theirs. This SOP reads the row and never edits it.

Adjacent component	The seam, stated so it is owned
#38 Bulk-file deployment mechanics	Owns staging and the T+1 verification diff. At activation, and only then, this component's output passes through there. Before activation this SOP stages nothing.
#21 DSP and Sponsored TV deferrals	The structural precedent for a deferral component held as a closed decision. Precedent only; this SOP borrows no content from it.
#1 Document governance, #2 A-to-Z System Map	Own component status vocabulary and routing. This SOP writes its state back and never re-declares its own status.
#44 Operator onboarding and certification	Consumes Element 11 of this SOP, due 2026-09-25. This SOP supplies ramp content and owns no certification test.
#4 and #40 Operating cadence	Own the calendar slots this SOP's weekly and quarterly procedures land on.

2. Trigger Conditions

Every trigger names the artifact that detects it and the maximum allowed lag from signal to procedure start. v4.0 contains no signal-to-start lag table, so the lag column is proposed and carried as amendment item AM-2 in Section 13. It is not a threshold quoted from v4.0 and must not be treated as one.

Trigger	Detecting artifact	Max lag (AM-2)	Procedure
Weekly deferral-state check	T7 MANIFEST row "D12 Hourly dayparting", column Status, read after the Monday workbook refresh completes at 09:00 per Data Architecture v1.0 Section 12 assembly cadence	Same day	P1
D12 status flip	T7 MANIFEST row "D12 Hourly dayparting" reads anything other than MISSING	1 business day	P8, then P1 suspends
Dayparting-shaped proposal	Any staged change reason string, any T6 DECISION LOG entry, or any verbal proposal that conditions a bid, budget, or modifier on time of day	Same day	P3, using P2
Time-conditioned automation rule found live	The weekly automation-platform rule diff owned by #37	Same day	P4
v4.0 increment	A new Criteria System export whose Section 9 L10 text or Section 5 G2 text differs from the version last read	5 business days	P5
AMC hourly-pattern evidence returned	The #20 result set answering the fifth query of Strategic Doctrine v1.0 Section 10	5 business days	P6
Quarterly deferral review	The calendar slot held by #40	On date	P7
Retirement evidence accepted	P7 record carrying a PPC Lead decision to retire the slot	5 business days	P9, then P10

3. Preconditions and Gates

Each gate below cites its manifest ID or its v4.0 gate ID and states the observable that proves it passed. A precondition that cannot be observed is treated as failed, not assumed.

Gate or manifest ID	Requirement	Observable that proves it passed
D12 Hourly dayparting (T7 MANIFEST)	The row exists and its Status is read this cycle. While Status is MISSING, every procedure in this SOP except P8 is legal and P8 is barred.	The Status cell of the D12 row, read after the 09:00 Monday refresh, and its date recorded in the P1 entry.
G8 Packet completeness (v4.0 Section 5 "Layer 0: Gates (G1-G12). Nothing Proceeds Until These Pass")	The manifest reports present, stale, or missing before analysis. A verdict whose gate fields are missing is downgraded to a flag; packets under 85% of gate-relevant fields produce flags only.	The T7 PACKET COMPLETENESS SCORE cell and the D12 row's own Status.

Gate or manifest ID	Requirement	Observable that proves it passed
G2 Sample sufficiency (v4.0 Section 5)	Bid reads need 15 clicks minimum over 7 days; CVR verdicts 100 clicks minimum; CTR verdicts 1,000 impressions minimum. Below threshold the verdict is EXTEND TEST, never a rate judgment.	The P2 computation output: clicks per hour bucket against the 15-click and 100-click minimums, shown per bucket, not per campaign.
G3 Budget truncation (v4.0 Section 5)	In-budget under 70% of the day marks every rate on that campaign TRUNCATED and the budget lever is evaluated before any bid conclusion is allowed.	T7 MANIFEST row A3 "Pacing buckets / in-budget %" and the campaign's in-budget percentage. Note that A3 currently reads PARTIAL with the day-of-week curve pending, which caps how far any hourly claim can be read even after truncation clears.
G9 Price stability (v4.0 Section 5)	A price change on the targeted variation within trailing 7 days marks CVR and ACOS reads PROVISIONAL.	T5 PRODUCT LOG, Event Type column, entries of type price move inside the read window.
G10 Event exclusion (v4.0 Section 5)	Event weeks and the 14 days following are excluded from trend verdicts.	T5 PRODUCT LOG Trend-Exclusion Flag column across the candidate evidence window.
G1 Validation quarantine (v4.0 Section 5)	Impossible metrics quarantine the row, produce a flag, and open a same-day data-refresh task.	Zero-flag columns computed in the workbook on any hour-level extract offered as evidence.
A6 Objective + per-term Expected CPC (T7 MANIFEST)	Reads PARTIAL, with Expected CPC held at the account reference until per-term values land.	The A6 Status cell. Any hourly claim priced against a per-term Expected CPC inherits this partial state and says so.

4. Step-Level Procedures

P1. Weekly deferral-state check

5 minutes

1. Monday, after Data Operations lands the Raw Pull Pack at 06:00 and the Power Query refresh completes the product workbooks by 09:00 (Data Architecture v1.0 Section 12, assembly cadence). Open the product workbook, tab T7 MANIFEST.
2. Locate the row whose Catalog ID + Field Group cell reads "D12 Hourly dayparting". Read four cells across: Source Feed, Refresh, Status, and Verdict Classes Blocked While Missing.
3. Expected reading while the deferral holds: Source Feed "Marketing Stream", Refresh "Deferred", Status "MISSING". Any other Status value stops this procedure and starts P8 within one business day.
4. Record one line in T6 DECISION LOG. Date = the cycle date. Level = product (see amendment item AM-7: the Level closed list carries no component-level value, so product is used as the nearest legal value and the target cell names the component). Target = "#45 deferral state". Rule / Scenario Fired = "v4.0 Section 9 L10". Change (prior to new) = "DEFERRED to DEFERRED" or the observed transition. Dollar Impact / wk = 0.00. Review Date = the next quarterly slot. Owner / Tier = PPC Lead.
5. Where the same product workbook set covers five brands, the check runs once per product workbook and the five results reconcile: a Status divergence across products is F2 the same day, because one feed cannot be present for one product and absent for another.

P2. Hour-bucket sufficiency computation

20 minutes

This is the instrument every hour-level claim passes through, whether it arrives from an operator, a platform recommendation, or an AMC result set. It answers one question with arithmetic instead of judgment: can an hour bucket on this campaign carry a read at all under v4.0 Section 5 G2.

1. Take the campaign's trailing 7-day clicks from T3 CAMPAIGN BOARD, or the term's clicks from T2 KEYWORD COMMAND where the claim is term-level. Use the same window the claim used; if the claim did not state its window, the claim fails here and routes to P3 branch D.
2. State the bucket count the claim assumes. Two are possible and v4.0 does not choose between them (amendment item AM-3): 24 buckets pooling each clock hour across all seven days, or 168 buckets holding day-of-week by hour separately. Compute both. The two answers differ by a factor of 7 and a claim that has not stated its bucketing has not stated its evidence.
3. Clicks per bucket per week = trailing 7-day clicks divided by bucket count.
4. Weeks to reach a bid read = 15 divided by clicks per bucket per week. Weeks to reach a CVR verdict = 100 divided by clicks per bucket per week. Both minimums are v4.0 Section 5 G2.
5. Compare against the L10 window. v4.0 Section 9 requires 4 weeks minimum of hour-level evidence per campaign. Where weeks-to-reach exceeds 4, the L10 window is necessary but not sufficient for this campaign and the claim cannot clear G2 inside it.
6. Compute the inverse for the record, because it is the number that tells the account which campaigns could ever be dayparted: required weekly clicks = minimum clicks multiplied by bucket count, divided by 4.

7. Write the computation into the triage record. It is six lines of arithmetic and it is the whole of the evidence question; a claim that cannot show these six lines has not been read, it has been believed.

P3. Dayparting-shaped proposal triage

25 minutes

A dayparting-shaped proposal is any proposal to condition a bid, a budget, a modifier, a campaign state, or an automation rule on time of day, whatever vocabulary it arrives in. The branches below are exhaustive; every proposal lands on exactly one.

Branch	Condition	Action
A	The proposal cites hour-level data and the campaign is TRUNCATED per G3 (in-budget under 70% of the day)	Stop. The hours carrying no spend carry no evidence. Route to the budget lever, which v4.0 Section 9 L1 places ahead of bid in the precedence order. Quote "FIX BUDGET (TRUNCATED)" from the KDR v3.0 Bands and Verdicts Reference closed vocabulary. Re-read only after a full validated window at in-budget 70% minimum. This is Worked Example B.
B	The proposal cites hour-level data, the campaign is not truncated, and P2 shows the bucket below the G2 minimum	Stop. Quote "EXTEND TEST" from v4.0 Section 13.1 "The Closed Verdict Vocabulary". Record the P2 output with the required weekly clicks figure so the same claim does not return next month unchanged. This is Worked Example A.
C	The proposal cites hour-level data, the campaign is not truncated, and P2 shows the bucket at or above the G2 minimum	Stop anyway, and state why in the record: the D12 row reads MISSING, so the data source is not the Marketing Stream curve L10 names. Sufficiency does not substitute for provenance. Route the observation to P6 as a candidate input for the AMC evidence file.
D	The proposal cites no window, no bucket count, or no click counts	Return it unread with the P2 template attached. A proposal without its window is not a thin claim, it is an unstated one.
E	The proposal is a platform-side recommendation or a vendor default carrying an hourly schedule	Treat as branch C plus F1: check whether it is already live as a rule, and if so route to P4 the same day.

1. Log the triage in T6 DECISION LOG with the branch letter in the Rule / Scenario Fired cell alongside the v4.0 citation, so the quarterly review can count how often each branch fires and which one is teaching the account something.
2. Where the proposal came from a role rather than a report, the response goes back to that role in writing with the six arithmetic lines. The arithmetic is the answer; the refusal alone teaches nothing and the proposal returns.

P4. Time-conditioned automation rule check

15 minutes

1. Read the weekly automation-platform rule diff produced by #37. This SOP does not run the diff and does not hold the register.
2. Scan for any rule whose trigger, schedule, or action window is conditioned on clock time rather than on a metric state. A bid rule that fires on a click count is not time-conditioned. A bid rule that fires between stated hours is.
3. For each one found, issue the exclusion reason to #37 in one sentence citing v4.0 Section 9 L10 and the D12 row status. The register entry, the suspension, and the reconciliation state are #37's to write.
4. Where the rule is already live and unregistered, that is #37's register failure and this SOP raises it there rather than acting on the rule directly. This SOP never edits an automation rule.
5. Record the count found and the count excluded in the P1 weekly line. A count above zero for two consecutive quarters is a finding for P7: the account keeps proposing the lever, which is evidence about the deferral itself.

P5. v4.0 increment re-read of L10

20 minutes

1. On any new Criteria System export, open Section 9 "Layer 4: The Lever Library (L1-L10). Order, Limits, Mechanics" and read the L10 row in full.
2. Diff four things against the version last read: the modifier bound, currently plus or minus 15%; the evidence window, currently 4 weeks minimum per campaign; the named feed, currently Marketing Stream; the named catalog ID, currently D12.
3. Read Section 5 G2 in the same pass, because the sufficiency minimums are what P2 computes against and a change there changes every stored required-weekly-clicks figure.

4. Read Section 17.1, adoption phase 4, which names "Marketing Stream dayparting (L10) when the feed lands" as a phase-4 deliverable. A change to the phase assignment changes this component's expected activation horizon and is reported to #1.
5. Where the increment adds any of the five missing rules listed as amendment items AM-3 to AM-6, this SOP is superseded rather than amended: the activation build starts. Where it does not, record the re-read with the date and the confirmed text.

P6. AMC hourly-pattern evidence intake

45 minutes

1. Trigger: #20 returns the result set answering the fifth of the five queries named in Strategic Doctrine v1.0 Section 10, which the Doctrine defines as "hourly-pattern evidence that either justifies or permanently retires the deferred dayparting slot."
2. Receive it as an intake artifact. This SOP does not write, run, or tune the query; the audience and query lifecycle belongs to #20 by its Ledger row.
3. Apply the gates in this order before reading a single hour: G1 on impossible values, G10 on event weeks and the 14 days following, G9 on any price move inside the window from T5 PRODUCT LOG, and G3 on every campaign inside the result set. An hourly curve assembled across a truncated campaign, an event window, or a price test is an artifact of those three things.
4. Run P2 against the surviving campaigns. Record which campaigns, if any, carry enough weekly clicks that a bucket could clear G2 inside the L10 4-week window. This list is the real output: it names the subset of the account for which dayparting could ever be a decision rather than a decoration.
5. Write the intake note with three findings stated as numbers, not adjectives: campaigns surviving the gates, campaigns clearing P2, and the spread between the best and worst bucket where both cleared. Attach it to the quarterly review file.
6. The intake note does not authorize a modifier. v4.0 Section 9 L10 names the Marketing Stream curve and the D12 row as the deployment condition, and AMC evidence is neither. It authorizes exactly two things: a better-informed P7, and a retirement case under P9.

P7. Quarterly deferral review

45 minutes

1. On the #40 calendar slot. Assemble four inputs: the trailing 13 P1 weekly lines, the P3 triage counts by branch, the P4 exclusion counts, and any P6 intake note.
2. Confirm the D12 row status has not changed unobserved: read it directly rather than trusting the last weekly line.
3. Answer three questions in writing. Has the feed status moved. Has the account tried to reach the lever by proxy, and how often. Has evidence arrived that would retire the slot.
4. Issue one of three states to the PPC Lead: hold the deferral, arm for activation, or retire. Hold is the default and requires no case. Arm requires the D12 status change. Retire requires a P6 intake note and moves to P9.
5. Write the review record and route the state to #1 for the register line and to #2 for the routing map. A quarter with no review record is F6.

P8. Activation gate

3 hours, then a build session

This procedure runs once, on the day the D12 row stops reading MISSING. It does not deploy anything. It establishes whether an execution SOP can now be written, and it is deliberately a gate rather than a launch.

1. Confirm the status change at source with #39, who own the T7 layer and the D12 row definition. A status change written by anyone else is F2.
2. Confirm feed coverage per campaign. v4.0 Section 9 L10 requires the 4-week evidence window per campaign, not per account. Record which campaigns have coverage from which date. Partial coverage arms partial scope, not full scope, and that is F10 if it is treated otherwise.
3. Start the 4-week clock and record its end date. No dayparting decision is legal before that date on any campaign, whatever the curve appears to show in week one.
4. Run P2 across every covered campaign at the end of week 4 and produce the eligible list. Campaigns whose buckets cannot clear G2 inside the window are out of scope permanently until their click volume changes, and the required weekly clicks figure is recorded against each so the boundary is a number rather than an opinion.
5. Check the five open rules. v4.0 must, by then, state the bucket granularity, the modifier step cadence, the re-read interval, the hold test, and the exit rule (amendment items AM-3 to AM-6). Where any is still absent, the gate closes with a written statement of which rule is missing, and the account holds. A modifier bound without a step cadence is not a lever, it is a range.
6. Where the five rules are present, hand the eligible list, the coverage dates, and the open-rule confirmation to the build session that writes the execution SOP, and hand the modifier-into-ceiling question to #28 as its worksheet input. This SOP is superseded on that document's release, not amended.

P9. Retirement path

30 minutes

1. Trigger: a P7 review issuing retire, on the strength of a P6 intake note. Strategic Doctrine v1.0 Section 10 names permanent retirement of the slot as one of two legitimate outcomes of the AMC evidence, the other being justification.
2. State the retirement case in numbers taken from the intake note: campaigns surviving the gates, campaigns clearing P2, and the observed spread between buckets. v4.0 states no evidence bar for retirement, which is amendment item AM-6, so the case is presented to the decision holder rather than resolved by a threshold.
3. The decision sits with the PPC Lead per the Authority Matrix rows named in Element 9. Retirement is a v4.0 change in effect, so the CEO row on v4.0 threshold changes applies to the L10 deletion itself, not to this SOP's recommendation.
4. On acceptance, route three artifacts: the retirement memo to #1 for the register status, the L10 deletion request to the v4.1 amendment list, and the D12 row disposition to #39.
5. Retirement closes this component the way #21 is closed: a dated decision with its evidence, no procedure owed. This SOP is then archived rather than maintained.

P10. State handoff at activation or retirement

30 minutes

1. Write the component state to #1 with the date and the deciding artifact. The register status vocabulary is #1's; this SOP supplies the evidence line, never the status word.
2. Return the updated Ledger row for #45 in copy-paste form to the Ledger close-out. Per the pack's standing rule, a session returns its row content and its findings; the close-out assigns the version number and any conflict identifier.
3. Hand Element 11 to #44 for the certification path, due 2026-09-25.
4. Release the #40 quarterly slot on retirement, or convert it to the execution SOP's review cadence on activation. A calendar slot left pointing at an archived document is how a review becomes theatre.

5. Verdict Mapping

This SOP issues zero verdicts of its own. v4.0 Section 13.1 "The Closed Verdict Vocabulary" contains eighteen actions and none of them is a dayparting action. The KDR v3.0 "CLOSED VERDICT VOCABULARY (Framework Section 17)" sheet contains twenty-five strings and none of them is a dayparting action. That absence is correct and it is the point: a dormant lever has no verdict because it authorizes no move. The table below maps the four situations this SOP meets to the strings it may quote, each with its source, and states the one gap that becomes an amendment item rather than a local string.

Verdict string, quoted verbatim	Source	When this SOP quotes it
EXTEND TEST	v4.0 Section 13.1 "The Closed Verdict Vocabulary" (primary scenarios G2, D9)	P3 branch B: an hour-level claim whose bucket sits below the G2 minimum. The definition in v4.0 reads "Evidence insufficient or confidence LOW; capped continuation to the sufficiency threshold", which is exactly the state of every hour bucket in the account today.
FIX BUDGET (TRUNCATED)	Keyword Decision Record Template v3.0, Bands and Verdicts Reference sheet, closed verdict vocabulary	P3 branch A: the campaign is truncated and the apparent hourly pattern is the budget running out. v4.0 Section 5 G3 supplies the 70% line and Section 9 L1 supplies the precedence.
FLAG: DATA REPAIR	v4.0 Section 13.1 "The Closed Verdict Vocabulary" (primary scenarios G1, G8, P8)	F2: the D12 row reads a status the feed table does not support, or an hourly extract fails G1 validation. v4.0 defines it as "Quarantine/manifest failure; no recommendation; repair task with owner and date."
QUARANTINE: DATA REFRESH	Keyword Decision Record Template v3.0, Bands and Verdicts Reference sheet, closed verdict vocabulary	The KDR-side equivalent where the finding lands on a keyword record rather than on the manifest.

The gap, stated rather than filled. The deferral states this SOP writes weekly, held, armed, and retired, are not verdicts in either closed vocabulary. v4.0 Section 13.1 states that "New verdicts enter only by V3 amendment," so these strings are carried as amendment item AM-1 for the v4.1 release and are recorded in T6 as state text rather than issued as verdicts. This follows the route the Ledger records for #43 and #46, which each logged a state vocabulary for admission rather than minting one locally.

Zero local thresholds. Two numbers govern this component and both are v4.0's: the plus or minus 15% modifier bound and the 4-week per-campaign evidence window, both from Section 9 L10. Three more are quoted where they gate a read: 15 clicks, 100 clicks, and 1,000 impressions from Section 5 G2, and 70% in-budget from Section 5 G3. Every other numeral in this document is either an arithmetic result computed on the page from those figures, a date, or a count of records.

6. Worked Examples

Both examples are illustrative. The anonymized live product economics spine was not attached to this session, so both use the standard fictional base from the PPC Economics Worked-Example Set v1.0: AOV $80.00, contribution $28.00 per order, break-even ACOS 35.0%, Target ACOS 17.5%, Maximum Acceptable 26.3%, CVR 5.2%. From that base, Max Profitable CPC = 28.00 multiplied by 0.052 = $1.46 and Target CPC = 80.00 multiplied by 0.175 multiplied by 0.052 = $0.73, both per the Set's case E1.

6.1 Worked Example A: the clean case. A proposal arrives and the arithmetic answers it.

An operator reads a seven-day hourly export and proposes cutting bids 15% between 00:00 and 06:00 on one single-keyword oversight exact campaign, on the observation that those hours "convert badly." The campaign took 40 clicks in the trailing 7 days and is not truncated: in-budget reads 91% of the day, clear of the G3 70% line. P3 routes to branch B by way of P2.

P2 computation, both bucketings, because v4.0 does not choose between them (AM-3).

- Trailing 7-day clicks: 40.
- 24-bucket pooling: 40 / 24 = 1.6667 clicks per bucket per week.
- 168-bucket day-of-week by hour: 40 / 168 = 0.2381 clicks per bucket per week.
- Weeks to a G2 bid read (15 clicks): 15 / 1.6667 = 9.0 weeks at 24 buckets; 15 / 0.2381 = 63.0 weeks at 168 buckets.
- Weeks to a G2 CVR verdict (100 clicks): 100 / 1.6667 = 60.0 weeks at 24 buckets; 100 / 0.2381 = 420.0 weeks at 168 buckets.
- What the L10 4-week window actually holds: 1.6667 x 4 = 6.67 clicks per bucket at 24 buckets; 0.2381 x 4 = 0.95 clicks per bucket at 168 buckets.
- Required weekly clicks to clear 15 clicks per bucket inside 4 weeks: 15 x 24 / 4 = 90 clicks/week at 24 buckets; 15 x 168 / 4 = 630 clicks/week at 168 buckets.
- Required weekly clicks to clear 100 clicks per bucket inside 4 weeks: 100 x 24 / 4 = 600 clicks/week at 24 buckets; 100 x 168 / 4 = 4,200 clicks/week at 168 buckets.
- What that costs at the training base: 630 clicks/week at Target CPC $0.73 = $459.90/week; 4,200 clicks/week = $3,066.00/week. At Max Profitable CPC $1.46 the same two figures are $919.80/week and $6,132.00/week.

The read. The proposal rests on 0.95 clicks per bucket at the bucketing that separates Saturday 02:00 from Tuesday 02:00, and 6.67 clicks per bucket at the bucketing that does not. Both sit under the G2 bid-read minimum of 15 clicks, and both sit far under the 100-click CVR minimum that the phrase "convert badly" actually requires. The verdict quoted is "EXTEND TEST" from v4.0 Section 13.1. No modifier is proposed, no rule is written, and the campaign's bids do not move on this basis.

What the record carries forward. The required weekly clicks figures are the durable output. This campaign would need 630 clicks per week to make a day-of-week hour bucket readable for a bid decision inside the L10 window, against 40 today, a factor of 15.75. The account now has a number for the question "which campaigns could ever be dayparted" rather than an argument, and P6 will run this same computation across the AMC result set to produce the eligible list. Note also that the answer changes by a factor of 7 depending on the bucketing, which is precisely why AM-3 is on the amendment list: two operators can reach opposite conclusions from the same export while both citing L10 correctly.

6.2 Worked Example B: the edge case. The pattern is real, the cause is not the hour.

A second campaign shows ten consecutive hours, 14:00 to 23:59, with zero impressions, zero clicks, and zero orders, seven days running. The pattern is stronger and cleaner than anything in Example A, and the proposal is to apply the full minus 15% L10 bound to those ten hours. P3 routes to branch A before P2 runs, because the gate order in Element 3 puts G3 ahead of the sufficiency question.

The observed day.

- Daily budget $45.00. Spend exhausts at 14:00, so in-budget = 14 / 24 = 58.3% of the day, under the G3 70% line: every rate on this campaign reads TRUNCATED.
- Burn rate = 45.00 / 14 = $3.2143 per hour.
- CPC $1.00, so clicks = 45.00 / 1.00 = 45.0 per day, all of them bought before 14:00.
- Orders at CVR 5.2% = 45.0 x 0.052 = 2.34. Sales = 2.34 x $80.00 = $187.20. ACOS = 45.00 / 187.20 = 24.0%, inside the 26.3% Maximum Acceptable line.
- Contribution = 2.34 x $28.00 = $65.52. Net contribution after ad spend = 65.52 - 45.00 = $20.52 per day.

The budget the G3 line implies.

- 70% of the day is 0.70 x 24 = 16.8 hours. At the observed burn rate: 16.8 x 3.2143 = $54.00 per day.

The corrected day, if the burn rate holds.

- Clicks = 54.00 / 1.00 = 54.0. Orders = 54.0 x 0.052 = 2.808. Sales = 2.808 x $80.00 = $224.64. ACOS = 54.00 / 224.64 = 24.0%, unchanged, because the constraint was budget and not efficiency.
- Contribution = 2.808 x $28.00 = $78.62. Net = 78.62 - 54.00 = $24.62 per day.
- Delta = 24.62 - 20.52 = $4.10 per day, $28.73 per week, on one campaign.

What goes wrong and how it is handled. The proposed minus 15% would have been applied to ten hours holding zero observations, on a campaign whose evening silence is an accounting fact about a $45.00 budget rather than a behavioural fact about shoppers. Worse, it would have been applied in the wrong lever order: v4.0 Section 9 L1 places budget ahead of bid and states that the sequence is never skipped forward, and G3 states plainly that "a starved campaign's ACOS and CTR describe its worst hours, not its potential." The verdict quoted is "FIX BUDGET (TRUNCATED)" from the KDR v3.0 Bands and Verdicts Reference. The budget lever runs, and the hour-level question is re-asked only after a full validated window at in-budget 70% minimum.

The second thing that goes wrong, and why the re-read still does not authorize anything. After the budget correction, an operator may reasonably expect the hourly curve to now be readable. It is not. The T7 MANIFEST row A3 "Pacing buckets / in-budget %" reads PARTIAL with the note that the day-of-week curve is pending, so the account cannot yet distinguish a genuine hourly shape from a day-of-week shape. And the D12 row still reads MISSING, so the provenance condition in v4.0 Section 9 L10 remains unmet regardless of how clean the sample becomes. P3 branch C applies: the observation routes into the P6 evidence file and the deferral holds. The $28.73 per week the budget lever recovered is the real yield of this triage, and it was found by refusing the dayparting frame rather than by accepting it.

7. Failure Modes and Escalation

Failure	Detection signal	Response
F1. Time-conditioned automation rule live	The #37 weekly rule diff returns a rule whose trigger or action window is conditioned on clock time rather than on a metric state	Issue the exclusion reason to #37 the same day citing v4.0 Section 9 L10 and the D12 status; #37 suspends and registers. This SOP never edits the rule. Motivating history: the account's automation register was seeded from a live platform inventory, which means rules predating the register exist by construction and a rule class nobody authorized is the normal case, not the exception.
F2. D12 marked present before the feed lands	The T7 D12 Status cell reads anything but MISSING while Data Architecture v1.0 Section 12's feed table still lists no Marketing Stream source, or the five product workbooks disagree on the D12 status	Quote "FLAG: DATA REPAIR" per v4.0 Section 13.1, open the repair task with owner and date per G8, and confirm at source with #39 before P8 runs a single step. A status cell is not a feed.
F3. Truncation read as an hourly pattern	An hour-level claim on a campaign whose in-budget reads under 70%	P3 branch A. Budget lever first per L1 precedence, then a full validated window before the hour question returns. This is the most likely failure in the account by a wide margin, because every truncated campaign in the set produces a clean, convincing, and entirely false evening pattern.
F4. AMC evidence used as deployment authority	A staged change reason string citing an AMC query as the basis for a bid, budget, or modifier conditioned on time of day	Blocked at staging by #38. AMC evidence authorizes a better P7 and a possible retirement case; it does not satisfy the provenance condition in L10, which names the Marketing Stream curve and D12 specifically.
F5. Bucket claim below sufficiency	P2 returns clicks per bucket under the G2 minimum for the read being claimed	Quote "EXTEND TEST" per v4.0 Section 13.1 and record the required weekly clicks figure. Without that figure the same claim returns next quarter with the same 40 clicks behind it.
F6. The deferral rots	No P7 review record in T6 across a full quarter, or four or more consecutive P1 lines missing	Escalate to the PPC Lead. A deferral nobody checks becomes an omission nobody owns, and the difference between the two is only visible in the log. The #40 calendar slot is the structural fix and its absence is itself the finding.
F7. Dayparting by proxy	T6 DECISION LOG shows a repeating clock-pattern in changes to T3 CAMPAIGN BOARD Daily Budget $, State, or Placement Mult TOS %, with no verdict behind them	Normalize and route to the owning SOP by number: budget and state changes to the campaign-type SOP, modifier changes to #28. Then record it here as a proxy attempt, because the count of proxy attempts is evidence for P7 about how much the account wants this lever.
F8. Activation without the window	P8 reaches step 4 with a coverage start date under 4 weeks old	Hold. The 4-week window is v4.0 Section 9 L10's and is not negotiable by enthusiasm on day three of a new feed. Record the end date and return on it.
F9. Contaminated first curve	The first Marketing Stream window overlaps an event week or the 14 days following (G10), a price move inside 7 days (G9), or a G1 quarantine	Exclude the affected days and extend the window rather than shortening the standard. A first curve built through an event is the account's own event, read back as consumer behaviour.
F10. Partial coverage treated as full	P8 step 2 shows the feed covering some campaigns and not others, and the eligible list is drawn account-wide	Scope to covered campaigns only. L10 states the evidence requirement per campaign; an account-level average of hourly curves is a shape no single campaign has.

Escalation. F1 and F7 escalate on discovery, because both are live spend conditioned on a lever the system has not authorized and v4.0's closing standard is that when spend cannot cite the rule that authorizes it, the spend is what changes. F2 escalates same day, because a false-present manifest row would let P8 open the activation gate on a feed that does not exist. F6 escalates at the quarter boundary to the PPC Lead. Every other failure is recorded and cleared inside the weekly cycle.

8. Outputs and Artifacts

Named record	Storage location	Retention	Downstream consumer
Deferral State Record (weekly)	T6 DECISION LOG, one line per product workbook per cycle	Life of the workbook; policy owed by #39 (AM-8)	P7 quarterly review; #1 register status
Hour-Bucket Sufficiency Sheet	Attached to the triage record that requested it	Until the campaign's click volume changes materially	P3 triage; P6 eligible list; P8 step 4
Proposal Triage Record	T6 DECISION LOG with the branch letter in the Rule / Scenario Fired cell	Life of the workbook (AM-8)	P7 branch counts; the proposing role, in writing
Schedule-Rule Exclusion Note	Issued to the #37 rule register; copy in the P1 weekly line	Held by #37 with the register entry	#37 register and reconciliation state
AMC Evidence Intake Note	Quarterly review file, attached to the #20 result set	Until superseded by a later intake	P7; P9 retirement case
Quarterly Deferral Review Record	T6 DECISION LOG plus the review file	Life of the component (AM-8)	#1 register line; #2 routing map; PPC Lead
Activation Packet (eligible list, coverage dates, open-rule confirmation)	Handed to the activation build session	Permanent; it is the execution SOP's input record	The execution SOP build; #28 ceiling worksheet input
Retirement Memo	Handed to #1 with the L10 deletion request	Permanent	#1 register; v4.1 amendment list; #39 for the D12 row disposition
Returned Ledger row for #45	Ledger close-out	Permanent, versioned by the close-out	Every subsequent build session citing this boundary

9. Skill-Conversion Block

Block	Specification
Input bindings	T7 MANIFEST row "D12 Hourly dayparting", columns Source Feed, Refresh, Status, Verdict Classes Blocked While Missing. T7 MANIFEST row A3 "Pacing buckets / in-budget %", column Status. T3 CAMPAIGN BOARD columns Daily Budget $, State, Bidding Strategy, Placement Mult TOS %, and trailing 7-day clicks. T2 KEYWORD COMMAND trailing 7-day clicks for term-level claims. T5 PRODUCT LOG columns Event Type and Trend-Exclusion Flag for the G9 and G10 checks. T6 DECISION LOG for the state, triage, and review lines. The #37 weekly rule diff. The #20 AMC result set for query 5 of Strategic Doctrine v1.0 Section 10. v4.0 Section 9 L10 text, Section 5 G2 and G3 text.
Deterministic logic	The D12 status read and its expected-value comparison. The P2 computation in full: clicks divided by bucket count at both 24 and 168; weeks-to-minimum against 15 and 100; the 4-week window comparison; the required-weekly-clicks inverse. The P3 branch selection, which is a strict ordering of G3 first, then P2, then provenance, with no discretionary branch. The F7 clock-pattern scan across T6. The G9, G10, and G1 exclusion filters on any candidate evidence window. None of this requires judgment and all of it is arithmetic or cell comparison.

Block	Specification
Human checkpoints, from the Authority Matrix v2.0 Section 1 rows that apply	AI-executable with audit: the P1 weekly read and its T6 line; the P2 computation; the P3 branch selection and its record; the F7 scan; the P6 gate filtering and eligible-list arithmetic. Human-mandatory, PPC Lead decides: the P7 quarterly state call, per the matrix row "Automation rules (#37)" where the PPC Lead holds D and the incoming operator holds E and R, and per "Structure creation (all SOPs)" where the PPC Lead holds D. Human-mandatory, PPC Lead decides with the AMC cycle: P6 acceptance of an intake note, per the matrix row "AMC cycle (#20)" where the PPC Lead holds D and the incoming operator holds E. Human-mandatory, CEO decides: the L10 deletion request in P9 and any activation that changes a v4.0 threshold, per the matrix row "v4.0 threshold changes" where the CEO holds D and the PPC Lead holds E and R. The vacancy rule applies unchanged: an unfilled role's D reverts to the PPC Lead and its E to the PPC Manager.
Cadence	Weekly: P1 on Monday after the 09:00 refresh, P4 on receipt of the #37 diff. On event: P3 same day, P5 within 5 business days of a v4.0 increment, P6 within 5 business days of an AMC return, P8 within 1 business day of a D12 status change. Quarterly: P7 on the #40 slot. Once, terminally: P9 and P10.
Outputs	The nine artifacts named in Element 8. Alarms F1 to F10 with the evidence line attached to each. Zero deployed changes: this component stages nothing through #38 before activation.

10. Version Governance

Owner: PPC Lead. This SOP increments on five component-specific triggers and on nothing generic.

1. Any v4.0 increment that touches Section 9 L10. A change to the plus or minus 15% bound, the 4-week window, the named feed, or the named catalog ID changes this document's subject matter directly. P5 runs and the increment is recorded here.
2. Any v4.0 increment that touches Section 5 G2 or G3. The sufficiency minimums are what P2 computes against and the 70% truncation line is what P3 branch A turns on. A change to either invalidates every stored required-weekly-clicks figure and forces a recomputation across the eligible list.
3. Any change to the T7 D12 row. A status, source feed, or refresh change is the activation trigger and is the one event that can supersede rather than amend this document.
4. Any AMC evidence return under #20 query 5. The intake changes what P7 is deciding against and, on a retirement finding, ends the component.
5. Any platform change introducing native hour-level controls. A new control surface would change what a dayparting decision even deploys into, and the ceiling arithmetic that #28 owns would change with it.

Quarterly review regardless, on the #40 slot, as part of P7. Registered in the Supersession Map. On activation this SOP is superseded in full by the execution build rather than amended, because a deferral control and an execution procedure are different documents with different readers. On retirement it is archived with its retirement memo and the component closes the way #21 is closed.

11. Operator Ramp

Built from the Authority Matrix v2.0 Section 1 rows that name this component's decisions: "Automation rules (#37)" (PPC Lead D, incoming operator E and R, PPC Manager E), "AMC cycle (#20)" (PPC Lead D, incoming operator E), "v4.0 threshold changes" (CEO D, PPC Lead E and R, incoming operator I), and "Structure creation (all SOPs)" (PPC Lead D, incoming operator E). The checkpoint holder throughout is the PPC Lead.

Week 1: observation and artifact familiarity

1. Locate the T7 MANIFEST D12 row unaided in all five product workbooks and state its four governing cells from memory of where they sit, not from memory of their values.
2. Read v4.0 Section 9 L10 and Section 5 G2 and G3 in the source file and point to the exact text for each of the five numbers this component uses.
3. Observe one P1 weekly check and one #37 rule diff read end to end. Produce no records.
4. Reproduce the P2 computation cold on a campaign of the checkpoint holder's choosing, both bucketings, all six lines, and get the required-weekly-clicks inverse right without prompting.

Week 2: supervised execution, checkpoint holder reviewing every record

1. Run P1 across all five product workbooks and write the T6 lines. The PPC Lead reviews each before it is committed, per the "Automation rules (#37)" row where the operator holds E and R under the Lead's D.
2. Run P4 against the live #37 diff and draft any exclusion reasons. The Lead reviews the reason text, because a reason string that does not cite L10 and the D12 status is not an exclusion, it is an opinion.

3. Triage two P3 proposals, at least one of which the checkpoint holder has seeded as a branch A truncation case, and defend the branch selection out loud.
4. Draft the quarterly review inputs from the trailing P1 lines without issuing a state call. The state call is the Lead's D and stays there.

Week 4: unsupervised execution, checkpoint holder auditing

1. Run the full weekly cycle unaided: P1, P4, and any P3 that arrives. The Lead audits the T6 lines after the fact rather than before.
2. Handle one live proposal end to end, including the written response carrying the six arithmetic lines back to the proposing role.
3. Assemble a P6 intake dry run against a synthetic result set, applying G1, G9, G10, and G3 in order and producing the eligible list. Acceptance of a real intake remains the Lead's D per the "AMC cycle (#20)" row.
4. Walk the P8 activation gate on paper, naming which of the five open rules would block it today and why.

Live-defense questions

These test the reasoning behind the thresholds, not recall of their values. An operator who can recite plus or minus 15% and 4 weeks has learned nothing this component needs.

1. A campaign shows a clean, seven-day-consistent collapse in evening performance. Give me the two conditions under which that pattern is real evidence, and the one that is by far the most likely explanation in this account. Then tell me which gate catches it and why that gate sits ahead of the sufficiency question rather than behind it.
2. Two operators read the same hourly export and reach opposite conclusions, both citing L10 correctly. Explain how that is possible, and tell me what has to change in v4.1 before it stops being possible.
3. An AMC result set arrives showing a strong and statistically clean hourly pattern across nine campaigns. Explain why that does not authorize a single modifier, and name the two things it does authorize.
4. The Marketing Stream feed lands on a Tuesday and the curve looks decisive by Friday. Give me the argument for acting, then dismantle it using the lever's own text.
5. Why is the required-weekly-clicks figure the durable output of P2 rather than the weeks-to-minimum figure, and what does it tell the account about which campaigns will never be dayparting candidates at all?
6. This component authorizes no spend and issues no verdicts. Make the case that it should exist as a document at all, in numbers, using the branch counts P3 produces.

12. Interface Contracts

Built from the Cross-Reference Ledger's Interfaces cell for row #45, "Triggered by the Marketing Stream hourly feed. Consumer of #20 AMC evidence," cross-checked against Master PPC Data Architecture v1.0 for the artifact names and tab locations. Every row states what happens when the handoff does not arrive.

ID	Sending owner	Receiving owner	Artifact	Timing	Consequence of non-arrival
I-1	#39 Data feed and manifest operations	#45	T7 MANIFEST row "D12 Hourly dayparting", Status cell	Weekly, readable by Monday 09:00 after the Power Query refresh	P1 runs against the last known status and the T6 line is marked stale with its age in cycles. Two consecutive stale cycles escalate to the PPC Lead, because a deferral held against an unread manifest is an assumption wearing a record's clothes.
I-2	#20 AMC operations	#45	Result set for query 5 of Strategic Doctrine v1.0 Section 10, hourly-pattern evidence	On the AMC cycle date set by #20	P6 does not run. P7 records evidence absent for the quarter and the deferral holds by default rather than by finding. Four consecutive quarters with no intake is itself a P7 finding routed to #20 and #1.

ID	Sending owner	Receiving owner	Artifact	Timing	Consequence of non-arrival
I-3	#37 Automation platform rules	#45	Weekly rule diff, filtered to time-conditioned rules	Weekly, with the Monday cycle	P4 cannot run. The exclusion set is recorded as unverified for that cycle rather than assumed unchanged, and the omission is named in the P1 line. An unverified exclusion set is the precondition for F1.
I-4	#45	#37 Automation platform rules	Exclusion reason string citing v4.0 Section 9 L10 and the D12 status	Same day as detection	#37 registers the rule with no exclusion basis, which lands in its own un documented-divergence class. The rule stays live and the account is dayparting without having decided to.
I-5	#45	#28 Placement and bid-modifier operations	At activation only: the hourly modifier as an additional multiplicand in the effective-CPC ceiling worksheet, plus the eligible campaign list	Once, at P8 handoff	Modifiers deploy against a ceiling computed without the hourly term. The worksheet understates the effective maximum CPC by the full hourly bound, which is the exact class of error the ceiling worksheet exists to prevent.
I-6	#12 Exact campaign operations	#45	Intraday truncation alarms on push campaigns, as the confound list for P3	Weekly, or on alarm	P3 branch A loses its primary detection input and truncation cases route to branch B, where the P2 arithmetic will stop them for the wrong reason. The proposal is refused correctly and the $28.73 per week budget finding in Worked Example B is never made.
I-7	#45	#1 Document governance and supersession	Component state line with its deciding artifact and date	On any state change; otherwise quarterly with P7	The register continues to show NOT BUILT against a live document, and the next Gap Register read counts this component as missing work that exists.
I-8	#45	#44 Operator onboarding and certification	Element 11 ramp content and the six live-defense questions	By 2026-09-25, the #44 due date	The certification path carries no dayparting module, and an operator can be certified while holding the belief that hourly patterns in this account are readable. That belief is what F3 costs money on.
I-9	#45	#40 Monthly, quarterly and annual cadence	The quarterly deferral review slot	Once, at this SOP's release	P7 has no scheduled trigger and F6 becomes the expected outcome rather than the exception. A deferral with no calendar slot is a deferral with no owner.

ID	Sending owner	Receiving owner	Artifact	Timing	Consequence of non-arrival
I-10	#45	Cross-Reference Ledger close-out	The returned row for component #45, in copy-paste form	At session close, before the next session opens	The next session citing this boundary builds against the seed row reading NOT BUILT, restates what this document owns, and the anti-drift rule 3 check fails silently at the corpus level.
I-11	#45	#38 Bulk-file deployment mechanics	At activation only: staged modifier sets	Not before activation	Not applicable while the deferral holds, and that is the contract: this component stages nothing. Any staged set citing #45 before activation is F4 and is blocked at the staging pre-flight.

13. Closing: Amendment Items and Findings

13.1 v4.1 amendment items discovered

Each is a number or rule this component needs and v4.0 does not contain. None is written into the document as a threshold.

ID	Item	Statement
AM-1	Deferral-state vocabulary	Held, armed, and retired are not in v4.0 Section 13.1 or the KDR v3.0 closed vocabulary. Admission runs through the V3 amendment route that Section 13.1 names. Same class as the state-vocabulary items the Ledger records for #43 and #46.
AM-2	Element 2 lag table	v4.0 contains no signal-to-start lag values. The Element 2 lags are proposed by this build. The Ledger records the same item owed for #20.
AM-3	Hour-bucket granularity	L10 says hour-level evidence without stating whether the bucket is 24 pooled clock hours or 168 day-of-week hours. The sufficiency arithmetic differs by a factor of 7 and two correct readings of L10 reach opposite verdicts. This is the highest-value item on the list.
AM-4	G2 application to buckets	G2's 15-click, 100-click, and 1,000-impression minimums are written for keyword and campaign windows. v4.0 does not state whether they apply per hour bucket. P2 applies them per bucket as the conservative reading and says so; v4.1 should confirm or replace that reading.
AM-5	Activation mechanics	L10 gives a bound and a window and nothing else. An execution SOP needs the step cadence, the re-read interval, the hold test, and the exit rule. Until v4.1 supplies them, P8 closes the gate rather than opening it.
AM-6	Retirement evidence bar	Strategic Doctrine v1.0 Section 10 permits permanent retirement of the slot on AMC evidence. v4.0 states no evidence bar for that decision, so P9 presents a case rather than clearing a threshold.
AM-7	T6 Level closed list, owner #39 and the workbook, not v4.1	The T6 DECISION LOG Level column's closed list is product, syntax, keyword, campaign, placement. A component-level state entry has no legal value and P1 currently uses product with the component named in Target.

ID	Item	Statement
AM-8	Retention policy, owner #39, not v4.1	The Ledger records retention policy lines as a pending fold-in under #39. Element 8 states retention as life of the workbook because no policy exists to cite.
AM-9	Lever ID for dayparting, owner v4.1 and #39	See finding N1 below. v4.1 should state L10 in one place and the instruments should follow it.
AM-10	Authority Matrix row for deferral components, owner #3	The matrix has no row naming deferral or activation decisions. Element 9 inherits from four adjacent rows, which works but is an inheritance rather than an assignment.

13.2 Conflicts found beyond C1 to C7

Per the pack's standing rule, this session returns findings and row content; the close-out assigns any conflict identifier and the Ledger version. No identifier is minted here.

1. N1. Lever ID drift, L8 against L10, for dayparting. v4.0 Section 9 "Layer 4: The Lever Library (L1-L10)" assigns dayparting to L10 and assigns L8 to format substitution, the cheap-surface pivot. Two instruments disagree: Master PPC Data Architecture v1.0 Section 8 states that the hourly dayparting columns are reserved but unbuilt until the Marketing Stream feed lands and cites L8; the PPC Command Workbook TEMPLATE v1.0 T7 MANIFEST row for D12 names the blocked class as L8 dayparting. Both instruments would route a dayparting reader to the format-substitution rule. v4.0 is the authority and L10 is correct. This is distinct from the gate-numbering conflict the Ledger already registers as C9, which concerns G-series identifiers rather than L-series.
2. N2. The version note's pack state does not reconcile with the attached Ledger. The attached note states Ledger v2.5 at 2026-08-08 with Wave 4 closed and Wave 5A opening. The attached Ledger's header line reads v3.3 at 2026-08-04 with Wave 1 the most recent return, while its own rows carry versions to 3.2 and dates to 2026-08-10 with Wave 5A and Wave 5B rows returned. Three different pack states are described by two documents in one session, and the Ledger header is stale against its own rows. Not resolved here; reported for the close-out.
3. Recurrence, not new: the Ledger's registered conflict C8 fired twice in this session. The Tier 2 items arrived as SOP-28 v1.0 and SOP-37 v1.0 while the Ledger records both at v2.0, at row versions 2.7 and 3.1. Handled the way the Ledger records for #20: both cited by number and one clause only, with no procedure number load-bearing anywhere in this document.

Inspiratek & Ecotero LLC | Confidential | Page 11 of 11